

```

1  2011/12/30 - HAProxy coding style - Willy Tarreau <w@wt.eu>
2  -----
3

```

4 A number of contributors are often embarrassed with coding style issues, they
5 don't always know if they're doing it right, especially since the coding style
6 has elvolved along the years. What is explained here is not necessarily what is
7 applied in the code, but new code should as much as possible conform to this
8 style. Coding style fixes happen when code is replaced. It is useless to send
9 patches to fix coding style only, they will be rejected, unless they belong to
10 a patch series which needs these fixes prior to get code changes. Also, please
11 avoid fixing coding style in the same patches as functional changes, they make
12 code review harder.

13 When modifying a file, you must accept the terms of the license of this file
14 which is recalled at the top of the file, or is explained in the LICENSE file,
15 or if not stated, defaults to LGPL version 2.1 or later for files in the
16 'include' directory, and GPL version 2 or later for all other files.

17 When adding a new file, you must add a copyright banner at the top of the
18 file with your real name, e-mail address and a reminder of the license.
19 Contributions under incompatible licenses or too restrictive licenses might
20 get rejected. If in doubt, please apply the principle above for existing files.

21 All code examples below will intentionally be prefixed with " | " to mark
22 where the code aligns with the first column, and tabs in this document will be
23 represented as a series of 8 spaces so that it displays the same everywhere.

1) Indentation and alignment

1.1) Indentation

Indentation and alignment are two completely different things that people often
get wrong. Indentation is used to mark a sub-level in the code. A sub-level
means that a block is executed in the context of another block (eg: a function
or a condition) :

```

34 | main(int argc, char **argv)
35 | {
36 |     int i;
37 |
38 |     if (argc < 2)
39 |         exit(1);
40 | }

```

In the example above, the code belongs to the main() function and the exit()
call belongs to the if statement. Indentation is made with tabs (\t, ASCII 9),
which allows any developer to configure their preferred editor to use their
own tab size and to still get the text properly indented. Exactly one tab is
used per sub-level. Tabs may only appear at the beginning of a line or after
another tab. It is illegal to put a tab after some text, as it mangles displays
in a different manner for different users (particularly when used to align
comments or values after a #define). If you're tempted to put a tab after some
text, then you're doing it wrong and you need alignment instead (see below).

Note that there are places where the code was not properly indented in the
past. In order to view it correctly, you may have to set your tab size to 8
characters.

1.2) Alignment

```

63 -----
64
65

```

66 Alignment is used to continue a line in a way to makes things easier to group
67 together. By definition, alignment is character-based, so it uses spaces. Tabs
68 would not work because for one tab there would not be as many characters on all
69 displays. For instance, the arguments in a function declaration may be broken
70 into multiple lines using alignment spaces :

```

71 | int http_header_match2(const char *hdr, const char *end,
72 |                        const char *name, int len)
73 | {
74 |     ...
75 |     ...
76 | }
77

```

In this example, the "const char *name" part is aligned with the first
character of the group it belongs to (list of function arguments). Placing it
here makes it obvious that it's one of the function's arguments. Multiple lines
are easy to handle this way. This is very common with long conditions too :

```

83 |     if ((len < eol - sol) &&
84 |         (sol[len] == ':') &&
85 |         (strncasecmp(sol, name, len) == 0)) {
86 |         ctx->del = len;
87 |     }
88

```

If we take again the example above marking tabs with "[-Tabs-]" and spaces
with "#", we get this :

```

91 | [-Tabs-]if ((len < eol - sol) &&
92 | [-Tabs-]####(sol[len] == ':') &&
93 | [-Tabs-]####(strncasecmp(sol, name, len) == 0)) {
94 | [-Tabs-] [-Tabs-]ctx->del = len;
95 | [-Tabs-]
96 | [-Tabs-]}
97

```

It is worth noting that some editors tend to confuse indentations and alignment.
Emacs is notoriously known for this brokenness, and is responsible for almost
all of the alignment mess. The reason is that Emacs only counts spaces, tries
to fill as many as possible with tabs and completes with spaces. Once you know
it, you just have to be careful, as alignment is not used much, so generally it
is just a matter of replacing the last tab with 8 spaces when this happens.

Indentation should be used everywhere there is a block or an opening brace. It
is not possible to have two consecutive closing braces on the same column, it
means that the innermost was not indented.

Right :

```

109 | main(int argc, char **argv)
110 | {
111 |     if (argc > 1) {
112 |         printf("Hello\n");
113 |     }
114 |     exit(0);
115 | }
116 |
117 | }
118

```

Wrong :

```

119 | main(int argc, char **argv)
120 | {
121 |     if (argc > 1) {
122 |         printf("Hello\n");
123 |     }
124 |     exit(0);
125 | }
126 |
127 | }
128

```

A special case applies to switch/case statements. Due to my editor's settings,
I've been used to align "case" with "switch" and to find it somewhat logical

131 since each of the "case" statements opens a sublevel belonging to the "switch"
 132 statement. But indenting "case" after "switch" is accepted too. However in any
 133 case, whatever follows the "case" statement must be indented, whether or not it
 134 contains braces :

```

135 | switch (*arg) {
136 | case 'A': {
137 |     int i;
138 |     for (i = 0; i < 10; i++)
139 |         printf("Please stop pressing 'A'!\n");
140 |     break;
141 | }
142 | case 'B':
143 |     printf("You pressed 'B'\n");
144 |     break;
145 | case 'C':
146 |     break;
147 | case 'D':
148 |     printf("You pressed 'C' or 'D'\n");
149 |     break;
150 | default:
151 |     printf("I don't know what you pressed\n");
152 | }
153 |
154 |

```

2) Braces

155 -----

156

157 Braces are used to delimit multiple-instruction blocks. In general it is
 158 preferred to avoid braces around single-instruction blocks as it reduces the
 159 number of lines :

160 Right :

```

161 | if (argc >= 2)
162 | {
163 |     exit(0);
164 | }
165 |
166 |

```

167 Wrong :

```

168 | if (argc >= 2) {
169 |     exit(0);
170 | }
171 |
172 |

```

173 But it is not that strict, it really depends on the context. It happens from
 174 time to time that single-instruction blocks are enclosed within braces because
 175 it makes the code more symmetrical, or more readable. Example :

```

176 | if (argc < 2) {
177 |     printf("Missing argument\n");
178 |     exit(1);
179 | } else {
180 |     exit(0);
181 | }
182 |
183 |

```

184 Braces are always needed to declare a function. A function's opening brace must
 185 be placed at the beginning of the next line :

186 Right :

```

187 |
188 |
189 | int main(int argc, char **argv)
190 | {
191 |     exit(0);
192 | }
193 |
194 |

```

195 Wrong :

```

196 | int main(int argc, char **argv) {
197 |     exit(0);
198 | }
199 |

```

200 Note that a large portion of the code still does not conforms to this rule, as
 201 it took years to me to adapt to this more common standard which I now tend to
 202 prefer, as it avoids visual confusion when function declarations are broken on
 203 multiple lines :

204 Right :

```

205 |
206 |
207 | int foo(const char *hdr, const char *end,
208 |         const char *name, const char *err,
209 |         int len)
210 | {
211 |     int i;
212 |
213 |

```

214 Wrong :

```

215 | int foo(const char *hdr, const char *end,
216 |         const char *name, const char *err,
217 |         int len) {
218 |     int i;
219 |
220 |

```

221 Braces should always be used where there might be an ambiguity with the code
 222 later. The most common example is the stacked "if" statement where an "else"
 223 may be added later at the wrong place breaking the code, but it also happens
 224 with comments or long arguments in function calls. In general, if a block is
 225 more than one line long, it should use braces.

226 Dangerous code waiting of a victim :

```

227 | if (argc < 2)
228 | {
229 |     /* ret must not be negative here */
230 |     if (ret < 0)
231 |         return -1;
232 | }
233 |

```

234 Wrong change :

```

235 | if (argc < 2)
236 | {
237 |     /* ret must not be negative here */
238 |     if (ret < 0)
239 |         return -1;
240 |     else
241 |         return 0;
242 | }
243 |

```

244 It will do this instead of what your eye seems to tell you :

```

245 | if (argc < 2)
246 | {
247 |     /* ret must not be negative here */
248 |     if (ret < 0)
249 |         return -1;
250 |     else
251 |         return 0;
252 | }
253 |

```

254 Right :

```

255 | if (argc < 2) {
256 |     /* ret must not be negative here */
257 |     if (ret < 0)
258 |         return -1;
259 | }
260 | else
261 |     return 0;
262 |

```

261 Similarly dangerous example :

```
262 | if (ret < 0)
263 |     /* ret must not be negative here */
264 |     complain();
265 |
266 |     init();
267
268 Wrong change to silent the annoying message :
```

```
269
270 | if (ret < 0)
271 |     /* ret must not be negative here */
272 |     //complain();
273 |     init();
274
275 ... which in fact means :
```

```
276
277 | if (ret < 0)
278 |     init();
279
280
281 3) Breaking lines
282 -----
283
```

284 There is no strict rule for line breaking. Some files try to stick to the 80
285 column limit, but given that various people use various tab sizes, it does not
286 make much sense. Also, code is sometimes easier to read with less lines, as it
287 represents less surface on the screen (since each new line adds its tabs and
288 spaces). The rule is to stick to the average line length of other lines. If you
289 are working in a file which fits in 80 columns, try to keep this goal in mind.
290 If you're in a function with 120-chars lines, there is no reason to add many
291 short lines, so you can make longer lines.

292
293 In general, opening a new block should lead to a new line. Similarly, multiple
294 instructions should be avoided on the same line. But some constructs make it
295 more readable when those are perfectly aligned :

296
297 A copy-paste bug in the following construct will be easier to spot :

```
298
299 | if (omult % idiv == 0) { omult /= idiv; idiv = 1; }
300 | if (idiv % omult == 0) { idiv /= omult; omult = 1; }
301 | if (imult % odiv == 0) { imult /= odiv; odiv = 1; }
302 | if (odiv % imult == 0) { odiv /= imult; imult = 1; }
303
304 than in this one :
```

```
305
306 | if (omult % idiv == 0) {
307 |     omult /= idiv;
308 |     idiv = 1;
309 | }
310 | if (idiv % omult == 0) {
311 |     idiv /= omult;
312 |     omult = 1;
313 | }
314 | if (imult % odiv == 0) {
315 |     imult /= odiv;
316 |     odiv = 1;
317 | }
318 | if (odiv % imult == 0) {
319 |     odiv /= imult;
320 |     imult = 1;
321 | }
```

322 What is important is not to mix styles. For instance there is nothing wrong
323 with having many one-line "case" statements as long as most of them are this
324 short like below :

```
326 | switch (*arg) {
327 | case 'A': ret = 1; break;
328 | case 'B': ret = 2; break;
329 | case 'C': ret = 4; break;
330 | case 'D': ret = 8; break;
331 | default : ret = 0; break;
332 | }
```

333
334 Otherwise, prefer to have the "case" statement on its own line as in the
335 example in section 1.2 about alignment. In any case, avoid to stack multiple
336 control statements on the same line, so that it will never be the needed to
337 add two tab levels at once :

```
338
339 Right :
340
341 | switch (*arg) {
342 | case 'A' :
343 |     if (ret < 0)
344 |         ret = 1;
345 |     break;
346 | default : ret = 0; break;
347 | }
```

```
348
349 Wrong :
350
351 | switch (*arg) {
352 | case 'A': if (ret < 0)
353 |         ret = 1;
354 |     break;
355 | default : ret = 0; break;
356 | }
```

```
357
358 Right :
359
360 | if (argc < 2)
361 |     if (ret < 0)
362 |         return -1;
363
364 or Right :
```

```
365
366 | if (argc < 2)
367 |     if (ret < 0) return -1;
368
369 but Wrong :
370
371 | if (argc < 2) if (ret < 0) return -1;
```

372
373 When complex conditions or expressions are broken into multiple lines, please
374 do ensure that alignment is perfectly appropriate, and group all main operators
375 on the same side (which you're free to choose as long as it does not change for
376 every block. Putting binary operators on the right side is preferred as it does
377 not mangle with alignment but various people have their preferences.

```
378
379 Right :
380
381 | if ((txn->flags & TX_NOT_FIRST) &&
382 |     ((req->flags & BF_FULL) ||
383 |      req->r < req->lr ||
384 |      req->r > req->data + req->size - global.tune.maxrewrite)) {
385 |     return 0;
386 | }
```

```
387
388 Right :
389
390
```

```

391 | if ((txn->flags & TX_NOT_FIRST)
392 |    && ((req->flags & BF_FULL)
393 |        || req->r < req->lr
394 |        || req->r > req->data + req->size - global.tune.maxrewrite)) {
395 |     return 0;
396 | }
397 |
398
399 Wrong :
400
401 | if ((txn->flags & TX_NOT_FIRST) &&
402 |     ((req->flags & BF_FULL) ||
403 |      req->r < req->lr
404 |      || req->r > req->data + req->size - global.tune.maxrewrite)) {
405 |     return 0;
406 | }
407

```

If it makes the result more readable, parenthesis may even be closed on their own line in order to align with the opening one. Note that should normally not be needed because such code would be too complex to be digged into.

The "else" statement may either be merged with the closing "if" brace or lie on its own line. The later is preferred but it adds one extra line to each control block which is annoying in short ones. However, if the "else" is followed by an "if", then it should really be on its own line and the rest of the if/else blocks must follow the same style.

```

417
418 Right :
419
420 | if (a < b) {
421 |     return a;
422 | }
423 | else {
424 |     return b;
425 | }
426
427 Right :
428
429 | if (a < b) {
430 |     return a;
431 | } else {
432 |     return b;
433 | }
434

```

```

435 Right :
436
437 | if (a < b) {
438 |     return a;
439 | }
440 | else if (a != b) {
441 |     return b;
442 | }
443 | else {
444 |     return 0;
445 | }
446

```

```

447 Wrong :
448
449 | if (a < b) {
450 |     return a;
451 | } else if (a != b) {
452 |     return b;
453 | } else {
454 |     return 0;
455 | }

```

```

456
457 Wrong :
458
459 | if (a < b) {
460 |     return a;
461 | }
462 | else if (a != b) {
463 |     return b;
464 | } else {
465 |     return 0;
466 | }
467

```

4) Spacing

Correctly spacing code is very important. When you have to spot a bug at 3am, you need it to be clear. When you expect other people to review your code, you want it to be clear and don't want them to get nervous when trying to find what you did.

Always place spaces around all binary or ternary operators, commas, as well as after semi-colons and opening braces if the line continues :

Right :

```

| int ret = 0;
| /* If (x >> 4) { x >= 4; ret += 4; } */
| ret += (x >> 4) ? (x >= 4, 4) : 0;
| val = ret + ((0xFFFFFAA50U >> (x << 1)) & 3) + 1;

```

Wrong :

```

| int ret=0;
| /* If (x>>4) {x>=4;ret+=4;} */
| ret+=(x>>4)?(x>=4,4):0;
| val=ret+((0xFFFFFAA50U>>(x<<1))&3)+1;

```

Never place spaces after unary operators (&, *, -, !, ~, ++, --) nor cast, as they might be confused with the binary counterpart, nor before commas or semicolons :

Right :

```

| bit = !!(~len++ ^ ~(unsigned char)*x);

```

Wrong :

```

| bit = ! ! (~len++ ^ - (unsigned char) * x) ;

```

Note that "sizeof" is a unary operator which is sometimes considered as a language keyword, but in no case it is a function. It does not require parenthesis so it is sometimes followed by spaces and sometimes not when there are no parenthesis. Most people do not really care as long as what is written is unambiguous.

Braces opening a block must be preceded by one space unless the brace is placed on the first column :

Right :

```

| if (argc < 2) {
| }

```

Wrong :

```

521 | if (argc < 2){
522 | }
523 | }
524
525 Do not add unneeded spaces inside parenthesis, they just make the code less
526 readable.
527
528 Right :
529
530 | if (x < 4 && (!y || !z))
531 |     break;
532
533 Wrong :
534
535 | if ( x < 4 && ( !y || !z ) )
536 |     break;
537
538 Language keywords must all be followed by a space. This is true for control
539 statements (do, for, while, if, else, return, switch, case), and for types
540 (int, char, unsigned). As an exception, the last type in a cast does not take
541 a space before the closing parenthesis). The "default" statement in a "switch"
542 construct is generally just followed by the colon. However the colon after a
543 "case" or "default" statement must be followed by a space.
544
545 Right :
546
547 | if (nbargs < 2) {
548 |     printf("Missing arg at %c\n", *(char *)ptr);
549 |     for (i = 0; i < 10; i++) beep();
550 |     return 0;
551 | }
552 | switch (*arg) {
553 |
554 Wrong :
555
556 | if (nbargs < 2){
557 |     printf("Missing arg at %c\n", *(char*)ptr);
558 |     for(i = 0; i < 10; i++)beep();
559 |     return 0;
560 | }
561 | switch(*arg) {
562 |
563 Function calls are different, the opening parenthesis is always coupled to the
564 function name without any space. But spaces are still needed after commas :
565
566 Right :
567
568 | if (!init(arqc, argv))
569 |     exit(1);
570
571 Wrong :
572
573 | if (!init (argc,argv))
574 |     exit(1);
575
576
577 5) Excess or lack of parenthesis
578 -----
579
580 Sometimes there are too many parenthesis in some formulas, sometimes there are
581 too few. There are a few rules of thumb for this. The first one is to respect
582 the compiler's advice. If it emits a warning and asks for more parenthesis to
583 avoid confusion, follow the advice at least to shut the warning. For instance,
584 the code below is quite ambiguous due to its alignment :

```

```

586 | if (var1 < 2 || var2 < 2 &&
587 |     var3 != var4) {
588 |     /* fail */
589 |     return -3;
590 | }
591
592 Note that this code does :
593
594 | if (var1 < 2 || (var2 < 2 && var3 != var4)) {
595 |     /* fail */
596 |     return -3;
597 | }
598
599 But maybe the author meant :
600
601 | if ((var1 < 2 || var2 < 2) && var3 != var4) {
602 |     /* fail */
603 |     return -3;
604 | }
605
606 A second rule to put parenthesis is that people don't always know operators
607 precedence too well. Most often they have no issue with operators of the same
608 category (eg: booleans, integers, bit manipulation, assignment) but once these
609 operators are mixed, it causes them all sort of issues. In this case, it is
610 wise to use parenthesis to avoid errors. One common error concerns the bit
611 shift operators because they're used to replace multiplies and divides but
612 don't have the same precedence :
613
614 The expression :
615
616 | x = y * 16 + 5;
617
618 becomes :
619
620 | x = y << 4 + 5;
621
622 which is wrong because it is equivalent to :
623
624 | x = y << (4 + 5);
625
626 while the following was desired instead :
627
628 | x = (y << 4) + 5;
629
630 It is generally fine to write boolean expressions based on comparisons without
631 any parenthesis. But on top of that, integer expressions and assignments should
632 then be protected. For instance, there is an error in the expression below
633 which should be safely rewritten :
634
635 Wrong :
636
637 | if (var1 > 2 && var1 < 10 ||
638 |     var1 > 2 + 256 && var2 < 10 + 256 ||
639 |     var1 > 2 + 1 << 16 && var2 < 10 + 2 << 16)
640 |     return 1;
641
642 Right (may remove a few parenthesis depending on taste) :
643
644 | if ((var1 > 2 && var1 < 10) ||
645 |     (var1 > (2 + 256) && var2 < (10 + 256)) ||
646 |     (var1 > (2 + (1 << 16)) && var2 < (10 + (1 << 16))))
647 |     return 1;
648
649 The "return" statement is not a function, so it takes no argument. It is a
650 control statement which is followed by the expression to be returned. It does

```

not need to be followed by parenthesis :

```
Wrong :
| int ret0()
| {
|     return 0;
| }
```

```
Right :
| int ret0()
| {
|     return 0;
| }
```

Parenthesis are also found in type casts. Type casting should be avoided as much as possible, especially when it concerns pointer types. Casting a pointer disables the compiler's type checking and is the best way to get caught doing wrong things with data not the size you expect. If you need to manipulate multiple data types, you can use a union instead. If the union is really not convenient and casts are easier, then try to isolate them as much as possible, for instance when initializing function arguments or in another function. Not proceeding this way causes huge risks of not using the proper pointer without any notification, which is especially true during copy-pastes.

Wrong :

```
| void *check_private_data(void *arg1, void *arg2)
| {
|     char *area;
```

```
|     if (*(int *)arg1 > 1000)
|         return NULL;
|     if (memcmp(*(const char *)arg2, "send(", 5) != 0)
|         return NULL;
|     area = malloc(*(int *)arg1);
|     if (!area)
|         return NULL;
|     memcpy(area, *(const char *)arg2 + 5, *(int *)arg1);
|     return area;
| }
```

Right :

```
| void *check_private_data(void *arg1, void *arg2)
| {
|     char *area;
|     int len = *(int *)arg1;
|     const char *msg = arg2;
|     if (len > 1000)
|         return NULL;
|     if (memcmp(msg, "send(", 5) != 0)
|         return NULL;
|     area = malloc(len);
|     if (!area)
|         return NULL;
|     memcpy(area, msg + 5, len);
|     return area;
| }
```

6) Ambiguous comparisons with zero or NULL

In C, '0' has no type, or it has the type of the variable it is assigned to. Comparing a variable or a return value with zero means comparing with the representation of zero for this variable's type. For a boolean, zero is false. For a pointer, zero is NULL. Very often, to make things shorter, it is fine to use the '!' unary operator to compare with zero, as it is shorter and easier to remind or understand than a plain '0'. Since the '!' operator is read "not", it helps read code faster when what follows it makes sense as a boolean, and it is often much more appropriate than a comparison with zero which makes an equal sign appear at an undesirable place. For instance :

```
| if (!isdigit(*c) && !isspace(*c))
|     break;
```

is easier to understand than :

```
| if (isdigit(*c) == 0 && isspace(*c) == 0)
|     break;
```

For a char this "not" operator can be reminded as "no remaining char", and the absence of comparison to zero implies existence of the tested entity, hence the simple strcpy() implementation below which automatically stops once the last zero is copied :

```
| void my_strcpy(char *d, const char *s)
| {
|     while ((*d++ = *s++));
| }
```

Note the double parenthesis in order to avoid the compiler telling us it looks like an equality test.

For a string or more generally any pointer, this test may be understood as an existence test or a validity test, as the only pointer which will fail to validate equality is the NULL pointer :

```
| area = malloc(1000);
| if (!area)
|     return -1;
```

However sometimes it can fool the reader. For instance, strcmp() precisely is one of such functions whose return value can make one think the opposite due to its name which may be understood as "if strings compare...". Thus it is strongly recommended to perform an explicit comparison with zero in such a case, and it makes sense considering that the comparison's operator is the same that is wanted to compare the strings (note that current config parser lacks a lot in this regards) :

```
strcmp(a, b) == 0 <=> a == b
strcmp(a, b) != 0 <=> a != b
strcmp(a, b) < 0 <=> a < b
strcmp(a, b) > 0 <=> a > b
```

Avoid this :

```
| if (strcmp(arg, "test"))
|     printf("this is not a test\n");
|
| if (!strcmp(arg, "test"))
|     printf("this is a test\n");
```

Prefer this :

```
| if (strcmp(arg, "test") != 0)
|     printf("this is not a test\n");
```

```

781 |
782 | if (strcmp(arg, "test") == 0)
783 |     printf("this is a test\n");
784 |

```

7) System call returns

```

785 -----
786
787
788

```

This is not directly a matter of coding style but more of bad habits. It is important to check for the correct value upon return of syscalls. The proper return code indicating an error is described in its man page. There is no reason to consider wider ranges than what is indicated. For instance, it is common to see such a thing :

```

794 | if ((fd = open(file, O_RDONLY)) < 0)
795 |     return -1;
796 |
797

```

This is wrong. The man page says that -1 is returned if an error occurred. It does not suggest that any other negative value will be an error. It is possible that a few such issues have been left in existing code. They are bugs for which fixes are accepted, even though they're currently harmless since open() is not known for returning negative values at the moment.

8) Declaring new types, names and values

```

805 -----
806
807

```

Please refrain from using "typedef" to declare new types, they only obfuscate the code. The reader never knows whether he's manipulating a scalar type or a struct. For instance it is not obvious why the following code fails to build :

```

811 | int delay_expired(timer_t exp, timer_us_t now)
812 | {
813 |     return now >= exp;
814 | }
815 |

```

With the types declared in another file this way :

```

819 | typedef unsigned int timer_t;
820 | typedef struct timeval timer_us_t;

```

This cannot work because we're comparing a scalar with a struct, which does not make sense. Without a typedef, the function would have been written this way without any ambiguity and would not have failed :

```

826 | int delay_expired(unsigned int exp, struct timeval *now)
827 | {
828 |     return now >= exp->tv_sec;
829 | }
830 |

```

Declaring special values may be done using enums. Enums are a way to define structured integer values which are related to each other. They are perfectly suited for state machines. While the first element is always assigned the zero value, not everybody knows that, especially people working with multiple languages all the day. For this reason it is recommended to explicitly force the first value even if it's zero. The last element should be followed by a comma if it is planned that new elements might later be added, this will make later patches shorter. Conversely, if the last element is placed in order to get the number of possible values, it must not be followed by a comma and must be preceded by a comment :

```

841 | enum {
842 |     first = 0,
843 |     second,
844 |     third,
845 |

```

```

846 |         fourth,
847 |     };
848 |
849 |
850 |     enum {
851 |         first = 0,
852 |         second,
853 |         third,
854 |         fourth,
855 |         /* nbvalues must always be placed last */
856 |         nbvalues
857 |     };
858 |

```

Structure names should be short enough not to mangle function declarations, and explicit enough to avoid confusion (which is the most important thing).

Wrong :

```

863 | struct request_args { /* arguments on the query string */
864 |     char *name;
865 |     char *value;
866 |     struct misc_args *next;
867 | };
868 |

```

Right :

```

870 | struct qs_args { /* arguments on the query string */
871 |     char *name;
872 |     char *value;
873 |     struct qs_args *next;
874 | };
875 |
876 |

```

When declaring new functions or structures, please do not use CamelCase, which is a style where upper and lower case are mixed in a single word. It causes a lot of confusion when words are composed from acronyms, because it's hard to stick to a rule. For instance, a function designed to generate an ISN (initial sequence number) for a TCP/IP connection could be called :

```

884 | - generateTcpiPISn()
885 | - generateTcpiPISn()
886 | - generateTcpiPISN()
887 | - generateTcpiPISN()
888 | - generateTCPIPISN()
889 | etc...

```

None is right, none is wrong, these are just preferences which might change along the code. Instead, please use an underscore to separate words. Lowercase is preferred for the words, but if acronyms are upcased it's not dramatic. The real advantage of this method is that it creates unambiguous levels even for short names.

Valid examples :

```

896 | - generate_tcpip_isn()
897 | - generate_tcp_ip_isn()
898 | - generate_TCPIP_ISN()
899 | - generate_TCP_IP_ISN()

```

Another example is easy to understand when 3 arguments are involved in naming the function :

Wrong (naming conflict) :

```

908 | /* returns A + B * C */
909 | int mulABC(int a, int b, int c)
910 |

```

```

911 | {
912 |     return a + b * c;
913 | }
914 |
915 | /* returns (A + B) * C */
916 | int mulABC(int a, int b, int c)
917 | {
918 |     return (a + b) * c;
919 | }
920 |
921 | Right (unambiguous naming) :
922 |
923 | /* returns A + B * C */
924 | int mul_a_bc(int a, int b, int c)
925 | {
926 |     return a + b * c;
927 | }
928 |
929 | /* returns (A + B) * C */
930 | int mul_ab_c(int a, int b, int c)
931 | {
932 |     return (a + b) * c;
933 | }
934 |
935 | Whenever you manipulate pointers, try to declare them as "const", as it will
936 | save you from many accidental misuses and will only cause warnings to be
937 | emitted when there is a real risk. In the examples below, it is possible to
938 | call my_strcpy() with a const string only in the first declaration. Note that
939 | people who ignore "const" are often the ones who cast a lot and who complain
940 | from segfaults when using strtok() !
941 |
942 | Right :
943 |
944 | void my_strcpy(char *d, const char *s)
945 | {
946 |     while ((*d++ = *s++));
947 | }
948 |
949 | void say_hello(char *dest)
950 | {
951 |     my_strcpy(dest, "hello\n");
952 | }
953 |
954 | Wrong :
955 |
956 | void my_strcpy(char *d, char *s)
957 | {
958 |     while ((*d++ = *s++));
959 | }
960 |
961 | void say_hello(char *dest)
962 | {
963 |     my_strcpy(dest, "hello\n");
964 | }
965 |
966 |
967 |
968 | 9) Getting macros right
969 | -----
970 |
971 | It is very common for macros to do the wrong thing when used in a way their
972 | author did not have in mind. For this reason, macros must always be named with
973 | uppercase letters only. This is the only way to catch the developer's eye when
974 | using them, so that he double-checks whether he's taking risks or not. First,
975 | macros must never ever be terminated by a semi-colon, or they will close the
976 | wrong block once in a while. For instance, the following will cause a build

```

```

976 | error before the "else" due to the double semi-colon :
977 |
978 | Wrong :
979 |
980 | #define WARN printf("warning\n");
981 | ...
982 |     if (a < 0)
983 |         WARN;
984 |     else
985 |         a--;
986 |
987 | Right :
988 |
989 | #define WARN printf("warning\n")
990 |
991 | If multiple instructions are needed, then use a do { } while (0) block, which
992 | is the only construct which respects *exactly* the semantics of a single
993 | instruction :
994 |
995 | #define WARN do { printf("warning\n"); log("warning\n"); } while (0)
996 | ...
997 |
998 |     if (a < 0)
999 |         WARN;
1000 |     else
1001 |         a--;
1002 |
1003 | Second, do not put unprotected control statements in macros, they will
1004 | definitely cause bugs :
1005 |
1006 | Wrong :
1007 |
1008 | #define WARN if (verbose) printf("warning\n")
1009 | ...
1010 |     if (a < 0)
1011 |         WARN;
1012 |     else
1013 |         a--;
1014 |
1015 | Which is equivalent to the undesired form below :
1016 |
1017 |     if (a < 0)
1018 |         if (verbose)
1019 |             printf("warning\n");
1020 |     else
1021 |         a--;
1022 |
1023 | Right way to do it :
1024 |
1025 | #define WARN do { if (verbose) printf("warning\n"); } while (0)
1026 | ...
1027 |     if (a < 0)
1028 |         WARN;
1029 |     else
1030 |         a--;
1031 |
1032 | Which is equivalent to :
1033 |
1034 |     if (a < 0)
1035 |         do { if (verbose) printf("warning\n"); } while (0);
1036 |     else
1037 |         a--;
1038 |
1039 | Macro parameters must always be surrounded by parenthesis, and must never be
1040 | duplicated in the same macro unless explicitly stated. Also, macros must not be

```


defined with operators without surrounding parenthesis. The MIN/MAX macros are a pretty common example of multiple misuses, but this happens as early as when using bit masks. Most often, in case of any doubt, try to use inline functions instead.

```
Wrong :
|
| #define MIN(a, b) a < b ? a : b
|
| /* returns 2 * min(a,b) + 1 */
1050 | int double_min_pl(int a, int b)
1051 | {
1052 |     return 2 * MIN(a, b) + 1;
1053 | }
1054 |
```

What this will do :

```
1058 | int double_min_pl(int a, int b)
1059 | {
1060 |     return 2 * a < b ? a : b + 1;
1061 | }
1062 |
```

Which is equivalent to :

```
1063 | int double_min_pl(int a, int b)
1064 | {
1065 |     return (2 * a) < b ? a : (b + 1);
1066 | }
1067 |
1068 |
```

The first thing to fix is to surround the macro definition with parenthesis to avoid this mistake :

```
1070 | #define MIN(a, b) (a < b ? a : b)
```

But this is still not enough, as can be seen in this example :

```
1076 |
1077 | /* compares either a or b with c */
1078 | int min_ab_c(int a, int b, int c)
1079 | {
1080 |     return MIN(a ? a : b, c);
1081 | }
1082 |
```

Which is equivalent to :

```
1083 | int min_ab_c(int a, int b, int c)
1084 | {
1085 |     return (a ? a : b < c ? a : b : c);
1086 | }
1087 |
1088 |
```

Which in turn means a totally different thing due to precedence :

```
1089 |
1090 | int min_ab_c(int a, int b, int c)
1091 | {
1092 |     return (a ? a : ((b < c) ? (a ? a : b) : c));
1093 | }
1094 |
1095 |
```

This can be fixed by surrounding *each* argument in the macro with parenthesis:

```
1098 | #define MIN(a, b) ((a) < (b) ? (a) : (b))
```

But this is still not enough, as can be seen in this example :

```
1100 | int min_apl_b(int a, int b)
1101 | {
1102 |     return MIN(++a, b);
1103 | }
1104 |
1105 |
```

```
1106 | }
1107 |
1108 | Which is equivalent to :
1109 |
1110 | int min_apl_b(int a, int b)
1111 | {
1112 |     return ((++a) < (b) ? (++a) : (b));
1113 | }
1114 |
```

Again, this is wrong because "a" is incremented twice if below b. The only way to fix this is to use a compound statement and to assign each argument exactly once to a local variable of the same type :

```
1118 | #define MIN(a, b) ({ typeof(a) __a = (a); typeof(b) __b = (b); \
1119 |     ((__a) < (__b) ? (__a) : (__b)); \
1120 | })
1121 |
```

At this point, using static inline functions is much cleaner if a single type is to be used :

```
1122 | static inline int min(int a, int b)
1123 | {
1124 |     return a < b ? a : b;
1125 | }
1126 |
1127 |
```

10) Includes

Includes are as much as possible listed in alphabetically ordered groups :

- the libc-standard includes (those without any path component)
- the includes more or less system-specific (sys/*, netinet/*, ...)
- includes from the local "common" subdirectory
- includes from the local "types" subdirectory
- includes from the local "proto" subdirectory

Each section is just visually delimited from the other ones using an empty line. The two first ones above may be merged into a single section depending on developer's preference. Please do not copy-paste include statements from other files. Having too many includes significantly increases build time and makes it hard to find which ones are needed later. Just include what you need and if possible in alphabetical order so that when something is missing, it becomes obvious where to look for it and where to add it.

All files should include <common/config.h> because this is where build options are prepared.

Header files are split in two directories ("types" and "proto") depending on what they provide. Types, structures, enums and #defines must go into the "types" directory. Function prototypes and inlined functions must go into the "proto" directory. This split is because of inlined functions which cross-reference types from other files, which cause a chicken-and-egg problem if the functions and types are declared at the same place.

All headers which do not depend on anything currently go to the "common" subdirectory, but are equally well placed into the "proto" directory. It is possible that one day the "common" directory will disappear.

Include files must be protected against multiple inclusion using the common #ifndef/#define/#endif trick with a tag derived from the include file and its location.

11) Comments

```

1171 Comments are preferably of the standard 'C' form using /* */. The C++ form "//"
1172 are tolerated for very short comments (eg: a word or two) but should be avoided
1173 as much as possible. Multi-line comments are made with each intermediate line
1174 starting with a star aligned with the first one, as in this example :
1175
1176
1177 | /*
1178 | * This is a multi-line
1179 | * comment.
1180 | */

```

```

1181
1182
1183
1184
1185

```

If multiple code lines need a short comment, try to align them so that you can have multi-line sentences. This is rarely needed, only for really complex constructs.

Do not tell what you're doing in comments, but explain why you're doing it if it seems not to be obvious. Also *do* indicate at the top of function what they accept and what they don't accept. For instance, strcpv() only accepts output buffers at least as large as the input buffer, and does not support any NULL pointer. There is nothing wrong with that if the caller knows it.

Wrong use of comments :

```

1191
1192
1193
1194 | int flsnz8(unsigned int x)
1195 | {
1196 |     int ret = 0;
1197 |     if (x >> 4) { x >=> 4; ret += 4; } /* initialize ret */
1198 |     return ret + ((0xFFFFAA50U >> (x << 1)) & 3) + 1; /* add 4 to ret if needed */
1199 | }
1200 | ...
1201 | bit = -len + (skip << 3) + 9; /* update bit */
1202
1203
1204

```

Right use of comments :

```

1205 | /* This function returns the positoin of the highest bit set in the lowest
1206 | * byte of <x>, between 0 and 7. It only works if <x> is non-null. It uses
1207 | * a 32-bit value as a lookup table to return one of 4 values for the
1208 | * highest 16 possible 4-bit values.
1209 | */
1210 | int flsnz8(unsigned int x)
1211 | {
1212 |     int ret = 0;
1213 |     if (x >> 4) { x >=> 4; ret += 4; }
1214 |     return ret + ((0xFFFFAA50U >> (x << 1)) & 3) + 1;
1215 | }
1216 | ...
1217 | bit = -len + (skip << 3) + 9; /* (skip << 3) + (8 - len), saves 1 cycle */
1218
1219

```

12) Use of assembly

```

1220 -----
1221
1222

```

There are many projects where use of assembly code is not welcome. There is no problem with use of assembly in haproxy, provided that :

- a) an alternate C-form is provided for architectures not covered
- b) the code is small enough and well commented enough to be maintained

It is important to take care of various incompatibilities between compiler versions, for instance regarding output and clobbered registers. There are a number of documentations on the subject on the net. Anyway if you are fiddling with assembly, you probably know that already.

Example :

```

1233 | /* gcc does not know when it can safely divide 64 bits by 32 bits. Use this
1234
1235

```

```

1236 |
1237 | * function when you know for sure that the result fits in 32 bits, because
1238 | * it is optimal on x86 and on 64bit processors.
1239 |
1240 | static inline unsigned int div64_32(unsigned long long o1, unsigned int o2)
1241 | {
1242 |     unsigned int result;
1243 |     #ifdef __i386__
1244 |         asm("divl %2"
1245 |             : "=a" (result)
1246 |             : "A"(o1), "r"(o2));
1247 |     #else
1248 |         result = o1 / o2;
1249 |     #endif
1250 |     return result;
1251 | }

```