

```
1 -----
2 HAPROXY
3 Configuration Manual
4 -----
5 version 1.5.15
6 willy tarreau
7 2015/11/01
8
9
10 This document covers the configuration language as implemented in the version
11 specified above. It does not provide any hint, example or advice. For such
12 documentation, please refer to the Reference Manual or the Architecture Manual.
13 The summary below is meant to help you search sections by name and navigate
14 through the document.
15
```

```
16 Note to documentation contributors :
17 This document is formatted with 80 columns per line, with even number of
18 spaces for indentation and without tabs. Please follow these rules strictly
19 so that it remains easily printable everywhere. If a line needs to be
20 printed verbatim and does not fit, please end each line with a backslash
21 ('\') and continue on next line, indented by two characters. It is also
22 sometimes useful to prefix all output lines (logs, console outs) with 3
23 closing angle brackets ('>>>') in order to help get the difference between
24 inputs and outputs when it can become ambiguous. If you add sections,
25 please update the summary below for easier searching.
26
```

Summary

```
27 -----
28
29
30
31 1. Quick reminder about HTTP
32 1.1. The HTTP transaction model
33 1.2. HTTP request
34 1.2.1. The Request line
35 1.2.2. The request headers
36 1.3. HTTP response
37 1.3.1. The Response line
38 1.3.2. The response headers
39
40 2. Configuring HAPROXY
41 2.1. Configuration file format
42 2.2. Time format
43 2.3. Examples
44
45 3. Global parameters
46 3.1. Process management and security
47 3.2. Performance tuning
48 3.3. Debugging
49 3.4. Userlists
50 3.5. Peers
51
52 4. Proxies
53 4.1. Proxy keywords matrix
54 4.2. Alphabetically sorted keywords reference
55
56 5. Bind and Server options
57 5.1. Bind options
58 5.2. Server and default-server options
59
60 6. HTTP header manipulation
61
62 7. Using ACLs and fetching samples
63 7.1. ACL basics
64 7.1.1. Matching booleans
65 7.1.2. Matching integers
```

```
66 7.1.3. Matching regular expressions (regexes)
67 7.1.4. Matching arbitrary data blocks
68 7.1.5. Matching IPv4 and IPv6 addresses
69 7.1.6. Using ACLs to form conditions
70 7.2. Fetching samples
71 7.3. Converters
72 7.3.1. Fetching samples from internal states
73 7.3.2. Fetching samples at Layer 4
74 7.3.3. Fetching samples at Layer 5
75 7.3.4. Fetching samples from buffer contents (Layer 6)
76 7.3.5. Fetching HTTP samples (Layer 7)
77 7.3.6. Pre-defined ACLs
78 7.4.
79
80 8. Logging
81 8.1. Log levels
82 8.2. Log formats
83 8.2.1. Default log format
84 8.2.2. TCP log format
85 8.2.3. HTTP log format
86 8.2.4. Custom log format
87 8.2.5. Error log format
88 8.3. Advanced logging options
89 8.3.1. Disabling logging of external tests
90 8.3.2. Logging before waiting for the session to terminate
91 8.3.3. Raising log level upon errors
92 8.3.4. Disabling logging of successful connections
93 8.4. Timing events
94 8.5. Session state at disconnection
95 8.6. Non-printable characters
96 8.7. Capturing HTTP cookies
97 8.8. Capturing HTTP headers
98 8.9. Examples of logs
99
100 9. Statistics and monitoring
101 9.1. CSV format
102 9.2. Unix Socket commands
103
104 1. Quick reminder about HTTP
105 -----
106
107 When haproxy is running in HTTP mode, both the request and the response are
108 fully analyzed and indexed, thus it becomes possible to build matching criteria
109 on almost anything found in the contents.
110
111 However, it is important to understand how HTTP requests and responses are
112 formed, and how HAPROXY decomposes them. It will then become easier to write
113 correct rules and to debug existing configurations.
114
115 1.1. The HTTP transaction model
116 -----
117
118 The HTTP protocol is transaction-driven. This means that each request will lead
119 to one and only one response. Traditionally, a TCP connection is established
120 from the client to the server, a request is sent by the client on the
121 connection, the server responds and the connection is closed. A new request
122 will involve a new connection :
123
124 [CON1] [REQ1] ... [RESP1] [CL01] [CON2] [REQ2] ... [RESP2] [CL02] ...
125
126 In this mode, called the "HTTP close" mode, there are as many connection
127 establishments as there are HTTP transactions. Since the connection is closed
128 by the server after the response, the client does not need to know the content
```

length.

Due to the transactional nature of the protocol, it was possible to improve it to avoid closing a connection between two subsequent transactions. In this mode however, it is mandatory that the server indicates the content length for each response so that the client does not wait indefinitely. For this, a special header is used: "Content-length". This mode is called the "keep-alive" mode :

```
[CON] [REQ1] ... [RESP1] [REQ2] ... [RESP2] [CLO] ...
```

Its advantages are a reduced latency between transactions, and less processing power required on the server side. It is generally better than the close mode, but not always because the clients often limit their concurrent connections to a smaller value.

A last improvement in the communications is the pipelining mode. It still uses keep-alive, but the client does not wait for the first response to send the second request. This is useful for fetching large number of images composing a page :

```
[CON] [REQ1] [REQ2] ... [RESP1] [RESP2] [CLO] ...
```

This can obviously have a tremendous benefit on performance because the network latency is eliminated between subsequent requests. Many HTTP agents do not correctly support pipelining since there is no way to associate a response with the corresponding request in HTTP. For this reason, it is mandatory for the server to reply in the exact same order as the requests were received.

By default HAProxy operates in keep-alive mode with regards to persistent connections: for each connection it processes each request and response, and leaves the connection idle on both sides between the end of a response and the start of a new request.

HAProxy supports 5 connection modes :

- keep alive : all requests and responses are processed (default)
- tunnel : only the first request and response are processed, everything else is forwarded with no analysis.
- passive close : tunnel with "Connection: close" added in both directions.
- server close : the server-facing connection is closed after the response.
- forced close : the connection is actively closed after end of response.

1.2. HTTP request

First, let's consider this HTTP request :

Line	Contents
number	
1	GET /serv/login.php?lang=en&profile=2 HTTP/1.1
2	Host: www.mydomain.com
3	User-agent: my small browser
4	Accept: image/jpeg, image/gif
5	Accept: image/png

1.2.1. The Request line

Line 1 is the "request line". It is always composed of 3 fields :

- a METHOD : GET
- a URI : /serv/login.php?lang=en&profile=2
- a version tag : HTTP/1.1

196 All of them are delimited by what the standard calls LWS (linear white spaces),
197 which are commonly spaces, but can also be tabs or line feeds/carriage returns
198 followed by spaces/tabs. The method itself cannot contain any colon (':') and
199 is limited to alphabetic letters. All those various combinations make it
200 desirable that HAProxy performs the splitting itself rather than leaving it to
201 the user to write a complex or inaccurate regular expression.

202 The URI itself can have several forms :

203
204 - A "relative URI" :

```
205 /serv/login.php?lang=en&profile=2
```

207 It is a complete URL without the host part. This is generally what is
208 received by servers, reverse proxies and transparent proxies.

209 - An "absolute URI", also called a "URL" :

```
210 http://192.168.0.12:8080/serv/login.php?lang=en&profile=2
```

211 It is composed of a "scheme" (the protocol name followed by '://'), a host
212 name or address, optionally a colon (':') followed by a port number, then
213 a relative URI beginning at the first slash ('/') after the address part.
214 This is generally what proxies receive, but a server supporting HTTP/1.1
215 must accept this form too.

216 - a star ('*') : this form is only accepted in association with the OPTIONS
217 method and is not relayable. It is used to inquiry a next hop's
218 capabilities.

219 - an address:port combination : 192.168.0.12:80

220 This is used with the CONNECT method, which is used to establish TCP
221 tunnels through HTTP proxies, generally for HTTPS, but sometimes for
222 other protocols too.

223 In a relative URI, two sub-parts are identified. The part before the question
224 mark is called the "path". It is typically the relative path to static objects
225 on the server. The part after the question mark is called the "query string".
226 It is mostly used with GET requests sent to dynamic scripts and is very
227 specific to the language, framework or application in use.

1.2.2. The request headers

238 The headers start at the second line. They are composed of a name at the
239 beginning of the line, immediately followed by a colon (':'). Traditionally,
240 an LWS is added after the colon but that's not required. Then come the values.
241 Multiple identical headers may be folded into one single line, delimiting the
242 values with commas, provided that their order is respected. This is commonly
243 encountered in the "Cookie:" field. A header may span over multiple lines if
244 the subsequent lines begin with an LWS. In the example in 1.2, lines 4 and 5
245 define a total of 3 values for the "Accept:" header.

246 Contrary to a common mis-conception, header names are not case-sensitive, and
247 their values are not either if they refer to other header names (such as the
248 "Connection:" header).

249 The end of the headers is indicated by the first empty line. People often say
250 that it's a double line feed, which is not exact, even if a double line feed
251 is one valid form of empty line.

252 Fortunately, HAProxy takes care of all these complex combinations when indexing
253 headers, checking values and counting them, so there is no reason to worry
254 about the way they could be written, but it is important not to accuse an

application of being buggy if it does unusual, valid things.

Important note:

As suggested by RFC2616, HAProxy normalizes headers by replacing line breaks in the middle of headers by LWS in order to join multi-line headers. This is necessary for proper analysis and helps less capable HTTP parsers to work correctly and not to be fooled by such complex constructs.

1.3. HTTP response

An HTTP response looks very much like an HTTP request. Both are called HTTP messages. Let's consider this HTTP response :

Line number	Contents
1	HTTP/1.1 200 OK
2	Content-length: 350
3	Content-Type: text/html

As a special case, HTTP supports so called "Informational responses" as status codes 1xx. These messages are special in that they don't convey any part of the response, they're just used as sort of a signaling message to ask a client to continue to post its request for instance. In the case of a status 100 response the requested information will be carried by the next non-100 response message following the informational one. This implies that multiple responses may be sent to a single request, and that this only works when keep-alive is enabled (1xx messages are HTTP/1.1 only). HAProxy handles these messages and is able to correctly forward and skip them, and only process the next non-100 response. As such, these messages are neither logged nor transformed, unless explicitly state otherwise. Status 101 messages indicate that the protocol is changing over the same connection and that haproxy must switch to tunnel mode, just as if a CONNECT had occurred. Then the Upgrade header would contain additional information about the type of protocol the connection is switching to.

1.3.1. The Response line

Line 1 is the "response line". It is always composed of 3 fields :

- a version tag : HTTP/1.1
- a status code : 200
- a reason : OK

The status code is always 3-digit. The first digit indicates a general status :

- 1xx = informational message to be skipped (eg: 100, 101)
- 2xx = OK, content is following (eg: 200, 206)
- 3xx = OK, no content following (eg: 302, 304)
- 4xx = error caused by the client (eg: 401, 403, 404)
- 5xx = error caused by the server (eg: 500, 502, 503)

Please refer to RFC2616 for the detailed meaning of all such codes. The "reason" field is just a hint, but is not parsed by clients. Anything can be found there, but it's a common practice to respect the well-established messages. It can be composed of one or multiple words, such as "OK", "Found", or "Authentication Required".

HAProxy may emit the following status codes by itself :

- | Code | When / reason |
|------|---|
| 200 | access to stats page, and when replying to monitoring requests |
| 301 | when performing a redirection, depending on the configured code |
| 302 | when performing a redirection, depending on the configured code |

- 326 when performing a redirection, depending on the configured code
- 307 when performing a redirection, depending on the configured code
- 308 when performing a redirection, depending on the configured code
- 400 for an invalid or too large request
- 329 when an authentication is required to perform the action (when accessing the stats page)
- 331 when a request is forbidden by a "block" ACL or "reqdeny" filter
- 332 when the request timeout strikes before the request is complete
- 333 when haproxy encounters an unrecoverable internal error, such as a memory allocation failure, which should never happen
- 335 when the server returns an empty, invalid or incomplete response, or when an "rspdeny" filter blocks the response.
- 336 when no server was available to handle the request, or in response to monitoring requests which match the "monitor fail" condition
- 338 when the response timeout strikes before the server responds
- 340
- 341 The error 4xx and 5xx codes above may be customized (see "errorloc" in section 4.2).
- 342
- 343
- 344
- 345
- 346
- 347

1.3.2. The response headers

Response headers work exactly like request headers, and as such, HAProxy uses the same parsing function for both. Please refer to paragraph 1.2.2 for more details.

2. Configuring HAProxy

2.1. Configuration file format

HAProxy's configuration process involves 3 major sources of parameters :

- the arguments from the command-line, which always take precedence
- the "global" section, which sets process-wide parameters
- the proxies sections which can take form of "defaults", "listen", "frontend" and "backend".

The configuration file syntax consists in lines beginning with a keyword referenced in this manual, optionally followed by one or several parameters delimited by spaces. If spaces have to be entered in strings, then they must be preceded by a backslash ('\') to be escaped. Backslashes also have to be escaped by doubling them.

2.2. Time format

Some parameters involve values representing time, such as timeouts. These values are generally expressed in milliseconds (unless explicitly stated otherwise) but may be expressed in any other unit by suffixing the unit to the numeric value. It is important to consider this because it will not be repeated for every keyword. Supported units are :

- us : microseconds. 1 microsecond = 1/1000000 second
- ms : milliseconds. 1 millisecond = 1/1000 second. This is the default.
- s : seconds. 1s = 1000ms
- m : minutes. 1m = 60s = 60000ms
- h : hours. 1h = 60m = 3600s = 3600000ms
- d : days. 1d = 24h = 1440m = 86400s = 86400000ms

2.3. Examples

```
-----
391
392
393
394
395 # Simple configuration for an HTTP proxy listening on port 80 on all
396 # interfaces and forwarding requests to a single backend "servers" with a
397 # single server "server1" listening on 127.0.0.1:8000
398 global
399     daemon
400     maxconn 256
401
402 defaults
403     mode http
404     timeout connect 5000ms
405     timeout client 50000ms
406     timeout server 50000ms
407
408 frontend http-in
409     bind *:80
410     default_backend servers
411
412 backend servers
413     server server1 127.0.0.1:8000 maxconn 32
414
415
416 # The same configuration defined with a single listen block. Shorter but
417 # less expressive, especially in HTTP mode.
418 global
419     daemon
420     maxconn 256
421
422 defaults
423     mode http
424     timeout connect 5000ms
425     timeout client 50000ms
426     timeout server 50000ms
427
428 listen http-in
429     bind *:80
430     server server1 127.0.0.1:8000 maxconn 32
431
432
433 Assuming haproxy is in $PATH, test these configurations in a shell with:
434
435 $ sudo haproxy -f configuration.conf -c
```

3. Global parameters

```
-----
437
438
439 Parameters in the "global" section are process-wide and often OS-specific. They
440 are generally set once for all and do not need being changed once correct. Some
441 of them have command-line equivalents.
```

The following keywords are supported in the "global" section :

* Process management and security

```
446
447 - ca-base
448 - chroot
449 - crt-base
450 - daemon
451 - gid
452 - group
453 - log
454 - log-send-hostname
455 - nbproc
```

```
456 - pidfile
457 - uid
458 - ulimit-n
459 - user
460 - stats
461 - ssl-server-verify
462 - node
463 - description
464 - unix-bind
465
466 * Performance tuning
467 - max-spread-checks
468 - maxconn
469 - maxconmrte
470 - maxcomprate
471 - maxcompcpuusage
472 - maxpipes
473 - maxsessrate
474 - maxsslconn
475 - maxsslrate
476 - noepoll
477 - nqueue
478 - nopoll
479 - nosplice
480 - nogetaddrinfo
481 - spread-checks
482 - tune.bufsize
483 - tune.chksize
484 - tune.comp.maxlevel
485 - tune.http.cookieleen
486 - tune.http.maxhdr
487 - tune.idletimer
488 - tune.maxaccept
489 - tune.maxpollevents
490 - tune.maxrewrite
491 - tune.pipesize
492 - tune.rcvbuf.client
493 - tune.rcvbuf.server
494 - tune.sndbuf.client
495 - tune.sndbuf.server
496 - tune.ssl.cachesize
497 - tune.ssl.lifetime
498 - tune.ssl.force-private-cache
499 - tune.ssl.maxrecord
500 - tune.ssl.default-dh-param
501 - tune.zlib.memlevel
502 - tune.zlib.windowsize
503
504 * Debugging
505 - debug
506 - quiet
507
508
509 3.1. Process management and security
510 -----
511
512 ca-base <dir>
513 Assigns a default directory to fetch SSL CA certificates and CRLs from when a
514 relative path is used with "ca-file" or "crl-file" directives. Absolute
515 locations specified in "ca-file" and "crl-file" prevail and ignore "ca-base".
516
517 chroot <jail dir>
518 Changes current directory to <jail dir> and performs a chroot() there before
519 dropping privileges. This increases the security level in case an unknown
520 vulnerability would be exploited, since it would make it very hard for the
```

attacker to exploit the system. This only works when the process is started with superuser privileges. It is important to ensure that <jail_dir> is both empty and unwritable to anyone.

cpu-map <"all"|"odd"|"even"|process_num> <cpu-set>...

On Linux 2.6 and above, it is possible to bind a process to a specific CPU set. This means that the process will never run on other CPUs. The "cpu-map" directive specifies CPU sets for process sets. The first argument is the process number to bind. This process must have a number between 1 and 32 or 64, depending on the machine's word size, and any process IDs above nbproc are ignored. It is possible to specify all processes at once using "all", only odd numbers using "odd" or even numbers using "even", just like with the "bind-process" directive. The second and forthcoming arguments are CPU sets. Each CPU set is either a unique number between 0 and 31 or 63 or a range with two such numbers delimited by a dash ('-'). Multiple CPU numbers or ranges may be specified, and the processes will be allowed to bind to all of them. Obviously, multiple "cpu-map" directives may be specified. Each "cpu-map" directive will replace the previous ones when they overlap.

crt-base <dir>

Assigns a default directory to fetch SSL certificates from when a relative path is used with "crtfile" directives. Absolute locations specified after "crtfile" prevail and ignore "crt-base".

daemon

Makes the process fork into background. This is the recommended mode of operation. It is equivalent to the command line "-D" argument. It can be disabled by the command line "-db" argument.

gid <number>

Changes the process' group ID to <number>. It is recommended that the group ID is dedicated to HAProxy or to a small set of similar daemons. HAProxy must be started with a user belonging to this group, or with superuser privileges. Note that if haproxy is started from a user having supplementary groups, it will only be able to drop these groups if started with superuser privileges. See also "group" and "uid".

group <group name>

Similar to "gid" but uses the GID of group name <group name> from /etc/group. See also "gid" and "user".

log <address> [len <length>] <facility> [max level [min level]]

Adds a global syslog server. Up to two global servers can be defined. They will receive logs for startups and exits, as well as all logs from proxies configured with "log global".

<address> can be one of:

- An IPv4 address optionally followed by a colon and a UDP port. If no port is specified, 514 is used by default (the standard syslog port).

- An IPv6 address followed by a colon and optionally a UDP port. If no port is specified, 514 is used by default (the standard syslog port).

- A filesystem path to a UNIX domain socket, keeping in mind considerations for chroot (be sure the path is accessible inside the chroot) and uid/gid (be sure the path is appropriately writable).

Any part of the address string may reference any number of environment variables by preceding their name with a dollar sign ('\$') and optionally enclosing them with braces ('{'}), similarly to what is done in Bourne shell.

<length> is an optional maximum line length. Log lines larger than this value will be truncated before being sent. The reason is that syslog servers act differently on log line length. All servers support the default value of 1024, but some servers simply drop larger lines while others do log them. If a server supports long lines, it may make sense to set this value here in order to avoid truncating long lines. Similarly, if a server drops long lines, it is preferable to truncate them before sending them. Accepted values are 80 to 65535 inclusive. The default value of 1024 is generally fine for all standard usages. Some specific cases of long captures or JSON-formatted logs may require larger values.

<facility> must be one of the 24 standard syslog facilities :

```
kern user mail daemon auth syslog lpr news
uucp cron auth2 ftp ntp audit alert cron2
local0 local1 local2 local3 local4 local5 local6 local7
```

An optional level can be specified to filter outgoing messages. By default, all messages are sent. If a maximum level is specified, only messages with a severity at least as important as this level will be sent. An optional minimum level can be specified. If it is set, logs emitted with a more severe level than this one will be capped to this level. This is used to avoid sending "emerg" messages on all terminals on some default syslog configurations. Eight levels are known :

```
emerg alert crit err warning notice info debug
```

log-send-hostname [<string>]

Sets the hostname field in the syslog header. If optional "string" parameter is set the header is set to the string contents, otherwise uses the hostname of the system. Generally used if one is not relaying logs through an intermediate syslog server or for simply customizing the hostname printed in the logs.

log-tag <string>

Sets the tag field in the syslog header to this string. It defaults to the program name as launched from the command line, which usually is "haproxy". Sometimes it can be useful to differentiate between multiple processes running on the same host.

nbproc <number>

Creates <number> processes when going daemon. This requires the "daemon" mode. By default, only one process is created, which is the recommended mode of operation. For systems limited to small sets of file descriptors per process, it may be needed to fork multiple daemons. USING MULTIPLE PROCESSES IS HARDER TO DEBUG AND IS REALLY DISCOURAGED. See also "daemon".

pidfile <pidfile>

Writes pids of all daemons into file <pidfile>. This option is equivalent to the "p" command line argument. The file must be accessible to the user starting the process. See also "daemon".

stats bind-process [all | odd | even | even | <number 1-64> [<number 1-64>]] ...

Limits the stats socket to a certain set of processes numbers. By default the stats socket is bound to all processes, causing a warning to be emitted when nbproc is greater than 1 because there is no way to select the target process when connecting. However, by using this setting, it becomes possible to pin the stats socket to a specific set of processes, typically the first one. The warning will automatically be disabled when this setting is used, whatever the number of processes used. The maximum process ID depends on the machine's word size (32 or 64). A better option consists in using the "process" setting of the "stats socket" line to force the process on each line.

650

651 **ssl-default-bind-ciphers <ciphers>**
 652 This setting is only available when support for OpenSSL was built in. It sets
 653 the default string describing the list of cipher algorithms ("cipher suite")
 654 that are negotiated during the SSL/TLS handshake for all "bind" lines which
 655 do not explicitly define theirs. The format of the string is defined in
 656 "man 1 ciphers" from OpenSSL man pages, and can be for instance a string such
 657 as "AES:ALL:!aNULL:!RC4:@STRENGTH" (without quotes). Please check the
 658 "bind" keyword for more information.

659 **ssl-default-bind-options [<option>]...**
 660 This setting is only available when support for OpenSSL was built in. It sets
 661 default ssl-options to force on all "bind" lines. Please check the "bind"
 662 keyword to see available options.

663 Example:

```
664 global
665     ssl-default-bind-options no-sslv3 no-tls-tickets
```

666 **ssl-default-server-ciphers <ciphers>**

667 This setting is only available when support for OpenSSL was built in. It
 668 sets the default string describing the list of cipher algorithms that are
 669 negotiated during the SSL/TLS handshake with the server, for all "server"
 670 lines which do not explicitly define theirs. The format of the string is
 671 defined in "man 1 ciphers". Please check the "server" keyword for more
 672 information.

673 **ssl-default-server-options [<option>]...**

674 This setting is only available when support for OpenSSL was built in. It sets
 675 default ssl-options to force on all "server" lines. Please check the "server"
 676 keyword to see available options.

677 **ssl-server-verify [none|required]**

678 The default behavior for SSL verify on servers side. If specified to 'none',
 679 servers certificates are not verified. The default is 'required' except if
 680 forced using cmdline option '-dV'.

681 **stats socket [<address:port>|<path>] [param*]**

682 Binds a UNIX socket to <path> or a TCPv4/v6 address to <address:port>.
 683 Connections to this socket will return various statistics outputs and even
 684 allow some commands to be issued to change some runtime settings. Please
 685 consult section 9.2 "Unix Socket commands" for more details.

686 All parameters supported by "bind" lines are supported, for instance to
 687 restrict access to some users or their access rights. Please consult
 688 section 5.1 for more information.

689 **stats timeout <timeout>, in milliseconds**

690 The default timeout on the stats socket is set to 10 seconds. It is possible
 691 to change this value with "stats timeout". The value must be passed in
 692 milliseconds, or be suffixed by a time unit among { us, ms, s, m, h, d }.

701 **stats maxconn <connections>**

702 By default, the stats socket is limited to 10 concurrent connections. It is
 703 possible to change this value with "stats maxconn".

704 **uid <number>**

705 Changes the process' user ID to <number>. It is recommended that the user ID
 706 is dedicated to HAProxy or to a small set of similar daemons. HAProxy must
 707 be started with superuser privileges in order to be able to switch to another
 708 one. See also "gid" and "user".

709 **ulimit-n <number>**

710 Sets the maximum number of per-process file-descriptors to <number>. By
 711 default, it is automatically computed, so it is recommended not to use this
 712 option.

```
716 unix-bind [ prefix <prefix> ] [ mode <mode> ] [ user <user> ] [ uid <uid> ]
717           [ group <group> ] [ gid <gid> ]
```

718 Fixes common settings to UNIX listening sockets declared in "bind" statements.
 719 This is mainly used to simplify declaration of those UNIX sockets and reduce
 720 the risk of errors, since those settings are most commonly required but are
 721 also process-specific. The <prefix> setting can be used to force all socket
 722 path to be relative to that directory. This might be needed to access another
 723 component's chroot. Note that those paths are resolved before haproxy chroots
 724 itself, so they are absolute. The <mode>, <user>, <uid>, <group> and <gid>
 725 all have the same meaning as their homonyms used by the "bind" statement. If
 726 both are specified, the "bind" statement has priority, meaning that the
 727 "unix-bind" settings may be seen as process-wide default settings.

728 user <user name>

729 Similar to "uid" but uses the UID of user name <user name> from /etc/passwd.
 730 See also "uid" and "group".

731 node <name>

732 Only letters, digits, hyphen and underscore are allowed, like in DNS names.

733 This statement is useful in HA configurations where two or more processes or
 734 servers share the same IP address. By setting a different node-name on all
 735 nodes, it becomes easy to immediately spot what server is handling the
 736 traffic.

737 description <text>

738 Add a text that describes the instance.

739 Please note that it is required to escape certain characters (# for example)
 740 and this text is inserted into a html page so you should avoid using
 741 "<" and ">" characters.

742 3.2. Performance tuning

743 -----

744 max-spread-checks <delay in milliseconds>

745 By default, haproxy tries to spread the start of health checks across the
 746 smallest health check interval of all the servers in a farm. The principle is
 747 to avoid hammering services running on the same server. But when using large
 748 check intervals (10 seconds or more), the last servers in the farm take some
 749 time before starting to be tested, which can be a problem. This parameter is
 750 used to enforce an upper bound on delay between the first and the last check,
 751 even if the servers' check intervals are larger. When servers run with
 752 shorter intervals, their intervals will be respected though.

753 maxconn <number>

754 Sets the maximum per-process number of concurrent connections to <number>. It
 755 is equivalent to the command-line argument "-n". Proxies will stop accepting
 756 connections when this limit is reached. The "ulimit-n" parameter is
 757 automatically adjusted according to this value. See also "ulimit-n". Note:
 758 the "select" poller cannot reliably use more than 1024 file descriptors on
 759 some platforms. If your platform only supports select and reports "select
 760 FAILED" on startup, you need to reduce maxconn until it works (slightly
 761 below 500 in general).

762 maxconnrate <number>

763 Sets the maximum per-process number of connections per second to <number>.
 764 Proxies will stop accepting connections when this limit is reached. It can be
 765 used to limit the global capacity regardless of each frontend capacity. It is
 766 important to note that this can only be used as a service protection measure,
 767 as there will not necessarily be a fair share between frontends when the
 768 limit is reached, so it's a good idea to also limit each frontend to some

781 value close to its expected share. Also, lowering tune.maxaccept can improve
782 fairness.
783

784 maxcomprate <number>

785 Sets the maximum per-process input compression rate to <number> kilobytes
786 per second. For each session, if the maximum is reached, the compression
787 level will be decreased during the session. If the maximum is reached at the
788 beginning of a session, the session will not compress at all. If the maximum
789 is not reached, the compression level will be increased up to
790 tune.comp.maxlevel. A value of zero means there is no limit, this is the
791 default value.
792

793 maxcompcpuusage <number>

794 Sets the maximum CPU usage HAProxy can reach before stopping the compression
795 for new requests or decreasing the compression level of current requests.
796 It works like 'maxcomprate' but measures CPU usage instead of incoming data
797 bandwidth. The value is expressed in percent of the CPU used by haproxy. In
798 case of multiple processes (nbproc > 1), each process manages its individual
799 usage. A value of 100 disable the limit. The default value is 100. Setting
800 a lower value will prevent the compression work from slowing the whole
801 process down and from introducing high latencies.
802

803 maxpipes <number>

804 Sets the maximum per-process number of pipes to <number>. Currently, pipes
805 are only used by kernel-based tcp splicing. Since a pipe contains two file
806 descriptors, the "ulimit-n" value will be increased accordingly. The default
807 value is maxconn/4, which seems to be more than enough for most heavy usages.
808 The splice code dynamically allocates and releases pipes, and can fall back
809 to standard copy, so setting this value too low may only impact performance.
810

811 maxessrate <number>

812 Sets the maximum per-process number of sessions per second to <number>.
813 Proxies will stop accepting connections when this limit is reached. It can be
814 used to limit the global capacity regardless of each frontend capacity. It is
815 important to note that this can only be used as a service protection measure,
816 as there will not necessarily be a fair share between frontends when the
817 limit is reached, so it's a good idea to also limit each frontend to some
818 value close to its expected share. Also, lowering tune.maxaccept can improve
819 fairness.
820

821 maxsslconn <number>

822 Sets the maximum per-process number of concurrent SSL connections to
823 <number>. By default there is no SSL-specific limit, which means that the
824 global maxconn setting will apply to all connections. Setting this limit
825 avoids having openssl use too much memory and crash when malloc returns NULL
826 (since it unfortunately does not reliably check for such conditions). Note
827 that the limit applies both to incoming and outgoing connections, so one
828 connection which is deciphered then ciphered accounts for 2 SSL connections.
829

830 maxsslrate <number>

831 Sets the maximum per-process number of SSL sessions per second to <number>.
832 SSL listeners will stop accepting connections when this limit is reached. It
833 can be used to limit the global SSL CPU usage regardless of each frontend
834 capacity. It is important to note that this can only be used as a service
835 protection measure, as there will not necessarily be a fair share between
836 frontends when the limit is reached, so it's a good idea to also limit each
837 frontend to some value close to its expected share. It is also important to
838 note that the sessions are accounted before they enter the SSL stack and not
839 after, which also protects the stack against bad handshakes. Also, lowering
840 tune.maxaccept can improve fairness.
841

842 maxlibmem <number>

843 Sets the maximum amount of RAM in megabytes per process usable by the zlib.
844 When the maximum amount is reached, future sessions will not compress as long
845 as RAM is unavailable. When sets to 0, there is no limit.

846 The default value is 0. The value is available in bytes on the UNIX socket
847 with "show info" on the line "MaxZlibMemUsage", the memory used by zlib is
848 "ZlibMemUsage" in bytes.
849

850 noepoll

851 Disables the use of the "epoll" event polling system on Linux. It is
852 equivalent to the command-line argument "-de". The next polling system
853 used will generally be "poll". See also "nopoll".
854

855 nokqueue

856 Disables the use of the "kqueue" event polling system on BSD. It is
857 equivalent to the command-line argument "-dk". The next polling system
858 used will generally be "poll". See also "nopoll".
859

860 nopoll

861 Disables the use of the "poll" event polling system. It is equivalent to the
862 command-line argument "-dp". The next polling system used will be "select".
863 It should never be needed to disable "poll" since it's available on all
864 platforms supported by HAProxy. See also "nokqueue" and "noepoll".
865

866 nosplice

867 Disables the use of kernel tcp splicing between sockets on Linux. It is
868 equivalent to the command line argument "-ds". Data will then be copied
869 using conventional and more portable recv/send calls. Kernel tcp splicing is
870 limited to some very recent instances of kernel 2.6. Most versions between
871 2.6.25 and 2.6.28 are buggy and will forward corrupted data, so they must not
872 be used. This option makes it easier to globally disable kernel splicing in
873 case of doubt. See also "option splice-auto", "option splice-request" and
874 "option splice-response".
875

876 nogetaddrinfo

877 Disables the use of getaddrinfo(3) for name resolving. It is equivalent to
878 the command line argument "-dg". Deprecated gethostbyname(3) will be used.
879

880 spread-checks <0..50, in percent>

881 Sometimes it is desirable to avoid sending agent and health checks to
882 servers at exact intervals, for instance when many logical servers are
883 located on the same physical server. With the help of this parameter, it
884 becomes possible to add some randomness in the check interval between 0
885 and +/- 50%. A value between 2 and 5 seems to show good results. The
886 default value remains at 0.
887

888 tune.bufsize <number>

889 Sets the buffer size to this size (in bytes). Lower values allow more
890 sessions to coexist in the same amount of RAM, and higher values allow some
891 applications with very large cookies to work. The default value is 16384 and
892 can be changed at build time. It is strongly recommended not to change this
893 from the default value, as very low values will break some services such as
894 statistics, and values larger than default size will increase memory usage,
895 possibly causing the system to run out of memory. At least the global maxconn
896 parameter should be decreased by the same factor as this one is increased.
897 If HTTP request is larger than (tune.bufsize - tune.maxrewrite), haproxy will
898 return HTTP 400 (Bad Request) error. Similarly if an HTTP response is larger
899 than this size, haproxy will return HTTP 502 (Bad Gateway).
900

901 tune.chksize <number>

902 Sets the check buffer size to this size (in bytes). Higher values may help
903 find string or regex patterns in very large pages, though doing so may imply
904 more memory and CPU usage. The default value is 16384 and can be changed at
905 build time. It is not recommended to change this value, but to use better
906 checks whenever possible.
907

908 tune.comp.maxlevel <number>

909 Sets the maximum compression level. The compression level affects CPU
910 usage during compression. This value affects CPU usage during compression.

Each session using compression initializes the compression algorithm with this value. The default value is 1.

`tune.http.cookie_len <number>`

Sets the maximum length of captured cookies. This is the maximum value that the "capture cookie xxx len yyy" will be allowed to take, and any upper value will automatically be truncated to this one. It is important not to set too high a value because all cookie captures still allocate this size whatever their configured value (they share a same pool). This value is per request per response, so the memory allocated is twice this value per connection. When not specified, the limit is set to 63 characters. It is recommended not to change this value.

`tune.http.maxhdr <number>`

Sets the maximum number of headers in a request. When a request comes with a number of headers greater than this value (including the first line), it is rejected with a "400 Bad Request" status code. Similarly, too large responses are blocked with "502 Bad Gateway". The default value is 101, which is enough for all usages, considering that the widely deployed Apache server uses the same limit. It can be useful to push this limit further to temporarily allow a buggy application to work by the time it gets fixed. Keep in mind that each new header consumes 32bits of memory for each session, so don't push this limit too high.

`tune.idletimer <timeout>`

Sets the duration after which haproxy will consider that an empty buffer is probably associated with an idle stream. This is used to optimally adjust some packet sizes while forwarding large and small data alternatively. The decision to use splice() or to send large buffers in SSL is modulated by this parameter. The value is in milliseconds between 0 and 65535. A value of zero means that haproxy will not try to detect idle streams. The default is 1000, which seems to correctly detect end user pauses (eg: read a page before clicking). There should be not reason for changing this value. Please check `tune.ssl.maxrecord` below.

`tune.maxaccept <number>`

Sets the maximum number of consecutive connections a process may accept in a row before switching to other work. In single process mode, higher numbers give better performance at high connection rates. However in multi-process modes, keeping a bit of fairness between processes generally is better to increase performance. This value applies individually to each listener, so that the number of processes a listener is bound to is taken into account. This value defaults to 64. In multi-process mode, it is divided by twice the number of processes the listener is bound to. Setting this value to -1 completely disables the limitation. It should normally not be needed to tweak this value.

`tune.maxpollevents <number>`

Sets the maximum amount of events that can be processed at once in a call to the polling system. The default value is adapted to the operating system. It has been noticed that reducing it below 200 tends to slightly decrease latency at the expense of network bandwidth, and increasing it above 200 tends to trade latency for slightly increased bandwidth.

`tune.maxrewrite <number>`

Sets the reserved buffer space to this size in bytes. The reserved space is used for header rewriting or appending. The first reads on sockets will never fill more than `bufsize-maxrewrite`. Historically it has defaulted to half of `bufsize`, though that does not make much sense since there are rarely large numbers of headers to add. Setting it too high prevents processing of large requests or responses. Setting it too low prevents addition of new headers to already large requests or to POST requests. It is generally wise to set it to about 1024. It is automatically readjusted to half of `bufsize` if it is larger than that. This means you don't have to worry about it when changing `bufsize`.

`tune.pipesize <number>`

Sets the kernel pipe buffer size to this size (in bytes). By default, pipes are the default size for the system. But sometimes when using TCP splicing, it can improve performance to increase pipe sizes, especially if it is suspected that pipes are not filled and that many calls to `splice()` are performed. This has an impact on the kernel's memory footprint, so this must not be changed if impacts are not understood.

`tune.rcvbuf.client <number>`

Forces the kernel socket receive buffer size on the client or the server side to the specified value in bytes. This value applies to all TCP/HTTP frontends and backends. It should normally never be set, and the default size (0) lets the kernel autotune this value depending on the amount of available memory. However it can sometimes help to set it to very low values (eg: 4096) in order to save kernel memory by preventing it from buffering too large amounts of received data. Lower values will significantly increase CPU usage though.

`tune.sndbuf.client <number>`

Forces the kernel socket send buffer size on the client or the server side to the specified value in bytes. This value applies to all TCP/HTTP frontends and backends. It should normally never be set, and the default size (0) lets the kernel autotune this value depending on the amount of available memory. However it can sometimes help to set it to very low values (eg: 4096) in order to save kernel memory by preventing it from buffering too large amounts of received data. Lower values will significantly increase CPU usage though. Another use case is to prevent write timeouts with extremely slow clients due to the kernel waiting for a large part of the buffer to be read before notifying haproxy again.

`tune.ssl.cachesize <number>`

Sets the size of the global SSL session cache, in a number of blocks. A block is large enough to contain an encoded session without peer certificate. An encoded session with peer certificate is stored in multiple blocks depending on the size of the peer certificate. A block uses approximately 200 bytes of memory. The default value may be forced at build time, otherwise defaults to 20000. When the cache is full, the most idle entries are purged and reassigned. Higher values reduce the occurrence of such a purge, hence the number of CPU-intensive SSL handshakes by ensuring that all users keep their session as long as possible. All entries are pre-allocated upon startup and are shared between all processes if "nbproc" is greater than 1. Setting this value to 0 disables the SSL session cache.

`tune.ssl.force-private-cache`

This boolean disables SSL session cache sharing between all processes. It should normally not be used since it will force many renegotiations due to clients hitting a random process. But it may be required on some operating systems where none of the SSL cache synchronization method may be used. In this case, adding a first layer of hash-based load balancing before the SSL layer might limit the impact of the lack of session sharing.

`tune.ssl.lifetime <timeout>`

Sets how long a cached SSL session may remain valid. This time is expressed in seconds and defaults to 300 (5 min). It is important to understand that it does not guarantee that sessions will last that long, because if the cache is full, the longest idle sessions will be purged despite their configured lifetime. The real usefulness of this setting is to prevent sessions from being used for too long.

`tune.ssl.maxrecord <number>`

Sets the maximum amount of bytes passed to `SSL_write()` at a time. Default value 0 means there is no limit. Over SSL/TLS, the client can decipher the data only once it has received a full record. With large records, it means

that clients might have to download up to 16kB of data before starting to process them. Limiting the value can improve page load times on browsers located over high latency or low bandwidth networks. It is suggested to find optimal values which fit into 1 or 2 TCP segments (generally 1448 bytes over Ethernet with TCP timestamps enabled, or 1460 when timestamps are disabled), keeping in mind that SSL/TLS add some overhead. Typical values of 1419 and 2859 gave good results during tests. Use "strace -e trace=write" to find the best value. Haproxy will automatically switch to this setting after an idle stream has been detected (see tune.idletimer above).

tune.ssl.default-dh-param <number>

Sets the maximum size of the Diffie-Hellman parameters used for generating the ephemeral/temporary Diffie-Hellman key in case of DHE key exchange. The final size will try to match the size of the server's RSA (or DSA) key (e.g. a 2048 bits temporary DH key for a 2048 bits RSA key), but will not exceed this maximum value. Default value is 1024. Only 1024 or higher values are allowed. Higher values will increase the CPU load, and values greater than 1024 bits are not supported by Java 7 and earlier clients. This value is not used if static Diffie-Hellman parameters are supplied via the certificate file.

tune.zlib.memlevel <number>

Sets the memLevel parameter in zlib initialization for each session. It defines how much memory should be allocated for the internal compression state. A value of 1 uses minimum memory but is slow and reduces compression ratio, a value of 9 uses maximum memory for optimal speed. Can be a value between 1 and 9. The default value is 8.

tune.zlib.windowSize <number>

Sets the window size (the size of the history buffer) as a parameter of the zlib initialization for each session. Larger values of this parameter result in better compression at the expense of memory usage. Can be a value between 8 and 15. The default value is 15.

3.3. Debugging

debug

Enables debug mode which dumps to stdout all exchanges, and disables forking into background. It is the equivalent of the command-line argument "-d". It should never be used in a production configuration since it may prevent full system startup.

quiet

Do not display any message during startup. It is equivalent to the command-line argument "-q".

3.4. Userlists

It is possible to control access to frontend/backend/listen sections or to http stats by allowing only authenticated and authorized users. To do this, it is required to create at least one userlist and to define users.

userlist <listname>

Creates new userlist with name <listname>. Many independent userlists can be used to store authentication & authorization data for independent customers.

group <groupname> [users <user>,<user>,...]

Adds group <groupname> to the current userlist. It is also possible to attach users to this group by using a comma separated list of names proceeded by "users" keyword.

user <username> [password]insecure-password <password>]

[groups <group>,<group>,...]

Adds user <username> to the current userlist. Both secure (encrypted) and

insecure (unencrypted) passwords can be used. Encrypted passwords are evaluated using the crypt(3) function so depending of the system's capabilities, different algorithms are supported. For example modern Glibc based Linux system supports MD5, SHA-256, SHA-512 and of course classic, DES-based method of encrypting passwords.

Example:

userlist L1

group G1 users tiger,scott

group G2 users xdb,scott

user tiger password \$6\$k6y3o.eP\$jlKBx9za9667qe4(...)xHSwRv6J.C0/D7cV91

user scott insecure-password elgato

user xdb insecure-password hello

userlist L2

group G1

group G2

user tiger password \$6\$k6y3o.eP\$jlKBx(...)xHSwRv6J.C0/D7cV91 groups G1

user scott insecure-password elgato groups G1,G2

user xdb insecure-password hello groups G2

Please note that both lists are functionally identical.

3.5. Peers

It is possible to synchronize server entries in stick tables between several haproxy instances over TCP connections in a multi-master fashion. Each instance pushes its local updates and insertions to remote peers. Server IDs are used to identify servers remotely, so it is important that configurations look similar or at least that the same IDs are forced on each server on all participants. Interrupted exchanges are automatically detected and recovered from the last known point. In addition, during a soft restart, the old process connects to the new one using such a TCP connection to push all its entries before the new process tries to connect to other peers. That ensures very fast replication during a reload, it typically takes a fraction of a second even for large tables.

peers <peersect>

Creates a new peer list with name <peersect>. It is an independent section, which is referenced by one or more stick-tables.

disabled

Disables a peers section. It disables both listening and any synchronization related to this section. This is provided to disable synchronization of stick tables without having to comment out all "peers" references.

enable

This re-enables a disabled peers section which was previously disabled.

peer <peername> <ip>:<port>

Defines a peer inside a peers section.

If <peername> is set to the local peer name (by default hostname, or forced using "-l" command line option), haproxy will listen for incoming remote peer connection on <ip>:<port>. Otherwise, <ip>:<port> defines where to connect to to join the remote peer, and <peername> is used at the protocol level to identify and validate the remote peer on the server side.

During a soft restart, local peer <ip>:<port> is used by the old instance to connect the new one and initiate a complete replication (teaching process).

It is strongly recommended to have the exact same peers declaration on all

peers and to only rely on the "-L" command line argument to change the local peer name. This makes it easier to maintain coherent configuration files across all peers.

Any part of the address string may reference any number of environment variables by preceding their name with a dollar sign ('\$') and optionally enclosing them with braces ('{}'), similarly to what is done in Bourne shell.

Example:

```
peers mypeers
peer haproxy1 192.168.0.1:1024
peer haproxy2 192.168.0.2:1024
peer haproxy3 10.2.0.1:1024

backend mybackend
mode tcp
balance roundrobin
stick-table type ip size 20k peers mypeers
stick on src

server srv1 192.168.0.30:80
server srv2 192.168.0.31:80
```

4. Proxies

Proxy configuration can be located in a set of sections :

```
- defaults <name>
- frontend <name>
- backend <name>
- listen <name>
```

A "defaults" section sets default parameters for all other sections following its declaration. Those default parameters are reset by the next "defaults" section. See below for the list of parameters which can be set in a "defaults" section. The name is optional but its use is encouraged for better readability.

A "frontend" section describes a set of listening sockets accepting client connections.

A "backend" section describes a set of servers to which the proxy will connect to forward incoming connections.

A "listen" section defines a complete proxy with its frontend and backend parts combined in one section. It is generally useful for TCP-only traffic.

All proxy names must be formed from upper and lower case letters, digits, '-' (dash), '_' (underscore), '.' (dot) and ':' (colon). ACL names are case-sensitive, which means that "www" and "WWW" are two different proxies.

Historically, all proxy names could overlap, it just caused troubles in the logs. Since the introduction of content switching, it is mandatory that two proxies with overlapping capabilities (frontend/backend) have different names. However, it is still permitted that a frontend and a backend share the same name, as this configuration seems to be commonly encountered.

Right now, two major proxy modes are supported : "tcp", also known as layer 4, and "http", also known as layer 7. In layer 4 mode, HAProxy simply forwards bidirectional traffic between two sides. In layer 7 mode, HAProxy analyzes the protocol, and can interact with it by allowing, blocking, switching, adding, modifying, or removing arbitrary contents in requests or responses, based on arbitrary criteria.

In HTTP mode, the processing applied to requests and responses flowing over

a connection depends in the combination of the frontend's HTTP options and the backend's. HAProxy supports 5 connection modes :

- KAL : keep alive ("option http-keep-alive") which is the default mode : all requests and responses are processed, and connections remain open but idle between responses and new requests.

- TUN: tunnel ("option http-tunnel") : this was the default mode for versions 1.0 to 1.5-dev21 : only the first request and response are processed, and everything else is forwarded with no analysis at all. This mode should not be used as it creates lots of trouble with logging and HTTP processing.

- PCL: passive close ("option httpclose") : exactly the same as tunnel mode, but with "Connection: close" appended in both directions to try to make both ends close after the first request/response exchange.

- SCL: server close ("option http-server-close") : the server-facing connection is closed after the end of the response is received, but the client-facing connection remains open.

- FCL: forced close ("option forceclose") : the connection is actively closed after the end of the response.

The effective mode that will be applied to a connection passing through a frontend and a backend can be determined by both proxy modes according to the following matrix, but in short, the modes are symmetric, keep-alive is the weakest option and force close is the strongest.

		Backend mode				
Frontend mode		KAL	TUN	PCL	SCL	FCL
	KAL	---	+	+	+	+
	TUN	+	---	+	+	+
	PCL	+	+	---	+	+
	SCL	+	+	+	---	+
	FCL	+	+	+	+	---

4.1. Proxy keywords matrix

The following list of keywords is supported. Most of them may only be used in a limited set of section types. Some of them are marked as "deprecated" because they are inherited from an old syntax which may be confusing or functionally limited, and there are new recommended keywords to replace them. Keywords marked with "(*)" can be optionally inverted using the "no" prefix, eg. "no option contstats". This makes sense when the option has been enabled by default and must be disabled for a specific instance. Such options may also be prefixed with "default" in order to restore default settings regardless of what has been specified in a previous "defaults" section.

keyword	defaults	frontend	listen	backend
acl	-	X	X	X
appsession	-	-	X	X
backlog	X	X	X	-
balance	X	-	X	X
bind	-	X	X	-

1301	bind-process	X	X	X	1366	option log-health-checks	(*)	X	-	X	X
1302	block	-	X	X	1367	option log-separate-errors	(*)	X	X	-	X
1303	capture cookie	-	X	X	1368	option logasap	(*)	X	X	-	X
1304	capture request header	-	X	X	1369	option mysql-check	(*)	X	X	X	X
1305	capture response header	-	X	X	1370	option nologger	(*)	X	X	X	X
1306	clitimeout	(deprecated)	X	X	1371	option originalto	(*)	X	X	X	X
1307	compression	(deprecated)	X	X	1372	option persist	(*)	X	-	X	X
1308	contimeout	(deprecated)	X	X	1373	option pgsql-check	(*)	X	-	X	X
1309	cookie	X	-	X	1374	option prefer-last-server	(*)	X	-	X	X
1310	default-server	X	-	X	1375	option redispatch	(*)	X	-	X	X
1311	default_backend	X	-	X	1376	option redis-check	(*)	X	-	X	X
1312	description	-	X	X	1377	option smtpchk	(*)	X	-	X	X
1313	disabled	-	X	X	1378	option socket-stats	(*)	X	X	-	X
1314	dispatch	-	-	X	1379	option splice-auto	(*)	X	X	X	X
1315	enabled	X	X	X	1380	option splice-request	(*)	X	X	X	X
1316	errorfile	X	X	X	1381	option splice-response	(*)	X	X	X	X
1317	errorloc	X	X	X	1382	option svtcpka	(*)	X	-	X	X
1318	errorloc302	X	X	X	1383	option ssl-hello-chk	(*)	X	-	X	X
1319	-- keyword	defaults	-	frontend - listen -- backend -	1384	-- keyword	defaults	-	frontend - listen -- backend -		
1320	errorloc303	X	X	X	1385	option tcp-check	(*)	X	-	X	X
1321	force-persist	-	X	X	1386	option tcp-smart-accept	(*)	X	X	-	X
1322	fullconn	-	X	X	1387	option tcp-smart-connect	(*)	X	-	X	X
1323	grace	X	X	X	1388	option tcpka	(*)	X	X	X	X
1324	hash-type	X	-	X	1389	option tcplog	(*)	X	X	X	X
1325	http-check disable-on-404	X	-	X	1390	option transparent	(*)	X	-	X	X
1326	http-check expect	-	-	X	1391	persist rdp-cookie	(*)	X	-	X	X
1327	http-check send-state	-	-	X	1392	rate-limit sessions	(*)	X	X	-	X
1328	http-request	-	X	X	1393	redirect	(deprecated)	X	X	X	X
1329	http-response	-	X	X	1394	redisp	(deprecated)	X	-	X	X
1330	http-send-name-header	-	-	X	1395	redispatch	(deprecated)	X	-	X	X
1331	id	-	X	X	1396	reqadd	(*)	X	X	X	X
1332	ignore-persist	-	X	X	1397	reqallow	(*)	X	X	X	X
1333	log	(*)	X	X	1398	reqdel	(*)	X	X	X	X
1334	log-format	X	X	X	1399	reqdeny	(*)	X	X	X	X
1335	max-keep-alive-queue	X	-	X	1400	reqallow	(*)	X	X	X	X
1336	maxconn	X	X	X	1401	reqidel	(*)	X	X	X	X
1337	mode	X	X	X	1402	reqideny	(*)	X	X	X	X
1338	monitor fail	-	X	X	1403	reqipass	(*)	X	X	X	X
1339	monitor-net	X	X	X	1404	reqirep	(*)	X	X	X	X
1340	monitor-uri	X	X	X	1405	reqisetbe	(*)	X	X	X	X
1341	option abortonclose	(*)	X	X	1406	reqitarpit	(*)	X	X	X	X
1342	option accept-invalid-http-request	(*)	X	X	1407	reqpass	(*)	X	X	X	X
1343	option accept-invalid-http-response	(*)	X	X	1408	reqrep	(*)	X	X	X	X
1344	option allbackups	(*)	X	X	1409	-- keyword	defaults	-	frontend - listen -- backend -		
1345	option checkcache	(*)	X	X	1410	reqsetbe	(*)	X	X	X	X
1346	option cliticpka	(*)	X	X	1411	reqtarpit	(*)	X	X	X	X
1347	option contstats	(*)	X	X	1412	retries	(*)	X	-	X	X
1348	option dontlog-normal	(*)	X	X	1413	rspadd	(*)	X	X	X	X
1349	option dontlognull	(*)	X	X	1414	rspdel	(*)	X	X	X	X
1350	option forceclose	(*)	X	X	1415	rspdeny	(*)	X	X	X	X
1351	-- keyword	defaults	-	frontend - listen -- backend -	1416	rspidel	(*)	X	X	X	X
1352	option forwardfor	X	X	X	1417	rspideny	(*)	X	X	X	X
1353	option http-ignore-probes	(*)	X	X	1418	rspirep	(*)	X	X	X	X
1354	option http-keep-alive	(*)	X	X	1419	rsprep	(*)	X	X	X	X
1355	option http-no-delay	(*)	X	X	1420	server	(*)	X	-	X	X
1356	option http-pretend-keepalive	(*)	X	X	1421	source	(*)	X	-	X	X
1357	option http-server-close	(*)	X	X	1422	srvtimeout	(*)	X	-	X	X
1358	option http-tunnel	(*)	X	X	1423	stats admin	(*)	X	X	X	X
1359	option http-use-proxy-header	(*)	X	X	1424	stats auth	(*)	X	X	X	X
1360	option httpchk	(*)	X	X	1425	stats enable	(*)	X	X	X	X
1361	option httpclose	(*)	X	X	1426	stats hide-version	(*)	X	X	X	X
1362	option httplog	(*)	X	X	1427	stats http-request	(*)	X	X	X	X
1363	option http_proxy	(*)	X	X	1428	stats realm	(*)	X	X	X	X
1364	option independent-streams	(*)	X	X	1429	stats refresh	(*)	X	X	X	X
1365	option ldap-check	X	-	X	1430	stats scope	(*)	X	X	X	X

```
1431 stats show-desc X X X X
1432 stats show-legends X X X X
1433 stats show-node X X X X
1434 stats uri X X X X
1435 -- keyword ----- defaults - frontend - listen -- backend -
1436 stick match - - - - -
1437 stick on - - - - -
1438 stick store-request - - - - -
1439 stick store-response - - - - -
1440 stick-table - - - - -
1441 tcp-check connect - - - - -
1442 tcp-check expect - - - - -
1443 tcp-check send - - - - -
1444 tcp-check send-binary - - - - -
1445 tcp-request connection - - - - -
1446 tcp-request content - - - - -
1447 tcp-request inspect-delay - - - - -
1448 tcp-response content - - - - -
1449 tcp-response inspect-delay - - - - -
1450 timeout check X - - - -
1451 timeout client X X X X
1452 timeout client-fin X X X X
1453 timeout clitimeout (deprecated) X X X X
1454 timeout connect (deprecated) X - - - -
1455 timeout contimeout X - - - -
1456 timeout http-keep-alive X X X X
1457 timeout http-request X X X X
1458 timeout queue X - - - -
1459 timeout server X - - - -
1460 timeout server-fin X - - - -
1461 timeout srvtimeout (deprecated) X X X X
1462 timeout tarpit X X X X
1463 timeout tunnel X - - - -
1464 transparent (deprecated) X - - - -
1465 unique-id-format X X X X
1466 unique-id-header X X X X
1467 use_backend - - - - -
1468 use_server - - - - -
1469 -----+-----+-----+-----+-----+-----+-----
1470 defaults frontend listen backend
1471 keyword
1472
1473
1474 4.2. Alphabetically sorted keywords reference
1475 -----
1476 This section provides a description of each keyword and its usage.
1477
1478 acl <aclname> <criterion> [flags] [operator] <value> ...
1479 Declare or complete an access list.
1480 May be used in sections : defaults | frontend | listen | backend
1481 no | yes | yes | yes
1482
1483 Example:
1484 acl invalid_src src 0.0.0/0/7 224.0.0.0/3
1485 acl invalid_src src_port 0.1023
1486 acl local_dst hdr(host) -i localhost
1487
1488 See section 7 about ACL usage.
1489
1490 appsession <cookie> len <length> timeout <holdtime>
1491 [request-learn] [prefix] [mode <path-parameters>|query-strings]
1492 Define session stickiness on an existing application cookie.
1493 May be used in sections : defaults | frontend | listen | backend
1494 no | no | yes | yes
1495
```

```
1496 Arguments :
1497 <cookie> this is the name of the cookie used by the application and which
1498 HAProxy will have to learn for each new session.
1499
1500 <length> this is the max number of characters that will be memorized and
1501 checked in each cookie value.
1502
1503 <holdtime> this is the time after which the cookie will be removed from
1504 memory if unused. If no unit is specified, this time is in
1505 milliseconds.
1506
1507 request-learn
1508 If this option is specified, then haproxy will be able to learn
1509 the cookie found in the request in case the server does not
1510 specify any in response. This is typically what happens with
1511 PHPSESSID cookies, or when haproxy's session expires before
1512 the application's session and the correct server is selected.
1513 It is recommended to specify this option to improve reliability.
1514
1515 prefix
1516 When this option is specified, haproxy will match on the cookie
1517 prefix (or URL parameter prefix). The appsession value is the
1518 data following this prefix.
1519
1520 Example :
1521 appsession ASPSESSIONID len 64 timeout 3h prefix
1522 This will match the cookie ASPSESSIONIDXXXX=XXXX,
1523 the appsession value will be XXXX=XXXX.
1524
1525 mode
1526 This option allows to change the URL parser mode.
1527 2 modes are currently supported :
1528 - path-parameters :
1529 The parser looks for the appsession in the path parameters
1530 part (each parameter is separated by a semi-colon), which is
1531 convenient for JSESSIONID for example.
1532 This is the default mode if the option is not set.
1533 - query-string :
1534 In this mode, the parser will look for the appsession in the
1535 query string.
1536
1537 When an application cookie is defined in a backend, HAProxy will check when
1538 the server sets such a cookie, and will store its value in a table, and
1539 associate it with the server's identifier. Up to <length> characters from
1540 the value will be retained. On each connection, haproxy will look for this
1541 cookie both in the "Cookie:" headers, and as a URL parameter (depending on
1542 the mode used). If a known value is found, the client will be directed to the
1543 server associated with this value. Otherwise, the load balancing algorithm is
1544 applied. Cookies are automatically removed from memory when they have been
1545 unused for a duration longer than <holdtime>.
1546
1547 The definition of an application cookie is limited to one per backend.
1548
1549 Note : Consider not using this feature in multi-process mode (nbproc > 1)
1550 unless you know what you do : memory is not shared between the
1551 processes, which can result in random behaviours.
1552
1553 Example :
1554 appsession JSESSIONID len 52 timeout 3h
1555
1556 See also : "cookie", "capture cookie", "balance", "stick", "stick-table",
1557 "ignore-persist", "nbproc" and "bind-process".
1558
1559 backlog <conns>
1560 Give hints to the system about the approximate listen backlog desired size
```

1561 May be used in sections : defaults | frontend | listen | backend
1562 yes | yes | yes | no
1563 Arguments :
1564 <conns> is the number of pending connections. Depending on the operating
1565 system, it may represent the number of already acknowledged
1566 connections, of non-acknowledged ones, or both.
1567

1568 In order to protect against SYN flood attacks, one solution is to increase
1569 the system's SYN backlog size. Depending on the system, sometimes it is just
1570 tunable via a system parameter, sometimes it is not adjustable at all, and
1571 sometimes the system relies on hints given by the application at the time of
1572 the listen() syscall. By default, HAPROXY passes the frontend's maxconn value
1573 to the listen() syscall. On systems which can make use of this value, it can
1574 sometimes be useful to be able to specify a different value, hence this
1575 backlog parameter.
1576

1577 On Linux 2.4, the parameter is ignored by the system. On Linux 2.6, it is
1578 used as a hint and the system accepts up to the smallest greater power of
1579 two, and never more than some limits (usually 32768).
1580

1581 See also : "maxconn" and the target operating system's tuning guide.
1582

1583 balance <algorithm> [<arguments>]

1584 balance url_param <param> [check_post]

1585 Define the load balancing algorithm to be used in a backend.

1586 May be used in sections : defaults | frontend | listen | backend

1587 defaults yes | no | yes | yes

1588 Arguments :
1589 <algorithm> is the algorithm used to select a server when doing load
1590 balancing. This only applies when no persistence information
1591 is available, or when a connection is redispatched to another
1592 server. <algorithm> may be one of the following :
1593
1594

1595 roundrobin

1596 Each server is used in turns, according to their weights.
1597 This is the smoothest and fairest algorithm when the server's
1598 processing time remains equally distributed. This algorithm
1599 is dynamic, which means that server weights may be adjusted
1600 on the fly for slow starts for instance. It is limited by
1601 design to 4095 active servers per backend. Note that in some
1602 large farms, when a server becomes up after having been down
1603 for a very short time, it may sometimes take a few hundreds
1604 requests for it to be re-integrated into the farm and start
1605 receiving traffic. This is normal, though very rare. It is
1606 indicated here in case you would have the chance to observe
1607 it, so that you don't worry.

1608 static-rr

1609 Each server is used in turns, according to their weights.
1610 This algorithm is as similar to roundrobin except that it is
1611 static, which means that changing a server's weight on the
1612 fly will have no effect. On the other hand, it has no design
1613 limitation on the number of servers, and when a server goes
1614 up, it is always immediately reintroduced into the farm, once
1615 the full map is recomputed. It also uses slightly less CPU to
1616 run (around -1%).
1617

1618 leastconn

1619 The server with the lowest number of connections receives the
1620 connection. Round-robin is performed within groups of servers
1621 of the same load to ensure that all servers will be used. Use
1622 of this algorithm is recommended where very long sessions are
1623 expected, such as LDAP, SQL, TSE, etc... but is not very well
1624 suited for protocols using short sessions such as HTTP. This
1625 algorithm is dynamic, which means that server weights may be
1626 adjusted on the fly for slow starts for instance.

1626 first

1627 The first server with available connection slots receives the
1628 connection. The servers are chosen from the lowest numeric
1629 identifier to the highest (see server parameter "id"), which
1630 defaults to the server's position in the farm. Once a server
1631 reaches its maxconn value, the next server is used. It does
1632 not make sense to use this algorithm without setting maxconn.
1633 The purpose of this algorithm is to always use the smallest
1634 number of servers so that extra servers can be powered off
1635 during non-intensive hours. This algorithm ignores the server
1636 weight, and brings more benefit to long session such as RDP
1637 or IMAP than HTTP, though it can be useful there too. In
1638 order to use this algorithm efficiently, it is recommended
1639 that a cloud controller regularly checks server usage to turn
1640 them off when unused, and regularly checks backend queue to
1641 turn new servers on when the queue inflates. Alternatively,
1642 using "http-check send-state" may inform servers on the load.
1643

1644 source

1645 The source IP address is hashed and divided by the total
1646 weight of the running servers to designate which server will
1647 receive the request. This ensures that the same client IP
1648 address will always reach the same server as long as no
1649 server goes down or up. If the hash result changes due to the
1650 number of running servers changing, many clients will be
1651 directed to a different server. This algorithm is generally
1652 used in TCP mode where no cookie may be inserted. It may also
1653 be used on the Internet to provide a best-effort stickiness
1654 to clients which refuse session cookies. This algorithm is
1655 static by default, which means that changing a server's
1656 weight on the fly will have no effect, but this can be
1657 changed using "hash-type".
1658

1659 uri

1660 This algorithm hashes either the left part of the URI (before
1661 the question mark) or the whole URI (if the "whole" parameter
1662 is present) and divides the hash value by the total weight of
1663 the running servers. The result designates which server will
1664 receive the request. This ensures that the same URI will
1665 always be directed to the same server as long as no server
1666 goes up or down. This is used with proxy caches and
1667 anti-varus proxies in order to maximize the cache hit rate.
1668 Note that this algorithm may only be used in an HTTP backend.
1669 This algorithm is static by default, which means that
1670 changing a server's weight on the fly will have no effect,
1671 but this can be changed using "hash-type".
1672

1673 This algorithm supports two optional parameters "len" and
1674 "depth", both followed by a positive integer number. These
1675 options may be helpful when it is needed to balance servers
1676 based on the beginning of the URI only. The "len" parameter
1677 indicates that the algorithm should only consider that many
1678 characters at the beginning of the URI to compute the hash.
1679 Note that having "len" set to 1 rarely makes sense since most
1680 URIs start with a leading "/".
1681

1682 The "depth" parameter indicates the maximum directory depth
1683 to be used to compute the hash. One level is counted for each
1684 slash in the request. If both parameters are specified, the
1685 evaluation stops when either is reached.
1686

1687 url_param

1688 The URL parameter specified in argument will be looked up in
1689 the query string of each HTTP GET request.

1690 If the modifier "check_post" is used, then an HTTP POST
1691 request entity will be searched for the parameter argument,
1692 when it is not found in a query string after a question mark
1693 ('?') in the URL. The message body will only start to be

analyzed once either the advertised amount of data has been received or the request buffer is full. In the unlikely event that chunked encoding is used, only the first chunk is scanned. Parameter values separated by a chunk boundary, may be randomly balanced if at all. This keyword used to support an optional `<max_wait>` parameter which is now ignored.

If the parameter is found followed by an equal sign ('=') and a value, then the value is hashed and divided by the total weight of the running servers. The result designates which server will receive the request.

This is used to track user identifiers in requests and ensure that a same user ID will always be sent to the same server as long as no server goes up or down. If no value is found or if the parameter is not found, then a round robin algorithm is applied. Note that this algorithm may only be used in an HTTP backend. This algorithm is static by default, which means that changing a server's weight on the fly will have no effect, but this can be changed using "hash-type".

`hdr(<name>)` The HTTP header `<name>` will be looked up in each HTTP request. Just as with the equivalent ACL `'hdr()'` function, the header name in parenthesis is not case sensitive. If the header is absent or if it does not contain any value, the roundrobin algorithm is applied instead.

An optional `'use_domain_only'` parameter is available, for reducing the hash algorithm to the main domain part with some specific headers such as 'Host'. For instance, in the Host value "haproxy.lwt.eu", only "lwt" will be considered.

This algorithm is static by default, which means that changing a server's weight on the fly will have no effect, but this can be changed using "hash-type".

`rdp-cookie`
`rdp-cookie(<name>)`

The RDP cookie `<name>` (or "msthash" if omitted) will be looked up and hashed for each incoming TCP request. Just as with the equivalent ACL `'req_rdp_cookie()'` function, the name is not case-sensitive. This mechanism is useful as a degraded persistence mode, as it makes it possible to always send the same user (or the same session ID) to the same server. If the cookie is not found, the normal roundrobin algorithm is used instead.

Note that for this to work, the frontend must ensure that an RDP cookie is already present in the request buffer. For this you must use `'tcp-request content accept'` rule combined with a `'req_rdp_cookie_cnt'` ACL.

This algorithm is static by default, which means that changing a server's weight on the fly will have no effect, but this can be changed using "hash-type".

See also the `rdp_cookie` pattern fetch function.

`<arguments>` is an optional list of arguments which may be needed by some algorithms. Right now, only `"url_param"` and `"uri"` support an optional argument.

The load balancing algorithm of a backend is set to roundrobin when no other algorithm, mode nor option have been set. The algorithm may only be set once for each backend.

Examples :

```
balance roundrobin
balance url_param userid
balance url_param session_id check_post 64
balance hdr(User-Agent)
balance hdr(host)
balance hdr(host) use_domain_only
```

Note: the following caveats and limitations on using the "check_post" extension with "url_param" must be considered :

- all POST requests are eligible for consideration, because there is no way to determine if the parameters will be found in the body or entity which may contain binary data. Therefore another method may be required to restrict consideration of POST requests that have no URL parameters in the body. (see acl reqideny http_end)
- using a `<max_wait>` value larger than the request buffer size does not make sense and is useless. The buffer size is set at build time, and defaults to 16 Kb.
- Content-Encoding is not supported, the parameter search will probably fail; and load balancing will fall back to Round Robin.
- Expect: 100-continue is not supported, load balancing will fall back to Round Robin.
- Transfer-Encoding (RFC2616 3.6.1) is only supported in the first chunk. If the entire parameter value is not present in the first chunk, the selection of server is undefined (actually, defined by how little actually appeared in the first chunk).

- This feature does not support generation of a 100, 411 or 501 response.

- In some cases, requesting "check_post" MAY attempt to scan the entire contents of a message body. Scanning normally terminates when linear white space or control characters are found, indicating the end of what might be a URL parameter list. This is probably not a concern with SGML type message bodies.

See also : "dispatch", "cookie", "appsession", "transparent", "hash-type" and "http_proxy".

```
bind [<address>]:<port_range> [, ...] [param*]
```

```
bind /<path> [, ...] [param*]
```

Define one or several listening addresses and/or ports in a frontend.

May be used in sections :

	defaults	frontend	listen	backend
	no	yes	yes	no

Arguments :

`<address>` is optional and can be a host name, an IPv4 address, an IPv6 address, or '*'. It designates the address the frontend will listen on. If unset, all IPv4 addresses of the system will be listened on. The same will apply for '*' or the system's special address "0.0.0.0". The IPv6 equivalent is '::'.
Optionally, an address family prefix may be used before the address to force the family regardless of the address format, which can be useful to specify a path to a unix socket with no slash ('/'). Currently supported prefixes are :
- 'ipv4q' -> address is always IPv4
- 'ipv6q' -> address is always IPv6
- 'unixq' -> address is a path to a local unix socket
- 'abnsq' -> address is in abstract namespace (linux only).
Note: since abstract sockets are not "rebindable", they

1821 do not cope well with multi-process mode during
 1822 soft-restart, so it is better to avoid them if
 1823 nbproc is greater than 1. The effect is that if the
 1824 new process fails to start, only one of the old ones
 1825 will be able to rebound to the socket.

1826 - 'fd@<n>' -> use file descriptor <n> inherited from the
 1827 parent. The fd must be bound and may or may not already
 1828 be listening.

1829 Any part of the address string may reference any number of
 1830 environment variables by preceding their name with a dollar
 1831 sign ('\$') and optionally enclosing them with braces ('{}'),
 1832 similarly to what is done in Bourne shell.

1833
 1834 <port_range> is either a unique TCP port, or a port range for which the
 1835 proxy will accept connections for the IP address specified
 1836 above. The port is mandatory for TCP listeners. Note that in
 1837 the case of an IPv6 address, the port is always the number
 1838 after the last colon (':'). A range can either be :
 1839 - a numerical port (ex: '80')
 1840 - a dash-delimited ports range explicitly stating the lower
 1841 and upper bounds (ex: '2000-2100') which are included in
 1842 the range.

1843 Particular care must be taken against port ranges, because
 1844 every <address:port> couple consumes one socket (= a file
 1845 descriptor), so it's easy to consume lots of descriptors
 1846 with a simple range, and to run out of sockets. Also, each
 1847 <address:port> couple must be used only once among all
 1848 instances running on a same system. Please note that binding
 1849 to ports lower than 1024 generally require particular
 1850 privileges to start the program, which are independent of
 1851 the 'uid' parameter.

1852
 1853 <path> is a UNIX socket path beginning with a slash ('/'). This is
 1854 alternative to the TCP listening port. Haproxy will then
 1855 receive UNIX connections on the socket located at this place.
 1856 The path must begin with a slash and by default is absolute.
 1857 It can be relative to the prefix defined by "unix-bind" in
 1858 the global section. Note that the total length of the prefix
 1859 followed by the socket path cannot exceed some system limits
 1860 for UNIX sockets, which commonly are set to 107 characters.

1861
 1862 <param*> is a list of parameters common to all sockets declared on the
 1863 same line. These numerous parameters depend on OS and build
 1864 options and have a complete section dedicated to them. Please
 1865 refer to section 5 for more details.

1866
 1867 It is possible to specify a list of address:port combinations delimited by
 1868 commas. The frontend will then listen on all of these addresses. There is no
 1869 fixed limit to the number of addresses and ports which can be listened on in
 1870 a frontend, as well as there is no limit to the number of "bind" statements
 1871 in a frontend.

1872 Example :

```
1873 listen http_proxy
1874 bind :80,443
1875 bind 10.0.0.1:10080,10.0.0.1:10443
1876 bind /var/run/ssl-frontend.sock user root mode 600 accept-proxy
1877
1878 listen http_https_proxy
1879 bind :80
1880 bind :443 ssl crt /etc/haproxy/site.pem
1881
1882 listen http_https_proxy_explicit
1883 bind ipv6::80
1884
1885
```

```
1886 bind ipv4@public ssl:443 ssl crt /etc/haproxy/site.pem
1887 bind unix@ssl-frontend.sock user root mode 600 accept-proxy
1888
1889 listen external_bind_app1
1890 bind fd@${FD_APP1}
1891
```

1892 Note: regarding Linux's abstract namespace sockets, HAProxy uses the whole
 1893 sun_path length is used for the address length. Some other programs
 1894 such as socat use the string length only by default. Pass the option
 1895 "unix-tightsocketlen=0" to any abstract socket definition in socat to
 1896 make it compatible with HAProxy's.

1897 See also : "source", "option forwardfor", "unix-bind" and the PROXY protocol
 1898 documentation, and section 5 about bind options.

1899 bind-process [all | odd | even | <number 1-64>[<-number 1-64>]] ...
 1900 Limit visibility of an instance to a certain set of processes numbers.

1901 May be used in sections : defaults | frontend | listen | backend
 1902 yes | yes | yes | yes | yes | yes

1903 Arguments : All process will see this instance. This is the default. It
 1904 all may be used to override a default value.

1905 odd This instance will be enabled on processes 1,3,5,...63. This
 1906 option may be combined with other numbers.

1907 even This instance will be enabled on processes 2,4,6,...64. This
 1908 option may be combined with other numbers. Do not use it
 1909 with less than 2 processes otherwise some instances might be
 1910 missing from all processes.

1911 number The instance will be enabled on this process number or range,
 1912 whose values must all be between 1 and 32 or 64 depending on
 1913 the machine's word size. If a proxy is bound to process
 1914 numbers greater than the configured global.nbproc, it will
 1915 either be forced to process #1 if a single process was
 1916 specified, or to all processes otherwise.

1917 This keyword limits binding of certain instances to certain processes. This
 1918 is useful in order not to have too many processes listening to the same
 1919 ports. For instance, on a dual-core machine, it might make sense to set
 1920 'nbproc 2' in the global section, then distributes the listeners among 'odd'
 1921 and 'even' instances.

1922 At the moment, it is not possible to reference more than 32 or 64 processes
 1923 using this keyword, but this should be more than enough for most setups.
 1924 Please note that 'all' really means all processes regardless of the machine's
 1925 word size, and is not limited to the first 32 or 64.

1926 Each "bind" line may further be limited to a subset of the proxy's processes,
 1927 please consult the "process" bind keyword in section 5.1.

1928 When a frontend has no explicit "bind-process" line, it tries to bind to all
 1929 the processes referenced by its "bind" lines. That means that frontends can
 1930 easily adapt to their listeners' processes.

1931 If some backends are referenced by frontends bound to other processes, the
 1932 backend automatically inherits the frontend's processes.

1933 Example :
 1934 listen app_ip1
 1935 bind 10.0.0.1:80
 1936 bind-process odd

1940

```
1951 Listen app_ip2
1952 bind 10.0.0.2:80
1953 bind-process even
1954
1955 listen management
1956 bind 10.0.0.3:80
1957 bind-process 1 2 3 4
1958
1959 listen management
1960 bind 10.0.0.4:80
1961 bind-process 1-4
1962
1963 See also : "nproc" in global section, and "process" in section 5.1.
1964
1965
1966 block { if | unless } <condition>
1967 Block a layer 7 request if/unless a condition is matched
1968 May be used in sections : defaults | frontend | listen | backend
1969 no | yes | yes | yes | yes
1970
1971 The HTTP request will be blocked very early in the layer 7 processing
1972 if/unless <condition> is matched. A 403 error will be returned if the request
1973 is blocked. The condition has to reference ACLs (see section 7). This is
1974 typically used to deny access to certain sensitive resources if some
1975 conditions are met or not met. There is no fixed limit to the number of
1976 "block" statements per instance.
1977
1978 Example:
1979 acl invalid_src src 0.0.0.0/7 224.0.0.0/3
1980 acl invalid_src src port 0:1023
1981 acl local_dst hdr(host) -i localhost
1982 block if invalid_src || local_dst
1983
1984 See section 7 about ACL usage.
1985
1986
1987 capture cookie <name> len <length>
1988 Capture and log a cookie in the request and in the response.
1989 May be used in sections : defaults | frontend | listen | backend
1990 no | yes | yes | yes | no
1991
1992 Arguments :
1993 <name> is the beginning of the name of the cookie to capture. In order
1994 to match the exact name, simply suffix the name with an equal
1995 sign ('='). The full name will appear in the logs, which is
1996 useful with application servers which adjust both the cookie name
1997 and value (eg: SESSIONXXXXX).
1998
1999 <length> is the maximum number of characters to report in the logs, which
2000 include the cookie name, the equal sign and the value, all in the
2001 standard "name=value" form. The string will be truncated on the
2002 right if it exceeds <length>.
2003
2004 Only the first cookie is captured. Both the "cookie" request headers and the
2005 "set-cookie" response headers are monitored. This is particularly useful to
2006 check for application bugs causing session crossing or stealing between
2007 users, because generally the user's cookies can only change on a login page.
2008
2009 When the cookie was not presented by the client, the associated log column
2010 will report "-". When a request does not cause a cookie to be assigned by the
2011 server, a "." is reported in the response column.
2012
2013 The capture is performed in the frontend only because it is necessary that
2014 the log format does not change for a given frontend depending on the
2015 backends. This may change in the future. Note that there can be only one
2016 "capture cookie" statement in a frontend. The maximum capture length is set
```

```
2016 by the global "tune.http.cookieien" setting and defaults to 63 characters. It
2017 is not possible to specify a capture in a "defaults" section.
2018
2019 Example:
2020 capture cookie ASPSESSION len 32
2021
2022 See also : "capture request header", "capture response header" as well as
2023 section 8 about logging.
2024
2025
2026 capture request header <name> len <length>
2027 Capture and log the last occurrence of the specified request header.
2028 May be used in sections : defaults | frontend | listen | backend
2029 no | yes | yes | yes | no
2030
2031 Arguments :
2032 <name> is the name of the header to capture. The header names are not
2033 case-sensitive, but it is a common practice to write them as they
2034 appear in the requests, with the first letter of each word in
2035 upper case. The header name will not appear in the logs, only the
2036 value is reported, but the position in the logs is respected.
2037
2038 <length> is the maximum number of characters to extract from the value and
2039 report in the logs. The string will be truncated on the right if
2040 it exceeds <length>.
2041
2042 The complete value of the last occurrence of the header is captured. The
2043 value will be added to the logs between braces ('{}'). If multiple headers
2044 are captured, they will be delimited by a vertical bar '|' and will appear
2045 in the same order they were declared in the configuration. Non-existent
2046 headers will be logged just as an empty string. Common uses for request
2047 header captures include the "Host" field in virtual hosting environments, the
2048 "Content-length" when uploads are supported, "User-agent" to quickly
2049 differentiate between real users and robots, and "X-Forwarded-For" in proxied
2050 environments to find where the request came from.
2051
2052 Note that when capturing headers such as "User-agent", some spaces may be
2053 logged, making the log analysis more difficult. Thus be careful about what
2054 you log if you know your log parser is not smart enough to rely on the
2055 braces.
2056
2057 There is no limit to the number of captured request headers nor to their
2058 length, though it is wise to keep them low to limit memory usage per session.
2059 In order to keep log format consistent for a same frontend, header captures
2060 can only be declared in a frontend. It is not possible to specify a capture
2061 in a "defaults" section.
2062
2063 Example:
2064 capture request header Host len 15
2065 capture request header X-Forwarded-For len 15
2066 capture request header Referer len 15
2067
2068 See also : "capture cookie", "capture response header" as well as section 8
2069 about logging.
2070
2071
2072 capture response header <name> len <length>
2073 Capture and log the last occurrence of the specified response header.
2074 May be used in sections : defaults | frontend | listen | backend
2075 no | yes | yes | yes | no
2076
2077 Arguments :
2078 <name> is the name of the header to capture. The header names are not
2079 case-sensitive, but it is a common practice to write them as they
2080 appear in the response, with the first letter of each word in
2081 upper case. The header name will not appear in the logs, only the
2082 value is reported, but the position in the logs is respected.
```


<length> is the maximum number of characters to extract from the value and report in the logs. The string will be truncated on the right if it exceeds <length>.

The complete value of the last occurrence of the header is captured. The result will be added to the logs between braces ('{}') after the captured request headers. If multiple headers are captured, they will be delimited by a vertical bar ('|') and will appear in the same order they were declared in the configuration. Non-existent headers will be logged just as an empty string. Common uses for response header captures include the "Content-length" header which indicates how many bytes are expected to be returned, the "Location" header to track redirections.

There is no limit to the number of captured response headers nor to their length, though it is wise to keep them low to limit memory usage per session. In order to keep log format consistent for a same frontend, header captures can only be declared in a frontend. It is not possible to specify a capture in a "defaults" section.

Example:
capture response header Content-length len 9
capture response header Location len 15

See also : "capture cookie", "capture request header" as well as section 8 about logging.

clitimeout <timeout> (deprecated)

Set the maximum inactivity time on the client side.

May be used in sections : defaults | frontend | listen | backend
yes | yes | yes | yes | no

Arguments :
<timeout> is the timeout value is specified in milliseconds by default, but can be in any other unit if the number is suffixed by the unit, as explained at the top of this document.

The inactivity timeout applies when the client is expected to acknowledge or send data. In HTTP mode, this timeout is particularly important to consider during the first phase, when the client sends the request, and during the response while it is reading data sent by the server. The value is specified in milliseconds by default, but can be in any other unit if the number is suffixed by the unit, as specified at the top of this document. In TCP mode (and to a lesser extent, in HTTP mode), it is highly recommended that the client timeout remains equal to the server timeout in order to avoid complex situations to debug. It is a good practice to cover one or several TCP packet losses by specifying timeouts that are slightly above multiples of 3 seconds (eg: 4 or 5 seconds).

This parameter is specific to frontends, but can be specified once for all in "defaults" sections. This is in fact one of the easiest solutions not to forget about it. An unspecified timeout results in an infinite timeout, which is not recommended. Such a usage is accepted and works but reports a warning during startup because it may results in accumulation of expired sessions in the system if the system's timeouts are not configured either.

This parameter is provided for compatibility but is currently deprecated. Please use "timeout client" instead.

See also : "timeout client", "timeout http-request", "timeout server", and "srvtimeout".

compression algo <algorithm> ...
compression type <mime type> ...
compression offload

Enable HTTP compression.

May be used in sections : defaults | frontend | listen | backend
yes | yes | yes | yes

Arguments :
algo is followed by the list of supported compression algorithms.
type is followed by the list of MIME types that will be compressed.
offload makes haproxy work as a compression offloader only (see notes).

The currently supported algorithms are :
identity this is mostly for debugging, and it was useful for developing the compression feature. Identity does not apply any change on data.

gzip applies gzip compression. This setting is only available when support for zlib was built in.

deflate same as gzip, but with deflate algorithm and zlib format.

Note that this algorithm has ambiguous support on many browsers and no support at all from recent ones. It is strongly recommended not to use it for anything else than experimentation. This setting is only available when support for zlib was built in.

Compression will be activated depending on the Accept-Encoding request header. With identity, it does not take care of that header. If backend servers support HTTP compression, these directives will be no-op: haproxy will see the compressed response and will not compress again. If backend servers do not support HTTP compression and there is Accept-Encoding header in request, haproxy will compress the matching response.

The "offload" setting makes haproxy remove the Accept-Encoding header to prevent backend servers from compressing responses. It is strongly recommended not to do this because this means that all the compression work will be done on the single point where haproxy is located. However in some deployment scenarios, haproxy may be installed in front of a buggy gateway with broken HTTP compression implementation which can't be turned off. In that case haproxy can be used to prevent that gateway from emitting invalid payloads. In this case, simply removing the header in the configuration does not work because it applies before the header is parsed, so that prevents haproxy from compressing. The "offload" setting should then be used for such scenarios. Note: for now, the "offload" setting is ignored when set in a defaults section.

Compression is disabled when:

- * the request does not advertise a supported compression algorithm in the "Accept-Encoding" header

- * the response message is not HTTP/1.1

- * HTTP status code is not 200

- * response header "Transfer-Encoding" contains "chunked" (Temporary Workaround)

- * response contain neither a "Content-Length" header nor a

- "Transfer-Encoding" whose last value is "chunked"

- * response contains a "Content-Type" header whose first value starts with "multipart"

- * the response contains the "no-transform" value in the "Cache-control" header

- * User-Agent matches "Mozilla/4" unless it is MSIE 6 with XP SP2, or MSIE 7

- * The response contains a "Content-Encoding" header, indicating that the response is already compressed (see compression offload)

Note: The compression does not rewrite Etag headers, and does not emit the Warning header.

2211 Examples :
 2212 compression algo gzip
 2213 compression type text/html text/plain
 2214
 2215 timeout <timeout> (deprecated)
 2216 Set the maximum time to wait for a connection attempt to a server to succeed.
 2217 May be used in sections : defaults | frontend | listen | backend
 2218 yes | no | yes | yes
 2219 Arguments :
 2220 <timeout> is the timeout value is specified in milliseconds by default, but
 2221 can be in any other unit if the number is suffixed by the unit,
 2222 as explained at the top of this document.
 2223
 2224 If the server is located on the same LAN as haproxy, the connection should be
 2225 immediate (less than a few milliseconds). Anyway, it is a good practice to
 2226 cover one or several TCP packet losses by specifying timeouts that are
 2227 slightly above multiples of 3 seconds (eg: 4 or 5 seconds). By default, the
 2228 connect timeout also presets the queue timeout to the same value if this one
 2229 has not been specified. Historically, the timeout was also used to set the
 2230 tarpit timeout in a listen section, which is not possible in a pure frontend.
 2231
 2232 This parameter is specific to backends, but can be specified once for all in
 2233 "defaults" sections. This is in fact one of the easiest solutions not to
 2234 forget about it. An unspecified timeout results in an infinite timeout, which
 2235 is not recommended. Such a usage is accepted and works but reports a warning
 2236 during startup because it may results in accumulation of failed sessions in
 2237 the system if the system's timeouts are not configured either.
 2238
 2239 This parameter is provided for backwards compatibility but is currently
 2240 deprecated. Please use "timeout connect", "timeout queue" or "timeout tarpit"
 2241 instead.
 2242
 2243 See also : "timeout connect", "timeout queue", "timeout tarpit",
 2244 "timeout server", "contimeout".
 2245
 2246 cookie <name> [rewrite | insert | prefix] [indirect] [nocache]
 2247 [postonly] [preserve] [httponly] [secure]
 2248 [domain <domain>] * [maxidle <idle>] [maxlife <life>]
 2249 Enable cookie-based persistence in a backend.
 2250 May be used in sections : defaults | frontend | listen | backend
 2251 yes | no | yes | yes
 2252 Arguments :
 2253 <name> is the name of the cookie which will be monitored, modified or
 2254 inserted in order to bring persistence. This cookie is sent to
 2255 the client via a "Set-Cookie" header in the response, and is
 2256 brought back by the client in a "Cookie" header in all requests.
 2257 Special care should be taken to choose a name which does not
 2258 conflict with any likely application cookie. Also, if the same
 2259 backends are subject to be used by the same clients (eg:
 2260 HTTP/HTTPS), care should be taken to use different cookie names
 2261 between all backends if persistence between them is not desired.
 2262
 2263 rewrite
 2264 This keyword indicates that the cookie will be provided by the
 2265 server and that haproxy will have to modify its value to set the
 2266 server's identifier in it. This mode is handy when the management
 2267 of complex combinations of "Set-cookie" and "Cache-control"
 2268 headers is left to the application. The application can then
 2269 decide whether or not it is appropriate to emit a persistence
 2270 cookie. Since all responses should be monitored, this mode only
 2271 works in HTTP close mode. Unless the application behaviour is
 2272 very complex and/or broken, it is advised not to start with this
 2273 mode for new deployments. This keyword is incompatible with
 2274 "insert" and "prefix".
 2275

2276 insert This keyword indicates that the persistence cookie will have to
 2277 be inserted by haproxy in server responses if the client did not
 2278 already have a cookie that would have permitted it to access this
 2279 server. When used without the "preserve" option, if the server
 2280 emits a cookie with the same name, it will be remove before
 2281 processing. For this reason, this mode can be used to upgrade
 2282 existing configurations running in the "rewrite" mode. The cookie
 2283 will only be a session cookie and will not be stored on the
 2284 client's disk. By default, unless the "indirect" option is added,
 2285 the server will see the cookies emitted by the client. Due to
 2286 caching effects, it is generally wise to add the "nocache" or
 2287 "postonly" keywords (see below). The "insert" keyword is not
 2288 compatible with "rewrite" and "prefix".
 2289
 2290 prefix
 2291 This keyword indicates that instead of relying on a dedicated
 2292 cookie for the persistence, an existing one will be completed.
 2293 This may be needed in some specific environments where the client
 2294 does not support more than one single cookie and the application
 2295 already needs it. In this case, whenever the server sets a cookie
 2296 named <name>, it will be prefixed with the server's identifier
 2297 and a delimiter. The prefix will be removed from all client
 2298 requests so that the server still finds the cookie it emitted.
 2299 Since all requests and responses are subject to being modified,
 2300 this mode requires the HTTP close mode. The "prefix" keyword is
 2301 not compatible with "rewrite" and "insert". Note: it is highly
 2302 recommended not to use "indirect" with "prefix", otherwise server
 2303 cookie updates would not be sent to clients.
 2304
 2305 indirect
 2306 When this option is specified, no cookie will be emitted to a
 2307 client which already has a valid one for the server which has
 2308 processed the request. If the server sets such a cookie itself,
 2309 it will be removed, unless the "preserve" option is also set. In
 2310 "insert" mode, this will additionally remove cookies from the
 2311 requests transmitted to the server, making the persistence
 2312 mechanism totally transparent from an application point of view.
 2313 Note: it is highly recommended not to use "indirect" with
 2314 "prefix", otherwise server cookie updates would not be sent to
 2315 clients.
 2316
 2317 nocache
 2318 This option is recommended in conjunction with the insert mode
 2319 when there is a cache between the client and HAProxy, as it
 2320 ensures that a cacheable response will be tagged non-cacheable if
 2321 a cookie needs to be inserted. This is important because if all
 2322 persistence cookies are added on a cacheable home page for
 2323 instance, then all customers will then fetch the page from an
 2324 outer cache and will all share the same persistence cookie,
 2325 leading to one server receiving much more traffic than others.
 2326 See also the "insert" and "postonly" options.
 2327
 2328 postonly
 2329 This option ensures that cookie insertion will only be performed
 2330 on responses to POST requests. It is an alternative to the
 2331 "nocache" option, because POST responses are not cacheable, so
 2332 this ensures that the persistence cookie will never get cached.
 2333 Since most sites do not need any sort of persistence before the
 2334 first POST which generally is a login request, this is a very
 2335 efficient method to optimize caching without risking to find a
 2336 persistence cookie in the cache.
 2337 See also the "insert" and "nocache" options.
 2338
 2339 preserve
 2340 This option may only be used with "insert" and/or "indirect". It
 2341 allows the server to emit the persistence cookie itself. In this
 2342 case, if a cookie is found in the response, haproxy will leave it
 2343 untouched. This is useful in order to end persistence after a
 2344 logout request for instance. For this, the server just has to

emit a cookie with an invalid value (eg: empty) or with a date in the past. By combining this mechanism with the "disable-on-404" check option, it is possible to perform a completely graceful shutdown because users will definitely leave the server after they logout.

httponly This option tells haproxy to add an "HttpOnly" cookie attribute when a cookie is inserted. This attribute is used so that a user agent doesn't share the cookie with non-HTTP components. Please check RFC6265 for more information on this attribute.

secure This option tells haproxy to add a "Secure" cookie attribute when a cookie is inserted. This attribute is used so that a user agent never emits this cookie over non-secure channels, which means that a cookie learned with this flag will be presented only over SSL/TLS connections. Please check RFC6265 for more information on this attribute.

domain This option allows to specify the domain at which a cookie is inserted. It requires exactly one parameter: a valid domain name. If the domain begins with a dot, the browser is allowed to use it for any host ending with that name. It is also possible to specify several domain names by invoking this option multiple times. Some browsers might have small limits on the number of domains, so be careful when doing that. For the record, sending 10 domains to MSIE 6 or Firefox 2 works as expected.

maxidle This option allows inserted cookies to be ignored after some idle time. It only works with insert-mode cookies. When a cookie is sent to the client, the date this cookie was emitted is sent too. Upon further presentations of this cookie, if the date is older than the delay indicated by the parameter (in seconds), it will be ignored. Otherwise, it will be refreshed if needed when the response is sent to the client. This is particularly useful to prevent users who never close their browsers from remaining for too long on the same server (eg: after a farm size change). When this option is set and a cookie has no date, it is always accepted, but gets refreshed in the response. This maintains the ability for admins to access their sites. Cookies that have a date in the future further than 24 hours are ignored. Doing so lets admins fix timezone issues without risking kicking users off the site.

maxlife This option allows inserted cookies to be ignored after some life time, whether they're in use or not. It only works with insert mode cookies. When a cookie is first sent to the client, the date this cookie was emitted is sent too. Upon further presentations of this cookie, if the date is older than the delay indicated by the parameter (in seconds), it will be ignored. If the cookie in the request has no date, it is accepted and a date will be set. Cookies that have a date in the future further than 24 hours are ignored. Doing so lets admins fix timezone issues without risking kicking users off the site. Contrary to maxidle, this value is not refreshed, only the first visit date counts. Both maxidle and maxlife may be used at the time. This is particularly useful to prevent users who never close their browsers from remaining for too long on the same server (eg: after a farm size change). This is stronger than the maxidle method in that it forces a redispatch after some absolute delay.

There can be only one persistence cookie per HTTP backend, and it can be declared in a defaults section. The value of the cookie will be the value indicated after the "cookie" keyword in a "server" statement. If no cookie is declared for a given server, the cookie is not set.

Examples :
cookie JSESSIONID prefix
cookie SRV insert indirect nocache
cookie SRV insert postonly indirect
cookie SRV insert indirect nocache maxidle 30m maxlife 8h

See also : "appsession", "balance source", "capture cookie", "server" and "ignore-persist".

default-server [param*]
Change default options for a server in a backend
May be used in sections : defaults | frontend | listen | backend
yes | no | yes | yes
Arguments:
<param> is a list of parameters for this server. The "default-server" keyword accepts an important number of options and has a complete section dedicated to it. Please refer to section 5 for more details.

Example :
default-server inter 1000 weight 13

See also: "server" and section 5 about server options

default_backend <backend>
Specify the backend to use when no "use_backend" rule has been matched.
May be used in sections : defaults | frontend | listen | backend
yes | yes | yes | no

Arguments :
<backend> is the name of the backend to use.

When doing content-switching between frontend and backends using the "use_backend" keyword, it is often useful to indicate which backend will be used when no rule has matched. It generally is the dynamic backend which will catch all undetermined requests.

Example :
use_backend dynamic if url_dyn
use_backend static if url_css url_img extension_img
default_backend dynamic

See also : "use_backend", "reqsetbe", "reqsetbe"

description <string>
Describe a listen, frontend or backend.
May be used in sections : defaults | frontend | listen | backend
no | yes | yes | yes
Arguments : string

Allows to add a sentence to describe the related object in the HAProxy HTML stats page. The description will be printed on the right of the object name it describes.
No need to backslash spaces in the <string> arguments.

disabled
Disable a proxy, frontend or backend.
May be used in sections : defaults | frontend | listen | backend
yes | yes | yes | yes
Arguments : none

The "disabled" keyword is used to disable an instance, mainly in order to liberate a listening port or to temporarily disable a service. The instance will still be created and its configuration will be checked, but it will be created in the "stopped" state and will appear as such in the statistics. It will not receive any traffic nor will it send any health-checks or logs. It is possible to disable many instances at once by adding the "disabled" keyword in a "defaults" section.

See also : "enabled"

```
dispatch <address>:<port>
```

Set a default server address

May be used in sections : defaults | frontend | listen | backend
no | no | yes | yes

Arguments :

<address> is the IPv4 address of the default server. Alternatively, a resolvable hostname is supported, but this name will be resolved during start-up.

<ports> is a mandatory port specification. All connections will be sent to this port, and it is not permitted to use port offsets as is possible with normal servers.

The "dispatch" keyword designates a default server for use when no other server can take the connection. In the past it was used to forward non persistent connections to an auxiliary load balancer. Due to its simple syntax, it has also been used for simple TCP relays. It is recommended not to use it for more clarity, and to use the "server" directive instead.

See also : "server"

enabled

Enable a proxy, frontend or backend.

May be used in sections : defaults | frontend | listen | backend
yes | yes | yes | yes

Arguments : none

The "enabled" keyword is used to explicitly enable an instance, when the defaults has been set to "disabled". This is very rarely used.

See also : "disabled"

```
errorfile <code> <file>
```

Return a file contents instead of errors generated by HAProxy

May be used in sections : defaults | frontend | listen | backend
yes | yes | yes | yes

Arguments :

<code> is the HTTP status code. Currently, HAProxy is capable of generating codes 200, 400, 403, 408, 500, 502, 503, and 504.

<file> designates a file containing the full HTTP response. It is recommended to follow the common practice of appending ".http" to the filename so that people do not confuse the response with HTML error pages, and to use absolute paths, since files are read before any chroot is performed.

It is important to understand that this keyword is not meant to rewrite errors returned by the server, but errors detected and returned by HAProxy. This is why the list of supported errors is limited to a small set.

Code 200 is emitted in response to requests matching a "monitor-uri" rule.

The files are returned verbatim on the TCP socket. This allows any trick such as redirections to another URL or site, as well as tricks to clean cookies, force enable or disable caching, etc... The package provides default error files returning the same contents as default errors.

The files should not exceed the configured buffer size (BUFSIZE), which generally is 8 or 16 kB, otherwise they will be truncated. It is also wise not to put any reference to local contents (eg: images) in order to avoid loops between the client and HAProxy when all servers are down, causing an error to be returned instead of an image. For better HTTP compliance, it is recommended that all header lines end with CR-LF and not LF alone.

The files are read at the same time as the configuration and kept in memory. For this reason, the errors continue to be returned even when the process is chrooted, and no file change is considered while the process is running. A simple method for developing those files consists in associating them to the 403 status code and interrogating a blocked URL.

See also : "errorloc", "errorloc302", "errorloc303"

Example :

```
errorfile 400 /etc/haproxy/errorfiles/400badreq.http
errorfile 408 /dev/null # workaround Chrome pre-connect bug
errorfile 403 /etc/haproxy/errorfiles/403forbid.http
errorfile 503 /etc/haproxy/errorfiles/503sorry.http
```

```
errorloc <code> <url>
```

```
errorloc302 <code> <url>
```

Return an HTTP redirection to a URL instead of errors generated by HAProxy

May be used in sections : defaults | frontend | listen | backend
yes | yes | yes | yes

Arguments :

<code> is the HTTP status code. Currently, HAProxy is capable of generating codes 200, 400, 403, 408, 500, 502, 503, and 504.

<url> it is the exact contents of the "Location" header. It may contain either a relative URI to an error page hosted on the same site, or an absolute URI designating an error page on another site. Special care should be given to relative URIs to avoid redirect loops if the URI itself may generate the same error (eg: 500).

It is important to understand that this keyword is not meant to rewrite errors returned by the server, but errors detected and returned by HAProxy. This is why the list of supported errors is limited to a small set.

Code 200 is emitted in response to requests matching a "monitor-uri" rule.

Note that both keyword return the HTTP 302 status code, which tells the client to fetch the designated URL using the same HTTP method. This can be quite problematic in case of non-GET methods such as POST, because the URL sent to the client might not be allowed for something other than GET. To workaround this problem, please use "errorloc303" which send the HTTP 303 status code, indicating to the client that the URL must be fetched with a GET request.

See also : "errorfile", "errorloc303"

```
errorloc303 <code> <url>
```

Return an HTTP redirection to a URL instead of errors generated by HAProxy

May be used in sections : defaults | frontend | listen | backend
yes | yes | yes | yes

Arguments :

Code 200 is emitted in response to requests matching a "monitor-uri" rule.

<code> is the HTTP status code. Currently, HAProxy is capable of generating codes 400, 403, 408, 500, 502, 503, and 504.

<url> it is the exact contents of the "Location" header. It may contain either a relative URI to an error page hosted on the same site, or an absolute URI designating an error page on another site. Special care should be given to relative URIs to avoid redirect loops if the URI itself may generate the same error (eg: 500).

It is important to understand that this keyword is not meant to rewrite errors returned by the server, but errors detected and returned by HAProxy. This is why the list of supported errors is limited to a small set.

Code 200 is emitted in response to requests matching a "monitor-uri" rule.

Note that both keyword return the HTTP 303 status code, which tells the client to fetch the designated URL using the same HTTP GET method. This solves the usual problems associated with "errorloc" and the 302 code. It is possible that some very old browsers designed before HTTP/1.1 do not support it, but no such problem has been reported till now.

See also : "errorfile", "errorloc", "errorloc302"

force-persist { if | unless } <condition>

Declare a condition to force persistence on down servers

May be used in sections : defaults | frontend | listen | backend
no | yes | yes | yes

By default, requests are not dispatched to down servers. It is possible to force this using "option persist", but it is unconditional and redispaches to a valid server if "option redispatch" is set. That leaves with very little possibilities to force some requests to reach a server which is artificially marked down for maintenance operations.

The "force-persist" statement allows one to declare various ACL-based conditions which, when met, will cause a request to ignore the down status of a server and still try to connect to it. That makes it possible to start a server, still replying an error to the health checks, and run a specially configured browser to test the service. Among the handy methods, one could use a specific source IP address, or a specific cookie. The cookie also has the advantage that it can easily be added/removed on the browser from a test page. Once the service is validated, it is then possible to open the service to the world by returning a valid response to health checks.

The forced persistence is enabled when an "if" condition is met, or unless an "unless" condition is met. The final redispatch is always disabled when this is used.

See also : "option redispatch", "ignore-persist", "persist", and section 7 about ACL usage.

fullconn <conns>

Specify at what backend load the servers will reach their maxconn

May be used in sections : defaults | frontend | listen | backend
yes | no | yes | yes

Arguments :

<conns> is the number of connections on the backend which will make the servers use the maximal number of connections.

When a server has a "maxconn" parameter specified, it means that its number of concurrent connections will never go higher. Additionally, if it has a "minconn" parameter, it indicates a dynamic limit following the backend's load. The server will then always accept at least <minconn> connections,

never more than <maxconn>, and the limit will be on the ramp between both values when the backend has less than <conns> concurrent connections. This makes it possible to limit the load on the servers during normal loads, but push it further for important loads without overloading the servers during exceptional loads.

Since it's hard to get this value right, haproxy automatically sets it to 10% of the sum of the maxconns of all frontends that may branch to this backend (based on "use_backend" and "default_backend" rules). That way it's safe to leave it unset. However, "use_backend" involving dynamic names are not counted since there is no way to know if they could match or not.

Example :

The servers will accept between 100 and 1000 concurrent connections each
and the maximum of 1000 will be reached when the backend reaches 10000
connections.
backend dynamic

fullconn 10000
server svr1 dyn1:80 minconn 100 maxconn 1000
server svr2 dyn2:80 minconn 100 maxconn 1000

See also : "maxconn", "server"

grace <time>

Maintain a proxy operational for some time after a soft stop

May be used in sections : defaults | frontend | listen | backend
yes | yes | yes | yes

Arguments :

<time> is the time (by default in milliseconds) for which the instance will remain operational with the frontend sockets still listening when a soft-stop is received via the SIGUSR1 signal.

This may be used to ensure that the services disappear in a certain order. This was designed so that frontends which are dedicated to monitoring by an external equipment fail immediately while other ones remain up for the time needed by the equipment to detect the failure.

Note that currently, there is very little benefit in using this parameter, and it may in fact complicate the soft-reconfiguration process more than simplify it.

hash-type <method> <function> <modifier>

Specify a method to use for mapping hashes to servers

May be used in sections : defaults | frontend | listen | backend
yes | no | yes | yes

Arguments :
<method> is the method used to select a server from the hash computed by the <function> :

map-based the hash table is a static array containing all alive servers. The hashes will be very smooth, will consider weights, but will be static in that weight changes while a server is up will be ignored. This means that there will be no slow start. Also, since a server is selected by its position in the array, most mappings are changed when the server count changes. This means that when a server goes up or down, or when a server is added to a farm, most connections will be redistributed to different servers. This can be inconvenient with caches for instance.

consistent the hash table is a tree filled with many occurrences of each server. The hash key is looked up in the tree and the closest server is chosen. This hash is dynamic, it supports changing

weights while the servers are up, so it is compatible with the slow start feature. It has the advantage that when a server goes up or down, only its associations are moved. When a server is added to the farm, only a few part of the mappings are redistributed, making it an ideal method for caches. However, due to its principle, the distribution will never be very smooth and it may sometimes be necessary to adjust a server's weight or its ID to get a more balanced distribution. In order to get the same distribution on multiple load balancers, it is important that all servers have the exact same IDs. Note: consistent hash uses sdbm and avalanche if no hash function is specified.

<function> is the hash function to be used :

sdbm this function was created initially for sdbm (a public-domain reimplementation of ndbm) database library. It was found to do well in scrambling bits, causing better distribution of the keys and fewer splits. It also happens to be a good general hashing function with good distribution, unless the total server weight is a multiple of 64, in which case applying the avalanche modifier may help.

djb2 this function was first proposed by Dan Bernstein many years ago on comp.lang.c. Studies have shown that for certain workload this function provides a better distribution than sdbm. It generally works well with text-based inputs though it can perform extremely poorly with numeric-only input or when the total server weight is a multiple of 33, unless the avalanche modifier is also used.

wt6 this function was designed for haproxy while testing other functions in the past. It is not as smooth as the other ones, but is much less sensible to the input data set or to the number of servers. It can make sense as an alternative to sdbm+avalanche or djb2+avalanche for consistent hashing or when hashing on numeric data such as a source IP address or a visitor identifier in a URL parameter.

<modifier> indicates an optional method applied after hashing the key :

avalanche This directive indicates that the result from the hash function above should not be used in its raw form but that a 4-byte full avalanche hash must be applied first. The purpose of this step is to mix the resulting bits from the previous hash in order to avoid any undesired effect when the input contains some limited values or when the number of servers is a multiple of one of the hash's components (64 for SPBM, 33 for DJB2). Enabling avalanche tends to make the result less predictable, but it's also not as smooth as when using the original function. Some testing might be needed with some workloads. This hash is one of the many proposed by Bob Jenkins.

The default hash type is "map-based" and is recommended for most usages. The default function is "sdbm", the selection of a function should be based on the range of the values being hashed.

See also : "balance", "server"

http-check disable-on-404

Enable a maintenance mode upon HTTP/404 response to health-checks

May be used in sections : defaults | frontend | listen | backend
yes | no | yes | yes

Arguments : none

When this option is set, a server which returns an HTTP code 404 will be excluded from further load-balancing, but will still receive persistent connections. This provides a very convenient method for Web administrators to perform a graceful shutdown of their servers. It is also important to note that a server which is detected as failed while it was in this mode will not generate an alert, just a notice. If the server responds 2xx or 3xx again, it will immediately be reinserted into the farm. The status on the stats page reports "NOLB" for a server in this mode. It is important to note that this option only works in conjunction with the "httpchk" option. If this option is used with "http-check expect", then it has precedence over it so that 404 responses will still be considered as soft-stop.

See also : "option httpchk", "http-check expect"

http-check expect [!] <match> <pattern>

Make HTTP health checks consider response contents or specific status codes

May be used in sections : defaults | frontend | listen | backend
defaults yes | no | yes | yes

Arguments : <match> is a keyword indicating how to look for a specific pattern in the response. The keyword may be one of "status", "rstatus", "string", or "rstring". The keyword may be preceded by an exclamation mark ("!") to negate the match. Spaces are allowed between the exclamation mark and the keyword. See below for more details on the supported keywords.

<pattern> is the pattern to look for. It may be a string or a regular expression. If the pattern contains spaces, they must be escaped with the usual backslash ('\').

By default, "option httpchk" considers that response statuses 2xx and 3xx are valid, and that others are invalid. When "http-check expect" is used, it defines what is considered valid or invalid. Only one "http-check" statement is supported in a backend. If a server fails to respond or times out, the check obviously fails. The available matches are :

status <string> : test the exact string match for the HTTP status code.

A health check response will be considered valid if the response's status code is exactly this string. If the "status" keyword is prefixed with "!", then the response will be considered invalid if the status code matches.

rstatus <regex> : test a regular expression for the HTTP status code.

A health check response will be considered valid if the response's status code matches the expression. If the "rstatus" keyword is prefixed with "!", then the response will be considered invalid if the status code matches. This is mostly used to check for multiple codes.

string <string> : test the exact string match in the HTTP response body.

A health check response will be considered valid if the response's body contains this exact string. If the "string" keyword is prefixed with "!", then the response will be considered invalid if the body contains this string. This can be used to look for a mandatory word at the end of a dynamic page, or to detect a failure when a specific error appears on the check page (eg: a stack trace).

rstring <regex> : test a regular expression on the HTTP response body.

A health check response will be considered valid if the response's body matches this expression. If the "rstring" keyword is prefixed with "!", then the response will be

considered invalid if the body matches the expression.
This can be used to look for a mandatory word at the end of a dynamic page, or to detect a failure when a specific error appears on the check page (eg: a stack trace).

It is important to note that the responses will be limited to a certain size defined by the global "tune.chksize" option, which defaults to 16384 bytes. Thus, too large responses may not contain the mandatory pattern when using "string" or "rstring". If a large response is absolutely required, it is possible to change the default max size by setting the global variable. However, it is worth keeping in mind that parsing very large responses can waste some CPU cycles, especially when regular expressions are used, and that it is always better to focus the checks on smaller resources.

Also "http-check expect" doesn't support HTTP keep-alive. Keep in mind that it will automatically append a "Connection: close" header, meaning that this header should not be present in the request provided by "option httpchk".

Last, if "http-check expect" is combined with "http-check disable-on-404", then this last one has precedence when the server responds with 404.

Examples :

```
# only accept status 200 as valid
http-check expect status 200
```

```
# consider SQL errors as errors
```

```
http-check expect ! string SQL\ Error
```

```
# consider status 5xx only as errors
```

```
http-check expect ! rstatus ^5
```

```
# check that we have a correct hexadecimal tag before /html
```

```
http-check expect rstring <!--tag:[0-9a-f]*</html>
```

See also : "option httpchk", "http-check disable-on-404"

http-check send-state

Enable emission of a state header with HTTP health checks

May be used in sections : defaults | frontend | listen | backend

```
yes | no | yes | yes
```

Arguments : none

When this option is set, haproxy will systematically send a special header "X-Haproxy-Server-State" with a list of parameters indicating to each server how they are seen by haproxy. This can be used for instance when a server is manipulated without access to haproxy and the operator needs to know whether haproxy still sees it up or not, or if the server is the last one in a farm.

The header is composed of fields delimited by semi-colons, the first of which is a word ("UP", "DOWN", "NOLB"), possibly followed by a number of valid checks on the total number before transition, just as appears in the stats interface. Next headers are in the form "<variable>=<value>", indicating in no specific order some values available in the stats interface :

- a variable "name", containing the name of the backend followed by a slash ("/") then the name of the server. This can be used when a server is checked in multiple backends.

- a variable "node" containing the name of the haproxy node, as set in the global "node" variable, otherwise the system's hostname if unspecified.

- a variable "weight" indicating the weight of the server, a slash ("/") and the total weight of the farm (just counting usable servers). This helps to know if other servers are available to handle the load when this one fails.

- a variable "scur" indicating the current number of concurrent connections on the server, followed by a slash ("/") then the total number of connections on all servers of the same backend.

- a variable "qcur" indicating the current number of requests in the server's queue.

Example of a header received by the application server :

```
>>> X-Haproxy-Server-State: UP 2/3; name=bck/srv2; node=lbl1; weight=1/2; \
      scur=13/22; qcur=0
```

See also : "option httpchk", "http-check disable-on-404"

```
http-request { allow | deny | tarpit | auth [realm <realm>] | redirect <rule> |
  add-header <name> <fmt> | set-header <name> <fmt> |
  del-header <name> | set-nice <nice> | set-log-level <level> |
  replace-header <name> <match-regex> <replace-fmt> |
  replace-value <name> <match-regex> <replace-fmt> |
  set-tos <tos> | set-mark <mark> |
  add-acl(<file name>) <key fmt> |
  del-acl(<file name>) <key fmt> |
  del-map(<file name>) <key fmt> |
  set-map(<file name>) <key fmt> <value fmt>
}
```

```
{ { if | unless } <condition> }
```

Access control for Layer 7 requests

May be used in sections: defaults | frontend | listen | backend

```
no | yes | yes | yes
```

The http-request statement defines a set of rules which apply to layer 7 processing. The rules are evaluated in their declaration order when they are met in a frontend, listen or backend section. Any rule may optionally be followed by an ACL-based condition, in which case it will only be evaluated if the condition is true.

The first keyword is the rule's action. Currently supported actions include :

- "allow" : this stops the evaluation of the rules and lets the request pass the check. No further "http-request" rules are evaluated.

- "deny" : this stops the evaluation of the rules and immediately rejects the request and emits an HTTP 403 error. No further "http-request" rules are evaluated.

- "tarpit" : this stops the evaluation of the rules and immediately blocks the request without responding for a delay specified by "timeout tarpit" or "timeout connect" if the former is not set. After that delay, if the client is still connected, an HTTP error 500 is returned so that the client does not suspect it has been tarpitted. Logs will report the flags "PT". The goal of the tarpit rule is to slow down robots during an attack when they're limited on the number of concurrent requests. It can be very efficient against very dumb robots, and will significantly reduce the load on firewalls compared to a "deny" rule. But when facing "correctly" developed robots, it can make things worse by forcing haproxy and the front firewall to support insane number of concurrent connections.

- "auth" : this stops the evaluation of the rules and immediately responds with an HTTP 401 or 407 error code to invite the user to present a valid user name and password. No further "http-request" rules are evaluated. An optional "realm" parameter is supported, it sets the authentication realm that is returned with the response (typically the application's name).

- "redirect" : this performs an HTTP redirection based on a redirect rule. This is exactly the same as the "redirect" statement except that it

inserts a redirect rule which can be processed in the middle of other "http-request" rules and that these rules use the "log-format" strings. See the "redirect" keyword for the rule's syntax.

- "add-header" appends an HTTP header field whose name is specified in <name> and whose value is defined by <fmt> which follows the log-format rules (see Custom Log Format in section 8.2.4). This is particularly useful to pass connection-specific information to the server (eg: the client's SSL certificate), or to combine several headers into one. This rule is not final, so it is possible to add other similar rules. Note that header addition is performed immediately, so one rule might reuse the resulting header from a previous rule.

- "set-header" does the same as "add-header" except that the header name is first removed if it existed. This is useful when passing security information to the server, where the header must not be manipulated by external users. Note that the new value is computed before the removal so it is possible to concatenate a value to an existing header.

- "del-header" removes all HTTP header fields whose name is specified in <name>.

- "replace-header" matches the regular expression in all occurrences of header field <name> according to <match-regex>, and replaces them with the <replace-fmt> argument. Format characters are allowed in replace-fmt and work like in <fmt> arguments in "add-header". The match is only case-sensitive. It is important to understand that this action only considers whole header lines, regardless of the number of values they may contain. This usage is suited to headers naturally containing commas in their value, such as If-Modified-Since and so on.

Example:

```
http-request replace-header Cookie foo=([^\;]*);(.* ) foo=\1;ip=%i;V2
```

applied to:

```
Cookie: foo=foobar; expires=Tue, 14-Jun-2016 01:40:45 GMT;
```

outputs:

```
Cookie: foo=foobar;ip=192.168.1.20; expires=Tue, 14-Jun-2016 01:40:45 GMT;
assuming the backend IP is 192.168.1.20
```

- "replace-value" works like "replace-header" except that it matches the regex against every comma-delimited value of the header field <name> instead of the entire header. This is suited for all headers which are allowed to carry more than one value. An example could be the Accept header.

Example:

```
http-request replace-value X-Forwarded-For ^192\.168\.(\.)*$ 172.16.\1
```

applied to:

```
X-Forwarded-For: 192.168.10.1, 192.168.13.24, 10.0.0.37
```

outputs:

```
X-Forwarded-For: 172.16.10.1, 172.16.13.24, 10.0.0.37
```

- "set-nice" sets the "nice" factor of the current request being processed. It only has effect against the other requests being processed at the same

time. The default value is 0, unless altered by the "nice" setting on the "bind" line. The accepted range is -1024..1024. The higher the value, the nicest the request will be. Lower values will make the request more important than other ones. This can be useful to improve the speed of some requests, or lower the priority of non-important requests. Using this setting without prior experimentation can cause some major slowdown.

- "set-log-level" is used to change the log level of the current request when a certain condition is met. Valid levels are the 8 syslog levels (see the "log" keyword) plus the special level "silent" which disables logging for this request. This rule is not final so the last matching rule wins. This rule can be useful to disable health checks coming from another equipment.

- "set-tos" is used to set the TOS or DSCP field value of packets sent to the client to the value passed in <tos> on platforms which support this. This value represents the whole 8 bits of the IP TOS field, and can be expressed both in decimal or hexadecimal format (prefixed by "0x"). Note that only the 6 higher bits are used in DSCP or TOS, and the two lower bits are always 0. This can be used to adjust some routing behaviour on border routers based on some information from the request. See RFC 2474, 2597, 3260 and 4594 for more information.

- "set-mark" is used to set the Netfilter MARK on all packets sent to the client to the value passed in <mark> on platforms which support it. This value is an unsigned 32 bit value which can be matched by netfilter and by the routing table. It can be expressed both in decimal or hexadecimal format (prefixed by "0x"). This can be useful to force certain packets to take a different route (for example a cheaper network path for bulk downloads). This works on Linux kernels 2.6.32 and above and requires admin privileges.

- "add-acl" is used to add a new entry into an ACL. The ACL must be loaded from a file (even a dummy empty file). The file name of the ACL to be updated is passed between parentheses. It takes one argument: <key fmt>, which follows log-format rules, to collect content of the new entry. It performs a lookup in the ACL before insertion, to avoid duplicated (or more) values. This lookup is done by a linear search and can be expensive with large lists! It is the equivalent of the "add acl" command from the stats socket, but can be triggered by an HTTP request.

- "del-acl" is used to delete an entry from an ACL. The ACL must be loaded from a file (even a dummy empty file). The file name of the ACL to be updated is passed between parentheses. It takes one argument: <key fmt>, which follows log-format rules, to collect content of the entry to delete. It is the equivalent of the "del acl" command from the stats socket, but can be triggered by an HTTP request.

- "del-map" is used to delete an entry from a MAP. The MAP must be loaded from a file (even a dummy empty file). The file name of the MAP to be updated is passed between parentheses. It takes one argument: <key fmt>, which follows log-format rules, to collect content of the entry to delete. It takes one argument: "file name" It is the equivalent of the "del map" command from the stats socket, but can be triggered by an HTTP request.

- "set-map" is used to add a new entry into a MAP. The MAP must be loaded from a file (even a dummy empty file). The file name of the MAP to be updated is passed between parentheses. It takes 2 arguments: <key fmt>, which follows log-format rules, used to collect MAP key, and <value fmt>, which follows log-format rules, used to collect content for the new entry. It performs a lookup in the MAP before insertion, to avoid duplicated (or more) values. This lookup is done by a linear search and can be expensive with large lists! It is the equivalent of the "set map" command from the stats socket, but can be triggered by an HTTP request.

There is no limit to the number of http-request statements per instance.

It is important to know that http-request rules are processed very early in the HTTP processing, just after "block" rules and before "redirect" or "reqrep" rules. That way, headers added by "add-header"/"set-header" are visible by almost all further ACL rules.

Example:

```
acl nagios src 192.168.129.3
acl local_net src 192.168.0.0/16
acl auth_ok http_auth(L1)
```

```
http-request allow if nagios
http-request allow if local_net auth_ok
http-request auth realm Gimme if local_net auth_ok
http-request deny
```

Example:

```
acl auth_ok http_auth_group(L1) G1
http-request auth unless auth_ok
```

Example:

```
http-request set-header X-Haproxy-Current-Date %T
http-request set-header X-SSL %ssl_fc
http-request set-header X-SSL-Session-ID %ssl_fc_session_id_hex]
http-request set-header X-SSL-Client-Verify %ssl_c_verify]
http-request set-header X-SSL-Client-DN %(+Q)[ssl_c_dn]
http-request set-header X-SSL-Client-CN %(+Q)[ssl_c_s_dn(cm)]
http-request set-header X-SSL-Issuer %(+Q)[ssl_c_i_dn]
http-request set-header X-SSL-Client-NotBefore %(+Q)[ssl_c_notbefore]
http-request set-header X-SSL-Client-NotAfter %(+Q)[ssl_c_notafter]
```

Example:

```
acl key req_hdr(X-Add-Acl-Key) -m found
acl del path /addacl
acl del path /delacl

acl myhost hdr(Host) -f myhost.lst
```

```
http-request add-acl(myhost.lst) %[req_hdr(X-Add-Acl-Key)] if key add
http-request del-acl(myhost.lst) %[req_hdr(X-Add-Acl-Key)] if key del
```

Example:

```
acl value req_hdr(X-Value) -m found
acl setmap path /setmap
acl delmap path /delmap

use_backend bk_appli if { hdr(Host),map_str(map.lst) -m found }

http-request set-map(map.lst) %[src] %[req_hdr(X-Value)] if setmap value
http-request del-map(map.lst) %[src] if delmap
```

See also : "stats http-request", section 3.4 about userlists and section 7 about ACL usage.

```
http-response { allow | deny | add-header <name> <fmt> | set-nice <nice> |
  set-header <name> <fmt> | del-header <name> |
  replace-header <name> <regex-match> <replace-fmt> |
  replace-value <name> <regex-match> <replace-fmt> |
  set-log-level <level> | set-mark <mark> | set-tos <tos> |
  add-acl(<file name>) <key fmt> |
  del-acl(<file name>) <key fmt> |
  del-map(<file name>) <key fmt> |
  set-map(<file name>) <key fmt> <value fmt>
}
```

```
[ { if | unless } <condition> ]
```

Access control for Layer 7 responses

May be used in sections: defaults | frontend | listen | backend
no | yes | yes | yes

The http-response statement defines a set of rules which apply to layer 7 processing. The rules are evaluated in their declaration order when they are met in a frontend, listen or backend section. Any rule may optionally be followed by an ACL-based condition, in which case it will only be evaluated if the condition is true. Since these rules apply on responses, the backend rules are applied first, followed by the frontend's rules.

The first keyword is the rule's action. Currently supported actions include :
- "allow" : this stops the evaluation of the rules and lets the response pass the check. No further "http-response" rules are evaluated for the current section.

- "deny" : this stops the evaluation of the rules and immediately rejects the response and emits an HTTP 502 error. No further "http-response" rules are evaluated.

- "add-header" appends an HTTP header field whose name is specified in <name> and whose value is defined by <fmt> which follows the log-format rules (see Custom Log Format in section 8.2.4). This may be used to send a cookie to a client for example, or to pass some internal information. This rule is not final, so it is possible to add other similar rules. Note that header addition is performed immediately, so one rule might reuse the resulting header from a previous rule.

- "set-header" does the same as "add-header" except that the header name is first removed if it existed. This is useful when passing security information to the server, where the header must not be manipulated by external users.

- "del-header" removes all HTTP header fields whose name is specified in <name>.

- "replace-header" matches the regular expression in all occurrences of header field <name> according to <match-regex>, and replaces them with the <replace-fmt> argument. Format characters are allowed in replace-fmt and work like in <fmt> arguments in "add-header". The match is only case-sensitive. It is important to understand that this action only considers whole header lines, regardless of the number of values they may contain. This usage is suited to headers naturally containing commas in their value, such as Set-Cookie, Expires and so on.

Example:

```
http-response replace-header Set-Cookie (C=[^;]*);(.* )\1:ip=%bi;\2
applied to:
```

```
Set-Cookie: C=1; expires=Tue, 14-Jun-2016 01:40:45 GMT
```

outputs:

```
Set-Cookie: C=1;ip=192.168.1.20; expires=Tue, 14-Jun-2016 01:40:45 GMT
assuming the backend IP is 192.168.1.20.
```

- "replace-value" works like "replace-header" except that it matches the regex against every comma-delimited value of the header field <name> instead of the entire header. This is suited for all headers which are allowed to carry more than one value. An example could be the Accept

```
3251 header.
3252
3253 Example:
3254 http-response replace-value Cache-control ^public$ private
3255
3256 applied to:
3257
3258 Cache-Control: max-age=3600, public
3259
3260 outputs:
3261
3262 Cache-Control: max-age=3600, private
3263
3264 - "set-nice" sets the "nice" factor of the current request being processed.
3265 It only has effect against the other requests being processed at the same
3266 time. The default value is 0, unless altered by the "nice" setting on the
3267 "bind" line. The accepted range is -1024..1024. The higher the value, the
3268 nicest the request will be. Lower values will make the request more
3269 important than other ones. This can be useful to improve the speed of
3270 some requests, or lower the priority of non-important requests. Using
3271 this setting without prior experimentation can cause some major slowdown.
3272
3273 - "set-log-level" is used to change the log level of the current request
3274 when a certain condition is met. Valid levels are the 8 syslog levels
3275 (see the "log" keyword) plus the special level "silent" which disables
3276 logging for this request. This rule is not final so the last matching
3277 rule wins. This rule can be useful to disable health checks coming from
3278 another equipment.
3279
3280 - "set-tos" is used to set the TOS or DSCP field value of packets sent to
3281 the client to the value passed in <tos> on platforms which support this.
3282 This value represents the whole 8 bits of the IP TOS field, and can be
3283 expressed both in decimal or hexadecimal format (prefixed by "0x"). Note
3284 that only the 6 higher bits are used in DSCP or TOS, and the two lower
3285 bits are always 0. This can be used to adjust some routing behaviour on
3286 border routers based on some information from the request. See RFC 2474,
3287 2597, 3260 and 4594 for more information.
3288
3289 - "set-mark" is used to set the Netfilter MARK on all packets sent to the
3290 client to the value passed in <mark> on platforms which support it. This
3291 value is an unsigned 32 bit value which can be matched by netfilter and
3292 by the routing table. It can be expressed both in decimal or hexadecimal
3293 format (prefixed by "0x"). This can be useful to force certain packets to
3294 take a different route (for example a cheaper network path for bulk
3295 downloads). This works on Linux kernels 2.6.32 and above and requires
3296 admin privileges.
3297
3298 - "add-acl" is used to add a new entry into an ACL. The ACL must be loaded
3299 from a file (even a dummy empty file). The file name of the ACL to be
3300 updated is passed between parentheses. It takes one argument: <key fmt>,
3301 which follows log-format rules, to collect content of the new entry. It
3302 performs a lookup in the ACL before insertion, to avoid duplicated (or
3303 more) values. This lookup is done by a linear search and can be expensive
3304 with large lists! It is the equivalent of the "add acl" command from the
3305 stats socket, but can be triggered by an HTTP response.
3306
3307 - "del-acl" is used to delete an entry from an ACL. The ACL must be loaded
3308 from a file (even a dummy empty file). The file name of the ACL to be
3309 updated is passed between parentheses. It takes one argument: <key fmt>,
3310 which follows log-format rules, to collect content of the entry to delete.
3311 It is the equivalent of the "del acl" command from the stats socket, but
3312 can be triggered by an HTTP response.
3313
3314 - "del-map" is used to delete an entry from a MAP. The MAP must be loaded
```

```
3316 from a file (even a dummy empty file). The file name of the MAP to be
3317 updated is passed between parentheses. It takes one argument: <key fmt>,
3318 which follows log-format rules, to collect content of the entry to delete.
3319 It takes one argument: "file name" It is the equivalent of the "del map"
3320 command from the stats socket, but can be triggered by an HTTP response.
3321
3322 - "set-map" is used to add a new entry into a MAP. The MAP must be loaded
3323 from a file (even a dummy empty file). The file name of the MAP to be
3324 updated is passed between parentheses. It takes 2 arguments: <key fmt>,
3325 which follows log-format rules, used to collect MAP key, and <value fmt>,
3326 which follows log-format rules, used to collect content for the new entry.
3327 It performs a lookup in the MAP before insertion, to avoid duplicated (or
3328 more) values. This lookup is done by a linear search and can be expensive
3329 with large lists! It is the equivalent of the "set map" command from the
3330 stats socket, but can be triggered by an HTTP response.
3331
3332 There is no limit to the number of http-response statements per instance.
3333
3334 It is important to know that http-response rules are processed very early in
3335 the HTTP processing, before "reqdel" or "reqrep" rules. That way, headers
3336 added by "add-header"/"set-header" are visible by almost all further ACL
3337 rules.
3338
3339 Example:
3340 acl key_acl res.hdr(X-Acl-Key) -m found
3341
3342 acl myhost hdr(Host) -f myhost.lst
3343
3344 http-response add-acl(myhost.lst) %[res.hdr(X-Acl-Key)] if key_acl
3345 http-response del-acl(myhost.lst) %[res.hdr(X-Acl-Key)] if key_acl
3346
3347 Example:
3348 acl value res.hdr(X-Value) -m found
3349
3350 use_backend bk_appli if { hdr(Host),map_str(map.lst) -m found }
3351
3352 http-response set-map(map.lst) %[src] %[res.hdr(X-Value)] if value
3353 http-response del-map(map.lst) %[src]
3354
3355 See also : "http-request", section 3.4 about userlists and section 7 about
3356 ACL usage.
3357
3358 http-send-name-header [<header>]
3359 Add the server name to a request. Use the header string given by <header>
3360
3361 May be used in sections: defaults | frontend | listen | backend
3362 yes | no | yes | yes
3363
3364 Arguments :
3365
3366 <header> The header string to use to send the server name
3367
3368 The "http-send-name-header" statement causes the name of the target
3369 server to be added to the headers of an HTTP request. The name
3370 is added with the header string proved.
3371
3372 See also : "server"
3373
3374 id <value>
3375 Set a persistent ID to a proxy.
3376 May be used in sections : defaults | frontend | listen | backend
3377 no | yes | yes | yes
3378
3379 Arguments : none
3380
```

3381 Set a persistent ID for the proxy. This ID must be unique and positive.
3382 An unused ID will automatically be assigned if unset. The first assigned
3383 value will be 1. This ID is currently only returned in statistics.
3384
3385
3386 ignore-persist { if | unless } <condition>
3387 Declare a condition to ignore persistence
3388 May be used in sections: defaults | frontend | listen | backend
3389 no | yes | yes | yes
3390
3391 By default, when cookie persistence is enabled, every requests containing
3392 the cookie are unconditionally persistent (assuming the target server is up
3393 and running).
3394
3395 The "ignore-persist" statement allows one to declare various ACL-based
3396 conditions which, when met, will cause a request to ignore persistence.
3397 This is sometimes useful to load balance requests for static files, which
3398 often don't require persistence. This can also be used to fully disable
3399 persistence for a specific User-Agent (for example, some web crawler bots).
3400
3401 Combined with "appsession", it can also help reduce HAProxy memory usage, as
3402 the appsession table won't grow if persistence is ignored.
3403
3404 The persistence is ignored when an "if" condition is met, or unless an
3405 "unless" condition is met.
3406
3407 See also : "force-persist", "cookie", and section 7 about ACL usage.
3408
3409
3410 log global
3411 log <address> [len <length>] <facility> [<level> [<minlevel>]]
3412 no log
3413 Enable per-instance logging of events and traffic.
3414 May be used in sections : defaults | frontend | listen | backend
3415 yes | yes | yes | yes
3416
3417 Prefix :
3418 no
3419 should be used when the logger list must be flushed. For example,
3420 if you don't want to inherit from the default logger list. This
3421 prefix does not allow arguments.
3422
3423 Arguments :
3424 global
3425 should be used when the instance's logging parameters are the
3426 same as the global ones. This is the most common usage. "global"
3427 replaces <address>, <facility> and <level> with those of the log
3428 entries found in the "global" section. Only one "log global"
3429 statement may be used per instance, and this form takes no other
3430 parameter.
3431
3432 <address>
3433 indicates where to send the logs. It takes the same format as
3434 for the "global" section's logs, and can be one of :
3435
3436 - An IPv4 address optionally followed by a colon (':') and a UDP
3437 port. If no port is specified, 514 is used by default (the
3438 standard syslog port).
3439
3440 - An IPv6 address followed by a colon (':') and optionally a UDP
3441 port. If no port is specified, 514 is used by default (the
3442 standard syslog port).
3443
3444 - A filesystem path to a UNIX domain socket, keeping in mind
3445 considerations for chroot (be sure the path is accessible
inside the chroot) and uid/gid (be sure the path is
appropriately writeable).

3446 Any part of the address string may reference any number of
3447 environment variables by preceding their name with a dollar
3448 sign ('\$') and optionally enclosing them with braces ('{') ,
3449 similarly to what is done in Bourne shell.
3450
3451 <length>
3452 is an optional maximum line length. Log lines larger than this
3453 value will be truncated before being sent. The reason is that
3454 syslog servers act differently on log line length. All servers
3455 support the default value of 1024, but some servers simply drop
3456 larger lines while others do log them. If a server supports long
3457 lines, it may make sense to set this value here in order to avoid
3458 truncating long lines. Similarly, if a server drops long lines,
3459 it is preferable to truncate them before sending them. Accepted
3460 values are 80 to 65535 inclusive. The default value of 1024 is
3461 generally fine for all standard usages. Some specific cases of
3462 long captures or JSON-formatted logs may require larger values.
3463
3464 <facility> must be one of the 24 standard syslog facilities :
3465 kern user mail daemon auth syslog lpr news
3466 uucp cron auth2 ftp ntp audit alert cron2
3467 local0 local1 local2 local3 local4 local5 local6 local7
3468
3469 <level>
3470 is optional and can be specified to filter outgoing messages. By
3471 default, all messages are sent. If a level is specified, only
3472 messages with a severity at least as important as this level
3473 will be sent. An optional minimum level can be specified. If it
3474 is set, logs emitted with a more severe level than this one will
3475 be capped to this level. This is used to avoid sending "emerg"
3476 messages on all terminals on some default syslog configurations.
3477 Eight levels are known :
3478 emerg alert crit err warning notice info debug
3479
3480 It is important to keep in mind that it is the frontend which decides what to
3481 log from a connection, and that in case of content switching, the log entries
3482 from the backend will be ignored. Connections are logged at level "info".
3483
3484 However, backend log declaration define how and where servers status changes
3485 will be logged. Level "notice" will be used to indicate a server going up,
3486 "warning" will be used for termination signals and definitive service
3487 termination, and "alert" will be used for when a server goes down.
3488
3489 Note : According to RFC3164, messages are truncated to 1024 bytes before
3490 being emitted.
3491
3492 Example :
3493 log global
3494 log 127.0.0.1:514 local0 notice # only send important events
3495 log 127.0.0.1:514 local0 notice # same but limit output level
3496 log \${LOCAL_SYSLOG}:514 local0 notice # send to local server
3497
3498 Log-format <string>
3499 Specifies the log format string to use for traffic logs
3500 May be used in sections: defaults | frontend | listen | backend
3501 yes | yes | yes | no
3502
3503 This directive specifies the log format string that will be used for all logs
3504 resulting from traffic passing through the frontend using this line. If the
3505 directive is used in a defaults section, all subsequent frontends will use
3506 the same log format. Please see section 8.2.4 which covers the log format
3507 string in depth.
3508
3509
3510

```

3511 max-keep-alive-queue <value>
3512 Set the maximum server queue size for maintaining keep-alive connections
3513 May be used in sections: defaults | frontend | listen | backend
3514 yes | no | yes | yes
3515
3516 HTTP keep-alive tries to reuse the same server connection whenever possible,
3517 but sometimes it can be counter-productive, for example if a server has a lot
3518 of connections while other ones are idle. This is especially true for static
3519 servers.
3520
3521 The purpose of this setting is to set a threshold on the number of queued
3522 connections at which haproxy stops trying to reuse the same server and prefers
3523 to find another one. The default value, -1, means there is no limit. A value
3524 of zero means that keep-alive requests will never be queued. For very close
3525 servers which can be reached with a low latency and which are not sensible to
3526 breaking keep-alive, a low value is recommended (eg: local static server can
3527 use a value of 10 or less). For remote servers suffering from a high latency,
3528 higher values might be needed to cover for the latency and/or the cost of
3529 picking a different server.
3530
3531 Note that this has no impact on responses which are maintained to the same
3532 server consecutively to a 401 response. They will still go to the same server
3533 even if they have to be queued.
3534
3535 See also : "option http-server-close", "option prefer-last-server", server
3536 "maxconn" and cookie persistence.
3537
3538 maxconn <conns>
3539 Fix the maximum number of concurrent connections on a frontend
3540 May be used in sections : defaults | frontend | listen | backend
3541 yes | yes | yes | no
3542
3543 Arguments :
3544 <conns> is the maximum number of concurrent connections the frontend will
3545 accept to serve. Excess connections will be queued by the system
3546 in the socket's listen queue and will be served once a connection
3547 closes.
3548
3549 If the system supports it, it can be useful on big sites to raise this limit
3550 very high so that haproxy manages connection queues, instead of leaving the
3551 clients with unanswered connection attempts. This value should not exceed the
3552 global maxconn. Also, keep in mind that a connection contains two buffers
3553 of 8kB each, as well as some other data resulting in about 17 kB of RAM being
3554 consumed per established connection. That means that a medium system equipped
3555 with 1GB of RAM can withstand around 40000-50000 concurrent connections if
3556 properly tuned.
3557
3558 Also, when <conns> is set to large values, it is possible that the servers
3559 are not sized to accept such loads, and for this reason it is generally wise
3560 to assign them some reasonable connection limits.
3561
3562 By default, this value is set to 2000.
3563
3564 See also : "server", global section's "maxconn", "fullconn"
3565
3566 mode { tcp|http|health }
3567 Set the running mode or protocol of the instance
3568 May be used in sections : defaults | frontend | listen | backend
3569 yes | yes | yes | yes
3570
3571 Arguments :
3572 tcp The instance will work in pure TCP mode. A full-duplex connection
3573 will be established between clients and servers, and no layer 7
3574 examination will be performed. This is the default mode. It
3575 should be used for SSL, SSH, SMTP, ...

```

```

3576
3577 http The instance will work in HTTP mode. The client request will be
3578 analyzed in depth before connecting to any server. Any request
3579 which is not RFC-compliant will be rejected. Layer 7 filtering,
3580 processing and switching will be possible. This is the mode which
3581 brings HAProxy most of its value.
3582
3583 health The instance will work in "health" mode. It will just reply "OK"
3584 to incoming connections and close the connection. Alternatively,
3585 If the "htpchk" option is set, "HTTP/1.0 200 OK" will be sent
3586 instead. Nothing will be logged in either case. This mode is used
3587 to reply to external components health checks. This mode is
3588 deprecated and should not be used anymore as it is possible to do
3589 the same and even better by combining TCP or HTTP modes with the
3590 "monitor" keyword.
3591
3592 When doing content switching, it is mandatory that the frontend and the
3593 backend are in the same mode (generally HTTP), otherwise the configuration
3594 will be refused.
3595
3596 Example :
3597 defaults http_instances
3598 mode http
3599
3600 See also : "monitor", "monitor-net"
3601
3602 monitor fail { if | unless } <condition>
3603 Add a condition to report a failure to a monitor HTTP request.
3604 May be used in sections : defaults | frontend | listen | backend
3605 no | yes | yes | no
3606
3607 Arguments :
3608 if <cond> the monitor request will fail if the condition is satisfied,
3609 and will succeed otherwise. The condition should describe a
3610 combined test which must induce a failure if all conditions
3611 are met, for instance a low number of servers both in a
3612 backend and its backup.
3613
3614 unless <cond> the monitor request will succeed only if the condition is
3615 satisfied, and will fail otherwise. Such a condition may be
3616 based on a test on the presence of a minimum number of active
3617 servers in a list of backends.
3618
3619 This statement adds a condition which can force the response to a monitor
3620 request to report a failure. By default, when an external component queries
3621 the URI dedicated to monitoring, a 200 response is returned. When one of the
3622 conditions above is met, haproxy will return 503 instead of 200. This is
3623 very useful to report a site failure to an external component which may base
3624 routing advertisements between multiple sites on the availability reported by
3625 haproxy. In this case, one would rely on an ACL involving the "nbsrv"
3626 criterion. Note that "monitor fail" only works in HTTP mode. Both status
3627 messages may be tweaked using "errorfile" or "errorloc" if needed.
3628
3629 Example:
3630 frontend www
3631 mode http
3632 acl site_dead nbsrv(dynamic) lt 2
3633 acl site_dead nbsrv(static) lt 2
3634 monitor-uri /site_alive
3635 monitor fail if site_dead
3636
3637 See also : "monitor-net", "monitor-uri", "errorfile", "errorloc"
3638
3639 monitor-net <source>
3640

```

3641 Declare a source network which is limited to monitor requests
 3642 May be used in sections : defaults | frontend | listen | backend
 3643 yes | yes | yes | no
 3644
 3645 Arguments :
 3646 <source> is the source IPv4 address or network which will only be able to
 3647 get monitor responses to any request. It can be either an IPv4
 3648 address, a host name, or an address followed by a slash ('/')
 3649 followed by a mask.

3650 In TCP mode, any connection coming from a source matching <source> will cause
 3651 the connection to be immediately closed without any log. This allows another
 3652 equipment to probe the port and verify that it is still listening, without
 3653 forwarding the connection to a remote server.
 3654

3655 In HTTP mode, a connection coming from a source matching <source> will be
 3656 accepted, the following response will be sent without waiting for a request,
 3657 then the connection will be closed : "HTTP/1.0 200 OK". This is normally
 3658 enough for any front-end HTTP probe to detect that the service is UP and
 3659 running without forwarding the request to a backend server. Note that this
 3660 response is sent in raw format, without any transformation. This is important
 3661 as it means that it will not be SSL-encrypted on SSL listeners.
 3662

3663 Monitor requests are processed very early, just after tcp-request connection
 3664 ACLs which are the only ones able to block them. These connections are short
 3665 lived and never wait for any data from the client. They cannot be logged, and
 3666 it is the intended purpose. They are only used to report HAProxy's health to
 3667 an upper component, nothing more. Please note that "monitor fail" rules do
 3668 not apply to connections intercepted by "monitor-net".

3669 Last, please note that only one "monitor-net" statement can be specified in
 3670 a frontend. If more than one is found, only the last one will be considered.

3671 Example :
 3672 # addresses .252 and .253 are just probing us.
 3673 frontend www
 3674 monitor-net 192.168.0.252/31
 3675

3676 See also : "monitor fail", "monitor-uri"

3677 monitor-uri <uri>

3678 Intercept a URI used by external components' monitor requests

3682 May be used in sections : defaults | frontend | listen | backend
 3683 yes | yes | yes | no
 3684

3685 Arguments :
 3686 <uri> is the exact URI which we want to intercept to return HAProxy's
 3687 health status instead of forwarding the request.
 3688

3689 When an HTTP request referencing <uri> will be received on a frontend,
 3690 HAProxy will not forward it nor log it, but instead will return either
 3691 "HTTP/1.0 200 OK" or "HTTP/1.0 503 Service unavailable", depending on failure
 3692 conditions defined with "monitor fail". This is normally enough for any
 3693 front-end HTTP probe to detect that the service is UP and running without
 3694 forwarding the request to a backend server. Note that the HTTP method, the
 3695 version and all headers are ignored, but the request must at least be valid
 3696 at the HTTP level. This keyword may only be used with an HTTP-mode frontend.
 3697

3698 Monitor requests are processed very early. It is not possible to block nor
 3699 divert them using ACLs. They cannot be logged either, and it is the intended
 3700 purpose. They are only used to report HAProxy's health to an upper component,
 3701 nothing more. However, it is possible to add any number of conditions using
 3702 "monitor fail" and ACLs so that the result can be adjusted to whatever check
 3703 can be imagined (most often the number of available servers in a backend).
 3704

3705 Example :

3706 # Use /haproxy_test to report haproxy's status
 3707 frontend www
 3708 mode http
 3709 monitor-uri /haproxy_test
 3710

3711 See also : "monitor fail", "monitor-net"

3712 option abortonclose

3714 no option abortonclose

3715 Enable or disable early dropping of aborted requests pending in queues.
 3716 May be used in sections : defaults | frontend | listen | backend
 3717 defaults yes | no | yes | yes
 3718

3719 Arguments : none

3720 In presence of very high loads, the servers will take some time to respond.
 3721 The per-instance connection queue will inflate, and the response time will
 3722 increase respective to the size of the queue times the average per-session
 3723 response time. When clients will wait for more than a few seconds, they will
 3724 often hit the "STOP" button on their browser, leaving a useless request in
 3725 the queue, and slowing down other users, and the servers as well, because the
 3726 request will eventually be served, then aborted at the first error
 3727 encountered while delivering the response.
 3728

3729 As there is no way to distinguish between a full STOP and a simple output
 3730 close on the client side, HTTP agents should be conservative and consider
 3731 that the client might only have closed its output channel while waiting for
 3732 the response. However, this introduces risks of congestion when lots of users
 3733 do the same, and is completely useless nowadays because probably no client at
 3734 all will close the session while waiting for the response. Some HTTP agents
 3735 support this behaviour (Squid, Apache, HAProxy), and others do not (TUX, most
 3736 hardware-based load balancers). So the probability for a closed input channel
 3737 to represent a user hitting the "STOP" button is close to 100%, and the risk
 3738 of being the single component to break rare but valid traffic is extremely
 3739 low, which adds to the temptation to be able to abort a session early while
 3740 still not served and not pollute the servers.
 3741

3742 In HAProxy, the user can choose the desired behaviour using the option
 3743 "abortonclose". By default (without the option) the behaviour is HTTP
 3744 compliant and aborted requests will be served. But when the option is
 3745 specified, a session with an incoming channel closed will be aborted while
 3746 it is still possible, either pending in the queue for a connection slot, or
 3747 during the connection establishment if the server has not yet acknowledged
 3748 the connection request. This considerably reduces the queue size and the load
 3749 on saturated servers when users are tempted to click on STOP, which in turn
 3750 reduces the response time for other users.
 3751

3752 If this option has been enabled in a "defaults" section, it can be disabled
 3753 in a specific instance by prepending the "no" keyword before it.
 3754

3755 See also : "timeout queue" and server's "maxconn" and "maxqueue" parameters

3756 option accept-invalid-http-request

3759 no option accept-invalid-http-request

3760 Enable or disable relaxing of HTTP request parsing
 3761 May be used in sections : defaults | frontend | listen | backend
 3762 defaults yes | yes | yes | no
 3763

3764 Arguments : none

3765 By default, HAProxy complies with RFC7230 in terms of message parsing. This
 3766 means that invalid characters in header names are not permitted and cause an
 3767 error to be returned to the client. This is the desired behaviour as such
 3768 forbidden characters are essentially used to build attacks exploiting server
 3769 weaknesses, and bypass security filtering. Sometimes, a buggy browser or
 3770

server will emit invalid header names for whatever reason (configuration, implementation) and the issue will not be immediately fixed. In such a case, it is possible to relax HAProxy's header name parser to accept any character even if that does not make sense, by specifying this option. Similarly, the list of characters allowed to appear in a URI is well defined by RFC3986, and chars 0-31, 32 (space), 34 ('"'), 60 ('<'), 62 ('>'), 92 ('\'), 94 ('\^'), 96 ('\`'), 123 ('{'), 124 ('|'), 125 ('}'), 127 (delete) and anything above are not allowed at all. HAProxy always blocks a number of them (0..32, 127). The remaining ones are blocked by default unless this option is enabled. This option also relaxes the test on the HTTP version format, it allows multiple digits for both the major and the minor version.

This option should never be enabled by default as it hides application bugs and open security breaches. It should only be deployed after a problem has been confirmed.

When this option is enabled, erroneous header names will still be accepted in requests, but the complete request will be captured in order to permit later analysis using the "show errors" request on the UNIX stats socket. Similarly, requests containing invalid chars in the URI part will be logged. Doing this also helps confirming that the issue has been solved.

If this option has been enabled in a "defaults" section, it can be disabled in a specific instance by prepending the "no" keyword before it.

See also : "option accept-invalid-http-response" and "show errors" on the stats socket.

option accept-invalid-http-response

no option accept-invalid-http-response

Enable or disable relaxing of HTTP response parsing

May be used in sections : defaults | frontend | listen | backend
yes | no | yes | yes

Arguments : none

By default, HAProxy complies with RFC7230 in terms of message parsing. This means that invalid characters in header names are not permitted and cause an error to be returned to the client. This is the desired behaviour as such forbidden characters are essentially used to build attacks exploiting server weaknesses, and bypass security filtering. Sometimes, a buggy browser or server will emit invalid header names for whatever reason (configuration, implementation) and the issue will not be immediately fixed. In such a case, it is possible to relax HAProxy's header name parser to accept any character even if that does not make sense, by specifying this option. This option also relaxes the test on the HTTP version format, it allows multiple digits for both the major and the minor version.

This option should never be enabled by default as it hides application bugs and open security breaches. It should only be deployed after a problem has been confirmed.

When this option is enabled, erroneous header names will still be accepted in responses, but the complete response will be captured in order to permit later analysis using the "show errors" request on the UNIX stats socket. Doing this also helps confirming that the issue has been solved.

If this option has been enabled in a "defaults" section, it can be disabled in a specific instance by prepending the "no" keyword before it.

See also : "option accept-invalid-http-request" and "show errors" on the stats socket.

option allbackups

no option allbackups

Use either all backup servers at a time or only the first one

May be used in sections : defaults | frontend | listen | backend
yes | no | yes | yes

Arguments : none

By default, the first operational backup server gets all traffic when normal servers are all down. Sometimes, it may be preferred to use multiple backups at once, because one will not be enough. When "option allbackups" is enabled, the load balancing will be performed among all backup servers when all normal ones are unavailable. The same load balancing algorithm will be used and the servers' weights will be respected. Thus, there will not be any priority order between the backup servers anymore.

This option is mostly used with static server farms dedicated to return a "sorry" page when an application is completely offline.

If this option has been enabled in a "defaults" section, it can be disabled in a specific instance by prepending the "no" keyword before it.

option checkcache

no option checkcache

Analyze all server responses and block responses with cacheable cookies

May be used in sections : defaults | frontend | listen | backend
yes | no | yes | yes

Arguments : none

Some high-level frameworks set application cookies everywhere and do not always let enough control to the developer to manage how the responses should be cached. When a session cookie is returned on a cacheable object, there is a high risk of session crossing or stealing between users traversing the same caches. In some situations, it is better to block the response than to let some sensitive session information go in the wild.

The option "checkcache" enables deep inspection of all server responses for strict compliance with HTTP specification in terms of cacheability. It carefully checks "Cache-control", "Pragma" and "Set-cookie" headers in server response to check if there's a risk of caching a cookie on a client-side proxy. When this option is enabled, the only responses which can be delivered to the client are :

- all those without 'Set-Cookie' header ;
- all those with a return code other than 200, 203, 206, 300, 301, 410, provided that the server has not set a "Cache-control: public" header ;
- all those that come from a POST request, provided that the server has not set a 'Cache-Control: public' header ;
- those with a 'Pragma: no-cache' header ;
- those with a 'Cache-control: private' header
- those with a 'Cache-control: no-store' header
- those with a 'Cache-control: max-age=0' header
- those with a 'Cache-control: s-maxage=0' header
- those with a 'Cache-control: no-cache' header
- those with a 'Cache-control: no-cache="set-cookie"' header
- those with a 'Cache-control: no-cache="set-cookie,' header (allowing other fields after set-cookie)

If a response doesn't respect these requirements, then it will be blocked just as if it was from an "rspdeny" filter, with an "HTTP 502 bad gateway". The session state shows "PH-" meaning that the proxy blocked the response during headers processing. Additionally, an alert will be sent in the logs so that admins are informed that there's something to be fixed.

Due to the high impact on the application, the application should be tested in depth with the option enabled before going to production. It is also a good practice to always activate it during tests, even if it is not used in

production, as it will report potentially dangerous application behaviours.

If this option has been enabled in a "defaults" section, it can be disabled in a specific instance by prepending the "no" keyword before it.

option cliticpka

no option cliticpka

Enable or disable the sending of TCP keepalive packets on the client side

May be used in sections : defaults | frontend | listen | backend

yes | yes | yes | no

Arguments : none

When there is a firewall or any session-aware component between a client and a server, and when the protocol involves very long sessions with long idle periods (eg: remote desktops), there is a risk that one of the intermediate components decides to expire a session which has remained idle for too long.

Enabling socket-level TCP keep-alives makes the system regularly send packets to the other end of the connection, leaving it active. The delay between keep-alive probes is controlled by the system only and depends both on the operating system and its tuning parameters.

It is important to understand that keep-alive packets are neither emitted nor received at the application level. It is only the network stacks which sees them. For this reason, even if one side of the proxy already uses keep-alives to maintain its connection alive, those keep-alive packets will not be forwarded to the other side of the proxy.

Please note that this has nothing to do with HTTP keep-alive.

Using option "cliticpka" enables the emission of TCP keep-alive probes on the client side of a connection, which should help when session expirations are noticed between HAProxy and a client.

If this option has been enabled in a "defaults" section, it can be disabled in a specific instance by prepending the "no" keyword before it.

See also : "option srvtcpka", "option tcpka"

option contstats

Enable continuous traffic statistics updates

May be used in sections : defaults | frontend | listen | backend

yes | yes | yes | no

Arguments : none

By default, counters used for statistics calculation are incremented only when a session finishes. It works quite well when serving small objects, but with big ones (for example large images or archives) or with A/V streaming, a graph generated from haproxy counters looks like a hedgehog. With this option enabled counters get incremented continuously, during a whole session. Recounting touches a hotpath directly so it is not enabled by default, as it has small performance impact (~0.5%).

option dontlog-normal

no option dontlog-normal

Enable or disable logging of normal, successful connections

May be used in sections : defaults | frontend | listen | backend

yes | yes | yes | no

Arguments : none

There are large sites dealing with several thousand connections per second and for which logging is a major pain. Some of them are even forced to turn

logs off and cannot debug production issues. Setting this option ensures that normal connections, those which experience no error, no timeout, no retry nor redispach, will not be logged. This leaves disk space for anomalies. In HTTP mode, the response status code is checked and return codes 5xx will still be logged.

It is strongly discouraged to use this option as most of the time, the key to complex issues is in the normal logs which will not be logged here. If you need to separate logs, see the "log-separate-errors" option instead.

See also : "log", "dontlognull", "log-separate-errors" and section 8 about logging.

option dontlognull

no option dontlognull

Enable or disable logging of null connections

May be used in sections : defaults | frontend | listen | backend

yes | yes | yes | no

Arguments : none

In certain environments, there are components which will regularly connect to various systems to ensure that they are still alive. It can be the case from another load balancer as well as from monitoring systems. By default, even a simple port probe or scan will produce a log. If those connections pollute the logs too much, it is possible to enable option "dontlognull" to indicate that a connection on which no data has been transferred will not be logged, which typically corresponds to those probes. Note that errors will still be returned to the client and accounted for in the stats. If this is not what is desired, option http-ignore-probes can be used instead.

It is generally recommended not to use this option in uncontrolled environments (eg: internet), otherwise scans and other malicious activities would not be logged.

If this option has been enabled in a "defaults" section, it can be disabled in a specific instance by prepending the "no" keyword before it.

See also : "log", "http-ignore-probes", "monitor-net", "monitor-uri", and section 8 about logging.

option forceclose

no option forceclose

Enable or disable active connection closing after response is transferred.

May be used in sections : defaults | frontend | listen | backend

yes | yes | yes | yes

Arguments : none

Some HTTP servers do not necessarily close the connections when they receive the "Connection: close" set by "option httpclose", and if the client does not close either, then the connection remains open till the timeout expires. This causes high number of simultaneous connections on the servers and shows high global session times in the logs.

When this happens, it is possible to use "option forceclose". It will actively close the outgoing server channel as soon as the server has finished to respond and release some resources earlier than with "option httpclose".

This option may also be combined with "option http-pretend-keepalive", which will disable sending of the "Connection: close" header, but will still cause the connection to be closed once the whole response is received.

This option disables and replaces any previous "option httpclose", "option http-server-close", "option http-keep-alive", or "option http-tunnel".

4031 If this option has been enabled in a "defaults" section, it can be disabled
 4032 in a specific instance by prepending the "no" keyword before it.
 4033
 4034

4035 See also : "option httpclose" and "option http-pretend-keepalive"
 4036
 4037

4038 option forwardfor [except <network>] [header <name>] [if-none]
 4039 Enable insertion of the X-Forwarded-For header to requests sent to servers
 4040 May be used in sections : defaults | frontend | listen | backend
 4041 yes | yes | yes | yes | yes | yes
 4042 Arguments :

4043 <network> : is an optional argument used to disable this option for sources
 4044 matching <network>
 4045 <name> : an optional argument to specify a different "X-Forwarded-For"
 4046 header name.
 4047

4048 Since HAProxy works in reverse-proxy mode, the servers see its IP address as
 4049 their client address. This is sometimes annoying when the client's IP address
 4050 is expected in server logs. To solve this problem, the well-known HTTP header
 4051 "X-Forwarded-For" may be added by HAProxy to all requests sent to the server.
 4052 This header contains a value representing the client's IP address. Since this
 4053 header is always appended at the end of the existing header list, the server
 4054 must be configured to always use the last occurrence of this header only. See
 4055 the server's manual to find how to enable use of this standard header. Note
 4056 that only the last occurrence of the header must be used, since it is really
 4057 possible that the client has already brought one.

4058 The keyword "header" may be used to supply a different header name to replace
 4059 the default "X-Forwarded-For". This can be useful where you might already
 4060 have a "X-Forwarded-For" header from a different application (eg: stunnel),
 4061 and you need preserve it. Also if your backend server doesn't use the
 4062 "X-Forwarded-For" header and requires different one (eg: Zeus Web Servers
 4063 require "X-Cluster-Client-IP").
 4064
 4065

4066 Sometimes, a same HAProxy instance may be shared between a direct client
 4067 access and a reverse-proxy access (for instance when an SSL reverse-proxy is
 4068 used to decrypt HTTPS traffic). It is possible to disable the addition of the
 4069 header for a known source address or network by adding the "except" keyword
 4070 followed by the network address. In this case, any source IP matching the
 4071 network will not cause an addition of this header. Most common uses are with
 4072 private networks or 127.0.0.1.

4073 Alternatively, the keyword "if-none" states that the header will only be
 4074 added if it is not present. This should only be used in perfectly trusted
 4075 environment, as this might cause a security issue if headers reaching haproxy
 4076 are under the control of the end-user.
 4077

4078 This option may be specified either in the frontend or in the backend. If at
 4079 least one of them uses it, the header will be added. Note that the backend's
 4080 setting of the header subargument takes precedence over the frontend's if
 4081 both are defined. In the case of the "if-none" argument, if at least one of
 4082 the frontend or the backend does not specify it, it wants the addition to be
 4083 mandatory, so it wins.
 4084
 4085

4086 Examples :
 4087 # Public HTTP address also used by stunnel on the same machine
 4088 frontend www
 4089 mode http
 4090 option forwardfor except 127.0.0.1 # stunnel already adds the header
 4091
 4092

4093 # Those servers want the IP Address in X-Client
 4094 backend www
 4095 mode http
 option forwardfor header X-Client

4096 See also : "option httpclose", "option http-server-close",
 4097 "option forceclose", "option http-keep-alive"
 4098
 4099

4100 option http-ignore-probes
 4101 no option http-ignore-probes
 4102 Enable or disable logging of null connections and request timeouts
 4103 May be used in sections : defaults | frontend | listen | backend
 4104 yes | yes | yes | yes | yes | no
 4105 Arguments : none
 4106
 4107

4108 Recently some browsers started to implement a "pre-connect" feature
 4109 consisting in speculatively connecting to some recently visited web sites
 4110 just in case the user would like to visit them. This results in many
 4111 connections being established to web sites, which end up in 408 Request
 4112 Timeout if the timeout strikes first, or 400 Bad Request when the browser
 4113 decides to close them first. These ones pollute the log and feed the error
 4114 counters. There was already "option dontlognull" but it's insufficient in
 4115 this case. Instead, this option does the following things :
 4116 - prevent any 400/408 message from being sent to the client if nothing
 4117 was received over a connection before it was closed ;
 4118 - prevent any log from being emitted in this situation ;
 4119 - prevent any error counter from being incremented
 4120
 4121

4122 That way the empty connection is silently ignored. Note that it is better
 4123 not to use this unless it is clear that it is needed, because it will hide
 4124 real problems. The most common reason for not receiving a request and seeing
 4125 a 408 is due to an MTU inconsistency between the client and an intermediary
 4126 element such as a VPN, which blocks too large packets. These issues are
 4127 generally seen with POST requests as well as GET with large cookies. The logs
 4128 are often the only way to detect them.

4129 If this option has been enabled in a "defaults" section, it can be disabled
 4130 in a specific instance by prepending the "no" keyword before it.
 4131

4132 See also : "log", "dontlognull", "errorfile", and section 8 about logging.
 4133
 4134

4135 option http-keep-alive
 4136 no option http-keep-alive
 4137 Enable or disable HTTP keep-alive from client to server
 4138 May be used in sections : defaults | frontend | listen | backend
 4139 yes | yes | yes | yes | yes | yes
 4140 Arguments : none
 4141
 4142

4143 By default HAProxy operates in keep-alive mode with regards to persistent
 4144 connections: for each connection it processes each request and response, and
 4145 leaves the connection idle on both sides between the end of a response and the
 4146 start of a new request. This mode may be changed by several options such as
 4147 "option http-server-close", "option forceclose", "option httpclose" or
 4148 "option http-tunnel". This option allows to set back the keep-alive mode,
 4149 which can be useful when another mode was used in a defaults section.

4150 Setting "option http-keep-alive" enables HTTP keep-alive mode on the client-
 4151 and server- sides. This provides the lowest latency on the client side (slow
 4152 network) and the fastest session reuse on the server side at the expense
 4153 of maintaining idle connections to the servers. In general, it is possible
 4154 with this option to achieve approximately twice the request rate that the
 4155 "http-server-close" option achieves on small objects. There are mainly two
 4156 situations where this option may be useful :

- when the server is non-HTTP compliant and authenticates the connection instead of requests (eg: NTLM authentication)

- when the cost of establishing the connection to the server is significant compared to the cost of retrieving the associated object from the server.

This last case can happen when the server is a fast static server of cache.

In this case, the server will need to be properly tuned to support high enough connection counts because connections will last until the client sends another request.

If the client request has to go to another backend or another server due to content switching or the load balancing algorithm, the idle connection will immediately be closed and a new one re-opened. Option "prefer-last-server" is available to try optimize server selection so that if the server currently attached to an idle connection is usable, it will be used.

In general it is preferred to use "option http-server-close" with application servers, and some static servers might benefit from "option http-keep-alive".

At the moment, logs will not indicate whether requests came from the same session or not. The accept date reported in the logs corresponds to the end of the previous request, and the request time corresponds to the time spent waiting for a new request. The keep-alive request time is still bound to the timeout defined by "timeout http-keep-alive" or "timeout http-request" if not set.

This option disables and replaces any previous "option httpclose", "option http-server-close", "option forceclose" or "option http-tunnel". When backend and frontend options differ, all of these 4 options have precedence over "option http-keep-alive".

See also : "option forceclose", "option http-server-close",
"option prefer-last-server", "option http-pretend-keepalive",
"option httpclose", and "I.I. The HTTP transaction model".

option http-no-delay

no option http-no-delay

Instruct the system to favor low interactive delays over performance in HTTP

May be used in sections : defaults | frontend | listen | backend
yes | yes | yes | yes

Arguments : none

In HTTP, each payload is unidirectional and has no notion of interactivity. Any agent is expected to queue data somewhat for a reasonably low delay.

There are some very rare server-to-server applications that abuse the HTTP

protocol and expect the payload phase to be highly interactive, with many

interleaved data chunks in both directions within a single request. This is

absolutely not supported by the HTTP specification and will not work across

most proxies or servers. When such applications attempt to do this through

haproxy, it works but they will experience high delays due to the network

optimizations which favor performance by instructing the system to wait for

enough data to be available in order to only send full packets. Typical

delays are around 200 ms per round trip. Note that this only happens with

abnormal uses. Normal uses such as CONNECT requests nor WebSockets are not

affected.

When "option http-no-delay" is present in either the frontend or the backend

used by a connection, all such optimizations will be disabled in order to

make the exchanges as fast as possible. Of course this offers no guarantee on

the functionality, as it may break at any other place. But if it works via

haproxy, it will work as fast as possible. This option should never be used

by default, and should never be used at all unless such a buggy application

is discovered. The impact of using this option is an increase of bandwidth

usage and CPU usage, which may significantly lower performance in high

latency environments.

option http-pretend-keepalive

no option http-pretend-keepalive

Define whether haproxy will announce keepalive to the server or not

May be used in sections : defaults | frontend | listen | backend
yes | yes | yes | yes

Arguments : none

When running with "option http-server-close" or "option forceclose", haproxy adds a "Connection: close" header to the request forwarded to the server. Unfortunately, when some servers see this header, they automatically refrain from using the chunked encoding for responses of unknown length, while this is totally unrelated. The immediate effect is that this prevents haproxy from maintaining the client connection alive. A second effect is that a client or a cache could receive an incomplete response without being aware of it, and consider the response complete.

By setting "option http-pretend-keepalive", haproxy will make the server believe it will keep the connection alive. The server will then not fall back to the abnormal undesired above. When haproxy gets the whole response, it will close the connection with the server just as it would do with the "forceclose" option. That way the client gets a normal response and the connection is correctly closed on the server side.

It is recommended not to enable this option by default, because most servers will more efficiently close the connection themselves after the last packet, and release its buffers slightly earlier. Also, the added packet on the network could slightly reduce the overall peak performance. However it is worth noting that when this option is enabled, haproxy will have slightly less work to do. So if haproxy is the bottleneck on the whole architecture, enabling this option might save a few CPU cycles.

This option may be set both in a frontend and in a backend. It is enabled if at least one of the frontend or backend holding a connection has it enabled. This option may be combined with "option httpclose", which will cause keepalive to be announced to the server and close to be announced to the client. This practice is discouraged though.

If this option has been enabled in a "defaults" section, it can be disabled in a specific instance by prepending the "no" keyword before it.

See also : "option forceclose", "option http-server-close", and
"option http-keep-alive"

option http-server-close

no option http-server-close

Enable or disable HTTP connection closing on the server side

May be used in sections : defaults | frontend | listen | backend
yes | yes | yes | yes

Arguments : none

By default HAProxy operates in keep-alive mode with regards to persistent connections: for each connection it processes each request and response, and leaves the connection idle on both sides between the end of a response and the start of a new request. This mode may be changed by several options such as "option http-server-close", "option forceclose", "option httpclose" or "option http-tunnel". Setting "option http-server-close" enables HTTP connection-close mode on the server side while keeping the ability to support HTTP keep-alive and pipelining on the client side. This provides the lowest latency on the client side (slow network) and the fastest session reuse on the server side to save server resources, similarly to "option forceclose". It also permits non-keepalive capable servers to be served in keep-alive mode to the clients if they conform to the requirements of RFC2616. Please note that some servers do not always conform to those requirements when they see

"Connection: close" in the request. The effect will be that keep-alive will never be used. A workaround consists in enabling "option http-pretend-keepalive".

At the moment, logs will not indicate whether requests came from the same session or not. The accept date reported in the logs corresponds to the end of the previous request, and the request time corresponds to the time spent waiting for a new request. The keep-alive request time is still bound to the timeout defined by "timeout http-keep-alive" or "timeout http-request" if not set.

This option may be set both in a frontend and in a backend. It is enabled if at least one of the frontend or backend holding a connection has it enabled. It disables and replaces any previous "option httpclose", "option forceclose", "option http-tunnel" or "option http-keep-alive". Please check section 4 ("Proxies") to see how this option combines with others when frontend and backend options differ.

If this option has been enabled in a "defaults" section, it can be disabled in a specific instance by prepending the "no" keyword before it.

See also : "option forceclose", "option http-pretend-keepalive", "option httpclose", "option http-keep-alive", and "1.1. The HTTP transaction model".

option http-tunnel

no option http-tunnel

Disable or enable HTTP connection processing after first transaction

May be used in sections : defaults | frontend | listen | backend

yes | yes | yes | yes

Arguments : none

By default HAProxy operates in keep-alive mode with regards to persistent connections: for each connection it processes each request and response, and leaves the connection idle on both sides between the end of a response and the start of a new request. This mode may be changed by several options such as "option http-server-close", "option forceclose", "option httpclose" or "option http-tunnel".

Option "http-tunnel" disables any HTTP processing past the first request and the first response. This is the mode which was used by default in versions 1.0 to 1.5-dev21. It is the mode with the lowest processing overhead, which is normally not needed anymore unless in very specific cases such as when using an in-house protocol that looks like HTTP but is not compatible, or just to log one request per client in order to reduce log size. Note that everything which works at the HTTP level, including header parsing/addition, cookie processing or content switching will only work for the first request and will be ignored after the first response.

If this option has been enabled in a "defaults" section, it can be disabled in a specific instance by prepending the "no" keyword before it.

See also : "option forceclose", "option http-server-close", "option httpclose", "option http-keep-alive", and "1.1. The HTTP transaction model".

option http-use-proxy-header

no option http-use-proxy-header

Make use of non-standard Proxy-Connection header instead of Connection

May be used in sections : defaults | frontend | listen | backend

yes | yes | yes | no

Arguments : none

While RFC2616 explicitly states that HTTP/1.1 agents must use the Connection header to indicate their wish of persistent or non-persistent connections, both browsers and proxies ignore this header for proxied connections and make use of the undocumented, non-standard Proxy-Connection header instead. The issue begins when trying to put a load balancer between browsers and such proxies, because there will be a difference between what haproxy understands and what the client and the proxy agree on.

By setting this option in a frontend, haproxy can automatically switch to use that non-standard header if it sees proxied requests. A proxied request is defined here as one where the URI begins with neither a '/' nor a '*'. The choice of header only affects requests passing through proxies making use of one of the "httpclose", "forceclose" and "http-server-close" options. Note that this option can only be specified in a frontend and will affect the request along its whole life.

Also, when this option is set, a request which requires authentication will automatically switch to use proxy authentication headers if it is itself a proxied request. That makes it possible to check or enforce authentication in front of an existing proxy.

This option should normally never be used, except in front of a proxy.

See also : "option httpclose", "option forceclose" and "option

http-server-close".

option httpchk

option httpchk <uri>

option httpchk <method> <uri>

option httpchk <method> <uri> <version>

Enable HTTP protocol to check on the servers health

May be used in sections : defaults | frontend | listen | backend

yes | no | yes | yes

Arguments :

<method> is the optional HTTP method used with the requests. When not set, the "OPTIONS" method is used, as it generally requires low server processing and is easy to filter out from the logs. Any method may be used, though it is not recommended to invent non-standard ones.

<uri> is the URI referenced in the HTTP requests. It defaults to " / " which is accessible by default on almost any server, but may be changed to any other URI. Query strings are permitted.

<version> is the optional HTTP version string. It defaults to "HTTP/1.0" but some servers might behave incorrectly in HTTP 1.0, so turning it to HTTP/1.1 may sometimes help. Note that the Host field is mandatory in HTTP/1.1, and as a trick, it is possible to pass it after "\r\n" following the version string.

By default, server health checks only consist in trying to establish a TCP connection. When "option httpchk" is specified, a complete HTTP request is sent once the TCP connection is established, and responses 2xx and 3xx are considered valid, while all other ones indicate a server failure, including the lack of any response.

The port and interval are specified in the server configuration.

This option does not necessarily require an HTTP backend, it also works with plain TCP backends. This is particularly useful to check simple scripts bound to some dedicated ports using the inetd daemon.

Examples :

Relay HTTPS traffic to Apache instance and check service availability

4420

4421 # using HTTP request "OPTIONS * HTTP/1.1" on port 80.

4422 backend https_relay

4423 mode tcp

4424 option httpchk OPTIONS * HTTP/1.1\r\nHost:\ www

4425 server apache1 192.168.1.1:443 check port 80

4426 See also : "option ssl-hello-chk", "option smtpchk", "option mysql-check",
4427 "option option psql-check", "http-check" and the "check", "port" and
4428 "inter" server options.
4429
4430

4431 option httpclose

4432 no option httpclose

4433 Enable or disable passive HTTP connection closing

4434 May be used in sections : defaults | frontend | listen | backend

4435 yes | yes | yes | yes

4436 Arguments : none

4437

4438

4439 By default HAProxy operates in keep-alive mode with regards to persistent
4440 connections: for each connection it processes each request and response, and
4441 leaves the connection idle on both sides between the end of a response and
4442 the start of a new request. This mode may be changed by several options such
4443 as "option http-server-close", "option forceclose", "option httpclose" or
4444 "option http-tunnel".
4445

4446 If "option httpclose" is set, HAProxy will work in HTTP tunnel mode and check
4447 if a "Connection: close" header is already set in each direction, and will
4448 add one if missing. Each end should react to this by actively closing the TCP
4449 connection after each transfer, thus resulting in a switch to the HTTP close
4450 mode. Any "Connection" header different from "close" will also be removed.
4451 Note that this option is deprecated since what it does is very cheap but not
4452 reliable. Using "option http-server-close" or "option forceclose" is strongly
4453 recommended instead.
4454

4455 It seldom happens that some servers incorrectly ignore this header and do not
4456 close the connection even though they reply "Connection: close". For this
4457 reason, they are not compatible with older HTTP 1.0 browsers. If this happens
4458 it is possible to use the "option forceclose" which actively closes the
4459 request connection once the server responds. Option "forceclose" also
4460 releases the server connection earlier because it does not have to wait for
4461 the client to acknowledge it.
4462

4463 This option may be set both in a frontend and in a backend. It is enabled if
4464 at least one of the frontend or backend holding a connection has it enabled.
4465 It disables and replaces any previous "option http-server-close",
4466 "option forceclose", "option http-keep-alive" or "option http-tunnel". Please
4467 check section 4 ("Proxies") to see how this option combines with others when
4468 frontend and backend options differ.
4469

4470 If this option has been enabled in a "defaults" section, it can be disabled
4471 in a specific instance by prepending the "no" keyword before it.

4472 See also : "option forceclose", "option http-server-close" and
4473 "1.1. The HTTP transaction model".
4474

4475 option httplog [clf]

4476 Enable logging of HTTP request, session state and timers

4477 May be used in sections : defaults | frontend | listen | backend

4478 yes | yes | yes | yes

4479 Arguments :

4480 clf if the "clf" argument is added, then the output format will be

4481 the CLF format instead of HAProxy's default HTTP format. You can
4482 use this when you need to feed HAProxy's logs through a specific
4483 log analyser which only support the CLF format and which is not
4484
4485

4486 extensible.

4487 By default, the log output format is very poor, as it only contains the
4488 source and destination addresses, and the instance name. By specifying
4489 "option httplog", each log line turns into a much richer format including,
4490 but not limited to, the HTTP request, the connection timers, the session
4491 status, the connections numbers, the captured headers and cookies, the
4492 frontend, backend and server name, and of course the source address and
4493 ports.
4494

4495 This option may be set either in the frontend or the backend.

4496 Specifying only "option httplog" will automatically clear the 'clf' mode
4497 if it was set by default.
4498

4499 See also : section 8 about logging.
4500

4501 option http_proxy

4502 no option http_proxy

4503 Enable or disable plain HTTP proxy mode

4504 May be used in sections : defaults | frontend | listen | backend

4505 yes | yes | yes | yes

4506 Arguments : none

4507

4508

4509

4510

4511

4512

4513

4514

4515

4516

4517

4518

4519

4520

4521

4522

4523

4524

4525

4526

4527

4528

4529

4530

4531

4532

4533

4534

4535

4536

4537

4538

4539

4540

4541

4542

4543

4544

4545

4546

4547

4548

4549

4550

4551 timeouts and low amounts of exchanged data such as telnet session. If the
4552 server suddenly disappears, the output data accumulates in the system's
4553 socket buffers, both timeouts are correctly refreshed, and there is no way
4554 to know the server does not receive them, so we don't timeout. However, when
4555 the underlying protocol always echoes sent data, it would be enough by itself
4556 to detect the issue using the read timeout. Note that this problem does not
4557 happen with more verbose protocols because data won't accumulate long in the
4558 socket buffers.

4559 When this option is set on the frontend, it will disable read timeout updates
4560 on data sent to the client. There probably is little use of this case. When
4561 the option is set on the backend, it will disable read timeout updates on
4562 data sent to the server. Doing so will typically break large HTTP posts from
4563 slow lines, so use it with caution.

4564
4565 Note: older versions used to call this setting "option independent-streams"
4566 with a spelling mistake. This spelling is still supported but
4567 deprecated.

4568
4569 See also : "timeout client", "timeout server" and "timeout tunnel"

option ldap-check

4573 Use LDAPv3 health checks for server testing

4574 May be used in sections : defaults | frontend | listen | backend
4575 yes | no | yes | yes
4576
4577 Arguments : none

4578
4579 It is possible to test that the server correctly talks LDAPv3 instead of just
4580 testing that it accepts the TCP connection. When this option is set, an
4581 LDAPv3 anonymous simple bind message is sent to the server, and the response
4582 is analyzed to find an LDAPv3 bind response message.

4583
4584 The server is considered valid only when the LDAP response contains success
4585 resultCode (http://tools.ietf.org/html/rfc4511#section-4.1.9).

4586
4587 Logging of bind requests is server dependent see your documentation how to
4588 configure it.

4589 Example :

4590 option ldap-check

4591
4592 See also : "option httpchk"

option log-health-checks

4596 no option log-health-checks

4597 Enable or disable logging of health checks status updates

4598 May be used in sections : defaults | frontend | listen | backend
4599 yes | no | yes | yes
4600
4601 Arguments : none

4602
4603 By default, failed health check are logged if server is UP and successful
4604 health checks are logged if server is DOWN, so the amount of additional
4605 information is limited.

4606
4607 When this option is enabled, any change of the health check status or to
4608 the server's health will be logged, so that it becomes possible to know
4609 that a server was failing occasional checks before crashing, or exactly when
4610 it failed to respond a valid HTTP status, then when the port started to
4611 reject connections, then when the server stopped responding at all.

4612
4613 Note that status changes not caused by health checks (eg: enable/disable on
4614 the CLI) are intentionally not logged by this option.

4616 See also: "option httpchk", "option ldap-check", "option mysql-check",
4617 "option pgsql-check", "option redis-check", "option smtpchk",
4618 "option tcp-check", "log" and section 8 about logging.

option log-separate-errors

4621 no option log-separate-errors

4622 Change log level for non-completely successful connections

4623 May be used in sections : defaults | frontend | listen | backend
4624 yes | yes | yes | no
4625
4626 Arguments : none

4627
4628 Sometimes looking for errors in logs is not easy. This option makes haproxy
4629 raise the level of logs containing potentially interesting information such
4630 as errors, timeouts, retries, redispatches, or HTTP status codes 5xx. The
4631 level changes from "info" to "err". This makes it possible to log them
4632 separately to a different file with most syslog daemons. Be careful not to
4633 remove them from the original file, otherwise you would lose ordering which
4634 provides very important information.

4635
4636 Using this option, large sites dealing with several thousand connections per
4637 second may log normal traffic to a rotating buffer and only archive smaller
4638 error logs.

4639
4640 See also : "log", "dontlognull", "dontlog-normal" and section 8 about
4641 logging.

option logasap

4644 no option logasap

4645 Enable or disable early logging of HTTP requests

4646 May be used in sections : defaults | frontend | listen | backend
4647 yes | yes | yes | no
4648
4649 Arguments : none

4650
4651 By default, HTTP requests are logged upon termination so that the total
4652 transfer time and the number of bytes appear in the logs. When large objects
4653 are being transferred, it may take a while before the request appears in the
4654 logs. Using "option logasap", the request gets logged as soon as the server
4655 sends the complete headers. The only missing information in the logs will be
4656 the total number of bytes which will indicate everything except the amount
4657 of data transferred, and the total time which will not take the transfer
4658 time into account. In such a situation, it's a good practice to capture the
4659 "Content-Length" response header so that the logs at least indicate how many
4660 bytes are expected to be transferred.

4661 Examples :

4662 listen http_proxy 0.0.0.0:80

4663 mode http

4664 option httplog

4665 option logasap

4666 log 192.168.2.200 local3

4667
4668 >>> Feb 6 12:14:14 localhost \

4669 haproxy[14389]: 10.0.1.2:33317 [06/Feb/2009:12:14:14.655] http-in \

4670 static/srv1 9/10/7/14/+30 200 +243 - - - - - 3/1/1/0 1/0 \

4671 "GET /image.iso HTTP/1.0"

4672
4673 See also : "option httplog", "capture response header", and section 8 about
4674 logging.

4675 option mysql-check [user <username> [post-41]]

4676 Use MySQL health checks for server testing

4677 May be used in sections : defaults | frontend | listen | backend

4680

4681 Arguments : yes | no | yes | yes
 4682 <username> This is the username which will be used when connecting to MySQL
 4683 server.
 4684 post-41 Send post v4.1 client compatible checks
 4685
 4686 If you specify a username, the check consists of sending two MySQL packet,
 4687 one Client Authentication packet, and one QUIT packet, to correctly close
 4688 MySQL session. We then parse the MySQL Handshake Initialisation packet and/or
 4689 Error packet. It is a basic but useful test which does not produce error nor
 4690 aborted connect on the server. However, it requires adding an authorization
 4691 in the MySQL table, like this :
 4692
 4693 USE mysql;
 4694 INSERT INTO user (Host,User) values ('<ip_of_haproxy>', '<username>');
 4695 FLUSH PRIVILEGES;
 4696
 4697 If you don't specify a username (it is deprecated and not recommended), the
 4698 check only consists in parsing the MySQL Handshake Initialisation packet or
 4699 Error packet, we don't send anything in this mode. It was reported that it
 4700 can generate lockout if check is too frequent and/or if there is not enough
 4701 traffic. In fact, you need in this case to check MySQL "max_connect_errors"
 4702 value as if a connection is established successfully within fewer than MySQL
 4703 "max_connect_errors" attempts after a previous connection was interrupted,
 4704 the error count for the host is cleared to zero. If HAProxy's server get
 4705 blocked, the "FLUSH HOSTS" statement is the only way to unblock it.
 4706
 4707 Remember that this does not check database presence nor database consistency.
 4708 To do this, you can use an external check with xinetd for example.
 4709
 4710 The check requires MySQL >=3.22, for older version, please use TCP check.
 4711
 4712 Most often, an incoming MySQL server needs to see the client's IP address for
 4713 various purposes, including IP privilege matching and connection logging.
 4714 When possible, it is often wise to masquerade the client's IP address when
 4715 connecting to the server using the "usescrc" argument of the "source" keyword,
 4716 which requires the ctproxy feature to be compiled in, and the MySQL server
 4717 to route the client via the machine hosting haproxy.
 4718
 4719 See also: "option httpchk"
 4720
 4721 option nolinger
 4722 no option nolinger
 4723 Enable or disable immediate session resource cleaning after close
 4724 May be used in sections: defaults | frontend | listen | backend
 4725 yes | yes | yes | yes
 4726
 4727 Arguments : none
 4728
 4729 When clients or servers abort connections in a dirty way (eg: they are
 4730 physically disconnected), the session timeouts triggers and the session is
 4731 closed. But it will remain in FIN_WAIT1 state for some time in the system,
 4732 using some resources and possibly limiting the ability to establish newer
 4733 connections.
 4734
 4735 When this happens, it is possible to activate "option nolinger" which forces
 4736 the system to immediately remove any socket's pending data on close. Thus,
 4737 the session is instantly purged from the system's tables. This usually has
 4738 side effects such as increased number of TCP resets due to old retransmits
 4739 getting immediately rejected. Some firewalls may sometimes complain about
 4740 this too.
 4741
 4742 For this reason, it is not recommended to use this option when not absolutely
 4743 needed. You know that you need it when you have thousands of FIN_WAIT1
 4744 sessions on your system (TIME_WAIT ones do not count).
 4745

4746 This option may be used both on frontends and backends, depending on the side
 4747 where it is required. Use it on the frontend for clients, and on the backend
 4748 for servers.
 4749
 4750 If this option has been enabled in a "defaults" section, it can be disabled
 4751 in a specific instance by prepending the "no" keyword before it.
 4752
 4753 option originalto [except <network>] [header <name>]
 4754 Enable insertion of the X-Original-To header to requests sent to servers
 4755 May be used in sections : defaults | frontend | listen | backend
 4756 yes | yes | yes | yes
 4757
 4758 Arguments :
 4759 <network> is an optional argument used to disable this option for sources
 4760 matching <network>
 4761 <name> an optional argument to specify a different "X-Original-To"
 4762 header name.
 4763
 4764 Since HAProxy can work in transparent mode, every request from a client can
 4765 be redirected to the proxy and HAProxy itself can proxy every request to a
 4766 complex SQUID environment and the destination host from SO_ORIGINAL_DST will
 4767 be lost. This is annoying when you want access rules based on destination ip
 4768 addresses. To solve this problem, a new HTTP header "X-Original-To" may be
 4769 added by HAProxy to all requests sent to the server. This header contains a
 4770 value representing the original destination IP address. Since this must be
 4771 configured to always use the last occurrence of this header only. Note that
 4772 only the last occurrence of the header must be used, since it is really
 4773 possible that the client has already brought one.
 4774
 4775 The keyword "header" may be used to supply a different header name to replace
 4776 the default "X-Original-To". This can be useful where you might already
 4777 have a "X-Original-To" header from a different application, and you need
 4778 preserve it. Also if your backend server doesn't use the "X-Original-To"
 4779 header and requires different one.
 4780
 4781 Sometimes, a same HAProxy instance may be shared between a direct client
 4782 access and a reverse-proxy access (for instance when an SSL reverse-proxy is
 4783 used to decrypt HTTPS traffic). It is possible to disable the addition of the
 4784 header for a known source address or network by adding the "except" keyword
 4785 followed by the network address. In this case, any source IP matching the
 4786 network will not cause an addition of this header. Most common uses are with
 4787 private networks or 127.0.0.1.
 4788
 4789 This option may be specified either in the frontend or in the backend. If at
 4790 least one of them uses it, the header will be added. Note that the backend's
 4791 setting of the header subargument takes precedence over the frontend's if
 4792 both are defined.
 4793
 4794 Examples :
 4795 # Original Destination address
 4796 frontend www
 4797 mode http
 4798 option originalto except 127.0.0.1
 4799
 4800 # Those servers want the IP Address in X-Client-Dst
 4801 backend www
 4802 mode http
 4803 option originalto header X-Client-Dst
 4804
 4805 See also : "option httpclose", "option http-server-close",
 4806 "option forceclose"
 4807
 4808 option persist
 4809
 4810

4811 no option persist
4812 Enable or disable forced persistence on down servers
4813 May be used in sections: defaults | frontend | listen | backend
4814 yes | no | yes | yes
4815 Arguments : none
4816
4817 When an HTTP request reaches a backend with a cookie which references a dead
4818 server, by default it is redispached to another server. It is possible to
4819 force the request to be sent to the dead server first using "option persist"
4820 if absolutely needed. A common use case is when servers are under extreme
4821 load and spend their time flapping. In this case, the users would still be
4822 directed to the server they opened the session on, in the hope they would be
4823 correctly served. It is recommended to use "option redispach" in conjunction
4824 with this option so that in the event it would not be possible to connect to
4825 the server at all (server definitely dead), the client would finally be
4826 redirected to another valid server.
4827
4828 If this option has been enabled in a "defaults" section, it can be disabled
4829 in a specific instance by prepending the "no" keyword before it.
4830
4831 See also : "option redispach", "retries", "force-persist"
4832
4833 option pgsqL-check [user <username>]
4834 Use PostgreSQL health checks for server testing
4835 May be used in sections : defaults | frontend | listen | backend
4836 yes | no | yes | yes
4837 Arguments :
4838 <username> This is the username which will be used when connecting to
4839 PostgreSQL server.
4840
4841 The check sends a PostgreSQL StartupMessage and waits for either
4842 Authentication request or ErrorResponse message. It is a basic but useful
4843 test which does not produce error nor aborted connect on the server.
4844 This check is identical with the "mysql-check".
4845
4846 See also: "option httpchk"
4847
4848
4849 option prefer-last-server
4850 no option prefer-last-server
4851 Allow multiple load balanced requests to remain on the same server
4852 May be used in sections: defaults | frontend | listen | backend
4853 yes | no | yes | yes
4854 Arguments : none
4855
4856 When the load balancing algorithm in use is not deterministic, and a previous
4857 request was sent to a server to which haproxy still holds a connection, it is
4858 sometimes desirable that subsequent requests on a same session go to the same
4859 server as much as possible. Note that this is different from persistence, as
4860 we only indicate a preference which haproxy tries to apply without any form
4861 of warranty. The real use is for keep-alive connections sent to servers. When
4862 this option is used, haproxy will try to reuse the same connection that is
4863 attached to the server instead of rebalancing to another server, causing a
4864 close of the connection. This can make sense for static file servers. It does
4865 not make much sense to use this in combination with hashing algorithms. Note,
4866 haproxy already automatically tries to stick to a server which sends a 401 or
4867 to a proxy which sends a 407 (authentication required). This is mandatory for
4868 use with the broken NTLM authentication challenge, and significantly helps in
4869 troubleshooting some faulty applications. Option prefer-last-server might be
4870 desirable in these environments as well, to avoid redistributing the traffic
4871 after every other response.
4872
4873 If this option has been enabled in a "defaults" section, it can be disabled
4874 in a specific instance by prepending the "no" keyword before it.
4875

4876 See also: "option http-keep-alive"
4877
4878
4879 option redispach
4880 no option redispach
4881 Enable or disable session redistribution in case of connection failure
4882 May be used in sections: defaults | frontend | listen | backend
4883 yes | no | yes | yes
4884 Arguments : none
4885
4886 In HTTP mode, if a server designated by a cookie is down, clients may
4887 definitely stick to it because they cannot flush the cookie, so they will not
4888 be able to access the service anymore.
4889
4890 Specifying "option redispach" will allow the proxy to break their
4891 persistence and redistribute them to a working server.
4892
4893 It also allows to retry last connection to another server in case of multiple
4894 connection failures. Of course, it requires having "retries" set to a nonzero
4895 value.
4896
4897 This form is the preferred form, which replaces both the "redispach" and
4898 "redispatch" keywords.
4899
4900 If this option has been enabled in a "defaults" section, it can be disabled
4901 in a specific instance by prepending the "no" keyword before it.
4902
4903 See also : "redispach", "retries", "force-persist"
4904
4905 option redis-check
4906 Use redis health checks for server testing
4907 May be used in sections : defaults | frontend | listen | backend
4908 yes | no | yes | yes
4909 Arguments : none
4910
4911 It is possible to test that the server correctly talks REDIS protocol instead
4912 of just testing that it accepts the TCP connection. When this option is set,
4913 a PING redis command is sent to the server, and the response is analyzed to
4914 find the "+PONG" response message.
4915
4916 Example :
4917 option redis-check
4918
4919 See also : "option httpchk"
4920
4921
4922 option smtpchk <hello> <domain>
4923 Use SMTP health checks for server testing
4924 May be used in sections : defaults | frontend | listen | backend
4925 yes | no | yes | yes
4926 Arguments :
4927 <hello> is an optional argument. It is the "hello" command to use. It can
4928 be either "HELO" (for SMTP) or "EHLO" (for ESTMP). All other
4929 values will be turned into the default command ("HELO").
4930
4931 <domain> is the domain name to present to the server. It may only be
4932 specified (and is mandatory) if the hello command has been
4933 specified. By default, "localhost" is used.
4934
4935 When "option smtpchk" is set, the health checks will consist in TCP
4936 connections followed by an SMTP command. By default, this command is
4937 "HELO localhost". The server's return code is analyzed and only return codes
4938 4939 4940

starting with a "2" will be considered as valid. All other responses, including a lack of response will constitute an error and will indicate a dead server.

This test is meant to be used with SMTP servers or relays. Depending on the request, it is possible that some servers do not log each connection attempt, so you may want to experiment to improve the behaviour. Using telnet on port 25 is often easier than adjusting the configuration.

Most often, an incoming SMTP server needs to see the client's IP address for various purposes, including spam filtering, anti-spoofing and logging. When possible, it is often wise to masquerade the client's IP address when connecting to the server using the "usesrc" argument of the "source" keyword, which requires the cttproxy feature to be compiled in.

Example :
option smtpchk HELO mydomain.org

See also : "option httpchk", "source"

option socket-stats
no option socket-stats

Enable or disable collecting & providing separate statistics for each socket.
May be used in sections : defaults | frontend | listen | backend
yes | yes | yes | no

Arguments : none

option splice-auto
no option splice-auto
Enable or disable automatic kernel acceleration on sockets in both directions
May be used in sections : defaults | frontend | listen | backend
yes | yes | yes | yes
Arguments : none

When this option is enabled either on a frontend or on a backend, haproxy will automatically evaluate the opportunity to use kernel tcp splicing to forward data between the client and the server, in either direction. Haproxy uses heuristics to estimate if kernel splicing might improve performance or not. Both directions are handled independently. Note that the heuristics used are not much aggressive in order to limit excessive use of splicing. This option requires splicing to be enabled at compile time, and may be globally disabled with the global option "nosplICE". Since splice uses pipes, using it requires that there are enough spare pipes.

Important note: kernel-based TCP splicing is a Linux-specific feature which first appeared in kernel 2.6.25. It offers kernel-based acceleration to transfer data between sockets without copying these data to user-space, thus providing noticeable performance gains and CPU cycles savings. Since many early implementations are buggy, corrupt data and/or are inefficient, this feature is not enabled by default, and it should be used with extreme care. While it is not possible to detect the correctness of an implementation. 2.6.29 is the first version offering a properly working implementation. In case of doubt, splicing may be globally disabled using the global "nosplICE" keyword.

Example :
option splice-auto

If this option has been enabled in a "defaults" section, it can be disabled in a specific instance by prepending the "no" keyword before it.

See also : "option splice-request", "option splice-response", and global options "nosplICE" and "maxpipes"

option splice-request
no option splice-request
Enable or disable automatic kernel acceleration on sockets for requests
May be used in sections : defaults | frontend | listen | backend
yes | yes | yes | yes
Arguments : none

When this option is enabled either on a frontend or on a backend, haproxy will use kernel tcp splicing whenever possible to forward data going from the client to the server. It might still use the recv/send scheme if there are no spare pipes left. This option requires splicing to be enabled at compile time, and may be globally disabled with the global option "nosplICE". Since splice uses pipes, using it requires that there are enough spare pipes.

Important note: see "option splice-auto" for usage limitations.

Example :
option splice-request

If this option has been enabled in a "defaults" section, it can be disabled in a specific instance by prepending the "no" keyword before it.

See also : "option splice-auto", "option splice-response", and global options "nosplICE" and "maxpipes"

option splice-response
no option splice-response
Enable or disable automatic kernel acceleration on sockets for responses
May be used in sections : defaults | frontend | listen | backend
yes | yes | yes | yes
Arguments : none

When this option is enabled either on a frontend or on a backend, haproxy will use kernel tcp splicing whenever possible to forward data going from the server to the client. It might still use the recv/send scheme if there are no spare pipes left. This option requires splicing to be enabled at compile time, and may be globally disabled with the global option "nosplICE". Since splice uses pipes, using it requires that there are enough spare pipes.

Important note: see "option splice-auto" for usage limitations.

Example :
option splice-response

If this option has been enabled in a "defaults" section, it can be disabled in a specific instance by prepending the "no" keyword before it.

See also : "option splice-auto", "option splice-request", and global options "nosplICE" and "maxpipes"

option srvtcpka
no option srvtcpka
Enable or disable the sending of TCP keepalive packets on the server side
May be used in sections : defaults | frontend | listen | backend
yes | no | yes | yes
Arguments : none

When there is a firewall or any session-aware component between a client and a server, and when the protocol involves very long sessions with long idle

periods (eg: remote desktops), there is a risk that one of the intermediate components decides to expire a session which has remained idle for too long.

Enabling socket-level TCP keep-alives makes the system regularly send packets to the other end of the connection, leaving it active. The delay between keep-alive probes is controlled by the system only and depends both on the operating system and its tuning parameters.

It is important to understand that keep-alive packets are neither emitted nor received at the application level. It is only the network stacks which sees them. For this reason, even if one side of the proxy already uses keep-alives to maintain its connection alive, those keep-alive packets will not be forwarded to the other side of the proxy.

Please note that this has nothing to do with HTTP keep-alive.

Using option "srvtcpka" enables the emission of TCP keep-alive probes on the server side of a connection, which should help when session expirations are noticed between HAProxy and a server.

If this option has been enabled in a "defaults" section, it can be disabled in a specific instance by prepending the "no" keyword before it.

See also : "option cli tcpka", "option tcpka"

option ssl-hello-chk

Use SSLv3 client hello health checks for server testing

May be used in sections : defaults | frontend | listen | backend

yes	no	yes	yes
-----	----	-----	-----

Arguments : none

When some SSL-based protocols are relayed in TCP mode through HAProxy, it is possible to test that the server correctly talks SSL instead of just testing that it accepts the TCP connection. When "option ssl-hello-chk" is set, pure SSLv3 client hello messages are sent once the connection is established to the server, and the response is analyzed to find an SSL server hello message. The server is considered valid only when the response contains this server hello message.

All servers tested till there correctly reply to SSLv3 client hello messages, and most servers tested do not even log the requests containing only hello messages, which is appreciable.

Note that this check works even when SSL support was not built into haproxy because it forges the SSL message. When SSL support is available, it is best to use native SSL health checks instead of this one.

See also: "option httpchk", "check-ssl"

option tcp-check

Perform health checks using tcp-check send/expect sequences

May be used in sections: defaults | frontend | listen | backend

yes	no	yes	yes
-----	----	-----	-----

This health check method is intended to be combined with "tcp-check" command lists in order to support send/expect types of health check sequences.

TCP checks currently support 4 modes of operations :

- no "tcp-check" directive : the health check only consists in a connection attempt, which remains the default mode.
- "tcp-check send" or "tcp-check send-binary" only is mentioned : this is used to send a string along with a connection opening. With some

protocols, it helps sending a "QUIT" message for example that prevents the server from logging a connection error for each health check. The check result will still be based on the ability to open the connection only.

- "tcp-check expect" only is mentioned : this is used to test a banner. The connection is opened and haproxy waits for the server to present some contents which must validate some rules. The check result will be based on the matching between the contents and the rules. This is suited for POP, IMAP, SMTP, FTP, SSH, TELNET.
 - both "tcp-check send" and "tcp-check expect" are mentioned : this is used to test a hello-type protocol. Haproxy sends a message, the server responds and its response is analysed. the check result will be based on the matching between the response contents and the rules. This is often suited for protocols which require a binding or a request/response model. LDAP, MySQL, Redis and SSL are example of such protocols, though they already all have their dedicated checks with a deeper understanding of the respective protocols.
- In this mode, many questions may be sent and many answers may be analysed.

Examples :

```
# perform a POP check (analyse only server's banner)
option tcp-check
tcp-check expect string +OK POP3 ready
```

```
# perform an IMAP check (analyse only server's banner)
option tcp-check
tcp-check expect string *OK IMAP4 ready
```

```
# look for the redis master server after ensuring it speaks well
# redis protocol, then it exits properly.
# (send a command then analyse the response 3 times)
option tcp-check
tcp-check send PING\r\n
```

```
tcp-check expect +PONG
tcp-check send info replication\r\n
```

```
tcp-check expect string role:master
tcp-check send QUIT\r\n
```

```
tcp-check expect string +OK
```

```
forge a HTTP request, then analyse the response
(send many headers before analyzing)
```

```
option tcp-check
tcp-check send HEAD /\ HTTP/1.1\r\n
```

```
tcp-check send Host:\ www.mydomain.com\r\n
```

```
tcp-check send User-Agent:\ HAProxy tcpcheck\r\n
```

```
tcp-check send \r\n
tcp-check expect rstring HTTP/1..\ (2..|3..)
```

See also : "tcp-check expect", "tcp-check send"

option tcp-smart-accept

no option tcp-smart-accept

Enable or disable the saving of one ACK packet during the accept sequence

May be used in sections : defaults | frontend | listen | backend

yes	yes	yes	no
-----	-----	-----	----

Arguments : none

When an HTTP connection request comes in, the system acknowledges it on behalf of HAProxy, then the client immediately sends its request, and the system acknowledges it too while it is notifying HAProxy about the new

connection. HAProxy then reads the request and responds. This means that we have one TCP ACK sent by the system for nothing, because the request could very well be acknowledged by HAProxy when it sends its response.

For this reason, in HTTP mode, HAProxy automatically asks the system to avoid sending this useless ACK on platforms which support it (currently at least Linux). It must not cause any problem, because the system will send it anyway after 40 ms if the response takes more time than expected to come.

During complex network debugging sessions, it may be desirable to disable this optimization because delayed ACKs can make troubleshooting more complex when trying to identify where packets are delayed. It is then possible to fall back to normal behaviour by specifying "no option tcp-smart-accept".

It is also possible to force it for non-HTTP proxies by simply specifying "option tcp-smart-accept". For instance, it can make sense with some services such as SMTP where the server speaks first.

It is recommended to avoid forcing this option in a defaults section. In case of doubt, consider setting it back to automatic values by prepending the "default" keyword before it, or disabling it using the "no" keyword.

See also : "option tcp-smart-connect"

option tcp-smart-connect

no option tcp-smart-connect

Enable or disable the saving of one ACK packet during the connect sequence
May be used in sections : defaults | frontend | listen | backend

defaults yes | no | yes | yes

Arguments : none

On certain systems (at least Linux), HAProxy can ask the kernel not to immediately send an empty ACK upon a connection request, but to directly send the buffer request instead. This saves one packet on the network and thus boosts performance. It can also be useful for some servers, because they immediately get the request along with the incoming connection.

This feature is enabled when "option tcp-smart-connect" is set in a backend. It is not enabled by default because it makes network troubleshooting more complex.

It only makes sense to enable it with protocols where the client speaks first such as HTTP. In other situations, if there is no data to send in place of the ACK, a normal ACK is sent.

If this option has been enabled in a "defaults" section, it can be disabled in a specific instance by prepending the "no" keyword before it.

See also : "option tcp-smart-accept"

option tcpka

Enable or disable the sending of TCP keepalive packets on both sides

May be used in sections : defaults | frontend | listen | backend

defaults yes | yes | yes | yes

Arguments : none

When there is a firewall or any session-aware component between a client and a server, and when the protocol involves very long sessions with long idle periods (eg: remote desktops), there is a risk that one of the intermediate components decides to expire a session which has remained idle for too long.

Enabling socket-level TCP keep-alives makes the system regularly send packets to the other end of the connection, leaving it active. The delay between

keep-alive probes is controlled by the system only and depends both on the operating system and its tuning parameters.

It is important to understand that keep-alive packets are neither emitted nor received at the application level. It is only the network stacks which sees them. For this reason, even if one side of the proxy already uses keep-alives to maintain its connection alive, those keep-alive packets will not be forwarded to the other side of the proxy.

Please note that this has nothing to do with HTTP keep-alive.

Using option "tcpka" enables the emission of TCP keep-alive probes on both the client and server sides of a connection. Note that this is meaningful only in "defaults" or "listen" sections. If this option is used in a frontend, only the client side will get keep-alives, and if this option is used in a backend, only the server side will get keep-alives. For this reason, it is strongly recommended to explicitly use "option cliptcpka" and "option srvtcpka" when the configuration is split between frontends and backends.

See also : "option cliptcpka", "option srvtcpka"

option tcplog

Enable advanced logging of TCP connections with session state and timers
May be used in sections : defaults | frontend | listen | backend

defaults yes | yes | yes | yes

Arguments : none

By default, the log output format is very poor, as it only contains the source and destination addresses, and the instance name. By specifying "option tcplog", each log line turns into a much richer format including, but not limited to, the connection timers, the session status, the connections numbers, the frontend, backend and server name, and of course the source address and ports. This option is useful for pure TCP proxies in order to find which of the client or server disconnects or times out. For normal HTTP proxies, it's better to use "option httplog" which is even more complete.

This option may be set either in the frontend or the backend.

See also : "option httplog", and section 8 about logging.

option transparent

no option transparent

Enable client-side transparent proxying

May be used in sections : defaults | frontend | listen | backend

defaults yes | no | yes | yes

Arguments : none

This option was introduced in order to provide layer 7 persistence to layer 3 load balancers. The idea is to use the OS's ability to redirect an incoming connection for a remote address to a local process (here HAProxy), and let this process know what address was initially requested. When this option is used, sessions without cookies will be forwarded to the original destination IP address of the incoming request (which should match that of another equipment), while requests with cookies will still be forwarded to the appropriate server.

Note that contrary to a common belief, this option does NOT make HAProxy present the client's IP to the server when establishing the connection.

See also: the "us-src" argument of the "source" keyword, and the "transparent" option of the "bind" keyword.

5300

5331 persist rdp-cookie
 5332 persist rdp-cookie(<name>)
 5333 Enable RDP cookie-based persistence
 5334 May be used in sections : defaults | frontend | listen | backend
 5335 yes | no | yes | yes
 5336
 5337 Arguments :
 5338 <name> is the optional name of the RDP cookie to check. If omitted, the
 5339 default cookie name "msts" will be used. There currently is no
 5340 valid reason to change this name.
 5341
 5342 This statement enables persistence based on an RDP cookie. The RDP cookie
 5343 contains all information required to find the server in the list of known
 5344 servers. So when this option is set in the backend, the request is analysed
 5345 and if an RDP cookie is found, it is decoded. If it matches a known server
 5346 which is still UP (or if "option persist" is set), then the connection is
 5347 forwarded to this server.
 5348
 5349 Note that this only makes sense in a TCP backend, but for this to work, the
 5350 frontend must have waited long enough to ensure that an RDP cookie is present
 5351 in the request buffer. This is the same requirement as with the "rdp-cookie"
 5352 load-balancing method. Thus it is highly recommended to put all statements in
 5353 a single "listen" section.
 5354
 5355 Also, it is important to understand that the terminal server will emit this
 5356 RDP cookie only if it is configured for "token redirection mode", which means
 5357 that the "ip address redirection" option is disabled.
 5358
 5359 Example :
 5360 listen tse-farm
 5361 bind :3389
 5362 # wait up to 5s for an RDP cookie in the request
 5363 tcp-request inspect-delay 5s
 5364 tcp-request content accept if RDP_COOKIE
 5365 # apply RDP cookie persistence
 5366 persist rdp-cookie
 5367 # if server is unknown, let's balance on the same cookie.
 5368 # alternatively, "balance leastconn" may be useful too.
 5369 balance rdp-cookie
 5370 server srv1 1.1.1.1:3389
 5371 server srv2 1.1.1.2:3389
 5372
 5373 See also : "balance rdp-cookie", "tcp-request", the "req_rdp_cookie" ACL and
 5374 the rdp_cookie pattern fetch function.
 5375

5376 rate-limit sessions <rate>
 5377 Set a limit on the number of new sessions accepted per second on a frontend
 5378 May be used in sections : defaults | frontend | listen | backend
 5379 yes | yes | yes | no
 5380
 5381 Arguments :
 5382 <rate> The <rate> parameter is an integer designating the maximum number
 5383 of new sessions per second to accept on the frontend.
 5384
 5385 When the frontend reaches the specified number of new sessions per second, it
 5386 stops accepting new connections until the rate drops below the limit again.
 5387 During this time, the pending sessions will be kept in the socket's backlog
 5388 (in system buffers) and haproxy will not even be aware that sessions are
 5389 pending. When applying very low limit on a highly loaded service, it may make
 5390 sense to increase the socket's backlog using the "backlog" keyword.
 5391
 5392 This feature is particularly efficient at blocking connection-based attacks
 5393 or service abuse on fragile servers. Since the session rate is measured every
 5394 millisecond, it is extremely accurate. Also, the limit applies immediately,
 5395 no delay is needed at all to detect the threshold.

5396
 5397 Example : limit the connection rate on SMTP to 10 per second max
 5398 listen smtp
 5399 mode tcp
 5400 bind :25
 5401 rate-limit sessions 10
 5402 server 127.0.0.1:1025
 5403
 5404 Note : when the maximum rate is reached, the frontend's status is not changed
 5405 but its sockets appear as "WAITING" in the statistics if the
 5406 "socket-stats" option is enabled.
 5407
 5408 See also : the "backlog" keyword and the "fe_sess_rate" ACL criterion.
 5409
 5410 redirect location <loc> [code <code>] <option> [{if | unless} <condition>]
 5411 redirect prefix <prfx> [code <code>] <option> [{if | unless} <condition>]
 5412 redirect scheme <sch> [code <code>] <option> [{if | unless} <condition>]
 5413 Return an HTTP redirection if/unless a condition is matched
 5414 May be used in sections : defaults | frontend | listen | backend
 5415 no | yes | yes | yes
 5416
 5417 If/unless the condition is matched, the HTTP request will lead to a redirect
 5418 response. If no condition is specified, the redirect applies unconditionally.
 5419
 5420 Arguments :
 5421 <loc> With "redirect location", the exact value in <loc> is placed into
 5422 the HTTP "Location" header. When used in an "http-request" rule,
 5423 <loc> value follows the log-format rules and can include some
 5424 dynamic values (see Custom Log Format in section 8.2.4).
 5425
 5426 With "redirect prefix", the "Location" header is built from the
 5427 concatenation of <prfx> and the complete URI path, including the
 5428 query string, unless the "drop-query" option is specified (see
 5429 below). As a special case, if <prfx> equals exactly "/", then
 5430 nothing is inserted before the original URI. It allows one to
 5431 redirect to the same URL (for instance, to insert a cookie). When
 5432 used in an "http-request" rule, <prfx> value follows the log-format
 5433 rules and can include some dynamic values (see Custom Log Format
 5434 in section 8.2.4).
 5435
 5436 With "redirect scheme", then the "Location" header is built by
 5437 concatenating <sch> with "://" then the first occurrence of the
 5438 "Host" header, and then the URI path, including the query string
 5439 unless the "drop-query" option is specified (see below). If no
 5440 path is found or if the path is "*", then "/" is used instead. If
 5441 no "Host" header is found, then an empty host component will be
 5442 returned, which most recent browsers interpret as redirecting to
 5443 the same host. This directive is mostly used to redirect HTTP to
 5444 HTTPS. When used in an "http-request" rule, <sch> value follows
 5445 the log-format rules and can include some dynamic values (see
 5446 Custom Log Format in section 8.2.4).
 5447
 5448 <code> The code is optional. It indicates which type of HTTP redirection
 5449 is desired. Only codes 301, 302, 303, 307 and 308 are supported,
 5450 with 302 used by default if no code is specified. 301 means
 5451 "Moved permanently", and a browser may cache the location. 302
 5452 means "Moved temporarily" and means that the browser should not
 5453 cache the redirection. 303 is equivalent to 302 except that the
 5454 browser will fetch the location with a GET method. 307 is just
 5455 like 302 but makes it clear that the same method must be reused.
 5456 Likewise, 308 replaces 301 if the same method must be used.
 5457
 5458 There are several options which can be specified to adjust the
 5459 expected behaviour of a redirection :
 5460

5461 - "drop-query"
 5462 When this keyword is used in a prefix-based redirection, then the
 5463 location will be set without any possible query-string, which is useful
 5464 for directing users to a non-secure page for instance. It has no effect
 5465 with a location-type redirect.
 5466

5467 - "append-slash"
 5468 This keyword may be used in conjunction with "drop-query" to redirect
 5469 users who use a URL not ending with a '/' to the same one with the '/'.
 5470 It can be useful to ensure that search engines will only see one URL.
 5471 For this, a return code 301 is preferred.
 5472

5473 - "set-cookie NAME[=value]"
 5474 A "Set-Cookie" header will be added with NAME (and optionally "=value")
 5475 to the response. This is sometimes used to indicate that a user has
 5476 been seen, for instance to protect against some types of DoS. No other
 5477 cookie option is added, so the cookie will be a session cookie. Note
 5478 that for a browser, a sole cookie name without an equal sign is
 5479 different from a cookie with an equal sign.
 5480

5481 - "clear-cookie NAME[=]"
 5482 A "Set-Cookie" header will be added with NAME (and optionally "=",) but
 5483 with the "Max-Age" attribute set to zero. This will tell the browser to
 5484 delete this cookie. It is useful for instance on logout pages. It is
 5485 important to note that clearing the cookie "NAME" will not remove a
 5486 cookie set with "NAME=value". You have to clear the cookie "NAME=" for
 5487 that, because the browser makes the difference.
 5488

5489 Example: move the login URL only to HTTPS.

```
5490 acl clear      dst_port 80
5491 acl secure     dst_port 8080
5492 acl login_page url_beg /login
5493 acl logout     url_beg /logout
5494 acl uid_given  url_reg /login?userid=[^&]+
5495 acl cookie_set hdr_sub(cookie) SEEN=1
```

```
5496 redirect prefix https://mysite.com set-cookie SEEN=1 if !cookie_set
5497 redirect prefix https://mysite.com if login_page !secure
5498 redirect prefix http://mysite.com drop-query if login_page !uid given
5499 redirect location http://mysite.com/ if !login_page secure
5500 redirect location / clear-cookie USERID= if logout
```

5501 Example: send redirects for request for articles without a '/'.
 5502

```
5503 acl missing_slash path_reg ~/article[/]*$
5504 redirect code 301 prefix / drop-query append-slash if missing_slash
```

5505 Example: redirect all HTTP traffic to HTTPS when SSL is handled by haproxy.
 5506

```
5507 redirect scheme https if !{ ssl_fc }
```

5508 Example: append 'www.' prefix in front of all hosts not having it
 5509

```
5510 http-request redirect code 301 location www. %[hdr(host)][req.uri] \
5511 unless { hdr_beg(host) -i www }
```

5512 See section 7 about ACL usage.

5513

5514 redisp (deprecated)

5515 redisp (deprecated)

5516 Enable or disable session redistribution in case of connection failure

5517 May be used in sections: defaults | frontend | listen | backend

5518 Arguments : none

5519 In HTTP mode, if a server designated by a cookie is down, clients may

5520

5521 definitely stick to it because they cannot flush the cookie, so they will not
 5522 be able to access the service anymore.

5523 Specifying "redispatch" will allow the proxy to break their persistence and
 5524 redistribute them to a working server.

5525 It also allows to retry last connection to another server in case of multiple
 5526 connection failures. Of course, it requires having "retries" set to a nonzero
 5527 value.

5528 This form is deprecated, do not use it in any new configuration, use the new
 5529 "option redispatch" instead.

5530 See also : "option redispatch"

```
5531 redadd <string> [{if | unless} <cond>]
```

5532 Add a header at the end of the HTTP request

5533 May be used in sections : defaults | frontend | listen | backend
 5534 no | yes | yes | yes
 5535 Arguments :
 5536 <string> is the complete line to be added. Any space or known delimiter
 5537 must be escaped using a backslash ('\'). Please refer to section
 5538 6 about HTTP header manipulation for more information.

5539 <cond> is an optional matching condition built from ACLs. It makes it
 5540 possible to ignore this rule when other conditions are not met.

5541 A new line consisting in <string> followed by a line feed will be added after
 5542 the last header of an HTTP request.

5543 Header transformations only apply to traffic which passes through HAProxy,
 5544 and not to traffic generated by HAProxy, such as health-checks or error
 5545 responses.

5546 Example : add "X-Proto: SSL" to requests coming via port 81
 5547 acl is-ssl dst_port 81
 5548 reqadd X-Proto:\ SSL if is-ssl

5549 See also: "rspadd", section 6 about HTTP header manipulation, and section 7
 5550 about ACLs.

5551 reqallow <search> [{if | unless} <cond>]
 5552 reqallow <search> [{if | unless} <cond>] (ignore case)
 5553 Definitely allow an HTTP request if a line matches a regular expression
 5554 May be used in sections : defaults | frontend | listen | backend
 5555 no | yes | yes | yes

5556 Arguments :
 5557 <search> is the regular expression applied to HTTP headers and to the
 5558 request line. This is an extended regular expression. Parenthesis
 5559 grouping is supported and no preliminary backslash is required.
 5560 Any space or known delimiter must be escaped using a backslash
 5561 ('\'). The pattern applies to a full line at a time. The
 5562 "reqallow" keyword strictly matches case while "reqallow"
 5563 ignores case.

5564 <cond> is an optional matching condition built from ACLs. It makes it
 5565 possible to ignore this rule when other conditions are not met.

5566 A request containing any line which matches extended regular expression
 5567 <search> will mark the request as allowed, even if any later test would
 5568 result in a deny. The test applies both to the request line and to request
 5569 headers. Keep in mind that URLs in request line are case-sensitive while
 5570 header names are not.

5591 It is easier, faster and more powerful to use ACLs to write access policies.
 5592 Reqdeny, reqallow and reqpass should be avoided in new designs.
 5593

5594 Example :

```
5595 # allow www.* but refuse *.local
5596 reqallow ^Host:\ www\.
5597 reqdeny ^Host:\ *.local
```

5598 See also: "reqdeny", "block", section 6 about HTTP header manipulation, and
 5599 section 7 about ACLs.

```
5600 reqdel <search> [(if | unless) <cond>] (ignore case)
5601 reqdel <search> [(if | unless) <cond>] (ignore case)
5602 Delete all headers matching a regular expression in an HTTP request
5603 May be used in sections : defaults | frontend | listen | backend
5604 no | yes | yes | yes | yes
5605 Arguments :
```

5606 **<search>** is the regular expression applied to HTTP headers and to the
 5607 request line. This is an extended regular expression. Parenthesis
 5608 grouping is supported and no preliminary backslash is required.
 5609 Any space or known delimiter must be escaped using a backslash
 5610 ('\\'). The pattern applies to a full line at a time. The "reqdel"
 5611 keyword strictly matches case while "reqidel" ignores case.

5612 **<cond>** is an optional matching condition built from ACLs. It makes it
 5613 possible to ignore this rule when other conditions are not met.
 5614 Any header line matching extended regular expression **<search>** in the request
 5615 will be completely deleted. Most common use of this is to remove unwanted
 5616 and/or dangerous headers or cookies from a request before passing it to the
 5617 next servers.

5618 Header transformations only apply to traffic which passes through HAProxy,
 5619 and not to traffic generated by HAProxy, such as health-checks or error
 5620 responses. Keep in mind that header names are not case-sensitive.

5621 Example :

```
5622 # remove X-Forwarded-For header and SERVER cookie
5623 reqdel ^X-Forwarded-For.*
5624 reqdel ^Cookie:.*SERVER=
```

5625 See also: "reqadd", "reqrep", "rspdel", section 6 about HTTP header
 5626 manipulation, and section 7 about ACLs.

```
5637 reqdeny <search> [(if | unless) <cond>]
5638 reqdeny <search> [(if | unless) <cond>] (ignore case)
5639 Deny an HTTP request if a line matches a regular expression
5640 May be used in sections : defaults | frontend | listen | backend
5641 no | yes | yes | yes | yes
5642 Arguments :
```

5643 **<search>** is the regular expression applied to HTTP headers and to the
 5644 request line. This is an extended regular expression. Parenthesis
 5645 grouping is supported and no preliminary backslash is required.
 5646 Any space or known delimiter must be escaped using a backslash
 5647 ('\\'). The pattern applies to a full line at a time. The
 5648 "reqdeny" keyword strictly matches case while "reqdeny" ignores
 5649 case.

5650 **<cond>** is an optional matching condition built from ACLs. It makes it
 5651 possible to ignore this rule when other conditions are not met.
 5652 A request containing any line which matches extended regular expression

5656 **<search>** will mark the request as denied, even if any later test would
 5657 result in an allow. The test applies both to the request line and to request
 5658 headers. Keep in mind that URLs in request line are case-sensitive while
 5659 header names are not.

5660 A denied request will generate an "HTTP 403 forbidden" response once the
 5661 complete request has been parsed. This is consistent with what is practiced
 5662 using ACLs.

5663 It is easier, faster and more powerful to use ACLs to write access policies.
 5664 Reqdeny, reqallow and reqpass should be avoided in new designs.

5665 Example :

```
5666 # refuse *.local, then allow www.*
5667 reqdeny ^Host:\ *.local
5668 reqallow ^Host:\ www\.
```

5669 See also: "reqallow", "rspdeny", "block", section 6 about HTTP header
 5670 manipulation, and section 7 about ACLs.

```
5671 reqpass <search> [(if | unless) <cond>]
5672 reqpass <search> [(if | unless) <cond>] (ignore case)
5673 Ignore any HTTP request line matching a regular expression in next rules
5674 May be used in sections : defaults | frontend | listen | backend
5675 no | yes | yes | yes | yes
```

5676 **Arguments :**
 5677 **<search>** is the regular expression applied to HTTP headers and to the
 5678 request line. This is an extended regular expression. Parenthesis
 5679 grouping is supported and no preliminary backslash is required.
 5680 Any space or known delimiter must be escaped using a backslash
 5681 ('\\'). The pattern applies to a full line at a time. The
 5682 "reqpass" keyword strictly matches case while "reqipass" ignores
 5683 case.

5684 **<cond>** is an optional matching condition built from ACLs. It makes it
 5685 possible to ignore this rule when other conditions are not met.

5686 A request containing any line which matches extended regular expression
 5687 **<search>** will skip next rules, without assigning any deny or allow verdict.
 5688 The test applies both to the request line and to request headers. Keep in
 5689 mind that URLs in request line are case-sensitive while header names are not.

5690 It is easier, faster and more powerful to use ACLs to write access policies.
 5691 Reqdeny, reqallow and reqpass should be avoided in new designs.

5692 Example :

```
5693 # refuse *.local, then allow www.*, but ignore "www.private.local"
5694 reqipass ^Host:\ www.private\.local
5695 reqdeny ^Host:\ *.local
5696 reqallow ^Host:\ www\.
```

5697 See also: "reqallow", "reqdeny", "block", section 6 about HTTP header
 5698 manipulation, and section 7 about ACLs.

```
5701 reqrep <search> <string> [(if | unless) <cond>]
5702 reqrep <search> <string> [(if | unless) <cond>] (ignore case)
5703 Replace a regular expression with a string in an HTTP request line
5704 May be used in sections : defaults | frontend | listen | backend
5705 no | yes | yes | yes | yes
```

5706 **Arguments :**
 5707 **<search>** is the regular expression applied to HTTP headers and to the
 5708 request line. This is an extended regular expression. Parenthesis
 5709 grouping is supported and no preliminary backslash is required.

5721 Any space or known delimiter must be escaped using a backslash
 5722 ('\\'). The pattern applies to a full line at a time. The "reqrep"
 5723 keyword strictly matches case while "requirep" ignores case.
 5724

5725 <string> is the complete line to be added. Any space or known delimiter
 5726 must be escaped using a backslash ('\\'). References to matched
 5727 pattern groups are possible using the common \N form, with N
 5728 being a single digit between 0 and 9. Please refer to section
 5729 6 about HTTP header manipulation for more information.

5730 <cond> is an optional matching condition built from ACLs. It makes it
 5731 possible to ignore this rule when other conditions are not met.
 5732
 5733

5734 Any line matching extended regular expression <search> in the request (both
 5735 the request line and header lines) will be completely replaced with <string>.
 5736 Most common use of this is to rewrite URLs or domain names in "Host" headers.
 5737

5738 Header transformations only apply to traffic which passes through HAProxy,
 5739 and not to traffic generated by HAProxy, such as health-checks or error
 5740 responses. Note that for increased readability, it is suggested to add enough
 5741 spaces between the request and the response. Keep in mind that URLs in
 5742 request line are case-sensitive while header names are not.

5743 Example :
 5744 # replace "/static/" with "/" at the beginning of any request path.
 5745 reqrep ^([\^:]*\)/static/(.*) \1/ /
 5746 # replace "www.mydomain.com" with "www" in the host name.
 5747 reqrep ^Host: www.mydomain.com Host: \ www
 5748

5749 See also: "reqadd", "reqdel", "rsprep", "tune.bufsize", section 6 about
 5750 HTTP header manipulation, and section 7 about ACLs.
 5751
 5752

5753 reqtarpit <search> [(if | unless) <cond>]
 5754 reqtarpit <search> [(if | unless) <cond>] (ignore case)
 5755 Tarpit an HTTP request containing a line matching a regular expression
 5756 May be used in sections : defaults | frontend | listen | backend
 5757 no | yes | yes | yes

5758 Arguments :
 5759 <search> is the regular expression applied to HTTP headers and to the
 5760 request line. This is an extended regular expression. Parenthesis
 5761 grouping is supported and no preliminary backslash is required.
 5762 Any space or known delimiter must be escaped using a backslash
 5763 ('\\'). The pattern applies to a full line at a time. The
 5764 "reqtarpit" keyword strictly matches case while "reqtarpit"
 5765 ignores case.
 5766

5767 <cond> is an optional matching condition built from ACLs. It makes it
 5768 possible to ignore this rule when other conditions are not met.
 5769
 5770

5771 A request containing any line which matches extended regular expression
 5772 <search> will be tarptitted, which means that it will connect to nowhere, will
 5773 be kept open for a pre-defined time, then will return an HTTP error 500 so
 5774 that the attacker does not suspect it has been tarptitted. The status 500 will
 5775 be reported in the logs, but the completion flags will indicate "PT". The
 5776 delay is defined by "timeout tarpit", or "timeout connect" if the former is
 5777 not set.
 5778

5779 The goal of the tarpit is to slow down robots attacking servers with
 5780 identifiable requests. Many robots limit their outgoing number of connections
 5781 and stay connected waiting for a reply which can take several minutes to
 5782 come. Depending on the environment and attack, it may be particularly
 5783 efficient at reducing the load on the network and firewalls.
 5784

5785 Examples :

5786 # ignore user-agents reporting any flavour of "Mozilla" or "MSIE", but
 5787 # block all others.
 5788 reqpass ^User-Agent:.*(Mozilla|MSIE)
 5789 reqtarpit ^User-Agent:

5790 # block bad guys
 5791 acl badguys src 10.1.0.3 172.16.13.20/28
 5792 reqtarpit . if badguys
 5793
 5794

5795 See also: "reqallow", "reqdeny", "reqpass", section 6 about HTTP header
 5796 manipulation, and section 7 about ACLs.
 5797
 5798

5799 retries <value>

5800 Set the number of retries to perform on a server after a connection failure
 5801 May be used in sections: defaults | frontend | listen | backend
 5802 yes | no | yes | yes

5803 Arguments :
 5804 <value> is the number of times a connection attempt should be retried on
 5805 a server when a connection either is refused or times out. The
 5806 default value is 3.
 5807

5808 It is important to understand that this value applies to the number of
 5809 connection attempts, not full requests. When a connection has effectively
 5810 been established to a server, there will be no more retry.
 5811

5812 In order to avoid immediate reconnections to a server which is restarting,
 5813 a turn-around timer of 1 second is applied before a retry occurs.
 5814

5815 When "option redispatch" is set, the last retry may be performed on another
 5816 server even if a cookie references a different server.
 5817

5818 See also : "option redispatch"

5819 rsppad <string> [(if | unless) <cond>]

5820 Add a header at the end of the HTTP response

5821 May be used in sections : defaults | frontend | listen | backend
 5822 no | yes | yes | yes

5823 Arguments :
 5824 <string> is the complete line to be added. Any space or known delimiter
 5825 must be escaped using a backslash ('\\'). Please refer to section
 5826 6 about HTTP header manipulation for more information.
 5827

5828 <cond> is an optional matching condition built from ACLs. It makes it
 5829 possible to ignore this rule when other conditions are not met.
 5830

5831 A new line consisting in <string> followed by a line feed will be added after
 5832 the last header of an HTTP response.
 5833

5834 Header transformations only apply to traffic which passes through HAProxy,
 5835 and not to traffic generated by HAProxy, such as health-checks or error
 5836 responses.
 5837

5838 See also: "reqadd", section 6 about HTTP header manipulation, and section 7
 5839 about ACLs.
 5840

5841 rspdel <search> [(if | unless) <cond>]

5842 rspdel <search> [(if | unless) <cond>] (ignore case)

5843 Delete all headers matching a regular expression in an HTTP response

5844 May be used in sections : defaults | frontend | listen | backend
 5845 no | yes | yes | yes

5846 Arguments :
 5847 <search> is the regular expression applied to HTTP headers and to the

5851 response line. This is an extended regular expression, so
5852 parenthesis grouping is supported and no preliminary backslash
5853 is required. Any space or known delimiter must be escaped using
5854 a backslash ('\'). The pattern applies to a full line at a time.
5855 The "rspdel" keyword strictly matches case while "rspdel"
5856 ignores case.

5857
5858 <cond> is an optional matching condition built from ACLs. It makes it
5859 possible to ignore this rule when other conditions are not met.

5860
5861 Any header line matching extended regular expression <search> in the response
5862 will be completely deleted. Most common use of this is to remove unwanted
5863 and/or sensitive headers or cookies from a response before passing it to the
5864 client.

5865 Header transformations only apply to traffic which passes through HAProxy,
5866 and not to traffic generated by HAProxy, such as health-checks or error
5867 responses. Keep in mind that header names are not case-sensitive.

5870 Example :

5871 # remove the Server header from responses
5872 rspdel ^Server:.*

5873
5874 See also: "rspadd", "rsprep", "reqdel", section 6 about HTTP header
5875 manipulation, and section 7 about ACLs.

5876
5877 rspdeny <search> [{if | unless} <cond>] (ignore case)
5878 rspdeny <search> [{if | unless} <cond>] (ignore case)

5879 Block an HTTP response if a line matches a regular expression
5880 May be used in sections : defaults | frontend | listen | backend
5881 no | yes | yes | yes | yes
5882 Arguments :
5883 <search> is the regular expression applied to HTTP headers and to the
5884 response line. This is an extended regular expression, so
5885 parenthesis grouping is supported and no preliminary backslash
5886 is required. Any space or known delimiter must be escaped using
5887 a backslash ('\'). The pattern applies to a full line at a time.
5888 The "rspdeny" keyword strictly matches case while "rspdeny"
5889 ignores case.

5890
5891 <cond> is an optional matching condition built from ACLs. It makes it
5892 possible to ignore this rule when other conditions are not met.

5893
5894 A response containing any line which matches extended regular expression
5895 <search> will mark the request as denied. The test applies both to the
5896 response line and to response headers. Keep in mind that header names are not
5897 case-sensitive.

5898
5899 Main use of this keyword is to prevent sensitive information leak and to
5900 block the response before it reaches the client. If a response is denied, it
5901 will be replaced with an HTTP 502 error so that the client never retrieves
5902 any sensitive data.

5903
5904 It is easier, faster and more powerful to use ACLs to write access policies.
5905 Rspdeny should be avoided in new designs.

5907 Example :

5908 # Ensure that no content type matching ms-word will leak
5909 rspdeny ^Content-type:.*ms-word

5910
5911 See also: "reqdeny", "acl", "block", section 6 about HTTP header manipulation
5912 and section 7 about ACLs.

5913
5914
5915

5916 rsprep <search> <string> [{if | unless} <cond>]
5917 rsprep <search> <string> [{if | unless} <cond>] (ignore case)
5918 Replace a regular expression with a string in an HTTP response line
5919 May be used in sections : defaults | frontend | listen | backend
5920 no | yes | yes | yes | yes

5921 Arguments :
5922 <search> is the regular expression applied to HTTP headers and to the
5923 response line. This is an extended regular expression, so
5924 parenthesis grouping is supported and no preliminary backslash
5925 is required. Any space or known delimiter must be escaped using
5926 a backslash ('\'). The pattern applies to a full line at a time.
5927 The "rsprep" keyword strictly matches case while "rsprep"
5928 ignores case.

5929
5930 <string> is the complete line to be added. Any space or known delimiter
5931 must be escaped using a backslash ('\'). References to matched
5932 pattern groups are possible using the common \N form, with N
5933 being a single digit between 0 and 9. Please refer to section
5934 6 about HTTP header manipulation for more information.

5935
5936 <cond> is an optional matching condition built from ACLs. It makes it
5937 possible to ignore this rule when other conditions are not met.

5938
5939 Any line matching extended regular expression <search> in the response (both
5940 the response line and header lines) will be completely replaced with
5941 <string>. Most common use of this is to rewrite Location headers.

5942
5943 Header transformations only apply to traffic which passes through HAProxy,
5944 and not to traffic generated by HAProxy, such as health-checks or error
5945 responses. Note that for increased readability, it is suggested to add enough
5946 spaces between the request and the response. Keep in mind that header names
5947 are not case-sensitive.

5948 Example :

5949 # replace "Location: 127.0.0.1:8080" with "Location: www.mydomain.com"
5950 rsprep ^Location:\ 127.0.0.1:8080 Location:\ www.mydomain.com

5951
5952 See also: "rspadd", "rspdel", "reqrep", section 6 about HTTP header
5953 manipulation, and section 7 about ACLs.

5954
5955 server <name> <address>[:[port]] [param*]

5956 Declare a server in a backend

5957 May be used in sections : defaults | frontend | listen | yes | yes
5958 no | no | yes | yes

5959 Arguments :
5960 <name> is the internal name assigned to this server. This name will
5961 appear in logs and alerts. If "http-send-name-header" is
5962 set, it will be added to the request header sent to the server.

5963
5964 <address> is the IPv4 or IPv6 address of the server. Alternatively, a
5965 resolvable hostname is supported, but this name will be resolved
5966 during start-up. Address "0.0.0.0" or "*" has a special meaning.
5967 It indicates that the connection will be forwarded to the same IP
5968 address as the one from the client connection. This is useful in
5969 transparent proxy architectures where the client's connection is
5970 intercepted and haproxy must forward to the original destination
5971 address. This is more or less what the "transparent" keyword does
5972 except that with a server it's possible to limit concurrency and
5973 to report statistics. Optionally, an address family prefix may be
5974 used before the address to force the family regardless of the
5975 address format, which can be useful to specify a path to a unix
5976 socket with no slash (/). Currently supported prefixes are :
5977 - 'ipv4q' -> address is always IPv4
5978 - 'ipv6q' -> address is always IPv6
5979
5980

```

5981 - 'unix@' -> address is a path to a local unix socket
5982 - 'absn@' -> address is in abstract namespace (Linux only)
5983 Any part of the address string may reference any number of
5984 environment variables by preceding their name with a dollar
5985 sign ('$') and optionally enclosing them with braces ('{ }'),
5986 similarly to what is done in Bourne shell.
5987
5988
5989

```

<port> is an optional port specification. If set, all connections will be sent to this port. If unset, the same port the client connected to will be used. The port may also be prefixed by a "+" or a "-". In this case, the server's port will be determined by adding this value to the client's port.

<param*> is a list of parameters for this server. The "server" keywords accepts an important number of options and has a complete section dedicated to it. Please refer to section 5 for more details.

Examples :

```

5998 server first 10.1.1.1:1080 cookie first check inter 1000
5999 server second 10.1.1.2:1080 cookie second check inter 1000
6000 server transp ipv4@
6001 server backup ${SRV_BACKUP}:1080 backup
6002 server ww1 dc1 ${LAN_DC1}.101:80
6003 server ww1_dc2 ${LAN_DC2}.101:80
6004
6005

```

Note: regarding Linux's abstract namespace sockets, HAProxy uses the whole sun_path length is used for the address length. Some other programs such as socat use the string length only by default. Pass the option "unix-tightsocketlen=0" to any abstract socket definition in socat to make it compatible with HAProxy's.

See also: "default-server", "http-send-name-header" and section 5 about server options

```

6016 source <addr>[:<port>] [usesrc { <addr2>[:<port2>] | client | clientip } ]
6017 source <addr>[:<port>] [usesrc { <addr2>[:<port2>] | hdr_ip(<hdr>[,<occ>]) } ]
6018 source <addr>[:<port>] [interface <name>]
6019 Set the source address for outgoing connections
6020 May be used in sections : defaults | frontend | listen | backend
6021                          yes | no | yes | yes
6022

```

Arguments :
<addr> is the IPv4 address HAProxy will bind to before connecting to a server. This address is also used as a source for health checks.

The default value of 0.0.0.0 means that the system will select the most appropriate address to reach its destination. Optionally an address family prefix may be used before the address to force the family regardless of the address format, which can be useful to specify a path to a unix socket with no slash ('/'). Currently supported prefixes are :

```

6032 - 'ipv4@' -> address is always IPv4
6033 - 'ipv6@' -> address is always IPv6
6034 - 'unix@' -> address is a path to a local unix socket
6035 - 'absn@' -> address is in abstract namespace (Linux only)
6036 Any part of the address string may reference any number of
6037 environment variables by preceding their name with a dollar
6038 sign ('$') and optionally enclosing them with braces ('{ }'),
6039 similarly to what is done in Bourne shell.
6040

```

<port> is an optional port. It is normally not needed but may be useful in some very specific contexts. The default value of zero means the system will select a free port. Note that port ranges are not supported in the backend. If you want to force port ranges, you have to specify them on each "server" line.

<addr2> is the IP address to present to the server when connections are forwarded in full transparent proxy mode. This is currently only supported on some patched Linux kernels. When this is currently specified, clients connecting to the server will be presented with this address, while health checks will still use the address <addr>.

<port2> is the optional port to present to the server when connections are forwarded in full transparent proxy mode (see <addr2> above). The default value of zero means the system will select a free port.

<hdr> is the name of a HTTP header in which to fetch the IP to bind to. This is the name of a comma-separated header list which can contain multiple IP addresses. By default, the last occurrence is used. This is designed to work with the X-Forwarded-For header and to automatically bind to the client's IP address as seen by previous proxy, typically Stunnel. In order to use another occurrence from the last one, please see the <occ> parameter below. When the header (or occurrence) is not found, no binding is performed so that the proxy's default IP address is used. Also keep in mind that the header name is case insensitive, as for any HTTP header.

<occ> is the occurrence number of a value to be used in a multi-value header. This is to be used in conjunction with "hdr_ip(<hdr>)", in order to specify which occurrence to use for the source IP address. Positive values indicate a position from the first occurrence, 1 being the first one. Negative values indicate positions relative to the last one, -1 being the last one. This is helpful for situations where an X-Forwarded-For header is set at the entry point of an infrastructure and must be used several proxy layers away. When this value is not specified, -1 is assumed. Passing a zero here disables the feature.

<name> is an optional interface name to which to bind to for outgoing traffic. On systems supporting this features (currently, only Linux), this allows one to bind all traffic to the server to this interface even if it is not the one the system would select based on routing tables. This should be used with extreme care. Note that using this option requires root privileges.

The "source" keyword is useful in complex environments where a specific address only is allowed to connect to the servers. It may be needed when a private address must be used through a public gateway for instance, and it is known that the system cannot determine the adequate source address by itself.

An extension which is available on certain patched Linux kernels may be used through the "usesrc" optional keyword. It makes it possible to connect to the servers with an IP address which does not belong to the system itself. This is called "full transparent proxy mode". For this to work, the destination servers have to route their traffic back to this address through the machine running HAProxy, and IP forwarding must generally be enabled on this machine.

In this "full transparent proxy" mode, it is possible to force a specific IP address to be presented to the servers. This is not much used in fact. A more common use is to tell HAProxy to present the client's IP address. For this, there are two methods :

- present the client's IP and port addresses. This is the most transparent mode, but it can cause problems when IP connection tracking is enabled on the machine, because a same connection may be seen twice with different states. However, this solution presents the huge advantage of not limiting the system to the 64k outgoing address+port couples, because all

of the client ranges may be used.

- present only the client's IP address and select a spare port. This solution is still quite elegant but slightly less transparent (downstream firewalls logs will not match upstream's). It also presents the downside of limiting the number of concurrent connections to the usual 64k ports. However, since the upstream and downstream ports are different, local IP connection tracking on the machine will not be upset by the reuse of the same session.

Note that depending on the transparent proxy technology used, it may be required to force the source address. In fact, ctpproxy version 2 requires an IP address in <addr> above, and does not support setting of "0.0.0.0" as the IP address because it creates NAT entries which much match the exact outgoing address. Tproxy version 4 and some other kernel patches which work in pure forwarding mode generally will not have this limitation.

This option sets the default source for all servers in the backend. It may also be specified in a "defaults" section. Finer source address specification is possible at the server level using the "source" server option. Refer to section 5 for more information.

In order to work, "usescr" requires root privileges.

Examples :

```
backend private
# Connect to the servers using our 192.168.1.200 source address
source 192.168.1.200
```

```
backend transparent_ssl1
# Connect to the SSL farm from the client's source address
source 192.168.1.200 usescr clientip
```

```
backend transparent_ssl2
# Connect to the SSL farm from the client's source address and port
# not recommended if IP conntrack is present on the local machine.
source 192.168.1.200 usescr client
```

```
backend transparent_ssl3
# Connect to the SSL farm from the client's source address. It
# is more conntrack-friendly.
source 192.168.1.200 usescr clientip
```

```
backend transparent_smtp
# Connect to the SMTP farm from the client's source address/port
# with Tproxy version 4.
source 0.0.0.0 usescr clientip
```

```
backend transparent_http
# Connect to the servers using the client's IP as seen by previous
# proxy.
source 0.0.0.0 usescr hdr_ip(x-forwarded-for,-1)
```

See also : the "source" server option in section 5, the Tproxy patches for the Linux kernel on www.balabit.com, the "bind" keyword.

srvtimeout <timeout> (deprecated)

Set the maximum inactivity time on the server side.

May be used in sections : defaults | frontend | listen | backend
yes | no | yes | yes

Arguments :

<timeout> is the timeout value specified in milliseconds by default, but can be in any other unit if the number is suffixed by the unit, as explained at the top of this document.

The inactivity timeout applies when the server is expected to acknowledge or send data. In HTTP mode, this timeout is particularly important to consider during the first phase of the server's response, when it has to send the headers, as it directly represents the server's processing time for the request. To find out what value to put there, it's often good to start with what would be considered as unacceptable response times, then check the logs to observe the response time distribution, and adjust the value accordingly.

The value is specified in milliseconds by default, but can be in any other unit if the number is suffixed by the unit, as specified at the top of this document. In TCP mode (and to a lesser extent, in HTTP mode), it is highly recommended that the client timeout remains equal to the server timeout in order to avoid complex situations to debug. Whatever the expected server response times, it is a good practice to cover at least one or several TCP packet losses by specifying timeouts that are slightly above multiples of 3 seconds (eg: 4 or 5 seconds minimum).

This parameter is specific to backends, but can be specified once for all in "defaults" sections. This is in fact one of the easiest solutions not to forget about it. An unspecified timeout results in an infinite timeout, which is not recommended. Such a usage is accepted and works but reports a warning during startup because it may results in accumulation of expired sessions in the system if the system's timeouts are not configured either.

This parameter is provided for compatibility but is currently deprecated. Please use "timeout server" instead.

See also : "timeout server", "timeout tunnel", "timeout client" and "clitimeout".

```
stats admin { if | unless } <cond>
```

Enable statistics admin level if/unless a condition is matched

May be used in sections : defaults | frontend | listen | backend
no | yes | yes | yes

This statement enables the statistics admin level if/unless a condition is matched.

The admin level allows to enable/disable servers from the web interface. By default, statistics page is read-only for security reasons.

Note : Consider not using this feature in multi-process mode (nbproc > 1) unless you know what you do : memory is not shared between the processes, which can result in random behaviours.

Currently, the POST request is limited to the buffer size minus the reserved buffer space, which means that if the list of servers is too long, the request won't be processed. It is recommended to alter few servers at a time.

Example :

```
# statistics admin level only for localhost
backend stats localhost
stats enable
stats admin if LOCALHOST
```

Example :

```
# statistics admin level always enabled because of the authentication
backend stats auth
stats enable
stats auth admin:Admin123
stats admin if TRUE
```


Example :
statistics admin level depends on the authenticated user
userlist stats-auth

```
group admin    users admin
user admin     insecure-password Admin123
group readonly users haproxy
user haproxy   insecure-password haproxy
```

```
backend stats_auth
stats enable
acl AUTH http_auth http_auth(stats-auth)
stats http-request auth unless AUTH
stats admin if AUTH_ADMIN
```

See also : "stats enable", "stats auth", "stats http-request", "noprocs",
"bind-process", section 3.4 about userlists and section 7 about
ACL usage.

stats auth <user>[:<passwd>

Enable statistics with authentication and grant access to an account
May be used in sections : defaults | frontend | listen | yes | yes
yes | yes | yes | yes

Arguments :
<user> is a user name to grant access to

<passwd> is the cleartext password associated to this user

This statement enables statistics with default settings, and restricts access
to declared users only. It may be repeated as many times as necessary to
allow as many users as desired. When a user tries to access the statistics
without a valid account, a "401 Forbidden" response will be returned so that
the browser asks the user to provide a valid user and password. The real
which will be returned to the browser is configurable using "stats realm".

Since the authentication method is HTTP Basic Authentication, the passwords
circulate in cleartext on the network. Thus, it was decided that the
configuration file would also use cleartext passwords to remind the users
that those ones should not be sensitive and not shared with any other account.

It is also possible to reduce the scope of the proxies which appear in the
report using "stats scope".

Though this statement alone is enough to enable statistics reporting, it is
recommended to set all other settings in order to avoid relying on default
unobvious parameters.

Example :
public access (limited to this backend only)
backend public_www
server srv1 192.168.0.1:80
stats enable
stats hide-version

```
stats scope .
stats uri /admin?stats
stats realm Haproxy\ Statistics
stats auth admin1:Admin123
stats auth admin2:Admin321
```

```
# internal monitoring access (unlimited)
backend private_monitoring
stats enable
stats uri /admin?stats
stats refresh 5s
```

See also : "stats enable", "stats realm", "stats scope", "stats uri"

stats enable

Enable statistics reporting with default settings
May be used in sections : defaults | frontend | listen | yes | yes
yes | yes | yes | yes

Arguments : none

This statement enables statistics reporting with default settings defined
at build time. Unless stated otherwise, these settings are used :

```
- stats uri : /haproxy?stats
- stats realm : "HAProxy Statistics"
- stats auth : no authentication
- stats scope : no restriction
```

Though this statement alone is enough to enable statistics reporting, it is
recommended to set all other settings in order to avoid relying on default
unobvious parameters.

Example :
public access (limited to this backend only)
backend public_www

```
server srv1 192.168.0.1:80
stats enable
stats hide-version
stats scope .
stats uri /admin?stats
stats realm Haproxy\ Statistics
stats auth admin1:Admin123
stats auth admin2:Admin321
```

```
# internal monitoring access (unlimited)
backend private_monitoring
stats enable
stats uri /admin?stats
stats refresh 5s
```

See also : "stats auth", "stats realm", "stats uri"

stats hide-version

Enable statistics and hide HAProxy version reporting
May be used in sections : defaults | frontend | listen | yes | yes
yes | yes | yes | yes

Arguments : none

By default, the stats page reports some useful status information along with
the statistics. Among them is HAProxy's version. However, it is generally
considered dangerous to report precise version to anyone, as it can help them
target known weaknesses with specific attacks. The "stats hide-version"
statement removes the version from the statistics report. This is recommended
for public sites or any site with a weak login/password.

Though this statement alone is enough to enable statistics reporting, it is
recommended to set all other settings in order to avoid relying on default
unobvious parameters.

Example :
public access (limited to this backend only)
backend public_www
server srv1 192.168.0.1:80
stats enable
stats hide-version

```
6371 stats scope , /admin?stats
6372 stats uri Haproxy\ Statistics
6373 stats realm admin1:AdMiN123
6374 stats auth admin2:AdMiN321
6375
6376 # internal monitoring access (unlimited)
6377 backend private_monitoring
6378 stats enable
6379 stats uri /admin?stats
6380 stats refresh 5s
6381
6382 See also : "stats auth", "stats enable", "stats realm", "stats uri"
6383
6384
6385 stats http-request { allow | deny | auth [realm <realm>] }
6386 [ { if | unless } <condition> ] }
6387 Access control for statistics
6388
6389 May be used in sections: defaults | frontend | listen | backend
6390 no | no | yes | yes | yes
6391
6392 As "http-request", these set of options allow to fine control access to
6393 statistics. Each option may be followed by if/unless and acl.
6394 First option with matched condition (or option without condition) is final.
6395 For "deny" a 403 error will be returned, for "allow" normal processing is
6396 performed, for "auth" a 401/407 error code is returned so the client
6397 should be asked to enter a username and password.
6398
6399 There is no fixed limit to the number of http-request statements per
6400 instance.
6401
6402 See also : "http-request", section 3.4 about userlists and section 7
6403 about ACL usage.
6404
6405
6406 stats realm <realm>
6407 Enable statistics and set authentication realm
6408 May be used in sections : defaults | frontend | listen | backend
6409 yes | yes | yes | yes
6410
6411 Arguments :
6412 <realm> is the name of the HTTP Basic Authentication realm reported to
6413 the browser. The browser uses it to display it in the pop-up
6414 inviting the user to enter a valid username and password.
6415
6416 The realm is read as a single word, so any spaces in it should be escaped
6417 using a backslash ('\').
6418
6419 This statement is useful only in conjunction with "stats auth" since it is
6420 only related to authentication.
6421
6422 Though this statement alone is enough to enable statistics reporting, it is
6423 recommended to set all other settings in order to avoid relying on default
6424 unobvious parameters.
6425
6426 Example :
6427 # public access (limited to this backend only)
6428 backend public_www
6429 server srv1 192.168.0.1:80
6430 stats enable
6431 stats hide-version
6432 stats scope , /admin?stats
6433 stats realm Haproxy\ Statistics
6434 stats auth admin1:AdMiN123
6435
```

```
6436 stats auth admin2:AdMiN321
6437
6438 # internal monitoring access (unlimited)
6439 backend private_monitoring
6440 stats enable
6441 stats uri /admin?stats
6442 stats refresh 5s
6443
6444 See also : "stats auth", "stats enable", "stats uri"
6445
6446
6447 stats refresh <delay>
6448 Enable statistics with automatic refresh
6449 May be used in sections : defaults | frontend | listen | backend
6450 yes | yes | yes | yes
6451
6452 Arguments :
6453 <delay> is the suggested refresh delay, specified in seconds, which will
6454 be returned to the browser consulting the report page. While the
6455 browser is free to apply any delay, it will generally respect it
6456 and refresh the page this every seconds. The refresh interval may
6457 be specified in any other non-default time unit, by suffixing the
6458 unit after the value, as explained at the top of this document.
6459
6460 This statement is useful on monitoring displays with a permanent page
6461 reporting the load balancer's activity. When set, the HTML report page will
6462 include a link "refresh"/"stop refresh" so that the user can select whether
6463 he wants automatic refresh of the page or not.
6464
6465 Though this statement alone is enough to enable statistics reporting, it is
6466 recommended to set all other settings in order to avoid relying on default
6467 unobvious parameters.
6468
6469 Example :
6470 # public access (limited to this backend only)
6471 backend public_www
6472 server srv1 192.168.0.1:80
6473 stats enable
6474 stats hide-version
6475 stats scope , /admin?stats
6476 stats uri Haproxy\ Statistics
6477 stats realm admin1:AdMiN123
6478 stats auth admin2:AdMiN321
6479
6480 # internal monitoring access (unlimited)
6481 backend private_monitoring
6482 stats enable
6483 stats uri /admin?stats
6484 stats refresh 5s
6485
6486 See also : "stats auth", "stats enable", "stats realm", "stats uri"
6487
6488
6489 stats scope { <name> | " " }
6490 Enable statistics and limit access scope
6491 May be used in sections : defaults | frontend | listen | backend
6492 yes | yes | yes | yes
6493
6494 Arguments :
6495 <name> is the name of a listen, frontend or backend section to be
6496 reported. The special name "." (a single dot) designates the
6497 section in which the statement appears.
6498
6499 When this statement is specified, only the sections enumerated with this
6500 statement will appear in the report. All other sections will be hidden. This
6501 statement may appear as many times as needed if multiple sections need to be
```

reported. Please note that the name checking is performed as simple string comparisons, and that it is never checked that a give section name really exists.

Though this statement alone is enough to enable statistics reporting, it is recommended to set all other settings in order to avoid relying on default unobvious parameters.

```
Example :
# public access (limited to this backend only)
backend public_www
server srv1 192.168.0.1:80
stats enable
stats hide-version
stats scope .
stats uri /admin?stats
stats realm Haproxy\ Statistics
stats auth admin1:Admin123
stats auth admin2:Admin321
```

```
# internal monitoring access (unlimited)
backend private_monitoring
stats enable
stats uri /admin?stats
stats refresh 5s
```

See also : "stats auth", "stats enable", "stats realm", "stats uri"

```
stats show-desc [ <desc> ]
```

Enable reporting of a description on the statistics page.

May be used in sections : defaults | frontend | listen | backend
yes | yes | yes | yes

<desc> is an optional description to be reported. If unspecified, the description from global section is automatically used instead.

This statement is useful for users that offer shared services to their customers, where node or description should be different for each customer.

Though this statement alone is enough to enable statistics reporting, it is recommended to set all other settings in order to avoid relying on default unobvious parameters. By default description is not shown.

```
Example :
# internal monitoring access (unlimited)
backend private_monitoring
stats enable
stats show-desc Master node for Europe, Asia, Africa
stats uri /admin?stats
stats refresh 5s
```

See also: "show-node", "stats enable", "stats uri" and "description" in global section.

```
stats show-legends
```

Enable reporting additional information on the statistics page

May be used in sections : defaults | frontend | listen | backend
yes | yes | yes | yes

Arguments : none

Enable reporting additional information on the statistics page :

- cap: capabilities (proxy)
- mode: one of tcp, http or health (proxy)

- id: SNMP ID (proxy, socket, server)
- IP (socket, server)
- cookie (backend, server)

Though this statement alone is enough to enable statistics reporting, it is recommended to set all other settings in order to avoid relying on default unobvious parameters. Default behaviour is not to show this information.

See also: "stats enable", "stats uri".

```
stats show-node [ <name> ]
```

Enable reporting of a host name on the statistics page.

May be used in sections : defaults | frontend | listen | backend
yes | yes | yes | yes

Arguments: <name> is an optional name to be reported. If unspecified, the node name from global section is automatically used instead.

This statement is useful for users that offer shared services to their customers, where node or description might be different on a stats page provided for each customer. Default behaviour is not to show host name.

Though this statement alone is enough to enable statistics reporting, it is recommended to set all other settings in order to avoid relying on default unobvious parameters.

Example:

```
# internal monitoring access (unlimited)
backend private_monitoring
stats enable
stats show-node Europe-1
stats uri /admin?stats
stats refresh 5s
```

See also: "show-desc", "stats enable", "stats uri", and "node" in global section.

```
stats uri <prefix>
```

Enable statistics and define the URI prefix to access them

May be used in sections : defaults | frontend | listen | backend
yes | yes | yes | yes

Arguments : <prefix> is the prefix of any URI which will be redirected to stats. This prefix may contain a question mark (?) to indicate part of a query string.

The statistics URI is intercepted on the relayed traffic, so it appears as a page within the normal application. It is strongly advised to ensure that the selected URI will never appear in the application, otherwise it will never be possible to reach it in the application.

The default URI compiled in haproxy is "/haproxy?stats", but this may be changed at build time, so it's better to always explicitly specify it here. It is generally a good idea to include a question mark in the URI so that intermediate proxies refrain from caching the results. Also, since any string beginning with the prefix will be accepted as a stats request, the question mark helps ensuring that no valid URI will begin with the same words.

It is sometimes very convenient to use "/" as the URI prefix, and put that statement in a "listen" instance of its own. That makes it easy to dedicate an address or a port to statistics only.

Though this statement alone is enough to enable statistics reporting, it is

recommended to set all other settings in order to avoid relying on default unobvious parameters.

Example :

```
# public access (limited to this backend only)
```

```
backend public_www
```

```
server srv1 192.168.0.1:80
```

```
stats enable
```

```
stats hide-version
```

```
stats scope
```

```
stats uri /admin?stats
```

```
stats realm Haproxy\ Statistics
```

```
stats auth admin1:Admin123
```

```
stats auth admin2:Admin321
```

```
# internal monitoring access (unlimited)
```

```
backend private_monitoring
```

```
stats enable
```

```
stats uri /admin?stats
```

```
stats refresh 5s
```

See also : "stats auth", "stats enable", "stats realm"

```
stick match <pattern> [table <table>] [{if | unless} <cond>]
```

Define a request pattern matching condition to stick a user to a server

May be used in sections : defaults | frontend | listen | backend

	no		no		yes		yes
--	----	--	----	--	-----	--	-----

Arguments :

<pattern> is a sample expression rule as described in section 7.3. It describes what elements of the incoming request or connection will be analysed in the hope to find a matching entry in a stickiness table. This rule is mandatory.

<table> is an optional stickiness table name. If unspecified, the same backend's table is used. A stickiness table is declared using the "stick-table" statement.

<cond> is an optional matching condition. It makes it possible to match on a certain criterion only when other conditions are met (or not met). For instance, it could be used to match on a source IP address except when a request passes through a known proxy, in which case we'd match on a header containing that IP address.

Some protocols or applications require complex stickiness rules and cannot always simply rely on cookies nor hashing. The "stick match" statement describes a rule to extract the stickiness criterion from an incoming request or connection. See section 7 for a complete list of possible patterns and transformation rules.

The table has to be declared using the "stick-table" statement. It must be of a type compatible with the pattern. By default it is the one which is present in the same backend. It is possible to share a table with other backends by referencing it using the "table" keyword. If another table is referenced, the server's ID inside the backends are used. By default, all server IDs start at 1 in each backend, so the server ordering is enough. But in case of doubt, it is highly recommended to force server IDs using their "id" setting.

It is possible to restrict the conditions where a "stick match" statement will apply, using "if" or "unless" followed by a condition. See section 7 for ACL based conditions.

There is no limit on the number of "stick match" statements. The first that applies and matches will cause the request to be directed to the same server

as was used for the request which created the entry. That way, multiple matches can be used as fallbacks.

The stick rules are checked after the persistence cookies, so they will not affect stickiness if a cookie has already been used to select a server. That way, it becomes very easy to insert cookies and match on IP addresses in order to maintain stickiness between HTTP and HTTPS.

Note : Consider not using this feature in multi-process mode (nbproc > 1) unless you know what you do : memory is not shared between the processes, which can result in random behaviours.

Example :

```
# forward SMTP users to the same server they just used for POP in the
```

```
# last 30 minutes
```

```
backend pop
```

```
mode tcp
```

```
balance roundrobin
```

```
stick store-request src
```

```
stick-table type ip size 200k expire 30m
```

```
server s1 192.168.1.1:110
```

```
server s2 192.168.1.1:110
```

```
backend smtp
```

```
mode tcp
```

```
balance roundrobin
```

```
stick match src table pop
```

```
server s1 192.168.1.1:25
```

```
server s2 192.168.1.1:25
```

See also : "stick-table", "stick on", "nbproc", "bind-process" and section 7 about ACLs and samples fetching.

```
stick on <pattern> [table <table>] [{if | unless} <condition>]
```

Define a request pattern to associate a user to a server

May be used in sections : defaults | frontend | listen | backend

	no		no		yes		yes
--	----	--	----	--	-----	--	-----

Note : This form is exactly equivalent to "stick match" followed by "stick store-request", all with the same arguments. Please refer to both keywords for details. It is only provided as a convenience for writing more maintainable configurations.

Note : Consider not using this feature in multi-process mode (nbproc > 1) unless you know what you do : memory is not shared between the processes, which can result in random behaviours.

Examples :

```
# The following form ...
```

```
stick on src table pop if !localhost
```

```
# ...is strictly equivalent to this one :
```

```
stick match src table pop if !localhost
```

```
stick store-request src table pop if !localhost
```

Use cookie persistence for HTTP, and stick on source address for HTTPS as # well as HTTP without cookie. Share the same table between both accesses.

```
backend http
```

```
mode http
```

```
balance roundrobin
```

```
stick on src table https
```

```
cookie SRV insert indirect nocache
```

```
server s1 192.168.1.1:80 cookie s1
```

6696

```
6761 server s2 192.168.1.1:80 cookie s2
6762
6763 backend https
6764     mode tcp
6765     balance roundrobin
6766     stick-table type ip size 200k expire 30m
6767     stick on src
6768     server s1 192.168.1.1:443
6769     server s2 192.168.1.1:443
6770
6771 See also : "stick match", "stick store-request", "nbproc" and "bind-process".
6772
6773
6774 stick store-request <pattern> [table <table>] [{if | unless} <condition>]
```

6775 Define a request pattern used to create an entry in a stickiness table

6776 May be used in sections : defaults | frontend | listen | backend

```
6777 no | no | yes | yes
```

6778 Arguments :

6779 <pattern> is a sample expression rule as described in section 7.3. It describes what elements of the incoming request or connection will be analysed, extracted and stored in the table once a server is selected.

6784 <table> is an optional stickiness table name. If unspecified, the same backend's table is used. A stickiness table is declared using the "stick-table" statement.

6789 <cond> is an optional storage condition. It makes it possible to store certain criteria only when some conditions are met (or not met). For instance, it could be used to store the source IP address except when the request passes through a known proxy, in which case we'd store a converted form of a header containing that IP address.

6795 Some protocols or applications require complex stickiness rules and cannot always simply rely on cookies nor hashing. The "stick store-request" statement describes a rule to decide what to extract from the request and when to do it, in order to store it into a stickiness table for further requests to match it using the "stick match" statement. Obviously the extracted part must make sense and have a chance to be matched in a further request. Storing a client's IP address for instance often makes sense. Storing an ID found in a URL parameter also makes sense. Storing a source port will almost never make any sense because it will be randomly matched. See section 7 for a complete list of possible patterns and transformation rules.

6806 The table has to be declared using the "stick-table" statement. It must be of a type compatible with the pattern. By default it is the one which is present in the same backend. It is possible to share a table with other backends by referencing it using the "table" keyword. If another table is referenced, the server's ID inside the backends are used. By default, all server IDs start at 1 in each backend, so the server ordering is enough. But in case of doubt, it is highly recommended to force server IDs using their "id" setting.

6814 It is possible to restrict the conditions where a "stick store-request" statement will apply, using "if" or "unless" followed by a condition. This condition will be evaluated while parsing the request, so any criteria can be used. See section 7 for ACL based conditions.

6819 There is no limit on the number of "stick store-request" statements, but there is a limit of 8 simultaneous stores per request or response. This makes it possible to store up to 8 criteria, all extracted from either the request or the response, regardless of the number of rules. Only the 8 first ones which match will be kept. Using this, it is possible to feed multiple tables at once in the hope to increase the chance to recognize a user on

6826 another protocol or access method. Using multiple store-request rules with the same table is possible and may be used to find the best criterion to rely on, by arranging the rules by decreasing preference order. Only the first extracted criterion for a given table will be stored. All subsequent store-request rules referencing the same table will be skipped and their ACLs will not be evaluated.

6832 The "store-request" rules are evaluated once the server connection has been established, so that the table will contain the real server that processed the request.

6837 Note : Consider not using this feature in multi-process mode (nbproc > 1) unless you know what you do : memory is not shared between the processes, which can result in random behaviours.

6841 Example :

```
6842 # forward SMTP users to the same server they just used for POP in the
6843 # last 30 minutes
6844 backend pop
```

```
6845     mode tcp
```

```
6846     balance roundrobin
```

```
6847     stick store-request src
```

```
6848     stick-table type ip size 200k expire 30m
```

```
6849     server s1 192.168.1.1:110
```

```
6850     server s2 192.168.1.1:110
```

```
6851 backend smtp
```

```
6852     mode tcp
```

```
6853     balance roundrobin
```

```
6854     stick match src table pop
```

```
6855     server s1 192.168.1.1:25
```

```
6856     server s2 192.168.1.1:25
```

6858 See also : "stick-table", "stick on", "nbproc", "bind-process" and section 7 about ACLs and sample fetching.

```
6862 stick-table type {ip | integer | string [len <length>] | binary [len <length>]}
6863     size <size> [expire <expire>] [nopurge] [peers <peersect>]
```

```
6864     [store <data type>]*
```

6865 Configure the stickiness table for the current section

6866 May be used in sections : defaults | frontend | listen | backend

```
6867 no | yes | yes | yes
6868
```

6869 Arguments :

6870 ip a table declared with "type ip" will only store IPv4 addresses. This form is very compact (about 50 bytes per entry) and allows very fast entry lookup and stores with almost no overhead. This is mainly used to store client source IP addresses.

6875 ipv6 a table declared with "type ipv6" will only store IPv6 addresses. This form is very compact (about 60 bytes per entry) and allows very fast entry lookup and stores with almost no overhead. This is mainly used to store client source IP addresses.

6877 integer a table declared with "type integer" will store 32bit integers which can represent a client identifier found in a request for instance.

6884 string a table declared with "type string" will store substrings of up to <len> characters. If the string provided by the pattern extractor is larger than <len>, it will be truncated before being stored. During matching, at most <len> characters will be compared between the string in the table and the extracted pattern. When not specified, the string is automatically limited

6891 to 32 characters.
6892
6893
6894
6895
6896
6897
6898
6899
6900
6901
6902
6903
6904
6905
6906
6907
6908
6909
6910
6911
6912
6913
6914
6915
6916
6917
6918
6919
6920
6921
6922
6923
6924
6925
6926
6927
6928
6929
6930
6931
6932
6933
6934
6935
6936
6937
6938
6939
6940
6941
6942
6943
6944
6945
6946
6947
6948
6949
6950
6951
6952
6953
6954
6955

binary

a table declared with "type binary" will store binary blocks of <len> bytes. If the block provided by the pattern extractor is larger than <len>, it will be truncated before being stored. If the block provided by the sample expression is shorter than <len>, it will be padded by 0. When not specified, the block is automatically limited to 32 bytes.

<length>

is the maximum number of characters that will be stored in a "string" type table (See type "string" above). Or the number of bytes of the block in "binary" type table. Be careful when changing this parameter as memory usage will proportionally increase.

<size>

is the maximum number of entries that can fit in the table. This value directly impacts memory usage. Count approximately 50 bytes per entry, plus the size of a string if any. The size supports suffixes "k", "m", "g" for 2^10, 2^20 and 2^30 factors.

[nopurge]

indicates that we refuse to purge older entries when the table is full. When not specified and the table is full when haproxy wants to store an entry in it, it will flush a few of the oldest entries in order to release some space for the new ones. This is most often the desired behaviour. In some specific cases, it be desirable to refuse new entries instead of purging the older ones. That may be the case when the amount of data to store is far above the hardware limits and we prefer not to offer access to new clients than to reject the ones already connected. When using this parameter, be sure to properly set the "expire" parameter (see below).

<peersect>

is the name of the peers section to use for replication. Entries which associate keys to server IDs are kept synchronized with the remote peers declared in this section. All entries are also automatically learned from the local peer (old process) during a soft restart.

NOTE : each peers section may be referenced only by tables belonging to the same unique process.

<expire>

defines the maximum duration of an entry in the table since it was last created, refreshed or matched. The expiration delay is defined using the standard time format, similarly as the various timeouts. The maximum duration is slightly above 24 days. See section 2.2 for more information. If this delay is not specified, the session won't automatically expire, but older entries will be removed once full. Be sure not to use the "nopurge" parameter if not expiration delay is specified.

<data_type>

is used to store additional information in the stick-table. This may be used by ACLs in order to control various criteria related to the activity of the client matching the stick-table. For each item specified here, the size of each entry will be inflated so that the additional data can fit. Several data types may be stored with an entry. Multiple data types may be specified after the "store" keyword, as a comma-separated list. Alternatively, it is possible to repeat the "store" keyword followed by one or several data types. Except for the "server_id" type which is automatically detected and enabled, all data types must be explicitly declared to be stored. If an ACL references a data type which is not stored, the ACL will simply not match. Some data types require an argument which must be passed just after the type between parenthesis. See below for the supported data types and their arguments.

6956
6957
6958
6959
6960
6961
6962
6963
6964
6965
6966
6967
6968
6969
6970
6971
6972
6973
6974
6975
6976
6977
6978
6979
6980
6981
6982
6983
6984
6985
6986
6987
6988
6989
6990
6991
6992
6993
6994
6995
6996
6997
6998
6999
7000
7001
7002
7003
7004
7005
7006
7007
7008
7009
7010
7011
7012
7013
7014
7015
7016
7017
7018
7019
7020

The data types that can be stored with an entry are the following :

- server_id : this is an integer which holds the numeric ID of the server a request was assigned to. It is used by the "stick match", "stick store", and "stick on" rules. It is automatically enabled when referenced.
- gpc0 : first General Purpose Counter. It is a positive 32-bit integer integer which may be used for anything. Most of the time it will be used to put a special tag on some entries, for instance to note that a specific behaviour was detected and must be known for future matches.
- gpc0_rate(<period>) : increment rate of the first General Purpose Counter over a period. It is a positive 32-bit integer integer which may be used for anything. Just like <gpc0>, it counts events, but instead of keeping a cumulative count, it maintains the rate at which the counter is incremented. Most of the time it will be used to measure the frequency of occurrence of certain events (eg: requests to a specific URL).
- conn_cnt : Connection Count. It is a positive 32-bit integer which counts the absolute number of connections received from clients which matched this entry. It does not mean the connections were accepted, just that they were received.
- conn_cur : Current Connections. It is a positive 32-bit integer which stores the concurrent connection counts for the entry. It is incremented once an incoming connection matches the entry, and decremented once the connection leaves. That way it is possible to know at any time the exact number of concurrent connections for an entry.
- conn_rate(<period>) : frequency counter (takes 12 bytes). It takes an integer parameter <period> which indicates in milliseconds the length of the period over which the average is measured. It reports the average incoming connection rate over that period, in connections per period. The result is an integer which can be matched using ACLs.
- sess_cnt : Session Count. It is a positive 32-bit integer which counts the absolute number of sessions received from clients which matched this entry. A session is a connection that was accepted by the layer 4 rules.
- sess_rate(<period>) : frequency counter (takes 12 bytes). It takes an integer parameter <period> which indicates in milliseconds the length of the period over which the average is measured. It reports the average incoming session rate over that period, in sessions per period. The result is an integer which can be matched using ACLs.
- http_req_cnt : HTTP request Count. It is a positive 32-bit integer which counts the absolute number of HTTP requests received from clients which matched this entry. It does not matter whether they are valid requests or not. Note that this is different from sessions when keep-alive is used on the client side.
- http_req_rate(<period>) : frequency counter (takes 12 bytes). It takes an integer parameter <period> which indicates in milliseconds the length of the period over which the average is measured. It reports the average HTTP request rate over that period in requests per period. The result is an integer which can be matched using ACLs. It does not matter whether they are valid requests or not. Note that this is different from sessions when keep-alive is used on the client side.
- http_err_cnt : HTTP Error Count. It is a positive 32-bit integer which counts the absolute number of HTTP requests errors induced by clients which matched this entry. Errors are counted on invalid and truncated requests, as well as on denied or tarpitted requests, and on failed authentications. If the server responds with 4xx, then the request is also counted as an error since it's an error triggered by the client

(eg: vulnerability scan).

- http_err_rate(<period>) : frequency counter (takes 12 bytes). It takes an integer parameter <period> which indicates in milliseconds the length of the period over which the average is measured. It reports the average HTTP request error rate over that period, in requests per period (see http_err_cnt above for what is accounted as an error). The result is an integer which can be matched using ACLs.
- bytes_in_cnt : client to server byte count. It is a positive 64-bit integer which counts the cumulated amount of bytes received from clients which matched this entry. Headers are included in the count. This may be used to limit abuse of upload features on photo or video servers.
- bytes_in_rate(<period>) : frequency counter (takes 12 bytes). It takes an integer parameter <period> which indicates in milliseconds the length of the period over which the average is measured. It reports the average incoming bytes rate over that period, in bytes per period. It may be used to detect users which upload too much and too fast. Warning: with large uploads, it is possible that the amount of uploaded data will be counted once upon termination, thus causing spikes in the average transfer speed instead of having a smooth one. This may partially be smoothed with "option contstats" though this is not perfect yet. Use of byte_in_cnt is recommended for better fairness.
- bytes_out_cnt : server to client byte count. It is a positive 64-bit integer which counts the cumulated amount of bytes sent to clients which matched this entry. Headers are included in the count. This may be used to limit abuse of bots sucking the whole site.
- bytes_out_rate(<period>) : frequency counter (takes 12 bytes). It takes an integer parameter <period> which indicates in milliseconds the length of the period over which the average is measured. It reports the average outgoing bytes rate over that period, in bytes per period. It may be used to detect users which download too much and too fast. Warning: with large transfers, it is possible that the amount of transferred data will be counted once upon termination, thus causing spikes in the average transfer speed instead of having a smooth one. This may partially be smoothed with "option contstats" though this is not perfect yet. Use of byte_out_cnt is recommended for better fairness.

There is only one stick-table per proxy. At the moment of writing this doc, it does not seem useful to have multiple tables per proxy. If this happens to be required, simply create a dummy backend with a stick-table in it and reference it.

It is important to understand that stickiness based on learning information has some limitations, including the fact that all learned associations are lost upon restart. In general it can be good as a complement but not always as an exclusive stickiness.

Last, memory requirements may be important when storing many data types. Indeed, storing all indicators above at once in each entry requires 116 bytes per entry, or 116 MB for a 1-million entries table. This is definitely not something that can be ignored.

Example:

- # Keep track of counters of up to 1 million IP addresses over 5 minutes
- # and store a general purpose counter and the average connection rate
- # computed over a sliding window of 30 seconds.
- stick-table type ip size 1m expire 5m store gpc0,conn_rate(30s)

See also : "stick match", "stick on", "stick store-request", section 2.2
about time format and section 7 about ACLs.

stick store-response <pattern> [table <table>] [if | unless] <condition>]

Define a request pattern used to create an entry in a stickiness table. May be used in sections : defaults | frontend | listen | backend

no | no | yes | yes

Arguments :

<pattern> is a sample expression rule as described in section 7.3. It describes what elements of the response or connection will be analysed, extracted and stored in the table once a server is selected.

<table> is an optional stickiness table name. If unspecified, the same backend's table is used. A stickiness table is declared using the "stick-table" statement.

<cond> is an optional storage condition. It makes it possible to store certain criteria only when some conditions are met (or not met). For instance, it could be used to store the SSL session ID only when the response is a SSL server hello.

Some protocols or applications require complex stickiness rules and cannot always simply rely on cookies nor hashing. The "stick store-response" statement describes a rule to decide what to extract from the response and when to do it, in order to store it into a stickiness table for further requests to match it using the "stick match" statement. Obviously the extracted part must make sense and have a chance to be matched in a further request. Storing an ID found in a header of a response makes sense. See section 7 for a complete list of possible patterns and transformation rules.

The table has to be declared using the "stick-table" statement. It must be of a type compatible with the pattern. By default it is the one which is present in the same backend. It is possible to share a table with other backends by referencing it using the "table" keyword. If another table is referenced, the server's ID inside the backends are used. By default, all server IDs start at 1 in each backend, so the server ordering is enough. But in case of doubt, it is highly recommended to force server IDs using their "id" setting.

It is possible to restrict the conditions where a "stick store-response" statement will apply, using "if" or "unless" followed by a condition. This condition will be evaluated while parsing the response, so any criteria can be used. See section 7 for ACL based conditions.

There is no limit on the number of "stick store-response" statements, but there is a limit of 8 simultaneous stores per request or response. This makes it possible to store up to 8 criteria, all extracted from either the request or the response, regardless of the number of rules. Only the 8 first ones which match will be kept. Using this, it is possible to feed multiple tables at once in the hope to increase the chance to recognize a user on another protocol or access method. Using multiple store-response rules with the same table is possible and may be used to find the best criterion to rely on, by arranging the rules by decreasing preference order. Only the first extracted criterion for a given table will be stored. All subsequent store-response rules referencing the same table will be skipped and their ACLs will not be evaluated. However, even if a store-request rule references a table, a store-response rule may also use the same table. This means that each table may learn exactly one element from the request and one element from the response at once.

The table will contain the real server that processed the request.

Example :

- # Learn SSL session ID from both request and response and create affinity.
- backend https

```
7151 mode tcp
7152 balance roundrobin
7153 # maximum SSL session ID length is 32 bytes.
7154 stick-table type binary len 32 size 30k expire 30m
7155
7156 acl clienthello req_ssl_hello_type 1
7157 acl serverhello rep_ssl_hello_type 2
7158
7159 # use tcp content accepts to detects ssl client and server hello.
7160 tcp-request inspect-delay 5s
7161 tcp-request content accept if clienthello
7162
7163 # no timeout on response inspect delay by default.
7164 tcp-response content accept if serverhello
7165
7166 # SSL session ID (SSLID) may be present on a client or server hello.
7167 # Its length is coded on 1 byte at offset 43 and its value starts
7168 # at offset 44.
7169
7170 # Match and learn on request if client hello.
7171 stick on payload_lv(43,1) if clienthello
7172
7173 # Learn on response if server hello.
7174 stick store-response payload_lv(43,1) if serverhello
7175
7176 server s1 192.168.1.1:443
7177 server s2 192.168.1.1:443
7178
```

See also : "stick-table", "stick on", and section 7 about ACLs and pattern extraction.

tcp-check connect [params*]

Opens a new connection

May be used in sections:

defaults	no	no	yes	yes	yes
frontend					
listen					
backend					

When an application lies on more than a single TCP port or when HAPROXY load-balance many services in a single backend, it makes sense to probe all the services individually before considering a server as operational.

When there are no TCP port configured on the server line neither server port directive, then the 'tcp-check connect port <port>' must be the first step of the sequence.

In a tcp-check ruleset a 'connect' is required, it is also mandatory to start the ruleset with a 'connect' rule. Purpose is to ensure admin know what they do.

Parameters :

They are optional and can be used to describe how HAProxy should open and use the TCP connection.

port if not set, check port or server port is used.

It tells HAProxy where to open the connection to.

<port> must be a valid TCP port source integer, from 1 to 65535.

send-proxy send a PROXY protocol string

ssl opens a ciphered connection

Examples:

check HTTP and HTTPS services on a server.

first open port 80 thanks to server line port directive, then

tcp-check opens port 443, ciphered and run a request on it:

```
7216 option tcp-check
7217 tcp-check connect
7218 tcp-check send GET \ HTTP/1.0\r\n
7219 tcp-check send Host:\ haproxy.lwt.eu\r\n
7220 tcp-check send \r\n
7221 tcp-check expect rstring (2..|3..)
7222 tcp-check connect port 443 ssl
7223 tcp-check send GET \ HTTP/1.0\r\n
7224 tcp-check send Host:\ haproxy.lwt.eu\r\n
7225 tcp-check send \r\n
7226 tcp-check expect rstring (2..|3..)
7227 server www 10.0.0.1 check port 80
7228
7229 # check both POP and IMAP from a single server:
7230 option tcp-check
7231 tcp-check connect port 110
7232 tcp-check expect string +OK POP3 ready
7233 tcp-check connect port 143
7234 tcp-check expect string * OK IMAP4 ready
7235 server mail 10.0.0.1 check
7236
7237
7238
7239
```

See also : "option tcp-check", "tcp-check send", "tcp-check expect"

tcp-check expect [!] <match> <pattern>

Specify data to be collected and analysed during a generic health check

May be used in sections:

defaults	no	no	yes	yes
frontend				
listen				
backend				

Arguments :

<match> is a keyword indicating how to look for a specific pattern in the response. The keyword may be one of "string", "rstring" or binary.

The keyword may be preceded by an exclamation mark ("!") to negate the match. Spaces are allowed between the exclamation mark and the keyword. See below for more details on the supported keywords.

<pattern> is the pattern to look for. It may be a string or a regular expression. If the pattern contains spaces, they must be escaped with the usual backslash ('\').

If the match is set to binary, then the pattern must be passed as a serie of hexadecimal digits in an even number. Each sequence of two digits will represent a byte. The hexadecimal digits may be used upper or lower case.

The available matches are intentionally similar to their http-check cousins :

string <string> : test the exact string matches in the response buffer.

A health check response will be considered valid if the response's buffer contains this exact string. If the "string" keyword is prefixed with "!", then the response will be considered invalid if the body contains this string. This can be used to look for a mandatory pattern in a protocol response, or to detect a failure when a specific error appears in a protocol banner.

rstring <regex> : test a regular expression on the response buffer.

A health check response will be considered valid if the response's buffer matches this expression. If the "rstring" keyword is prefixed with "!", then the response will be considered invalid if the body matches the expression.

binary <hexstring> : test the exact string in its hexadecimal form matches

7281 in the response buffer. A health check response will
7282 be considered valid if the response's buffer contains
7283 this exact hexadecimal string.
7284 Purpose is to match data on binary protocols.
7285

7286 It is important to note that the responses will be limited to a certain size
7287 defined by the global "tune.chksize" option, which defaults to 16384 bytes.
7288 Thus, too large responses may not contain the mandatory pattern when using
7289 "string", "rstring" or binary. If a large response is absolutely required, it
7290 is possible to change the default max size by setting the global variable.
7291 However, it is worth keeping in mind that parsing very large responses can
7292 waste some CPU cycles, especially when regular expressions are used, and that
7293 it is always better to focus the checks on smaller resources. Also, in its
7294 current state, the check will not find any string nor regex past a null
7295 character in the response. Similarly it is not possible to request matching
7296 the null character.

7297 Examples :
7298 # perform a POP check
7299 option tcp-check
7300 tcp-check expect string +OK\ POP3\ ready
7301
7302 # perform an IMAP check
7303 option tcp-check
7304 tcp-check expect string *\ OK\ IMAP4\ ready
7305
7306 # look for the redis master server
7307 option tcp-check
7308 tcp-check send PING\r\n
7309 tcp-check expect +PONG
7310 tcp-check send info\ replication\r\n
7311 tcp-check expect string role:master
7312 tcp-check send QUIT\r\n
7313 tcp-check expect string +OK
7314
7315

7316 See also : "option tcp-check", "tcp-check connect", "tcp-check send",
7317 "tcp-check send-binary", "http-check expect", tune.chksize
7318
7319

7320 tcp-check send <data>

7321 Specify a string to be sent as a question during a generic health check
7322 May be used in sections: defaults | frontend | listen | backend
7323 no | no | yes | yes
7324

7325 <data> : the data to be sent as a question during a generic health check
7326 session. For now, <data> must be a string.
7327
7328

7329 Examples :

7330 # look for the redis master server
7331 option tcp-check
7332 tcp-check send info\ replication\r\n
7333 tcp-check expect string role:master
7334

7335 See also : "option tcp-check", "tcp-check connect", "tcp-check expect",
7336 "tcp-check send-binary", tune.chksize
7337
7338

7339 tcp-check send-binary <hexastring>

7340 Specify an hexa digits string to be sent as a binary question during a raw
7341 tcp health check
7342 May be used in sections: defaults | frontend | listen | backend
7343 no | no | yes | yes
7344

7345 <data> : the data to be sent as a question during a generic health check

7346 session. For now, <data> must be a string.
7347 <hexastring> : test the exact string in its hexadecimal form matches in the
7348 response buffer. A health check response will be considered
7349 valid if the response's buffer contains this exact
7350 hexadecimal string.
7351 Purpose is to send binary data to ask on binary protocols.
7352

7353 Examples :

7354 # redis check in binary
7355 option tcp-check
7356 tcp-check send-binary 50494e470d0a # PING\r\n
7357 tcp-check expect binary 2b504f4e47 # +PONG
7358
7359

7360 See also : "option tcp-check", "tcp-check connect", "tcp-check expect",
7361 "tcp-check send", tune.chksize
7362
7363

7364 tcp-request connection <action> [(if | unless) <condition>]

7365 Perform an action on an incoming connection depending on a layer 4 condition
7366 May be used in sections: defaults | frontend | listen | backend
7367 no | yes | yes | no
7368

7369 Arguments : defines the action to perform if the condition applies. Valid
7370 <action> actions include : "accept", "reject", "track-scl", "track-scl",
7371 "track-sc2", and "expect-proxy". See below for more details.
7372
7373

7374 <condition> is a standard layer4-only ACL-based condition (see section 7).

7375 Immediately after acceptance of a new incoming connection, it is possible to
7376 evaluate some conditions to decide whether this connection must be accepted
7377 or dropped or have its counters tracked. Those conditions cannot make use of
7378 any data contents because the connection has not been read from yet, and the
7379 buffers are not yet allocated. This is used to selectively and very quickly
7380 accept or drop connections from various sources with a very low overhead. If
7381 some contents need to be inspected in order to take the decision, the
7382 "tcp-request content" statements must be used instead.
7383

7384 The "tcp-request connection" rules are evaluated in their exact declaration
7385 order. If no rule matches or if there is no rule, the default action is to
7386 accept the incoming connection. There is no specific limit to the number of
7387 rules which may be inserted.
7388

7389 Five types of actions are supported :

7390 - accept :
7391 accepts the connection if the condition is true (when used with "if")
7392 or false (when used with "unless"). The first such rule executed ends
7393 the rules evaluation.
7394

7395 - reject :
7396 rejects the connection if the condition is true (when used with "if")
7397 or false (when used with "unless"). The first such rule executed ends
7398 the rules evaluation. Rejected connections do not even become a
7399 session, which is why they are accounted separately for in the stats,
7400 as "denied connections". They are not considered for the session
7401 rate-limit and are not logged either. The reason is that these rules
7402 should only be used to filter extremely high connection rates such as
7403 the ones encountered during a massive DDoS attack. Under these extreme
7404 conditions, the simple action of logging each event would make the
7405 system collapse and would considerably lower the filtering capacity. If
7406 logging is absolutely desired, then "tcp-request content" rules should
7407 be used instead.
7408

7409 - expect-proxy layer4 :
7410 configures the client-facing connection to receive a PROXY protocol

header before any byte is read from the socket. This is equivalent to having the "accept-proxy" keyword on the "bind" line, except that using the TCP rule allows the PROXY protocol to be accepted only for certain IP address ranges using an ACL. This is convenient when multiple layers of load balancers are passed through by traffic coming from public hosts.

- capture <sample> len <length> :

This only applies to "tcp-request content" rules. It captures sample expression <sample> from the request buffer, and converts it to a string of at most <len> characters. The resulting string is stored into the next request "capture" slot, so it will possibly appear next to some captured HTTP headers. It will then automatically appear in the logs, and it will be possible to extract it using sample fetch rules to feed it into headers or anything. The length should be limited given that this size will be allocated for each capture during the whole session life. Since it applies to Please check section 7.3 (Fetching samples) and "capture request header" for more information.

- { track-sc0 | track-scl | track-sc2 } <key> [table <table>] : enables tracking of sticky counters from current connection. These rules do not stop evaluation and do not change default action. Two sets of counters may be simultaneously tracked by the same connection. The first "track-sc0" rule executed enables tracking of the counters of the specified table as the first set. The first "track-scl" rule executed enables tracking of the counters of the specified table as the second set. The first "track-sc2" rule executed enables tracking of the counters of the specified table as the third set. It is a recommended practice to use the first set of counters for the per-frontend counters and the second set for the per-backend ones. But this is just a guideline, all may be used everywhere.

These actions take one or two arguments :

<key> is mandatory, and is a sample expression rule as described in section 7.3. It describes what elements of the incoming request or connection will be analysed, extracted, combined, and used to select which table entry to update the counters. Note that "tcp-request connection" cannot use content-based fetches.

<table> is an optional table to be used instead of the default one, which is the stick-table declared in the current proxy. All the counters for the matches and updates for the key will then be performed in that table until the session ends.

Once a "track-sc*" rule is executed, the key is looked up in the table and if it is not found, an entry is allocated for it. Then a pointer to that entry is kept during all the session's life, and this entry's counters are updated as often as possible, every time the session's counters are updated, and also systematically when the session ends. Counters are only updated for events that happen after the tracking has been started. For example, connection counters will not be updated when tracking layer 7 information, since the connection event happens before layer7 information is extracted.

If the entry tracks concurrent connection counters, one connection is counted for as long as the entry is tracked, and the entry will not expire during that time. Tracking counters also provides a performance advantage over just checking the keys, because only one table lookup is performed for all ACL checks that make use of it.

Note that the "if/unless" condition is optional. If no condition is set on the action, it is simply performed unconditionally. That can be useful for "track-sc*" actions as well as for changing the default action to a reject.

Example: accept all connections from white-listed hosts, reject too fast connection without counting them, and track accepted connections. This results in connection rate being capped from abusive sources.

```
tcp-request connection accept if { src -f /etc/haproxy/whitelist.lst }
tcp-request connection reject if { src_conn_rate gt 10 }
tcp-request connection track-sc0 src
```

Example: accept all connections from white-listed hosts, count all other connections and reject too fast ones. This results in abusive ones being blocked as long as they don't slow down.

```
tcp-request connection accept if { src -f /etc/haproxy/whitelist.lst }
tcp-request connection track-sc0 src
tcp-request connection reject if { sc0_conn_rate gt 10 }
```

Example: enable the PROXY protocol for traffic coming from all known proxies.

```
tcp-request connection expect-proxy layer4 if { src -f proxies.lst }
```

See section 7 about ACL usage.

See also : "tcp-request content", "stick-table"

tcp-request content <action> [(if | unless) <condition>]

Perform an action on a new session depending on a layer 4-7 condition

May be used in sections : defaults | frontend | listen | backend

no | yes | yes | yes

Arguments : defines the action to perform if the condition applies. Valid actions include : "accept", "reject", "track-sc0", "track-scl", "track-sc2" and "capture". See "tcp-request connection" above for their signification.

<condition> is a standard layer 4-7 ACL-based condition (see section 7).

A request's contents can be analysed at an early stage of request processing called "TCP content inspection". During this stage, ACL-based rules are evaluated every time the request contents are updated, until either an "accept" or a "reject" rule matches, or the TCP request inspection delay expires with no matching rule.

The first difference between these rules and "tcp-request connection" rules is that "tcp-request content" rules can make use of contents to take a decision. Most often, these decisions will consider a protocol recognition or validity. The second difference is that content-based rules can be used in both frontends and backends. In case of HTTP keep-alive with the client, all tcp-request content rules are evaluated again, so haproxy keeps a record of what sticky counters were assigned by a "tcp-request connection" versus a "tcp-request content" rule, and flushes all the content-related ones after processing an HTTP request, so that they may be evaluated again by the rules being evaluated again for the next request. This is of particular importance when the rule tracks some L7 information or when it is conditioned by an L7-based ACL, since tracking may change between requests.

Content-based rules are evaluated in their exact declaration order. If no rule matches or if there is no rule, the default action is to accept the contents. There is no specific limit to the number of rules which may be inserted.

Four types of actions are supported :

- accept : the request is accepted
- reject : the request is rejected and the connection is closed
- capture : the specified sample expression is captured

```
7541 - { track-sc0 | track-scl | track-sc2 } <key> [table <table>]
7542
7543 They have the same meaning as their counter-parts in "tcp-request connection"
7544 so please refer to that section for a complete description.
7545
7546 While there is nothing mandatory about it, it is recommended to use the
7547 track-sc0 in "tcp-request connection" rules, track-scl for "tcp-request
7548 content" rules in the frontend, and track-sc2 for "tcp-request content"
7549 rules in the backend, because that makes the configuration more readable
7550 and easier to troubleshoot, but this is just a guideline and all counters
7551 may be used everywhere.
7552
7553 Note that the "if/unless" condition is optional. If no condition is set on
7554 the action, it is simply performed unconditionally. That can be useful for
7555 "track-sc*" actions as well as for changing the default action to a reject.
7556
7557 It is perfectly possible to match layer 7 contents with "tcp-request content"
7558 rules, since HTTP-specific ACL matches are able to preliminarily parse the
7559 contents of a buffer before extracting the required data. If the buffered
7560 contents do not parse as a valid HTTP message, then the ACL does not match.
7561 The parser which is involved there is exactly the same as for all other HTTP
7562 processing, so there is no risk of parsing something differently. In an HTTP
7563 backend connected to from an HTTP frontend, it is guaranteed that HTTP
7564 contents will always be immediately present when the rule is evaluated first.
7565
7566 Tracking layer7 information is also possible provided that the information
7567 are present when the rule is processed. The rule processing engine is able to
7568 wait until the inspect delay expires when the data to be tracked is not yet
7569 available.
7570
7571 Example:
7572 # Accept HTTP requests containing a Host header saying "example.com"
7573 # and reject everything else.
7574 acl is_host_com hdr(Host) -i example.com
7575 tcp-request inspect-delay 30s
7576 tcp-request content accept if is_host_com
7577 tcp-request content reject
7578
7579 Example:
7580 # reject SMTP connection if client speaks first
7581 tcp-request inspect-delay 30s
7582 acl content_present req_len gt 0
7583 tcp-request content reject if content_present
7584
7585 # Forward HTTPS connection only if client speaks
7586 tcp-request inspect-delay 30s
7587 acl content_present req_len gt 0
7588 tcp-request content accept if content_present
7589 tcp-request content reject
7590
7591 Example:
7592 # Track the last IP from X-Forwarded-For
7593 tcp-request inspect-delay 10s
7594 tcp-request content track-sc0 hdr(x-forwarded-for,-1)
7595
7596 Example:
7597 # track request counts per "base" (concatenation of Host+URL)
7598 tcp-request inspect-delay 10s
7599 tcp-request content track-sc0 base table req-rate
7600
7601 Example: track per-frontend and per-backend counters, block abusers at the
7602 frontend when the backend detects abuse.
7603
7604 frontend http
7605     # Use General Purpose Counter 0 in SC0 as a global abuse counter
```

```
7606 # protecting all our sites
7607 stick-table type ip size 1m expire 5m store gpc0
7608 tcp-request connection track-sc0 src
7609 tcp-request connection reject if { sc0_get_gpc0 gt 0 }
7610 ...
7611 use_backend http_dynamic if { path_end .php }
7612
7613 backend http_dynamic
7614 # if a source makes too fast requests to this dynamic site (tracked
7615 # by SCL), block it globally in the frontend.
7616 stick-table type ip size 1m expire 5m store 1m http_req_rate(10s)
7617 acl click_too_fast scl_http_req_rate gt 10
7618 acl mark_as_abuser sc0_inc_gpc0 gt 0
7619 tcp-request content track-scl src
7620 tcp-request content reject if click_too_fast mark_as_abuser
7621
7622 See section 7 about ACL usage.
7623
7624 See also : "tcp-request connection", "tcp-request inspect-delay"
7625
7626 tcp-request inspect-delay <timeout>
7627 Set the maximum allowed time to wait for data during content inspection
7628 May be used in sections : defaults | frontend | listen | backend
7629                          no | yes | yes | yes | yes
7630 Arguments :
7631 <timeout> is the timeout value specified in milliseconds by default, but
7632 can be in any other unit if the number is suffixed by the unit,
7633 as explained at the top of this document.
7634
7635 People using haproxy primarily as a TCP relay are often worried about the
7636 risk of passing any type of protocol to a server without any analysis. In
7637 order to be able to analyze the request contents, we must first withhold
7638 the data then analyze them. This statement simply enables withholding of
7639 data for at most the specified amount of time.
7640
7641 TCP content inspection applies very early when a connection reaches a
7642 frontend, then very early when the connection is forwarded to a backend. This
7643 means that a connection may experience a first delay in the frontend and a
7644 second delay in the backend if both have tcp-request rules.
7645
7646 Note that when performing content inspection, haproxy will evaluate the whole
7647 rules for every new chunk which gets in, taking into account the fact that
7648 those data are partial. If no rule matches before the aforementioned delay,
7649 a last check is performed upon expiration, this time considering that the
7650 contents are definitive. If no delay is set, haproxy will not wait at all
7651 and will immediately apply a verdict based on the available information.
7652 Obviously this is unlikely to be very useful and might even be racy, so such
7653 setups are not recommended.
7654
7655 As soon as a rule matches, the request is released and continues as usual. If
7656 the timeout is reached and no rule matches, the default policy will be to let
7657 it pass through unaffected.
7658
7659 For most protocols, it is enough to set it to a few seconds, as most clients
7660 send the full request immediately upon connection. Add 3 or more seconds to
7661 cover TCP retransmits but that's all. For some protocols, it may make sense
7662 to use large values, for instance to ensure that the client never talks
7663 before the server (eg: SMTP), or to wait for a client to talk before passing
7664 data to the server (eg: SSL). Note that the client timeout must cover at
7665 least the inspection delay, otherwise it will expire first. If the client
7666 closes the connection or if the buffer is full, the delay immediately expires
7667 since the contents will not be able to change anymore.
7668
7669 See also : "tcp-request content accept", "tcp-request content reject",
7670
```

7671 "timeout client".
7672
7673
7674 tcp-response content <action> [{if | unless} <condition>]
7675 Perform an action on a session response depending on a layer 4-7 condition
7676 May be used in sections : defaults | frontend | listen | backend
7677 no | no | yes | yes
7678
7679 Arguments :
7680 <action> defines the action to perform if the condition applies. Valid
7681 actions include : "accept", "close", "reject".
7682
7683 <condition> is a standard layer 4-7 ACL-based condition (see section 7).
7684
7685 Response contents can be analysed at an early stage of response processing
7686 called "TCP content inspection". During this stage, ACL-based rules are
7687 evaluated every time the response contents are updated, until either an
7688 "accept", "close" or a "reject" rule matches, or a TCP response inspection
7689 delay is set and expires with no matching rule.
7690
7691 Most often, these decisions will consider a protocol recognition or validity.
7692
7693 Content-based rules are evaluated in their exact declaration order. If no
7694 rule matches or if there is no rule, the default action is to accept the
7695 contents. There is no specific limit to the number of rules which may be
7696 inserted.
7697
7698 Two types of actions are supported :
7699 - accept :
7700 accepts the response if the condition is true (when used with "if")
7701 or false (when used with "unless"). The first such rule executed ends
7702 the rules evaluation.
7703
7704 - close :
7705 immediately closes the connection with the server if the condition is
7706 true (when used with "if"), or false (when used with "unless"). The
7707 first such rule executed ends the rules evaluation. The main purpose of
7708 this action is to force a connection to be finished between a client
7709 and a server after an exchange when the application protocol expects
7710 some long time outs to elapse first. The goal is to eliminate idle
7711 connections which take significant resources on servers with certain
7712 protocols.
7713
7714 - reject :
7715 rejects the response if the condition is true (when used with "if")
7716 or false (when used with "unless"). The first such rule executed ends
7717 the rules evaluation. Rejected session are immediately closed.
7718
7719 Note that the "if/unless" condition is optional. If no condition is set on
7720 the action, it is simply performed unconditionally. That can be useful for
7721 for changing the default action to a reject.
7722
7723 It is perfectly possible to match layer 7 contents with "tcp-response
7724 content" rules, but then it is important to ensure that a full response has
7725 been buffered, otherwise no contents will match. In order to achieve this,
7726 the best solution involves detecting the HTTP protocol during the inspection
7727 period.
7728
7729 See section 7 about ACL usage.
7730
7731 See also : "tcp-request content", "tcp-response inspect-delay"
7732
7733 tcp-response inspect-delay <timeout>
7734 Set the maximum allowed time to wait for a response during content inspection
7735 May be used in sections : defaults | frontend | listen | backend

7736 Arguments :
7737 no | no | yes | yes
7738 <timeout> is the timeout value specified in milliseconds by default, but
7739 can be in any other unit if the number is suffixed by the unit,
7740 as explained at the top of this document.
7741
7742 See also : "tcp-response content", "tcp-request inspect-delay".
7743
7744 timeout check <timeout>
7745 Set additional check timeout, but only after a connection has been already
7746 established.
7747
7748 May be used in sections: defaults | frontend | listen | backend
7749 yes | no | yes | yes
7750
7751 Arguments:
7752 <timeout> is the timeout value specified in milliseconds by default, but
7753 can be in any other unit if the number is suffixed by the unit,
7754 as explained at the top of this document.
7755
7756 If set, haproxy uses min("timeout connect", "inter") as a connect timeout
7757 for check and "timeout check" as an additional read timeout. The "min" is
7758 used so that people running with *very* long "timeout connect" (eg. those
7759 who needed this due to the queue or tarpit) do not slow down their checks.
7760 (Please also note that there is no valid reason to have such long connect
7761 timeouts, because "timeout queue" and "timeout tarpit" can always be used to
7762 avoid that).
7763
7764 If "timeout check" is not set haproxy uses "inter" for complete check
7765 timeout (connect + read) exactly like all <1.3.15 version.
7766
7767 In most cases check request is much simpler and faster to handle than normal
7768 requests and people may want to kick out laggy servers so this timeout should
7769 be smaller than "timeout server".
7770
7771 This parameter is specific to backends, but can be specified once for all in
7772 "defaults" sections. This is in fact one of the easiest solutions not to
7773 forget about it.
7774
7775 See also: "timeout connect", "timeout queue", "timeout server",
7776 "timeout tarpit".
7777
7778 timeout client <timeout>
7779 timeout cli timeout <timeout> (deprecated)
7780 Set the maximum inactivity time on the client side.
7781 May be used in sections : defaults | frontend | listen | backend
7782 yes | yes | yes | no
7783
7784 Arguments :
7785 <timeout> is the timeout value specified in milliseconds by default, but
7786 can be in any other unit if the number is suffixed by the unit,
7787 as explained at the top of this document.
7788
7789 The inactivity timeout applies when the client is expected to acknowledge or
7790 send data. In HTTP mode, this timeout is particularly important to consider
7791 during the first phase, when the client sends the request, and during the
7792 response while it is reading data sent by the server. The value is specified
7793 in milliseconds by default, but can be in any other unit if the number is
7794 suffixed by the unit, as specified at the top of this document. In TCP mode
7795 (and to a lesser extent, in HTTP mode), it is highly recommended that the
7796 client timeout remains equal to the server timeout in order to avoid complex
7797 situations to debug. It is a good practice to cover one or several TCP packet
7798 losses by specifying timeouts that are slightly above multiples of 3 seconds
7799 (eg: 4 or 5 seconds). If some long-lived sessions are mixed with short-lived
7800 sessions (eg: WebSocket and HTTP), it's worth considering "timeout tunnel",

7801 which overrides "timeout client" and "timeout server" for tunnels, as well as
7802 "timeout client-fin" for half-closed connections.

7803
7804 This parameter is specific to frontends, but can be specified once for all in
7805 "defaults" sections. This is in fact one of the easiest solutions not to
7806 forget about it. An unspecified timeout results in an infinite timeout, which
7807 is not recommended. Such a usage is accepted and works but reports a warning
7808 during startup because it may results in accumulation of expired sessions in
7809 the system if the system's timeouts are not configured either.

7810
7811 This parameter replaces the old, deprecated "clitimeout". It is recommended
7812 to use it to write new configurations. The form "timeout clitimeout" is
7813 provided only by backwards compatibility but its use is strongly discouraged.

7814
7815 See also : "clitimeout", "timeout server", "timeout tunnel".

7816 timeout client-fin <timeout>

7817 Set the inactivity timeout on the client side for half-closed connections.
7818 May be used in sections : defaults | frontend | listen | backend

7819 yes | yes | yes | no

7820 Arguments :

7821 <timeout> is the timeout value specified in milliseconds by default, but
7822 can be in any other unit if the number is suffixed by the unit,
7823 as explained at the top of this document.

7824
7825 The inactivity timeout applies when the client is expected to acknowledge or
7826 send data while one direction is already shut down. This timeout is different
7827 from "timeout client" in that it only applies to connections which are closed
7828 in one direction. This is particularly useful to avoid keeping connections in
7829 FIN_WAIT state for too long when clients do not disconnect cleanly. This
7830 problem is particularly common long connections such as RDP or WebSocket.
7831 Note that this timeout can override "timeout tunnel" when a connection shuts
7832 down in one direction.

7833
7834 This parameter is specific to frontends, but can be specified once for all in
7835 "defaults" sections. By default it is not set, so half-closed connections
7836 will use the other timeouts (timeout.client or timeout.tunnel).

7837
7838 See also : "timeout client", "timeout server-fin", and "timeout tunnel".

7839 timeout connect <timeout>

7840 timeout contimeout <timeout> (deprecated)

7841 Set the maximum time to wait for a connection attempt to a server to succeed.

7842 May be used in sections : defaults | frontend | listen | backend
7843 yes | no | yes | yes

7844 Arguments :

7845 <timeout> is the timeout value specified in milliseconds by default, but
7846 can be in any other unit if the number is suffixed by the unit,
7847 as explained at the top of this document.

7848
7849 If the server is located on the same LAN as haproxy, the connection should be
7850 immediate (less than a few milliseconds). Anyway, it is a good practice to
7851 cover one or several TCP packet losses by specifying timeouts that are
7852 slightly above multiples of 3 seconds (eg: 4 or 5 seconds). By default, the
7853 connect timeout also presets both queue and tarpit timeouts to the same value
7854 if these have not been specified.

7855
7856 This parameter is specific to backends, but can be specified once for all in
7857 "defaults" sections. This is in fact one of the easiest solutions not to
7858 forget about it. An unspecified timeout results in an infinite timeout, which
7859 is not recommended. Such a usage is accepted and works but reports a warning
7860 during startup because it may results in accumulation of failed sessions in
7861 the system if the system's timeouts are not configured either.

7867 This parameter replaces the old, deprecated "contimeout". It is recommended
7868 to use it to write new configurations. The form "timeout contimeout" is
7869 provided only by backwards compatibility but its use is strongly discouraged.

7870
7871 See also: "timeout check", "timeout queue", "timeout server", "contimeout",
7872 "timeout tarpit".

7873 timeout http-keep-alive <timeout>

7874 Set the maximum allowed time to wait for a new HTTP request to appear
7875 May be used in sections : defaults | frontend | listen | backend
7876 yes | yes | yes | yes

7877 Arguments :
7878 <timeout> is the timeout value specified in milliseconds by default, but
7879 can be in any other unit if the number is suffixed by the unit,
7880 as explained at the top of this document.

7881
7882 By default, the time to wait for a new request in case of keep-alive is set
7883 by "timeout http-request". However this is not always convenient because some
7884 people want very short keep-alive timeouts in order to release connections
7885 faster, and others prefer to have larger ones but still have short timeouts
7886 once the request has started to present itself.

7887
7888 The "http-keep-alive" timeout covers these needs. It will define how long to
7889 wait for a new HTTP request to start coming after a response was sent. Once
7890 the first byte of request has been seen, the "http-request" timeout is used
7891 to wait for the complete request to come. Note that empty lines prior to a
7892 new request do not refresh the timeout and are not counted as a new request.

7893
7894 There is also another difference between the two timeouts : when a connection
7895 expires during timeout http-keep-alive, no error is returned, the connection
7896 just closes. If the connection expires in "http-request" while waiting for a
7897 connection to complete, a HTTP 408 error is returned.

7900
7901 In general it is optimal to set this value to a few tens to hundreds of
7902 milliseconds, to allow users to fetch all objects of a page at once but
7903 without waiting for further clicks. Also, if set to a very small value (eg:
7904 1 millisecond) it will probably only accept pipelined requests but not the
7905 non-pipelined ones. It may be a nice trade-off for very large sites running
7906 with tens to hundreds of thousands of clients.

7907
7908 If this parameter is not set, the "http-request" timeout applies, and if both
7909 are not set, "timeout client" still applies at the lower level. It should be
7910 set in the frontend to take effect, unless the frontend is in TCP mode, in
7911 which case the HTTP backend's timeout will be used.

7912
7913 See also : "timeout http-request", "timeout client".

7914 timeout http-request <timeout>

7915 Set the maximum allowed time to wait for a complete HTTP request
7916 May be used in sections : defaults | frontend | listen | backend
7917 yes | yes | yes | yes

7918 Arguments :
7919 <timeout> is the timeout value specified in milliseconds by default, but
7920 can be in any other unit if the number is suffixed by the unit,
7921 as explained at the top of this document.

7922
7923 In order to offer DoS protection, it may be required to lower the maximum
7924 accepted time to receive a complete HTTP request without affecting the client
7925 timeout. This helps protecting against established connections on which
7926 nothing is sent. The client timeout cannot offer a good protection against
7927 this abuse because it is an inactivity timeout, which means that if the
7928 attacker sends one character every now and then, the timeout will not
7929
7930

7931 trigger. With the HTTP request timeout, no matter what speed the client
7932 types, the request will be aborted if it does not complete in time. When the
7933 timeout expires, an HTTP 408 response is sent to the client to inform it
7934 about the problem, and the connection is closed. The logs will report
7935 termination codes "CR". Some recent browsers are having problems with this
7936 standard, well-documented behaviour, so it might be needed to hide the 408
7937 code using "option http-ignore-probes" or "errorfile 408/dev/null". See
7938 more details in the explanations of the "CR" termination code in section 8.5.
7939

7940 Note that this timeout only applies to the header part of the request, and
7941 not to any data. As soon as the empty line is received, this timeout is not
7942 used anymore. It is used again on keep-alive connections to wait for a second
7943 request if "timeout http-keep-alive" is not set.
7944

7945 Generally it is enough to set it to a few seconds, as most clients send the
7946 full request immediately upon connection. Add 3 or more seconds to cover TCP
7947 retransmits but that's all. Setting it to very low values (eg: 50 ms) will
7948 generally work on local networks as long as there are no packet losses. This
7949 will prevent people from sending bare HTTP requests using telnet.

7950 If this parameter is not set, the client timeout still applies between each
7951 chunk of the incoming request. It should be set in the frontend to take
7952 effect, unless the frontend is in TCP mode, in which case the HTTP backend's
7953 timeout will be used.
7954

7955 See also : "errorfile", "http-ignore-probes", "timeout http-keep-alive", and
7956 "timeout client".
7957

7958 timeout queue <timeout>
7959

7960 Set the maximum time to wait in the queue for a connection slot to be free
7961 May be used in sections : defaults | frontend | listen | backend
7962 yes | no | yes | yes

7963 Arguments :
7964 <timeout> is the timeout value specified in milliseconds by default, but
7965 can be in any other unit if the number is suffixed by the unit,
7966 as explained at the top of this document.
7967

7968 When a server's maxconn is reached, connections are left pending in a queue
7969 which may be server-specific or global to the backend. In order not to wait
7970 indefinitely, a timeout is applied to requests pending in the queue. If the
7971 timeout is reached, it is considered that the request will almost never be
7972 served, so it is dropped and a 503 error is returned to the client.
7973

7974 The "timeout queue" statement allows to fix the maximum time for a request to
7975 be left pending in a queue. If unspecified, the same value as the backend's
7976 connection timeout ("timeout connect") is used, for backwards compatibility
7977 with older versions with no "timeout queue" parameter.
7978

7979 See also : "timeout connect", "contimeout".
7980

7981 timeout server <timeout>
7982

7983 timeout srvtimeout <timeout> (deprecated)
7984

7985 Set the maximum inactivity time on the server side.
7986

7987 May be used in sections : defaults | frontend | listen | backend
7988 yes | no | yes | yes

7989 Arguments :

7990 <timeout> is the timeout value specified in milliseconds by default, but
7991 can be in any other unit if the number is suffixed by the unit,
7992 as explained at the top of this document.

7993 The inactivity timeout applies when the server is expected to acknowledge or
7994 send data. In HTTP mode, this timeout is particularly important to consider
7995 during the first phase of the server's response, when it has to send the

7996 headers, as it directly represents the server's processing time for the
7997 request. To find out what value to put there, it's often good to start with
7998 what would be considered as unacceptable response times, then check the logs
7999 to observe the response time distribution, and adjust the value accordingly.
8000

8001 The value is specified in milliseconds by default, but can be in any other
8002 unit if the number is suffixed by the unit, as specified at the top of this
8003 document. In TCP mode (and to a lesser extent, in HTTP mode), it is highly
8004 recommended that the client timeout remains equal to the server timeout in
8005 order to avoid complex situations to debug. Whatever the expected server
8006 response times, it is a good practice to cover at least one or several TCP
8007 packet losses by specifying timeouts that are slightly above multiples of 3
8008 seconds (eg: 4 or 5 seconds minimum). If some long-lived sessions are mixed
8009 with short-lived sessions (eg: WebSocket and HTTP), it's worth considering
8010 "timeout tunnel", which overrides "timeout client" and "timeout server" for
8011 tunnels.
8012

8013 This parameter is specific to backends, but can be specified once for all in
8014 "defaults" sections. This is in fact one of the easiest solutions not to
8015 forget about it. An unspecified timeout results in an infinite timeout, which
8016 is not recommended. Such a usage is accepted and works but reports a warning
8017 during startup because it may results in accumulation of expired sessions in
8018 the system if the system's timeouts are not configured either.
8019

8020 This parameter replaces the old, deprecated "srvtimeout". It is recommended
8021 to use it to write new configurations. The form "timeout srvtimeout" is
8022 provided only by backwards compatibility but its use is strongly discouraged.
8023

8024 See also : "srvtimeout", "timeout client" and "timeout tunnel".
8025

8026 timeout server-fin <timeout>
8027

8028 Set the inactivity timeout on the server side for half-closed connections.
8029

8030 May be used in sections : defaults | frontend | listen | backend
8031 yes | no | yes | yes

8032 Arguments :

8033 <timeout> is the timeout value specified in milliseconds by default, but
8034 can be in any other unit if the number is suffixed by the unit,
8035 as explained at the top of this document.

8036 The inactivity timeout applies when the server is expected to acknowledge or
8037 send data while one direction is already shut down. This timeout is different
8038 from "timeout server" in that it only applies to connections which are closed
8039 in one direction. This is particularly useful to avoid keeping connections in
8040 FIN_WAIT state for too long when a remote server does not disconnect cleanly.
8041 This problem is particularly common long connections such as RDP or WebSocket.
8042 Note that this timeout can override "timeout tunnel" when a connection shuts
8043 down in one direction. This setting was provided for completeness, but in most
8044 situations, it should not be needed.

8045 This parameter is specific to backends, but can be specified once for all in
8046 "defaults" sections. By default it is not set, so half-closed connections
8047 will use the other timeouts (timeout.server or timeout.tunnel).
8048

8049 See also : "timeout client-fin", "timeout server", and "timeout tunnel".
8050

8051 timeout tarpit <timeout>
8052

8053 Set the duration for which tarpitted connections will be maintained
8054

8055 May be used in sections : defaults | frontend | listen | backend
8056 yes | yes | yes | yes

8057 Arguments :

8058 <timeout> is the tarpit duration specified in milliseconds by default, but
8059 can be in any other unit if the number is suffixed by the unit,
8060 as explained at the top of this document.

When a connection is tarptitted using "reqtarpit", it is maintained open with no activity for a certain amount of time, then closed. "timeout tarpit" defines how long it will be maintained open.

The value is specified in milliseconds by default, but can be in any other unit if the number is suffixed by the unit, as specified at the top of this document. If unspecified, the same value as the backend's connection timeout ("timeout connect") is used, for backwards compatibility with older versions with no "timeout tarpit" parameter.

See also : "timeout connect", "contimeout".

timeout tunnel <timeout>

Set the maximum inactivity time on the client and server side for tunnels.

May be used in sections : defaults | frontend | listen | backend
yes | no | yes | yes

Arguments :

<timeout> is the timeout value specified in milliseconds by default, but can be in any other unit if the number is suffixed by the unit, as explained at the top of this document.

The tunnel timeout applies when a bidirectional connection is established between a client and a server, and the connection remains inactive in both directions. This timeout supersedes both the client and server timeouts once the connection becomes a tunnel. In TCP, this timeout is used as soon as no analyser remains attached to either connection (eg: tcp content rules are accepted). In HTTP, this timeout is used when a connection is upgraded (eg: when switching to the WebSocket protocol, or forwarding a CONNECT request to a proxy), or after the first response when no keepalive/close option is specified.

Since this timeout is usually used in conjunction with long-lived connections, it usually is a good idea to also set "timeout client-fin" to handle the situation where a client suddenly disappears from the net and does not acknowledge a close, or sends a shutdown and does not acknowledge pending data anymore. This can happen in lossy networks where firewalls are present, and is detected by the presence of large amounts of sessions in a FIN_WAIT state.

The value is specified in milliseconds by default, but can be in any other unit if the number is suffixed by the unit, as specified at the top of this document. Whatever the expected normal idle time, it is a good practice to cover at least one or several TCP packet losses by specifying timeouts that are slightly above multiples of 3 seconds (eg: 4 or 5 seconds minimum).

This parameter is specific to backends, but can be specified once for all in "defaults" sections. This is in fact one of the easiest solutions not to forget about it.

Example :

```
defaults http
  option http-server-close
  timeout connect 5s
  timeout client 30s
  timeout client-fin 30s
  timeout server 30s
  timeout tunnel 1h # timeout to use with WebSocket and CONNECT
```

See also : "timeout client", "timeout client-fin", "timeout server".

transparent (deprecated)

Enable client-side transparent proxying

May be used in sections : defaults | frontend | listen | backend
yes | no | yes | yes

Arguments : none

This keyword was introduced in order to provide layer 7 persistence to layer 3 load balancers. The idea is to use the OS's ability to redirect an incoming connection for a remote address to a local process (here HAProxy), and let this process know what address was initially requested. When this option is used, sessions without cookies will be forwarded to the original destination IP address of the incoming request (which should match that of another equipment), while requests with cookies will still be forwarded to the appropriate server.

The "transparent" keyword is deprecated, use "option transparent" instead.

Note that contrary to a common belief, this option does NOT make HAProxy present the client's IP to the server when establishing the connection.

See also: "option transparent"

unique-id-format <string>

Generate a unique ID for each request.

May be used in sections : defaults | frontend | listen | backend
defaults | yes | yes | yes | no

Arguments :

<string> is a log-format string.

This keyword creates a ID for each request using the custom log format. A unique ID is useful to trace a request passing through many components of a complex infrastructure. The newly created ID may also be logged using the %ID tag the log-format string.

The format should be composed from elements that are guaranteed to be unique when combined together. For instance, if multiple haproxy instances are involved, it might be important to include the node name. It is often needed to log the incoming connection's source and destination addresses and ports. Note that since multiple requests may be performed over the same connection, including a request counter may help differentiate them. Similarly, a timestamp may protect against a rollover of the counter. Logging the process ID will avoid collisions after a service restart.

It is recommended to use hexadecimal notation for many fields since it makes them more compact and saves space in logs.

Example:

unique-id-format %{+X}o\ %ci:%cp_%fi:%fp_%Ts_%rt:%pid

will generate:

7F000001:8296_7F00001E:1F90_4F7B0A69_0003:790A

See also: "unique-id-header"

unique-id-header <name>

Add a unique ID header in the HTTP request.

May be used in sections : defaults | frontend | listen | backend
defaults | yes | yes | yes | no

Arguments :

<name> is the name of the header.

Add a unique-id header in the HTTP request sent to the server, using the unique-id-format. It can't work if the unique-id-format doesn't exist.

Example:

```
8191 unique-id-format %(+X)%\ %ci:%cp_%fi:%fp_%Ts_%rt:%pid
8192 unique-id-header X-Unique-ID
8193 will generate:
8194
8195 X-Unique-ID: 7F000001:8296_7F00001E:1F90_4F7B0A69_0003:790A
8196
8197 See also: "unique-id-format"
8198
8200 use_backend <backend> [{if | unless} <condition>]
8201 Switch to a specific backend if/unless an ACL-based condition is matched.
8202 May be used in sections : defaults | frontend | listen | backend
8203
8204 Arguments :
8205 no | yes | yes | no
8206
8207 <backend> is the name of a valid backend or "listen" section, or a
8208 "log-format" string resolving to a backend name.
8209
8210 <condition> is a condition composed of ACLs, as described in section 7. If
8211 it is omitted, the rule is unconditionally applied.
8212
8213 When doing content-switching, connections arrive on a frontend and are then
8214 dispatched to various backends depending on a number of conditions. The
8215 relation between the conditions and the backends is described with the
8216 "use_backend" keyword. While it is normally used with HTTP processing, it can
8217 also be used in pure TCP, either without content using stateless ACLs (eg:
8218 source address validation) or combined with a "tcp-request" rule to wait for
8219 some payload.
8220
8221 There may be as many "use_backend" rules as desired. All of these rules are
8222 evaluated in their declaration order, and the first one which matches will
8223 assign the backend.
8224
8225 In the first form, the backend will be used if the condition is met. In the
8226 second form, the backend will be used if the condition is not met. If no
8227 condition is valid, the backend defined with "default_backend" will be used.
8228 If no default backend is defined, either the servers in the same section are
8229 used (in case of a "listen" section) or, in case of a frontend, no server is
8230 used and a 503 service unavailable response is returned.
8231
8232 Note that it is possible to switch from a TCP frontend to an HTTP backend. In
8233 this case, either the frontend has already checked that the protocol is HTTP,
8234 and backend processing will immediately follow, or the backend will wait for
8235 a complete HTTP request to get in. This feature is useful when a frontend
8236 must decode several protocols on a unique port, one of them being HTTP.
8237
8238 When <backend> is a simple name, it is resolved at configuration time, and an
8239 error is reported if the specified backend does not exist. If <backend> is
8240 a log-format string instead, no check may be done at configuration time, so
8241 the backend name is resolved dynamically at run time. If the resulting
8242 backend name does not correspond to any valid backend, no other rule is
8243 evaluated, and the default_backend directive is applied instead. Note that
8244 when using dynamic backend names, it is highly recommended to use a prefix
8245 that no other backend uses in order to ensure that an unauthorized backend
8246 cannot be forced from the request.
8247
8248 It is worth mentioning that "use_backend" rules with an explicit name are
8249 used to detect the association between frontends and backends to compute the
8250 backend's "fullconn" setting. This cannot be done for dynamic names.
8251
8252 See also: "default_backend", "tcp-request", "fullconn", "log-format", and
8253 section 7 about ACLs.
8254
8255 use-server <server> if <condition>
```

```
8256 use-server <server> unless <condition>
8257 Only use a specific server if/unless an ACL-based condition is matched.
8258 May be used in sections : defaults | frontend | listen | backend
8259
8260 Arguments :
8261 no | no | yes | yes
8262
8263 <server> is the name of a valid server in the same backend section.
8264
8265 <condition> is a condition composed of ACLs, as described in section 7.
8266
8267 By default, connections which arrive to a backend are load-balanced across
8268 the available servers according to the configured algorithm, unless a
8269 persistence mechanism such as a cookie is used and found in the request.
8270
8271 Sometimes it is desirable to forward a particular request to a specific
8272 server without having to declare a dedicated backend for this server. This
8273 can be achieved using the "use-server" rules. These rules are evaluated after
8274 the "redirect" rules and before evaluating cookies, and they have precedence
8275 on them. There may be as many "use-server" rules as desired. All of these
8276 rules are evaluated in their declaration order, and the first one which
8277 matches will assign the server.
8278
8279 If a rule designates a server which is down, and "option persist" is not used
8280 and no force-persist rule was validated, it is ignored and evaluation goes on
8281 with the next rules until one matches.
8282
8283 In the first form, the server will be used if the condition is met. In the
8284 second form, the server will be used if the condition is not met. If no
8285 condition is valid, the processing continues and the server will be assigned
8286 according to other persistence mechanisms.
8287
8288 Note that even if a rule is matched, cookie processing is still performed but
8289 does not assign the server. This allows prefixed cookies to have their prefix
8290 stripped.
8291
8292 The "use-server" statement works both in HTTP and TCP mode. This makes it
8293 suitable for use with content-based inspection. For instance, a server could
8294 be selected in a farm according to the TLS SNI field. And if these servers
8295 have their weight set to zero, they will not be used for other traffic.
8296
8297 Example :
8298 # intercept incoming TLS requests based on the SNI field
8299 use-server www if { req_ssl_sni -i www.example.com }
8300 server www 192.168.0.1:443 weight 0
8301 use-server mail if { req_ssl_sni -i mail.example.com }
8302 server mail 192.168.0.1:587 weight 0
8303 use-server imap if { req_ssl_sni -i imap.example.com }
8304 server imap 192.168.0.1:993 weight 0
8305 # all the rest is forwarded to this server
8306 server default 192.168.0.2:443 check
8307
8308 See also: "use_backend", section 5 about server and section 7 about ACLs.
8309
8310 5. Bind and Server options
8311 -----
8312
8313 The "bind", "server" and "default-server" keywords support a number of settings
8314 depending on some build options and on the system hAProxy was built on. These
8315 settings generally each consist in one word sometimes followed by a value,
8316 written on the same line as the "bind" or "server" line. All these options are
8317 described in this section.
8318
8319 5.1. Bind options
8320 -----
```


8321 The "bind" keyword supports a certain number of settings which are all passed
8322 as arguments on the same line. The order in which those arguments appear makes
8323 no importance, provided that they appear after the bind address. All of these
8324 parameters are optional. Some of them consist in a single words (booleans),
8325 while other ones expect a value after them. In this case, the value must be
8326 provided immediately after the setting name.
8327

8328 The currently supported settings are the following ones.
8329
8330

8331 **accept-proxy**

8332 Enforces the use of the PROXY protocol over any connection accepted by any of
8333 the sockets declared on the same line. Versions 1 and 2 of the PROXY protocol
8334 are supported and correctly detected. The PROXY protocol dictates the layer
8335 3/4 addresses of the incoming connection to be used everywhere an address is
8336 used, with the only exception of "tcp-request connection" rules which will
8337 only see the real connection address. Logs will reflect the addresses
8338 indicated in the protocol, unless it is violated, in which case the real
8339 address will still be used. This keyword combined with support from external
8340 components can be used as an efficient and reliable alternative to the
8341 X-Forwarded-For mechanism which is not always reliable and not even always
8342 usable. See also "tcp-request connection expect-proxy" for a finer-grained
8343 setting of which client is allowed to use the protocol.
8344

8345 **alpn <protocols>**

8346 This enables the TLS ALPN extension and advertises the specified protocol
8347 list as supported on top of ALPN. The protocol list consists in a comma-
8348 delimited list of protocol names, for instance: "http/1.1,http/1.0" (without
8349 quotes). This requires that the SSL library is build with support for TLS
8350 extensions enabled (check with haproxy -vv). The ALPN extension replaces the
8351 initial NPN extension.

8352 **backlog <backlog>**

8353 Sets the socket's backlog to this value. If unspecified, the frontend's
8354 backlog is used instead, which generally defaults to the maxconn value.
8355
8356

8357 **edchke <named curves>**

8358 This setting is only available when support for OpenSSL was built in. It sets
8359 the named curve (RFC 4492) used to generate ECDH ephemeral keys. By default,
8360 used named curve is prime256v1.
8361

8362 **ca-file <cafile>**

8363 This setting is only available when support for OpenSSL was built in. It
8364 designates a PEM file from which to load CA certificates used to verify
8365 client's certificate.
8366

8367 **ca-ignore-err [all]<errorIDs>,...]**

8368 This setting is only available when support for OpenSSL was built in.
8369 Sets a comma separated list of errorIDs to ignore during verify at depth > 0.
8370 If set to 'all', all errors are ignored. SSL handshake is not aborted if an
8371 error is ignored.
8372

8373 **ciphers <ciphers>**

8374 This setting is only available when support for OpenSSL was built in. It sets
8375 the string describing the list of cipher algorithms ("cipher suite") that are
8376 negotiated during the SSL/TLS handshake. The format of the string is defined
8377 in "man 1 ciphers" from OpenSSL man pages, and can be for instance a string
8378 such as "AES:ALL:!aNULL:!eNULL:RC4:@STRENGTH" (without quotes).
8379

8380 **crl-file <crlfile>**

8381 This setting is only available when support for OpenSSL was built in. It
8382 designates a PEM file from which to load certificate revocation list used
8383 to verify client's certificate.
8384

8385 **crt <cert>**

8386
8387
8388
8389
8390
8391
8392
8393
8394
8395
8396
8397
8398
8399
8400
8401
8402
8403
8404
8405
8406
8407
8408
8409
8410
8411
8412
8413
8414
8415
8416
8417
8418
8419
8420
8421
8422
8423
8424
8425
8426
8427
8428
8429
8430
8431
8432
8433
8434
8435
8436
8437
8438
8439
8440
8441
8442
8443
8444
8445
8446
8447
8448
8449
8450

This setting is only available when support for OpenSSL was built in. It
designates a PEM file containing both the required certificates and any
associated private keys. This file can be built by concatenating multiple
PEM files into one (e.g. cat cert.pem key.pem > combined.pem). If your CA
requires an intermediate certificate, this can also be concatenated into this
file.

If the OpenSSL used supports Diffie-Hellman, parameters present in this file
are loaded.

If a directory name is used instead of a PEM file, then all files found in
that directory will be loaded in alphabetic order unless their name ends with
'-issuer' or '-ocsp' (reserved extensions). This directive may be specified
multiple times in order to load certificates from multiple files or
directories. The certificates will be presented to clients who provide a valid
TLS Server Name Indication field matching one of their CN or alt subjects.
Wildcards are supported, where a wildcard character '*' is used instead of the
first hostname component (eg: *.example.org matches www.example.org but not
www.sub.example.org).

If no SNI is provided by the client or if the SSL library does not support
TLS extensions, or if the client provides an SNI hostname which does not
match any certificate, then the first loaded certificate will be presented.
This means that when loading certificates from a directory, it is highly
recommended to load the default one first as a file or to ensure that it will
always be the first one in the directory.

Note that the same cert may be loaded multiple times without side effects.

Some CAs (such as Godaddy) offer a drop down list of server types that do not
include HAProxy when obtaining a certificate. If this happens be sure to
choose a webserver that the CA believes requires an intermediate CA (for
Godaddy, selection Apache Tomcat will get the correct bundle, but many
others, e.g. nginx, result in a wrong bundle that will not work for some
clients).

For each PEM file, haproxy checks for the presence of file at the same path
suffixed by ".ocsp". If such file is found, support for the TLS Certificate
Status Request extension (also known as "OCSP stapling") is automatically
enabled. The content of this file is optional. If not empty, it must contain
a valid OCSP Response in DER format. In order to be valid an OCSP Response
must comply with the following rules: it has to indicate a good status,
it has to be a single response for the certificate of the PEM file, and it
has to be valid at the moment of addition. If these rules are not respected
the OCSP Response is ignored and a warning is emitted. In order to identify
which certificate an OCSP Response applies to, the issuer's certificate is
necessary. If the issuer's certificate is not found in the PEM file, it will
be loaded from a file at the same path as the PEM file suffixed by ".issuer"
if it exists otherwise it will fail with an error.

crt-ignore-err <errors>

This setting is only available when support for OpenSSL was built in. Sets a
comma separated list of errorIDs to ignore during verify at depth == 0. If
set to 'all', all errors are ignored. SSL handshake is not aborted if an error
is ignored.

crt-list <file>

This setting is only available when support for OpenSSL was built in. It
designates a list of PEM file with an optional list of SNI filter per
certificate, with the following format for each line :

```
<crtfile> [[!]<sni-filter> ...]
```

Wildcards are supported in the SNI filter. Negative filter are also supported,
only useful in combination with a wildcard filter to exclude a particular SNI.

8451 The certificates will be presented to clients who provide a valid TLS Server
8452 Name Indication field matching one of the SNI filters. If no SNI filter is
8453 specified, the CN and alt subjects are used. This directive may be specified
8454 multiple times. See the "crt" option for more information. The default
8455 certificate is still needed to meet OpenSSL expectations. If it is not used,
8456 the 'strict-sni' option may be used.

8457 defer-accept

8458 Is an optional keyword which is supported only on certain Linux kernels. It
8459 states that a connection will only be accepted once some data arrive on it,
8460 or at worst after the first retransmit. This should be used only on protocols
8461 for which the client talks first (eg: HTTP). It can slightly improve
8462 performance by ensuring that most of the request is already available when
8463 the connection is accepted. On the other hand, it will not be able to detect
8464 connections which don't talk. It is important to note that this option is
8465 broken in all kernels up to 2.6.31, as the connection is never accepted until
8466 the client talks. This can cause issues with front firewalls which would see
8467 an established connection while the proxy will only see it in SYN_RECV. This
8468 option is only supported on TCPv4/TCPv6 sockets and ignored by other ones.

8470 force-sslv3

8471 This option enforces use of SSLv3 only on SSL connections instantiated from
8472 this listener. SSLv3 is generally less expensive than the TLS counterparts
8473 for high connection rates. This option is also available on global statement
8474 "ssl-default-bind-options". See also "no-tlsv*" and "no-sslv3".

8476 force-tlsv10

8477 This option enforces use of TLSv1.0 only on SSL connections instantiated from
8478 this listener. This option is also available on global statement
8479 "ssl-default-bind-options". See also "no-tlsv*" and "no-sslv3".

8481 force-tlsv11

8482 This option enforces use of TLSv1.1 only on SSL connections instantiated from
8483 this listener. This option is also available on global statement
8484 "ssl-default-bind-options". See also "no-tlsv*", and "no-sslv3".

8486 force-tlsv12

8487 This option enforces use of TLSv1.2 only on SSL connections instantiated from
8488 this listener. This option is also available on global statement
8489 "ssl-default-bind-options". See also "no-tlsv*", and "no-sslv3".

8491 gid <gid>

8492 Sets the group of the UNIX sockets to the designated system gid. It can also
8493 be set by default in the global section's "unix-bind" statement. Note that
8494 some platforms simply ignore this. This setting is equivalent to the "group"
8495 setting except that the group ID is used instead of its name. This setting is
8496 ignored by non UNIX sockets.

8498 group <group>

8499 Sets the group of the UNIX sockets to the designated system group. It can
8500 also be set by default in the global section's "unix-bind" statement. Note
8501 that some platforms simply ignore this. This setting is equivalent to the
8502 "gid" setting except that the group name is used instead of its gid. This
8503 setting is ignored by non UNIX sockets.

8505 id <id>

8506 Fixes the socket ID. By default, socket IDs are automatically assigned, but
8507 sometimes it is more convenient to fix them to ease monitoring. This value
8508 must be strictly positive and unique within the listener/frontend. This
8509 option can only be used when defining only a single socket.

8510 interface <interface>

8511 Restricts the socket to a specific interface. When specified, only packets
8512 received from that particular interface are processed by the socket. This is
8513 currently only supported on Linux. The interface must be a primary system

8516 interface, not an aliased interface. It is also possible to bind multiple
8517 frontends to the same address if they are bound to different interfaces. Note
8518 that binding to a network interface requires root privileges. This parameter
8519 is only compatible with TCPv4/TCPv6 sockets.

8521 level <level>

8522 This setting is used with the stats sockets only to restrict the nature of
8523 the commands that can be issued on the socket. It is ignored by other
8524 sockets. <level> can be one of :
8525 - "user" is the least privileged level ; only non-sensitive stats can be
8526 read, and no change is allowed. It would make sense on systems where it
8527 is not easy to restrict access to the socket.
8528 - "operator" is the default level and fits most common uses. All data can
8529 be read, and only non-sensitive changes are permitted (eg: clear max
8530 counters).
8531 - "admin" should be used with care, as everything is permitted (eg: clear
8532 all counters).

8533 maxconn <maxconn>

8534 Limits the sockets to this number of concurrent connections. Extraneous
8535 connections will remain in the system's backlog until a connection is
8536 released. If unspecified, the limit will be the same as the frontend's
8537 maxconn. Note that in case of port ranges or multiple addresses, the same
8538 value will be applied to each socket. This setting enables different
8539 limitations on expensive sockets, for instance SSL entries which may easily
8540 eat all memory.

8542 mode <mode>

8543 Sets the octal mode used to define access permissions on the UNIX socket. It
8544 can also be set by default in the global section's "unix-bind" statement.
8545 Note that some platforms simply ignore this. This setting is ignored by non
8546 UNIX sockets.

8548 mss <maxseq>

8549 Sets the TCP Maximum Segment Size (MSS) value to be advertised on incoming
8550 connections. This can be used to force a lower MSS for certain specific
8551 ports, for instance for connections passing through a VPN. Note that this
8552 relies on a kernel feature which is theoretically supported under Linux but
8553 was buggy in all versions prior to 2.6.28. It may or may not work on other
8554 operating systems. It may also not change the advertised value but change the
8555 effective size of outgoing segments. The commonly advertised value for IPv4
8556 over Ethernet networks is 1460 = 1500(MTU) - 40(IP+TCP). If this value is
8557 positive, it will be used as the advertised MSS. If it is negative, it will
8558 indicate by how much to reduce the incoming connection's advertised MSS for
8559 outgoing segments. This parameter is only compatible with TCP v4/v6 sockets.

8561 name <name>

8562 Sets an optional name for these sockets, which will be reported on the stats
8563 page.

8564 nice <nice>

8565 Sets the 'niceness' of connections initiated from the socket. Value must be
8566 in the range -1024..1024 inclusive, and defaults to zero. Positive values
8567 means that such connections are more friendly to others and easily offer
8568 their place in the scheduler. On the opposite, negative values mean that
8569 connections want to run with a higher priority than others. The difference
8570 only happens under high loads when the system is close to saturation.
8571 Negative values are appropriate for low-latency or administration services,
8572 and high values are generally recommended for CPU intensive tasks such as SSL
8573 processing or bulk transfers which are less sensible to latency. For example,
8574 it may make sense to use a positive value for an SMTP socket and a negative
8575 one for an RDP socket.

8578 no-sslv3

8579 This setting is only available when support for OpenSSL was built in. It

disables support for SSLv3 on any sockets instantiated from the listener when SSL is supported. Note that SSLv2 is forced disabled in the code and cannot be enabled using any configuration option. This option is also available on global statement "ssl-default-bind-options". See also "force-tls*", and "force-sslv3".

no-tls-tickets

This setting is only available when support for OpenSSL was built in. It disables the stateless session resumption (RFC 5077 TLS Ticket extension) and force to use stateful session resumption. Stateless session resumption is more expensive in CPU usage. This option is also available on global statement "ssl-default-bind-options".

no-tlsv10

This setting is only available when support for OpenSSL was built in. It disables support for TLSv1.0 on any sockets instantiated from the listener when SSL is supported. Note that SSLv2 is forced disabled in the code and cannot be enabled using any configuration option. This option is also available on global statement "ssl-default-bind-options". See also "force-tlsv*", and "force-sslv3".

no-tlsv11

This setting is only available when support for OpenSSL was built in. It disables support for TLSv1.1 on any sockets instantiated from the listener when SSL is supported. Note that SSLv2 is forced disabled in the code and cannot be enabled using any configuration option. This option is also available on global statement "ssl-default-bind-options". See also "force-tlsv*", and "force-sslv3".

no-tlsv12

This setting is only available when support for OpenSSL was built in. It disables support for TLSv1.2 on any sockets instantiated from the listener when SSL is supported. Note that SSLv2 is forced disabled in the code and cannot be enabled using any configuration option. This option is also available on global statement "ssl-default-bind-options". See also "force-tlsv*", and "force-sslv3".

npn <protocols>

This enables the NPN TLS extension and advertises the specified protocol list as supported on top of NPN. The protocol list consists in a comma-delimited list of protocol names, for instance: "http/1.1,http/1.0" (without quotes). This requires that the SSL library is build with support for TLS extensions enabled (check with haproxy -vv). Note that the NPN extension has been replaced with the ALPN extension (see the "alpn" keyword).

```
process [ all | odd | even | <number 1-64> [ -<number 1-64> ] ]
```

This restricts the list of processes on which this listener is allowed to run. It does not enforce any process but eliminates those which do not match. If the frontend uses a "bind-process" setting, the intersection between the two is applied. If in the end the listener is not allowed to run on any remaining process, a warning is emitted, and the listener will either run on the first process of the listener if a single process was specified, or on all of its processes if multiple processes were specified. For the unlikely case where several ranges are needed, this directive may be repeated. The main purpose of this directive is to be used with the stats sockets and have one different socket per process. The second purpose is to have multiple bind lines sharing the same IP:port but not the same process in a listener, so that the system can distribute the incoming connections into multiple queues and allow a smoother inter-process load balancing. Currently Linux 3.9 and above is known for supporting this. See also "bind-process" and "nproc".

ssl

This setting is only available when support for OpenSSL was built in. It enables SSL deciphering on connections instantiated from this listener. A certificate is necessary (see "crt" above). All contents in the buffers will

appear in clear text, so that ACLs and HTTP processing will only have access to deciphered contents.

strict-sni

This setting is only available when support for OpenSSL was built in. The SSL/TLS negotiation is allow only if the client provided an SNI which match a certificate. The default certificate is not used. See the "crt" option for more information.

tfo

Is an optional keyword which is supported only on Linux kernels >= 3.7. It enables TCP Fast Open on the listening socket, which means that clients which support this feature will be able to send a request and receive a response during the 3-way handshake starting from second connection, thus saving one round-trip after the first connection. This only makes sense with protocols that use high connection rates and where each round trip matters. This can possibly cause issues with many firewalls which do not accept data on SYN packets, so this option should only be enabled once well tested. This option is only supported on TCPv4/TCPv6 sockets and ignored by other ones. You may need to build HAProxy with USE_TFO=1 if your libc doesn't define TCP_FASTOPEN.

transparent

Is an optional keyword which is supported only on certain Linux kernels. It indicates that the addresses will be bound even if they do not belong to the local machine, and that packets targeting any of these addresses will be intercepted just as if the addresses were locally configured. This normally requires that IP forwarding is enabled. Caution! do not use this with the default address '*', as it would redirect any traffic for the specified port. This keyword is available only when HAProxy is built with USE_LINUX_TPROXY=1. This parameter is only compatible with TCPv4 and TCPv6 sockets, depending on kernel version. Some distribution kernels include backports of the feature, so check for support with your vendor.

v4v6

Is an optional keyword which is supported only on most recent systems including Linux kernels >= 2.4.21. It is used to bind a socket to both IPv4 and IPv6 when it uses the default address. Doing so is sometimes necessary on systems which bind to IPv6 only by default. It has no effect on non-IPv6 sockets, and is overridden by the "v6only" option.

v6only

Is an optional keyword which is supported only on most recent systems including Linux kernels >= 2.4.21. It is used to bind a socket to IPv6 only when it uses the default address. Doing so is sometimes preferred to doing it system-wide as it is per-listener. It has no effect on non-IPv6 sockets and has precedence over the "v4v6" option.

uid <uid>

Sets the owner of the UNIX sockets to the designated system uid. It can also be set by default in the global section's "unix-bind" statement. Note that some platforms simply ignore this. This setting is equivalent to the "user" setting except that the user numeric ID is used instead of its name. This setting is ignored by non UNIX sockets.

user <user>

Sets the owner of the UNIX sockets to the designated system user. It can also be set by default in the global section's "unix-bind" statement. Note that some platforms simply ignore this. This setting is equivalent to the "uid" setting except that the user name is used instead of its uid. This setting is ignored by non UNIX sockets.

verify [none|optional|required]

This setting is only available when support for OpenSSL was built in. If set to 'none', client certificate is not requested. This is the default. In other

cases, a client certificate is requested. If the client does not provide a certificate after the request and if 'verify' is set to 'required', then the handshake is aborted, while it would have succeeded if set to 'optional'. The certificate provided by the client is always verified using CAs from 'ca-file' and optional CRLs from 'crl-file'. On verify failure the handshake is aborted, regardless of the 'verify' option, unless the error code exactly matches one of those listed with 'ca-ignore-err' or 'crt-ignore-err'.

5.2. Server and default-server options

The "server" and "default-server" keywords support a certain number of settings which are all passed as arguments on the server line. The order in which those arguments appear does not count, and they are all optional. Some of those settings are single words (booleans) while others expect one or several values after them. In this case, the values must immediately follow the setting name. Except default-server, all those settings must be specified after the server's address if they are used:

```
server <name> <address>[:port] [settings ...]
default-server [settings ...]
```

The currently supported settings are the following ones.

addr <ipv4|ipv6>

Using the "addr" parameter, it becomes possible to use a different IP address to send health-checks. On some servers, it may be desirable to dedicate an IP address to specific component able to perform complex tests which are more suitable to health-checks than the application. This parameter is ignored if the "check" parameter is not set. See also the "port" parameter.

Supported in default-server: No

agent-check

Enable an auxiliary agent check which is run independently of a regular health check. An agent health check is performed by making a TCP connection to the port set by the "agent-port" parameter and reading an ASCII string. The string is made of a series of words delimited by spaces, tabs or commas in any order, optionally terminated by '\r' and/or '\n', each consisting of :

- An ASCII representation of a positive integer percentage, e.g. "75%". Values in this format will set the weight proportional to the initial weight of a server as configured when haproxy starts. Note that a zero weight is reported on the stats page as "DRAIN" since it has the same effect on the server (it's removed from the LB farm).

- The word "ready". This will turn the server's administrative state to the READY mode, thus cancelling any DRAIN or MAINT state

- The word "drain". This will turn the server's administrative state to the DRAIN mode, thus it will not accept any new connections other than those that are accepted via persistence.

- The word "maint". This will turn the server's administrative state to the MAINT mode, thus it will not accept any new connections at all, and health checks will be stopped.

- The words "down", "failed", or "stopped", optionally followed by a description string after a sharp ('#'). All of these mark the server's operating state as DOWN, but since the word itself is reported on the stats page, the difference allows an administrator to know if the situation was expected or not : the service may intentionally be stopped, may appear up but fail some validity tests, or may be seen as down (eg: missing process, or port not responding).

- The word "up" sets back the server's operating state as UP if health checks also report that the service is accessible.

Parameters which are not advertised by the agent are not changed. For example, an agent might be designed to monitor CPU usage and only report a relative weight and never interact with the operating status. Similarly, an agent could be designed as an end-user interface with 3 radio buttons allowing an administrator to change only the administrative state. However, it is important to consider that only the agent may revert its own actions, so if a server is set to DRAIN mode or to DOWN state using the agent, the agent must implement the other equivalent actions to bring the service into operations again.

Failure to connect to the agent is not considered an error as connectivity is tested by the regular health check which is enabled by the "check" parameter. Warning though, it is not a good idea to stop an agent after it reports "down", since only an agent reporting "up" will be able to turn the server up again. Note that the CLI on the Unix stats socket is also able to force an agent's result in order to workaround a bogus agent if needed.

Requires the "agent-port" parameter to be set. See also the "agent-inter" parameter.

Supported in default-server: No

agent-inter <delay>

The "agent-inter" parameter sets the interval between two agent checks to <delay> milliseconds. If left unspecified, the delay defaults to 2000 ms.

Just as with every other time-based parameter, it may be entered in any other explicit unit among { us, ms, s, m, h, d }. The "agent-inter" parameter also serves as a timeout for agent checks "timeout check" is not set. In order to reduce "resonance" effects when multiple servers are hosted on the same hardware, the agent and health checks of all servers are started with a small time offset between them. It is also possible to add some random noise in the agent and health checks interval using the global "spread-checks" keyword. This makes sense for instance when a lot of backends use the same servers.

See also the "agent-check" and "agent-port" parameters.

Supported in default-server: Yes

agent-port <port>

The "agent-port" parameter sets the TCP port used for agent checks.

See also the "agent-check" and "agent-inter" parameters.

Supported in default-server: Yes

backup

When "backup" is present on a server line, the server is only used in load balancing when all other non-backup servers are unavailable. Requests coming with a persistence cookie referencing the server will always be served though. By default, only the first operational backup server is used, unless the "allbackups" option is set in the backend. See also the "allbackups" option.

Supported in default-server: No

ca-file <cafile>

This setting is only available when support for OpenSSL was built in. It designates a PEM file from which to load CA certificates used to verify server's certificate.

8840

8841 Supported in default-server: No
8842
8843 check
8844 This option enables health checks on the server. By default, a server is
8845 always considered available. If "check" is set, the server is available when
8846 accepting periodic TCP connections, to ensure that it is really able to serve
8847 requests. The default address and port to send the tests to are those of the
8848 server, and the default source is the same as the one defined in the
8849 backend. It is possible to change the address using the "addr:" parameter, the
8850 port using the "port" parameter, the source address using the "source"
8851 address, and the interval and timers using the "inter", "rise" and "fall"
8852 parameters. The request method is define in the backend using the "httpchk",
8853 "smtpchk", "mysql-check", "pgsql-check" and "ssl-hello-chk" options. Please
8854 refer to those options and parameters for more information.
8855
8856 Supported in default-server: No
8857
8858 check-send-proxy
8859 This option forces emission of a PROXY protocol line with outgoing health
8860 checks, regardless of whether the server uses send-proxy or not for the
8861 normal traffic. By default, the PROXY protocol is enabled for health checks
8862 if it is already enabled for normal traffic and if no "port" nor "addr"
8863 directive is present. However, if such a directive is present, the
8864 "check-send-proxy" option needs to be used to force the use of the
8865 protocol. See also the "send-proxy" option for more information.
8866
8867 Supported in default-server: No
8868
8869 check-ssl
8870 This option forces encryption of all health checks over SSL, regardless of
8871 whether the server uses SSL or not for the normal traffic. This is generally
8872 used when an explicit "port" or "addr" directive is specified and SSL health
8873 checks are not inherited. It is important to understand that this option
8874 inserts an SSL transport layer below the checks, so that a simple TCP connect
8875 check becomes an SSL connect, which replaces the old ssl-hello-chk. The most
8876 common use is to send HTTPS checks by combining "httpchk" with SSL checks.
8877 All SSL settings are common to health checks and traffic (eg: ciphers).
8878 See the "ssl" option for more information.
8879
8880 Supported in default-server: No
8881
8882 ciphers <ciphers>
8883 This option sets the string describing the list of cipher algorithms that is
8884 is negotiated during the SSL/TLS handshake with the server. The format of the
8885 string is defined in "man 1 ciphers". When SSL is used to communicate with
8886 servers on the local network, it is common to see a weaker set of algorithms
8887 than what is used over the internet. Doing so reduces CPU usage on both the
8888 server and haproxy while still keeping it compatible with deployed software.
8889 Some algorithms such as RC4-SHA1 are reasonably cheap. If no security at all
8890 is needed and just connectivity, using DES can be appropriate.
8891
8892 Supported in default-server: No
8893
8894 cookie <value>
8895 The "cookie" parameter sets the cookie value assigned to the server to
8896 <value>. This value will be checked in incoming requests, and the first
8897 operational server possessing the same value will be selected. In return, in
8898 cookie insertion or rewrite modes, this value will be assigned to the cookie
8899 sent to the client. There is nothing wrong in having several servers sharing
8900 the same cookie value, and it is in fact somewhat common between normal and
8901 backup servers. See also the "cookie" keyword in backend section.
8902
8903 Supported in default-server: No
8904
8905 crl-file <crlfile>

8906 This setting is only available when support for OpenSSL was built in. It
8907 designates a PEM file from which to load certificate revocation list used
8908 to verify server's certificate.
8909
8910 Supported in default-server: No
8911
8912 crt <cert>
8913 This setting is only available when support for OpenSSL was built in.
8914 It designates a PEM file from which to load both a certificate and the
8915 associated private key. This file can be built by concatenating both PEM
8916 files into one. This certificate will be sent if the server send a client
8917 certificate request.
8918
8919 Supported in default-server: No
8920
8921 disabled
8922 The "disabled" keyword starts the server in the "disabled" state. That means
8923 that it is marked down in maintenance mode, and no connection other than the
8924 ones allowed by persist mode will reach it. It is very well suited to setup
8925 new servers, because normal traffic will never reach them, while it is still
8926 possible to test the service by making use of the force-persist mechanism.
8927
8928 Supported in default-server: No
8929
8930 error-limit <count>
8931 If health observing is enabled, the "error-limit" parameter specifies the
8932 number of consecutive errors that triggers event selected by the "on-error"
8933 option. By default it is set to 10 consecutive errors.
8934
8935 Supported in default-server: Yes
8936
8937 See also the "check", "error-limit" and "on-error".
8938
8939 fall <count>
8940 The "fall" parameter states that a server will be considered as dead after
8941 <count> consecutive unsuccessful health checks. This value defaults to 3 if
8942 unspecified. See also the "check", "inter" and "rise" parameters.
8943
8944 Supported in default-server: Yes
8945
8946 force-sslv3
8947 This option enforces use of SSLv3 only when SSL is used to communicate with
8948 the server. SSLv3 is generally less expensive than the TLS counterparts for
8949 high connection rates. This option is also available on global statement
8950 "ssl-default-server-options". See also "no-tlsv*", "no-sslv3".
8951
8952 Supported in default-server: No
8953
8954 force-tlsv10
8955 This option enforces use of TLSv1.0 only when SSL is used to communicate with
8956 the server. This option is also available on global statement
8957 "ssl-default-server-options". See also "no-tlsv*", "no-sslv3".
8958
8959 Supported in default-server: No
8960
8961 force-tlsv11
8962 This option enforces use of TLSv1.1 only when SSL is used to communicate with
8963 the server. This option is also available on global statement
8964 "ssl-default-server-options". See also "no-tlsv*", "no-sslv3".
8965
8966 Supported in default-server: No
8967
8968 force-tlsv12
8969 This option enforces use of TLSv1.2 only when SSL is used to communicate with
8970 the server. This option is also available on global statement

8971 "ssl-default-server-options". See also "no-tlsv*", "no-sslv3".
8972
8973 Supported in default-server: No
8974
8975 id <value>
8976 Set a persistent ID for the server. This ID must be positive and unique for
8977 the proxy. An unused ID will automatically be assigned if unset. The first
8978 assigned value will be 1. This ID is currently only returned in statistics.
8979
8980 Supported in default-server: No
8981
8982 inter <delay>
8983 fastinter <delay>
8984 downinter <delay>
8985 The "inter" parameter sets the interval between two consecutive health checks
8986 to <delay> milliseconds. If left unspecified, the delay defaults to 2000 ms.
8987 It is also possible to use "fastinter" and "downinter" to optimize delays
8988 between checks depending on the server state :
8989
8990 Server state | Interval used
8991 -----+-----
8992 UP 100% (non-transitional) | "inter"
8993 -----+-----
8994 Transitionally UP (going down), |
8995 Transitionally DOWN (going up), | "fastinter" if set, "inter" otherwise.
8996 or yet unchecked. |
8997 -----+-----
8998 DOWN 100% (non-transitional) | "downinter" if set, "inter" otherwise.
8999 -----+-----
9000
9001 Just as with every other time-based parameter, they can be entered in any
9002 other explicit unit among { us, ms, s, m, h, d }. The "inter" parameter also
9003 serves as a timeout for health checks sent to servers if "timeout check" is
9004 not set. In order to reduce "resonance" effects when multiple servers are
9005 hosted on the same hardware, the agent and health checks of all servers
9006 are started with a small time offset between them. It is also possible to
9007 add some random noise in the agent and health checks interval using the
9008 global "spread-checks" keyword. This makes sense for instance when a lot
9009 of backends use the same servers.
9010
9011 Supported in default-server: Yes
9012
9013 maxconn <maxconn>
9014 The "maxconn" parameter specifies the maximal number of concurrent
9015 connections that will be sent to this server. If the number of incoming
9016 concurrent requests goes higher than this value, they will be queued, waiting
9017 for a connection to be released. This parameter is very important as it can
9018 save fragile servers from going down under extreme loads. If a "minconn"
9019 parameter is specified, the limit becomes dynamic. The default value is "0"
9020 which means unlimited. See also the "minconn" and "maxqueue" parameters, and
9021 the backend's "fullconn" keyword.
9022
9023 Supported in default-server: Yes
9024
9025 maxqueue <maxqueue>
9026 The "maxqueue" parameter specifies the maximal number of connections which
9027 will wait in the queue for this server. If this limit is reached, next
9028 requests will be redispached to other servers instead of indefinitely
9029 waiting to be served. This will break persistence but may allow people to
9030 quickly re-log in when the server they try to connect to is dying. The
9031 default value is "0" which means the queue is unlimited. See also the
9032 "maxconn" and "minconn" parameters.
9033
9034 Supported in default-server: Yes
9035

9036 minconn <minconn>
9037 When the "minconn" parameter is set, the maxconn limit becomes a dynamic
9038 limit following the backend's load. The server will always accept at least
9039 <minconn> connections, never more than <maxconn>, and the limit will be on
9040 the ramp between both values when the backend has less than <fullconn>
9041 concurrent connections. This makes it possible to limit the load on the
9042 server during normal loads, but push it further for important loads without
9043 overloading the server during exceptional loads. See also the "maxconn"
9044 and "maxqueue" parameters, as well as the "fullconn" backend keyword.
9045
9046 Supported in default-server: Yes
9047
9048 no-sslv3
9049 This option disables support for SSLv3 when SSL is used to communicate with
9050 the server. Note that SSLv2 is disabled in the code and cannot be enabled
9051 using any configuration option. See also "force-sslv3", "force-tlsv*".
9052
9053 Supported in default-server: No
9054
9055 no-tls-tickets
9056 This setting is only available when support for OpenSSL was built in. It
9057 disables the stateless session resumption (RFC 5077 TLS Ticket
9058 extension) and force to use stateful session resumption. Stateless
9059 session resumption is more expensive in CPU usage for servers. This option
9060 is also available on global statement "ssl-default-server-options".
9061
9062 Supported in default-server: No
9063
9064 no-tlsv10
9065 This option disables support for TLSv1.0 when SSL is used to communicate with
9066 the server. Note that SSLv2 is disabled in the code and cannot be enabled
9067 using any configuration option. TLSv1 is more expensive than SSLv3 so it
9068 often makes sense to disable it when communicating with local servers. This
9069 option is also available on global statement "ssl-default-server-options".
9070 See also "force-sslv3", "force-tlsv*".
9071
9072 Supported in default-server: No
9073
9074 no-tlsv11
9075 This option disables support for TLSv1.1 when SSL is used to communicate with
9076 the server. Note that SSLv2 is disabled in the code and cannot be enabled
9077 using any configuration option. TLSv1 is more expensive than SSLv3 so it
9078 often makes sense to disable it when communicating with local servers. This
9079 option is also available on global statement "ssl-default-server-options".
9080 See also "force-sslv3", "force-tlsv*".
9081
9082 Supported in default-server: No
9083
9084 no-tlsv12
9085 This option disables support for TLSv1.2 when SSL is used to communicate with
9086 the server. Note that SSLv2 is disabled in the code and cannot be enabled
9087 using any configuration option. TLSv1 is more expensive than SSLv3 so it
9088 often makes sense to disable it when communicating with local servers. This
9089 option is also available on global statement "ssl-default-server-options".
9090 See also "force-sslv3", "force-tlsv*".
9091
9092 Supported in default-server: No
9093
9094 non-stick
9095 Never add connections allocated to this sever to a stick-table.
9096 This may be used in conjunction with backup to ensure that
9097 stick-table persistence is disabled for backup servers.
9098
9099 Supported in default-server: No
9100

9101 observe <mode>
9102 This option enables health adjusting based on observing communication with
9103 the server. By default this functionality is disabled and enabling it also
9104 requires to enable health checks. There are two supported modes: "layer4" and
9105 "layer7". In layer4 mode, only successful/unsuccessful tcp connections are
9106 significant. In layer7, which is only allowed for http proxies, responses
9107 received from server are verified, like valid/wrong http code, unparseable
9108 headers, a timeout, etc. Valid status codes include 100 to 499, 501 and 505.
9109
9110 Supported in default-server: No
9111
9112 See also the "check", "on-error" and "error-limit".
9113
9114 on-error <mode>
9115 Select what should happen when enough consecutive errors are detected.
9116 Currently, four modes are available:
9117 - fastinter: force fastinter
9118 - fail-check: simulate a failed check, also forces fastinter (default)
9119 - sudden-death: simulate a pre-fatal failed health check, one more failed
9120 check will mark a server down, forces fastinter
9121 - mark-down: mark the server immediately down and force fastinter
9122
9123 Supported in default-server: Yes
9124
9125 See also the "check", "observe" and "error-limit".
9126
9127 on-marked-down <action>
9128 Modify what occurs when a server is marked down.
9129 Currently one action is available:
9130 - shutdown-sessions: Shutdown peer sessions. When this setting is enabled,
9131 all connections to the server are immediately terminated when the server
9132 goes down. It might be used if the health check detects more complex cases
9133 than a simple connection status, and long timeouts would cause the service
9134 to remain unresponsive for too long a time. For instance, a health check
9135 might detect that a database is stuck and that there's no chance to reuse
9136 existing connections anymore. Connections killed this way are logged with
9137 a 'D' termination code (for "Down").
9138
9139 Actions are disabled by default
9140
9141 Supported in default-server: Yes
9142
9143 on-marked-up <action>
9144 Modify what occurs when a server is marked up.
9145 Currently one action is available:
9146 - shutdown-backup-sessions: Shutdown sessions on all backup servers. This is
9147 done only if the server is not in backup state and if it is not disabled
9148 (it must have an effective weight > 0). This can be used sometimes to force
9149 an active server to take all the traffic back after recovery when dealing
9150 with long sessions (eg: LDAP, SQL, ...). Doing this can cause more trouble
9151 than it tries to solve (eg: incomplete transactions), so use this feature
9152 with extreme care. Sessions killed because a server comes up are logged
9153 with an 'U' termination code (for "Up").
9154
9155 Actions are disabled by default
9156
9157 Supported in default-server: Yes
9158
9159 port <port>
9160 Using the "port" parameter, it becomes possible to use a different port to
9161 send health-checks. On some servers, it may be desirable to dedicate a port
9162 to a specific component able to perform complex tests which are more suitable
9163 to health-checks than the application. It is common to run a simple script in
9164 inetd for instance. This parameter is ignored if the "check" parameter is not
9165 set. See also the "addr" parameter.

9166 Supported in default-server: Yes
9167
9168
9169 redir <prefix>
9170 The "redir" parameter enables the redirection mode for all GET and HEAD
9171 requests addressing this server. This means that instead of having HAProxy
9172 forward the request to the server, it will send an "HTTP 302" response with
9173 the "Location" header composed of this prefix immediately followed by the
9174 requested URI beginning at the leading '/' of the path component. That means
9175 that no trailing slash should be used after <prefix>. All invalid requests
9176 will be rejected, and all non-GET or HEAD requests will be normally served by
9177 the server. Note that since the response is completely forged, no header
9178 mangling nor cookie insertion is possible in the response. However, cookies in
9179 requests are still analysed, making this solution completely usable to direct
9180 users to a remote location in case of local disaster. Main use consists in
9181 increasing bandwidth for static servers by having the clients directly
9182 connect to them. Note: never use a relative location here, it would cause a
9183 loop between the client and HAProxy!
9184
9185 Example : server srv1 192.168.1.1:80 redir http://image1.mydomain.com check
9186
9187 Supported in default-server: No
9188
9189 rise <count>
9190 The "rise" parameter states that a server will be considered as operational
9191 after <count> consecutive successful health checks. This value defaults to 2
9192 if unspecified. See also the "check", "inter" and "fall" parameters.
9193
9194 Supported in default-server: Yes
9195
9196 send-proxy
9197 The "send-proxy" parameter enforces use of the PROXY protocol over any
9198 connection established to this server. The PROXY protocol informs the other
9199 end about the layer 3/4 addresses of the incoming connection, so that it can
9200 know the client's address or the public address it accessed to, whatever the
9201 upper layer protocol. For connections accepted by an "accept-proxy" listener,
9202 the advertised address will be used. Only TCPv4 and TCPv6 address families
9203 are supported. Other families such as Unix sockets, will report an UNKNOWN
9204 family. Servers using this option can fully be chained to another instance of
9205 haproxy listening with an "accept-proxy" setting. This setting must not be
9206 used if the server isn't aware of the protocol. When health checks are sent
9207 to the server, the PROXY protocol is automatically used when this option is
9208 set, unless there is an explicit "port" or "addr" directive, in which case an
9209 explicit "check-send-proxy" directive would also be needed to use the PROXY
9210 protocol. See also the "accept-proxy" option of the "bind" keyword.
9211
9212 Supported in default-server: No
9213
9214 send-proxy-v2
9215 The "send-proxy-v2" parameter enforces use of the PROXY protocol version 2
9216 over any connection established to this server. The PROXY protocol informs
9217 the other end about the layer 3/4 addresses of the incoming connection, so
9218 that it can know the client's address or the public address it accessed to,
9219 whatever the upper layer protocol. This setting must not be used if the
9220 server isn't aware of this version of the protocol. See also the "send-proxy"
9221 option of the "bind" keyword.
9222
9223 Supported in default-server: No
9224
9225 send-proxy-v2-ssl
9226 The "send-proxy-v2-ssl" parameter enforces use of the PROXY protocol version
9227 2 over any connection established to this server. The PROXY protocol informs
9228 the other end about the layer 3/4 addresses of the incoming connection, so
9229 that it can know the client's address or the public address it accessed to,
9230 whatever the upper layer protocol. In addition, the SSL information extension

9231 of the PROXY protocol is added to the PROXY protocol header. This setting
9232 must not be used if the server isn't aware of this version of the protocol.
9233 See also the "send-proxy-v2" option of the "bind" keyword.
9234
9235 Supported in default-server: No
9236
9237 send-proxy-v2-ssl-cn
9238 The "send-proxy-v2-ssl" parameter enforces use of the PROXY protocol version
9239 2 over any connection established to this server. The PROXY protocol informs
9240 the other end about the layer 3/4 addresses of the incoming connection, so
9241 that it can know the client's address or the public address it accessed to,
9242 whatever the upper layer protocol. In addition, the SSL information extension
9243 of the PROXY protocol, along with the Common Name from the subject of
9244 the client certificate (if any), is added to the PROXY protocol header. This
9245 setting must not be used if the server isn't aware of this version of the
9246 protocol. See also the "send-proxy-v2" option of the "bind" keyword.
9247
9248 Supported in default-server: No
9249
9250 slowstart <start_time_in_ms>
9251 The "slowstart" parameter for a server accepts a value in milliseconds which
9252 indicates after how long a server which has just come back up will run at
9253 full speed. Just as with every other time-based parameter, it can be entered
9254 in any other explicit unit among { us, ms, s, m, h, d }. The speed grows
9255 linearly from 0 to 100% during this time. The limitation applies to two
9256 parameters :
9257 - maxconn: the number of connections accepted by the server will grow from 1
9258 to 100% of the usual dynamic limit defined by (minconn,maxconn,fullconn).
9259
9260 - weight: when the backend uses a dynamic weighted algorithm, the weight
9261 grows linearly from 1 to 100%. In this case, the weight is updated at every
9262 health-check. For this reason, it is important that the "inter" parameter
9263 is smaller than the "slowstart", in order to maximize the number of steps.
9264
9265 The slowstart never applies when haproxy starts, otherwise it would cause
9266 trouble to running servers. It only applies when a server has been previously
9267 seen as failed.
9268
9269 Supported in default-server: Yes
9270
9271 source <addr>[:<pl>[:<ph>]] [usesrc { <addr2>[:<port2>] | client | clientip }]
9272 source <addr>[:<port>] [usesrc { <addr2>[:<port2>] | hdr_ip(<hdr>[,<occ>]) }]
9273 source <addr>[:<pl>[:<ph>]] [interface <name>] ...
9274 The "source" parameter sets the source address which will be used when
9275 connecting to the server. It follows the exact same parameters and principle
9276 as the backend "source" keyword, except that it only applies to the server
9277 referencing it. Please consult the "source" keyword for details.
9278
9279 Additionally, the "source" statement on a server line allows one to specify a
9280 source port range by indicating the lower and higher bounds delimited by a
9281 dash ('-'). Some operating systems might require a valid IP address when a
9282 source port range is specified. It is permitted to have the same IP/range for
9283 several servers. Doing so makes it possible to bypass the maximum of 64k
9284 total concurrent connections. The limit will then reach 64k connections per
9285 server.
9286
9287 Supported in default-server: No
9288
9289 ssl
9290 This option enables SSL ciphering on outgoing connections to the server. It
9291 is critical to verify server certificates using "verify" when using SSL to
9292 connect to servers, otherwise the communication is prone to trivial man in
9293 the-middle attacks rendering SSL useless. When this option is used, health
9294 checks are automatically sent in SSL too unless there is a "port" or an
9295

9296 "addr" directive indicating the check should be sent to a different location.
9297 See the "check-ssl" option to force SSL health checks.
9298
9299 Supported in default-server: No
9300
9301 track [<proxy>]/<server>
9302 This option enables ability to set the current state of the server by tracking
9303 another one. It is possible to track a server which itself tracks another
9304 server, provided that at the end of the chain, a server has health checks
9305 enabled. If <proxy> is omitted the current one is used. If disable-on-404 is
9306 used, it has to be enabled on both proxies.
9307
9308 Supported in default-server: No
9309
9310 verify [none|required]
9311 This setting is only available when support for OpenSSL was built in. If set
9312 to 'none', server certificate is not verified. In the other case, The
9313 certificate provided by the server is verified using CAs from 'ca-file'
9314 and optional CRLs from 'crl-file'. If 'ssl server verify' is not specified
9315 in global section, this is the default. On verify failure the handshake
9316 is aborted. It is critically important to verify server certificates when
9317 using SSL to connect to servers, otherwise the communication is prone to
9318 trivial man-in-the-middle attacks rendering SSL totally useless.
9319
9320 Supported in default-server: No
9321
9322 verifyhost <hostname>
9323 This setting is only available when support for OpenSSL was built in, and
9324 only takes effect if 'verify required' is also specified. When set, the
9325 hostnames in the subject and subjectAltNameNames of the certificate
9326 provided by the server are checked. If none of the hostnames in the
9327 certificate match the specified hostname, the handshake is aborted. The
9328 hostnames in the server-provided certificate may include wildcards.
9329
9330 Supported in default-server: No
9331
9332 weight <weight>
9333 The "weight" parameter is used to adjust the server's weight relative to
9334 other servers. All servers will receive a load proportional to their weight
9335 relative to the sum of all weights, so the higher the weight, the higher the
9336 load. The default weight is 1, and the maximal value is 256. A value of 0
9337 means the server will not participate in load-balancing but will still accept
9338 persistent connections. If this parameter is used to distribute the load
9339 according to server's capacity, it is recommended to start with values which
9340 can both grow and shrink, for instance between 10 and 100 to leave enough
9341 room above and below for later adjustments.
9342
9343 Supported in default-server: Yes
9344
9345
9346
9347 6. HTTP header manipulation
9348 -----
9349 In HTTP mode, it is possible to rewrite, add or delete some of the request and
9350 response headers based on regular expressions. It is also possible to block a
9351 request or a response if a particular header matches a regular expression,
9352 which is enough to stop most elementary protocol attacks, and to protect
9353 against information leak from the internal network.
9354
9355 If HAProxy encounters an "Informational Response" (status code 1xx), it is able
9356 to process all rsp* rules which can allow, deny, rewrite or delete a header,
9357 but it will refuse to add a header to any such messages as this is not
9358 HTTP-compliant. The reason for still processing headers in such responses is to
9359 stop and/or fix any possible information leak which may happen, for instance
9360 because another downstream equipment would unconditionally add a header, or if

9361 a server name appears there. When such messages are seen, normal processing
 9362 still occurs on the next non-informational messages.
 9363

9364 This section covers common usage of the following keywords, described in detail
 9365 in section 4.2 :
 9366

```

9367 - reqadd <string>
9368 - reqallow <search>
9369 - reqallow <search>
9370 - reqdel <search>
9371 - reqdel <search>
9372 - reqdeny <search>
9373 - reqdeny <search>
9374 - reqpass <search>
9375 - reqpass <search>
9376 - reqrep <search> <replace>
9377 - reqrep <search> <replace>
9378 - reqtarpit <search>
9379 - reqtarpit <search>
9380 - rspadd <string>
9381 - rspdel <search>
9382 - rspdel <search>
9383 - rspdeny <search>
9384 - rspdeny <search>
9385 - rsprep <search> <replace>
9386 - rsprep <search> <replace>
9387
```

9388 With all these keywords, the same conventions are used. The <search> parameter
 9389 is a POSIX extended regular expression (regex) which supports grouping through
 9390 parenthesis (without the backslash). Spaces and other delimiters must be
 9391 prefixed with a backslash ('\') to avoid confusion with a field delimiter.
 9392 Other characters may be prefixed with a backslash to change their meaning :

```

9393 \t for a tab
9394 \r for a carriage return (CR)
9395 \n for a new line (LF)
9396 \ to mark a space and differentiate it from a delimiter
9397 \# to mark a sharp and differentiate it from a comment
9398 \ to use a backslash in a regex
9399 \\\\ to use a backslash in the text (*2 for regex, *2 for haproxy)
9400 \xxx to write the ASCII hex code XX as in the C language
9401
9402
```

9403 The <replace> parameter contains the string to be used to replace the largest
 9404 portion of text matching the regex. It can make use of the special characters
 9405 above, and can reference a substring which is delimited by parenthesis in the
 9406 regex, by writing a backslash ('\') immediately followed by one digit from 0 to
 9407 9 indicating the group position (0 designating the entire line). This practice
 9408 is very common to users of the "sed" program.

9409 The <string> parameter represents the string which will systematically be added
 9410 after the last header line. It can also use special character sequences above.

9411 Notes related to these keywords :

9412 -----
 9413 - these keywords are not always convenient to allow/deny based on header
 9414 contents. It is strongly recommended to use ACLs with the "block" keyword
 9415 instead, resulting in far more flexible and manageable rules.

9416 - lines are always considered as a whole. It is not possible to reference
 9417 a header name only or a value only. This is important because of the way
 9418 headers are written (notably the number of spaces after the colon).

9419 - the first line is always considered as a header, which makes it possible to
 9420 rewrite or filter HTTP requests URIs or response codes, but in turn makes
 9421 it harder to distinguish between headers and request line. The regex prefix
 9422

9426 ^([^\t]*[^\t] matches any HTTP method followed by a space, and the prefix
 9427 ^[^\t:]*: matches any header name followed by a colon.
 9428

9429 - for performances reasons, the number of characters added to a request or to
 9430 a response is limited at build time to values between 1 and 4 kB. This
 9431 should normally be far more than enough for most usages. If it is too short
 9432 on occasional usages, it is possible to gain some space by removing some
 9433 useless headers before adding new ones.
 9434

9435 - keywords beginning with "reqi" and "rspi" are the same as their counterpart
 9436 without the 'i' letter except that they ignore case when matching patterns.
 9437
 9438 - when a request passes through a frontend then a backend, all req* rules
 9439 from the frontend will be evaluated, then all req* rules from the backend
 9440 will be evaluated. The reverse path is applied to responses.
 9441

9442 - req* statements are applied after "block" statements, so that "block" is
 9443 always the first one, but before "use_backend" in order to permit rewriting
 9444 before switching.
 9445

9446 7. Using ACLs and fetching samples

9447 -----

9448 Haproxy is capable of extracting data from request or response streams, from
 9449 client or server information, from tables, environmental information etc...
 9450 The action of extracting such data is called fetching a sample. Once retrieved,
 9451 these samples may be used for various purposes such as a key to a stick-table,
 9452 but most common usages consist in matching them against predefined constant
 9453 data called patterns.
 9454

9455 7.1. ACL basics

9456 -----

9457 The use of Access Control Lists (ACL) provides a flexible solution to perform
 9458 content switching and generally to take decisions based on content extracted
 9459 from the request, the response or any environmental status. The principle is
 9460 simple :

```

9461 - extract a data sample from a stream, table or the environment
9462 - optionally apply some format conversion to the extracted sample
9463 - apply one or multiple pattern matching methods on this sample
9464 - perform actions only when a pattern matches the sample
9465
```

9466 The actions generally consist in blocking a request, selecting a backend, or
 9467 adding a header.

9468 In order to define a test, the "acl" keyword is used. The syntax is :

```

9469 acl <aclname> <criterion> [flags] [operator] [<value>] ...
9470
```

9471 This creates a new ACL <aclname> or completes an existing one with new tests.
 9472 Those tests apply to the portion of request/response specified in <criterion>
 9473 and may be adjusted with optional flags [flags]. Some criteria also support
 9474 an operator which may be specified before the set of values. Optionally some
 9475 conversion operators may be applied to the sample, and they will be specified
 9476 as a comma-delimited list of keywords just after the first keyword. The values
 9477 are of the type supported by the criterion, and are separated by spaces.

9478 ACL names must be formed from upper and lower case letters, digits, '-' (dash),
 9479 '.' (underscore), ':' (dot) and ':' (colon). ACL names are case-sensitive,
 9480 which means that "my_acl" and "My_Acl" are two different ACLs.
 9481

9482 There is no enforced limit to the number of ACLs. The unused ones do not affect
 9483

performance, they just consume a small amount of memory.

The criterion generally is the name of a sample fetch method, or one of its ACL specific declinations. The default test method is implied by the output type of this sample fetch method. The ACL declinations can describe alternate matching methods of a same sample fetch method. The sample fetch methods are the only ones supporting a conversion.

Sample fetch methods return data which can be of the following types :

- boolean
- integer (signed or unsigned)
- IPv4 or IPv6 address
- string
- data block

Converters transform any of these data into any of these. For example, some converters might convert a string to a lower-case string while other ones would turn a string to an IPv4 address, or apply a netmask to an IP address. The resulting sample is of the type of the last converter applied to the list, which defaults to the type of the sample fetch method.

Each sample or converter returns data of a specific type, specified with its keyword in this documentation. When an ACL is declared using a standard sample fetch method, certain types automatically involved a default matching method which are summarized in the table below :

+	-----+	-----+
	Sample or converter	Default
	output type	matching method
+	-----+	-----+
	boolean	bool
+	-----+	-----+
	integer	int
+	-----+	-----+
	ip	ip
+	-----+	-----+
	string	str
+	-----+	-----+
	binary	none, use "-m"
+	-----+	-----+

Note that in order to match a binary samples, it is mandatory to specify a matching method, see below.

The ACL engine can match these types against patterns of the following types :

- boolean
- integer or integer range
- IP address / network
- string (exact, substring, suffix, prefix, subidir, domain)
- regular expression
- hex block

The following ACL flags are currently supported :

- i : ignore case during matching of all subsequent patterns.
- f : load patterns from a file.
- m : use a specific pattern matching method
- n : forbid the DNS resolutions
- M : load the file pointed by -f like a map file.
- u : force the unique id of the ACL
- : force end of flags. Useful when a string looks like one of the flags.

The "-f" flag is followed by the name of a file from which all lines will be read as individual values. It is even possible to pass multiple "-f" arguments if the patterns are to be loaded from multiple files. Empty lines as well as

lines beginning with a sharp ('#') will be ignored. All leading spaces and tabs will be stripped. If it is absolutely necessary to insert a valid pattern beginning with a sharp, just prefix it with a space so that it is not taken for a comment. Depending on the data type and match method, haproxy may load the lines into a binary tree, allowing very fast lookups. This is true for IPv4 and exact string matching. In this case, duplicates will automatically be removed.

The "-M" flag allows an ACL to use a map file. If this flag is set, the file is parsed as two column file. The first column contains the patterns used by the ACL, and the second column contain the samples. The sample can be used later by a map. This can be useful in some rare cases where an ACL would just be used to check for the existence of a pattern in a map before a mapping is applied.

The "-u" flag forces the unique id of the ACL. This unique id is used with the socket interface to identify ACL and dynamically change its values. Note that a file is always identified by its name even if an id is set.

Also, note that the "-i" flag applies to subsequent entries and not to entries loaded from files preceding it. For instance :

```
acl valid-ua hdr(user-agent) -f exact-ua.lst -i -f generic-ua.lst test
```

In this example, each line of "exact-ua.lst" will be exactly matched against the "user-agent" header of the request. Then each line of "generic-ua" will be case-insensitively matched. Then the word "test" will be insensitively matched as well.

The "-m" flag is used to select a specific pattern matching method on the input sample. All ACL-specific criteria imply a pattern matching method and generally do not need this flag. However, this flag is useful with generic sample fetch methods to describe how they're going to be matched against the patterns. This is required for sample fetches which return data type for which there is no obvious matching method (eg: string or binary). When "-m" is specified and followed by a pattern matching method name, this method is used instead of the default one for the criterion. This makes it possible to match contents in ways that were not initially planned, or with sample fetch methods which return a string. The matching method also affects the way the patterns are parsed.

The "-n" flag forbids the dns resolutions. It is used with the load of ip files. By default, if the parser cannot parse ip address it considers that the parsed string is maybe a domain name and try dns resolution. The flag "-n" disable this resolution. It is useful for detecting malformed ip lists. Note that if the DNS server is not reachable, the haproxy configuration parsing may last many minutes waiting fir the timeout. During this time no error messages are displayed. The flag "-n" disable this behavior. Note also that during the runtime, this function is disabled for the dynamic acl modifications.

There are some restrictions however. Not all methods can be used with all sample fetch methods. Also, if "-m" is used in conjunction with "-f", it must be placed first. The pattern matching method must be one of the following :

- "found" : only check if the requested sample could be found in the stream, but do not compare it against any pattern. It is recommended not to pass any pattern to avoid confusion. This matching method is particularly useful to detect presence of certain contents such as headers, cookies, etc... even if they are empty and without comparing them to anything nor counting them.

- "bool" : check the value as a boolean. It can only be applied to fetches which return a boolean or integer value, and takes no pattern. Value zero or false does not match, all other values do match.

- "int" : match the value as an integer. It can be used with integer and boolean samples. Boolean false is integer 0, true is integer 1.

```
- "ip"      : match the value as an IPv4 or IPv6 address. It is compatible
              with IP address samples only, so it is implied and never needed.
- "bin"     : match the contents against an hexadecimal string representing a
              binary sequence. This may be used with binary or string samples.
- "len"     : match the sample's length as an integer. This may be used with
              binary or string samples.
- "str"     : exact match : match the contents against a string. This may be
              used with binary or string samples.
- "sub"     : substring match : check that the contents contain at least one of
              the provided string patterns. This may be used with binary or
              string samples.
- "reg"     : regex match : match the contents against a list of regular
              expressions. This may be used with binary or string samples.
- "beg"     : prefix match : check that the contents begin like the provided
              string patterns. This may be used with binary or string samples.
- "end"     : suffix match : check that the contents end like the provided
              string patterns. This may be used with binary or string samples.
- "dir"     : subdir match : check that a slash-delimited portion of the
              contents exactly matches one of the provided string patterns.
              This may be used with binary or string samples.
- "dom"     : domain match : check that a dot-delimited portion of the contents
              exactly match one of the provided string patterns. This may be
              used with binary or string samples.
```

For example, to quickly detect the presence of cookie "JSESSIONID" in an HTTP request, it is possible to do :

```
acl jsess_present cook(JSESSIONID) -m found
```

In order to apply a regular expression on the 500 first bytes of data in the buffer, one would use the following acl :

```
acl script_tag payload(0,500) -m reg -i <script>
```

On systems where the regex library is much slower when using "-i", it is possible to convert the sample to lowercase before matching, like this :

```
acl script_tag payload(0,500),lower -m reg <script>
```

All ACL-specific criteria imply a default matching method. Most often, these criteria are composed by concatenating the name of the original sample fetch method and the matching method. For example, "hdr_beg" applies the "beg" match to samples retrieved using the "hdr" fetch method. Since all ACL-specific criteria rely on a sample fetch method, it is always possible instead to use the original sample fetch method and the explicit matching method using "-m".

If an alternate match is specified using "-m" on an ACL-specific criterion, the matching method is simply applied to the underlying sample fetch method. For example, all ACLs below are exact equivalent :

```
acl short_form  hdr_beg(host)          www.
acl alternate1  hdr_beg(host) -m beg www.
acl alternate2  hdr_dom(host) -m beg www.
acl alternate3  hdr(host)      -m beg www.
```

The table below summarizes the compatibility matrix between sample or converter types and the pattern types to fetch against. It indicates for each compatible combination the name of the matching method to be used, surrounded with angle brackets ">" and "<," when the method is the default one and will work by default without "-m".

	Input sample type									
	pattern type	boolean	integer	ip	string	binary	none (presence only)	found	found	found
	none (boolean value)	> bool <	bool		bool		none	boolean value	> bool <	bool
	integer (value)	int > int <	int		int		integer (value)	int > int <	int	
	integer (length)	len	len		len		integer (length)	len	len	
	IP address	> ip <	ip		ip		IP address	> ip <	ip	
	exact string	str	str		str	> str <	exact string	str	str	> str <
	prefix	beg	beg		beg		prefix	beg	beg	
	suffix	end	end		end		suffix	end	end	
	substring	sub	sub		sub		substring	sub	sub	
	subdir	dir	dir		dir		subdir	dir	dir	
	domain	dom	dom		dom		domain	dom	dom	
	regex	reg	reg		reg		regex	reg	reg	
	hex block				bin		hex block			

7.1.1. Matching booleans

In order to match a boolean, no value is needed and all values are ignored. Boolean matching is used by default for all fetch methods of type "boolean". When boolean matching is used, the fetched value is returned as-is, which means that a boolean "true" will always match and a boolean "false" will never match.

Boolean matching may also be enforced using "-m bool" on fetch methods which return an integer value. Then, integer value 0 is converted to the boolean "false" and all other values are converted to "true".

7.1.2. Matching integers

Integer matching applies by default to integer fetch methods. It can also be enforced on boolean fetches using "-m int". In this case, "false" is converted to the integer 0, and "true" is converted to the integer 1.

Integer matching also supports integer ranges and operators. Note that integer matching only applies to positive values. A range is a value expressed with a lower and an upper bound separated with a colon, both of which may be omitted.

For instance, "1024:65535" is a valid range to represent a range of unprivileged ports, and "1024:" would also work. "0:1023" is a valid

representation of privileged ports, and ":1023" would also work.

As a special case, some ACL functions support decimal numbers which are in fact two integers separated by a dot. This is used with some version checks for instance. All integer properties apply to those decimal numbers, including ranges and operators.

For an easier usage, comparison operators are also supported. Note that using operators with ranges does not make much sense and is strongly discouraged. Similarly, it does not make much sense to perform order comparisons with a set of values.

Available operators for integer matching are :

```
eq : true if the tested value equals at least one value
ge : true if the tested value is greater than or equal to at least one value
gt : true if the tested value is greater than at least one value
le : true if the tested value is less than or equal to at least one value
lt : true if the tested value is less than at least one value
```

For instance, the following ACL matches any negative Content-Length header :

```
acl negative-length hdr_val(content-length) lt 0
```

This one matches SSL versions between 3.0 and 3.1 (inclusive) :

```
acl sslv3 req_ssl_ver 3:3.1
```

7.1.3. Matching strings

String matching applies to string or binary fetch methods, and exists in 6 different forms :

- exact match (-m str) : the extracted string must exactly match the patterns ;
- substring match (-m sub) : the patterns are looked up inside the extracted string, and the ACL matches if any of them is found inside ;
- prefix match (-m beg) : the patterns are compared with the beginning of the extracted string, and the ACL matches if any of them matches.
- suffix match (-m end) : the patterns are compared with the end of the extracted string, and the ACL matches if any of them matches.
- subdir match (-m sub) : the patterns are looked up inside the extracted string, delimited with slashes ("/"), and the ACL matches if any of them matches.
- domain match (-m dom) : the patterns are looked up inside the extracted string, delimited with dots ("."), and the ACL matches if any of them matches.

String matching applies to verbatim strings as they are passed, with the exception of the backslash ("\") which makes it possible to escape some characters such as the space. If the "-i" flag is passed before the first string, then the matching will be performed ignoring the case. In order to match the string "-i", either set it second, or pass the "--" flag before the first string. Same applies of course to match the string "--".

7.1.4. Matching regular expressions (regexes)

Just like with string matching, regex matching applies to verbatim strings as they are passed, with the exception of the backslash ("\") which makes it possible to escape some characters such as the space. If the "-i" flag is passed before the first regex, then the matching will be performed ignoring the case. In order to match the string "-i", either set it second, or pass the "--" flag before the first string. Same principle applies of course to match the string "--".

7.1.5. Matching arbitrary data blocks

It is possible to match some extracted samples against a binary block which may not safely be represented as a string. For this, the patterns must be passed as a series of hexadecimal digits in an even number, when the match method is set to binary. Each sequence of two digits will represent a byte. The hexadecimal digits may be used upper or lower case.

Example :

```
# match "Hello\n" in the input stream (\x48 \x65 \x6c \x6f \x0a)
acl hello payload(0,6) -m bin 48656c6f0a
```

7.1.6. Matching IPv4 and IPv6 addresses

IPv4 addresses values can be specified either as plain addresses or with a netmask appended, in which case the IPv4 address matches whenever it is within the network. Plain addresses may also be replaced with a resolvable host name, but this practice is generally discouraged as it makes it more difficult to read and debug configurations. If hostnames are used, you should at least ensure that they are present in /etc/hosts so that the configuration does not depend on any random DNS match at the moment the configuration is parsed.

IPv6 may be entered in their usual form, with or without a netmask appended. Only bit counts are accepted for IPv6 netmasks. In order to avoid any risk of trouble with randomly resolved IP addresses, host names are never allowed in IPv6 patterns.

Haproxy is also able to match IPv4 addresses with IPv6 addresses in the following situations :

- tested address is IPv4, pattern address is IPv4, the match applies in IPv4 using the supplied mask if any.
- tested address is IPv6, pattern address is IPv6, the match applies in IPv6 using the supplied mask if any.
- tested address is IPv6, pattern address is IPv4, the match applies in IPv4 using the pattern's mask if the IPv6 address matches with 2002::IPV4:::IPV4 or ::ffff:IPV4, otherwise it fails.
- tested address is IPv4, pattern address is IPv6, the IPv4 address is first converted to IPv6 by prefixing ::ffff: in front of it, then the match is applied in IPv6 using the supplied IPv6 mask.

7.2. Using ACLs to form conditions

Some actions are only performed upon a valid condition. A condition is a combination of ACLs with operators. 3 operators are supported :

- AND (implicit)
- OR (explicit with the "or" keyword or the "||" operator)
- Negation with the exclamation mark ("!")

9800

A condition is formed as a disjunctive form:

```
[!acl1 [!acl2 ... [!acln { or [!acl1 [!acl2 ... [!acln } ...
```

Such conditions are generally used after an "if" or "unless" statement, indicating when the condition will trigger the action.

For instance, to block HTTP requests to the "*" URL with methods other than "OPTIONS", as well as POST requests without content-length, and GET or HEAD requests with a content-length greater than 0, and finally every request which is not either GET/HEAD/POST/OPTIONS !

```
acl missing_cl hdr_cnt(Content-length) eq 0
block if HTTP_URL_STAR !METH_OPTIONS || METH_POST missing_cl
block if METH_GET HTTP_CONTENT
block unless METH_GET or METH_POST or METH_OPTIONS
```

To select a different backend for requests to static contents on the "www" site and to every request on the "img", "video", "download" and "ftp" hosts :

```
acl url_static path_beg          /static/images /img /css
acl url_static path_end         .gif .png .jpg .css .js
acl host_www    hdr_beg(host) -i www
acl host_static hdr_beg(host) -i img. video. download. ftp.
```

```
# now use backend "static" for all static-only hosts, and for static urls
# of host "www". Use backend "www" for the rest.
use_backend static if host_static or host_www url_static
use_backend www    if host_www
```

It is also possible to form rules using "anonymous ACLs". Those are unnamed ACL expressions that are built on the fly without needing to be declared. They must be enclosed between braces, with a space before and after each brace (because the braces must be seen as independent words). Example :

The following rule :

```
acl missing_cl hdr_cnt(Content-length) eq 0
block if METH_POST missing_cl
```

Can also be written that way :

```
block if METH_POST { hdr_cnt(Content-length) eq 0 }
```

It is generally not recommended to use this construct because it's a lot easier to leave errors in the configuration when written that way. However, for very simple rules matching only one source IP address for instance, it can make more sense to use them than to declare ACLs with random names. Another example of good use is the following :

With named ACLs :

```
acl site_dead nbsrv(dynamic) lt 2
acl site_dead nbsrv(static)  lt 2
monitor fail if site_dead
```

With anonymous ACLs :

```
monitor fail if { nbsrv(dynamic) lt 2 } || { nbsrv(static) lt 2 }
```

See section 4.2 for detailed help on the "block" and "use_backend" keywords.

7.3. Fetching samples

Historically, sample fetch methods were only used to retrieve data to match against patterns using ACLs. With the arrival of stick-tables, a new class of sample fetch methods was created, most often sharing the same syntax as their ACL counterpart. These sample fetch methods are also known as "fetches". As of now, ACLs and fetches have converged. All ACL fetch methods have been made available as fetch methods, and ACLs may use any sample fetch method as well.

This section details all available sample fetch methods and their output type. Some sample fetch methods have deprecated aliases that are used to maintain compatibility with existing configurations. They are then explicitly marked as deprecated and should not be used in new setups.

The ACL derivatives are also indicated when available, with their respective matching methods. These ones all have a well defined default pattern matching method, so it is never necessary (though allowed) to pass the "-m" option to indicate how the sample will be matched using ACLs.

As indicated in the sample type versus matching compatibility matrix above, when using a generic sample fetch method in an ACL, the "-m" option is mandatory unless the sample type is one of boolean, integer, IPv4 or IPv6. When the same keyword exists as an ACL keyword and as a standard fetch method, the ACL engine will automatically pick the ACL-only one by default.

Some of these keywords support one or multiple mandatory arguments, and one or multiple optional arguments. These arguments are strongly typed and are checked when the configuration is parsed so that there is no risk of running with an incorrect argument (eg: an unresolved backend name). Fetch function arguments are passed between parenthesis and are delimited by commas. When an argument is optional, it will be indicated below between square brackets ('[']). When all arguments are optional, the parenthesis may be omitted.

Thus, the syntax of a standard sample fetch method is one of the following :

```
- name
- name(arg1)
- name(arg1,arg2)
```

7.3.1. Converters

Sample fetch methods may be combined with transformations to be applied on top of the fetched sample (also called "converters"). These combinations form what is called "sample expressions" and the result is a "sample". Initially this was only supported by "stick on" and "stick store-request" directives but this has now be extended to all places where samples may be used (acls, log-format, unique-id-format, add-header, ...).

These transformations are enumerated as a series of specific keywords after the sample fetch method. These keywords may equally be appended immediately after the fetch keyword's argument, delimited by a comma. These keywords can also support some arguments (eg: a netmask) which must be passed in parenthesis.

The currently available list of transformation keywords include :

base64

Converts a binary input sample to a base64 string. It is used to log or transfer binary content in a way that can be reliably transferred (eg: an SSL ID can be copied in a header).

hex

Converts a binary input sample to an hex string containing two hex digits per input byte. It is used to log or transfer hex dumps of some binary input data in a way that can be reliably transferred (eg: an SSL ID can be copied in a header).

`http_date([-offsets])` Converts an integer supposed to contain a date since epoch to a string representing this date in a format suitable for use in HTTP header fields. If an offset value is specified, then it is a number of seconds that is added to the date before the conversion is operated. This is particularly useful to emit Date header fields, Expires values in responses when combined with a positive offset, or Last-Modified values when the offset is negative.

ipmask(<mask>)

Apply a mask to an IPv4 address, and use the result for lookups and storage. This can be used to make all hosts within a certain mask to share the same table entries and as such use the same server. The mask can be passed in dotted form (eg: 255.255.255.0) or in CIDR form (eg: 24).

```
language(<value>[, <default>])
```

Returns the value with the highest q-factor from a list as extracted from the "accept-language" header using "req.fhdr". Values with no q-factor have a q-factor of 1. Values with a q-factor of 0 are dropped. Only values which belong to the list of semi-colon delimited <values> will be considered. The argument <values> syntax is "lang1;lang2;...;J|J". If no value matches the given list and a default value is provided, it is returned. Note that language names may have a variant after a dash ('-'). If this variant is present in the list, it will be matched, but if it is not, only the base language is checked. The match is case-sensitive, and the output string is always one of those provided in arguments. The ordering of arguments is meaningless, only the ordering of the values in the request counts, as the first value among multiple sharing the same q-factor is used.

Example :

```
# this configuration switches to the backend matching a
# given language based on the request :
```

```

acl es req. fhdr(accept-language), language(es;fr;en) -m str fr
acl fr req. fhdr(accept-language), language(es;fr;en) -m str fr
acl en req. fhdr(accept-language), language(es;fr;en) -m str en
use_backend spanish if es
use_backend french if fr
use_backend english if en
default backend choose your language

```

Lower

Convert a string sample to lower case. This can only be placed after a string sample fetch function or after a transformation keyword returning a string type. The result is of type string.

```
map(<map_file>[,<default_value>])
map_match_type(<map_file>[,<default_value>])
map_match_type(<map_file>[,<default_value>])
map_match_type(<map_file>[,<default_value>])
```

Search the input value from `<map_file>` using the `<match_type>` matching method, and return the associated value converted to the type `<output_type>`. If the input value cannot be found in the `<map_file>`, the converter returns the `<default_value>`. If the `<default_value>` is not set, the converter fails and acts as if no input value could be fetched. If the `<match_type>` is not set, it defaults to "str". Likewise, if the `<output_type>` is not set, it defaults to "str". For convenience, the "map" keyword is an alias for "map_str" and maps a string to another string.

It is important to avoid overlapping between the keys : IP addresses and strings are stored in trees, so the first of the finest match will be used. Other keys are stored in lists, so the first matching occurrence will be used.

The following array contains the list of all map functions available sorted by input type, match type and output type.

input type	match method	output type	str	output type	int	output type	ip
str	str	map_str		map_str_int		map_str_ip	
str	beg	map_beg		map_beg_int		map_end_ip	
str	sub	map_sub		map_sub_int		map_sub_ip	
str	dir	map_dir		map_dir_int		map_dir_ip	
str	dom	map_dom		map_dom_int		map_dom_ip	
str	end	map_end		map_end_int		map_end_ip	
str	reg	map_reg		map_reg_int		map_reg_ip	
int	int	map_int		map_int_int		map_int_ip	
ip	ip	map_ip		map_ip_int		map_ip_ip	

The file contains one key + value per line. Lines which start with '#' are ignored, just like empty lines. Leading tabs and spaces are stripped. The key is then the first "word" (series of non-space/tabs characters), and the value is what follows this series of space/tab till the end of the line excluding trailing spaces/tabs.

Example :

[illegible]

upper

Convert a string sample to upper case. This can only be placed after a string sample fetch function or after a transformation keyword returning a string type. The result is of type string.

7.3.2. Fetching samples from internal states

A first set of sample fetch methods applies to internal information which does not even relate to any client information. These ones are sometimes used with "monitor-fail" directives to report an internal status to external watchers. The sample fetch methods described in this section are usable anywhere.

always false : boolean

Always returns the boolean "false" value. It may be used with ACLs as a temporary replacement for another one when adjusting configurations.

always true : boolean

Always returns the boolean "true" value. It may be used with ACLs as a temporary replacement for another one when adjusting configurations.

```
avg queue(['<backend>']) : integer
```

Returns the total number of queued connections of the designated backend divided by the number of active servers. The current backend is used if no backend is specified. This is very similar to "queue" except that the size of

the farm is considered, in order to give a more accurate measurement of the time it may take for a new connection to be processed. The main usage is with ACL to return a sorry page to new users when it becomes certain they will get a degraded service, or to pass to the backend servers in a header so that they decide to work in degraded mode or to disable some functions to speed up the processing a bit. Note that in the event there would not be any active server anymore, twice the number of queued connections would be considered as the measured value. This is a fair estimate, as we expect one server to get back soon anyway, but we still prefer to send new traffic to another backend if in better shape. See also the "queue", "be_conn", and "be_sess_rate" sample fetches.

```
be_conn(<backend>]) : integer
```

Applies to the number of currently established connections on the backend, possibly including the connection being evaluated. If no backend name is specified, the current one is used. But it is also possible to check another backend. It can be used to use a specific farm when the nominal one is full. See also the "fe_conn", "queue" and "be_sess_rate" criteria.

```
be_sess_rate(<backends>]) : integer
```

Returns an integer value corresponding to the sessions creation rate on the backend, in number of new sessions per second. This is used with ACLs to switch to an alternate backend when an expensive or fragile one reaches too high a session rate, or to limit abuse of service (eg. prevent sucking of an online dictionary). It can also be useful to add this element to logs using a log-format directive.

Example :

```
# Redirect to an error page if the dictionary is requested too often
backend dynamic
mode http
acl being_scanned be_sess_rate gt 100
redirect location /denied.html if being_scanned
```

```
connslots(<backend>]) : integer
```

Returns an integer value corresponding to the number of connection slots still available in the backend, by totaling the maximum amount of connections on all servers and the maximum queue size. This is probably only used with ACLs.

The basic idea here is to be able to measure the number of connection "slots" still available (connection + queue), so that anything beyond that (intended usage; see "use_backend" keyword) can be redirected to a different backend.

'connslots' = number of available server connection slots, + number of available server queue slots.

Note that while "fe_conn" may be used, "connslots" comes in especially useful when you have a case of traffic going to one single ip, splitting into multiple backends (perhaps using ACLs to do name-based load balancing) and you want to be able to differentiate between different backends, and their available "connslots". Also, whereas "nbsrv" only measures servers that are actually *down*, this fetch is more fine-grained and looks into the number of available connection slots as well. See also "queue" and "avg_queue".

OTHER CAVEATS AND NOTES: at this point in time, the code does not take care of dynamic connections. Also, if any of the server maxconn, or maxqueue is 0, then this fetch clearly does not make sense, in which case the value returned will be -1.

```
date(<offset>]) : integer
```

Returns the current date as the epoch (number of seconds since 01/01/1970). If an offset value is specified, then it is a number of seconds that is added to the current date before returning the value. This is particularly useful to compute relative dates, as both positive and negative offsets are allowed.

It is useful combined with the http_date converter.

Example :

```
# set an expires header to now+1 hour in every response
http-response set-header Expires %[date(3600),http_date]
```

```
env(<name>) : string
```

Returns a string containing the value of environment variable <name>. As a reminder, environment variables are per-process and are sampled when the process starts. This can be useful to pass some information to a next hop server, or with ACLs to take specific action when the process is started a certain way.

Examples :

```
# Pass the Via header to next hop with the local hostname in it
http-request add-header Via 1.1\ %[env(HOSTNAME)]

# reject cookie-less requests when the STOP environment variable is set
http-request deny if !{ cook(SESSIONID) -m found } { env(STOP) -m found }
```

```
fe_conn(<frontend>]) : integer
```

Returns the number of currently established connections on the frontend, possibly including the connection being evaluated. If no frontend name is specified, the current one is used. But it is also possible to check another frontend. It can be used to return a sorry page before hard-blocking, or to use a specific backend to drain new requests when the farm is considered full. This is mostly used with ACLs but can also be used to pass some statistics to servers in HTTP headers. See also the "dst_conn", "be_conn", "fe_sess_rate" fetches.

```
fe_sess_rate(<frontend>]) : integer
```

Returns an integer value corresponding to the sessions creation rate on the frontend, in number of new sessions per second. This is used with ACLs to limit the incoming session rate to an acceptable range in order to prevent abuse of service at the earliest moment, for example when combined with other layer 4 ACLs in order to force the clients to wait a bit for the rate to go down below the limit. It can also be useful to add this element to logs using a log-format directive. See also the "rate-limit sessions" directive for use in frontends.

Example :

```
# This frontend limits incoming mails to 10/s with a max of 100
# concurrent connections. We accept any connection below 10/s, and
# force excess clients to wait for 100 ms. Since clients are limited to
# 100 max, there cannot be more than 10 incoming mails per second.
frontend mail
bind :25
mode tcp
maxconn 100
acl too_fast fe_sess_rate ge 10
tcp-request inspect-delay 100ms
tcp-request content accept if ! too_fast
tcp-request content accept if WAIT_END
```

```
nbproc : integer
```

Returns an integer value corresponding to the number of processes that were started (it equals the global "nbproc" setting). This is useful for logging and debugging purposes.

```
nbsrv(<backend>]) : integer
```

Returns an integer value corresponding to the number of usable servers of either the current backend or the named backend. This is mostly used with ACLs but can also be useful when added to logs. This is normally used to switch to an alternate backend when the number of servers is too low to

```

10271 to handle some load. It is useful to report a failure when combined with
10272 "monitor fail".
10273
10274 proc : integer
10275 Returns an integer value corresponding to the position of the process calling
10276 the function, between 1 and global.nbproc. This is useful for logging and
10277 debugging purposes.
10278
10279 queue(<backends>)) : integer
10280 Returns the total number of queued connections of the designated backend,
10281 including all the connections in server queues. If no backend name is
10282 specified, the current one is used, but it is also possible to check another
10283 one. This is useful with ACLs or to pass statistics to backend servers. This
10284 can be used to take actions when queuing goes above a known level, generally
10285 indicating a surge of traffic or a massive slowdown on the servers. One
10286 possible action could be to reject new users but still accept old ones. See
10287 also the "avg_queue", "be_conn", and "be_sess_rate" fetches.
10288
10289 rand(<range>)) : integer
10290 Returns a random integer value within a range of <range> possible values,
10291 starting at zero. If the range is not specified, it defaults to 2^32, which
10292 gives numbers between 0 and 4294967295. It can be useful to pass some values
10293 needed to take some routing decisions for example, or just for debugging
10294 purposes. This random must not be used for security purposes.
10295
10296 srv_conn(<backends>[/<server>]) : integer
10297 Returns an integer value corresponding to the number of currently established
10298 connections on the designated server, possibly including the connection being
10299 evaluated. If <backends> is omitted, then the server is looked up in the
10300 current backend. It can be used to use a specific farm when one server is
10301 full, or to inform the server about our view of the number of active
10302 connections with it. See also the "fe_conn", "be_conn" and "queue" fetch
10303 methods.
10304
10305 srv_is_up(<backends>[/<server>]) : boolean
10306 Returns true when the designated server is UP, and false when it is either
10307 DOWN or in maintenance mode. If <backends> is omitted, then the server is
10308 looked up in the current backend. It is mainly used to take action based on
10309 an external status reported via a health check (eg: a geographical site's
10310 availability). Another possible use which is more of a hack consists in
10311 using dummy servers as boolean variables that can be enabled or disabled from
10312 the CLI, so that rules depending on those ACLs can be tweaked in realtime.
10313
10314 srv_sess_rate(<backends>[/<server>]) : integer
10315 Returns an integer corresponding to the sessions creation rate on the
10316 designated server, in number of new sessions per second. If <backends> is
10317 omitted, then the server is looked up in the current backend. This is mostly
10318 used with ACLs but can make sense with logs too. This is used to switch to an
10319 alternate backend when an expensive or fragile one reaches too high a session
10320 rate, or to limit abuse of service (eg. prevent latent requests from
10321 overloading servers).
10322
10323 Example :
10324 # Redirect to a separate back
10325 acl srv1_full srv_sess_rate(be1/srv1) gt 50
10326 acl srv2_full srv_sess_rate(be1/srv2) gt 50
10327 use_backend be2 if srv1_full or srv2_full
10328
10329 stopping : boolean
10330 Returns TRUE if the process calling the function is currently stopping. This
10331 can be useful for logging, or for relaxing certain checks or helping close
10332 certain connections upon graceful shutdown.
10333
10334 table_avl(<table>)) : integer
10335 Returns the total number of available entries in the current proxy's

```

```

10336 stick-table or in the designated stick-table. See also table_cnt.
10337
10338 table_cnt(<table>)) : integer
10339 Returns the total number of entries currently in use in the current proxy's
10340 stick-table or in the designated stick-table. See also src_conn_cnt and
10341 table_avl for other entry counting methods.
10342
10343
10344 7.3.3. Fetching samples at Layer 4
10345 -----
10346
10347 The layer 4 usually describes just the transport layer which in haproxy is
10348 closest to the connection, where no content is yet made available. The fetch
10349 methods described here are usable as low as the "tcp-request connection" rule
10350 sets unless they require some future information. Those generally include
10351 TCP/IP addresses and ports, as well as elements from stick-tables related to
10352 the incoming connection. For retrieving a value from a sticky counters, the
10353 counter number can be explicitly set as 0, 1, or 2 using the pre-defined
10354 "sc0", "sc1", or "sc2" prefix, or it can be specified as the first integer
10355 argument when using the "sc_" prefix. An optional table may be specified with
10356 the "sc*" form, in which case the currently tracked key will be looked up into
10357 this alternate table instead of the table currently being tracked.
10358
10359 be_id : integer
10360 Returns an integer containing the current backend's id. It can be used in
10361 frontends with responses to check which backend processed the request.
10362
10363 dst : ip
10364 This is the destination IPv4 address of the connection on the client side,
10365 which is the address the client connected to. It can be useful when running
10366 in transparent mode. It is of type IP and works on both IPv4 and IPv6 tables.
10367 On IPv6 tables, IPv4 address is mapped to its IPv6 equivalent, according to
10368 RFC 4291.
10369
10370 dst_conn : integer
10371 Returns an integer value corresponding to the number of currently established
10372 connections on the same socket including the one being evaluated. It is
10373 normally used with ACLs but can as well be used to pass the information to
10374 servers in an HTTP header or in logs. It can be used to either return a sorry
10375 page before hard-blocking, or to use a specific backend to drain new requests
10376 when the socket is considered saturated. This offers the ability to assign
10377 different limits to different listening ports or addresses. See also the
10378 "fe_conn" and "be_conn" fetches.
10379
10380 dst_port : integer
10381 Returns an integer value corresponding to the destination TCP port of the
10382 connection on the client side, which is the port the client connected to.
10383 This might be used when running in transparent mode, when assigning dynamic
10384 ports to some clients for a whole application session, to stick all users to
10385 a same server, or to pass the destination port information to a server using
10386 an HTTP header.
10387
10388 fe_id : integer
10389 Returns an integer containing the current frontend's id. It can be used in
10390 backends to check from which backend it was called, or to stick all users
10391 coming via a same frontend to the same server.
10392
10393 sc_bytes_in_rate(<ctr>[/<table>]) : integer
10394 sc0_bytes_in_rate(<table>)) : integer
10395 sc1_bytes_in_rate(<table>)) : integer
10396 sc2_bytes_in_rate(<table>)) : integer
10397 Returns the average client-to-server bytes rate from the currently tracked
10398 counters, measured in amount of bytes over the period configured in the
10399 table. See also src_bytes_in_rate.
10400

```



```
10401 sc_bytes_out_rate(<ctr>[,<table>]) : integer
10402 sc0_bytes_out_rate([<table>]) : integer
10403 sc1_bytes_out_rate([<table>]) : integer
10404 sc2_bytes_out_rate([<table>]) : integer
10405 Returns the average server-to-client bytes rate from the currently tracked
10406 counters, measured in amount of bytes over the period configured in the
10407 table. See also src_bytes_out_rate.
10408
10409 sc_clr_gpc0(<ctr>[,<table>]) : integer
10410 sc0_clr_gpc0([<table>]) : integer
10411 sc1_clr_gpc0([<table>]) : integer
10412 sc2_clr_gpc0([<table>]) : integer
10413 Clears the first General Purpose Counter associated to the currently tracked
10414 counters, and returns its previous value. Before the first invocation, the
10415 stored value is zero, so first invocation will always return zero. This is
10416 typically used as a second ACL in an expression in order to mark a connection
10417 when a first ACL was verified :
10418
10419 # block if 5 consecutive requests continue to come faster than 10 sess
10420 # per second, and reset the counter as soon as the traffic slows down.
10421 sc1_clr_gpc0(sc0_http_req_rate gt 10
10422 acl kill sc0_inc_gpc0 gt 5
10423 acl save sc0_clr_gpc0 ge 0
10424 tcp-request connection accept if !abuse save
10425
10426
10427 sc_conn_cnt(<ctr>[,<table>]) : integer
10428 sc0_conn_cnt([<table>]) : integer
10429 sc1_conn_cnt([<table>]) : integer
10430 sc2_conn_cnt([<table>]) : integer
10431 Returns the cumulated number of incoming connections from currently tracked
10432 counters. See also src_conn_cnt.
10433
10434 sc_conn_cur(<ctr>[,<table>]) : integer
10435 sc0_conn_cur([<table>]) : integer
10436 sc1_conn_cur([<table>]) : integer
10437 sc2_conn_cur([<table>]) : integer
10438 Returns the current amount of concurrent connections tracking the same
10439 tracked counters. This number is automatically incremented when tracking
10440 begins and decremented when tracking stops. See also src_conn_cur.
10441
10442 sc_conn_rate(<ctr>[,<table>]) : integer
10443 sc0_conn_rate([<table>]) : integer
10444 sc1_conn_rate([<table>]) : integer
10445 sc2_conn_rate([<table>]) : integer
10446 Returns the average connection rate from the currently tracked counters,
10447 measured in amount of connections over the period configured in the table.
10448 See also src_conn_rate.
10449
10450 sc_get_gpc0(<ctr>[,<table>]) : integer
10451 sc0_get_gpc0([<table>]) : integer
10452 sc1_get_gpc0([<table>]) : integer
10453 sc2_get_gpc0([<table>]) : integer
10454 Returns the value of the first General Purpose Counter associated to the
10455 currently tracked counters. See also src_get_gpc0 and sc/sc0/sc1/sc2_inc_gpc0.
10456
10457 sc_gpc0_rate(<ctr>[,<table>]) : integer
10458 sc0_gpc0_rate([<table>]) : integer
10459 sc1_gpc0_rate([<table>]) : integer
10460 sc2_gpc0_rate([<table>]) : integer
10461 Returns the average increment rate of the first General Purpose Counter
10462 associated to the currently tracked counters. It reports the frequency
10463 which the gpc0 counter was incremented over the configured period. See also
10464 src_gpc0_rate, sc/sc0/sc1/sc2_get_gpc0, and sc/sc0/sc1/sc2_inc_gpc0. Note
10465 that the "gpc0_rate" counter must be stored in the stick-table for a value to
```

```
10466 be returned, as "gpc0" only holds the event count.
10467
10468 sc_http_err_cnt(<ctr>[,<table>]) : integer
10469 sc0_http_err_cnt([<table>]) : integer
10470 sc1_http_err_cnt([<table>]) : integer
10471 sc2_http_err_cnt([<table>]) : integer
10472 Returns the cumulated number of HTTP errors from the currently tracked
10473 counters. This includes the both request errors and 4xx error responses.
10474 See also src_http_err_cnt.
10475
10476 sc_http_err_rate(<ctr>[,<table>]) : integer
10477 sc0_http_err_rate([<table>]) : integer
10478 sc1_http_err_rate([<table>]) : integer
10479 sc2_http_err_rate([<table>]) : integer
10480 Returns the average rate of HTTP errors from the currently tracked counters,
10481 measured in amount of errors over the period configured in the table. This
10482 includes the both request errors and 4xx error responses. See also
10483 src_http_err_rate.
10484
10485 sc_http_req_cnt(<ctr>[,<table>]) : integer
10486 sc0_http_req_cnt([<table>]) : integer
10487 sc1_http_req_cnt([<table>]) : integer
10488 sc2_http_req_cnt([<table>]) : integer
10489 Returns the cumulated number of HTTP requests from the currently tracked
10490 counters. This includes every started request, valid or not. See also
10491 src_http_req_cnt.
10492
10493 sc_http_req_rate(<ctr>[,<table>]) : integer
10494 sc0_http_req_rate([<table>]) : integer
10495 sc1_http_req_rate([<table>]) : integer
10496 sc2_http_req_rate([<table>]) : integer
10497 Returns the average rate of HTTP requests from the currently tracked
10498 counters, measured in amount of requests over the period configured in
10499 the table. This includes every started request, valid or not. See also
10500 src_http_req_rate.
10501
10502 sc_inc_gpc0(<ctr>[,<table>]) : integer
10503 sc0_inc_gpc0([<table>]) : integer
10504 sc1_inc_gpc0([<table>]) : integer
10505 sc2_inc_gpc0([<table>]) : integer
10506 Increments the first General Purpose Counter associated to the currently
10507 tracked counters, and returns its new value. Before the first invocation,
10508 the stored value is zero, so first invocation will increase it to 1 and will
10509 return 1. This is typically used as a second ACL in an expression in order
10510 to mark a connection when a first ACL was verified :
10511
10512 acl abuse sc0_http_req_rate gt 10
10513 acl kill sc0_inc_gpc0 gt 0
10514 tcp-request connection reject if abuse kill
10515
10516 sc_kbytes_in(<ctr>[,<table>]) : integer
10517 sc0_kbytes_in([<table>]) : integer
10518 sc1_kbytes_in([<table>]) : integer
10519 sc2_kbytes_in([<table>]) : integer
10520 Returns the total amount of client-to-server data from the currently tracked
10521 counters, measured in kilobytes. The test is currently performed on 32-bit
10522 integers, which limits values to 4 terabytes. See also src_kbytes_in.
10523
10524 sc_kbytes_out(<ctr>[,<table>]) : integer
10525 sc0_kbytes_out([<table>]) : integer
10526 sc1_kbytes_out([<table>]) : integer
10527 sc2_kbytes_out([<table>]) : integer
10528 Returns the total amount of server-to-client data from the currently tracked
10529 counters, measured in kilobytes. The test is currently performed on 32-bit
10530 integers, which limits values to 4 terabytes. See also src_kbytes_out.
```

```
10531 sc_sess_cnt(<ctr>[,<table>]) : integer
10532 sc0_sess_cnt([<table>]) : integer
10533 sc1_sess_cnt([<table>]) : integer
10534 sc2_sess_cnt([<table>]) : integer
10535 Returns the cumulated number of incoming connections that were transformed
10536 into sessions, which means that they were accepted by a "tcp-request
10537 connection" rule, from the currently tracked counters. A backend may count
10538 more sessions than connections because each connection could result in many
10539 backend sessions if some HTTP keep-alive is performed over the connection
10540 with the client. See also src_sess_cnt.
10541
10542
10543 sc_sess_rate(<ctr>[,<table>]) : integer
10544 sc0_sess_rate([<table>]) : integer
10545 sc1_sess_rate([<table>]) : integer
10546 sc2_sess_rate([<table>]) : integer
10547 Returns the average session rate from the currently tracked counters,
10548 measured in amount of sessions over the period configured in the table. A
10549 session is a connection that got past the early "tcp-request connection"
10550 rules. A backend may count more sessions than connections because each
10551 connection could result in many backend sessions if some HTTP keep-alive is
10552 performed over the connection with the client. See also src_sess_rate.
10553
10554
10555 sc_tracked(<ctr>[,<table>]) : boolean
10556 sc0_tracked([<table>]) : boolean
10557 sc1_tracked([<table>]) : boolean
10558 sc2_tracked([<table>]) : boolean
10559 Returns true if the designated session counter is currently being tracked by
10560 the current session. This can be useful when deciding whether or not we want
10561 to set some values in a header passed to the server.
10562
10563
10564 sc_trackers(<ctr>[,<table>]) : integer
10565 sc0_trackers([<table>]) : integer
10566 sc1_trackers([<table>]) : integer
10567 sc2_trackers([<table>]) : integer
10568 Returns the current amount of concurrent connections tracking the same
10569 tracked counters. This number is automatically incremented when tracking
10570 begins and decremented when tracking stops. It differs from sc0_conn_cur in
10571 that it does not rely on any stored information but on the table's reference
10572 count (the "use" value which is returned by "show table" on the CLI). This
10573 may sometimes be more suited for layer7 tracking. It can be used to tell a
10574 server how many concurrent connections there are from a given address for
10575 example.
10576
10577
10578 so_id : integer
10579 Returns an integer containing the current listening socket's id. It is useful
10580 in frontends involving many "bind" lines, or to stick all users coming via a
10581 same socket to the same server.
10582
10583
10584 src : ip
10585 This is the source IPv4 address of the client of the session. It is of type
10586 IP and works on both IPv4 and IPv6 tables. On IPv6 tables, IPv4 addresses are
10587 mapped to their IPv6 equivalent, according to RFC 4291. Note that it is the
10588 TCP-level source address which is used, and not the address of a client
10589 behind a proxy. However if the "accept-proxy" bind directive is used, it can
10590 be the address of a client behind another PROXY-protocol compatible component
10591 for all rule sets except "tcp-request connection" which sees the real address.
10592
10593
10594 Example:
10595 # add an HTTP header in requests with the originating address' country
10596 http-request set-header X-Country %[src,map_ip(geoip.lst)]
10597
10598
10599 src_bytes_in_rate([<table>]) : integer
10600 Returns the average bytes rate from the incoming connection's source address
10601 in the current proxy's stick-table or in the designated stick-table, measured
```

```
10596 in amount of bytes over the period configured in the table. If the address is
10597 not found, zero is returned. See also sc/sc0/sc1/sc2_bytes_in_rate.
10598
10599
10600 src_bytes_out_rate([<table>]) : integer
10601 Returns the average bytes rate to the incoming connection's source address in
10602 the current proxy's stick-table or in the designated stick-table, measured in
10603 amount of bytes over the period configured in the table. If the address is
10604 not found, zero is returned. See also sc/sc0/sc1/sc2_bytes_out_rate.
10605
10606
10607 src_clr_gpc0([<table>]) : integer
10608 Clears the first General Purpose Counter associated to the incoming
10609 connection's source address in the current proxy's stick-table or in the
10610 designated stick-table, and returns its previous value. If the address is not
10611 found, an entry is created and 0 is returned. This is typically used as a
10612 second ACL in an expression in order to mark a connection when a first ACL
10613 was verified :
10614
10615 # block if 5 consecutive requests continue to come faster than 10 sess
10616 # per second, and reset the counter as soon as the traffic slows down.
10617 acl abuse src_http_req_rate gt 10
10618 acl kill src_inc_gpc0 gt 5
10619 acl save src_clr_gpc0 ge 0
10620 tcp-request connection accept if !abuse save
10621 tcp-request connection reject if abuse kill
10622
10623
10624 src_conn_cnt([<table>]) : integer
10625 Returns the cumulated number of connections initiated from the current
10626 incoming connection's source address in the current proxy's stick-table or in
10627 the designated stick-table. If the address is not found, zero is returned.
10628 See also sc/sc0/sc1/sc2_conn_cnt.
10629
10630
10631 src_conn_cur([<table>]) : integer
10632 Returns the current amount of concurrent connections initiated from the
10633 current incoming connection's source address in the current proxy's
10634 stick-table or in the designated stick-table. If the address is not found,
10635 zero is returned. See also sc/sc0/sc1/sc2_conn_cur.
10636
10637
10638 src_conn_rate([<table>]) : integer
10639 Returns the average connection rate from the incoming connection's source
10640 address in the current proxy's stick-table or in the designated stick-table,
10641 measured in amount of connections over the period configured in the table. If
10642 the address is not found, zero is returned. See also sc/sc0/sc1/sc2_conn_rate.
10643
10644
10645 src_get_gpc0([<table>]) : integer
10646 Returns the value of the first General Purpose Counter associated to the
10647 incoming connection's source address in the current proxy's stick-table or in
10648 the designated stick-table. If the address is not found, zero is returned.
10649 See also sc/sc0/sc1/sc2_get_gpc0 and src_inc_gpc0.
10650
10651
10652 src_gpc0_rate([<table>]) : integer
10653 Returns the average increment rate of the first General Purpose Counter
10654 associated to the incoming connection's source address in the current proxy's
10655 stick-table or in the designated stick-table. It reports the frequency
10656 which the gpc0 counter was incremented over the configured period. See also
10657 sc/sc0/sc1/sc2_gpc0_rate, src_get_gpc0, and sc/sc0/sc1/sc2_inc_gpc0. Note
10658 that the "gpc0_rate" counter must be stored in the stick-table for a value to
10659 be returned, as "gpc0" only holds the event count.
10660
10661
10662 src_http_err_cnt([<table>]) : integer
10663 Returns the cumulated number of HTTP errors from the incoming connection's
10664 source address in the current proxy's stick-table or in the designated
10665 stick-table. This includes the both request errors and 4xx error responses.
10666 See also sc/sc0/sc1/sc2_http_err_cnt. If the address is not found, zero is
10667 returned.
10668
10669
10670
```

```
10661 src_http_err_rate(<table>)) : integer
10662 Returns the average rate of HTTP errors from the incoming connection's source
10663 address in the current proxy's stick-table or in the designated stick-table,
10664 measured in amount of errors over the period configured in the table. This
10665 includes the both request errors and 4xx error responses. If the address is
10666 not found, zero is returned. See also sc/sc0/sc1/sc2_http_err_rate.
10667
10668 src_http_req_cnt(<table>)) : integer
10669 Returns the cumulated number of HTTP requests from the incoming connection's
10670 source address in the current proxy's stick-table or in the designated stick-
10671 table. This includes every started request, valid or not. If the address is
10672 not found, zero is returned. See also sc/sc0/sc1/sc2_http_req_cnt.
10673
10674 src_http_req_rate(<table>)) : integer
10675 Returns the average rate of HTTP requests from the incoming connection's
10676 source address in the current proxy's stick-table or in the designated stick-
10677 table, measured in amount of requests over the period configured in the
10678 table. This includes every started request, valid or not. If the address is
10679 not found, zero is returned. See also sc/sc0/sc1/sc2_http_req_rate.
10680
10681 src_inc_gpc0(<table>)) : integer
10682 Increments the first General Purpose Counter associated to the incoming
10683 connection's source address in the current proxy's stick-table or in the
10684 designated stick-table, and returns its new value. If the address is not
10685 found, an entry is created and 1 is returned. See also sc0/sc2/inc_gpc0.
10686 This is typically used as a second ACL in an expression in order to mark a
10687 connection when a first ACL was verified :
10688
10689 acl abuse src_http_req_rate gt 10
10690 acl kill src_inc_gpc0 gt 0
10691 tcp-request connection reject if abuse kill
10692
10693 src_kbytes_in(<table>)) : integer
10694 Returns the total amount of data received from the incoming connection's
10695 source address in the current proxy's stick-table or in the designated
10696 stick-table, measured in kilobytes. If the address is not found, zero is
10697 returned. The test is currently performed on 32-bit integers, which limits
10698 values to 4 terabytes. See also sc/sc0/sc1/sc2_kbytes_in.
10699
10700 src_kbytes_out(<table>)) : integer
10701 Returns the total amount of data sent to the incoming connection's source
10702 address in the current proxy's stick-table or in the designated stick-table,
10703 measured in kilobytes. If the address is not found, zero is returned. The
10704 test is currently performed on 32-bit integers, which limits values to 4
10705 terabytes. See also sc/sc0/sc1/sc2_kbytes_out.
10706
10707 src_port : integer
10708 Returns an integer value corresponding to the TCP source port of the
10709 connection on the client side, which is the port the client connected from.
10710 Usage of this function is very limited as modern protocols do not care much
10711 about source ports nowadays.
10712
10713 src_sess_cnt(<table>)) : integer
10714 Returns the cumulated number of connections initiated from the incoming
10715 connection's source IPv4 address in the current proxy's stick-table or in the
10716 designated stick-table, that were transformed into sessions, which means that
10717 they were accepted by "tcp-request" rules. If the address is not found, zero
10718 is returned. See also sc/sc0/sc1/sc2_sess_cnt.
10719
10720 src_sess_rate(<table>)) : integer
10721 Returns the average session rate from the incoming connection's source
10722 address in the current proxy's stick-table or in the designated stick-table,
10723 measured in amount of sessions over the period configured in the table. A
10724 session is a connection that went past the early "tcp-request" rules. If the
10725 address is not found, zero is returned. See also sc/sc0/sc1/sc2_sess_rate.
```

```
10726 src_updt_conn_cnt(<table>)) : integer
10727 Creates or updates the entry associated to the incoming connection's source
10728 address in the current proxy's stick-table or in the designated stick-table.
10729 This table must be configured to store the "conn_cnt" data type, otherwise
10730 the match will be ignored. The current count is incremented by one, and the
10731 expiration timer refreshed. The updated count is returned, so this match
10732 can't return zero. This was used to reject service abusers based on their
10733 source address. Note: it is recommended to use the more complete "track-sc*"
10734 actions in "tcp-request" rules instead.
10735
10736 Example :
10737 # This frontend limits incoming SSH connections to 3 per 10 second for
10738 # each source address, and rejects excess connections until a 10 second
10739 # silence is observed. At most 20 addresses are tracked.
10740 listen ssh
10741 bind :22
10742 mode tcp
10743 maxconn 100
10744 stick-table type ip size 20 expire 10s store conn_cnt
10745 tcp-request content reject if { src_updt_conn_cnt gt 3 }
10746 server local 127.0.0.1:22
10747
10748 srv_id : integer
10749 Returns an integer containing the server's id when processing the response.
10750 While it's almost only used with ACLs, it may be used for logging or
10751 debugging.
10752
10753 7.3.4. Fetching samples at Layer 5
10754 -----
10755
10756 The layer 5 usually describes just the session layer which in haproxy is
10757 closest to the session once all the connection handshakes are finished, but
10758 when no content is yet made available. The fetch methods described here are
10759 usable as low as the "tcp-request content" rule sets unless they require some
10760 future information. Those generally include the results of SSL negotiations.
10761
10762 ssl_bc : boolean
10763 Returns true when the back connection was made via an SSL/TLS transport
10764 layer and is locally deciphered. This means the outgoing connection was made
10765 other a server with the "ssl" option.
10766
10767 ssl_bc_alg_keysize : integer
10768 Returns the symmetric cipher key size supported in bits when the outgoing
10769 connection was made over an SSL/TLS transport layer.
10770
10771 ssl_bc_cipher : string
10772 Returns the name of the used cipher when the outgoing connection was made
10773 over an SSL/TLS transport layer.
10774
10775 ssl_bc_protocol : string
10776 Returns the name of the used protocol when the outgoing connection was made
10777 over an SSL/TLS transport layer.
10778
10779 ssl_bc_unique_id : binary
10780 When the outgoing connection was made over an SSL/TLS transport layer,
10781 returns the TLS unique ID as defined in RFC5929 section 3. The unique ID
10782 can be encoded to base64 using the converter: "ssl_bc_unique_id,base64".
10783
10784 ssl_bc_session_id : binary
10785 Returns the SSL ID of the back connection when the outgoing connection was
10786 made over an SSL/TLS transport layer. It is useful to log if we want to know
10787 if session was reused or not.
10788
10789
```

10791 **ssl_bc_use_keysize** : integer
10792 Returns the symmetric cipher key size used in bits when the outgoing
10793 connection was made over an SSL/TLS transport layer.
10794
10795 **ssl_c_ca_err** : integer
10796 When the incoming connection was made over an SSL/TLS transport layer,
10797 returns the ID of the first error detected during verification of the client
10798 certificate at depth > 0, or 0 if no error was encountered during this
10799 verification process. Please refer to your SSL library's documentation to
10800 find the exhaustive list of error codes.
10801
10802 **ssl_c_ca_err_depth** : integer
10803 When the incoming connection was made over an SSL/TLS transport layer,
10804 returns the depth in the CA chain of the first error detected during the
10805 verification of the client certificate. If no error is encountered, 0 is
10806 returned.
10807
10808 **ssl_c_der** : binary
10809 Returns the DER formatted certificate presented by the client when the
10810 incoming connection was made over an SSL/TLS transport layer. When used for
10811 an ACL, the value(s) to match against can be passed in hexadecimal form.
10812
10813 **ssl_c_err** : integer
10814 When the incoming connection was made over an SSL/TLS transport layer,
10815 returns the ID of the first error detected during verification at depth 0, or
10816 0 if no error was encountered during this verification process. Please refer
10817 to your SSL library's documentation to find the exhaustive list of error
10818 codes.
10819
10820 **ssl_c_i_dn(<entry>[,<occ>]])** : string
10821 When the incoming connection was made over an SSL/TLS transport layer,
10822 returns the full distinguished name of the issuer of the certificate
10823 presented by the client when no <entry> is specified, or the value of the
10824 first given entry found from the beginning of the DN. If a positive/negative
10825 occurrence number is specified as the optional second argument, it returns
10826 the value of the nth given entry value from the beginning/end of the DN.
10827 For instance, "ssl_c_i_dn(OU,2)" the second organization unit, and
10828 "ssl_c_i_dn(CN)" retrieves the common name.
10829
10830 **ssl_c_key_alg** : string
10831 Returns the name of the algorithm used to generate the key of the certificate
10832 presented by the client when the incoming connection was made over an SSL/TLS
10833 transport layer.
10834
10835 **ssl_c_notafter** : string
10836 Returns the end date presented by the client as a formatted string
10837 YYYMMDDhhmmss[Z] when the incoming connection was made over an SSL/TLS
10838 transport layer.
10839
10840 **ssl_c_notbefore** : string
10841 Returns the start date presented by the client as a formatted string
10842 YYYMMDDhhmmss[Z] when the incoming connection was made over an SSL/TLS
10843 transport layer.
10844
10845 **ssl_c_s_dn(<entry>[,<occ>]])** : string
10846 When the incoming connection was made over an SSL/TLS transport layer,
10847 returns the full distinguished name of the subject of the certificate
10848 presented by the client when no <entry> is specified, or the value of the
10849 first given entry found from the beginning of the DN. If a positive/negative
10850 occurrence number is specified as the optional second argument, it returns
10851 the value of the nth given entry value from the beginning/end of the DN.
10852 For instance, "ssl_c_s_dn(OU,2)" the second organization unit, and
10853 "ssl_c_s_dn(CN)" retrieves the common name.
10854
10855 **ssl_c_serial** : binary

10856 Returns the serial of the certificate presented by the client when the
10857 incoming connection was made over an SSL/TLS transport layer. When used for
10858 an ACL, the value(s) to match against can be passed in hexadecimal form.
10859
10860 **ssl_c_shal** : binary
10861 Returns the SHA-1 fingerprint of the certificate presented by the client when
10862 the incoming connection was made over an SSL/TLS transport layer. This can be
10863 used to stick a client to a server, or to pass this information to a server.
10864 Note that the output is binary, so if you want to pass that signature to the
10865 server, you need to encode it in hex or base64, such as in the example below:
10866
10867 http-request set-header X-SSL-Client-SHA1 %[ssl_c_shal.hex]
10868
10869 **ssl_c_sig_alg** : string
10870 Returns the name of the algorithm used to sign the certificate presented by
10871 the client when the incoming connection was made over an SSL/TLS transport
10872 layer.
10873
10874 **ssl_c_used** : boolean
10875 Returns true if current SSL session uses a client certificate even if current
10876 connection uses SSL session resumption. See also "ssl_fc_has_crt".
10877
10878 **ssl_c_verify** : integer
10879 Returns the verify result error ID when the incoming connection was made over
10880 an SSL/TLS transport layer, otherwise zero if no error is encountered. Please
10881 refer to your SSL library's documentation for an exhaustive list of error
10882 codes.
10883
10884 **ssl_c_version** : integer
10885 Returns the version of the certificate presented by the client when the
10886 incoming connection was made over an SSL/TLS transport layer.
10887
10888 **ssl_f_der** : binary
10889 Returns the DER formatted certificate presented by the frontend when the
10890 incoming connection was made over an SSL/TLS transport layer. When used for
10891 an ACL, the value(s) to match against can be passed in hexadecimal form.
10892
10893 **ssl_f_i_dn(<entry>[,<occ>]])** : string
10894 When the incoming connection was made over an SSL/TLS transport layer,
10895 returns the full distinguished name of the issuer of the certificate
10896 presented by the frontend when no <entry> is specified, or the value of the
10897 first given entry found from the beginning of the DN. If a positive/negative
10898 occurrence number is specified as the optional second argument, it returns
10899 the value of the nth given entry value from the beginning/end of the DN.
10900 For instance, "ssl_f_i_dn(OU,2)" the second organization unit, and
10901 "ssl_f_i_dn(CN)" retrieves the common name.
10902
10903 **ssl_f_key_alg** : string
10904 Returns the name of the algorithm used to generate the key of the certificate
10905 presented by the frontend when the incoming connection was made over an
10906 SSL/TLS transport layer.
10907
10908 **ssl_f_notafter** : string
10909 Returns the end date presented by the frontend as a formatted string
10910 YYYMMDDhhmmss[Z] when the incoming connection was made over an SSL/TLS
10911 transport layer.
10912
10913 **ssl_f_notbefore** : string
10914 Returns the start date presented by the frontend as a formatted string
10915 YYYMMDDhhmmss[Z] when the incoming connection was made over an SSL/TLS
10916 transport layer.
10917
10918 **ssl_f_s_dn(<entry>[,<occ>]])** : string
10919 When the incoming connection was made over an SSL/TLS transport layer,
10920 returns the full distinguished name of the subject of the certificate

presented by the frontend when no <entry> is specified, or the value of the first given entry found from the beginning of the DN. If a positive/negative occurrence number is specified as the optional second argument, it returns the value of the nth given entry value from the beginning/end of the DN. For instance, "ssl_f_s_dn(0U,2)" the second organization unit, and "ssl_f_s_dn(CN)" retrieves the common name.

ssl_f_serial : binary

Returns the serial of the certificate presented by the frontend when the incoming connection was made over an SSL/TLS transport layer. When used for an ACL, the value(s) to match against can be passed in hexadecimal form.

ssl_f_sha1 : binary

Returns the SHA-1 fingerprint of the certificate presented by the frontend when the incoming connection was made over an SSL/TLS transport layer. This can be used to know which certificate was chosen using SNI.

ssl_f_sig_alg : string

Returns the name of the algorithm used to sign the certificate presented by the frontend when the incoming connection was made over an SSL/TLS transport layer.

ssl_f_version : integer

Returns the version of the certificate presented by the frontend when the incoming connection was made over an SSL/TLS transport layer.

ssl_fc : boolean

Returns true when the front connection was made via an SSL/TLS transport layer and is locally deciphered. This means it has matched a socket declared with a "bind" line having the "ssl" option.

Example :

```
# This passes "X-Proto: https" to servers when client connects over SSL
listen http-https
    bind :80
    bind :443 ssl crt /etc/haproxy.pem
    http-request add-header X-Proto https if { ssl_fc }
```

ssl_fc_alg_keysize : integer

Returns the symmetric cipher key size supported in bits when the incoming connection was made over an SSL/TLS transport layer.

ssl_fc_alpn : string

This extracts the Application Layer Protocol Negotiation field from an incoming connection made via a TLS transport layer and locally deciphered by haproxy. The result is a string containing the protocol name advertised by the client. The SSL library must have been built with support for TLS extensions enabled (check haproxy -vv). Note that the TLS ALPN extension is not advertised unless the "alpn" keyword on the "bind" line specifies a protocol list. Also, nothing forces the client to pick a protocol from this list, any other one may be requested. The TLS ALPN extension is meant to replace the TLS NPN extension. See also "ssl_fc_npn".

ssl_fc_cipher : string

Returns the name of the used cipher when the incoming connection was made over an SSL/TLS transport layer.

ssl_fc_has_crt : boolean

Returns true if a client certificate is present in an incoming connection over SSL/TLS transport layer. Useful if 'verify' statement is set to 'optional'. Note: on SSL session resumption with Session ID or TLS ticket, client certificate is not present in the current connection but may be retrieved from the cache or the ticket. So prefer "ssl_c_used" if you want to check if current SSL session uses a client certificate.

ssl_fc_has_sni : boolean

This checks for the presence of a Server Name Indication TLS extension (SNI) in an incoming connection was made over an SSL/TLS transport layer. Returns true when the incoming connection presents a TLS SNI field. This requires that the SSL library is build with support for TLS extensions enabled (check haproxy -vv).

ssl_fc_npn : string

This extracts the Next Protocol Negotiation field from an incoming connection made via a TLS transport layer and locally deciphered by haproxy. The result is a string containing the protocol name advertised by the client. The SSL library must have been built with support for TLS extensions enabled (check haproxy -vv). Note that the TLS NPN extension is not advertised unless the "npn" keyword on the "bind" line specifies a protocol list. Also, nothing forces the client to pick a protocol from this list, any other one may be requested. Please note that the TLS NPN extension was replaced with ALPN.

ssl_fc_protocol : string

Returns the name of the used protocol when the incoming connection was made over an SSL/TLS transport layer.

ssl_fc_unique_id : binary

When the incoming connection was made over an SSL/TLS transport layer, returns the TLS unique ID as defined in RFC5929 section 3. The unique id can be encoded to base64 using the converter: "ssl_bc_unique_id,base64".

ssl_fc_session_id : binary

Returns the SSL ID of the front connection when the incoming connection was made over an SSL/TLS transport layer. It is useful to stick a given client to a server. It is important to note that some browsers refresh their session ID every few minutes.

ssl_fc_sni : string

This extracts the Server Name Indication TLS extension (SNI) field from an incoming connection made via an SSL/TLS transport layer and locally deciphered by haproxy. The result (when present) typically is a string matching the HTTPS host name (253 chars or less). The SSL library must have been built with support for TLS extensions enabled (check haproxy -vv).

This fetch is different from "req_ssl_sni" above in that it applies to the connection being deciphered by haproxy and not to SSL contents being blindly forwarded. See also "ssl_fc_sni_end" and "ssl_fc_sni_reg" below. This requires that the SSL library is build with support for TLS extensions enabled (check haproxy -vv).

ACL derivatives :

ssl_fc_sni_end : suffix match

ssl_fc_sni_reg : regex match

ssl_fc_use_keysize : integer

Returns the symmetric cipher key size used in bits when the incoming connection was made over an SSL/TLS transport layer.

7.3.5. Fetching samples from buffer contents (Layer 6)

Fetching samples from buffer contents is a bit different from the previous sample fetches above because the sampled data are ephemeral. These data can only be used when they're available and will be lost when they're forwarded. For this reason, samples fetched from buffer contents during a request cannot be used in a response for example. Even while the data are being fetched, they can change. Sometimes it is necessary to set some delays or combine multiple sample fetch methods to ensure that the expected data are complete and usable, for example through TCP request content inspection. Please see the "tcp-request

content" keyword for more detailed information on the subject.

payload(<offset>,<length>) : binary (deprecated)

This is an alias for "req.payload" when used in the context of a request (eg: "stick on", "stick match"), and for "res.payload" when used in the context of a response such as in "stick store response".

payload_lv(<offset>,<length>[,<offset2>]) : binary (deprecated)

This is an alias for "req.payload_lv" when used in the context of a request (eg: "stick on", "stick match"), and for "res.payload_lv" when used in the context of a response such as in "stick store response".

req.len : integer

req.len : integer (deprecated)

Returns an integer value corresponding to the number of bytes present in the request buffer. This is mostly used in ACL. It is important to understand that this test does not return false as long as the buffer is changing. This means that a check with equality to zero will almost always immediately match at the beginning of the session, while a test for more data will wait for that data to come in and return false only when haproxy is certain that no more data will come in. This test was designed to be used with TCP request content inspection.

req.payload(<offset>,<length>) : binary

This extracts a binary block of <length> bytes and starting at byte <offset> in the request buffer. As a special case, if the <length> argument is zero, the whole buffer from <offset> to the end is extracted. This can be used with ACLs in order to check for the presence of some content in a buffer at any location.

ACL alternatives :

payload(<offset>,<length>) : hex binary match

req.payload_lv(<offset>,<length>[,<offset2>]) : binary

This extracts a binary block whose size is specified at <offset> for <length> bytes, and which starts at <offset2> if specified or just after the length in the request buffer. The <offset2> parameter also supports relative offsets if prepended with a '+' or '-' sign.

ACL alternatives :

payload_lv(<offset>,<length>[,<offset2>]) : hex binary match

Example : please consult the example from the "stick store-response" keyword.

req.proto_http : boolean

Returns true when data in the request buffer look like HTTP and correctly parses as such. It is the same parser as the common HTTP request parser which is used so there should be no surprises. The test does not match until the request is complete, failed or timed out. This test may be used to report the protocol in TCP logs, but the biggest use is to block TCP request analysis until a complete HTTP request is present in the buffer, for example to track a header.

Example:

```
# track request counts per "base" (concatenation of Host+URL)
tcp-request inspect-delay 10s
tcp-request content reject if !HTTP
tcp-request content track-sc0 base table req-rate
```

req.rdp_cookie(<lname>]) : string

rdp_cookie(<lname>]) : string (deprecated)

When the request buffer looks like the RDP protocol, extracts the RDP cookie <lname>, or any cookie if unspecified. The parser only checks for the first cookie, as illustrated in the RDP protocol specification. The cookie name is

case insensitive. Generally the "MSTS" cookie name will be used, as it can contain the user name of the client connecting to the server if properly configured on the client. The "MSTHASH" cookie is often used as well for session stickiness to servers.

This differs from "balance rdp-cookie" in that any balancing algorithm may be used and thus the distribution of clients to backend servers is not linked to a hash of the RDP cookie. It is envisaged that using a balancing algorithm such as "balance roundrobin" or "balance leastconn" will lead to a more even distribution of clients to backend servers than the hash used by "balance rdp-cookie".

ACL derivatives :

req_rdp_cookie(<lname>]) : exact string match

Example :

```
listen tse-farm
bind 0.0.0.0:3389
# wait up to 5s for an RDP cookie in the request
tcp-request inspect-delay 5s
tcp-request content accept if RDP_COOKIE
# apply RDP cookie persistence
persist rdp-cookie
# Persist based on the msthash cookie
# This is only useful makes sense if
# balance rdp-cookie is not used
stick-table type string size 204800
stick on req_rdp_cookie(msthash)
server srv1 1.1.1.1:3389
server srv1 1.1.1.2:3389
```

See also : "balance rdp-cookie", "persist rdp-cookie", "tcp-request" and the "req_rdp_cookie" ACL.

req_rdp_cookie_cnt(<lname>) : integer

rdp_cookie_cnt(<lname>) : integer (deprecated)

Tries to parse the request buffer as RDP protocol, then returns an integer corresponding to the number of RDP cookies found. If an optional cookie name is passed, only cookies matching this name are considered. This is mostly used in ACL.

ACL derivatives :

req_rdp_cookie_cnt(<lname>]) : integer match

req_ssl_hello_type : integer

req_ssl_hello_type : integer (deprecated)

Returns an integer value containing the type of the SSL hello message found in the request buffer if the buffer contains data that parse as a complete SSL (v3 or superior) client hello message. Note that this only applies to raw contents found in the request buffer and not to contents deciphered via an SSL data layer, so this will not work with "bind" lines having the "ssl" option. This is mostly used in ACL to detect presence of an SSL hello message that is supposed to contain an SSL session ID usable for stickiness.

req_ssl_sni : string

req_ssl_sni : string (deprecated)

Returns a string containing the value of the Server Name TLS extension sent by a client in a TLS stream passing through the request buffer if the buffer contains data that parse as a complete SSL (v3 or superior) client hello message. Note that this only applies to raw contents found in the request buffer and not to contents deciphered via an SSL data layer, so this will not work with "bind" lines having the "ssl" option. SNI normally contains the name of the host the client tries to connect to (for recent browsers). SNI is useful for allowing or denying access to certain hosts when SSL/TLS is used by the client. This test was designed to be used with TCP request content

```
11181 inspection. If content switching is needed, it is recommended to first wait
11182 for a complete client hello (type 1), like in the example below. See also
11183 "ssl_fc_sni".
11184
11185 ACL derivatives :
11186   req_ssl_sni : exact string match
11187
11188 Examples :
11189   # Wait for a client hello for at most 5 seconds
11190   tcp-request inspect-delay 5s
11191   tcp-request content accept if { req_ssl_hello_type 1 }
11192   use backend bk allow if { req_ssl_sni -f allowed_sites }
11193   default_backend bk_sorry_page
11194
11195 res_ssl_hello_type : integer
11196 rep_ssl_hello_type : integer (deprecated)
11197 Returns an integer value containing the type of the SSL hello message found
11198 in the response buffer if the buffer contains data that parses as a complete
11199 SSL (v3 or superior) hello message. Note that this only applies to raw
11200 contents found in the response buffer and not to contents deciphered via an
11201 SSL data layer, so this will not work with "server" lines having the "ssl"
11202 option. This is mostly used in ACL to detect presence of an SSL hello message
11203 that is supposed to contain an SSL session ID usable for stickiness.
11204
11205 req_ssl_ver : integer
11206 rep_ssl_ver : integer (deprecated)
11207 Returns an integer value containing the version of the SSL/TLS protocol of a
11208 stream present in the request buffer. Both SSLv2 hello messages and SSLv3
11209 messages are supported. TLSv1 is announced as SSL version 3.1. The value is
11210 composed of the major version multiplied by 65536, added to the minor
11211 version. Note that this only applies to raw contents found in the request
11212 buffer and not to contents deciphered via an SSL data layer, so this will not
11213 work with "bind" lines having the "ssl" option. The ACL version of the test
11214 matches against a decimal notation in the form MAJOR.MINOR (eg: 3.1). This
11215 fetch is mostly used in ACL.
11216
11217 ACL derivatives :
11218   req_ssl_ver : decimal match
11219
11220 res.len : integer
11221 Returns an integer value corresponding to the number of bytes present in the
11222 response buffer. This is mostly used in ACL. It is important to understand
11223 that this test does not return false as long as the buffer is changing. This
11224 means that a check with equality to zero will almost always immediately match
11225 at the beginning of the session, while a test for more data will wait for
11226 that data to come in and return false only when haproxy is certain that no
11227 more data will come in. This test was designed to be used with TCP response
11228 content inspection.
11229
11230 res.payload(<offset>,<length>) : binary
11231 This extracts a binary block of <length> bytes and starting at byte <offset>
11232 in the response buffer. As a special case, if the <length> argument is zero,
11233 the whole buffer from <offset> to the end is extracted. This can be used
11234 with ACLs in order to check for the presence of some content in a buffer at
11235 any location.
11236
11237 res.payload_lv(<offset1>,<length>[,<offset2>]) : binary
11238 This extracts a binary block whose size is specified at <offset1> for <length>
11239 bytes, and which starts at <offset2> if specified or just after the length in
11240 the response buffer. The <offset2> parameter also supports relative offsets
11241 if prepended with a '+' or '-' sign.
11242
11243 Example : please consult the example from the "stick store-response" keyword.
11244
11245 wait_end : boolean
```

```
11246 This fetch either returns true when the inspection period is over, or does
11247 not fetch. It is only used in ACLs, in conjunction with content analysis to
11248 avoid returning a wrong verdict early. It may also be used to delay some
11249 actions, such as a delayed reject for some special addresses. Since it either
11250 stops the rules evaluation or immediately returns true, it is recommended to
11251 use this acl as the last one in a rule. Please note that the default ACL
11252 "WAIT_END" is always usable without prior declaration. This test was designed
11253 to be used with TCP request content inspection.
11254
11255 Examples :
11256   # delay every incoming request by 2 seconds
11257   tcp-request inspect-delay 2s
11258   tcp-request content accept if WAIT_END
11259
11260   # don't immediately tell bad guys they are rejected
11261   tcp-request inspect-delay 10s
11262   acl goodguys src 10.0.0.0/24
11263   acl badguys src 10.0.1.0/24
11264   tcp-request content accept if goodguys
11265   tcp-request content reject if badguys WAIT_END
11266   tcp-request content reject
11267
11268 7.3.6. Fetching HTTP samples (Layer 7)
11269 -----
11270
11271 It is possible to fetch samples from HTTP contents, requests and responses.
11272 This application layer is also called layer 7. It is only possible to fetch the
11273 data in this section when a full HTTP request or response has been parsed from
11274 its respective request or response buffer. This is always the case with all
11275 HTTP specific rules and for sections running with "mode http". When using TCP
11276 content inspection, it may be necessary to support an inspection delay in order
11277 to let the request or response come in first. These fetches may require a bit
11278 more CPU resources than the layer 4 ones, but not much since the request and
11279 response are indexed.
11280
11281 base : string
11282 This returns the concatenation of the first Host header and the path part of
11283 the request, which starts at the first slash and ends before the question
11284 mark. It can be useful in virtual hosted environments to detect URL abuses as
11285 well as to improve shared caches efficiency. Using this with a limited size
11286 stick table also allows one to collect statistics about most commonly
11287 requested objects by host/path. With ACLs it can allow simple content
11288 switching rules involving the host and the path at the same time, such as
11289 "www.example.com/favicon.ico". See also "path" and "url".
11290
11291 ACL derivatives :
11292   base : exact string match
11293   base_beg : prefix match
11294   base_dir : subdir match
11295   base_dom : domain match
11296   base_end : suffix match
11297   base_len : length match
11298   base_reg : regex match
11299   base_sub : substring match
11300
11301 base32 : integer
11302 This returns a 32-bit hash of the value returned by the "base" fetch method
11303 above. This is useful to track per-URL activity on high traffic sites without
11304 having to store all URLs. Instead a shorter hash is stored, saving a lot of
11305 memory. The output type is an unsigned integer.
11306
11307 base32+src : binary
11308 This returns the concatenation of the base32 fetch above and the src fetch
11309 below. The resulting type is of type binary, with a size of 8 or 20 bytes
11310
```

depending on the source address family. This can be used to track per-IP, per-URL counters.

`capture.req.hdr(<idx>) : string`
 This extracts the content of the header captured by the "capture request header", `idx` is the position of the capture keyword in the configuration. The first entry is an index of 0. See also: "capture request header".

`capture.req.method : string`
 This extracts the METHOD of an HTTP request. It can be used in both request and response. Unlike "method", it can be used in both request and response because it's allocated.

`capture.req.uri : string`
 This extracts the request's URI, which starts at the first slash and ends before the first space in the request (without the host part). Unlike "path" and "url", it can be used in both request and response because it's allocated.

`capture.req.ver : string`
 This extracts the request's HTTP version and returns either "HTTP/1.0" or "HTTP/1.1". Unlike "req.ver", it can be used in both request, response, and logs because it relies on a persistent flag.

`capture.res.hdr(<idx>) : string`
 This extracts the content of the header captured by the "capture response header", `idx` is the position of the capture keyword in the configuration. The first entry is an index of 0.

See also: "capture response header"

`capture.res.ver : string`
 This extracts the response's HTTP version and returns either "HTTP/1.0" or "HTTP/1.1". Unlike "res.ver", it can be used in logs because it relies on a persistent flag.

`req.cookie(<[<name>]) : string`
 Cook(<[<name>]) : string (deprecated)
 This extracts the last occurrence of the cookie name <name> on a "Cookie" header line from the request, and returns its value as string. If no name is specified, the first cookie value is returned. When used with ACLs, all matching cookies are evaluated. Spaces around the name and the value are ignored as requested by the Cookie header specification (RFC6265). The cookie name is case-sensitive. Empty cookies are valid, so an empty cookie may very well return an empty value if it is present. Use the "found" match to detect presence. Use the `res.cookie()` variant for response cookies sent by the server.

ACL derivatives :
`cook(<[<name>])` : exact string match
`cook_beg(<[<name>])` : prefix match
`cook_dir(<[<name>])` : subdir match
`cook_dom(<[<name>])` : domain match
`cook_end(<[<name>])` : suffix match
`cook_len(<[<name>])` : length match
`cook_reg(<[<name>])` : regex match
`cook_sub(<[<name>])` : substring match

`req.cookie_cnt(<[<name>]) : integer`
`cook_val(<[<name>]) : integer (deprecated)`
 Returns an integer value representing the number of occurrences of the cookie <name> in the request, or all cookies if <name> is not specified.

`req.cookie_val(<[<name>]) : integer`
`cook_val(<[<name>]) : integer (deprecated)`
 This extracts the last occurrence of the cookie name <name> on a "Cookie" header line from the request, and converts its value to an integer which is

returned. If no name is specified, the first cookie value is returned. When used in ACLs, all matching names are iterated over until a value matches.

`cookie(<[<name>]) : string (deprecated)`
 This extracts the last occurrence of the cookie name <name> on a "Cookie" header line from the request, or a "Set-Cookie" header from the response, and returns its value as a string. A typical use is to get multiple clients sharing a same profile use the same server. This can be similar to what "appsession" does with the "request-learn" statement, but with support for multi-peer synchronization and state keeping across restarts. If no name is specified, the first cookie value is returned. This fetch should not be used anymore and should be replaced by `req.cookie()` or `res.cookie()` instead as it ambiguously uses the direction based on the context where it is used. See also : "appsession".

`hdr(<[<name>],<occ>]) : string`
 This is equivalent to `req.hdr()` when used on requests, and to `res.hdr()` when used on responses. Please refer to these respective fetches for more details. In case of doubt about the fetch direction, please use the explicit ones. Note that contrary to the `hdr()` sample fetch method, the `hdr_*` ACL keywords unambiguously apply to the request headers.

`req.fhdr(<name>[,<occ>]) : string`
 This extracts the last occurrence of header <name> in an HTTP request. When used from an ACL, all occurrences are iterated over until a match is found. Optionally, a specific occurrence might be specified as a position number. Positive values indicate a position from the first occurrence, with 1 being the first one. Negative values indicate positions relative to the last one, with -1 being the last one. It differs from `req.hdr()` in that any commas present in the value are returned and are not used as delimiters. This is sometimes useful with headers such as User-Agent.

`req.hdr_cnt(<[<name>]) : integer`
 Returns an integer value representing the number of occurrences of request header field name <name>, or the total number of header fields if <name> is not specified. Contrary to its `req.hdr_cnt()` cousin, this function returns the number of full line headers and does not stop on commas.

`req.hdr(<[<name>],<occ>]) : string`
 This extracts the last occurrence of header <name> in an HTTP request. When used from an ACL, all occurrences are iterated over until a match is found. Optionally, a specific occurrence might be specified as a position number. Positive values indicate a position from the first occurrence, with 1 being the first one. Negative values indicate positions relative to the last one, with -1 being the last one. A typical use is with the X-Forwarded-For header once converted to IP, associated with an IP stick-table. The function considers any comma as a delimiter for distinct values. If full-line headers are desired instead, use `req.fhdr()`. Please carefully check RFC2616 to know how certain headers are supposed to be parsed. Also, some of them are case insensitive (eg: Connection).

ACL derivatives :
`hdr(<[<name>],<occ>])` : exact string match
`hdr_beg(<[<name>],<occ>])` : prefix match
`hdr_dir(<[<name>],<occ>])` : subdir match
`hdr_dom(<[<name>],<occ>])` : domain match
`hdr_end(<[<name>],<occ>])` : suffix match
`hdr_len(<[<name>],<occ>])` : length match
`hdr_reg(<[<name>],<occ>])` : regex match
`hdr_sub(<[<name>],<occ>])` : substring match

`req.hdr_cnt(<[<name>]) : integer`
`hdr_cnt(<[<header>]) : integer (deprecated)`
 Returns an integer value representing the number of occurrences of request header field name <name>, or the total number of header field values if

<name> is not specified. It is important to remember that one header line may count as several headers if it has several values. The function considers any comma as a delimiter for distinct values. If full-line headers are desired to instead, req.hndr_cnt() should be used instead. With ACLs, it can be used to detect presence, absence or abuse of a specific header, as well as to block request smuggling attacks by rejecting requests which contain more than one of certain headers. See "req.hdr" for more information on header matching.

```
req.hdr_ip(<name>[,<occ>]) : ip
```

```
hdr_ip(<name>[,<occ>]) : ip (deprecated)
```

This extracts the last occurrence of header <name> in an HTTP request, converts it to an IPv4 or IPv6 address and returns this address. When used with ACLs, all occurrences are checked, and if <name> is omitted, every value of every header is checked. Optionally, a specific occurrence might be specified as a position number. Positive values indicate a position from the first occurrence, with 1 being the first one. Negative values indicate positions relative to the last one, with -1 being the last one. A typical use is with the X-Forwarded-For and X-Client-IP headers.

```
req.hdr_val(<name>[,<occ>]) : integer
```

```
hdr_val(<name>[,<occ>]) : integer (deprecated)
```

This extracts the last occurrence of header <name> in an HTTP request, and converts it to an integer value. When used with ACLs, all occurrences are checked, and if <name> is omitted, every value of every header is checked. Optionally, a specific occurrence might be specified as a position number. Positive values indicate a position from the first occurrence, with 1 being the first one. Negative values indicate positions relative to the last one, with -1 being the last one. A typical use is with the X-Forwarded-For header.

```
http_auth(<userlist>) : boolean
```

Returns a boolean indicating whether the authentication data received from the client match a username & password stored in the specified userlist. This fetch function is not really useful outside of ACLs. Currently only http basic auth is supported.

```
http_auth_group(<userlist>) : string
```

Returns a string corresponding to the user name found in the authentication data received from the client if both the user name and password are valid according to the specified userlist. The main purpose is to use it in ACLs where it is then checked whether the user belongs to any group within a list. This fetch function is not really useful outside of ACLs. Currently only http basic auth is supported.

```
ACL derivatives :
```

```
http_auth_group(<userlist>) : group ...
```

Returns true when the user extracted from the request and whose password is valid according to the specified userlist belongs to at least one of the groups.

```
http_first_req : boolean
```

Returns true when the request being processed is the first one of the connection. This can be used to add or remove headers that may be missing from some requests when a request is not the first one, or to help grouping requests in the logs.

```
method : integer + string
```

Returns an integer value corresponding to the method in the HTTP request. For example, "GET" equals 1 (check sources to establish the matching). Value 9 means "other method" and may be converted to a string extracted from the stream. This should not be used directly as a sample, this is only meant to be used from ACLs, which transparently convert methods from patterns to these integer + string values. Some predefined ACL already check for most common methods.

```
ACL derivatives :
```

```
method : case insensitive method match
```

```
Example :
```

```
# only accept GET and HEAD requests
acl valid_method method GET HEAD
http-request deny if ! valid_method
```

```
path : string
```

This extracts the request's URL path, which starts at the first slash and ends before the question mark (without the host part). A typical use is with prefetch-capable caches, and with portals which need to aggregate multiple information from databases and keep them in caches. Note that with outgoing caches, it would be wiser to use "url" instead. With ACLs, it's typically used to match exact file names (eg: "/login.php"), or directory parts using the derivative forms. See also the "url" and "base" fetch methods.

```
ACL derivatives :
```

```
path : exact string match
path_beg : prefix match
path_dir : subdir match
path_dom : domain match
path_end : suffix match
path_len : length match
path_reg : regex match
path_sub : substring match
```

```
req.ver : string
```

```
req_ver : string (deprecated)
```

Returns the version string from the HTTP request, for example "1.1". This can be useful for logs, but is mostly there for ACL. Some predefined ACL already check for versions 1.0 and 1.1.

```
ACL derivatives :
```

```
req_ver : exact string match
```

```
res.comp : boolean
```

Returns the boolean "true" value if the response has been compressed by HAPROXY, otherwise returns boolean "false". This may be used to add information in the logs.

```
res.comp_algo : string
```

Returns a string containing the name of the algorithm used if the response was compressed by HAPROXY, for example : "deflate". This may be used to add some information in the logs.

```
res.cookie(<name>]) : string
```

```
scookie(<name>]) : string (deprecated)
```

This extracts the last occurrence of the cookie name <name> on a "Set-Cookie" header line from the response, and returns its value as string. If no name is specified, the first cookie value is returned.

```
ACL derivatives :
```

```
scookie(<name>) : exact string match
```

```
res.cookie_cnt(<name>]) : integer
```

```
scookie_cnt(<name>]) : integer (deprecated)
```

Returns an integer value representing the number of occurrences of the cookie <name> in the response, or all cookies if <name> is not specified. This is mostly useful when combined with ACLs to detect suspicious responses.

```
res.cookie_val(<name>]) : integer
```

```
scookie_val(<name>]) : integer (deprecated)
```

This extracts the last occurrence of the cookie name <name> on a "Set-Cookie" header line from the response, and converts its value to an integer which is returned. If no name is specified, the first cookie value is returned.

```
11570
```

```
11571 res.fhdr([<name>[,<occ>]]) : string
11572 This extracts the last occurrence of header <name> in an HTTP response, or of
11573 the last header if no <name> is specified. Optionally, a specific occurrence
11574 might be specified as a position number. Positive values indicate a position
11575 from the first occurrence, with 1 being the first one. Negative values
11576 indicate positions relative to the last one, with -1 being the last one. It
11577 differs from res.hdr() in that any commas present in the value are returned
11578 and are not used as delimiters. If this is not desired, the res.hdr() fetch
11579 should be used instead. This is sometimes useful with headers such as Date or
11580 Expires.
11581
11582
11583 res.fhdr_cnt([<name>]) : integer
11584 Returns an integer value representing the number of occurrences of response
11585 header field name <name>, or the total number of header fields if <name> is
11586 not specified. Contrary to its res.hdr_cnt() cousin, this function returns
11587 the number of full line headers and does not stop on commas. If this is not
11588 desired, the res.hdr_cnt() fetch should be used instead.
11589
11590 res.hdr([<name>[,<occ>]]) : string
11591 shdr([<name>[,<occ>]]) : string (deprecated)
11592 This extracts the last occurrence of header <name> in an HTTP response, or of
11593 the last header if no <name> is specified. Optionally, a specific occurrence
11594 might be specified as a position number. Positive values indicate a position
11595 from the first occurrence, with 1 being the first one. Negative values
11596 indicate positions relative to the last one, with -1 being the last one. This
11597 can be useful to learn some data into a stick-table. The function considers
11598 any comma as a delimiter for distinct values. If this is not desired, the
11599 res.fhdr() fetch should be used instead.
11600
11601 ACL derivatives :
11602 shdr([<name>[,<occ>]]) : exact string match
11603 shdr_beg([<name>[,<occ>]]) : prefix match
11604 shdr_dir([<name>[,<occ>]]) : subdir match
11605 shdr_dom([<name>[,<occ>]]) : domain match
11606 shdr_end([<name>[,<occ>]]) : suffix match
11607 shdr_len([<name>[,<occ>]]) : length match
11608 shdr_reg([<name>[,<occ>]]) : regex match
11609 shdr_sub([<name>[,<occ>]]) : substring match
11610
11611 res.hdr_cnt([<name>]) : integer
11612 shdr_cnt([<name>]) : integer (deprecated)
11613 Returns an integer value representing the number of occurrences of response
11614 header field name <name>, or the total number of header fields if <name> is
11615 not specified. The function considers any comma as a delimiter for distinct
11616 values. If this is not desired, the res.fhdr_cnt() fetch should be used
11617 instead.
11618
11619 res.hdr_ip([<name>[,<occ>]]) : ip
11620 shdr_ip([<name>[,<occ>]]) : ip (deprecated)
11621 This extracts the last occurrence of header <name> in an HTTP response,
11622 convert it to an IPv4 or IPv6 address and returns this address. Optionally, a
11623 specific occurrence might be specified as a position number. Positive values
11624 indicate a position from the first occurrence, with 1 being the first one.
11625 Negative values indicate positions relative to the last one, with -1 being
11626 the last one. This can be useful to learn some data into a stick table.
11627
11628 res.hdr_val([<name>[,<occ>]]) : integer
11629 shdr_val([<name>[,<occ>]]) : integer (deprecated)
11630 This extracts the last occurrence of header <name> in an HTTP response, and
11631 converts it to an integer value. Optionally, a specific occurrence might be
11632 specified as a position number. Positive values indicate a position from the
11633 first occurrence, with 1 being the first one. Negative values indicate
11634 positions relative to the last one, with -1 being the last one. This can be
11635 useful to learn some data into a stick table.
```

```
11636 res.ver : string
11637 resp_ver : string (deprecated)
11638 Returns the version string from the HTTP response, for example "1.1". This
11639 can be useful for logs, but is mostly there for ACL.
11640
11641 ACL derivatives :
11642 resp_ver : exact string match
11643
11644 set-cookie([<name>]) : string (deprecated)
11645 This extracts the last occurrence of the cookie name <name> on a "Set-Cookie"
11646 header line from the response and uses the corresponding value to match. This
11647 can be comparable to what "appsession" does with default options, but with
11648 support for multi-peer synchronization and state keeping across restarts.
11649
11650 This fetch function is deprecated and has been superseded by the "res.cook"
11651 fetch. This keyword will disappear soon.
11652
11653 See also : "appsession"
11654
11655 status : integer
11656 Returns an integer containing the HTTP status code in the HTTP response, for
11657 example, 302. It is mostly used within ACLs and integer ranges, for example,
11658 to remove any Location header if the response is not a 3xx.
11659
11660 url : string
11661 This extracts the request's URL as presented in the request. A typical use is
11662 with prefetch-capable caches, and with portals which need to aggregate
11663 multiple information from databases and keep them in caches. With ACLs, using
11664 "path" is preferred over using "url", because clients may send a full URL as
11665 is normally done with proxies. The only real use is to match "*" which does
11666 not match in "path", and for which there is already a predefined ACL. See
11667 also "path" and "base".
11668
11669 ACL derivatives :
11670 url : exact string match
11671 url_beg : prefix match
11672 url_dir : subdir match
11673 url_dom : domain match
11674 url_end : suffix match
11675 url_len : length match
11676 url_reg : regex match
11677 url_sub : substring match
11678
11679 url_ip : ip
11680 This extracts the IP address from the request's URL when the host part is
11681 presented as an IP address. Its use is very limited. For instance, a
11682 monitoring system might use this field as an alternative for the source IP in
11683 order to test what path a given source address would follow, or to force an
11684 entry in a table for a given source address. With ACLs it can be used to
11685 restrict access to certain systems through a proxy, for example when combined
11686 with option "http_proxy".
11687
11688 url_port : integer
11689 This extracts the port part from the request's URL. Note that if the port is
11690 not specified in the request, port 80 is assumed. With ACLs it can be used to
11691 restrict access to certain systems through a proxy, for example when combined
11692 with option "http_proxy".
11693
11694 urlp([<name>[,<delim>]]) : string
11695 url_param([<name>[,<delim>]]) : string
11696 This extracts the first occurrence of the parameter <name> in the query
11697 string, which begins after either '?' or <delim>, and which ends before '&',
11698 ';' or <delim>. The parameter name is case-sensitive. The result is a string
11699 corresponding to the value of the parameter <name> as presented in the
11700
```

request (no URL decoding is performed). This can be used for session stickiness based on a client ID, to extract an application cookie passed as a URL parameter, or in ACLs to apply some checks. Note that the ACL version of this fetch do not iterate over multiple parameters and stop at the first one as well.

```
ACL derivatives :
  urlp(<name>[,<delim>]) : exact string match
  urlp_beg(<names>[,<delim>]) : prefix match
  urlp_dir(<names>[,<delim>]) : subdir match
  urlp_dom(<names>[,<delim>]) : domain match
  urlp_end(<names>[,<delim>]) : suffix match
  urlp_len(<names>[,<delim>]) : length match
  urlp_reg(<names>[,<delim>]) : regex match
  urlp_sub(<names>[,<delim>]) : substring match
```

```
Example :
# match http://example.com/foo?PHPSESSIONID=some_id
stick on urlp(PHPSESSIONID)
# match http://example.com/foo;JSESSIONID=some_id
stick on urlp(JSESSIONID,;)
```

```
urlp_val(<name>[,<delim>]) : integer
See "urlp" above. This one extracts the URL parameter <name> in the request and converts it to an integer value. This can be used for session stickiness based on a user ID for example, or with ACLs to match a page number or price.
```

7.4. Pre-defined ACLs

Some predefined ACLs are hard-coded so that they do not have to be declared in every frontend which needs them. They all have their names in upper case in order to avoid confusion. Their equivalence is provided below.

ACL name	Equivalent to	Usage
FALSE	always_false	never match
HTTP	req_proto.http	match if protocol is valid HTTP
HTTP_1_0	req_ver 1.0	match HTTP version 1.0
HTTP_1_1	req_ver 1.1	match HTTP version 1.1
HTTP_CONTENT	hdr_val(content-length) gt 0	match an existing content-length
HTTP_URL_ABS	url_reg ^[/:]*://	match absolute URL with scheme
HTTP_URL_SLASH	url_beg /	match URL beginning with "/"
HTTP_URL_STAR	url	match URL equal to "*"
LOCALHOST	src 127.0.0.1/8	match connection from local host
METH_CONNECT	method CONNECT	match HTTP CONNECT method
METH_GET	method GET	match HTTP GET or HEAD method
METH_HEAD	method HEAD	match HTTP HEAD method
METH_OPTIONS	method OPTIONS	match HTTP OPTIONS method
METH_POST	method POST	match HTTP POST method
METH_TRACE	method TRACE	match HTTP TRACE method
RDP_COOKIE	req_rdp.cookie_cnt gt 0	match presence of an RDP cookie
REQ_CONTENT	req_len gt 0	match data in the request buffer
TRUE	always_true	always match
WAIT_END	wait_end	wait for end of content analysis

8. Logging

One of HAProxy's strong points certainly lies in its precise logs. It probably provides the finest level of information available for such a product, which is

very important for troubleshooting complex environments. Standard information provided in logs include client ports, TCP/HTTP state timers, precise session state at termination and precise termination cause, information about decisions to direct traffic to a server, and of course the ability to capture arbitrary headers.

In order to improve administrators reactivity, it offers a great transparency about encountered problems, both internal and external, and it is possible to send logs to different sources at the same time with different level filters :

- global process-level logs (system errors, start/stop, etc..)
- per-instance system and internal errors (lack of resource, bugs, ...)
- per-instance external troubles (servers up/down, max connections)
- per-instance activity (client connections), either at the establishment or at the termination.

The ability to distribute different levels of logs to different log servers allow several production teams to interact and to fix their problems as soon as possible. For example, the system team might monitor system-wide errors, while the application team might be monitoring the up/down for their servers in real time, and the security team might analyze the activity logs with one hour delay.

8.1. Log levels

TCP and HTTP connections can be logged with information such as the date, time, source IP address, destination address, connection duration, response times, HTTP request, HTTP return code, number of bytes transmitted, conditions in which the session ended, and even exchanged cookies values. For example track a particular user's problems. All messages may be sent to up to two syslog servers. Check the "log" keyword in section 4.2 for more information about log facilities.

8.2. Log formats

HAProxy supports 5 log formats. Several fields are common between these formats and will be detailed in the following sections. A few of them may vary slightly with the configuration, due to indicators specific to certain options. The supported formats are as follows :

- the default format, which is very basic and very rarely used. It only provides very basic information about the incoming connection at the moment it is accepted : source IP:port, destination IP:port, and frontend-name. This mode will eventually disappear so it will not be described to great extents.
- the TCP format, which is more advanced. This format is enabled when "option tcplog" is set on the frontend. HAProxy will then usually wait for the connection to terminate before logging. This format provides much richer information, such as timers, connection counts, queue size, etc... This format is recommended for pure TCP proxies.

- the HTTP format, which is the most advanced for HTTP proxying. This format is enabled when "option httplog" is set on the frontend. It provides the same information as the TCP format with some HTTP-specific fields such as the request, the status code, and captures of headers and cookies. This format is recommended for HTTP proxies.

- the CLF HTTP format, which is equivalent to the HTTP format, but with the fields arranged in the same order as the CLF format. In this mode, all timers, captures, flags, etc... appear one per field after the end of the

11831 common fields, in the same order they appear in the standard HTTP format.
11832
11833 - the custom log format, allows you to make your own log line.
11834

11835 Next sections will go deeper into details for each of these formats. Format
11836 specification will be performed on a "field" basis. Unless stated otherwise, a
11837 field is a portion of text delimited by any number of spaces. Since syslog
11838 servers are susceptible of inserting fields at the beginning of a line, it is
11839 always assumed that the first field is the one containing the process name and
11840 identifier.
11841

11842 Note : Since log lines may be quite long, the log examples in sections below
11843 might be broken into multiple lines. The example log lines will be
11844 prefixed with 3 closing angle brackets ('>>>') and each time a log is
11845 broken into multiple lines, each non-final line will end with a
11846 backslash ('\') and the next line will start indented by two characters.
11847
11848

11849 8.2.1. Default log format
11850 -----
11851

11852 This format is used when no specific option is set. The log is emitted as soon
11853 as the connection is accepted. One should note that this currently is the only
11854 format which logs the request's destination IP and ports.
11855

11856 Example :
11857 Listen www
11858 mode http
11859 log global
11860 server srv1 127.0.0.1:8000
11861

11862 >>> Feb 6 12:12:09 localhost \
11863 haproxy[14385]: Connect from 10.0.1.2:33312 to 10.0.3.31:8012 \
11864 (www/HTTP)
11865

11866 Field Format Extract from the example above
11867 1 process_name '[' pid ':'
11868 2 'Connect from', haproxy[14385]:
11869 3 source_ip ':' source_port Connect from
11870 4 'to' 10.0.1.2:33312
11871 5 destination_ip ':' destination_port to
11872 6 '(' frontend_name '/' mode ')' 10.0.3.31:8012
11873 (www/HTTP)

11874 Detailed fields description :
11875 - "source_ip" is the IP address of the client which initiated the connection.
11876 - "source_port" is the TCP port of the client which initiated the connection.
11877 - "destination_ip" is the IP address the client connected to.
11878 - "destination_port" is the TCP port the client connected to.
11879 - "frontend_name" is the name of the frontend (or listener) which received
11880 and processed the connection.
11881 - "mode" is the mode the frontend is operating (TCP or HTTP).
11882

11883 In case of a UNIX socket, the source and destination addresses are marked as
11884 "unix:" and the ports reflect the internal ID of the socket which accepted the
11885 connection (the same ID as reported in the stats).
11886

11887 It is advised not to use this deprecated format for newer installations as it
11888 will eventually disappear.
11889

11890 8.2.2. TCP log format
11891 -----
11892

11893 The TCP format is used when "option tcplog" is specified in the frontend, and
11894 is the recommended format for pure TCP proxies. It provides a lot of precious
11895

11896 information for troubleshooting. Since this format includes timers and byte
11897 counts, the log is normally emitted at the end of the session. It can be
11898 emitted earlier if "option logasap" is specified, which makes sense in most
11899 environments with long sessions such as remote terminals. Sessions which match
11900 the "monitor" rules are never logged. It is also possible not to emit logs for
11901 sessions for which no data were exchanged between the client and the server, by
11902 specifying "option dontlognull" in the frontend. Successful connections will
11903 not be logged if "option dontlog-normal" is specified in the frontend. A few
11904 fields may slightly vary depending on some configuration options, those are
11905 marked with a star (*) after the field name below.
11906

11907 Example :
11908 frontend fnt
11909 mode tcp
11910 option tcplog
11911 log global
11912 default_backend bck
11913
11914 backend bck
11915 server srv1 127.0.0.1:8000
11916

11917 >>> Feb 6 12:12:56 localhost \
11918 haproxy[14387]: 10.0.1.2:33313 [06/Feb/2009:12:12:51.443] fnt \
11919 bck/srv1 0/0/5007 212 -- 0/0/0/0/3 0/0
11920

11921 Field Format Extract from the example above
11922 1 process_name '[' pid ':'
11923 2 client_ip ':' client_port haproxy[14387]:
11924 3 '[' accept_date ']' 10.0.1.2:33313
11925 4 frontend_name '[06/Feb/2009:12:12:51.443]
11926 5 backend_name '/' server_name fnt
11927 6 Tw '/' Tc '/' Tt* bck/srv1
11928 7 bytes_read* 0/0/5007
11929 8 termination_state 212
11930 9 actconn '/' feconn '/' beconn '/' srv_conn '/' retries* --
11931 10 srv_queue '/' backend_queue 0/0/0/0/3
11932 0/0

11933 Detailed fields description :

11934 - "client_ip" is the IP address of the client which initiated the TCP
11935 connection to haproxy. If the connection was accepted on a UNIX socket
11936 instead, the IP address would be replaced with the word "unix". Note that
11937 when the connection is accepted on a socket configured with "accept-proxy"
11938 and the PROXY protocol is correctly used, then the logs will reflect the
11939 forwarded connection's information.
11940

11941 - "client_port" is the TCP port of the client which initiated the connection.
11942 If the connection was accepted on a UNIX socket instead, the port would be
11943 replaced with the ID of the accepting socket, which is also reported in the
11944 stats interface.
11945

11946 - "accept_date" is the exact date when the connection was received by haproxy
11947 (which might be very slightly different from the date observed on the
11948 network if there was some queuing in the system's backlog). This is usually
11949 the same date which may appear in any upstream firewall's log.
11950

11951 - "frontend_name" is the name of the frontend (or listener) which received
11952 and processed the connection.
11953

11954 - "backend_name" is the name of the backend (or listener) which was selected
11955 to manage the connection to the server. This will be the same as the
11956 frontend if no switching rule has been applied, which is common for TCP
11957 applications.
11958

11959 - "server_name" is the name of the last server to which the connection was
11960 sent, which might differ from the first one if there were connection errors
11961

and a redispatch occurred. Note that this server belongs to the backend which processed the request. If the connection was aborted before reaching a server, "<NOSRV>" is indicated instead of a server name.

- "Tm" is the total time in milliseconds spent waiting in the various queues. It can be "-1" if the connection was aborted before reaching the queue. See "timers" below for more details.

- "Tc" is the total time in milliseconds spent waiting for the connection to establish to the final server, including retries. It can be "-1" if the connection was aborted before a connection could be established. See "timers" below for more details.

- "Tt" is the total time in milliseconds elapsed between the accept and the last close. It covers all possible processing. There is one exception, if "option logasap" was specified, then the time counting stops at the moment the log is emitted. In this case, a '+' sign is prepended before the value, indicating that the final one will be larger. See "timers" below for more details.

- "bytes_read" is the total number of bytes transmitted from the server to the client when the log is emitted. If "option logasap" is specified, the this value will be prefixed with a '+' sign indicating that the final one may be larger. Please note that this value is a 64-bit counter, so log analysis tools must be able to handle it without overflowing.

- "termination_state" is the condition the session was in when the session ended. This indicates the session state, which side caused the end of session to happen, and for what reason (timeout, error, ...). The normal flags should be "-", indicating the session was closed by either end with no data remaining in buffers. See below 'Session state at disconnection' for more details.

- "actconn" is the total number of concurrent connections on the process when the session was logged. It is useful to detect when some per-process system limits have been reached. For instance, if actconn is close to 512 when multiple connection errors occur, chances are high that the system limits the process to use a maximum of 1024 file descriptors and that all of them are used. See section 3 "Global parameters" to find how to tune the system.

- "feconn" is the total number of concurrent connections on the frontend when the session was logged. It is useful to estimate the amount of resource required to sustain high loads, and to detect when the frontend's "maxconn" has been reached. Most often when this value increases by huge jumps, it is because there is congestion on the backend servers, but sometimes it can be caused by a denial of service attack.

- "beconn" is the total number of concurrent connections handled by the backend when the session was logged. It includes the total number of concurrent connections active on servers as well as the number of connections pending in queues. It is useful to estimate the amount of additional servers needed to support high loads for a given application. Most often when this value increases by huge jumps, it is because there is congestion on the backend servers, but sometimes it can be caused by a denial of service attack.

- "srv_conn" is the total number of concurrent connections still active on the server when the session was logged. It can never exceed the server's configured "maxconn" parameter. If this value is very often close or equal to the server's "maxconn", it means that traffic regulation is involved a lot, meaning that either the server's maxconn value is too low, or that there aren't enough servers to process the load with an optimal response time. When only one of the server's "srv_conn" is high, it usually means that this server has some trouble causing the connections to take longer to be processed than on other servers.

- "retries" is the number of connection retries experienced by this session when trying to connect to the server. It must normally be zero, unless a server is being stopped at the same moment the connection was attempted. Frequent retries generally indicate either a network problem between haproxy and the server, or a misconfigured system backlog on the server preventing new connections from being queued. This field may optionally be prefixed with a '+' sign, indicating that the session has experienced a redispatch after the maximal retry count has been reached on the initial server. In this case, the server name appearing in the log is the one the connection was redispatched to, and not the first one, though both may sometimes be the same in case of hashing for instance. So as a general rule of thumb, when a '+' is present in front of the retry count, this count should not be attributed to the logged server.

- "srv_queue" is the total number of requests which were processed before this one in the server queue. It is zero when the request has not gone through the server queue. It makes it possible to estimate the approximate server's response time by dividing the time spent in queue by the number of requests in the queue. It is worth noting that if a session experiences a redispatch and passes through two server queues, their positions will be cumulated. A request should not pass through both the server queue and the backend queue unless a redispatch occurs.

- "backend_queue" is the total number of requests which were processed before this one in the backend's global queue. It is zero when the request has not gone through the global queue. It makes it possible to estimate the average queue length, which easily translates into a number of missing servers when divided by a server's "maxconn" parameter. It is worth noting that if a session experiences a redispatch, it may pass twice in the backend's queue, and then both positions will be cumulated. A request should not pass through both the server queue and the backend queue unless a redispatch occurs.

8.2.3. HTTP log format

The HTTP format is the most complete and the best suited for HTTP proxies. It is enabled by when "option httplog" is specified in the frontend. It provides the same level of information as the TCP format with additional features which are specific to the HTTP protocol. Just like the TCP format, the log is usually emitted at the end of the session, unless "option logasap" is specified, which generally only makes sense for download sites. A session which matches the "monitor" rules will never be logged. It is also possible not to log sessions for which no data were sent by the client by specifying "option dontlognull" in the frontend. Successful connections will not be logged if "option dontlog-normal" is specified in the frontend.

Most fields are shared with the TCP log, some being different. A few fields may slightly vary depending on some configuration options. Those ones are marked with a star (*) after the field name below.

Example :

```
frontend http-in
mode http
option httplog
log global
default_backend bck

backend static
server srv1 127.0.0.1:8000
```

>>> Feb 6 12:14:14 localhost \
haproxy[14389]: 10.0.1.2:33317 [06/Feb/2009:12:14:14.655] http-in \

```
12091 static/srv1 10/0/30/69/109 200 2750 - - ---- 1/1/1/1/0 0/0 {lwt.eu} \
12092 {} "GET /index.html HTTP/1.1"
12093
12094 Format
12095 1 process_name '[' pid '];'
12096 2 client_ip ',' client_port
12097 3 '[' accept_date ','
12098 4 frontend_name '/' server_name http-in
12099 5 backend_name '/' backend_queue static/srv1
12100 6 Tq '/' Tw '/' Tc '/' Tr '/' Tt* 10/0/30/69/109
12101 7 status_code 200
12102 8 bytes_read* 2750
12103 9 captured_request_cookie -
12104 10 captured_response_cookie -
12105 11 termination_state ----
12106 12 actconn '/' feconn '/' beconn '/' retries* 1/1/1/1/0
12107 13 srv_queue '/' backend_queue 0/0
12108 14 '{' captured_request_headers* '}' {haproxy.lwt.eu}
12109 15 '{' captured_response_headers* '}' {}
12110 16 ',"' http_request ',' "GET /index.html HTTP/1.1"
12111
```

Detailed fields description :

- "client_ip" is the IP address of the client which initiated the TCP connection to haproxy. If the connection was accepted on a UNIX socket instead, the IP address would be replaced with the word "unix". Note that when the connection is accepted on a socket configured with "accept-proxy" and the PROXY protocol is correctly used, then the logs will reflect the forwarded connection's information.
- "client_port" is the TCP port of the client which initiated the connection. If the connection was accepted on a UNIX socket instead, the port would be replaced with the ID of the accepting socket, which is also reported in the stats interface.
- "accept_date" is the exact date when the TCP connection was received by haproxy (which might be very slightly different from the date observed on the network if there was some queuing in the system's backlog). This is usually the same date which may appear in any upstream firewall's log. This does not depend on the fact that the client has sent the request or not.
- "frontend_name" is the name of the frontend (or listener) which received and processed the connection.
- "backend_name" is the name of the backend (or listener) which was selected to manage the connection to the server. This will be the same as the frontend if no switching rule has been applied.
- "server_name" is the name of the last server to which the connection was sent, which might differ from the first one if there were connection errors and a redispatch occurred. Note that this server belongs to the backend which processed the request. If the request was aborted before reaching a server, "<NOSRV>" is indicated instead of a server name. If the request was intercepted by the stats subsystem, "<STATS>" is indicated instead.
- "Tq" is the total time in milliseconds spent waiting for the client to send a full HTTP request, not counting data. It can be "-1" if the connection was aborted before a complete request could be received. It should always be very small because a request generally fits in one single packet. Large times here generally indicate network trouble between the client and haproxy. See "Timers" below for more details.
- "Tw" is the total time in milliseconds spent waiting in the various queues. It can be "-1" if the connection was aborted before reaching the queue. See "Timers" below for more details.

- "Tc" is the total time in milliseconds spent waiting for the connection to establish to the final server, including retries. It can be "-1" if the request was aborted before a connection could be established. See "Timers" below for more details.
- "Tr" is the total time in milliseconds spent waiting for the server to send a full HTTP response, not counting data. It can be "-1" if the request was aborted before a complete response could be received. It generally matches the server's processing time for the request, though it may be altered by the amount of data sent by the client to the server. Large times here on "GET" requests generally indicate an overloaded server. See "Timers" below for more details.
- "Tt" is the total time in milliseconds elapsed between the accept and the last close. It covers all possible processing. There is one exception, if "option logasap" was specified, then the time counting stops at the moment the log is emitted. In this case, a '+' sign is prepended before the value, indicating that the final one will be larger. See "Timers" below for more details.
- "status_code" is the HTTP status code returned to the client. This status is generally set by the server, but it might also be set by haproxy when the server cannot be reached or when its response is blocked by haproxy.
- "bytes_read" is the total number of bytes transmitted to the client when the log is emitted. This does include HTTP headers. If "option logasap" is specified, the this value will be prefixed with a '+' sign indicating that the final one may be larger. Please note that this value is a 64-bit counter, so log analysis tools must be able to handle it without overflowing.
- "captured_request_cookie" is an optional "name=value" entry indicating that the client had this cookie in the request. The cookie name and its maximum length are defined by the "capture cookie" statement in the frontend configuration. The field is a single dash ('-') when the option is not set. Only one cookie may be captured, it is generally used to track session ID exchanges between a client and a server to detect session crossing between clients due to application bugs. For more details, please consult the section "Capturing HTTP headers and cookies" below.
- "captured_response_cookie" is an optional "name=value" entry indicating that the server has returned a cookie with its response. The cookie name and its maximum length are defined by the "capture cookie" statement in the frontend configuration. The field is a single dash ('-') when the option is not set. Only one cookie may be captured, it is generally used to track session ID exchanges between a client and a server to detect session crossing between clients due to application bugs. For more details, please consult the section "Capturing HTTP headers and cookies" below.
- "termination_state" is the condition the session was in when the session ended. This indicates the session state, which side caused the end of session to happen, for what reason (timeout, error, ...), just like in TCP logs, and information about persistence operations on cookies in the last two characters. The normal flags should begin with "---", indicating the session was closed by either end with no data remaining in buffers. See below "Session state at disconnection" for more details.
- "actconn" is the total number of concurrent connections on the process when the session was logged. It is useful to detect when some per-process system limits have been reached. For instance, if actconn is close to 512 or 1024 when multiple connection errors occur, chances are high that the system limits the process to use a maximum of 1024 file descriptors and that all of them are used. See section 3 "Global parameters" to find how to tune the system.

12351	%b\ %sc\ %rc\ %sq\ %bq\ %CC\ %CS\ %hrl\ %hsl
12352	
12353	and the default TCP format is defined this way :
12354	
12355	log-format %ci:%p\ [%t]\ %ft\ %b/%s\ %Tw/%Tc/%Tt\ %B\ %ts\ \
12356	%ac/%Tc/%bc/%sc/%rc\ %sq/%bq
12357	
12358	Please refer to the table below for currently defined variables :
12359	
12360	
12361	R var field name (8.2.2 and 8.2.3 for description) type
12362	
12363	%o special variable, apply flags on all next var
12364	
12365	%B bytes_read (from server to client) numeric
12366	H %CC captured_request_cookie string
12367	H %CS captured_response_cookie string
12368	H %H hostname string
12369	%ID unique-id string
12370	%ST status_code numeric
12371	%T gmt_date_time date
12372	%Tc Tc numeric
12373	%Tl local_date_time date
12374	%Tq Tq numeric
12375	H %Tr Tr numeric
12376	%Ts timestamp numeric
12377	%Tt Tt numeric
12378	%Tw Tw numeric
12379	%U bytes_uploaded (from client to server) numeric
12380	%ac actconn numeric
12381	%b backend_name string
12382	%bc beconn (backend concurrent connections) numeric
12383	%bi backend_source_ip (connecting address) IP
12384	%bp backend_source_port (connecting address) numeric
12385	%bq backend_queue numeric
12386	%ci client_ip (accepted address) IP
12387	%cp client_port (accepted address) numeric
12388	%f frontend_name string
12389	%fc feconn (frontend concurrent connections) numeric
12390	%fi frontend_ip (accepting address) IP
12391	%fp frontend_port (accepting address) numeric
12392	%ft frontend_name_transport ('_' suffix for SSL) string
12393	%hr captured_request_headers default style string
12394	%hrl captured_request_headers CLF style string
12395	%hs captured_response_headers default style string
12396	%hsl captured_response_headers CLF style string
12397	%ms accept date milliseconds (left-padded with 0) numeric
12398	%pid PID numeric
12399	H %r http_request string
12400	%rc retries numeric
12401	%rt request_counter (HTTP req or TCP session) numeric
12402	%s server_name string
12403	%sc srv_conn (server concurrent connections) numeric
12404	%si server_IP (target address) IP
12405	%sp server_port (target address) numeric
12406	%sq srv_queue numeric
12407	S %ssl ssl_ciphers (ex: AES-SHA) string
12408	S %ssl_v ssl_version (ex: TLSv1) string
12409	%t date_time (with millisecond resolution) date
12410	%ts termination_state string
12411	H %tsc termination_state with cookie status string
12412	

R = Restrictions : H = mode http only ; S = SSL only

12416
12417
12418
12419

8.2.5. Error log format

When an incoming connection fails due to an SSL handshake or an invalid PROXY protocol header, haproxy will log the event using a shorter, fixed line format. By default, logs are emitted at the LOG_INFO level, unless the option "log-separate-errors" is set in the backend, in which case the LOG_ERR level will be used. Connections on which no data are exchanged (eg: probes) are not logged if the "dontlognull" option is set.

12420
12421
12422
12423
12424
12425

12427 The format looks like this :

```
12429 >>> Dec 3 18:27:14 localhost \
12430 haproxy[6103]: 127.0.0.1:56059 [03/Dec/2012:17:35:10.380] frrt/fl: \
12431 Connection error during SSL handshake
```

Field	Format	Extract from the example above
1	process_name ['pid ':'']	haproxy[61031]
2	client_ip ':' client_port	127.0.0.1:56059
3	['accept_date ']	[03/Dec/2012:17:35:10.380]
4	frontend_name "/" bind_name "	frtfrk
5	message	Connection error during SSL handshake

These fields just provide minimal information to help debugging connection failures.

8.3. Advanced logging options

Some advanced logging options are often looked for but are not easy to find out just by looking at the various options. Here is an entry point for the few options which can enable better logging. Please refer to the keywords reference for more information about their usage.

8.3.1. Disabling logging of external tests

It is quite common to have some monitoring tools perform health checks on haproxy. Sometimes it will be a layer 3 load-balancer such as LVS or any commercial load-balancer, and sometimes it will simply be a more complete monitoring system such as Nagios. When the tests are very frequent, users often ask how to disable logging for those checks. There are three possibilities :

- if connections come from everywhere and are just TCP probes, it is often desired to simply disable logging of connections without data exchange, by setting "option dontlognull" in the frontend. It also disables logging of port scans, which may or may not be desired.
- if the connection come from a known source network, use "monitor-net" to declare this network as monitoring only. Any host in this network will then only be able to perform health checks, and their requests will not be logged. This is generally appropriate to designate a list of equipment such as other load-balancers.
- if the tests are performed on a known URI, use "monitor-uri" to declare this URI as dedicated to monitoring. Any host sending this request will only get the result of a health-check, and the request will not be logged.

12477 8.3.2. Logging before waiting for the session to terminate

12479
12480

12481 The problem with logging at end of connection is that you have no clue about
12482 what is happening during very long sessions, such as remote terminal sessions
12483 or large file downloads. This problem can be worked around by specifying
12484 "option logasap" in the frontend. Haproxy will then log as soon as possible,
12485 just before data transfer begins. This means that in case of TCP, it will still
12486 log the connection status to the server, and in case of HTTP, it will log just
12487 after processing the server headers. In this case, the number of bytes reported
12488 is the number of header bytes sent to the client. In order to avoid confusion
12489 with normal logs, the total time field and the number of bytes are prefixed
12490 with a '+' sign which means that real numbers are certainly larger.

8.3.3. Raising log level upon errors

12491 -----
12492
12493 Sometimes it is more convenient to separate normal traffic from errors logs,
12494 for instance in order to ease error monitoring from log files. When the option
12495 "log-separate-errors" is used, connections which experience errors, timeouts,
12496 retries, redispatches or HTTP status codes 5xx will see their syslog level
12497 raised from "info" to "err". This will help a syslog daemon store the log in
12498 a separate file. It is very important to keep the errors in the normal traffic
12499 file too, so that log ordering is not altered. You should also be careful if
12500 you already have configured your syslog daemon to store all logs higher than
12501 "notice" in an "admin" file, because the "err" level is higher than "notice".

8.3.4. Disabling logging of successful connections

12502 -----

12503 Although this may sound strange at first, some large sites have to deal with
12504 multiple thousands of logs per second and are experiencing difficulties keeping
12505 them intact for a long time or detecting errors within them. If the option
12506 "dontlog-normal" is set on the frontend, all normal connections will not be
12507 logged. In this regard, a normal connection is defined as one without any
12508 error, timeout, retry nor redispatch. In HTTP, the status code is checked too,
12509 and a response with a status 5xx is not considered normal and will be logged
12510 too. Of course, doing is is really discouraged as it will remove most of the
12511 useful information from the logs. Do this only if you have no other
12512 alternative.

8.4. Timing events

12513 -----

12514 Timers provide a great help in troubleshooting network problems. All values are
12515 reported in milliseconds (ms). These timers should be used in conjunction with
12516 the session termination flags. In TCP mode with "option tcplog" set on the
12517 frontend, 3 control points are reported under the form "Tw/Tc/Tt", and in HTTP
12518 mode, 5 control points are reported under the form "Tq/Tw/Tc/Tr/Tt" :

12519 - Tq: total time to get the client request (HTTP mode only). It's the time
12520 elapsed between the moment the client connection was accepted and the
12521 moment the proxy received the last HTTP header. The value "-1" indicates
12522 that the end of headers (empty line) has never been seen. This happens when
12523 the client closes prematurely or times out.

12524 - Tw: total time spent in the queues waiting for a connection slot. It
12525 accounts for backend queue as well as the server queues, and depends on the
12526 queue size, and the time needed for the server to complete previous
12527 requests. The value "-1" means that the request was killed before reaching
12528 the queue, which is generally what happens with invalid or denied requests.

12529 - Tc: total time to establish the TCP connection to the server. It's the time
12530 elapsed between the moment the proxy sent the connection request, and the
12531 moment it was acknowledged by the server, or between the TCP SYN packet and

12546 the matching SYN/ACK packet in return. The value "-1" means that the
12547 connection never established.

12548 - Tr: server response time (HTTP mode only). It's the time elapsed between
12549 the moment the TCP connection was established to the server and the moment
12550 the server sent its complete response headers. It purely shows its request
12551 processing time, without the network overhead due to the data transmission.
12552 It is worth noting that when the client has data to send to the server, for
12553 instance during a POST request, the time already runs, and this can distort
12554 apparent response time. For this reason, it's generally wise not to trust
12555 too much this field for POST requests initiated from clients behind an
12556 untrusted network. A value of "-1" here means that the last the response
12557 header (empty line) was never seen, most likely because the server timeout
12558 stroke before the server managed to process the request.

12559 - Tt: total session duration time, between the moment the proxy accepted it
12560 and the moment both ends were closed. The exception is when the "logasap"
12561 option is specified. In this case, it only equals (Tq+Tw+Tc+Tr), and is
12562 prefixed with a '+' sign. From this field, we can deduce "Td", the data
12563 transmission time, by subtracting other timers when valid :

$$12564 Td = Tt - (Tq + Tw + Tc + Tr)$$

12565 Timers with "-1" values have to be excluded from this equation. In TCP
12566 mode, "Tq" and "Tr" have to be excluded too. Note that "Tt" can never be
12567 negative.

12568 These timers provide precious indications on trouble causes. Since the TCP
12569 protocol defines retransmit delays of 3, 6, 12... seconds, we know for sure
12570 that timers close to multiples of 3s are nearly always related to lost packets
12571 due to network problems (wires, negotiation, congestion). Moreover, if "Tt" is
12572 close to a timeout value specified in the configuration, it often means that a
12573 session has been aborted on timeout.

Most common cases :

12574 - If "Tq" is close to 3000, a packet has probably been lost between the
12575 client and the proxy. This is very rare on local networks but might happen
12576 when clients are on far remote networks and send large requests. It may
12577 happen that values larger than usual appear here without any network cause.
12578 Sometimes, during an attack or just after a resource starvation has ended,
12579 haproxy may accept thousands of connections in a few milliseconds. The time
12580 spent accepting these connections will inevitably slightly delay processing
12581 of other connections, and it can happen that request times in the order of
12582 a few tens of milliseconds are measured after a few thousands of new
12583 connections have been accepted at once. Setting "option http-server-close"
12584 may display larger request times since "Tq" also measures the time spent
12585 waiting for additional requests.

12586 - If "Tc" is close to 3000, a packet has probably been lost between the
12587 server and the proxy during the server connection phase. This value should
12588 always be very low, such as 1 ms on local networks and less than a few tens
12589 of ms on remote networks.

12590 - If "Tr" is nearly always lower than 3000 except some rare values which seem
12591 to be the average majored by 3000, there are probably some packets lost
12592 between the proxy and the server.

12593 - If "Tt" is large even for small byte counts, it generally is because
12594 neither the client nor the server decides to close the connection, for
12595 instance because both have agreed on a keep-alive connection mode. In order
12596 to solve this issue, it will be needed to specify "option httpclose" on
12597 either the frontend or the backend. If the problem persists, it means that
12598 the server ignores the "close" connection mode and expects the client to
12599 close. Then it will be required to use "option forceclose". Having the

12611 smallest possible 'Tt' is important when connection regulation is used with
12612 the "maxconn" option on the servers, since no new connection will be sent
12613 to the server until another one is released.
12614

12615 Other noticeable HTTP log cases ('xx' means any value to be ignored) :

12616 Tq/Tw/Tc/Tr/+Tt The "option logasap" is present on the frontend and the log
12617 was emitted before the data phase. All the timers are valid
12618 except "Tt" which is shorter than reality.
12619

12620 -1/xx/xx/xx/Tt The client was not able to send a complete request in time
12621 or it aborted too early. Check the session termination flags
12622 then "timeout http-request" and "timeout client" settings.
12623

12624 Tq/-1/xx/xx/Tt It was not possible to process the request, maybe because
12625 servers were out of order, because the request was invalid
12626 or forbidden by ACL rules. Check the session termination
12627 flags.
12628

12629 Tq/Tw/-1/xx/Tt The connection could not establish on the server. Either it
12630 actively refused it or it timed out after Tt-(Tq+Tw) ms.
12631 Check the session termination flags, then check the
12632 "timeout connect" setting. Note that the tarpit action might
12633 return similar-looking patterns, with "Tw" equal to the time
12634 the client connection was maintained open.
12635

12636 Tq/Tw/Tc/-1/Tt The server has accepted the connection but did not return
12637 a complete response in time, or it closed its connection
12638 unexpectedly after Tt-(Tq+Tw+Tc) ms. Check the session
12639 termination flags, then check the "timeout server" setting.
12640

8.5. Session state at disconnection

12641 TCP and HTTP logs provide a session termination indicator in the
12642 "termination state" field, just before the number of active connections. It is
12643 2-characters long in TCP mode, and is extended to 4 characters in HTTP mode,
12644 each of which has a special meaning :

12645 - On the first character, a code reporting the first event which caused the
12646 session to terminate :

12647 C : the TCP session was unexpectedly aborted by the client.

12648 S : the TCP session was unexpectedly aborted by the server, or the
12649 server explicitly refused it.

12650 P : the session was prematurely aborted by the proxy, because of a
12651 connection limit enforcement, because a DENY filter was matched,
12652 because of a security check which detected and blocked a dangerous
12653 error in server response which might have caused information leak
(eg: cacheable cookie).
12654

12655 L : the session was locally processed by haproxy and was not passed to
12656 a server. This is what happens for stats and redirects.
12657

12658 R : a resource on the proxy has been exhausted (memory, sockets, source
12659 ports, ...). Usually, this appears during the connection phase, and
12660 system logs should contain a copy of the precise error. If this
12661 happens, it must be considered as a very serious anomaly which
12662 should be fixed as soon as possible by any means.
12663

12664 I : an internal error was identified by the proxy during a self-check.
12665 This should NEVER happen, and you are encouraged to report any log
12666

12667 containing this, because this would almost certainly be a bug. It
12668 would be wise to preventively restart the process after such an
12669 event too, in case it would be caused by memory corruption.

12670 D : the session was killed by haproxy because the server was detected
12671 as down and was configured to kill all connections when going down.

12672 U : the session was killed by haproxy on this backup server because an
12673 active server was detected as up and was configured to kill all
12674 backup connections when going up.

12675 K : the session was actively killed by an admin operating on haproxy.

12676 c : the client-side timeout expired while waiting for the client to
12677 send or receive data.

12678 s : the server-side timeout expired while waiting for the server to
12679 send or receive data.

12680 - : normal session completion, both the client and the server closed
12681 with nothing left in the buffers.

12682 - on the second character, the TCP or HTTP session state when it was closed :
12683 R : the proxy was waiting for a complete, valid REQUEST from the client
(HTTP mode only). Nothing was sent to any server.

12684 Q : the proxy was waiting in the QUEUE for a connection slot. This can
12685 only happen when servers have a 'maxconn' parameter set. It can
12686 also happen in the global queue after a redispatch consecutive to
12687 a failed attempt to connect to a dying server. If no redispatch is
12688 reported, then no connection attempt was made to any server.

12689 C : the proxy was waiting for the CONNECTION to establish on the
12690 server. The server might at most have noticed a connection attempt.

12691 H : the proxy was waiting for complete, valid response HEADERS from the
12692 server (HTTP only).

12693 D : the session was in the DATA phase.

12694 L : the proxy was still transmitting LAST data to the client while the
12695 server had already finished. This one is very rare as it can only
12696 happen when the client dies while receiving the last packets.

12697 T : the request was tarptimed. It has been held open with the client
12698 during the whole "timeout tarpit" duration or until the client
12699 closed, both of which will be reported in the "Tw" timer.

12700 - : normal session completion after end of data transfer.

12701 - the third character tells whether the persistence cookie was provided by
12702 the client (only in HTTP mode) :

12703 N : the client provided NO cookie. This is usually the case for new
12704 visitors, so counting the number of occurrences of this flag in the
12705 logs generally indicate a valid trend for the site frequentation.

12706 I : the client provided an INVALID cookie matching no known server.
12707 This might be caused by a recent configuration change, mixed
12708 cookies between HTTP/HTTPS sites, persistence conditionally
12709 ignored, or an attack.

12710 D : the client provided a cookie designating a server which was DOWN,
12711 so either "option persist" was used and the client was sent to
12712

12741 this server, or it was not set and the client was redispached to
12742 another server.
12743

12744 V : the client provided a VALID cookie, and was sent to the associated
12745 server.
12746

12747 E : the client provided a valid cookie, but with a last date which was
12748 older than what is allowed by the "maxidle" cookie parameter, so
12749 the cookie is consider EXPIRED and is ignored. The request will be
12750 redispached just as if there was no cookie.
12751

12752 O : the client provided a valid cookie, but with a first date which was
12753 older than what is allowed by the "maxlife" cookie parameter, so
12754 the cookie is consider too OLD and is ignored. The request will be
12755 redispached just as if there was no cookie.
12756

12757 U : a cookie was present but was not used to select the server because
12758 some other server selection mechanism was used instead (typically a
12759 "use-server" rule).
12760

12761 - : does not apply (no cookie set in configuration).
12762

12763 - the last character reports what operations were performed on the persistence
12764 cookie returned by the server (only in HTTP mode) :
12765

12766 N : NO cookie was provided by the server, and none was inserted either.
12767

12768 I : no cookie was provided by the server, and the proxy INSERTED one.
12769 Note that in "cookie insert" mode, if the server provides a cookie,
12770 it will still be overwritten and reported as "I" here.
12771

12772 U : the proxy UPDATED the last date in the cookie that was presented by
12773 the client. This can only happen in insert mode with "maxidle". It
12774 happens every time there is activity at a different date than the
12775 date indicated in the cookie. If any other change happens, such as
12776 a redispach, then the cookie will be marked as inserted instead.
12777

12778 P : a cookie was PROVIDED by the server and transmitted as-is.
12779

12780 R : the cookie provided by the server was REMWRITTEN by the proxy, which
12781 happens in "cookie rewrite" or "cookie prefix" modes.
12782

12783 D : the cookie provided by the server was DELETED by the proxy.
12784

12785 - : does not apply (no cookie set in configuration).
12786

12787 The combination of the two first flags gives a lot of information about what
12788 was happening when the session terminated, and why it did terminate. It can be
12789 helpful to detect server saturation, network troubles, local system resource
12790 starvation, attacks, etc...
12791

12792 The most common termination flags combinations are indicated below. They are
12793 alphabetically sorted, with the lowercase set just after the upper case for
12794 easier finding and understanding.
12795

12796 Flags Reason

12797 -- Normal termination.
12798

12799 CC The client aborted before the connection could be established to the
12800 server. This can happen when haproxy tries to connect to a recently
12801 dead (or unchecked) server, and the client aborts while haproxy is
12802 waiting for the server to respond or for "timeout connect" to expire.
12803

12804 CD The client unexpectedly aborted during data transfer. This can be
12805

12806 caused by a browser crash, by an intermediate equipment between the
12807 client and haproxy which decided to actively break the connection,
12808 by network routing issues between the client and haproxy, or by a
12809 keep-alive session between the server and the client terminated first
12810 by the client.
12811

12812 CD The client did not send nor acknowledge any data for as long as the
12813 "timeout client" delay. This is often caused by network failures on
12814 the client side, or the client simply leaving the net uncleanly.
12815

12816 CH The client aborted while waiting for the server to start responding.
12817 It might be the server taking too long to respond or the client
12818 clicking the 'Stop' button too fast.
12819

12820 cH The "timeout client" stroke while waiting for client data during a
12821 POST request. This is sometimes caused by too large TCP MSS values
12822 for PPoE networks which cannot transport full-sized packets. It can
12823 also happen when client timeout is smaller than server timeout and
12824 the server takes too long to respond.
12825

12826 CQ The client aborted while its session was queued, waiting for a server
12827 with enough empty slots to accept it. It might be that either all the
12828 servers were saturated or that the assigned server was taking too
12829 long a time to respond.
12830

12831 CR The client aborted before sending a full HTTP request. Most likely
12832 the request was typed by hand using a telnet client, and aborted
12833 too early. The HTTP status code is likely a 400 here. Sometimes this
12834 might also be caused by an IDS killing the connection between haproxy
12835 and the client. "option http-ignore-probes" can be used to ignore
12836 connections without any data transfer.
12837

12838 cR The "timeout http-request" stroke before the client sent a full HTTP
12839 request. This is sometimes caused by too large TCP MSS values on the
12840 client side for PPoE networks which cannot transport full-sized
12841 packets, or by clients sending requests by hand and not typing fast
12842 enough, or forgetting to enter the empty line at the end of the
12843 request. The HTTP status code is likely a 408 here. Note: recently,
12844 some browsers started to implement a "pre-connect" feature consisting
12845 in speculatively connecting to some recently visited web sites just
12846 in case the user would like to visit them. This results in many
12847 connections being established to web sites, which end up in 408
12848 Request Timeout if the timeout strikes first, or 400 Bad Request when
12849 the browser decides to close them first. These ones pollute the log
12850 and feed the error counters. Some versions of some browsers have even
12851 been reported to display the error code. It is possible to work
12852 around the undesirable effects of this behaviour by adding "option
12853 http-ignore-probes" in the frontend, resulting in connections with
12854 zero data transfer to be totally ignored. This will definitely hide
12855 the errors of people experiencing connectivity issues though.
12856

12857 CT The client aborted while its session was tarpitted. It is important to
12858 check if this happens on valid requests, in order to be sure that no
12859 wrong tarpit rules have been written. If a lot of them happen, it
12860 might make sense to lower the "timeout tarpit" value to something
12861 closer to the average reported "Tw" timer, in order not to consume
12862 resources for just a few attackers.
12863

12864 LR The request was intercepted and locally handled by haproxy. Generally
12865 it means that this was a redirect or a stats request.
12866

12867 SC The server or an equipment between it and haproxy explicitly refused
12868 the TCP connection (the proxy received a TCP RST or an ICMP message
12869 in return). Under some circumstances, it can also be the network
12870 stack telling the proxy that the server is unreachable (eg: no route,

12871 or no ARP response on local network). When this happens in HTTP mode,
12872 the status code is likely a 502 or 503 here.
12873

12874 sC The "timeout connect" stroke before a connection to the server could
12875 complete. When this happens in HTTP mode, the status code is likely a
12876 503 or 504 here.
12877

12878 SD The connection to the server died with an error during the data
12879 transfer. This usually means that haproxy has received an RST from
12880 the server or an ICMP message from an intermediate equipment while
12881 exchanging data with the server. This can be caused by a server crash
12882 or by a network issue on an intermediate equipment.
12883

12884 sD The server did not send nor acknowledge any data for as long as the
12885 "timeout server" setting during the data phase. This is often caused
12886 by too short timeouts on L4 equipments before the server (firewalls,
12887 load-balancers, ...), as well as keep-alive sessions maintained
12888 between the client and the server expiring first on haproxy.
12889

12890 SH The server aborted before sending its full HTTP response headers, or
12891 it crashed while processing the request. Since a server aborting at
12892 this moment is very rare, it would be wise to inspect its logs to
12893 control whether it crashed and why. The logged request may indicate a
12894 small set of faulty requests, demonstrating bugs in the application.
12895 Sometimes this might also be caused by an IDS killing the connection
12896 between haproxy and the server.
12897

12898 sH The "timeout server" stroke before the server could return its
12899 response headers. This is the most common anomaly, indicating too
12900 long transactions, probably caused by server or database saturation.
12901 The immediate workaround consists in increasing the "timeout server"
12902 setting, but it is important to keep in mind that the user experience
12903 will suffer from these long response times. The only long term
12904 solution is to fix the application.
12905

12906 sQ The session spent too much time in queue and has been expired. See
12907 the "timeout queue" and "timeout connect" settings to find out how to
12908 fix this if it happens too often. If it often happens massively in
12909 short periods, it may indicate general problems on the affected
12910 servers due to I/O or database congestion, or saturation caused by
12911 external attacks.
12912

12913 PC The proxy refused to establish a connection to the server because the
12914 process' socket limit has been reached while attempting to connect.
12915 The global "maxconn" parameter may be increased in the configuration
12916 so that it does not happen anymore. This status is very rare and
12917 might happen when the global "ulimit-n" parameter is forced by hand.
12918

12919 PD The proxy blocked an incorrectly formatted chunked encoded message in
12920 a request or a response, after the server has emitted its headers. In
12921 most cases, this will indicate an invalid message from the server to
12922 the client. Haproxy supports chunk sizes of up to 2GB - 1 (2147483647
12923 bytes). Any larger size will be considered as an error.
12924

12925 PH The proxy blocked the server's response, because it was invalid,
12926 incomplete, dangerous (cache control), or matched a security filter.
12927 In any case, an HTTP 502 error is sent to the client. One possible
12928 cause for this error is an invalid syntax in an HTTP header name
12929 containing unauthorized characters. It is also possible but quite
12930 rare, that the proxy blocked a chunked-encoding request from the
12931 client due to an invalid syntax, before the server responded. In this
12932 case, an HTTP 400 error is sent to the client and reported in the
12933 logs.
12934

12935 PR The proxy blocked the client's HTTP request, either because of an

12936 invalid HTTP syntax, in which case it returned an HTTP 400 error to
12937 the client, or because a deny filter matched, in which case it
12938 returned an HTTP 403 error.
12939

12940 PT The proxy blocked the client's request and has tarptitted its
12941 connection before returning it a 500 server error. Nothing was sent
12942 to the server. The connection was maintained open for as long as
12943 reported by the "Tw" timer field.
12944

12945 RC A local resource has been exhausted (memory, sockets, source ports)
12946 preventing the connection to the server from establishing. The error
12947 logs will tell precisely what was missing. This is very rare and can
12948 only be solved by proper system tuning.
12949

12950 The combination of the two last flags gives a lot of information about how
12951 persistence was handled by the client, the server and by haproxy. This is very
12952 important to troubleshoot disconnections, when users complain they have to
12953 re-authenticate. The commonly encountered flags are :

12954 -- Persistence cookie is not enabled.

12955 NN No cookie was provided by the client, none was inserted in the
12956 response. For instance, this can be in insert mode with "postonly"
12957 set on a GET request.
12958

12959 II A cookie designating an invalid server was provided by the client,
12960 a valid one was inserted in the response. This typically happens when
12961 a "server" entry is removed from the configuration, since its cookie
12962 value can be presented by a client when no other server knows it.
12963

12964 NI No cookie was provided by the client, one was inserted in the
12965 response. This typically happens for first requests from every user
12966 in "insert" mode, which makes it an easy way to count real users.
12967

12968 VN A cookie was provided by the client, none was inserted in the
12969 response. This happens for most responses for which the client has
12970 already got a cookie.
12971

12972 VU A cookie was provided by the client, with a last visit date which is
12973 not completely up-to-date, so an updated cookie was provided in
12974 response. This can also happen if there was no date at all, or if
12975 there was a date but the "maxidle" parameter was not set, so that the
12976 cookie can be switched to unlimited time.
12977

12978 EI A cookie was provided by the client, with a last visit date which is
12979 too old for the "maxidle" parameter, so the cookie was ignored and a
12980 new cookie was inserted in the response.
12981

12982 OI A cookie was provided by the client, with a first visit date which is
12983 too old for the "maxlife" parameter, so the cookie was ignored and a
12984 new cookie was inserted in the response.
12985

12986 DI The server designated by the cookie was down, a new server was
12987 selected and a new cookie was emitted in the response.
12988

12989 VI The server designated by the cookie was not marked dead but could not
12990 be reached. A redispatch happened and selected another one, which was
12991 then advertised in the response.
12992

8.6. Non-printable characters

12996 In order not to cause trouble to log analysis tools or terminals during log
12997 consulting, non-printable characters are not sent as-is into log files, but are
12998
12999
13000

converted to the two-digits hexadecimal representation of their ASCII code, prefixed by the character '#'. The only characters that can be logged without being escaped are comprised between 32 and 126 (inclusive). Obviously, the escape character '#' itself is also encoded to avoid any ambiguity ("23"). It is the same for the character '"' which becomes "22", as well as '{', '|' and '}' when logging headers.

Note that the space character (' ') is not encoded in headers, which can cause issues for tools relying on space count to locate fields. A typical header containing spaces is "User-Agent".

Last, it has been observed that some syslog daemons such as syslog-ng escape the quote ('') with a backslash ('\'). The reverse operation can safely be performed since no quote may appear anywhere else in the logs.

8.7. Capturing HTTP cookies

Cookie capture simplifies the tracking a complete user session. This can be achieved using the "capture cookie" statement in the frontend. Please refer to section 4.2 for more details. Only one cookie can be captured, and the same cookie will simultaneously be checked in the request ("Cookie:" header) and in the response ("Set-Cookie:" header). The respective values will be reported in the HTTP logs at the "captured_request_cookie" and "captured_response_cookie" locations (see section 8.2.3 about HTTP log format). When either cookie is not seen, a dash ('-') replaces the value. This way, it's easy to detect when a user switches to a new session for example, because the server will reassign it a new cookie. It is also possible to detect if a server unexpectedly sets a wrong cookie to a client, leading to session crossing.

Examples :

```
# capture the first cookie whose name starts with "ASPSESSION"
capture cookie ASPSESSION len 32

# capture the first cookie whose name is exactly "vgnvisitor"
capture cookie vgnvisitor= len 32
```

8.8. Capturing HTTP headers

Header captures are useful to track unique request identifiers set by an upper proxy, virtual host names, user-agents, POST content-length, referrers, etc. In the response, one can search for information about the response length, how the server asked the cache to behave, or an object location during a redirection.

Header captures are performed using the "capture request header" and "capture response header" statements in the frontend. Please consult their definition in section 4.2 for more details.

It is possible to include both request headers and response headers at the same time. Non-existent headers are logged as empty strings, and if one header appears more than once, only its last occurrence will be logged. Request headers are grouped within braces '{' and '}' in the same order as they were declared, and delimited with a vertical bar '|', without any space. Response headers follow the same representation, but are displayed after a space following the request headers block. These blocks are displayed just before the HTTP request in the logs.

As a special case, it is possible to specify an HTTP header capture in a TCP frontend. The purpose is to enable logging of headers which will be parsed in an HTTP backend if the request is then switched to this HTTP backend.

Example :

```
# This instance chains to the outgoing proxy
listen proxy-out
mode http
option httplog
option logasap
log global
server cache1 192.168.1.1:3128

# log the name of the virtual server
capture request header Host len 20

# log the amount of data uploaded during a POST
capture request header Content-Length len 10

# log the beginning of the referrer
capture request header Referer len 20

# server name (useful for outgoing proxies only)
capture response header Server len 20

# logging the content-length is useful with "option logasap"
capture response header Content-Length len 10

# log the expected cache behaviour on the response
capture response header Cache-Control len 8

# the Via header will report the next proxy's name
capture response header Via len 20

# log the URL location during a redirection
capture response header Location len 20
```

```
>>> Aug 9 20:26:09 localhost \
haproxy[2022]: 127.0.0.1:34014 [09/Aug/2004:20:26:09] proxy-out \
proxy-out/cache1 0/0/0/162/+162 200 - - ---- 0/0/0/0/0 0/0 \
{fr.adserver.yahoo.co}|http://fr.f416.mail.} {1864|private|} \
"GET http://fr.adserver.yahoo.com/"
```

```
>>> Aug 9 20:30:46 localhost \
haproxy[2022]: 127.0.0.1:34020 [09/Aug/2004:20:30:46] proxy-out \
proxy-out/cache1 0/0/0/182/+182 200 +279 - - ---- 0/0/0/0/0 0/0 \
{w.ods.org|} {Formilux/0.1.8|3495|}| \
"GET http://trafic.lwt.eu/ HTTP/1.1"
```

```
>>> Aug 9 20:30:46 localhost \
haproxy[2022]: 127.0.0.1:34028 [09/Aug/2004:20:30:46] proxy-out \
proxy-out/cache1 0/0/2/126/+128 301 +223 - - ---- 0/0/0/0/0 0/0 \
{www.sytadin.equipement.gouv.fr}|http://trafic.lwt.eu/} \
{Apache|230|}|http://www.sytadin.} \
"GET http://www.sytadin.equipement.gouv.fr/ HTTP/1.1"
```

8.9. Examples of logs

These are real-world examples of logs accompanied with an explanation. Some of them have been made up by hand. The syslog part has been removed for better reading. Their sole purpose is to explain how to decipher them.

```
>>> haproxy[674]: 127.0.0.1:33318 [15/Oct/2003:08:31:57.130] px-http \
px-http/srv1 6559/0/7/147/6723 200 243 - - ---- 5/3/3/1/0 0/0 \
"HEAD / HTTP/1.0"
```

```
=> long request (6.5s) entered by hand through 'telnet'. The server replied
in 147 ms, and the session ended normally ('-----')
```

```
13131 >>> haproxy[674]: 127.0.0.1:33319 [15/Oct/2003:08:31:57.149] px-http \
13132 px-http/srv1 6559/1230/7/147/6870 200 243 - - ---- 324/239/239/99/0 \
13133 0/9 "HEAD / HTTP/1.0"
13134
13135 => Idem, but the request was queued in the global queue behind 9 other
13136 requests, and waited there for 1230 ms.
13137
13138 >>> haproxy[674]: 127.0.0.1:33320 [15/Oct/2003:08:32:17.654] px-http \
13139 px-http/srv1 9/0/7/14/+30 200 +243 - - ---- 3/3/3/1/0 0/0 \
13140 "GET /image.iso HTTP/1.0"
13141
13142 => request for a long data transfer. The "logasap" option was specified, so
13143 the log was produced just before transferring data. The server replied in
13144 14 ms, 243 bytes of headers were sent to the client, and total time from
13145 accept to first data byte is 30 ms.
13146
13147 >>> haproxy[674]: 127.0.0.1:33320 [15/Oct/2003:08:32:17.925] px-http \
13148 px-http/srv1 9/0/7/14/30 502 243 - - PH-- 3/2/2/0/0 0/0 \
13149 "GET /cgi-bin/bug.cgi? HTTP/1.0"
13150
13151 => the proxy blocked a server response either because of an "rspdeny" or
13152 "rspidenvy" filter, or because the response was improperly formatted and
13153 not HTTP-compliant, or because it blocked sensitive information which
13154 risked being cached. In this case, the response is replaced with a "502
13155 bad gateway". The flags ("PH--") tell us that it was haproxy who decided
13156 to return the 502 and not the server.
13157
13158 >>> haproxy[18113]: 127.0.0.1:34548 [15/Oct/2003:15:18:55.798] px-http \
13159 px-http/<N0SRV> -1/-1/-1/-1/8490 -1 0 - - CR-- 2/2/2/0/0 0/0 ""
13160
13161 => the client never completed its request and aborted itself ("C---") after
13162 8.5s, while the proxy was waiting for the request headers ("R--").
13163 Nothing was sent to any server.
13164
13165 >>> haproxy[18113]: 127.0.0.1:34549 [15/Oct/2003:15:19:06.103] px-http \
13166 px-http/<N0SRV> -1/-1/-1/-1/50001 408 0 - - cR-- 2/2/2/0/0 0/0 ""
13167
13168 => The client never completed its request, which was aborted by the
13169 time-out ("C---") after 50s, while the proxy was waiting for the request
13170 headers ("R--"). Nothing was sent to any server, but the proxy could
13171 send a 408 return code to the client.
13172
13173 >>> haproxy[18989]: 127.0.0.1:34550 [15/Oct/2003:15:24:28.312] px-tcp \
13174 px-tcp/srv1 0/0/5007 0 cD 0/0/0/0/0 0/0
13175
13176 => This log was produced with "option tcplog". The client timed out after
13177 5 seconds ("C-----").
13178
13179 >>> haproxy[18989]: 10.0.0.1:34552 [15/Oct/2003:15:26:31.462] px-http \
13180 px-http/srv1 3183/-1/-1/-1/11215 503 0 - - SC-- 205/202/202/115/3 \
13181 0/0 "HEAD / HTTP/1.0"
13182
13183 => The request took 3s to complete (probably a network problem), and the
13184 connection to the server failed ('SC-') after 4 attempts of 2 seconds
13185 (config says 'retries 3'), and no redispatch (otherwise we would have
13186 seen '+/3'). Status code 503 was returned to the client. There were 115
13187 connections on this server, 202 connections on this proxy, and 205 on
13188 the global process. It is possible that the server refused the
13189 connection because of too many already established.
13190
13191
13192
13193
13194
13195
```

9. Statistics and monitoring

```
13196 It is possible to query HAProxy about its status. The most commonly used
13197 mechanism is the HTTP statistics page. This page also exposes an alternative
13198 CSV output format for monitoring tools. The same format is provided on the
13199 Unix socket.
13200
13201
13202
13203
13204
13205
13206
13207
13208
13209
13210
13211
13212
13213
13214
13215
13216
13217
13218
13219
13220
13221
13222
13223
13224
13225
13226
13227
13228
13229
13230
13231
13232
13233
13234
13235
13236
13237
13238
13239
13240
13241
13242
13243
13244
13245
13246
13247
13248
13249
13250
13251
13252
13253
13254
13255
13256
13257
13258
13259
13260
```

The statistics may be consulted either from the unix socket or from the HTTP page. Both means provide a CSV format whose fields follow. The first line begins with a sharp ('#') and has one word per comma-delimited field which represents the title of the column. All other lines starting at the second one use a classical CSV format using a comma as the delimiter, and the double quote (") as an optional text delimiter, but only if the enclosed text is ambiguous (if it contains a quote or a comma). The double-quote character (") in the text is doubled (""), which is the format that most tools recognize. Please do not insert any column before these ones in order not to break tools which use hard-coded column positions.

In brackets after each field name are the types which may have a value for that field. The types are L (Listeners), F (Frontends), B (Backends), and S (Servers).

0. pxname [LFBS]: proxy name
1. svname [LFBS]: service name (FRONTEND for frontend, BACKEND for backend, any name for server/listener)
2. qcur [..BS]: current queued requests. For the backend this reports the number queued without a server assigned.
3. qmax [..BS]: max value of qcur
4. scur [LFBS]: current sessions
5. smax [LFBS]: max sessions
6. slim [LFBS]: configured session limit
7. stot [LFBS]: cumulative number of connections
8. bin [LFBS]: bytes in
9. bout [LFBS]: bytes out
10. dreq [LFB..]: requests denied because of security concerns.
 - For tcp this is because of a matched tcp-request content rule.
 - For http this is because of a matched http-request or tarpit rule.
11. dresp [LFBS]: responses denied because of security concerns.
 - For http this is because of a matched http-request rule, or "option checkcache".
12. ereq [LF..]: request errors. Some of the possible causes are:
 - early termination from the client, before the request has been sent.
 - read error from the client
 - client timeout
 - client closed connection
 - various bad requests from the client.
 - request was tarpitted.
13. econ [..BS]: number of requests that encountered an error trying to connect to a backend server. The backend stat is the sum of the stat for all servers of that backend, plus any connection errors not associated with a particular server (such as the backend having no active servers).
14. eresp [..BS]: response errors. srv_abrt will be counted here also. Some other errors are:
 - write error on the client socket (won't be counted for the server stat)
 - failure applying filters to the response.
15. wrctr [..BS]: number of times a connection to a server was retried.
16. wredis [..BS]: number of times a request was redispatched to another server. The server value counts the number of times that server was switched away from.
17. status [LFBS]: status (UP/DOWN/NOLB/MAINT(via)..)
18. weight [..BS]: total weight (backend), server weight (server)
19. act [..BS]: number of active servers (backend), server is active (server)

```
13261 20. bck [..BS]: number of backup servers (backend), server is backup (server)
13262 21. chkfail [..S]: number of failed checks. (Only counts checks failed when
13263 the server is up.)
13264 22. chkdown [..BS]: number of UP->DOWN transitions. The backend counter counts
13265 transitions to the whole backend being down, rather than the sum of the
13266 counters for each server.
13267 23. lastchg [..BS]: number of seconds since the last UP->DOWN transition
13268 24. downtime [..BS]: total downtime (in seconds). The value for the backend
13269 is the downtime for the whole backend, not the sum of the server downtime.
13270 25. qlimit [..S]: configured maxqueue for the server, or nothing in the
13271 value is 0 (default, meaning no limit)
13272 26. pid [LFBS]: process id (0 for first instance, 1 for second, ...)
13273 27. iid [LFBS]: unique proxy id
13274 28. sid [L..S]: server id (unique inside a proxy)
13275 29. throttle [..S]: current throttle percentage for the server, when
13276 slowstart is active, or no value if not in slowstart.
13277 30. lbtot [..BS]: total number of times a server was selected, either for new
13278 sessions, or when re-dispatching. The server counter is the number
13279 of times that server was selected.
13280 31. tracked [..S]: id of proxy/server if tracking is enabled.
13281 32. type [LFBS]: (0=frontend, 1=backend, 2=server, 3=socket/listener)
13282 33. rate [LFBS]: number of sessions per second over last elapsed second
13283 34. rate_lim [F..]: configured limit on new sessions per second
13284 35. rate_max [LFBS]: max number of new sessions per second
13285 36. check_status [..S]: status of last health check, one of:
13286 UNK -> unknown
13287 INIT -> initializing
13288 SOCKERR -> socket error
13289 L4OK -> check passed on layer 4, no upper layers testing enabled
13290 L4TOUIT -> layer 1-4 timeout
13291 L4CON -> layer 1-4 connection problem, for example
13292 L6OK -> check passed on layer 6
13293 L6TOUIT -> layer 6 (SSL) timeout
13294 L6RSP -> layer 6 invalid response - protocol error
13295 L7OK -> check passed on layer 7
13296 L7OKC -> check conditionally passed on layer 7, for example 404 with
13297 disable-on-404
13298 L7TOUIT -> layer 7 (HTTP/SMTP) timeout
13299 L7RSP -> layer 7 invalid response - protocol error
13300 L7STS -> layer 7 response error, for example HTTP 5xx
13301
13302 37. check_code [..S]: layer5-7 code, if available
13303 38. check_duration [..S]: time in ms took to finish last health check
13304 39. hrsp_lxx [FBS]: http responses with lxx code
13305 40. hrsp_2xx [FBS]: http responses with 2xx code
13306 41. hrsp_3xx [FBS]: http responses with 3xx code
13307 42. hrsp_4xx [FBS]: http responses with 4xx code
13308 43. hrsp_5xx [FBS]: http responses with 5xx code
13309 44. hrsp_other [FBS]: http responses with other codes (protocol error)
13310 45. hanafail [..S]: failed health checks details
13311 46. req_rate [F..]: HTTP requests per second over last elapsed second
13312 47. req_rate_max [F..]: max number of HTTP requests per second observed
13313 48. req_tot [F..]: total number of HTTP requests received
13314 49. cli_abrt [..BS]: number of data transfers aborted by the client
13315 50. srv_abrt [..BS]: number of data transfers aborted by the server
13316 (inc. in eresp)
13317 51. comp.in [F.B.]: number of HTTP response bytes fed to the compressor
13318 52. comp.out [F.B.]: number of HTTP response bytes emitted by the compressor
13319 53. comp_byp [F.B.]: number of bytes that bypassed the HTTP compressor
13320 (CPU/BW limit)
13321 54. comp_rsp [F.B.]: number of HTTP responses that were compressed
13322 55. lastsess [..BS]: number of seconds since last session assigned to
13323 server/backend
13324 56. last_chk [..S]: last health check contents or textual error
13325 57. last_agt [..S]: last agent check contents or textual error
```

```
13326 58. qtime [..BS]: the average queue time in ms over the 1024 last requests
13327 59. ctime [..BS]: the average connect time in ms over the 1024 last requests
13328 60. rtime [..BS]: the average response time in ms over the 1024 last requests
13329 (0 for TCP)
13330 61. ttime [..BS]: the average total session time in ms over the 1024 last
13331 requests
13332
13333 9.2. Unix Socket commands
13334 -----
13335
13336 The stats socket is not enabled by default. In order to enable it, it is
13337 necessary to add one line in the global section of the haproxy configuration.
13338 A second line is recommended to set a larger timeout, always appreciated when
13339 issuing commands by hand :
13340
13341 global
13342 stats socket /var/run/haproxy.sock mode 600 level admin
13343 stats timeout 2m
13344
13345 It is also possible to add multiple instances of the stats socket by repeating
13346 the line, and make them listen to a TCP port instead of a UNIX socket. This is
13347 never done by default because this is dangerous, but can be handy in some
13348 situations :
13349
13350 global
13351 stats socket /var/run/haproxy.sock mode 600 level admin
13352 stats socket ipv4@192.168.0.1:9999 level admin
13353 stats timeout 2m
13354
13355 To access the socket, an external utility such as "socat" is required. Socat is a
13356 swiss-army knife to connect anything to anything. We use it to connect terminals
13357 to the socket, or a couple of stdin/stdout pipes to it for scripts. The two main
13358 syntaxes we'll use are the following :
13359
13360 # socat /var/run/haproxy.sock stdio
13361 # socat /var/run/haproxy.sock readline
13362
13363 The first one is used with scripts. It is possible to send the output of a
13364 script to haproxy, and pass haproxy's output to another script. That's useful
13365 for retrieving counters or attack traces for example.
13366
13367 The second one is only useful for issuing commands by hand. It has the benefit
13368 that the terminal is handled by the readline library which supports line
13369 editing and history, which is very convenient when issuing repeated commands
13370 (eg: watch a counter).
13371
13372 The socket supports two operation modes :
13373 - interactive
13374 - non-interactive
13375
13376 The non-interactive mode is the default when socat connects to the socket. In
13377 this mode, a single line may be sent. It is processed as a whole, responses are
13378 sent back, and the connection closes after the end of the response. This is the
13379 mode that scripts and monitoring tools use. It is possible to send multiple
13380 commands in this mode, they need to be delimited by a semi-colon (;). For
13381 example :
13382
13383 # echo "show info;show stat;show table" | socat /var/run/haproxy stdio
13384
13385 The interactive mode displays a prompt ('>') and waits for commands to be
13386 entered on the line, then processes them, and displays the prompt again to wait
13387 for a new command. This mode is entered via the "prompt" command which must be
13388 sent on the first line in non-interactive mode. The mode is a flip switch, if
13389 "prompt" is sent in interactive mode, it is disabled and the connection closes
13390
```

13391 after processing the last command of the same line.

13392
13393 For this reason, when debugging by hand, it's quite common to start with the
13394 "prompt" command :

13395
13396 # socat /var/run/haproxy readline
13397 prompt
13398 > show info
13399 ...
13400 >

13401
13402 Since multiple commands may be issued at once, haproxy uses the empty line as a
13403 delimiter to mark an end of output for each command, and takes care of ensuring
13404 that no command can emit an empty line on output. A script can thus easily
13405 parse the output even when multiple commands were pipelined on a single line.

13406
13407 It is important to understand that when multiple haproxy processes are started
13408 on the same sockets, any process may pick up the request and will output its
13409 own stats.

13410
13411 The list of commands currently supported on the stats socket is provided below.
13412 If an unknown command is sent, haproxy displays the usage message which reminds
13413 all supported commands. Some commands support a more complex syntax, generally
13414 it will explain what part of the command is invalid when this happens.

13415 add acl <acl> <pattern>

13416 Add an entry into the acl <acl>. <acl> is the #<id> or the <file> returned by
13417 "show acl". This command does not verify if the entry already exists. This
13418 command cannot be used if the reference <acl> is a file also used with a map.
13419 In this case, you must use the command "add map" in place of "add acl".

13420 add map <map> <key> <value>

13421 Add an entry into the map <map> to associate the value <value> to the key
13422 <key>. This command does not verify if the entry already exists. It is
13423 mainly used to fill a map after a clear operation. Note that if the reference
13424 <map> is a file and is shared with a map, this map will contain also a new
13425 pattern entry.

13426 clear counters

13427 Clear the max values of the statistics counters in each proxy (frontend &
13428 backend) and in each server. The cumulated counters are not affected. This
13429 can be used to get clean counters after an incident, without having to
13430 restart nor to clear traffic counters. This command is restricted and can
13431 only be issued on sockets configured for levels "operator" or "admin".

13432 clear counters all

13433 Clear all statistics counters in each proxy (frontend & backend) and in each
13434 server. This has the same effect as restarting. This command is restricted
13435 and can only be issued on sockets configured for level "admin".

13436 clear acl <acl>

13437 Remove all entries from the acl <acl>. <acl> is the #<id> or the <file>
13438 returned by "show acl". Note that if the reference <acl> is a file and is
13439 shared with a map, this map will be also cleared.

13440 clear map <map>

13441 Remove all entries from the map <map>. <map> is the #<id> or the <file>
13442 returned by "show map". Note that if the reference <map> is a file and is
13443 shared with a acl, this acl will be also cleared.

13444 clear table <table> [data.<type> <operator> <value>] [key <key>]

13445 Remove entries from the stick-table <table>.

13450
13451 This is typically used to unblock some users complaining they have been
13452 abusively denied access to a service, but this can also be used to clear some

13456 stickiness entries matching a server that is going to be replaced (see "show
13457 table" below for details). Note that sometimes, removal of an entry will be
13458 refused because it is currently tracked by a session. Retrying a few seconds
13459 later after the session ends is usual enough.

13460
13461 In the case where no options arguments are given all entries will be removed.

13462
13463 When the "data." form is used entries matching a filter applied using the
13464 stored data (see "stick-table" in section 4.2) are removed. A stored data
13465 type must be specified in <type>, and this data type must be stored in the
13466 table otherwise an error is reported. The data is compared according to
13467 <operator> with the 64-bit integer <value>. Operators are the same as with
13468 the ACLs :

- 13469 - eq : match entries whose data is equal to this value
- 13470 - ne : match entries whose data is not equal to this value
- 13471 - le : match entries whose data is less than or equal to this value
- 13472 - ge : match entries whose data is greater than or equal to this value
- 13473 - lt : match entries whose data is less than this value
- 13474 - gt : match entries whose data is greater than this value

13475
13476 When the key form is used the entry <key> is removed. The key must be of the
13477 same type as the table, which currently is limited to IPv4, IPv6, integer and
13478 string.

13479 Example :
13480 \$ echo "show table http_proxy" | socat stdio /tmp/socket1

13481 >>> # table: http_proxy, type: ip, size:204800, used:2

13482 >>> 0x80e6a4c: key=127.0.0.1 use=0 exp=3594729 gpc0=0 conn_rate(30000)=1 \

13483 bytes_out_rate(60000)=187

13484 >>> 0x80e6a80: key=127.0.0.2 use=0 exp=3594740 gpc0=1 conn_rate(30000)=10 \

13485 bytes_out_rate(60000)=191

13486 \$ echo "clear table http_proxy key 127.0.0.1" | socat stdio /tmp/socket1

13487 >>> # table: http_proxy, type: ip, size:204800, used:1

13488 >>> 0x80e6a80: key=127.0.0.2 use=0 exp=3594740 gpc0=1 conn_rate(30000)=10 \

13489 bytes_out_rate(60000)=191

13490 \$ echo "show table http_proxy data.gpc0 eq 1" | socat stdio /tmp/socket1

13491 >>> # table: http_proxy, type: ip, size:204800, used:1

13492 del acl <acl> [-<key> #-<ref>]

13493 Delete all the acl entries from the acl <acl> corresponding to the key <key>.

13494 <acl> is the #<id> or the <file> returned by "show acl". If the <ref> is used,

13495 this command delete only the listed reference. The reference can be found with

13496 listing the content of the acl. Note that if the reference <acl> is a file and

13497 is shared with a map, the entry will be also deleted in the map.

13500 del map <map> [-<key> #-<ref>]

13501 Delete all the map entries from the map <map> corresponding to the key <key>.

13502 <map> is the #<id> or the <file> returned by "show map". If the <ref> is used,

13503 this command delete only the listed reference. The reference can be found with

13504 listing the content of the map. Note that if the reference <map> is a file and

13505 is shared with a acl, the entry will be also deleted in the map.

13506 disable agent <backend>[-<server>]

13507 Mark the auxiliary agent check as temporarily stopped.

13508
13509 In the case where an agent check is being run as a auxiliary check, due

13510 to the agent-check parameter of a server directive, new checks are only

13511 initialised when the agent is in the enabled. Thus, disable agent will

13512 prevent any new agent checks from begin initiated until the agent

13513 re-enabled using enable agent.

13514
13515

When an agent is disabled the processing of an auxiliary agent check that was initiated while the agent was set as enabled is as follows: All results that would alter the weight, specifically "drain" or a weight returned by the agent, are ignored. The processing of agent check is otherwise unchanged.

The motivation for this feature is to allow the weight changing effects of the agent checks to be paused to allow the weight of a server to be configured using set weight without being overridden by the agent.

This command is restricted and can only be issued on sockets configured for level "admin".

disable frontend <frontend>

Mark the frontend as temporarily stopped. This corresponds to the mode which is used during a soft restart : the frontend releases the port but can be enabled again if needed. This should be used with care as some non-Linux OSes are unable to enable it back. This is intended to be used in environments where stopping a proxy is not even imaginable but a misconfigured proxy must be fixed. That way it's possible to release the port and bind it into another process to restore operations. The frontend will appear with status "STOP" on the stats page.

The frontend may be specified either by its name or by its numeric ID, prefixed with a sharp ('#').

This command is restricted and can only be issued on sockets configured for level "admin".

disable health <backend>/<server>

Mark the primary health check as temporarily stopped. This will disable sending of health checks, and the last health check result will be ignored. The server will be in unchecked state and considered UP unless an auxiliary agent check forces it down.

This command is restricted and can only be issued on sockets configured for level "admin".

disable server <backend>/<server>

Mark the server DOWN for maintenance. In this mode, no more checks will be performed on the server until it leaves maintenance. If the server is tracked by other servers, those servers will be set to DOWN during the maintenance.

In the statistics page, a server DOWN for maintenance will appear with a "MAINT" status, its tracking servers with the "MAINT(via)" one.

Both the backend and the server may be specified either by their name or by their numeric ID, prefixed with a sharp ('#').

This command is restricted and can only be issued on sockets configured for level "admin".

enable agent <backend>/<server>
Resume auxiliary agent check that was temporarily stopped.

See "disable agent" for details of the effect of temporarily starting and stopping an auxiliary agent.

This command is restricted and can only be issued on sockets configured for level "admin".

enable frontend <frontend>

Resume a frontend which was temporarily stopped. It is possible that some of

the listening ports won't be able to bind anymore (eg: if another process took them since the 'disable frontend' operation). If this happens, an error is displayed. Some operating systems might not be able to resume a frontend which was disabled.

The frontend may be specified either by its name or by its numeric ID, prefixed with a sharp ('#').

This command is restricted and can only be issued on sockets configured for level "admin".

enable health <backend>/<server>

Resume a primary health check that was temporarily stopped. This will enable sending of health checks again. Please see "disable health" for details.

This command is restricted and can only be issued on sockets configured for level "admin".

enable server <backend>/<server>

If the server was previously marked as DOWN for maintenance, this marks the server UP and checks are re-enabled.

Both the backend and the server may be specified either by their name or by their numeric ID, prefixed with a sharp ('#').

This command is restricted and can only be issued on sockets configured for level "admin".

get map <map> <value>

get acl <acl> <value>

Lookup the value <value> in the map <map> or in the ACL <acl>. <map> or <acl> are the #<id> or the <file> returned by "show map" or "show acl". This command returns all the matching patterns associated with this map. This is useful for debugging maps and ACLs. The output format is composed by one line per matching type. Each line is composed by space-delimited series of words.

The first two words are:

<match method>: The match method applied. It can be "found", "bool", "int", "ip", "bin", "len", "str", "beg", "sub", "dir", "dom", "end" or "reg".

<match result>: The result. Can be "match" or "no-match".

The following words are returned only if the pattern matches an entry.

<index type>: "tree" or "list". The internal lookup algorithm.

<cases>: "case-insensitive" or "case-sensitive". The interpretation of the case.

<entry matched>: match=<entry>. Return the matched pattern. It is useful with regular expressions.

The two last word are used to show the returned value and its type. With the "acl" case, the pattern doesn't exist.

return=nothing: No return because there are no "map".
return=<values>: The value returned in the string format.
return=cannot-display: The value cannot be converted as string.

type=<type>: The type of the returned sample.

get weight <backend>/<server>

Report the current weight and the initial weight of server <server> in

13650

13651 backend <backend> or an error if either doesn't exist. The initial weight is
13652 the one that appears in the configuration file. Both are normally equal
13653 unless the current weight has been changed. Both the backend and the server
13654 may be specified either by their name or by their numeric ID, prefixed with a
13655 sharp ('#').

13656 help

13657 Print the list of known keywords and their basic usage. The same help screen
13658 is also displayed for unknown commands.

13660 prompt

13661 Toggle the prompt at the beginning of the line and enter or leave interactive
13662 mode. In interactive mode, the connection is not closed after a command
13663 completes. Instead, the prompt will appear again, indicating the user that
13664 the interpreter is waiting for a new command. The prompt consists in a right
13665 angle bracket followed by a space "> ". This mode is particularly convenient
13666 when one wants to periodically check information such as stats or errors.
13667 It is also a good idea to enter interactive mode before issuing a "help"
13668 command.

13670 quit

13671 Close the connection when in interactive mode.

13673 set map <map> [<key>|<ref>] <value>

13674 Modify the value corresponding to each key <key> in a map <map>. <map> is the
13675 #<id> or <file> returned by "show map". If the <ref> is used in place of
13676 <key>, only the entry pointed by <ref> is changed. The new value is <value>.

13678 set maxconn frontend <frontend> <value>

13679 Dynamically change the specified frontend's maxconn setting. Any positive
13680 value is allowed including zero, but setting values larger than the global
13681 maxconn does not make much sense. If the limit is increased and connections
13682 were pending, they will immediately be accepted. If it is lowered to a value
13683 below the current number of connections, new connections acceptance will be
13684 delayed until the threshold is reached. The frontend might be specified by
13685 either its name or its numeric ID prefixed with a sharp ('#').

13687 set maxconn global <maxconn>

13688 Dynamically change the global maxconn setting within the range defined by the
13689 initial global maxconn setting. If it is increased and connections were
13690 pending, they will immediately be accepted. If it is lowered to a value below
13691 the current number of connections, new connections acceptance will be
13692 delayed until the threshold is reached. A value of zero restores the initial
13693 setting.

13695 set rate-limit connections global <value>

13696 Change the process-wide connection rate limit, which is set by the global
13697 'maxconnrate' setting. A value of zero disables the limitation. This limit
13698 applies to all frontends and the change has an immediate effect. The value
13699 is passed in number of connections per second.

13701 set rate-limit http-compression global <value>

13702 Change the maximum input compression rate, which is set by the global
13703 'maxcompress' setting. A value of zero disables the limitation. The value is
13704 passed in number of kilobytes per second. The value is available in the "show
13705 info" on the line "CompressBpsRateLim" in bytes.

13707 set rate-limit sessions global <value>

13708 Change the process-wide session rate limit, which is set by the global
13709 'maxessrate' setting. A value of zero disables the limitation. This limit
13710 applies to all frontends and the change has an immediate effect. The value
13711 is passed in number of sessions per second.

13713 set rate-limit ssl-sessions global <value>

13714 Change the process-wide SSL session rate limit, which is set by the global

13716 'maxsslrate' setting. A value of zero disables the limitation. This limit
13717 applies to all frontends and the change has an immediate effect. The value
13718 is passed in number of sessions per second sent to the SSL stack. It applies
13719 before the handshake in order to protect the stack against handshake abuses.

13720 set server <backend>/<server> agent [up | down]

13721 Force a server's agent to a new state. This can be useful to immediately
13722 switch a server's state regardless of some slow agent checks for example.
13723 Note that the change is propagated to tracking servers if any.

13725 set server <backend>/<server> health [up | stopping | down]

13726 Force a server's health to a new state. This can be useful to immediately
13727 switch a server's state regardless of some slow health checks for example.
13728 Note that the change is propagated to tracking servers if any.

13730 set server <backend>/<server> state [ready | drain | maint]

13731 Force a server's administrative state to a new state. This can be useful to
13732 disable load balancing and/or any traffic to a server. Setting the state to
13733 "ready" puts the server in normal mode, and the command is the equivalent of
13734 the "enable server" command. Setting the state to "maint" disables any traffic
13735 to the server as well as any health checks. This is the equivalent of the
13736 "disable server" command. Setting the mode to "drain" only removes the server
13737 from load balancing but still allows it to be checked and to accept new
13738 persistent connections. Changes are propagated to tracking servers if any.

13739 set server <backend>/<server> weight <weight>[%]

13740 Change a server's weight to the value passed in argument. This is the exact
13741 equivalent of the "set weight" command below.

13744 set ssl ocsdp-response <response>

13745 This command is used to update an OSCP Response for a certificate (see "crt"
13746 on "bind" lines). Same controls are performed as during the initial loading of
13747 the response. The <response> must be passed as a base64 encoded string of the
13748 DER encoded response from the OSCP server.

13750 Example:

```
13751 openssl ocsdp -issuer issuer.pem -cert server.pem \  
13752 -host ocsdp.issuer.com:80 -respout resp.der" | \  
13753 echo "set ssl ocsdp-response $(base64 -w 10000 resp.der)" | \  
13754 socat stdio /var/run/haproxy.stat
```

13756 set table <table> key <key> [data.<data_type> <value>]*

13757 Create or update a stick-table entry in the table. If the key is not present,
13758 an entry is inserted. See stick-table in section 4.2 to find all possible
13759 values for <data_type>. The most likely use consists in dynamically entering
13760 entries for source IP addresses, with a flag in gpc0 to dynamically block an
13761 IP address or affect its quality of service. It is possible to pass multiple
13762 data_types in a single call.

13764 set timeout cli <delay>

13765 Change the CLI interface timeout for current connection. This can be useful
13766 during long debugging sessions where the user needs to constantly inspect
13767 some indicators without being disconnected. The delay is passed in seconds.

13770 set weight <backend>/<server> <weight>[%]

13771 Change a server's weight to the value passed in argument. If the value ends
13772 with the '%' sign, then the new weight will be relative to the initially
13773 configured weight. Absolute weights are permitted between 0 and 256.
13774 Relative weights must be positive with the resulting absolute weight is
13775 capped at 256. Servers which are part of a farm running a static
13776 load-balancing algorithm have stricter limitations because the weight
13777 cannot change once set. Thus for these servers, the only accepted values
13778 are 0 and 100% (or 0 and the initial weight). Changes take effect
13779 immediately, though certain LB algorithms require a certain amount of
13780 requests to consider changes. A typical usage of this command is to

disable a server during an update by setting its weight to zero, then to enable it again after the update by setting it back to 100%. This command is restricted and can only be issued on sockets configured for level "admin". Both the backend and the server may be specified either by their name or by their numeric ID, prefixed with a sharp ('#').

show errors [<id>]

Dump last known request and response errors collected by frontends and backends. If <id> is specified, the limit the dump to errors concerning either frontend or backend whose ID is <id>. This command is restricted and can only be issued on sockets configured for levels "operator" or "admin".

The errors which may be collected are the last request and response errors caused by protocol violations, often due to invalid characters in header names. The report precisely indicates what exact character violated the protocol. Other important information such as the exact date the error was detected, frontend and backend names, the server name (when known), the internal session ID and the source address which has initiated the session are reported too.

All characters are returned, and non-printable characters are encoded. The most common ones (\t = 9, \n = 10, \r = 13 and \e = 27) are encoded as one letter following a backslash. The backslash itself is encoded as '\\' to avoid confusion. Other non-printable characters are encoded '\xNN' where NN is the two-digits hexadecimal representation of the character's ASCII code.

Lines are prefixed with the position of their first character, starting at 0 for the beginning of the buffer. At most one input line is printed per line, and large lines will be broken into multiple consecutive output lines so that the output never goes beyond 79 characters wide. It is easy to detect if a line was broken, because it will not end with '\n' and the next line's offset will be followed by a '+' sign, indicating it is a continuation of previous line.

Example :

```
$ echo "show errors" | socat stdio /tmp/sock1
>>> [04/Mar/2009:15:46:56.081] backend http-in (#2) : invalid response
src 127.0.0.1, session #54, frontend fe-eth0 (#1), server s2 (#1)
response length 213 bytes, error at position 23:
```

```
00000 HTTP/1.0 200 OK\r\n
00017 header/bizarre:blah\r\n
00038 Location: blah\r\n
00054 Long-line: this is a very long line which should b
00104+ e broken into multiple lines on the output buffer,
00154+ otherwise it would be too large to print in a ter
00204+ minal\r\n
00211 \r\n
```

In the example above, we see that the backend "http-in" which has internal ID 2 has blocked an invalid response from its server s2 which has internal ID 1. The request was on session 54 initiated by source 127.0.0.1 and received by frontend fe-eth0 whose ID is 1. The total response length was 213 bytes when the error was detected, and the error was at byte 23. This is the slash '/' in header name "header/bizarre", which is not a valid HTTP character for a header name.

show info

Dump info about haproxy status on current process.

show map [<map>]

Dump info about map converters. Without argument, the list of all available maps is returned. If a <map> is specified, its contents are dumped. <map> is

the #<id> or <file>. The first column is a unique identifier. It can be used as reference for the operation "del map" and "set map". The second column is the pattern and the third column is the sample if available. The data returned are not directly a list of available maps, but are the list of all patterns composing any map. Many of these patterns can be shared with ACL.

show acl [<acl>]

Dump info about acl converters. Without argument, the list of all available acls is returned. If a <acl> is specified, its contents are dumped. <acl> if the #<id> or <file>. The dump format is the same than the map even for the sample value. The data returned are not a list of available ACL, but are the list of all patterns composing any ACL. Many of these patterns can be shared with maps.

show pools

Dump the status of internal memory pools. This is useful to track memory usage when suspecting a memory leak for example. It does exactly the same as the SIGQUIT when running in foreground except that it does not flush the pools.

show sess

Dump all known sessions. Avoid doing this on slow connections as this can be huge. This command is restricted and can only be issued on sockets configured for levels "operator" or "admin".

show sess <id>

Display a lot of internal information about the specified session identifier. This identifier is the first field at the beginning of the lines in the dumps of "show sess" (it corresponds to the session pointer). Those information are useless to most users but may be used by haproxy developers to troubleshoot a complex bug. The output format is intentionally not documented so that it can freely evolve depending on demands. You may find a description of all fields returned in src/dumpstats.c

The special id "all" dumps the states of all sessions, which must be avoided as much as possible as it is highly CPU intensive and can take a lot of time.

show stat [<iid> <type> <sid>]

Dump statistics in the CSV format. By passing <id>, <type> and <sid>, it is possible to dump only selected items :

- <iid> is a proxy ID, -1 to dump everything

- <type> selects the type of dumpable objects : 1 for frontends, 2 for backends, 4 for servers, -1 for everything. These values can be 0Red, for example:

```
1 + 2 = 3 -> frontend + backend.
1 + 2 + 4 = 7 -> frontend + backend + server.
```

- <sid> is a server ID, -1 to dump everything from the selected proxy.

Example :

```
$ echo "show info;show stat" | socat stdio unix-connect:/tmp/sock1
>>> Name: HAProxy
Version: 1.4-dev2.49
Release_date: 2009/09/23
Nbproc: 1
Process_num: 1
(...)
# pxname,svname,qcur,qmax,scur,smax,slim,stat,bout,dreq, (...)
stats.FRONTEND,,,0,0,1000,0,0,0,0,0,,,,,OPEN,,,,,,,,,1,1,0, (...)
stats.BACKEND,0,0,0,1000,0,0,0,0,0,0,0,0,0,0,0,UP,0,0,0,250, (...)
(...)
www1.BACKEND,0,0,0,1000,0,0,0,0,0,0,0,0,0,0,0,UP,1,1,0,0,250, (...)
$
```

Here, two commands have been issued at once. That way it's easy to find which process the stats apply to in multi-process mode. Notice the empty line after the information output which marks the end of the first block. A similar empty line appears at the end of the second block (stats) so that the reader knows the output has not been truncated.

show table

Dump general information on all known stick-tables. Their name is returned (the name of the proxy which holds them), their type (currently zero, always IP), their size in maximum possible number of entries, and the number of entries currently in use.

```
Example :
$ echo "show table" | socat stdio /tmp/sockl1
>>> # table: front_pub, type: ip, size:204800, used:171454
>>> # table: back_rdp, type: ip, size:204800, used:0
```

show table <name> [data.<type> <operator> <value>] [[key <key>] Dump contents of stick-table <name>. In this mode, a first line of generic information about the table is reported as with "show table", then all entries are dumped. Since this can be quite heavy, it is possible to specify a filter in order to specify what entries to display.

When the "data." form is used the filter applies to the stored data (see "stick-table" in section 4.2). A stored data type must be specified in <type>, and this data type must be stored in the table otherwise an error is reported. The data is compared according to <operator> with the 64-bit integer <value>. Operators are the same as with the ACLs :

- eq : match entries whose data is equal to this value
- ne : match entries whose data is not equal to this value
- le : match entries whose data is less than or equal to this value
- ge : match entries whose data is greater than or equal to this value
- lt : match entries whose data is less than this value
- gt : match entries whose data is greater than this value

When the key form is used the entry <key> is shown. The key must be of the same type as the table, which currently is limited to IPv4, IPv6, integer, and string.

```
Example :
$ echo "show table http_proxy" | socat stdio /tmp/sockl1
>>> # table: http_proxy, type: ip, size:204800, used:2
>>> 0x80e6a4c: key=127.0.0.1 use=0 exp=3594729 gpc0=0 conn_rate(30000)=1 \
bytes_out_rate(60000)=187
>>> 0x80e6a80: key=127.0.0.2 use=0 exp=3594740 gpc0=1 conn_rate(30000)=10 \
bytes_out_rate(60000)=191

$ echo "show table http_proxy data.gpc0 gt 0" | socat stdio /tmp/sockl1
>>> # table: http_proxy, type: ip, size:204800, used:2
>>> 0x80e6a80: key=127.0.0.2 use=0 exp=3594740 gpc0=1 conn_rate(30000)=10 \
bytes_out_rate(60000)=191

$ echo "show table http_proxy data.conn_rate gt 5" | \
socat stdio /tmp/sockl1
>>> # table: http_proxy, type: ip, size:204800, used:2
>>> 0x80e6a80: key=127.0.0.2 use=0 exp=3594740 gpc0=1 conn_rate(30000)=10 \
bytes_out_rate(60000)=191

$ echo "show table http_proxy key 127.0.0.2" | \
socat stdio /tmp/sockl1
>>> # table: http_proxy, type: ip, size:204800, used:2
>>> 0x80e6a80: key=127.0.0.2 use=0 exp=3594740 gpc0=1 conn_rate(30000)=10 \
bytes_out_rate(60000)=191
```

When the data criterion applies to a dynamic value dependent on time such as a bytes rate, the value is dynamically computed during the evaluation of the entry in order to decide whether it has to be dumped or not. This means that such a filter could match for some time then not match anymore because as time goes, the average event rate drops.

It is possible to use this to extract lists of IP addresses abusing the service, in order to monitor them or even blacklist them in a firewall.

```
Example :
$ echo "show table http_proxy data.gpc0 gt 0" \
| socat stdio /tmp/sockl1 \
| fgrep 'key=' | cut -d' ' -f2 | cut -d= -f2 > abusers-ip.txt
( or | awk '/key/{ print a[split($2,a,"=")]; }' )
```

shutdown frontend <frontend>

Completely delete the specified frontend. All the ports it was bound to will be released. It will not be possible to enable the frontend anymore after this operation. This is intended to be used in environments where stopping a proxy is not even imaginable but a misconfigured proxy must be fixed. That way it's possible to release the port and bind it into another process to restore operations. The frontend will not appear at all on the stats page once it is terminated.

The frontend may be specified either by its name or by its numeric ID, prefixed with a sharp ('#').

This command is restricted and can only be issued on sockets configured for level "admin".

shutdown session <id>

Immediately terminate the session matching the specified session identifier. This identifier is the first field at the beginning of the lines in the dumps of "show sess" (it corresponds to the session pointer). This can be used to terminate a long-running session without waiting for a timeout or when an endless transfer is ongoing. Such terminated sessions are reported with a 'K' flag in the logs.

shutdown sessions server <backend>/<server>

Immediately terminate all the sessions attached to the specified server. This can be used to terminate long-running sessions after a server is put into maintenance mode, for instance. Such terminated sessions are reported with a 'K' flag in the logs.

```
/* Local variables:
* fill-column: 79
* End:
*/
```