```
                          ----------------------
                                HAProxy
                           Configuration Manual
                          ----------------------
                               version 1.6
                              willy tarreau
                               2015/12/21


This document covers the configuration language as implemented in the version
specified above. It does not provide any hint, example or advice. For such
documentation, please refer to the Reference Manual or the Architecture Manual.
The summary below is meant to help you search sections by name and navigate
through the document.

Note to documentation contributors :
    This document is formatted with 80 columns per line, with even number of
    spaces for indentation and without tabs. Please follow these rules strictly
    so that it remains easily printable everywhere. If a line needs to be
    printed verbatim and does not fit, please end each line with a backslash
    ('\') and continue on next line, indented by two characters. It is also
    sometimes useful to prefix all output lines (logs, console outs) with 3
    closing angle brackets ('>>>') in order to help get the difference between
    inputs and outputs when it can become ambiguous. If you add sections,
    please update the summary below for easier searching.


Summary
-------

1.    Quick reminder about HTTP
1.1.      The HTTP transaction model
1.2.      HTTP request
1.2.1.        The Request line
1.2.2.        The request headers
1.3.      HTTP response
1.3.1.        The Response line
1.3.2.        The response headers

2.    Configuring HAProxy
2.1.      Configuration file format
2.2.      Quoting and escaping
2.3.      Environment variables
2.4.      Time format
2.5.      Examples

3.    Global parameters
3.1.      Process management and security
3.2.      Performance tuning
3.3.      Debugging
3.4.      Userlists
3.5.      Peers
3.6.      Mailers

4.    Proxies
4.1.      Proxy keywords matrix
4.2.      Alphabetically sorted keywords reference

5.    Bind and Server options
5.1.      Bind options
5.2.      Server and default-server options
5.3.      Server DNS resolution
5.3.1.      Global overview
5.3.2.      The resolvers section

```



```
6.    HTTP header manipulation

7.    Using ACLs and fetching samples
7.1.      ACL basics
7.1.1.      Matching booleans
7.1.2.      Matching integers
7.1.3.      Matching strings
7.1.4.      Matching regular expressions (regexes)
7.1.5.      Matching arbitrary data blocks
7.1.6.      Matching IPv4 and IPv6 addresses
7.2.      Using ACLs to form conditions
7.3.      Fetching samples
7.3.1.        Converters
7.3.2.        Fetching samples from internal states
7.3.3.        Fetching samples at Layer 4
7.3.4.        Fetching samples at Layer 5
7.3.5.        Fetching samples from buffer contents (Layer 6)
7.3.6.        Fetching HTTP samples (Layer 7)
7.4.      Pre-defined ACLs

8.    Logging
8.1.      Log levels
8.2.      Log formats
8.2.1.        Default log format
8.2.2.        TCP log format
8.2.3.        HTTP log format
8.2.4.        Custom log format
8.2.5.        Error log format
8.3.      Advanced logging options
8.3.1.        Disabling logging of external tests
8.3.2.        Logging before waiting for the session to terminate
8.3.3.        Raising log level upon errors
8.3.4.        Disabling logging of successful connections
8.4.      Timing events
8.5.      Session state at disconnection
8.6.      Non-printable characters
8.7.      Capturing HTTP cookies
8.8.      Capturing HTTP headers
8.9.      Examples of logs


1. Quick reminder about HTTP
----------------------------

When haproxy is running in HTTP mode, both the request and the response are
fully analyzed and indexed, thus it becomes possible to build matching criteria
on almost anything found in the contents.

However, it is important to understand how HTTP requests and responses are
formed, and how HAProxy decomposes them. It will then become easier to write
correct rules and to debug existing configurations.


1.1. The HTTP transaction model
-------------------------------

The HTTP protocol is transaction-driven. This means that each request will lead
to one and only one response. Traditionally, a TCP connection is established
from the client to the server, a request is sent by the client on the
connection, the server responds and the connection is closed. A new request
will involve a new connection :

  [CON1] [REQ1] ... [RESP1] [CLO1] [CON2] [REQ2] ... [RESP2] [CLO2] ...

In this mode, called the "HTTP close" mode, there are as many connection
```

establishments as there are HTTP transactions. Since the connection is closed by the server after the response, the client does not need to know the content length.

Due to the transactional nature of the protocol, it was possible to improve it to avoid closing a connection between two subsequent transactions. In this mode however, it is mandatory that the server indicates the content length for each response so that the client does not wait indefinitely. For this, a special header is used: "Content-length". This mode is called the "keep-alive" mode :

```
  [CON] [REQ1] ... [RESP1] [REQ2] ... [RESP2] [CLO] ...
```

Its advantages are a reduced latency between transactions, and less processing power required on the server side. It is generally better than the close mode, but not always because the clients often limit their concurrent connections to a smaller value.

A last improvement in the communications is the pipelining mode. It still uses keep-alive, but the client does not wait for the first response to send the second request. This is useful for fetching large number of images composing a page :

```
  [CON] [REQ1] [REQ2] ... [RESP1] [RESP2] [CLO] ...
```

This can obviously have a tremendous benefit on performance because the network latency is eliminated between subsequent requests. Many HTTP agents do not correctly support pipelining since there is no way to associate a response with the corresponding request in HTTP. For this reason, it is mandatory for the server to reply in the exact same order as the requests were received.

By default HAProxy operates in keep-alive mode with regards to persistent connections: for each connection it processes each request and response, and leaves the connection idle on both sides between the end of a response and the start of a new request.

HAProxy supports 5 connection modes :
  - keep alive   : all requests and responses are processed (default)
  - tunnel   : only the first request and response are processed,
                    everything else is forwarded with no analysis.
  - passive close : tunnel with "Connection: close" added in both directions.
  - server close : the server-facing connection is closed after the response.
  - forced close : the connection is actively closed after end of response.

1.2. HTTP request

First, let's consider this HTTP request :

Line number	Contents
1	GET /serv/login.php?lang=en&profile=2 HTTP/1.1
2	Host: www.mydomain.com
3	User-agent: my small browser
4	Accept: image/jpeg, image/gif
5	Accept: image/png

1.2.1. The Request line

Line 1 is the "request line". It is always composed of 3 fields :

  - a METHOD   : GET
  - a URI   : /serv/login.php?lang=en&profile=2

  - a version tag : HTTP/1.1

All of them are delimited by what the standard calls LWS (linear white spaces), which are commonly spaces, but can also be tabs or line feeds/carriage returns followed by spaces/tabs. The method itself cannot contain any colon (':') and is limited to alphabetic letters. All those various combinations make it desirable that HAProxy performs the splitting itself rather than leaving it to the user to write a complex or inaccurate regular expression.

The URI itself can have several forms :

  - A "relative URI" :

      /serv/login.php?lang=en&profile=2

    It is a complete URL without the host part. This is generally what is
    received by servers, reverse proxies and transparent proxies.

  - An "absolute URI", also called a "URL" :

      http://192.168.0.12:8080/serv/login.php?lang=en&profile=2

    It is composed of a "scheme" (the protocol name followed by '://'), a host
    name or address, optionally a colon (':') followed by a port number, then
    a relative URI beginning at the first slash ('/') after the address part.
    This is generally what proxies receive, but a server supporting HTTP/1.1
    must accept this form too.

  - a star ('*') : this form is only accepted in association with the OPTIONS
    method and is not relayable. It is used to inquiry a next hop's
    capabilities.

  - an address:port combination : 192.168.0.12:80
    This is used with the CONNECT method, which is used to establish TCP
    tunnels through HTTP proxies, generally for HTTPS, but sometimes for
    other protocols too.

In a relative URI, two sub-parts are identified. The part before the question mark is called the "path". It is typically the relative path to static objects on the server. The part after the question mark is called the "query string". It is mostly used with GET requests sent to dynamic scripts and is very specific to the language, framework or application in use.

1.2.2. The request headers

The headers start at the second line. They are composed of a name at the beginning of the line, immediately followed by a colon (':'). Traditionally, an LWS is added after the colon but that's not required. Then come the values. Multiple identical headers may be folded into one single line, delimiting the values with commas, provided that their order is respected. This is commonly encountered in the "Cookie:" field. A header may span over multiple lines if the subsequent lines begin with an LWS. In the example in 1.2, lines 4 and 5 define a total of 3 values for the "Accept:" header.

Contrary to a common mis-conception, header names are not case-sensitive, and their values are not either if they refer to other header names (such as the "Connection:" header).

The end of the headers is indicated by the first empty line. People often say that it's a double line feed, which is not exact, even if a double line feed is one valid form of empty line.

Fortunately, HAProxy takes care of all these complex combinations when indexing

headers, checking values and counting them, so there is no reason to worry
about the way they could be written, but it is important not to accuse an
application of being buggy if it does unusual, valid things.

Important note:
   As suggested by RFC2616, HAProxy normalizes headers by replacing line breaks
   in the middle of headers by LWS in order to join multi-line headers. This
   is necessary for proper analysis and helps less capable HTTP parsers to work
   correctly and not to be fooled by such complex constructs.

1.3. HTTP response

An HTTP response looks very much like an HTTP request. Both are called HTTP
messages. Let's consider this HTTP response :

  Line   Contents
  number
     1   HTTP/1.1 200 OK
     2   Content-length: 350
     3   Content-Type: text/html

As a special case, HTTP supports so called "Informational responses" as status
codes 1xx. These messages are special in that they don't convey any part of the
response, they're just used as sort of a signaling message to ask a client to
continue to post its request for instance. In the case of a status 100 response
the requested information will be carried by the next non-100 response message
following the informational one. This implies that multiple responses may be
sent to a single request, and that this only works when keep-alive is enabled
(1xx messages are HTTP/1.1 only). HAProxy handles these messages and is able to
correctly forward and skip them, and only process the next non-100 response. As
such, these messages are neither logged nor transformed, unless explicitly
state otherwise. Status 101 messages indicate that the protocol is changing
over the same connection and that haproxy must switch to tunnel mode, just as
if a CONNECT had occurred. Then the Upgrade header would contain additional
information about the type of protocol the connection is switching to.

1.3.1. The Response line

Line 1 is the "response line". It is always composed of 3 fields :

  - a version tag : HTTP/1.1
  - a status code : 200
  - a reason   : OK

The status code is always 3-digit. The first digit indicates a general status :
 - 1xx = informational message to be skipped (eg: 100, 101)
 - 2xx = OK, content is following   (eg: 200, 206)
 - 3xx = OK, no content following   (eg: 302, 304)
 - 4xx = error caused by the client (eg: 401, 403, 404)
 - 5xx = error caused by the server (eg: 500, 502, 503)

Please refer to RFC2616 for the detailed meaning of all such codes. The
"reason" field is just a hint, but is not parsed by clients. Anything can be
found there, but it's a common practice to respect the well-established
messages. It can be composed of one or multiple words, such as "OK", "Found",
or "Authentication Required".

Haproxy may emit the following status codes by itself :

  Code When / reason
   200 access to stats page, and when replying to monitoring requests

   301 when performing a redirection, depending on the configured code
   302 when performing a redirection, depending on the configured code
   303 when performing a redirection, depending on the configured code
   307 when performing a redirection, depending on the configured code
   308 when performing a redirection, depending on the configured code
   400 for an invalid or too large request
   401 when an authentication is required to perform the action (when
        accessing the stats page)
   403 when a request is forbidden by a "block" ACL or "reqdeny" filter
   408 when the request timeout strikes before the request is complete
   500 when haproxy encounters an unrecoverable internal error, such as a
        memory allocation failure, which should never happen
   502 when the server returns an empty, invalid or incomplete response, or
        when an "rspdeny" filter blocks the response.
   503 when no server was available to handle the request, or in response to
        monitoring requests which match the "monitor fail" condition
   504 when the response timeout strikes before the server responds

The error 4xx and 5xx codes above may be customized (see "errorloc" in section
4.2).

1.3.2. The response headers

Response headers work exactly like request headers, and as such, HAProxy uses
the same parsing function for both. Please refer to paragraph 1.2.2 for more
details.

2. Configuring HAProxy

2.1. Configuration file format

HAProxy's configuration process involves 3 major sources of parameters :

  - the arguments from the command-line, which always take precedence
  - the "global" section, which sets process-wide parameters
  - the proxies sections which can take form of "defaults", "listen",
    "frontend" and "backend".

The configuration file syntax consists in lines beginning with a keyword
referenced in this manual, optionally followed by one or several parameters
delimited by spaces.

2.2. Quoting and escaping

HAProxy's configuration introduces a quoting and escaping system similar to
many programming languages. The configuration file supports 3 types: escaping
with a backslash, weak quoting with double quotes, and strong quoting with
single quotes.

If spaces have to be entered in strings, then they must be escaped by preceding
them by a backslash ('\') or by quoting them. Backslashes also have to be
escaped by doubling or strong quoting them.

Escaping is achieved by preceding a special character by a backslash ('\'):

  \   to mark a space and differentiate it from a delimiter
  \#   to mark a hash and differentiate it from a comment
  \\   to use a backslash

```
   \'   to use a single quote and differentiate it from strong quoting
   \"   to use a double quote and differentiate it from weak quoting

Weak quoting is achieved by using double quotes (""). Weak quoting prevents
the interpretation of:

       space as a parameter separator
   '   single quote as a strong quoting delimiter
   #   hash as a comment start

Weak quoting permits the interpretation of variables, if you want to use a non
-interpreted dollar within a double quoted string, you should escape it with a
backslash ("\$"), it does not work outside weak quoting.

Interpretation of escaping and special characters are not prevented by weak
quoting.

Strong quoting is achieved by using single quotes (''). Inside single quotes,
nothing is interpreted, it's the efficient way to quote regexes.

Quoted and escaped strings are replaced in memory by their interpreted
equivalent, it allows you to perform concatenation.

  Example:
      # those are equivalents:
      log-format %{+Q}o\ %t\ %s\ %{-Q}r
      log-format "%{+Q}o %t %s %{-Q}r"
      log-format '%{+Q}o %t %s %{-Q}r'
      log-format "%{+Q}o %t"' %s %{-Q}r'
      log-format "%{+Q}o %t"' %s'\ %{-Q}r

      # those are equivalents:
      reqrep "^([^\ :]*)\ /static/(.*)"     \1\ /\2
      reqrep "^([^ :]*)\ /static/(.*)"     '\1 /\2'
      reqrep "^([^ :]*)\ /static/(.*)"     "\1 /\2"
      reqrep "^([^ :]*)\ /static/(.*)"     "\1\ /\2"


2.3. Environment variables
--------------------------

HAProxy's configuration supports environment variables. Those variables are
interpreted only within double quotes. Variables are expanded during the
configuration parsing. Variable names must be preceded by a dollar ("$") and
optionally enclosed with braces ("{}") similarly to what is done in Bourne
shell. Variable names can contain alphanumerical characters or the character
underscore ("_") but should not start with a digit.

  Example:

      bind "fd@${FD_APP1}"

      log "${LOCAL_SYSLOG}:514" local0 notice   # send to local server

      user "$HAPROXY_USER"


2.4. Time format
----------------

Some parameters involve values representing time, such as timeouts. These
values are generally expressed in milliseconds (unless explicitly stated
otherwise) but may be expressed in any other unit by suffixing the unit to the
numeric value. It is important to consider this because it will not be repeated
for every keyword. Supported units are :
```



```

   - us : microseconds. 1 microsecond = 1/1000000 second
   - ms : milliseconds. 1 millisecond = 1/1000 second. This is the default.
   - s  : seconds. 1s = 1000ms
   - m  : minutes. 1m = 60s = 60000ms
   - h  : hours.   1h = 60m = 3600s = 3600000ms
   - d  : days.    1d = 24h = 1440m = 86400s = 86400000ms


2.4. Examples
-------------

    # Simple configuration for an HTTP proxy listening on port 80 on all
    # interfaces and forwarding requests to a single backend "servers" with a
    # single server "server1" listening on 127.0.0.1:8000
    global
        daemon
        maxconn 256

    defaults
        mode http
        timeout connect 5000ms
        timeout client 50000ms
        timeout server 50000ms

    frontend http-in
        bind *:80
        default_backend servers

    backend servers
        server server1 127.0.0.1:8000 maxconn 32


    # The same configuration defined with a single listen block. Shorter but
    # less expressive, especially in HTTP mode.
    global
        daemon
        maxconn 256

    defaults
        mode http
        timeout connect 5000ms
        timeout client 50000ms
        timeout server 50000ms

    listen http-in
        bind *:80
        server server1 127.0.0.1:8000 maxconn 32


Assuming haproxy is in $PATH, test these configurations in a shell with:

    $ sudo haproxy -f configuration.conf -c


3. Global parameters
--------------------

Parameters in the "global" section are process-wide and often OS-specific. They
are generally set once for all and do not need being changed once correct. Some
of them have command-line equivalents.

The following keywords are supported in the "global" section :

 * Process management and security
```

```
    - ca-base
    - chroot
    - crt-base
    - cpu-map
    - daemon
    - description
    - deviceatlas-json-file
    - deviceatlas-log-level
    - deviceatlas-separator
    - deviceatlas-properties-cookie
    - external-check
    - gid
    - group
    - log
    - log-tag
    - log-send-hostname
    - lua-load
    - nbproc
    - node
    - pidfile
    - uid
    - ulimit-n
    - user
    - stats
    - ssl-default-bind-ciphers
    - ssl-default-bind-options
    - ssl-default-server-ciphers
    - ssl-default-server-options
    - ssl-dh-param-file
    - ssl-server-verify
    - unix-bind
    - 51degrees-data-file
    - 51degrees-property-name-list
    - 51degrees-property-separator
    - 51degrees-cache-size

  * Performance tuning
    - max-spread-checks
    - maxconn
    - maxconnrate
    - maxcomprate
    - maxcompcpuusage
    - maxpipes
    - maxsessrate
    - maxsslconn
    - maxsslrate
    - maxzlibmem
    - noepoll
    - nokqueue
    - nopoll
    - nosplice
    - nogetaddrinfo
    - spread-checks
    - server-state-base
    - server-state-file
    - tune.buffers.limit
    - tune.buffers.reserve
    - tune.bufsize
    - tune.chksize
    - tune.comp.maxlevel
    - tune.http.cookielen
    - tune.http.maxhdr
    - tune.idletimer
    - tune.lua.forced-yield
    - tune.lua.maxmem
```



```
    - tune.lua.session-timeout
    - tune.lua.task-timeout
    - tune.lua.service-timeout
    - tune.maxaccept
    - tune.maxpollevents
    - tune.maxrewrite
    - tune.pattern.cache-size
    - tune.pipesize
    - tune.rcvbuf.client
    - tune.rcvbuf.server
    - tune.sndbuf.client
    - tune.sndbuf.server
    - tune.ssl.cachesize
    - tune.ssl.lifetime
    - tune.ssl.force-private-cache
    - tune.ssl.maxrecord
    - tune.ssl.default-dh-param
    - tune.ssl.ssl-ctx-cache-size
    - tune.vars.global-max-size
    - tune.vars.reqres-max-size
    - tune.vars.sess-max-size
    - tune.vars.txn-max-size
    - tune.zlib.memlevel
    - tune.zlib.windowsize

  * Debugging
    - debug
    - quiet


3.1. Process management and security
------------------------------------

ca-base <dir>
  Assigns a default directory to fetch SSL CA certificates and CRLs from when a
  relative path is used with "ca-file" or "crl-file" directives. Absolute
  locations specified in "ca-file" and "crl-file" prevail and ignore "ca-base".

chroot <jail dir>
  Changes current directory to <jail dir> and performs a chroot() there before
  dropping privileges. This increases the security level in case an unknown
  vulnerability would be exploited, since it would make it very hard for the
  attacker to exploit the system. This only works when the process is started
  with superuser privileges. It is important to ensure that <jail_dir> is both
  empty and unwritable to anyone.

cpu-map <"all"|"odd"|"even"|process_num> <cpu-set>...
  On Linux 2.6 and above, it is possible to bind a process to a specific CPU
  set. This means that the process will never run on other CPUs. The "cpu-map"
  directive specifies CPU sets for process sets. The first argument is the
  process number to bind. This process must have a number between 1 and 32 or
  64, depending on the machine's word size, and any process IDs above nbproc
  are ignored. It is possible to specify all processes at once using "all",
  only odd numbers using "odd" or even numbers using "even", just like with the
  "bind-process" directive. The second and forthcoming arguments are CPU sets.
  Each CPU set is either a unique number between 0 and 31 or 63 or a range with
  two such numbers delimited by a dash ('-'). Multiple CPU numbers or ranges
  may be specified, and the processes will be allowed to bind to all of them.
  Obviously, multiple "cpu-map" directives may be specified. Each "cpu-map"
  directive will replace the previous ones when they overlap.

crt-base <dir>
  Assigns a default directory to fetch SSL certificates from when a relative
  path is used with "crtfile" directives. Absolute locations specified after
  "crtfile" prevail and ignore "crt-base".
```

```
651
652  daemon
653    Makes the process fork into background. This is the recommended mode of
654    operation. It is equivalent to the command line "-D" argument. It can be
655    disabled by the command line "-db" argument.
656
657  deviceatlas-json-file <path>
658    Sets the path of the DeviceAtlas JSON data file to be loaded by the API.
659    The path must be a valid JSON data file and accessible by Haproxy process.
660
661  deviceatlas-log-level <value>
662    Sets the level of informations returned by the API. This directive is
663    optional and set to 0 by default if not set.
664
665  deviceatlas-separator <char>
666    Sets the character separator for the API properties results. This directive
667    is optional and set to | by default if not set.
668
669  deviceatlas-properties-cookie <name>
670    Sets the client cookie's name used for the detection if the DeviceAtlas
671    Client-side component was used during the request. This directive is optional
672    and set to DAPROPS by default if not set.
673
674  external-check
675    Allows the use of an external agent to perform health checks.
676    This is disabled by default as a security precaution.
677    See "option external-check".
678
679  gid <number>
680    Changes the process' group ID to <number>. It is recommended that the group
681    ID is dedicated to HAProxy or to a small set of similar daemons. HAProxy must
682    be started with a user belonging to this group, or with superuser privileges.
683    Note that if haproxy is started from a user having supplementary groups, it
684    will only be able to drop these groups if started with superuser privileges.
685    See also "group" and "uid".
686
687  group <group name>
688    Similar to "gid" but uses the GID of group name <group name> from /etc/group.
689    See also "gid" and "user".
690
691  log <address> [len <length>] [format <format>] <facility> [max level [min level]]
692    Adds a global syslog server. Up to two global servers can be defined. They
693    will receive logs for startups and exits, as well as all logs from proxies
694    configured with "log global".
695
696    <address> can be one of:
697
698          - An IPv4 address optionally followed by a colon and a UDP port. If
699            no port is specified, 514 is used by default (the standard syslog
700            port).
701
702          - An IPv6 address followed by a colon and optionally a UDP port. If
703            no port is specified, 514 is used by default (the standard syslog
704            port).
705
706          - A filesystem path to a UNIX domain socket, keeping in mind
707            considerations for chroot (be sure the path is accessible inside
708            the chroot) and uid/gid (be sure the path is appropriately
709            writeable).
710
711          You may want to reference some environment variables in the address
712          parameter, see section 2.3 about environment variables.
713
714    <length> is an optional maximum line length. Log lines larger than this value
715             will be truncated before being sent. The reason is that syslog
```



```
716             servers act differently on log line length. All servers support the
717             default value of 1024, but some servers simply drop larger lines
718             while others do log them. If a server supports long lines, it may
719             make sense to set this value here in order to avoid truncating long
720             lines. Similarly, if a server drops long lines, it is preferable to
721             truncate them before sending them. Accepted values are 80 to 65535
722             inclusive. The default value of 1024 is generally fine for all
723             standard usages. Some specific cases of long captures or
724             JSON-formated logs may require larger values.
725
726    <format> is the log format used when generating syslog messages. It may be
727             one of the following :
728
729      rfc3164   The RFC3164 syslog message format. This is the default.
730                (https://tools.ietf.org/html/rfc3164)
731
732      rfc5424   The RFC5424 syslog message format.
733                (https://tools.ietf.org/html/rfc5424)
734
735    <facility> must be one of the 24 standard syslog facilities :
736
737            kern   user   mail   daemon auth   syslog lpr    news
738            uucp   cron   auth2  ftp    ntp    audit  alert  cron2
739            local0 local1 local2 local3 local4 local5 local6 local7
740
741    An optional level can be specified to filter outgoing messages. By default,
742    all messages are sent. If a maximum level is specified, only messages with a
743    severity at least as important as this level will be sent. An optional minimum
744    level can be specified. If it is set, logs emitted with a more severe level
745    than this one will be capped to this level. This is used to avoid sending
746    "emerg" messages on all terminals on some default syslog configurations.
747    Eight levels are known :
748
749            emerg  alert  crit   err    warning notice info  debug
750
751  log-send-hostname [<string>]
752    Sets the hostname field in the syslog header. If optional "string" parameter
753    is set the header is set to the string contents, otherwise uses the hostname
754    of the system. Generally used if one is not relaying logs through an
755    intermediate syslog server or for simply customizing the hostname printed in
756    the logs.
757
758  log-tag <string>
759    Sets the tag field in the syslog header to this string. It defaults to the
760    program name as launched from the command line, which usually is "haproxy".
761    Sometimes it can be useful to differentiate between multiple processes
762    running on the same host. See also the per-proxy "log-tag" directive.
763
764  lua-load <file>
765    This global directive loads and executes a Lua file. This directive can be
766    used multiple times.
767
768  nbproc <number>
769    Creates <number> processes when going daemon. This requires the "daemon"
770    mode. By default, only one process is created, which is the recommended mode
771    of operation. For systems limited to small sets of file descriptors per
772    process, it may be needed to fork multiple daemons. USING MULTIPLE PROCESSES
773    IS HARDER TO DEBUG AND IS REALLY DISCOURAGED. See also "daemon".
774
775  pidfile <pidfile>
776    Writes pids of all daemons into file <pidfile>. This option is equivalent to
777    the "-p" command line argument. The file must be accessible to the user
778    starting the process. See also "daemon".
779
780  stats bind-process [ all | odd | even | <number 1-64>[-<number 1-64>] ] ...
```

```
   Limits the stats socket to a certain set of processes numbers. By default the
   stats socket is bound to all processes, causing a warning to be emitted when
   nbproc is greater than 1 because there is no way to select the target process
   when connecting. However, by using this setting, it becomes possible to pin
   the stats socket to a specific set of processes, typically the first one. The
   warning will automatically be disabled when this setting is used, whatever
   the number of processes used. The maximum process ID depends on the machine's
   word size (32 or 64). A better option consists in using the "process" setting
   of the "stats socket" line to force the process on each line.

server-state-base <directory>
   Specifies the directory prefix to be prepended in front of all servers state
   file names which do not start with a '/'. See also "server-state-file",
   "load-server-state-from-file" and "server-state-file-name".

server-state-file <file>
   Specifies the path to the file containing state of servers. If the path starts
   with a slash ('/'), it is considered absolute, otherwise it is considered
   relative to the directory specified using "server-state-base" (if set) or to
   the current directory. Before reloading HAProxy, it is possible to save the
   servers' current state using the stats command "show servers state". The
   output of this command must be written in the file pointed by <file>. When
   starting up, before handling traffic, HAProxy will read, load and apply state
   for each server found in the file and available in its current running
   configuration. See also "server-state-base" and "show servers state",
   "load-server-state-from-file" and "server-state-file-name"

ssl-default-bind-ciphers <ciphers>
   This setting is only available when support for OpenSSL was built in. It sets
   the default string describing the list of cipher algorithms ("cipher suite")
   that are negotiated during the SSL/TLS handshake for all "bind" lines which
   do not explicitly define theirs. The format of the string is defined in
   "man 1 ciphers" from OpenSSL man pages, and can be for instance a string such
   as "AES:ALL:!aNULL:!eNULL:+RC4:@STRENGTH" (without quotes). Please check the
   "bind" keyword for more information.

ssl-default-bind-options [<option>]...
   This setting is only available when support for OpenSSL was built in. It sets
   default ssl-options to force on all "bind" lines. Please check the "bind"
   keyword to see available options.

   Example:
         global
            ssl-default-bind-options no-sslv3 no-tls-tickets

ssl-default-server-ciphers <ciphers>
   This setting is only available when support for OpenSSL was built in. It
   sets the default string describing the list of cipher algorithms that are
   negotiated during the SSL/TLS handshake with the server, for all "server"
   lines which do not explicitly define theirs. The format of the string is
   defined in "man 1 ciphers". Please check the "server" keyword for more
   information.

ssl-default-server-options [<option>]...
   This setting is only available when support for OpenSSL was built in. It sets
   default ssl-options to force on all "server" lines. Please check the "server"
   keyword to see available options.

ssl-dh-param-file <file>
   This setting is only available when support for OpenSSL was built in. It sets
   the default DH parameters that are used during the SSL/TLS handshake when
   ephemeral Diffie-Hellman (DHE) key exchange is used, for all "bind" lines
   which do not explicitely define theirs. It will be overridden by custom DH
   parameters found in a bind certificate file if any. If custom DH parameters
   are not specified either by using ssl-dh-param-file or by setting them
```



```
   directly in the certificate file, pre-generated DH parameters of the size
   specified by tune.ssl.default-dh-param will be used. Custom parameters are
   known to be more secure and therefore their use is recommended.
   Custom DH parameters may be generated by using the OpenSSL command
   "openssl dhparam <size>", where size should be at least 2048, as 1024-bit DH
   parameters should not be considered secure anymore.

ssl-server-verify [none|required]
   The default behavior for SSL verify on servers side. If specified to 'none',
   servers certificates are not verified. The default is 'required' except if
   forced using cmdline option '-dV'.

stats socket [<address:port>|<path>] [param*]
   Binds a UNIX socket to <path> or a TCPv4/v6 address to <address:port>.
   Connections to this socket will return various statistics outputs and even
   allow some commands to be issued to change some runtime settings. Please
   consult section 9.2 "Unix Socket commands" of Management Guide for more
   details.

   All parameters supported by "bind" lines are supported, for instance to
   restrict access to some users or their access rights. Please consult
   section 5.1 for more information.

stats timeout <timeout, in milliseconds>
   The default timeout on the stats socket is set to 10 seconds. It is possible
   to change this value with "stats timeout". The value must be passed in
   milliseconds, or be suffixed by a time unit among { us, ms, s, m, h, d }.

stats maxconn <connections>
   By default, the stats socket is limited to 10 concurrent connections. It is
   possible to change this value with "stats maxconn".

uid <number>
   Changes the process' user ID to <number>. It is recommended that the user ID
   is dedicated to HAProxy or to a small set of similar daemons. HAProxy must
   be started with superuser privileges in order to be able to switch to another
   one. See also "gid" and "user".

ulimit-n <number>
   Sets the maximum number of per-process file-descriptors to <number>. By
   default, it is automatically computed, so it is recommended not to use this
   option.

unix-bind [ prefix <prefix> ] [ mode <mode> ] [ user <user> ] [ uid <uid> ]
          [ group <group> ] [ gid <gid> ]

   Fixes common settings to UNIX listening sockets declared in "bind" statements.
   This is mainly used to simplify declaration of those UNIX sockets and reduce
   the risk of errors, since those settings are most commonly required but are
   also process-specific. The <prefix> setting can be used to force all socket
   path to be relative to that directory. This might be needed to access another
   component's chroot. Note that those paths are resolved before haproxy chroots
   itself, so they are absolute. The <mode>, <user>, <uid>, <group> and <gid>
   all have the same meaning as their homonyms used by the "bind" statement. If
   both are specified, the "bind" statement has priority, meaning that the
   "unix-bind" settings may be seen as process-wide default settings.

user <user name>
   Similar to "uid" but uses the UID of user name <user name> from /etc/passwd.
   See also "uid" and "group".

node <name>
   Only letters, digits, hyphen and underscore are allowed, like in DNS names.

   This statement is useful in HA configurations where two or more processes or
```

servers share the same IP address. By setting a different node-name on all
nodes, it becomes easy to immediately spot what server is handling the
traffic.

description <text>
Add a text that describes the instance.

Please note that it is required to escape certain characters (# for example)
and this text is inserted into a html page so you should avoid using
"<" and ">" characters.

51degrees-data-file <file path>
The path of the 51Degrees data file to provide device detection services. The
file should be unzipped and accessible by HAProxy with relevavnt permissions.

Please note that this option is only available when haproxy has been
compiled with USE_51DEGREES.

51degrees-property-name-list [<string>]
A list of 51Degrees property names to be load from the dataset. A full list
of names is available on the 51Degrees website:
https://51degrees.com/resources/property-dictionary

Please note that this option is only available when haproxy has been
compiled with USE_51DEGREES.

51degrees-property-separator <char>
A char that will be appended to every property value in a response header
containing 51Degrees results. If not set that will be set as ','.

Please note that this option is only available when haproxy has been
compiled with USE_51DEGREES.

51degrees-cache-size <number>
Sets the size of the 51Degrees converter cache to <number> entries. This
is an LRU cache which reminds previous device detections and their results.
By default, this cache is disabled.

Please note that this option is only available when haproxy has been
compiled with USE_51DEGREES.

3.2. Performance tuning

max-spread-checks <delay in milliseconds>
By default, haproxy tries to spread the start of health checks across the
smallest health check interval of all the servers in a farm. The principle is
to avoid hammering services running on the same server. But when using large
check intervals (10 seconds or more), the last servers in the farm take some
time before starting to be tested, which can be a problem. This parameter is
used to enforce an upper bound on delay between the first and the last check,
even if the servers' check intervals are larger. When servers run with
shorter intervals, their intervals will be respected though.

maxconn <number>
Sets the maximum per-process number of concurrent connections to <number>. It
is equivalent to the command-line argument "-n". Proxies will stop accepting
connections when this limit is reached. The "ulimit-n" parameter is
automatically adjusted according to this value. See also "ulimit-n". Note:
the "select" poller cannot reliably use more than 1024 file descriptors on
some platforms. If your platform only supports select and reports "select
FAILED" on startup, you need to reduce maxconn until it works (slightly
below 500 in general). If this value is not set, it will default to the value
set in DEFAULT_MAXCONN at build time (reported in haproxy -vv) if no memory

limit is enforced, or will be computed based on the memory limit, the buffer
size, memory allocated to compression, SSL cache size, and use or not of SSL
and the associated maxsslconn (which can also be automatic).

maxconnrate <number>
Sets the maximum per-process number of connections per second to <number>.
Proxies will stop accepting connections when this limit is reached. It can be
used to limit the global capacity regardless of each frontend capacity. It is
important to note that this can only be used as a service protection measure,
as there will not necessarily be a fair share between frontends when the
limit is reached, so it's a good idea to also limit each frontend to some
value close to its expected share. Also, lowering tune.maxaccept can improve
fairness.

maxcomprate <number>
Sets the maximum per-process input compression rate to <number> kilobytes
per second. For each session, if the maximum is reached, the compression
level will be decreased during the session. If the maximum is reached at the
beginning of a session, the session will not compress at all. If the maximum
is not reached, the compression level will be increased up to
tune.comp.maxlevel. A value of zero means there is no limit, this is the
default value.

maxcompcpuusage <number>
Sets the maximum CPU usage HAProxy can reach before stopping the compression
for new requests or decreasing the compression level of current requests.
It works like 'maxcomprate' but measures CPU usage instead of incoming data
bandwidth. The value is expressed in percent of the CPU used by haproxy. In
case of multiple processes (nbproc > 1), each process manages its individual
usage. A value of 100 disable the limit. The default value is 100. Setting
a lower value will prevent the compression work from slowing the whole
process down and from introducing high latencies.

maxpipes <number>
Sets the maximum per-process number of pipes to <number>. Currently, pipes
are only used by kernel-based tcp splicing. Since a pipe contains two file
descriptors, the "ulimit-n" value will be increased accordingly. The default
value is maxconn/4, which seems to be more than enough for most heavy usages.
The splice code dynamically allocates and releases pipes, and can fall back
to standard copy, so setting this value too low may only impact performance.

maxsessrate <number>
Sets the maximum per-process number of sessions per second to <number>.
Proxies will stop accepting connections when this limit is reached. It can be
used to limit the global capacity regardless of each frontend capacity. It is
important to note that this can only be used as a service protection measure,
as there will not necessarily be a fair share between frontends when the
limit is reached, so it's a good idea to also limit each frontend to some
value close to its expected share. Also, lowering tune.maxaccept can improve
fairness.

maxsslconn <number>
Sets the maximum per-process number of concurrent SSL connections to
<number>. By default there is no SSL-specific limit, which means that the
global maxconn setting will apply to all connections. Setting this limit
avoids having openssl use too much memory and crash when malloc returns NULL
(since it unfortunately does not reliably check for such conditions). Note
that the limit applies both to incoming and outgoing connections, so one
connection which is deciphered then ciphered accounts for 2 SSL connections.
If this value is not set, but a memory limit is enforced, this value will be
automatically computed based on the memory limit, maxconn, the buffer size,
memory allocated to compression, SSL cache size, and use of SSL in either
frontends, backends or both. If neither maxconn nor maxsslconn are specified
when there is a memory limit, haproxy will automatically adjust these values
so that 100% of the connections can be made over SSL with no risk, and will

```
1041    consider the sides where it is enabled (frontend, backend, both).
1042
1043  maxsslrate <number>
1044    Sets the maximum per-process number of SSL sessions per second to <number>.
1045    SSL listeners will stop accepting connections when this limit is reached. It
1046    can be used to limit the global SSL CPU usage regardless of each frontend
1047    capacity. It is important to note that this can only be used as a service
1048    protection measure, as there will not necessarily be a fair share between
1049    frontends when the limit is reached, so it's a good idea to also limit each
1050    frontend to some value close to its expected share. It is also important to
1051    note that the sessions are accounted before they enter the SSL stack and not
1052    after, which also protects the stack against bad handshakes. Also, lowering
1053    tune.maxaccept can improve fairness.
1054
1055  maxzlibmem <number>
1056    Sets the maximum amount of RAM in megabytes per process usable by the zlib.
1057    When the maximum amount is reached, future sessions will not compress as long
1058    as RAM is unavailable. When sets to 0, there is no limit.
1059    The default value is 0. The value is available in bytes on the UNIX socket
1060    with "show info" on the line "MaxZlibMemUsage", the memory used by zlib is
1061    "ZlibMemUsage" in bytes.
1062
1063  noepoll
1064    Disables the use of the "epoll" event polling system on Linux. It is
1065    equivalent to the command-line argument "-de". The next polling system
1066    used will generally be "poll". See also "nopoll".
1067
1068  nokqueue
1069    Disables the use of the "kqueue" event polling system on BSD. It is
1070    equivalent to the command-line argument "-dk". The next polling system
1071    used will generally be "poll". See also "nopoll".
1072
1073  nopoll
1074    Disables the use of the "poll" event polling system. It is equivalent to the
1075    command-line argument "-dp". The next polling system used will be "select".
1076    It should never be needed to disable "poll" since it's available on all
1077    platforms supported by HAProxy. See also "nokqueue" and "noepoll".
1078
1079  nosplice
1080    Disables the use of kernel tcp splicing between sockets on Linux. It is
1081    equivalent to the command line argument "-dS".  Data will then be copied
1082    using conventional and more portable recv/send calls. Kernel tcp splicing is
1083    limited to some very recent instances of kernel 2.6. Most versions between
1084    2.6.25 and 2.6.28 are buggy and will forward corrupted data, so they must not
1085    be used. This option makes it easier to globally disable kernel splicing in
1086    case of doubt. See also "option splice-auto", "option splice-request" and
1087    "option splice-response".
1088
1089  nogetaddrinfo
1090    Disables the use of getaddrinfo(3) for name resolving. It is equivalent to
1091    the command line argument "-dG". Deprecated gethostbyname(3) will be used.
1092
1093  spread-checks <0..50, in percent>
1094    Sometimes it is desirable to avoid sending agent and health checks to
1095    servers at exact intervals, for instance when many logical servers are
1096    located on the same physical server. With the help of this parameter, it
1097    becomes possible to add some randomness in the check interval between 0
1098    and +/- 50%. A value between 2 and 5 seems to show good results. The
1099    default value remains at 0.
1100
1101  tune.buffers.limit <number>
1102    Sets a hard limit on the number of buffers which may be allocated per process.
1103    The default value is zero which means unlimited. The minimum non-zero value
1104    will always be greater than "tune.buffers.reserve" and should ideally always
1105    be about twice as large. Forcing this value can be particularly useful to
```



```
1106    limit the amount of memory a process may take, while retaining a sane
1107    behaviour. When this limit is reached, sessions which need a buffer wait for
1108    another one to be released by another session. Since buffers are dynamically
1109    allocated and released, the waiting time is very short and not perceptible
1110    provided that limits remain reasonable. In fact sometimes reducing the limit
1111    may even increase performance by increasing the CPU cache's efficiency. Tests
1112    have shown good results on average HTTP traffic with a limit to 1/10 of the
1113    expected global maxconn setting, which also significantly reduces memory
1114    usage. The memory savings come from the fact that a number of connections
1115    will not allocate 2*tune.bufsize. It is best not to touch this value unless
1116    advised to do so by an haproxy core developer.
1117
1118  tune.buffers.reserve <number>
1119    Sets the number of buffers which are pre-allocated and reserved for use only
1120    during memory shortage conditions resulting in failed memory allocations. The
1121    minimum value is 2 and is also the default. There is no reason a user would
1122    want to change this value, it's mostly aimed at haproxy core developers.
1123
1124  tune.bufsize <number>
1125    Sets the buffer size to this size (in bytes). Lower values allow more
1126    sessions to coexist in the same amount of RAM, and higher values allow some
1127    applications with very large cookies to work. The default value is 16384 and
1128    can be changed at build time. It is strongly recommended not to change this
1129    from the default value, as very low values will break some services such as
1130    statistics, and values larger than default size will increase memory usage,
1131    possibly causing the system to run out of memory. At least the global maxconn
1132    parameter should be decreased by the same factor as this one is increased.
1133    If HTTP request is larger than (tune.bufsize - tune.maxrewrite), haproxy will
1134    return HTTP 400 (Bad Request) error. Similarly if an HTTP response is larger
1135    than this size, haproxy will return HTTP 502 (Bad Gateway).
1136
1137  tune.chksize <number>
1138    Sets the check buffer size to this size (in bytes). Higher values may help
1139    find string or regex patterns in very large pages, though doing so may imply
1140    more memory and CPU usage. The default value is 16384 and can be changed at
1141    build time. It is not recommended to change this value, but to use better
1142    checks whenever possible.
1143
1144  tune.comp.maxlevel <number>
1145    Sets the maximum compression level. The compression level affects CPU
1146    usage during compression. This value affects CPU usage during compression.
1147    Each session using compression initializes the compression algorithm with
1148    this value. The default value is 1.
1149
1150  tune.http.cookielen <number>
1151    Sets the maximum length of captured cookies. This is the maximum value that
1152    the "capture cookie xxx len yyy" will be allowed to take, and any upper value
1153    will automatically be truncated to this one. It is important not to set too
1154    high a value because all cookie captures still allocate this size whatever
1155    their configured value (they share a same pool). This value is per request
1156    per response, so the memory allocated is twice this value per connection.
1157    When not specified, the limit is set to 63 characters. It is recommended not
1158    to change this value.
1159
1160  tune.http.maxhdr <number>
1161    Sets the maximum number of headers in a request. When a request comes with a
1162    number of headers greater than this value (including the first line), it is
1163    rejected with a "400 Bad Request" status code. Similarly, too large responses
1164    are blocked with "502 Bad Gateway". The default value is 101, which is enough
1165    for all usages, considering that the widely deployed Apache server uses the
1166    same limit. It can be useful to push this limit further to temporarily allow
1167    a buggy application to work by the time it gets fixed. Keep in mind that each
1168    new header consumes 32bits of memory for each session, so don't push this
1169    limit too high.
1170
```

```
tune.idletimer <timeout>
  Sets the duration after which haproxy will consider that an empty buffer is
  probably associated with an idle stream. This is used to optimally adjust
  some packet sizes while forwarding large and small data alternatively. The
  decision to use splice() or to send large buffers in SSL is modulated by this
  parameter. The value is in milliseconds between 0 and 65535. A value of zero
  means that haproxy will not try to detect idle streams. The default is 1000,
  which seems to correctly detect end user pauses (eg: read a page before
  clicking). There should be not reason for changing this value. Please check
  tune.ssl.maxrecord below.

tune.lua.forced-yield <number>
  This directive forces the Lua engine to execute a yield each <number> of
  instructions executed. This permits interruptng a long script and allows the
  HAProxy scheduler to process other tasks like accepting connections or
  forwarding traffic. The default value is 10000 instructions. If HAProxy often
  executes some Lua code but more reactivity is required, this value can be
  lowered. If the Lua code is quite long and its result is absolutely required
  to process the data, the <number> can be increased.

tune.lua.maxmem
  Sets the maximum amount of RAM in megabytes per process usable by Lua. By
  default it is zero which means unlimited. It is important to set a limit to
  ensure that a bug in a script will not result in the system running out of
  memory.

tune.lua.session-timeout <timeout>
  This is the execution timeout for the Lua sessions. This is useful for
  preventing infinite loops or spending too much time in Lua. This timeout
  counts only the pure Lua runtime. If the Lua does a sleep, the sleep is
  not taked in account. The default timeout is 4s.

tune.lua.task-timeout <timeout>
  Purpose is the same as "tune.lua.session-timeout", but this timeout is
  dedicated to the tasks. By default, this timeout isn't set because a task may
  remain alive during of the lifetime of HAProxy. For example, a task used to
  check servers.

tune.lua.service-timeout <timeout>
  This is the execution timeout for the Lua services. This is useful for
  preventing infinite loops or spending too much time in Lua. This timeout
  counts only the pure Lua runtime. If the Lua does a sleep, the sleep is
  not taked in account. The default timeout is 4s.

tune.maxaccept <number>
  Sets the maximum number of consecutive connections a process may accept in a
  row before switching to other work. In single process mode, higher numbers
  give better performance at high connection rates. However in multi-process
  modes, keeping a bit of fairness between processes generally is better to
  increase performance. This value applies individually to each listener, so
  that the number of processes a listener is bound to is taken into account.
  This value defaults to 64. In multi-process mode, it is divided by twice
  the number of processes the listener is bound to. Setting this value to -1
  completely disables the limitation. It should normally not be needed to tweak
  this value.

tune.maxpollevents <number>
  Sets the maximum amount of events that can be processed at once in a call to
  the polling system. The default value is adapted to the operating system. It
  has been noticed that reducing it below 200 tends to slightly decrease
  latency at the expense of network bandwidth, and increasing it above 200
  tends to trade latency for slightly increased bandwidth.

tune.maxrewrite <number>
  Sets the reserved buffer space to this size in bytes. The reserved space is
```



```
  used for header rewriting or appending. The first reads on sockets will never
  fill more than bufsize-maxrewrite. Historically it has defaulted to half of
  bufsize, though that does not make much sense since there are rarely large
  numbers of headers to add. Setting it too high prevents processing of large
  requests or responses. Setting it too low prevents addition of new headers
  to already large requests or to POST requests. It is generally wise to set it
  to about 1024. It is automatically readjusted to half of bufsize if it is
  larger than that. This means you don't have to worry about it when changing
  bufsize.

tune.pattern.cache-size <number>
  Sets the size of the pattern lookup cache to <number> entries. This is an LRU
  cache which reminds previous lookups and their results. It is used by ACLs
  and maps on slow pattern lookups, namely the ones using the "sub", "reg",
  "dir", "dom", "end", "bin" match methods as well as the case-insensitive
  strings. It applies to pattern expressions which means that it will be able
  to memorize the result of a lookup among all the patterns specified on a
  configuration line (including all those loaded from files). It automatically
  invalidates entries which are updated using HTTP actions or on the CLI. The
  default cache size is set to 10000 entries, which limits its footprint to
  about 5 MB on 32-bit systems and 8 MB on 64-bit systems. There is a very low
  risk of collision in this cache, which is in the order of the size of the
  cache divided by 2^64. Typically, at 10000 requests per second with the
  default cache size of 10000 entries, there's 1% chance that a brute force
  attack could cause a single collision after 60 years, or 0.1% after 6 years.
  This is considered much lower than the risk of a memory corruption caused by
  aging components. If this is not acceptable, the cache can be disabled by
  setting this parameter to 0.

tune.pipesize <number>
  Sets the kernel pipe buffer size to this size (in bytes). By default, pipes
  are the default size for the system. But sometimes when using TCP splicing,
  it can improve performance to increase pipe sizes, especially if it is
  suspected that pipes are not filled and that many calls to splice() are
  performed. This has an impact on the kernel's memory footprint, so this must
  not be changed if impacts are not understood.

tune.rcvbuf.client <number>
tune.rcvbuf.server <number>
  Forces the kernel socket receive buffer size on the client or the server side
  to the specified value in bytes. This value applies to all TCP/HTTP frontends
  and backends. It should normally never be set, and the default size (0) lets
  the kernel autotune this value depending on the amount of available memory.
  However it can sometimes help to set it to very low values (eg: 4096) in
  order to save kernel memory by preventing it from buffering too large amounts
  of received data. Lower values will significantly increase CPU usage though.

tune.sndbuf.client <number>
tune.sndbuf.server <number>
  Forces the kernel socket send buffer size on the client or the server side to
  the specified value in bytes. This value applies to all TCP/HTTP frontends
  and backends. It should normally never be set, and the default size (0) lets
  the kernel autotune this value depending on the amount of available memory.
  However it can sometimes help to set it to very low values (eg: 4096) in
  order to save kernel memory by preventing it from buffering too large amounts
  of received data. Lower values will significantly increase CPU usage though.
  Another use case is to prevent write timeouts with extremely slow clients due
  to the kernel waiting for a large part of the buffer to be read before
  notifying haproxy again.

tune.ssl.cachesize <number>
  Sets the size of the global SSL session cache, in a number of blocks. A block
  is large enough to contain an encoded session without peer certificate.
  An encoded session with peer certificate is stored in multiple blocks
  depending on the size of the peer certificate. A block uses approximately
```

```
   200 bytes of memory. The default value may be forced at build time, otherwise
   defaults to 20000.  When the cache is full, the most idle entries are purged
   and reassigned. Higher values reduce the occurrence of such a purge, hence
   the number of CPU-intensive SSL handshakes by ensuring that all users keep
   their session as long as possible. All entries are pre-allocated upon startup
   and are shared between all processes if "nbproc" is greater than 1. Setting
   this value to 0 disables the SSL session cache.

tune.ssl.force-private-cache
   This boolean disables SSL session cache sharing between all processes. It
   should normally not be used since it will force many renegotiations due to
   clients hitting a random process. But it may be required on some operating
   systems where none of the SSL cache synchronization method may be used. In
   this case, adding a first layer of hash-based load balancing before the SSL
   layer might limit the impact of the lack of session sharing.

tune.ssl.lifetime <timeout>
   Sets how long a cached SSL session may remain valid. This time is expressed
   in seconds and defaults to 300 (5 min). It is important to understand that it
   does not guarantee that sessions will last that long, because if the cache is
   full, the longest idle sessions will be purged despite their configured
   lifetime. The real usefulness of this setting is to prevent sessions from
   being used for too long.

tune.ssl.maxrecord <number>
   Sets the maximum amount of bytes passed to SSL_write() at a time. Default
   value 0 means there is no limit. Over SSL/TLS, the client can decipher the
   data only once it has received a full record. With large records, it means
   that clients might have to download up to 16kB of data before starting to
   process them. Limiting the value can improve page load times on browsers
   located over high latency or low bandwidth networks. It is suggested to find
   optimal values which fit into 1 or 2 TCP segments (generally 1448 bytes over
   Ethernet with TCP timestamps enabled, or 1460 when timestamps are disabled),
   keeping in mind that SSL/TLS add some overhead. Typical values of 1419 and
   2859 gave good results during tests. Use "strace -e trace=write" to find the
   best value. Haproxy will automatically switch to this setting after an idle
   stream has been detected (see tune.idletimer above).

tune.ssl.default-dh-param <number>
   Sets the maximum size of the Diffie-Hellman parameters used for generating
   the ephemeral/temporary Diffie-Hellman key in case of DHE key exchange. The
   final size will try to match the size of the server's RSA (or DSA) key (e.g,
   a 2048 bits temporary DH key for a 2048 bits RSA key), but will not exceed
   this maximum value. Default value if 1024. Only 1024 or higher values are
   allowed. Higher values will increase the CPU load, and values greater than
   1024 bits are not supported by Java 7 and earlier clients. This value is not
   used if static Diffie-Hellman parameters are supplied either directly
   in the certificate file or by using the ssl-dh-param-file parameter.

tune.ssl.ssl-ctx-cache-size <number>
   Sets the size of the cache used to store generated certificates to <number>
   entries. This is a LRU cache. Because generating a SSL certificate
   dynamically is expensive, they are cached. The default cache size is set to
   1000 entries.

tune.vars.global-max-size <size>
tune.vars.reqres-max-size <size>
tune.vars.sess-max-size <size>
tune.vars.txn-max-size <size>
   These four tunes helps to manage the allowed amount of memory used by the
   variables system. "global" limits the memory for all the systems. "sess" limit
   the memory by session, "txn" limits the memory by transaction and "reqres"
   limits the memory for each request or response processing. during the
   accounting, "sess" embbed "txn" and "txn" embed "reqres".
```



```
   By example, we considers that "tune.vars.sess-max-size" is fixed to 100,
   "tune.vars.txn-max-size" is fixed to 100, "tune.vars.reqres-max-size" is
   also fixed to 100. If we create a variable "txn.var" that contains 100 bytes,
   we cannot create any more variable in the other contexts.

tune.zlib.memlevel <number>
   Sets the memLevel parameter in zlib initialization for each session. It
   defines how much memory should be allocated for the internal compression
   state. A value of 1 uses minimum memory but is slow and reduces compression
   ratio, a value of 9 uses maximum memory for optimal speed.  Can be a value
   between 1 and 9. The default value is 8.

tune.zlib.windowsize <number>
   Sets the window size (the size of the history buffer) as a parameter of the
   zlib initialization for each session. Larger values of this parameter result
   in better compression at the expense of memory usage.  Can be a value between
   8 and 15.  The default value is 15.

3.3. Debugging
--------------

debug
   Enables debug mode which dumps to stdout all exchanges, and disables forking
   into background. It is the equivalent of the command-line argument "-d". It
   should never be used in a production configuration since it may prevent full
   system startup.

quiet
   Do not display any message during startup. It is equivalent to the command-
   line argument "-q".


3.4. Userlists
--------------
It is possible to control access to frontend/backend/listen sections or to
http stats by allowing only authenticated and authorized users. To do this,
it is required to create at least one userlist and to define users.

userlist <listname>
   Creates new userlist with name <listname>. Many independent userlists can be
   used to store authentication & authorization data for independent customers.

group <groupname> [users <user>,<user>,(...)]
   Adds group <groupname> to the current userlist. It is also possible to
   attach users to this group by using a comma separated list of names
   proceeded by "users" keyword.

user <username> [password|insecure-password <password>]
                [groups <group>,<group>,(...)]
   Adds user <username> to the current userlist. Both secure (encrypted) and
   insecure (unencrypted) passwords can be used. Encrypted passwords are
   evaluated using the crypt(3) function so depending of the system's
   capabilities, different algorithms are supported. For example modern Glibc
   based Linux system supports MD5, SHA-256, SHA-512 and of course classic,
   DES-based method of encrypting passwords.


   Example:
         userlist L1
           group G1 users tiger,scott
           group G2 users xdb,scott

           user tiger password $6$k6y3o.eP$JlKBx9za9667qe4(...)xHSwRv6J.C0/D7cV91
           user scott insecure-password elgato
           user xdb insecure-password hello
```

```
1431
1432          userlist L2
1433            group G1
1434            group G2
1435
1436            user tiger password $6$k6y3o.eP$JlKBx(...)xHSwRv6J.C0/D7cV91 groups G1
1437            user scott insecure-password elgato groups G1,G2
1438            user xdb insecure-password hello groups G2
1439
1440    Please note that both lists are functionally identical.
1441
1442
1443  3.5. Peers
1444  ----------
1445  It is possible to propagate entries of any data-types in stick-tables between
1446  several haproxy instances over TCP connections in a multi-master fashion. Each
1447  instance pushes its local updates and insertions to remote peers. The pushed
1448  values overwrite remote ones without aggregation. Interrupted exchanges are
1449  automatically detected and recovered from the last known point.
1450  In addition, during a soft restart, the old process connects to the new one
1451  using such a TCP connection to push all its entries before the new process
1452  tries to connect to other peers. That ensures very fast replication during a
1453  reload, it typically takes a fraction of a second even for large tables.
1454  Note that Server IDs are used to identify servers remotely, so it is important
1455  that configurations look similar or at least that the same IDs are forced on
1456  each server on all participants.
1457
1458  peers <peersect>
1459    Creates a new peer list with name <peersect>. It is an independent section,
1460    which is referenced by one or more stick-tables.
1461
1462  disabled
1463    Disables a peers section. It disables both listening and any synchronization
1464    related to this section. This is provided to disable synchronization of stick
1465    tables without having to comment out all "peers" references.
1466
1467  enable
1468    This re-enables a disabled peers section which was previously disabled.
1469
1470  peer <peername> <ip>:<port>
1471    Defines a peer inside a peers section.
1472    If <peername> is set to the local peer name (by default hostname, or forced
1473    using "-L" command line option), haproxy will listen for incoming remote peer
1474    connection on <ip>:<port>. Otherwise, <ip>:<port> defines where to connect to
1475    to join the remote peer, and <peername> is used at the protocol level to
1476    identify and validate the remote peer on the server side.
1477
1478    During a soft restart, local peer <ip>:<port> is used by the old instance to
1479    connect the new one and initiate a complete replication (teaching process).
1480
1481    It is strongly recommended to have the exact same peers declaration on all
1482    peers and to only rely on the "-L" command line argument to change the local
1483    peer name. This makes it easier to maintain coherent configuration files
1484    across all peers.
1485
1486    You may want to reference some environment variables in the address
1487    parameter, see section 2.3 about environment variables.
1488
1489    Example:
1490      peers mypeers
1491          peer haproxy1 192.168.0.1:1024
1492          peer haproxy2 192.168.0.2:1024
1493          peer haproxy3 10.2.0.1:1024
1494
1495      backend mybackend
```



```
1496          mode tcp
1497          balance roundrobin
1498          stick-table type ip size 20k peers mypeers
1499          stick on src
1500
1501          server srv1 192.168.0.30:80
1502          server srv2 192.168.0.31:80
1503
1504
1505  3.6. Mailers
1506  ------------
1507  It is possible to send email alerts when the state of servers changes.
1508  If configured email alerts are sent to each mailer that is configured
1509  in a mailers section. Email is sent to mailers using SMTP.
1510
1511  mailers <mailersect>
1512    Creates a new mailer list with the name <mailersect>. It is an
1513    independent section which is referenced by one or more proxies.
1514
1515  mailer <mailername> <ip>:<port>
1516    Defines a mailer inside a mailers section.
1517
1518    Example:
1519      mailers mymailers
1520          mailer smtp1 192.168.0.1:587
1521          mailer smtp2 192.168.0.2:587
1522
1523      backend mybackend
1524          mode tcp
1525          balance roundrobin
1526
1527          email-alert mailers mymailers
1528          email-alert from test1@horms.org
1529          email-alert to test2@horms.org
1530
1531          server srv1 192.168.0.30:80
1532          server srv2 192.168.0.31:80
1533
1534
1535  4. Proxies
1536  ----------
1537
1538  Proxy configuration can be located in a set of sections :
1539   - defaults [<name>]
1540   - frontend <name>
1541   - backend  <name>
1542   - listen   <name>
1543
1544  A "defaults" section sets default parameters for all other sections following
1545  its declaration. Those default parameters are reset by the next "defaults"
1546  section. See below for the list of parameters which can be set in a "defaults"
1547  section. The name is optional but its use is encouraged for better readability.
1548
1549  A "frontend" section describes a set of listening sockets accepting client
1550  connections.
1551
1552  A "backend" section describes a set of servers to which the proxy will connect
1553  to forward incoming connections.
1554
1555  A "listen" section defines a complete proxy with its frontend and backend
1556  parts combined in one section. It is generally useful for TCP-only traffic.
1557
1558  All proxy names must be formed from upper and lower case letters, digits,
1559  '-' (dash), '_' (underscore) , '.' (dot) and ':' (colon). ACL names are
1560  case-sensitive, which means that "www" and "WWW" are two different proxies.
```

Historically, all proxy names could overlap, it just caused troubles in the logs. Since the introduction of content switching, it is mandatory that two proxies with overlapping capabilities (frontend/backend) have different names. However, it is still permitted that a frontend and a backend share the same name, as this configuration seems to be commonly encountered.

Right now, two major proxy modes are supported : "tcp", also known as layer 4, and "http", also known as layer 7. In layer 4 mode, HAProxy simply forwards bidirectional traffic between two sides. In layer 7 mode, HAProxy analyzes the protocol, and can interact with it by allowing, blocking, switching, adding, modifying, or removing arbitrary contents in requests or responses, based on arbitrary criteria.

In HTTP mode, the processing applied to requests and responses flowing over a connection depends in the combination of the frontend's HTTP options and the backend's. HAProxy supports 5 connection modes :

  - KAL : keep alive ("option http-keep-alive") which is the default mode : all
    requests and responses are processed, and connections remain open but idle
    between responses and new requests.

  - TUN: tunnel ("option http-tunnel") : this was the default mode for versions
    1.0 to 1.5-dev21 : only the first request and response are processed, and
    everything else is forwarded with no analysis at all. This mode should not
    be used as it creates lots of trouble with logging and HTTP processing.

  - PCL: passive close ("option httpclose") : exactly the same as tunnel mode,
    but with "Connection: close" appended in both directions to try to make
    both ends close after the first request/response exchange.

  - SCL: server close ("option http-server-close") : the server-facing
    connection is closed after the end of the response is received, but the
    client-facing connection remains open.

  - FCL: forced close ("option forceclose") : the connection is actively closed
    after the end of the response.

The effective mode that will be applied to a connection passing through a frontend and a backend can be determined by both proxy modes according to the following matrix, but in short, the modes are symmetric, keep-alive is the weakest option and force close is the strongest.

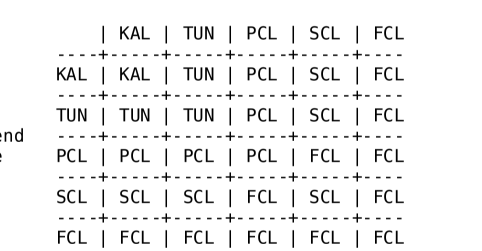

4.1. Proxy keywords matrix

The following list of keywords is supported. Most of them may only be used in a limited set of section types. Some of them are marked as "deprecated" because they are inherited from an old syntax which may be confusing or functionally

limited, and there are new recommended keywords to replace them. Keywords marked with "(*)" can be optionally inverted using the "no" prefix, eg. "no option contstats". This makes sense when the option has been enabled by default and must be disabled for a specific instance. Such options may also be prefixed with "default" in order to restore default settings regardless of what has been specified in a previous "defaults" section.

keyword		defaults	frontend	listen	backend
acl		-	X	X	X
appsession		-	-	-	-
backlog		X	X	X	-
balance		X	-	X	X
bind		-	X	X	-
bind-process		X	X	X	X
block		-	X	X	X
capture cookie		-	X	X	-
capture request header		-	X	X	-
capture response header		-	X	X	-
clitimeout	(deprecated)	X	X	X	-
compression		X	X	X	X
contimeout	(deprecated)	X	-	X	X
cookie		X	-	X	X
declare capture		-	X	X	-
default-server		X	-	X	X
default_backend		X	X	X	-
description		-	X	X	X
disabled		X	X	X	X
dispatch		-	-	X	X
email-alert from		X	X	X	X
email-alert level		X	X	X	X
email-alert mailers		X	X	X	X
email-alert myhostname		X	X	X	X
email-alert to		X	X	X	X
enabled		X	X	X	X
errorfile		X	X	X	X
errorloc		X	X	X	X
errorloc302		X	X	X	X
-- keyword -------------------------		defaults	- frontend	- listen	-- backend -
errorloc303		X	X	X	X
force-persist		-	X	X	X
fullconn		X	-	X	X
grace		X	X	X	X
hash-type		X	-	X	X
http-check disable-on-404		X	-	X	X
http-check expect		-	-	X	X
http-check send-state		X	-	X	X
http-request		-	X	X	X
http-response		-	X	X	X
http-reuse		X	-	X	X
http-send-name-header		-	-	X	X
id		-	X	X	X
ignore-persist		-	X	X	X
load-server-state-from-file		X	-	X	X
log	(*)	X	X	X	X
log-format		X	X	X	-
log-format-sd		X	X	X	-
log-tag		X	X	X	X
max-keep-alive-queue		X	-	X	X
maxconn		X	X	X	-
mode		X	X	X	X
monitor fail		-	X	X	-
monitor-net		X	X	X	-
monitor-uri		X	X	X	-

```
1691  option abortonclose                  (*)  X          -          X          X
1692  option accept-invalid-http-request   (*)  X          X          X          -
1693  option accept-invalid-http-response  (*)  X          -          X          X
1694  option allbackups                    (*)  X          -          X          X
1695  option checkcache                    (*)  X          -          X          X
1696  option clitcpka                      (*)  X          X          X          -
1697  option contstats                     (*)  X          X          X          -
1698  option dontlog-normal                (*)  X          X          X          -
1699  option dontlognull                   (*)  X          X          X          -
1700  option forceclose                    (*)  X          X          X          X
1701  -- keyword -------------------------- defaults - frontend - listen -- backend -
1702  option forwardfor                         X          X          X          X
1703  option http-buffer-request           (*)  X          X          X          X
1704  option http-ignore-probes            (*)  X          X          X          -
1705  option http-keep-alive               (*)  X          X          X          X
1706  option http-no-delay                 (*)  X          X          X          X
1707  option http-pretend-keepalive        (*)  X          X          X          X
1708  option http-server-close             (*)  X          X          X          X
1709  option http-tunnel                   (*)  X          X          X          X
1710  option http-use-proxy-header         (*)  X          X          X          -
1711  option httpchk                            X          -          X          X
1712  option httpclose                     (*)  X          X          X          X
1713  option httplog                            X          X          X          X
1714  option http_proxy                    (*)  X          X          X          X
1715  option independent-streams           (*)  X          X          X          X
1716  option ldap-check                         X          -          X          X
1717  option external-check                     X          -          X          X
1718  option log-health-checks             (*)  X          -          X          X
1719  option log-separate-errors           (*)  X          X          X          -
1720  option logasap                       (*)  X          X          X          -
1721  option mysql-check                        X          -          X          X
1722  option nolinger                      (*)  X          X          X          X
1723  option originalto                         X          X          X          X
1724  option persist                       (*)  X          -          X          X
1725  option pgsql-check                        X          -          X          X
1726  option prefer-last-server            (*)  X          -          X          X
1727  option redispatch                    (*)  X          -          X          X
1728  option redis-check                        X          -          X          X
1729  option smtpchk                            X          -          X          X
1730  option socket-stats                  (*)  X          X          X          -
1731  option splice-auto                   (*)  X          X          X          X
1732  option splice-request                (*)  X          X          X          X
1733  option splice-response               (*)  X          X          X          X
1734  option srvtcpka                      (*)  X          -          X          X
1735  option ssl-hello-chk                      X          -          X          X
1736  -- keyword -------------------------- defaults - frontend - listen -- backend -
1737  option tcp-check                          X          -          X          X
1738  option tcp-smart-accept              (*)  X          X          X          -
1739  option tcp-smart-connect             (*)  X          -          X          X
1740  option tcpka                              X          X          X          X
1741  option tcplog                             X          X          X          X
1742  option transparent                   (*)  X          -          X          X
1743  external-check command                    X          -          X          X
1744  external-check path                       X          -          X          X
1745  persist rdp-cookie                        X          -          X          X
1746  rate-limit sessions                       X          X          X          -
1747  redirect                                  -          X          X          X
1748  redisp                     (deprecated)   X          -          X          X
1749  redispatch                 (deprecated)   X          -          X          X
1750  reqadd                                    -          X          X          X
1751  reqallow                                  -          X          X          X
1752  reqdel                                    -          X          X          X
1753  reqdeny                                   -          X          X          X
1754  reqiallow                                 -          X          X          X
1755  reqidel                                   -          X          X          X
```



```
1756  reqideny                                  -          X          X          X
1757  reqipass                                  -          X          X          X
1758  reqirep                                   -          X          X          X
1759  reqitarpit                                -          X          X          X
1760  reqpass                                   -          X          X          X
1761  reqrep                                    -          X          X          X
1762  -- keyword -------------------------- defaults - frontend - listen -- backend -
1763  reqtarpit                                 -          X          X          X
1764  retries                                   X          -          X          X
1765  rspadd                                    -          X          X          X
1766  rspdel                                    -          X          X          X
1767  rspdeny                                   -          X          X          X
1768  rspidel                                   -          X          X          X
1769  rspideny                                  -          X          X          X
1770  rspirep                                   -          X          X          X
1771  rsprep                                    -          X          X          X
1772  server                                    -          -          X          X
1773  server-state-file-name                    X          -          X          X
1774  source                                    X          -          X          X
1775  srvtimeout                 (deprecated)   X          -          X          X
1776  stats admin                               -          X          X          X
1777  stats auth                                X          X          X          X
1778  stats enable                              X          X          X          X
1779  stats hide-version                        X          X          X          X
1780  stats http-request                        -          X          X          X
1781  stats realm                               X          X          X          X
1782  stats refresh                             X          X          X          X
1783  stats scope                               X          X          X          X
1784  stats show-desc                           X          X          X          X
1785  stats show-legends                        X          X          X          X
1786  stats show-node                           X          X          X          X
1787  stats uri                                 X          X          X          X
1788  -- keyword -------------------------- defaults - frontend - listen -- backend -
1789  stick match                               -          -          X          X
1790  stick on                                  -          -          X          X
1791  stick store-request                       -          -          X          X
1792  stick store-response                      -          -          X          X
1793  stick-table                               -          -          X          X
1794  tcp-check connect                         -          -          X          X
1795  tcp-check expect                          -          -          X          X
1796  tcp-check send                            -          -          X          X
1797  tcp-check send-binary                     -          -          X          X
1798  tcp-request connection                    -          X          X          -
1799  tcp-request content                       -          X          X          X
1800  tcp-request inspect-delay                 -          X          X          X
1801  tcp-response content                      -          -          X          X
1802  tcp-response inspect-delay                -          -          X          X
1803  timeout check                             X          -          X          X
1804  timeout client                            X          X          X          -
1805  timeout client-fin                        X          X          X          -
1806  timeout clitimeout         (deprecated)   X          X          X          -
1807  timeout connect                           X          -          X          X
1808  timeout contimeout         (deprecated)   X          -          X          X
1809  timeout http-keep-alive                   X          X          X          X
1810  timeout http-request                      X          X          X          X
1811  timeout queue                             X          -          X          X
1812  timeout server                            X          -          X          X
1813  timeout server-fin                        X          -          X          X
1814  timeout srvtimeout         (deprecated)   X          -          X          X
1815  timeout tarpit                            X          X          X          X
1816  timeout tunnel                            X          -          X          X
1817  transparent                (deprecated)   X          -          X          X
1818  unique-id-format                          X          X          X          -
1819  unique-id-header                          X          X          X          -
1820  use_backend                               -          X          X          -
```

```
use-server                                -          -         X         X
----------------------------------+----------+----------+---------+---------
 keyword                          | defaults | frontend | listen  | backend
```

4.2. Alphabetically sorted keywords reference

This section provides a description of each keyword and its usage.

```
acl <aclname> <criterion> [flags] [operator] <value> ...
  Declare or complete an access list.
  May be used in sections :   defaults | frontend | listen | backend
                                 no    |    yes   |   yes  |   yes
  Example:
        acl invalid_src  src          0.0.0.0/7 224.0.0.0/3
        acl invalid_src  src_port     0:1023
        acl local_dst    hdr(host) -i localhost

  See section 7 about ACL usage.


appsession <cookie> len <length> timeout <holdtime>
           [request-learn] [prefix] [mode <path-parameters|query-string>]
  Define session stickiness on an existing application cookie.
  May be used in sections :   defaults | frontend | listen | backend
                                 no    |    no    |   yes  |   yes
  Arguments :
    <cookie>   this is the name of the cookie used by the application and which
               HAProxy will have to learn for each new session.

    <length>   this is the max number of characters that will be memorized and
               checked in each cookie value.

    <holdtime> this is the time after which the cookie will be removed from
               memory if unused. If no unit is specified, this time is in
               milliseconds.

    request-learn
               If this option is specified, then haproxy will be able to learn
               the cookie found in the request in case the server does not
               specify any in response. This is typically what happens with
               PHPSESSID cookies, or when haproxy's session expires before
               the application's session and the correct server is selected.
               It is recommended to specify this option to improve reliability.

    prefix     When this option is specified, haproxy will match on the cookie
               prefix (or URL parameter prefix). The appsession value is the
               data following this prefix.

               Example :
               appsession ASPSESSIONID len 64 timeout 3h prefix

               This will match the cookie ASPSESSIONIDXXXX=XXXXX,
               the appsession value will be XXXX=XXXXX.

    mode       This option allows to change the URL parser mode.
               2 modes are currently supported :
               - path-parameters :
                 The parser looks for the appsession in the path parameters
                 part (each parameter is separated by a semi-colon), which is
                 convenient for JSESSIONID for example.
                 This is the default mode if the option is not set.
               - query-string :
```



```
                 In this mode, the parser will look for the appsession in the
                 query string.

  As of version 1.6, appsessions was removed. It is more flexible and more
  convenient to use stick-tables instead, and stick-tables support multi-master
  replication and data conservation across reloads, which appsessions did not.

  See also : "cookie", "capture cookie", "balance", "stick", "stick-table",
             "ignore-persist", "nbproc" and "bind-process".


backlog <conns>
  Give hints to the system about the approximate listen backlog desired size
  May be used in sections :   defaults | frontend | listen | backend
                                 yes   |    yes   |   yes  |   no
  Arguments :
    <conns>   is the number of pending connections. Depending on the operating
              system, it may represent the number of already acknowledged
              connections, of non-acknowledged ones, or both.

  In order to protect against SYN flood attacks, one solution is to increase
  the system's SYN backlog size. Depending on the system, sometimes it is just
  tunable via a system parameter, sometimes it is not adjustable at all, and
  sometimes the system relies on hints given by the application at the time of
  the listen() syscall. By default, HAProxy passes the frontend's maxconn value
  to the listen() syscall. On systems which can make use of this value, it can
  sometimes be useful to be able to specify a different value, hence this
  backlog parameter.

  On Linux 2.4, the parameter is ignored by the system. On Linux 2.6, it is
  used as a hint and the system accepts up to the smallest greater power of
  two, and never more than some limits (usually 32768).

  See also : "maxconn" and the target operating system's tuning guide.


balance <algorithm> [ <arguments> ]
balance url_param <param> [check_post]
  Define the load balancing algorithm to be used in a backend.
  May be used in sections :   defaults | frontend | listen | backend
                                 yes   |    no    |   yes  |   yes
  Arguments :
    <algorithm> is the algorithm used to select a server when doing load
                balancing. This only applies when no persistence information
                is available, or when a connection is redispatched to another
                server. <algorithm> may be one of the following :

       roundrobin  Each server is used in turns, according to their weights.
                   This is the smoothest and fairest algorithm when the server's
                   processing time remains equally distributed. This algorithm
                   is dynamic, which means that server weights may be adjusted
                   on the fly for slow starts for instance. It is limited by
                   design to 4095 active servers per backend. Note that in some
                   large farms, when a server becomes up after having been down
                   for a very short time, it may sometimes take a few hundreds
                   requests for it to be re-integrated into the farm and start
                   receiving traffic. This is normal, though very rare. It is
                   indicated here in case you would have the chance to observe
                   it, so that you don't worry.

       static-rr   Each server is used in turns, according to their weights.
                   This algorithm is as similar to roundrobin except that it is
                   static, which means that changing a server's weight on the
                   fly will have no effect. On the other hand, it has no design
                   limitation on the number of servers, and when a server goes
```

up, it is always immediately reintroduced into the farm, once the full map is recomputed. It also uses slightly less CPU to run (around -1%).

leastconn The server with the lowest number of connections receives the connection. Round-robin is performed within groups of servers of the same load to ensure that all servers will be used. Use of this algorithm is recommended where very long sessions are expected, such as LDAP, SQL, TSE, etc... but is not very well suited for protocols using short sessions such as HTTP. This algorithm is dynamic, which means that server weights may be adjusted on the fly for slow starts for instance.

first The first server with available connection slots receives the connection. The servers are chosen from the lowest numeric identifier to the highest (see server parameter "id"), which defaults to the server's position in the farm. Once a server reaches its maxconn value, the next server is used. It does not make sense to use this algorithm without setting maxconn. The purpose of this algorithm is to always use the smallest number of servers so that extra servers can be powered off during non-intensive hours. This algorithm ignores the server weight, and brings more benefit to long session such as RDP or IMAP than HTTP, though it can be useful there too. In order to use this algorithm efficiently, it is recommended that a cloud controller regularly checks server usage to turn them off when unused, and regularly checks backend queue to turn new servers on when the queue inflates. Alternatively, using "http-check send-state" may inform servers on the load.

source The source IP address is hashed and divided by the total weight of the running servers to designate which server will receive the request. This ensures that the same client IP address will always reach the same server as long as no server goes down or up. If the hash result changes due to the number of running servers changing, many clients will be directed to a different server. This algorithm is generally used in TCP mode where no cookie may be inserted. It may also be used on the Internet to provide a best-effort stickiness to clients which refuse session cookies. This algorithm is static by default, which means that changing a server's weight on the fly will have no effect, but this can be changed using "hash-type".

uri This algorithm hashes either the left part of the URI (before the question mark) or the whole URI (if the "whole" parameter is present) and divides the hash value by the total weight of the running servers. The result designates which server will receive the request. This ensures that the same URI will always be directed to the same server as long as no server goes up or down. This is used with proxy caches and anti-virus proxies in order to maximize the cache hit rate. Note that this algorithm may only be used in an HTTP backend. This algorithm is static by default, which means that changing a server's weight on the fly will have no effect, but this can be changed using "hash-type".

This algorithm supports two optional parameters "len" and "depth", both followed by a positive integer number. These options may be helpful when it is needed to balance servers based on the beginning of the URI only. The "len" parameter indicates that the algorithm should only consider that many characters at the beginning of the URI to compute the hash. Note that having "len" set to 1 rarely makes sense since most URIs start with a leading "/".

The "depth" parameter indicates the maximum directory depth to be used to compute the hash. One level is counted for each slash in the request. If both parameters are specified, the evaluation stops when either is reached.

url_param The URL parameter specified in argument will be looked up in the query string of each HTTP GET request.

If the modifier "check_post" is used, then an HTTP POST request entity will be searched for the parameter argument, when it is not found in a query string after a question mark ('?') in the URL. The message body will only start to be analyzed once either the advertised amount of data has been received or the request buffer is full. In the unlikely event that chunked encoding is used, only the first chunk is scanned. Parameter values separated by a chunk boundary, may be randomly balanced if at all. This keyword used to support an optional <max_wait> parameter which is now ignored.

If the parameter is found followed by an equal sign ('=') and a value, then the value is hashed and divided by the total weight of the running servers. The result designates which server will receive the request.

This is used to track user identifiers in requests and ensure that a same user ID will always be sent to the same server as long as no server goes up or down. If no value is found or if the parameter is not found, then a round robin algorithm is applied. Note that this algorithm may only be used in an HTTP backend. This algorithm is static by default, which means that changing a server's weight on the fly will have no effect, but this can be changed using "hash-type".

hdr(<name>) The HTTP header <name> will be looked up in each HTTP request. Just as with the equivalent ACL 'hdr()' function, the header name in parenthesis is not case sensitive. If the header is absent or if it does not contain any value, the roundrobin algorithm is applied instead.

An optional 'use_domain_only' parameter is available, for reducing the hash algorithm to the main domain part with some specific headers such as 'Host'. For instance, in the Host value "haproxy.1wt.eu", only "1wt" will be considered.

This algorithm is static by default, which means that changing a server's weight on the fly will have no effect, but this can be changed using "hash-type".

rdp-cookie
rdp-cookie(<name>)
The RDP cookie <name> (or "mstshash" if omitted) will be looked up and hashed for each incoming TCP request. Just as with the equivalent ACL 'req_rdp_cookie()' function, the name is not case-sensitive. This mechanism is useful as a degraded persistence mode, as it makes it possible to always send the same user (or the same session ID) to the same server. If the cookie is not found, the normal roundrobin algorithm is used instead.

Note that for this to work, the frontend must ensure that an RDP cookie is already present in the request buffer. For this you must use 'tcp-request content accept' rule combined with a 'req_rdp_cookie_cnt' ACL.

```
                  This algorithm is static by default, which means that
                  changing a server's weight on the fly will have no effect,
                  but this can be changed using "hash-type".

                  See also the rdp_cookie pattern fetch function.

    <arguments> is an optional list of arguments which may be needed by some
                algorithms. Right now, only "url_param" and "uri" support an
                optional argument.

  The load balancing algorithm of a backend is set to roundrobin when no other
  algorithm, mode nor option have been set. The algorithm may only be set once
  for each backend.

  Examples :
        balance roundrobin
        balance url_param userid
        balance url_param session_id check_post 64
        balance hdr(User-Agent)
        balance hdr(host)
        balance hdr(Host) use_domain_only

  Note: the following caveats and limitations on using the "check_post"
  extension with "url_param" must be considered :

    - all POST requests are eligible for consideration, because there is no way
      to determine if the parameters will be found in the body or entity which
      may contain binary data. Therefore another method may be required to
      restrict consideration of POST requests that have no URL parameters in
      the body. (see acl reqideny http_end)

    - using a <max_wait> value larger than the request buffer size does not
      make sense and is useless. The buffer size is set at build time, and
      defaults to 16 kB.

    - Content-Encoding is not supported, the parameter search will probably
      fail; and load balancing will fall back to Round Robin.

    - Expect: 100-continue is not supported, load balancing will fall back to
      Round Robin.

    - Transfer-Encoding (RFC2616 3.6.1) is only supported in the first chunk.
      If the entire parameter value is not present in the first chunk, the
      selection of server is undefined (actually, defined by how little
      actually appeared in the first chunk).

    - This feature does not support generation of a 100, 411 or 501 response.

    - In some cases, requesting "check_post" MAY attempt to scan the entire
      contents of a message body. Scanning normally terminates when linear
      white space or control characters are found, indicating the end of what
      might be a URL parameter list. This is probably not a concern with SGML
      type message bodies.

  See also : "dispatch", "cookie", "transparent", "hash-type" and "http_proxy".


bind [<address>]:<port_range> [, ...] [param*]
bind /<path> [, ...] [param*]
  Define one or several listening addresses and/or ports in a frontend.
  May be used in sections :   defaults | frontend | listen | backend
                                 no    |    yes   |   yes  |   no
  Arguments :
    <address>     is optional and can be a host name, an IPv4 address, an IPv6
                  address, or '*'. It designates the address the frontend will
```



```
                  listen on. If unset, all IPv4 addresses of the system will be
                  listened on. The same will apply for '*' or the system's
                  special address "0.0.0.0". The IPv6 equivalent is '::'.
                  Optionally, an address family prefix may be used before the
                  address to force the family regardless of the address format,
                  which can be useful to specify a path to a unix socket with
                  no slash ('/'). Currently supported prefixes are :
                    - 'ipv4@'  -> address is always IPv4
                    - 'ipv6@'  -> address is always IPv6
                    - 'unix@'  -> address is a path to a local unix socket
                    - 'abns@'  -> address is in abstract namespace (Linux only).
                      Note: since abstract sockets are not "rebindable", they
                            do not cope well with multi-process mode during
                            soft-restart, so it is better to avoid them if
                            nbproc is greater than 1. The effect is that if the
                            new process fails to start, only one of the old ones
                            will be able to rebind to the socket.
                    - 'fd@<n>' -> use file descriptor <n> inherited from the
                      parent. The fd must be bound and may or may not already
                      be listening.
                  You may want to reference some environment variables in the
                  address parameter, see section 2.3 about environment
                  variables.

    <port_range>  is either a unique TCP port, or a port range for which the
                  proxy will accept connections for the IP address specified
                  above. The port is mandatory for TCP listeners. Note that in
                  the case of an IPv6 address, the port is always the number
                  after the last colon (':'). A range can either be :
                   - a numerical port (ex: '80')
                   - a dash-delimited ports range explicitly stating the lower
                     and upper bounds (ex: '2000-2100') which are included in
                     the range.

                  Particular care must be taken against port ranges, because
                  every <address:port> couple consumes one socket (= a file
                  descriptor), so it's easy to consume lots of descriptors
                  with a simple range, and to run out of sockets. Also, each
                  <address:port> couple must be used only once among all
                  instances running on a same system. Please note that binding
                  to ports lower than 1024 generally require particular
                  privileges to start the program, which are independent of
                  the 'uid' parameter.

    <path>        is a UNIX socket path beginning with a slash ('/'). This is
                  alternative to the TCP listening port. Haproxy will then
                  receive UNIX connections on the socket located at this place.
                  The path must begin with a slash and by default is absolute.
                  It can be relative to the prefix defined by "unix-bind" in
                  the global section. Note that the total length of the prefix
                  followed by the socket path cannot exceed some system limits
                  for UNIX sockets, which commonly are set to 107 characters.

    <param*>      is a list of parameters common to all sockets declared on the
                  same line. These numerous parameters depend on OS and build
                  options and have a complete section dedicated to them. Please
                  refer to section 5 to for more details.

  It is possible to specify a list of address:port combinations delimited by
  commas. The frontend will then listen on all of these addresses. There is no
  fixed limit to the number of addresses and ports which can be listened on in
  a frontend, as well as there is no limit to the number of "bind" statements
  in a frontend.

  Example :
```

```
        listen http_proxy
            bind :80,:443
            bind 10.0.0.1:10080,10.0.0.1:10443
            bind /var/run/ssl-frontend.sock user root mode 600 accept-proxy

        listen http_https_proxy
            bind :80
            bind :443 ssl crt /etc/haproxy/site.pem

        listen http_https_proxy_explicit
            bind ipv6@:80
            bind ipv4@public_ssl:443 ssl crt /etc/haproxy/site.pem
            bind unix@ssl-frontend.sock user root mode 600 accept-proxy

        listen external_bind_app1
            bind "fd@${FD_APP1}"

    Note: regarding Linux's abstract namespace sockets, HAProxy uses the whole
          sun_path length is used for the address length. Some other programs
          such as socat use the string length only by default. Pass the option
          ",unix-tightsocklen=0" to any abstract socket definition in socat to
          make it compatible with HAProxy's.

    See also : "source", "option forwardfor", "unix-bind" and the PROXY protocol
               documentation, and section 5 about bind options.


  bind-process [ all | odd | even | <number 1-64>[-<number 1-64>] ] ...
    Limit visibility of an instance to a certain set of processes numbers.
    May be used in sections :   defaults | frontend | listen | backend
                                   yes   |    yes   |   yes  |   yes
    Arguments :
      all           All process will see this instance. This is the default. It
                    may be used to override a default value.

      odd           This instance will be enabled on processes 1,3,5,...63. This
                    option may be combined with other numbers.

      even          This instance will be enabled on processes 2,4,6,...64. This
                    option may be combined with other numbers. Do not use it
                    with less than 2 processes otherwise some instances might be
                    missing from all processes.

      number        The instance will be enabled on this process number or range,
                    whose values must all be between 1 and 32 or 64 depending on
                    the machine's word size. If a proxy is bound to process
                    numbers greater than the configured global.nbproc, it will
                    either be forced to process #1 if a single process was
                    specified, or to all processes otherwise.

    This keyword limits binding of certain instances to certain processes. This
    is useful in order not to have too many processes listening to the same
    ports. For instance, on a dual-core machine, it might make sense to set
    'nbproc 2' in the global section, then distributes the listeners among 'odd'
    and 'even' instances.

    At the moment, it is not possible to reference more than 32 or 64 processes
    using this keyword, but this should be more than enough for most setups.
    Please note that 'all' really means all processes regardless of the machine's
    word size, and is not limited to the first 32 or 64.

    Each "bind" line may further be limited to a subset of the proxy's processes,
    please consult the "process" bind keyword in section 5.1.

    When a frontend has no explicit "bind-process" line, it tries to bind to all
```



```
    the processes referenced by its "bind" lines. That means that frontends can
    easily adapt to their listeners' processes.

    If some backends are referenced by frontends bound to other processes, the
    backend automatically inherits the frontend's processes.

    Example :
          listen app_ip1
              bind 10.0.0.1:80
              bind-process odd

          listen app_ip2
              bind 10.0.0.2:80
              bind-process even

          listen management
              bind 10.0.0.3:80
              bind-process 1 2 3 4

          listen management
              bind 10.0.0.4:80
              bind-process 1-4

    See also : "nbproc" in global section, and "process" in section 5.1.


  block { if | unless } <condition>
    Block a layer 7 request if/unless a condition is matched
    May be used in sections :   defaults | frontend | listen | backend
                                   no    |    yes   |   yes  |   yes

    The HTTP request will be blocked very early in the layer 7 processing
    if/unless <condition> is matched. A 403 error will be returned if the request
    is blocked. The condition has to reference ACLs (see section 7). This is
    typically used to deny access to certain sensitive resources if some
    conditions are met or not met. There is no fixed limit to the number of
    "block" statements per instance.

    Example:
          acl invalid_src  src          0.0.0.0/7 224.0.0.0/3
          acl invalid_src  src_port     0:1023
          acl local_dst    hdr(host) -i localhost
          block if invalid_src || local_dst

    See section 7 about ACL usage.


  capture cookie <name> len <length>
    Capture and log a cookie in the request and in the response.
    May be used in sections :   defaults | frontend | listen | backend
                                   no    |    yes   |   yes  |   no
    Arguments :
      <name>    is the beginning of the name of the cookie to capture. In order
                to match the exact name, simply suffix the name with an equal
                sign ('='). The full name will appear in the logs, which is
                useful with application servers which adjust both the cookie name
                and value (eg: ASPSESSIONXXXXX).

      <length>  is the maximum number of characters to report in the logs, which
                include the cookie name, the equal sign and the value, all in the
                standard "name=value" form. The string will be truncated on the
                right if it exceeds <length>.

    Only the first cookie is captured. Both the "cookie" request headers and the
    "set-cookie" response headers are monitored. This is particularly useful to
```

```
2341    check for application bugs causing session crossing or stealing between
2342    users, because generally the user's cookies can only change on a login page.
2343
2344    When the cookie was not presented by the client, the associated log column
2345    will report "-". When a request does not cause a cookie to be assigned by the
2346    server, a "-" is reported in the response column.
2347
2348    The capture is performed in the frontend only because it is necessary that
2349    the log format does not change for a given frontend depending on the
2350    backends. This may change in the future. Note that there can be only one
2351    "capture cookie" statement in a frontend. The maximum capture length is set
2352    by the global "tune.http.cookielen" setting and defaults to 63 characters. It
2353    is not possible to specify a capture in a "defaults" section.
2354
2355    Example:
2356          capture cookie ASPSESSION len 32
2357
2358    See also : "capture request header", "capture response header" as well as
2359               section 8 about logging.
2360
2361
2362  capture request header <name> len <length>
2363    Capture and log the last occurrence of the specified request header.
2364    May be used in sections :   defaults | frontend | listen | backend
2365                                   no   |    yes   |   yes  |   no
2366    Arguments :
2367      <name>    is the name of the header to capture. The header names are not
2368                case-sensitive, but it is a common practice to write them as they
2369                appear in the requests, with the first letter of each word in
2370                upper case. The header name will not appear in the logs, only the
2371                value is reported, but the position in the logs is respected.
2372
2373      <length>  is the maximum number of characters to extract from the value and
2374                report in the logs. The string will be truncated on the right if
2375                it exceeds <length>.
2376
2377    The complete value of the last occurrence of the header is captured. The
2378    value will be added to the logs between braces ('{}'). If multiple headers
2379    are captured, they will be delimited by a vertical bar ('|') and will appear
2380    in the same order they were declared in the configuration. Non-existent
2381    headers will be logged just as an empty string. Common uses for request
2382    header captures include the "Host" field in virtual hosting environments, the
2383    "Content-length" when uploads are supported, "User-agent" to quickly
2384    differentiate between real users and robots, and "X-Forwarded-For" in proxied
2385    environments to find where the request came from.
2386
2387    Note that when capturing headers such as "User-agent", some spaces may be
2388    logged, making the log analysis more difficult. Thus be careful about what
2389    you log if you know your log parser is not smart enough to rely on the
2390    braces.
2391
2392    There is no limit to the number of captured request headers nor to their
2393    length, though it is wise to keep them low to limit memory usage per session.
2394    In order to keep log format consistent for a same frontend, header captures
2395    can only be declared in a frontend. It is not possible to specify a capture
2396    in a "defaults" section.
2397
2398    Example:
2399          capture request header Host len 15
2400          capture request header X-Forwarded-For len 15
2401          capture request header Referer len 15
2402
2403    See also : "capture cookie", "capture response header" as well as section 8
2404               about logging.
2405
```



```
2406
2407  capture response header <name> len <length>
2408    Capture and log the last occurrence of the specified response header.
2409    May be used in sections :   defaults | frontend | listen | backend
2410                                   no   |    yes   |   yes  |   no
2411    Arguments :
2412      <name>    is the name of the header to capture. The header names are not
2413                case-sensitive, but it is a common practice to write them as they
2414                appear in the response, with the first letter of each word in
2415                upper case. The header name will not appear in the logs, only the
2416                value is reported, but the position in the logs is respected.
2417
2418      <length>  is the maximum number of characters to extract from the value and
2419                report in the logs. The string will be truncated on the right if
2420                it exceeds <length>.
2421
2422    The complete value of the last occurrence of the header is captured. The
2423    result will be added to the logs between braces ('{}') after the captured
2424    request headers. If multiple headers are captured, they will be delimited by
2425    a vertical bar ('|') and will appear in the same order they were declared in
2426    the configuration. Non-existent headers will be logged just as an empty
2427    string. Common uses for response header captures include the "Content-length"
2428    header which indicates how many bytes are expected to be returned, the
2429    "Location" header to track redirections.
2430
2431    There is no limit to the number of captured response headers nor to their
2432    length, though it is wise to keep them low to limit memory usage per session.
2433    In order to keep log format consistent for a same frontend, header captures
2434    can only be declared in a frontend. It is not possible to specify a capture
2435    in a "defaults" section.
2436
2437    Example:
2438          capture response header Content-length len 9
2439          capture response header Location len 15
2440
2441    See also : "capture cookie", "capture request header" as well as section 8
2442               about logging.
2443
2444
2445  clitimeout <timeout> (deprecated)
2446    Set the maximum inactivity time on the client side.
2447    May be used in sections :   defaults | frontend | listen | backend
2448                                   yes   |    yes   |   yes  |   no
2449    Arguments :
2450      <timeout> is the timeout value is specified in milliseconds by default, but
2451                can be in any other unit if the number is suffixed by the unit,
2452                as explained at the top of this document.
2453
2454    The inactivity timeout applies when the client is expected to acknowledge or
2455    send data. In HTTP mode, this timeout is particularly important to consider
2456    during the first phase, when the client sends the request, and during the
2457    response while it is reading data sent by the server. The value is specified
2458    in milliseconds by default, but can be in any other unit if the number is
2459    suffixed by the unit, as specified at the top of this document. In TCP mode
2460    (and to a lesser extent, in HTTP mode), it is highly recommended that the
2461    client timeout remains equal to the server timeout in order to avoid complex
2462    situations to debug. It is a good practice to cover one or several TCP packet
2463    losses by specifying timeouts that are slightly above multiples of 3 seconds
2464    (eg: 4 or 5 seconds).
2465
2466    This parameter is specific to frontends, but can be specified once for all in
2467    "defaults" sections. This is in fact one of the easiest solutions not to
2468    forget about it. An unspecified timeout results in an infinite timeout, which
2469    is not recommended. Such a usage is accepted and works but reports a warning
2470    during startup because it may results in accumulation of expired sessions in
```

```
   the system if the system's timeouts are not configured either.

   This parameter is provided for compatibility but is currently deprecated.
   Please use "timeout client" instead.

   See also : "timeout client", "timeout http-request", "timeout server", and
              "srvtimeout".

compression algo <algorithm> ...
compression type <mime type> ...
compression offload
  Enable HTTP compression.
  May be used in sections :   defaults | frontend | listen | backend
                                 yes   |    yes   |   yes  |   yes
  Arguments :
    algo     is followed by the list of supported compression algorithms.
    type     is followed by the list of MIME types that will be compressed.
    offload  makes haproxy work as a compression offloader only (see notes).

   The currently supported algorithms are :
     identity     this is mostly for debugging, and it was useful for developing
                  the compression feature. Identity does not apply any change on
                  data.

     gzip         applies gzip compression. This setting is only available when
                  support for zlib or libslz was built in.

     deflate      same as "gzip", but with deflate algorithm and zlib format.
                  Note that this algorithm has ambiguous support on many
                  browsers and no support at all from recent ones. It is
                  strongly recommended not to use it for anything else than
                  experimentation. This setting is only available when support
                  for zlib or libslz was built in.

     raw-deflate  same as "deflate" without the zlib wrapper, and used as an
                  alternative when the browser wants "deflate". All major
                  browsers understand it and despite violating the standards,
                  it is known to work better than "deflate", at least on MSIE
                  and some versions of Safari. Do not use it in conjunction
                  with "deflate", use either one or the other since both react
                  to the same Accept-Encoding token. This setting is only
                  available when support for zlib or libslz was built in.

   Compression will be activated depending on the Accept-Encoding request
   header. With identity, it does not take care of that header.
   If backend servers support HTTP compression, these directives
   will be no-op: haproxy will see the compressed response and will not
   compress again. If backend servers do not support HTTP compression and
   there is Accept-Encoding header in request, haproxy will compress the
   matching response.

   The "offload" setting makes haproxy remove the Accept-Encoding header to
   prevent backend servers from compressing responses. It is strongly
   recommended not to do this because this means that all the compression work
   will be done on the single point where haproxy is located. However in some
   deployment scenarios, haproxy may be installed in front of a buggy gateway
   with broken HTTP compression implementation which can't be turned off.
   In that case haproxy can be used to prevent that gateway from emitting
   invalid payloads. In this case, simply removing the header in the
   configuration does not work because it applies before the header is parsed,
   so that prevents haproxy from compressing. The "offload" setting should
   then be used for such scenarios. Note: for now, the "offload" setting is
   ignored when set in a defaults section.

   Compression is disabled when:
```



```
     * the request does not advertise a supported compression algorithm in the
       "Accept-Encoding" header
     * the response message is not HTTP/1.1
     * HTTP status code is not 200
     * response header "Transfer-Encoding" contains "chunked" (Temporary
       Workaround)
     * response contain neither a "Content-Length" header nor a
       "Transfer-Encoding" whose last value is "chunked"
     * response contains a "Content-Type" header whose first value starts with
       "multipart"
     * the response contains the "no-transform" value in the "Cache-control"
       header
     * User-Agent matches "Mozilla/4" unless it is MSIE 6 with XP SP2, or MSIE 7
       and later
     * The response contains a "Content-Encoding" header, indicating that the
       response is already compressed (see compression offload)

  Note: The compression does not rewrite Etag headers, and does not emit the
        Warning header.

  Examples :
        compression algo gzip
        compression type text/html text/plain

contimeout <timeout> (deprecated)
  Set the maximum time to wait for a connection attempt to a server to succeed.
  May be used in sections :   defaults | frontend | listen | backend
                                 yes   |    no    |   yes  |   yes
  Arguments :
    <timeout> is the timeout value is specified in milliseconds by default, but
              can be in any other unit if the number is suffixed by the unit,
              as explained at the top of this document.

  If the server is located on the same LAN as haproxy, the connection should be
  immediate (less than a few milliseconds). Anyway, it is a good practice to
  cover one or several TCP packet losses by specifying timeouts that are
  slightly above multiples of 3 seconds (eg: 4 or 5 seconds). By default, the
  connect timeout also presets the queue timeout to the same value if this one
  has not been specified. Historically, the contimeout was also used to set the
  tarpit timeout in a listen section, which is not possible in a pure frontend.

  This parameter is specific to backends, but can be specified once for all in
  "defaults" sections. This is in fact one of the easiest solutions not to
  forget about it. An unspecified timeout results in an infinite timeout, which
  is not recommended. Such a usage is accepted and works but reports a warning
  during startup because it may results in accumulation of failed sessions in
  the system if the system's timeouts are not configured either.

  This parameter is provided for backwards compatibility but is currently
  deprecated. Please use "timeout connect", "timeout queue" or "timeout tarpit"
  instead.

  See also : "timeout connect", "timeout queue", "timeout tarpit",
             "timeout server", "contimeout".


cookie <name> [ rewrite | insert | prefix ] [ indirect ] [ nocache ]
              [ postonly ] [ preserve ] [ httponly ] [ secure ]
              [ domain <domain> ]* [ maxidle <idle> ] [ maxlife <life> ]
  Enable cookie-based persistence in a backend.
  May be used in sections :   defaults | frontend | listen | backend
                                 yes   |    no    |   yes  |   yes
  Arguments :
    <name>    is the name of the cookie which will be monitored, modified or
              inserted in order to bring persistence. This cookie is sent to
```

the client via a "Set-Cookie" header in the response, and is
brought back by the client in a "Cookie" header in all requests.
Special care should be taken to choose a name which does not
conflict with any likely application cookie. Also, if the same
backends are subject to be used by the same clients (eg:
HTTP/HTTPS), care should be taken to use different cookie names
between all backends if persistence between them is not desired.

rewrite This keyword indicates that the cookie will be provided by the
server and that haproxy will have to modify its value to set the
server's identifier in it. This mode is handy when the management
of complex combinations of "Set-cookie" and "Cache-control"
headers is left to the application. The application can then
decide whether or not it is appropriate to emit a persistence
cookie. Since all responses should be monitored, this mode only
works in HTTP close mode. Unless the application behaviour is
very complex and/or broken, it is advised not to start with this
mode for new deployments. This keyword is incompatible with
"insert" and "prefix".

insert This keyword indicates that the persistence cookie will have to
be inserted by haproxy in server responses if the client did not

already have a cookie that would have permitted it to access this
server. When used without the "preserve" option, if the server
emits a cookie with the same name, it will be remove before
processing. For this reason, this mode can be used to upgrade
existing configurations running in the "rewrite" mode. The cookie
will only be a session cookie and will not be stored on the
client's disk. By default, unless the "indirect" option is added,
the server will see the cookies emitted by the client. Due to
caching effects, it is generally wise to add the "nocache" or
"postonly" keywords (see below). The "insert" keyword is not
compatible with "rewrite" and "prefix".

prefix This keyword indicates that instead of relying on a dedicated
cookie for the persistence, an existing one will be completed.
This may be needed in some specific environments where the client
does not support more than one single cookie and the application
already needs it. In this case, whenever the server sets a cookie
named <name>, it will be prefixed with the server's identifier
and a delimiter. The prefix will be removed from all client
requests so that the server still finds the cookie it emitted.
Since all requests and responses are subject to being modified,
this mode requires the HTTP close mode. The "prefix" keyword is
not compatible with "rewrite" and "insert". Note: it is highly
recommended not to use "indirect" with "prefix", otherwise server
cookie updates would not be sent to clients.

indirect When this option is specified, no cookie will be emitted to a
client which already has a valid one for the server which has
processed the request. If the server sets such a cookie itself,
it will be removed, unless the "preserve" option is also set. In
"insert" mode, this will additionally remove cookies from the
requests transmitted to the server, making the persistence
mechanism totally transparent from an application point of view.
Note: it is highly recommended not to use "indirect" with
"prefix", otherwise server cookie updates would not be sent to
clients.

nocache This option is recommended in conjunction with the insert mode
when there is a cache between the client and HAProxy, as it
ensures that a cacheable response will be tagged non-cacheable if
a cookie needs to be inserted. This is important because if all
persistence cookies are added on a cacheable home page for

instance, then all customers will then fetch the page from an
outer cache and will all share the same persistence cookie,
leading to one server receiving much more traffic than others.
See also the "insert" and "postonly" options.

postonly This option ensures that cookie insertion will only be performed
on responses to POST requests. It is an alternative to the
"nocache" option, because POST responses are not cacheable, so
this ensures that the persistence cookie will never get cached.
Since most sites do not need any sort of persistence before the
first POST which generally is a login request, this is a very
efficient method to optimize caching without risking to find a
persistence cookie in the cache.
See also the "insert" and "nocache" options.

preserve This option may only be used with "insert" and/or "indirect". It
allows the server to emit the persistence cookie itself. In this
case, if a cookie is found in the response, haproxy will leave it
untouched. This is useful in order to end persistence after a
logout request for instance. For this, the server just has to
emit a cookie with an invalid value (eg: empty) or with a date in
the past. By combining this mechanism with the "disable-on-404"
check option, it is possible to perform a completely graceful
shutdown because users will definitely leave the server after
they logout.

httponly This option tells haproxy to add an "HttpOnly" cookie attribute
when a cookie is inserted. This attribute is used so that a
user agent doesn't share the cookie with non-HTTP components.
Please check RFC6265 for more information on this attribute.

secure This option tells haproxy to add a "Secure" cookie attribute when
a cookie is inserted. This attribute is used so that a user agent
never emits this cookie over non-secure channels, which means
that a cookie learned with this flag will be presented only over
SSL/TLS connections. Please check RFC6265 for more information on
this attribute.

domain This option allows to specify the domain at which a cookie is
inserted. It requires exactly one parameter: a valid domain
name. If the domain begins with a dot, the browser is allowed to
use it for any host ending with that name. It is also possible to
specify several domain names by invoking this option multiple
times. Some browsers might have small limits on the number of
domains, so be careful when doing that. For the record, sending
10 domains to MSIE 6 or Firefox 2 works as expected.

maxidle This option allows inserted cookies to be ignored after some idle
time. It only works with insert-mode cookies. When a cookie is
sent to the client, the date this cookie was emitted is sent too.
Upon further presentations of this cookie, if the date is older
than the delay indicated by the parameter (in seconds), it will
be ignored. Otherwise, it will be refreshed if needed when the
response is sent to the client. This is particularly useful to
prevent users who never close their browsers from remaining for
too long on the same server (eg: after a farm size change). When
this option is set and a cookie has no date, it is always
accepted, but gets refreshed in the response. This maintains the
ability for admins to access their sites. Cookies that have a
date in the future further than 24 hours are ignored. Doing so
lets admins fix timezone issues without risking kicking users off
the site.

maxlife This option allows inserted cookies to be ignored after some life
time, whether they're in use or not. It only works with insert

```
2731                mode cookies. When a cookie is first sent to the client, the date
2732                this cookie was emitted is sent too. Upon further presentations
2733                of this cookie, if the date is older than the delay indicated by
2734                the parameter (in seconds), it will be ignored. If the cookie in
2735                the request has no date, it is accepted and a date will be set.
2736                Cookies that have a date in the future further than 24 hours are
2737                ignored. Doing so lets admins fix timezone issues without risking
2738                kicking users off the site. Contrary to maxidle, this value is
2739                not refreshed, only the first visit date counts. Both maxidle and
2740                maxlife may be used at the time. This is particularly useful to
2741                prevent users who never close their browsers from remaining for
2742                too long on the same server (eg: after a farm size change). This
2743                is stronger than the maxidle method in that it forces a
2744                redispatch after some absolute delay.
2745
2746    There can be only one persistence cookie per HTTP backend, and it can be
2747    declared in a defaults section. The value of the cookie will be the value
2748    indicated after the "cookie" keyword in a "server" statement. If no cookie
2749    is declared for a given server, the cookie is not set.
2750
2751    Examples :
2752          cookie JSESSIONID prefix
2753          cookie SRV insert indirect nocache
2754          cookie SRV insert postonly indirect
2755          cookie SRV insert indirect nocache maxidle 30m maxlife 8h
2756
2757    See also : "balance source", "capture cookie", "server" and "ignore-persist".
2758
2759
2760  declare capture [ request | response ] len <length>
2761    Declares a capture slot.
2762    May be used in sections :   defaults | frontend | listen | backend
2763                                   no    |    yes   |   yes  |   no
2764    Arguments:
2765      <length> is the length allowed for the capture.
2766
2767    This declaration is only available in the frontend or listen section, but the
2768    reserved slot can be used in the backends. The "request" keyword allocates a
2769    capture slot for use in the request, and "response" allocates a capture slot
2770    for use in the response.
2771
2772    See also: "capture-req", "capture-res" (sample converters),
2773              "capture.req.hdr", "capture.res.hdr" (sample fetches),
2774              "http-request capture" and "http-response capture".
2775
2776
2777  default-server [param*]
2778    Change default options for a server in a backend
2779    May be used in sections :   defaults | frontend | listen | backend
2780                                   yes   |    no    |   yes  |   yes
2781    Arguments:
2782      <param*>  is a list of parameters for this server. The "default-server"
2783                keyword accepts an important number of options and has a complete
2784                section dedicated to it. Please refer to section 5 for more
2785                details.
2786
2787    Example :
2788          default-server inter 1000 weight 13
2789
2790    See also: "server" and section 5 about server options
2791
2792
2793  default_backend <backend>
2794    Specify the backend to use when no "use_backend" rule has been matched.
2795    May be used in sections :   defaults | frontend | listen | backend
```



```
2796                                   yes   |    yes   |   yes  |   no
2797    Arguments :
2798      <backend> is the name of the backend to use.
2799
2800    When doing content-switching between frontend and backends using the
2801    "use_backend" keyword, it is often useful to indicate which backend will be
2802    used when no rule has matched. It generally is the dynamic backend which
2803    will catch all undetermined requests.
2804
2805    Example :
2806
2807          use_backend     dynamic  if  url_dyn
2808          use_backend     static   if  url_css url_img extension_img
2809          default_backend dynamic
2810
2811    See also : "use_backend"
2812
2813
2814  description <string>
2815    Describe a listen, frontend or backend.
2816    May be used in sections :   defaults | frontend | listen | backend
2817                                   no    |    yes   |   yes  |   yes
2818    Arguments : string
2819
2820    Allows to add a sentence to describe the related object in the HAProxy HTML
2821    stats page. The description will be printed on the right of the object name
2822    it describes.
2823    No need to backslash spaces in the <string> arguments.
2824
2825
2826  disabled
2827    Disable a proxy, frontend or backend.
2828    May be used in sections :   defaults | frontend | listen | backend
2829                                   yes   |    yes   |   yes  |   yes
2830    Arguments : none
2831
2832    The "disabled" keyword is used to disable an instance, mainly in order to
2833    liberate a listening port or to temporarily disable a service. The instance
2834    will still be created and its configuration will be checked, but it will be
2835    created in the "stopped" state and will appear as such in the statistics. It
2836    will not receive any traffic nor will it send any health-checks or logs. It
2837    is possible to disable many instances at once by adding the "disabled"
2838    keyword in a "defaults" section.
2839
2840    See also : "enabled"
2841
2842
2843  dispatch <address>:<port>
2844    Set a default server address
2845    May be used in sections :   defaults | frontend | listen | backend
2846                                   no    |    no    |   yes  |   yes
2847    Arguments :
2848
2849      <address> is the IPv4 address of the default server. Alternatively, a
2850                resolvable hostname is supported, but this name will be resolved
2851                during start-up.
2852
2853      <ports>   is a mandatory port specification. All connections will be sent
2854                to this port, and it is not permitted to use port offsets as is
2855                possible with normal servers.
2856
2857    The "dispatch" keyword designates a default server for use when no other
2858    server can take the connection. In the past it was used to forward non
2859    persistent connections to an auxiliary load balancer. Due to its simple
2860    syntax, it has also been used for simple TCP relays. It is recommended not to
```

```
   use it for more clarity, and to use the "server" directive instead.

   See also : "server"


enabled
   Enable a proxy, frontend or backend.
   May be used in sections :   defaults | frontend | listen | backend
                                  yes   |    yes   |   yes  |   yes
   Arguments : none

   The "enabled" keyword is used to explicitly enable an instance, when the
   defaults has been set to "disabled". This is very rarely used.

   See also : "disabled"


errorfile <code> <file>
   Return a file contents instead of errors generated by HAProxy
   May be used in sections :   defaults | frontend | listen | backend
                                  yes   |    yes   |   yes  |   yes
   Arguments :
     <code>    is the HTTP status code. Currently, HAProxy is capable of
               generating codes 200, 400, 403, 405, 408, 429, 500, 502, 503, and
               504.

     <file>    designates a file containing the full HTTP response. It is
               recommended to follow the common practice of appending ".http" to
               the filename so that people do not confuse the response with HTML
               error pages, and to use absolute paths, since files are read
               before any chroot is performed.

   It is important to understand that this keyword is not meant to rewrite
   errors returned by the server, but errors detected and returned by HAProxy.
   This is why the list of supported errors is limited to a small set.

   Code 200 is emitted in response to requests matching a "monitor-uri" rule.

   The files are returned verbatim on the TCP socket. This allows any trick such
   as redirections to another URL or site, as well as tricks to clean cookies,
   force enable or disable caching, etc... The package provides default error
   files returning the same contents as default errors.

   The files should not exceed the configured buffer size (BUFSIZE), which
   generally is 8 or 16 kB, otherwise they will be truncated. It is also wise
   not to put any reference to local contents (eg: images) in order to avoid
   loops between the client and HAProxy when all servers are down, causing an
   error to be returned instead of an image. For better HTTP compliance, it is
   recommended that all header lines end with CR-LF and not LF alone.

   The files are read at the same time as the configuration and kept in memory.
   For this reason, the errors continue to be returned even when the process is
   chrooted, and no file change is considered while the process is running. A
   simple method for developing those files consists in associating them to the
   403 status code and interrogating a blocked URL.

   See also : "errorloc", "errorloc302", "errorloc303"

   Example :
         errorfile 400 /etc/haproxy/errorfiles/400badreq.http
         errorfile 408 /dev/null  # workaround Chrome pre-connect bug
         errorfile 403 /etc/haproxy/errorfiles/403forbid.http
         errorfile 503 /etc/haproxy/errorfiles/503sorry.http
```



```
errorloc <code> <url>
errorloc302 <code> <url>
   Return an HTTP redirection to a URL instead of errors generated by HAProxy
   May be used in sections :   defaults | frontend | listen | backend
                                  yes   |    yes   |   yes  |   yes
   Arguments :
     <code>    is the HTTP status code. Currently, HAProxy is capable of
               generating codes 200, 400, 403, 408, 500, 502, 503, and 504.

     <url>     it is the exact contents of the "Location" header. It may contain
               either a relative URI to an error page hosted on the same site,
               or an absolute URI designating an error page on another site.
               Special care should be given to relative URIs to avoid redirect
               loops if the URI itself may generate the same error (eg: 500).

   It is important to understand that this keyword is not meant to rewrite
   errors returned by the server, but errors detected and returned by HAProxy.
   This is why the list of supported errors is limited to a small set.

   Code 200 is emitted in response to requests matching a "monitor-uri" rule.

   Note that both keyword return the HTTP 302 status code, which tells the
   client to fetch the designated URL using the same HTTP method. This can be
   quite problematic in case of non-GET methods such as POST, because the URL
   sent to the client might not be allowed for something other than GET. To
   workaround this problem, please use "errorloc303" which send the HTTP 303
   status code, indicating to the client that the URL must be fetched with a GET
   request.

   See also : "errorfile", "errorloc303"


errorloc303 <code> <url>
   Return an HTTP redirection to a URL instead of errors generated by HAProxy
   May be used in sections :   defaults | frontend | listen | backend
                                  yes   |    yes   |   yes  |   yes
   Arguments :
     <code>    is the HTTP status code. Currently, HAProxy is capable of
               generating codes 400, 403, 408, 500, 502, 503, and 504.

     <url>     it is the exact contents of the "Location" header. It may contain
               either a relative URI to an error page hosted on the same site,
               or an absolute URI designating an error page on another site.
               Special care should be given to relative URIs to avoid redirect
               loops if the URI itself may generate the same error (eg: 500).

   It is important to understand that this keyword is not meant to rewrite
   errors returned by the server, but errors detected and returned by HAProxy.
   This is why the list of supported errors is limited to a small set.

   Code 200 is emitted in response to requests matching a "monitor-uri" rule.

   Note that both keyword return the HTTP 303 status code, which tells the
   client to fetch the designated URL using the same HTTP GET method. This
   solves the usual problems associated with "errorloc" and the 302 code. It is
   possible that some very old browsers designed before HTTP/1.1 do not support
   it, but no such problem has been reported till now.

   See also : "errorfile", "errorloc", "errorloc302"


email-alert from <emailaddr>
   Declare the from email address to be used in both the envelope and header
   of email alerts.  This is the address that email alerts are sent from.
   May be used in sections:    defaults | frontend | listen | backend
```

```
                                  yes   |    yes   |   yes  |   yes

   Arguments :

     <emailaddr> is the from email address to use when sending email alerts

   Also requires "email-alert mailers" and "email-alert to" to be set
   and if so sending email alerts is enabled for the proxy.

   See also : "email-alert level", "email-alert mailers",
              "email-alert myhostname", "email-alert to", section 3.6 about
              mailers.


email-alert level <level>
  Declare the maximum log level of messages for which email alerts will be
  sent. This acts as a filter on the sending of email alerts.
  May be used in sections:    defaults | frontend | listen | backend
                                  yes   |    yes   |   yes  |   yes

  Arguments :

    <level> One of the 8 syslog levels:
              emerg alert crit err warning notice info  debug
            The above syslog levels are ordered from lowest to highest.

  By default level is alert

  Also requires "email-alert from", "email-alert mailers" and
  "email-alert to" to be set and if so sending email alerts is enabled
  for the proxy.

  Alerts are sent when :

  * An un-paused server is marked as down and <level> is alert or lower
  * A paused server is marked as down and <level> is notice or lower
  * A server is marked as up or enters the drain state and <level>
    is notice or lower
  * "option log-health-checks" is enabled, <level> is info or lower,
     and a health check status update occurs

  See also : "email-alert from", "email-alert mailers",
             "email-alert myhostname", "email-alert to",
             section 3.6 about mailers.


email-alert mailers <mailersect>
  Declare the mailers to be used when sending email alerts
  May be used in sections:    defaults | frontend | listen | backend
                                  yes   |    yes   |   yes  |   yes

  Arguments :

    <mailersect> is the name of the mailers section to send email alerts.

  Also requires "email-alert from" and "email-alert to" to be set
  and if so sending email alerts is enabled for the proxy.

  See also : "email-alert from", "email-alert level", "email-alert myhostname",
             "email-alert to", section 3.6 about mailers.


email-alert myhostname <hostname>
  Declare the to hostname address to be used when communicating with
  mailers.
```



```
  May be used in sections:    defaults | frontend | listen | backend
                                  yes   |    yes   |   yes  |   yes

  Arguments :

    <hostname> is the hostname to use when communicating with mailers

  By default the systems hostname is used.

  Also requires "email-alert from", "email-alert mailers" and
  "email-alert to" to be set and if so sending email alerts is enabled
  for the proxy.

  See also : "email-alert from", "email-alert level", "email-alert mailers",
             "email-alert to", section 3.6 about mailers.


email-alert to <emailaddr>
  Declare both the recipent address in the envelope and to address in the
  header of email alerts. This is the address that email alerts are sent to.
  May be used in sections:    defaults | frontend | listen | backend
                                  yes   |    yes   |   yes  |   yes

  Arguments :

    <emailaddr> is the to email address to use when sending email alerts

  Also requires "email-alert mailers" and "email-alert to" to be set
  and if so sending email alerts is enabled for the proxy.

  See also : "email-alert from", "email-alert level", "email-alert mailers",
             "email-alert myhostname", section 3.6 about mailers.


force-persist { if | unless } <condition>
  Declare a condition to force persistence on down servers
  May be used in sections:    defaults | frontend | listen | backend
                                  no   |    yes   |   yes  |   yes

  By default, requests are not dispatched to down servers. It is possible to
  force this using "option persist", but it is unconditional and redispatches
  to a valid server if "option redispatch" is set. That leaves with very little
  possibilities to force some requests to reach a server which is artificially
  marked down for maintenance operations.

  The "force-persist" statement allows one to declare various ACL-based
  conditions which, when met, will cause a request to ignore the down status of
  a server and still try to connect to it. That makes it possible to start a
  server, still replying an error to the health checks, and run a specially
  configured browser to test the service. Among the handy methods, one could
  use a specific source IP address, or a specific cookie. The cookie also has
  the advantage that it can easily be added/removed on the browser from a test
  page. Once the service is validated, it is then possible to open the service
  to the world by returning a valid response to health checks.

  The forced persistence is enabled when an "if" condition is met, or unless an
  "unless" condition is met. The final redispatch is always disabled when this
  is used.

  See also : "option redispatch", "ignore-persist", "persist",
             and section 7 about ACL usage.


fullconn <conns>
  Specify at what backend load the servers will reach their maxconn
```

```
   May be used in sections :   defaults | frontend | listen | backend
                                  yes   |    no    |   yes  |   yes
   Arguments :
     <conns>   is the number of connections on the backend which will make the
               servers use the maximal number of connections.

   When a server has a "maxconn" parameter specified, it means that its number
   of concurrent connections will never go higher. Additionally, if it has a
   "minconn" parameter, it indicates a dynamic limit following the backend's
   load. The server will then always accept at least <minconn> connections,
   never more than <maxconn>, and the limit will be on the ramp between both
   values when the backend has less than <conns> concurrent connections. This
   makes it possible to limit the load on the servers during normal loads, but
   push it further for important loads without overloading the servers during
   exceptional loads.

   Since it's hard to get this value right, haproxy automatically sets it to
   10% of the sum of the maxconns of all frontends that may branch to this
   backend (based on "use_backend" and "default_backend" rules). That way it's
   safe to leave it unset. However, "use_backend" involving dynamic names are
   not counted since there is no way to know if they could match or not.

   Example :
      # The servers will accept between 100 and 1000 concurrent connections each
      # and the maximum of 1000 will be reached when the backend reaches 10000
      # connections.
      backend dynamic
         fullconn   10000
         server     srv1   dyn1:80 minconn 100 maxconn 1000
         server     srv2   dyn2:80 minconn 100 maxconn 1000

   See also : "maxconn", "server"


grace <time>
   Maintain a proxy operational for some time after a soft stop
   May be used in sections :   defaults | frontend | listen | backend
                                  yes   |    yes   |   yes  |   yes
   Arguments :
     <time>    is the time (by default in milliseconds) for which the instance
               will remain operational with the frontend sockets still listening
               when a soft-stop is received via the SIGUSR1 signal.

   This may be used to ensure that the services disappear in a certain order.
   This was designed so that frontends which are dedicated to monitoring by an
   external equipment fail immediately while other ones remain up for the time
   needed by the equipment to detect the failure.

   Note that currently, there is very little benefit in using this parameter,
   and it may in fact complicate the soft-reconfiguration process more than
   simplify it.


hash-type <method> <function> <modifier>
   Specify a method to use for mapping hashes to servers
   May be used in sections :   defaults | frontend | listen | backend
                                  yes   |    no    |   yes  |   yes
   Arguments :
     <method> is the method used to select a server from the hash computed by
              the <function> :

        map-based   the hash table is a static array containing all alive servers.
                    The hashes will be very smooth, will consider weights, but
                    will be static in that weight changes while a server is up
                    will be ignored. This means that there will be no slow start.
```



```
                    Also, since a server is selected by its position in the array,
                    most mappings are changed when the server count changes. This
                    means that when a server goes up or down, or when a server is
                    added to a farm, most connections will be redistributed to
                    different servers. This can be inconvenient with caches for
                    instance.

        consistent  the hash table is a tree filled with many occurrences of each
                    server. The hash key is looked up in the tree and the closest
                    server is chosen. This hash is dynamic, it supports changing
                    weights while the servers are up, so it is compatible with the
                    slow start feature. It has the advantage that when a server
                    goes up or down, only its associations are moved. When a
                    server is added to the farm, only a few part of the mappings
                    are redistributed, making it an ideal method for caches.
                    However, due to its principle, the distribution will never be
                    very smooth and it may sometimes be necessary to adjust a
                    server's weight or its ID to get a more balanced distribution.
                    In order to get the same distribution on multiple load
                    balancers, it is important that all servers have the exact
                    same IDs. Note: consistent hash uses sdbm and avalanche if no
                    hash function is specified.

     <function> is the hash function to be used :

        sdbm   this function was created initially for sdbm (a public-domain
               reimplementation of ndbm) database library. It was found to do
               well in scrambling bits, causing better distribution of the keys
               and fewer splits. It also happens to be a good general hashing
               function with good distribution, unless the total server weight
               is a multiple of 64, in which case applying the avalanche
               modifier may help.

        djb2   this function was first proposed by Dan Bernstein many years ago
               on comp.lang.c. Studies have shown that for certain workload this
               function provides a better distribution than sdbm. It generally
               works well with text-based inputs though it can perform extremely
               poorly with numeric-only input or when the total server weight is
               a multiple of 33, unless the avalanche modifier is also used.

        wt6    this function was designed for haproxy while testing other
               functions in the past. It is not as smooth as the other ones, but
               is much less sensible to the input data set or to the number of
               servers. It can make sense as an alternative to sdbm+avalanche or
               djb2+avalanche for consistent hashing or when hashing on numeric
               data such as a source IP address or a visitor identifier in a URL
               parameter.

        crc32  this is the most common CRC32 implementation as used in Ethernet,
               gzip, PNG, etc. It is slower than the other ones but may provide
               a better distribution or less predictable results especially when
               used on strings.

     <modifier> indicates an optional method applied after hashing the key :

        avalanche   This directive indicates that the result from the hash
                    function above should not be used in its raw form but that
                    a 4-byte full avalanche hash must be applied first. The
                    purpose of this step is to mix the resulting bits from the
                    previous hash in order to avoid any undesired effect when
                    the input contains some limited values or when the number of
                    servers is a multiple of one of the hash's components (64
                    for SDBM, 33 for DJB2). Enabling avalanche tends to make the
                    result less predictable, but it's also not as smooth as when
                    using the original function. Some testing might be needed
```

```
                    with some workloads. This hash is one of the many proposed
                    by Bob Jenkins.

  The default hash type is "map-based" and is recommended for most usages. The
  default function is "sdbm", the selection of a function should be based on
  the range of the values being hashed.

  See also : "balance", "server"


http-check disable-on-404
  Enable a maintenance mode upon HTTP/404 response to health-checks
  May be used in sections :   defaults | frontend | listen | backend
                                 yes   |    no    |   yes  |   yes
  Arguments : none

  When this option is set, a server which returns an HTTP code 404 will be
  excluded from further load-balancing, but will still receive persistent
  connections. This provides a very convenient method for Web administrators
  to perform a graceful shutdown of their servers. It is also important to note
  that a server which is detected as failed while it was in this mode will not
  generate an alert, just a notice. If the server responds 2xx or 3xx again, it
  will immediately be reinserted into the farm. The status on the stats page
  reports "NOLB" for a server in this mode. It is important to note that this
  option only works in conjunction with the "httpchk" option. If this option
  is used with "http-check expect", then it has precedence over it so that 404
  responses will still be considered as soft-stop.

  See also : "option httpchk", "http-check expect"


http-check expect [!] <match> <pattern>
  Make HTTP health checks consider response contents or specific status codes
  May be used in sections :   defaults | frontend | listen | backend
                                 yes   |    no    |   yes  |   yes
  Arguments :
    <match>   is a keyword indicating how to look for a specific pattern in the
              response. The keyword may be one of "status", "rstatus",
              "string", or "rstring". The keyword may be preceded by an
              exclamation mark ("!") to negate the match. Spaces are allowed
              between the exclamation mark and the keyword. See below for more
              details on the supported keywords.

    <pattern> is the pattern to look for. It may be a string or a regular
              expression. If the pattern contains spaces, they must be escaped
              with the usual backslash ('\').

  By default, "option httpchk" considers that response statuses 2xx and 3xx
  are valid, and that others are invalid. When "http-check expect" is used,
  it defines what is considered valid or invalid. Only one "http-check"
  statement is supported in a backend. If a server fails to respond or times
  out, the check obviously fails. The available matches are :

    status <string> : test the exact string match for the HTTP status code.
                      A health check response will be considered valid if the
                      response's status code is exactly this string. If the
                      "status" keyword is prefixed with "!", then the response
                      will be considered invalid if the status code matches.

    rstatus <regex> : test a regular expression for the HTTP status code.
                      A health check response will be considered valid if the
                      response's status code matches the expression. If the
                      "rstatus" keyword is prefixed with "!", then the response
                      will be considered invalid if the status code matches.
                      This is mostly used to check for multiple codes.
```



```

    string <string> : test the exact string match in the HTTP response body.
                      A health check response will be considered valid if the
                      response's body contains this exact string. If the
                      "string" keyword is prefixed with "!", then the response
                      will be considered invalid if the body contains this
                      string. This can be used to look for a mandatory word at
                      the end of a dynamic page, or to detect a failure when a
                      specific error appears on the check page (eg: a stack
                      trace).

    rstring <regex> : test a regular expression on the HTTP response body.
                      A health check response will be considered valid if the
                      response's body matches this expression. If the "rstring"
                      keyword is prefixed with "!", then the response will be
                      considered invalid if the body matches the expression.
                      This can be used to look for a mandatory word at the end
                      of a dynamic page, or to detect a failure when a specific
                      error appears on the check page (eg: a stack trace).

  It is important to note that the responses will be limited to a certain size
  defined by the global "tune.chksize" option, which defaults to 16384 bytes.
  Thus, too large responses may not contain the mandatory pattern when using
  "string" or "rstring". If a large response is absolutely required, it is
  possible to change the default max size by setting the global variable.
  However, it is worth keeping in mind that parsing very large responses can
  waste some CPU cycles, especially when regular expressions are used, and that
  it is always better to focus the checks on smaller resources.

  Also "http-check expect" doesn't support HTTP keep-alive. Keep in mind that it
  will automatically append a "Connection: close" header, meaning that this
  header should not be present in the request provided by "option httpchk".

  Last, if "http-check expect" is combined with "http-check disable-on-404",
  then this last one has precedence when the server responds with 404.

  Examples :
        # only accept status 200 as valid
        http-check expect status 200

        # consider SQL errors as errors
        http-check expect ! string SQL\ Error

        # consider status 5xx only as errors
        http-check expect ! rstatus ^5

        # check that we have a correct hexadecimal tag before /html
        http-check expect rstring <!--tag:[0-9a-f]*</html>

  See also : "option httpchk", "http-check disable-on-404"


http-check send-state
  Enable emission of a state header with HTTP health checks
  May be used in sections :   defaults | frontend | listen | backend
                                 yes   |    no    |   yes  |   yes
  Arguments : none

  When this option is set, haproxy will systematically send a special header
  "X-Haproxy-Server-State" with a list of parameters indicating to each server
  how they are seen by haproxy. This can be used for instance when a server is
  manipulated without access to haproxy and the operator needs to know whether
  haproxy still sees it up or not, or if the server is the last one in a farm.

  The header is composed of fields delimited by semi-colons, the first of which
```

```
   is a word ("UP", "DOWN", "NOLB"), possibly followed by a number of valid
   checks on the total number before transition, just as appears in the stats
   interface. Next headers are in the form "<variable>=<value>", indicating in
   no specific order some values available in the stats interface :
     - a variable "address", containing the address of the backend server.
       This corresponds to the <address> field in the server declaration. For
       unix domain sockets, it will read "unix".

     - a variable "port", containing the port of the backend server. This
       corresponds to the <port> field in the server declaration. For unix
       domain sockets, it will read "unix".

     - a variable "name", containing the name of the backend followed by a slash
       ("/") then the name of the server. This can be used when a server is
       checked in multiple backends.

     - a variable "node" containing the name of the haproxy node, as set in the
       global "node" variable, otherwise the system's hostname if unspecified.

     - a variable "weight" indicating the weight of the server, a slash ("/")
       and the total weight of the farm (just counting usable servers). This
       helps to know if other servers are available to handle the load when this
       one fails.

     - a variable "scur" indicating the current number of concurrent connections
       on the server, followed by a slash ("/") then the total number of
       connections on all servers of the same backend.

     - a variable "qcur" indicating the current number of requests in the
       server's queue.

   Example of a header received by the application server :
     >>>  X-Haproxy-Server-State: UP 2/3; name=bck/srv2; node=lb1; weight=1/2; \
          scur=13/22; qcur=0

   See also : "option httpchk", "http-check disable-on-404"

http-request { allow | deny | tarpit | auth [realm <realm>] | redirect <rule> |
               add-header <name> <fmt> | set-header <name> <fmt> |
               capture <sample> [ len <length> | id <id> ] |
               del-header <name> | set-nice <nice> | set-log-level <level> |
               replace-header <name> <match-regex> <replace-fmt> |
               replace-value <name> <match-regex> <replace-fmt> |
               set-method <fmt> | set-path <fmt> | set-query <fmt> |
               set-uri <fmt> | set-tos <tos> | set-mark <mark> |
               add-acl(<file name>) <key fmt> |
               del-acl(<file name>) <key fmt> |
               del-map(<file name>) <key fmt> |
               set-map(<file name>) <key fmt> <value fmt> |
               set-var(<var name>) <expr> |
               { track-sc0 | track-sc1 | track-sc2 } <key> [table <table>] |
               sc-inc-gpc0(<sc-id>) |
               sc-set-gpt0(<sc-id>) <int> |
               silent-drop |
              }
              [ { if | unless } <condition> ]
   Access control for Layer 7 requests

   May be used in sections:   defaults | frontend | listen | backend
                                  no    |    yes   |   yes  |   yes

   The http-request statement defines a set of rules which apply to layer 7
   processing. The rules are evaluated in their declaration order when they are
   met in a frontend, listen or backend section. Any rule may optionally be
   followed by an ACL-based condition, in which case it will only be evaluated
```



```
   if the condition is true.

   The first keyword is the rule's action. Currently supported actions include :
     - "allow" : this stops the evaluation of the rules and lets the request
       pass the check. No further "http-request" rules are evaluated.

     - "deny" : this stops the evaluation of the rules and immediately rejects
       the request and emits an HTTP 403 error. No further "http-request" rules
       are evaluated.

     - "tarpit" : this stops the evaluation of the rules and immediately blocks
       the request without responding for a delay specified by "timeout tarpit"
       or "timeout connect" if the former is not set. After that delay, if the
       client is still connected, an HTTP error 500 is returned so that the
       client does not suspect it has been tarpitted. Logs will report the flags
       "PT". The goal of the tarpit rule is to slow down robots during an attack
       when they're limited on the number of concurrent requests. It can be very
       efficient against very dumb robots, and will significantly reduce the
       load on firewalls compared to a "deny" rule. But when facing "correctly"
       developed robots, it can make things worse by forcing haproxy and the
       front firewall to support insane number of concurrent connections. See
       also the "silent-drop" action below.

     - "auth" : this stops the evaluation of the rules and immediately responds
       with an HTTP 401 or 407 error code to invite the user to present a valid
       user name and password. No further "http-request" rules are evaluated. An
       optional "realm" parameter is supported, it sets the authentication realm
       that is returned with the response (typically the application's name).

     - "redirect" : this performs an HTTP redirection based on a redirect rule.
       This is exactly the same as the "redirect" statement except that it
       inserts a redirect rule which can be processed in the middle of other
       "http-request" rules and that these rules use the "log-format" strings.
       See the "redirect" keyword for the rule's syntax.

     - "add-header" appends an HTTP header field whose name is specified in
       <name> and whose value is defined by <fmt> which follows the log-format
       rules (see Custom Log Format in section 8.2.4). This is particularly
       useful to pass connection-specific information to the server (eg: the
       client's SSL certificate), or to combine several headers into one. This
       rule is not final, so it is possible to add other similar rules. Note
       that header addition is performed immediately, so one rule might reuse
       the resulting header from a previous rule.

     - "set-header" does the same as "add-header" except that the header name
       is first removed if it existed. This is useful when passing security
       information to the server, where the header must not be manipulated by
       external users. Note that the new value is computed before the removal so
       it is possible to concatenate a value to an existing header.

     - "del-header" removes all HTTP header fields whose name is specified in
       <name>.

     - "replace-header" matches the regular expression in all occurrences of
       header field <name> according to <match-regex>, and replaces them with
       the <replace-fmt> argument. Format characters are allowed in replace-fmt
       and work like in <fmt> arguments in "add-header". The match is only
       case-sensitive. It is important to understand that this action only
       considers whole header lines, regardless of the number of values they
       may contain. This usage is suited to headers naturally containing commas
       in their value, such as If-Modified-Since and so on.

       Example:

         http-request replace-header Cookie foo=([^;]*);(.*) foo=\1;ip=%bi;\2
```

```
3511
3512        applied to:
3513
3514          Cookie: foo=foobar; expires=Tue, 14-Jun-2016 01:40:45 GMT;
3515
3516        outputs:
3517
3518          Cookie: foo=foobar;ip=192.168.1.20; expires=Tue, 14-Jun-2016 01:40:45 GMT;
3519
3520        assuming the backend IP is 192.168.1.20
3521
3522      - "replace-value" works like "replace-header" except that it matches the
3523        regex against every comma-delimited value of the header field <name>
3524        instead of the entire header. This is suited for all headers which are
3525        allowed to carry more than one value. An example could be the Accept
3526        header.
3527
3528        Example:
3529
3530          http-request replace-value X-Forwarded-For ^192\.168\.(.*)$ 172.16.\1
3531
3532        applied to:
3533
3534          X-Forwarded-For: 192.168.10.1, 192.168.13.24, 10.0.0.37
3535
3536        outputs:
3537
3538          X-Forwarded-For: 172.16.10.1, 172.16.13.24, 10.0.0.37
3539
3540      - "set-method" rewrites the request method with the result of the
3541        evaluation of format string <fmt>. There should be very few valid reasons
3542        for having to do so as this is more likely to break something than to fix
3543        it.
3544
3545      - "set-path" rewrites the request path with the result of the evaluation of
3546        format string <fmt>. The query string, if any, is left intact. If a
3547        scheme and authority is found before the path, they are left intact as
3548        well. If the request doesn't have a path ("*"), this one is replaced with
3549        the format. This can be used to prepend a directory component in front of
3550        a path for example. See also "set-query" and "set-uri".
3551
3552        Example :
3553            # prepend the host name before the path
3554            http-request set-path /%[hdr(host)]%[path]
3555
3556      - "set-query" rewrites the request's query string which appears after the
3557        first question mark ("?") with the result of the evaluation of format
3558        string <fmt>. The part prior to the question mark is left intact. If the
3559        request doesn't contain a question mark and the new value is not empty,
3560        then one is added at the end of the URI, followed by the new value. If
3561        a question mark was present, it will never be removed even if the value
3562        is empty. This can be used to add or remove parameters from the query
3563        string. See also "set-query" and "set-uri".
3564
3565        Example :
3566            # replace "%3D" with "=" in the query string
3567            http-request set-query %[query,regsub(%3D,=,g)]
3568
3569      - "set-uri" rewrites the request URI with the result of the evaluation of
3570        format string <fmt>. The scheme, authority, path and query string are all
3571        replaced at once. This can be used to rewrite hosts in front of proxies,
3572        or to perform complex modifications to the URI such as moving parts
3573        between the path and the query string. See also "set-path" and
3574        "set-query".
3575
```



```
3576      - "set-nice" sets the "nice" factor of the current request being processed.
3577        It only has effect against the other requests being processed at the same
3578        time. The default value is 0, unless altered by the "nice" setting on the
3579        "bind" line. The accepted range is -1024..1024. The higher the value, the
3580        nicest the request will be. Lower values will make the request more
3581        important than other ones. This can be useful to improve the speed of
3582        some requests, or lower the priority of non-important requests. Using
3583        this setting without prior experimentation can cause some major slowdown.
3584
3585      - "set-log-level" is used to change the log level of the current request
3586        when a certain condition is met. Valid levels are the 8 syslog levels
3587        (see the "log" keyword) plus the special level "silent" which disables
3588        logging for this request. This rule is not final so the last matching
3589        rule wins. This rule can be useful to disable health checks coming from
3590        another equipment.
3591
3592      - "set-tos" is used to set the TOS or DSCP field value of packets sent to
3593        the client to the value passed in <tos> on platforms which support this.
3594        This value represents the whole 8 bits of the IP TOS field, and can be
3595        expressed both in decimal or hexadecimal format (prefixed by "0x"). Note
3596        that only the 6 higher bits are used in DSCP or TOS, and the two lower
3597        bits are always 0. This can be used to adjust some routing behaviour on
3598        border routers based on some information from the request. See RFC 2474,
3599        2597, 3260 and 4594 for more information.
3600
3601      - "set-mark" is used to set the Netfilter MARK on all packets sent to the
3602        client to the value passed in <mark> on platforms which support it. This
3603        value is an unsigned 32 bit value which can be matched by netfilter and
3604        by the routing table. It can be expressed both in decimal or hexadecimal
3605        format (prefixed by "0x"). This can be useful to force certain packets to
3606        take a different route (for example a cheaper network path for bulk
3607        downloads). This works on Linux kernels 2.6.32 and above and requires
3608        admin privileges.
3609
3610      - "add-acl" is used to add a new entry into an ACL. The ACL must be loaded
3611        from a file (even a dummy empty file). The file name of the ACL to be
3612        updated is passed between parentheses. It takes one argument: <key fmt>,
3613        which follows log-format rules, to collect content of the new entry. It
3614        performs a lookup in the ACL before insertion, to avoid duplicated (or
3615        more) values. This lookup is done by a linear search and can be expensive
3616        with large lists! It is the equivalent of the "add acl" command from the
3617        stats socket, but can be triggered by an HTTP request.
3618
3619      - "del-acl" is used to delete an entry from an ACL. The ACL must be loaded
3620        from a file (even a dummy empty file). The file name of the ACL to be
3621        updated is passed between parentheses. It takes one argument: <key fmt>,
3622        which follows log-format rules, to collect content of the entry to delete.
3623        It is the equivalent of the "del acl" command from the stats socket, but
3624        can be triggered by an HTTP request.
3625
3626      - "del-map" is used to delete an entry from a MAP. The MAP must be loaded
3627        from a file (even a dummy empty file). The file name of the MAP to be
3628        updated is passed between parentheses. It takes one argument: <key fmt>,
3629        which follows log-format rules, to collect content of the entry to delete.
3630        It takes one argument: "file name" It is the equivalent of the "del map"
3631        command from the stats socket, but can be triggered by an HTTP request.
3632
3633      - "set-map" is used to add a new entry into a MAP. The MAP must be loaded
3634        from a file (even a dummy empty file). The file name of the MAP to be
3635        updated is passed between parentheses. It takes 2 arguments: <key fmt>,
3636        which follows log-format rules, used to collect MAP key, and <value fmt>,
3637        which follows log-format rules, used to collect content for the new entry.
3638        It performs a lookup in the MAP before insertion, to avoid duplicated (or
3639        more) values. This lookup is done by a linear search and can be expensive
3640        with large lists! It is the equivalent of the "set map" command from the
```

```
3641        stats socket, but can be triggered by an HTTP request.
3642
3643      - capture <sample> [ len <length> | id <id> ] :
3644        captures sample expression <sample> from the request buffer, and converts
3645        it to a string of at most <len> characters. The resulting string is
3646        stored into the next request "capture" slot, so it will possibly appear
3647        next to some captured HTTP headers. It will then automatically appear in
3648        the logs, and it will be possible to extract it using sample fetch rules
3649        to feed it into headers or anything. The length should be limited given
3650        that this size will be allocated for each capture during the whole
3651        session life. Please check section 7.3 (Fetching samples) and "capture
3652        request header" for more information.
3653
3654        If the keyword "id" is used instead of "len", the action tries to store
3655        the captured string in a previously declared capture slot. This is useful
3656        to run captures in backends. The slot id can be declared by a previous
3657        directive "http-request capture" or with the "declare capture" keyword.
3658        If the slot <id> doesn't exist, then HAProxy fails parsing the
3659        configuration to prevent unexpected behavior at run time.
3660
3661      - { track-sc0 | track-sc1 | track-sc2 } <key> [table <table>] :
3662        enables tracking of sticky counters from current request. These rules
3663        do not stop evaluation and do not change default action. Three sets of
3664        counters may be simultaneously tracked by the same connection. The first
3665        "track-sc0" rule executed enables tracking of the counters of the
3666        specified table as the first set. The first "track-sc1" rule executed
3667        enables tracking of the counters of the specified table as the second
3668        set. The first "track-sc2" rule executed enables tracking of the
3669        counters of the specified table as the third set. It is a recommended
3670        practice to use the first set of counters for the per-frontend counters
3671        and the second set for the per-backend ones. But this is just a
3672        guideline, all may be used everywhere.
3673
3674        These actions take one or two arguments :
3675          <key>   is mandatory, and is a sample expression rule as described
3676                  in section 7.3. It describes what elements of the incoming
3677                  request or connection will be analysed, extracted, combined,
3678                  and used to select which table entry to update the counters.
3679
3680          <table> is an optional table to be used instead of the default one,
3681                  which is the stick-table declared in the current proxy. All
3682                  the counters for the matches and updates for the key will
3683                  then be performed in that table until the session ends.
3684
3685        Once a "track-sc*" rule is executed, the key is looked up in the table
3686        and if it is not found, an entry is allocated for it. Then a pointer to
3687        that entry is kept during all the session's life, and this entry's
3688        counters are updated as often as possible, every time the session's
3689        counters are updated, and also systematically when the session ends.
3690        Counters are only updated for events that happen after the tracking has
3691        been started. As an exception, connection counters and request counters
3692        are systematically updated so that they reflect useful information.
3693
3694        If the entry tracks concurrent connection counters, one connection is
3695        counted for as long as the entry is tracked, and the entry will not
3696        expire during that time. Tracking counters also provides a performance
3697        advantage over just checking the keys, because only one table lookup is
3698        performed for all ACL checks that make use of it.
3699
3700      - sc-set-gpt0(<sc-id>) <int> :
3701        This action sets the GPT0 tag according to the sticky counter designated
3702        by <sc-id> and the value of <int>. The expected result is a boolean. If
3703        an error occurs, this action silently fails and the actions evaluation
3704        continues.
3705
```



```
3706      - sc-inc-gpc0(<sc-id>):
3707        This action increments the GPC0 counter according with the sticky counter
3708        designated by <sc-id>. If an error occurs, this action silently fails and
3709        the actions evaluation continues.
3710
3711      - set-var(<var-name>) <expr> :
3712        Is used to set the contents of a variable. The variable is declared
3713        inline.
3714
3715          <var-name> The name of the variable starts by an indication about its
3716                     scope. The allowed scopes are:
3717                       "sess" : the variable is shared with all the session,
3718                       "txn"  : the variable is shared with all the transaction
3719                                (request and response)
3720                       "req"  : the variable is shared only during the request
3721                                processing
3722                       "res"  : the variable is shared only during the response
3723                                processing.
3724                     This prefix is followed by a name. The separator is a '.'.
3725                     The name may only contain characters 'a-z', 'A-Z', '0-9',
3726                     and '_'.
3727
3728             <expr>  Is a standard HAProxy expression formed by a sample-fetch
3729                     followed by some converters.
3730
3731        Example:
3732
3733           http-request set-var(req.my_var) req.fhdr(user-agent),lower
3734
3735      - set-src <expr> :
3736        Is used to set the source IP address to the value of specified
3737        expression. Useful when a proxy in front of HAProxy rewrites source IP,
3738        but provides the correct IP in a HTTP header; or you want to mask
3739        source IP for privacy.
3740
3741           <expr>  Is a standard HAProxy expression formed by a sample-fetch
3742                   followed by some converters.
3743
3744        Example:
3745
3746           http-request set-src hdr(x-forwarded-for)
3747           http-request set-src src,ipmask(24)
3748
3749        When set-src is successful, the source port is set to 0.
3750
3751      - "silent-drop" : this stops the evaluation of the rules and makes the
3752        client-facing connection suddenly disappear using a system-dependant way
3753        that tries to prevent the client from being notified. The effect it then
3754        that the client still sees an established connection while there's none
3755        on HAProxy. The purpose is to achieve a comparable effect to "tarpit"
3756        except that it doesn't use any local resource at all on the machine
3757        running HAProxy. It can resist much higher loads than "tarpit", and slow
3758        down stronger attackers. It is important to undestand the impact of using
3759        this mechanism. All stateful equipments placed between the client and
3760        HAProxy (firewalls, proxies, load balancers) will also keep the
3761        established connection for a long time and may suffer from this action.
3762        On modern Linux systems running with enough privileges, the TCP_REPAIR
3763        socket option is used to block the emission of a TCP reset. On other
3764        systems, the socket's TTL is reduced to 1 so that the TCP reset doesn't
3765        pass the first router, though it's still delivered to local networks. Do
3766        not use it unless you fully understand how it works.
3767
3768    There is no limit to the number of http-request statements per instance.
3769
3770    It is important to know that http-request rules are processed very early in
```

```
   the HTTP processing, just after "block" rules and before "reqdel" or "reqrep"
   or "reqadd" rules. That way, headers added by "add-header"/"set-header" are
   visible by almost all further ACL rules.

   Using "reqadd"/"reqdel"/"reqrep" to manipulate request headers is discouraged
   in newer versions (>= 1.5). But if you need to use regular expression to
   delete headers, you can still use "reqdel". Also please use
   "http-request deny/allow/tarpit" instead of "reqdeny"/"reqpass"/"reqtarpit".

   Example:
         acl nagios src 192.168.129.3
         acl local_net src 192.168.0.0/16
         acl auth_ok http_auth(L1)

         http-request allow if nagios
         http-request allow if local_net auth_ok
         http-request auth realm Gimme if local_net auth_ok
         http-request deny

   Example:
         acl auth_ok http_auth_group(L1) G1
         http-request auth unless auth_ok

   Example:
         http-request set-header X-Haproxy-Current-Date %T
         http-request set-header X-SSL                  %[ssl_fc]
         http-request set-header X-SSL-Session_ID       %[ssl_fc_session_id,hex]
         http-request set-header X-SSL-Client-Verify    %[ssl_c_verify]
         http-request set-header X-SSL-Client-DN        %{+Q}[ssl_c_s_dn]
         http-request set-header X-SSL-Client-CN        %{+Q}[ssl_c_s_dn(cn)]
         http-request set-header X-SSL-Issuer           %{+Q}[ssl_c_i_dn]
         http-request set-header X-SSL-Client-NotBefore %{+Q}[ssl_c_notbefore]
         http-request set-header X-SSL-Client-NotAfter  %{+Q}[ssl_c_notafter]

   Example:
         acl key req.hdr(X-Add-Acl-Key) -m found
         acl add path /addacl
         acl del path /delacl

         acl myhost hdr(Host) -f myhost.lst

         http-request add-acl(myhost.lst) %[req.hdr(X-Add-Acl-Key)] if key add
         http-request del-acl(myhost.lst) %[req.hdr(X-Add-Acl-Key)] if key del

   Example:
         acl value  req.hdr(X-Value) -m found
         acl setmap path /setmap
         acl delmap path /delmap

         use_backend bk_appli if { hdr(Host),map_str(map.lst) -m found }

         http-request set-map(map.lst) %[src] %[req.hdr(X-Value)] if setmap value
         http-request del-map(map.lst) %[src]                     if delmap

   See also : "stats http-request", section 3.4 about userlists and section 7
              about ACL usage.

http-response { allow | deny | add-header <name> <fmt> | set-nice <nice> |
                capture <sample> id <id> | redirect <rule> |
                set-header <name> <fmt> | del-header <name> |
                replace-header <name> <regex-match> <replace-fmt> |
                replace-value <name> <regex-match> <replace-fmt> |
                set-status <status> |
                set-log-level <level> | set-mark <mark> | set-tos <tos> |
                add-acl(<file name>) <key fmt> |
```



```
                del-acl(<file name>) <key fmt> |
                del-map(<file name>) <key fmt> |
                set-map(<file name>) <key fmt> <value fmt> |
                set-var(<var-name>) <expr> |
                sc-inc-gpc0(<sc-id>) |
                sc-set-gpt0(<sc-id>) <int> |
                silent-drop |
              }
              [ { if | unless } <condition> ]
   Access control for Layer 7 responses

   May be used in sections:   defaults | frontend | listen | backend
                                 no    |    yes   |   yes  |   yes

   The http-response statement defines a set of rules which apply to layer 7
   processing. The rules are evaluated in their declaration order when they are
   met in a frontend, listen or backend section. Any rule may optionally be
   followed by an ACL-based condition, in which case it will only be evaluated
   if the condition is true. Since these rules apply on responses, the backend
   rules are applied first, followed by the frontend's rules.

   The first keyword is the rule's action. Currently supported actions include :
     - "allow" : this stops the evaluation of the rules and lets the response
       pass the check. No further "http-response" rules are evaluated for the
       current section.

     - "deny" : this stops the evaluation of the rules and immediately rejects
       the response and emits an HTTP 502 error. No further "http-response"
       rules are evaluated.

     - "add-header" appends an HTTP header field whose name is specified in
       <name> and whose value is defined by <fmt> which follows the log-format
       rules (see Custom Log Format in section 8.2.4). This may be used to send
       a cookie to a client for example, or to pass some internal information.
       This rule is not final, so it is possible to add other similar rules.
       Note that header addition is performed immediately, so one rule might
       reuse the resulting header from a previous rule.

     - "set-header" does the same as "add-header" except that the header name
       is first removed if it existed. This is useful when passing security
       information to the server, where the header must not be manipulated by
       external users.

     - "del-header" removes all HTTP header fields whose name is specified in
       <name>.

     - "replace-header" matches the regular expression in all occurrences of
       header field <name> according to <match-regex>, and replaces them with
       the <replace-fmt> argument. Format characters are allowed in replace-fmt
       and work like in <fmt> arguments in "add-header". The match is only
       case-sensitive. It is important to understand that this action only
       considers whole header lines, regardless of the number of values they
       may contain. This usage is suited to headers naturally containing commas
       in their value, such as Set-Cookie, Expires and so on.

       Example:

         http-response replace-header Set-Cookie (C=[^;]*);(.*) \1;ip=%bi;\2

       applied to:

         Set-Cookie: C=1; expires=Tue, 14-Jun-2016 01:40:45 GMT

       outputs:
```

```
3901          Set-Cookie: C=1;ip=192.168.1.20; expires=Tue, 14-Jun-2016 01:40:45 GMT
3902
3903        assuming the backend IP is 192.168.1.20.
3904
3905      - "replace-value" works like "replace-header" except that it matches the
3906        regex against every comma-delimited value of the header field <name>
3907        instead of the entire header. This is suited for all headers which are
3908        allowed to carry more than one value. An example could be the Accept
3909        header.
3910
3911        Example:
3912
3913          http-response replace-value Cache-control ^public$ private
3914
3915        applied to:
3916
3917          Cache-Control: max-age=3600, public
3918
3919        outputs:
3920
3921          Cache-Control: max-age=3600, private
3922
3923      - "set-status" replaces the response status code with <status> which must
3924        be an integer between 100 and 999. Note that the reason is automatically
3925        adapted to the new code.
3926
3927        Example:
3928
3929          # return "431 Request Header Fields Too Large"
3930          http-response set-status 431
3931
3932      - "set-nice" sets the "nice" factor of the current request being processed.
3933        It only has effect against the other requests being processed at the same
3934        time. The default value is 0, unless altered by the "nice" setting on the
3935        "bind" line. The accepted range is -1024..1024. The higher the value, the
3936        nicest the request will be. Lower values will make the request more
3937        important than other ones. This can be useful to improve the speed of
3938        some requests, or lower the priority of non-important requests. Using
3939        this setting without prior experimentation can cause some major slowdown.
3940
3941      - "set-log-level" is used to change the log level of the current request
3942        when a certain condition is met. Valid levels are the 8 syslog levels
3943        (see the "log" keyword) plus the special level "silent" which disables
3944        logging for this request. This rule is not final so the last matching
3945        rule wins. This rule can be useful to disable health checks coming from
3946        another equipment.
3947
3948      - "set-tos" is used to set the TOS or DSCP field value of packets sent to
3949        the client to the value passed in <tos> on platforms which support this.
3950        This value represents the whole 8 bits of the IP TOS field, and can be
3951        expressed both in decimal or hexadecimal format (prefixed by "0x"). Note
3952        that only the 6 higher bits are used in DSCP or TOS, and the two lower
3953        bits are always 0. This can be used to adjust some routing behaviour on
3954        border routers based on some information from the request. See RFC 2474,
3955        2597, 3260 and 4594 for more information.
3956
3957      - "set-mark" is used to set the Netfilter MARK on all packets sent to the
3958        client to the value passed in <mark> on platforms which support it. This
3959        value is an unsigned 32 bit value which can be matched by netfilter and
3960        by the routing table. It can be expressed both in decimal or hexadecimal
3961        format (prefixed by "0x"). This can be useful to force certain packets to
3962        take a different route (for example a cheaper network path for bulk
3963        downloads). This works on Linux kernels 2.6.32 and above and requires
3964        admin privileges.
3965
```



```
3966      - "add-acl" is used to add a new entry into an ACL. The ACL must be loaded
3967        from a file (even a dummy empty file). The file name of the ACL to be
3968        updated is passed between parentheses. It takes one argument: <key fmt>,
3969        which follows log-format rules, to collect content of the new entry. It
3970        performs a lookup in the ACL before insertion, to avoid duplicated (or
3971        more) values. This lookup is done by a linear search and can be expensive
3972        with large lists! It is the equivalent of the "add acl" command from the
3973        stats socket, but can be triggered by an HTTP response.
3974
3975      - "del-acl" is used to delete an entry from an ACL. The ACL must be loaded
3976        from a file (even a dummy empty file). The file name of the ACL to be
3977        updated is passed between parentheses. It takes one argument: <key fmt>,
3978        which follows log-format rules, to collect content of the entry to delete.
3979        It is the equivalent of the "del acl" command from the stats socket, but
3980        can be triggered by an HTTP response.
3981
3982      - "del-map" is used to delete an entry from a MAP. The MAP must be loaded
3983        from a file (even a dummy empty file). The file name of the MAP to be
3984        updated is passed between parentheses. It takes one argument: <key fmt>,
3985        which follows log-format rules, to collect content of the entry to delete.
3986        It takes one argument: "file name" It is the equivalent of the "del map"
3987        command from the stats socket, but can be triggered by an HTTP response.
3988
3989      - "set-map" is used to add a new entry into a MAP. The MAP must be loaded
3990        from a file (even a dummy empty file). The file name of the MAP to be
3991        updated is passed between parentheses. It takes 2 arguments: <key fmt>,
3992        which follows log-format rules, used to collect MAP key, and <value fmt>,
3993        which follows log-format rules, used to collect content for the new entry.
3994        It performs a lookup in the MAP before insertion, to avoid duplicated (or
3995        more) values. This lookup is done by a linear search and can be expensive
3996        with large lists! It is the equivalent of the "set map" command from the
3997        stats socket, but can be triggered by an HTTP response.
3998
3999      - capture <sample> id <id> :
4000        captures sample expression <sample> from the response buffer, and converts
4001        it to a string. The resulting string is stored into the next request
4002        "capture" slot, so it will possibly appear next to some captured HTTP
4003        headers. It will then automatically appear in the logs, and it will be
4004        possible to extract it using sample fetch rules to feed it into headers or
4005        anything. Please check section 7.3 (Fetching samples) and "capture
4006        response header" for more information.
4007
4008        The keyword "id" is the id of the capture slot which is used for storing
4009        the string. The capture slot must be defined in an associated frontend.
4010        This is useful to run captures in backends. The slot id can be declared by
4011        a previous directive "http-response capture" or with the "declare capture"
4012        keyword.
4013        If the slot <id> doesn't exist, then HAProxy fails parsing the
4014        configuration to prevent unexpected behavior at run time.
4015
4016      - "redirect" : this performs an HTTP redirection based on a redirect rule.
4017        This supports a format string similarly to "http-request redirect" rules,
4018        with the exception that only the "location" type of redirect is possible
4019        on the response. See the "redirect" keyword for the rule's syntax. When
4020        a redirect rule is applied during a response, connections to the server
4021        are closed so that no data can be forwarded from the server to the client.
4022
4023      - set-var(<var-name>) expr:
4024        Is used to set the contents of a variable. The variable is declared
4025        inline.
4026
4027          <var-name> The name of the variable starts by an indication about its
4028                     scope. The allowed scopes are:
4029                       "sess" : the variable is shared with all the session,
4030                       "txn"  : the variable is shared with all the transaction
```

```
                                (request and response)
                        "req"  : the variable is shared only during the request
                                 processing
                        "res"  : the variable is shared only during the response
                                 processing.
                    This prefix is followed by a name. The separator is a '.'.
                    The name may only contain characters 'a-z', 'A-Z', '0-9',
                    and '_'.

          <expr>    Is a standard HAProxy expression formed by a sample-fetch
                    followed by some converters.

       Example:

          http-response set-var(sess.last_redir) res.hdr(location)

     - sc-set-gpt0(<sc-id>) <int> :
       This action sets the GPT0 tag according to the sticky counter designated
       by <sc-id> and the value of <int>. The expected result is a boolean. If
       an error occurs, this action silently fails and the actions evaluation
       continues.

     - sc-inc-gpc0(<sc-id>):
       This action increments the GPC0 counter according with the sticky counter
       designated by <sc-id>. If an error occurs, this action silently fails and
       the actions evaluation continues.

     - "silent-drop" : this stops the evaluation of the rules and makes the
       client-facing connection suddenly disappear using a system-dependant way
       that tries to prevent the client from being notified. The effect it then
       that the client still sees an established connection while there's none
       on HAProxy. The purpose is to achieve a comparable effect to "tarpit"
       except that it doesn't use any local resource at all on the machine
       running HAProxy. It can resist much higher loads than "tarpit", and slow
       down stronger attackers. It is important to undestand the impact of using
       this mechanism. All stateful equipments placed between the client and
       HAProxy (firewalls, proxies, load balancers) will also keep the
       established connection for a long time and may suffer from this action.
       On modern Linux systems running with enough privileges, the TCP_REPAIR
       socket option is used to block the emission of a TCP reset. On other
       systems, the socket's TTL is reduced to 1 so that the TCP reset doesn't
       pass the first router, though it's still delivered to local networks. Do
       not use it unless you fully understand how it works.

   There is no limit to the number of http-response statements per instance.

   It is important to know that http-response rules are processed very early in
   the HTTP processing, before "rspdel" or "rsprep" or "rspadd" rules. That way,
   headers added by "add-header"/"set-header" are visible by almost all further ACL
   rules.

   Using "rspadd"/"rspdel"/"rsprep" to manipulate request headers is discouraged
   in newer versions (>= 1.5). But if you need to use regular expression to
   delete headers, you can still use "rspdel". Also please use
   "http-response deny" instead of "rspdeny".

   Example:
          acl key_acl res.hdr(X-Acl-Key) -m found

          acl myhost hdr(Host) -f myhost.lst

          http-response add-acl(myhost.lst) %[res.hdr(X-Acl-Key)] if key_acl
          http-response del-acl(myhost.lst) %[res.hdr(X-Acl-Key)] if key_acl

   Example:
```



```
          acl value  res.hdr(X-Value) -m found

          use_backend bk_appli if { hdr(Host),map_str(map.lst) -m found }

          http-response set-map(map.lst) %[src] %[res.hdr(X-Value)] if value
          http-response del-map(map.lst) %[src]                     if ! value

   See also : "http-request", section 3.4 about userlists and section 7 about
              ACL usage.


http-reuse { never | safe | aggressive | always }
  Declare how idle HTTP connections may be shared between requests

  May be used in sections:   defaults | frontend | listen | backend
                                yes    |    no    |   yes  |   yes

  By default, a connection established between haproxy and the backend server
  belongs to the session that initiated it. The downside is that between the
  response and the next request, the connection remains idle and is not used.
  In many cases for performance reasons it is desirable to make it possible to
  reuse these idle connections to serve other requests from different sessions.
  This directive allows to tune this behaviour.

  The argument indicates the desired connection reuse strategy :

    - "never"  : idle connections are never shared between sessions. This is
                 the default choice. It may be enforced to cancel a different
                 strategy inherited from a defaults section or for
                 troubleshooting. For example, if an old bogus application
                 considers that multiple requests over the same connection come
                 from the same client and it is not possible to fix the
                 application, it may be desirable to disable connection sharing
                 in a single backend. An example of such an application could
                 be an old haproxy using cookie insertion in tunnel mode and
                 not checking any request past the first one.

    - "safe"   : this is the recommended strategy. The first request of a
                 session is always sent over its own connection, and only
                 subsequent requests may be dispatched over other existing
                 connections. This ensures that in case the server closes the
                 connection when the request is being sent, the browser can
                 decide to silently retry it. Since it is exactly equivalent to
                 regular keep-alive, there should be no side effects.

    - "aggressive" : this mode may be useful in webservices environments where
                 all servers are not necessarily known and where it would be
                 appreciable to deliver most first requests over existing
                 connections. In this case, first requests are only delivered
                 over existing connections that have been reused at least once,
                 proving that the server correctly supports connection reuse.
                 It should only be used when it's sure that the client can
                 retry a failed request once in a while and where the benefit
                 of aggressive connection reuse significantly outweights the
                 downsides of rare connection failures.

    - "always" : this mode is only recommended when the path to the server is
                 known for never breaking existing connections quickly after
                 releasing them. It allows the first request of a session to be
                 sent to an existing connection. This can provide a significant
                 performance increase over the "safe" strategy when the backend
                 is a cache farm, since such components tend to show a
                 consistent behaviour and will benefit from the connection
                 sharing. It is recommended that the "http-keep-alive" timeout
                 remains low in this mode so that no dead connections remain
```

```
                  usable. In most cases, this will lead to the same performance
                  gains as "aggressive" but with more risks. It should only be
                  used when it improves the situation over "aggressive".

  When http connection sharing is enabled, a great care is taken to respect the
  connection properties and compatiblities. Specifically :
    - connections made with "usesrc" followed by a client-dependant value
      ("client", "clientip", "hdr_ip") are marked private and never shared ;

    - connections sent to a server with a TLS SNI extension are marked private
      and are never shared ;

    - connections receiving a status code 401 or 407 expect some authentication
      to be sent in return. Due to certain bogus authentication schemes (such
      as NTLM) relying on the connection, these connections are marked private
      and are never shared ;

  No connection pool is involved, once a session dies, the last idle connection
  it was attached to is deleted at the same time. This ensures that connections
  may not last after all sessions are closed.

  Note: connection reuse improves the accuracy of the "server maxconn" setting,
  because almost no new connection will be established while idle connections
  remain available. This is particularly true with the "always" strategy.

  See also : "option http-keep-alive", "server maxconn"


http-send-name-header [<header>]
  Add the server name to a request. Use the header string given by <header>

  May be used in sections:   defaults | frontend | listen | backend
                               yes    |    no    |   yes  |   yes

  Arguments :

    <header>  The header string to use to send the server name

  The "http-send-name-header" statement causes the name of the target
  server to be added to the headers of an HTTP request.  The name
  is added with the header string proved.

  See also : "server"

id <value>
  Set a persistent ID to a proxy.
  May be used in sections :   defaults | frontend | listen | backend
                                  no   |    yes   |   yes  |   yes
  Arguments : none

  Set a persistent ID for the proxy. This ID must be unique and positive.
  An unused ID will automatically be assigned if unset. The first assigned
  value will be 1. This ID is currently only returned in statistics.


ignore-persist { if | unless } <condition>
  Declare a condition to ignore persistence
  May be used in sections:    defaults | frontend | listen | backend
                                  no   |    yes   |   yes  |   yes

  By default, when cookie persistence is enabled, every requests containing
  the cookie are unconditionally persistent (assuming the target server is up
  and running).

  The "ignore-persist" statement allows one to declare various ACL-based
```



```
  conditions which, when met, will cause a request to ignore persistence.
  This is sometimes useful to load balance requests for static files, which
  often don't require persistence. This can also be used to fully disable
  persistence for a specific User-Agent (for example, some web crawler bots).

  The persistence is ignored when an "if" condition is met, or unless an
  "unless" condition is met.

  See also : "force-persist", "cookie", and section 7 about ACL usage.

load-server-state-from-file { global | local | none }
  Allow seamless reload of HAProxy
  May be used in sections:    defaults | frontend | listen | backend
                                  yes   |    no   |   yes  |   yes

  This directive points HAProxy to a file where server state from previous
  running process has been saved. That way, when starting up, before handling
  traffic, the new process can apply old states to servers exactly has if no
  reload occured. The purpose of the "load-server-state-from-file" directive is
  to tell haproxy which file to use. For now, only 2 arguments to either prevent
  loading state or load states from a file containing all backends and servers.
  The state file can be generated by running the command "show servers state"
  over the stats socket and redirect output.

  The format of the file is versionned and is very specific. To understand it,
  please read the documentation of the "show servers state" command (chapter
  9.2 of Management Guide).

  Arguments:
    global     load the content of the file pointed by the global directive
               named "server-state-file".

    local      load the content of the file pointed by the directive
               "server-state-file-name" if set. If not set, then the backend
               name is used as a file name.

    none       don't load any stat for this backend

  Notes:
    - server's IP address is not updated unless DNS resolution is enabled on
      the server. It means that if a server IP address has been changed using
      the stat socket, this information won't be re-applied after reloading.

    - server's weight is applied from previous running process unless it has
      has changed between previous and new configuration files.

  Example 1:

  Minimal configuration:

    global
     stats socket /tmp/socket
     server-state-file /tmp/server_state

    defaults
     load-server-state-from-file global

    backend bk
     server s1 127.0.0.1:22 check weight 11
     server s2 127.0.0.1:22 check weight 12

  Then one can run :

    socat /tmp/socket - <<< "show servers state" > /tmp/server_state

```

```
   Content of the file /tmp/server_state would be like this:

     1
     # <field names skipped for the doc example>
     1 bk 1 s1 127.0.0.1 2 0 11 11 4 6 3 4 6 0 0
     1 bk 2 s2 127.0.0.1 2 0 12 12 4 6 3 4 6 0 0

   Example 2:

   Minimal configuration:

     global
      stats socket /tmp/socket
      server-state-base /etc/haproxy/states

     defaults
      load-server-state-from-file local

     backend bk
      server s1 127.0.0.1:22 check weight 11
      server s2 127.0.0.1:22 check weight 12

   Then one can run :

     socat /tmp/socket - <<< "show servers state bk" > /etc/haproxy/states/bk

   Content of the file /etc/haproxy/states/bk would be like this:

     1
     # <field names skipped for the doc example>
     1 bk 1 s1 127.0.0.1 2 0 11 11 4 6 3 4 6 0 0
     1 bk 2 s2 127.0.0.1 2 0 12 12 4 6 3 4 6 0 0

   See also: "server-state-file", "server-state-file-name", and
   "show servers state"


log global
log <address> [len <length>] <facility> [<level> [<minlevel>]]
no log
  Enable per-instance logging of events and traffic.
  May be used in sections :   defaults | frontend | listen | backend
                                  yes   |    yes   |   yes  |   yes

  Prefix :
    no         should be used when the logger list must be flushed. For example,
               if you don't want to inherit from the default logger list. This
               prefix does not allow arguments.

  Arguments :
    global     should be used when the instance's logging parameters are the
               same as the global ones. This is the most common usage. "global"
               replaces <address>, <facility> and <level> with those of the log
               entries found in the "global" section. Only one "log global"
               statement may be used per instance, and this form takes no other
               parameter.

    <address>  indicates where to send the logs. It takes the same format as
               for the "global" section's logs, and can be one of :

               - An IPv4 address optionally followed by a colon (':') and a UDP
                 port. If no port is specified, 514 is used by default (the
                 standard syslog port).

               - An IPv6 address followed by a colon (':') and optionally a UDP
```



```
                 port. If no port is specified, 514 is used by default (the
                 standard syslog port).

               - A filesystem path to a UNIX domain socket, keeping in mind
                 considerations for chroot (be sure the path is accessible
                 inside the chroot) and uid/gid (be sure the path is
                 appropriately writeable).

              You may want to reference some environment variables in the
              address parameter, see section 2.3 about environment variables.

    <length>   is an optional maximum line length. Log lines larger than this
               value will be truncated before being sent. The reason is that
               syslog servers act differently on log line length. All servers
               support the default value of 1024, but some servers simply drop
               larger lines while others do log them. If a server supports long
               lines, it may make sense to set this value here in order to avoid
               truncating long lines. Similarly, if a server drops long lines,
               it is preferable to truncate them before sending them. Accepted
               values are 80 to 65535 inclusive. The default value of 1024 is
               generally fine for all standard usages. Some specific cases of
               long captures or JSON-formated logs may require larger values.

    <facility> must be one of the 24 standard syslog facilities :

                 kern   user   mail   daemon auth   syslog lpr    news
                 uucp   cron   auth2  ftp    ntp    audit  alert  cron2
                 local0 local1 local2 local3 local4 local5 local6 local7

    <level>    is optional and can be specified to filter outgoing messages. By
               default, all messages are sent. If a level is specified, only
               messages with a severity at least as important as this level
               will be sent. An optional minimum level can be specified. If it
               is set, logs emitted with a more severe level than this one will
               be capped to this level. This is used to avoid sending "emerg"
               messages on all terminals on some default syslog configurations.
               Eight levels are known :

                 emerg  alert  crit   err    warning notice info  debug

  It is important to keep in mind that it is the frontend which decides what to
  log from a connection, and that in case of content switching, the log entries
  from the backend will be ignored. Connections are logged at level "info".

  However, backend log declaration define how and where servers status changes
  will be logged. Level "notice" will be used to indicate a server going up,
  "warning" will be used for termination signals and definitive service
  termination, and "alert" will be used for when a server goes down.

  Note : According to RFC3164, messages are truncated to 1024 bytes before
         being emitted.

  Example :
    log global
    log 127.0.0.1:514 local0 notice         # only send important events
    log 127.0.0.1:514 local0 notice notice  # same but limit output level
    log "${LOCAL_SYSLOG}:514" local0 notice   # send to local server


log-format <string>
  Specifies the log format string to use for traffic logs
  May be used in sections:    defaults | frontend | listen | backend
                                 yes   |    yes   |   yes  |    no

  This directive specifies the log format string that will be used for all logs
```

```
   resulting from traffic passing through the frontend using this line. If the
   directive is used in a defaults section, all subsequent frontends will use
   the same log format. Please see section 8.2.4 which covers the log format
   string in depth.

log-format-sd <string>
   Specifies the RFC5424 structured-data log format string
   May be used in sections:    defaults | frontend | listen | backend
                                  yes   |    yes   |   yes  |    no

   This directive specifies the RFC5424 structured-data log format string that
   will be used for all logs resulting from traffic passing through the frontend
   using this line. If the directive is used in a defaults section, all
   subsequent frontends will use the same log format. Please see section 8.2.4
   which covers the log format string in depth.

   See https://tools.ietf.org/html/rfc5424#section-6.3 for more information
   about the RFC5424 structured-data part.

   Note : This log format string will be used only for loggers that have set
          log format to "rfc5424".

   Example :
     log-format-sd [exampleSDID@1234\ bytes=\"%B\"\ status=\"%ST\"]


log-tag <string>
   Specifies the log tag to use for all outgoing logs
   May be used in sections:    defaults | frontend | listen | backend
                                  yes   |    yes   |   yes  |   yes

   Sets the tag field in the syslog header to this string. It defaults to the
   log-tag set in the global section, otherwise the program name as launched
   from the command line, which usually is "haproxy". Sometimes it can be useful
   to differentiate between multiple processes running on the same host, or to
   differentiate customer instances running in the same process. In the backend,
   logs about servers up/down will use this tag. As a hint, it can be convenient
   to set a log-tag related to a hosted customer in a defaults section then put
   all the frontends and backends for that customer, then start another customer
   in a new defaults section. See also the global "log-tag" directive.

max-keep-alive-queue <value>
   Set the maximum server queue size for maintaining keep-alive connections
   May be used in sections:    defaults | frontend | listen | backend
                                  yes   |     no   |   yes  |   yes

   HTTP keep-alive tries to reuse the same server connection whenever possible,
   but sometimes it can be counter-productive, for example if a server has a lot
   of connections while other ones are idle. This is especially true for static
   servers.

   The purpose of this setting is to set a threshold on the number of queued
   connections at which haproxy stops trying to reuse the same server and prefers
   to find another one. The default value, -1, means there is no limit. A value
   of zero means that keep-alive requests will never be queued. For very close
   servers which can be reached with a low latency and which are not sensible to
   breaking keep-alive, a low value is recommended (eg: local static server can
   use a value of 10 or less). For remote servers suffering from a high latency,
   higher values might be needed to cover for the latency and/or the cost of
   picking a different server.

   Note that this has no impact on responses which are maintained to the same
   server consecutively to a 401 response. They will still go to the same server
   even if they have to be queued.

```



```
   See also : "option http-server-close", "option prefer-last-server", server
              "maxconn" and cookie persistence.


maxconn <conns>
   Fix the maximum number of concurrent connections on a frontend
   May be used in sections :   defaults | frontend | listen | backend
                                  yes   |    yes   |   yes  |   no
   Arguments :
     <conns>   is the maximum number of concurrent connections the frontend will
               accept to serve. Excess connections will be queued by the system
               in the socket's listen queue and will be served once a connection
               closes.

   If the system supports it, it can be useful on big sites to raise this limit
   very high so that haproxy manages connection queues, instead of leaving the
   clients with unanswered connection attempts. This value should not exceed the
   global maxconn. Also, keep in mind that a connection contains two buffers
   of 8kB each, as well as some other data resulting in about 17 kB of RAM being
   consumed per established connection. That means that a medium system equipped
   with 1GB of RAM can withstand around 40000-50000 concurrent connections if
   properly tuned.

   Also, when <conns> is set to large values, it is possible that the servers
   are not sized to accept such loads, and for this reason it is generally wise
   to assign them some reasonable connection limits.

   By default, this value is set to 2000.

   See also : "server", global section's "maxconn", "fullconn"


mode { tcp|http|health }
   Set the running mode or protocol of the instance
   May be used in sections :   defaults | frontend | listen | backend
                                  yes   |    yes   |   yes  |   yes
   Arguments :
     tcp       The instance will work in pure TCP mode. A full-duplex connection
               will be established between clients and servers, and no layer 7
               examination will be performed. This is the default mode. It
               should be used for SSL, SSH, SMTP, ...

     http      The instance will work in HTTP mode. The client request will be
               analyzed in depth before connecting to any server. Any request
               which is not RFC-compliant will be rejected. Layer 7 filtering,
               processing and switching will be possible. This is the mode which
               brings HAProxy most of its value.

     health    The instance will work in "health" mode. It will just reply "OK"
               to incoming connections and close the connection. Alternatively,
               If the "httpchk" option is set, "HTTP/1.0 200 OK" will be sent
               instead. Nothing will be logged in either case. This mode is used
               to reply to external components health checks. This mode is
               deprecated and should not be used anymore as it is possible to do
               the same and even better by combining TCP or HTTP modes with the
               "monitor" keyword.

   When doing content switching, it is mandatory that the frontend and the
   backend are in the same mode (generally HTTP), otherwise the configuration
   will be refused.

   Example :
      defaults http_instances
          mode http

```

```
4551    See also : "monitor", "monitor-net"
4552
4553
4554  monitor fail { if | unless } <condition>
4555    Add a condition to report a failure to a monitor HTTP request.
4556    May be used in sections :   defaults | frontend | listen | backend
4557                                   no    |    yes   |   yes  |   no
4558    Arguments :
4559      if <cond>     the monitor request will fail if the condition is satisfied,
4560                    and will succeed otherwise. The condition should describe a
4561                    combined test which must induce a failure if all conditions
4562                    are met, for instance a low number of servers both in a
4563                    backend and its backup.
4564
4565      unless <cond> the monitor request will succeed only if the condition is
4566                    satisfied, and will fail otherwise. Such a condition may be
4567                    based on a test on the presence of a minimum number of active
4568                    servers in a list of backends.
4569
4570    This statement adds a condition which can force the response to a monitor
4571    request to report a failure. By default, when an external component queries
4572    the URI dedicated to monitoring, a 200 response is returned. When one of the
4573    conditions above is met, haproxy will return 503 instead of 200. This is
4574    very useful to report a site failure to an external component which may base
4575    routing advertisements between multiple sites on the availability reported by
4576    haproxy. In this case, one would rely on an ACL involving the "nbsrv"
4577    criterion. Note that "monitor fail" only works in HTTP mode. Both status
4578    messages may be tweaked using "errorfile" or "errorloc" if needed.
4579
4580    Example:
4581       frontend www
4582          mode http
4583          acl site_dead nbsrv(dynamic) lt 2
4584          acl site_dead nbsrv(static)  lt 2
4585          monitor-uri   /site_alive
4586          monitor fail  if site_dead
4587
4588    See also : "monitor-net", "monitor-uri", "errorfile", "errorloc"
4589
4590
4591  monitor-net <source>
4592    Declare a source network which is limited to monitor requests
4593    May be used in sections :   defaults | frontend | listen | backend
4594                                   yes   |    yes   |   yes  |   no
4595    Arguments :
4596      <source>  is the source IPv4 address or network which will only be able to
4597                get monitor responses to any request. It can be either an IPv4
4598                address, a host name, or an address followed by a slash ('/')
4599                followed by a mask.
4600
4601    In TCP mode, any connection coming from a source matching <source> will cause
4602    the connection to be immediately closed without any log. This allows another
4603    equipment to probe the port and verify that it is still listening, without
4604    forwarding the connection to a remote server.
4605
4606    In HTTP mode, a connection coming from a source matching <source> will be
4607    accepted, the following response will be sent without waiting for a request,
4608    then the connection will be closed : "HTTP/1.0 200 OK". This is normally
4609    enough for any front-end HTTP probe to detect that the service is UP and
4610    running without forwarding the request to a backend server. Note that this
4611    response is sent in raw format, without any transformation. This is important
4612    as it means that it will not be SSL-encrypted on SSL listeners.
4613
4614    Monitor requests are processed very early, just after tcp-request connection
4615    ACLs which are the only ones able to block them. These connections are short
```



```
4616    lived and never wait for any data from the client. They cannot be logged, and
4617    it is the intended purpose. They are only used to report HAProxy's health to
4618    an upper component, nothing more. Please note that "monitor fail" rules do
4619    not apply to connections intercepted by "monitor-net".
4620
4621    Last, please note that only one "monitor-net" statement can be specified in
4622    a frontend. If more than one is found, only the last one will be considered.
4623
4624    Example :
4625      # addresses .252 and .253 are just probing us.
4626      frontend www
4627          monitor-net 192.168.0.252/31
4628
4629    See also : "monitor fail", "monitor-uri"
4630
4631
4632  monitor-uri <uri>
4633    Intercept a URI used by external components' monitor requests
4634    May be used in sections :   defaults | frontend | listen | backend
4635                                   yes   |    yes   |   yes  |   no
4636    Arguments :
4637      <uri>     is the exact URI which we want to intercept to return HAProxy's
4638                health status instead of forwarding the request.
4639
4640    When an HTTP request referencing <uri> will be received on a frontend,
4641    HAProxy will not forward it nor log it, but instead will return either
4642    "HTTP/1.0 200 OK" or "HTTP/1.0 503 Service unavailable", depending on failure
4643    conditions defined with "monitor fail". This is normally enough for any
4644    front-end HTTP probe to detect that the service is UP and running without
4645    forwarding the request to a backend server. Note that the HTTP method, the
4646    version and all headers are ignored, but the request must at least be valid
4647    at the HTTP level. This keyword may only be used with an HTTP-mode frontend.
4648
4649    Monitor requests are processed very early. It is not possible to block nor
4650    divert them using ACLs. They cannot be logged either, and it is the intended
4651    purpose. They are only used to report HAProxy's health to an upper component,
4652    nothing more. However, it is possible to add any number of conditions using
4653    "monitor fail" and ACLs so that the result can be adjusted to whatever check
4654    can be imagined (most often the number of available servers in a backend).
4655
4656    Example :
4657      # Use /haproxy_test to report haproxy's status
4658      frontend www
4659          mode http
4660          monitor-uri /haproxy_test
4661
4662    See also : "monitor fail", "monitor-net"
4663
4664
4665  option abortonclose
4666  no option abortonclose
4667    Enable or disable early dropping of aborted requests pending in queues.
4668    May be used in sections :   defaults | frontend | listen | backend
4669                                   yes   |    no    |   yes  |   yes
4670    Arguments : none
4671
4672    In presence of very high loads, the servers will take some time to respond.
4673    The per-instance connection queue will inflate, and the response time will
4674    increase respective to the size of the queue times the average per-session
4675    response time. When clients will wait for more than a few seconds, they will
4676    often hit the "STOP" button on their browser, leaving a useless request in
4677    the queue, and slowing down other users, and the servers as well, because the
4678    request will eventually be served, then aborted at the first error
4679    encountered while delivering the response.
4680
```

```
   As there is no way to distinguish between a full STOP and a simple output
   close on the client side, HTTP agents should be conservative and consider
   that the client might only have closed its output channel while waiting for
   the response. However, this introduces risks of congestion when lots of users
   do the same, and is completely useless nowadays because probably no client at
   all will close the session while waiting for the response. Some HTTP agents
   support this behaviour (Squid, Apache, HAProxy), and others do not (TUX, most
   hardware-based load balancers). So the probability for a closed input channel
   to represent a user hitting the "STOP" button is close to 100%, and the risk
   of being the single component to break rare but valid traffic is extremely
   low, which adds to the temptation to be able to abort a session early while
   still not served and not pollute the servers.

   In HAProxy, the user can choose the desired behaviour using the option
   "abortonclose". By default (without the option) the behaviour is HTTP
   compliant and aborted requests will be served. But when the option is
   specified, a session with an incoming channel closed will be aborted while
   it is still possible, either pending in the queue for a connection slot, or
   during the connection establishment if the server has not yet acknowledged
   the connection request. This considerably reduces the queue size and the load
   on saturated servers when users are tempted to click on STOP, which in turn
   reduces the response time for other users.

   If this option has been enabled in a "defaults" section, it can be disabled
   in a specific instance by prepending the "no" keyword before it.

   See also : "timeout queue" and server's "maxconn" and "maxqueue" parameters


option accept-invalid-http-request
no option accept-invalid-http-request
  Enable or disable relaxing of HTTP request parsing
  May be used in sections :   defaults | frontend | listen | backend
                                 yes   |    yes   |   yes  |   no
  Arguments : none

  By default, HAProxy complies with RFC7230 in terms of message parsing. This
  means that invalid characters in header names are not permitted and cause an
  error to be returned to the client. This is the desired behaviour as such
  forbidden characters are essentially used to build attacks exploiting server
  weaknesses, and bypass security filtering. Sometimes, a buggy browser or
  server will emit invalid header names for whatever reason (configuration,
  implementation) and the issue will not be immediately fixed. In such a case,
  it is possible to relax HAProxy's header name parser to accept any character
  even if that does not make sense, by specifying this option. Similarly, the
  list of characters allowed to appear in a URI is well defined by RFC3986, and
  chars 0-31, 32 (space), 34 ('"'), 60 ('<'), 62 ('>'), 92 ('\'), 94 ('^'), 96
  ('`'), 123 ('{'), 124 ('|'), 125 ('}'), 127 (delete) and anything above are
  not allowed at all. Haproxy always blocks a number of them (0..32, 127). The
  remaining ones are blocked by default unless this option is enabled. This
  option also relaxes the test on the HTTP version, it allows HTTP/0.9 requests
  to pass through (no version specified) and multiple digits for both the major
  and the minor version.

  This option should never be enabled by default as it hides application bugs
  and open security breaches. It should only be deployed after a problem has
  been confirmed.

  When this option is enabled, erroneous header names will still be accepted in
  requests, but the complete request will be captured in order to permit later
  analysis using the "show errors" request on the UNIX stats socket. Similarly,
  requests containing invalid chars in the URI part will be logged. Doing this
  also helps confirming that the issue has been solved.

  If this option has been enabled in a "defaults" section, it can be disabled
```



```
  in a specific instance by prepending the "no" keyword before it.

  See also : "option accept-invalid-http-response" and "show errors" on the
             stats socket.


option accept-invalid-http-response
no option accept-invalid-http-response
  Enable or disable relaxing of HTTP response parsing
  May be used in sections :   defaults | frontend | listen | backend
                                 yes   |    no    |   yes  |   yes
  Arguments : none

  By default, HAProxy complies with RFC7230 in terms of message parsing. This
  means that invalid characters in header names are not permitted and cause an
  error to be returned to the client. This is the desired behaviour as such
  forbidden characters are essentially used to build attacks exploiting server
  weaknesses, and bypass security filtering. Sometimes, a buggy browser or
  server will emit invalid header names for whatever reason (configuration,
  implementation) and the issue will not be immediately fixed. In such a case,
  it is possible to relax HAProxy's header name parser to accept any character
  even if that does not make sense, by specifying this option. This option also
  relaxes the test on the HTTP version format, it allows multiple digits for
  both the major and the minor version.

  This option should never be enabled by default as it hides application bugs
  and open security breaches. It should only be deployed after a problem has
  been confirmed.

  When this option is enabled, erroneous header names will still be accepted in
  responses, but the complete response will be captured in order to permit
  later analysis using the "show errors" request on the UNIX stats socket.
  Doing this also helps confirming that the issue has been solved.

  If this option has been enabled in a "defaults" section, it can be disabled
  in a specific instance by prepending the "no" keyword before it.

  See also : "option accept-invalid-http-request" and "show errors" on the
             stats socket.


option allbackups
no option allbackups
  Use either all backup servers at a time or only the first one
  May be used in sections :   defaults | frontend | listen | backend
                                 yes   |    no    |   yes  |   yes
  Arguments : none

  By default, the first operational backup server gets all traffic when normal
  servers are all down. Sometimes, it may be preferred to use multiple backups
  at once, because one will not be enough. When "option allbackups" is enabled,
  the load balancing will be performed among all backup servers when all normal
  ones are unavailable. The same load balancing algorithm will be used and the
  servers' weights will be respected. Thus, there will not be any priority
  order between the backup servers anymore.

  This option is mostly used with static server farms dedicated to return a
  "sorry" page when an application is completely offline.

  If this option has been enabled in a "defaults" section, it can be disabled
  in a specific instance by prepending the "no" keyword before it.


option checkcache
no option checkcache
```

```
   Analyze all server responses and block responses with cacheable cookies
   May be used in sections :   defaults | frontend | listen | backend
                                  yes   |     no   |   yes  |   yes
   Arguments : none

   Some high-level frameworks set application cookies everywhere and do not
   always let enough control to the developer to manage how the responses should
   be cached. When a session cookie is returned on a cacheable object, there is a
   high risk of session crossing or stealing between users traversing the same
   caches. In some situations, it is better to block the response than to let
   some sensitive session information go in the wild.

   The option "checkcache" enables deep inspection of all server responses for
   strict compliance with HTTP specification in terms of cacheability. It
   carefully checks "Cache-control", "Pragma" and "Set-cookie" headers in server
   response to check if there's a risk of caching a cookie on a client-side
   proxy. When this option is enabled, the only responses which can be delivered
   to the client are :
     - all those without "Set-Cookie" header ;
     - all those with a return code other than 200, 203, 206, 300, 301, 410,
       provided that the server has not set a "Cache-control: public" header ;
     - all those that come from a POST request, provided that the server has not
       set a 'Cache-Control: public' header ;
     - those with a 'Pragma: no-cache' header
     - those with a 'Cache-control: private' header
     - those with a 'Cache-control: no-store' header
     - those with a 'Cache-control: max-age=0' header
     - those with a 'Cache-control: s-maxage=0' header
     - those with a 'Cache-control: no-cache' header
     - those with a 'Cache-control: no-cache="set-cookie"' header
     - those with a 'Cache-control: no-cache="set-cookie,' header
       (allowing other fields after set-cookie)

   If a response doesn't respect these requirements, then it will be blocked
   just as if it was from an "rspdeny" filter, with an "HTTP 502 bad gateway".
   The session state shows "PH--" meaning that the proxy blocked the response
   during headers processing. Additionally, an alert will be sent in the logs so
   that admins are informed that there's something to be fixed.

   Due to the high impact on the application, the application should be tested
   in depth with the option enabled before going to production. It is also a
   good practice to always activate it during tests, even if it is not used in
   production, as it will report potentially dangerous application behaviours.

   If this option has been enabled in a "defaults" section, it can be disabled
   in a specific instance by prepending the "no" keyword before it.


option clitcpka
no option clitcpka
   Enable or disable the sending of TCP keepalive packets on the client side
   May be used in sections :   defaults | frontend | listen | backend
                                  yes   |    yes   |   yes  |   no
   Arguments : none

   When there is a firewall or any session-aware component between a client and
   a server, and when the protocol involves very long sessions with long idle
   periods (eg: remote desktops), there is a risk that one of the intermediate
   components decides to expire a session which has remained idle for too long.

   Enabling socket-level TCP keep-alives makes the system regularly send packets
   to the other end of the connection, leaving it active. The delay between
   keep-alive probes is controlled by the system only and depends both on the
   operating system and its tuning parameters.

```



```
   It is important to understand that keep-alive packets are neither emitted nor
   received at the application level. It is only the network stacks which sees
   them. For this reason, even if one side of the proxy already uses keep-alives
   to maintain its connection alive, those keep-alive packets will not be
   forwarded to the other side of the proxy.

   Please note that this has nothing to do with HTTP keep-alive.

   Using option "clitcpka" enables the emission of TCP keep-alive probes on the
   client side of a connection, which should help when session expirations are
   noticed between HAProxy and a client.

   If this option has been enabled in a "defaults" section, it can be disabled
   in a specific instance by prepending the "no" keyword before it.

   See also : "option srvtcpka", "option tcpka"


option contstats
   Enable continuous traffic statistics updates
   May be used in sections :   defaults | frontend | listen | backend
                                  yes   |    yes   |   yes  |   no
   Arguments : none

   By default, counters used for statistics calculation are incremented
   only when a session finishes. It works quite well when serving small
   objects, but with big ones (for example large images or archives) or
   with A/V streaming, a graph generated from haproxy counters looks like
   a hedgehog. With this option enabled counters get incremented continuously,
   during a whole session. Recounting touches a hotpath directly so
   it is not enabled by default, as it has small performance impact (~0.5%).


option dontlog-normal
no option dontlog-normal
   Enable or disable logging of normal, successful connections
   May be used in sections :   defaults | frontend | listen | backend
                                  yes   |    yes   |   yes  |   no
   Arguments : none

   There are large sites dealing with several thousand connections per second
   and for which logging is a major pain. Some of them are even forced to turn
   logs off and cannot debug production issues. Setting this option ensures that
   normal connections, those which experience no error, no timeout, no retry nor
   redispatch, will not be logged. This leaves disk space for anomalies. In HTTP
   mode, the response status code is checked and return codes 5xx will still be
   logged.

   It is strongly discouraged to use this option as most of the time, the key to
   complex issues is in the normal logs which will not be logged here. If you
   need to separate logs, see the "log-separate-errors" option instead.

   See also : "log", "dontlognull", "log-separate-errors" and section 8 about
              logging.


option dontlognull
no option dontlognull
   Enable or disable logging of null connections
   May be used in sections :   defaults | frontend | listen | backend
                                  yes   |    yes   |   yes  |   no
   Arguments : none

   In certain environments, there are components which will regularly connect to
   various systems to ensure that they are still alive. It can be the case from
```

```
   another load balancer as well as from monitoring systems. By default, even a
   simple port probe or scan will produce a log. If those connections pollute
   the logs too much, it is possible to enable option "dontlognull" to indicate
   that a connection on which no data has been transferred will not be logged,
   which typically corresponds to those probes. Note that errors will still be
   returned to the client and accounted for in the stats. If this is not what is
   desired, option http-ignore-probes can be used instead.

   It is generally recommended not to use this option in uncontrolled
   environments (eg: internet), otherwise scans and other malicious activities
   would not be logged.

   If this option has been enabled in a "defaults" section, it can be disabled
   in a specific instance by prepending the "no" keyword before it.

   See also : "log", "http-ignore-probes", "monitor-net", "monitor-uri", and
              section 8 about logging.


option forceclose
no option forceclose
  Enable or disable active connection closing after response is transferred.
  May be used in sections :   defaults | frontend | listen | backend
                                 yes   |    yes   |   yes  |   yes
  Arguments : none

  Some HTTP servers do not necessarily close the connections when they receive
  the "Connection: close" set by "option httpclose", and if the client does not
  close either, then the connection remains open till the timeout expires. This
  causes high number of simultaneous connections on the servers and shows high
  global session times in the logs.

  When this happens, it is possible to use "option forceclose". It will
  actively close the outgoing server channel as soon as the server has finished
  to respond and release some resources earlier than with "option httpclose".

  This option may also be combined with "option http-pretend-keepalive", which
  will disable sending of the "Connection: close" header, but will still cause
  the connection to be closed once the whole response is received.

  This option disables and replaces any previous "option httpclose", "option
  http-server-close", "option http-keep-alive", or "option http-tunnel".

  If this option has been enabled in a "defaults" section, it can be disabled
  in a specific instance by prepending the "no" keyword before it.

  See also : "option httpclose" and "option http-pretend-keepalive"


option forwardfor [ except <network> ] [ header <name> ] [ if-none ]
  Enable insertion of the X-Forwarded-For header to requests sent to servers
  May be used in sections :   defaults | frontend | listen | backend
                                 yes   |    yes   |   yes  |   yes
  Arguments :
    <network> is an optional argument used to disable this option for sources
              matching <network>
    <name>    an optional argument to specify a different "X-Forwarded-For"
              header name.

  Since HAProxy works in reverse-proxy mode, the servers see its IP address as
  their client address. This is sometimes annoying when the client's IP address
  is expected in server logs. To solve this problem, the well-known HTTP header
  "X-Forwarded-For" may be added by HAProxy to all requests sent to the server.
  This header contains a value representing the client's IP address. Since this
  header is always appended at the end of the existing header list, the server
```



```
  must be configured to always use the last occurrence of this header only. See
  the server's manual to find how to enable use of this standard header. Note
  that only the last occurrence of the header must be used, since it is really
  possible that the client has already brought one.

  The keyword "header" may be used to supply a different header name to replace
  the default "X-Forwarded-For". This can be useful where you might already
  have a "X-Forwarded-For" header from a different application (eg: stunnel),
  and you need preserve it. Also if your backend server doesn't use the
  "X-Forwarded-For" header and requires different one (eg: Zeus Web Servers
  require "X-Cluster-Client-IP").

  Sometimes, a same HAProxy instance may be shared between a direct client
  access and a reverse-proxy access (for instance when an SSL reverse-proxy is
  used to decrypt HTTPS traffic). It is possible to disable the addition of the
  header for a known source address or network by adding the "except" keyword
  followed by the network address. In this case, any source IP matching the
  network will not cause an addition of this header. Most common uses are with
  private networks or 127.0.0.1.

  Alternatively, the keyword "if-none" states that the header will only be
  added if it is not present. This should only be used in perfectly trusted
  environment, as this might cause a security issue if headers reaching haproxy
  are under the control of the end-user.

  This option may be specified either in the frontend or in the backend. If at
  least one of them uses it, the header will be added. Note that the backend's
  setting of the header subargument takes precedence over the frontend's if
  both are defined. In the case of the "if-none" argument, if at least one of
  the frontend or the backend does not specify it, it wants the addition to be
  mandatory, so it wins.

  Examples :
    # Public HTTP address also used by stunnel on the same machine
    frontend www
        mode http
        option forwardfor except 127.0.0.1  # stunnel already adds the header

    # Those servers want the IP Address in X-Client
    backend www
        mode http
        option forwardfor header X-Client

  See also : "option httpclose", "option http-server-close",
             "option forceclose", "option http-keep-alive"


option http-buffer-request
no option http-buffer-request
  Enable or disable waiting for whole HTTP request body before proceeding
  May be used in sections :   defaults | frontend | listen | backend
                                 yes   |    yes   |   yes  |   yes
  Arguments : none

  It is sometimes desirable to wait for the body of an HTTP request before
  taking a decision. This is what is being done by "balance url_param" for
  example. The first use case is to buffer requests from slow clients before
  connecting to the server. Another use case consists in taking the routing
  decision based on the request body's contents. This option placed in a
  frontend or backend forces the HTTP processing to wait until either the whole
  body is received, or the request buffer is full, or the first chunk is
  complete in case of chunked encoding. It can have undesired side effects with
  some applications abusing HTTP by expecting unbufferred transmissions between
  the frontend and the backend, so this should definitely not be used by
  default.
```

```
5071
5072    See also : "option http-no-delay", "timeout http-request"
5073
5074
5075  option http-ignore-probes
5076  no option http-ignore-probes
5077    Enable or disable logging of null connections and request timeouts
5078    May be used in sections :   defaults | frontend | listen | backend
5079                                   yes   |    yes   |   yes  |   no
5080    Arguments : none
5081
5082    Recently some browsers started to implement a "pre-connect" feature
5083    consisting in speculatively connecting to some recently visited web sites
5084    just in case the user would like to visit them. This results in many
5085    connections being established to web sites, which end up in 408 Request
5086    Timeout if the timeout strikes first, or 400 Bad Request when the browser
5087    decides to close them first. These ones pollute the log and feed the error
5088    counters. There was already "option dontlognull" but it's insufficient in
5089    this case. Instead, this option does the following things :
5090       - prevent any 400/408 message from being sent to the client if nothing
5091         was received over a connection before it was closed ;
5092       - prevent any log from being emitted in this situation ;
5093       - prevent any error counter from being incremented
5094
5095    That way the empty connection is silently ignored. Note that it is better
5096    not to use this unless it is clear that it is needed, because it will hide
5097    real problems. The most common reason for not receiving a request and seeing
5098    a 408 is due to an MTU inconsistency between the client and an intermediary
5099    element such as a VPN, which blocks too large packets. These issues are
5100    generally seen with POST requests as well as GET with large cookies. The logs
5101    are often the only way to detect them.
5102
5103    If this option has been enabled in a "defaults" section, it can be disabled
5104    in a specific instance by prepending the "no" keyword before it.
5105
5106    See also : "log", "dontlognull", "errorfile", and section 8 about logging.
5107
5108
5109  option http-keep-alive
5110  no option http-keep-alive
5111    Enable or disable HTTP keep-alive from client to server
5112    May be used in sections :   defaults | frontend | listen | backend
5113                                   yes   |    yes   |   yes  |   yes
5114    Arguments : none
5115
5116    By default HAProxy operates in keep-alive mode with regards to persistent
5117    connections: for each connection it processes each request and response, and
5118    leaves the connection idle on both sides between the end of a response and the
5119    start of a new request. This mode may be changed by several options such as
5120    "option http-server-close", "option forceclose", "option httpclose" or
5121    "option http-tunnel". This option allows to set back the keep-alive mode,
5122    which can be useful when another mode was used in a defaults section.
5123
5124    Setting "option http-keep-alive" enables HTTP keep-alive mode on the client-
5125    and server- sides. This provides the lowest latency on the client side (slow
5126    network) and the fastest session reuse on the server side at the expense
5127    of maintaining idle connections to the servers. In general, it is possible
5128    with this option to achieve approximately twice the request rate that the
5129    "http-server-close" option achieves on small objects. There are mainly two
5130    situations where this option may be useful :
5131
5132      - when the server is non-HTTP compliant and authenticates the connection
5133        instead of requests (eg: NTLM authentication)
5134
5135      - when the cost of establishing the connection to the server is significant
```



```
5136        compared to the cost of retrieving the associated object from the server.
5137
5138    This last case can happen when the server is a fast static server of cache.
5139    In this case, the server will need to be properly tuned to support high enough
5140    connection counts because connections will last until the client sends another
5141    request.
5142
5143    If the client request has to go to another backend or another server due to
5144    content switching or the load balancing algorithm, the idle connection will
5145    immediately be closed and a new one re-opened. Option "prefer-last-server" is
5146    available to try optimize server selection so that if the server currently
5147    attached to an idle connection is usable, it will be used.
5148
5149    In general it is preferred to use "option http-server-close" with application
5150    servers, and some static servers might benefit from "option http-keep-alive".
5151
5152    At the moment, logs will not indicate whether requests came from the same
5153    session or not. The accept date reported in the logs corresponds to the end
5154    of the previous request, and the request time corresponds to the time spent
5155    waiting for a new request. The keep-alive request time is still bound to the
5156    timeout defined by "timeout http-keep-alive" or "timeout http-request" if
5157    not set.
5158
5159    This option disables and replaces any previous "option httpclose", "option
5160    http-server-close", "option forceclose" or "option http-tunnel". When backend
5161    and frontend options differ, all of these 4 options have precedence over
5162    "option http-keep-alive".
5163
5164    See also : "option forceclose", "option http-server-close",
5165               "option prefer-last-server", "option http-pretend-keepalive",
5166               "option httpclose", and "1.1. The HTTP transaction model".
5167
5168
5169  option http-no-delay
5170  no option http-no-delay
5171    Instruct the system to favor low interactive delays over performance in HTTP
5172    May be used in sections :   defaults | frontend | listen | backend
5173                                   yes   |    yes   |   yes  |   yes
5174    Arguments : none
5175
5176    In HTTP, each payload is unidirectional and has no notion of interactivity.
5177    Any agent is expected to queue data somewhat for a reasonably low delay.
5178    There are some very rare server-to-server applications that abuse the HTTP
5179    protocol and expect the payload phase to be highly interactive, with many
5180    interleaved data chunks in both directions within a single request. This is
5181    absolutely not supported by the HTTP specification and will not work across
5182    most proxies or servers. When such applications attempt to do this through
5183    haproxy, it works but they will experience high delays due to the network
5184    optimizations which favor performance by instructing the system to wait for
5185    enough data to be available in order to only send full packets. Typical
5186    delays are around 200 ms per round trip. Note that this only happens with
5187    abnormal uses. Normal uses such as CONNECT requests nor WebSockets are not
5188    affected.
5189
5190    When "option http-no-delay" is present in either the frontend or the backend
5191    used by a connection, all such optimizations will be disabled in order to
5192    make the exchanges as fast as possible. Of course this offers no guarantee on
5193    the functionality, as it may break at any other place. But if it works via
5194    HAProxy, it will work as fast as possible. This option should never be used
5195    by default, and should never be used at all unless such a buggy application
5196    is discovered. The impact of using this option is an increase of bandwidth
5197    usage and CPU usage, which may significantly lower performance in high
5198    latency environments.
5199
5200    See also : "option http-buffer-request"
```

```
5201
5202
5203  option http-pretend-keepalive
5204  no option http-pretend-keepalive
5205    Define whether haproxy will announce keepalive to the server or not
5206    May be used in sections :   defaults | frontend | listen | backend
5207                                   yes   |    yes   |   yes  |   yes
5208    Arguments : none
5209
5210    When running with "option http-server-close" or "option forceclose", haproxy
5211    adds a "Connection: close" header to the request forwarded to the server.
5212    Unfortunately, when some servers see this header, they automatically refrain
5213    from using the chunked encoding for responses of unknown length, while this
5214    is totally unrelated. The immediate effect is that this prevents haproxy from
5215    maintaining the client connection alive. A second effect is that a client or
5216    a cache could receive an incomplete response without being aware of it, and
5217    consider the response complete.
5218
5219    By setting "option http-pretend-keepalive", haproxy will make the server
5220    believe it will keep the connection alive. The server will then not fall back
5221    to the abnormal undesired above. When haproxy gets the whole response, it
5222    will close the connection with the server just as it would do with the
5223    "forceclose" option. That way the client gets a normal response and the
5224    connection is correctly closed on the server side.
5225
5226    It is recommended not to enable this option by default, because most servers
5227    will more efficiently close the connection themselves after the last packet,
5228    and release its buffers slightly earlier. Also, the added packet on the
5229    network could slightly reduce the overall peak performance. However it is
5230    worth noting that when this option is enabled, haproxy will have slightly
5231    less work to do. So if haproxy is the bottleneck on the whole architecture,
5232    enabling this option might save a few CPU cycles.
5233
5234    This option may be set both in a frontend and in a backend. It is enabled if
5235    at least one of the frontend or backend holding a connection has it enabled.
5236    This option may be combined with "option httpclose", which will cause
5237    keepalive to be announced to the server and close to be announced to the
5238    client. This practice is discouraged though.
5239
5240    If this option has been enabled in a "defaults" section, it can be disabled
5241    in a specific instance by prepending the "no" keyword before it.
5242
5243    See also : "option forceclose", "option http-server-close", and
5244               "option http-keep-alive"
5245
5246
5247  option http-server-close
5248  no option http-server-close
5249    Enable or disable HTTP connection closing on the server side
5250    May be used in sections :   defaults | frontend | listen | backend
5251                                   yes   |    yes   |   yes  |   yes
5252    Arguments : none
5253
5254    By default HAProxy operates in keep-alive mode with regards to persistent
5255    connections: for each connection it processes each request and response, and
5256    leaves the connection idle on both sides between the end of a response and
5257    the start of a new request. This mode may be changed by several options such
5258    as "option http-server-close", "option forceclose", "option httpclose" or
5259    "option http-tunnel". Setting "option http-server-close" enables HTTP
5260    connection-close mode on the server side while keeping the ability to support
5261    HTTP keep-alive and pipelining on the client side.  This provides the lowest
5262    latency on the client side (slow network) and the fastest session reuse on
5263    the server side to save server resources, similarly to "option forceclose".
5264    It also permits non-keepalive capable servers to be served in keep-alive mode
5265    to the clients if they conform to the requirements of RFC2616. Please note
```



```
5266    that some servers do not always conform to those requirements when they see
5267    "Connection: close" in the request. The effect will be that keep-alive will
5268    never be used. A workaround consists in enabling "option
5269    http-pretend-keepalive".
5270
5271    At the moment, logs will not indicate whether requests came from the same
5272    session or not. The accept date reported in the logs corresponds to the end
5273    of the previous request, and the request time corresponds to the time spent
5274    waiting for a new request. The keep-alive request time is still bound to the
5275    timeout defined by "timeout http-keep-alive" or "timeout http-request" if
5276    not set.
5277
5278    This option may be set both in a frontend and in a backend. It is enabled if
5279    at least one of the frontend or backend holding a connection has it enabled.
5280    It disables and replaces any previous "option httpclose", "option forceclose",
5281    "option http-tunnel" or "option http-keep-alive". Please check section 4
5282    ("Proxies") to see how this option combines with others when frontend and
5283    backend options differ.
5284
5285    If this option has been enabled in a "defaults" section, it can be disabled
5286    in a specific instance by prepending the "no" keyword before it.
5287
5288    See also : "option forceclose", "option http-pretend-keepalive",
5289               "option httpclose", "option http-keep-alive", and
5290               "1.1. The HTTP transaction model".
5291
5292
5293  option http-tunnel
5294  no option http-tunnel
5295    Disable or enable HTTP connection processing after first transaction
5296    May be used in sections :   defaults | frontend | listen | backend
5297                                   yes   |    yes   |   yes  |   yes
5298    Arguments : none
5299
5300    By default HAProxy operates in keep-alive mode with regards to persistent
5301    connections: for each connection it processes each request and response, and
5302    leaves the connection idle on both sides between the end of a response and
5303    the start of a new request. This mode may be changed by several options such
5304    as "option http-server-close", "option forceclose", "option httpclose" or
5305    "option http-tunnel".
5306
5307    Option "http-tunnel" disables any HTTP processing past the first request and
5308    the first response. This is the mode which was used by default in versions
5309    1.0 to 1.5-dev21. It is the mode with the lowest processing overhead, which
5310    is normally not needed anymore unless in very specific cases such as when
5311    using an in-house protocol that looks like HTTP but is not compatible, or
5312    just to log one request per client in order to reduce log size. Note that
5313    everything which works at the HTTP level, including header parsing/addition,
5314    cookie processing or content switching will only work for the first request
5315    and will be ignored after the first response.
5316
5317    If this option has been enabled in a "defaults" section, it can be disabled
5318    in a specific instance by prepending the "no" keyword before it.
5319
5320    See also : "option forceclose", "option http-server-close",
5321               "option httpclose", "option http-keep-alive", and
5322               "1.1. The HTTP transaction model".
5323
5324
5325  option http-use-proxy-header
5326  no option http-use-proxy-header
5327    Make use of non-standard Proxy-Connection header instead of Connection
5328    May be used in sections :   defaults | frontend | listen | backend
5329                                   yes   |    yes   |   yes  |   no
5330    Arguments : none
```

While RFC2616 explicitly states that HTTP/1.1 agents must use the
Connection header to indicate their wish of persistent or non-persistent
connections, both browsers and proxies ignore this header for proxied
connections and make use of the undocumented, non-standard Proxy-Connection
header instead. The issue begins when trying to put a load balancer between
browsers and such proxies, because there will be a difference between what
haproxy understands and what the client and the proxy agree on.

By setting this option in a frontend, haproxy can automatically switch to use
that non-standard header if it sees proxied requests. A proxied request is
defined here as one where the URI begins with neither a '/' nor a '*'. The
choice of header only affects requests passing through proxies making use of
one of the "httpclose", "forceclose" and "http-server-close" options. Note
that this option can only be specified in a frontend and will affect the
request along its whole life.

Also, when this option is set, a request which requires authentication will
automatically switch to use proxy authentication headers if it is itself a
proxied request. That makes it possible to check or enforce authentication in
front of an existing proxy.

This option should normally never be used, except in front of a proxy.

See also : "option httpclose", "option forceclose" and "option
           http-server-close".

option httpchk
option httpchk <uri>
option httpchk <method> <uri>
option httpchk <method> <uri> <version>
  Enable HTTP protocol to check on the servers health
  May be used in sections :   defaults | frontend | listen | backend
                                 yes | no | yes | yes
  Arguments :
    <method> is the optional HTTP method used with the requests. When not set,
              the "OPTIONS" method is used, as it generally requires low server
              processing and is easy to filter out from the logs. Any method
              may be used, though it is not recommended to invent non-standard
              ones.

    <uri>   is the URI referenced in the HTTP requests. It defaults to " / "
              which is accessible by default on almost any server, but may be
              changed to any other URI. Query strings are permitted.

    <version> is the optional HTTP version string. It defaults to "HTTP/1.0"
              but some servers might behave incorrectly in HTTP 1.0, so turning
              it to HTTP/1.1 may sometimes help. Note that the Host field is
              mandatory in HTTP/1.1, and as a trick, it is possible to pass it
              after "\r\n" following the version string.

  By default, server health checks only consist in trying to establish a TCP
  connection. When "option httpchk" is specified, a complete HTTP request is
  sent once the TCP connection is established, and responses 2xx and 3xx are
  considered valid, while all other ones indicate a server failure, including
  the lack of any response.

  The port and interval are specified in the server configuration.

  This option does not necessarily require an HTTP backend, it also works with
  plain TCP backends. This is particularly useful to check simple scripts bound
  to some dedicated ports using the inetd daemon.

  Examples :

      # Relay HTTPS traffic to Apache instance and check service availability
      # using HTTP request "OPTIONS * HTTP/1.1" on port 80.
      backend https_relay
          mode tcp
          option httpchk OPTIONS * HTTP/1.1\r\nHost:\ www
          server apache1 192.168.1.1:443 check port 80

  See also : "option ssl-hello-chk", "option smtpchk", "option mysql-check",
             "option pgsql-check", "http-check" and the "check", "port" and
             "inter" server options.

option httpclose
no option httpclose
  Enable or disable passive HTTP connection closing
  May be used in sections :   defaults | frontend | listen | backend
                                 yes | yes | yes | yes
  Arguments : none

  By default HAProxy operates in keep-alive mode with regards to persistent
  connections: for each connection it processes each request and response, and
  leaves the connection idle on both sides between the end of a response and
  the start of a new request. This mode may be changed by several options such
  as "option http-server-close", "option forceclose", "option httpclose" or
  "option http-tunnel".

  If "option httpclose" is set, HAProxy will work in HTTP tunnel mode and check
  if a "Connection: close" header is already set in each direction, and will
  add one if missing. Each end should react to this by actively closing the TCP
  connection after each transfer, thus resulting in a switch to the HTTP close
  mode. Any "Connection" header different from "close" will also be removed.
  Note that this option is deprecated since what it does is very cheap but not
  reliable. Using "option http-server-close" or "option forceclose" is strongly
  recommended instead.

  It seldom happens that some servers incorrectly ignore this header and do not
  close the connection even though they reply "Connection: close". For this
  reason, they are not compatible with older HTTP 1.0 browsers. If this happens
  it is possible to use the "option forceclose" which actively closes the
  request connection once the server responds. Option "forceclose" also
  releases the server connection earlier because it does not have to wait for
  the client to acknowledge it.

  This option may be set both in a frontend and in a backend. It is enabled if
  at least one of the frontend or backend holding a connection has it enabled.
  It disables and replaces any previous "option http-server-close",
  "option forceclose", "option http-keep-alive" or "option http-tunnel". Please
  check section 4 ("Proxies") to see how this option combines with others when
  frontend and backend options differ.

  If this option has been enabled in a "defaults" section, it can be disabled
  in a specific instance by prepending the "no" keyword before it.

  See also : "option forceclose", "option http-server-close" and
             "1.1. The HTTP transaction model".

option httplog [clf]
  Enable logging of HTTP request, session state and timers
  May be used in sections :   defaults | frontend | listen | backend
                                 yes | yes | yes | yes
  Arguments :
    clf   if the "clf" argument is added, then the output format will be
              the CLF format instead of HAProxy's default HTTP format. You can
              use this when you need to feed HAProxy's logs through a specific

```
5461                log analyser which only support the CLF format and which is not
5462                extensible.
5463
5464     By default, the log output format is very poor, as it only contains the
5465     source and destination addresses, and the instance name. By specifying
5466     "option httplog", each log line turns into a much richer format including,
5467     but not limited to, the HTTP request, the connection timers, the session
5468     status, the connections numbers, the captured headers and cookies, the
5469     frontend, backend and server name, and of course the source address and
5470     ports.
5471
5472     This option may be set either in the frontend or the backend.
5473
5474     Specifying only "option httplog" will automatically clear the 'clf' mode
5475     if it was set by default.
5476
5477     See also :  section 8 about logging.
5478
5479
5480   option http_proxy
5481   no option http_proxy
5482     Enable or disable plain HTTP proxy mode
5483     May be used in sections :   defaults | frontend | listen | backend
5484                                    yes   |    yes   |   yes  |   yes
5485     Arguments : none
5486
5487     It sometimes happens that people need a pure HTTP proxy which understands
5488     basic proxy requests without caching nor any fancy feature. In this case,
5489     it may be worth setting up an HAProxy instance with the "option http_proxy"
5490     set. In this mode, no server is declared, and the connection is forwarded to
5491     the IP address and port found in the URL after the "http://" scheme.
5492
5493     No host address resolution is performed, so this only works when pure IP
5494     addresses are passed. Since this option's usage perimeter is rather limited,
5495     it will probably be used only by experts who know they need exactly it. Last,
5496     if the clients are susceptible of sending keep-alive requests, it will be
5497     needed to add "option httpclose" to ensure that all requests will correctly
5498     be analyzed.
5499
5500     If this option has been enabled in a "defaults" section, it can be disabled
5501     in a specific instance by prepending the "no" keyword before it.
5502
5503     Example :
5504       # this backend understands HTTP proxy requests and forwards them directly.
5505       backend direct_forward
5506           option httpclose
5507           option http_proxy
5508
5509     See also : "option httpclose"
5510
5511
5512   option independent-streams
5513   no option independent-streams
5514     Enable or disable independent timeout processing for both directions
5515     May be used in sections :   defaults | frontend | listen | backend
5516                                    yes   |    yes   |   yes  |  yes
5517     Arguments : none
5518
5519     By default, when data is sent over a socket, both the write timeout and the
5520     read timeout for that socket are refreshed, because we consider that there is
5521     activity on that socket, and we have no other means of guessing if we should
5522     receive data or not.
5523
5524     While this default behaviour is desirable for almost all applications, there
5525     exists a situation where it is desirable to disable it, and only refresh the
```



```
5526     read timeout if there are incoming data. This happens on sessions with large
5527     timeouts and low amounts of exchanged data such as telnet session. If the
5528     server suddenly disappears, the output data accumulates in the system's
5529     socket buffers, both timeouts are correctly refreshed, and there is no way
5530     to know the server does not receive them, so we don't timeout. However, when
5531     the underlying protocol always echoes sent data, it would be enough by itself
5532     to detect the issue using the read timeout. Note that this problem does not
5533     happen with more verbose protocols because data won't accumulate long in the
5534     socket buffers.
5535
5536     When this option is set on the frontend, it will disable read timeout updates
5537     on data sent to the client. There probably is little use of this case. When
5538     the option is set on the backend, it will disable read timeout updates on
5539     data sent to the server. Doing so will typically break large HTTP posts from
5540     slow lines, so use it with caution.
5541
5542     Note: older versions used to call this setting "option independent-streams"
5543           with a spelling mistake. This spelling is still supported but
5544           deprecated.
5545
5546     See also : "timeout client", "timeout server" and "timeout tunnel"
5547
5548
5549   option ldap-check
5550     Use LDAPv3 health checks for server testing
5551     May be used in sections :   defaults | frontend | listen | backend
5552                                    yes   |    no    |   yes  |   yes
5553     Arguments : none
5554
5555     It is possible to test that the server correctly talks LDAPv3 instead of just
5556     testing that it accepts the TCP connection. When this option is set, an
5557     LDAPv3 anonymous simple bind message is sent to the server, and the response
5558     is analyzed to find an LDAPv3 bind response message.
5559
5560     The server is considered valid only when the LDAP response contains success
5561     resultCode (http://tools.ietf.org/html/rfc4511#section-4.1.9).
5562
5563     Logging of bind requests is server dependent see your documentation how to
5564     configure it.
5565
5566     Example :
5567           option ldap-check
5568
5569     See also : "option httpchk"
5570
5571
5572   option external-check
5573     Use external processes for server health checks
5574     May be used in sections :   defaults | frontend | listen | backend
5575                                    yes   |    no    |   yes  |   yes
5576
5577     It is possible to test the health of a server using an external command.
5578     This is achieved by running the executable set using "external-check
5579     command".
5580
5581     Requires the "external-check" global to be set.
5582
5583     See also : "external-check", "external-check command", "external-check path"
5584
5585
5586   option log-health-checks
5587   no option log-health-checks
5588     Enable or disable logging of health checks status updates
5589     May be used in sections :   defaults | frontend | listen | backend
5590                                    yes   |    no    |   yes  |  yes
```

```
   Arguments : none

   By default, failed health check are logged if server is UP and successful
   health checks are logged if server is DOWN, so the amount of additional
   information is limited.

   When this option is enabled, any change of the health check status or to
   the server's health will be logged, so that it becomes possible to know
   that a server was failing occasional checks before crashing, or exactly when
   it failed to respond a valid HTTP status, then when the port started to
   reject connections, then when the server stopped responding at all.

   Note that status changes not caused by health checks (eg: enable/disable on
   the CLI) are intentionally not logged by this option.

   See also: "option httpchk", "option ldap-check", "option mysql-check",
             "option pgsql-check", "option redis-check", "option smtpchk",
             "option tcp-check", "log" and section 8 about logging.


option log-separate-errors
no option log-separate-errors
  Change log level for non-completely successful connections
  May be used in sections :   defaults | frontend | listen | backend
                                 yes   |    yes   |   yes  |   no
  Arguments : none

  Sometimes looking for errors in logs is not easy. This option makes haproxy
  raise the level of logs containing potentially interesting information such
  as errors, timeouts, retries, redispatches, or HTTP status codes 5xx. The
  level changes from "info" to "err". This makes it possible to log them
  separately to a different file with most syslog daemons. Be careful not to
  remove them from the original file, otherwise you would lose ordering which
  provides very important information.

  Using this option, large sites dealing with several thousand connections per
  second may log normal traffic to a rotating buffer and only archive smaller
  error logs.

  See also : "log", "dontlognull", "dontlog-normal" and section 8 about
             logging.


option logasap
no option logasap
  Enable or disable early logging of HTTP requests
  May be used in sections :   defaults | frontend | listen | backend
                                 yes   |    yes   |   yes  |   no
  Arguments : none

  By default, HTTP requests are logged upon termination so that the total
  transfer time and the number of bytes appear in the logs. When large objects
  are being transferred, it may take a while before the request appears in the
  logs. Using "option logasap", the request gets logged as soon as the server
  sends the complete headers. The only missing information in the logs will be
  the total number of bytes which will indicate everything except the amount
  of data transferred, and the total time which will not take the transfer
  time into account. In such a situation, it's a good practice to capture the
  "Content-Length" response header so that the logs at least indicate how many
  bytes are expected to be transferred.

  Examples :
      listen http_proxy 0.0.0.0:80
          mode http
          option httplog
```



```
          option logasap
          log 192.168.2.200 local3

    >>> Feb  6 12:14:14 localhost \
          haproxy[14389]: 10.0.1.2:33317 [06/Feb/2009:12:14:14.655] http-in \
          static/srv1 9/10/7/14/+30 200 +243 - - ---- 3/1/1/1/0 1/0 \
          "GET /image.iso HTTP/1.0"

  See also : "option httplog", "capture response header", and section 8 about
             logging.


option mysql-check [ user <username> [ post-41 ] ]
  Use MySQL health checks for server testing
  May be used in sections :   defaults | frontend | listen | backend
                                 yes   |    no    |   yes  |   yes
  Arguments :
    <username> This is the username which will be used when connecting to MySQL
               server.
    post-41    Send post v4.1 client compatible checks

  If you specify a username, the check consists of sending two MySQL packet,
  one Client Authentication packet, and one QUIT packet, to correctly close
  MySQL session. We then parse the MySQL Handshake Initialisation packet and/or
  Error packet. It is a basic but useful test which does not produce error nor
  aborted connect on the server. However, it requires adding an authorization
  in the MySQL table, like this :

      USE mysql;
      INSERT INTO user (Host,User) values ('<ip_of_haproxy>','<username>');
      FLUSH PRIVILEGES;

  If you don't specify a username (it is deprecated and not recommended), the
  check only consists in parsing the Mysql Handshake Initialisation packet or
  Error packet, we don't send anything in this mode. It was reported that it
  can generate lockout if check is too frequent and/or if there is not enough
  traffic. In fact, you need in this case to check MySQL "max_connect_errors"
  value as if a connection is established successfully within fewer than MySQL
  "max_connect_errors" attempts after a previous connection was interrupted,
  the error count for the host is cleared to zero. If HAProxy's server get
  blocked, the "FLUSH HOSTS" statement is the only way to unblock it.

  Remember that this does not check database presence nor database consistency.
  To do this, you can use an external check with xinetd for example.

  The check requires MySQL >=3.22, for older version, please use TCP check.

  Most often, an incoming MySQL server needs to see the client's IP address for
  various purposes, including IP privilege matching and connection logging.
  When possible, it is often wise to masquerade the client's IP address when
  connecting to the server using the "usesrc" argument of the "source" keyword,
  which requires the transparent proxy feature to be compiled in, and the MySQL
  server to route the client via the machine hosting haproxy.

  See also: "option httpchk"


option nolinger
no option nolinger
  Enable or disable immediate session resource cleaning after close
  May be used in sections:    defaults | frontend | listen | backend
                                 yes   |    yes   |   yes  |   yes
  Arguments : none

  When clients or servers abort connections in a dirty way (eg: they are
```

```
   physically disconnected), the session timeouts triggers and the session is
   closed. But it will remain in FIN_WAIT1 state for some time in the system,
   using some resources and possibly limiting the ability to establish newer
   connections.

   When this happens, it is possible to activate "option nolinger" which forces
   the system to immediately remove any socket's pending data on close. Thus,
   the session is instantly purged from the system's tables. This usually has
   side effects such as increased number of TCP resets due to old retransmits
   getting immediately rejected. Some firewalls may sometimes complain about
   this too.

   For this reason, it is not recommended to use this option when not absolutely
   needed. You know that you need it when you have thousands of FIN_WAIT1
   sessions on your system (TIME_WAIT ones do not count).

   This option may be used both on frontends and backends, depending on the side
   where it is required. Use it on the frontend for clients, and on the backend
   for servers.

   If this option has been enabled in a "defaults" section, it can be disabled
   in a specific instance by prepending the "no" keyword before it.


option originalto [ except <network> ] [ header <name> ]
  Enable insertion of the X-Original-To header to requests sent to servers
  May be used in sections :   defaults | frontend | listen | backend
                                 yes   |    yes   |   yes  |   yes
  Arguments :
    <network> is an optional argument used to disable this option for sources
              matching <network>
    <name>    an optional argument to specify a different "X-Original-To"
              header name.

   Since HAProxy can work in transparent mode, every request from a client can
   be redirected to the proxy and HAProxy itself can proxy every request to a
   complex SQUID environment and the destination host from SO_ORIGINAL_DST will
   be lost. This is annoying when you want access rules based on destination ip
   addresses. To solve this problem, a new HTTP header "X-Original-To" may be
   added by HAProxy to all requests sent to the server. This header contains a
   value representing the original destination IP address. Since this must be
   configured to always use the last occurrence of this header only. Note that
   only the last occurrence of the header must be used, since it is really
   possible that the client has already brought one.

   The keyword "header" may be used to supply a different header name to replace
   the default "X-Original-To". This can be useful where you might already
   have a "X-Original-To" header from a different application, and you need
   preserve it. Also if your backend server doesn't use the "X-Original-To"
   header and requires different one.

   Sometimes, a same HAProxy instance may be shared between a direct client
   access and a reverse-proxy access (for instance when an SSL reverse-proxy is
   used to decrypt HTTPS traffic). It is possible to disable the addition of the
   header for a known source address or network by adding the "except" keyword
   followed by the network address. In this case, any source IP matching the
   network will not cause an addition of this header. Most common uses are with
   private networks or 127.0.0.1.

   This option may be specified either in the frontend or in the backend. If at
   least one of them uses it, the header will be added. Note that the backend's
   setting of the header subargument takes precedence over the frontend's if
   both are defined.

   Examples :
```



```
     # Original Destination address
     frontend www
         mode http
         option originalto except 127.0.0.1

     # Those servers want the IP Address in X-Client-Dst
     backend www
         mode http
         option originalto header X-Client-Dst

  See also : "option httpclose", "option http-server-close",
             "option forceclose"


option persist
no option persist
  Enable or disable forced persistence on down servers
  May be used in sections:    defaults | frontend | listen | backend
                                 yes   |    no    |   yes  |   yes
  Arguments : none

  When an HTTP request reaches a backend with a cookie which references a dead
  server, by default it is redispatched to another server. It is possible to
  force the request to be sent to the dead server first using "option persist"
  if absolutely needed. A common use case is when servers are under extreme
  load and spend their time flapping. In this case, the users would still be
  directed to the server they opened the session on, in the hope they would be
  correctly served. It is recommended to use "option redispatch" in conjunction
  with this option so that in the event it would not be possible to connect to
  the server at all (server definitely dead), the client would finally be
  redirected to another valid server.

  If this option has been enabled in a "defaults" section, it can be disabled
  in a specific instance by prepending the "no" keyword before it.

  See also : "option redispatch", "retries", "force-persist"


option pgsql-check [ user <username> ]
  Use PostgreSQL health checks for server testing
  May be used in sections :   defaults | frontend | listen | backend
                                 yes   |    no    |   yes  |   yes
  Arguments :
    <username> This is the username which will be used when connecting to
               PostgreSQL server.

  The check sends a PostgreSQL StartupMessage and waits for either
  Authentication request or ErrorResponse message. It is a basic but useful
  test which does not produce error nor aborted connect on the server.
  This check is identical with the "mysql-check".

  See also: "option httpchk"


option prefer-last-server
no option prefer-last-server
  Allow multiple load balanced requests to remain on the same server
  May be used in sections:    defaults | frontend | listen | backend
                                 yes   |    no    |   yes  |   yes
  Arguments : none

  When the load balancing algorithm in use is not deterministic, and a previous
  request was sent to a server to which haproxy still holds a connection, it is
  sometimes desirable that subsequent requests on a same session go to the same
  server as much as possible. Note that this is different from persistence, as
```

```
   we only indicate a preference which haproxy tries to apply without any form
   of warranty. The real use is for keep-alive connections sent to servers. When
   this option is used, haproxy will try to reuse the same connection that is
   attached to the server instead of rebalancing to another server, causing a
   close of the connection. This can make sense for static file servers. It does
   not make much sense to use this in combination with hashing algorithms. Note,
   haproxy already automatically tries to stick to a server which sends a 401 or
   to a proxy which sends a 407 (authentication required). This is mandatory for
   use with the broken NTLM authentication challenge, and significantly helps in
   troubleshooting some faulty applications. Option prefer-last-server might be
   desirable in these environments as well, to avoid redistributing the traffic
   after every other response.

   If this option has been enabled in a "defaults" section, it can be disabled
   in a specific instance by prepending the "no" keyword before it.

   See also: "option http-keep-alive"


option redispatch
option redispatch <interval>
no option redispatch
  Enable or disable session redistribution in case of connection failure
  May be used in sections:    defaults | frontend | listen | backend
                                 yes   |    no    |   yes  |   yes
  Arguments :
    <interval> The optional integer value that controls how often redispatches
               occur when retrying connections. Positive value P indicates a
               redispatch is desired on every Pth retry, and negative value
               N indicate a redispath is desired on the Nth retry prior to the
               last retry. For example, the default of -1 preserves the
               historical behaviour of redispatching on the last retry, a
               positive value of 1 would indicate a redispatch on every retry,
               and a positive value of 3 would indicate a redispatch on every
               third retry. You can disable redispatches with a value of 0.


   In HTTP mode, if a server designated by a cookie is down, clients may
   definitely stick to it because they cannot flush the cookie, so they will not
   be able to access the service anymore.

   Specifying "option redispatch" will allow the proxy to break their
   persistence and redistribute them to a working server.

   It also allows to retry connections to another server in case of multiple
   connection failures. Of course, it requires having "retries" set to a nonzero
   value.

   This form is the preferred form, which replaces both the "redispatch" and
   "redisp" keywords.

   If this option has been enabled in a "defaults" section, it can be disabled
   in a specific instance by prepending the "no" keyword before it.

   See also : "redispatch", "retries", "force-persist"


option redis-check
  Use redis health checks for server testing
  May be used in sections :   defaults | frontend | listen | backend
                                 yes   |    no    |   yes  |   yes
  Arguments : none

  It is possible to test that the server correctly talks REDIS protocol instead
  of just testing that it accepts the TCP connection. When this option is set,
```



```
  a PING redis command is sent to the server, and the response is analyzed to
  find the "+PONG" response message.

  Example :
        option redis-check

  See also : "option httpchk"


option smtpchk
option smtpchk <hello> <domain>
  Use SMTP health checks for server testing
  May be used in sections :   defaults | frontend | listen | backend
                                 yes   |    no    |   yes  |   yes
  Arguments :
    <hello>   is an optional argument. It is the "hello" command to use. It can
              be either "HELO" (for SMTP) or "EHLO" (for ESTMP). All other
              values will be turned into the default command ("HELO").

    <domain>  is the domain name to present to the server. It may only be
              specified (and is mandatory) if the hello command has been
              specified. By default, "localhost" is used.

  When "option smtpchk" is set, the health checks will consist in TCP
  connections followed by an SMTP command. By default, this command is
  "HELO localhost". The server's return code is analyzed and only return codes
  starting with a "2" will be considered as valid. All other responses,
  including a lack of response will constitute an error and will indicate a
  dead server.

  This test is meant to be used with SMTP servers or relays. Depending on the
  request, it is possible that some servers do not log each connection attempt,
  so you may want to experiment to improve the behaviour. Using telnet on port
  25 is often easier than adjusting the configuration.

  Most often, an incoming SMTP server needs to see the client's IP address for
  various purposes, including spam filtering, anti-spoofing and logging. When
  possible, it is often wise to masquerade the client's IP address when
  connecting to the server using the "usesrc" argument of the "source" keyword,
  which requires the transparent proxy feature to be compiled in.

  Example :
        option smtpchk HELO mydomain.org

  See also : "option httpchk", "source"


option socket-stats
no option socket-stats

  Enable or disable collecting & providing separate statistics for each socket.
  May be used in sections :   defaults | frontend | listen | backend
                                 yes   |    yes   |   yes  |   no

  Arguments : none


option splice-auto
no option splice-auto
  Enable or disable automatic kernel acceleration on sockets in both directions
  May be used in sections :   defaults | frontend | listen | backend
                                 yes   |    yes   |   yes  |   yes
  Arguments : none

  When this option is enabled either on a frontend or on a backend, haproxy
```

```
   will automatically evaluate the opportunity to use kernel tcp splicing to
   forward data between the client and the server, in either direction. Haproxy
   uses heuristics to estimate if kernel splicing might improve performance or
   not. Both directions are handled independently. Note that the heuristics used
   are not much aggressive in order to limit excessive use of splicing. This
   option requires splicing to be enabled at compile time, and may be globally
   disabled with the global option "nosplice". Since splice uses pipes, using it
   requires that there are enough spare pipes.

   Important note: kernel-based TCP splicing is a Linux-specific feature which
   first appeared in kernel 2.6.25. It offers kernel-based acceleration to
   transfer data between sockets without copying these data to user-space, thus
   providing noticeable performance gains and CPU cycles savings. Since many
   early implementations are buggy, corrupt data and/or are inefficient, this
   feature is not enabled by default, and it should be used with extreme care.
   While it is not possible to detect the correctness of an implementation,
   2.6.29 is the first version offering a properly working implementation. In
   case of doubt, splicing may be globally disabled using the global "nosplice"
   keyword.

   Example :
         option splice-auto

   If this option has been enabled in a "defaults" section, it can be disabled
   in a specific instance by prepending the "no" keyword before it.

   See also : "option splice-request", "option splice-response", and global
              options "nosplice" and "maxpipes"


option splice-request
no option splice-request
  Enable or disable automatic kernel acceleration on sockets for requests
  May be used in sections :   defaults | frontend | listen | backend
                                 yes   |    yes   |   yes  |   yes
  Arguments : none

  When this option is enabled either on a frontend or on a backend, haproxy
  will use kernel tcp splicing whenever possible to forward data going from
  the client to the server. It might still use the recv/send scheme if there
  are no spare pipes left. This option requires splicing to be enabled at
  compile time, and may be globally disabled with the global option "nosplice".
  Since splice uses pipes, using it requires that there are enough spare pipes.

  Important note: see "option splice-auto" for usage limitations.

  Example :
        option splice-request

  If this option has been enabled in a "defaults" section, it can be disabled
  in a specific instance by prepending the "no" keyword before it.

  See also : "option splice-auto", "option splice-response", and global options
             "nosplice" and "maxpipes"


option splice-response
no option splice-response
  Enable or disable automatic kernel acceleration on sockets for responses
  May be used in sections :   defaults | frontend | listen | backend
                                 yes   |    yes   |   yes  |   yes
  Arguments : none

  When this option is enabled either on a frontend or on a backend, haproxy
  will use kernel tcp splicing whenever possible to forward data going from
```



```
  the server to the client. It might still use the recv/send scheme if there
  are no spare pipes left. This option requires splicing to be enabled at
  compile time, and may be globally disabled with the global option "nosplice".
  Since splice uses pipes, using it requires that there are enough spare pipes.

  Important note: see "option splice-auto" for usage limitations.

  Example :
        option splice-response

  If this option has been enabled in a "defaults" section, it can be disabled
  in a specific instance by prepending the "no" keyword before it.

  See also : "option splice-auto", "option splice-request", and global options
             "nosplice" and "maxpipes"


option srvtcpka
no option srvtcpka
  Enable or disable the sending of TCP keepalive packets on the server side
  May be used in sections :   defaults | frontend | listen | backend
                                 yes   |    no    |   yes  |   yes
  Arguments : none

  When there is a firewall or any session-aware component between a client and
  a server, and when the protocol involves very long sessions with long idle
  periods (eg: remote desktops), there is a risk that one of the intermediate
  components decides to expire a session which has remained idle for too long.

  Enabling socket-level TCP keep-alives makes the system regularly send packets
  to the other end of the connection, leaving it active. The delay between
  keep-alive probes is controlled by the system only and depends both on the
  operating system and its tuning parameters.

  It is important to understand that keep-alive packets are neither emitted nor
  received at the application level. It is only the network stacks which sees
  them. For this reason, even if one side of the proxy already uses keep-alives
  to maintain its connection alive, those keep-alive packets will not be
  forwarded to the other side of the proxy.

  Please note that this has nothing to do with HTTP keep-alive.

  Using option "srvtcpka" enables the emission of TCP keep-alive probes on the
  server side of a connection, which should help when session expirations are
  noticed between HAProxy and a server.

  If this option has been enabled in a "defaults" section, it can be disabled
  in a specific instance by prepending the "no" keyword before it.

  See also : "option clitcpka", "option tcpka"


option ssl-hello-chk
  Use SSLv3 client hello health checks for server testing
  May be used in sections :   defaults | frontend | listen | backend
                                 yes   |    no    |   yes  |   yes
  Arguments : none

  When some SSL-based protocols are relayed in TCP mode through HAProxy, it is
  possible to test that the server correctly talks SSL instead of just testing
  that it accepts the TCP connection. When "option ssl-hello-chk" is set, pure
  SSLv3 client hello messages are sent once the connection is established to
  the server, and the response is analyzed to find an SSL server hello message.
  The server is considered valid only when the response contains this server
  hello message.
```

```
6111
6112    All servers tested till there correctly reply to SSLv3 client hello messages,
6113    and most servers tested do not even log the requests containing only hello
6114    messages, which is appreciable.
6115
6116    Note that this check works even when SSL support was not built into haproxy
6117    because it forges the SSL message. When SSL support is available, it is best
6118    to use native SSL health checks instead of this one.
6119
6120    See also: "option httpchk", "check-ssl"
6121
6122
6123  option tcp-check
6124    Perform health checks using tcp-check send/expect sequences
6125    May be used in sections:   defaults | frontend | listen | backend
6126                                  yes    |    no    |   yes  |   yes
6127
6128    This health check method is intended to be combined with "tcp-check" command
6129    lists in order to support send/expect types of health check sequences.
6130
6131    TCP checks currently support 4 modes of operations :
6132      - no "tcp-check" directive : the health check only consists in a connection
6133        attempt, which remains the default mode.
6134
6135      - "tcp-check send" or "tcp-check send-binary" only is mentioned : this is
6136        used to send a string along with a connection opening. With some
6137        protocols, it helps sending a "QUIT" message for example that prevents
6138        the server from logging a connection error for each health check. The
6139        check result will still be based on the ability to open the connection
6140        only.
6141
6142      - "tcp-check expect" only is mentioned : this is used to test a banner.
6143        The connection is opened and haproxy waits for the server to present some
6144        contents which must validate some rules. The check result will be based
6145        on the matching between the contents and the rules. This is suited for
6146        POP, IMAP, SMTP, FTP, SSH, TELNET.
6147
6148      - both "tcp-check send" and "tcp-check expect" are mentioned : this is
6149        used to test a hello-type protocol. Haproxy sends a message, the server
6150        responds and its response is analysed. the check result will be based on
6151        the matching between the response contents and the rules. This is often
6152        suited for protocols which require a binding or a request/response model.
6153        LDAP, MySQL, Redis and SSL are example of such protocols, though they
6154        already all have their dedicated checks with a deeper understanding of
6155        the respective protocols.
6156        In this mode, many questions may be sent and many answers may be
6157        analysed.
6158
6159      A fifth mode can be used to insert comments in different steps of the
6160      script.
6161
6162      For each tcp-check rule you create, you can add a "comment" directive,
6163      followed by a string. This string will be reported in the log and stderr
6164      in debug mode. It is useful to make user-friendly error reporting.
6165      The "comment" is of course optional.
6166
6167
6168    Examples :
6169          # perform a POP check (analyse only server's banner)
6170          option tcp-check
6171          tcp-check expect string +OK\ POP3\ ready comment POP\ protocol
6172
6173          # perform an IMAP check (analyse only server's banner)
6174          option tcp-check
6175          tcp-check expect string *\ OK\ IMAP4\ ready comment IMAP\ protocol
```



```
6176
6177          # look for the redis master server after ensuring it speaks well
6178          # redis protocol, then it exits properly.
6179          # (send a command then analyse the response 3 times)
6180          option tcp-check
6181          tcp-check comment PING\ phase
6182          tcp-check send PING\r\n
6183          tcp-check expect string +PONG
6184          tcp-check comment role\ check
6185          tcp-check send info\ replication\r\n
6186          tcp-check expect string role:master
6187          tcp-check comment QUIT\ phase
6188          tcp-check send QUIT\r\n
6189          tcp-check expect string +OK
6190
6191          forge a HTTP request, then analyse the response
6192          (send many headers before analyzing)
6193          option tcp-check
6194          tcp-check comment forge\ and\ send\ HTTP\ request
6195          tcp-check send HEAD\ /\ HTTP/1.1\r\n
6196          tcp-check send Host:\ www.mydomain.com\r\n
6197          tcp-check send User-Agent:\ HAProxy\ tcpcheck\r\n
6198          tcp-check send \r\n
6199          tcp-check expect rstring HTTP/1\..\ (2..|3..) comment check\ HTTP\ response
6200
6201
6202    See also : "tcp-check expect", "tcp-check send"
6203
6204
6205  option tcp-smart-accept
6206  no option tcp-smart-accept
6207    Enable or disable the saving of one ACK packet during the accept sequence
6208    May be used in sections :   defaults | frontend | listen | backend
6209                                   yes   |    yes   |   yes  |    no
6210    Arguments : none
6211
6212    When an HTTP connection request comes in, the system acknowledges it on
6213    behalf of HAProxy, then the client immediately sends its request, and the
6214    system acknowledges it too while it is notifying HAProxy about the new
6215    connection. HAProxy then reads the request and responds. This means that we
6216    have one TCP ACK sent by the system for nothing, because the request could
6217    very well be acknowledged by HAProxy when it sends its response.
6218
6219    For this reason, in HTTP mode, HAProxy automatically asks the system to avoid
6220    sending this useless ACK on platforms which support it (currently at least
6221    Linux). It must not cause any problem, because the system will send it anyway
6222    after 40 ms if the response takes more time than expected to come.
6223
6224    During complex network debugging sessions, it may be desirable to disable
6225    this optimization because delayed ACKs can make troubleshooting more complex
6226    when trying to identify where packets are delayed. It is then possible to
6227    fall back to normal behaviour by specifying "no option tcp-smart-accept".
6228
6229    It is also possible to force it for non-HTTP proxies by simply specifying
6230    "option tcp-smart-accept". For instance, it can make sense with some services
6231    such as SMTP where the server speaks first.
6232
6233    It is recommended to avoid forcing this option in a defaults section. In case
6234    of doubt, consider setting it back to automatic values by prepending the
6235    "default" keyword before it, or disabling it using the "no" keyword.
6236
6237    See also : "option tcp-smart-connect"
6238
6239
6240  option tcp-smart-connect
```

```
no option tcp-smart-connect
  Enable or disable the saving of one ACK packet during the connect sequence
  May be used in sections :   defaults | frontend | listen | backend
                                 yes   |    no    |   yes  |   yes
  Arguments : none

  On certain systems (at least Linux), HAProxy can ask the kernel not to
  immediately send an empty ACK upon a connection request, but to directly
  send the buffer request instead. This saves one packet on the network and
  thus boosts performance. It can also be useful for some servers, because they
  immediately get the request along with the incoming connection.

  This feature is enabled when "option tcp-smart-connect" is set in a backend.
  It is not enabled by default because it makes network troubleshooting more
  complex.

  It only makes sense to enable it with protocols where the client speaks first
  such as HTTP. In other situations, if there is no data to send in place of
  the ACK, a normal ACK is sent.

  If this option has been enabled in a "defaults" section, it can be disabled
  in a specific instance by prepending the "no" keyword before it.

  See also : "option tcp-smart-accept"


option tcpka
  Enable or disable the sending of TCP keepalive packets on both sides
  May be used in sections :   defaults | frontend | listen | backend
                                 yes   |    yes   |   yes  |   yes
  Arguments : none

  When there is a firewall or any session-aware component between a client and
  a server, and when the protocol involves very long sessions with long idle
  periods (eg: remote desktops), there is a risk that one of the intermediate
  components decides to expire a session which has remained idle for too long.

  Enabling socket-level TCP keep-alives makes the system regularly send packets
  to the other end of the connection, leaving it active. The delay between
  keep-alive probes is controlled by the system only and depends both on the
  operating system and its tuning parameters.

  It is important to understand that keep-alive packets are neither emitted nor
  received at the application level. It is only the network stacks which sees
  them. For this reason, even if one side of the proxy already uses keep-alives
  to maintain its connection alive, those keep-alive packets will not be
  forwarded to the other side of the proxy.

  Please note that this has nothing to do with HTTP keep-alive.

  Using option "tcpka" enables the emission of TCP keep-alive probes on both
  the client and server sides of a connection. Note that this is meaningful
  only in "defaults" or "listen" sections. If this option is used in a
  frontend, only the client side will get keep-alives, and if this option is
  used in a backend, only the server side will get keep-alives. For this
  reason, it is strongly recommended to explicitly use "option clitcpka" and
  "option srvtcpka" when the configuration is split between frontends and
  backends.

  See also : "option clitcpka", "option srvtcpka"


option tcplog
  Enable advanced logging of TCP connections with session state and timers
  May be used in sections :   defaults | frontend | listen | backend
```



```
                                 yes   |    yes   |   yes  |   yes
  Arguments : none

  By default, the log output format is very poor, as it only contains the
  source and destination addresses, and the instance name. By specifying
  "option tcplog", each log line turns into a much richer format including, but
  not limited to, the connection timers, the session status, the connections
  numbers, the frontend, backend and server name, and of course the source
  address and ports. This option is useful for pure TCP proxies in order to
  find which of the client or server disconnects or times out. For normal HTTP
  proxies, it's better to use "option httplog" which is even more complete.

  This option may be set either in the frontend or the backend.

  See also :  "option httplog", and section 8 about logging.


option transparent
no option transparent
  Enable client-side transparent proxying
  May be used in sections :   defaults | frontend | listen | backend
                                 yes   |    no    |   yes  |   yes
  Arguments : none

  This option was introduced in order to provide layer 7 persistence to layer 3
  load balancers. The idea is to use the OS's ability to redirect an incoming
  connection for a remote address to a local process (here HAProxy), and let
  this process know what address was initially requested. When this option is
  used, sessions without cookies will be forwarded to the original destination
  IP address of the incoming request (which should match that of another
  equipment), while requests with cookies will still be forwarded to the
  appropriate server.

  Note that contrary to a common belief, this option does NOT make HAProxy
  present the client's IP to the server when establishing the connection.

  See also: the "usesrc" argument of the "source" keyword, and the
            "transparent" option of the "bind" keyword.


external-check command <command>
  Executable to run when performing an external-check
  May be used in sections :   defaults | frontend | listen | backend
                                 yes   |    no    |   yes  |   yes

  Arguments :
    <command> is the external command to run

  The arguments passed to the to the command are:

  <proxy_address> <proxy_port> <server_address> <server_port>

  The <proxy_address> and <proxy_port> are derived from the first listener
  that is either IPv4, IPv6 or a UNIX socket. In the case of a UNIX socket
  listener the proxy_address will be the path of the socket and the
  <proxy_port> will be the string "NOT_USED". In a backend section, it's not
  possible to determine a listener, and both <proxy_address> and <proxy_port>
  will have the string value "NOT_USED".

  Some values are also provided through environment variables.

  Environment variables :
    HAPROXY_PROXY_ADDR      The first bind address if available (or empty if not
                            applicable, for example in a "backend" section).
```

```
      HAPROXY_PROXY_ID        The backend id.

      HAPROXY_PROXY_NAME      The backend name.

      HAPROXY_PROXY_PORT      The first bind port if available (or empty if not
                              applicable, for example in a "backend" section or
                              for a UNIX socket).

      HAPROXY_SERVER_ADDR     The server address.

      HAPROXY_SERVER_CURCONN  The current number of connections on the server.

      HAPROXY_SERVER_ID       The server id.

      HAPROXY_SERVER_MAXCONN  The server max connections.

      HAPROXY_SERVER_NAME     The server name.

      HAPROXY_SERVER_PORT     The server port if available (or empty for a UNIX
                              socket).

      PATH                    The PATH environment variable used when executing
                              the command may be set using "external-check path".

    If the command executed and exits with a zero status then the check is
    considered to have passed, otherwise the check is considered to have
    failed.

    Example :
          external-check command /bin/true

    See also : "external-check", "option external-check", "external-check path"


  external-check path <path>
    The value of the PATH environment variable used when running an external-check
    May be used in sections :   defaults | frontend | listen | backend
                                   yes   |    no    |   yes  |   yes

    Arguments :
      <path> is the path used when executing external command to run

    The default path is "".

    Example :
          external-check path "/usr/bin:/bin"

    See also : "external-check", "option external-check",
               "external-check command"


  persist rdp-cookie
  persist rdp-cookie(<name>)
    Enable RDP cookie-based persistence
    May be used in sections :   defaults | frontend | listen | backend
                                   yes   |    no    |   yes  |   yes
    Arguments :
      <name>    is the optional name of the RDP cookie to check. If omitted, the
                default cookie name "msts" will be used. There currently is no
                valid reason to change this name.

    This statement enables persistence based on an RDP cookie. The RDP cookie
    contains all information required to find the server in the list of known
    servers. So when this option is set in the backend, the request is analysed
    and if an RDP cookie is found, it is decoded. If it matches a known server
```



```
    which is still UP (or if "option persist" is set), then the connection is
    forwarded to this server.

    Note that this only makes sense in a TCP backend, but for this to work, the
    frontend must have waited long enough to ensure that an RDP cookie is present
    in the request buffer. This is the same requirement as with the "rdp-cookie"
    load-balancing method. Thus it is highly recommended to put all statements in
    a single "listen" section.

    Also, it is important to understand that the terminal server will emit this
    RDP cookie only if it is configured for "token redirection mode", which means
    that the "IP address redirection" option is disabled.

    Example :
          listen tse-farm
              bind :3389
              # wait up to 5s for an RDP cookie in the request
              tcp-request inspect-delay 5s
              tcp-request content accept if RDP_COOKIE
              # apply RDP cookie persistence
              persist rdp-cookie
              # if server is unknown, let's balance on the same cookie.
              # alternatively, "balance leastconn" may be useful too.
              balance rdp-cookie
              server srv1 1.1.1.1:3389
              server srv2 1.1.1.2:3389

    See also : "balance rdp-cookie", "tcp-request", the "req_rdp_cookie" ACL and
    the rdp_cookie pattern fetch function.


  rate-limit sessions <rate>
    Set a limit on the number of new sessions accepted per second on a frontend
    May be used in sections :   defaults | frontend | listen | backend
                                   yes   |    yes   |   yes  |   no
    Arguments :
      <rate>    The <rate> parameter is an integer designating the maximum number
                of new sessions per second to accept on the frontend.

    When the frontend reaches the specified number of new sessions per second, it
    stops accepting new connections until the rate drops below the limit again.
    During this time, the pending sessions will be kept in the socket's backlog
    (in system buffers) and haproxy will not even be aware that sessions are
    pending. When applying very low limit on a highly loaded service, it may make
    sense to increase the socket's backlog using the "backlog" keyword.

    This feature is particularly efficient at blocking connection-based attacks
    or service abuse on fragile servers. Since the session rate is measured every
    millisecond, it is extremely accurate. Also, the limit applies immediately,
    no delay is needed at all to detect the threshold.

    Example : limit the connection rate on SMTP to 10 per second max
          listen smtp
              mode tcp
              bind :25
              rate-limit sessions 10
              server 127.0.0.1:1025

    Note : when the maximum rate is reached, the frontend's status is not changed
           but its sockets appear as "WAITING" in the statistics if the
           "socket-stats" option is enabled.

    See also : the "backlog" keyword and the "fe_sess_rate" ACL criterion.
```

```
redirect location <loc> [code <code>] <option> [{if | unless} <condition>]
redirect prefix   <pfx> [code <code>] <option> [{if | unless} <condition>]
redirect scheme   <sch> [code <code>] <option> [{if | unless} <condition>]
  Return an HTTP redirection if/unless a condition is matched
  May be used in sections :   defaults | frontend | listen | backend
                                 no    |    yes   |   yes  |   yes

  If/unless the condition is matched, the HTTP request will lead to a redirect
  response. If no condition is specified, the redirect applies unconditionally.

  Arguments :
    <loc>     With "redirect location", the exact value in <loc> is placed into
              the HTTP "Location" header. When used in an "http-request" rule,
              <loc> value follows the log-format rules and can include some
              dynamic values (see Custom Log Format in section 8.2.4).

    <pfx>     With "redirect prefix", the "Location" header is built from the
              concatenation of <pfx> and the complete URI path, including the
              query string, unless the "drop-query" option is specified (see
              below). As a special case, if <pfx> equals exactly "/", then
              nothing is inserted before the original URI. It allows one to
              redirect to the same URL (for instance, to insert a cookie). When
              used in an "http-request" rule, <pfx> value follows the log-format
              rules and can include some dynamic values (see Custom Log Format
              in section 8.2.4).

    <sch>     With "redirect scheme", then the "Location" header is built by
              concatenating <sch> with "://" then the first occurrence of the
              "Host" header, and then the URI path, including the query string
              unless the "drop-query" option is specified (see below). If no
              path is found or if the path is "*", then "/" is used instead. If
              no "Host" header is found, then an empty host component will be
              returned, which most recent browsers interpret as redirecting to
              the same host. This directive is mostly used to redirect HTTP to
              HTTPS. When used in an "http-request" rule, <sch> value follows
              the log-format rules and can include some dynamic values (see
              Custom Log Format in section 8.2.4).

    <code>    The code is optional. It indicates which type of HTTP redirection
              is desired. Only codes 301, 302, 303, 307 and 308 are supported,
              with 302 used by default if no code is specified. 301 means
              "Moved permanently", and a browser may cache the Location. 302
              means "Moved temporarily" and means that the browser should not
              cache the redirection. 303 is equivalent to 302 except that the
              browser will fetch the location with a GET method. 307 is just
              like 302 but makes it clear that the same method must be reused.
              Likewise, 308 replaces 301 if the same method must be used.

    <option>  There are several options which can be specified to adjust the
              expected behaviour of a redirection :

      - "drop-query"
        When this keyword is used in a prefix-based redirection, then the
        location will be set without any possible query-string, which is useful
        for directing users to a non-secure page for instance. It has no effect
        with a location-type redirect.

      - "append-slash"
        This keyword may be used in conjunction with "drop-query" to redirect
        users who use a URL not ending with a '/' to the same one with the '/'.
        It can be useful to ensure that search engines will only see one URL.
        For this, a return code 301 is preferred.

      - "set-cookie NAME[=value]"
        A "Set-Cookie" header will be added with NAME (and optionally "=value")
```



```
        to the response. This is sometimes used to indicate that a user has
        been seen, for instance to protect against some types of DoS. No other
        cookie option is added, so the cookie will be a session cookie. Note
        that for a browser, a sole cookie name without an equal sign is
        different from a cookie with an equal sign.

      - "clear-cookie NAME[=]"
        A "Set-Cookie" header will be added with NAME (and optionally "="), but
        with the "Max-Age" attribute set to zero. This will tell the browser to
        delete this cookie. It is useful for instance on logout pages. It is
        important to note that clearing the cookie "NAME" will not remove a
        cookie set with "NAME=value". You have to clear the cookie "NAME=" for
        that, because the browser makes the difference.

  Example: move the login URL only to HTTPS.
        acl clear      dst_port  80
        acl secure     dst_port  8080
        acl login_page url_beg   /login
        acl logout     url_beg   /logout
        acl uid_given  url_reg   /login?userid=[^&]+
        acl cookie_set hdr_sub(cookie) SEEN=1

        redirect prefix   https://mysite.com set-cookie SEEN=1 if !cookie_set
        redirect prefix   https://mysite.com           if login_page !secure
        redirect prefix   http://mysite.com drop-query if login_page !uid_given
        redirect location http://mysite.com/           if !login_page secure
        redirect location / clear-cookie USERID=       if logout

  Example: send redirects for request for articles without a '/'.
        acl missing_slash path_reg ^/article/[^/]*$
        redirect code 301 prefix / drop-query append-slash if missing_slash

  Example: redirect all HTTP traffic to HTTPS when SSL is handled by haproxy.
        redirect scheme https if !{ ssl_fc }

  Example: append 'www.' prefix in front of all hosts not having it
        http-request redirect code 301 location www.%[hdr(host)]%[req.uri] \
          unless { hdr_beg(host) -i www }

  See section 7 about ACL usage.


redisp (deprecated)
redispatch (deprecated)
  Enable or disable session redistribution in case of connection failure
  May be used in sections:    defaults | frontend | listen | backend
                                 yes   |    no    |   yes  |   yes
  Arguments : none

  In HTTP mode, if a server designated by a cookie is down, clients may
  definitely stick to it because they cannot flush the cookie, so they will not
  be able to access the service anymore.

  Specifying "redispatch" will allow the proxy to break their persistence and
  redistribute them to a working server.

  It also allows to retry last connection to another server in case of multiple
  connection failures. Of course, it requires having "retries" set to a nonzero
  value.

  This form is deprecated, do not use it in any new configuration, use the new
  "option redispatch" instead.

  See also : "option redispatch"
```

```
6631
6632  reqadd  <string> [{if | unless} <cond>]
6633    Add a header at the end of the HTTP request
6634    May be used in sections :   defaults | frontend | listen | backend
6635                                   no    |    yes   |   yes  |   yes
6636    Arguments :
6637      <string>  is the complete line to be added. Any space or known delimiter
6638                must be escaped using a backslash ('\'). Please refer to section
6639                6 about HTTP header manipulation for more information.
6640
6641      <cond>    is an optional matching condition built from ACLs. It makes it
6642                possible to ignore this rule when other conditions are not met.
6643
6644    A new line consisting in <string> followed by a line feed will be added after
6645    the last header of an HTTP request.
6646
6647    Header transformations only apply to traffic which passes through HAProxy,
6648    and not to traffic generated by HAProxy, such as health-checks or error
6649    responses.
6650
6651    Example : add "X-Proto: SSL" to requests coming via port 81
6652       acl is-ssl  dst_port       81
6653       reqadd      X-Proto:\ SSL  if is-ssl
6654
6655    See also: "rspadd", "http-request", section 6 about HTTP header manipulation,
6656              and section 7 about ACLs.
6657
6658
6659  reqallow  <search> [{if | unless} <cond>]
6660  reqiallow <search> [{if | unless} <cond>] (ignore case)
6661    Definitely allow an HTTP request if a line matches a regular expression
6662    May be used in sections :   defaults | frontend | listen | backend
6663                                   no    |    yes   |   yes  |   yes
6664    Arguments :
6665      <search>  is the regular expression applied to HTTP headers and to the
6666                request line. This is an extended regular expression. Parenthesis
6667                grouping is supported and no preliminary backslash is required.
6668                Any space or known delimiter must be escaped using a backslash
6669                ('\'). The pattern applies to a full line at a time. The
6670                "reqallow" keyword strictly matches case while "reqiallow"
6671                ignores case.
6672
6673      <cond>    is an optional matching condition built from ACLs. It makes it
6674                possible to ignore this rule when other conditions are not met.
6675
6676    A request containing any line which matches extended regular expression
6677    <search> will mark the request as allowed, even if any later test would
6678    result in a deny. The test applies both to the request line and to request
6679    headers. Keep in mind that URLs in request line are case-sensitive while
6680    header names are not.
6681
6682    It is easier, faster and more powerful to use ACLs to write access policies.
6683    Reqdeny, reqallow and reqpass should be avoided in new designs.
6684
6685    Example :
6686       # allow www.* but refuse *.local
6687       reqiallow ^Host:\ www\.
6688       reqideny  ^Host:\ .*\.local
6689
6690    See also: "reqdeny", "block", "http-request",  section 6 about HTTP header
6691              manipulation, and section 7 about ACLs.
6692
6693
6694  reqdel  <search> [{if | unless} <cond>]
6695  reqidel <search> [{if | unless} <cond>]  (ignore case)
```



```
6696    Delete all headers matching a regular expression in an HTTP request
6697    May be used in sections :   defaults | frontend | listen | backend
6698                                   no    |    yes   |   yes  |   yes
6699    Arguments :
6700      <search>  is the regular expression applied to HTTP headers and to the
6701                request line. This is an extended regular expression. Parenthesis
6702                grouping is supported and no preliminary backslash is required.
6703                Any space or known delimiter must be escaped using a backslash
6704                ('\'). The pattern applies to a full line at a time. The "reqdel"
6705                keyword strictly matches case while "reqidel" ignores case.
6706
6707      <cond>    is an optional matching condition built from ACLs. It makes it
6708                possible to ignore this rule when other conditions are not met.
6709
6710    Any header line matching extended regular expression <search> in the request
6711    will be completely deleted. Most common use of this is to remove unwanted
6712    and/or dangerous headers or cookies from a request before passing it to the
6713    next servers.
6714
6715    Header transformations only apply to traffic which passes through HAProxy,
6716    and not to traffic generated by HAProxy, such as health-checks or error
6717    responses. Keep in mind that header names are not case-sensitive.
6718
6719    Example :
6720       # remove X-Forwarded-For header and SERVER cookie
6721       reqidel ^X-Forwarded-For:.*
6722       reqidel ^Cookie:.*SERVER=
6723
6724    See also: "reqadd", "reqrep", "rspdel", "http-request", section 6 about
6725              HTTP header manipulation, and section 7 about ACLs.
6726
6727
6728  reqdeny  <search> [{if | unless} <cond>]
6729  reqideny <search> [{if | unless} <cond>]  (ignore case)
6730    Deny an HTTP request if a line matches a regular expression
6731    May be used in sections :   defaults | frontend | listen | backend
6732                                   no    |    yes   |   yes  |   yes
6733    Arguments :
6734      <search>  is the regular expression applied to HTTP headers and to the
6735                request line. This is an extended regular expression. Parenthesis
6736                grouping is supported and no preliminary backslash is required.
6737                Any space or known delimiter must be escaped using a backslash
6738                ('\'). The pattern applies to a full line at a time. The
6739                "reqdeny" keyword strictly matches case while "reqideny" ignores
6740                case.
6741
6742      <cond>    is an optional matching condition built from ACLs. It makes it
6743                possible to ignore this rule when other conditions are not met.
6744
6745    A request containing any line which matches extended regular expression
6746    <search> will mark the request as denied, even if any later test would
6747    result in an allow. The test applies both to the request line and to request
6748    headers. Keep in mind that URLs in request line are case-sensitive while
6749    header names are not.
6750
6751    A denied request will generate an "HTTP 403 forbidden" response once the
6752    complete request has been parsed. This is consistent with what is practiced
6753    using ACLs.
6754
6755    It is easier, faster and more powerful to use ACLs to write access policies.
6756    Reqdeny, reqallow and reqpass should be avoided in new designs.
6757
6758    Example :
6759       # refuse *.local, then allow www.*
6760       reqideny  ^Host:\ .*\.local
```

```
      reqiallow ^Host:\ www\.

   See also: "reqallow", "rspdeny", "block", "http-request", section 6 about
             HTTP header manipulation, and section 7 about ACLs.


reqpass  <search> [{if | unless} <cond>]
reqipass <search> [{if | unless} <cond>]  (ignore case)
  Ignore any HTTP request line matching a regular expression in next rules
  May be used in sections :   defaults | frontend | listen | backend
                                 no    |    yes   |   yes  |   yes
  Arguments :
    <search>  is the regular expression applied to HTTP headers and to the
              request line. This is an extended regular expression. Parenthesis
              grouping is supported and no preliminary backslash is required.
              Any space or known delimiter must be escaped using a backslash
              ('\'). The pattern applies to a full line at a time. The
              "reqpass" keyword strictly matches case while "reqipass" ignores
              case.

    <cond>    is an optional matching condition built from ACLs. It makes it
              possible to ignore this rule when other conditions are not met.

  A request containing any line which matches extended regular expression
  <search> will skip next rules, without assigning any deny or allow verdict.
  The test applies both to the request line and to request headers. Keep in
  mind that URLs in request line are case-sensitive while header names are not.

  It is easier, faster and more powerful to use ACLs to write access policies.
  Reqdeny, reqallow and reqpass should be avoided in new designs.

  Example :
     # refuse *.local, then allow www.*, but ignore "www.private.local"
     reqipass  ^Host:\ www.private\.local
     reqideny  ^Host:\ .*\.local
     reqiallow ^Host:\ www\.

  See also: "reqallow", "reqdeny", "block", "http-request", section 6 about
            HTTP header manipulation, and section 7 about ACLs.


reqrep  <search> <string> [{if | unless} <cond>]
reqirep <search> <string> [{if | unless} <cond>]   (ignore case)
  Replace a regular expression with a string in an HTTP request line
  May be used in sections :   defaults | frontend | listen | backend
                                 no    |    yes   |   yes  |   yes
  Arguments :
    <search>  is the regular expression applied to HTTP headers and to the
              request line. This is an extended regular expression. Parenthesis
              grouping is supported and no preliminary backslash is required.
              Any space or known delimiter must be escaped using a backslash
              ('\'). The pattern applies to a full line at a time. The "reqrep"
              keyword strictly matches case while "reqirep" ignores case.

    <string>  is the complete line to be added. Any space or known delimiter
              must be escaped using a backslash ('\'). References to matched
              pattern groups are possible using the common \N form, with N
              being a single digit between 0 and 9. Please refer to section
              6 about HTTP header manipulation for more information.

    <cond>    is an optional matching condition built from ACLs. It makes it
              possible to ignore this rule when other conditions are not met.

  Any line matching extended regular expression <search> in the request (both
  the request line and header lines) will be completely replaced with <string>.
```



```
  Most common use of this is to rewrite URLs or domain names in "Host" headers.

  Header transformations only apply to traffic which passes through HAProxy,
  and not to traffic generated by HAProxy, such as health-checks or error
  responses. Note that for increased readability, it is suggested to add enough
  spaces between the request and the response. Keep in mind that URLs in
  request line are case-sensitive while header names are not.

  Example :
     # replace "/static/" with "/" at the beginning of any request path.
     reqrep ^([^\ :]*)\ /static/(.*)     \1\ /\2
     # replace "www.mydomain.com" with "www" in the host name.
     reqirep ^Host:\ www.mydomain.com   Host:\ www

  See also: "reqadd", "reqdel", "rsprep", "tune.bufsize", "http-request",
            section 6 about HTTP header manipulation, and section 7 about ACLs.


reqtarpit  <search> [{if | unless} <cond>]
reqitarpit <search> [{if | unless} <cond>]  (ignore case)
  Tarpit an HTTP request containing a line matching a regular expression
  May be used in sections :   defaults | frontend | listen | backend
                                 no    |    yes   |   yes  |   yes
  Arguments :
    <search>  is the regular expression applied to HTTP headers and to the
              request line. This is an extended regular expression. Parenthesis
              grouping is supported and no preliminary backslash is required.
              Any space or known delimiter must be escaped using a backslash
              ('\'). The pattern applies to a full line at a time. The
              "reqtarpit" keyword strictly matches case while "reqitarpit"
              ignores case.

    <cond>    is an optional matching condition built from ACLs. It makes it
              possible to ignore this rule when other conditions are not met.

  A request containing any line which matches extended regular expression
  <search> will be tarpitted, which means that it will connect to nowhere, will
  be kept open for a pre-defined time, then will return an HTTP error 500 so
  that the attacker does not suspect it has been tarpitted. The status 500 will
  be reported in the logs, but the completion flags will indicate "PT". The
  delay is defined by "timeout tarpit", or "timeout connect" if the former is
  not set.

  The goal of the tarpit is to slow down robots attacking servers with
  identifiable requests. Many robots limit their outgoing number of connections
  and stay connected waiting for a reply which can take several minutes to
  come. Depending on the environment and attack, it may be particularly
  efficient at reducing the load on the network and firewalls.

  Examples :
     # ignore user-agents reporting any flavour of "Mozilla" or "MSIE", but
     # block all others.
     reqipass   ^User-Agent:\.*(Mozilla|MSIE)
     reqitarpit ^User-Agent:

     # block bad guys
     acl badguys src 10.1.0.3 172.16.13.20/28
     reqitarpit . if badguys

  See also: "reqallow", "reqdeny", "reqpass", "http-request",  section 6
            about HTTP header manipulation, and section 7 about ACLs.


retries <value>
  Set the number of retries to perform on a server after a connection failure
```

```
   May be used in sections:    defaults | frontend | listen | backend
                                  yes   |    no    |   yes  |   yes
   Arguments :
     <value>   is the number of times a connection attempt should be retried on
               a server when a connection either is refused or times out. The
               default value is 3.

   It is important to understand that this value applies to the number of
   connection attempts, not full requests. When a connection has effectively
   been established to a server, there will be no more retry.

   In order to avoid immediate reconnections to a server which is restarting,
   a turn-around timer of min("timeout connect", one second) is applied before
   a retry occurs.

   When "option redispatch" is set, the last retry may be performed on another
   server even if a cookie references a different server.

   See also : "option redispatch"


rspadd <string> [{if | unless} <cond>]
  Add a header at the end of the HTTP response
  May be used in sections :   defaults | frontend | listen | backend
                                 no    |    yes   |   yes  |   yes
  Arguments :
    <string>  is the complete line to be added. Any space or known delimiter
              must be escaped using a backslash ('\'). Please refer to section
              6 about HTTP header manipulation for more information.

    <cond>    is an optional matching condition built from ACLs. It makes it
              possible to ignore this rule when other conditions are not met.

  A new line consisting in <string> followed by a line feed will be added after
  the last header of an HTTP response.

  Header transformations only apply to traffic which passes through HAProxy,
  and not to traffic generated by HAProxy, such as health-checks or error
  responses.

  See also: "rspdel" "reqadd", "http-response", section 6 about HTTP header
            manipulation, and section 7 about ACLs.


rspdel  <search> [{if | unless} <cond>]
rspidel <search> [{if | unless} <cond>]  (ignore case)
  Delete all headers matching a regular expression in an HTTP response
  May be used in sections :   defaults | frontend | listen | backend
                                 no    |    yes   |   yes  |   yes
  Arguments :
    <search>  is the regular expression applied to HTTP headers and to the
              response line. This is an extended regular expression, so
              parenthesis grouping is supported and no preliminary backslash
              is required. Any space or known delimiter must be escaped using
              a backslash ('\'). The pattern applies to a full line at a time.
              The "rspdel" keyword strictly matches case while "rspidel"
              ignores case.

    <cond>    is an optional matching condition built from ACLs. It makes it
              possible to ignore this rule when other conditions are not met.

  Any header line matching extended regular expression <search> in the response
  will be completely deleted. Most common use of this is to remove unwanted
  and/or sensitive headers or cookies from a response before passing it to the
  client.
```



```

  Header transformations only apply to traffic which passes through HAProxy,
  and not to traffic generated by HAProxy, such as health-checks or error
  responses. Keep in mind that header names are not case-sensitive.

  Example :
     # remove the Server header from responses
     rspidel ^Server:.*

  See also: "rspadd", "rsprep", "reqdel", "http-response",  section 6 about
            HTTP header manipulation, and section 7 about ACLs.


rspdeny  <search> [{if | unless} <cond>]
rspideny <search> [{if | unless} <cond>]  (ignore case)
  Block an HTTP response if a line matches a regular expression
  May be used in sections :   defaults | frontend | listen | backend
                                 no    |    yes   |   yes  |   yes
  Arguments :
    <search>  is the regular expression applied to HTTP headers and to the
              response line. This is an extended regular expression, so
              parenthesis grouping is supported and no preliminary backslash
              is required. Any space or known delimiter must be escaped using
              a backslash ('\'). The pattern applies to a full line at a time.
              The "rspdeny" keyword strictly matches case while "rspideny"
              ignores case.

    <cond>    is an optional matching condition built from ACLs. It makes it
              possible to ignore this rule when other conditions are not met.

  A response containing any line which matches extended regular expression
  <search> will mark the request as denied. The test applies both to the
  response line and to response headers. Keep in mind that header names are not
  case-sensitive.

  Main use of this keyword is to prevent sensitive information leak and to
  block the response before it reaches the client. If a response is denied, it
  will be replaced with an HTTP 502 error so that the client never retrieves
  any sensitive data.

  It is easier, faster and more powerful to use ACLs to write access policies.
  Rspdeny should be avoided in new designs.

  Example :
     # Ensure that no content type matching ms-word will leak
     rspideny  ^Content-type:\.*/ms-word

  See also: "reqdeny", "acl", "block", "http-response",  section 6 about
            HTTP header manipulation and section 7 about ACLs.


rsprep  <search> <string> [{if | unless} <cond>]
rspirep <search> <string> [{if | unless} <cond>]  (ignore case)
  Replace a regular expression with a string in an HTTP response line
  May be used in sections :   defaults | frontend | listen | backend
                                 no    |    yes   |   yes  |   yes
  Arguments :
    <search>  is the regular expression applied to HTTP headers and to the
              response line. This is an extended regular expression, so
              parenthesis grouping is supported and no preliminary backslash
              is required. Any space or known delimiter must be escaped using
              a backslash ('\'). The pattern applies to a full line at a time.
              The "rsprep" keyword strictly matches case while "rspirep"
              ignores case.

```

```
      <string>  is the complete line to be added. Any space or known delimiter
                must be escaped using a backslash ('\'). References to matched
                pattern groups are possible using the common \N form, with N
                being a single digit between 0 and 9. Please refer to section
                6 about HTTP header manipulation for more information.

      <cond>    is an optional matching condition built from ACLs. It makes it
                possible to ignore this rule when other conditions are not met.

    Any line matching extended regular expression <search> in the response (both
    the response line and header lines) will be completely replaced with
    <string>. Most common use of this is to rewrite Location headers.

    Header transformations only apply to traffic which passes through HAProxy,
    and not to traffic generated by HAProxy, such as health-checks or error
    responses. Note that for increased readability, it is suggested to add enough
    spaces between the request and the response. Keep in mind that header names
    are not case-sensitive.

    Example :
       # replace "Location: 127.0.0.1:8080" with "Location: www.mydomain.com"
       rspirep ^Location:\ 127.0.0.1:8080    Location:\ www.mydomain.com

    See also: "rspadd", "rspdel", "reqrep", "http-response", section 6 about
              HTTP header manipulation, and section 7 about ACLs.


server <name> <address>[:[port]] [param*]
  Declare a server in a backend
  May be used in sections :   defaults | frontend | listen | backend
                                 no    |    no    |   yes  |   yes
  Arguments :
    <name>    is the internal name assigned to this server. This name will
              appear in logs and alerts.  If "http-send-name-header" is
              set, it will be added to the request header sent to the server.

    <address> is the IPv4 or IPv6 address of the server. Alternatively, a
              resolvable hostname is supported, but this name will be resolved
              during start-up. Address "0.0.0.0" or "*" has a special meaning.
              It indicates that the connection will be forwarded to the same IP
              address as the one from the client connection. This is useful in
              transparent proxy architectures where the client's connection is
              intercepted and haproxy must forward to the original destination
              address. This is more or less what the "transparent" keyword does
              except that with a server it's possible to limit concurrency and
              to report statistics. Optionally, an address family prefix may be
              used before the address to force the family regardless of the
              address format, which can be useful to specify a path to a unix
              socket with no slash ('/'). Currently supported prefixes are :
                    - 'ipv4@'  -> address is always IPv4
                    - 'ipv6@'  -> address is always IPv6
                    - 'unix@'  -> address is a path to a local unix socket
                    - 'abns@'  -> address is in abstract namespace (Linux only)
              You may want to reference some environment variables in the
              address parameter, see section 2.3 about environment
              variables.

    <port>    is an optional port specification. If set, all connections will
              be sent to this port. If unset, the same port the client
              connected to will be used. The port may also be prefixed by a "+"
              or a "-". In this case, the server's port will be determined by
              adding this value to the client's port.

    <param*>  is a list of parameters for this server. The "server" keywords
              accepts an important number of options and has a complete section
```



```
              dedicated to it. Please refer to section 5 for more details.

    Examples :
          server first  10.1.1.1:1080 cookie first  check inter 1000
          server second 10.1.1.2:1080 cookie second check inter 1000
          server transp ipv4@
          server backup "${SRV_BACKUP}:1080" backup
          server www1_dc1 "${LAN_DC1}.101:80"
          server www1_dc2 "${LAN_DC2}.101:80"

    Note: regarding Linux's abstract namespace sockets, HAProxy uses the whole
          sun_path length is used for the address length. Some other programs
          such as socat use the string length only by default. Pass the option
          ",unix-tightsocklen=0" to any abstract socket definition in socat to
          make it compatible with HAProxy's.

    See also: "default-server", "http-send-name-header" and section 5 about
              server options

server-state-file-name [<file>]
  Set the server state file to read, load and apply to servers available in
  this backend. It only applies when the directive "load-server-state-from-file"
  is set to "local". When <file> is not provided or if this directive is not
  set, then backend name is used. If <file> starts with a slash '/', then it is
  considered as an absolute path. Otherwise, <file> is concatenated to the
  global directive "server-state-file-base".

  Example: the minimal configuration below would make HAProxy look for the
           state server file '/etc/haproxy/states/bk':

    global
      server-state-file-base /etc/haproxy/states

   backend bk
      load-server-state-from-file

  See also: "server-state-file-base", "load-server-state-from-file", and
  "show servers state"

source <addr>[:<port>] [usesrc { <addr2>[:<port2>] | client | clientip } ]
source <addr>[:<port>] [usesrc { <addr2>[:<port2>] | hdr_ip(<hdr>[,<occ>]) } ]
source <addr>[:<port>] [interface <name>]
  Set the source address for outgoing connections
  May be used in sections :   defaults | frontend | listen | backend
                                 yes   |    no    |   yes  |   yes
  Arguments :
    <addr>    is the IPv4 address HAProxy will bind to before connecting to a
              server. This address is also used as a source for health checks.

              The default value of 0.0.0.0 means that the system will select
              the most appropriate address to reach its destination. Optionally
              an address family prefix may be used before the address to force
              the family regardless of the address format, which can be useful
              to specify a path to a unix socket with no slash ('/'). Currently
              supported prefixes are :
                - 'ipv4@' -> address is always IPv4
                - 'ipv6@' -> address is always IPv6
                - 'unix@' -> address is a path to a local unix socket
                - 'abns@' -> address is in abstract namespace (Linux only)
              You may want to reference some environment variables in the
              address parameter, see section 2.3 about environment variables.

    <port>    is an optional port. It is normally not needed but may be useful
              in some very specific contexts. The default value of zero means
              the system will select a free port. Note that port ranges are not
```

supported in the backend. If you want to force port ranges, you
have to specify them on each "server" line.

<addr2> is the IP address to present to the server when connections are
forwarded in full transparent proxy mode. This is currently only
supported on some patched Linux kernels. When this address is
specified, clients connecting to the server will be presented
with this address, while health checks will still use the address
<addr>.

<port2> is the optional port to present to the server when connections
are forwarded in full transparent proxy mode (see <addr2> above).
The default value of zero means the system will select a free
port.

<hdr> is the name of a HTTP header in which to fetch the IP to bind to.
This is the name of a comma-separated header list which can
contain multiple IP addresses. By default, the last occurrence is
used. This is designed to work with the X-Forwarded-For header
and to automatically bind to the client's IP address as seen
by previous proxy, typically Stunnel. In order to use another
occurrence from the last one, please see the <occ> parameter
below. When the header (or occurrence) is not found, no binding
is performed so that the proxy's default IP address is used. Also
keep in mind that the header name is case insensitive, as for any
HTTP header.

<occ> is the occurrence number of a value to be used in a multi-value
header. This is to be used in conjunction with "hdr_ip(<hdr>)",
in order to specify which occurrence to use for the source IP
address. Positive values indicate a position from the first
occurrence, 1 being the first one. Negative values indicate
positions relative to the last one, -1 being the last one. This
is helpful for situations where an X-Forwarded-For header is set
at the entry point of an infrastructure and must be used several
proxy layers away. When this value is not specified, -1 is
assumed. Passing a zero here disables the feature.

<name> is an optional interface name to which to bind to for outgoing
traffic. On systems supporting this features (currently, only
Linux), this allows one to bind all traffic to the server to
this interface even if it is not the one the system would select
based on routing tables. This should be used with extreme care.
Note that using this option requires root privileges.

The "source" keyword is useful in complex environments where a specific
address only is allowed to connect to the servers. It may be needed when a
private address must be used through a public gateway for instance, and it is
known that the system cannot determine the adequate source address by itself.

An extension which is available on certain patched Linux kernels may be used
through the "usesrc" optional keyword. It makes it possible to connect to the
servers with an IP address which does not belong to the system itself. This
is called "full transparent proxy mode". For this to work, the destination
servers have to route their traffic back to this address through the machine
running HAProxy, and IP forwarding must generally be enabled on this machine.

In this "full transparent proxy" mode, it is possible to force a specific IP
address to be presented to the servers. This is not much used in fact. A more
common use is to tell HAProxy to present the client's IP address. For this,
there are two methods :

- present the client's IP and port addresses. This is the most transparent
mode, but it can cause problems when IP connection tracking is enabled on
the machine, because a same connection may be seen twice with different

states. However, this solution presents the huge advantage of not
limiting the system to the 64k outgoing address+port couples, because all
of the client ranges may be used.

- present only the client's IP address and select a spare port. This
solution is still quite elegant but slightly less transparent (downstream
firewalls logs will not match upstream's). It also presents the downside
of limiting the number of concurrent connections to the usual 64k ports.
However, since the upstream and downstream ports are different, local IP
connection tracking on the machine will not be upset by the reuse of the
same session.

This option sets the default source for all servers in the backend. It may
also be specified in a "defaults" section. Finer source address specification
is possible at the server level using the "source" server option. Refer to
section 5 for more information.

In order to work, "usesrc" requires root privileges.

Examples :

```
backend private
    # Connect to the servers using our 192.168.1.200 source address
    source 192.168.1.200

backend transparent_ssl1
    # Connect to the SSL farm from the client's source address
    source 192.168.1.200 usesrc clientip

backend transparent_ssl2
    # Connect to the SSL farm from the client's source address and port
    # not recommended if IP conntrack is present on the local machine.
    source 192.168.1.200 usesrc client

backend transparent_ssl3
    # Connect to the SSL farm from the client's source address. It
    # is more conntrack-friendly.
    source 192.168.1.200 usesrc clientip

backend transparent_smtp
    # Connect to the SMTP farm from the client's source address/port
    # with Tproxy version 4.
    source 0.0.0.0 usesrc clientip

backend transparent_http
    # Connect to the servers using the client's IP as seen by previous
    # proxy.
    source 0.0.0.0 usesrc hdr_ip(x-forwarded-for,-1)
```

See also : the "source" server option in section 5, the Tproxy patches for
the Linux kernel on www.balabit.com, the "bind" keyword.

srvtimeout <timeout> (deprecated)
Set the maximum inactivity time on the server side.
May be used in sections : defaults | frontend | listen | backend
yes | no | yes | yes
Arguments :
<timeout> is the timeout value specified in milliseconds by default, but
can be in any other unit if the number is suffixed by the unit,
as explained at the top of this document.

The inactivity timeout applies when the server is expected to acknowledge or
send data. In HTTP mode, this timeout is particularly important to consider
during the first phase of the server's response, when it has to send the
headers, as it directly represents the server's processing time for the

```
   request. To find out what value to put there, it's often good to start with
   what would be considered as unacceptable response times, then check the logs
   to observe the response time distribution, and adjust the value accordingly.

   The value is specified in milliseconds by default, but can be in any other
   unit if the number is suffixed by the unit, as specified at the top of this
   document. In TCP mode (and to a lesser extent, in HTTP mode), it is highly
   recommended that the client timeout remains equal to the server timeout in
   order to avoid complex situations to debug. Whatever the expected server
   response times, it is a good practice to cover at least one or several TCP
   packet losses by specifying timeouts that are slightly above multiples of 3
   seconds (eg: 4 or 5 seconds minimum).

   This parameter is specific to backends, but can be specified once for all in
   "defaults" sections. This is in fact one of the easiest solutions not to
   forget about it. An unspecified timeout results in an infinite timeout, which
   is not recommended. Such a usage is accepted and works but reports a warning
   during startup because it may results in accumulation of expired sessions in
   the system if the system's timeouts are not configured either.

   This parameter is provided for compatibility but is currently deprecated.
   Please use "timeout server" instead.

   See also : "timeout server", "timeout tunnel", "timeout client" and
              "clitimeout".


stats admin { if | unless } <cond>
  Enable statistics admin level if/unless a condition is matched
  May be used in sections :   defaults | frontend | listen | backend
                                 no    |    yes   |   yes  |   yes

   This statement enables the statistics admin level if/unless a condition is
   matched.

   The admin level allows to enable/disable servers from the web interface. By
   default, statistics page is read-only for security reasons.

   Note : Consider not using this feature in multi-process mode (nbproc > 1)
          unless you know what you do : memory is not shared between the
          processes, which can result in random behaviours.

   Currently, the POST request is limited to the buffer size minus the reserved
   buffer space, which means that if the list of servers is too long, the
   request won't be processed. It is recommended to alter few servers at a
   time.

   Example :
     # statistics admin level only for localhost
     backend stats_localhost
         stats enable
         stats admin if LOCALHOST

   Example :
     # statistics admin level always enabled because of the authentication
     backend stats_auth
         stats enable
         stats auth  admin:AdMiN123
         stats admin if TRUE

   Example :
     # statistics admin level depends on the authenticated user
     userlist stats-auth
         group admin    users admin
         user  admin    insecure-password AdMiN123
```



```
         group readonly users haproxy
         user  haproxy  insecure-password haproxy

     backend stats_auth
         stats enable
         acl AUTH       http_auth(stats-auth)
         acl AUTH_ADMIN http_auth_group(stats-auth) admin
         stats http-request auth unless AUTH
         stats admin if AUTH_ADMIN

  See also : "stats enable", "stats auth", "stats http-request", "nbproc",
             "bind-process", section 3.4 about userlists and section 7 about
             ACL usage.


stats auth <user>:<passwd>
  Enable statistics with authentication and grant access to an account
  May be used in sections :   defaults | frontend | listen | backend
                                 yes   |    yes   |   yes  |   yes
  Arguments :
    <user>    is a user name to grant access to

    <passwd>  is the cleartext password associated to this user

  This statement enables statistics with default settings, and restricts access
  to declared users only. It may be repeated as many times as necessary to
  allow as many users as desired. When a user tries to access the statistics
  without a valid account, a "401 Forbidden" response will be returned so that
  the browser asks the user to provide a valid user and password. The real
  which will be returned to the browser is configurable using "stats realm".

  Since the authentication method is HTTP Basic Authentication, the passwords
  circulate in cleartext on the network. Thus, it was decided that the
  configuration file would also use cleartext passwords to remind the users
  that those ones should not be sensitive and not shared with any other account.

  It is also possible to reduce the scope of the proxies which appear in the
  report using "stats scope".

  Though this statement alone is enough to enable statistics reporting, it is
  recommended to set all other settings in order to avoid relying on default
  unobvious parameters.

  Example :
    # public access (limited to this backend only)
    backend public_www
        server srv1 192.168.0.1:80
        stats enable
        stats hide-version
        stats scope   .
        stats uri     /admin?stats
        stats realm   Haproxy\ Statistics
        stats auth    admin1:AdMiN123
        stats auth    admin2:AdMiN321

    # internal monitoring access (unlimited)
    backend private_monitoring
        stats enable
        stats uri     /admin?stats
        stats refresh 5s

  See also : "stats enable", "stats realm", "stats scope", "stats uri"


stats enable
```

```
   Enable statistics reporting with default settings
   May be used in sections :   defaults | frontend | listen | backend
                                   yes   |    yes   |   yes  |   yes
   Arguments : none

   This statement enables statistics reporting with default settings defined
   at build time. Unless stated otherwise, these settings are used :
     - stats uri   : /haproxy?stats
     - stats realm : "HAProxy Statistics"
     - stats auth  : no authentication
     - stats scope : no restriction

   Though this statement alone is enough to enable statistics reporting, it is
   recommended to set all other settings in order to avoid relying on default
   unobvious parameters.

   Example :
     # public access (limited to this backend only)
     backend public_www
         server srv1 192.168.0.1:80
         stats enable
         stats hide-version
         stats scope   .
         stats uri     /admin?stats
         stats realm   Haproxy\ Statistics
         stats auth    admin1:AdMiN123
         stats auth    admin2:AdMiN321

     # internal monitoring access (unlimited)
     backend private_monitoring
         stats enable
         stats uri     /admin?stats
         stats refresh 5s

   See also : "stats auth", "stats realm", "stats uri"


stats hide-version
   Enable statistics and hide HAProxy version reporting
   May be used in sections :   defaults | frontend | listen | backend
                                   yes   |    yes   |   yes  |   yes
   Arguments : none

   By default, the stats page reports some useful status information along with
   the statistics. Among them is HAProxy's version. However, it is generally
   considered dangerous to report precise version to anyone, as it can help them
   target known weaknesses with specific attacks. The "stats hide-version"
   statement removes the version from the statistics report. This is recommended
   for public sites or any site with a weak login/password.

   Though this statement alone is enough to enable statistics reporting, it is
   recommended to set all other settings in order to avoid relying on default
   unobvious parameters.

   Example :
     # public access (limited to this backend only)
     backend public_www
         server srv1 192.168.0.1:80
         stats enable
         stats hide-version
         stats scope   .
         stats uri     /admin?stats
         stats realm   Haproxy\ Statistics
         stats auth    admin1:AdMiN123
         stats auth    admin2:AdMiN321
```



```

     # internal monitoring access (unlimited)
     backend private_monitoring
         stats enable
         stats uri     /admin?stats
         stats refresh 5s

   See also : "stats auth", "stats enable", "stats realm", "stats uri"


stats http-request { allow | deny | auth [realm <realm>] }
            [ { if | unless } <condition> ]
   Access control for statistics

   May be used in sections:   defaults | frontend | listen | backend
                                  no    |    no    |   yes  |   yes

   As "http-request", these set of options allow to fine control access to
   statistics. Each option may be followed by if/unless and acl.
   First option with matched condition (or option without condition) is final.
   For "deny" a 403 error will be returned, for "allow" normal processing is
   performed, for "auth" a 401/407 error code is returned so the client
   should be asked to enter a username and password.

   There is no fixed limit to the number of http-request statements per
   instance.

   See also : "http-request", section 3.4 about userlists and section 7
              about ACL usage.


stats realm <realm>
   Enable statistics and set authentication realm
   May be used in sections :   defaults | frontend | listen | backend
                                   yes   |    yes   |   yes  |   yes
   Arguments :
     <realm>   is the name of the HTTP Basic Authentication realm reported to
               the browser. The browser uses it to display it in the pop-up
               inviting the user to enter a valid username and password.

   The realm is read as a single word, so any spaces in it should be escaped
   using a backslash ('\').

   This statement is useful only in conjunction with "stats auth" since it is
   only related to authentication.

   Though this statement alone is enough to enable statistics reporting, it is
   recommended to set all other settings in order to avoid relying on default
   unobvious parameters.

   Example :
     # public access (limited to this backend only)
     backend public_www
         server srv1 192.168.0.1:80
         stats enable
         stats hide-version
         stats scope   .
         stats uri     /admin?stats
         stats realm   Haproxy\ Statistics
         stats auth    admin1:AdMiN123
         stats auth    admin2:AdMiN321

     # internal monitoring access (unlimited)
     backend private_monitoring
         stats enable
```

```
7541          stats uri     /admin?stats
7542          stats refresh 5s
7543
7544    See also : "stats auth", "stats enable", "stats uri"
7545
7546
7547  stats refresh <delay>
7548    Enable statistics with automatic refresh
7549    May be used in sections :   defaults | frontend | listen | backend
7550                                  yes   |    yes   |   yes  |   yes
7551    Arguments :
7552      <delay>   is the suggested refresh delay, specified in seconds, which will
7553                be returned to the browser consulting the report page. While the
7554                browser is free to apply any delay, it will generally respect it
7555                and refresh the page this every seconds. The refresh interval may
7556                be specified in any other non-default time unit, by suffixing the
7557                unit after the value, as explained at the top of this document.
7558
7559    This statement is useful on monitoring displays with a permanent page
7560    reporting the load balancer's activity. When set, the HTML report page will
7561    include a link "refresh"/"stop refresh" so that the user can select whether
7562    he wants automatic refresh of the page or not.
7563
7564    Though this statement alone is enough to enable statistics reporting, it is
7565    recommended to set all other settings in order to avoid relying on default
7566    unobvious parameters.
7567
7568    Example :
7569      # public access (limited to this backend only)
7570      backend public_www
7571          server srv1 192.168.0.1:80
7572          stats enable
7573          stats hide-version
7574          stats scope   .
7575          stats uri     /admin?stats
7576          stats realm   Haproxy\ Statistics
7577          stats auth    admin1:AdMiN123
7578          stats auth    admin2:AdMiN321
7579
7580      # internal monitoring access (unlimited)
7581      backend private_monitoring
7582          stats enable
7583          stats uri     /admin?stats
7584          stats refresh 5s
7585
7586    See also : "stats auth", "stats enable", "stats realm", "stats uri"
7587
7588
7589  stats scope { <name> | "." }
7590    Enable statistics and limit access scope
7591    May be used in sections :   defaults | frontend | listen | backend
7592                                  yes   |    yes   |   yes  |   yes
7593    Arguments :
7594      <name>    is the name of a listen, frontend or backend section to be
7595                reported. The special name "." (a single dot) designates the
7596                section in which the statement appears.
7597
7598    When this statement is specified, only the sections enumerated with this
7599    statement will appear in the report. All other ones will be hidden. This
7600    statement may appear as many times as needed if multiple sections need to be
7601    reported. Please note that the name checking is performed as simple string
7602    comparisons, and that it is never checked that a give section name really
7603    exists.
7604
7605    Though this statement alone is enough to enable statistics reporting, it is
```



```
7606    recommended to set all other settings in order to avoid relying on default
7607    unobvious parameters.
7608
7609    Example :
7610      # public access (limited to this backend only)
7611      backend public_www
7612          server srv1 192.168.0.1:80
7613          stats enable
7614          stats hide-version
7615          stats scope   .
7616          stats uri     /admin?stats
7617          stats realm   Haproxy\ Statistics
7618          stats auth    admin1:AdMiN123
7619          stats auth    admin2:AdMiN321
7620
7621      # internal monitoring access (unlimited)
7622      backend private_monitoring
7623          stats enable
7624          stats uri     /admin?stats
7625          stats refresh 5s
7626
7627    See also : "stats auth", "stats enable", "stats realm", "stats uri"
7628
7629
7630  stats show-desc [ <desc> ]
7631    Enable reporting of a description on the statistics page.
7632    May be used in sections :   defaults | frontend | listen | backend
7633                                  yes   |    yes   |   yes  |   yes
7634
7635      <desc>    is an optional description to be reported. If unspecified, the
7636                description from global section is automatically used instead.
7637
7638    This statement is useful for users that offer shared services to their
7639    customers, where node or description should be different for each customer.
7640
7641    Though this statement alone is enough to enable statistics reporting, it is
7642    recommended to set all other settings in order to avoid relying on default
7643    unobvious parameters.  By default description is not shown.
7644
7645    Example :
7646      # internal monitoring access (unlimited)
7647      backend private_monitoring
7648          stats enable
7649          stats show-desc Master node for Europe, Asia, Africa
7650          stats uri       /admin?stats
7651          stats refresh   5s
7652
7653    See also: "show-node", "stats enable", "stats uri" and "description" in
7654              global section.
7655
7656
7657  stats show-legends
7658    Enable reporting additional information on the statistics page
7659    May be used in sections :   defaults | frontend | listen | backend
7660                                  yes   |    yes   |   yes  |   yes
7661    Arguments : none
7662
7663    Enable reporting additional information on the statistics page :
7664      - cap: capabilities (proxy)
7665      - mode: one of tcp, http or health (proxy)
7666      - id: SNMP ID (proxy, socket, server)
7667      - IP (socket, server)
7668      - cookie (backend, server)
7669
7670    Though this statement alone is enough to enable statistics reporting, it is
```

```
   recommended to set all other settings in order to avoid relying on default
   unobvious parameters.  Default behaviour is not to show this information.

   See also: "stats enable", "stats uri".


stats show-node [ <name> ]
   Enable reporting of a host name on the statistics page.
   May be used in sections :   defaults | frontend | listen | backend
                                   yes   |    yes   |   yes  |   yes
   Arguments:
     <name>    is an optional name to be reported. If unspecified, the
               node name from global section is automatically used instead.

   This statement is useful for users that offer shared services to their
   customers, where node or description might be different on a stats page
   provided for each customer.  Default behaviour is not to show host name.

   Though this statement alone is enough to enable statistics reporting, it is
   recommended to set all other settings in order to avoid relying on default
   unobvious parameters.

   Example:
     # internal monitoring access (unlimited)
     backend private_monitoring
         stats enable
         stats show-node Europe-1
         stats uri       /admin?stats
         stats refresh   5s

   See also: "show-desc", "stats enable", "stats uri", and "node" in global
             section.


stats uri <prefix>
   Enable statistics and define the URI prefix to access them
   May be used in sections :   defaults | frontend | listen | backend
                                   yes   |    yes   |   yes  |   yes
   Arguments :
     <prefix>  is the prefix of any URI which will be redirected to stats. This
               prefix may contain a question mark ('?') to indicate part of a
               query string.

   The statistics URI is intercepted on the relayed traffic, so it appears as a
   page within the normal application. It is strongly advised to ensure that the
   selected URI will never appear in the application, otherwise it will never be
   possible to reach it in the application.

   The default URI compiled in haproxy is "/haproxy?stats", but this may be
   changed at build time, so it's better to always explicitly specify it here.
   It is generally a good idea to include a question mark in the URI so that
   intermediate proxies refrain from caching the results. Also, since any string
   beginning with the prefix will be accepted as a stats request, the question
   mark helps ensuring that no valid URI will begin with the same words.

   It is sometimes very convenient to use "/" as the URI prefix, and put that
   statement in a "listen" instance of its own. That makes it easy to dedicate
   an address or a port to statistics only.

   Though this statement alone is enough to enable statistics reporting, it is
   recommended to set all other settings in order to avoid relying on default
   unobvious parameters.

   Example :
     # public access (limited to this backend only)
```



```
     backend public_www
         server srv1 192.168.0.1:80
         stats enable
         stats hide-version
         stats scope   .
         stats uri     /admin?stats
         stats realm   Haproxy\ Statistics
         stats auth    admin1:AdMiN123
         stats auth    admin2:AdMiN321

     # internal monitoring access (unlimited)
     backend private_monitoring
         stats enable
         stats uri     /admin?stats
         stats refresh 5s

   See also : "stats auth", "stats enable", "stats realm"


stick match <pattern> [table <table>] [{if | unless} <cond>]
   Define a request pattern matching condition to stick a user to a server
   May be used in sections :   defaults | frontend | listen | backend
                                   no    |    no    |   yes  |   yes

   Arguments :
     <pattern>  is a sample expression rule as described in section 7.3. It
                describes what elements of the incoming request or connection
                will be analysed in the hope to find a matching entry in a
                stickiness table. This rule is mandatory.

     <table>    is an optional stickiness table name. If unspecified, the same
                backend's table is used. A stickiness table is declared using
                the "stick-table" statement.

     <cond>     is an optional matching condition. It makes it possible to match
                on a certain criterion only when other conditions are met (or
                not met). For instance, it could be used to match on a source IP
                address except when a request passes through a known proxy, in
                which case we'd match on a header containing that IP address.

   Some protocols or applications require complex stickiness rules and cannot
   always simply rely on cookies nor hashing. The "stick match" statement
   describes a rule to extract the stickiness criterion from an incoming request
   or connection. See section 7 for a complete list of possible patterns and
   transformation rules.

   The table has to be declared using the "stick-table" statement. It must be of
   a type compatible with the pattern. By default it is the one which is present
   in the same backend. It is possible to share a table with other backends by
   referencing it using the "table" keyword. If another table is referenced,
   the server's ID inside the backends are used. By default, all server IDs
   start at 1 in each backend, so the server ordering is enough. But in case of
   doubt, it is highly recommended to force server IDs using their "id" setting.

   It is possible to restrict the conditions where a "stick match" statement
   will apply, using "if" or "unless" followed by a condition. See section 7 for
   ACL based conditions.

   There is no limit on the number of "stick match" statements. The first that
   applies and matches will cause the request to be directed to the same server
   as was used for the request which created the entry. That way, multiple
   matches can be used as fallbacks.

   The stick rules are checked after the persistence cookies, so they will not
   affect stickiness if a cookie has already been used to select a server. That
```

```
   way, it becomes very easy to insert cookies and match on IP addresses in
   order to maintain stickiness between HTTP and HTTPS.

   Note : Consider not using this feature in multi-process mode (nbproc > 1)
          unless you know what you do : memory is not shared between the
          processes, which can result in random behaviours.

   Example :
     # forward SMTP users to the same server they just used for POP in the
     # last 30 minutes
     backend pop
         mode tcp
         balance roundrobin
         stick store-request src
         stick-table type ip size 200k expire 30m
         server s1 192.168.1.1:110
         server s2 192.168.1.1:110

     backend smtp
         mode tcp
         balance roundrobin
         stick match src table pop
         server s1 192.168.1.1:25
         server s2 192.168.1.1:25

   See also : "stick-table", "stick on", "nbproc", "bind-process" and section 7
              about ACLs and samples fetching.


stick on <pattern> [table <table>] [{if | unless} <condition>]
  Define a request pattern to associate a user to a server
  May be used in sections :   defaults | frontend | listen | backend
                                 no    |    no    |   yes  |   yes

   Note : This form is exactly equivalent to "stick match" followed by
          "stick store-request", all with the same arguments. Please refer
          to both keywords for details. It is only provided as a convenience
          for writing more maintainable configurations.

   Note : Consider not using this feature in multi-process mode (nbproc > 1)
          unless you know what you do : memory is not shared between the
          processes, which can result in random behaviours.

   Examples :
     # The following form ...
     stick on src table pop if !localhost

     # ...is strictly equivalent to this one :
     stick match src table pop if !localhost
     stick store-request src table pop if !localhost


     # Use cookie persistence for HTTP, and stick on source address for HTTPS as
     # well as HTTP without cookie. Share the same table between both accesses.
     backend http
         mode http
         balance roundrobin
         stick on src table https
         cookie SRV insert indirect nocache
         server s1 192.168.1.1:80 cookie s1
         server s2 192.168.1.1:80 cookie s2

     backend https
         mode tcp
         balance roundrobin
```



```
         stick-table type ip size 200k expire 30m
         stick on src
         server s1 192.168.1.1:443
         server s2 192.168.1.1:443

   See also : "stick match", "stick store-request", "nbproc" and "bind-process".


stick store-request <pattern> [table <table>] [{if | unless} <condition>]
  Define a request pattern used to create an entry in a stickiness table
  May be used in sections :   defaults | frontend | listen | backend
                                 no    |    no    |   yes  |   yes

  Arguments :
    <pattern>  is a sample expression rule as described in section 7.3. It
               describes what elements of the incoming request or connection
               will be analysed, extracted and stored in the table once a
               server is selected.

    <table>    is an optional stickiness table name. If unspecified, the same
               backend's table is used. A stickiness table is declared using
               the "stick-table" statement.

    <cond>     is an optional storage condition. It makes it possible to store
               certain criteria only when some conditions are met (or not met).
               For instance, it could be used to store the source IP address
               except when the request passes through a known proxy, in which
               case we'd store a converted form of a header containing that IP
               address.

  Some protocols or applications require complex stickiness rules and cannot
  always simply rely on cookies nor hashing. The "stick store-request" statement
  describes a rule to decide what to extract from the request and when to do
  it, in order to store it into a stickiness table for further requests to
  match it using the "stick match" statement. Obviously the extracted part must
  make sense and have a chance to be matched in a further request. Storing a
  client's IP address for instance often makes sense. Storing an ID found in a
  URL parameter also makes sense. Storing a source port will almost never make
  any sense because it will be randomly matched. See section 7 for a complete
  list of possible patterns and transformation rules.

  The table has to be declared using the "stick-table" statement. It must be of
  a type compatible with the pattern. By default it is the one which is present
  in the same backend. It is possible to share a table with other backends by
  referencing it using the "table" keyword. If another table is referenced,
  the server's ID inside the backends are used. By default, all server IDs
  start at 1 in each backend, so the server ordering is enough. But in case of
  doubt, it is highly recommended to force server IDs using their "id" setting.

  It is possible to restrict the conditions where a "stick store-request"
  statement will apply, using "if" or "unless" followed by a condition. This
  condition will be evaluated while parsing the request, so any criteria can be
  used. See section 7 for ACL based conditions.

  There is no limit on the number of "stick store-request" statements, but
  there is a limit of 8 simultaneous stores per request or response. This
  makes it possible to store up to 8 criteria, all extracted from either the
  request or the response, regardless of the number of rules. Only the 8 first
  ones which match will be kept. Using this, it is possible to feed multiple
  tables at once in the hope to increase the chance to recognize a user on
  another protocol or access method. Using multiple store-request rules with
  the same table is possible and may be used to find the best criterion to rely
  on, by arranging the rules by decreasing preference order. Only the first
  extracted criterion for a given table will be stored. All subsequent store-
  request rules referencing the same table will be skipped and their ACLs will
```

```
   not be evaluated.

   The "store-request" rules are evaluated once the server connection has been
   established, so that the table will contain the real server that processed
   the request.

   Note : Consider not using this feature in multi-process mode (nbproc > 1)
          unless you know what you do : memory is not shared between the
          processes, which can result in random behaviours.

   Example :
     # forward SMTP users to the same server they just used for POP in the
     # last 30 minutes
     backend pop
         mode tcp
         balance roundrobin
         stick store-request src
         stick-table type ip size 200k expire 30m
         server s1 192.168.1.1:110
         server s2 192.168.1.1:110

     backend smtp
         mode tcp
         balance roundrobin
         stick match src table pop
         server s1 192.168.1.1:25
         server s2 192.168.1.1:25

   See also : "stick-table", "stick on", "nbproc", "bind-process" and section 7
              about ACLs and sample fetching.


stick-table type {ip | integer | string [len <length>] | binary [len <length>]}
            size <size> [expire <expire>] [nopurge] [peers <peersect>]
            [store <data_type>]*
   Configure the stickiness table for the current section
   May be used in sections :   defaults | frontend | listen | backend
                                  no    |    yes   |   yes  |   yes

   Arguments :
     ip         a table declared with "type ip" will only store IPv4 addresses.
                This form is very compact (about 50 bytes per entry) and allows
                very fast entry lookup and stores with almost no overhead. This
                is mainly used to store client source IP addresses.

     ipv6       a table declared with "type ipv6" will only store IPv6 addresses.
                This form is very compact (about 60 bytes per entry) and allows
                very fast entry lookup and stores with almost no overhead. This
                is mainly used to store client source IP addresses.

     integer    a table declared with "type integer" will store 32bit integers
                which can represent a client identifier found in a request for
                instance.

     string     a table declared with "type string" will store substrings of up
                to <len> characters. If the string provided by the pattern
                extractor is larger than <len>, it will be truncated before
                being stored. During matching, at most <len> characters will be
                compared between the string in the table and the extracted
                pattern. When not specified, the string is automatically limited
                to 32 characters.

     binary     a table declared with "type binary" will store binary blocks
                of <len> bytes. If the block provided by the pattern
                extractor is larger than <len>, it will be truncated before
```



```
                being stored. If the block provided by the sample expression
                is shorter than <len>, it will be padded by 0. When not
                specified, the block is automatically limited to 32 bytes.

     <length>   is the maximum number of characters that will be stored in a
                "string" type table (See type "string" above). Or the number
                of bytes of the block in "binary" type table. Be careful when
                changing this parameter as memory usage will proportionally
                increase.

     <size>     is the maximum number of entries that can fit in the table. This
                value directly impacts memory usage. Count approximately
                50 bytes per entry, plus the size of a string if any. The size
                supports suffixes "k", "m", "g" for 2^10, 2^20 and 2^30 factors.

     [nopurge]  indicates that we refuse to purge older entries when the table
                is full. When not specified and the table is full when haproxy
                wants to store an entry in it, it will flush a few of the oldest
                entries in order to release some space for the new ones. This is
                most often the desired behaviour. In some specific cases, it
                be desirable to refuse new entries instead of purging the older
                ones. That may be the case when the amount of data to store is
                far above the hardware limits and we prefer not to offer access
                to new clients than to reject the ones already connected. When
                using this parameter, be sure to properly set the "expire"
                parameter (see below).

     <peersect> is the name of the peers section to use for replication. Entries
                which associate keys to server IDs are kept synchronized with
                the remote peers declared in this section. All entries are also
                automatically learned from the local peer (old process) during a
                soft restart.

                NOTE : each peers section may be referenced only by tables
                       belonging to the same unique process.

     <expire>   defines the maximum duration of an entry in the table since it
                was last created, refreshed or matched. The expiration delay is
                defined using the standard time format, similarly as the various
                timeouts. The maximum duration is slightly above 24 days. See
                section 2.2 for more information. If this delay is not specified,
                the session won't automatically expire, but older entries will
                be removed once full. Be sure not to use the "nopurge" parameter
                if not expiration delay is specified.

    <data_type> is used to store additional information in the stick-table. This
                may be used by ACLs in order to control various criteria related
                to the activity of the client matching the stick-table. For each
                item specified here, the size of each entry will be inflated so
                that the additional data can fit. Several data types may be
                stored with an entry. Multiple data types may be specified after
                the "store" keyword, as a comma-separated list. Alternatively,
                it is possible to repeat the "store" keyword followed by one or
                several data types. Except for the "server_id" type which is
                automatically detected and enabled, all data types must be
                explicitly declared to be stored. If an ACL references a data
                type which is not stored, the ACL will simply not match. Some
                data types require an argument which must be passed just after
                the type between parenthesis. See below for the supported data
                types and their arguments.

   The data types that can be stored with an entry are the following :
     - server_id : this is an integer which holds the numeric ID of the server a
       request was assigned to. It is used by the "stick match", "stick store",
       and "stick on" rules. It is automatically enabled when referenced.
```

- gpc0 : first General Purpose Counter. It is a positive 32-bit integer integer which may be used for anything. Most of the time it will be used to put a special tag on some entries, for instance to note that a specific behaviour was detected and must be known for future matches.

- gpc0_rate(<period>) : increment rate of the first General Purpose Counter over a period. It is a positive 32-bit integer integer which may be used for anything. Just like <gpc0>, it counts events, but instead of keeping a cumulative count, it maintains the rate at which the counter is incremented. Most of the time it will be used to measure the frequency of occurrence of certain events (eg: requests to a specific URL).

- conn_cnt : Connection Count. It is a positive 32-bit integer which counts the absolute number of connections received from clients which matched this entry. It does not mean the connections were accepted, just that they were received.

- conn_cur : Current Connections. It is a positive 32-bit integer which stores the concurrent connection counts for the entry. It is incremented once an incoming connection matches the entry, and decremented once the connection leaves. That way it is possible to know at any time the exact number of concurrent connections for an entry.

- conn_rate(<period>) : frequency counter (takes 12 bytes). It takes an integer parameter <period> which indicates in milliseconds the length of the period over which the average is measured. It reports the average incoming connection rate over that period, in connections per period. The result is an integer which can be matched using ACLs.

- sess_cnt : Session Count. It is a positive 32-bit integer which counts the absolute number of sessions received from clients which matched this entry. A session is a connection that was accepted by the layer 4 rules.

- sess_rate(<period>) : frequency counter (takes 12 bytes). It takes an integer parameter <period> which indicates in milliseconds the length of the period over which the average is measured. It reports the average incoming session rate over that period, in sessions per period. The result is an integer which can be matched using ACLs.

- http_req_cnt : HTTP request Count. It is a positive 32-bit integer which counts the absolute number of HTTP requests received from clients which matched this entry. It does not matter whether they are valid requests or not. Note that this is different from sessions when keep-alive is used on the client side.

- http_req_rate(<period>) : frequency counter (takes 12 bytes). It takes an integer parameter <period> which indicates in milliseconds the length of the period over which the average is measured. It reports the average HTTP request rate over that period, in requests per period. The result is an integer which can be matched using ACLs. It does not matter whether they are valid requests or not. Note that this is different from sessions when keep-alive is used on the client side.

- http_err_cnt : HTTP Error Count. It is a positive 32-bit integer which counts the absolute number of HTTP requests errors induced by clients which matched this entry. Errors are counted on invalid and truncated requests, as well as on denied or tarpitted requests, and on failed authentications. If the server responds with 4xx, then the request is also counted as an error since it's an error triggered by the client (eg: vulnerability scan).

- http_err_rate(<period>) : frequency counter (takes 12 bytes). It takes an integer parameter <period> which indicates in milliseconds the length of the period over which the average is measured. It reports the average

HTTP request error rate over that period, in requests per period (see http_err_cnt above for what is accounted as an error). The result is an integer which can be matched using ACLs.

- bytes_in_cnt : client to server byte count. It is a positive 64-bit integer which counts the cumulated amount of bytes received from clients which matched this entry. Headers are included in the count. This may be used to limit abuse of upload features on photo or video servers.

- bytes_in_rate(<period>) : frequency counter (takes 12 bytes). It takes an integer parameter <period> which indicates in milliseconds the length of the period over which the average is measured. It reports the average incoming bytes rate over that period, in bytes per period. It may be used to detect users which upload too much and too fast. Warning: with large uploads, it is possible that the amount of uploaded data will be counted once upon termination, thus causing spikes in the average transfer speed instead of having a smooth one. This may partially be smoothed with "option contstats" though this is not perfect yet. Use of byte_in_cnt is recommended for better fairness.

- bytes_out_cnt : server to client byte count. It is a positive 64-bit integer which counts the cumulated amount of bytes sent to clients which matched this entry. Headers are included in the count. This may be used to limit abuse of bots sucking the whole site.

- bytes_out_rate(<period>) : frequency counter (takes 12 bytes). It takes an integer parameter <period> which indicates in milliseconds the length of the period over which the average is measured. It reports the average outgoing bytes rate over that period, in bytes per period. It may be used to detect users which download too much and too fast. Warning: with large transfers, it is possible that the amount of transferred data will be counted once upon termination, thus causing spikes in the average transfer speed instead of having a smooth one. This may partially be smoothed with "option contstats" though this is not perfect yet. Use of byte_out_cnt is recommended for better fairness.

There is only one stick-table per proxy. At the moment of writing this doc, it does not seem useful to have multiple tables per proxy. If this happens to be required, simply create a dummy backend with a stick-table in it and reference it.

It is important to understand that stickiness based on learning information has some limitations, including the fact that all learned associations are lost upon restart. In general it can be good as a complement but not always as an exclusive stickiness.

Last, memory requirements may be important when storing many data types. Indeed, storing all indicators above at once in each entry requires 116 bytes per entry, or 116 MB for a 1-million entries table. This is definitely not something that can be ignored.

Example:
```
      # Keep track of counters of up to 1 million IP addresses over 5 minutes
      # and store a general purpose counter and the average connection rate
      # computed over a sliding window of 30 seconds.
      stick-table type ip size 1m expire 5m store gpc0,conn_rate(30s)
```

See also : "stick match", "stick on", "stick store-request", section 2.2 about time format and section 7 about ACLs.

stick store-response <pattern> [table <table>] [{if | unless} <condition>]
Define a request pattern used to create an entry in a stickiness table
May be used in sections : defaults | frontend | listen | backend
no | no | yes | yes

```
8191
8192    Arguments :
8193      <pattern>  is a sample expression rule as described in section 7.3. It
8194                 describes what elements of the response or connection will
8195                 be analysed, extracted and stored in the table once a
8196                 server is selected.
8197
8198      <table>    is an optional stickiness table name. If unspecified, the same
8199                 backend's table is used. A stickiness table is declared using
8200                 the "stick-table" statement.
8201
8202      <cond>     is an optional storage condition. It makes it possible to store
8203                 certain criteria only when some conditions are met (or not met).
8204                 For instance, it could be used to store the SSL session ID only
8205                 when the response is a SSL server hello.
8206
8207    Some protocols or applications require complex stickiness rules and cannot
8208    always simply rely on cookies nor hashing. The "stick store-response"
8209    statement  describes a rule to decide what to extract from the response and
8210    when to do it, in order to store it into a stickiness table for further
8211    requests to match it using the "stick match" statement. Obviously the
8212    extracted part must make sense and have a chance to be matched in a further
8213    request. Storing an ID found in a header of a response makes sense.
8214    See section 7 for a complete list of possible patterns and transformation
8215    rules.
8216
8217    The table has to be declared using the "stick-table" statement. It must be of
8218    a type compatible with the pattern. By default it is the one which is present
8219    in the same backend. It is possible to share a table with other backends by
8220    referencing it using the "table" keyword. If another table is referenced,
8221    the server's ID inside the backends are used. By default, all server IDs
8222    start at 1 in each backend, so the server ordering is enough. But in case of
8223    doubt, it is highly recommended to force server IDs using their "id" setting.
8224
8225    It is possible to restrict the conditions where a "stick store-response"
8226    statement will apply, using "if" or "unless" followed by a condition. This
8227    condition will be evaluated while parsing the response, so any criteria can
8228    be used. See section 7 for ACL based conditions.
8229
8230    There is no limit on the number of "stick store-response" statements, but
8231    there is a limit of 8 simultaneous stores per request or response. This
8232    makes it possible to store up to 8 criteria, all extracted from either the
8233    request or the response, regardless of the number of rules. Only the 8 first
8234    ones which match will be kept. Using this, it is possible to feed multiple
8235    tables at once in the hope to increase the chance to recognize a user on
8236    another protocol or access method. Using multiple store-response rules with
8237    the same table is possible and may be used to find the best criterion to rely
8238    on, by arranging the rules by decreasing preference order. Only the first
8239    extracted criterion for a given table will be stored. All subsequent store-
8240    response rules referencing the same table will be skipped and their ACLs will
8241    not be evaluated. However, even if a store-request rule references a table, a
8242    store-response rule may also use the same table. This means that each table
8243    may learn exactly one element from the request and one element from the
8244    response at once.
8245
8246    The table will contain the real server that processed the request.
8247
8248    Example :
8249      # Learn SSL session ID from both request and response and create affinity.
8250      backend https
8251          mode tcp
8252          balance roundrobin
8253          # maximum SSL session ID length is 32 bytes.
8254          stick-table type binary len 32 size 30k expire 30m
8255
```



```
8256          acl clienthello req_ssl_hello_type 1
8257          acl serverhello rep_ssl_hello_type 2
8258
8259          # use tcp content accepts to detects ssl client and server hello.
8260          tcp-request inspect-delay 5s
8261          tcp-request content accept if clienthello
8262
8263          # no timeout on response inspect delay by default.
8264          tcp-response content accept if serverhello
8265
8266          # SSL session ID (SSLID) may be present on a client or server hello.
8267          # Its length is coded on 1 byte at offset 43 and its value starts
8268          # at offset 44.
8269
8270          # Match and learn on request if client hello.
8271          stick on payload_lv(43,1) if clienthello
8272
8273          # Learn on response if server hello.
8274          stick store-response payload_lv(43,1) if serverhello
8275
8276          server s1 192.168.1.1:443
8277          server s2 192.168.1.1:443
8278
8279    See also : "stick-table", "stick on", and section 7 about ACLs and pattern
8280               extraction.
8281
8282
8283  tcp-check connect [params*]
8284    Opens a new connection
8285    May be used in sections:   defaults | frontend | listen | backend
8286                                  no     |    no    |   yes  |   yes
8287
8288    When an application lies on more than a single TCP port or when HAProxy
8289    load-balance many services in a single backend, it makes sense to probe all
8290    the services individually before considering a server as operational.
8291
8292    When there are no TCP port configured on the server line neither server port
8293    directive, then the 'tcp-check connect port <port>' must be the first step
8294    of the sequence.
8295
8296    In a tcp-check ruleset a 'connect' is required, it is also mandatory to start
8297    the ruleset with a 'connect' rule. Purpose is to ensure admin know what they
8298    do.
8299
8300    Parameters :
8301      They are optional and can be used to describe how HAProxy should open and
8302      use the TCP connection.
8303
8304      port      if not set, check port or server port is used.
8305                It tells HAProxy where to open the connection to.
8306                <port> must be a valid TCP port source integer, from 1 to 65535.
8307
8308      send-proxy   send a PROXY protocol string
8309
8310      ssl          opens a ciphered connection
8311
8312      Examples:
8313           # check HTTP and HTTPs services on a server.
8314           # first open port 80 thanks to server line port directive, then
8315           # tcp-check opens port 443, ciphered and run a request on it:
8316           option tcp-check
8317           tcp-check connect
8318           tcp-check send GET\ /\ HTTP/1.0\r\n
8319           tcp-check send Host:\ haproxy.1wt.eu\r\n
8320           tcp-check send \r\n
```

```
          tcp-check expect rstring (2..|3..)
          tcp-check connect port 443 ssl
          tcp-check send GET\ /\ HTTP/1.0\r\n
          tcp-check send Host:\ haproxy.1wt.eu\r\n
          tcp-check send \r\n
          tcp-check expect rstring (2..|3..)
          server www 10.0.0.1 check port 80

          # check both POP and IMAP from a single server:
          option tcp-check
          tcp-check connect port 110
          tcp-check expect string +OK\ POP3\ ready
          tcp-check connect port 143
          tcp-check expect string *\ OK\ IMAP4\ ready
          server mail 10.0.0.1 check

   See also : "option tcp-check", "tcp-check send", "tcp-check expect"


tcp-check expect [!] <match> <pattern>
  Specify data to be collected and analysed during a generic health check
  May be used in sections:   defaults | frontend | listen | backend
                               no     |    no    |   yes  |   yes

  Arguments :
    <match>   is a keyword indicating how to look for a specific pattern in the
              response. The keyword may be one of "string", "rstring" or
              binary.
              The keyword may be preceded by an exclamation mark ("!") to negate
              the match. Spaces are allowed between the exclamation mark and the
              keyword. See below for more details on the supported keywords.

    <pattern> is the pattern to look for. It may be a string or a regular
              expression. If the pattern contains spaces, they must be escaped
              with the usual backslash ('\').
              If the match is set to binary, then the pattern must be passed as
              a serie of hexadecimal digits in an even number. Each sequence of
              two digits will represent a byte. The hexadecimal digits may be
              used upper or lower case.


   The available matches are intentionally similar to their http-check cousins :

     string <string> : test the exact string matches in the response buffer.
                       A health check response will be considered valid if the
                       response's buffer contains this exact string. If the
                       "string" keyword is prefixed with "!", then the response
                       will be considered invalid if the body contains this
                       string. This can be used to look for a mandatory pattern
                       in a protocol response, or to detect a failure when a
                       specific error appears in a protocol banner.

     rstring <regex> : test a regular expression on the response buffer.
                       A health check response will be considered valid if the
                       response's buffer matches this expression. If the
                       "rstring" keyword is prefixed with "!", then the response
                       will be considered invalid if the body matches the
                       expression.

     binary <hexstring> : test the exact string in its hexadecimal form matches
                          in the response buffer. A health check response will
                          be considered valid if the response's buffer contains
                          this exact hexadecimal string.
                          Purpose is to match data on binary protocols.
```



```
   It is important to note that the responses will be limited to a certain size
   defined by the global "tune.chksize" option, which defaults to 16384 bytes.
   Thus, too large responses may not contain the mandatory pattern when using
   "string", "rstring" or binary. If a large response is absolutely required, it
   is possible to change the default max size by setting the global variable.
   However, it is worth keeping in mind that parsing very large responses can
   waste some CPU cycles, especially when regular expressions are used, and that
   it is always better to focus the checks on smaller resources. Also, in its
   current state, the check will not find any string nor regex past a null
   character in the response. Similarly it is not possible to request matching
   the null character.

   Examples :
          # perform a POP check
          option tcp-check
          tcp-check expect string +OK\ POP3\ ready

          # perform an IMAP check
          option tcp-check
          tcp-check expect string *\ OK\ IMAP4\ ready

          # look for the redis master server
          option tcp-check
          tcp-check send PING\r\n
          tcp-check expect string +PONG
          tcp-check send info\ replication\r\n
          tcp-check expect string role:master
          tcp-check send QUIT\r\n
          tcp-check expect string +OK


   See also : "option tcp-check", "tcp-check connect", "tcp-check send",
              "tcp-check send-binary", "http-check expect", tune.chksize


tcp-check send <data>
  Specify a string to be sent as a question during a generic health check
  May be used in sections:   defaults | frontend | listen | backend
                               no     |    no    |   yes  |   yes

    <data> : the data to be sent as a question during a generic health check
             session. For now, <data> must be a string.

  Examples :
          # look for the redis master server
          option tcp-check
          tcp-check send info\ replication\r\n
          tcp-check expect string role:master

  See also : "option tcp-check", "tcp-check connect", "tcp-check expect",
             "tcp-check send-binary", tune.chksize


tcp-check send-binary <hexastring>
  Specify an hexa digits string to be sent as a binary question during a raw
  tcp health check
  May be used in sections:   defaults | frontend | listen | backend
                               no     |    no    |   yes  |   yes

    <data> : the data to be sent as a question during a generic health check
             session. For now, <data> must be a string.
    <hexastring> : test the exact string in its hexadecimal form matches in the
                   response buffer. A health check response will be considered
                   valid if the response's buffer contains this exact
                   hexadecimal string.
```

```
                    Purpose is to send binary data to ask on binary protocols.

   Examples :
         # redis check in binary
         option tcp-check
         tcp-check send-binary 50494e470d0a # PING\r\n
         tcp-check expect binary 2b504F4e47 # +PONG


   See also : "option tcp-check", "tcp-check connect", "tcp-check expect",
              "tcp-check send", tune.chksize


tcp-request connection <action> [{if | unless} <condition>]
  Perform an action on an incoming connection depending on a layer 4 condition
  May be used in sections :   defaults | frontend | listen | backend
                                 no    |    yes   |   yes  |   no
  Arguments :
    <action>    defines the action to perform if the condition applies. See
                below.

    <condition> is a standard layer4-only ACL-based condition (see section 7).

  Immediately after acceptance of a new incoming connection, it is possible to
  evaluate some conditions to decide whether this connection must be accepted
  or dropped or have its counters tracked. Those conditions cannot make use of
  any data contents because the connection has not been read from yet, and the
  buffers are not yet allocated. This is used to selectively and very quickly
  accept or drop connections from various sources with a very low overhead. If
  some contents need to be inspected in order to take the decision, the
  "tcp-request content" statements must be used instead.

  The "tcp-request connection" rules are evaluated in their exact declaration
  order. If no rule matches or if there is no rule, the default action is to
  accept the incoming connection. There is no specific limit to the number of
  rules which may be inserted.

  Four types of actions are supported :
    - accept :
        accepts the connection if the condition is true (when used with "if")
        or false (when used with "unless"). The first such rule executed ends
        the rules evaluation.

    - reject :
        rejects the connection if the condition is true (when used with "if")
        or false (when used with "unless"). The first such rule executed ends
        the rules evaluation. Rejected connections do not even become a
        session, which is why they are accounted separately for in the stats,
        as "denied connections". They are not considered for the session
        rate-limit and are not logged either. The reason is that these rules
        should only be used to filter extremely high connection rates such as
        the ones encountered during a massive DDoS attack. Under these extreme
        conditions, the simple action of logging each event would make the
        system collapse and would considerably lower the filtering capacity. If
        logging is absolutely desired, then "tcp-request content" rules should
        be used instead.

    - expect-proxy layer4 :
        configures the client-facing connection to receive a PROXY protocol
        header before any byte is read from the socket. This is equivalent to
        having the "accept-proxy" keyword on the "bind" line, except that using
        the TCP rule allows the PROXY protocol to be accepted only for certain
        IP address ranges using an ACL. This is convenient when multiple layers
        of load balancers are passed through by traffic coming from public
        hosts.
```



```

    - capture <sample> len <length> :
        This only applies to "tcp-request content" rules. It captures sample
        expression <sample> from the request buffer, and converts it to a
        string of at most <len> characters. The resulting string is stored into
        the next request "capture" slot, so it will possibly appear next to
        some captured HTTP headers. It will then automatically appear in the
        logs, and it will be possible to extract it using sample fetch rules to
        feed it into headers or anything. The length should be limited given
        that this size will be allocated for each capture during the whole
        session life. Please check section 7.3 (Fetching samples) and "capture
        request header" for more information.

    - { track-sc0 | track-sc1 | track-sc2 } <key> [table <table>] :
        enables tracking of sticky counters from current connection. These
        rules do not stop evaluation and do not change default action. 3 sets
        of counters may be simultaneously tracked by the same connection. The
        first "track-sc0" rule executed enables tracking of the counters of the
        specified table as the first set. The first "track-sc1" rule executed
        enables tracking of the counters of the specified table as the second
        set. The first "track-sc2" rule executed enables tracking of the
        counters of the specified table as the third set. It is a recommended
        practice to use the first set of counters for the per-frontend counters
        and the second set for the per-backend ones. But this is just a
        guideline, all may be used everywhere.

        These actions take one or two arguments :
          <key>   is mandatory, and is a sample expression rule as described
                  in section 7.3. It describes what elements of the incoming
                  request or connection will be analysed, extracted, combined,
                  and used to select which table entry to update the counters.
                  Note that "tcp-request connection" cannot use content-based
                  fetches.

         <table>  is an optional table to be used instead of the default one,
                  which is the stick-table declared in the current proxy. All
                  the counters for the matches and updates for the key will
                  then be performed in that table until the session ends.

        Once a "track-sc*" rule is executed, the key is looked up in the table
        and if it is not found, an entry is allocated for it. Then a pointer to
        that entry is kept during all the session's life, and this entry's
        counters are updated as often as possible, every time the session's
        counters are updated, and also systematically when the session ends.
        Counters are only updated for events that happen after the tracking has
        been started. For example, connection counters will not be updated when
        tracking layer 7 information, since the connection event happens before
        layer7 information is extracted.

        If the entry tracks concurrent connection counters, one connection is
        counted for as long as the entry is tracked, and the entry will not
        expire during that time. Tracking counters also provides a performance
        advantage over just checking the keys, because only one table lookup is
        performed for all ACL checks that make use of it.

    - sc-inc-gpc0(<sc-id>):
        The "sc-inc-gpc0" increments the GPC0 counter according to the sticky
        counter designated by <sc-id>. If an error occurs, this action silently
        fails and the actions evaluation continues.

    - sc-set-gpt0(<sc-id>) <int>:
        This action sets the GPT0 tag according to the sticky counter designated
        by <sc-id> and the value of <int>. The expected result is a boolean. If
        an error occurs, this action silently fails and the actions evaluation
        continues.
```

```
8581
8582      - "silent-drop" :
8583          This stops the evaluation of the rules and makes the client-facing
8584          connection suddenly disappear using a system-dependant way that tries
8585          to prevent the client from being notified. The effect it then that the
8586          client still sees an established connection while there's none on
8587          HAProxy. The purpose is to achieve a comparable effect to "tarpit"
8588          except that it doesn't use any local resource at all on the machine
8589          running HAProxy. It can resist much higher loads than "tarpit", and
8590          slow down stronger attackers. It is important to undestand the impact
8591          of using this mechanism. All stateful equipments placed between the
8592          client and HAProxy (firewalls, proxies, load balancers) will also keep
8593          the established connection for a long time and may suffer from this
8594          action.  On modern Linux systems running with enough privileges, the
8595          TCP_REPAIR socket option is used to block the emission of a TCP
8596          reset. On other systems, the socket's TTL is reduced to 1 so that the
8597          TCP reset doesn't pass the first router, though it's still delivered to
8598          local networks. Do not use it unless you fully understand how it works.
8599
8600    Note that the "if/unless" condition is optional. If no condition is set on
8601    the action, it is simply performed unconditionally. That can be useful for
8602    "track-sc*" actions as well as for changing the default action to a reject.
8603
8604    Example: accept all connections from white-listed hosts, reject too fast
8605             connection without counting them, and track accepted connections.
8606             This results in connection rate being capped from abusive sources.
8607
8608          tcp-request connection accept if { src -f /etc/haproxy/whitelist.lst }
8609          tcp-request connection reject if { src_conn_rate gt 10 }
8610          tcp-request connection track-sc0 src
8611
8612    Example: accept all connections from white-listed hosts, count all other
8613             connections and reject too fast ones. This results in abusive ones
8614             being blocked as long as they don't slow down.
8615
8616          tcp-request connection accept if { src -f /etc/haproxy/whitelist.lst }
8617          tcp-request connection track-sc0 src
8618          tcp-request connection reject if { sc0_conn_rate gt 10 }
8619
8620    Example: enable the PROXY protocol for traffic coming from all known proxies.
8621
8622          tcp-request connection expect-proxy layer4 if { src -f proxies.lst }
8623
8624    See section 7 about ACL usage.
8625
8626    See also : "tcp-request content", "stick-table"
8627
8628
8629  tcp-request content <action> [{if | unless} <condition>]
8630    Perform an action on a new session depending on a layer 4-7 condition
8631    May be used in sections :   defaults | frontend | listen | backend
8632                                   no    |    yes   |   yes  |   yes
8633    Arguments :
8634      <action>    defines the action to perform if the condition applies. See
8635                  below.
8636
8637      <condition> is a standard layer 4-7 ACL-based condition (see section 7).
8638
8639    A request's contents can be analysed at an early stage of request processing
8640    called "TCP content inspection". During this stage, ACL-based rules are
8641    evaluated every time the request contents are updated, until either an
8642    "accept" or a "reject" rule matches, or the TCP request inspection delay
8643    expires with no matching rule.
8644
8645    The first difference between these rules and "tcp-request connection" rules
```



```
8646    is that "tcp-request content" rules can make use of contents to take a
8647    decision. Most often, these decisions will consider a protocol recognition or
8648    validity. The second difference is that content-based rules can be used in
8649    both frontends and backends. In case of HTTP keep-alive with the client, all
8650    tcp-request content rules are evaluated again, so haproxy keeps a record of
8651    what sticky counters were assigned by a "tcp-request connection" versus a
8652    "tcp-request content" rule, and flushes all the content-related ones after
8653    processing an HTTP request, so that they may be evaluated again by the rules
8654    being evaluated again for the next request. This is of particular importance
8655    when the rule tracks some L7 information or when it is conditioned by an
8656    L7-based ACL, since tracking may change between requests.
8657
8658    Content-based rules are evaluated in their exact declaration order. If no
8659    rule matches or if there is no rule, the default action is to accept the
8660    contents. There is no specific limit to the number of rules which may be
8661    inserted.
8662
8663    Several types of actions are supported :
8664      - accept : the request is accepted
8665      - reject : the request is rejected and the connection is closed
8666      - capture : the specified sample expression is captured
8667      - { track-sc0 | track-sc1 | track-sc2 } <key> [table <table>]
8668      - sc-inc-gpc0(<sc-id>)
8669      - set-gpt0(<sc-id>) <int>
8670      - set-var(<var-name>) <expr>
8671      - silent-drop
8672
8673    They have the same meaning as their counter-parts in "tcp-request connection"
8674    so please refer to that section for a complete description.
8675
8676    While there is nothing mandatory about it, it is recommended to use the
8677    track-sc0 in "tcp-request connection" rules, track-sc1 for "tcp-request
8678    content" rules in the frontend, and track-sc2 for "tcp-request content"
8679    rules in the backend, because that makes the configuration more readable
8680    and easier to troubleshoot, but this is just a guideline and all counters
8681    may be used everywhere.
8682
8683    Note that the "if/unless" condition is optional. If no condition is set on
8684    the action, it is simply performed unconditionally. That can be useful for
8685    "track-sc*" actions as well as for changing the default action to a reject.
8686
8687    It is perfectly possible to match layer 7 contents with "tcp-request content"
8688    rules, since HTTP-specific ACL matches are able to preliminarily parse the
8689    contents of a buffer before extracting the required data. If the buffered
8690    contents do not parse as a valid HTTP message, then the ACL does not match.
8691    The parser which is involved there is exactly the same as for all other HTTP
8692    processing, so there is no risk of parsing something differently. In an HTTP
8693    backend connected to from an HTTP frontend, it is guaranteed that HTTP
8694    contents will always be immediately present when the rule is evaluated first.
8695
8696    Tracking layer7 information is also possible provided that the information
8697    are present when the rule is processed. The rule processing engine is able to
8698    wait until the inspect delay expires when the data to be tracked is not yet
8699    available.
8700
8701    The "set-var" is used to set the content of a variable. The variable is
8702    declared inline.
8703
8704      <var-name> The name of the variable starts by an indication about its scope.
8705                 The allowed scopes are:
8706                   "sess" : the variable is shared with all the session,
8707                   "txn"  : the variable is shared with all the transaction
8708                            (request and response)
8709                   "req"  : the variable is shared only during the request
8710                            processing
```

```
                  "res"  : the variable is shared only during the response
                           processing.
               This prefix is followed by a name. The separator is a '.'.
               The name may only contain characters 'a-z', 'A-Z', '0-9' and '_'.

  <expr>       Is a standard HAProxy expression formed by a sample-fetch
               followed by some converters.

Example:

      tcp-request content set-var(sess.my_var) src

Example:
      # Accept HTTP requests containing a Host header saying "example.com"
      # and reject everything else.
      acl is_host_com hdr(Host) -i example.com
      tcp-request inspect-delay 30s
      tcp-request content accept if is_host_com
      tcp-request content reject

Example:
      # reject SMTP connection if client speaks first
      tcp-request inspect-delay 30s
      acl content_present req_len gt 0
      tcp-request content reject if content_present

      # Forward HTTPS connection only if client speaks
      tcp-request inspect-delay 30s
      acl content_present req_len gt 0
      tcp-request content accept if content_present
      tcp-request content reject

Example:
      # Track the last IP from X-Forwarded-For
      tcp-request inspect-delay 10s
      tcp-request content track-sc0 hdr(x-forwarded-for,-1)

Example:
      # track request counts per "base" (concatenation of Host+URL)
      tcp-request inspect-delay 10s
      tcp-request content track-sc0 base table req-rate

Example: track per-frontend and per-backend counters, block abusers at the
         frontend when the backend detects abuse.

      frontend http
          # Use General Purpose Couter 0 in SC0 as a global abuse counter
          # protecting all our sites
          stick-table type ip size 1m expire 5m store gpc0
          tcp-request connection track-sc0 src
          tcp-request connection reject if { sc0_get_gpc0 gt 0 }
          ...
          use_backend http_dynamic if { path_end .php }

      backend http_dynamic
          # if a source makes too fast requests to this dynamic site (tracked
          # by SC1), block it globally in the frontend.
          stick-table type ip size 1m expire 5m store http_req_rate(10s)
          acl click_too_fast sc1_http_req_rate gt 10
          acl mark_as_abuser sc0_inc_gpc0 gt 0
          tcp-request content track-sc1 src
          tcp-request content reject if click_too_fast mark_as_abuser

See section 7 about ACL usage.
```



```
  See also : "tcp-request connection", "tcp-request inspect-delay"


tcp-request inspect-delay <timeout>
  Set the maximum allowed time to wait for data during content inspection
  May be used in sections :   defaults | frontend | listen | backend
                                 no    |    yes   |   yes  |   yes
  Arguments :
    <timeout> is the timeout value specified in milliseconds by default, but
              can be in any other unit if the number is suffixed by the unit,
              as explained at the top of this document.

  People using haproxy primarily as a TCP relay are often worried about the
  risk of passing any type of protocol to a server without any analysis. In
  order to be able to analyze the request contents, we must first withhold
  the data then analyze them. This statement simply enables withholding of
  data for at most the specified amount of time.

  TCP content inspection applies very early when a connection reaches a
  frontend, then very early when the connection is forwarded to a backend. This
  means that a connection may experience a first delay in the frontend and a
  second delay in the backend if both have tcp-request rules.

  Note that when performing content inspection, haproxy will evaluate the whole
  rules for every new chunk which gets in, taking into account the fact that
  those data are partial. If no rule matches before the aforementioned delay,
  a last check is performed upon expiration, this time considering that the
  contents are definitive. If no delay is set, haproxy will not wait at all
  and will immediately apply a verdict based on the available information.
  Obviously this is unlikely to be very useful and might even be racy, so such
  setups are not recommended.

  As soon as a rule matches, the request is released and continues as usual. If
  the timeout is reached and no rule matches, the default policy will be to let
  it pass through unaffected.

  For most protocols, it is enough to set it to a few seconds, as most clients
  send the full request immediately upon connection. Add 3 or more seconds to
  cover TCP retransmits but that's all. For some protocols, it may make sense
  to use large values, for instance to ensure that the client never talks
  before the server (eg: SMTP), or to wait for a client to talk before passing
  data to the server (eg: SSL). Note that the client timeout must cover at
  least the inspection delay, otherwise it will expire first. If the client
  closes the connection or if the buffer is full, the delay immediately expires
  since the contents will not be able to change anymore.

  See also : "tcp-request content accept", "tcp-request content reject",
             "timeout client".


tcp-response content <action> [{if | unless} <condition>]
  Perform an action on a session response depending on a layer 4-7 condition
  May be used in sections :   defaults | frontend | listen | backend
                                 no    |    no    |   yes  |   yes
  Arguments :
    <action>    defines the action to perform if the condition applies. See
                below.

    <condition> is a standard layer 4-7 ACL-based condition (see section 7).

  Response contents can be analysed at an early stage of response processing
  called "TCP content inspection". During this stage, ACL-based rules are
  evaluated every time the response contents are updated, until either an
  "accept", "close" or a "reject" rule matches, or a TCP response inspection
  delay is set and expires with no matching rule.
```

```
8841
8842    Most often, these decisions will consider a protocol recognition or validity.
8843
8844    Content-based rules are evaluated in their exact declaration order. If no
8845    rule matches or if there is no rule, the default action is to accept the
8846    contents. There is no specific limit to the number of rules which may be
8847    inserted.
8848
8849    Several types of actions are supported :
8850      - accept :
8851          accepts the response if the condition is true (when used with "if")
8852          or false (when used with "unless"). The first such rule executed ends
8853          the rules evaluation.
8854
8855      - close :
8856          immediately closes the connection with the server if the condition is
8857          true (when used with "if"), or false (when used with "unless"). The
8858          first such rule executed ends the rules evaluation. The main purpose of
8859          this action is to force a connection to be finished between a client
8860          and a server after an exchange when the application protocol expects
8861          some long time outs to elapse first. The goal is to eliminate idle
8862          connections which take significant resources on servers with certain
8863          protocols.
8864
8865      - reject :
8866          rejects the response if the condition is true (when used with "if")
8867          or false (when used with "unless"). The first such rule executed ends
8868          the rules evaluation. Rejected session are immediately closed.
8869
8870      - set-var(<var-name>) <expr>
8871          Sets a variable.
8872
8873      - sc-inc-gpc0(<sc-id>):
8874          This action increments the GPC0 counter according to the sticky
8875          counter designated by <sc-id>. If an error occurs, this action fails
8876          silently and the actions evaluation continues.
8877
8878      - sc-set-gpt0(<sc-id>) <int> :
8879          This action sets the GPT0 tag according to the sticky counter designated
8880          by <sc-id> and the value of <int>. The expected result is a boolean. If
8881          an error occurs, this action silently fails and the actions evaluation
8882          continues.
8883
8884      - "silent-drop" :
8885          This stops the evaluation of the rules and makes the client-facing
8886          connection suddenly disappear using a system-dependant way that tries
8887          to prevent the client from being notified. The effect it then that the
8888          client still sees an established connection while there's none on
8889          HAProxy. The purpose is to achieve a comparable effect to "tarpit"
8890          except that it doesn't use any local resource at all on the machine
8891          running HAProxy. It can resist much higher loads than "tarpit", and
8892          slow down stronger attackers. It is important to undestand the impact
8893          of using this mechanism. All stateful equipments placed between the
8894          client and HAProxy (firewalls, proxies, load balancers) will also keep
8895          the established connection for a long time and may suffer from this
8896          action.  On modern Linux systems running with enough privileges, the
8897          TCP_REPAIR socket option is used to block the emission of a TCP
8898          reset. On other systems, the socket's TTL is reduced to 1 so that the
8899          TCP reset doesn't pass the first router, though it's still delivered to
8900          local networks. Do not use it unless you fully understand how it works.
8901
8902    Note that the "if/unless" condition is optional. If no condition is set on
8903    the action, it is simply performed unconditionally. That can be useful for
8904    for changing the default action to a reject.
8905
```



```
8906    It is perfectly possible to match layer 7 contents with "tcp-response
8907    content" rules, but then it is important to ensure that a full response has
8908    been buffered, otherwise no contents will match. In order to achieve this,
8909    the best solution involves detecting the HTTP protocol during the inspection
8910    period.
8911
8912    The "set-var" is used to set the content of a variable. The variable is
8913    declared inline.
8914
8915      <var-name> The name of the variable starts by an indication about its scope.
8916                 The allowed scopes are:
8917                   "sess" : the variable is shared with all the session,
8918                   "txn"  : the variable is shared with all the transaction
8919                            (request and response)
8920                   "req"  : the variable is shared only during the request
8921                            processing
8922                   "res"  : the variable is shared only during the response
8923                            processing.
8924                 This prefix is followed by a name. The separator is a '.'.
8925                 The name may only contain characters 'a-z', 'A-Z', '0-9' and '_'.
8926
8927      <expr>     Is a standard HAProxy expression formed by a sample-fetch
8928                 followed by some converters.
8929
8930    Example:
8931
8932         tcp-request content set-var(sess.my_var) src
8933
8934    See section 7 about ACL usage.
8935
8936    See also : "tcp-request content", "tcp-response inspect-delay"
8937
8938
8939  tcp-response inspect-delay <timeout>
8940    Set the maximum allowed time to wait for a response during content inspection
8941    May be used in sections :   defaults | frontend | listen | backend
8942                                   no    |    no    |   yes  |   yes
8943    Arguments :
8944      <timeout> is the timeout value specified in milliseconds by default, but
8945                can be in any other unit if the number is suffixed by the unit,
8946                as explained at the top of this document.
8947
8948    See also : "tcp-response content", "tcp-request inspect-delay".
8949
8950
8951  timeout check <timeout>
8952    Set additional check timeout, but only after a connection has been already
8953    established.
8954
8955    May be used in sections:    defaults | frontend | listen | backend
8956                                   yes   |    no    |   yes  |   yes
8957    Arguments:
8958      <timeout> is the timeout value specified in milliseconds by default, but
8959                can be in any other unit if the number is suffixed by the unit,
8960                as explained at the top of this document.
8961
8962    If set, haproxy uses min("timeout connect", "inter") as a connect timeout
8963    for check and "timeout check" as an additional read timeout. The "min" is
8964    used so that people running with *very* long "timeout connect" (eg. those
8965    who needed this due to the queue or tarpit) do not slow down their checks.
8966    (Please also note that there is no valid reason to have such long connect
8967    timeouts, because "timeout queue" and "timeout tarpit" can always be used to
8968    avoid that).
8969
8970    If "timeout check" is not set haproxy uses "inter" for complete check
```

```
   timeout (connect + read) exactly like all <1.3.15 version.

   In most cases check request is much simpler and faster to handle than normal
   requests and people may want to kick out laggy servers so this timeout should
   be smaller than "timeout server".

   This parameter is specific to backends, but can be specified once for all in
   "defaults" sections. This is in fact one of the easiest solutions not to
   forget about it.

   See also: "timeout connect", "timeout queue", "timeout server",
             "timeout tarpit".


timeout client <timeout>
timeout clitimeout <timeout> (deprecated)
  Set the maximum inactivity time on the client side.
  May be used in sections :   defaults | frontend | listen | backend
                                 yes   |    yes   |   yes  |   no
  Arguments :
    <timeout> is the timeout value specified in milliseconds by default, but
              can be in any other unit if the number is suffixed by the unit,
              as explained at the top of this document.

   The inactivity timeout applies when the client is expected to acknowledge or
   send data. In HTTP mode, this timeout is particularly important to consider
   during the first phase, when the client sends the request, and during the
   response while it is reading data sent by the server. The value is specified
   in milliseconds by default, but can be in any other unit if the number is
   suffixed by the unit, as specified at the top of this document. In TCP mode
   (and to a lesser extent, in HTTP mode), it is highly recommended that the
   client timeout remains equal to the server timeout in order to avoid complex
   situations to debug. It is a good practice to cover one or several TCP packet
   losses by specifying timeouts that are slightly above multiples of 3 seconds
   (eg: 4 or 5 seconds). If some long-lived sessions are mixed with short-lived
   sessions (eg: WebSocket and HTTP), it's worth considering "timeout tunnel",
   which overrides "timeout client" and "timeout server" for tunnels, as well as
   "timeout client-fin" for half-closed connections.

   This parameter is specific to frontends, but can be specified once for all in
   "defaults" sections. This is in fact one of the easiest solutions not to
   forget about it. An unspecified timeout results in an infinite timeout, which
   is not recommended. Such a usage is accepted and works but reports a warning
   during startup because it may results in accumulation of expired sessions in
   the system if the system's timeouts are not configured either.

   This parameter replaces the old, deprecated "clitimeout". It is recommended
   to use it to write new configurations. The form "timeout clitimeout" is
   provided only by backwards compatibility but its use is strongly discouraged.

   See also : "clitimeout", "timeout server", "timeout tunnel".


timeout client-fin <timeout>
  Set the inactivity timeout on the client side for half-closed connections.
  May be used in sections :   defaults | frontend | listen | backend
                                 yes   |    yes   |   yes  |   no
  Arguments :
    <timeout> is the timeout value specified in milliseconds by default, but
              can be in any other unit if the number is suffixed by the unit,
              as explained at the top of this document.

   The inactivity timeout applies when the client is expected to acknowledge or
   send data while one direction is already shut down. This timeout is different
   from "timeout client" in that it only applies to connections which are closed
```



```
   in one direction. This is particularly useful to avoid keeping connections in
   FIN_WAIT state for too long when clients do not disconnect cleanly. This
   problem is particularly common long connections such as RDP or WebSocket.
   Note that this timeout can override "timeout tunnel" when a connection shuts
   down in one direction.

   This parameter is specific to frontends, but can be specified once for all in
   "defaults" sections. By default it is not set, so half-closed connections
   will use the other timeouts (timeout.client or timeout.tunnel).

   See also : "timeout client", "timeout server-fin", and "timeout tunnel".


timeout connect <timeout>
timeout contimeout <timeout> (deprecated)
  Set the maximum time to wait for a connection attempt to a server to succeed.
  May be used in sections :   defaults | frontend | listen | backend
                                 yes   |    no    |   yes  |   yes
  Arguments :
    <timeout> is the timeout value specified in milliseconds by default, but
              can be in any other unit if the number is suffixed by the unit,
              as explained at the top of this document.

   If the server is located on the same LAN as haproxy, the connection should be
   immediate (less than a few milliseconds). Anyway, it is a good practice to
   cover one or several TCP packet losses by specifying timeouts that are
   slightly above multiples of 3 seconds (eg: 4 or 5 seconds). By default, the
   connect timeout also presets both queue and tarpit timeouts to the same value
   if these have not been specified.

   This parameter is specific to backends, but can be specified once for all in
   "defaults" sections. This is in fact one of the easiest solutions not to
   forget about it. An unspecified timeout results in an infinite timeout, which
   is not recommended. Such a usage is accepted and works but reports a warning
   during startup because it may results in accumulation of failed sessions in
   the system if the system's timeouts are not configured either.

   This parameter replaces the old, deprecated "contimeout". It is recommended
   to use it to write new configurations. The form "timeout contimeout" is
   provided only by backwards compatibility but its use is strongly discouraged.

   See also: "timeout check", "timeout queue", "timeout server", "contimeout",
             "timeout tarpit".


timeout http-keep-alive <timeout>
  Set the maximum allowed time to wait for a new HTTP request to appear
  May be used in sections :   defaults | frontend | listen | backend
                                 yes   |    yes   |   yes  |   yes
  Arguments :
    <timeout> is the timeout value specified in milliseconds by default, but
              can be in any other unit if the number is suffixed by the unit,
              as explained at the top of this document.

   By default, the time to wait for a new request in case of keep-alive is set
   by "timeout http-request". However this is not always convenient because some
   people want very short keep-alive timeouts in order to release connections
   faster, and others prefer to have larger ones but still have short timeouts
   once the request has started to present itself.

   The "http-keep-alive" timeout covers these needs. It will define how long to
   wait for a new HTTP request to start coming after a response was sent. Once
   the first byte of request has been seen, the "http-request" timeout is used
   to wait for the complete request to come. Note that empty lines prior to a
   new request do not refresh the timeout and are not counted as a new request.
```

There is also another difference between the two timeouts : when a connection expires during timeout http-keep-alive, no error is returned, the connection just closes. If the connection expires in "http-request" while waiting for a connection to complete, a HTTP 408 error is returned.

In general it is optimal to set this value to a few tens to hundreds of milliseconds, to allow users to fetch all objects of a page at once but without waiting for further clicks. Also, if set to a very small value (eg: 1 millisecond) it will probably only accept pipelined requests but not the non-pipelined ones. It may be a nice trade-off for very large sites running with tens to hundreds of thousands of clients.

If this parameter is not set, the "http-request" timeout applies, and if both are not set, "timeout client" still applies at the lower level. It should be set in the frontend to take effect, unless the frontend is in TCP mode, in which case the HTTP backend's timeout will be used.

See also : "timeout http-request", "timeout client".

timeout http-request <timeout>
  Set the maximum allowed time to wait for a complete HTTP request
  May be used in sections :   defaults | frontend | listen | backend
                                 yes | yes | yes | yes
  Arguments :
    <timeout> is the timeout value specified in milliseconds by default, but
              can be in any other unit if the number is suffixed by the unit,
              as explained at the top of this document.

  In order to offer DoS protection, it may be required to lower the maximum accepted time to receive a complete HTTP request without affecting the client timeout. This helps protecting against established connections on which nothing is sent. The client timeout cannot offer a good protection against this abuse because it is an inactivity timeout, which means that if the attacker sends one character every now and then, the timeout will not trigger. With the HTTP request timeout, no matter what speed the client types, the request will be aborted if it does not complete in time. When the timeout expires, an HTTP 408 response is sent to the client to inform it about the problem, and the connection is closed. The logs will report termination codes "cR". Some recent browsers are having problems with this standard, well-documented behaviour, so it might be needed to hide the 408 code using "option http-ignore-probes" or "errorfile 408 /dev/null". See more details in the explanations of the "cR" termination code in section 8.5.

  By default, this timeout only applies to the header part of the request, and not to any data. As soon as the empty line is received, this timeout is not used anymore. When combined with "option http-buffer-request", this timeout also applies to the body of the request..
  It is used again on keep-alive connections to wait for a second request if "timeout http-keep-alive" is not set.

  Generally it is enough to set it to a few seconds, as most clients send the full request immediately upon connection. Add 3 or more seconds to cover TCP retransmits but that's all. Setting it to very low values (eg: 50 ms) will generally work on local networks as long as there are no packet losses. This will prevent people from sending bare HTTP requests using telnet.

  If this parameter is not set, the client timeout still applies between each chunk of the incoming request. It should be set in the frontend to take effect, unless the frontend is in TCP mode, in which case the HTTP backend's timeout will be used.

  See also : "errorfile", "http-ignore-probes", "timeout http-keep-alive", and
             "timeout client", "option http-buffer-request".

timeout queue <timeout>
  Set the maximum time to wait in the queue for a connection slot to be free
  May be used in sections :   defaults | frontend | listen | backend
                                 yes | no | yes | yes
  Arguments :
    <timeout> is the timeout value specified in milliseconds by default, but
              can be in any other unit if the number is suffixed by the unit,
              as explained at the top of this document.

  When a server's maxconn is reached, connections are left pending in a queue which may be server-specific or global to the backend. In order not to wait indefinitely, a timeout is applied to requests pending in the queue. If the timeout is reached, it is considered that the request will almost never be served, so it is dropped and a 503 error is returned to the client.

  The "timeout queue" statement allows to fix the maximum time for a request to be left pending in a queue. If unspecified, the same value as the backend's connection timeout ("timeout connect") is used, for backwards compatibility with older versions with no "timeout queue" parameter.

  See also : "timeout connect", "contimeout".

timeout server <timeout>
timeout srvtimeout <timeout> (deprecated)
  Set the maximum inactivity time on the server side.
  May be used in sections :   defaults | frontend | listen | backend
                                 yes | no | yes | yes
  Arguments :
    <timeout> is the timeout value specified in milliseconds by default, but
              can be in any other unit if the number is suffixed by the unit,
              as explained at the top of this document.

  The inactivity timeout applies when the server is expected to acknowledge or send data. In HTTP mode, this timeout is particularly important to consider during the first phase of the server's response, when it has to send the headers, as it directly represents the server's processing time for the request. To find out what value to put there, it's often good to start with what would be considered as unacceptable response times, then check the logs to observe the response time distribution, and adjust the value accordingly.

  The value is specified in milliseconds by default, but can be in any other unit if the number is suffixed by the unit, as specified at the top of this document. In TCP mode (and to a lesser extent, in HTTP mode), it is highly recommended that the client timeout remains equal to the server timeout in order to avoid complex situations to debug. Whatever the expected server response times, it is a good practice to cover at least one or several TCP packet losses by specifying timeouts that are slightly above multiples of 3 seconds (eg: 4 or 5 seconds minimum). If some long-lived sessions are mixed with short-lived sessions (eg: WebSocket and HTTP), it's worth considering "timeout tunnel", which overrides "timeout client" and "timeout server" for tunnels.

  This parameter is specific to backends, but can be specified once for all in "defaults" sections. This is in fact one of the easiest solutions not to forget about it. An unspecified timeout results in an infinite timeout, which is not recommended. Such a usage is accepted and works but reports a warning during startup because it may results in accumulation of expired sessions in the system if the system's timeouts are not configured either.

  This parameter replaces the old, deprecated "srvtimeout". It is recommended to use it to write new configurations. The form "timeout srvtimeout" is provided only by backwards compatibility but its use is strongly discouraged.

```
9231
9232    See also : "srvtimeout", "timeout client" and "timeout tunnel".
9233
9234
9235  timeout server-fin <timeout>
9236    Set the inactivity timeout on the server side for half-closed connections.
9237    May be used in sections :   defaults | frontend | listen | backend
9238                                   yes   |    no    |   yes  |   yes
9239    Arguments :
9240      <timeout> is the timeout value specified in milliseconds by default, but
9241                can be in any other unit if the number is suffixed by the unit,
9242                as explained at the top of this document.
9243
9244    The inactivity timeout applies when the server is expected to acknowledge or
9245    send data while one direction is already shut down. This timeout is different
9246    from "timeout server" in that it only applies to connections which are closed
9247    in one direction. This is particularly useful to avoid keeping connections in
9248    FIN_WAIT state for too long when a remote server does not disconnect cleanly.
9249    This problem is particularly common long connections such as RDP or WebSocket.
9250    Note that this timeout can override "timeout tunnel" when a connection shuts
9251    down in one direction. This setting was provided for completeness, but in most
9252    situations, it should not be needed.
9253
9254    This parameter is specific to backends, but can be specified once for all in
9255    "defaults" sections. By default it is not set, so half-closed connections
9256    will use the other timeouts (timeout.server or timeout.tunnel).
9257
9258    See also : "timeout client-fin", "timeout server", and "timeout tunnel".
9259
9260
9261  timeout tarpit <timeout>
9262    Set the duration for which tarpitted connections will be maintained
9263    May be used in sections :   defaults | frontend | listen | backend
9264                                   yes   |    yes   |   yes  |   yes
9265    Arguments :
9266      <timeout> is the tarpit duration specified in milliseconds by default, but
9267                can be in any other unit if the number is suffixed by the unit,
9268                as explained at the top of this document.
9269
9270    When a connection is tarpitted using "reqtarpit", it is maintained open with
9271    no activity for a certain amount of time, then closed. "timeout tarpit"
9272    defines how long it will be maintained open.
9273
9274    The value is specified in milliseconds by default, but can be in any other
9275    unit if the number is suffixed by the unit, as specified at the top of this
9276    document. If unspecified, the same value as the backend's connection timeout
9277    ("timeout connect") is used, for backwards compatibility with older versions
9278    with no "timeout tarpit" parameter.
9279
9280    See also : "timeout connect", "contimeout".
9281
9282
9283  timeout tunnel <timeout>
9284    Set the maximum inactivity time on the client and server side for tunnels.
9285    May be used in sections :   defaults | frontend | listen | backend
9286                                   yes   |    no    |   yes  |   yes
9287    Arguments :
9288      <timeout> is the timeout value specified in milliseconds by default, but
9289                can be in any other unit if the number is suffixed by the unit,
9290                as explained at the top of this document.
9291
9292    The tunnel timeout applies when a bidirectional connection is established
9293    between a client and a server, and the connection remains inactive in both
9294    directions. This timeout supersedes both the client and server timeouts once
9295    the connection becomes a tunnel. In TCP, this timeout is used as soon as no
```



```
9296    analyser remains attached to either connection (eg: tcp content rules are
9297    accepted). In HTTP, this timeout is used when a connection is upgraded (eg:
9298    when switching to the WebSocket protocol, or forwarding a CONNECT request
9299    to a proxy), or after the first response when no keepalive/close option is
9300    specified.
9301
9302    Since this timeout is usually used in conjunction with long-lived connections,
9303    it usually is a good idea to also set "timeout client-fin" to handle the
9304    situation where a client suddenly disappears from the net and does not
9305    acknowledge a close, or sends a shutdown and does not acknowledge pending
9306    data anymore. This can happen in lossy networks where firewalls are present,
9307    and is detected by the presence of large amounts of sessions in a FIN_WAIT
9308    state.
9309
9310    The value is specified in milliseconds by default, but can be in any other
9311    unit if the number is suffixed by the unit, as specified at the top of this
9312    document. Whatever the expected normal idle time, it is a good practice to
9313    cover at least one or several TCP packet losses by specifying timeouts that
9314    are slightly above multiples of 3 seconds (eg: 4 or 5 seconds minimum).
9315
9316    This parameter is specific to backends, but can be specified once for all in
9317    "defaults" sections. This is in fact one of the easiest solutions not to
9318    forget about it.
9319
9320    Example :
9321          defaults http
9322              option http-server-close
9323              timeout connect 5s
9324              timeout client 30s
9325              timeout client-fin 30s
9326              timeout server 30s
9327              timeout tunnel  1h    # timeout to use with WebSocket and CONNECT
9328
9329    See also : "timeout client", "timeout client-fin", "timeout server".
9330
9331
9332  transparent (deprecated)
9333    Enable client-side transparent proxying
9334    May be used in sections :   defaults | frontend | listen | backend
9335                                   yes   |    no    |   yes  |   yes
9336    Arguments : none
9337
9338    This keyword was introduced in order to provide layer 7 persistence to layer
9339    3 load balancers. The idea is to use the OS's ability to redirect an incoming
9340    connection for a remote address to a local process (here HAProxy), and let
9341    this process know what address was initially requested. When this option is
9342    used, sessions without cookies will be forwarded to the original destination
9343    IP address of the incoming request (which should match that of another
9344    equipment), while requests with cookies will still be forwarded to the
9345    appropriate server.
9346
9347    The "transparent" keyword is deprecated, use "option transparent" instead.
9348
9349    Note that contrary to a common belief, this option does NOT make HAProxy
9350    present the client's IP to the server when establishing the connection.
9351
9352    See also: "option transparent"
9353
9354  unique-id-format <string>
9355    Generate a unique ID for each request.
9356    May be used in sections :   defaults | frontend | listen | backend
9357                                   yes   |    yes   |   yes  |   no
9358    Arguments :
9359      <string>   is a log-format string.
9360
```

```
   This keyword creates a ID for each request using the custom log format. A
   unique ID is useful to trace a request passing through many components of
   a complex infrastructure. The newly created ID may also be logged using the
   %ID tag the log-format string.

   The format should be composed from elements that are guaranteed to be
   unique when combined together. For instance, if multiple haproxy instances
   are involved, it might be important to include the node name. It is often
   needed to log the incoming connection's source and destination addresses
   and ports. Note that since multiple requests may be performed over the same
   connection, including a request counter may help differentiate them.
   Similarly, a timestamp may protect against a rollover of the counter.
   Logging the process ID will avoid collisions after a service restart.

   It is recommended to use hexadecimal notation for many fields since it
   makes them more compact and saves space in logs.

   Example:

         unique-id-format %{+X}o\ %ci:%cp_%fi:%fp_%Ts_%rt:%pid

         will generate:

               7F000001:8296_7F00001E:1F90_4F7B0A69_0003:790A

   See also: "unique-id-header"

unique-id-header <name>
  Add a unique ID header in the HTTP request.
  May be used in sections :   defaults | frontend | listen | backend
                                  yes  |    yes   |   yes  |   no
  Arguments :
    <name>   is the name of the header.

  Add a unique-id header in the HTTP request sent to the server, using the
  unique-id-format. It can't work if the unique-id-format doesn't exist.

  Example:

         unique-id-format %{+X}o\ %ci:%cp_%fi:%fp_%Ts_%rt:%pid
         unique-id-header X-Unique-ID

         will generate:

            X-Unique-ID: 7F000001:8296_7F00001E:1F90_4F7B0A69_0003:790A

    See also: "unique-id-format"

use_backend <backend> [{if | unless} <condition>]
  Switch to a specific backend if/unless an ACL-based condition is matched.
  May be used in sections :   defaults | frontend | listen | backend
                                  no   |    yes   |   yes  |   no
  Arguments :
    <backend>   is the name of a valid backend or "listen" section, or a
                "log-format" string resolving to a backend name.

    <condition> is a condition composed of ACLs, as described in section 7. If
                it is omitted, the rule is unconditionally applied.

  When doing content-switching, connections arrive on a frontend and are then
  dispatched to various backends depending on a number of conditions. The
  relation between the conditions and the backends is described with the
  "use_backend" keyword. While it is normally used with HTTP processing, it can
  also be used in pure TCP, either without content using stateless ACLs (eg:
  source address validation) or combined with a "tcp-request" rule to wait for
```



```
  some payload.

  There may be as many "use_backend" rules as desired. All of these rules are
  evaluated in their declaration order, and the first one which matches will
  assign the backend.

  In the first form, the backend will be used if the condition is met. In the
  second form, the backend will be used if the condition is not met. If no
  condition is valid, the backend defined with "default_backend" will be used.
  If no default backend is defined, either the servers in the same section are
  used (in case of a "listen" section) or, in case of a frontend, no server is
  used and a 503 service unavailable response is returned.

  Note that it is possible to switch from a TCP frontend to an HTTP backend. In
  this case, either the frontend has already checked that the protocol is HTTP,
  and backend processing will immediately follow, or the backend will wait for
  a complete HTTP request to get in. This feature is useful when a frontend
  must decode several protocols on a unique port, one of them being HTTP.

  When <backend> is a simple name, it is resolved at configuration time, and an
  error is reported if the specified backend does not exist. If <backend> is
  a log-format string instead, no check may be done at configuration time, so
  the backend name is resolved dynamically at run time. If the resulting
  backend name does not correspond to any valid backend, no other rule is
  evaluated, and the default_backend directive is applied instead. Note that
  when using dynamic backend names, it is highly recommended to use a prefix
  that no other backend uses in order to ensure that an unauthorized backend
  cannot be forced from the request.

  It is worth mentioning that "use_backend" rules with an explicit name are
  used to detect the association between frontends and backends to compute the
  backend's "fullconn" setting. This cannot be done for dynamic names.

  See also: "default_backend", "tcp-request", "fullconn", "log-format", and
            section 7 about ACLs.


use-server <server> if <condition>
use-server <server> unless <condition>
  Only use a specific server if/unless an ACL-based condition is matched.
  May be used in sections :   defaults | frontend | listen | backend
                                  no   |    no    |   yes  |   yes
  Arguments :
    <server>    is the name of a valid server in the same backend section.

    <condition> is a condition composed of ACLs, as described in section 7.

  By default, connections which arrive to a backend are load-balanced across
  the available servers according to the configured algorithm, unless a
  persistence mechanism such as a cookie is used and found in the request.

  Sometimes it is desirable to forward a particular request to a specific
  server without having to declare a dedicated backend for this server. This
  can be achieved using the "use-server" rules. These rules are evaluated after
  the "redirect" rules and before evaluating cookies, and they have precedence
  on them. There may be as many "use-server" rules as desired. All of these
  rules are evaluated in their declaration order, and the first one which
  matches will assign the server.

  If a rule designates a server which is down, and "option persist" is not used
  and no force-persist rule was validated, it is ignored and evaluation goes on
  with the next rules until one matches.

  In the first form, the server will be used if the condition is met. In the
  second form, the server will be used if the condition is not met. If no
```

```
   condition is valid, the processing continues and the server will be assigned
   according to other persistence mechanisms.

   Note that even if a rule is matched, cookie processing is still performed but
   does not assign the server. This allows prefixed cookies to have their prefix
   stripped.

   The "use-server" statement works both in HTTP and TCP mode. This makes it
   suitable for use with content-based inspection. For instance, a server could
   be selected in a farm according to the TLS SNI field. And if these servers
   have their weight set to zero, they will not be used for other traffic.

   Example :
      # intercept incoming TLS requests based on the SNI field
      use-server www if { req_ssl_sni -i www.example.com }
      server     www 192.168.0.1:443 weight 0
      use-server mail if { req_ssl_sni -i mail.example.com }
      server     mail 192.168.0.1:587 weight 0
      use-server imap if { req_ssl_sni -i imap.example.com }
      server     mail 192.168.0.1:993 weight 0
      # all the rest is forwarded to this server
      server  default 192.168.0.2:443 check

   See also: "use_backend", section 5 about server and section 7 about ACLs.


5. Bind and Server options
--------------------------

The "bind", "server" and "default-server" keywords support a number of settings
depending on some build options and on the system HAProxy was built on. These
settings generally each consist in one word sometimes followed by a value,
written on the same line as the "bind" or "server" line. All these options are
described in this section.


5.1. Bind options
-----------------

The "bind" keyword supports a certain number of settings which are all passed
as arguments on the same line. The order in which those arguments appear makes
no importance, provided that they appear after the bind address. All of these
parameters are optional. Some of them consist in a single words (booleans),
while other ones expect a value after them. In this case, the value must be
provided immediately after the setting name.

The currently supported settings are the following ones.

accept-proxy
  Enforces the use of the PROXY protocol over any connection accepted by any of
  the sockets declared on the same line. Versions 1 and 2 of the PROXY protocol
  are supported and correctly detected. The PROXY protocol dictates the layer
  3/4 addresses of the incoming connection to be used everywhere an address is
  used, with the only exception of "tcp-request connection" rules which will
  only see the real connection address. Logs will reflect the addresses
  indicated in the protocol, unless it is violated, in which case the real
  address will still be used.  This keyword combined with support from external
  components can be used as an efficient and reliable alternative to the
  X-Forwarded-For mechanism which is not always reliable and not even always
  usable. See also "tcp-request connection expect-proxy" for a finer-grained
  setting of which client is allowed to use the protocol.

alpn <protocols>
  This enables the TLS ALPN extension and advertises the specified protocol
  list as supported on top of ALPN. The protocol list consists in a comma-
```



```
  delimited list of protocol names, for instance: "http/1.1,http/1.0" (without
  quotes). This requires that the SSL library is build with support for TLS
  extensions enabled (check with haproxy -vv). The ALPN extension replaces the
  initial NPN extension.

backlog <backlog>
  Sets the socket's backlog to this value. If unspecified, the frontend's
  backlog is used instead, which generally defaults to the maxconn value.

ecdhe <named curve>
  This setting is only available when support for OpenSSL was built in. It sets
  the named curve (RFC 4492) used to generate ECDH ephemeral keys. By default,
  used named curve is prime256v1.

ca-file <cafile>
  This setting is only available when support for OpenSSL was built in. It
  designates a PEM file from which to load CA certificates used to verify
  client's certificate.

ca-ignore-err [all|<errorID>,...]
  This setting is only available when support for OpenSSL was built in.
  Sets a comma separated list of errorIDs to ignore during verify at depth > 0.
  If set to 'all', all errors are ignored. SSL handshake is not aborted if an
  error is ignored.

ca-sign-file <cafile>
  This setting is only available when support for OpenSSL was built in. It
  designates a PEM file containing both the CA certificate and the CA private
  key used to create and sign server's certificates. This is a mandatory
  setting when the dynamic generation of certificates is enabled. See
  'generate-certificates' for details.

ca-sign-passphrase <passphrase>
  This setting is only available when support for OpenSSL was built in. It is
  the CA private key passphrase. This setting is optional and used only when
  the dynamic generation of certificates is enabled. See
  'generate-certificates' for details.

ciphers <ciphers>
  This setting is only available when support for OpenSSL was built in. It sets
  the string describing the list of cipher algorithms ("cipher suite") that are
  negotiated during the SSL/TLS handshake. The format of the string is defined
  in "man 1 ciphers" from OpenSSL man pages, and can be for instance a string
  such as "AES:ALL:!aNULL:!eNULL:+RC4:@STRENGTH" (without quotes).

crl-file <crlfile>
  This setting is only available when support for OpenSSL was built in. It
  designates a PEM file from which to load certificate revocation list used
  to verify client's certificate.

crt <cert>
  This setting is only available when support for OpenSSL was built in. It
  designates a PEM file containing both the required certificates and any
  associated private keys. This file can be built by concatenating multiple
  PEM files into one (e.g. cat cert.pem key.pem > combined.pem). If your CA
  requires an intermediate certificate, this can also be concatenated into this
  file.

  If the OpenSSL used supports Diffie-Hellman, parameters present in this file
  are loaded.

  If a directory name is used instead of a PEM file, then all files found in
  that directory will be loaded in alphabetic order unless their name ends with
  '.issuer', '.ocsp' or '.sctl' (reserved extensions). This directive may be
  specified multiple times in order to load certificates from multiple files or
```

```
   directories. The certificates will be presented to clients who provide a
   valid TLS Server Name Indication field matching one of their CN or alt
   subjects.  Wildcards are supported, where a wildcard character '*' is used
   instead of the first hostname component (eg: *.example.org matches
   www.example.org but not www.sub.example.org).

   If no SNI is provided by the client or if the SSL library does not support
   TLS extensions, or if the client provides an SNI hostname which does not
   match any certificate, then the first loaded certificate will be presented.
   This means that when loading certificates from a directory, it is highly
   recommended to load the default one first as a file or to ensure that it will
   always be the first one in the directory.

   Note that the same cert may be loaded multiple times without side effects.

   Some CAs (such as Godaddy) offer a drop down list of server types that do not
   include HAProxy when obtaining a certificate. If this happens be sure to
   choose a webserver that the CA believes requires an intermediate CA (for
   Godaddy, selection Apache Tomcat will get the correct bundle, but many
   others, e.g. nginx, result in a wrong bundle that will not work for some
   clients).

   For each PEM file, haproxy checks for the presence of file at the same path
   suffixed by ".ocsp". If such file is found, support for the TLS Certificate
   Status Request extension (also known as "OCSP stapling") is automatically
   enabled. The content of this file is optional. If not empty, it must contain
   a valid OCSP Response in DER format. In order to be valid an OCSP Response
   must comply with the following rules: it has to indicate a good status,
   it has to be a single response for the certificate of the PEM file, and it
   has to be valid at the moment of addition. If these rules are not respected
   the OCSP Response is ignored and a warning is emitted. In order to  identify
   which certificate an OCSP Response applies to, the issuer's certificate is
   necessary. If the issuer's certificate is not found in the PEM file, it will
   be loaded from a file at the same path as the PEM file suffixed by ".issuer"
   if it exists otherwise it will fail with an error.

   For each PEM file, haproxy also checks for the presence of file at the same
   path suffixed by ".sctl". If such file is found, support for Certificate
   Transparency (RFC6962) TLS extension is enabled. The file must contain a
   valid Signed Certificate Timestamp List, as described in RFC. File is parsed
   to check basic syntax, but no signatures are verified.

 crt-ignore-err <errors>
   This setting is only available when support for OpenSSL was built in. Sets a
   comma separated list of errorIDs to ignore during verify at depth == 0.  If
   set to 'all', all errors are ignored. SSL handshake is not aborted if an error
   is ignored.

 crt-list <file>
   This setting is only available when support for OpenSSL was built in. It
   designates a list of PEM file with an optional list of SNI filter per
   certificate, with the following format for each line :

         <crtfile> [[!]<snifilter> ...]

   Wildcards are supported in the SNI filter. Negative filter are also supported,
   only useful in combination with a wildcard filter to exclude a particular SNI.
   The certificates will be presented to clients who provide a valid TLS Server
   Name Indication field matching one of the SNI filters. If no SNI filter is
   specified, the CN and alt subjects are used. This directive may be specified
   multiple times. See the "crt" option for more information. The default
   certificate is still needed to meet OpenSSL expectations. If it is not used,
   the 'strict-sni' option may be used.

 defer-accept
```



```
   Is an optional keyword which is supported only on certain Linux kernels. It
   states that a connection will only be accepted once some data arrive on it,
   or at worst after the first retransmit. This should be used only on protocols
   for which the client talks first (eg: HTTP). It can slightly improve
   performance by ensuring that most of the request is already available when
   the connection is accepted. On the other hand, it will not be able to detect
   connections which don't talk. It is important to note that this option is
   broken in all kernels up to 2.6.31, as the connection is never accepted until
   the client talks. This can cause issues with front firewalls which would see
   an established connection while the proxy will only see it in SYN_RECV. This
   option is only supported on TCPv4/TCPv6 sockets and ignored by other ones.

 force-sslv3
   This option enforces use of SSLv3 only on SSL connections instantiated from
   this listener. SSLv3 is generally less expensive than the TLS counterparts
   for high connection rates. This option is also available on global statement
   "ssl-default-bind-options". See also "no-tlsv*" and "no-sslv3".

 force-tlsv10
   This option enforces use of TLSv1.0 only on SSL connections instantiated from
   this listener. This option is also available on global statement
   "ssl-default-bind-options". See also "no-tlsv*" and "no-sslv3".

 force-tlsv11
   This option enforces use of TLSv1.1 only on SSL connections instantiated from
   this listener. This option is also available on global statement
   "ssl-default-bind-options". See also "no-tlsv*", and "no-sslv3".

 force-tlsv12
   This option enforces use of TLSv1.2 only on SSL connections instantiated from
   this listener. This option is also available on global statement
   "ssl-default-bind-options". See also "no-tlsv*", and "no-sslv3".

 generate-certificates
   This setting is only available when support for OpenSSL was built in. It
   enables the dynamic SSL certificates generation. A CA certificate and its
   private key are necessary (see 'ca-sign-file'). When HAProxy is configured as
   a transparent forward proxy, SSL requests generate errors because of a common
   name mismatch on the certificate presented to the client. With this option
   enabled, HAProxy will try to forge a certificate using the SNI hostname
   indicated by the client. This is done only if no certificate matches the SNI
   hostname (see 'crt-list'). If an error occurs, the default certificate is
   used, else the 'strict-sni' option is set.
   It can also be used when HAProxy is configured as a reverse proxy to ease the
   deployment of an architecture with many backends.

   Creating a SSL certificate is an expensive operation, so a LRU cache is used
   to store forged certificates (see 'tune.ssl.ssl-ctx-cache-size'). It
   increases the HAProxy's memroy footprint to reduce latency when the same
   certificate is used many times.

 gid <gid>
   Sets the group of the UNIX sockets to the designated system gid. It can also
   be set by default in the global section's "unix-bind" statement. Note that
   some platforms simply ignore this. This setting is equivalent to the "group"
   setting except that the group ID is used instead of its name. This setting is
   ignored by non UNIX sockets.

 group <group>
   Sets the group of the UNIX sockets to the designated system group. It can
   also be set by default in the global section's "unix-bind" statement. Note
   that some platforms simply ignore this. This setting is equivalent to the
   "gid" setting except that the group name is used instead of its gid. This
   setting is ignored by non UNIX sockets.
```

```
9751  id <id>
9752    Fixes the socket ID. By default, socket IDs are automatically assigned, but
9753    sometimes it is more convenient to fix them to ease monitoring. This value
9754    must be strictly positive and unique within the listener/frontend. This
9755    option can only be used when defining only a single socket.
9756
9757  interface <interface>
9758    Restricts the socket to a specific interface. When specified, only packets
9759    received from that particular interface are processed by the socket. This is
9760    currently only supported on Linux. The interface must be a primary system
9761    interface, not an aliased interface. It is also possible to bind multiple
9762    frontends to the same address if they are bound to different interfaces. Note
9763    that binding to a network interface requires root privileges. This parameter
9764    is only compatible with TCPv4/TCPv6 sockets.
9765
9766  level <level>
9767    This setting is used with the stats sockets only to restrict the nature of
9768    the commands that can be issued on the socket. It is ignored by other
9769    sockets. <level> can be one of :
9770    - "user" is the least privileged level ; only non-sensitive stats can be
9771      read, and no change is allowed. It would make sense on systems where it
9772      is not easy to restrict access to the socket.
9773    - "operator" is the default level and fits most common uses. All data can
9774      be read, and only non-sensitive changes are permitted (eg: clear max
9775      counters).
9776    - "admin" should be used with care, as everything is permitted (eg: clear
9777      all counters).
9778
9779  maxconn <maxconn>
9780    Limits the sockets to this number of concurrent connections. Extraneous
9781    connections will remain in the system's backlog until a connection is
9782    released. If unspecified, the limit will be the same as the frontend's
9783    maxconn. Note that in case of port ranges or multiple addresses, the same
9784    value will be applied to each socket. This setting enables different
9785    limitations on expensive sockets, for instance SSL entries which may easily
9786    eat all memory.
9787
9788  mode <mode>
9789    Sets the octal mode used to define access permissions on the UNIX socket. It
9790    can also be set by default in the global section's "unix-bind" statement.
9791    Note that some platforms simply ignore this. This setting is ignored by non
9792    UNIX sockets.
9793
9794  mss <maxseg>
9795    Sets the TCP Maximum Segment Size (MSS) value to be advertised on incoming
9796    connections. This can be used to force a lower MSS for certain specific
9797    ports, for instance for connections passing through a VPN. Note that this
9798    relies on a kernel feature which is theoretically supported under Linux but
9799    was buggy in all versions prior to 2.6.28. It may or may not work on other
9800    operating systems. It may also not change the advertised value but change the
9801    effective size of outgoing segments. The commonly advertised value for TCPv4
9802    over Ethernet networks is 1460 = 1500(MTU) - 40(IP+TCP). If this value is
9803    positive, it will be used as the advertised MSS. If it is negative, it will
9804    indicate by how much to reduce the incoming connection's advertised MSS for
9805    outgoing segments. This parameter is only compatible with TCP v4/v6 sockets.
9806
9807  name <name>
9808    Sets an optional name for these sockets, which will be reported on the stats
9809    page.
9810
9811  namespace <name>
9812    On Linux, it is possible to specify which network namespace a socket will
9813    belong to. This directive makes it possible to explicitly bind a listener to
9814    a namespace different from the default one. Please refer to your operating
9815    system's documentation to find more details about network namespaces.
```



```
9816
9817  nice <nice>
9818    Sets the 'niceness' of connections initiated from the socket. Value must be
9819    in the range -1024..1024 inclusive, and defaults to zero. Positive values
9820    means that such connections are more friendly to others and easily offer
9821    their place in the scheduler. On the opposite, negative values mean that
9822    connections want to run with a higher priority than others. The difference
9823    only happens under high loads when the system is close to saturation.
9824    Negative values are appropriate for low-latency or administration services,
9825    and high values are generally recommended for CPU intensive tasks such as SSL
9826    processing or bulk transfers which are less sensible to latency. For example,
9827    it may make sense to use a positive value for an SMTP socket and a negative
9828    one for an RDP socket.
9829
9830  no-sslv3
9831    This setting is only available when support for OpenSSL was built in. It
9832    disables support for SSLv3 on any sockets instantiated from the listener when
9833    SSL is supported. Note that SSLv2 is forced disabled in the code and cannot
9834    be enabled using any configuration option. This option is also available on
9835    global statement "ssl-default-bind-options". See also "force-tls*",
9836    and "force-sslv3".
9837
9838  no-tls-tickets
9839    This setting is only available when support for OpenSSL was built in. It
9840    disables the stateless session resumption (RFC 5077 TLS Ticket
9841    extension) and force to use stateful session resumption. Stateless
9842    session resumption is more expensive in CPU usage. This option is also
9843    available on global statement "ssl-default-bind-options".
9844
9845  no-tlsv10
9846    This setting is only available when support for OpenSSL was built in. It
9847    disables support for TLSv1.0 on any sockets instantiated from the listener
9848    when SSL is supported. Note that SSLv2 is forced disabled in the code and
9849    cannot be enabled using any configuration option. This option is also
9850    available on global statement "ssl-default-bind-options". See also
9851    "force-tlsv*", and "force-sslv3".
9852
9853  no-tlsv11
9854    This setting is only available when support for OpenSSL was built in. It
9855    disables support for TLSv1.1 on any sockets instantiated from the listener
9856    when SSL is supported. Note that SSLv2 is forced disabled in the code and
9857    cannot be enabled using any configuration option. This option is also
9858    available on global statement "ssl-default-bind-options". See also
9859    "force-tlsv*", and "force-sslv3".
9860
9861  no-tlsv12
9862    This setting is only available when support for OpenSSL was built in. It
9863    disables support for TLSv1.2 on any sockets instantiated from the listener
9864    when SSL is supported. Note that SSLv2 is forced disabled in the code and
9865    cannot be enabled using any configuration option. This option is also
9866    available on global statement "ssl-default-bind-options". See also
9867    "force-tlsv*", and "force-sslv3".
9868
9869  npn <protocols>
9870    This enables the NPN TLS extension and advertises the specified protocol list
9871    as supported on top of NPN. The protocol list consists in a comma-delimited
9872    list of protocol names, for instance: "http/1.1,http/1.0" (without quotes).
9873    This requires that the SSL library is build with support for TLS extensions
9874    enabled (check with haproxy -vv). Note that the NPN extension has been
9875    replaced with the ALPN extension (see the "alpn" keyword).
9876
9877  process [ all | odd | even | <number 1-64>[-<number 1-64>] ]
9878    This restricts the list of processes on which this listener is allowed to
9879    run. It does not enforce any process but eliminates those which do not match.
9880    If the frontend uses a "bind-process" setting, the intersection between the
```

```
9881    two is applied. If in the end the listener is not allowed to run on any
9882    remaining process, a warning is emitted, and the listener will either run on
9883    the first process of the listener if a single process was specified, or on
9884    all of its processes if multiple processes were specified. For the unlikely
9885    case where several ranges are needed, this directive may be repeated. The
9886    main purpose of this directive is to be used with the stats sockets and have
9887    one different socket per process. The second purpose is to have multiple bind
9888    lines sharing the same IP:port but not the same process in a listener, so
9889    that the system can distribute the incoming connections into multiple queues
9890    and allow a smoother inter-process load balancing. Currently Linux 3.9 and
9891    above is known for supporting this. See also "bind-process" and "nbproc".
9892
9893  ssl
9894    This setting is only available when support for OpenSSL was built in. It
9895    enables SSL deciphering on connections instantiated from this listener. A
9896    certificate is necessary (see "crt" above). All contents in the buffers will
9897    appear in clear text, so that ACLs and HTTP processing will only have access
9898    to deciphered contents.
9899
9900  strict-sni
9901    This setting is only available when support for OpenSSL was built in. The
9902    SSL/TLS negotiation is allow only if the client provided an SNI which match
9903    a certificate. The default certificate is not used.
9904    See the "crt" option for more information.
9905
9906  tcp-ut <delay>
9907    Sets the TCP User Timeout for all incoming connections instanciated from this
9908    listening socket. This option is available on Linux since version 2.6.37. It
9909    allows haproxy to configure a timeout for sockets which contain data not
9910    receiving an acknoledgement for the configured delay. This is especially
9911    useful on long-lived connections experiencing long idle periods such as
9912    remote terminals or database connection pools, where the client and server
9913    timeouts must remain high to allow a long period of idle, but where it is
9914    important to detect that the client has disappeared in order to release all
9915    resources associated with its connection (and the server's session). The
9916    argument is a delay expressed in milliseconds by default. This only works
9917    for regular TCP connections, and is ignored for other protocols.
9918
9919  tfo
9920    Is an optional keyword which is supported only on Linux kernels >= 3.7. It
9921    enables TCP Fast Open on the listening socket, which means that clients which
9922    support this feature will be able to send a request and receive a response
9923    during the 3-way handshake starting from second connection, thus saving one
9924    round-trip after the first connection. This only makes sense with protocols
9925    that use high connection rates and where each round trip matters. This can
9926    possibly cause issues with many firewalls which do not accept data on SYN
9927    packets, so this option should only be enabled once well tested. This option
9928    is only supported on TCPv4/TCPv6 sockets and ignored by other ones. You may
9929    need to build HAProxy with USE_TFO=1 if your libc doesn't define
9930    TCP_FASTOPEN.
9931
9932  tls-ticket-keys <keyfile>
9933    Sets the TLS ticket keys file to load the keys from. The keys need to be 48
9934    bytes long, encoded with base64 (ex. openssl rand -base64 48). Number of keys
9935    is specified by the TLS_TICKETS_NO build option (default 3) and at least as
9936    many keys need to be present in the file. Last TLS_TICKETS_NO keys will be
9937    used for decryption and the penultimate one for encryption. This enables easy
9938    key rotation by just appending new key to the file and reloading the process.
9939    Keys must be periodically rotated (ex. every 12h) or Perfect Forward Secrecy
9940    is compromised. It is also a good idea to keep the keys off any permanent
9941    storage such as hard drives (hint: use tmpfs and don't swap those files).
9942    Lifetime hint can be changed using tune.ssl.timeout.
9943
9944  transparent
9945    Is an optional keyword which is supported only on certain Linux kernels. It
```



```
9946    indicates that the addresses will be bound even if they do not belong to the
9947    local machine, and that packets targeting any of these addresses will be
9948    intercepted just as if the addresses were locally configured. This normally
9949    requires that IP forwarding is enabled. Caution! do not use this with the
9950    default address '*', as it would redirect any traffic for the specified port.
9951    This keyword is available only when HAProxy is built with USE_LINUX_TPROXY=1.
9952    This parameter is only compatible with TCPv4 and TCPv6 sockets, depending on
9953    kernel version. Some distribution kernels include backports of the feature,
9954    so check for support with your vendor.
9955
9956  v4v6
9957    Is an optional keyword which is supported only on most recent systems
9958    including Linux kernels >= 2.4.21. It is used to bind a socket to both IPv4
9959    and IPv6 when it uses the default address. Doing so is sometimes necessary
9960    on systems which bind to IPv6 only by default. It has no effect on non-IPv6
9961    sockets, and is overridden by the "v6only" option.
9962
9963  v6only
9964    Is an optional keyword which is supported only on most recent systems
9965    including Linux kernels >= 2.4.21. It is used to bind a socket to IPv6 only
9966    when it uses the default address. Doing so is sometimes preferred to doing it
9967    system-wide as it is per-listener. It has no effect on non-IPv6 sockets and
9968    has precedence over the "v4v6" option.
9969
9970  uid <uid>
9971    Sets the owner of the UNIX sockets to the designated system uid. It can also
9972    be set by default in the global section's "unix-bind" statement. Note that
9973    some platforms simply ignore this. This setting is equivalent to the "user"
9974    setting except that the user numeric ID is used instead of its name. This
9975    setting is ignored by non UNIX sockets.
9976
9977  user <user>
9978    Sets the owner of the UNIX sockets to the designated system user. It can also
9979    be set by default in the global section's "unix-bind" statement. Note that
9980    some platforms simply ignore this. This setting is equivalent to the "uid"
9981    setting except that the user name is used instead of its uid. This setting is
9982    ignored by non UNIX sockets.
9983
9984  verify [none|optional|required]
9985    This setting is only available when support for OpenSSL was built in. If set
9986    to 'none', client certificate is not requested. This is the default. In other
9987    cases, a client certificate is requested. If the client does not provide a
9988    certificate after the request and if 'verify' is set to 'required', then the
9989    handshake is aborted, while it would have succeeded if set to 'optional'. The
9990    certificate provided by the client is always verified using CAs from
9991    'ca-file' and optional CRLs from 'crl-file'. On verify failure the handshake
9992    is aborted, regardless of the 'verify' option, unless the error code exactly
9993    matches one of those listed with 'ca-ignore-err' or 'crt-ignore-err'.
9994
9995  5.2. Server and default-server options
9996  ------------------------------------
9997
9998  The "server" and "default-server" keywords support a certain number of settings
9999  which are all passed as arguments on the server line. The order in which those
10000 arguments appear does not count, and they are all optional. Some of those
10001 settings are single words (booleans) while others expect one or several values
10002 after them. In this case, the values must immediately follow the setting name.
10003 Except default-server, all those settings must be specified after the server's
10004 address if they are used:
10005
10006   server <name> <address>[:port] [settings ...]
10007   default-server [settings ...]
10008
10009 The currently supported settings are the following ones.
10010
```

```
addr <ipv4|ipv6>
  Using the "addr" parameter, it becomes possible to use a different IP address
  to send health-checks. On some servers, it may be desirable to dedicate an IP
  address to specific component able to perform complex tests which are more
  suitable to health-checks than the application. This parameter is ignored if
  the "check" parameter is not set. See also the "port" parameter.

  Supported in default-server: No

agent-check
  Enable an auxiliary agent check which is run independently of a regular
  health check. An agent health check is performed by making a TCP connection
  to the port set by the "agent-port" parameter and reading an ASCII string.
  The string is made of a series of words delimited by spaces, tabs or commas
  in any order, optionally terminated by '\r' and/or '\n', each consisting of :

  - An ASCII representation of a positive integer percentage, e.g. "75%".
    Values in this format will set the weight proportional to the initial
    weight of a server as configured when haproxy starts. Note that a zero
    weight is reported on the stats page as "DRAIN" since it has the same
    effect on the server (it's removed from the LB farm).

  - The word "ready". This will turn the server's administrative state to the
    READY mode, thus cancelling any DRAIN or MAINT state

  - The word "drain". This will turn the server's administrative state to the
    DRAIN mode, thus it will not accept any new connections other than those
    that are accepted via persistence.

  - The word "maint". This will turn the server's administrative state to the
    MAINT mode, thus it will not accept any new connections at all, and health
    checks will be stopped.

  - The words "down", "failed", or "stopped", optionally followed by a
    description string after a sharp ('#'). All of these mark the server's
    operating state as DOWN, but since the word itself is reported on the stats
    page, the difference allows an administrator to know if the situation was
    expected or not : the service may intentionally be stopped, may appear up
    but fail some validity tests, or may be seen as down (eg: missing process,
    or port not responding).

  - The word "up" sets back the server's operating state as UP if health checks
    also report that the service is accessible.

  Parameters which are not advertised by the agent are not changed. For
  example, an agent might be designed to monitor CPU usage and only report a
  relative weight and never interact with the operating status. Similarly, an
  agent could be designed as an end-user interface with 3 radio buttons
  allowing an administrator to change only the administrative state. However,
  it is important to consider that only the agent may revert its own actions,
  so if a server is set to DRAIN mode or to DOWN state using the agent, the
  agent must implement the other equivalent actions to bring the service into
  operations again.

  Failure to connect to the agent is not considered an error as connectivity
  is tested by the regular health check which is enabled by the "check"
  parameter. Warning though, it is not a good idea to stop an agent after it
  reports "down", since only an agent reporting "up" will be able to turn the
  server up again. Note that the CLI on the Unix stats socket is also able to
  force an agent's result in order to workaround a bogus agent if needed.

  Requires the "agent-port" parameter to be set. See also the "agent-inter"
  parameter.

  Supported in default-server: No
```



```

agent-inter <delay>
  The "agent-inter" parameter sets the interval between two agent checks
  to <delay> milliseconds. If left unspecified, the delay defaults to 2000 ms.

  Just as with every other time-based parameter, it may be entered in any
  other explicit unit among { us, ms, s, m, h, d }. The "agent-inter"
  parameter also serves as a timeout for agent checks "timeout check" is
  not set. In order to reduce "resonance" effects when multiple servers are
  hosted on the same hardware, the agent and health checks of all servers
  are started with a small time offset between them. It is also possible to
  add some random noise in the agent and health checks interval using the
  global "spread-checks" keyword. This makes sense for instance when a lot
  of backends use the same servers.

  See also the "agent-check" and "agent-port" parameters.

  Supported in default-server: Yes

agent-port <port>
  The "agent-port" parameter sets the TCP port used for agent checks.

  See also the "agent-check" and "agent-inter" parameters.

  Supported in default-server: Yes

backup
  When "backup" is present on a server line, the server is only used in load
  balancing when all other non-backup servers are unavailable. Requests coming
  with a persistence cookie referencing the server will always be served
  though. By default, only the first operational backup server is used, unless
  the "allbackups" option is set in the backend. See also the "allbackups"
  option.

  Supported in default-server: No

ca-file <cafile>
  This setting is only available when support for OpenSSL was built in. It
  designates a PEM file from which to load CA certificates used to verify
  server's certificate.

  Supported in default-server: No

check
  This option enables health checks on the server. By default, a server is
  always considered available. If "check" is set, the server is available when
  accepting periodic TCP connections, to ensure that it is really able to serve
  requests. The default address and port to send the tests to are those of the
  server, and the default source is the same as the one defined in the
  backend. It is possible to change the address using the "addr" parameter, the
  port using the "port" parameter, the source address using the "source"
  address, and the interval and timers using the "inter", "rise" and "fall"
  parameters. The request method is define in the backend using the "httpchk",
  "smtpchk", "mysql-check", "pgsql-check" and "ssl-hello-chk" options. Please
  refer to those options and parameters for more information.

  Supported in default-server: No

check-send-proxy
  This option forces emission of a PROXY protocol line with outgoing health
  checks, regardless of whether the server uses send-proxy or not for the
  normal traffic. By default, the PROXY protocol is enabled for health checks
  if it is already enabled for normal traffic and if no "port" nor "addr"
  directive is present. However, if such a directive is present, the
  "check-send-proxy" option needs to be used to force the use of the
```

protocol. See also the "send-proxy" option for more information.

Supported in default-server: No

check-ssl
This option forces encryption of all health checks over SSL, regardless of whether the server uses SSL or not for the normal traffic. This is generally used when an explicit "port" or "addr" directive is specified and SSL health checks are not inherited. It is important to understand that this option inserts an SSL transport layer below the checks, so that a simple TCP connect check becomes an SSL connect, which replaces the old ssl-hello-chk. The most common use is to send HTTPS checks by combining "httpchk" with SSL checks. All SSL settings are common to health checks and traffic (eg: ciphers). See the "ssl" option for more information.

Supported in default-server: No

ciphers <ciphers>
This option sets the string describing the list of cipher algorithms that is is negotiated during the SSL/TLS handshake with the server. The format of the string is defined in "man 1 ciphers". When SSL is used to communicate with servers on the local network, it is common to see a weaker set of algorithms than what is used over the internet. Doing so reduces CPU usage on both the server and haproxy while still keeping it compatible with deployed software. Some algorithms such as RC4-SHA1 are reasonably cheap. If no security at all is needed and just connectivity, using DES can be appropriate.

Supported in default-server: No

cookie <value>
The "cookie" parameter sets the cookie value assigned to the server to <value>. This value will be checked in incoming requests, and the first operational server possessing the same value will be selected. In return, in cookie insertion or rewrite modes, this value will be assigned to the cookie sent to the client. There is nothing wrong in having several servers sharing the same cookie value, and it is in fact somewhat common between normal and backup servers. See also the "cookie" keyword in backend section.

Supported in default-server: No

crl-file <crlfile>
This setting is only available when support for OpenSSL was built in. It designates a PEM file from which to load certificate revocation list used to verify server's certificate.

Supported in default-server: No

crt <cert>
This setting is only available when support for OpenSSL was built in. It designates a PEM file from which to load both a certificate and the associated private key. This file can be built by concatenating both PEM files into one. This certificate will be sent if the server send a client certificate request.

Supported in default-server: No

disabled
The "disabled" keyword starts the server in the "disabled" state. That means that it is marked down in maintenance mode, and no connection other than the ones allowed by persist mode will reach it. It is very well suited to setup new servers, because normal traffic will never reach them, while it is still possible to test the service by making use of the force-persist mechanism.

Supported in default-server: No

error-limit <count>
If health observing is enabled, the "error-limit" parameter specifies the number of consecutive errors that triggers event selected by the "on-error" option. By default it is set to 10 consecutive errors.

Supported in default-server: Yes

See also the "check", "error-limit" and "on-error".

fall <count>
The "fall" parameter states that a server will be considered as dead after <count> consecutive unsuccessful health checks. This value defaults to 3 if unspecified. See also the "check", "inter" and "rise" parameters.

Supported in default-server: Yes

force-sslv3
This option enforces use of SSLv3 only when SSL is used to communicate with the server. SSLv3 is generally less expensive than the TLS counterparts for high connection rates. This option is also available on global statement "ssl-default-server-options". See also "no-tlsv*", "no-sslv3".

Supported in default-server: No

force-tlsv10
This option enforces use of TLSv1.0 only when SSL is used to communicate with the server. This option is also available on global statement "ssl-default-server-options". See also "no-tlsv*", "no-sslv3".

Supported in default-server: No

force-tlsv11
This option enforces use of TLSv1.1 only when SSL is used to communicate with the server. This option is also available on global statement "ssl-default-server-options". See also "no-tlsv*", "no-sslv3".

Supported in default-server: No

force-tlsv12
This option enforces use of TLSv1.2 only when SSL is used to communicate with the server. This option is also available on global statement "ssl-default-server-options". See also "no-tlsv*", "no-sslv3".

Supported in default-server: No

id <value>
Set a persistent ID for the server. This ID must be positive and unique for the proxy. An unused ID will automatically be assigned if unset. The first assigned value will be 1. This ID is currently only returned in statistics.

Supported in default-server: No

inter <delay>
fastinter <delay>
downinter <delay>
The "inter" parameter sets the interval between two consecutive health checks to <delay> milliseconds. If left unspecified, the delay defaults to 2000 ms. It is also possible to use "fastinter" and "downinter" to optimize delays between checks depending on the server state :

Server state	Interval used
UP 100% (non-transitional)	"inter"
Transitionally UP (going down "fall"),	"fastinter" if set,

```
      Transitionally DOWN (going up "rise"), | "inter" otherwise.
      or yet unchecked.                      |
     ----------------------------------------+---------------------------------
      DOWN 100% (non-transitional)           | "downinter" if set,
                                             | "inter" otherwise.
     ----------------------------------------+---------------------------------

   Just as with every other time-based parameter, they can be entered in any
   other explicit unit among { us, ms, s, m, h, d }. The "inter" parameter also
   serves as a timeout for health checks sent to servers if "timeout check" is
   not set. In order to reduce "resonance" effects when multiple servers are
   hosted on the same hardware, the agent and health checks of all servers
   are started with a small time offset between them. It is also possible to
   add some random noise in the agent and health checks interval using the
   global "spread-checks" keyword. This makes sense for instance when a lot
   of backends use the same servers.

   Supported in default-server: Yes

maxconn <maxconn>
   The "maxconn" parameter specifies the maximal number of concurrent
   connections that will be sent to this server. If the number of incoming
   concurrent requests goes higher than this value, they will be queued, waiting
   for a connection to be released. This parameter is very important as it can
   save fragile servers from going down under extreme loads. If a "minconn"
   parameter is specified, the limit becomes dynamic. The default value is "0"
   which means unlimited. See also the "minconn" and "maxqueue" parameters, and
   the backend's "fullconn" keyword.

   Supported in default-server: Yes

maxqueue <maxqueue>
   The "maxqueue" parameter specifies the maximal number of connections which
   will wait in the queue for this server. If this limit is reached, next
   requests will be redispatched to other servers instead of indefinitely
   waiting to be served. This will break persistence but may allow people to
   quickly re-log in when the server they try to connect to is dying. The
   default value is "0" which means the queue is unlimited. See also the
   "maxconn" and "minconn" parameters.

   Supported in default-server: Yes

minconn <minconn>
   When the "minconn" parameter is set, the maxconn limit becomes a dynamic
   limit following the backend's load. The server will always accept at least
   <minconn> connections, never more than <maxconn>, and the limit will be on
   the ramp between both values when the backend has less than <fullconn>
   concurrent connections. This makes it possible to limit the load on the
   server during normal loads, but push it further for important loads without
   overloading the server during exceptional loads. See also the "maxconn"
   and "maxqueue" parameters, as well as the "fullconn" backend keyword.

   Supported in default-server: Yes

namespace <name>
   On Linux, it is possible to specify which network namespace a socket will
   belong to. This directive makes it possible to explicitly bind a server to
   a namespace different from the default one. Please refer to your operating
   system's documentation to find more details about network namespaces.

no-ssl-reuse
   This option disables SSL session reuse when SSL is used to communicate with
   the server. It will force the server to perform a full handshake for every
   new connection. It's probably only useful for benchmarking, troubleshooting,
   and for paranoid users.
```



```

   Supported in default-server: No

no-sslv3
   This option disables support for SSLv3 when SSL is used to communicate with
   the server. Note that SSLv2 is disabled in the code and cannot be enabled
   using any configuration option. See also "force-sslv3", "force-tlsv*".

   Supported in default-server: No

no-tls-tickets
   This setting is only available when support for OpenSSL was built in. It
   disables the stateless session resumption (RFC 5077 TLS Ticket
   extension) and force to use stateful session resumption. Stateless
   session resumption is more expensive in CPU usage for servers. This option
   is also available on global statement "ssl-default-server-options".

   Supported in default-server: No

no-tlsv10
   This option disables support for TLSv1.0 when SSL is used to communicate with
   the server. Note that SSLv2 is disabled in the code and cannot be enabled
   using any configuration option. TLSv1 is more expensive than SSLv3 so it
   often makes sense to disable it when communicating with local servers. This
   option is also available on global statement "ssl-default-server-options".
   See also "force-sslv3", "force-tlsv*".

   Supported in default-server: No

no-tlsv11
   This option disables support for TLSv1.1 when SSL is used to communicate with
   the server. Note that SSLv2 is disabled in the code and cannot be enabled
   using any configuration option. TLSv1 is more expensive than SSLv3 so it
   often makes sense to disable it when communicating with local servers. This
   option is also available on global statement "ssl-default-server-options".
   See also "force-sslv3", "force-tlsv*".

   Supported in default-server: No

no-tlsv12
   This option disables support for TLSv1.2 when SSL is used to communicate with
   the server. Note that SSLv2 is disabled in the code and cannot be enabled
   using any configuration option. TLSv1 is more expensive than SSLv3 so it
   often makes sense to disable it when communicating with local servers. This
   option is also available on global statement "ssl-default-server-options".
   See also "force-sslv3", "force-tlsv*".

   Supported in default-server: No

non-stick
   Never add connections allocated to this sever to a stick-table.
   This may be used in conjunction with backup to ensure that
   stick-table persistence is disabled for backup servers.

   Supported in default-server: No

observe <mode>
   This option enables health adjusting based on observing communication with
   the server. By default this functionality is disabled and enabling it also
   requires to enable health checks. There are two supported modes: "layer4" and
   "layer7". In layer4 mode, only successful/unsuccessful tcp connections are
   significant. In layer7, which is only allowed for http proxies, responses
   received from server are verified, like valid/wrong http code, unparsable
   headers, a timeout, etc. Valid status codes include 100 to 499, 501 and 505.

```

```
10401     Supported in default-server: No
10402
10403     See also the "check", "on-error" and "error-limit".
10404
10405   on-error <mode>
10406     Select what should happen when enough consecutive errors are detected.
10407     Currently, four modes are available:
10408     - fastinter: force fastinter
10409     - fail-check: simulate a failed check, also forces fastinter (default)
10410     - sudden-death: simulate a pre-fatal failed health check, one more failed
10411       check will mark a server down, forces fastinter
10412     - mark-down: mark the server immediately down and force fastinter
10413
10414     Supported in default-server: Yes
10415
10416     See also the "check", "observe" and "error-limit".
10417
10418   on-marked-down <action>
10419     Modify what occurs when a server is marked down.
10420     Currently one action is available:
10421     - shutdown-sessions: Shutdown peer sessions. When this setting is enabled,
10422       all connections to the server are immediately terminated when the server
10423       goes down. It might be used if the health check detects more complex cases
10424       than a simple connection status, and long timeouts would cause the service
10425       to remain unresponsive for too long a time. For instance, a health check
10426       might detect that a database is stuck and that there's no chance to reuse
10427       existing connections anymore. Connections killed this way are logged with
10428       a 'D' termination code (for "Down").
10429
10430     Actions are disabled by default
10431
10432     Supported in default-server: Yes
10433
10434   on-marked-up <action>
10435     Modify what occurs when a server is marked up.
10436     Currently one action is available:
10437     - shutdown-backup-sessions: Shutdown sessions on all backup servers. This is
10438       done only if the server is not in backup state and if it is not disabled
10439       (it must have an effective weight > 0). This can be used sometimes to force
10440       an active server to take all the traffic back after recovery when dealing
10441       with long sessions (eg: LDAP, SQL, ...). Doing this can cause more trouble
10442       than it tries to solve (eg: incomplete transactions), so use this feature
10443       with extreme care. Sessions killed because a server comes up are logged
10444       with an 'U' termination code (for "Up").
10445
10446     Actions are disabled by default
10447
10448     Supported in default-server: Yes
10449
10450   port <port>
10451     Using the "port" parameter, it becomes possible to use a different port to
10452     send health-checks. On some servers, it may be desirable to dedicate a port
10453     to a specific component able to perform complex tests which are more suitable
10454     to health-checks than the application. It is common to run a simple script in
10455     inetd for instance. This parameter is ignored if the "check" parameter is not
10456     set. See also the "addr" parameter.
10457
10458     Supported in default-server: Yes
10459
10460   redir <prefix>
10461     The "redir" parameter enables the redirection mode for all GET and HEAD
10462     requests addressing this server. This means that instead of having HAProxy
10463     forward the request to the server, it will send an "HTTP 302" response with
10464     the "Location" header composed of this prefix immediately followed by the
10465     requested URI beginning at the leading '/' of the path component. That means
```



```
10466     that no trailing slash should be used after <prefix>. All invalid requests
10467     will be rejected, and all non-GET or HEAD requests will be normally served by
10468     the server. Note that since the response is completely forged, no header
10469     mangling nor cookie insertion is possible in the response. However, cookies in
10470     requests are still analysed, making this solution completely usable to direct
10471     users to a remote location in case of local disaster. Main use consists in
10472     increasing bandwidth for static servers by having the clients directly
10473     connect to them. Note: never use a relative location here, it would cause a
10474     loop between the client and HAProxy!
10475
10476     Example :  server srv1 192.168.1.1:80 redir http://image1.mydomain.com check
10477
10478     Supported in default-server: No
10479
10480   rise <count>
10481     The "rise" parameter states that a server will be considered as operational
10482     after <count> consecutive successful health checks. This value defaults to 2
10483     if unspecified. See also the "check", "inter" and "fall" parameters.
10484
10485     Supported in default-server: Yes
10486
10487   resolve-prefer <family>
10488     When DNS resolution is enabled for a server and multiple IP addresses from
10489     different families are returned, HAProxy will prefer using an IP address
10490     from the family mentioned in the "resolve-prefer" parameter.
10491     Available families: "ipv4" and "ipv6"
10492
10493     Default value: ipv6
10494
10495     Supported in default-server: Yes
10496
10497     Example: server s1 app1.domain.com:80 resolvers mydns resolve-prefer ipv6
10498
10499   resolvers <id>
10500     Points to an existing "resolvers" section to resolve current server's
10501     hostname.
10502     In order to be operational, DNS resolution requires that health check is
10503     enabled on the server. Actually, health checks triggers the DNS resolution.
10504     You must precise one 'resolvers' parameter on each server line where DNS
10505     resolution is required.
10506
10507     Supported in default-server: No
10508
10509     Example: server s1 app1.domain.com:80 check resolvers mydns
10510
10511     See also chapter 5.3
10512
10513   send-proxy
10514     The "send-proxy" parameter enforces use of the PROXY protocol over any
10515     connection established to this server. The PROXY protocol informs the other
10516     end about the layer 3/4 addresses of the incoming connection, so that it can
10517     know the client's address or the public address it accessed to, whatever the
10518     upper layer protocol. For connections accepted by an "accept-proxy" listener,
10519     the advertised address will be used. Only TCPv4 and TCPv6 address families
10520     are supported. Other families such as Unix sockets, will report an UNKNOWN
10521     family. Servers using this option can fully be chained to another instance of
10522     haproxy listening with an "accept-proxy" setting. This setting must not be
10523     used if the server isn't aware of the protocol. When health checks are sent
10524     to the server, the PROXY protocol is automatically used when this option is
10525     set, unless there is an explicit "port" or "addr" directive, in which case an
10526     explicit "check-send-proxy" directive would also be needed to use the PROXY
10527     protocol. See also the "accept-proxy" option of the "bind" keyword.
10528
10529     Supported in default-server: No
10530
```

```
10531  send-proxy-v2
10532    The "send-proxy-v2" parameter enforces use of the PROXY protocol version 2
10533    over any connection established to this server. The PROXY protocol informs
10534    the other end about the layer 3/4 addresses of the incoming connection, so
10535    that it can know the client's address or the public address it accessed to,
10536    whatever the upper layer protocol. This setting must not be used if the
10537    server isn't aware of this version of the protocol. See also the "send-proxy"
10538    option of the "bind" keyword.
10539
10540    Supported in default-server: No
10541
10542  send-proxy-v2-ssl
10543    The "send-proxy-v2-ssl" parameter enforces use of the PROXY protocol version
10544    2 over any connection established to this server. The PROXY protocol informs
10545    the other end about the layer 3/4 addresses of the incoming connection, so
10546    that it can know the client's address or the public address it accessed to,
10547    whatever the upper layer protocol. In addition, the SSL information extension
10548    of the PROXY protocol is added to the PROXY protocol header. This setting
10549    must not be used if the server isn't aware of this version of the protocol.
10550    See also the "send-proxy-v2" option of the "bind" keyword.
10551
10552    Supported in default-server: No
10553
10554  send-proxy-v2-ssl-cn
10555    The "send-proxy-v2-ssl" parameter enforces use of the PROXY protocol version
10556    2 over any connection established to this server. The PROXY protocol informs
10557    the other end about the layer 3/4 addresses of the incoming connection, so
10558    that it can know the client's address or the public address it accessed to,
10559    whatever the upper layer protocol. In addition, the SSL information extension
10560    of the PROXY protocol, along along with the Common Name from the subject of
10561    the client certificate (if any), is added to the PROXY protocol header. This
10562    setting must not be used if the server isn't aware of this version of the
10563    protocol. See also the "send-proxy-v2" option of the "bind" keyword.
10564
10565    Supported in default-server: No
10566
10567  slowstart <start_time_in_ms>
10568    The "slowstart" parameter for a server accepts a value in milliseconds which
10569    indicates after how long a server which has just come back up will run at
10570    full speed. Just as with every other time-based parameter, it can be entered
10571    in any other explicit unit among { us, ms, s, m, h, d }. The speed grows
10572    linearly from 0 to 100% during this time. The limitation applies to two
10573    parameters :
10574
10575    - maxconn: the number of connections accepted by the server will grow from 1
10576      to 100% of the usual dynamic limit defined by (minconn,maxconn,fullconn).
10577
10578    - weight: when the backend uses a dynamic weighted algorithm, the weight
10579      grows linearly from 1 to 100%. In this case, the weight is updated at every
10580      health-check. For this reason, it is important that the "inter" parameter
10581      is smaller than the "slowstart", in order to maximize the number of steps.
10582
10583    The slowstart never applies when haproxy starts, otherwise it would cause
10584    trouble to running servers. It only applies when a server has been previously
10585    seen as failed.
10586
10587    Supported in default-server: Yes
10588
10589  sni <expression>
10590    The "sni" parameter evaluates the sample fetch expression, converts it to a
10591    string and uses the result as the host name sent in the SNI TLS extension to
10592    the server. A typical use case is to send the SNI received from the client in
10593    a bridged HTTPS scenario, using the "ssl_fc_sni" sample fetch for the
10594    expression, though alternatives such as req.hdr(host) can also make sense.
10595
```



```
10596    Supported in default-server: no
10597
10598  source <addr>[:<pl>[-<ph>]] [usesrc { <addr2>[:<port2>] | client | clientip } ]
10599  source <addr>[:<port>] [usesrc { <addr2>[:<port2>] | hdr_ip(<hdr>[,<occ>]) } ]
10600  source <addr>[:<pl>[-<ph>]] [interface <name>] ...
10601    The "source" parameter sets the source address which will be used when
10602    connecting to the server. It follows the exact same parameters and principle
10603    as the backend "source" keyword, except that it only applies to the server
10604    referencing it. Please consult the "source" keyword for details.
10605
10606    Additionally, the "source" statement on a server line allows one to specify a
10607    source port range by indicating the lower and higher bounds delimited by a
10608    dash ('-'). Some operating systems might require a valid IP address when a
10609    source port range is specified. It is permitted to have the same IP/range for
10610    several servers. Doing so makes it possible to bypass the maximum of 64k
10611    total concurrent connections. The limit will then reach 64k connections per
10612    server.
10613
10614    Supported in default-server: No
10615
10616  ssl
10617    This option enables SSL ciphering on outgoing connections to the server. It
10618    is critical to verify server certificates using "verify" when using SSL to
10619    connect to servers, otherwise the communication is prone to trivial man in
10620    the-middle attacks rendering SSL useless. When this option is used, health
10621    checks are automatically sent in SSL too unless there is a "port" or an
10622    "addr" directive indicating the check should be sent to a different location.
10623    See the "check-ssl" option to force SSL health checks.
10624
10625    Supported in default-server: No
10626
10627  tcp-ut <delay>
10628    Sets the TCP User Timeout for all outgoing connections to this server. This
10629    option is available on Linux since version 2.6.37. It allows haproxy to
10630    configure a timeout for sockets which contain data not receiving an
10631    ackноledgement for the configured delay. This is especially useful on
10632    long-lived connections experiencing long idle periods such as remote
10633    terminals or database connection pools, where the client and server timeouts
10634    must remain high to allow a long period of idle, but where it is important to
10635    detect that the server has disappeared in order to release all resources
10636    associated with its connection (and the client's session). One typical use
10637    case is also to force dead server connections to die when health checks are
10638    too slow or during a soft reload since health checks are then disabled. The
10639    argument is a delay expressed in milliseconds by default. This only works for
10640    regular TCP connections, and is ignored for other protocols.
10641
10642  track [<proxy>/]<server>
10643    This option enables ability to set the current state of the server by tracking
10644    another one. It is possible to track a server which itself tracks another
10645    server, provided that at the end of the chain, a server has health checks
10646    enabled. If <proxy> is omitted the current one is used. If disable-on-404 is
10647    used, it has to be enabled on both proxies.
10648
10649    Supported in default-server: No
10650
10651  verify [none|required]
10652    This setting is only available when support for OpenSSL was built in. If set
10653    to 'none', server certificate is not verified. In the other case, The
10654    certificate provided by the server is verified using CAs from 'ca-file'
10655    and optional CRLs from 'crl-file'. If 'ssl_server_verify' is not specified
10656    in global  section, this is the default. On verify failure the handshake
10657    is aborted. It is critically important to verify server certificates when
10658    using SSL to connect to servers, otherwise the communication is prone to
10659    trivial man-in-the-middle attacks rendering SSL totally useless.
10660
```

```
    Supported in default-server: No

verifyhost <hostname>
  This setting is only available when support for OpenSSL was built in, and
  only takes effect if 'verify required' is also specified. When set, the
  hostnames in the subject and subjectAlternateNames of the certificate
  provided by the server are checked. If none of the hostnames in the
  certificate match the specified hostname, the handshake is aborted. The
  hostnames in the server-provided certificate may include wildcards.

  Supported in default-server: No

weight <weight>
  The "weight" parameter is used to adjust the server's weight relative to
  other servers. All servers will receive a load proportional to their weight
  relative to the sum of all weights, so the higher the weight, the higher the
  load. The default weight is 1, and the maximal value is 256. A value of 0
  means the server will not participate in load-balancing but will still accept
  persistent connections. If this parameter is used to distribute the load
  according to server's capacity, it is recommended to start with values which
  can both grow and shrink, for instance between 10 and 100 to leave enough
  room above and below for later adjustments.

  Supported in default-server: Yes


5.3. Server IP address resolution using DNS
-------------------------------------------

HAProxy allows using a host name on the server line to retrieve its IP address
using name servers. By default, HAProxy resolves the name when parsing the
configuration file, at startup and cache the result for the process' life.
This is not sufficient in some cases, such as in Amazon where a server's IP
can change after a reboot or an ELB Virtual IP can change based on current
workload.
This chapter describes how HAProxy can be configured to process server's name
resolution at run time.
Whether run time server name resolution has been enable or not, HAProxy will
carry on doing the first resolution when parsing the configuration.

Bear in mind that DNS resolution is triggered by health checks. This makes
health checks mandatory to allow DNS resolution.


5.3.1. Global overview
----------------------

As we've seen in introduction, name resolution in HAProxy occurs at two
different steps of the process life:

  1. when starting up, HAProxy parses the server line definition and matches a
     host name. It uses libc functions to get the host name resolved. This
     resolution relies on /etc/resolv.conf file.

  2. at run time, when HAProxy gets prepared to run a health check on a server,
     it verifies if the current name resolution is still considered as valid.
     If not, it processes a new resolution, in parallel of the health check.

A few other events can trigger a name resolution at run time:
  - when a server's health check ends up in a connection timeout: this may be
    because the server has a new IP address. So we need to trigger a name
    resolution to know this new IP.

A few things important to notice:
  - all the name servers are queried in the mean time. HAProxy will process the
```



```
    first valid response.

  - a resolution is considered as invalid (NX, timeout, refused), when all the
    servers return an error.


5.3.2. The resolvers section
----------------------------

This section is dedicated to host information related to name resolution in
HAProxy.
There can be as many as resolvers section as needed. Each section can contain
many name servers.

When multiple name servers are configured in a resolvers section, then HAProxy
uses the first valid response. In case of invalid responses, only the last one
is treated. Purpose is to give the chance to a slow server to deliver a valid
answer after a fast faulty or outdated server.

When each server returns a different error type, then only the last error is
used by HAProxy to decide what type of behavior to apply.

Two types of behavior can be applied:
 1. stop DNS resolution
 2. replay the DNS query with a new query type
      In such case, the following types are applied in this exact order:
        1. ANY query type
        2. query type corresponding to family pointed by resolve-prefer
           server's parameter
        3. remaining family type

HAProxy stops DNS resolution when the following errors occur:
 - invalid DNS response packet
 - wrong name in the query section of the response
 - NX domain
 - Query refused by server
 - CNAME not pointing to an IP address

HAProxy tries a new query type when the following errors occur:
 - no Answer records in the response
 - DNS response truncated
 - Error in DNS response
 - No expected DNS records found in the response
 - name server timeout

For example, with 2 name servers configured in a resolvers section:
 - first response is valid and is applied directly, second response is ignored
 - first response is invalid and second one is valid, then second response is
   applied;
 - first response is a NX domain and second one a truncated response, then
   HAProxy replays the query with a new type;
 - first response is truncated and second one is a NX Domain, then HAProxy
   stops resolution.


resolvers <resolvers id>
  Creates a new name server list labelled <resolvers id>

A resolvers section accept the following parameters:

nameserver <id> <ip>:<port>
  DNS server description:
    <id>   : label of the server, should be unique
    <ip>   : IP address of the server
    <port> : port where the DNS service actually runs
```

```
10791
10792  hold <status> <period>
10793    Defines <period> during which the last name resolution should be kept based
10794    on last resolution <status>
10795      <status> : last name resolution status. Only "valid" is accepted for now.
10796      <period> : interval between two successive name resolution when the last
10797                 answer was in <status>. It follows the HAProxy time format.
10798                 <period> is in milliseconds by default.
10799
10800    Default value is 10s for "valid".
10801
10802    Note: since the name resolution is triggered by the health checks, a new
10803          resolution is triggered after <period> modulo the <inter> parameter of
10804          the healch check.
10805
10806  resolve_retries <nb>
10807    Defines the number <nb> of queries to send to resolve a server name before
10808    giving up.
10809    Default value: 3
10810
10811    A retry occurs on name server timeout or when the full sequence of DNS query
10812    type failover is over and we need to start up from the default ANY query
10813    type.
10814
10815  timeout <event> <time>
10816    Defines timeouts related to name resolution
10817       <event> : the event on which the <time> timeout period applies to.
10818                 events available are:
10819                 - retry: time between two DNS queries, when no response have
10820                          been received.
10821                          Default value: 1s
10822       <time>  : time related to the event. It follows the HAProxy time format.
10823                 <time> is expressed in milliseconds.
10824
10825  Example of a resolvers section (with default values):
10826
10827     resolvers mydns
10828       nameserver dns1 10.0.0.1:53
10829       nameserver dns2 10.0.0.2:53
10830       resolve_retries       3
10831       timeout retry         1s
10832       hold valid           10s
10833
10834
10835  6. HTTP header manipulation
10836  ---------------------------
10837
10838  In HTTP mode, it is possible to rewrite, add or delete some of the request and
10839  response headers based on regular expressions. It is also possible to block a
10840  request or a response if a particular header matches a regular expression,
10841  which is enough to stop most elementary protocol attacks, and to protect
10842  against information leak from the internal network.
10843
10844  If HAProxy encounters an "Informational Response" (status code 1xx), it is able
10845  to process all rsp* rules which can allow, deny, rewrite or delete a header,
10846  but it will refuse to add a header to any such messages as this is not
10847  HTTP-compliant. The reason for still processing headers in such responses is to
10848  stop and/or fix any possible information leak which may happen, for instance
10849  because another downstream equipment would unconditionally add a header, or if
10850  a server name appears there. When such messages are seen, normal processing
10851  still occurs on the next non-informational messages.
10852
10853  This section covers common usage of the following keywords, described in detail
10854  in section 4.2 :
10855
```



```
10856     - reqadd     <string>
10857     - reqallow   <search>
10858     - reqiallow  <search>
10859     - reqdel     <search>
10860     - reqidel    <search>
10861     - reqdeny    <search>
10862     - reqideny   <search>
10863     - reqpass    <search>
10864     - reqipass   <search>
10865     - reqrep     <search> <replace>
10866     - reqirep    <search> <replace>
10867     - reqtarpit  <search>
10868     - reqitarpit <search>
10869     - rspadd     <string>
10870     - rspdel     <search>
10871     - rspidel    <search>
10872     - rspdeny    <search>
10873     - rspideny   <search>
10874     - rsprep     <search> <replace>
10875     - rspirep    <search> <replace>
10876
10877  With all these keywords, the same conventions are used. The <search> parameter
10878  is a POSIX extended regular expression (regex) which supports grouping through
10879  parenthesis (without the backslash). Spaces and other delimiters must be
10880  prefixed with a backslash ('\') to avoid confusion with a field delimiter.
10881  Other characters may be prefixed with a backslash to change their meaning :
10882
10883    \t   for a tab
10884    \r   for a carriage return (CR)
10885    \n   for a new line (LF)
10886    \    to mark a space and differentiate it from a delimiter
10887    \#   to mark a sharp and differentiate it from a comment
10888    \\   to use a backslash in a regex
10889    \\\\ to use a backslash in the text (*2 for regex, *2 for haproxy)
10890    \xXX to write the ASCII hex code XX as in the C language
10891
10892  The <replace> parameter contains the string to be used to replace the largest
10893  portion of text matching the regex. It can make use of the special characters
10894  above, and can reference a substring which is delimited by parenthesis in the
10895  regex, by writing a backslash ('\') immediately followed by one digit from 0 to
10896  9 indicating the group position (0 designating the entire line). This practice
10897  is very common to users of the "sed" program.
10898
10899  The <string> parameter represents the string which will systematically be added
10900  after the last header line. It can also use special character sequences above.
10901
10902  Notes related to these keywords :
10903  ---------------------------------
10904    - these keywords are not always convenient to allow/deny based on header
10905      contents. It is strongly recommended to use ACLs with the "block" keyword
10906      instead, resulting in far more flexible and manageable rules.
10907
10908    - lines are always considered as a whole. It is not possible to reference
10909      a header name only or a value only. This is important because of the way
10910      headers are written (notably the number of spaces after the colon).
10911
10912    - the first line is always considered as a header, which makes it possible to
10913      rewrite or filter HTTP requests URIs or response codes, but in turn makes
10914      it harder to distinguish between headers and request line. The regex prefix
10915      ^[^\ \t]*[\ \t] matches any HTTP method followed by a space, and the prefix
10916      ^[^ \t:]*: matches any header name followed by a colon.
10917
10918    - for performances reasons, the number of characters added to a request or to
10919      a response is limited at build time to values between 1 and 4 kB. This
10920      should normally be far more than enough for most usages. If it is too short
```

```
10921      on occasional usages, it is possible to gain some space by removing some
10922      useless headers before adding new ones.
10923
10924    - keywords beginning with "reqi" and "rspi" are the same as their counterpart
10925      without the 'i' letter except that they ignore case when matching patterns.
10926
10927    - when a request passes through a frontend then a backend, all req* rules
10928      from the frontend will be evaluated, then all req* rules from the backend
10929      will be evaluated. The reverse path is applied to responses.
10930
10931    - req* statements are applied after "block" statements, so that "block" is
10932      always the first one, but before "use_backend" in order to permit rewriting
10933      before switching.
10934
10935
10936  7. Using ACLs and fetching samples
10937  ----------------------------------
10938
10939  Haproxy is capable of extracting data from request or response streams, from
10940  client or server information, from tables, environmental information etc...
10941  The action of extracting such data is called fetching a sample. Once retrieved,
10942  these samples may be used for various purposes such as a key to a stick-table,
10943  but most common usages consist in matching them against predefined constant
10944  data called patterns.
10945
10946
10947  7.1. ACL basics
10948  ---------------
10949
10950  The use of Access Control Lists (ACL) provides a flexible solution to perform
10951  content switching and generally to take decisions based on content extracted
10952  from the request, the response or any environmental status. The principle is
10953  simple :
10954
10955    - extract a data sample from a stream, table or the environment
10956    - optionally apply some format conversion to the extracted sample
10957    - apply one or multiple pattern matching methods on this sample
10958    - perform actions only when a pattern matches the sample
10959
10960  The actions generally consist in blocking a request, selecting a backend, or
10961  adding a header.
10962
10963  In order to define a test, the "acl" keyword is used. The syntax is :
10964
10965     acl <aclname> <criterion> [flags] [operator] [<value>] ...
10966
10967  This creates a new ACL <aclname> or completes an existing one with new tests.
10968  Those tests apply to the portion of request/response specified in <criterion>
10969  and may be adjusted with optional flags [flags]. Some criteria also support
10970  an operator which may be specified before the set of values. Optionally some
10971  conversion operators may be applied to the sample, and they will be specified
10972  as a comma-delimited list of keywords just after the first keyword. The values
10973  are of the type supported by the criterion, and are separated by spaces.
10974
10975  ACL names must be formed from upper and lower case letters, digits, '-' (dash),
10976  '_' (underscore) , '.' (dot) and ':' (colon). ACL names are case-sensitive,
10977  which means that "my_acl" and "My_Acl" are two different ACLs.
10978
10979  There is no enforced limit to the number of ACLs. The unused ones do not affect
10980  performance, they just consume a small amount of memory.
10981
10982  The criterion generally is the name of a sample fetch method, or one of its ACL
10983  specific declinations. The default test method is implied by the output type of
10984  this sample fetch method. The ACL declinations can describe alternate matching
10985  methods of a same sample fetch method. The sample fetch methods are the only
```



```
10986  ones supporting a conversion.
10987
10988  Sample fetch methods return data which can be of the following types :
10989    - boolean
10990    - integer (signed or unsigned)
10991    - IPv4 or IPv6 address
10992    - string
10993    - data block
10994
10995  Converters transform any of these data into any of these. For example, some
10996  converters might convert a string to a lower-case string while other ones
10997  would turn a string to an IPv4 address, or apply a netmask to an IP address.
10998  The resulting sample is of the type of the last converter applied to the list,
10999  which defaults to the type of the sample fetch method.
11000
11001  Each sample or converter returns data of a specific type, specified with its
11002  keyword in this documentation. When an ACL is declared using a standard sample
11003  fetch method, certain types automatically involved a default matching method
11004  which are summarized in the table below :
11005
11006     +---------------------+-----------------+
11007     | Sample or converter | Default         |
11008     |    output type      | matching method |
11009     +---------------------+-----------------+
11010     | boolean             | bool            |
11011     +---------------------+-----------------+
11012     | integer             | int             |
11013     +---------------------+-----------------+
11014     | ip                  | ip              |
11015     +---------------------+-----------------+
11016     | string              | str             |
11017     +---------------------+-----------------+
11018     | binary              | none, use "-m"  |
11019     +---------------------+-----------------+
11020
11021  Note that in order to match a binary samples, it is mandatory to specify a
11022  matching method, see below.
11023
11024  The ACL engine can match these types against patterns of the following types :
11025    - boolean
11026    - integer or integer range
11027    - IP address / network
11028    - string (exact, substring, suffix, prefix, subdir, domain)
11029    - regular expression
11030    - hex block
11031
11032  The following ACL flags are currently supported :
11033
11034     -i : ignore case during matching of all subsequent patterns.
11035     -f : load patterns from a file.
11036     -m : use a specific pattern matching method
11037     -n : forbid the DNS resolutions
11038     -M : load the file pointed by -f like a map file.
11039     -u : force the unique id of the ACL
11040     -- : force end of flags. Useful when a string looks like one of the flags.
11041
11042  The "-f" flag is followed by the name of a file from which all lines will be
11043  read as individual values. It is even possible to pass multiple "-f" arguments
11044  if the patterns are to be loaded from multiple files. Empty lines as well as
11045  lines beginning with a sharp ('#') will be ignored. All leading spaces and tabs
11046  will be stripped. If it is absolutely necessary to insert a valid pattern
11047  beginning with a sharp, just prefix it with a space so that it is not taken for
11048  a comment. Depending on the data type and match method, haproxy may load the
11049  lines into a binary tree, allowing very fast lookups. This is true for IPv4 and
11050  exact string matching. In this case, duplicates will automatically be removed.
```

The "-M" flag allows an ACL to use a map file. If this flag is set, the file is parsed as two column file. The first column contains the patterns used by the ACL, and the second column contain the samples. The sample can be used later by a map. This can be useful in some rare cases where an ACL would just be used to check for the existence of a pattern in a map before a mapping is applied.

The "-u" flag forces the unique id of the ACL. This unique id is used with the socket interface to identify ACL and dynamically change its values. Note that a file is always identified by its name even if an id is set.

Also, note that the "-i" flag applies to subsequent entries and not to entries loaded from files preceding it. For instance :

```
    acl valid-ua hdr(user-agent) -f exact-ua.lst -i -f generic-ua.lst test
```

In this example, each line of "exact-ua.lst" will be exactly matched against the "user-agent" header of the request. Then each line of "generic-ua" will be case-insensitively matched. Then the word "test" will be insensitively matched as well.

The "-m" flag is used to select a specific pattern matching method on the input sample. All ACL-specific criteria imply a pattern matching method and generally do not need this flag. However, this flag is useful with generic sample fetch methods to describe how they're going to be matched against the patterns. This is required for sample fetches which return data type for which there is no obvious matching method (eg: string or binary). When "-m" is specified and followed by a pattern matching method name, this method is used instead of the default one for the criterion. This makes it possible to match contents in ways that were not initially planned, or with sample fetch methods which return a string. The matching method also affects the way the patterns are parsed.

The "-n" flag forbids the dns resolutions. It is used with the load of ip files. By default, if the parser cannot parse ip address it considers that the parsed string is maybe a domain name and try dns resolution. The flag "-n" disable this resolution. It is useful for detecting malformed ip lists. Note that if the DNS server is not reachable, the haproxy configuration parsing may last many minutes waiting fir the timeout. During this time no error messages are displayed. The flag "-n" disable this behavior. Note also that during the runtime, this function is disabled for the dynamic acl modifications.

There are some restrictions however. Not all methods can be used with all sample fetch methods. Also, if "-m" is used in conjunction with "-f", it must be placed first. The pattern matching method must be one of the following :

```
  - "found" : only check if the requested sample could be found in the stream,
              but do not compare it against any pattern. It is recommended not
              to pass any pattern to avoid confusion. This matching method is
              particularly useful to detect presence of certain contents such
              as headers, cookies, etc... even if they are empty and without
              comparing them to anything nor counting them.

  - "bool"  : check the value as a boolean. It can only be applied to fetches
              which return a boolean or integer value, and takes no pattern.
              Value zero or false does not match, all other values do match.

  - "int"   : match the value as an integer. It can be used with integer and
              boolean samples. Boolean false is integer 0, true is integer 1.

  - "ip"    : match the value as an IPv4 or IPv6 address. It is compatible
              with IP address samples only, so it is implied and never needed.

  - "bin"   : match the contents against an hexadecimal string representing a
              binary sequence. This may be used with binary or string samples.
```



```
  - "len"   : match the sample's length as an integer. This may be used with
              binary or string samples.

  - "str"   : exact match : match the contents against a string. This may be
              used with binary or string samples.

  - "sub"   : substring match : check that the contents contain at least one of
              the provided string patterns. This may be used with binary or
              string samples.

  - "reg"   : regex match : match the contents against a list of regular
              expressions. This may be used with binary or string samples.

  - "beg"   : prefix match : check that the contents begin like the provided
              string patterns. This may be used with binary or string samples.

  - "end"   : suffix match : check that the contents end like the provided
              string patterns. This may be used with binary or string samples.

  - "dir"   : subdir match : check that a slash-delimited portion of the
              contents exactly matches one of the provided string patterns.
              This may be used with binary or string samples.

  - "dom"   : domain match : check that a dot-delimited portion of the contents
              exactly match one of the provided string patterns. This may be
              used with binary or string samples.
```

For example, to quickly detect the presence of cookie "JSESSIONID" in an HTTP request, it is possible to do :

```
    acl jsess_present cook(JSESSIONID) -m found
```

In order to apply a regular expression on the 500 first bytes of data in the buffer, one would use the following acl :

```
    acl script_tag payload(0,500) -m reg -i <script>
```

On systems where the regex library is much slower when using "-i", it is possible to convert the sample to lowercase before matching, like this :

```
    acl script_tag payload(0,500),lower -m reg <script>
```

All ACL-specific criteria imply a default matching method. Most often, these criteria are composed by concatenating the name of the original sample fetch method and the matching method. For example, "hdr_beg" applies the "beg" match to samples retrieved using the "hdr" fetch method. Since all ACL-specific criteria rely on a sample fetch method, it is always possible instead to use the original sample fetch method and the explicit matching method using "-m".

If an alternate match is specified using "-m" on an ACL-specific criterion, the matching method is simply applied to the underlying sample fetch method. For example, all ACLs below are exact equivalent :

```
    acl short_form  hdr_beg(host)        www.
    acl alternate1  hdr_beg(host) -m beg www.
    acl alternate2  hdr_dom(host) -m beg www.
    acl alternate3  hdr(host)     -m beg www.
```

The table below summarizes the compatibility matrix between sample or converter types and the pattern types to fetch against. It indicates for each compatible combination the name of the matching method to be used, surrounded with angle brackets ">" and "<" when the method is the default one and will work by default without "-m".

```
                            +-----------------------------------------------+
                            |                Input sample type              |
      +---------------------+---------+---------+---------+---------+---------+
      |     pattern type    | boolean | integer |   ip    | string  | binary  |
      +---------------------+---------+---------+---------+---------+---------+
      | none (presence only) |  found  |  found  |  found  |  found  |  found  |
      +---------------------+---------+---------+---------+---------+---------+
      | none (boolean value) |> bool <|   bool  |         |   bool  |         |
      +---------------------+---------+---------+---------+---------+---------+
      | integer (value)     |   int   |> int  <|   int   |   int   |         |
      +---------------------+---------+---------+---------+---------+---------+
      | integer (length)    |   len   |   len   |   len   |   len   |   len   |
      +---------------------+---------+---------+---------+---------+---------+
      | IP address          |         |         |>   ip  <|    ip   |    ip   |
      +---------------------+---------+---------+---------+---------+---------+
      | exact string        |   str   |   str   |   str   |> str  <|   str   |
      +---------------------+---------+---------+---------+---------+---------+
      | prefix              |   beg   |   beg   |   beg   |   beg   |   beg   |
      +---------------------+---------+---------+---------+---------+---------+
      | suffix              |   end   |   end   |   end   |   end   |   end   |
      +---------------------+---------+---------+---------+---------+---------+
      | substring           |   sub   |   sub   |   sub   |   sub   |   sub   |
      +---------------------+---------+---------+---------+---------+---------+
      | subdir              |   dir   |   dir   |   dir   |   dir   |   dir   |
      +---------------------+---------+---------+---------+---------+---------+
      | domain              |   dom   |   dom   |   dom   |   dom   |   dom   |
      +---------------------+---------+---------+---------+---------+---------+
      | regex               |   reg   |   reg   |   reg   |   reg   |   reg   |
      +---------------------+---------+---------+---------+---------+---------+
      | hex block           |         |         |         |   bin   |   bin   |
      +---------------------+---------+---------+---------+---------+---------+
```

7.1.1. Matching booleans

In order to match a boolean, no value is needed and all values are ignored.
Boolean matching is used by default for all fetch methods of type "boolean".
When boolean matching is used, the fetched value is returned as-is, which means
that a boolean "true" will always match and a boolean "false" will never match.

Boolean matching may also be enforced using "-m bool" on fetch methods which
return an integer value. Then, integer value 0 is converted to the boolean
"false" and all other values are converted to "true".

7.1.2. Matching integers

Integer matching applies by default to integer fetch methods. It can also be
enforced on boolean fetches using "-m int". In this case, "false" is converted
to the integer 0, and "true" is converted to the integer 1.

Integer matching also supports integer ranges and operators. Note that integer
matching only applies to positive values. A range is a value expressed with a
lower and an upper bound separated with a colon, both of which may be omitted.

For instance, "1024:65535" is a valid range to represent a range of
unprivileged ports, and "1024:" would also work. "0:1023" is a valid
representation of privileged ports, and ":1023" would also work.

As a special case, some ACL functions support decimal numbers which are in fact
two integers separated by a dot. This is used with some version checks for
instance. All integer properties apply to those decimal numbers, including
ranges and operators.

For an easier usage, comparison operators are also supported. Note that using
operators with ranges does not make much sense and is strongly discouraged.
Similarly, it does not make much sense to perform order comparisons with a set
of values.

Available operators for integer matching are :

  eq : true if the tested value equals at least one value
  ge : true if the tested value is greater than or equal to at least one value
  gt : true if the tested value is greater than at least one value
  le : true if the tested value is less than or equal to at least one value
  lt : true if the tested value is less than at least one value

For instance, the following ACL matches any negative Content-Length header :

```
  acl negative-length hdr_val(content-length) lt 0
```

This one matches SSL versions between 3.0 and 3.1 (inclusive) :

```
  acl sslv3 req_ssl_ver 3:3.1
```

7.1.3. Matching strings

String matching applies to string or binary fetch methods, and exists in 6
different forms :

  - exact match   (-m str) : the extracted string must exactly match the
    patterns ;

  - substring match (-m sub) : the patterns are looked up inside the
    extracted string, and the ACL matches if any of them is found inside ;

  - prefix match   (-m beg) : the patterns are compared with the beginning of
    the extracted string, and the ACL matches if any of them matches.

  - suffix match   (-m end) : the patterns are compared with the end of the
    extracted string, and the ACL matches if any of them matches.

  - subdir match   (-m sub) : the patterns are looked up inside the extracted
    string, delimited with slashes ("/"), and the ACL matches if any of them
    matches.

  - domain match   (-m dom) : the patterns are looked up inside the extracted
    string, delimited with dots ("."), and the ACL matches if any of them
    matches.

String matching applies to verbatim strings as they are passed, with the
exception of the backslash ("\") which makes it possible to escape some
characters such as the space. If the "-i" flag is passed before the first
string, then the matching will be performed ignoring the case. In order
to match the string "-i", either set it second, or pass the "--" flag
before the first string. Same applies of course to match the string "--".

7.1.4. Matching regular expressions (regexes)

Just like with string matching, regex matching applies to verbatim strings as
they are passed, with the exception of the backslash ("\") which makes it
possible to escape some characters such as the space. If the "-i" flag is
passed before the first regex, then the matching will be performed ignoring
the case. In order to match the string "-i", either set it second, or pass

the "--" flag before the first string. Same principle applies of course to
match the string "--".

7.1.5. Matching arbitrary data blocks

It is possible to match some extracted samples against a binary block which may
not safely be represented as a string. For this, the patterns must be passed as
a series of hexadecimal digits in an even number, when the match method is set
to binary. Each sequence of two digits will represent a byte. The hexadecimal
digits may be used upper or lower case.

Example :
```
    # match "Hello\n" in the input stream (\x48 \x65 \x6c \x6c \x6f \x0a)
    acl hello payload(0,6) -m bin 48656c6c6f0a
```

7.1.6. Matching IPv4 and IPv6 addresses

IPv4 addresses values can be specified either as plain addresses or with a
netmask appended, in which case the IPv4 address matches whenever it is
within the network. Plain addresses may also be replaced with a resolvable
host name, but this practice is generally discouraged as it makes it more
difficult to read and debug configurations. If hostnames are used, you should
at least ensure that they are present in /etc/hosts so that the configuration
does not depend on any random DNS match at the moment the configuration is
parsed.

IPv6 may be entered in their usual form, with or without a netmask appended.
Only bit counts are accepted for IPv6 netmasks. In order to avoid any risk of
trouble with randomly resolved IP addresses, host names are never allowed in
IPv6 patterns.

HAProxy is also able to match IPv4 addresses with IPv6 addresses in the
following situations :

- tested address is IPv4, pattern address is IPv4, the match applies
  in IPv4 using the supplied mask if any.
- tested address is IPv6, pattern address is IPv6, the match applies
  in IPv6 using the supplied mask if any.
- tested address is IPv6, pattern address is IPv4, the match applies in IPv4
  using the pattern's mask if the IPv6 address matches with 2002:IPV4::,
  ::IPV4 or ::ffff:IPV4, otherwise it fails.
- tested address is IPv4, pattern address is IPv6, the IPv4 address is first
  converted to IPv6 by prefixing ::ffff: in front of it, then the match is
  applied in IPv6 using the supplied IPv6 mask.

7.2. Using ACLs to form conditions

Some actions are only performed upon a valid condition. A condition is a
combination of ACLs with operators. 3 operators are supported :

- AND (implicit)
- OR (explicit with the "or" keyword or the "||" operator)
- Negation with the exclamation mark ("!")

A condition is formed as a disjunctive form:

```
  [!]acl1 [!]acl2 ... [!]acln  { or [!]acl1 [!]acl2 ... [!]acln } ...
```

Such conditions are generally used after an "if" or "unless" statement,
indicating when the condition will trigger the action.

For instance, to block HTTP requests to the "*" URL with methods other than
"OPTIONS", as well as POST requests without content-length, and GET or HEAD
requests with a content-length greater than 0, and finally every request which
is not either GET/HEAD/POST/OPTIONS !

```
   acl missing_cl hdr_cnt(Content-length) eq 0
   block if HTTP_URL_STAR !METH_OPTIONS || METH_POST missing_cl
   block if METH_GET HTTP_CONTENT
   block unless METH_GET or METH_POST or METH_OPTIONS
```

To select a different backend for requests to static contents on the "www" site
and to every request on the "img", "video", "download" and "ftp" hosts :

```
   acl url_static  path_beg         /static /images /img /css
   acl url_static  path_end         .gif .png .jpg .css .js
   acl host_www    hdr_beg(host) -i www
   acl host_static hdr_beg(host) -i img. video. download. ftp.

   # now use backend "static" for all static-only hosts, and for static urls
   # of host "www". Use backend "www" for the rest.
   use_backend static if host_static or host_www url_static
   use_backend www    if host_www
```

It is also possible to form rules using "anonymous ACLs". Those are unnamed ACL
expressions that are built on the fly without needing to be declared. They must
be enclosed between braces, with a space before and after each brace (because
the braces must be seen as independent words). Example :

   The following rule :

```
       acl missing_cl hdr_cnt(Content-length) eq 0
       block if METH_POST missing_cl
```

   Can also be written that way :

```
       block if METH_POST { hdr_cnt(Content-length) eq 0 }
```

It is generally not recommended to use this construct because it's a lot easier
to leave errors in the configuration when written that way. However, for very
simple rules matching only one source IP address for instance, it can make more
sense to use them than to declare ACLs with random names. Another example of
good use is the following :

   With named ACLs :

```
        acl site_dead nbsrv(dynamic) lt 2
        acl site_dead nbsrv(static)  lt 2
        monitor fail  if site_dead
```

   With anonymous ACLs :

```
        monitor fail if { nbsrv(dynamic) lt 2 } || { nbsrv(static) lt 2 }
```

See section 4.2 for detailed help on the "block" and "use_backend" keywords.

7.3. Fetching samples

Historically, sample fetch methods were only used to retrieve data to match
against patterns using ACLs. With the arrival of stick-tables, a new class of
sample fetch methods was created, most often sharing the same syntax as their
ACL counterpart. These sample fetch methods are also known as "fetches". As
of now, ACLs and fetches have converged. All ACL fetch methods have been made

available as fetch methods, and ACLs may use any sample fetch method as well.

This section details all available sample fetch methods and their output type.
Some sample fetch methods have deprecated aliases that are used to maintain
compatibility with existing configurations. They are then explicitly marked as
deprecated and should not be used in new setups.

The ACL derivatives are also indicated when available, with their respective
matching methods. These ones all have a well defined default pattern matching
method, so it is never necessary (though allowed) to pass the "-m" option to
indicate how the sample will be matched using ACLs.

As indicated in the sample type versus matching compatibility matrix above,
when using a generic sample fetch method in an ACL, the "-m" option is
mandatory unless the sample type is one of boolean, integer, IPv4 or IPv6. When
the same keyword exists as an ACL keyword and as a standard fetch method, the
ACL engine will automatically pick the ACL-only one by default.

Some of these keywords support one or multiple mandatory arguments, and one or
multiple optional arguments. These arguments are strongly typed and are checked
when the configuration is parsed so that there is no risk of running with an
incorrect argument (eg: an unresolved backend name). Fetch function arguments
are passed between parenthesis and are delimited by commas. When an argument
is optional, it will be indicated below between square brackets ('[ ]'). When
all arguments are optional, the parenthesis may be omitted.

Thus, the syntax of a standard sample fetch method is one of the following :
   - name
   - name(arg1)
   - name(arg1,arg2)

7.3.1. Converters

Sample fetch methods may be combined with transformations to be applied on top
of the fetched sample (also called "converters"). These combinations form what
is called "sample expressions" and the result is a "sample". Initially this
was only supported by "stick on" and "stick store-request" directives but this
has now be extended to all places where samples may be used (acls, log-format,
unique-id-format, add-header, ...).

These transformations are enumerated as a series of specific keywords after the
sample fetch method. These keywords may equally be appended immediately after
the fetch keyword's argument, delimited by a comma. These keywords can also
support some arguments (eg: a netmask) which must be passed in parenthesis.

A certain category of converters are bitwise and arithmetic operators which
support performing basic operations on integers. Some bitwise operations are
supported (and, or, xor, cpl) and some arithmetic operations are supported
(add, sub, mul, div, mod, neg). Some comparators are provided (odd, even, not,
bool) which make it possible to report a match without having to write an ACL.

The currently available list of transformation keywords include :

add(<value>)
  Adds <value> to the input value of type signed integer, and returns the
  result as a signed integer. <value> can be a numeric value or a variable
  name. The name of the variable starts by an indication about its scope. The
  allowed scopes are:
    "sess" : the variable is shared with all the session,
    "txn" : the variable is shared with all the transaction (request and
             response),
    "req" : the variable is shared only during the request processing,
    "res" : the variable is shared only during the response processing.

  This prefix is followed by a name. The separator is a '.'. The name may only
  contain characters 'a-z', 'A-Z', '0-9' and '_'.

and(<value>)
  Performs a bitwise "AND" between <value> and the input value of type signed
  integer, and returns the result as an signed integer. <value> can be a
  numeric value or a variable name. The name of the variable starts by an
  indication about its scope. The allowed scopes are:
    "sess" : the variable is shared with all the session,
    "txn" : the variable is shared with all the transaction (request and
             response),
    "req" : the variable is shared only during the request processing,
    "res" : the variable is shared only during the response processing.
  This prefix is followed by a name. The separator is a '.'. The name may only
  contain characters 'a-z', 'A-Z', '0-9' and '_'.

base64
  Converts a binary input sample to a base64 string. It is used to log or
  transfer binary content in a way that can be reliably transferred (eg:
  an SSL ID can be copied in a header).

bool
  Returns a boolean TRUE if the input value of type signed integer is
  non-null, otherwise returns FALSE. Used in conjunction with and(), it can be
  used to report true/false for bit testing on input values (eg: verify the
  presence of a flag).

bytes(<offset>[,<length>])
  Extracts some bytes from an input binary sample. The result is a binary
  sample starting at an offset (in bytes) of the original sample and
  optionnaly truncated at the given length.

cpl
  Takes the input value of type signed integer, applies a ones-complement
  (flips all bits) and returns the result as an signed integer.

crc32([<avalanche>])
  Hashes a binary input sample into an unsigned 32-bit quantity using the CRC32
  hash function. Optionally, it is possible to apply a full avalanche hash
  function to the output if the optional <avalanche> argument equals 1. This
  converter uses the same functions as used by the various hash-based load
  balancing algorithms, so it will provide exactly the same results. It is
  provided for compatibility with other software which want a CRC32 to be
  computed on some input keys, so it follows the most common implementation as
  found in Ethernet, Gzip, PNG, etc... It is slower than the other algorithms
  but may provide a better or at least less predictable distribution. It must
  not be used for security purposes as a 32-bit hash is trivial to break. See
  also "djb2", "sdbm", "wt6" and the "hash-type" directive.

da-csv-conv(<prop>[,<prop>*])
  Asks the DeviceAtlas converter to identify the User Agent string passed on
  input, and to emit a string made of the concatenation of the properties
  enumerated in argument, delimited by the separator defined by the global
  keyword "deviceatlas-property-separator", or by default the pipe character
  ('|'). There's a limit of 5 different properties imposed by the haproxy
  configuration language.

  Example:
    frontend www
      bind *:8881
      default_backend servers
      http-request set-header X-DeviceAtlas-Data %[req.fhdr(User-Agent),da-csv(prima…

debug
  This converter is used as debug tool. It dumps on screen the content and the

```
   type of the input sample. The sample is returned as is on its output. This
   converter only exists when haproxy was built with debugging enabled.

  div(<value>)
    Divides the input value of type signed integer by <value>, and returns the
    result as an signed integer. If <value> is null, the largest unsigned
    integer is returned (typically 2^63-1). <value> can be a numeric value or a
    variable name. The name of the variable starts by an indication about it
    scope. The scope allowed are:
      "sess" : the variable is shared with all the session,
      "txn"  : the variable is shared with all the transaction (request and
               response),
      "req"  : the variable is shared only during the request processing,
      "res"  : the variable is shared only during the response processing.
    This prefix is followed by a name. The separator is a '.'. The name may only
    contain characters 'a-z', 'A-Z', '0-9' and '_'.

  djb2([<avalanche>])
    Hashes a binary input sample into an unsigned 32-bit quantity using the DJB2
    hash function. Optionally, it is possible to apply a full avalanche hash
    function to the output if the optional <avalanche> argument equals 1. This
    converter uses the same functions as used by the various hash-based load
    balancing algorithms, so it will provide exactly the same results. It is
    mostly intended for debugging, but can be used as a stick-table entry to
    collect rough statistics. It must not be used for security purposes as a
    32-bit hash is trivial to break. See also "crc32", "sdbm", "wt6" and the
    "hash-type" directive.

  even
    Returns a boolean TRUE if the input value of type signed integer is even
    otherwise returns FALSE. It is functionally equivalent to "not,and(1),bool".

  field(<index>,<delimiters>)
    Extracts the substring at the given index considering given delimiters from
    an input string. Indexes start at 1 and delimiters are a string formatted
    list of chars.

  hex
    Converts a binary input sample to an hex string containing two hex digits per
    input byte. It is used to log or transfer hex dumps of some binary input data
    in a way that can be reliably transferred (eg: an SSL ID can be copied in a
    header).

  http_date([<offset>])
    Converts an integer supposed to contain a date since epoch to a string
    representing this date in a format suitable for use in HTTP header fields. If
    an offset value is specified, then it is a number of seconds that is added to
    the date before the conversion is operated. This is particularly useful to
    emit Date header fields, Expires values in responses when combined with a
    positive offset, or Last-Modified values when the offset is negative.

  in_table(<table>)
    Uses the string representation of the input sample to perform a look up in
    the specified table. If the key is not found in the table, a boolean false
    is returned. Otherwise a boolean true is returned. This can be used to verify
    the presence of a certain key in a table tracking some elements (eg: whether
    or not a source IP address or an Authorization header was already seen).

  ipmask(<mask>)
    Apply a mask to an IPv4 address, and use the result for lookups and storage.
    This can be used to make all hosts within a certain mask to share the same
    table entries and as such use the same server. The mask can be passed in
    dotted form (eg: 255.255.255.0) or in CIDR form (eg: 24).

  json([<input-code>])
```



```
    Escapes the input string and produces an ASCII ouput string ready to use as a
    JSON string. The converter tries to decode the input string according to the
    <input-code> parameter. It can be "ascii", "utf8", "utf8s", "utf8"" or
    "utf8ps". The "ascii" decoder never fails. The "utf8" decoder detects 3 types
    of errors:
     - bad UTF-8 sequence (lone continuation byte, bad number of continuation
       bytes, ...)
     - invalid range (the decoded value is within a UTF-8 prohibited range),
     - code overlong (the value is encoded with more bytes than necessary).

    The UTF-8 JSON encoding can produce a "too long value" error when the UTF-8
    character is greater than 0xffff because the JSON string escape specification
    only authorizes 4 hex digits for the value encoding. The UTF-8 decoder exists
    in 4 variants designated by a combination of two suffix letters : "p" for
    "permissive" and "s" for "silently ignore". The behaviors of the decoders
    are :
     - "ascii"  : never fails ;
     - "utf8"   : fails on any detected errors ;
     - "utf8s"  : never fails, but removes characters corresponding to errors ;
     - "utf8p"  : accepts and fixes the overlong errors, but fails on any other
                  error ;
     - "utf8ps" : never fails, accepts and fixes the overlong errors, but removes
                  characters corresponding to the other errors.

    This converter is particularly useful for building properly escaped JSON for
    logging to servers which consume JSON-formated traffic logs.

    Example:
       capture request header user-agent len 150
       capture request header Host len 15
       log-format {"ip":"%[src]","user-agent":"%[capture.req.hdr(1),json]"}

    Input request from client 127.0.0.1:
       GET / HTTP/1.0
       User-Agent: Very "Ugly" UA 1/2

    Output log:
       {"ip":"127.0.0.1","user-agent":"Very \"Ugly\" UA 1\/2"}

  language(<value>[,<default>])
    Returns the value with the highest q-factor from a list as extracted from the
    "accept-language" header using "req.fhdr". Values with no q-factor have a
    q-factor of 1. Values with a q-factor of 0 are dropped. Only values which
    belong to the list of semi-colon delimited <values> will be considered. The
    argument <value> syntax is "lang[;lang[;lang[;...]]]". If no value matches the
    given list and a default value is provided, it is returned. Note that language
    names may have a variant after a dash ('-'). If this variant is present in the
    list, it will be matched, but if it is not, only the base language is checked.
    The match is case-sensitive, and the output string is always one of those
    provided in arguments.  The ordering of arguments is meaningless, only the
    ordering of the values in the request counts, as the first value among
    multiple sharing the same q-factor is used.

    Example :

      # this configuration switches to the backend matching a
      # given language based on the request :

      acl es req.fhdr(accept-language),language(es;fr;en) -m str es
      acl fr req.fhdr(accept-language),language(es;fr;en) -m str fr
      acl en req.fhdr(accept-language),language(es;fr;en) -m str en
      use_backend spanish if es
      use_backend french  if fr
      use_backend english if en
      default_backend choose_your_language
```

```
11701
11702  lower
11703    Convert a string sample to lower case. This can only be placed after a string
11704    sample fetch function or after a transformation keyword returning a string
11705    type. The result is of type string.
11706
11707  ltime(<format>[,<offset>])
11708    Converts an integer supposed to contain a date since epoch to a string
11709    representing this date in local time using a format defined by the <format>
11710    string using strftime(3). The purpose is to allow any date format to be used
11711    in logs. An optional <offset> in seconds may be applied to the input date
11712    (positive or negative). See the strftime() man page for the format supported
11713    by your operating system. See also the utime converter.
11714
11715    Example :
11716
11717        # Emit two colons, one with the local time and another with ip:port
11718        # Eg:  20140710162350 127.0.0.1:57325
11719        log-format %[date,ltime(%Y%m%d%H%M%S)]\ %ci:%cp
11720
11721  map(<map_file>[,<default_value>])
11722  map_<match_type>(<map_file>[,<default_value>])
11723  map_<match_type>_<output_type>(<map_file>[,<default_value>])
11724    Search the input value from <map_file> using the <match_type> matching method,
11725    and return the associated value converted to the type <output_type>. If the
11726    input value cannot be found in the <map_file>, the converter returns the
11727    <default_value>. If the <default_value> is not set, the converter fails and
11728    acts as if no input value could be fetched. If the <match_type> is not set, it
11729    defaults to "str". Likewise, if the <output_type> is not set, it defaults to
11730    "str". For convenience, the "map" keyword is an alias for "map_str" and maps a
11731    string to another string.
11732
11733    It is important to avoid overlapping between the keys : IP addresses and
11734    strings are stored in trees, so the first of the finest match will be used.
11735    Other keys are stored in lists, so the first matching occurrence will be used.
11736
11737    The following array contains the list of all map functions avalaible sorted by
11738    input type, match type and output type.
11739
11740    input type | match method | output type str | output type int | output type ip
11741    -----------+--------------+-----------------+-----------------+---------------
11742      str      | str          | map_str         | map_str_int     | map_str_ip
11743    -----------+--------------+-----------------+-----------------+---------------
11744      str      | beg          | map_beg         | map_beg_int     | map_end_ip
11745    -----------+--------------+-----------------+-----------------+---------------
11746      str      | sub          | map_sub         | map_sub_int     | map_sub_ip
11747    -----------+--------------+-----------------+-----------------+---------------
11748      str      | dir          | map_dir         | map_dir_int     | map_dir_ip
11749    -----------+--------------+-----------------+-----------------+---------------
11750      str      | dom          | map_dom         | map_dom_int     | map_dom_ip
11751    -----------+--------------+-----------------+-----------------+---------------
11752      str      | end          | map_end         | map_end_int     | map_end_ip
11753    -----------+--------------+-----------------+-----------------+---------------
11754      str      | reg          | map_reg         | map_reg_int     | map_reg_ip
11755    -----------+--------------+-----------------+-----------------+---------------
11756      int      | int          | map_int         | map_int_int     | map_int_ip
11757    -----------+--------------+-----------------+-----------------+---------------
11758      ip       | ip           | map_ip          | map_ip_int      | map_ip_ip
11759    -----------+--------------+-----------------+-----------------+---------------
11760
11761    The file contains one key + value per line. Lines which start with '#' are
11762    ignored, just like empty lines. Leading tabs and spaces are stripped. The key
11763    is then the first "word" (series of non-space/tabs characters), and the value
11764    is what follows this series of space/tab till the end of the line excluding
11765    trailing spaces/tabs.
```



```
11766
11767    Example :
11768
11769       # this is a comment and is ignored
11770          2.22.246.0/23    United Kingdom      \n
11771       <-><----------><--><------------><---->
11772        |       |       |         |        `- trailing spaces ignored
11773        |       |       |         `---------- value
11774        |       |       `-------------------- middle spaces ignored
11775        |       `---------------------------- key
11776        `------------------------------------ leading spaces ignored
11777
11778  mod(<value>)
11779    Divides the input value of type signed integer by <value>, and returns the
11780    remainder as an signed integer. If <value> is null, then zero is returned.
11781    <value> can be a numeric value or a variable name. The name of the variable
11782    starts by an indication about its scope. The allowed scopes are:
11783      "sess" : the variable is shared with all the session,
11784      "txn"  : the variable is shared with all the transaction (request and
11785               response),
11786      "req"  : the variable is shared only during the request processing,
11787      "res"  : the variable is shared only during the response processing.
11788    This prefix is followed by a name. The separator is a '.'. The name may only
11789    contain characters 'a-z', 'A-Z', '0-9' and '_'.
11790
11791  mul(<value>)
11792    Multiplies the input value of type signed integer by <value>, and returns
11793    the product as an signed integer. In case of overflow, the largest possible
11794    value for the sign is returned so that the operation doesn't wrap around.
11795    <value> can be a numeric value or a variable name. The name of the variable
11796    starts by an indication about its scope. The allowed scopes are:
11797      "sess" : the variable is shared with all the session,
11798      "txn"  : the variable is shared with all the transaction (request and
11799               response),
11800      "req"  : the variable is shared only during the request processing,
11801      "res"  : the variable is shared only during the response processing.
11802    This prefix is followed by a name. The separator is a '.'. The name may only
11803    contain characters 'a-z', 'A-Z', '0-9' and '_'.
11804
11805  neg
11806    Takes the input value of type signed integer, computes the opposite value,
11807    and returns the remainder as an signed integer. 0 is identity. This operator
11808    is provided for reversed subtracts : in order to subtract the input from a
11809    constant, simply perform a "neg,add(value)".
11810
11811  not
11812    Returns a boolean FALSE if the input value of type signed integer is
11813    non-null, otherwise returns TRUE. Used in conjunction with and(), it can be
11814    used to report true/false for bit testing on input values (eg: verify the
11815    absence of a flag).
11816
11817  odd
11818    Returns a boolean TRUE if the input value of type signed integer is odd
11819    otherwise returns FALSE. It is functionally equivalent to "and(1),bool".
11820
11821  or(<value>)
11822    Performs a bitwise "OR" between <value> and the input value of type signed
11823    integer, and returns the result as an signed integer. <value> can be a
11824    numeric value or a variable name. The name of the variable starts by an
11825    indication about its scope. The allowed scopes are:
11826      "sess" : the variable is shared with all the session,
11827      "txn"  : the variable is shared with all the transaction (request and
11828               response),
11829      "req"  : the variable is shared only during the request processing,
11830      "res"  : the variable is shared only during the response processing.
```

```
   This prefix is followed by a name. The separator is a '.'. The name may only
   contain characters 'a-z', 'A-Z', '0-9' and '_'.

regsub(<regex>,<subst>[,<flags>])
   Applies a regex-based substitution to the input string. It does the same
   operation as the well-known "sed" utility with "s/<regex>/<subst>/". By
   default it will replace in the input string the first occurrence of the
   largest part matching the regular expression <regex> with the substitution
   string <subst>. It is possible to replace all occurrences instead by adding
   the flag "g" in the third argument <flags>. It is also possible to make the
   regex case insensitive by adding the flag "i" in <flags>. Since <flags> is a
   string, it is made up from the concatenation of all desired flags. Thus if
   both "i" and "g" are desired, using "gi" or "ig" will have the same effect.
   It is important to note that due to the current limitations of the
   configuration parser, some characters such as closing parenthesis or comma
   are not possible to use in the arguments. The first use of this converter is
   to replace certain characters or sequence of characters with other ones.

   Example :

      # de-duplicate "/" in header "x-path".
      # input:  x-path: /////a///b/c/xzxyz/
      # output: x-path: /a/b/c/xzxyz/
      http-request set-header x-path %[hdr(x-path),regsub(/+,/,g)]

capture-req(<id>)
   Capture the string entry in the request slot <id> and returns the entry as
   is. If the slot doesn't exist, the capture fails silently.

   See also: "declare capture", "http-request capture",
             "http-response capture", "capture.req.hdr" and
             "capture.res.hdr" (sample fetches).

capture-res(<id>)
   Capture the string entry in the response slot <id> and returns the entry as
   is. If the slot doesn't exist, the capture fails silently.

   See also: "declare capture", "http-request capture",
             "http-response capture", "capture.req.hdr" and
             "capture.res.hdr" (sample fetches).

sdbm([<avalanche>])
   Hashes a binary input sample into an unsigned 32-bit quantity using the SDBM
   hash function. Optionally, it is possible to apply a full avalanche hash
   function to the output if the optional <avalanche> argument equals 1. This
   converter uses the same functions as used by the various hash-based load
   balancing algorithms, so it will provide exactly the same results. It is
   mostly intended for debugging, but can be used as a stick-table entry to
   collect rough statistics. It must not be used for security purposes as a
   32-bit hash is trivial to break. See also "crc32", "djb2", "wt6" and the
   "hash-type" directive.

set-var(<var name>)
   Sets a variable with the input content and return the content on the output as
   is. The variable keep the value and the associated input type. The name of the
   variable starts by an indication about it scope. The scope allowed are:
     "sess" : the variable is shared with all the session,
     "txn"  : the variable is shared with all the transaction (request and
              response),
     "req"  : the variable is shared only during the request processing,
     "res"  : the variable is shared only during the response processing.
   This prefix is followed by a name. The separator is a '.'. The name may only
   contain characters 'a-z', 'A-Z', '0-9' and '_'.

sub(<value>)
```



```
   Subtracts <value> from the input value of type signed integer, and returns
   the result as an signed integer. Note: in order to subtract the input from
   a constant, simply perform a "neg,add(value)". <value> can be a numeric value
   or a variable name. The name of the variable starts by an indication about its
   scope. The allowed scopes are:
     "sess" : the variable is shared with all the session,
     "txn"  : the variable is shared with all the transaction (request and
              response),
     "req"  : the variable is shared only during the request processing,
     "res"  : the variable is shared only during the response processing.
   This prefix is followed by a name. The separator is a '.'. The name may only
   contain characters 'a-z', 'A-Z', '0-9' and '_'.

table_bytes_in_rate(<table>)
   Uses the string representation of the input sample to perform a look up in
   the specified table. If the key is not found in the table, integer value zero
   is returned. Otherwise the converter returns the average client-to-server
   bytes rate associated with the input sample in the designated table, measured
   in amount of bytes over the period configured in the table. See also the
   sc_bytes_in_rate sample fetch keyword.


table_bytes_out_rate(<table>)
   Uses the string representation of the input sample to perform a look up in
   the specified table. If the key is not found in the table, integer value zero
   is returned. Otherwise the converter returns the average server-to-client
   bytes rate associated with the input sample in the designated table, measured
   in amount of bytes over the period configured in the table. See also the
   sc_bytes_out_rate sample fetch keyword.

table_conn_cnt(<table>)
   Uses the string representation of the input sample to perform a look up in
   the specified table. If the key is not found in the table, integer value zero
   is returned. Otherwise the converter returns the cumulated amount of incoming
   connections associated with the input sample in the designated table. See
   also the sc_conn_cnt sample fetch keyword.

table_conn_cur(<table>)
   Uses the string representation of the input sample to perform a look up in
   the specified table. If the key is not found in the table, integer value zero
   is returned. Otherwise the converter returns the current amount of concurrent
   tracked connections associated with the input sample in the designated table.
   See also the sc_conn_cur sample fetch keyword.

table_conn_rate(<table>)
   Uses the string representation of the input sample to perform a look up in
   the specified table. If the key is not found in the table, integer value zero
   is returned. Otherwise the converter returns the average incoming connection
   rate associated with the input sample in the designated table. See also the
   sc_conn_rate sample fetch keyword.

table_gpt0(<table>)
   Uses the string representation of the input sample to perform a look up in
   the specified table. If the key is not found in the table, boolean value zero
   is returned. Otherwise the converter returns the current value of the first
   general purpose tag associated with the input sample in the designated table.
   See also the sc_get_gpt0 sample fetch keyword.

table_gpc0(<table>)
   Uses the string representation of the input sample to perform a look up in
   the specified table. If the key is not found in the table, integer value zero
   is returned. Otherwise the converter returns the current value of the first
   general purpose counter associated with the input sample in the designated
   table. See also the sc_get_gpc0 sample fetch keyword.
```

```
11961  table_gpc0_rate(<table>)
11962    Uses the string representation of the input sample to perform a look up in
11963    the specified table. If the key is not found in the table, integer value zero
11964    is returned. Otherwise the converter returns the frequency which the gpc0
11965    counter was incremented over the configured period in the table, associated
11966    with the input sample in the designated table. See also the sc_get_gpc0_rate
11967    sample fetch keyword.
11968
11969  table_http_err_cnt(<table>)
11970    Uses the string representation of the input sample to perform a look up in
11971    the specified table. If the key is not found in the table, integer value zero
11972    is returned. Otherwise the converter returns the cumulated amount of HTTP
11973    errors associated with the input sample in the designated table. See also the
11974    sc_http_err_cnt sample fetch keyword.
11975
11976  table_http_err_rate(<table>)
11977    Uses the string representation of the input sample to perform a look up in
11978    the specified table. If the key is not found in the table, integer value zero
11979    is returned. Otherwise the average rate of HTTP errors associated with the
11980    input sample in the designated table, measured in amount of errors over the
11981    period configured in the table. See also the sc_http_err_rate sample fetch
11982    keyword.
11983
11984  table_http_req_cnt(<table>)
11985    Uses the string representation of the input sample to perform a look up in
11986    the specified table. If the key is not found in the table, integer value zero
11987    is returned. Otherwise the converter returns the cumulated amount of HTTP
11988    requests associated with the input sample in the designated table. See also
11989    the sc_http_req_cnt sample fetch keyword.
11990
11991  table_http_req_rate(<table>)
11992    Uses the string representation of the input sample to perform a look up in
11993    the specified table. If the key is not found in the table, integer value zero
11994    is returned. Otherwise the average rate of HTTP requests associated with the
11995    input sample in the designated table, measured in amount of requests over the
11996    period configured in the table. See also the sc_http_req_rate sample fetch
11997    keyword.
11998
11999  table_kbytes_in(<table>)
12000    Uses the string representation of the input sample to perform a look up in
12001    the specified table. If the key is not found in the table, integer value zero
12002    is returned. Otherwise the converter returns the cumulated amount of client-
12003    to-server data associated with the input sample in the designated table,
12004    measured in kilobytes. The test is currently performed on 32-bit integers,
12005    which limits values to 4 terabytes. See also the sc_kbytes_in sample fetch
12006    keyword.
12007
12008  table_kbytes_out(<table>)
12009    Uses the string representation of the input sample to perform a look up in
12010    the specified table. If the key is not found in the table, integer value zero
12011    is returned. Otherwise the converter returns the cumulated amount of server-
12012    to-client data associated with the input sample in the designated table,
12013    measured in kilobytes. The test is currently performed on 32-bit integers,
12014    which limits values to 4 terabytes. See also the sc_kbytes_out sample fetch
12015    keyword.
12016
12017  table_server_id(<table>)
12018    Uses the string representation of the input sample to perform a look up in
12019    the specified table. If the key is not found in the table, integer value zero
12020    is returned. Otherwise the converter returns the server ID associated with
12021    the input sample in the designated table. A server ID is associated to a
12022    sample by a "stick" rule when a connection to a server succeeds. A server ID
12023    zero means that no server is associated with this key.
12024
12025  table_sess_cnt(<table>)
```



```
12026    Uses the string representation of the input sample to perform a look up in
12027    the specified table. If the key is not found in the table, integer value zero
12028    is returned. Otherwise the converter returns the cumulated amount of incoming
12029    sessions associated with the input sample in the designated table. Note that
12030    a session here refers to an incoming connection being accepted by the
12031    "tcp-request connection" rulesets. See also the sc_sess_cnt sample fetch
12032    keyword.
12033
12034  table_sess_rate(<table>)
12035    Uses the string representation of the input sample to perform a look up in
12036    the specified table. If the key is not found in the table, integer value zero
12037    is returned. Otherwise the converter returns the average incoming session
12038    rate associated with the input sample in the designated table. Note that a
12039    session here refers to an incoming connection being accepted by the
12040    "tcp-request connection" rulesets. See also the sc_sess_rate sample fetch
12041    keyword.
12042
12043  table_trackers(<table>)
12044    Uses the string representation of the input sample to perform a look up in
12045    the specified table. If the key is not found in the table, integer value zero
12046    is returned. Otherwise the converter returns the current amount of concurrent
12047    connections tracking the same key as the input sample in the designated
12048    table. It differs from table_conn_cur in that it does not rely on any stored
12049    information but on the table's reference count (the "use" value which is
12050    returned by "show table" on the CLI). This may sometimes be more suited for
12051    layer7 tracking. It can be used to tell a server how many concurrent
12052    connections there are from a given address for example. See also the
12053    sc_trackers sample fetch keyword.
12054
12055  upper
12056    Convert a string sample to upper case. This can only be placed after a string
12057    sample fetch function or after a transformation keyword returning a string
12058    type. The result is of type string.
12059
12060  url_dec
12061    Takes an url-encoded string provided as input and returns the decoded
12062    version as output. The input and the output are of type string.
12063
12064  utime(<format>[,<offset>])
12065    Converts an integer supposed to contain a date since epoch to a string
12066    representing this date in UTC time using a format defined by the <format>
12067    string using strftime(3). The purpose is to allow any date format to be used
12068    in logs. An optional <offset> in seconds may be applied to the input date
12069    (positive or negative). See the strftime() man page for the format supported
12070    by your operating system. See also the ltime converter.
12071
12072    Example :
12073
12074        # Emit two colons, one with the UTC time and another with ip:port
12075        # Eg:  20140710162350 127.0.0.1:57325
12076        log-format %[date,utime(%Y%m%d%H%M%S)]\ %ci:%cp
12077
12078  word(<index>,<delimiters>)
12079    Extracts the nth word considering given delimiters from an input string.
12080    Indexes start at 1 and delimiters are a string formatted list of chars.
12081
12082  wt6([<avalanche>])
12083    Hashes a binary input sample into an unsigned 32-bit quantity using the WT6
12084    hash function. Optionally, it is possible to apply a full avalanche hash
12085    function to the output if the optional <avalanche> argument equals 1. This
12086    converter uses the same functions as used by the various hash-based load
12087    balancing algorithms, so it will provide exactly the same results. It is
12088    mostly intended for debugging, but can be used as a stick-table entry to
12089    collect rough statistics. It must not be used for security purposes as a
12090    32-bit hash is trivial to break. See also "crc32", "djb2", "sdbm", and the
```

```
   "hash-type" directive.

xor(<value>)
  Performs a bitwise "XOR" (exclusive OR) between <value> and the input value
  of type signed integer, and returns the result as an signed integer.
  <value> can be a numeric value or a variable name. The name of the variable
  starts by an indication about its scope. The allowed scopes are:
    "sess" : the variable is shared with all the session,
    "txn"  : the variable is shared with all the transaction (request and
             response),
    "req"  : the variable is shared only during the request processing,
    "res"  : the variable is shared only during the response processing.
  This prefix is followed by a name. The separator is a '.'. The name may only
  contain characters 'a-z', 'A-Z', '0-9' and '_'.


7.3.2. Fetching samples from internal states
--------------------------------------------

A first set of sample fetch methods applies to internal information which does
not even relate to any client information. These ones are sometimes used with
"monitor-fail" directives to report an internal status to external watchers.
The sample fetch methods described in this section are usable anywhere.

always_false : boolean
  Always returns the boolean "false" value. It may be used with ACLs as a
  temporary replacement for another one when adjusting configurations.

always_true : boolean
  Always returns the boolean "true" value. It may be used with ACLs as a
  temporary replacement for another one when adjusting configurations.

avg_queue([<backend>]) : integer
  Returns the total number of queued connections of the designated backend
  divided by the number of active servers. The current backend is used if no
  backend is specified. This is very similar to "queue" except that the size of
  the farm is considered, in order to give a more accurate measurement of the
  time it may take for a new connection to be processed. The main usage is with
  ACL to return a sorry page to new users when it becomes certain they will get
  a degraded service, or to pass to the backend servers in a header so that
  they decide to work in degraded mode or to disable some functions to speed up
  the processing a bit. Note that in the event there would not be any active
  server anymore, twice the number of queued connections would be considered as
  the measured value. This is a fair estimate, as we expect one server to get
  back soon anyway, but we still prefer to send new traffic to another backend
  if in better shape. See also the "queue", "be_conn", and "be_sess_rate"
  sample fetches.

be_conn([<backend>]) : integer
  Applies to the number of currently established connections on the backend,
  possibly including the connection being evaluated. If no backend name is
  specified, the current one is used. But it is also possible to check another
  backend. It can be used to use a specific farm when the nominal one is full.
  See also the "fe_conn", "queue" and "be_sess_rate" criteria.

be_sess_rate([<backend>]) : integer
  Returns an integer value corresponding to the sessions creation rate on the
  backend, in number of new sessions per second. This is used with ACLs to
  switch to an alternate backend when an expensive or fragile one reaches too
  high a session rate, or to limit abuse of service (eg. prevent sucking of an
  online dictionary). It can also be useful to add this element to logs using a
  log-format directive.

  Example :
        # Redirect to an error page if the dictionary is requested too often
```



```
        backend dynamic
            mode http
            acl being_scanned be_sess_rate gt 100
            redirect location /denied.html if being_scanned

bin(<hexa>) : bin
  Returns a binary chain. The input is the hexadecimal representation
  of the string.

bool(<bool>) : bool
  Returns a boolean value. <bool> can be 'true', 'false', '1' or '0'.
  'false' and '0' are the same. 'true' and '1' are the same.

connslots([<backend>]) : integer
  Returns an integer value corresponding to the number of connection slots
  still available in the backend, by totaling the maximum amount of
  connections on all servers and the maximum queue size. This is probably only
  used with ACLs.

  The basic idea here is to be able to measure the number of connection "slots"
  still available (connection + queue), so that anything beyond that (intended
  usage; see "use_backend" keyword) can be redirected to a different backend.

  'connslots' = number of available server connection slots, + number of
  available server queue slots.

  Note that while "fe_conn" may be used, "connslots" comes in especially
  useful when you have a case of traffic going to one single ip, splitting into
  multiple backends (perhaps using ACLs to do name-based load balancing) and
  you want to be able to differentiate between different backends, and their
  available "connslots".  Also, whereas "nbsrv" only measures servers that are
  actually *down*, this fetch is more fine-grained and looks into the number of
  available connection slots as well. See also "queue" and "avg_queue".

  OTHER CAVEATS AND NOTES: at this point in time, the code does not take care
  of dynamic connections. Also, if any of the server maxconn, or maxqueue is 0,
  then this fetch clearly does not make sense, in which case the value returned
  will be -1.

date([<offset>]) : integer
  Returns the current date as the epoch (number of seconds since 01/01/1970).
  If an offset value is specified, then it is a number of seconds that is added
  to the current date before returning the value. This is particularly useful
  to compute relative dates, as both positive and negative offsets are allowed.
  It is useful combined with the http_date converter.

  Example :

     # set an expires header to now+1 hour in every response
     http-response set-header Expires %[date(3600),http_date]

env(<name>) : string
  Returns a string containing the value of environment variable <name>. As a
  reminder, environment variables are per-process and are sampled when the
  process starts. This can be useful to pass some information to a next hop
  server, or with ACLs to take specific action when the process is started a
  certain way.

  Examples :
      # Pass the Via header to next hop with the local hostname in it
      http-request add-header Via 1.1\ %[env(HOSTNAME)]

      # reject cookie-less requests when the STOP environment variable is set
      http-request deny if !{ cook(SESSIONID) -m found } { env(STOP) -m found }
```

```
fe_conn([<frontend>]) : integer
  Returns the number of currently established connections on the frontend,
  possibly including the connection being evaluated. If no frontend name is
  specified, the current one is used. But it is also possible to check another
  frontend. It can be used to return a sorry page before hard-blocking, or to
  use a specific backend to drain new requests when the farm is considered
  full.  This is mostly used with ACLs but can also be used to pass some
  statistics to servers in HTTP headers. See also the "dst_conn", "be_conn",
  "fe_sess_rate" fetches.

fe_sess_rate([<frontend>]) : integer
  Returns an integer value corresponding to the sessions creation rate on the
  frontend, in number of new sessions per second. This is used with ACLs to
  limit the incoming session rate to an acceptable range in order to prevent
  abuse of service at the earliest moment, for example when combined with other
  layer 4 ACLs in order to force the clients to wait a bit for the rate to go
  down below the limit. It can also be useful to add this element to logs using
  a log-format directive. See also the "rate-limit sessions" directive for use
  in frontends.

  Example :
        # This frontend limits incoming mails to 10/s with a max of 100
        # concurrent connections. We accept any connection below 10/s, and
        # force excess clients to wait for 100 ms. Since clients are limited to
        # 100 max, there cannot be more than 10 incoming mails per second.
        frontend mail
            bind :25
            mode tcp
            maxconn 100
            acl too_fast fe_sess_rate ge 10
            tcp-request inspect-delay 100ms
            tcp-request content accept if ! too_fast
            tcp-request content accept if WAIT_END

int(<integer>) : signed integer
  Returns a signed integer.

ipv4(<ipv4>) : ipv4
  Returns an ipv4.

ipv6(<ipv6>) : ipv6
  Returns an ipv6.

meth(<method>) : method
  Returns a method.

nbproc : integer
  Returns an integer value corresponding to the number of processes that were
  started (it equals the global "nbproc" setting). This is useful for logging
  and debugging purposes.

nbsrv([<backend>]) : integer
  Returns an integer value corresponding to the number of usable servers of
  either the current backend or the named backend. This is mostly used with
  ACLs but can also be useful when added to logs. This is normally used to
  switch to an alternate backend when the number of servers is too low to
  to handle some load. It is useful to report a failure when combined with
  "monitor fail".

proc : integer
  Returns an integer value corresponding to the position of the process calling
  the function, between 1 and global.nbproc. This is useful for logging and
  debugging purposes.

queue([<backend>]) : integer
```



```
  Returns the total number of queued connections of the designated backend,
  including all the connections in server queues. If no backend name is
  specified, the current one is used, but it is also possible to check another
  one. This is useful with ACLs or to pass statistics to backend servers. This
  can be used to take actions when queuing goes above a known level, generally
  indicating a surge of traffic or a massive slowdown on the servers. One
  possible action could be to reject new users but still accept old ones. See
  also the "avg_queue", "be_conn", and "be_sess_rate" fetches.

rand([<range>]) : integer
  Returns a random integer value within a range of <range> possible values,
  starting at zero. If the range is not specified, it defaults to 2^32, which
  gives numbers between 0 and 4294967295. It can be useful to pass some values
  needed to take some routing decisions for example, or just for debugging
  purposes. This random must not be used for security purposes.

srv_conn([<backend>/]<server>) : integer
  Returns an integer value corresponding to the number of currently established
  connections on the designated server, possibly including the connection being
  evaluated. If <backend> is omitted, then the server is looked up in the
  current backend. It can be used to use a specific farm when one server is
  full, or to inform the server about our view of the number of active
  connections with it. See also the "fe_conn", "be_conn" and "queue" fetch
  methods.

srv_is_up([<backend>/]<server>) : boolean
  Returns true when the designated server is UP, and false when it is either
  DOWN or in maintenance mode. If <backend> is omitted, then the server is
  looked up in the current backend. It is mainly used to take action based on
  an external status reported via a health check (eg: a geographical site's
  availability). Another possible use which is more of a hack consists in
  using dummy servers as boolean variables that can be enabled or disabled from
  the CLI, so that rules depending on those ACLs can be tweaked in realtime.

srv_sess_rate([<backend>/]<server>) : integer
  Returns an integer corresponding to the sessions creation rate on the
  designated server, in number of new sessions per second. If <backend> is
  omitted, then the server is looked up in the current backend. This is mostly
  used with ACLs but can make sense with logs too. This is used to switch to an
  alternate backend when an expensive or fragile one reaches too high a session
  rate, or to limit abuse of service (eg. prevent latent requests from
  overloading servers).

  Example :
        # Redirect to a separate back
        acl srv1_full srv_sess_rate(be1/srv1) gt 50
        acl srv2_full srv_sess_rate(be1/srv2) gt 50
        use_backend be2 if srv1_full or srv2_full

stopping : boolean
  Returns TRUE if the process calling the function is currently stopping. This
  can be useful for logging, or for relaxing certain checks or helping close
  certain connections upon graceful shutdown.

str(<string>) : string
  Returns a string.

table_avl([<table>]) : integer
  Returns the total number of available entries in the current proxy's
  stick-table or in the designated stick-table. See also table_cnt.

table_cnt([<table>]) : integer
  Returns the total number of entries currently in use in the current proxy's
  stick-table or in the designated stick-table. See also src_conn_cnt and
  table_avl for other entry counting methods.
```

```
12351
12352  var(<var-name>) : undefined
12353    Returns a variable with the stored type. If the variable is not set, the
12354    sample fetch fails. The name of the variable starts by an indication about its
12355    scope. The scope allowed are:
12356      "sess" : the variable is shared with all the session,
12357      "txn"  : the variable is shared with all the transaction (request and
12358               response),
12359      "req"  : the variable is shared only during the request processing,
12360      "res"  : the variable is shared only during the response processing.
12361    This prefix is followed by a name. The separator is a '.'. The name may only
12362    contain characters 'a-z', 'A-Z', '0-9' and '_'.
12363
12364
12365  7.3.3. Fetching samples at Layer 4
12366  ----------------------------------
12367
12368  The layer 4 usually describes just the transport layer which in haproxy is
12369  closest to the connection, where no content is yet made available. The fetch
12370  methods described here are usable as low as the "tcp-request connection" rule
12371  sets unless they require some future information. Those generally include
12372  TCP/IP addresses and ports, as well as elements from stick-tables related to
12373  the incoming connection. For retrieving a value from a sticky counters, the
12374  counter number can be explicitly set as 0, 1, or 2 using the pre-defined
12375  "sc0_", "sc1_", or "sc2_" prefix, or it can be specified as the first integer
12376  argument when using the "sc_" prefix. An optional table may be specified with
12377  the "sc*" form, in which case the currently tracked key will be looked up into
12378  this alternate table instead of the table currently being tracked.
12379
12380  be_id : integer
12381    Returns an integer containing the current backend's id. It can be used in
12382    frontends with responses to check which backend processed the request.
12383
12384  dst : ip
12385    This is the destination IPv4 address of the connection on the client side,
12386    which is the address the client connected to. It can be useful when running
12387    in transparent mode. It is of type IP and works on both IPv4 and IPv6 tables.
12388    On IPv6 tables, IPv4 address is mapped to its IPv6 equivalent, according to
12389    RFC 4291.
12390
12391  dst_conn : integer
12392    Returns an integer value corresponding to the number of currently established
12393    connections on the same socket including the one being evaluated. It is
12394    normally used with ACLs but can as well be used to pass the information to
12395    servers in an HTTP header or in logs. It can be used to either return a sorry
12396    page before hard-blocking, or to use a specific backend to drain new requests
12397    when the socket is considered saturated. This offers the ability to assign
12398    different limits to different listening ports or addresses. See also the
12399    "fe_conn" and "be_conn" fetches.
12400
12401  dst_port : integer
12402    Returns an integer value corresponding to the destination TCP port of the
12403    connection on the client side, which is the port the client connected to.
12404    This might be used when running in transparent mode, when assigning dynamic
12405    ports to some clients for a whole application session, to stick all users to
12406    a same server, or to pass the destination port information to a server using
12407    an HTTP header.
12408
12409  fe_id : integer
12410    Returns an integer containing the current frontend's id. It can be used in
12411    backends to check from which backend it was called, or to stick all users
12412    coming via a same frontend to the same server.
12413
12414  sc_bytes_in_rate(<ctr>[,<table>]) : integer
12415  sc0_bytes_in_rate([<table>]) : integer
```



```
12416  sc1_bytes_in_rate([<table>]) : integer
12417  sc2_bytes_in_rate([<table>]) : integer
12418    Returns the average client-to-server bytes rate from the currently tracked
12419    counters, measured in amount of bytes over the period configured in the
12420    table. See also src_bytes_in_rate.
12421
12422  sc_bytes_out_rate(<ctr>[,<table>]) : integer
12423  sc0_bytes_out_rate([<table>]) : integer
12424  sc1_bytes_out_rate([<table>]) : integer
12425  sc2_bytes_out_rate([<table>]) : integer
12426    Returns the average server-to-client bytes rate from the currently tracked
12427    counters, measured in amount of bytes over the period configured in the
12428    table. See also src_bytes_out_rate.
12429
12430  sc_clr_gpc0(<ctr>[,<table>]) : integer
12431  sc0_clr_gpc0([<table>]) : integer
12432  sc1_clr_gpc0([<table>]) : integer
12433  sc2_clr_gpc0([<table>]) : integer
12434    Clears the first General Purpose Counter associated to the currently tracked
12435    counters, and returns its previous value. Before the first invocation, the
12436    stored value is zero, so first invocation will always return zero. This is
12437    typically used as a second ACL in an expression in order to mark a connection
12438    when a first ACL was verified :
12439
12440          # block if 5 consecutive requests continue to come faster than 10 sess
12441          # per second, and reset the counter as soon as the traffic slows down.
12442          acl abuse sc0_http_req_rate gt 10
12443          acl kill  sc0_inc_gpc0 gt 5
12444          acl save  sc0_clr_gpc0 ge 0
12445          tcp-request connection accept if !abuse save
12446          tcp-request connection reject if abuse kill
12447
12448  sc_conn_cnt(<ctr>[,<table>]) : integer
12449  sc0_conn_cnt([<table>]) : integer
12450  sc1_conn_cnt([<table>]) : integer
12451  sc2_conn_cnt([<table>]) : integer
12452    Returns the cumulated number of incoming connections from currently tracked
12453    counters. See also src_conn_cnt.
12454
12455  sc_conn_cur(<ctr>[,<table>]) : integer
12456  sc0_conn_cur([<table>]) : integer
12457  sc1_conn_cur([<table>]) : integer
12458  sc2_conn_cur([<table>]) : integer
12459    Returns the current amount of concurrent connections tracking the same
12460    tracked counters. This number is automatically incremented when tracking
12461    begins and decremented when tracking stops. See also src_conn_cur.
12462
12463  sc_conn_rate(<ctr>[,<table>]) : integer
12464  sc0_conn_rate([<table>]) : integer
12465  sc1_conn_rate([<table>]) : integer
12466  sc2_conn_rate([<table>]) : integer
12467    Returns the average connection rate from the currently tracked counters,
12468    measured in amount of connections over the period configured in the table.
12469    See also src_conn_rate.
12470
12471  sc_get_gpc0(<ctr>[,<table>]) : integer
12472  sc0_get_gpc0([<table>]) : integer
12473  sc1_get_gpc0([<table>]) : integer
12474  sc2_get_gpc0([<table>]) : integer
12475    Returns the value of the first General Purpose Counter associated to the
12476    currently tracked counters. See also src_get_gpc0 and sc/sc0/sc1/sc2_inc_gpc0.
12477
12478  sc_get_gpt0(<ctr>[,<table>]) : integer
12479  sc0_get_gpt0([<table>]) : integer
12480  sc1_get_gpt0([<table>]) : integer
```

```
sc2_get_gpt0([<table>]) : integer
  Returns the value of the first General Purpose Tag associated to the
  currently tracked counters. See also src_get_gpt0.

sc_gpc0_rate(<ctr>[,<table>]) : integer
sc0_gpc0_rate([<table>]) : integer
sc1_gpc0_rate([<table>]) : integer
sc2_gpc0_rate([<table>]) : integer
  Returns the average increment rate of the first General Purpose Counter
  associated to the currently tracked counters. It reports the frequency
  which the gpc0 counter was incremented over the configured period. See also
  src_gpc0_rate, sc/sc0/sc1/sc2_get_gpc0, and sc/sc0/sc1/sc2_inc_gpc0. Note
  that the "gpc0_rate" counter must be stored in the stick-table for a value to
  be returned, as "gpc0" only holds the event count.

sc_http_err_cnt(<ctr>[,<table>]) : integer
sc0_http_err_cnt([<table>]) : integer
sc1_http_err_cnt([<table>]) : integer
sc2_http_err_cnt([<table>]) : integer
  Returns the cumulated number of HTTP errors from the currently tracked
  counters. This includes the both request errors and 4xx error responses.
  See also src_http_err_cnt.

sc_http_err_rate(<ctr>[,<table>]) : integer
sc0_http_err_rate([<table>]) : integer
sc1_http_err_rate([<table>]) : integer
sc2_http_err_rate([<table>]) : integer
  Returns the average rate of HTTP errors from the currently tracked counters,
  measured in amount of errors over the period configured in the table. This
  includes the both request errors and 4xx error responses. See also
  src_http_err_rate.

sc_http_req_cnt(<ctr>[,<table>]) : integer
sc0_http_req_cnt([<table>]) : integer
sc1_http_req_cnt([<table>]) : integer
sc2_http_req_cnt([<table>]) : integer
  Returns the cumulated number of HTTP requests from the currently tracked
  counters. This includes every started request, valid or not. See also
  src_http_req_cnt.

sc_http_req_rate(<ctr>[,<table>]) : integer
sc0_http_req_rate([<table>]) : integer
sc1_http_req_rate([<table>]) : integer
sc2_http_req_rate([<table>]) : integer
  Returns the average rate of HTTP requests from the currently tracked
  counters, measured in amount of requests over the period configured in
  the table. This includes every started request, valid or not. See also
  src_http_req_rate.

sc_inc_gpc0(<ctr>[,<table>]) : integer
sc0_inc_gpc0([<table>]) : integer
sc1_inc_gpc0([<table>]) : integer
sc2_inc_gpc0([<table>]) : integer
  Increments the first General Purpose Counter associated to the currently
  tracked counters, and returns its new value. Before the first invocation,
  the stored value is zero, so first invocation will increase it to 1 and will
  return 1. This is typically used as a second ACL in an expression in order
  to mark a connection when a first ACL was verified :

        acl abuse sc0_http_req_rate gt 10
        acl kill  sc0_inc_gpc0 gt 0
        tcp-request connection reject if abuse kill

sc_kbytes_in(<ctr>[,<table>]) : integer
sc0_kbytes_in([<table>]) : integer
```



```
sc1_kbytes_in([<table>]) : integer
sc2_kbytes_in([<table>]) : integer
  Returns the total amount of client-to-server data from the currently tracked
  counters, measured in kilobytes. The test is currently performed on 32-bit
  integers, which limits values to 4 terabytes. See also src_kbytes_in.

sc_kbytes_out(<ctr>[,<table>]) : integer
sc0_kbytes_out([<table>]) : integer
sc1_kbytes_out([<table>]) : integer
sc2_kbytes_out([<table>]) : integer
  Returns the total amount of server-to-client data from the currently tracked
  counters, measured in kilobytes. The test is currently performed on 32-bit
  integers, which limits values to 4 terabytes. See also src_kbytes_out.

sc_sess_cnt(<ctr>[,<table>]) : integer
sc0_sess_cnt([<table>]) : integer
sc1_sess_cnt([<table>]) : integer
sc2_sess_cnt([<table>]) : integer
  Returns the cumulated number of incoming connections that were transformed
  into sessions, which means that they were accepted by a "tcp-request
  connection" rule, from the currently tracked counters. A backend may count
  more sessions than connections because each connection could result in many
  backend sessions if some HTTP keep-alive is performed over the connection
  with the client. See also src_sess_cnt.

sc_sess_rate(<ctr>[,<table>]) : integer
sc0_sess_rate([<table>]) : integer
sc1_sess_rate([<table>]) : integer
sc2_sess_rate([<table>]) : integer
  Returns the average session rate from the currently tracked counters,
  measured in amount of sessions over the period configured in the table. A
  session is a connection that got past the early "tcp-request connection"
  rules. A backend may count more sessions than connections because each
  connection could result in many backend sessions if some HTTP keep-alive is
  performed over the connection with the client. See also src_sess_rate.

sc_tracked(<ctr>[,<table>]) : boolean
sc0_tracked([<table>]) : boolean
sc1_tracked([<table>]) : boolean
sc2_tracked([<table>]) : boolean
  Returns true if the designated session counter is currently being tracked by
  the current session. This can be useful when deciding whether or not we want
  to set some values in a header passed to the server.

sc_trackers(<ctr>[,<table>]) : integer
sc0_trackers([<table>]) : integer
sc1_trackers([<table>]) : integer
sc2_trackers([<table>]) : integer
  Returns the current amount of concurrent connections tracking the same
  tracked counters. This number is automatically incremented when tracking
  begins and decremented when tracking stops. It differs from sc0_conn_cur in
  that it does not rely on any stored information but on the table's reference
  count (the "use" value which is returned by "show table" on the CLI). This
  may sometimes be more suited for layer7 tracking. It can be used to tell a
  server how many concurrent connections there are from a given address for
  example.

so_id : integer
  Returns an integer containing the current listening socket's id. It is useful
  in frontends involving many "bind" lines, or to stick all users coming via a
  same socket to the same server.

src : ip
  This is the source IPv4 address of the client of the session.  It is of type
  IP and works on both IPv4 and IPv6 tables. On IPv6 tables, IPv4 addresses are
```

```
12611    mapped to their IPv6 equivalent, according to RFC 4291. Note that it is the
12612    TCP-level source address which is used, and not the address of a client
12613    behind a proxy. However if the "accept-proxy" bind directive is used, it can
12614    be the address of a client behind another PROXY-protocol compatible component
12615    for all rule sets except "tcp-request connection" which sees the real address.
12616
12617    Example:
12618         # add an HTTP header in requests with the originating address' country
12619         http-request set-header X-Country %[src,map_ip(geoip.lst)]
12620
12621  src_bytes_in_rate([<table>]) : integer
12622    Returns the average bytes rate from the incoming connection's source address
12623    in the current proxy's stick-table or in the designated stick-table, measured
12624    in amount of bytes over the period configured in the table. If the address is
12625    not found, zero is returned. See also sc/sc0/sc1/sc2_bytes_in_rate.
12626
12627  src_bytes_out_rate([<table>]) : integer
12628    Returns the average bytes rate to the incoming connection's source address in
12629    the current proxy's stick-table or in the designated stick-table, measured in
12630    amount of bytes over the period configured in the table. If the address is
12631    not found, zero is returned. See also sc/sc0/sc1/sc2_bytes_out_rate.
12632
12633  src_clr_gpc0([<table>]) : integer
12634    Clears the first General Purpose Counter associated to the incoming
12635    connection's source address in the current proxy's stick-table or in the
12636    designated stick-table, and returns its previous value. If the address is not
12637    found, an entry is created and 0 is returned. This is typically used as a
12638    second ACL in an expression in order to mark a connection when a first ACL
12639    was verified :
12640
12641          # block if 5 consecutive requests continue to come faster than 10 sess
12642          # per second, and reset the counter as soon as the traffic slows down.
12643          acl abuse src_http_req_rate gt 10
12644          acl kill  src_inc_gpc0 gt 5
12645          acl save  src_clr_gpc0 ge 0
12646          tcp-request connection accept if !abuse save
12647          tcp-request connection reject if abuse kill
12648
12649  src_conn_cnt([<table>]) : integer
12650    Returns the cumulated number of connections initiated from the current
12651    incoming connection's source address in the current proxy's stick-table or in
12652    the designated stick-table. If the address is not found, zero is returned.
12653    See also sc/sc0/sc1/sc2_conn_cnt.
12654
12655  src_conn_cur([<table>]) : integer
12656    Returns the current amount of concurrent connections initiated from the
12657    current incoming connection's source address in the current proxy's
12658    stick-table or in the designated stick-table. If the address is not found,
12659    zero is returned. See also sc/sc0/sc1/sc2_conn_cur.
12660
12661  src_conn_rate([<table>]) : integer
12662    Returns the average connection rate from the incoming connection's source
12663    address in the current proxy's stick-table or in the designated stick-table,
12664    measured in amount of connections over the period configured in the table. If
12665    the address is not found, zero is returned. See also sc/sc0/sc1/sc2_conn_rate.
12666
12667  src_get_gpc0([<table>]) : integer
12668    Returns the value of the first General Purpose Counter associated to the
12669    incoming connection's source address in the current proxy's stick-table or in
12670    the designated stick-table. If the address is not found, zero is returned.
12671    See also sc/sc0/sc1/sc2_get_gpc0 and src_inc_gpc0.
12672
12673  src_get_gpt0([<table>]) : integer
12674    Returns the value of the first General Purpose Tag associated to the
12675    incoming connection's source address in the current proxy's stick-table or in
```



```
12676    the designated stick-table. If the address is not found, zero is returned.
12677    See also sc/sc0/sc1/sc2_get_gpt0.
12678
12679  src_gpc0_rate([<table>]) : integer
12680    Returns the average increment rate of the first General Purpose Counter
12681    associated to the incoming connection's source address in the current proxy's
12682    stick-table or in the designated stick-table. It reports the frequency
12683    which the gpc0 counter was incremented over the configured period. See also
12684    sc/sc0/sc1/sc2_gpc0_rate, src_get_gpc0, and sc/sc0/sc1/sc2_inc_gpc0. Note
12685    that the "gpc0_rate" counter must be stored in the stick-table for a value to
12686    be returned, as "gpc0" only holds the event count.
12687
12688  src_http_err_cnt([<table>]) : integer
12689    Returns the cumulated number of HTTP errors from the incoming connection's
12690    source address in the current proxy's stick-table or in the designated
12691    stick-table. This includes the both request errors and 4xx error responses.
12692    See also sc/sc0/sc1/sc2_http_err_cnt. If the address is not found, zero is
12693    returned.
12694
12695  src_http_err_rate([<table>]) : integer
12696    Returns the average rate of HTTP errors from the incoming connection's source
12697    address in the current proxy's stick-table or in the designated stick-table,
12698    measured in amount of errors over the period configured in the table. This
12699    includes the both request errors and 4xx error responses. If the address is
12700    not found, zero is returned. See also sc/sc0/sc1/sc2_http_err_rate.
12701
12702  src_http_req_cnt([<table>]) : integer
12703    Returns the cumulated number of HTTP requests from the incoming connection's
12704    source address in the current proxy's stick-table or in the designated stick-
12705    table. This includes every started request, valid or not. If the address is
12706    not found, zero is returned. See also sc/sc0/sc1/sc2_http_req_cnt.
12707
12708  src_http_req_rate([<table>]) : integer
12709    Returns the average rate of HTTP requests from the incoming connection's
12710    source address in the current proxy's stick-table or in the designated stick-
12711    table, measured in amount of requests over the period configured in the
12712    table. This includes every started request, valid or not. If the address is
12713    not found, zero is returned. See also sc/sc0/sc1/sc2_http_req_rate.
12714
12715  src_inc_gpc0([<table>]) : integer
12716    Increments the first General Purpose Counter associated to the incoming
12717    connection's source address in the current proxy's stick-table or in the
12718    designated stick-table, and returns its new value. If the address is not
12719    found, an entry is created and 1 is returned. See also sc0/sc2/sc2_inc_gpc0.
12720    This is typically used as a second ACL in an expression in order to mark a
12721    connection when a first ACL was verified :
12722
12723          acl abuse src_http_req_rate gt 10
12724          acl kill  src_inc_gpc0 gt 0
12725          tcp-request connection reject if abuse kill
12726
12727  src_kbytes_in([<table>]) : integer
12728    Returns the total amount of data received from the incoming connection's
12729    source address in the current proxy's stick-table or in the designated
12730    stick-table, measured in kilobytes. If the address is not found, zero is
12731    returned. The test is currently performed on 32-bit integers, which limits
12732    values to 4 terabytes. See also sc/sc0/sc1/sc2_kbytes_in.
12733
12734  src_kbytes_out([<table>]) : integer
12735    Returns the total amount of data sent to the incoming connection's source
12736    address in the current proxy's stick-table or in the designated stick-table,
12737    measured in kilobytes. If the address is not found, zero is returned. The
12738    test is currently performed on 32-bit integers, which limits values to 4
12739    terabytes. See also sc/sc0/sc1/sc2_kbytes_out.
12740
```

```
src_port : integer
  Returns an integer value corresponding to the TCP source port of the
  connection on the client side, which is the port the client connected from.
  Usage of this function is very limited as modern protocols do not care much
  about source ports nowadays.

src_sess_cnt([<table>]) : integer
  Returns the cumulated number of connections initiated from the incoming
  connection's source IPv4 address in the current proxy's stick-table or in the
  designated stick-table, that were transformed into sessions, which means that
  they were accepted by "tcp-request" rules. If the address is not found, zero
  is returned. See also sc/sc0/sc1/sc2_sess_cnt.

src_sess_rate([<table>]) : integer
  Returns the average session rate from the incoming connection's source
  address in the current proxy's stick-table or in the designated stick-table,
  measured in amount of sessions over the period configured in the table. A
  session is a connection that went past the early "tcp-request" rules. If the
  address is not found, zero is returned. See also sc/sc0/sc1/sc2_sess_rate.

src_updt_conn_cnt([<table>]) : integer
  Creates or updates the entry associated to the incoming connection's source
  address in the current proxy's stick-table or in the designated stick-table.
  This table must be configured to store the "conn_cnt" data type, otherwise
  the match will be ignored. The current count is incremented by one, and the
  expiration timer refreshed. The updated count is returned, so this match
  can't return zero. This was used to reject service abusers based on their
  source address. Note: it is recommended to use the more complete "track-sc*"
  actions in "tcp-request" rules instead.

  Example :
        # This frontend limits incoming SSH connections to 3 per 10 second for
        # each source address, and rejects excess connections until a 10 second
        # silence is observed. At most 20 addresses are tracked.
        listen ssh
            bind :22
            mode tcp
            maxconn 100
            stick-table type ip size 20 expire 10s store conn_cnt
            tcp-request content reject if { src_updt_conn_cnt gt 3 }
            server local 127.0.0.1:22

srv_id : integer
  Returns an integer containing the server's id when processing the response.
  While it's almost only used with ACLs, it may be used for logging or
  debugging.


7.3.4. Fetching samples at Layer 5
----------------------------------

The layer 5 usually describes just the session layer which in haproxy is
closest to the session once all the connection handshakes are finished, but
when no content is yet made available. The fetch methods described here are
usable as low as the "tcp-request content" rule sets unless they require some
future information. Those generally include the results of SSL negotiations.

ssl_bc : boolean
  Returns true when the back connection was made via an SSL/TLS transport
  layer and is locally deciphered. This means the outgoing connection was made
  other a server with the "ssl" option.

ssl_bc_alg_keysize : integer
  Returns the symmetric cipher key size supported in bits when the outgoing
  connection was made over an SSL/TLS transport layer.
```



```

ssl_bc_cipher : string
  Returns the name of the used cipher when the outgoing connection was made
  over an SSL/TLS transport layer.

ssl_bc_protocol : string
  Returns the name of the used protocol when the outgoing connection was made
  over an SSL/TLS transport layer.

ssl_bc_unique_id : binary
  When the outgoing connection was made over an SSL/TLS transport layer,
  returns the TLS unique ID as defined in RFC5929 section 3. The unique id
  can be encoded to base64 using the converter: "ssl_bc_unique_id,base64".

ssl_bc_session_id : binary
  Returns the SSL ID of the back connection when the outgoing connection was
  made over an SSL/TLS transport layer. It is useful to log if we want to know
  if session was reused or not.

ssl_bc_use_keysize : integer
  Returns the symmetric cipher key size used in bits when the outgoing
  connection was made over an SSL/TLS transport layer.

ssl_c_ca_err : integer
  When the incoming connection was made over an SSL/TLS transport layer,
  returns the ID of the first error detected during verification of the client
  certificate at depth > 0, or 0 if no error was encountered during this
  verification process. Please refer to your SSL library's documentation to
  find the exhaustive list of error codes.

ssl_c_ca_err_depth : integer
  When the incoming connection was made over an SSL/TLS transport layer,
  returns the depth in the CA chain of the first error detected during the
  verification of the client certificate. If no error is encountered, 0 is
  returned.

ssl_c_der : binary
  Returns the DER formatted certificate presented by the client when the
  incoming connection was made over an SSL/TLS transport layer. When used for
  an ACL, the value(s) to match against can be passed in hexadecimal form.

ssl_c_err : integer
  When the incoming connection was made over an SSL/TLS transport layer,
  returns the ID of the first error detected during verification at depth 0, or
  0 if no error was encountered during this verification process. Please refer
  to your SSL library's documentation to find the exhaustive list of error
  codes.

ssl_c_i_dn([<entry>[,<occ>]]) : string
  When the incoming connection was made over an SSL/TLS transport layer,
  returns the full distinguished name of the issuer of the certificate
  presented by the client when no <entry> is specified, or the value of the
  first given entry found from the beginning of the DN. If a positive/negative
  occurrence number is specified as the optional second argument, it returns
  the value of the nth given entry value from the beginning/end of the DN.
  For instance, "ssl_c_i_dn(OU,2)" the second organization unit, and
  "ssl_c_i_dn(CN)" retrieves the common name.

ssl_c_key_alg : string
  Returns the name of the algorithm used to generate the key of the certificate
  presented by the client when the incoming connection was made over an SSL/TLS
  transport layer.

ssl_c_notafter : string
  Returns the end date presented by the client as a formatted string
```

```
   YYMMDDhhmmss[Z] when the incoming connection was made over an SSL/TLS
   transport layer.

ssl_c_notbefore : string
   Returns the start date presented by the client as a formatted string
   YYMMDDhhmmss[Z] when the incoming connection was made over an SSL/TLS
   transport layer.

ssl_c_s_dn([<entry>[,<occ>]]) : string
   When the incoming connection was made over an SSL/TLS transport layer,
   returns the full distinguished name of the subject of the certificate
   presented by the client when no <entry> is specified, or the value of the
   first given entry found from the beginning of the DN. If a positive/negative
   occurrence number is specified as the optional second argument, it returns
   the value of the nth given entry value from the beginning/end of the DN.
   For instance, "ssl_c_s_dn(OU,2)" the second organization unit, and
   "ssl_c_s_dn(CN)" retrieves the common name.

ssl_c_serial : binary
   Returns the serial of the certificate presented by the client when the
   incoming connection was made over an SSL/TLS transport layer. When used for
   an ACL, the value(s) to match against can be passed in hexadecimal form.

ssl_c_sha1 : binary
   Returns the SHA-1 fingerprint of the certificate presented by the client when
   the incoming connection was made over an SSL/TLS transport layer. This can be
   used to stick a client to a server, or to pass this information to a server.
   Note that the output is binary, so if you want to pass that signature to the
   server, you need to encode it in hex or base64, such as in the example below:

      http-request set-header X-SSL-Client-SHA1 %[ssl_c_sha1,hex]

ssl_c_sig_alg : string
   Returns the name of the algorithm used to sign the certificate presented by
   the client when the incoming connection was made over an SSL/TLS transport
   layer.

ssl_c_used : boolean
   Returns true if current SSL session uses a client certificate even if current
   connection uses SSL session resumption. See also "ssl_fc_has_crt".

ssl_c_verify : integer
   Returns the verify result error ID when the incoming connection was made over
   an SSL/TLS transport layer, otherwise zero if no error is encountered. Please
   refer to your SSL library's documentation for an exhaustive list of error
   codes.

ssl_c_version : integer
   Returns the version of the certificate presented by the client when the
   incoming connection was made over an SSL/TLS transport layer.

ssl_f_der : binary
   Returns the DER formatted certificate presented by the frontend when the
   incoming connection was made over an SSL/TLS transport layer. When used for
   an ACL, the value(s) to match against can be passed in hexadecimal form.

ssl_f_i_dn([<entry>[,<occ>]]) : string
   When the incoming connection was made over an SSL/TLS transport layer,
   returns the full distinguished name of the issuer of the certificate
   presented by the frontend when no <entry> is specified, or the value of the
   first given entry found from the beginning of the DN. If a positive/negative
   occurrence number is specified as the optional second argument, it returns
   the value of the nth given entry value from the beginning/end of the DN.
   For instance, "ssl_f_i_dn(OU,2)" the second organization unit, and
   "ssl_f_i_dn(CN)" retrieves the common name.
```



```

ssl_f_key_alg : string
   Returns the name of the algorithm used to generate the key of the certificate
   presented by the frontend when the incoming connection was made over an
   SSL/TLS transport layer.

ssl_f_notafter : string
   Returns the end date presented by the frontend as a formatted string
   YYMMDDhhmmss[Z] when the incoming connection was made over an SSL/TLS
   transport layer.

ssl_f_notbefore : string
   Returns the start date presented by the frontend as a formatted string
   YYMMDDhhmmss[Z] when the incoming connection was made over an SSL/TLS
   transport layer.

ssl_f_s_dn([<entry>[,<occ>]]) : string
   When the incoming connection was made over an SSL/TLS transport layer,
   returns the full distinguished name of the subject of the certificate
   presented by the frontend when no <entry> is specified, or the value of the
   first given entry found from the beginning of the DN. If a positive/negative
   occurrence number is specified as the optional second argument, it returns
   the value of the nth given entry value from the beginning/end of the DN.
   For instance, "ssl_f_s_dn(OU,2)" the second organization unit, and
   "ssl_f_s_dn(CN)" retrieves the common name.

ssl_f_serial : binary
   Returns the serial of the certificate presented by the frontend when the
   incoming connection was made over an SSL/TLS transport layer. When used for
   an ACL, the value(s) to match against can be passed in hexadecimal form.

ssl_f_sha1 : binary
   Returns the SHA-1 fingerprint of the certificate presented by the frontend
   when the incoming connection was made over an SSL/TLS transport layer. This
   can be used to know which certificate was chosen using SNI.

ssl_f_sig_alg : string
   Returns the name of the algorithm used to sign the certificate presented by
   the frontend when the incoming connection was made over an SSL/TLS transport
   layer.

ssl_f_version : integer
   Returns the version of the certificate presented by the frontend when the
   incoming connection was made over an SSL/TLS transport layer.

ssl_fc : boolean
   Returns true when the front connection was made via an SSL/TLS transport
   layer and is locally deciphered. This means it has matched a socket declared
   with a "bind" line having the "ssl" option.

   Example :
         # This passes "X-Proto: https" to servers when client connects over SSL
         listen http-https
             bind :80
             bind :443 ssl crt /etc/haproxy.pem
             http-request add-header X-Proto https if { ssl_fc }

ssl_fc_alg_keysize : integer
   Returns the symmetric cipher key size supported in bits when the incoming
   connection was made over an SSL/TLS transport layer.

ssl_fc_alpn : string
   This extracts the Application Layer Protocol Negotiation field from an
   incoming connection made via a TLS transport layer and locally deciphered by
   haproxy. The result is a string containing the protocol name advertised by
```

```
13001     the client. The SSL library must have been built with support for TLS
13002     extensions enabled (check haproxy -vv). Note that the TLS ALPN extension is
13003     not advertised unless the "alpn" keyword on the "bind" line specifies a
13004     protocol list. Also, nothing forces the client to pick a protocol from this
13005     list, any other one may be requested. The TLS ALPN extension is meant to
13006     replace the TLS NPN extension. See also "ssl_fc_npn".
13007
13008   ssl_fc_cipher : string
13009     Returns the name of the used cipher when the incoming connection was made
13010     over an SSL/TLS transport layer.
13011
13012   ssl_fc_has_crt : boolean
13013     Returns true if a client certificate is present in an incoming connection over
13014     SSL/TLS transport layer. Useful if 'verify' statement is set to 'optional'.
13015     Note: on SSL session resumption with Session ID or TLS ticket, client
13016     certificate is not present in the current connection but may be retrieved
13017     from the cache or the ticket. So prefer "ssl_c_used" if you want to check if
13018     current SSL session uses a client certificate.
13019
13020   ssl_fc_has_sni : boolean
13021     This checks for the presence of a Server Name Indication TLS extension (SNI)
13022     in an incoming connection was made over an SSL/TLS transport layer. Returns
13023     true when the incoming connection presents a TLS SNI field. This requires
13024     that the SSL library is build with support for TLS extensions enabled (check
13025     haproxy -vv).
13026
13027   ssl_fc_is_resumed: boolean
13028     Returns true if the SSL/TLS session has been resumed through the use of
13029     SSL session cache or TLS tickets.
13030
13031   ssl_fc_npn : string
13032     This extracts the Next Protocol Negotiation field from an incoming connection
13033     made via a TLS transport layer and locally deciphered by haproxy. The result
13034     is a string containing the protocol name advertised by the client. The SSL
13035     library must have been built with support for TLS extensions enabled (check
13036     haproxy -vv). Note that the TLS NPN extension is not advertised unless the
13037     "npn" keyword on the "bind" line specifies a protocol list. Also, nothing
13038     forces the client to pick a protocol from this list, any other one may be
13039     requested. Please note that the TLS NPN extension was replaced with ALPN.
13040
13041   ssl_fc_protocol : string
13042     Returns the name of the used protocol when the incoming connection was made
13043     over an SSL/TLS transport layer.
13044
13045   ssl_fc_unique_id : binary
13046     When the incoming connection was made over an SSL/TLS transport layer,
13047     returns the TLS unique ID as defined in RFC5929 section 3. The unique id
13048     can be encoded to base64 using the converter: "ssl_bc_unique_id,base64".
13049
13050   ssl_fc_session_id : binary
13051     Returns the SSL ID of the front connection when the incoming connection was
13052     made over an SSL/TLS transport layer. It is useful to stick a given client to
13053     a server. It is important to note that some browsers refresh their session ID
13054     every few minutes.
13055
13056   ssl_fc_sni : string
13057     This extracts the Server Name Indication TLS extension (SNI) field from an
13058     incoming connection made via an SSL/TLS transport layer and locally
13059     deciphered by haproxy. The result (when present) typically is a string
13060     matching the HTTPS host name (253 chars or less). The SSL library must have
13061     been built with support for TLS extensions enabled (check haproxy -vv).
13062
13063     This fetch is different from "req_ssl_sni" above in that it applies to the
13064     connection being deciphered by haproxy and not to SSL contents being blindly
13065     forwarded. See also "ssl_fc_sni_end" and "ssl_fc_sni_reg" below. This
```



```
13066     requires that the SSL library is build with support for TLS extensions
13067     enabled (check haproxy -vv).
13068
13069     ACL derivatives :
13070       ssl_fc_sni_end : suffix match
13071       ssl_fc_sni_reg : regex match
13072
13073   ssl_fc_use_keysize : integer
13074     Returns the symmetric cipher key size used in bits when the incoming
13075     connection was made over an SSL/TLS transport layer.
13076
13077
13078   7.3.5. Fetching samples from buffer contents (Layer 6)
13079   ------------------------------------------------------
13080
13081   Fetching samples from buffer contents is a bit different from the previous
13082   sample fetches above because the sampled data are ephemeral. These data can
13083   only be used when they're available and will be lost when they're forwarded.
13084   For this reason, samples fetched from buffer contents during a request cannot
13085   be used in a response for example. Even while the data are being fetched, they
13086   can change. Sometimes it is necessary to set some delays or combine multiple
13087   sample fetch methods to ensure that the expected data are complete and usable,
13088   for example through TCP request content inspection. Please see the "tcp-request
13089   content" keyword for more detailed information on the subject.
13090
13091   payload(<offset>,<length>) : binary (deprecated)
13092     This is an alias for "req.payload" when used in the context of a request (eg:
13093     "stick on", "stick match"), and for "res.payload" when used in the context of
13094     a response such as in "stick store response".
13095
13096   payload_lv(<offset1>,<length>[,<offset2>]) : binary (deprecated)
13097     This is an alias for "req.payload_lv" when used in the context of a request
13098     (eg: "stick on", "stick match"), and for "res.payload_lv" when used in the
13099     context of a response such as in "stick store response".
13100
13101   req.len : integer
13102   req_len : integer (deprecated)
13103     Returns an integer value corresponding to the number of bytes present in the
13104     request buffer. This is mostly used in ACL. It is important to understand
13105     that this test does not return false as long as the buffer is changing. This
13106     means that a check with equality to zero will almost always immediately match
13107     at the beginning of the session, while a test for more data will wait for
13108     that data to come in and return false only when haproxy is certain that no
13109     more data will come in. This test was designed to be used with TCP request
13110     content inspection.
13111
13112   req.payload(<offset>,<length>) : binary
13113     This extracts a binary block of <length> bytes and starting at byte <offset>
13114     in the request buffer. As a special case, if the <length> argument is zero,
13115     the the whole buffer from <offset> to the end is extracted. This can be used
13116     with ACLs in order to check for the presence of some content in a buffer at
13117     any location.
13118
13119     ACL alternatives :
13120       payload(<offset>,<length>) : hex binary match
13121
13122   req.payload_lv(<offset1>,<length>[,<offset2>]) : binary
13123     This extracts a binary block whose size is specified at <offset1> for <length>
13124     bytes, and which starts at <offset2> if specified or just after the length in
13125     the request buffer. The <offset2> parameter also supports relative offsets if
13126     prepended with a '+' or '-' sign.
13127
13128     ACL alternatives :
13129       payload_lv(<offset1>,<length>[,<offset2>]) : hex binary match
13130
```

```
13131     Example : please consult the example from the "stick store-response" keyword.
13132
13133   req.proto_http : boolean
13134   req_proto_http : boolean (deprecated)
13135     Returns true when data in the request buffer look like HTTP and correctly
13136     parses as such. It is the same parser as the common HTTP request parser which
13137     is used so there should be no surprises. The test does not match until the
13138     request is complete, failed or timed out. This test may be used to report the
13139     protocol in TCP logs, but the biggest use is to block TCP request analysis
13140     until a complete HTTP request is present in the buffer, for example to track
13141     a header.
13142
13143     Example:
13144           # track request counts per "base" (concatenation of Host+URL)
13145           tcp-request inspect-delay 10s
13146           tcp-request content reject if !HTTP
13147           tcp-request content track-sc0 base table req-rate
13148
13149   req.rdp_cookie([<name>]) : string
13150   rdp_cookie([<name>]) : string (deprecated)
13151     When the request buffer looks like the RDP protocol, extracts the RDP cookie
13152     <name>, or any cookie if unspecified. The parser only checks for the first
13153     cookie, as illustrated in the RDP protocol specification. The cookie name is
13154     case insensitive. Generally the "MSTS" cookie name will be used, as it can
13155     contain the user name of the client connecting to the server if properly
13156     configured on the client. The "MSTSHASH" cookie is often used as well for
13157     session stickiness to servers.
13158
13159     This differs from "balance rdp-cookie" in that any balancing algorithm may be
13160     used and thus the distribution of clients to backend servers is not linked to
13161     a hash of the RDP cookie. It is envisaged that using a balancing algorithm
13162     such as "balance roundrobin" or "balance leastconn" will lead to a more even
13163     distribution of clients to backend servers than the hash used by "balance
13164     rdp-cookie".
13165
13166     ACL derivatives :
13167       req_rdp_cookie([<name>]) : exact string match
13168
13169     Example :
13170      listen tse-farm
13171          bind 0.0.0.0:3389
13172          # wait up to 5s for an RDP cookie in the request
13173          tcp-request inspect-delay 5s
13174          tcp-request content accept if RDP_COOKIE
13175          # apply RDP cookie persistence
13176          persist rdp-cookie
13177          # Persist based on the mstshash cookie
13178          # This is only useful makes sense if
13179          # balance rdp-cookie is not used
13180          stick-table type string size 204800
13181          stick on req.rdp_cookie(mstshash)
13182          server srv1 1.1.1.1:3389
13183          server srv1 1.1.1.2:3389
13184
13185     See also : "balance rdp-cookie", "persist rdp-cookie", "tcp-request" and the
13186     "req_rdp_cookie" ACL.
13187
13188   req.rdp_cookie_cnt([name]) : integer
13189   rdp_cookie_cnt([name]) : integer (deprecated)
13190     Tries to parse the request buffer as RDP protocol, then returns an integer
13191     corresponding to the number of RDP cookies found. If an optional cookie name
13192     is passed, only cookies matching this name are considered. This is mostly
13193     used in ACL.
13194
13195     ACL derivatives :
```



```
13196       req_rdp_cookie_cnt([<name>]) : integer match
13197
13198   req.ssl_ec_ext : boolean
13199     Returns a boolean identifying if client sent the Supported Elliptic Curves
13200     Extension as defined in RFC4492, section 5.1. within the SSL ClientHello
13201     message. This can be used to present ECC compatible clients with EC
13202     certificate and to use RSA for all others, on the same IP address. Note that
13203     this only applies to raw contents found in the request buffer and not to
13204     contents deciphered via an SSL data layer, so this will not work with "bind"
13205     lines having the "ssl" option.
13206
13207   req.ssl_hello_type : integer
13208   req_ssl_hello_type : integer (deprecated)
13209     Returns an integer value containing the type of the SSL hello message found
13210     in the request buffer if the buffer contains data that parse as a complete
13211     SSL (v3 or superior) client hello message. Note that this only applies to raw
13212     contents found in the request buffer and not to contents deciphered via an
13213     SSL data layer, so this will not work with "bind" lines having the "ssl"
13214     option. This is mostly used in ACL to detect presence of an SSL hello message
13215     that is supposed to contain an SSL session ID usable for stickiness.
13216
13217   req.ssl_sni : string
13218   req_ssl_sni : string (deprecated)
13219     Returns a string containing the value of the Server Name TLS extension sent
13220     by a client in a TLS stream passing through the request buffer if the buffer
13221     contains data that parse as a complete SSL (v3 or superior) client hello
13222     message. Note that this only applies to raw contents found in the request
13223     buffer and not to contents deciphered via an SSL data layer, so this will not
13224     work with "bind" lines having the "ssl" option. SNI normally contains the
13225     name of the host the client tries to connect to (for recent browsers). SNI is
13226     useful for allowing or denying access to certain hosts when SSL/TLS is used
13227     by the client. This test was designed to be used with TCP request content
13228     inspection. If content switching is needed, it is recommended to first wait
13229     for a complete client hello (type 1), like in the example below. See also
13230     "ssl_fc_sni".
13231
13232     ACL derivatives :
13233       req_ssl_sni : exact string match
13234
13235     Examples :
13236        # Wait for a client hello for at most 5 seconds
13237        tcp-request inspect-delay 5s
13238        tcp-request content accept if { req_ssl_hello_type 1 }
13239        use_backend bk_allow if { req_ssl_sni -f allowed_sites }
13240        default_backend bk_sorry_page
13241
13242   req.ssl_st_ext : integer
13243     Returns 0 if the client didn't send a SessionTicket TLS Extension (RFC5077)
13244     Returns 1 if the client sent SessionTicket TLS Extension
13245     Returns 2 if the client also sent non-zero length TLS SessionTicket
13246     Note that this only applies to raw contents found in the request buffer and
13247     not to contents deciphered via an SSL data layer, so this will not work with
13248     "bind" lines having the "ssl" option. This can for example be used to detect
13249     whether the client sent a SessionTicket or not and stick it accordingly, if
13250     no SessionTicket then stick on SessionID or don't stick as there's no server
13251     side state is there when SessionTickets are in use.
13252
13253   req.ssl_ver : integer
13254   req_ssl_ver : integer (deprecated)
13255     Returns an integer value containing the version of the SSL/TLS protocol of a
13256     stream present in the request buffer. Both SSLv2 hello messages and SSLv3
13257     messages are supported. TLSv1 is announced as SSL version 3.1. The value is
13258     composed of the major version multiplied by 65536, added to the minor
13259     version. Note that this only applies to raw contents found in the request
13260     buffer and not to contents deciphered via an SSL data layer, so this will not
```

```
13261    work with "bind" lines having the "ssl" option. The ACL version of the test
13262    matches against a decimal notation in the form MAJOR.MINOR (eg: 3.1). This
13263    fetch is mostly used in ACL.
13264
13265    ACL derivatives :
13266      req_ssl_ver : decimal match
13267
13268  res.len : integer
13269    Returns an integer value corresponding to the number of bytes present in the
13270    response buffer. This is mostly used in ACL. It is important to understand
13271    that this test does not return false as long as the buffer is changing. This
13272    means that a check with equality to zero will almost always immediately match
13273    at the beginning of the session, while a test for more data will wait for
13274    that data to come in and return false only when haproxy is certain that no
13275    more data will come in. This test was designed to be used with TCP response
13276    content inspection.
13277
13278  res.payload(<offset>,<length>) : binary
13279    This extracts a binary block of <length> bytes and starting at byte <offset>
13280    in the response buffer. As a special case, if the <length> argument is zero,
13281    the the whole buffer from <offset> to the end is extracted. This can be used
13282    with ACLs in order to check for the presence of some content in a buffer at
13283    any location.
13284
13285  res.payload_lv(<offset1>,<length>[,<offset2>]) : binary
13286    This extracts a binary block whose size is specified at <offset1> for <length>
13287    bytes, and which starts at <offset2> if specified or just after the length in
13288    the response buffer. The <offset2> parameter also supports relative offsets
13289    if prepended with a '+' or '-' sign.
13290
13291    Example : please consult the example from the "stick store-response" keyword.
13292
13293  res.ssl_hello_type : integer
13294  rep_ssl_hello_type : integer (deprecated)
13295    Returns an integer value containing the type of the SSL hello message found
13296    in the response buffer if the buffer contains data that parses as a complete
13297    SSL (v3 or superior) hello message. Note that this only applies to raw
13298    contents found in the response buffer and not to contents deciphered via an
13299    SSL data layer, so this will not work with "server" lines having the "ssl"
13300    option. This is mostly used in ACL to detect presence of an SSL hello message
13301    that is supposed to contain an SSL session ID usable for stickiness.
13302
13303  wait_end : boolean
13304    This fetch either returns true when the inspection period is over, or does
13305    not fetch. It is only used in ACLs, in conjunction with content analysis to
13306    avoid returning a wrong verdict early.  It may also be used to delay some
13307    actions, such as a delayed reject for some special addresses. Since it either
13308    stops the rules evaluation or immediately returns true, it is recommended to
13309    use this acl as the last one in a rule.  Please note that the default ACL
13310    "WAIT_END" is always usable without prior declaration. This test was designed
13311    to be used with TCP request content inspection.
13312
13313    Examples :
13314       # delay every incoming request by 2 seconds
13315       tcp-request inspect-delay 2s
13316       tcp-request content accept if WAIT_END
13317
13318       # don't immediately tell bad guys they are rejected
13319       tcp-request inspect-delay 10s
13320       acl goodguys src 10.0.0.0/24
13321       acl badguys  src 10.0.1.0/24
13322       tcp-request content accept if goodguys
13323       tcp-request content reject if badguys WAIT_END
13324       tcp-request content reject
13325
```



```
13326
13327  7.3.6. Fetching HTTP samples (Layer 7)
13328  --------------------------------------
13329
13330  It is possible to fetch samples from HTTP contents, requests and responses.
13331  This application layer is also called layer 7. It is only possible to fetch the
13332  data in this section when a full HTTP request or response has been parsed from
13333  its respective request or response buffer. This is always the case with all
13334  HTTP specific rules and for sections running with "mode http". When using TCP
13335  content inspection, it may be necessary to support an inspection delay in order
13336  to let the request or response come in first. These fetches may require a bit
13337  more CPU resources than the layer 4 ones, but not much since the request and
13338  response are indexed.
13339
13340  base : string
13341    This returns the concatenation of the first Host header and the path part of
13342    the request, which starts at the first slash and ends before the question
13343    mark. It can be useful in virtual hosted environments to detect URL abuses as
13344    well as to improve shared caches efficiency. Using this with a limited size
13345    stick table also allows one to collect statistics about most commonly
13346    requested objects by host/path. With ACLs it can allow simple content
13347    switching rules involving the host and the path at the same time, such as
13348    "www.example.com/favicon.ico". See also "path" and "uri".
13349
13350    ACL derivatives :
13351      base     : exact string match
13352      base_beg : prefix match
13353      base_dir : subdir match
13354      base_dom : domain match
13355      base_end : suffix match
13356      base_len : length match
13357      base_reg : regex match
13358      base_sub : substring match
13359
13360  base32 : integer
13361    This returns a 32-bit hash of the value returned by the "base" fetch method
13362    above. This is useful to track per-URL activity on high traffic sites without
13363    having to store all URLs. Instead a shorter hash is stored, saving a lot of
13364    memory. The output type is an unsigned integer. The hash function used is
13365    SDBM with full avalanche on the output. Technically, base32 is exactly equal
13366    to "base,sdbm(1)".
13367
13368  base32+src : binary
13369    This returns the concatenation of the base32 fetch above and the src fetch
13370    below. The resulting type is of type binary, with a size of 8 or 20 bytes
13371    depending on the source address family. This can be used to track per-IP,
13372    per-URL counters.
13373
13374  capture.req.hdr(<idx>) : string
13375    This extracts the content of the header captured by the "capture request
13376    header", idx is the position of the capture keyword in the configuration.
13377    The first entry is an index of 0. See also: "capture request header".
13378
13379  capture.req.method : string
13380    This extracts the METHOD of an HTTP request. It can be used in both request
13381    and response. Unlike "method", it can be used in both request and response
13382    because it's allocated.
13383
13384  capture.req.uri : string
13385    This extracts the request's URI, which starts at the first slash and ends
13386    before the first space in the request (without the host part). Unlike "path"
13387    and "url", it can be used in both request and response because it's
13388    allocated.
13389
13390  capture.req.ver : string
```

```
13391     This extracts the request's HTTP version and returns either "HTTP/1.0" or
13392     "HTTP/1.1". Unlike "req.ver", it can be used in both request, response, and
13393     logs because it relies on a persistent flag.
13394
13395   capture.res.hdr(<idx>) : string
13396     This extracts the content of the header captured by the "capture response
13397     header", idx is the position of the capture keyword in the configuration.
13398     The first entry is an index of 0.
13399     See also: "capture response header"
13400
13401   capture.res.ver : string
13402     This extracts the response's HTTP version and returns either "HTTP/1.0" or
13403     "HTTP/1.1". Unlike "res.ver", it can be used in logs because it relies on a
13404     persistent flag.
13405
13406   req.body : binary
13407     This returns the HTTP request's available body as a block of data. It
13408     requires that the request body has been buffered made available using
13409     "option http-buffer-request". In case of chunked-encoded body, currently only
13410     the first chunk is analyzed.
13411
13412   req.body_param([<name>) : string
13413     This fetch assumes that the body of the POST request is url-encoded. The user
13414     can check if the "content-type" contains the value
13415     "application/x-www-form-urlencoded". This extracts the first occurrence of the
13416     parameter <name> in the body, which ends before '&'. The parameter name is
13417     case-sensitive. If no name is given, any parameter will match, and the first
13418     one will be returned. The result is a string corresponding to the value of the
13419     parameter <name> as presented in the request body (no URL decoding is
13420     performed). Note that the ACL version of this fetch iterates over multiple
13421     parameters and will iteratively report all parameters values if no name is
13422     given.
13423
13424   req.body_len : integer
13425     This returns the length of the HTTP request's available body in bytes. It may
13426     be lower than the advertised length if the body is larger than the buffer. It
13427     requires that the request body has been buffered made available using
13428     "option http-buffer-request".
13429
13430   req.body_size : integer
13431     This returns the advertised length of the HTTP request's body in bytes. It
13432     will represent the advertised Content-Length header, or the size of the first
13433     chunk in case of chunked encoding. In order to parse the chunks, it requires
13434     that the request body has been buffered made available using
13435     "option http-buffer-request".
13436
13437   req.cook([<name>]) : string
13438   cook([<name>]) : string (deprecated)
13439     This extracts the last occurrence of the cookie name <name> on a "Cookie"
13440     header line from the request, and returns its value as string. If no name is
13441     specified, the first cookie value is returned. When used with ACLs, all
13442     matching cookies are evaluated. Spaces around the name and the value are
13443     ignored as requested by the Cookie header specification (RFC6265). The cookie
13444     name is case-sensitive. Empty cookies are valid, so an empty cookie may very
13445     well return an empty value if it is present. Use the "found" match to detect
13446     presence. Use the res.cook() variant for response cookies sent by the server.
13447
13448     ACL derivatives :
13449       cook([<name>])     : exact string match
13450       cook_beg([<name>]) : prefix match
13451       cook_dir([<name>]) : subdir match
13452       cook_dom([<name>]) : domain match
13453       cook_end([<name>]) : suffix match
13454       cook_len([<name>]) : length match
13455       cook_reg([<name>]) : regex match
```



```
13456       cook_sub([<name>]) : substring match
13457
13458   req.cook_cnt([<name>]) : integer
13459   cook_cnt([<name>]) : integer (deprecated)
13460     Returns an integer value representing the number of occurrences of the cookie
13461     <name> in the request, or all cookies if <name> is not specified.
13462
13463   req.cook_val([<name>]) : integer
13464   cook_val([<name>]) : integer (deprecated)
13465     This extracts the last occurrence of the cookie name <name> on a "Cookie"
13466     header line from the request, and converts its value to an integer which is
13467     returned. If no name is specified, the first cookie value is returned. When
13468     used in ACLs, all matching names are iterated over until a value matches.
13469
13470   cookie([<name>]) : string (deprecated)
13471     This extracts the last occurrence of the cookie name <name> on a "Cookie"
13472     header line from the request, or a "Set-Cookie" header from the response, and
13473     returns its value as a string. A typical use is to get multiple clients
13474     sharing a same profile use the same server. This can be similar to what
13475     "appsession" did with the "request-learn" statement, but with support for
13476     multi-peer synchronization and state keeping across restarts. If no name is
13477     specified, the first cookie value is returned. This fetch should not be used
13478     anymore and should be replaced by req.cook() or res.cook() instead as it
13479     ambiguously uses the direction based on the context where it is used.
13480
13481   hdr([<name>[,<occ>]]) : string
13482     This is equivalent to req.hdr() when used on requests, and to res.hdr() when
13483     used on responses. Please refer to these respective fetches for more details.
13484     In case of doubt about the fetch direction, please use the explicit ones.
13485     Note that contrary to the hdr() sample fetch method, the hdr_* ACL keywords
13486     unambiguously apply to the request headers.
13487
13488   req.fhdr(<name>[,<occ>]) : string
13489     This extracts the last occurrence of header <name> in an HTTP request. When
13490     used from an ACL, all occurrences are iterated over until a match is found.
13491     Optionally, a specific occurrence might be specified as a position number.
13492     Positive values indicate a position from the first occurrence, with 1 being
13493     the first one. Negative values indicate positions relative to the last one,
13494     with -1 being the last one. It differs from req.hdr() in that any commas
13495     present in the value are returned and are not used as delimiters. This is
13496     sometimes useful with headers such as User-Agent.
13497
13498   req.fhdr_cnt([<name>]) : integer
13499     Returns an integer value representing the number of occurrences of request
13500     header field name <name>, or the total number of header fields if <name> is
13501     not specified. Contrary to its req.hdr_cnt() cousin, this function returns
13502     the number of full line headers and does not stop on commas.
13503
13504   req.hdr([<name>[,<occ>]]) : string
13505     This extracts the last occurrence of header <name> in an HTTP request. When
13506     used from an ACL, all occurrences are iterated over until a match is found.
13507     Optionally, a specific occurrence might be specified as a position number.
13508     Positive values indicate a position from the first occurrence, with 1 being
13509     the first one. Negative values indicate positions relative to the last one,
13510     with -1 being the last one. A typical use is with the X-Forwarded-For header
13511     once converted to IP, associated with an IP stick-table. The function
13512     considers any comma as a delimiter for distinct values. If full-line headers
13513     are desired instead, use req.fhdr(). Please carefully check RFC2616 to know
13514     how certain headers are supposed to be parsed. Also, some of them are case
13515     insensitive (eg: Connection).
13516
13517     ACL derivatives :
13518       hdr([<name>[,<occ>]])     : exact string match
13519       hdr_beg([<name>[,<occ>]]) : prefix match
13520       hdr_dir([<name>[,<occ>]]) : subdir match
```

```
13521      hdr_dom([<name>[,<occ>]]) : domain match
13522      hdr_end([<name>[,<occ>]]) : suffix match
13523      hdr_len([<name>[,<occ>]]) : length match
13524      hdr_reg([<name>[,<occ>]]) : regex match
13525      hdr_sub([<name>[,<occ>]]) : substring match
13526
13527  req.hdr_cnt([<name>]) : integer
13528  hdr_cnt([<header>]) : integer (deprecated)
13529    Returns an integer value representing the number of occurrences of request
13530    header field name <name>, or the total number of header field values if
13531    <name> is not specified. It is important to remember that one header line may
13532    count as several headers if it has several values. The function considers any
13533    comma as a delimiter for distinct values. If full-line headers are desired
13534    instead, req.fhdr_cnt() should be used instead. With ACLs, it can be used to
13535    detect presence, absence or abuse of a specific header, as well as to block
13536    request smuggling attacks by rejecting requests which contain more than one
13537    of certain headers. See "req.hdr" for more information on header matching.
13538
13539  req.hdr_ip([<name>[,<occ>]]) : ip
13540  hdr_ip([<name>[,<occ>]]) : ip (deprecated)
13541    This extracts the last occurrence of header <name> in an HTTP request,
13542    converts it to an IPv4 or IPv6 address and returns this address. When used
13543    with ACLs, all occurrences are checked, and if <name> is omitted, every value
13544    of every header is checked. Optionally, a specific occurrence might be
13545    specified as a position number. Positive values indicate a position from the
13546    first occurrence, with 1 being the first one.  Negative values indicate
13547    positions relative to the last one, with -1 being the last one. A typical use
13548    is with the X-Forwarded-For and X-Client-IP headers.
13549
13550  req.hdr_val([<name>[,<occ>]]) : integer
13551  hdr_val([<name>[,<occ>]]) : integer (deprecated)
13552    This extracts the last occurrence of header <name> in an HTTP request, and
13553    converts it to an integer value. When used with ACLs, all occurrences are
13554    checked, and if <name> is omitted, every value of every header is checked.
13555    Optionally, a specific occurrence might be specified as a position number.
13556    Positive values indicate a position from the first occurrence, with 1 being
13557    the first one. Negative values indicate positions relative to the last one,
13558    with -1 being the last one. A typical use is with the X-Forwarded-For header.
13559
13560  http_auth(<userlist>) : boolean
13561    Returns a boolean indicating whether the authentication data received from
13562    the client match a username & password stored in the specified userlist. This
13563    fetch function is not really useful outside of ACLs. Currently only http
13564    basic auth is supported.
13565
13566  http_auth_group(<userlist>) : string
13567    Returns a string corresponding to the user name found in the authentication
13568    data received from the client if both the user name and password are valid
13569    according to the specified userlist. The main purpose is to use it in ACLs
13570    where it is then checked whether the user belongs to any group within a list.
13571    This fetch function is not really useful outside of ACLs. Currently only http
13572    basic auth is supported.
13573
13574    ACL derivatives :
13575      http_auth_group(<userlist>) : group ...
13576      Returns true when the user extracted from the request and whose password is
13577      valid according to the specified userlist belongs to at least one of the
13578      groups.
13579
13580  http_first_req : boolean
13581    Returns true when the request being processed is the first one of the
13582    connection. This can be used to add or remove headers that may be missing
13583    from some requests when a request is not the first one, or to help grouping
13584    requests in the logs.
13585
```



```
13586  method : integer + string
13587    Returns an integer value corresponding to the method in the HTTP request. For
13588    example, "GET" equals 1 (check sources to establish the matching). Value 9
13589    means "other method" and may be converted to a string extracted from the
13590    stream. This should not be used directly as a sample, this is only meant to
13591    be used from ACLs, which transparently convert methods from patterns to these
13592    integer + string values. Some predefined ACL already check for most common
13593    methods.
13594
13595    ACL derivatives :
13596      method : case insensitive method match
13597
13598    Example :
13599        # only accept GET and HEAD requests
13600        acl valid_method method GET HEAD
13601        http-request deny if ! valid_method
13602
13603  path : string
13604    This extracts the request's URL path, which starts at the first slash and
13605    ends before the question mark (without the host part). A typical use is with
13606    prefetch-capable caches, and with portals which need to aggregate multiple
13607    information from databases and keep them in caches. Note that with outgoing
13608    caches, it would be wiser to use "url" instead. With ACLs, it's typically
13609    used to match exact file names (eg: "/login.php"), or directory parts using
13610    the derivative forms. See also the "url" and "base" fetch methods.
13611
13612    ACL derivatives :
13613      path     : exact string match
13614      path_beg : prefix match
13615      path_dir : subdir match
13616      path_dom : domain match
13617      path_end : suffix match
13618      path_len : length match
13619      path_reg : regex match
13620      path_sub : substring match
13621
13622  query : string
13623    This extracts the request's query string, which starts after the first
13624    question mark. If no question mark is present, this fetch returns nothing. If
13625    a question mark is present but nothing follows, it returns an empty string.
13626    This means it's possible to easily know whether a query string is present
13627    using the "found" matching method. This fetch is the completemnt of "path"
13628    which stops before the question mark.
13629
13630  req.hdr_names([<delim>]) : string
13631    This builds a string made from the concatenation of all header names as they
13632    appear in the request when the rule is evaluated. The default delimiter is
13633    the comma (',') but it may be overridden as an optional argument <delim>. In
13634    this case, only the first character of <delim> is considered.
13635
13636  req.ver : string
13637  req_ver : string (deprecated)
13638    Returns the version string from the HTTP request, for example "1.1". This can
13639    be useful for logs, but is mostly there for ACL. Some predefined ACL already
13640    check for versions 1.0 and 1.1.
13641
13642    ACL derivatives :
13643      req_ver : exact string match
13644
13645  res.comp : boolean
13646    Returns the boolean "true" value if the response has been compressed by
13647    HAProxy, otherwise returns boolean "false". This may be used to add
13648    information in the logs.
13649
13650  res.comp_algo : string
```

```
13651     Returns a string containing the name of the algorithm used if the response
13652     was compressed by HAProxy, for example : "deflate". This may be used to add
13653     some information in the logs.
13654
13655   res.cook([<name>]) : string
13656   scook([<name>]) : string (deprecated)
13657     This extracts the last occurrence of the cookie name <name> on a "Set-Cookie"
13658     header line from the response, and returns its value as string. If no name is
13659     specified, the first cookie value is returned.
13660
13661     ACL derivatives :
13662       scook([<name>] : exact string match
13663
13664   res.cook_cnt([<name>]) : integer
13665   scook_cnt([<name>]) : integer (deprecated)
13666     Returns an integer value representing the number of occurrences of the cookie
13667     <name> in the response, or all cookies if <name> is not specified. This is
13668     mostly useful when combined with ACLs to detect suspicious responses.
13669
13670   res.cook_val([<name>]) : integer
13671   scook_val([<name>]) : integer (deprecated)
13672     This extracts the last occurrence of the cookie name <name> on a "Set-Cookie"
13673     header line from the response, and converts its value to an integer which is
13674     returned. If no name is specified, the first cookie value is returned.
13675
13676   res.fhdr([<name>[,<occ>]]) : string
13677     This extracts the last occurrence of header <name> in an HTTP response, or of
13678     the last header if no <name> is specified. Optionally, a specific occurrence
13679     might be specified as a position number. Positive values indicate a position
13680     from the first occurrence, with 1 being the first one. Negative values
13681     indicate positions relative to the last one, with -1 being the last one. It
13682     differs from res.hdr() in that any commas present in the value are returned
13683     and are not used as delimiters. If this is not desired, the res.hdr() fetch
13684     should be used instead. This is sometimes useful with headers such as Date or
13685     Expires.
13686
13687   res.fhdr_cnt([<name>]) : integer
13688     Returns an integer value representing the number of occurrences of response
13689     header field name <name>, or the total number of header fields if <name> is
13690     not specified. Contrary to its res.hdr_cnt() cousin, this function returns
13691     the number of full line headers and does not stop on commas. If this is not
13692     desired, the res.hdr_cnt() fetch should be used instead.
13693
13694   res.hdr([<name>[,<occ>]]) : string
13695   shdr([<name>[,<occ>]]) : string (deprecated)
13696     This extracts the last occurrence of header <name> in an HTTP response, or of
13697     the last header if no <name> is specified. Optionally, a specific occurrence
13698     might be specified as a position number. Positive values indicate a position
13699     from the first occurrence, with 1 being the first one. Negative values
13700     indicate positions relative to the last one, with -1 being the last one. This
13701     can be useful to learn some data into a stick-table. The function considers
13702     any comma as a delimiter for distinct values. If this is not desired, the
13703     res.fhdr() fetch should be used instead.
13704
13705     ACL derivatives :
13706       shdr([<name>[,<occ>]])     : exact string match
13707       shdr_beg([<name>[,<occ>]]) : prefix match
13708       shdr_dir([<name>[,<occ>]]) : subdir match
13709       shdr_dom([<name>[,<occ>]]) : domain match
13710       shdr_end([<name>[,<occ>]]) : suffix match
13711       shdr_len([<name>[,<occ>]]) : length match
13712       shdr_reg([<name>[,<occ>]]) : regex match
13713       shdr_sub([<name>[,<occ>]]) : substring match
13714
13715   res.hdr_cnt([<name>]) : integer
```



```
13716   shdr_cnt([<name>]) : integer (deprecated)
13717     Returns an integer value representing the number of occurrences of response
13718     header field name <name>, or the total number of header fields if <name> is
13719     not specified. The function considers any comma as a delimiter for distinct
13720     values. If this is not desired, the res.fhdr_cnt() fetch should be used
13721     instead.
13722
13723   res.hdr_ip([<name>[,<occ>]]) : ip
13724   shdr_ip([<name>[,<occ>]]) : ip (deprecated)
13725     This extracts the last occurrence of header <name> in an HTTP response,
13726     convert it to an IPv4 or IPv6 address and returns this address. Optionally, a
13727     specific occurrence might be specified as a position number. Positive values
13728     indicate a position from the first occurrence, with 1 being the first one.
13729     Negative values indicate positions relative to the last one, with -1 being
13730     the last one. This can be useful to learn some data into a stick table.
13731
13732   res.hdr_names([<delim>]) : string
13733     This builds a string made from the concatenation of all header names as they
13734     appear in the response when the rule is evaluated. The default delimiter is
13735     the comma (',') but it may be overridden as an optional argument <delim>. In
13736     this case, only the first character of <delim> is considered.
13737
13738   res.hdr_val([<name>[,<occ>]]) : integer
13739   shdr_val([<name>[,<occ>]]) : integer (deprecated)
13740     This extracts the last occurrence of header <name> in an HTTP response, and
13741     converts it to an integer value. Optionally, a specific occurrence might be
13742     specified as a position number. Positive values indicate a position from the
13743     first occurrence, with 1 being the first one. Negative values indicate
13744     positions relative to the last one, with -1 being the last one. This can be
13745     useful to learn some data into a stick table.
13746
13747   res.ver : string
13748   resp_ver : string (deprecated)
13749     Returns the version string from the HTTP response, for example "1.1". This
13750     can be useful for logs, but is mostly there for ACL.
13751
13752     ACL derivatives :
13753       resp_ver : exact string match
13754
13755   set-cookie([<name>]) : string (deprecated)
13756     This extracts the last occurrence of the cookie name <name> on a "Set-Cookie"
13757     header line from the response and uses the corresponding value to match. This
13758     can be comparable to what "appsession" did with default options, but with
13759     support for multi-peer synchronization and state keeping across restarts.
13760
13761     This fetch function is deprecated and has been superseded by the "res.cook"
13762     fetch. This keyword will disappear soon.
13763
13764   status : integer
13765     Returns an integer containing the HTTP status code in the HTTP response, for
13766     example, 302. It is mostly used within ACLs and integer ranges, for example,
13767     to remove any Location header if the response is not a 3xx.
13768
13769   url : string
13770     This extracts the request's URL as presented in the request. A typical use is
13771     with prefetch-capable caches, and with portals which need to aggregate
13772     multiple information from databases and keep them in caches. With ACLs, using
13773     "path" is preferred over using "url", because clients may send a full URL as
13774     is normally done with proxies. The only real use is to match "*" which does
13775     not match in "path", and for which there is already a predefined ACL. See
13776     also "path" and "base".
13777
13778     ACL derivatives :
13779       url     : exact string match
13780       url_beg : prefix match
```

```
     url_dir : subdir match
     url_dom : domain match
     url_end : suffix match
     url_len : length match
     url_reg : regex match
     url_sub : substring match

url_ip : ip
  This extracts the IP address from the request's URL when the host part is
  presented as an IP address. Its use is very limited. For instance, a
  monitoring system might use this field as an alternative for the source IP in
  order to test what path a given source address would follow, or to force an
  entry in a table for a given source address. With ACLs it can be used to
  restrict access to certain systems through a proxy, for example when combined
  with option "http_proxy".

url_port : integer
  This extracts the port part from the request's URL. Note that if the port is
  not specified in the request, port 80 is assumed. With ACLs it can be used to
  restrict access to certain systems through a proxy, for example when combined
  with option "http_proxy".

urlp([<name>[,<delim>]]) : string
url_param([<name>[,<delim>]]) : string
  This extracts the first occurrence of the parameter <name> in the query
  string, which begins after either '?' or <delim>, and which ends before '&',
  ';' or <delim>. The parameter name is case-sensitive. If no name is given,
  any parameter will match, and the first one will be returned. The result is
  a string corresponding to the value of the parameter <name> as presented in
  the request (no URL decoding is performed). This can be used for session
  stickiness based on a client ID, to extract an application cookie passed as a
  URL parameter, or in ACLs to apply some checks. Note that the ACL version of
  this fetch iterates over multiple parameters and will iteratively report all
  parameters values if no name is given

  ACL derivatives :
    urlp(<name>[,<delim>])     : exact string match
    urlp_beg(<name>[,<delim>]) : prefix match
    urlp_dir(<name>[,<delim>]) : subdir match
    urlp_dom(<name>[,<delim>]) : domain match
    urlp_end(<name>[,<delim>]) : suffix match
    urlp_len(<name>[,<delim>]) : length match
    urlp_reg(<name>[,<delim>]) : regex match
    urlp_sub(<name>[,<delim>]) : substring match


  Example :
      # match http://example.com/foo?PHPSESSIONID=some_id
      stick on urlp(PHPSESSIONID)
      # match http://example.com/foo;JSESSIONID=some_id
      stick on urlp(JSESSIONID,;)

urlp_val([<name>[,<delim>])] : integer
  See "urlp" above. This one extracts the URL parameter <name> in the request
  and converts it to an integer value. This can be used for session stickiness
  based on a user ID for example, or with ACLs to match a page number or price.


7.4. Pre-defined ACLs
---------------------

Some predefined ACLs are hard-coded so that they do not have to be declared in
every frontend which needs them. They all have their names in upper case in
order to avoid confusion. Their equivalence is provided below.

```



```
ACL name          Equivalent to                Usage
--------------+----------------------------+--------------------------------
FALSE           always_false                 never match
HTTP            req_proto_http               match if protocol is valid HTTP
HTTP_1.0        req_ver 1.0                  match HTTP version 1.0
HTTP_1.1        req_ver 1.1                  match HTTP version 1.1
HTTP_CONTENT    hdr_val(content-length) gt 0 match an existing content-length
HTTP_URL_ABS    url_reg ^[^/:]*://           match absolute URL with scheme
HTTP_URL_SLASH  url_beg /                    match URL beginning with "/"
HTTP_URL_STAR   url     *                    match URL equal to "*"
LOCALHOST       src 127.0.0.1/8              match connection from local host
METH_CONNECT    method  CONNECT              match HTTP CONNECT method
METH_GET        method  GET HEAD             match HTTP GET or HEAD method
METH_HEAD       method  HEAD                 match HTTP HEAD method
METH_OPTIONS    method  OPTIONS              match HTTP OPTIONS method
METH_POST       method  POST                 match HTTP POST method
METH_TRACE      method  TRACE                match HTTP TRACE method
RDP_COOKIE      req_rdp_cookie_cnt gt 0      match presence of an RDP cookie
REQ_CONTENT     req_len gt 0                 match data in the request buffer
TRUE            always_true                  always match
WAIT_END        wait_end                     wait for end of content analysis
--------------+----------------------------+--------------------------------


8. Logging
----------

One of HAProxy's strong points certainly lies is its precise logs. It probably
provides the finest level of information available for such a product, which is
very important for troubleshooting complex environments. Standard information
provided in logs include client ports, TCP/HTTP state timers, precise session
state at termination and precise termination cause, information about decisions
to direct traffic to a server, and of course the ability to capture arbitrary
headers.

In order to improve administrators reactivity, it offers a great transparency
about encountered problems, both internal and external, and it is possible to
send logs to different sources at the same time with different level filters :

  - global process-level logs (system errors, start/stop, etc..)
  - per-instance system and internal errors (lack of resource, bugs, ...)
  - per-instance external troubles (servers up/down, max connections)
  - per-instance activity (client connections), either at the establishment or
    at the termination.
  - per-request control of log-level, eg:
        http-request set-log-level silent if sensitive_request

The ability to distribute different levels of logs to different log servers
allow several production teams to interact and to fix their problems as soon
as possible. For example, the system team might monitor system-wide errors,
while the application team might be monitoring the up/down for their servers in
real time, and the security team might analyze the activity logs with one hour
delay.


8.1. Log levels
---------------

TCP and HTTP connections can be logged with information such as the date, time,
source IP address, destination address, connection duration, response times,
HTTP request, HTTP return code, number of bytes transmitted, conditions
in which the session ended, and even exchanged cookies values. For example
track a particular user's problems. All messages may be sent to up to two
syslog servers. Check the "log" keyword in section 4.2 for more information
about log facilities.
```

8.2. Log formats

HAProxy supports 5 log formats. Several fields are common between these formats and will be detailed in the following sections. A few of them may vary slightly with the configuration, due to indicators specific to certain options. The supported formats are as follows :

- the default format, which is very basic and very rarely used. It only provides very basic information about the incoming connection at the moment it is accepted : source IP:port, destination IP:port, and frontend-name. This mode will eventually disappear so it will not be described to great extents.

- the TCP format, which is more advanced. This format is enabled when "option tcplog" is set on the frontend. HAProxy will then usually wait for the connection to terminate before logging. This format provides much richer information, such as timers, connection counts, queue size, etc... This format is recommended for pure TCP proxies.

- the HTTP format, which is the most advanced for HTTP proxying. This format is enabled when "option httplog" is set on the frontend. It provides the same information as the TCP format with some HTTP-specific fields such as the request, the status code, and captures of headers and cookies. This format is recommended for HTTP proxies.

- the CLF HTTP format, which is equivalent to the HTTP format, but with the fields arranged in the same order as the CLF format. In this mode, all timers, captures, flags, etc... appear one per field after the end of the common fields, in the same order they appear in the standard HTTP format.

- the custom log format, allows you to make your own log line.

Next sections will go deeper into details for each of these formats. Format specification will be performed on a "field" basis. Unless stated otherwise, a field is a portion of text delimited by any number of spaces. Since syslog servers are susceptible of inserting fields at the beginning of a line, it is always assumed that the first field is the one containing the process name and identifier.

Note : Since log lines may be quite long, the log examples in sections below might be broken into multiple lines. The example log lines will be prefixed with 3 closing angle brackets ('>>>') and each time a log is broken into multiple lines, each non-final line will end with a backslash ('\') and the next line will start indented by two characters.

8.2.1. Default log format

This format is used when no specific option is set. The log is emitted as soon as the connection is accepted. One should note that this currently is the only format which logs the request's destination IP and ports.

```
  Example :
        listen www
            mode http
            log global
            server srv1 127.0.0.1:8000

    >>> Feb  6 12:12:09 localhost \
          haproxy[14385]: Connect from 10.0.1.2:33312 to 10.0.3.31:8012 \
          (www/HTTP)
```



```
  Field   Format                                Extract from the example above
      1   process_name '[' pid ']:'                            haproxy[14385]:
      2   'Connect from'                                          Connect from
      3   source_ip ':' source_port                            10.0.1.2:33312
      4   'to'                                                             to
      5   destination_ip ':' destination_port                  10.0.3.31:8012
      6   '(' frontend_name '/' mode ')'                           (www/HTTP)

Detailed fields description :
  - "source_ip" is the IP address of the client which initiated the connection.
  - "source_port" is the TCP port of the client which initiated the connection.
  - "destination_ip" is the IP address the client connected to.
  - "destination_port" is the TCP port the client connected to.
  - "frontend_name" is the name of the frontend (or listener) which received
    and processed the connection.
  - "mode is the mode the frontend is operating (TCP or HTTP).

In case of a UNIX socket, the source and destination addresses are marked as
"unix:" and the ports reflect the internal ID of the socket which accepted the
connection (the same ID as reported in the stats).

It is advised not to use this deprecated format for newer installations as it
will eventually disappear.
```

8.2.2. TCP log format

The TCP format is used when "option tcplog" is specified in the frontend, and is the recommended format for pure TCP proxies. It provides a lot of precious information for troubleshooting. Since this format includes timers and byte counts, the log is normally emitted at the end of the session. It can be emitted earlier if "option logasap" is specified, which makes sense in most environments with long sessions such as remote terminals. Sessions which match the "monitor" rules are never logged. It is also possible not to emit logs for sessions for which no data were exchanged between the client and the server, by specifying "option dontlognull" in the frontend. Successful connections will not be logged if "option dontlog-normal" is specified in the frontend. A few fields may slightly vary depending on some configuration options, those are marked with a star ('*') after the field name below.

```
  Example :
        frontend fnt
            mode tcp
            option tcplog
            log global
            default_backend bck

        backend bck
            server srv1 127.0.0.1:8000

    >>> Feb  6 12:12:56 localhost \
          haproxy[14387]: 10.0.1.2:33313 [06/Feb/2009:12:12:51.443] fnt \
          bck/srv1 0/0/5007 212 -- 0/0/0/0/3 0/0

  Field   Format                                Extract from the example above
      1   process_name '[' pid ']:'                            haproxy[14387]:
      2   client_ip ':' client_port                             10.0.1.2:33313
      3   '[' accept_date ']'                      [06/Feb/2009:12:12:51.443]
      4   frontend_name                                                   fnt
      5   backend_name '/' server_name                                bck/srv1
      6   Tw '/' Tc '/' Tt*                                          0/0/5007
      7   bytes_read*                                                      212
      8   termination_state                                                 --
```

```
14041        9   actconn '/' feconn '/' beconn '/' srv_conn '/' retries*    0/0/0/0/3
14042       10   srv_queue '/' backend_queue                                  0/0
14043
14044  Detailed fields description :
14045    - "client_ip" is the IP address of the client which initiated the TCP
14046      connection to haproxy. If the connection was accepted on a UNIX socket
14047      instead, the IP address would be replaced with the word "unix". Note that
14048      when the connection is accepted on a socket configured with "accept-proxy"
14049      and the PROXY protocol is correctly used, then the logs will reflect the
14050      forwarded connection's information.
14051
14052    - "client_port" is the TCP port of the client which initiated the connection.
14053      If the connection was accepted on a UNIX socket instead, the port would be
14054      replaced with the ID of the accepting socket, which is also reported in the
14055      stats interface.
14056
14057    - "accept_date" is the exact date when the connection was received by haproxy
14058      (which might be very slightly different from the date observed on the
14059      network if there was some queuing in the system's backlog). This is usually
14060      the same date which may appear in any upstream firewall's log.
14061
14062    - "frontend_name" is the name of the frontend (or listener) which received
14063      and processed the connection.
14064
14065    - "backend_name" is the name of the backend (or listener) which was selected
14066      to manage the connection to the server. This will be the same as the
14067      frontend if no switching rule has been applied, which is common for TCP
14068      applications.
14069
14070    - "server_name" is the name of the last server to which the connection was
14071      sent, which might differ from the first one if there were connection errors
14072      and a redispatch occurred. Note that this server belongs to the backend
14073      which processed the request. If the connection was aborted before reaching
14074      a server, "<NOSRV>" is indicated instead of a server name.
14075
14076    - "Tw" is the total time in milliseconds spent waiting in the various queues.
14077      It can be "-1" if the connection was aborted before reaching the queue.
14078      See "Timers" below for more details.
14079
14080    - "Tc" is the total time in milliseconds spent waiting for the connection to
14081      establish to the final server, including retries. It can be "-1" if the
14082      connection was aborted before a connection could be established. See
14083      "Timers" below for more details.
14084
14085    - "Tt" is the total time in milliseconds elapsed between the accept and the
14086      last close. It covers all possible processing. There is one exception, if
14087      "option logasap" was specified, then the time counting stops at the moment
14088      the log is emitted. In this case, a '+' sign is prepended before the value,
14089      indicating that the final one will be larger. See "Timers" below for more
14090      details.
14091
14092    - "bytes_read" is the total number of bytes transmitted from the server to
14093      the client when the log is emitted. If "option logasap" is specified, the
14094      this value will be prefixed with a '+' sign indicating that the final one
14095      may be larger. Please note that this value is a 64-bit counter, so log
14096      analysis tools must be able to handle it without overflowing.
14097
14098    - "termination_state" is the condition the session was in when the session
14099      ended. This indicates the session state, which side caused the end of
14100      session to happen, and for what reason (timeout, error, ...). The normal
14101      flags should be "--", indicating the session was closed by either end with
14102      no data remaining in buffers. See below "Session state at disconnection"
14103      for more details.
14104
14105    - "actconn" is the total number of concurrent connections on the process when
```



```
14106      the session was logged. It is useful to detect when some per-process system
14107      limits have been reached. For instance, if actconn is close to 512 when
14108      multiple connection errors occur, chances are high that the system limits
14109      the process to use a maximum of 1024 file descriptors and that all of them
14110      are used. See section 3 "Global parameters" to find how to tune the system.
14111
14112    - "feconn" is the total number of concurrent connections on the frontend when
14113      the session was logged. It is useful to estimate the amount of resource
14114      required to sustain high loads, and to detect when the frontend's "maxconn"
14115      has been reached. Most often when this value increases by huge jumps, it is
14116      because there is congestion on the backend servers, but sometimes it can be
14117      caused by a denial of service attack.
14118
14119    - "beconn" is the total number of concurrent connections handled by the
14120      backend when the session was logged. It includes the total number of
14121      concurrent connections active on servers as well as the number of
14122      connections pending in queues. It is useful to estimate the amount of
14123      additional servers needed to support high loads for a given application.
14124      Most often when this value increases by huge jumps, it is because there is
14125      congestion on the backend servers, but sometimes it can be caused by a
14126      denial of service attack.
14127
14128    - "srv_conn" is the total number of concurrent connections still active on
14129      the server when the session was logged. It can never exceed the server's
14130      configured "maxconn" parameter. If this value is very often close or equal
14131      to the server's "maxconn", it means that traffic regulation is involved a
14132      lot, meaning that either the server's maxconn value is too low, or that
14133      there aren't enough servers to process the load with an optimal response
14134      time. When only one of the server's "srv_conn" is high, it usually means
14135      that this server has some trouble causing the connections to take longer to
14136      be processed than on other servers.
14137
14138    - "retries" is the number of connection retries experienced by this session
14139      when trying to connect to the server. It must normally be zero, unless a
14140      server is being stopped at the same moment the connection was attempted.
14141      Frequent retries generally indicate either a network problem between
14142      haproxy and the server, or a misconfigured system backlog on the server
14143      preventing new connections from being queued. This field may optionally be
14144      prefixed with a '+' sign, indicating that the session has experienced a
14145      redispatch after the maximal retry count has been reached on the initial
14146      server. In this case, the server name appearing in the log is the one the
14147      connection was redispatched to, and not the first one, though both may
14148      sometimes be the same in case of hashing for instance. So as a general rule
14149      of thumb, when a '+' is present in front of the retry count, this count
14150      should not be attributed to the logged server.
14151
14152    - "srv_queue" is the total number of requests which were processed before
14153      this one in the server queue. It is zero when the request has not gone
14154      through the server queue. It makes it possible to estimate the approximate
14155      server's response time by dividing the time spent in queue by the number of
14156      requests in the queue. It is worth noting that if a session experiences a
14157      redispatch and passes through two server queues, their positions will be
14158      cumulated. A request should not pass through both the server queue and the
14159      backend queue unless a redispatch occurs.
14160
14161    - "backend_queue" is the total number of requests which were processed before
14162      this one in the backend's global queue. It is zero when the request has not
14163      gone through the global queue. It makes it possible to estimate the average
14164      queue length, which easily translates into a number of missing servers when
14165      divided by a server's "maxconn" parameter. It is worth noting that if a
14166      session experiences a redispatch, it may pass twice in the backend's queue,
14167      and then both positions will be cumulated. A request should not pass
14168      through both the server queue and the backend queue unless a redispatch
14169      occurs.
14170
```

8.2.3. HTTP log format

The HTTP format is the most complete and the best suited for HTTP proxies. It is enabled by when "option httplog" is specified in the frontend. It provides the same level of information as the TCP format with additional features which are specific to the HTTP protocol. Just like the TCP format, the log is usually emitted at the end of the session, unless "option logasap" is specified, which generally only makes sense for download sites. A session which matches the "monitor" rules will never logged. It is also possible not to log sessions for which no data were sent by the client by specifying "option dontlognull" in the frontend. Successful connections will not be logged if "option dontlog-normal" is specified in the frontend.

Most fields are shared with the TCP log, some being different. A few fields may slightly vary depending on some configuration options. Those ones are marked with a star ('*') after the field name below.

```
  Example :
        frontend http-in
            mode http
            option httplog
            log global
            default_backend bck

        backend static
            server srv1 127.0.0.1:8000

    >>> Feb  6 12:14:14 localhost \
          haproxy[14389]: 10.0.1.2:33317 [06/Feb/2009:12:14:14.655] http-in \
          static/srv1 10/0/30/69/109 200 2750 - - ---- 1/1/1/1/0 0/0 {1wt.eu} \
          {} "GET /index.html HTTP/1.1"

  Field   Format                                Extract from the example above
      1   process_name '[' pid ']:'                            haproxy[14389]:
      2   client_ip ':' client_port                             10.0.1.2:33317
      3   '[' accept_date ']'                      [06/Feb/2009:12:14:14.655]
      4   frontend_name                                             http-in
      5   backend_name '/' server_name                           static/srv1
      6   Tq '/' Tw '/' Tc '/' Tr '/' Tt*                     10/0/30/69/109
      7   status_code                                                   200
      8   bytes_read*                                                  2750
      9   captured_request_cookie                                         -
     10   captured_response_cookie                                        -
     11   termination_state                                            ----
     12   actconn '/' feconn '/' beconn '/' srv_conn '/' retries*    1/1/1/1/0
     13   srv_queue '/' backend_queue                                   0/0
     14   '{' captured_request_headers* '}'                  {haproxy.1wt.eu}
     15   '{' captured_response_headers* '}'                             {}
     16   '"' http_request '"'                      "GET /index.html HTTP/1.1"


Detailed fields description :
  - "client_ip" is the IP address of the client which initiated the TCP
    connection to haproxy. If the connection was accepted on a UNIX socket
    instead, the IP address would be replaced with the word "unix". Note that
    when the connection is accepted on a socket configured with "accept-proxy"
    and the PROXY protocol is correctly used, then the logs will reflect the
    forwarded connection's information.

  - "client_port" is the TCP port of the client which initiated the connection.
    If the connection was accepted on a UNIX socket instead, the port would be
    replaced with the ID of the accepting socket, which is also reported in the
    stats interface.
```



```
  - "accept_date" is the exact date when the TCP connection was received by
    haproxy (which might be very slightly different from the date observed on
    the network if there was some queuing in the system's backlog). This is
    usually the same date which may appear in any upstream firewall's log. This
    does not depend on the fact that the client has sent the request or not.

  - "frontend_name" is the name of the frontend (or listener) which received
    and processed the connection.

  - "backend_name" is the name of the backend (or listener) which was selected
    to manage the connection to the server. This will be the same as the
    frontend if no switching rule has been applied.

  - "server_name" is the name of the last server to which the connection was
    sent, which might differ from the first one if there were connection errors
    and a redispatch occurred. Note that this server belongs to the backend
    which processed the request. If the request was aborted before reaching a
    server, "<NOSRV>" is indicated instead of a server name. If the request was
    intercepted by the stats subsystem, "<STATS>" is indicated instead.

  - "Tq" is the total time in milliseconds spent waiting for the client to send
    a full HTTP request, not counting data. It can be "-1" if the connection
    was aborted before a complete request could be received. It should always
    be very small because a request generally fits in one single packet. Large
    times here generally indicate network trouble between the client and
    haproxy. See "Timers" below for more details.

  - "Tw" is the total time in milliseconds spent waiting in the various queues.
    It can be "-1" if the connection was aborted before reaching the queue.
    See "Timers" below for more details.

  - "Tc" is the total time in milliseconds spent waiting for the connection to
    establish to the final server, including retries. It can be "-1" if the
    request was aborted before a connection could be established. See "Timers"
    below for more details.

  - "Tr" is the total time in milliseconds spent waiting for the server to send
    a full HTTP response, not counting data. It can be "-1" if the request was
    aborted before a complete response could be received. It generally matches
    the server's processing time for the request, though it may be altered by
    the amount of data sent by the client to the server. Large times here on
    "GET" requests generally indicate an overloaded server. See "Timers" below
    for more details.

  - "Tt" is the total time in milliseconds elapsed between the accept and the
    last close. It covers all possible processing. There is one exception, if
    "option logasap" was specified, then the time counting stops at the moment
    the log is emitted. In this case, a '+' sign is prepended before the value,
    indicating that the final one will be larger. See "Timers" below for more
    details.

  - "status_code" is the HTTP status code returned to the client. This status
    is generally set by the server, but it might also be set by haproxy when
    the server cannot be reached or when its response is blocked by haproxy.

  - "bytes_read" is the total number of bytes transmitted to the client when
    the log is emitted. This does include HTTP headers. If "option logasap" is
    specified, the this value will be prefixed with a '+' sign indicating that
    the final one may be larger. Please note that this value is a 64-bit
    counter, so log analysis tools must be able to handle it without
    overflowing.

  - "captured_request_cookie" is an optional "name=value" entry indicating that
    the client had this cookie in the request. The cookie name and its maximum
```

length are defined by the "capture cookie" statement in the frontend
configuration. The field is a single dash ('-') when the option is not
set. Only one cookie may be captured, it is generally used to track session
ID exchanges between a client and a server to detect session crossing
between clients due to application bugs. For more details, please consult
the section "Capturing HTTP headers and cookies" below.

- "captured_response_cookie" is an optional "name=value" entry indicating
that the server has returned a cookie with its response. The cookie name
and its maximum length are defined by the "capture cookie" statement in the
frontend configuration. The field is a single dash ('-') when the option is
not set. Only one cookie may be captured, it is generally used to track
session ID exchanges between a client and a server to detect session
crossing between clients due to application bugs. For more details, please
consult the section "Capturing HTTP headers and cookies" below.

- "termination_state" is the condition the session was in when the session
ended. This indicates the session state, which side caused the end of
session to happen, for what reason (timeout, error, ...), just like in TCP
logs, and information about persistence operations on cookies in the last
two characters. The normal flags should begin with "--", indicating the
session was closed by either end with no data remaining in buffers. See
below "Session state at disconnection" for more details.

- "actconn" is the total number of concurrent connections on the process when
the session was logged. It is useful to detect when some per-process system
limits have been reached. For instance, if actconn is close to 512 or 1024
when multiple connection errors occur, chances are high that the system
limits the process to use a maximum of 1024 file descriptors and that all
of them are used. See section 3 "Global parameters" to find how to tune the
system.

- "feconn" is the total number of concurrent connections on the frontend when
the session was logged. It is useful to estimate the amount of resource
required to sustain high loads, and to detect when the frontend's "maxconn"
has been reached. Most often when this value increases by huge jumps, it is
because there is congestion on the backend servers, but sometimes it can be
caused by a denial of service attack.

- "beconn" is the total number of concurrent connections handled by the
backend when the session was logged. It includes the total number of
concurrent connections active on servers as well as the number of
connections pending in queues. It is useful to estimate the amount of
additional servers needed to support high loads for a given application.
Most often when this value increases by huge jumps, it is because there is
congestion on the backend servers, but sometimes it can be caused by a
denial of service attack.

- "srv_conn" is the total number of concurrent connections still active on
the server when the session was logged. It can never exceed the server's
configured "maxconn" parameter. If this value is very often close or equal
to the server's "maxconn", it means that traffic regulation is involved a
lot, meaning that either the server's maxconn value is too low, or that
there aren't enough servers to process the load with an optimal response
time. When only one of the server's "srv_conn" is high, it usually means
that this server has some trouble causing the requests to take longer to be
processed than on other servers.

- "retries" is the number of connection retries experienced by this session
when trying to connect to the server. It must normally be zero, unless a
server is being stopped at the same moment the connection was attempted.
Frequent retries generally indicate either a network problem between
haproxy and the server, or a misconfigured system backlog on the server
preventing new connections from being queued. This field may optionally be
prefixed with a '+' sign, indicating that the session has experienced a

redispatch after the maximal retry count has been reached on the initial
server. In this case, the server name appearing in the log is the one the
connection was redispatched to, and not the first one, though both may
sometimes be the same in case of hashing for instance. So as a general rule
of thumb, when a '+' is present in front of the retry count, this count
should not be attributed to the logged server.

- "srv_queue" is the total number of requests which were processed before
this one in the server queue. It is zero when the request has not gone
through the server queue. It makes it possible to estimate the approximate
server's response time by dividing the time spent in queue by the number of
requests in the queue. It is worth noting that if a session experiences a
redispatch and passes through two server queues, their positions will be
cumulated. A request should not pass through both the server queue and the
backend queue unless a redispatch occurs.

- "backend_queue" is the total number of requests which were processed before
this one in the backend's global queue. It is zero when the request has not
gone through the global queue. It makes it possible to estimate the average
queue length, which easily translates into a number of missing servers when
divided by a server's "maxconn" parameter. It is worth noting that if a
session experiences a redispatch, it may pass twice in the backend's queue,
and then both positions will be cumulated. A request should not pass
through both the server queue and the backend queue unless a redispatch
occurs.

- "captured_request_headers" is a list of headers captured in the request due
to the presence of the "capture request header" statement in the frontend.
Multiple headers can be captured, they will be delimited by a vertical bar
('|'). When no capture is enabled, the braces do not appear, causing a
shift of remaining fields. It is important to note that this field may
contain spaces, and that using it requires a smarter log parser than when
it's not used. Please consult the section "Capturing HTTP headers and
cookies" below for more details.

- "captured_response_headers" is a list of headers captured in the response
due to the presence of the "capture response header" statement in the
frontend. Multiple headers can be captured, they will be delimited by a
vertical bar ('|'). When no capture is enabled, the braces do not appear,
causing a shift of remaining fields. It is important to note that this
field may contain spaces, and that using it requires a smarter log parser
than when it's not used. Please consult the section "Capturing HTTP headers
and cookies" below for more details.

- "http_request" is the complete HTTP request line, including the method,
request and HTTP version string. Non-printable characters are encoded (see
below the section "Non-printable characters"). This is always the last
field, and it is always delimited by quotes and is the only one which can
contain quotes. If new fields are added to the log format, they will be
added before this field. This field might be truncated if the request is
huge and does not fit in the standard syslog buffer (1024 characters). This
is the reason why this field must always remain the last one.

8.2.4. Custom log format

The directive log-format allows you to customize the logs in http mode and tcp
mode. It takes a string as argument.

HAproxy understands some log format variables. % precedes log format variables.
Variables can take arguments using braces ('{}'), and multiple arguments are
separated by commas within the braces. Flags may be added or removed by
prefixing them with a '+' or '-' sign.

Special variable "%o" may be used to propagate its flags to all other
variables on the same format string. This is particularly handy with quoted
string formats ("Q").

If a variable is named between square brackets ('[' .. ']') then it is used
as a sample expression rule (see section 7.3). This it useful to add some
less common information such as the client's SSL certificate's DN, or to log
the key that would be used to store an entry into a stick table.

Note: spaces must be escaped. A space character is considered as a separator.
In order to emit a verbatim '%', it must be preceded by another '%' resulting
in '%%'. HAProxy will automatically merge consecutive separators.

Flags are :
  * Q: quote a string
  * X: hexadecimal representation (IPs, Ports, %Ts, %rt, %pid)

  Example:

    log-format %T\ %t\ Some\ Text
    log-format %{+Q}o\ %t\ %s\ %{-Q}r

At the moment, the default HTTP format is defined this way :

    log-format %ci:%cp\ [%t]\ %ft\ %b/%s\ %Tq/%Tw/%Tc/%Tr/%Tt\ %ST\ %B\ %CC\ \
                %CS\ %tsc\ %ac/%fc/%bc/%sc/%rc\ %sq/%bq\ %hr\ %hs\ %{+Q}r

the default CLF format is defined this way :

    log-format %{+Q}o\ %{-Q}ci\ -\ -\ [%T]\ %r\ %ST\ %B\ \"\"\ \"\"\ %cp\ \
                %ms\ %ft\ %b\ %s\ \%Tq\ %Tw\ %Tc\ %Tr\ %Tt\ %tsc\ %ac\ %fc\ \
                %bc\ %sc\ %rc\ %sq\ %bq\ %CC\ %CS\ \%hrl\ %hsl

and the default TCP format is defined this way :

    log-format %ci:%cp\ [%t]\ %ft\ %b/%s\ %Tw/%Tc/%Tt\ %B\ %ts\ \
                %ac/%fc/%bc/%sc/%rc\ %sq/%bq

Please refer to the table below for currently defined variables :

R	var	field name (8.2.2 and 8.2.3 for description)	type
	%o	special variable, apply flags on all next var	
	%B	bytes_read (from server to client)	numeric
H	%CC	captured_request_cookie	string
H	%CS	captured_response_cookie	string
	%H	hostname	string
H	%HM	HTTP method (ex: POST)	string
H	%HP	HTTP request URI without query string (path)	string
H	%HQ	HTTP request URI query string (ex: ?bar=baz)	string
H	%HU	HTTP request URI (ex: /foo?bar=baz)	string
H	%HV	HTTP version (ex: HTTP/1.0)	string
	%ID	unique-id	string
	%ST	status_code	numeric
	%T	gmt_date_time	date
	%Tc	Tc	numeric
	%Tl	local_date_time	date
H	%Tq	Tq	numeric
H	%Tr	Tr	numeric
	%Ts	timestamp	numeric
	%Tt	Tt	numeric
	%Tw	Tw	numeric
	%U	bytes_uploaded (from client to server)	numeric

		%ac	actconn	numeric
		%b	backend_name	string
		%bc	beconn (backend concurrent connections)	numeric
		%bi	backend_source_ip (connecting address)	IP
		%bp	backend_source_port (connecting address)	numeric
		%bq	backend_queue	numeric
		%ci	client_ip (accepted address)	IP
		%cp	client_port (accepted address)	numeric
		%f	frontend_name	string
		%fc	feconn (frontend concurrent connections)	numeric
		%fi	frontend_ip (accepting address)	IP
		%fp	frontend_port (accepting address)	numeric
		%ft	frontend_name_transport ('~' suffix for SSL)	string
		%lc	frontend_log_counter	numeric
		%hr	captured_request_headers default style	string
		%hrl	captured_request_headers CLF style	string list
		%hs	captured_response_headers default style	string
		%hsl	captured_response_headers CLF style	string list
		%ms	accept date milliseconds (left-padded with 0)	numeric
		%pid	PID	numeric
H		%r	http_request	string
		%rc	retries	numeric
		%rt	request_counter (HTTP req or TCP session)	numeric
		%s	server_name	string
		%sc	srv_conn (server concurrent connections)	numeric
		%si	server_IP (target address)	IP
		%sp	server_port (target address)	numeric
		%sq	srv_queue	numeric
S		%sslc	ssl_ciphers (ex: AES-SHA)	string
S		%sslv	ssl_version (ex: TLSv1)	string
		%t	date_time (with millisecond resolution)	date
		%ts	termination_state	string
H		%tsc	termination_state with cookie status	string

  R = Restrictions : H = mode http only ; S = SSL only

8.2.5. Error log format

When an incoming connection fails due to an SSL handshake or an invalid PROXY
protocol header, haproxy will log the event using a shorter, fixed line format.
By default, logs are emitted at the LOG_INFO level, unless the option
"log-separate-errors" is set in the backend, in which case the LOG_ERR level
will be used. Connections on which no data are exchanged (eg: probes) are not
logged if the "dontlognull" option is set.

The format looks like this :

    >>> Dec 3 18:27:14 localhost \
          haproxy[6103]: 127.0.0.1:56059 [03/Dec/2012:17:35:10.380] frt/f1: \
          Connection error during SSL handshake

  Field   Format   Extract from the example above
      1   process_name '[' pid ']:'   haproxy[6103]:
      2   client_ip ':' client_port   127.0.0.1:56059
      3   '[' accept_date ']'   [03/Dec/2012:17:35:10.380]
      4   frontend_name "/" bind_name ":"   frt/f1:
      5   message   Connection error during SSL handshake

These fields just provide minimal information to help debugging connection
failures.

8.3. Advanced logging options

Some advanced logging options are often looked for but are not easy to find out just by looking at the various options. Here is an entry point for the few options which can enable better logging. Please refer to the keywords reference for more information about their usage.

8.3.1. Disabling logging of external tests

It is quite common to have some monitoring tools perform health checks on haproxy. Sometimes it will be a layer 3 load-balancer such as LVS or any commercial load-balancer, and sometimes it will simply be a more complete monitoring system such as Nagios. When the tests are very frequent, users often ask how to disable logging for those checks. There are three possibilities :

  - if connections come from everywhere and are just TCP probes, it is often
    desired to simply disable logging of connections without data exchange, by
    setting "option dontlognull" in the frontend. It also disables logging of
    port scans, which may or may not be desired.

  - if the connection come from a known source network, use "monitor-net" to
    declare this network as monitoring only. Any host in this network will then
    only be able to perform health checks, and their requests will not be
    logged. This is generally appropriate to designate a list of equipment
    such as other load-balancers.

  - if the tests are performed on a known URI, use "monitor-uri" to declare
    this URI as dedicated to monitoring. Any host sending this request will
    only get the result of a health-check, and the request will not be logged.

8.3.2. Logging before waiting for the session to terminate

The problem with logging at end of connection is that you have no clue about what is happening during very long sessions, such as remote terminal sessions or large file downloads. This problem can be worked around by specifying "option logasap" in the frontend. Haproxy will then log as soon as possible, just before data transfer begins. This means that in case of TCP, it will still log the connection status to the server, and in case of HTTP, it will log just after processing the server headers. In this case, the number of bytes reported is the number of header bytes sent to the client. In order to avoid confusion with normal logs, the total time field and the number of bytes are prefixed with a '+' sign which means that real numbers are certainly larger.

8.3.3. Raising log level upon errors

Sometimes it is more convenient to separate normal traffic from errors logs, for instance in order to ease error monitoring from log files. When the option "log-separate-errors" is used, connections which experience errors, timeouts, retries, redispatches or HTTP status codes 5xx will see their syslog level raised from "info" to "err". This will help a syslog daemon store the log in a separate file. It is very important to keep the errors in the normal traffic file too, so that log ordering is not altered. You should also be careful if you already have configured your syslog daemon to store all logs higher than "notice" in an "admin" file, because the "err" level is higher than "notice".

8.3.4. Disabling logging of successful connections

Although this may sound strange at first, some large sites have to deal with multiple thousands of logs per second and are experiencing difficulties keeping them intact for a long time or detecting errors within them. If the option "dontlog-normal" is set on the frontend, all normal connections will not be logged. In this regard, a normal connection is defined as one without any error, timeout, retry nor redispatch. In HTTP, the status code is checked too, and a response with a status 5xx is not considered normal and will be logged too. Of course, doing is is really discouraged as it will remove most of the useful information from the logs. Do this only if you have no other alternative.

8.4. Timing events

Timers provide a great help in troubleshooting network problems. All values are reported in milliseconds (ms). These timers should be used in conjunction with the session termination flags. In TCP mode with "option tcplog" set on the frontend, 3 control points are reported under the form "Tw/Tc/Tt", and in HTTP mode, 5 control points are reported under the form "Tq/Tw/Tc/Tr/Tt" :

  - Tq: total time to get the client request (HTTP mode only). It's the time
    elapsed between the moment the client connection was accepted and the
    moment the proxy received the last HTTP header. The value "-1" indicates
    that the end of headers (empty line) has never been seen. This happens when
    the client closes prematurely or times out.

  - Tw: total time spent in the queues waiting for a connection slot. It
    accounts for backend queue as well as the server queues, and depends on the
    queue size, and the time needed for the server to complete previous
    requests. The value "-1" means that the request was killed before reaching
    the queue, which is generally what happens with invalid or denied requests.

  - Tc: total time to establish the TCP connection to the server. It's the time
    elapsed between the moment the proxy sent the connection request, and the
    moment it was acknowledged by the server, or between the TCP SYN packet and
    the matching SYN/ACK packet in return. The value "-1" means that the
    connection never established.

  - Tr: server response time (HTTP mode only). It's the time elapsed between
    the moment the TCP connection was established to the server and the moment
    the server sent its complete response headers. It purely shows its request
    processing time, without the network overhead due to the data transmission.
    It is worth noting that when the client has data to send to the server, for
    instance during a POST request, the time already runs, and this can distort
    apparent response time. For this reason, it's generally wise not to trust
    too much this field for POST requests initiated from clients behind an
    untrusted network. A value of "-1" here means that the last the response
    header (empty line) was never seen, most likely because the server timeout
    stroke before the server managed to process the request.

  - Tt: total session duration time, between the moment the proxy accepted it
    and the moment both ends were closed. The exception is when the "logasap"
    option is specified. In this case, it only equals (Tq+Tw+Tc+Tr), and is
    prefixed with a '+' sign. From this field, we can deduce "Td", the data
    transmission time, by subtracting other timers when valid :

        Td = Tt - (Tq + Tw + Tc + Tr)

    Timers with "-1" values have to be excluded from this equation. In TCP
    mode, "Tq" and "Tr" have to be excluded too. Note that "Tt" can never be
    negative.

These timers provide precious indications on trouble causes. Since the TCP

protocol defines retransmit delays of 3, 6, 12... seconds, we know for sure
that timers close to multiples of 3s are nearly always related to lost packets
due to network problems (wires, negotiation, congestion). Moreover, if "Tt" is
close to a timeout value specified in the configuration, it often means that a
session has been aborted on timeout.

Most common cases :

  - If "Tq" is close to 3000, a packet has probably been lost between the
    client and the proxy. This is very rare on local networks but might happen
    when clients are on far remote networks and send large requests. It may
    happen that values larger than usual appear here without any network cause.
    Sometimes, during an attack or just after a resource starvation has ended,
    haproxy may accept thousands of connections in a few milliseconds. The time
    spent accepting these connections will inevitably slightly delay processing
    of other connections, and it can happen that request times in the order of
    a few tens of milliseconds are measured after a few thousands of new
    connections have been accepted at once. Setting "option http-server-close"
    may display larger request times since "Tq" also measures the time spent
    waiting for additional requests.

  - If "Tc" is close to 3000, a packet has probably been lost between the
    server and the proxy during the server connection phase. This value should
    always be very low, such as 1 ms on local networks and less than a few tens
    of ms on remote networks.

  - If "Tr" is nearly always lower than 3000 except some rare values which seem
    to be the average majored by 3000, there are probably some packets lost
    between the proxy and the server.

  - If "Tt" is large even for small byte counts, it generally is because
    neither the client nor the server decides to close the connection, for
    instance because both have agreed on a keep-alive connection mode. In order
    to solve this issue, it will be needed to specify "option httpclose" on
    either the frontend or the backend. If the problem persists, it means that
    the server ignores the "close" connection mode and expects the client to
    close. Then it will be required to use "option forceclose". Having the
    smallest possible 'Tt' is important when connection regulation is used with
    the "maxconn" option on the servers, since no new connection will be sent
    to the server until another one is released.

Other noticeable HTTP log cases ('xx' means any value to be ignored) :

  Tq/Tw/Tc/Tr/+Tt The "option logasap" is present on the frontend and the log
                   was emitted before the data phase. All the timers are valid
                   except "Tt" which is shorter than reality.

  -1/xx/xx/xx/Tt   The client was not able to send a complete request in time
                   or it aborted too early. Check the session termination flags
                   then "timeout http-request" and "timeout client" settings.

  Tq/-1/xx/xx/Tt   It was not possible to process the request, maybe because
                   servers were out of order, because the request was invalid
                   or forbidden by ACL rules. Check the session termination
                   flags.

  Tq/Tw/-1/xx/Tt   The connection could not establish on the server. Either it
                   actively refused it or it timed out after Tt-(Tq+Tw) ms.
                   Check the session termination flags, then check the
                   "timeout connect" setting. Note that the tarpit action might
                   return similar-looking patterns, with "Tw" equal to the time
                   the client connection was maintained open.

  Tq/Tw/Tc/-1/Tt   The server has accepted the connection but did not return
                   a complete response in time, or it closed its connection

                   unexpectedly after Tt-(Tq+Tw+Tc) ms. Check the session
                   termination flags, then check the "timeout server" setting.

8.5. Session state at disconnection

TCP and HTTP logs provide a session termination indicator in the
"termination_state" field, just before the number of active connections. It is
2-characters long in TCP mode, and is extended to 4 characters in HTTP mode,
each of which has a special meaning :

  - On the first character, a code reporting the first event which caused the
    session to terminate :

        C : the TCP session was unexpectedly aborted by the client.

        S : the TCP session was unexpectedly aborted by the server, or the
            server explicitly refused it.

        P : the session was prematurely aborted by the proxy, because of a
            connection limit enforcement, because a DENY filter was matched,
            because of a security check which detected and blocked a dangerous
            error in server response which might have caused information leak
            (eg: cacheable cookie).

        L : the session was locally processed by haproxy and was not passed to
            a server. This is what happens for stats and redirects.

        R : a resource on the proxy has been exhausted (memory, sockets, source
            ports, ...). Usually, this appears during the connection phase, and
            system logs should contain a copy of the precise error. If this
            happens, it must be considered as a very serious anomaly which
            should be fixed as soon as possible by any means.

        I : an internal error was identified by the proxy during a self-check.
            This should NEVER happen, and you are encouraged to report any log
            containing this, because this would almost certainly be a bug. It
            would be wise to preventively restart the process after such an
            event too, in case it would be caused by memory corruption.

        D : the session was killed by haproxy because the server was detected
            as down and was configured to kill all connections when going down.

        U : the session was killed by haproxy on this backup server because an
            active server was detected as up and was configured to kill all
            backup connections when going up.

        K : the session was actively killed by an admin operating on haproxy.

        c : the client-side timeout expired while waiting for the client to
            send or receive data.

        s : the server-side timeout expired while waiting for the server to
            send or receive data.

        - : normal session completion, both the client and the server closed
            with nothing left in the buffers.

  - on the second character, the TCP or HTTP session state when it was closed :

        R : the proxy was waiting for a complete, valid REQUEST from the client
            (HTTP mode only). Nothing was sent to any server.

        Q : the proxy was waiting in the QUEUE for a connection slot. This can

```
14821              only happen when servers have a 'maxconn' parameter set. It can
14822              also happen in the global queue after a redispatch consecutive to
14823              a failed attempt to connect to a dying server. If no redispatch is
14824              reported, then no connection attempt was made to any server.
14825
14826          C : the proxy was waiting for the CONNECTION to establish on the
14827              server. The server might at most have noticed a connection attempt.
14828
14829          H : the proxy was waiting for complete, valid response HEADERS from the
14830              server (HTTP only).
14831
14832          D : the session was in the DATA phase.
14833
14834          L : the proxy was still transmitting LAST data to the client while the
14835              server had already finished. This one is very rare as it can only
14836              happen when the client dies while receiving the last packets.
14837
14838          T : the request was tarpitted. It has been held open with the client
14839              during the whole "timeout tarpit" duration or until the client
14840              closed, both of which will be reported in the "Tw" timer.
14841
14842          - : normal session completion after end of data transfer.
14843
14844      - the third character tells whether the persistence cookie was provided by
14845        the client (only in HTTP mode) :
14846
14847          N : the client provided NO cookie. This is usually the case for new
14848              visitors, so counting the number of occurrences of this flag in the
14849              logs generally indicate a valid trend for the site frequentation.
14850
14851          I : the client provided an INVALID cookie matching no known server.
14852              This might be caused by a recent configuration change, mixed
14853              cookies between HTTP/HTTPS sites, persistence conditionally
14854              ignored, or an attack.
14855
14856          D : the client provided a cookie designating a server which was DOWN,
14857              so either "option persist" was used and the client was sent to
14858              this server, or it was not set and the client was redispatched to
14859              another server.
14860
14861          V : the client provided a VALID cookie, and was sent to the associated
14862              server.
14863
14864          E : the client provided a valid cookie, but with a last date which was
14865              older than what is allowed by the "maxidle" cookie parameter, so
14866              the cookie is consider EXPIRED and is ignored. The request will be
14867              redispatched just as if there was no cookie.
14868
14869          O : the client provided a valid cookie, but with a first date which was
14870              older than what is allowed by the "maxlife" cookie parameter, so
14871              the cookie is consider too OLD and is ignored. The request will be
14872              redispatched just as if there was no cookie.
14873
14874          U : a cookie was present but was not used to select the server because
14875              some other server selection mechanism was used instead (typically a
14876              "use-server" rule).
14877
14878          - : does not apply (no cookie set in configuration).
14879
14880      - the last character reports what operations were performed on the persistence
14881        cookie returned by the server (only in HTTP mode) :
14882
14883          N : NO cookie was provided by the server, and none was inserted either.
14884
14885          I : no cookie was provided by the server, and the proxy INSERTED one.
```



```
14886              Note that in "cookie insert" mode, if the server provides a cookie,
14887              it will still be overwritten and reported as "I" here.
14888
14889          U : the proxy UPDATED the last date in the cookie that was presented by
14890              the client. This can only happen in insert mode with "maxidle". It
14891              happens every time there is activity at a different date than the
14892              date indicated in the cookie. If any other change happens, such as
14893              a redispatch, then the cookie will be marked as inserted instead.
14894
14895          P : a cookie was PROVIDED by the server and transmitted as-is.
14896
14897          R : the cookie provided by the server was REWRITTEN by the proxy, which
14898              happens in "cookie rewrite" or "cookie prefix" modes.
14899
14900          D : the cookie provided by the server was DELETED by the proxy.
14901
14902          - : does not apply (no cookie set in configuration).
14903
14904  The combination of the two first flags gives a lot of information about what
14905  was happening when the session terminated, and why it did terminate. It can be
14906  helpful to detect server saturation, network troubles, local system resource
14907  starvation, attacks, etc...
14908
14909  The most common termination flags combinations are indicated below. They are
14910  alphabetically sorted, with the lowercase set just after the upper case for
14911  easier finding and understanding.
14912
14913    Flags   Reason
14914
14915       --   Normal termination.
14916
14917       CC   The client aborted before the connection could be established to the
14918            server. This can happen when haproxy tries to connect to a recently
14919            dead (or unchecked) server, and the client aborts while haproxy is
14920            waiting for the server to respond or for "timeout connect" to expire.
14921
14922       CD   The client unexpectedly aborted during data transfer. This can be
14923            caused by a browser crash, by an intermediate equipment between the
14924            client and haproxy which decided to actively break the connection,
14925            by network routing issues between the client and haproxy, or by a
14926            keep-alive session between the server and the client terminated first
14927            by the client.
14928
14929       cD   The client did not send nor acknowledge any data for as long as the
14930            "timeout client" delay. This is often caused by network failures on
14931            the client side, or the client simply leaving the net uncleanly.
14932
14933       CH   The client aborted while waiting for the server to start responding.
14934            It might be the server taking too long to respond or the client
14935            clicking the 'Stop' button too fast.
14936
14937       cH   The "timeout client" stroke while waiting for client data during a
14938            POST request. This is sometimes caused by too large TCP MSS values
14939            for PPPoE networks which cannot transport full-sized packets. It can
14940            also happen when client timeout is smaller than server timeout and
14941            the server takes too long to respond.
14942
14943       CQ   The client aborted while its session was queued, waiting for a server
14944            with enough empty slots to accept it. It might be that either all the
14945            servers were saturated or that the assigned server was taking too
14946            long a time to respond.
14947
14948       CR   The client aborted before sending a full HTTP request. Most likely
14949            the request was typed by hand using a telnet client, and aborted
14950            too early. The HTTP status code is likely a 400 here. Sometimes this
```

might also be caused by an IDS killing the connection between haproxy and the client. "option http-ignore-probes" can be used to ignore connections without any data transfer.

cR The "timeout http-request" stroke before the client sent a full HTTP request. This is sometimes caused by too large TCP MSS values on the client side for PPPoE networks which cannot transport full-sized packets, or by clients sending requests by hand and not typing fast enough, or forgetting to enter the empty line at the end of the request. The HTTP status code is likely a 408 here. Note: recently, some browsers started to implement a "pre-connect" feature consisting in speculatively connecting to some recently visited web sites just in case the user would like to visit them. This results in many connections being established to web sites, which end up in 408 Request Timeout if the timeout strikes first, or 400 Bad Request when the browser decides to close them first. These ones pollute the log and feed the error counters. Some versions of some browsers have even been reported to display the error code. It is possible to work around the undesirable effects of this behaviour by adding "option http-ignore-probes" in the frontend, resulting in connections with zero data transfer to be totally ignored. This will definitely hide the errors of people experiencing connectivity issues though.

CT The client aborted while its session was tarpitted. It is important to check if this happens on valid requests, in order to be sure that no wrong tarpit rules have been written. If a lot of them happen, it might make sense to lower the "timeout tarpit" value to something closer to the average reported "Tw" timer, in order not to consume resources for just a few attackers.

LR The request was intercepted and locally handled by haproxy. Generally it means that this was a redirect or a stats request.

SC The server or an equipment between it and haproxy explicitly refused the TCP connection (the proxy received a TCP RST or an ICMP message in return). Under some circumstances, it can also be the network stack telling the proxy that the server is unreachable (eg: no route, or no ARP response on local network). When this happens in HTTP mode, the status code is likely a 502 or 503 here.

sC The "timeout connect" stroke before a connection to the server could complete. When this happens in HTTP mode, the status code is likely a 503 or 504 here.

SD The connection to the server died with an error during the data transfer. This usually means that haproxy has received an RST from the server or an ICMP message from an intermediate equipment while exchanging data with the server. This can be caused by a server crash or by a network issue on an intermediate equipment.

sD The server did not send nor acknowledge any data for as long as the "timeout server" setting during the data phase. This is often caused by too short timeouts on L4 equipments before the server (firewalls, load-balancers, ...), as well as keep-alive sessions maintained between the client and the server expiring first on haproxy.

SH The server aborted before sending its full HTTP response headers, or it crashed while processing the request. Since a server aborting at this moment is very rare, it would be wise to inspect its logs to control whether it crashed and why. The logged request may indicate a small set of faulty requests, demonstrating bugs in the application. Sometimes this might also be caused by an IDS killing the connection between haproxy and the server.

sH The "timeout server" stroke before the server could return its

response headers. This is the most common anomaly, indicating too long transactions, probably caused by server or database saturation. The immediate workaround consists in increasing the "timeout server" setting, but it is important to keep in mind that the user experience will suffer from these long response times. The only long term solution is to fix the application.

sQ The session spent too much time in queue and has been expired. See the "timeout queue" and "timeout connect" settings to find out how to fix this if it happens too often. If it often happens massively in short periods, it may indicate general problems on the affected servers due to I/O or database congestion, or saturation caused by external attacks.

PC The proxy refused to establish a connection to the server because the process' socket limit has been reached while attempting to connect. The global "maxconn" parameter may be increased in the configuration so that it does not happen anymore. This status is very rare and might happen when the global "ulimit-n" parameter is forced by hand.

PD The proxy blocked an incorrectly formatted chunked encoded message in a request or a response, after the server has emitted its headers. In most cases, this will indicate an invalid message from the server to the client. Haproxy supports chunk sizes of up to 2GB - 1 (2147483647 bytes). Any larger size will be considered as an error.

PH The proxy blocked the server's response, because it was invalid, incomplete, dangerous (cache control), or matched a security filter. In any case, an HTTP 502 error is sent to the client. One possible cause for this error is an invalid syntax in an HTTP header name containing unauthorized characters. It is also possible but quite rare, that the proxy blocked a chunked-encoding request from the client due to an invalid syntax, before the server responded. In this case, an HTTP 400 error is sent to the client and reported in the logs.

PR The proxy blocked the client's HTTP request, either because of an invalid HTTP syntax, in which case it returned an HTTP 400 error to the client, or because a deny filter matched, in which case it returned an HTTP 403 error.

PT The proxy blocked the client's request and has tarpitted its connection before returning it a 500 server error. Nothing was sent to the server. The connection was maintained open for as long as reported by the "Tw" timer field.

RC A local resource has been exhausted (memory, sockets, source ports) preventing the connection to the server from establishing. The error logs will tell precisely what was missing. This is very rare and can only be solved by proper system tuning.

The combination of the two last flags gives a lot of information about how persistence was handled by the client, the server and by haproxy. This is very important to troubleshoot disconnections, when users complain they have to re-authenticate. The commonly encountered flags are :

-- Persistence cookie is not enabled.

NN No cookie was provided by the client, none was inserted in the response. For instance, this can be in insert mode with "postonly" set on a GET request.

II A cookie designating an invalid server was provided by the client, a valid one was inserted in the response. This typically happens when a "server" entry is removed from the configuration, since its cookie

value can be presented by a client when no other server knows it.

NI No cookie was provided by the client, one was inserted in the response. This typically happens for first requests from every user in "insert" mode, which makes it an easy way to count real users.

VN A cookie was provided by the client, none was inserted in the response. This happens for most responses for which the client has already got a cookie.

VU A cookie was provided by the client, with a last visit date which is not completely up-to-date, so an updated cookie was provided in response. This can also happen if there was no date at all, or if there was a date but the "maxidle" parameter was not set, so that the cookie can be switched to unlimited time.

EI A cookie was provided by the client, with a last visit date which is too old for the "maxidle" parameter, so the cookie was ignored and a new cookie was inserted in the response.

OI A cookie was provided by the client, with a first visit date which is too old for the "maxlife" parameter, so the cookie was ignored and a new cookie was inserted in the response.

DI The server designated by the cookie was down, a new server was selected and a new cookie was emitted in the response.

VI The server designated by the cookie was not marked dead but could not be reached. A redispatch happened and selected another one, which was then advertised in the response.

8.6. Non-printable characters

In order not to cause trouble to log analysis tools or terminals during log consulting, non-printable characters are not sent as-is into log files, but are converted to the two-digits hexadecimal representation of their ASCII code, prefixed by the character '#'. The only characters that can be logged without being escaped are comprised between 32 and 126 (inclusive). Obviously, the escape character '#' itself is also encoded to avoid any ambiguity ("#23"). It is the same for the character '"' which becomes "#22", as well as '{', '|' and '}' when logging headers.

Note that the space character (' ') is not encoded in headers, which can cause issues for tools relying on space count to locate fields. A typical header containing spaces is "User-Agent".

Last, it has been observed that some syslog daemons such as syslog-ng escape the quote ('"') with a backslash ('\'). The reverse operation can safely be performed since no quote may appear anywhere else in the logs.

8.7. Capturing HTTP cookies

Cookie capture simplifies the tracking a complete user session. This can be achieved using the "capture cookie" statement in the frontend. Please refer to section 4.2 for more details. Only one cookie can be captured, and the same cookie will simultaneously be checked in the request ("Cookie:" header) and in the response ("Set-Cookie:" header). The respective values will be reported in the HTTP logs at the "captured_request_cookie" and "captured_response_cookie" locations (see section 8.2.3 about HTTP log format). When either cookie is not seen, a dash ('-') replaces the value. This way, it's easy to detect when a user switches to a new session for example, because the server will reassign it

a new cookie. It is also possible to detect if a server unexpectedly sets a wrong cookie to a client, leading to session crossing.

Examples :

```
# capture the first cookie whose name starts with "ASPSESSION"
capture cookie ASPSESSION len 32

# capture the first cookie whose name is exactly "vgnvisitor"
capture cookie vgnvisitor= len 32
```

8.8. Capturing HTTP headers

Header captures are useful to track unique request identifiers set by an upper proxy, virtual host names, user-agents, POST content-length, referrers, etc. In the response, one can search for information about the response length, how the server asked the cache to behave, or an object location during a redirection.

Header captures are performed using the "capture request header" and "capture response header" statements in the frontend. Please consult their definition in section 4.2 for more details.

It is possible to include both request headers and response headers at the same time. Non-existent headers are logged as empty strings, and if one header appears more than once, only its last occurrence will be logged. Request headers are grouped within braces '{' and '}' in the same order as they were declared, and delimited with a vertical bar '|' without any space. Response headers follow the same representation, but are displayed after a space following the request headers block. These blocks are displayed just before the HTTP request in the logs.

As a special case, it is possible to specify an HTTP header capture in a TCP frontend. The purpose is to enable logging of headers which will be parsed in an HTTP backend if the request is then switched to this HTTP backend.

Example :

```
# This instance chains to the outgoing proxy
listen proxy-out
    mode http
    option httplog
    option logasap
    log global
    server cache1 192.168.1.1:3128

    # log the name of the virtual server
    capture request  header Host len 20

    # log the amount of data uploaded during a POST
    capture request  header Content-Length len 10

    # log the beginning of the referrer
    capture request  header Referer len 20

    # server name (useful for outgoing proxies only)
    capture response header Server len 20

    # logging the content-length is useful with "option logasap"
    capture response header Content-Length len 10

    # log the expected cache behaviour on the response
    capture response header Cache-Control len 8

    # the Via header will report the next proxy's name
    capture response header Via len 20
```

```
15211
15212              # log the URL location during a redirection
15213              capture response header Location len 20
15214
15215      >>> Aug  9 20:26:09 localhost \
15216            haproxy[2022]: 127.0.0.1:34014 [09/Aug/2004:20:26:09] proxy-out \
15217            proxy-out/cache1 0/0/0/162/+162 200 +350 - - ---- 0/0/0/0/0 0/0 \
15218            {fr.adserver.yahoo.co||http://fr.f416.mail.} {|864|private||} \
15219            "GET http://fr.adserver.yahoo.com/"
15220
15221      >>> Aug  9 20:30:46 localhost \
15222            haproxy[2022]: 127.0.0.1:34020 [09/Aug/2004:20:30:46] proxy-out \
15223            proxy-out/cache1 0/0/0/182/+182 200 +279 - - ---- 0/0/0/0/0 0/0 \
15224            {w.ods.org||} {Formilux/0.1.8|3495|||} \
15225            "GET http://trafic.1wt.eu/ HTTP/1.1"
15226
15227      >>> Aug  9 20:30:46 localhost \
15228            haproxy[2022]: 127.0.0.1:34028 [09/Aug/2004:20:30:46] proxy-out \
15229            proxy-out/cache1 0/0/2/126/+128 301 +223 - - ---- 0/0/0/0/0 0/0 \
15230            {www.sytadin.equipement.gouv.fr||http://trafic.1wt.eu/} \
15231            {Apache|230|||http://www.sytadin.} \
15232            "GET http://www.sytadin.equipement.gouv.fr/ HTTP/1.1"
15233
15234
15235  8.9. Examples of logs
15236  ---------------------
15237
15238  These are real-world examples of logs accompanied with an explanation. Some of
15239  them have been made up by hand. The syslog part has been removed for better
15240  reading. Their sole purpose is to explain how to decipher them.
15241
15242      >>> haproxy[674]: 127.0.0.1:33318 [15/Oct/2003:08:31:57.130] px-http \
15243            px-http/srv1 6559/0/7/147/6723 200 243 - - ---- 5/3/3/1/0 0/0 \
15244            "HEAD / HTTP/1.0"
15245
15246      => long request (6.5s) entered by hand through 'telnet'. The server replied
15247         in 147 ms, and the session ended normally ('----')
15248
15249      >>> haproxy[674]: 127.0.0.1:33319 [15/Oct/2003:08:31:57.149] px-http \
15250            px-http/srv1 6559/1230/7/147/6870 200 243 - - ---- 324/239/239/99/0 \
15251            0/9 "HEAD / HTTP/1.0"
15252
15253      => Idem, but the request was queued in the global queue behind 9 other
15254         requests, and waited there for 1230 ms.
15255
15256      >>> haproxy[674]: 127.0.0.1:33320 [15/Oct/2003:08:32:17.654] px-http \
15257            px-http/srv1 9/0/7/14/+30 200 +243 - - ---- 3/3/3/1/0 0/0 \
15258            "GET /image.iso HTTP/1.0"
15259
15260      => request for a long data transfer. The "logasap" option was specified, so
15261         the log was produced just before transferring data. The server replied in
15262         14 ms, 243 bytes of headers were sent to the client, and total time from
15263         accept to first data byte is 30 ms.
15264
15265      >>> haproxy[674]: 127.0.0.1:33320 [15/Oct/2003:08:32:17.925] px-http \
15266            px-http/srv1 9/0/7/14/30 502 243 - - PH-- 3/2/2/0/0 0/0 \
15267            "GET /cgi-bin/bug.cgi? HTTP/1.0"
15268
15269      => the proxy blocked a server response either because of an "rspdeny" or
15270         "rspideny" filter, or because the response was improperly formatted and
15271         not HTTP-compliant, or because it blocked sensitive information which
15272         risked being cached. In this case, the response is replaced with a "502
15273         bad gateway". The flags ("PH--") tell us that it was haproxy who decided
15274         to return the 502 and not the server.
15275
```



```
15276      >>> haproxy[18113]: 127.0.0.1:34548 [15/Oct/2003:15:18:55.798] px-http \
15277            px-http/<NOSRV> -1/-1/-1/-1/8490 -1 0 - - CR-- 2/2/2/0/0 0/0 ""
15278
15279      => the client never completed its request and aborted itself ("C---") after
15280         8.5s, while the proxy was waiting for the request headers ("-R--").
15281         Nothing was sent to any server.
15282
15283      >>> haproxy[18113]: 127.0.0.1:34549 [15/Oct/2003:15:19:06.103] px-http \
15284           px-http/<NOSRV> -1/-1/-1/-1/50001 408 0 - - cR-- 2/2/2/0/0 0/0 ""
15285
15286      => The client never completed its request, which was aborted by the
15287         time-out ("c---") after 50s, while the proxy was waiting for the request
15288         headers ("-R--").  Nothing was sent to any server, but the proxy could
15289         send a 408 return code to the client.
15290
15291      >>> haproxy[18989]: 127.0.0.1:34550 [15/Oct/2003:15:24:28.312] px-tcp \
15292            px-tcp/srv1 0/0/5007 0 cD 0/0/0/0/0 0/0
15293
15294      => This log was produced with "option tcplog". The client timed out after
15295         5 seconds ("c----").
15296
15297      >>> haproxy[18989]: 10.0.0.1:34552 [15/Oct/2003:15:26:31.462] px-http \
15298            px-http/srv1 3183/-1/-1/-1/11215 503 0 - - SC-- 205/202/202/115/3 \
15299            0/0 "HEAD / HTTP/1.0"
15300
15301      => The request took 3s to complete (probably a network problem), and the
15302         connection to the server failed ('SC--') after 4 attempts of 2 seconds
15303         (config says 'retries 3'), and no redispatch (otherwise we would have
15304         seen "/+3"). Status code 503 was returned to the client. There were 115
15305         connections on this server, 202 connections on this proxy, and 205 on
15306         the global process. It is possible that the server refused the
15307         connection because of too many already established.
15308
15309
15310  /*
15311   * Local variables:
15312   *  fill-column: 79
15313   * End:
15314   */
```