

```
1
2
3
4
5
6
7
8
9
10
11
12
13
14
15
16
17
18
19
20
21
22
23
24
25
26
27
28
29
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48
49
50
51
52
53
54
55
56
57
58
59
60
61
62
63
64
65
```

HAProxy
Configuration Manual

version 1.7
willy tarreau
2015/12/20

This document covers the configuration language as implemented in the version specified above. It does not provide any hint, example or advice. For such documentation, please refer to the Reference Manual or the Architecture Manual. The summary below is meant to help you search sections by name and navigate through the document.

Note to documentation contributors :

This document is formatted with 80 columns per line, with even number of spaces for indentation and without tabs. Please follow these rules strictly so that it remains easily printable everywhere. If a line needs to be printed verbatim and does not fit, please end each line with a backslash ('\') and continue on next line, indented by two characters. It is also sometimes useful to prefix all output lines (logs, console outs) with 3 closing angle brackets ('>>>') in order to help get the difference between inputs and outputs when it can become ambiguous. If you add sections, please update the summary below for easier searching.

Summary

- 1. Quick reminder about HTTP
 - 1.1. The HTTP transaction model
 - 1.2. HTTP request
 - 1.2.1. The Request line
 - 1.2.2. The request headers
 - 1.3. HTTP response
 - 1.3.1. The Response line
 - 1.3.2. The response headers
- 2. Configuring HAProxy
 - 2.1. Configuration file format
 - 2.2. Quoting and escaping
 - 2.3. Environment variables
 - 2.4. Time format
 - 2.5. Examples
- 3. Global parameters
 - 3.1. Process management and security
 - 3.2. Performance tuning
 - 3.3. Debugging
 - 3.4. Userlists
 - 3.5. Peers
 - 3.6. Mailers
- 4. Proxies
 - 4.1. Proxy keywords matrix
 - 4.2. Alphabetically sorted keywords reference
- 5. Bind and Server options
 - 5.1. Bind options
 - 5.2. Server and default-server options
 - 5.3. Server DNS resolution
 - 5.3.1. Global overview
 - 5.3.2. The resolvers section

```
66
67
68
69
70
71
72
73
74
75
76
77
78
79
80
81
82
83
84
85
86
87
88
89
90
91
92
93
94
95
96
97
98
99
100
101
102
103
104
105
106
107
108
109
110
111
112
113
114
115
116
117
118
119
120
121
122
123
124
125
126
127
128
129
130
```

6. HTTP header manipulation

7. Using ACLs and fetching samples

7.1. ACL basics

- 7.1.1. Matching booleans
- 7.1.2. Matching integers
- 7.1.3. Matching strings
- 7.1.4. Matching regular expressions (regexes)
- 7.1.5. Matching arbitrary data blocks
- 7.1.6. Matching IPv4 and IPv6 addresses

7.2. Using ACLs to form conditions

7.3. Fetching samples

7.3.1. Converters

7.3.2. Fetching samples from internal states

7.3.3. Fetching samples at Layer 4

7.3.4. Fetching samples at Layer 5

7.3.5. Fetching samples from buffer contents (Layer 6)

7.3.6. Fetching HTTP samples (Layer 7)

7.4. Pre-defined ACLs

8. Logging

- 8.1. Log levels
- 8.2. Log formats
 - 8.2.1. Default log format
 - 8.2.2. TCP log format
 - 8.2.3. HTTP log format
 - 8.2.4. Custom log format
 - 8.2.5. Error log format
- 8.3. Advanced logging options
 - 8.3.1. Disabling logging of external tests
 - 8.3.2. Logging before waiting for the session to terminate
 - 8.3.3. Raising log level upon errors
 - 8.3.4. Disabling logging of successful connections
- 8.4. Timing events
- 8.5. Session state at disconnection
- 8.6. Non-printable characters
- 8.7. Capturing HTTP cookies
- 8.8. Capturing HTTP headers
- 8.9. Examples of logs

1. Quick reminder about HTTP

When haproxy is running in HTTP mode, both the request and the response are fully analyzed and indexed, thus it becomes possible to build matching criteria on almost anything found in the contents.

However, it is important to understand how HTTP requests and responses are formed, and how HAProxy decomposes them. It will then become easier to write correct rules and to debug existing configurations.

1.1. The HTTP transaction model

The HTTP protocol is transaction-driven. This means that each request will lead to one and only one response. Traditionally, a TCP connection is established from the client to the server, a request is sent by the client on the connection, the server responds and the connection is closed. A new request will involve a new connection :

[CON1] [REQ1] ... [RESP1] [CL01] [CON2] [REQ2] ... [RESP2] [CL02] ...

In this mode, called the "HTTP close" mode, there are as many connection

131 establishments as there are HTTP transactions. Since the connection is closed
132 by the server after the response, the client does not need to know the content
133 length.
134

135 Due to the transactional nature of the protocol, it was possible to improve it
136 to avoid closing a connection between two subsequent transactions. In this mode
137 however, it is mandatory that the server indicates the content length for each
138 response so that the client does not wait indefinitely. For this, a special
139 header is used: "Content-length". This mode is called the "keep-alive" mode :
140

141 [CON] [REQ1] ... [RESP1] [REQ2] ... [RESP2] [CLO] ...
142

143 Its advantages are a reduced latency between transactions, and less processing
144 power required on the server side. It is generally better than the close mode,
145 but not always because the clients often limit their concurrent connections to
146 a smaller value.
147

148 A last improvement in the communications is the pipelining mode. It still uses
149 keep-alive, but the client does not wait for the first response to send the
150 second request. This is useful for fetching large number of images composing a
151 page :
152

153 [CON] [REQ1] [REQ2] ... [RESP1] [RESP2] [CLO] ...
154

155 This can obviously have a tremendous benefit on performance because the network
156 latency is eliminated between subsequent requests. Many HTTP agents do not
157 correctly support pipelining since there is no way to associate a response with
158 the corresponding request in HTTP. For this reason, it is mandatory for the
159 server to reply in the exact same order as the requests were received.
160

161 By default HAProxy operates in keep-alive mode with regards to persistent
162 connections: for each connection it processes each request and response, and
163 leaves the connection idle on both sides between the end of a response and the
164 start of a new request.
165

166 HAProxy supports 5 connection modes :

- 167 - keep alive : all requests and responses are processed (default)
- 168 - tunnel : only the first request and response are processed,
169 everything else is forwarded with no analysis.
- 170 - passive close : tunnel with "Connection: close" added in both directions.
- 171 - server close : the server-facing connection is closed after the response.
- 172 - forced close : the connection is actively closed after end of response.
173

174 1.2. HTTP request

175 -----
176

177 First, let's consider this HTTP request :

178	Line	Contents
179	number	
180	1	GET /serv/login.php?lang=en&profile=2 HTTP/1.1
181	2	Host: www.mydomain.com
182	3	User-agent: my small browser
183	4	Accept: image/jpeg, image/gif
184	5	Accept: image/png
185		
186		
187		
188		

189 1.2.1. The Request line

190 -----
191

192 Line 1 is the "request line". It is always composed of 3 fields :

- 193 - a METHOD : GET
- 194 - a URI : /serv/login.php?lang=en&profile=2
- 195

196 - a version tag : HTTP/1.1
197

198 All of them are delimited by what the standard calls LWS (linear white spaces),
199 which are commonly spaces, but can also be tabs or line feeds/carriage returns
200 followed by spaces/tabs. The method itself cannot contain any colon (':') and
201 is limited to alphabetic letters. All those various combinations make it
202 desirable that HAProxy performs the splitting itself rather than leaving it to
203 the user to write a complex or inaccurate regular expression.
204

205 The URI itself can have several forms :

206 - A "relative URI" :
207

208 /serv/login.php?lang=en&profile=2
209

210 It is a complete URL without the host part. This is generally what is
211 received by servers, reverse proxies and transparent proxies.
212

213 - An "absolute URI", also called a "URL" :
214

215 http://192.168.0.12:8080/serv/login.php?lang=en&profile=2
216

217 It is composed of a "scheme" (the protocol name followed by '://'), a host
218 name or address, optionally a colon (':') followed by a port number, then
219 a relative URI beginning at the first slash ('/') after the address part.
220 This is generally what proxies receive, but a server supporting HTTP/1.1
221 must accept this form too.
222

223 - a star ('*') : this form is only accepted in association with the OPTIONS
224 method and is not relayable. It is used to inquiry a next hop's
225 capabilities.
226

227 - an address:port combination : 192.168.0.12:80
228

229 This is used with the CONNECT method, which is used to establish TCP
230 tunnels through HTTP proxies, generally for HTTPS, but sometimes for
231 other protocols too.
232

233 In a relative URI, two sub-parts are identified. The part before the question
234 mark is called the "path". It is typically the relative path to static objects
235 on the server. The part after the question mark is called the "query string".
236 It is mostly used with GET requests sent to dynamic scripts and is very
237 specific to the language, framework or application in use.
238

239 1.2.2. The request headers

240 -----
241

242 The headers start at the second line. They are composed of a name at the
243 beginning of the line, immediately followed by a colon (':'). Traditionally,
244 an LWS is added after the colon but that's not required. Then come the values.
245 Multiple identical headers may be folded into one single line, delimiting the
246 values with commas, provided that their order is respected. This is commonly
247 encountered in the "Cookie:" field. A header may span over multiple lines if
248 the subsequent lines begin with an LWS. In the example in 1.2, lines 4 and 5
249 define a total of 3 values for the "Accept:" header.
250

251 Contrary to a common mis-conception, header names are not case-sensitive, and
252 their values are not either if they refer to other header names (such as the
253 "Connection:" header).
254

255 The end of the headers is indicated by the first empty line. People often say
256 that it's a double line feed, which is not exact, even if a double line feed
257 is one valid form of empty line.
258

259 Fortunately, HAProxy takes care of all these complex combinations when indexing
260

headers, checking values and counting them, so there is no reason to worry about the way they could be written, but it is important not to accuse an application of being buggy if it does unusual, valid things.

Important note:

As suggested by RFC2616, HAProxy normalizes headers by replacing line breaks in the middle of headers by LWS in order to join multi-line headers. This is necessary for proper analysis and helps less capable HTTP parsers to work correctly and not to be fooled by such complex constructs.

1.3. HTTP response

An HTTP response looks very much like an HTTP request. Both are called HTTP messages. Let's consider this HTTP response :

Line number	Contents
1	HTTP/1.1 200 OK
2	Content-length: 350
3	Content-Type: text/html

As a special case, HTTP supports so called "Informational responses" as status codes 1xx. These messages are special in that they don't convey any part of the response, they're just used as sort of a signaling message to ask a client to continue to post its request for instance. In the case of a status 100 response the requested information will be carried by the next non-100 response message following the informational one. This implies that multiple responses may be sent to a single request, and that this only works when keep-alive is enabled (1xx messages are HTTP/1.1 only). HAProxy handles these messages and is able to correctly forward and skip them, and only process the next non-100 response. As such, these messages are neither logged nor transformed, unless explicitly state otherwise. Status 101 messages indicate that the protocol is changing over the same connection and that haproxy must switch to tunnel mode, just as if a CONNECT had occurred. Then the Upgrade header would contain additional information about the type of protocol the connection is switching to.

1.3.1. The Response line

Line 1 is the "response line". It is always composed of 3 fields :

- a version tag : HTTP/1.1
- a status code : 200
- a reason : OK

The status code is always 3-digit. The first digit indicates a general status :

- 1xx = informational message to be skipped (eg: 100, 101)
- 2xx = OK, content is following (eg: 200, 206)
- 3xx = OK, no content following (eg: 302, 304)
- 4xx = error caused by the client (eg: 401, 403, 404)
- 5xx = error caused by the server (eg: 500, 502, 503)

Please refer to RFC2616 for the detailed meaning of all such codes. The "reason" field is just a hint, but is not parsed by clients. Anything can be found there, but it's a common practice to respect the well-established messages. It can be composed of one or multiple words, such as "OK", "Found", or "Authentication Required".

HAProxy may emit the following status codes by itself :

Code	When / reason
200	access to stats page, and when replying to monitoring requests

326	301 when performing a redirection, depending on the configured code
327	302 when performing a redirection, depending on the configured code
328	303 when performing a redirection, depending on the configured code
329	307 when performing a redirection, depending on the configured code
330	308 when performing a redirection, depending on the configured code
331	400 for an invalid or too large request
332	401 when an authentication is required to perform the action (when accessing the stats page)
333	403 when a request is forbidden by a "block" ACL or "reqdeny" filter
334	408 when the request timeout strikes before the request is complete
335	500 when haproxy encounters an unrecoverable internal error, such as a memory allocation failure, which should never happen
336	502 when the server returns an empty, invalid or incomplete response, or when an "rspdeny" filter blocks the response.
337	503 when no server was available to handle the request, or in response to monitoring requests which match the "monitor fail" condition
338	504 when the response timeout strikes before the server responds
339	
340	
341	
342	
343	The error 4xx and 5xx codes above may be customized (see "errorloc" in section 4.2).
344	
345	
346	
347	
348	1.3.2. The response headers
349	-----
350	
351	Response headers work exactly like request headers, and as such, HAProxy uses the same parsing function for both. Please refer to paragraph 1.2.2 for more details.
352	
353	
354	
355	
356	2. Configuring HAProxy
357	-----
358	
359	2.1. Configuration file format
360	-----
361	
362	HAProxy's configuration process involves 3 major sources of parameters :
363	
364	- the arguments from the command-line, which always take precedence
365	- the "global" section, which sets process-wide parameters
366	- the proxies sections which can take form of "defaults", "listen", "frontend" and "backend".
367	
368	
369	The configuration file syntax consists in lines beginning with a keyword
370	referenced in this manual, optionally followed by one or several parameters delimited by spaces.
371	
372	
373	
374	2.2. Quoting and escaping
375	-----
376	
377	HAProxy's configuration introduces a quoting and escaping system similar to many programming languages. The configuration file supports 3 types: escaping with a backslash, weak quoting with double quotes, and strong quoting with single quotes.
378	
379	
380	
381	
382	If spaces have to be entered in strings, then they must be escaped by preceding them by a backslash ('\') or by quoting them. Backslashes also have to be escaped by doubling or strong quoting them.
383	
384	
385	Escaping is achieved by preceding a special character by a backslash ('\')
386	
387	\ to mark a space and differentiate it from a delimiter
388	\# to mark a hash and differentiate it from a comment
389	\\ to use a backslash
390	

```

391 \ ' to use a single quote and differentiate it from strong quoting
392 \ ' to use a double quote and differentiate it from weak quoting
393
394 Weak quoting is achieved by using double quotes (""), Weak quoting prevents
395 the interpretation of:
396
397     space as a parameter separator
398     ' single quote as a strong quoting delimiter
399     # hash as a comment start
400
401 Weak quoting permits the interpretation of variables, if you want to use a non
402 -interpreted dollar within a double quoted string, you should escape it with a
403 backslash ("\$"), it does not work outside weak quoting.
404
405 Interpretation of escaping and special characters are not prevented by weak
406 quoting.
407
408 Strong quoting is achieved by using single quotes ('). Inside single quotes,
409 nothing is interpreted, it's the efficient way to quote regexes.
410
411 Quoted and escaped strings are replaced in memory by their interpreted
412 equivalent, it allows you to perform concatenation.
413
414 Example:
415     # those are equivalents:
416     log-format "%{+Q}o\ %t\ %s\ %{-Q}r"
417     log-format "%{+Q}o %t %s %{-Q}r"
418     log-format "%{+Q}o %t %s %{-Q}r"
419     log-format "%{+Q}o %t" %s %{-Q}r"
420     log-format "%{+Q}o %t" %s\ %{-Q}r
421
422     # those are equivalents:
423     reqrep "\^([\^:]*)\ /static/(.*)" "\1 />2"
424     reqrep "\^([\^:]*)\ /static/(.*)" "\1 />2"
425     reqrep "\^([\^:]*)\ /static/(.*)" "\1 />2"
426     reqrep "\^([\^:]*)\ /static/(.*)" "\1 />2"
427
428
429
430
431
432
433
434
435
436
437
438
439
440
441
442
443
444
445
446
447
448
449
450
451
452
453
454
455

```

2.3. Environment variables

HAProxy's configuration supports environment variables. Those variables are interpreted only within double quotes. Variables are expanded during the configuration parsing. Variable names must be preceded by a dollar ("\$\$") and optionally enclosed with braces ("{}") similarly to what is done in Bourne shell. Variable names can contain alphanumerical characters or the character underscore ("_") but should not start with a digit.

Example:

```

bind "fd${FD_APP1}"
log "${LOCAL_SYSL0G}:514" local0 notice # send to local server
user "$HAPROXY_USER"

```

2.4. Time format

Some parameters involve values representing time, such as timeouts. These values are generally expressed in milliseconds (unless explicitly stated otherwise) but may be expressed in any other unit by suffixing the unit to the numeric value. It is important to consider this because it will not be repeated for every keyword. Supported units are :

```

456 - us : microseconds. 1 microsecond = 1/1000000 second
457 - ms : milliseconds. 1 millisecond = 1/1000 second. This is the default.
458 - s : seconds. 1s = 1000ms
459 - m : minutes. 1m = 60s = 60000ms
460 - h : hours. 1h = 60m = 3600s = 3600000ms
461 - d : days. 1d = 24h = 1440m = 86400s = 86400000ms
462
463
464
465
466
467
468
469
470
471
472
473
474
475
476
477
478
479
480
481
482
483
484
485
486
487
488
489
490
491
492
493
494
495
496
497
498
499
500
501
502
503
504
505
506
507
508
509
510
511
512
513
514
515
516
517
518
519
520

```

2.4. Examples

Simple configuration for an HTTP proxy listening on port 80 on all
interfaces and forwarding requests to a single backend "servers" with a
single server "server1" listening on 127.0.0.1:8000

```

global
    daemon
    maxconn 256

defaults
    mode http
    timeout connect 5000ms
    timeout client 5000ms
    timeout server 5000ms

frontend http-in
    bind *:80
    default_backend servers

backend servers
    server server1 127.0.0.1:8000 maxconn 32

```

The same configuration defined with a single listen block. Shorter but
less expressive, especially in HTTP mode.

```

global
    daemon
    maxconn 256

defaults
    mode http
    timeout connect 5000ms
    timeout client 5000ms
    timeout server 5000ms

listen http-in
    bind *:80
    server server1 127.0.0.1:8000 maxconn 32

```

Assuming haproxy is in \$PATH, test these configurations in a shell with:

```

$ sudo haproxy -f configuration.conf -c

```

3. Global parameters

Parameters in the "global" section are process-wide and often OS-specific. They are generally set once for all and do not need being changed once correct. Some of them have command-line equivalents.

The following keywords are supported in the "global" section :

- * Process management and security

```
521 - ca-base
522 - chroot
523 - crt-base
524 - cpu-map
525 - daemon
526 - description
527 - deviceatlas-json-file
528 - deviceatlas-log-level
529 - deviceatlas-separator
530 - deviceatlas-properties-cookie
531 - external-check
532 - gid
533 - group
534 - log
535 - log-tag
536 - log-send-hostname
537 - lua-load
538 - nbproc
539 - node
540 - pidfile
541 - uid
542 - ulimit-n
543 - user
544 - stats
545 - ssl-default-bind-ciphers
546 - ssl-default-bind-options
547 - ssl-default-server-ciphers
548 - ssl-default-server-options
549 - ssl-dh-param-file
550 - ssl-server-verify
551 - unix-bind
552 - 51degrees-data-file
553 - 51degrees-property-name-list
554 - 51degrees-property-separator
555 - 51degrees-cache-size
556
557 * Performance tuning
558 - max-spread-checks
559 - maxconn
560 - maxconnrate
561 - maxcomprate
562 - maxcompcpuusage
563 - maxpipes
564 - maxsserate
565 - maxsslconn
566 - maxssirate
567 - maxzlibmem
568 - nopoll
569 - nokqueue
570 - nopoll
571 - nosplice
572 - nogetaddrinfo
573 - spread-checks
574 - server-state-base
575 - server-state-file
576 - tune.buffer.limit
577 - tune.buffer.reserve
578 - tune.bufsize
579 - tune.chksize
580 - tune.comp.maxlevel
581 - tune.http.cookieielen
582 - tune.http.maxhdr
583 - tune.idletimer
584 - tune.lua.forced-yield
585 - tune.lua.maxmem

586 - tune.lua.session-timeout
587 - tune.lua.task-timeout
588 - tune.lua.service-timeout
589 - tune.maxaccept
590 - tune.maxpollvents
591 - tune.maxrewrite
592 - tune.pattern.cache-size
593 - tune.pipesize
594 - tune.rcvbuf.client
595 - tune.rcvbuf.server
596 - tune.recv_enough
597 - tune.sndbuf.client
598 - tune.sndbuf.server
599 - tune.ssl.cachesize
600 - tune.ssl.lifetime
601 - tune.ssl.force-private-cache
602 - tune.ssl.maxrecord
603 - tune.ssl.default-dh-param
604 - tune.ssl.ctx-cache-size
605 - tune.vars.global-max-size
606 - tune.vars.reqres-max-size
607 - tune.vars.sess-max-size
608 - tune.vars.txn-max-size
609 - tune.zlib.memlevel
610 - tune.zlib.windowsize
611
612 * Debugging
613 - debug
614 - quiet
615
616
617 3.1. Process management and security
618 -----
619
620 ca-base <dir>
621   Assigns a default directory to fetch SSL CA certificates and CRLs from when a
622   relative path is used with "ca-file" or "crl-file" directives. Absolute
623   locations specified in "ca-file" and "crl-file" prevail and ignore "ca-base".
624
625 chroot <jail dir>
626   Changes current directory to <jail dir> and performs a chroot() there before
627   dropping privileges. This increases the security level in case an unknown
628   vulnerability would be exploited, since it would make it very hard for the
629   attacker to exploit the system. This only works when the process is started
630   with superuser privileges. It is important to ensure that <jail_dir> is both
631   empty and unwritable to anyone.
632
633 cpu-map <"all"|"odd"|"even"|process_num> <cpu-set>....
634   On Linux 2.6 and above, it is possible to bind a process to a specific CPU
635   set. This means that the process will never run on other CPUs. The "cpu-map"
636   directive specifies CPU sets for process sets. The first argument is the
637   process number to bind. This process must have a number between 1 and 32 or
638   64, depending on the machine's word size, and any process IDs above nbproc
639   are ignored. It is possible to specify all processes at once using "all",
640   only odd numbers using "odd" or even numbers using "even", just like with the
641   "bind-process" directive. The second and forthcoming arguments are CPU sets.
642   Each CPU set is either a unique number between 0 and 31 or 63 or a range with
643   two such numbers delimited by a dash ('-'). Multiple CPU numbers or ranges
644   may be specified, and the processes will be allowed to bind to all of them.
645   Obviously, multiple "cpu-map" directives may be specified. Each "cpu-map"
646   directive will replace the previous ones when they overlap.
647
648 crt-base <dir>
649   Assigns a default directory to fetch SSL certificates from when a relative
650   path is used with "crtfile" directives. Absolute locations specified after
```

651 "crtfile" prevail and ignore "crt-base".
652
653 daemon
654 Makes the process fork into background. This is the recommended mode of
655 operation. It is equivalent to the command line "-D" argument. It can be
656 disabled by the command line "-db" argument.
657
658 deviceatlas-json-file <path>
659 Sets the path of the DeviceAtlas JSON data file to be loaded by the API.
660 The path must be a valid JSON data file and accessible by Haproxy process.
661
662 deviceatlas-log-level <value>
663 Sets the level of informations returned by the API. This directive is
664 optional and set to 0 by default if not set.
665
666 deviceatlas-separator <char>
667 Sets the character separator for the API properties results. This directive
668 is optional and set to | by default if not set.
669
670 deviceatlas-properties-cookie <name>
671 Sets the client cookie's name used for the detection if the DeviceAtlas
672 Client-side component was used during the request. This directive is optional
673 and set to DAPROPS by default if not set.
674
675 external-check
676 Allows the use of an external agent to perform health checks.
677 This is disabled by default as a security precaution.
678 See "option external-check".
679
680 gid <number>
681 Changes the process' group ID to <number>. It is recommended that the group
682 ID is dedicated to Haproxy or to a small set of similar daemons. Haproxy must
683 be started with a user belonging to this group, or with superuser privileges.
684 Note that if haproxy is started from a user having supplementary groups, it
685 will only be able to drop these groups if started with superuser privileges.
686 See also "group" and "uid".
687
688 group <group name>
689 Similar to "gid" but uses the GID of group name <group name> from /etc/group.
690 See also "gid" and "user".
691
692 log <address> [len <length>] [format <format>] <facility> [max level [min level]]
693 Adds a global syslog server. Up to two global servers can be defined. They
694 will receive logs for startups and exits, as well as all logs from proxies
695 configured with "log global".
696
697 <address> can be one of:
698
699 - An IPv4 address optionally followed by a colon and a UDP port. If
700 no port is specified, 514 is used by default (the standard syslog
701 port).
702
703 - An IPv6 address followed by a colon and optionally a UDP port. If
704 no port is specified, 514 is used by default (the standard syslog
705 port).
706
707 - A filesystem path to a UNIX domain socket, keeping in mind
708 considerations for chroot (be sure the path is accessible inside
709 the chroot) and uid/gid (be sure the path is appropriately
710 writable).
711
712 You may want to reference some environment variables in the address
713 parameter, see section 2.3 about environment variables.
714
715 <length> is an optional maximum line length. Log lines larger than this value

716 will be truncated before being sent. The reason is that syslog
717 servers act differently on log line length. All servers support the
718 default value of 1024, but some servers simply drop larger lines
719 while others do log them. If a server supports long lines, it may
720 make sense to set this value here in order to avoid truncating long
721 lines. Similarly, if a server drops long lines, it is preferable to
722 truncate them before sending them. Accepted values are 80 to 65535
723 inclusive. The default value of 1024 is generally fine for all
724 standard usages. Some specific cases of long captures or
725 JSON-formatted logs may require larger values.
726
727 <format> is the log format used when generating syslog messages. It may be
728 one of the following :
729
730 rfc3164 The RFC3164 syslog message format. This is the default.
731 (https://tools.ietf.org/html/rfc3164)
732
733 rfc5424 The RFC5424 syslog message format.
734 (https://tools.ietf.org/html/rfc5424)
735
736 <facility> must be one of the 24 standard syslog facilities :
737
738 kern user mail daemon auth syslog lpr news
739 uucp cron auth2 ftp ntp audit alert cron2
740 local0 local1 local2 local3 local4 local5 local6 local7
741
742 An optional level can be specified to filter outgoing messages. By default,
743 all messages are sent. If a maximum level is specified, only messages with a
744 severity at least as important as this level will be sent. An optional minimum
745 level can be specified. If it is set, logs emitted with a more severe level
746 than this one will be capped to this level. This is used to avoid sending
747 "emerg" messages on all terminals on some default syslog configurations.
748 Eight levels are known :
749
750 emerg alert crit err warning notice info debug
751
752 log-send-hostname [<string>]
753 Sets the hostname field in the syslog header. If optional "string" parameter
754 is set the header is set to the string contents, otherwise uses the hostname
755 of the system. Generally used if one is not relaying logs through an
756 intermediate syslog server or for simply customizing the hostname printed in
757 the logs.
758
759 log-tag <string>
760 Sets the tag field in the syslog header to this string. It defaults to the
761 program name as launched from the command line, which usually is "haproxy".
762 Sometimes it can be useful to differentiate between multiple processes
763 running on the same host. See also the per-proxy "log-tag" directive.
764
765 lua-load <file>
766 This global directive loads and executes a Lua file. This directive can be
767 used multiple times.
768
769 nbproc <number>
770 Creates <number> processes when going daemon. This requires the "daemon"
771 mode. By default, only one process is created, which is the recommended mode
772 of operation. For systems limited to small sets of file descriptors per
773 process, it may be needed to fork multiple daemons. USING MULTIPLE PROCESSES
774 IS HARDER TO DEBUG AND IS REALLY DISCOURAGED. See also "daemon".
775
776 pidfile <pidfile>
777 Writes pids of all daemons into file <pidfile>. This option is equivalent to
778 the "-p" command line argument. The file must be accessible to the user
779 starting the process. See also "daemon".
780

781 stats bind-process [all | odd | even | <number 1-64> [-<number 1-64>]] ...
782 Limits the stats socket to a certain set of processes numbers. By default the
783 stats socket is bound to all processes, causing a warning to be emitted when
784 nproc is greater than 1 because there is no way to select the target process
785 when connecting. However, by using this setting, it becomes possible to pin
786 the stats socket to a specific set of processes, typically the first one. The
787 warning will automatically be disabled when this setting is used, whatever
788 the number of processes used. The maximum process ID depends on the machine's
789 word size (32 or 64). A better option consists in using the "process" setting
790 of the "stats socket" line to force the process on each line.
791
792 server-state-base <directory>
793 Specifies the directory prefix to be prepended in front of all servers state
794 file names which do not start with a '/'. See also "server-state-file",
795 "load-server-state-from-file" and "server-state-file-name".
796
797 server-state-file <file>
798 Specifies the path to the file containing state of servers. If the path starts
799 with a slash ('/'), it is considered absolute, otherwise it is considered
800 relative to the directory specified using "server-state-base" (if set) or to
801 the current directory. Before reloading HAProxy, it is possible to save the
802 servers' current state using the stats command "show servers state". The
803 output of this command must be written in the file pointed by <file>. When
804 starting up, before handling traffic, HAProxy will read, load and apply state
805 for each server found in the file and available in its current running
806 configuration. See also "server-state-base" and "show servers state",
807 "load-server-state-from-file" and "server-state-file-name"
808
809 ssl-default-bind-ciphers <ciphers>
810 This setting is only available when support for OpenSSL was built in. It sets
811 the default string describing the list of cipher algorithms ("cipher suite")
812 that are negotiated during the SSL/TLS handshake for all "bind" lines which
813 do not explicitly define theirs. The format of the string is defined in
814 "man 1 ciphers" from OpenSSL man pages, and can be for instance a string such
815 as "AES:ALL:!NULL:!RC4:@STRENGTH" (without quotes). Please check the
816 "bind" keyword for more information.
817
818 ssl-default-bind-options [<option>] ...
819 This setting is only available when support for OpenSSL was built in. It sets
820 default ssl-options to force on all "bind" lines. Please check the "bind"
821 keyword to see available options.
822
823 Example:
824 global
825 ssl-default-bind-options no-sslv3 no-tls-tickets
826
827 ssl-default-server-ciphers <ciphers>
828 This setting is only available when support for OpenSSL was built in. It
829 sets the default string describing the list of cipher algorithms that are
830 negotiated during the SSL/TLS handshake with the server, for all "server"
831 lines which do not explicitly define theirs. The format of the string is
832 defined in "man 1 ciphers". Please check the "server" keyword for more
833 information.
834
835 ssl-default-server-options [<option>] ...
836 This setting is only available when support for OpenSSL was built in. It sets
837 default ssl-options to force on all "server" lines. Please check the "server"
838 keyword to see available options.
839
840 ssl-dh-param-file <file>
841 This setting is only available when support for OpenSSL was built in. It sets
842 the default DH parameters that are used during the SSL/TLS handshake when
843 ephemeral Diffie-Hellman (DHE) key exchange is used, for all "bind" lines
844 which do not explicitly define theirs. It will be overridden by custom DH
845 parameters found in a bind certificate file if any. If custom DH parameters

846 are not specified either by using ssl-dh-param-file or by setting them
847 directly in the certificate file, pre-generated DH parameters of the size
848 specified by tune.ssl.default-dh-param will be used. Custom parameters are
849 known to be more secure and therefore their use is recommended.
850 Custom DH parameters may be generated by using the OpenSSL command
851 "openssl dhparam <size>", where size should be at least 2048, as 1024-bit DH
852 parameters should not be considered secure anymore.
853
854 ssl-server-verify [none|required]
855 The default behavior for SSL verify on servers side. If specified to 'none',
856 servers certificates are not verified. The default is 'required' except if
857 forced using cmdline option '-dv'.
858
859 stats socket [<address:port> | <path>] [param*]
860 Binds a UNIX socket to <path> or a TCPv4/v6 address to <address:port>.
861 Connections to this socket will return various statistics outputs and even
862 allow some commands to be issued to change some runtime settings. Please
863 consult section 9.2 "Unix Socket commands" of Management Guide for more
864 details.
865
866 All parameters supported by "bind" lines are supported, for instance to
867 restrict access to some users or their access rights. Please consult
868 section 5.1 for more information.
869
870 stats timeout <timeout, in milliseconds>
871 The default timeout on the stats socket is set to 10 seconds. It is possible
872 to change this value with "stats timeout". The value must be passed in
873 milliseconds, or be suffixed by a time unit among { us, ms, s, m, h }.
874
875 stats maxconn <connections>
876 By default, the stats socket is limited to 10 concurrent connections. It is
877 possible to change this value with "stats maxconn".
878
879 uid <number>
880 Changes the process' user ID to <number>. It is recommended that the user ID
881 is dedicated to HAProxy or to a small set of similar daemons. HAProxy must
882 be started with superuser privileges in order to be able to switch to another
883 one. See also "gid" and "user".
884
885 ulimit-n <number>
886 Sets the maximum number of per-process file-descriptors to <number>. By
887 default, it is automatically computed, so it is recommended not to use this
888 option.
889
890 unix-bind [prefix <prefix>] [mode <mode>] [user <user>] [uid <uid>]
891 [group <group>] [gid <gid>]
892
893 Fixes common settings to UNIX listening sockets declared in "bind" statements.
894 This is mainly used to simplify declaration of those UNIX sockets and reduce
895 the risk of errors, since those settings are most commonly required but are
896 also process-specific. The <prefix> setting can be used to force all socket
897 path to be relative to that directory. This might be needed to access another
898 component's chroot. Note that those paths are resolved before haproxy chroots
899 itself, so they are absolute. The <mode>, <user>, <uid>, <group> and <gid>
900 all have the same meaning as their homonyms used by the "bind" statement. If
901 both are specified, the "bind" statement has priority, meaning that the
902 "unix-bind" settings may be seen as process-wide default settings.
903
904 user <user name>
905 Similar to "uid" but uses the UID of user name <user name> from /etc/passwd.
906 See also "uid" and "group".
907
908 node <name>
909 Only letters, digits, hyphen and underscore are allowed, like in DNS names.
910

This statement is useful in HA configurations where two or more processes or servers share the same IP address. By setting a different node-name on all nodes, it becomes easy to immediately spot what server is handling the traffic.

description <text>

Add a text that describes the instance.

Please note that it is required to escape certain characters (# for example) and this text is inserted into a html page so you should avoid using "<" and ">" characters.

51degrees-data-file <file path>

The path of the 51Degrees data file to provide device detection services. The file should be unzipped and accessible by HAProxy with relevant permissions.

Please note that this option is only available when haproxy has been compiled with USE_51DEGREES.

51degrees-property-name-list [<string>]

A list of 51Degrees property names to be load from the dataset. A full list of names is available on the 51Degrees website:
<https://51degrees.com/resources/property-dictionary>

Please note that this option is only available when haproxy has been compiled with USE_51DEGREES.

51degrees-property-separator <char>

A char that will be appended to every property value in a response header containing 51Degrees results. If not set that will be set as ','.

Please note that this option is only available when haproxy has been compiled with USE_51DEGREES.

51degrees-cache-size <number>

Sets the size of the 51Degrees converter cache to <number> entries. This is an LRU cache which reminds previous device detections and their results. By default, this cache is disabled.

Please note that this option is only available when haproxy has been compiled with USE_51DEGREES.

3.2. Performance tuning

max-spread-checks <delay in milliseconds>

By default, haproxy tries to spread the start of health checks across the smallest health check interval of all the servers in a farm. The principle is to avoid hammering services running on the same server. But when using large check intervals (10 seconds or more), the last servers in the farm take some time before starting to be tested, which can be a problem. This parameter is used to enforce an upper bound on delay between the first and the last check, even if the servers' check intervals are larger. When servers run with shorter intervals, their intervals will be respected though.

maxconn <number>

Sets the maximum per-process number of concurrent connections to <number>. It is equivalent to the command-line argument "-n". Proxies will stop accepting connections when this limit is reached. The "ulimit-n" parameter is automatically adjusted according to this value. See also "ulimit-n". Note: the "select" poller cannot reliably use more than 1024 file descriptors on some platforms. If your platform only supports select and reports "select FAILED" on startup, you need to reduce maxconn until it works (slightly below 500 in general). If this value is not set, it will default to the value

set in DEFAULT_MAXCONN at build time (reported in haproxy -vv) if no memory limit is enforced, or will be computed based on the memory limit, the buffer size, memory allocated to compression, SSL cache size, and use or not of SSL and the associated maxsslconn (which can also be automatic).

maxconnrate <number>

Sets the maximum per-process number of connections per second to <number>. Proxies will stop accepting connections when this limit is reached. It can be used to limit the global capacity regardless of each frontend capacity. It is important to note that this can only be used as a service protection measure, as there will not necessarily be a fair share between frontends when the limit is reached, so it's a good idea to also limit each frontend to some value close to its expected share. Also, lowering tune.maxaccept can improve fairness.

maxcomprate <number>

Sets the maximum per-process input compression rate to <number> kilobytes per second. For each session, if the maximum is reached, the compression level will be decreased during the session. If the maximum is reached at the beginning of a session, the session will not compress at all. If the maximum is not reached, the compression level will be increased up to tune.comp.maxlevel. A value of zero means there is no limit, this is the default value.

maxcompcpuusage <number>

Sets the maximum CPU usage HAProxy can reach before stopping the compression for new requests or decreasing the compression level of current requests. It works like 'maxcomprate' but measures CPU usage instead of incoming data bandwidth. The value is expressed in percent of the CPU used by haproxy. In case of multiple processes (nproc > 1), each process manages its individual usage. A value of 100 disable the limit. The default value is 100. Setting a lower value will prevent the compression work from slowing the whole process down and from introducing high latencies.

maxpipes <number>

Sets the maximum per-process number of pipes to <number>. Currently, pipes are only used by kernel-based tcp splicing. Since a pipe contains two file descriptors, the "ulimit-n" value will be increased accordingly. The default value is maxconn/4, which seems to be more than enough for most heavy usages. The splice code dynamically allocates and releases pipes, and can fall back to standard copy, so setting this value too low may only impact performance.

maxssrrate <number>

Sets the maximum per-process number of sessions per second to <number>. Proxies will stop accepting connections when this limit is reached. It can be used to limit the global capacity regardless of each frontend capacity. It is important to note that this can only be used as a service protection measure, as there will not necessarily be a fair share between frontends when the limit is reached, so it's a good idea to also limit each frontend to some value close to its expected share. Also, lowering tune.maxaccept can improve fairness.

maxsslconn <number>

Sets the maximum per-process number of concurrent SSL connections to <number>. By default there is no SSL-specific limit, which means that the global maxconn setting will apply to all connections. Setting this limit avoids having openssl use too much memory and crash when malloc returns NULL (since it unfortunately does not reliably check for such conditions). Note that the limit applies both to incoming and outgoing connections, so one connection which is deciphered then ciphered accounts for 2 SSL connections. If this value is not set, but a memory limit is enforced, this value will be automatically computed based on the memory limit, maxconn, the buffer size, memory allocated to compression, SSL cache size, and use of SSL in either frontends, backends or both. If neither maxconn nor maxsslconn are specified when there is a memory limit, haproxy will automatically adjust these values

so that 100% of the connections can be made over SSL with no risk, and will consider the sides where it is enabled (frontend, backend, both).

maxsslrate <number>

Sets the maximum per-process number of SSL sessions per second to <number>. SSL listeners will stop accepting connections when this limit is reached. It can be used to limit the global SSL CPU usage regardless of each frontend capacity. It is important to note that this can only be used as a service protection measure, as there will not necessarily be a fair share between frontends when the limit is reached, so it's a good idea to also limit each frontend to some value close to its expected share. It is also important to note that the sessions are accounted before they enter the SSL stack and not after, which also protects the stack against bad handshakes. Also, lowering tune.maxaccept can improve fairness.

maxlibmem <number>

Sets the maximum amount of RAM in megabytes per process usable by the zlib. When the maximum amount is reached, future sessions will not compress as long as RAM is unavailable. When sets to 0, there is no limit. The default value is 0. The value is available in bytes on the UNIX socket with "show info" on the line "MaxLibMemUsage", the memory used by zlib is "ZlibMemUsage" in bytes.

noepoll

Disables the use of the "epoll" event polling system on Linux. It is equivalent to the command-line argument "-de". The next polling system used will generally be "poll". See also "nopoll".

nokqueue

Disables the use of the "kqueue" event polling system on BSD. It is equivalent to the command-line argument "-dk". The next polling system used will generally be "poll". See also "nopoll".

nopoll

Disables the use of the "poll" event polling system. It is equivalent to the command-line argument "-dp". The next polling system used will be "select". It should never be needed to disable "poll" since it's available on all platforms supported by HAProxy. See also "nokqueue" and "noepoll".

nossplice

Disables the use of kernel tcp splicing between sockets on Linux. It is equivalent to the command line argument "-ds". Data will then be copied using conventional and more portable recv/send calls. Kernel tcp splicing is limited to some very recent instances of kernel 2.6. Most versions between 2.6.25 and 2.6.28 are buggy and will forward corrupted data, so they must not be used. This option makes it easier to globally disable kernel splicing in case of doubt. See also "option splice-auto", "option splice-request" and "option splice-response".

nogetaddrinfo

Disables the use of getaddrinfo(3) for name resolving. It is equivalent to the command line argument "-dg". Deprecated gethostbyname(3) will be used.

spread-checks <0..50, in percent>

Sometimes it is desirable to avoid sending agent and health checks to servers at exact intervals, for instance when many logical servers are located on the same physical server. With the help of this parameter, it becomes possible to add some randomness in the check interval between 0 and +/- 50%. A value between 2 and 5 seems to show good results. The default value remains at 0.

tune.buffer.limit <number>

Sets a hard limit on the number of buffers which may be allocated per process. The default value is zero which means unlimited. The minimum non-zero value will always be greater than "tune.buffer.reserve" and should ideally always

be about twice as large. Forcing this value can be particularly useful to limit the amount of memory a process may take, while retaining a sane behaviour. When this limit is reached, sessions which need a buffer wait for another one to be released by another session. Since buffers are dynamically allocated and released, the waiting time is very short and not perceptible provided that limits remain reasonable. In fact sometimes reducing the limit may even increase performance by increasing the CPU cache's efficiency. Tests have shown good results on average HTTP traffic with a limit to 1/10 of the expected global maxconn setting, which also significantly reduces memory usage. The memory savings come from the fact that a number of connections will not allocate 2*tune.bufsize. It is best not to touch this value unless advised to do so by an haproxy core developer.

tune.buffer.reserve <number>

Sets the number of buffers which are pre-allocated and reserved for use only during memory shortage conditions resulting in failed memory allocations. The minimum value is 2 and is also the default. There is no reason a user would want to change this value, it's mostly aimed at haproxy core developers.

tune.bufsize <number>

Sets the buffer size to this size (in bytes). Lower values allow more sessions to coexist in the same amount of RAM, and higher values allow some applications with very large cookies to work. The default value is 16384 and can be changed at build time. It is strongly recommended not to change this from the default value, as very low values will break some services such as statistics, and values larger than default size will increase memory usage, possibly causing the system to run out of memory. At least the global maxconn parameter should be decreased by the same factor as this one is increased. If HTTP request is larger than (tune.bufsize - tune.maxrewrite), haproxy will return HTTP 400 (Bad Request) error. Similarly if an HTTP response is larger than this size, haproxy will return HTTP 502 (Bad Gateway).

tune.chksize <number>

Sets the check buffer size to this size (in bytes). Higher values may help find string or regex patterns in very large pages, though doing so may imply more memory and CPU usage. The default value is 16384 and can be changed at build time. It is not recommended to change this value, but to use better checks whenever possible.

tune.comp.maxlevel <number>

Sets the maximum compression level. The compression level affects CPU usage during compression. This value affects CPU usage during compression. Each session using compression initializes the compression algorithm with this value. The default value is 1.

tune.http.cookie.len <number>

Sets the maximum length of captured cookies. This is the maximum value that the "capture cookie xxx len yyy" will be allowed to take, and any upper value will automatically be truncated to this one. It is important not to set too high a value because all cookie captures still allocate this size whatever their configured value (they share a same pool). This value is per request per response, so the memory allocated is twice this value per connection. When not specified, the limit is set to 63 characters. It is recommended not to change this value.

tune.http.maxhdr <number>

Sets the maximum number of headers in a request. When a request comes with a number of headers greater than this value (including the first line), it is rejected with a "400 Bad Request" status code. Similarly, too large responses are blocked with "502 Bad Gateway". The default value is 101, which is enough for all usages, considering that the widely deployed Apache server uses the same limit. It can be useful to push this limit further to temporarily allow a buggy application to work by the time it gets fixed. Keep in mind that each new header consumes 32bits of memory for each session, so don't push this limit too high.

1171 tune.idletimer <timeout>
1172 Sets the duration after which haproxy will consider that an empty buffer is
1173 probably associated with an idle stream. This is used to optimally adjust
1174 some packet sizes while forwarding large and small data alternatively. The
1175 decision to use splice() or to send large buffers in SSL is modulated by this
1176 parameter. The value is in milliseconds between 0 and 65535. A value of zero
1177 means that haproxy will not try to detect idle streams. The default is 1000,
1178 which seems to correctly detect end user pauses (eg: read a page before
1179 clicking). There should be not reason for changing this value. Please check
1180 tune.ssl.maxrecord below.
1181
1182
1183 tune.lua.forced-yield <number>
1184 This directive forces the Lua engine to execute a yield each <number> of
1185 instructions executed. This permits interrupting a long script and allows the
1186 HAProxy scheduler to process other tasks like accepting connections or
1187 forwarding traffic. The default value is 10000 instructions. If HAProxy often
1188 executes some Lua code but more reactivity is required, this value can be
1189 lowered. If the Lua code is quite long and its result is absolutely required
1190 to process the data, the <number> can be increased.
1191
1192 tune.lua.maxmem
1193 Sets the maximum amount of RAM in megabytes per process usable by Lua. By
1194 default it is zero which means unlimited. It is important to set a limit to
1195 ensure that a bug in a script will not result in the system running out of
1196 memory.
1197
1198 tune.lua.session-timeout <timeout>
1199 This is the execution timeout for the Lua sessions. This is useful for
1200 preventing infinite loops or spending too much time in Lua. This timeout
1201 counts only the pure Lua runtime. If the Lua does a sleep, the sleep is
1202 not taked in account. The default timeout is 4s.
1203
1204 tune.lua.task-timeout <timeout>
1205 Purpose is the same as "tune.lua.session-timeout", but this timeout is
1206 dedicated to the tasks. By default, this timeout isn't set because a task may
1207 remain alive during of the lifetime of HAProxy. For example, a task used to
1208 check servers.
1209
1210 tune.lua.service-timeout <timeout>
1211 This is the execution timeout for the Lua services. This is useful for
1212 preventing infinite loops or spending too much time in Lua. This timeout
1213 counts only the pure Lua runtime. If the Lua does a sleep, the sleep is
1214 not taked in account. The default timeout is 4s.
1215
1216 tune.maxaccept <number>
1217 Sets the maximum number of consecutive connections a process may accept in a
1218 row before switching to other work. In single process mode, higher numbers
1219 give better performance at high connection rates. However in multi-process
1220 modes, keeping a bit of fairness between processes generally is better to
1221 increase performance. This value applies individually to each listener, so
1222 that the number of processes a listener is bound to is taken into account.
1223 This value defaults to 64. In multi-process mode, it is divided by twice
1224 the number of processes the listener is bound to. Setting this value to -1
1225 completely disables the limitation. It should normally not be needed to tweak
1226 this value.
1227
1228 tune.maxpollvents <number>
1229 Sets the maximum amount of events that can be processed at once in a call to
1230 the polling system. The default value is adapted to the operating system. It
1231 has been noticed that reducing it below 200 tends to slightly decrease
1232 latency at the expense of network bandwidth, and increasing it above 200
1233 tends to trade latency for slightly increased bandwidth.
1234
1235 tune.maxrewrite <number>

1236 Sets the reserved buffer space to this size in bytes. The reserved space is
1237 used for header rewriting or appending. The first reads on sockets will never
1238 fill more than bufsize-maxrewrite. Historically it has defaulted to half of
1239 bufsize, though that does not make much sense since there are rarely large
1240 numbers of headers to add. Setting it too high prevents processing of large
1241 requests or responses. Setting it too low prevents addition of new headers
1242 to already large requests or to POST requests. It is generally wise to set it
1243 to about 1024. It is automatically readjusted to half of bufsize if it is
1244 larger than that. This means you don't have to worry about it when changing
1245 bufsize.
1246
1247 tune.pattern.cache-size <number>
1248 Sets the size of the pattern lookup cache to <number> entries. This is an LRU
1249 cache which reminds previous lookups and their results. It is used by ACLs
1250 and maps on slow pattern lookups, namely the ones using the "sub", "reg",
1251 "dir", "dom", "end", "bin" match methods as well as the case-insensitive
1252 strings. It applies to pattern expressions which means that it will be able
1253 to memorize the result of a lookup among all the patterns specified on a
1254 configuration line (including all those loaded from files). It automatically
1255 invalidates entries which are updated using HTTP actions or on the CLI. The
1256 default cache size is set to 10000 entries, which limits its footprint to
1257 about 5 MB on 32-bit systems and 8 MB on 64-bit systems. There is a very low
1258 risk of collision in this cache, which is in the order of the size of the
1259 cache divided by 2^64. Typically, at 10000 requests per second with the
1260 default cache size of 10000 entries, there's 1% chance that a brute force
1261 attack could cause a single collision after 60 years, or 0.1% after 6 years.
1262 This is considered much lower than the risk of a memory corruption caused by
1263 aging components. If this is not acceptable, the cache can be disabled by
1264 setting this parameter to 0.
1265
1266 tune.pipesize <number>
1267 Sets the kernel pipe buffer size to this size (in bytes). By default, pipes
1268 are the default size for the system. But sometimes when using TCP splicing,
1269 it can improve performance to increase pipe sizes, especially if it is
1270 suspected that pipes are not filled and that many calls to splice() are
1271 performed. This has an impact on the kernel's memory footprint, so this must
1272 not be changed if impacts are not understood.
1273
1274 tune.rcvbuf.client <number>
1275 tune.rcvbuf.server <number>
1276 Forces the kernel socket receive buffer size on the client or the server side
1277 to the specified value in bytes. This value applies to all TCP/HTTP frontends
1278 and backends. It should normally never be set, and the default size (0) lets
1279 the kernel autotune this value depending on the amount of available memory.
1280 However it can sometimes help to set it to very low values (eg: 4096) in
1281 order to save kernel memory by preventing it from buffering too large amounts
1282 of received data. Lower values will significantly increase CPU usage though.
1283
1284 tune.recv_enough <number>
1285 Haproxy uses some hints to detect that a short read indicates the end of the
1286 socket buffers. One of them is that a read returns more than <recv_enough>
1287 bytes, which defaults to 10136 (7 segments of 1448 each). This default value
1288 may be changed by this setting to better deal with workloads involving lots
1289 of short messages such as telnet or SSH sessions.
1290
1291 tune.sndbuf.client <number>
1292 tune.sndbuf.server <number>
1293 Forces the kernel socket send buffer size on the client or the server side to
1294 the specified value in bytes. This value applies to all TCP/HTTP frontends
1295 and backends. It should normally never be set, and the default size (0) lets
1296 the kernel autotune this value depending on the amount of available memory.
1297 However it can sometimes help to set it to very low values (eg: 4096) in
1298 order to save kernel memory by preventing it from buffering too large amounts
1299 of received data. Lower values will significantly increase CPU usage though.
1300
1301 Another use case is to prevent write timeouts with extremely slow clients due

to the kernel waiting for a large part of the buffer to be read before notifying haproxy again.

tune.ssl.cachesize <number>

Sets the size of the global SSL session cache, in a number of blocks. A block is large enough to contain an encoded session without peer certificate. An encoded session with peer certificate is stored in multiple blocks depending on the size of the peer certificate. A block uses approximately 200 bytes of memory. The default value may be forced at build time, otherwise defaults to 20000. When the cache is full, the most idle entries are purged and reassigned. Higher values reduce the occurrence of such a purge, hence the number of CPU-intensive SSL handshakes by ensuring that all users keep their session as long as possible. All entries are pre-allocated upon startup and are shared between all processes if "nproc" is greater than 1. Setting this value to 0 disables the SSL session cache.

tune.ssl.force-private-cache

This boolean disables SSL session cache sharing between all processes. It should normally not be used since it will force many renegotiations due to clients hitting a random process. But it may be required on some operating systems where none of the SSL cache synchronization method may be used. In this case, adding a first layer of hash-based load balancing before the SSL layer might limit the impact of the lack of session sharing.

tune.ssl.lifetime <timeout>

Sets how long a cached SSL session may remain valid. This time is expressed in seconds and defaults to 300 (5 min). It is important to understand that it does not guarantee that sessions will last that long, because if the cache is full, the longest idle sessions will be purged despite their configured lifetime. The real usefulness of this setting is to prevent sessions from being used for too long.

tune.ssl.maxrecord <number>

Sets the maximum amount of bytes passed to SSL_write() at a time. Default value 0 means there is no limit. Over SSL/TLS, the client can decipher the data only once it has received a full record. With large records, it means that clients might have to download up to 16kB of data before starting to process them. Limiting the value can improve page load times on browsers located over high latency or low bandwidth networks. It is suggested to find optimal values which fit into 1 or 2 TCP segments (generally 1448 bytes over Ethernet with TCP timestamps enabled, or 1460 when timestamps are disabled), keeping in mind that SSL/TLS add some overhead. Typical values of 1419 and 2859 gave good results during tests. Use "strace -e trace=write" to find the best value. Haproxy will automatically switch to this setting after an idle stream has been detected (see tune.idletimer above).

tune.ssl.default-dh-param <number>

Sets the maximum size of the Diffie-Hellman parameters used for generating the ephemeral/temporary Diffie-Hellman key in case of DHE key exchange. The final size will try to match the size of the server's RSA (or DSA) key (e.g, a 2048 bits temporary DH key for a 2048 bits RSA key), but will not exceed this maximum value. Default value if 1024. Only 1024 or higher values are allowed. Higher values will increase the CPU load, and values greater than 1024 bits are not supported by Java 7 and earlier clients. This value is not used if static Diffie-Hellman parameters are supplied either directly in the certificate file or by using the ssl-dh-param-file parameter.

tune.ssl.ctx-cache-size <number>

Sets the size of the cache used to store generated certificates to <number> entries. This is a LRU cache. Because generating a SSL certificate dynamically is expensive, they are cached. The default cache size is set to 1000 entries.

tune.vars.global-max-size <size>

tune.vars.reqres-max-size <size>

tune.vars.ssess-max-size <size>

tune.vars.txn-max-size <size>

These four tunes helps to manage the allowed amount of memory used by the variables system. "global" limits the memory for all the systems. "sess" limit the memory by session, "txn" limits the memory by transaction and "reqres" limits the memory for each request or response processing. during the accounting, "sess" embed "txn" and "txn" embed "reqres".

By example, we considers that "tune.vars.ssess-max-size" is fixed to 100, "tune.vars.txn-max-size" is fixed to 100, "tune.vars.reqres-max-size" is also fixed to 100. If we create a variable "txn.var" that contains 100 bytes, we cannot create any more variable in the other contexts.

tune.zlib.memlevel <number>

Sets the memLevel parameter in zlib initialization for each session. It defines how much memory should be allocated for the internal compression state. A value of 1 uses minimum memory but is slow and reduces compression ratio, a value of 9 uses maximum memory for optimal speed. Can be a value between 1 and 9. The default value is 8.

tune.zlib.windowsize <number>

Sets the window size (the size of the history buffer) as a parameter of the zlib initialization for each session. Larger values of this parameter result in better compression at the expense of memory usage. Can be a value between 8 and 15. The default value is 15.

3.3. Debugging

debug

Enables debug mode which dumps to stdout all exchanges, and disables forking into background. It is the equivalent of the command-line argument "-d". It should never be used in a production configuration since it may prevent full system startup.

quiet

Do not display any message during startup. It is equivalent to the command-line argument "-q".

3.4. Userlists

It is possible to control access to frontend/backend/listen sections or to http stats by allowing only authenticated and authorized users. To do this, it is required to create at least one userlist and to define users.

userlist <listname>

Creates new userlist with name <listname>. Many independent userlists can be used to store authentication & authorization data for independent customers.

group <groupname> [users <user>,<user>,(...)]

Adds group <groupname> to the current userlist. It is also possible to attach users to this group by using a comma separated list of names proceeded by "users" keyword.

user <username> [password|insecure-password <password>]

[groups <group>,<group>,(...)]

Adds user <username> to the current userlist. Both secure (encrypted) and insecure (unencrypted) passwords can be used. Encrypted passwords are evaluated using the crypt(3) function so depending of the system's capabilities, different algorithms are supported. For example modern Glibc based Linux system supports MD5, SHA-256, SHA-512 and of course classic, DES-based method of encrypting passwords.

1430

```
1431
1432
1433
1434
1435
1436
1437
1438
1439
1440
1441
1442
1443
1444
1445
1446
1447
1448
1449
1450
1451
1452
1453
1454
1455
1456
1457
1458
1459
1460
1461
1462
1463
1464
1465
1466
1467
1468
1469
1470
1471
1472
1473
1474
1475
1476
1477
1478
1479
1480
1481
1482
1483
1484
1485
1486
1487
1488
1489
1490
1491
1492
1493
1494
1495
```

Example:

```
userlist L1
group G1 users tiger,scott
group G2 users xdb,scott

user tiger password $6$k6y3o.eP$JlKBx9za9667qe4(...)xH$wRv6J.C0/D7cV91
user scott insecure-password elgato
user xdb insecure-password hello

userlist L2
group G1
group G2

user tiger password $6$k6y3o.eP$JlKBx(...)xH$wRv6J.C0/D7cV91 groups G1
user scott insecure-password elgato groups G1,G2
user xdb insecure-password hello groups G2
```

Please note that both lists are functionally identical.

3.5. Peers

It is possible to propagate entries of any data-types in stick-tables between several haproxy instances over TCP connections in a multi-master fashion. Each instance pushes its local updates and insertions to remote peers. The pushed values overwrite remote ones without aggregation. Interrupted exchanges are automatically detected and recovered from the last known point.

In addition, during a soft restart, the old process connects to the new one using such a TCP connection to push all its entries before the new process tries to connect to other peers. That ensures very fast replication during a reload, it typically takes a fraction of a second even for large tables.

Note that Server IDs are used to identify servers remotely, so it is important that configurations look similar or at least that the same IDs are forced on each server on all participants.

peers <peersect>
Creates a new peer list with name <peersect>. It is an independent section, which is referenced by one or more stick-tables.

disabled
Disables a peers section. It disables both listening and any synchronization related to this section. This is provided to disable synchronization of stick tables without having to comment out all "peers" references.

enable
This re-enables a disabled peers section which was previously disabled.

peer <peername> <ip>:<port>
Defines a peer inside a peers section.
If <peername> is set to the local peer name (by default hostname, or forced using "-L" command line option), haproxy will listen for incoming remote peer connection on <ip>:<port>. Otherwise, <ip>:<port> defines where to connect to join the remote peer, and <peername> is used at the protocol level to identify and validate the remote peer on the server side.

During a soft restart, local peer <ip>:<port> is used by the old instance to connect the new one and initiate a complete replication (teaching process).

It is strongly recommended to have the exact same peers declaration on all peers and to only rely on the "-L" command line argument to change the local peer name. This makes it easier to maintain coherent configuration files across all peers.

You may want to reference some environment variables in the address parameter, see section 2.3 about environment variables.

```
1496
1497
1498
1499
1500
1501
1502
1503
1504
1505
1506
1507
1508
1509
1510
1511
1512
1513
1514
1515
1516
1517
1518
1519
1520
1521
1522
1523
1524
1525
1526
1527
1528
1529
1530
1531
1532
1533
1534
1535
1536
1537
1538
1539
1540
1541
1542
1543
1544
1545
1546
1547
1548
1549
1550
1551
1552
1553
1554
1555
1556
1557
1558
1559
1560
```

Example:

```
peers mypeers
peer haproxy1 192.168.0.1:1024
peer haproxy2 192.168.0.2:1024
peer haproxy3 10.2.0.1:1024

backend mybackend
mode tcp
balance roundrobin
stick-table type ip size 20k peers mypeers
stick on src

server srv1 192.168.0.30:80
server srv2 192.168.0.31:80
```

3.6. Mailers

It is possible to send email alerts when the state of servers changes.
If configured email alerts are sent to each mailer that is configured in a mailers section. Email is sent to mailers using SMTP.

mailers <mailersect>
Creates a new mailer list with the name <mailersect>. It is an independent section which is referenced by one or more proxies.

mailer <mailername> <ip>:<port>
Defines a mailer inside a mailers section.

Example:

```
mailers mymailers
mailer smtp1 192.168.0.1:587
mailer smtp2 192.168.0.2:587

backend mybackend
mode tcp
balance roundrobin

email-alert mailers mymailers
email-alert from test1@horms.org
email-alert to test2@horms.org

server srv1 192.168.0.30:80
server srv2 192.168.0.31:80
```

4. Proxies

Proxy configuration can be located in a set of sections :

- defaults [<name>]
- frontend <name>
- backend <name>
- listen <name>

A "defaults" section sets default parameters for all other sections following its declaration. Those default parameters are reset by the next "defaults" section. See below for the list of parameters which can be set in a "defaults" section. The name is optional but its use is encouraged for better readability.

A "frontend" section describes a set of listening sockets accepting client connections.

A "backend" section describes a set of servers to which the proxy will connect

to forward incoming connections.

A "listen" section defines a complete proxy with its frontend and backend parts combined in one section. It is generally useful for TCP-only traffic.

All proxy names must be formed from upper and lower case letters, digits, '-' (dash), '_' (underscore), '.' (dot) and ':' (colon). ACL names are case-sensitive, which means that "www" and "WWW" are two different proxies.

Historically, all proxy names could overlap, it just caused troubles in the logs. Since the introduction of content switching, it is mandatory that two proxies with overlapping capabilities (frontend/backend) have different names. However, it is still permitted that a frontend and a backend share the same name, as this configuration seems to be commonly encountered.

Right now, two major proxy modes are supported : "tcp", also known as layer 4, and "http", also known as layer 7. In layer 4 mode, HAProxy simply forwards bidirectional traffic between two sides. In layer 7 mode, HAProxy analyzes the protocol, and can interact with it by allowing, blocking, switching, adding, modifying, or removing arbitrary contents in requests or responses, based on arbitrary criteria.

In HTTP mode, the processing applied to requests and responses flowing over a connection depends in the combination of the frontend's HTTP options and the backend's. HAProxy supports 5 connection modes :

- KAL : keep alive ("option http-keep-alive") which is the default mode : all requests and responses are processed, and connections remain open but idle between responses and new requests.

- TUN: tunnel ("option http-tunnel"): this was the default mode for versions 1.0 to 1.5-dev21 : only the first request and response are processed, and everything else is forwarded with no analysis at all. This mode should not be used as it creates lots of trouble with logging and HTTP processing.

- PCL: passive close ("option httpclose") : exactly the same as tunnel mode, but with "Connection: close" appended in both directions to try to make both ends close after the first request/response exchange.

- `SCl: server close ("option http-server-close")` : the server-facing connection is closed after the end of the response is received, but the client-facing connection remains open.

- FCL: forced close ("option forceclose") : the connection is actively closed after the end of the response.

The effective mode that will be applied to a connection passing through a frontend and a backend can be determined by both proxy modes according to the following matrix, but in short, the modes are symmetric, keep-alive is the weakest option and force close is the strongest.

Backend mode

| KAL | TUN | PCL | SCL | FCL |

	KAL	KAL	TUN	PCL	SCL	FCL
1	0.78	0.69	0.69	0.69	0.69	0.69
2	0.78	0.69	0.69	0.69	0.69	0.69
3	0.78	0.69	0.69	0.69	0.69	0.69
4	0.78	0.69	0.69	0.69	0.69	0.69
5	0.78	0.69	0.69	0.69	0.69	0.69
6	0.78	0.69	0.69	0.69	0.69	0.69
7	0.78	0.69	0.69	0.69	0.69	0.69
8	0.78	0.69	0.69	0.69	0.69	0.69
9	0.78	0.69	0.69	0.69	0.69	0.69
10	0.78	0.69	0.69	0.69	0.69	0.69
11	0.78	0.69	0.69	0.69	0.69	0.69
12	0.78	0.69	0.69	0.69	0.69	0.69
13	0.78	0.69	0.69	0.69	0.69	0.69
14	0.78	0.69	0.69	0.69	0.69	0.69
15	0.78	0.69	0.69	0.69	0.69	0.69
16	0.78	0.69	0.69	0.69	0.69	0.69
17	0.78	0.69	0.69	0.69	0.69	0.69
18	0.78	0.69	0.69	0.69	0.69	0.69
19	0.78	0.69	0.69	0.69	0.69	0.69
20	0.78	0.69	0.69	0.69	0.69	0.69
21	0.78	0.69	0.69	0.69	0.69	0.69
22	0.78	0.69	0.69	0.69	0.69	0.69
23	0.78	0.69	0.69	0.69	0.69	0.69
24	0.78	0.69	0.69	0.69	0.69	0.69
25	0.78	0.69	0.69	0.69	0.69	0.69
26	0.78	0.69	0.69	0.69	0.69	0.69
27	0.78	0.69	0.69	0.69	0.69	0.69
28	0.78	0.69	0.69	0.69	0.69	0.69
29	0.78	0.69	0.69	0.69	0.69	0.69
30	0.78	0.69	0.69	0.69	0.69	0.69
31	0.78	0.69	0.69	0.69	0.69	0.69
32	0.78	0.69	0.69	0.69	0.69	0.69
33	0.78	0.69	0.69	0.69	0.69	0.69
34	0.78	0.69	0.69	0.69	0.69	0.69
35	0.78	0.69	0.69	0.69	0.69	0.69
36	0.78	0.69	0.69	0.69	0.69	0.69
37	0.78	0.69	0.69	0.69	0.69	0.69
38	0.78	0.69	0.69	0.69	0.69	0.69
39	0.78	0.69	0.69	0.69	0.69	0.69
40	0.78	0.69	0.69	0.69	0.69	0.69
41	0.78	0.69	0.69	0.69	0.69	0.69
42	0.78	0.69	0.69	0.69	0.69	0.69
43	0.78	0.69	0.69	0.69	0.69	0.69
44	0.78	0.69	0.69	0.69	0.69	0.69
45	0.78	0.69	0.69	0.69	0.69	0.69
46	0.78	0.69	0.69	0.69	0.69	0.69
47	0.78	0.69	0.69	0.69	0.69	0.69
48	0.78	0.69	0.69	0.69	0.69	0.69
49	0.78	0.69	0.69	0.69	0.69	0.69
50	0.78	0.69	0.69	0.69	0.69	0.69
51	0.78	0.69	0.69	0.69	0.69	0.69
52	0.78	0.69	0.69	0.69	0.69	0.69
53	0.78	0.69	0.69	0.69	0.69	0.69
54	0.78	0.69	0.69	0.69	0.69	0.69
55	0.78	0.69	0.69	0.69	0.69	0.69
56	0.78	0.69	0.69	0.69	0.69	0.69
57	0.78	0.69	0.69	0.69	0.69	0.69
58	0.78	0.69	0.69	0.69	0.69	0.69
59	0.78	0.69	0.69	0.69	0.69	0.69
60	0.78	0.69	0.69	0.69	0.69	0.69
61	0.78	0.				

TUN	TUN	TUN	PCL	SCL	FCL
-----	-----	-----	-----	-----	-----

	PCL	PCL	PCL	PCL	PCL	FCL	FCL
1							
2							
3							
4							
5							
6							
7							
8							
9							
10							
11							
12							
13							
14							
15							
16							
17							
18							
19							
20							
21							
22							
23							
24							
25							
26							
27							
28							
29							
30							
31							
32							
33							
34							
35							
36							
37							
38							
39							
40							
41							
42							
43							
44							
45							
46							
47							
48							
49							
50							
51							
52							
53							
54							
55							
56							
57							
58							
59							
60							
61							
62							
63							
64							
65							
66							
67							
68							
69							
70							
71							
72							
73							
74							
75							
76							
77							
78							
79							
80							
81				</			

SCL		SCL		SCL		FCL		FCL
-----	--	-----	--	-----	--	-----	--	-----

	FCL	FCL	FCL	FCL	FCL

1266	1267	1268	1269	1270	1271	1272	1273	1274	1275	1276	1277	1278	1279	1280	1281	1282	1283	1284	1285	1286	1287	1288	1289	1290	1291	1292	1293	1294	1295	1296	1297	1298	1299	1300	1301	1302	1303	1304	1305	1306	1307	1308	1309	1310	1311	1312	1313	1314	1315	1316	1317	1318	1319	1320	1321	1322	1323	1324	1325	1326	1327	1328	1329	1330	1331	1332	1333	1334	1335	1336	1337	1338	1339	1340	1341	1342	1343	1344	1345	1346	1347	1348	1349	1350	1351	1352	1353	1354	1355	1356	1357	1358	1359	1360	1361	1362	1363	1364	1365	1366	1367	1368	1369	1370	1371	1372	1373	1374	1375	1376	1377	1378	1379	1380	1381	1382	1383	1384	1385	1386	1387	1388	1389	1390	1391	1392	1393	1394	1395	1396	1397	1398	1399	1400	1401	1402	1403	1404	1405	1406	1407	1408	1409	1410	1411	1412	1413	1414	1415	1416	1417	1418	1419	1420	1421	1422	1423	1424	1425	1426	1427	1428	1429	1430	1431	1432	1433	1434	1435	1436	1437	1438	1439	1440	1441	1442	1443	1444	1445	1446	1447	1448	1449	1450	1451	1452	1453	1454	1455	1456	1457	1458	1459	1460	1461	1462	1463	1464	1465	1466	1467	1468	1469	1470	1471	1472	1473	1474	1475	1476	1477	1478	1479	1480	1481	1482	1483	1484	1485	1486	1487	1488	1489	1490	1491	1492	1493	1494	1495	1496	1497	1498	1499	1500	1501	1502	1503	1504	1505	1506	1507	1508	1509	1510	1511	1512	1513	1514	1515	1516	1517	1518	1519	1520	1521	1522	1523	1524	1525	1526	1527	1528	1529	1530	1531	1532	1533	1534	1535	1536	1537	1538	1539	1540	1541	1542	1543	1544	1545	1546	1547	1548	1549	1550	1551	1552	1553	1554	1555	1556	1557	1558	1559	1560	1561	1562	1563	1564	1565	1566	1567	1568	1569	1570	1571	1572	1573	1574	1575	1576	1577	1578	1579	1580	1581	1582	1583	1584	1585	1586	1587	1588	1589	1590	1591	1592	1593	1594	1595	1596	1597	1598	1599	1600	1601	1602	1603	1604	1605	1606	1607	1608	1609	1610	1611	1612	1613	1614	1615	1616	1617	1618	1619	1620	1621	1622	1623	1624	1625	1626	1627	1628	1629	1630	1631	1632	1633	1634	1635	1636	1637	1638	1639	1640	1641	1642	1643	1644	1645	1646	1647	1648	1649	1650	1651	1652	1653	1654	1655	1656	1657	1658	1659	1660	1661	1662	1663	1664	1665	1666	1667	1668	1669	1670	1671	1672	1673	1674	1675	1676	1677	1678	1679	1680	1681	1682	1683	1684	1685	1686	1687	1688	1689	1690	1691	1692	1693	1694	1695	1696	1697	1698	1699	1700	1701	1702	1703	1704	1705	1706	1707	1708	1709	1710	1711	1712	1713	1714	1715	1716	1717	1718	1719
------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------

4.1. Proxy keywords matrix

The following list of keywords is supported. Most of them may only be used in a limited set of section types. Some of them are marked as "deprecated" because they are inherited from an old syntax which may be confusing or functionally limited, and there are new recommended keywords to replace them. Keywords marked with "(*)" can be optionally inverted using the "no" prefix, eg. "no option constants". This makes sense when the option has been enabled by default and must be disabled for a specific instance. Such options may also be prefixed with "default" in order to restore default settings regardless of what has been specified in a previous "defaults" section.

keyword	defaults	frontend	listen	backend
acl	-	X	X	X
appsession	-	-	-	-
backlog	X	X	X	X
balance	X	-	X	-
bind	-	X	X	-
bind-process	X	X	X	X
block	-	X	X	X
capture cookie	-	X	X	-
capture request header	-	X	X	-
capture response header	-	X	X	-
clitimeout	(deprecated)	X	X	-
compression	X	X	X	X
contimeout	(deprecated)	X	-	X
cookie	X	-	X	X
declare capture	-	X	X	-
default-server	X	-	X	X
default-backend	X	X	X	-
description	-	X	X	X
disabled	X	X	X	X
dispatch	-	-	X	X
email-alert from	X	X	X	X
email-alert level	X	X	X	X
email-alert mailers	X	X	X	X
email-alert myhostname	X	X	X	X
email-alert to	X	X	X	X
enabled	X	X	X	X
errorfile	X	X	X	X
errorloc	X	X	X	X
errorloc302	X	X	X	X
-- keyword	defaults	frontend	listen	-- backend
errorloc303	X	X	X	X
force-persist	-	X	X	X
fullconn	X	-	X	X
grace	X	X	X	X
hash-type	X	-	X	X
http-check disable-on-404	X	-	X	X
http-check expect	-	-	X	X
http-check send-state	X	-	X	X
http-request	-	X	X	X
http-response	-	-	X	X
http-reuse	X	-	X	X
http-send-name-header	-	-	X	X
id	-	X	X	X
ignore-persist	-	X	X	X
load-server-state-from-file	-	X	X	X
log	X	-	X	X
log-format	(*)	X	X	-

1821	timeout server-fin		X	-	X	X	X
1822	timeout srvtimeout	(deprecated)	X	-	X	X	X
1823	timeout tarpit		X	X	X	X	X
1824	timeout tunnel		X	-	X	X	X
1825	transparent	(deprecated)	X	-	X	X	X
1826	unique-id-format		X	X	X	X	-
1827	unique-id-header		X	X	X	X	-
1828	use_backend		-	X	X	X	-
1829	use-server		-	-	X	X	-
1830							X
1831	keyword	defaults	frontend	listen			backend
1832							
1833							
1834							
1835	4.2. Alphabetically sorted keywords reference						
1836							
1837	This section provides a description of each keyword and its usage.						
1838							
1839							
1840	acl <aclname> <criterion> [flags] [operator] <value> ...						
1841	Declare or complete an access list.						
1842	May be used in sections :	defaults	frontend	listen	backend		
1843		no	yes	yes	yes		
1844	Example:						
1845	acl invalid_src src	0.0.0.0/7	224.0.0.0/3				
1846	acl invalid_src src_port	0:1023					
1847	acl local_dst hdr(host) -i localhost						
1848							
1849	See section 7 about ACL usage.						
1850							
1851							
1852	appsession <cookie> len <length> timeout <holdtime>						
1853	[request-learn] [prefix] [mode <path-parameters[query-string]>]						
1854	Define session stickiness on an existing application cookie.						
1855	May be used in sections :	defaults	frontend	listen	backend		
1856		no	no	yes	yes		
1857	Arguments :						
1858	<cookie>	this is the name of the cookie used by the application and which HAProxy will have to learn for each new session.					
1859							
1860							
1861	<length>	this is the max number of characters that will be memorized and checked in each cookie value.					
1862							
1863							
1864	<holdtime>	this is the time after which the cookie will be removed from memory if unused. If no unit is specified, this time is in milliseconds.					
1865							
1866							
1867							
1868	request-learn						
1869							
1870	If this option is specified, then haproxy will be able to learn the cookie found in the request in case the server does not specify any in response. This is typically what happens with PHPSESSID cookies, or when haproxy's session expires before the application's session and the correct server is selected. It is recommended to specify this option to improve reliability						
1871							
1872							
1873							
1874							
1875							
1876	prefix	When this option is specified, haproxy will match on the cookie prefix (or URL parameter prefix). The appsession value is the data following this prefix.					
1877							
1878							
1879							
1880	Example :						
1881	appsession ASPSESSIONID len 64 timeout 3h prefix						
1882							
1883	This will match the cookie ASPSESSIONIDXXX=XXXX,						
1884	the appsession value will be XXXX=XXXXX.						
1885							

```

1886 mode
1887 This option allows to change the URL parser mode.
1888 2 modes are currently supported :
1889 - path-parameters :
1890 The parser looks for the appsession in the path parameters
1891 part (each parameter is separated by a semi-colon), which is
1892 convenient for JSESSIONID for example.
1893 This is the default mode if the option is not set.
1894 - query-string :
1895 In this mode, the parser will look for the appsession in the
1896 query string.
1897
1898 As of version 1.6, appsessions was removed. It is more flexible and more
1899 convenient to use stick-tables instead, and stick-tables support multi-master
1900 replication and data conservation across reloads, which appsessions did not.
1901
1902 See also : "cookie", "capture cookie", "balance", "stick", "stick-table",
1903 "ignore-persist", "nbproc" and "bind-process".
1904
1905 backlog <conns>
1906 Give hints to the system about the approximate listen backlog desired size
1907 May be used in sections : defaults | frontend | listen | backend
1908 yes | yes | yes | no
1909
1910 Arguments :
1911 <conns> is the number of pending connections. Depending on the operating
1912 system, it may represent the number of already acknowledged
1913 connections, of non-acknowledged ones, or both.
1914
1915 In order to protect against SYN flood attacks, one solution is to increase
1916 the system's SYN backlog size. Depending on the system, sometimes it is just
1917 tunable via a system parameter, sometimes it is not adjustable at all, and
1918 sometimes the system relies on hints given by the application at the time of
1919 the listen() syscall. By default, HAProxy passes the frontend's maxconn value
1920 to the listen() syscall. On systems which can make use of this value, it can
1921 sometimes be useful to be able to specify a different value, hence this
1922 backlog parameter.
1923
1924 On Linux 2.4, the parameter is ignored by the system. On Linux 2.6, it is
1925 used as a hint and the system accepts up to the smallest greater power of
1926 two, and never more than some limits (usually 32768).
1927
1928 See also : "maxconn" and the target operating system's tuning guide.
1929
1930 balance <algorithm> [ <arguments> ]
1931 balance url_param <param> [check_post]
1932 Define the load balancing algorithm to be used in a backend.
1933 May be used in sections : defaults | frontend | listen | backend
1934 yes | no | yes | yes
1935
1936 Arguments :
1937 <algorithm> is the algorithm used to select a server when doing load
1938 balancing. This only applies when no persistence information
1939 is available, or when a connection is redispatched to another
1940 server. <algorithm> may be one of the following :
1941
1942 roundrobin Each server is used in turns, according to their weights.
1943 This is the smoothest and fairest algorithm when the server's
1944 processing time remains equally distributed. This algorithm
1945 is dynamic, which means that server weights may be adjusted
1946 on the fly for slow starts for instance. It is limited by
1947 design to 4095 active servers per backend. Note that in some
1948 large farms, when a server becomes up after having been down
1949 for a very short time, it may sometimes take a few hundreds
1950 requests for it to be re-integrated into the farm and start
1951 receiving traffic. This is normal, though very rare. It is

```

indicated here in case you would have the chance to observe it, so that you don't worry.

static-rr

Each server is used in turns, according to their weights. This algorithm is as similar to roundrobin except that it is static, which means that changing a server's weight on the fly will have no effect. On the other hand, it has no design limitation on the number of servers, and when a server goes up, it is always immediately reintroduced into the farm, once the full map is recomputed. It also uses slightly less CPU to run (around -1%).

leastconn

The server with the lowest number of connections receives the connection. Round-robin is performed within groups of servers of the same load to ensure that all servers will be used. Use of this algorithm is recommended where very long sessions are expected, such as LDAP, SQL, TSE, etc... but is not very well suited for protocols using short sessions such as HTTP. This algorithm is dynamic, which means that server weights may be adjusted on the fly for slow starts for instance.

first

The first server with available connection slots receives the connection. The servers are chosen from the lowest numeric identifier to the highest (see server parameter "id"), which defaults to the server's position in the farm. Once a server reaches its maxconn value, the next server is used. It does not make sense to use this algorithm without setting maxconn. The purpose of this algorithm is to always use the smallest number of servers so that extra servers can be powered off during non-intensive hours. This algorithm ignores the server weight, and brings more benefit to long session such as RDP or IMAP than HTTP, though it can be useful there too. In order to use this algorithm efficiently, it is recommended that a cloud controller regularly checks server usage to turn them off when unused, and regularly checks backend queue to turn new servers on when the queue inflates. Alternatively, using "http-check send-state" may inform servers on the load.

source

The source IP address is hashed and divided by the total weight of the running servers to designate which server will receive the request. This ensures that the same client IP address will always reach the same server as long as no server goes down or up. If the hash result changes due to the number of running servers changing, many clients will be directed to a different server. This algorithm is generally used in TCP mode where no cookie may be inserted. It may also be used on the Internet to provide a best-effort stickiness to clients which refuse session cookies. This algorithm is static by default, which means that changing a server's weight on the fly will have no effect, but this can be changed using "hash-type".

uri

This algorithm hashes either the left part of the URI (before the question mark) or the whole URI (if the "whole" parameter is present) and divides the hash value by the total weight of the running servers. The result designates which server will receive the request. This ensures that the same URI will always be directed to the same server as long as no server goes up or down. This is used with proxy caches and anti-virus proxies in order to maximize the cache hit rate. Note that this algorithm may only be used in an HTTP backend. This algorithm is static by default, which means that changing a server's weight on the fly will have no effect, but this can be changed using "hash-type".

This algorithm supports two optional parameters "len" and "depth", both followed by a positive integer number. These options may be helpful when it is needed to balance servers based on the beginning of the URI only. The "len" parameter indicates that the algorithm should only consider that many characters at the beginning of the URI to compute the hash. Note that having "len" set to 1 rarely makes sense since most URIs start with a leading "/".

The "depth" parameter indicates the maximum directory depth to be used to compute the hash. One level is counted for each slash in the request. If both parameters are specified, the evaluation stops when either is reached.

The URL parameter specified in argument will be looked up in the query string of each HTTP GET request.

url_param

If the modifier "check_post" is used, then an HTTP POST request entity will be searched for the parameter argument, when it is not found in a query string after a question mark ('?') in the URL. The message body will only start to be analyzed once either the advertised amount of data has been received or the request buffer is full. In the unlikely event that chunked encoding is used, only the first chunk is scanned. Parameter values separated by a chunk boundary, may be randomly balanced if at all. This keyword used to support an optional <max_wait> parameter which is now ignored.

If the parameter is found followed by an equal sign ('=') and a value, then the value is hashed and divided by the total weight of the running servers. The result designates which server will receive the request.

This is used to track user identifiers in requests and ensure that a same user ID will always be sent to the same server as long as no server goes up or down. If no value is found or if the parameter is not found, then a round robin algorithm is applied. Note that this algorithm may only be used in an HTTP backend. This algorithm is static by default, which means that changing a server's weight on the fly will have no effect, but this can be changed using "hash-type".

hdr(<name>)

The HTTP header <name> will be looked up in each HTTP request. Just as with the equivalent ACL 'hdr()' function, the header name in parenthesis is not case sensitive. If the header is absent or if it does not contain any value, the roundrobin algorithm is applied instead.

An optional 'use_domain_only' parameter is available, for reducing the hash algorithm to the main domain part with some specific headers such as 'Host'. For instance, in the Host value 'haproxy.lwt.eu', only 'lwt' will be considered.

This algorithm is static by default, which means that changing a server's weight on the fly will have no effect, but this can be changed using "hash-type".

rdp-cookie

rdp-cookie<name>

The RDP cookie <name> (or "msthash" if omitted) will be looked up and hashed for each incoming TCP request. Just as with the equivalent ACL 'req_rdp_cookie()' function, the name is not case-sensitive. This mechanism is useful as a degraded persistence mode, as it makes it possible to always send the same user (or the same session ID) to the same server. If the

cookie is not found, the normal roundrobin algorithm is used instead.

Note that for this to work, the frontend must ensure that an RDP cookie is already present in the request buffer. For this you must use 'tcp-request content accept' rule combined with a 'req_rdp_cookie_cnt' ACL.

This algorithm is static by default, which means that changing a server's weight on the fly will have no effect, but this can be changed using "hash-type".

See also the rdp_cookie pattern fetch function.

<arguments> is an optional list of arguments which may be needed by some algorithms. Right now, only "url_param" and "uri" support an optional argument.

The load balancing algorithm of a backend is set to roundrobin when no other algorithm, mode nor option have been set. The algorithm may only be set once for each backend.

Examples :

```
balance roundrobin
balance url_param userid
balance url_param session_id check_post 64
balance hdr(User-Agent)
balance hdr(host)
balance hdr(Host) use_domain_only
```

Note: the following caveats and limitations on using the "check_post" extension with "url_param" must be considered :

- all POST requests are eligible for consideration, because there is no way to determine if the parameters will be found in the body or entity which may contain binary data. Therefore another method may be required to restrict consideration of POST requests that have no URL parameters in the body. (see acl reqidney http_end)
- using a <max.wait> value larger than the request buffer size does not make sense and is useless. The buffer size is set at build time, and defaults to 16 kB.
- Content-Encoding is not supported, the parameter search will probably fail; and load balancing will fall back to Round Robin.
- Expect: 100-continue is not supported, load balancing will fall back to Round Robin.
- Transfer-Encoding (RFC2616 3.6.1) is only supported in the first chunk. If the entire parameter value is not present in the first chunk, the selection of server is undefined (actually, defined by how little actually appeared in the first chunk).
- This feature does not support generation of a 100, 411 or 501 response.
- In some cases, requesting "check_post" MAY attempt to scan the entire contents of a message body. Scanning normally terminates when linear white space or control characters are found, indicating the end of what might be a URL parameter list. This is probably not a concern with SQL type message bodies.

See also : "dispatch", "cookie", "transparent", "hash-type" and "http_proxy".

```
bind [<address>]:<port_range> [, ...] [param*]
```

```
bind /<path> [, ...] [param*]
```

Define one or several listening addresses and/or ports in a frontend.

May be used in sections :

defaults	frontend	listen	backend
no	yes	yes	no

Arguments :

<address>

is optional and can be a host name, an IPv4 address, an IPv6 address, or '*'. It designates the address the frontend will listen on. If unset, all IPv4 addresses of the system will be listened on. The same will apply for '*' or the system's special address '0.0.0.0'. The IPv6 equivalent is '::'.

Optionally, an address family prefix may be used before the address to force the family regardless of the address format, which can be useful to specify a path to a unix socket with no slash ('/'). Currently supported prefixes are :

- 'ipv4@' -> address is always IPv4

- 'ipv6@' -> address is always IPv6

- 'unix@' -> address is a path to a local unix socket

- 'abns@' -> address is in abstract namespace (linux only). Note: since abstract sockets are not "rebindable", they do not cope well with multi-process mode during

soft-restart, so it is better to avoid them if

nbproc is greater than 1. The effect is that if the

new process fails to start, only one of the old ones

will be able to rebind to the socket.

- 'fd@<n>' -> use file descriptor <n> inherited from the

parent. The fd must be bound and may or may not already

be listening.

You may want to reference some environment variables in the

address parameter, see section 2.3 about environment

variables.

<port_range>

is either a unique TCP port, or a port range for which the

proxy will accept connections for the IP address specified

above. The port is mandatory for TCP listeners. Note that in

the case of an IPv6 address, the port is always the number

after the last colon (':'). A range can either be :

- a numerical port (ex: '80')

- a dash-delimited ports range explicitly stating the lower

and upper bounds (ex: '2000-2100') which are included in

the range.

Particular care must be taken against port ranges, because every <address:port> couple consumes one socket (= a file

descriptor), so it's easy to consume lots of descriptors

with a simple range, and to run out of sockets. Also, each

<address:port> couple must be used only once among all

instances running on a same system. Please note that binding

to ports lower than 1024 generally require particular

privileges to start the program, which are independent of

the 'uid' parameter.

<path>

is a UNIX socket path beginning with a slash ('/'). This is

alternative to the TCP listening port. Haproxy will then

receive UNIX connections on the socket located at this place.

The path must begin with a slash and by default is absolute.

It can be relative to the prefix defined by "unix-bind" in

the global section. Note that the total length of the prefix

followed by the socket path cannot exceed some system limits

for UNIX sockets, which commonly are set to 107 characters.

<param*>

is a list of parameters common to all sockets declared on the

same line. These numerous parameters depend on OS and build

options and have a complete section dedicated to them. Please

refer to section 5 to for more details.

It is possible to specify a list of address:port combinations delimited by commas. The frontend will then listen on all of these addresses. There is no fixed limit to the number of addresses and ports which can be listened on in a frontend, as well as there is no limit to the number of "bind" statements in a frontend.

Example :

```
listen http_proxy
bind :80,:443
bind 10.0.0.1:10080,10.0.0.1:10443
bind /var/run/ssl-frontend.sock user root mode 600 accept-proxy

listen http_https_proxy
bind :80
bind :443 ssl crt /etc/haproxy/site.pem

listen http_https_proxy_explicit
bind ipv6::80
bind ipv4@public_ssl:443 ssl crt /etc/haproxy/site.pem
bind unix@ssl-frontend.sock user root mode 600 accept-proxy

listen external_bind_app1
bind "fd@${FD_APP1}"
```

Note: regarding Linux's abstract namespace sockets, HAProxy uses the whole `sun_path` length is used for the address length. Some other programs such as `socket` use the string length only by default. Pass the option `"unix-tightsocketlen=0"` to any abstract socket definition in `socket` to make it compatible with HAProxy's.

See also : "source", "option forwardfor", "unix-bind" and the PROXY protocol documentation, and section 5 about bind options.

`bind-process [all | odd | even | <number 1-64> [-<number 1-64>] ...`
Limit visibility of an instance to a certain set of processes numbers.

May be used in sections : defaults | frontend | listen | backend
yes | yes | yes | yes

Arguments :

all All process will see this instance. This is the default. It may be used to override a default value.

odd This instance will be enabled on processes 1,3,5,...63. This option may be combined with other numbers.

even This instance will be enabled on processes 2,4,6,...64. This option may be combined with other numbers. Do not use it with less than 2 processes otherwise some instances might be missing from all processes.

number The instance will be enabled on this process number or range, whose values must all be between 1 and 32 or 64 depending on the machine's word size. If a proxy is bound to process numbers greater than the configured `global.nbproc`, it will either be forced to process #1 if a single process was specified, or to all processes otherwise.

This keyword limits binding of certain instances to certain processes. This is useful in order not to have too many processes listening to the same ports. For instance, on a dual-core machine, it might make sense to set 'nbproc 2' in the global section, then distributes the listeners among 'odd' and 'even' instances.

At the moment, it is not possible to reference more than 32 or 64 processes

using this keyword, but this should be more than enough for most setups. Please note that 'all' really means all processes regardless of the machine's word size, and is not limited to the first 32 or 64.

Each "bind" line may further be limited to a subset of the proxy's processes, please consult the "process" bind keyword in section 5.1.

When a frontend has no explicit "bind-process" line, it tries to bind to all the processes referenced by its "bind" lines. That means that frontends can easily adapt to their listeners' processes.

If some backends are referenced by frontends bound to other processes, the backend automatically inherits the frontend's processes.

Example :

```
listen app_ip1
bind 10.0.0.1:80
bind-process odd

listen app_ip2
bind 10.0.0.2:80
bind-process even

listen management
bind 10.0.0.3:80
bind-process 1 2 3 4

listen management
bind 10.0.0.4:80
bind-process 1-4
```

See also : "nbproc" in global section, and "process" in section 5.1.

`block { if | unless } <condition>`

Block a layer 7 request if/unless a condition is matched

May be used in sections : defaults | frontend | listen | backend
no | yes | yes | yes

The HTTP request will be blocked very early in the layer 7 processing if/unless `<condition>` is matched. A 403 error will be returned if the request is blocked. The condition has to reference ACLs (see section 7). This is typically used to deny access to certain sensitive resources if some conditions are met or not met. There is no fixed limit to the number of "block" statements per instance.

Example:

```
acl invalid_src src 0.0.0.0/7 224.0.0.0/3
acl invalid_src src_port 0:1023
acl local_dst hdr(host) -i localhost
block if invalid_src || local_dst
```

See section 7 about ACL usage.

`capture cookie <name> len <length>`

Capture and log a cookie in the request and in the response.

May be used in sections : defaults | frontend | listen | backend
no | yes | yes | no

Arguments :

<name> is the beginning of the name of the cookie to capture. In order to match the exact name, simply suffix the name with an equal sign ('='). The full name will appear in the logs, which is useful with application servers which adjust both the cookie name and value (eg: `ASPSESSIONXXXX`).

2341 `<length>` is the maximum number of characters to report in the logs, which
2342 include the cookie name, the equal sign and the value, all in the
2343 standard "name=value" form. The string will be truncated on the
2344 right if it exceeds `<length>`.
2345

2346 Only the first cookie is captured. Both the "cookie" request headers and the
2347 "set-cookie" response headers are monitored. This is particularly useful to
2348 check for application bugs causing session crossing or stealing between
2349 users, because generally the user's cookies can only change on a login page.

2351 When the cookie was not presented by the client, the associated log column
2352 will report "-". When a request does not cause a cookie to be assigned by the
2353 server, a "-" is reported in the response column.
2354

2355 The capture is performed in the frontend only because it is necessary that
2356 the log format does not change for a given frontend depending on the
2357 backends. This may change in the future. Note that there can be only one
2358 "capture cookie" statement in a frontend. The maximum capture length is set
2359 by the global "tune.http.cookie_len" setting and defaults to 63 characters. It
2360 is not possible to specify a capture in a "defaults" section.
2361

2362 Example:
2363 capture cookie ASPSESSION len 32
2364

2365 See also : "capture request header", "capture response header" as well as
2366 section 8 about logging.
2367

2368 capture request header <name> len <length>

2370 Capture and log the last occurrence of the specified request header.

2371 May be used in sections : defaults | frontend | listen | backend
2372 no | yes | yes | no
2373

2374 Arguments :
2375 <name> is the name of the header to capture. The header names are not
2376 case-sensitive, but it is a common practice to write them as they
2377 appear in the requests, with the first letter of each word in
2378 upper case. The header name will not appear in the logs, only the
2379 value is reported, but the position in the logs is respected.
2380

2381 <length> is the maximum number of characters to extract from the value and
2382 report in the logs. The string will be truncated on the right if
2383 it exceeds `<length>`.
2384

2385 The complete value of the last occurrence of the header is captured. The
2386 value will be added to the logs between braces ({}). If multiple headers
2387 are captured, they will be delimited by a vertical bar (|) and will appear
2388 in the same order they were declared in the configuration. Non-existent
2389 headers will be logged just as an empty string. Common uses for request
2390 header captures include the "Host" field in virtual hosting environments, the
2391 "Content-length" when uploads are supported, "User-agent" to quickly
2392 differentiate between real users and robots, and "X-Forwarded-For" in proxied
2393 environments to find where the request came from.
2394

2395 Note that when capturing headers such as "User-agent", some spaces may be
2396 logged, making the log analysis more difficult. Thus be careful about what
2397 you log if you know your log parser is not smart enough to rely on the
2398 braces.
2399

2400 There is no limit to the number of captured request headers nor to their
2401 length, though it is wise to keep them low to limit memory usage per session.
2402 In order to keep log format consistent for a same frontend, header captures
2403 can only be declared in a frontend. It is not possible to specify a capture
2404 in a "defaults" section.
2405

2406 Example:

2407 capture request header Host len 15
2408 capture request header X-Forwarded-For len 15
2409 capture request header Referer len 15
2410

2411 See also : "capture cookie", "capture response header" as well as section 8
2412 about logging.
2413

2414 capture response header <name> len <length>

2415 Capture and log the last occurrence of the specified response header.

2416 May be used in sections : defaults | frontend | listen | backend
2417 no | yes | yes | no
2418

2419 Arguments :
2420 <name> is the name of the header to capture. The header names are not
2421 case-sensitive, but it is a common practice to write them as they
2422 appear in the response, with the first letter of each word in
2423 upper case. The header name will not appear in the logs, only the
2424 value is reported, but the position in the logs is respected.
2425

2426 <length> is the maximum number of characters to extract from the value and
2427 report in the logs. The string will be truncated on the right if
2428 it exceeds `<length>`.
2429

2430 The complete value of the last occurrence of the header is captured. The
2431 result will be added to the logs between braces ({}). If multiple headers
2432 are captured, they will be delimited by a vertical bar (|) and will appear
2433 in the same order they were declared in the configuration. Non-existent
2434 headers will be logged just as an empty string. Common uses for response
2435 header captures include the "Content-length" header which indicates how many
2436 bytes are expected to be returned, the "Location" header to track redirections.
2437

2438 There is no limit to the number of captured response headers nor to their
2439 length, though it is wise to keep them low to limit memory usage per session.
2440 In order to keep log format consistent for a same frontend, header captures
2441 can only be declared in a frontend. It is not possible to specify a capture
2442 in a "defaults" section.
2443

2444 Example:

2445 capture response header Content-length len 9
2446 capture response header Location len 15
2447

2448 See also : "capture cookie", "capture request header" as well as section 8
2449 about logging.
2450

2451 cliptimeout <timeout> (deprecated)

2452 Set the maximum inactivity time on the client side.

2453 May be used in sections : defaults | frontend | listen | backend
2454 yes | yes | yes | no
2455

2456 Arguments :
2457 <timeout> is the timeout value is specified in milliseconds by default, but
2458 can be in any other unit if the number is suffixed by the unit,
2459 as explained at the top of this document.
2460

2461 The inactivity timeout applies when the client is expected to acknowledge or
2462 send data. In HTTP mode, this timeout is particularly important to consider
2463 during the first phase, when the client sends the request, and during the
2464 response while it is reading data sent by the server. The value is specified
2465 in milliseconds by default, but can be in any other unit if the number is
2466 suffixed by the unit, as specified at the top of this document. In TCP mode
2467 (and to a lesser extent, in HTTP mode), it is highly recommended that the
2468 client timeout remains equal to the server timeout in order to avoid complex
2469 situations to debug. It is a good practice to cover one or several TCP packet
2470

losses by specifying timeouts that are slightly above multiples of 3 seconds (eg: 4 or 5 seconds).

This parameter is specific to frontends, but can be specified once for all in "defaults" sections. This is in fact one of the easiest solutions not to forget about it. An unspecified timeout results in an infinite timeout, which is not recommended. Such a usage is accepted and works but reports a warning during startup because it may results in accumulation of expired sessions in the system if the system's timeouts are not configured either.

This parameter is provided for compatibility but is currently deprecated. Please use "timeout client" instead.

See also : "timeout client", "timeout http-request", "timeout server", and "srvtimeout".

compression algo <algorithm> ...
compression type <mime type> ...
compression offload

Enable HTTP compression.

May be used in sections : defaults | frontend | listen | backend
yes | yes | yes | yes

Arguments :
algo is followed by the list of supported compression algorithms.
type is followed by the list of MIME types that will be compressed.
offload makes haproxy work as a compression offloader only (see notes).

The currently supported algorithms are :

identity this is mostly for debugging, and it was useful for developing the compression feature. Identity does not apply any change on data.

gzip applies gzip compression. This setting is only available when support for zlib was built in.

deflate same as "gzip", but with deflate algorithm and zlib format. Note that this algorithm has ambiguous support on many browsers and no support at all from recent ones. It is strongly recommended not to use it for anything else than experimentation. This setting is only available when support for zlib was built in.

raw-deflate same as "deflate" without the zlib wrapper, and used as an alternative when the browser wants "deflate". All major browsers understand it and despite violating the standards, it is known to work better than "deflate", at least on MSIE and some versions of Safari. Do not use it in conjunction with "deflate", use either one or the other since both react to the same Accept-Encoding token. This setting is only available when support for zlib was built in.

Compression will be activated depending on the Accept-Encoding request header. With identity, it does not take care of that header. If backend servers support HTTP compression, these directives will be no-op: haproxy will see the compressed response and will not compress again. If backend servers do not support HTTP compression and there is Accept-Encoding header in request, haproxy will compress the matching response.

The "offload" setting makes haproxy remove the Accept-Encoding header to prevent backend servers from compressing responses. It is strongly recommended not to do this because this means that all the compression work will be done on the single point where haproxy is located. However in some deployment scenarios, haproxy may be installed in front of a buggy gateway with broken HTTP compression implementation which can't be turned off.

In that case haproxy can be used to prevent that gateway from emitting invalid payloads. In this case, simply removing the header in the configuration does not work because it applies before the header is parsed, so that prevents haproxy from compressing. The "offload" setting should then be used for such scenarios. Note: for now, the "offload" setting is ignored when set in a defaults section.

Compression is disabled when:

- * the request does not advertise a supported compression algorithm in the "Accept-Encoding" header
- * the response message is not HTTP/1.1
- * HTTP status code is not 200
- * response header "Transfer-Encoding" contains "chunked" (Temporary Workaround)

* response contain neither a "Content-Length" header nor a

"Transfer-Encoding" whose last value is "chunked"

* response contains a "Content-Type" header whose first value starts with

"multipart"

* the response contains the "no-transform" value in the "Cache-control"

header

- * User-Agent matches "Mozilla/4" unless it is MSIE 6 with XP SP2, or MSIE 7 and later

* The response contains a "Content-Encoding" header, indicating that the response is already compressed (see compression offload)

Note: The compression does not rewrite Etag headers, and does not emit the Warning header.

Examples :

compression algo gzip
compression type text/html text/plain

contimeout <timeout> (deprecated)

Set the maximum time to wait for a connection attempt to a server to succeed.

May be used in sections : defaults | frontend | listen | backend
yes | no | yes | yes

Arguments :

<timeout> is the timeout value is specified in milliseconds by default, but can be in any other unit if the number is suffixed by the unit, as explained at the top of this document.

If the server is located on the same LAN as haproxy, the connection should be immediate (less than a few milliseconds). Anyway, it is a good practice to cover one or several TCP packet losses by specifying timeouts that are slightly above multiples of 3 seconds (eg: 4 or 5 seconds). By default, the connect timeout also presets the queue timeout to the same value if this one has not been specified. Historically, the contimeout was also used to set the tarpit timeout in a listen section, which is not possible in a pure frontend.

This parameter is specific to backends, but can be specified once for all in "defaults" sections. This is in fact one of the easiest solutions not to forget about it. An unspecified timeout results in an infinite timeout, which is not recommended. Such a usage is accepted and works but reports a warning during startup because it may results in accumulation of failed sessions in the system if the system's timeouts are not configured either.

This parameter is provided for backwards compatibility but is currently deprecated. Please use "timeout connect", "timeout queue" or "timeout tarpit" instead.

See also : "timeout connect", "timeout queue", "timeout tarpit",
"timeout server", "contimeout".

cookie <name> [rewrite | insert | prefix] [indirect] [nocache]

```

2601 [ postonly ] [ preserve ] [ httponly ] [ secure ]
2602 [ domain <domain> ] * [ maxidle <idle> ] [ maxlife <life> ]
2603 Enable cookie-based persistence in a backend.
2604 May be used in sections : defaults | frontend | listen | backend
2605 yes | no | yes | yes
2606 Arguments :
2607 <name> is the name of the cookie which will be monitored, modified or
2608 inserted in order to bring persistence. This cookie is sent to
2609 the client via a "Set-Cookie" header in the response, and is
2610 brought back by the client in a "Cookie" header in all requests.
2611 Special care should be taken to choose a name which does not
2612 conflict with any likely application cookie. Also, if the same
2613 backends are subject to be used by the same clients (eg:
2614 HTTP/HTTPS), care should be taken to use different cookie names
2615 between all backends if persistence between them is not desired.
2616
2617 rewrite This keyword indicates that the cookie will be provided by the
2618 server and that haproxy will have to modify its value to set the
2619 server's identifier in it. This mode is handy when the management
2620 of complex combinations of "Set-cookie" and "Cache-control"
2621 headers is left to the application. The application can then
2622 decide whether or not it is appropriate to emit a persistence
2623 cookie. Since all responses should be monitored, this mode only
2624 works in HTTP close mode. Unless the application behaviour is
2625 very complex and/or broken, it is advised not to start with this
2626 mode for new deployments. This keyword is incompatible with
2627 "insert" and "prefix".
2628
2629 insert This keyword indicates that the persistence cookie will have to
2630 be inserted by haproxy in server responses if the client did not
2631 already have a cookie that would have permitted it to access this
2632 server. When used without the "preserve" option, if the server
2633 emits a cookie with the same name, it will be remove before
2634 processing. For this reason, this mode can be used to upgrade
2635 existing configurations running in the "rewrite" mode. The cookie
2636 will only be a session cookie and will not be stored on the
2637 client's disk. By default, unless the "indirect" option is added,
2638 the server will see the cookies emitted by the client. Due to
2639 caching effects, it is generally wise to add the "nocache" or
2640 "postonly" keywords (see below). The "insert" keyword is not
2641 compatible with "rewrite" and "prefix".
2642
2643 prefix This keyword indicates that instead of relying on a dedicated
2644 cookie for the persistence, an existing one will be completed.
2645 This may be needed in some specific environments where the client
2646 does not support more than one single cookie and the application
2647 already needs it. In this case, whenever the server sets a cookie
2648 named <name>, it will be prefixed with the server's identifier
2649 and a delimiter. The prefix will be removed from all client
2650 requests so that the server still finds the cookie it emitted.
2651 Since all requests and responses are subject to being modified,
2652 this mode requires the HTTP close mode. The "prefix" keyword is
2653 not compatible with "rewrite" and "insert". Note: it is highly
2654 recommended not to use "indirect" with "prefix", otherwise server
2655 cookie updates would not be sent to clients.
2656
2657 indirect When this option is specified, no cookie will be emitted to a
2658 client which already has a valid one for the server which has
2659 processed the request. If the server sets such a cookie itself,
2660 it will be removed, unless the "preserve" option is also set. In
2661 "insert" mode, this will additionally remove cookies from the
2662 requests transmitted to the server, making the persistence
2663 mechanism totally transparent from an application point of view.
2664 Note: it is highly recommended not to use "indirect" with
2665

```

```

2666 "prefix", otherwise server cookie updates would not be sent to
2667 clients.
2668
2669 nocache This option is recommended in conjunction with the insert mode
2670 when there is a cache between the client and HAProxy, as it
2671 ensures that a cacheable response will be tagged non-cacheable if
2672 a cookie needs to be inserted. This is important because if all
2673 persistence cookies are added on a cacheable home page for
2674 instance, then all customers will then fetch the page from an
2675 outer cache and will all share the same persistence cookie,
2676 leading to one server receiving much more traffic than others.
2677 See also the "insert" and "postonly" options.
2678
2679 postonly This option ensures that cookie insertion will only be performed
2680 on responses to POST requests. It is an alternative to the
2681 "nocache" option, because POST responses are not cacheable, so
2682 this ensures that the persistence cookie will never get cached.
2683 Since most sites do not need any sort of persistence before the
2684 first POST which generally is a login request, this is a very
2685 efficient method to optimize caching without risking to find a
2686 persistence cookie in the cache.
2687 See also the "insert" and "nocache" options.
2688
2689 preserve This option may only be used with "insert" and/or "indirect". It
2690 allows the server to emit the persistence cookie itself. In this
2691 case, if a cookie is found in the response, haproxy will leave it
2692 untouched. This is useful in order to end persistence after a
2693 logout request for instance. For this, the server just has to
2694 emit a cookie with an invalid value (eg: empty) or with a date in
2695 the past. By combining this mechanism with the "disable-on-404"
2696 check option, it is possible to perform a completely graceful
2697 shutdown because users will definitely leave the server after
2698 they logout.
2699
2700 httponly This option tells haproxy to add an "HttpOnly" cookie attribute
2701 when a cookie is inserted. This attribute is used so that a
2702 user agent doesn't share the cookie with non-HTTP components.
2703 Please check RFC6265 for more information on this attribute.
2704
2705 secure This option tells haproxy to add a "Secure" cookie attribute when
2706 a cookie is inserted. This attribute is used so that a user agent
2707 never emits this cookie over non-secure channels, which means
2708 that a cookie learned with this flag will be presented only over
2709 SSL/TLS connections. Please check RFC6265 for more information on
2710 this attribute.
2711
2712 domain This option allows to specify the domain at which a cookie is
2713 inserted. It requires exactly one parameter: a valid domain
2714 name. If the domain begins with a dot, the browser is allowed to
2715 use it for any host ending with that name. It is also possible to
2716 specify several domain names by invoking this option multiple
2717 times. Some browsers might have small limits on the number of
2718 domains, so be careful when doing that. For the record, sending
2719 10 domains to MSIE 6 or Firefox 2 works as expected.
2720
2721 maxidle This option allows inserted cookies to be ignored after some idle
2722 time. It only works with insert-mode cookies. When a cookie is
2723 sent to the client, the date this cookie was emitted is sent too.
2724 Upon further presentations of this cookie, if the date is older
2725 than the delay indicated by the parameter (in seconds), it will
2726 be ignored. Otherwise, it will be refreshed if needed when the
2727 response is sent to the client. This is particularly useful to
2728 prevent users who never close their browsers from remaining for
2729 too long on the same server (eg: after a farm size change). When
2730 this option is set and a cookie has no date, it is always
2731

```

accepted, but gets refreshed in the response. This maintains the ability for admins to access their sites. Cookies that have a date in the future further than 24 hours are ignored. Doing so lets admins fix timezone issues without risking kicking users off the site.

maxlife This option allows inserted cookies to be ignored after some life time, whether they're in use or not. It only works with insert mode cookies. When a cookie is first sent to the client, the date this cookie was emitted is sent too. Upon further presentations of this cookie, if the date is older than the delay indicated by the parameter (in seconds), it will be ignored. If the cookie in the request has no date, it is accepted and a date will be set. Cookies that have a date in the future further than 24 hours are ignored. Doing so lets admins fix timezone issues without risking kicking users off the site. Contrary to maxidle, this value is not refreshed, only the first visit date counts. Both maxidle and maxlife may be used at the time. This is particularly useful to prevent users who never close their browsers from remaining for too long on the same server (eg: after a farm size change). This is stronger than the maxidle method in that it forces a redispatch after some absolute delay.

There can be only one persistence cookie per HTTP backend, and it can be declared in a defaults section. The value of the cookie will be the value indicated after the "cookie" keyword in a "server" statement. If no cookie is declared for a given server, the cookie is not set.

Examples :

```
cookie JSESSIONID prefix
cookie SRV insert indirect nocache
cookie SRV insert postonly indirect
cookie SRV insert indirect nocache maxidle 30m maxlife 8h
```

See also : "balance source", "capture cookie", "server" and "ignore-persist".

declare capture [request | response] len <length>

Declares a capture slot.

May be used in sections : defaults | frontend | listen | backend
no | yes | yes | no

Arguments:

<length> is the length allowed for the capture.

This declaration is only available in the frontend or listen section, but the reserved slot can be used in the backends. The "request" keyword allocates a capture slot for use in the request, and "response" allocates a capture slot for use in the response.

See also: "capture-req", "capture-res" (sample converters),
"capture.req.hdr", "capture.res.hdr" (sample fetches),
"http-request capture" and "http-response capture".

default-server [param*]

Change default options for a server in a backend

May be used in sections : defaults | frontend | listen | backend
yes | no | yes | yes

Arguments:

<param*> is a list of parameters for this server. The "default-server" keyword accepts an important number of options and has a complete section dedicated to it. Please refer to section 5 for more details.

Example :

default-server inter 1000 weight 13

See also: "server" and section 5 about server options

default_backend <backend>

Specify the backend to use when no "use_backend" rule has been matched. May be used in sections : defaults | frontend | listen | backend
yes | yes | yes | no

Arguments :

<backend> is the name of the backend to use.

When doing content-switching between frontend and backends using the "use_backend" keyword, it is often useful to indicate which backend will be used when no rule has matched. It generally is the dynamic backend which will catch all undetermined requests.

Example :

```
use_backend dynamic if url_dyn
use_backend static if url_css url_img extension_img
default_backend dynamic
```

See also : "use_backend"

description <string>

Describe a listen, frontend or backend.

May be used in sections : defaults | frontend | listen | backend
no | yes | yes | yes

Arguments : string

Allows to add a sentence to describe the related object in the HAProxy HTML stats page. The description will be printed on the right of the object name it describes.

No need to backslash spaces in the <string> arguments.

disabled

Disable a proxy, frontend or backend.

May be used in sections : defaults | frontend | listen | backend
yes | yes | yes | yes

Arguments : none

The "disabled" keyword is used to disable an instance, mainly in order to liberate a listening port or to temporarily disable a service. The instance will still be created and its configuration will be checked, but it will be created in the "stopped" state and will appear as such in the statistics. It will not receive any traffic nor will it send any health-checks or logs. It is possible to disable many instances at once by adding the "disabled" keyword in a "defaults" section.

See also : "enabled"

dispatch <address>:<port>

Set a default server address

May be used in sections : defaults | frontend | listen | backend
no | no | yes | yes

Arguments :

<address> is the IPv4 address of the default server. Alternatively, a resolvable hostname is supported, but this name will be resolved during start-up.

<ports> is a mandatory port specification. All connections will be sent to this port, and it is not permitted to use port offsets as is possible with normal servers.

The "dispatch" keyword designates a default server for use when no other server can take the connection. In the past it was used to forward non persistent connections to an auxiliary load balancer. Due to its simple syntax, it has also been used for simple TCP relays. It is recommended not to use it for more clarity, and to use the "server" directive instead.

See also : "server"

enabled

Enable a proxy, frontend or backend.

May be used in sections : defaults | frontend | listen | backend
yes | yes | yes | yes

Arguments : none

The "enabled" keyword is used to explicitly enable an instance, when the defaults has been set to "disabled". This is very rarely used.

See also : "disabled"

errorfile <code> <file>

Return a file contents instead of errors generated by HAProxy

May be used in sections : defaults | frontend | listen | backend
yes | yes | yes | yes

Arguments :

<code> is the HTTP status code. Currently, HAProxy is capable of generating codes 200, 400, 403, 405, 408, 429, 500, 502, 503, and 504.

<file> designates a file containing the full HTTP response. It is recommended to follow the common practice of appending ".http" to the filename so that people do not confuse the response with HTML error pages, and to use absolute paths, since files are read before any chroot is performed.

It is important to understand that this keyword is not meant to rewrite errors returned by the server, but errors detected and returned by HAProxy. This is why the list of supported errors is limited to a small set.

Code 200 is emitted in response to requests matching a "monitor-uri" rule.

The files are returned verbatim on the TCP socket. This allows any trick such as redirections to another URL or site, as well as tricks to clean cookies, force enable or disable caching, etc... The package provides default error files returning the same contents as default errors.

The files should not exceed the configured buffer size (BUFSIZE), which generally is 8 or 16 kB, otherwise they will be truncated. It is also wise not to put any reference to local contents (eg: images) in order to avoid loops between the client and HAProxy when all servers are down, causing an error to be returned instead of an image. For better HTTP compliance, it is recommended that all header lines end with CR-LF and not LF alone.

The files are read at the same time as the configuration and kept in memory. For this reason, the errors continue to be returned even when the process is chrooted, and no file change is considered while the process is running. A simple method for developing those files consists in associating them to the 403 status code and interrogating a blocked URL.

See also : "errorloc", "errorloc302", "errorloc303"

Example :

```
errorfile 400 /etc/haproxy/errorfiles/400badreq.http
errorfile 408 /dev/null # workaround Chrome pre-connect bug
errorfile 403 /etc/haproxy/errorfiles/403forbid.http
errorfile 503 /etc/haproxy/errorfiles/503sorry.http
```

errorloc <code> <url>

errorloc302 <code> <url>

Return an HTTP redirection to a URL instead of errors generated by HAProxy
May be used in sections : defaults | frontend | listen | backend
yes | yes | yes | yes

Arguments :

<code> is the HTTP status code. Currently, HAProxy is capable of generating codes 200, 400, 403, 408, 500, 502, 503, and 504.

<url> it is the exact contents of the "Location" header. It may contain either a relative URI to an error page hosted on the same site, or an absolute URI designating an error page on another site. Special care should be given to relative URIs to avoid redirect loops if the URI itself may generate the same error (eg: 500).

It is important to understand that this keyword is not meant to rewrite errors returned by the server, but errors detected and returned by HAProxy. This is why the list of supported errors is limited to a small set.

Code 200 is emitted in response to requests matching a "monitor-uri" rule.

Note that both keyword return the HTTP 302 status code, which tells the client to fetch the designated URL using the same HTTP method. This can be quite problematic in case of non-GET methods such as POST, because the URL sent to the client might not be allowed for something other than GET. To workaround this problem, please use "errorloc303" which send the HTTP 303 status code, indicating to the client that the URL must be fetched with a GET request.

See also : "errorfile", "errorloc303"

errorloc303 <code> <url>

Return an HTTP redirection to a URL instead of errors generated by HAProxy
May be used in sections : defaults | frontend | listen | backend
yes | yes | yes | yes

Arguments :
<code> is the HTTP status code. Currently, HAProxy is capable of generating codes 400, 403, 408, 500, 502, 503, and 504.

<url> it is the exact contents of the "Location" header. It may contain either a relative URI to an error page hosted on the same site, or an absolute URI designating an error page on another site. Special care should be given to relative URIs to avoid redirect loops if the URI itself may generate the same error (eg: 500).

It is important to understand that this keyword is not meant to rewrite errors returned by the server, but errors detected and returned by HAProxy. This is why the list of supported errors is limited to a small set.

Code 200 is emitted in response to requests matching a "monitor-uri" rule.

Note that both keyword return the HTTP 303 status code, which tells the client to fetch the designated URL using the same HTTP GET method. This solves the usual problems associated with "errorloc" and the 302 code. It is possible that some very old browsers designed before HTTP/1.1 do not support it, but no such problem has been reported till now.

2991
2992
2993
2994
2995
2996
2997
2998
2999
3000
3001
3002
3003
3004
3005
3006
3007
3008
3009
3010
3011
3012
3013
3014
3015
3016
3017
3018
3019
3020
3021
3022
3023
3024
3025
3026
3027
3028
3029
3030
3031
3032
3033
3034
3035
3036
3037
3038
3039
3040
3041
3042
3043
3044
3045
3046
3047
3048
3049
3050
3051
3052
3053
3054
3055

See also : "errorfile", "errorloc", "errorloc302"

email-alert from <emailaddr>
Declare the from email address to be used in both the envelope and header of email alerts. This is the address that email alerts are sent from.
May be used in sections: defaults | frontend | listen | backend
yes | yes | yes | yes

Arguments :

<emailaddr> is the from email address to use when sending email alerts

Also requires "email-alert mailers" and "email-alert to" to be set and if so sending email alerts is enabled for the proxy.

See also : "email-alert level", "email-alert mailers",
"email-alert myhostname", "email-alert to", section 3.6 about mailers.

email-alert level <level>
Declare the maximum log level of messages for which email alerts will be sent. This acts as a filter on the sending of email alerts.
May be used in sections: defaults | frontend | listen | backend
yes | yes | yes | yes

Arguments :

<level> One of the 8 syslog levels:
emerg alert crit err warning notice info debug
The above syslog levels are ordered from lowest to highest.

By default level is alert

Also requires "email-alert from", "email-alert mailers" and "email-alert to" to be set and if so sending email alerts is enabled for the proxy.

Alerts are sent when :

* An un-paused server is marked as down and <level> is alert or lower
* A paused server is marked as down and <level> is notice or lower
* A server is marked as up or enters the drain state and <level> is notice or lower
* "option log-health-checks" is enabled, <level> is info or lower, and a health check status update occurs

See also : "email-alert from", "email-alert mailers",
"email-alert myhostname", "email-alert to",
section 3.6 about mailers.

email-alert mailers <mailersect>
Declare the mailers to be used when sending email alerts
May be used in sections: defaults | frontend | listen | backend
yes | yes | yes | yes

Arguments :

<mailersect> is the name of the mailers section to send email alerts.

Also requires "email-alert from" and "email-alert to" to be set and if so sending email alerts is enabled for the proxy.

3056
3057
3058
3059
3060
3061
3062
3063
3064
3065
3066
3067
3068
3069
3070
3071
3072
3073
3074
3075
3076
3077
3078
3079
3080
3081
3082
3083
3084
3085
3086
3087
3088
3089
3090
3091
3092
3093
3094
3095
3096
3097
3098
3099
3100
3101
3102
3103
3104
3105
3106
3107
3108
3109
3110
3111
3112
3113
3114
3115
3116
3117
3118
3119
3120

See also : "email-alert from", "email-alert level", "email-alert myhostname",
"email-alert to", section 3.6 about mailers.

email-alert myhostname <hostname>
Declare the to hostname address to be used when communicating with mailers.
May be used in sections: defaults | frontend | listen | backend
yes | yes | yes | yes

Arguments :

<emailaddr> is the to email address to use when sending email alerts

By default the systems hostname is used.

Also requires "email-alert from", "email-alert mailers" and "email-alert to" to be set and if so sending email alerts is enabled for the proxy.

See also : "email-alert from", "email-alert level", "email-alert mailers",
"email-alert to", section 3.6 about mailers.

email-alert to <emailaddr>
Declare both the recipient address in the envelope and to address in the header of email alerts. This is the address that email alerts are sent to.
May be used in sections: defaults | frontend | listen | backend
yes | yes | yes | yes

Arguments :

<emailaddr> is the to email address to use when sending email alerts

Also requires "email-alert mailers" and "email-alert to" to be set and if so sending email alerts is enabled for the proxy.

See also : "email-alert from", "email-alert level", "email-alert mailers",
"email-alert myhostname", section 3.6 about mailers.

force-persist { if | unless } <condition>
Declare a condition to force persistence on down servers
May be used in sections: defaults | frontend | listen | backend
no | yes | yes | yes

By default, requests are not dispatched to down servers. It is possible to force this using "option persist", but it is unconditional and redispatches to a valid server if "option redispatch" is set. That leaves with very little possibilities to force some requests to reach a server which is artificially marked down for maintenance operations.

The "force-persist" statement allows one to declare various ACL-based conditions which, when met, will cause a request to ignore the down status of a server and still try to connect to it. That makes it possible to start a server, still replying an error to the health checks, and run a specially configured browser to test the service. Among the handy methods, one could use a specific source IP address, or a specific cookie. The cookie also has the advantage that it can easily be added/removed on the browser from a test page. Once the service is validated, it is then possible to open the service to the world by returning a valid response to health checks.

The forced persistence is enabled when an "if" condition is met, or unless an "unless" condition is met. The final redispatch is always disabled when this


```

3121 is used.
3122
3123 See also : "option redispatch", "ignore-persist", "persist",
3124 and section 7 about ACL usage.
3125
3126 fullconn <conns>
3127 Specify at what backend load the servers will reach their maxconn
3128 May be used in sections : defaults | frontend | listen | backend
3129 yes | no | yes | yes
3130 Arguments :
3131 <conns> is the number of connections on the backend which will make the
3132 servers use the maximal number of connections.
3133
3134 When a server has a "maxconn" parameter specified, it means that its number
3135 of concurrent connections will never go higher. Additionally, if it has a
3136 "minconn" parameter, it indicates a dynamic limit following the backend's
3137 load. The server will then always accept at least <minconn> connections,
3138 never more than <maxconn>, and the limit will be on the ramp between both
3139 values when the backend has less than <conns> concurrent connections. This
3140 makes it possible to limit the load on the servers during normal loads, but
3141 push it further for important loads without overloading the servers during
3142 exceptional loads.
3143
3144 Since it's hard to get this value right, haproxy automatically sets it to
3145 10% of the sum of the maxconns of all frontends that may branch to this
3146 backend (based on "use_backend" and "default_backend" rules). That way it's
3147 safe to leave it unset. However, "use_backend" involving dynamic names are
3148 not counted since there is no way to know if they could match or not.
3149
3150 Example :
3151 # The servers will accept between 100 and 1000 concurrent connections each
3152 # and the maximum of 1000 will be reached when the backend reaches 10000
3153 # connections.
3154 backend dynamic
3155     fullconn 10000
3156     server sv1 dyn1:80 minconn 100 maxconn 1000
3157     server sv2 dyn2:80 minconn 100 maxconn 1000
3158
3159 See also : "maxconn", "server"
3160
3161 grace <time>
3162 Maintain a proxy operational for some time after a soft stop
3163 May be used in sections : defaults | frontend | listen | backend
3164 yes | yes | yes | yes
3165
3166 Arguments :
3167 <time> is the time (by default in milliseconds) for which the instance
3168 will remain operational with the frontend sockets still listening
3169 when a soft-stop is received via the SIGUSR1 signal.
3170
3171 This may be used to ensure that the services disappear in a certain order.
3172 This was designed so that frontends which are dedicated to monitoring by an
3173 external equipment fail immediately while other ones remain up for the time
3174 needed by the equipment to detect the failure.
3175
3176 Note that currently, there is very little benefit in using this parameter,
3177 and it may in fact complicate the soft-reconfiguration process more than
3178 simplify it.
3179
3180 hash-type <method> <function> <modifier>
3181 Specify a method to use for mapping hashes to servers
3182 May be used in sections : defaults | frontend | listen | backend
3183 yes | no | yes | yes
3184
3185

```

```

3186 Arguments :
3187 <method> is the method used to select a server from the hash computed by
3188 the <function> :
3189
3190 map-based the hash table is a static array containing all alive servers.
3191 The hashes will be very smooth, will consider weights, but
3192 will be static in that weight changes while a server is up
3193 will be ignored. This means that there will be no slow start.
3194 Also, since a server is selected by its position in the array,
3195 most mappings are changed when the server count changes. This
3196 means that when a server goes up or down, or when a server is
3197 added to a farm, most connections will be redistributed to
3198 different servers. This can be inconvenient with caches for
3199 instance.
3200
3201 consistent the hash table is a tree filled with many occurrences of each
3202 server. The hash key is looked up in the tree and the closest
3203 server is chosen. This hash is dynamic, it supports changing
3204 weights while the servers are up, so it is compatible with the
3205 slow start feature. It has the advantage that when a server
3206 goes up or down, only its associations are moved. When a
3207 server is added to the farm, only a few part of the mappings
3208 are redistributed, making it an ideal method for caches.
3209 However, due to its principle, the distribution will never be
3210 very smooth and it may sometimes be necessary to adjust a
3211 server's weight or its ID to get a more balanced distribution.
3212 In order to get the same distribution on multiple load
3213 balancers, it is important that all servers have the exact
3214 same IDs. Note: consistent hash uses sdbm and avalanche if no
3215 hash function is specified.
3216
3217 <function> is the hash function to be used :
3218
3219 sdbm this function was created initially for sdbm (a public-domain
3220 reimplementation of ndbm) database library. It was found to do
3221 well in scrambling bits, causing better distribution of the keys
3222 and fewer splits. It also happens to be a good general hashing
3223 function with good distribution, unless the total server weight
3224 is a multiple of 64, in which case applying the avalanche
3225 modifier may help.
3226
3227 djb2 this function was first proposed by Dan Bernstein many years ago
3228 on comp.lang.c. Studies have shown that for certain workload this
3229 function provides a better distribution than sdbm. It generally
3230 works well with text-based inputs though it can perform extremely
3231 poorly with numeric-only input or when the total server weight is
3232 a multiple of 33, unless the avalanche modifier is also used.
3233
3234 wt6 this function was designed for haproxy while testing other
3235 functions in the past. It is not as smooth as the other ones, but
3236 is much less sensible to the input data set or to the number of
3237 servers. It can make sense as an alternative to sdbm-avalanche or
3238 djb2-avalanche for consistent hashing or when hashing on numeric
3239 data such as a source IP address or a visitor identifier in a URL
3240 parameter.
3241
3242 crc32 this is the most common CRC32 implementation as used in Ethernet,
3243 gzip, PNG, etc. It is slower than the other ones but may provide
3244 a better distribution or less predictable results especially when
3245 used on strings.
3246
3247 <modifier> indicates an optional method applied after hashing the key :
3248
3249 avalanche This directive indicates that the result from the hash
3250 function above should not be used in its raw form but that

```

3251 a 4-byte full avalanche hash must be applied first. The
3252 purpose of this step is to mix the resulting bits from the
3253 previous hash in order to avoid any undesired effect when
3254 the input contains some limited values or when the number of
3255 servers is a multiple of one of the hash's components (64
3256 for SDBM, 33 for DJB2). Enabling avalanche tends to make the
3257 result less predictable, but it's also not as smooth as when
3258 using the original function. Some testing might be needed
3259 with some workloads. This hash is one of the many proposed
3260 by Bob Jenkins.

3261 The default hash type is "map-based" and is recommended for most usages. The
3262 default function is "sdbm", the selection of a function should be based on
3263 the range of the values being hashed.

3264 See also : "balance", "server"

3265 http-check disable-on-404

3266 Enable a maintenance mode upon HTTP/404 response to health-checks

3267 May be used in sections : defaults | frontend | listen | backend
3268 yes | no | yes | yes

3269 Arguments : none

3270 When this option is set, a server which returns an HTTP code 404 will be
3271 excluded from further load-balancing, but will still receive persistent
3272 connections. This provides a very convenient method for Web administrators
3273 to perform a graceful shutdown of their servers. It is also important to note
3274 that a server which is detected as failed while it was in this mode will not
3275 generate an alert, just a notice. If the server responds 2xx or 3xx again, it
3276 will immediately be reinserted into the farm. The status on the stats page
3277 reports "NOLB" for a server in this mode. It is important to note that this
3278 option only works in conjunction with the "httpchk" option. If this option
3279 is used with "http-check expect", then it has precedence over it so that 404
3280 responses will still be considered as soft-stop.

3281 See also : "option httpchk", "http-check expect"

3282 http-check expect [!] <match> <pattern>

3283 Make HTTP health checks consider response contents or specific status codes

3284 May be used in sections : defaults | frontend | listen | backend
3285 yes | no | yes | yes

3286 Arguments :

3287 <match> is a keyword indicating how to look for a specific pattern in the
3288 response. The keyword may be one of "status", "rstatus",
3289 "string", or "rstring". The keyword may be preceded by an
3290 exclamation mark ("!") to negate the match. Spaces are allowed
3291 between the exclamation mark and the keyword. See below for more
3292 details on the supported keywords.

3293 <pattern> is the pattern to look for. It may be a string or a regular
3294 expression. If the pattern contains spaces, they must be escaped
3295 with the usual backslash ('\').

3296 By default, "option httpchk" considers that response statuses 2xx and 3xx
3297 are valid, and that others are invalid. When "http-check expect" is used,
3298 it defines what is considered valid or invalid. Only one "http-check"
3299 statement is supported in a backend. If a server fails to respond or times
3300 out, the check obviously fails. The available matches are :

3301 status <string> : test the exact string match for the HTTP status code.
3302 A health check response will be considered valid if the
3303 response's status code is exactly this string. If the
3304 "status" keyword is prefixed with "!", then the response
3305

3316 will be considered invalid if the status code matches.
3317
3318 rstatus <regex> : test a regular expression for the HTTP status code.
3319 A health check response will be considered valid if the
3320 response's status code matches the expression. If the
3321 "rstatus" keyword is prefixed with "!", then the response
3322 will be considered invalid if the status code matches.
3323 This is mostly used to check for multiple codes.
3324
3325 string <string> : test the exact string match in the HTTP response body.
3326 A health check response will be considered valid if the
3327 response's body contains this exact string. If the
3328 "string" keyword is prefixed with "!", then the response
3329 will be considered invalid if the body contains this
3330 string. This can be used to look for a mandatory word at
3331 the end of a dynamic page, or to detect a failure when a
3332 specific error appears on the check page (eg: a stack
3333 trace).
3334
3335 rstring <regex> : test a regular expression on the HTTP response body.
3336 A health check response will be considered valid if the
3337 response's body matches this expression. If the "rstring"
3338 keyword is prefixed with "!", then the response will be
3339 considered invalid if the body matches the expression.
3340 This can be used to look for a mandatory word at the end
3341 of a dynamic page, or to detect a failure when a specific
3342 error appears on the check page (eg: a stack trace).
3343
3344 It is important to note that the responses will be limited to a certain size
3345 defined by the global "tune.chksize" option, which defaults to 16384 bytes.
3346 Thus, too large responses may not contain the mandatory pattern when using
3347 "string" or "rstring". If a large response is absolutely required, it is
3348 possible to change the default max size by setting the global variable.
3349 However, it is worth keeping in mind that parsing very large responses can
3350 waste some CPU cycles, especially when regular expressions are used, and that
3351 it is always better to focus the checks on smaller resources.
3352
3353 Also "http-check expect" doesn't support HTTP keep-alive. Keep in mind that it
3354 will automatically append a "Connection: close" header, meaning that this
3355 header should not be present in the request provided by "option httpchk".
3356
3357 Last, if "http-check expect" is combined with "http-check disable-on-404",
3358 then this last one has precedence when the server responds with 404.

3359 Examples :
3360 # only accept status 200 as valid
3361 http-check expect status 200

3362 # consider SQL errors as errors
3363 http-check expect ! string SQL Error

3364 # consider status 5xx only as errors
3365 http-check expect ! rstatus ^5

3366 # check that we have a correct hexadecimal tag before /html
3367 http-check expect rstring <!--tag:[0-9a-f]*-->/html>

3368 See also : "option httpchk", "http-check disable-on-404"

3369 http-check send-state

3370 Enable emission of a state header with HTTP health checks

3371 May be used in sections : defaults | frontend | listen | backend
3372 yes | no | yes | yes

3373 Arguments : none

```

3381 When this option is set, haproxy will systematically send a special header
3382 "X-Haproxy-Server-State" with a list of parameters indicating to each server
3383 how they are seen by haproxy. This can be used for instance when a server is
3384 manipulated without access to haproxy and the operator needs to know whether
3385 haproxy still sees it up or not, or if the server is the last one in a farm.
3386
3387 The header is composed of fields delimited by semi-colons, the first of which
3388 is a word ("UP", "DOWN", "NOLB"), possibly followed by a number of valid
3389 checks on the total number before transition, just as appears in the stats
3390 interface. Next headers are in the form "<variable>=<value>", indicating in
3391 no specific order some values available in the stats interface :
3392 - a variable "address", containing the address of the backend server.
3393 This corresponds to the <address> field in the server declaration. For
3394 unix domain sockets, it will read "unix".
3395
3396 - a variable "port", containing the port of the backend server. This
3397 corresponds to the <port> field in the server declaration. For unix
3398 domain sockets, it will read "unix".
3399
3400 - a variable "name", containing the name of the backend followed by a slash
3401 ("/") then the name of the server. This can be used when a server is
3402 checked in multiple backends.
3403
3404 - a variable "node" containing the name of the haproxy node, as set in the
3405 global "node" variable, otherwise the system's hostname if unspecified.
3406
3407 - a variable "weight" indicating the weight of the server, a slash ("/")
3408 and the total weight of the farm (just counting usable servers). This
3409 helps to know if other servers are available to handle the load when this
3410 one fails.
3411
3412 - a variable "scur" indicating the current number of concurrent connections
3413 on the server, followed by a slash ("/") then the total number of
3414 connections on all servers of the same backend.
3415
3416 - a variable "qcur" indicating the current number of requests in the
3417 server's queue.
3418
3419 Example of a header received by the application server :
3420 >>> X-Haproxy-Server-State: UP 2/3; name=bck/srv2; node=lb1; weight=1/2; \
3421      scur=13/22; qcur=0
3422
3423 See also : "option httpchk", "http-check disable-on-404"
3424
3425 http-request { allow | deny | tarpit | auth [realm <realm>] | redirect <rule> |
3426      add-header <name> <fmt> | set-header <name> <fmt> |
3427      capture <sample> [ len <length> | id <id> ] |
3428      del-header <name> | set-nice <nice> | set-log-level <level> |
3429      replace-header <name> <match-regex> <replace-fmt> |
3430      replace-value <name> <match-regex> <replace-fmt> |
3431      set-method <fmt> | set-path <fmt> | set-query <fmt> |
3432      set-uri <fmt> | set-tos <tos> | set-mark <mark> |
3433      add-acl(<file name>) <key fmt> |
3434      del-acl(<file name>) <key fmt> |
3435      del-map(<file name>) <key fmt> |
3436      set-map(<file name>) <key fmt> <value fmt> |
3437      set-var(<var name>) <expr> |
3438      { track-sc0 | track-scl | track-sc2 } <key> [table <tables>] |
3439      sc-inc-gpc0(<sc-id>) |
3440      sc-set-gpt0(<sc-id>) <int> |
3441      silent-drop |
3442      { if | unless } <condition> |
3443      }
3444 Access control for Layer 7 requests
3445

```

```

3446 May be used in sections: defaults | frontend | listen | backend
3447 no | yes | yes | yes
3448
3449 The http-request statement defines a set of rules which apply to layer 7
3450 processing. The rules are evaluated in their declaration order when they are
3451 met in a frontend, listen or backend section. Any rule may optionally be
3452 followed by an ACL-based condition, in which case it will only be evaluated
3453 if the condition is true.
3454
3455 The first keyword is the rule's action. Currently supported actions include :
3456 - "allow" : this stops the evaluation of the rules and lets the request
3457 pass the check. No further "http-request" rules are evaluated.
3458
3459 - "deny" : this stops the evaluation of the rules and immediately rejects
3460 the request and emits an HTTP 403 error. No further "http-request" rules
3461 are evaluated.
3462
3463 - "tarpit" : this stops the evaluation of the rules and immediately blocks
3464 the request without responding for a delay specified by "timeout tarpit"
3465 or "timeout connect" if the former is not set. After that delay, if the
3466 client is still connected, an HTTP error 500 is returned so that the
3467 client does not suspect it has been tarptitted. Logs will report the flags
3468 "PT". The goal of the tarpit rule is to slow down robots during an attack
3469 when they're limited on the number of concurrent requests. It can be very
3470 efficient against very dumb robots, and will significantly reduce the
3471 load on firewalls compared to a "deny" rule. But when facing "correctly"
3472 developed robots, it can make things worse by forcing haproxy and the
3473 front firewall to support insane number of concurrent connections. See
3474 also the "silent-drop" action below.
3475
3476 - "auth" : this stops the evaluation of the rules and immediately responds
3477 with an HTTP 401 or 407 error code to invite the user to present a valid
3478 user name and password. No further "http-request" rules are evaluated. An
3479 optional "realm" parameter is supported, it sets the authentication realm
3480 that is returned with the response (typically the application's name).
3481
3482 - "redirect" : this performs an HTTP redirection based on a redirect rule.
3483 This is exactly the same as the "redirect" statement except that it
3484 inserts a redirect rule which can be processed in the middle of other
3485 "http-request" rules and that these rules use the "log-format" strings.
3486 See the "redirect" keyword for the rule's syntax.
3487
3488 - "add-header" appends an HTTP header field whose name is specified in
3489 <name> and whose value is defined by <fmt> which follows the log-format
3490 rules (see Custom Log Format in section 8.2.4). This is particularly
3491 useful to pass connection-specific information to the server (eg: the
3492 client's SSL certificate), or to combine several headers into one. This
3493 rule is not final, so it is possible to add other similar rules. Note
3494 that header addition is performed immediately, so one rule might reuse
3495 the resulting header from a previous rule.
3496
3497 - "set-header" does the same as "add-header" except that the header name
3498 is first removed if it existed. This is useful when passing security
3499 information to the server, where the header must not be manipulated by
3500 external users. Note that the new value is computed before the removal so
3501 it is possible to concatenate a value to an existing header.
3502
3503 - "del-header" removes all HTTP header fields whose name is specified in
3504 <name>.
3505
3506 - "replace-header" matches the regular expression in all occurrences of
3507 header field <name> according to <match-regex>, and replaces them with
3508 the <replace-fmt> argument. Format characters are allowed in replace-fmt
3509 and work like in <fmt> arguments in "add-header". The match is only
3510

```

case-sensitive. It is important to understand that this action only considers whole header lines, regardless of the number of values they may contain. This usage is suited to headers naturally containing commas in their value, such as If-Modified-Since and so on.

Example:

```
http-request replace-header Cookie foo=[^:]*;(.*) foo=\1;ip=%bi;\2
```

applied to:

```
Cookie: foo=foobar; expires=Tue, 14-Jun-2016 01:40:45 GMT;
```

outputs:

```
Cookie: foo=foobar;ip=192.168.1.20; expires=Tue, 14-Jun-2016 01:40:45 GMT;
```

assuming the backend IP is 192.168.1.20

- "replace-value" works like "replace-header" except that it matches the regex against every comma-delimited value of the header field <name> instead of the entire header. This is suited for all headers which are allowed to carry more than one value. An example could be the Accept header.

Example:

```
http-request replace-value X-Forwarded-For ^192\.168\.(.*)$ 172.16.\1
```

applied to:

```
X-Forwarded-For: 192.168.10.1, 192.168.13.24, 10.0.0.37
```

outputs:

```
X-Forwarded-For: 172.16.10.1, 172.16.13.24, 10.0.0.37
```

- "set-method" rewrites the request method with the result of the evaluation of format string <fmt>. There should be very few valid reasons for having to do so as this is more likely to break something than to fix it.

- "set-path" rewrites the request path with the result of the evaluation of format string <fmt>. The query string, if any, is left intact. If a scheme and authority is found before the path, they are left intact as well. If the request doesn't have a path ("**"), this one is replaced with the format. This can be used to prepend a directory component in front of a path for example. See also "set-query" and "set-uri".

Example :

```
# prepend the host name before the path
http-request set-path /[hdr(host)]%[path]
```

- "set-query" rewrites the request's query string which appears after the first question mark ("?") with the result of the evaluation of format string <fmt>. The part prior to the question mark is left intact. If the request doesn't contain a question mark and the new value is not empty, then one is added at the end of the URI, followed by the new value. If a question mark was present, it will never be removed even if the value is empty. This can be used to add or remove parameters from the query string. See also "set-query" and "set-uri".

Example :

```
# replace "%3D" with "=" in the query string
http-request set-query %[query,regsub(%3D,=,g)]
```

- "set-uri" rewrites the request URI with the result of the evaluation of format string <fmt>. The scheme, authority, path and query string are all replaced at once. This can be used to rewrite hosts in front of proxies, or to perform complex modifications to the URI such as moving parts between the path and the query string. See also "set-path" and "set-query".

- "set-nice" sets the "nice" factor of the current request being processed. It only has effect against the other requests being processed at the same time. The default value is 0, unless altered by the "nice" setting on the "bind" line. The accepted range is -1024..1024. The higher the value, the nicest the request will be. Lower values will make the request more important than other ones. This can be useful to improve the speed of some requests, or lower the priority of non-important requests. Using this setting without prior experimentation can cause some major slowdown.

- "set-log-level" is used to change the log level of the current request when a certain condition is met. Valid levels are the 8 syslog levels (see the "log" keyword) plus the special level "silent" which disables logging for this request. This rule is not final so the last matching rule wins. This rule can be useful to disable health checks coming from another equipment.

- "set-tos" is used to set the TOS or DSCP field value of packets sent to the client to the value passed in <tos> on platforms which support this. This value represents the whole 8 bits of the IP TOS field, and can be expressed both in decimal or hexadecimal format (prefixed by "0x"). Note that only the 6 higher bits are used in DSCP or TOS, and the two lower bits are always 0. This can be used to adjust some routing behaviour on border routers based on some information from the request. See RFC 2474, 2597, 3260 and 4594 for more information.

- "set-mark" is used to set the Netfilter MARK on all packets sent to the client to the value passed in <mark> on platforms which support it. This value is an unsigned 32 bit value which can be matched by netfilter and by the routing table. It can be expressed both in decimal or hexadecimal format (prefixed by "0x"). This can be useful to force certain packets to take a different route (for example a cheaper network path for bulk downloads). This works on Linux kernels 2.6.32 and above and requires admin privileges.

- "add-acl" is used to add a new entry into an ACL. The ACL must be loaded from a file (even a dummy empty file). The file name of the ACL to be updated is passed between parentheses. It takes one argument: <key fmt>, which follows log-format rules, to collect content of the new entry. It performs a lookup in the ACL before insertion, to avoid duplicated (or more) values. This lookup is done by a linear search and can be expensive with large lists! It is the equivalent of the "add acl" command from the stats socket, but can be triggered by an HTTP request.

- "del-acl" is used to delete an entry from an ACL. The ACL must be loaded from a file (even a dummy empty file). The file name of the ACL to be updated is passed between parentheses. It takes one argument: <key fmt>, which follows log-format rules, to collect content of the entry to delete. It is the equivalent of the "del acl" command from the stats socket, but can be triggered by an HTTP request.

- "del-map" is used to delete an entry from a MAP. The MAP must be loaded from a file (even a dummy empty file). The file name of the MAP to be updated is passed between parentheses. It takes one argument: <key fmt>, which follows log-format rules, to collect content of the entry to delete. It takes one argument: "file name" It is the equivalent of the "del map" command from the stats socket, but can be triggered by an HTTP request.

3640

- "set-map" is used to add a new entry into a MAP. The MAP must be loaded from a file (even a dummy empty file). The file name of the MAP to be updated is passed between parentheses. It takes 2 arguments: <key fmt>, which follows log-format rules, used to collect MAP key, and <value fmt>, which follows log-format rules, used to collect content for the new entry. It performs a lookup in the MAP before insertion, to avoid duplicated (or more) values. This lookup is done by a linear search and can be expensive with large lists! It is the equivalent of the "set map" command from the stats socket, but can be triggered by an HTTP request.

- capture <sample> [len <length> | id <id>] : captures sample expression <sample> from the request buffer, and converts it to a string of at most <len> characters. The resulting string is stored into the next request "capture" slot, so it will possibly appear next to some captured HTTP headers. It will then automatically appear in the logs, and it will be possible to extract it using sample fetch rules to feed it into headers or anything. The length should be limited given that this size will be allocated for each capture during the whole session life. Please check section 7.3 (Fetching samples) and "capture request header" for more information..

If the keyword "id" is used instead of "len", the action tries to store the captured string in a previously declared capture slot. This is useful to run captures in backends. The slot id can be declared by a previous directive "http-request capture" or with the "declare capture" keyword. If the slot <id> doesn't exist, then HAProxy fails parsing the configuration to prevent unexpected behavior at run time.

- { track-sc0 | track-scl | track-sc2 } <key> [table <table>] : enables tracking of sticky counters from current request. These rules do not stop evaluation and do not change default action. Three sets of counters may be simultaneously tracked by the same connection. The first "track-sc0" rule executed enables tracking of the counters of the specified table as the first set. The first "track-scl" rule executed enables tracking of the counters of the specified table as the second set. The first "track-sc2" rule executed enables tracking of the counters of the specified table as the third set. It is a recommended practice to use the first set of counters for the per-frontend counters and the second set for the per-backend ones. But this is just a guideline, all may be used everywhere.

These actions take one or two arguments :

<key> is mandatory, and is a sample expression rule as described in section 7.3. It describes what elements of the incoming request or connection will be analysed, extracted, combined, and used to select which table entry to update the counters.

<table> is an optional table to be used instead of the default one, which is the stick-table declared in the current proxy. All the counters for the matches and updates for the key will then be performed in that table until the session ends.

Once a "track-sc*" rule is executed, the key is looked up in the table and if it is not found, an entry is allocated for it. Then a pointer to that entry is kept during all the session's life, and this entry's counters are updated as often as possible, every time the session's counters are updated, and also systematically when the session ends. Counters are only updated for events that happen after the tracking has been started. As an exception, connection counters and request counters are systematically updated so that they reflect useful information.

If the entry tracks concurrent connection counters, one connection is counted for as long as the entry is tracked, and the entry will not expire during that time. Tracking counters also provides a performance advantage over just checking the keys, because only one table lookup is

performed for all ACL checks that make use of it.

- sc-set-gpt0(<sc-id>) <int> : This action sets the GPT0 tag according to the sticky counter designated by <sc-id> and the value of <int>. The expected result is a boolean. If an error occurs, this action silently fails and the actions evaluation continues.

- sc-inc-gpc0(<sc-id>) : This action increments the GPC0 counter according with the sticky counter designated by <sc-id>. If an error occurs, this action silently fails and the actions evaluation continues.

- set-var(<var-name>) <expr> : Is used to set the contents of a variable. The variable is declared inline.

<var-name> The name of the variable starts by an indication about its scope. The allowed scopes are:

"sess" : the variable is shared with all the session.

"txn" : the variable is shared with all the transaction (request and response)

"req" : the variable is shared only during the request processing

"res" : the variable is shared only during the response processing.

This prefix is followed by a name. The separator is a '.',

The name may only contain characters 'a-z', 'A-Z', '0-9', and '_'.

<expr> Is a standard HAProxy expression formed by a sample-fetch followed by some converters.

Example:

```
http-request set-var(req.my_var) req.fhdr(user-agent),lower
```

- set-src <expr> :

Is used to set the source IP address to the value of specified expression. Useful when a proxy in front of HAProxy rewrites source IP, but provides the correct IP in a HTTP header; or you want to mask source IP for privacy.

<expr> Is a standard HAProxy expression formed by a sample-fetch followed by some converters.

Example:

```
http-request set-src hdr(x-forwarded-for)
```

```
http-request set-src src.ipmask(24)
```

When set-src is successful, the source port is set to 0.

- "silent-drop" : this stops the evaluation of the rules and makes the client-facing connection suddenly disappear using a system-dependant way that tries to prevent the client from being notified. The effect it then that the client still sees an established connection while there's none on HAProxy. The purpose is to achieve a comparable effect to "tarpit" except that it doesn't use any local resource at all on the machine running HAProxy. It can resist much higher loads than "tarpit", and slow down stronger attackers. It is important to understand the impact of using this mechanism. All stateful equipments placed between the client and HAProxy (firewalls, proxies, load balancers) will also keep the established connection for a long time and may suffer from this action. On modern Linux systems running with enough privileges, the TCP_REPAIR

socket option is used to block the emission of a TCP reset. On other systems, the socket's TTL is reduced to 1 so that the TCP reset doesn't pass the first router, though it's still delivered to local networks. Do not use it unless you fully understand how it works.

There is no limit to the number of http-request statements per instance.

It is important to know that http-request rules are processed very early in the HTTP processing, just after "block" rules and before "reqdel" or "reqrep" or "reqadd" rules. That way, headers added by "add-header"/"set-header" are visible by almost all further ACL rules.

Using "reqadd"/"reqdel"/"reqrep" to manipulate request headers is discouraged in newer versions (≥ 1.5). But if you need to use regular expression to delete headers, you can still use "reqdel". Also please use "http-request deny/allow/tarpit" instead of "reqdeny"/"reqpass"/"reqtarpit".

Example:

```
acl nagios src 192.168.129.3
acl local_net src 192.168.0.0/16
acl auth_ok http_auth(L1)
```

```
http-request allow if nagios
http-request allow if local_net auth_ok
http-request auth realm Gimme if local_net auth_ok
http-request deny
```

Example:

```
acl auth_ok http_auth_group(L1) G1
http-request auth unless auth_ok
```

Example:

```
http-request set-header X-Haproxy-Current-Date %T
http-request set-header X-SSL %[ssl_fc]
http-request set-header X-SSL-Session-ID %[ssl_fc_session_id_hex]
http-request set-header X-SSL-Client-Verify %[ssl_c_verify]
http-request set-header X-SSL-Client-DN %({+Q})[ssl_c_s_dn]
http-request set-header X-SSL-Client-CN %({+Q})[ssl_c_s_dn(cn)]
http-request set-header X-SSL-Issuer %({+Q})[ssl_c_i_dn]
http-request set-header X-SSL-Client-NotBefore %({+Q})[ssl_c_notbefore]
http-request set-header X-SSL-Client-NotAfter %({+Q})[ssl_c_notafter]
```

Example:

```
acl key req,hdr(X-Add-Acl-Key) -m found
acl add path /addacl
acl del path /delacl
```

```
acl myhost hdr(Host) -f myhost.lst
```

```
http-request add-acl(myhost.lst) %[req,hdr(X-Add-Acl-Key)] if key add
http-request del-acl(myhost.lst) %[req,hdr(X-Add-Acl-Key)] if key del
```

Example:

```
acl value req,hdr(X-Value) -m found
acl setmap path /setmap
acl delmap path /delmap

use_backend bk_appli if { hdr(Host),map_str(map.lst) -m found }

http-request set-map(map.lst) %[src] %[req,hdr(X-Value)] if setmap value
http-request del-map(map.lst) %[src] if delmap
```

See also : "stats http-request", section 3.4 about userlists and section 7 about ACL usage.

```
http-response { allow | deny | add-header <name> <fmt> | set-nice <nice> |
capture <sample> id <id> | redirect <rules> |
set-header <name> <fmt> | del-header <name> |
replace-header <name> <regex-match> <replace-fmt> |
replace-value <name> <regex-match> <replace-fmt> |
set-status <status> |
set-log-level <level> | set-mark <mark> | set-tos <tos> |
add-acl(<file name>) <key fmt> |
del-acl(<file name>) <key fmt> |
del-map(<file name>) <key fmt> |
set-map(<file name>) <key fmt> <value fmt> |
set-var(<var-name>) <expr> |
sc-inc-gp0(<sc-id>) |
sc-set-gp0(<sc-id>) <int> |
silent-drop |
}
[ { if | unless } <condition> ]
```

Access control for Layer 7 responses

May be used in sections: defaults no | yes | listen | backend
no | yes | yes | yes

The http-response statement defines a set of rules which apply to layer 7 processing. The rules are evaluated in their declaration order when they are met in a frontend, listen or backend section. Any rule may optionally be followed by an ACL-based condition, in which case it will only be evaluated if the condition is true. Since these rules apply on responses, the backend rules are applied first, followed by the frontend's rules.

The first keyword is the rule's action. Currently supported actions include :
- "allow" : this stops the evaluation of the rules and lets the response pass the check. No further "http-response" rules are evaluated for the current section.

- "deny" : this stops the evaluation of the rules and immediately rejects the response and emits an HTTP 502 error. No further "http-response" rules are evaluated.

- "add-header" appends an HTTP header field whose name is specified in <name> and whose value is defined by <fmt> which follows the log-format rules (see Custom Log Format in section 8.2.4). This may be used to send a cookie to a client for example, or to pass some internal information. This rule is not final, so it is possible to add other similar rules. Note that header addition is performed immediately, so one rule might reuse the resulting header from a previous rule.

- "set-header" does the same as "add-header" except that the header name is first removed if it existed. This is useful when passing security information to the server, where the header must not be manipulated by external users.

- "del-header" removes all HTTP header fields whose name is specified in <name>.

- "replace-header" matches the regular expression in all occurrences of header field <name> according to <match-regex>, and replaces them with the <replace-fmt> argument. Format characters are allowed in replace-fmt and work like in <fmt> arguments in "add-header". The match is only case-sensitive. It is important to understand that this action only considers whole header lines, regardless of the number of values they may contain. This usage is suited to headers naturally containing commas in their value, such as Set-Cookie, Expires and so on.

Example:

3900

```
3901 http-response replace-header Set-Cookie (C=[^:]*);(.*) \1:ip=%bi;\2
3902 applied to:
3903
3904 Set-Cookie: C=1; expires=Tue, 14-Jun-2016 01:40:45 GMT
3905
3906 outputs:
3907
3908 Set-Cookie: C=1;ip=192.168.1.20; expires=Tue, 14-Jun-2016 01:40:45 GMT
3909 assuming the backend IP is 192.168.1.20.
3910
3911 - "replace-value" works like "replace-header" except that it matches the
3912 regex against every comma-delimited value of the header field <name>
3913 instead of the entire header. This is suited for all headers which are
3914 allowed to carry more than one value. An example could be the Accept
3915 header.
3916
3917 Example:
3918
3919 http-response replace-value Cache-control ^public$ private
3920 applied to:
3921
3922 Cache-Control: max-age=3600, public
3923
3924 outputs:
3925
3926 Cache-Control: max-age=3600, private
3927
3928 - "set-status" replaces the response status code with <status> which must
3929 be an integer between 100 and 999. Note that the reason is automatically
3930 adapted to the new code.
3931
3932 Example:
3933
3934 # return "431 Request Header Fields Too Large"
3935 http-response set-status 431
3936
3937 - "set-nice" sets the "nice" factor of the current request being processed.
3938 It only has effect against the other requests being processed at the same
3939 time. The default value is 0, unless altered by the "nice" setting on the
3940 "bind" line. The accepted range is -1024..1024. The higher the value, the
3941 nicest the request will be. Lower values will make the request more
3942 important than other ones. This can be useful to improve the speed of
3943 some requests, or lower the priority of non-important requests. Using
3944 this setting without prior experimentation can cause some major slowdown.
3945
3946 - "set-log-level" is used to change the log level of the current request
3947 when a certain condition is met. Valid levels are the 8 syslog levels
3948 (see the "log" keyword) plus the special level "silent" which disables
3949 logging for this request. This rule is not final so the last matching
3950 rule wins. This rule can be useful to disable health checks coming from
3951 another equipment.
3952
3953 - "set-tos" is used to set the TOS or DSCP field value of packets sent to
3954 the client to the value passed in <tos> on platforms which support this.
3955 This value represents the whole 8 bits of the IP TOS field, and can be
3956 expressed both in decimal or hexadecimal format (prefixed by "0x"). Note
3957 that only the 6 higher bits are used in DSCP or TOS, and the two lower
3958 bits are always 0. This can be used to adjust some routing behaviour on
3959 border routers based on some information from the request. See RFC 2474,
3960 2597, 3260 and 4594 for more information.
3961
3962 - "set-mark" is used to set the Netfilter MARK on all packets sent to the
```

```
3966 client to the value passed in <mark> on platforms which support it. This
3967 value is an unsigned 32 bit value which can be matched by netfilter and
3968 by the routing table. It can be expressed both in decimal or hexadecimal
3969 format (prefixed by "0x"). This can be useful to force certain packets to
3970 take a different route (for example a cheaper network path for bulk
3971 downloads). This works on Linux kernels 2.6.32 and above and requires
3972 admin privileges.
3973
3974 - "add-acl" is used to add a new entry into an ACL. The ACL must be loaded
3975 from a file (even a dummy empty file). The file name of the ACL to be
3976 updated is passed between parentheses. It takes one argument: <key fmt>,
3977 which follows log-format rules, to collect content of the new entry. It
3978 performs a lookup in the ACL before insertion, to avoid duplicated (or
3979 more) values. This lookup is done by a linear search and can be expensive
3980 with large lists! It is the equivalent of the "add acl" command from the
3981 stats socket, but can be triggered by an HTTP response.
3982
3983 - "del-acl" is used to delete an entry from an ACL. The ACL must be loaded
3984 from a file (even a dummy empty file). The file name of the ACL to be
3985 updated is passed between parentheses. It takes one argument: <key fmt>,
3986 which follows log-format rules, to collect content of the entry to delete.
3987 It is the equivalent of the "del acl" command from the stats socket, but
3988 can be triggered by an HTTP response.
3989
3990 - "del-map" is used to delete an entry from a MAP. The MAP must be loaded
3991 from a file (even a dummy empty file). The file name of the MAP to be
3992 updated is passed between parentheses. It takes one argument: <key fmt>,
3993 which follows log-format rules, to collect content of the entry to delete.
3994 It takes one argument: "file name" It is the equivalent of the "del map"
3995 command from the stats socket, but can be triggered by an HTTP response.
3996
3997 - "set-map" is used to add a new entry into a MAP. The MAP must be loaded
3998 from a file (even a dummy empty file). The file name of the MAP to be
3999 updated is passed between parentheses. It takes 2 arguments: <key fmt>,
4000 which follows log-format rules, used to collect MAP key, and <value fmt>,
4001 which follows log-format rules, used to collect content for the new entry.
4002 It performs a lookup in the MAP before insertion, to avoid duplicated (or
4003 more) values. This lookup is done by a linear search and can be expensive
4004 with large lists! It is the equivalent of the "set map" command from the
4005 stats socket, but can be triggered by an HTTP response.
4006
4007 - capture <sample> id <id> :
4008 captures sample expression <sample> from the response buffer, and converts
4009 it to a string. The resulting string is stored into the next request
4010 "capture" slot, so it will possibly appear next to some captured HTTP
4011 headers. It will then automatically appear in the logs, and it will be
4012 possible to extract it using sample fetch rules to feed it into headers or
4013 anything. Please check section 7.3 (Fetching samples) and "capture
4014 response header" for more information.
4015
4016 The keyword "id" is the id of the capture slot which is used for storing
4017 the string. The capture slot must be defined in an associated frontend.
4018 This is useful to run captures in backends. The slot id can be declared by
4019 a previous directive "http-response capture" or with the "declare capture"
4020 keyword.
4021
4022 If the slot <id> doesn't exist, then HAProxy fails parsing the
4023 configuration to prevent unexpected behavior at run time.
4024
4025 - "redirect" : this performs an HTTP redirection based on a redirect rule.
4026 This supports a format string similarly to "http-request redirect" rules,
4027 with the exception that only the "location" type of redirect is possible
4028 on the response. See the "redirect" keyword for the rule's syntax. When
4029 a redirect rule is applied during a response, connections to the server
4030 are closed so that no data can be forwarded from the server to the client.
```

```

4031 - set-var(<var-name>) expr:
4032 Is used to set the contents of a variable. The variable is declared
4033 inline.
4034
4035 <var-name> The name of the variable starts by an indication about its
4036 scope. The allowed scopes are:
4037 "sess" : the variable is shared with all the session,
4038 "txn" : the variable is shared with all the transaction
4039 (request and response)
4040 "req" : the variable is shared only during the request
4041 processing
4042 "res" : the variable is shared only during the response
4043 processing.
4044 This prefix is followed by a name. The separator is a '::'.
4045 The name may only contain characters 'a-z', 'A-Z', '0-9',
4046 and '_'.
4047
4048 <expr> Is a standard HAProxy expression formed by a sample-fetch
4049 followed by some converters.
4050
4051 Example:
4052 http-response set-var(sess.last_redir) res.hdr(location)
4053
4054 - sc-set-gpt0(<sc-id>) <int> :
4055 This action sets the GPT0 tag according to the sticky counter designated
4056 by <sc-id> and the value of <int>. The expected result is a boolean. If
4057 an error occurs, this action silently fails and the actions evaluation
4058 continues.
4059
4060 - sc-inc-gpc0(<sc-id>):
4061 This action increments the GPC0 counter according with the sticky counter
4062 designated by <sc-id>. If an error occurs, this action silently fails and
4063 the actions evaluation continues.
4064
4065 - "silent-drop" : this stops the evaluation of the rules and makes the
4066 client-facing connection suddenly disappear using a system-dependant way
4067 that tries to prevent the client from being notified. The effect it then
4068 has is that the client still sees an established connection while there's none
4069 on HAProxy. The purpose is to achieve a comparable effect to "tarpit"
4070 except that it doesn't use any local resource at all on the machine
4071 running HAProxy. It can resist much higher loads than "tarpit", and slow
4072 down stronger attackers. It is important to understand the impact of using
4073 this mechanism. All stateful equipments placed between the client and
4074 HAProxy (firewalls, proxies, load balancers) will also keep the
4075 established connection for a long time and may suffer from this action.
4076 On modern Linux systems running with enough privileges, the TCP REPAIR
4077 socket option is used to block the emission of a TCP reset. On other
4078 systems, the socket's TTL is reduced to 1 so that the TCP reset doesn't
4079 pass the first router, though it's still delivered to local networks. Do
4080 not use it unless you fully understand how it works.
4081
4082 There is no limit to the number of http-response statements per instance.
4083
4084 It is important to know that http-response rules are processed very early in
4085 the HTTP processing, before "rspdel" or "rsprep" or "rspadd" rules. That way,
4086 headers added by "add-header"/"set-header" are visible by almost all further
4087 rules.
4088
4089 Using "rspadd"/"rspdel"/"rsprep" to manipulate request headers is discouraged
4090 in newer versions (>= 1.5). But if you need to use regular expression to
4091 delete headers, you can still use "rspdel". Also please use
4092 "http-response deny" instead of "rspdeny".
4093
4094 Example:
4095

```

```

4096 acl key_acl res.hdr(X-Acl-Key) -m found
4097
4098 acl myhost hdr(Host) -f myhost.lst
4099
4100 http-response add-acl(myhost.lst) %[res.hdr(X-Acl-Key)] if key_acl
4101 http-response del-acl(myhost.lst) %[res.hdr(X-Acl-Key)] if key_acl
4102
4103 Example:
4104 acl value res.hdr(X-Value) -m found
4105
4106 use_backend bk_appli if { hdr(Host).map_str(map.lst) -m found }
4107
4108 http-response set-map(map.lst) %[src] %[res.hdr(X-Value)] if value
4109 http-response del-map(map.lst) %[src] if ! value
4110
4111 See also : "http-request", section 3.4 about userlists and section 7 about
4112 ACL usage.
4113
4114 http-reuse { never | safe | aggressive | always }
4115 Declare how idle HTTP connections may be shared between requests
4116
4117 May be used in sections: defaults | frontend | listen | backend
4118 yes | no | yes | yes
4119
4120 By default, a connection established between haproxy and the backend server
4121 belongs to the session that initiated it. The downside is that between the
4122 response and the next request, the connection remains idle and is not used.
4123 In many cases for performance reasons it is desirable to make it possible to
4124 reuse these idle connections to serve other requests from different sessions.
4125 This directive allows to tune this behaviour.
4126
4127 The argument indicates the desired connection reuse strategy :
4128
4129 - "never" : idle connections are never shared between sessions. This is
4130 the default choice. It may be enforced to cancel a different
4131 strategy inherited from a defaults section or for
4132 troubleshooting. For example, if an old bogus application
4133 considers that multiple requests over the same connection come
4134 from the same client and it is not possible to fix the
4135 application, it may be desirable to disable connection sharing
4136 in a single backend. An example of such an application could
4137 be an old haproxy using cookie insertion in tunnel mode and
4138 not checking any request past the first one.
4139
4140 - "safe" : this is the recommended strategy. The first request of a
4141 session is always sent over its own connection, and only
4142 subsequent requests may be dispatched over other existing
4143 connections. This ensures that in case the server closes the
4144 connection when the request is being sent, the browser can
4145 decide to silently retry it. Since it is exactly equivalent to
4146 regular keep-alive, there should be no side effects.
4147
4148 - "aggressive" : this mode may be useful in websockets environments where
4149 all servers are not necessarily known and where it would be
4150 appreciable to deliver most first requests over existing
4151 connections. In this case, first requests are only delivered
4152 over existing connections that have been reused at least once,
4153 proving that the server correctly supports connection reuse.
4154 It should only be used when it's sure that the client can
4155 retry a failed request once in a while and where the benefit
4156 of aggressive connection reuse significantly outweighs the
4157 downsides of rare connection failures.
4158
4159 - "always" : this mode is only recommended when the path to the server is
4160

```


known for never breaking existing connections quickly after releasing them. It allows the first request of a session to be sent to an existing connection. This can provide a significant performance increase over the "safe" strategy when the backend is a cache farm, since such components tend to show a consistent behaviour and will benefit from the connection sharing. It is recommended that the "http-keep-alive" timeout remains low in this mode so that no dead connections remain usable. In most cases, this will lead to the same performance gains as "aggressive" but with more risks. It should only be used when it improves the situation over "aggressive".

When http connection sharing is enabled, a great care is taken to respect the connection properties and compatibilities. Specifically :

- connections made with "userc" followed by a client-dependant value ("client", "clientip", "hdr_ip") are marked private and never shared ;
- connections sent to a server with a TLS SNI extension are marked private and are never shared ;
- connections receiving a status code 401 or 407 expect some authentication to be sent in return. Due to certain bogus authentication schemes (such as NTLM) relying on the connection, these connections are marked private and are never shared ;

No connection pool is involved, once a session dies, the last idle connection it was attached to is deleted at the same time. This ensures that connections may not last after all sessions are closed.

Note: connection reuse improves the accuracy of the "server maxconn" setting, because almost no new connection will be established while idle connections remain available. This is particularly true with the "always" strategy.

See also : "option http-keep-alive", "server maxconn"

http-send-name-header [<header>]

Add the server name to a request. Use the header string given by <header>

May be used in sections: defaults yes | no | yes | listen | backend

Arguments :

<header> The header string to use to send the server name

The "http-send-name-header" statement causes the name of the target server to be added to the headers of an HTTP request. The name is added with the header string proved.

See also : "server"

id <value>

Set a persistent ID to a proxy.

May be used in sections : defaults | frontend | listen | backend

Arguments : none

Set a persistent ID for the proxy. This ID must be unique and positive. An unused ID will automatically be assigned if unset. The first assigned value will be 1. This ID is currently only returned in statistics.

ignore-persist { if | unless } <condition>
Declare a condition to ignore persistence

May be used in sections: defaults no | frontend | listen | backend

By default, when cookie persistence is enabled, every requests containing the cookie are unconditionally persistent (assuming the target server is up and running).

The "ignore-persist" statement allows one to declare various ACL-based conditions which, when met, will cause a request to ignore persistence. This is sometimes useful to load balance requests for static files, which often don't require persistence. This can also be used to fully disable persistence for a specific User-Agent (for example, some web crawler bots).

The persistence is ignored when an "if" condition is met, or unless an "unless" condition is met.

See also : "force-persist", "cookie", and section 7 about ACL usage.

load-server-state-from-file { global | local | none }

Allow seamless reload of HAProxy

May be used in sections: defaults | frontend | listen | backend

This directive points HAProxy to a file where server state from previous running process has been saved. That way, when starting up, before handling traffic, the new process can apply old states to servers exactly has if no reload occurred. The purpose of the "load-server-state-from-file" directive is to tell haproxy which file to use. For now, only 2 arguments to either prevent loading state or load states from a file containing all backends and servers. The state file can be generated by running the command "show servers state" over the stats socket and redirect output.

The format of the file is versionned and is very specific. To understand it, please read the documentation of the "show servers state" command (chapter 9.2 of Management Guide).

Arguments:

global load the content of the file pointed by the global directive named "server-state-file".

local load the content of the file pointed by the directive "server-state-file-name" if set. If not set, then the backend name is used as a file name.

none don't load any stat for this backend

Notes:

- server's IP address is not updated unless DNS resolution is enabled on the server. It means that if a server IP address has been changed using the stat socket, this information won't be re-applied after reloading.

- server's weight is applied from previous running process unless it has been changed between previous and new configuration files.

Example 1:

Minimal configuration:

global
stats socket /tmp/socket
server-state-file /tmp/server_state

defaults
load-server-state-from-file global

```
4291 backend bk
4292 server s1 127.0.0.1:22 check weight 11
4293 server s2 127.0.0.1:22 check weight 12
4294
4295 Then one can run :
4296
4297 socat /tmp/socket - <<< "show servers state" > /tmp/server_state
4298
4299 Content of the file /tmp/server_state would be like this:
4300
4301 1
4302 # <field names skipped for the doc example>
4303 1 bk 1 s1 127.0.0.1 2 0 11 11 4 6 3 4 6 0 0
4304 1 bk 2 s2 127.0.0.1 2 0 12 12 4 6 3 4 6 0 0
4305
4306 Example 2:
4307
4308 Minimal configuration:
4309
4310 global
4311 stats socket /tmp/socket
4312 server-state-base /etc/haproxy/states
4313
4314 defaults
4315 load-server-state-from-file local
4316
4317 backend bk
4318 server s1 127.0.0.1:22 check weight 11
4319 server s2 127.0.0.1:22 check weight 12
4320
4321 Then one can run :
4322
4323 socat /tmp/socket - <<< "show servers state bk" > /etc/haproxy/states/bk
4324
4325 Content of the file /etc/haproxy/states/bk would be like this:
4326
4327 1
4328 # <field names skipped for the doc example>
4329 1 bk 1 s1 127.0.0.1 2 0 11 11 4 6 3 4 6 0 0
4330 1 bk 2 s2 127.0.0.1 2 0 12 12 4 6 3 4 6 0 0
4331
4332 See also: "server-state-file", "server-state-file-name", and
4333 "show servers state"
4334
4335 log global
4336 log <address> [len <length>] <facility> [<level> [<minlevel>]]
4337 no log
4338
4339 Enable per-instance logging of events and traffic.
4340 May be used in sections : defaults | frontend | listen | backend
4341 yes | yes | yes | yes | yes | yes
4342
4343 Prefix :
4344 no
4345
4346 should be used when the logger list must be flushed. For example,
4347 if you don't want to inherit from the default logger list. This
4348 prefix does not allow arguments.
4349
4350 Arguments :
4351 global
4352
4353 should be used when the instance's logging parameters are the
4354 same as the global ones. This is the most common usage. "global"
4355 replaces <address>, <facility> and <level> with those of the log
4356 entries found in the "global" section. Only one "log global"
4357 statement may be used per instance, and this form takes no other
4358 parameter.
```

```
4356 <address> indicates where to send the logs. It takes the same format as
4357 for the "global" section's logs, and can be one of :
4358
4359 - An IPv4 address optionally followed by a colon (':') and a UDP
4360 port. If no port is specified, 514 is used by default (the
4361 standard syslog port).
4362
4363 - An IPv6 address followed by a colon (':') and optionally a UDP
4364 port. If no port is specified, 514 is used by default (the
4365 standard syslog port).
4366
4367 - A filesystem path to a UNIX domain socket, keeping in mind
4368 considerations for chroot (be sure the path is accessible
4369 inside the chroot) and uid/gid (be sure the path is
4370 appropriately writable).
4371
4372 You may want to reference some environment variables in the
4373 address parameter, see section 2.3 about environment variables.
4374
4375 <length> is an optional maximum line length. Log lines larger than this
4376 value will be truncated before being sent. The reason is that
4377 syslog servers act differently on log line length. All servers
4378 support the default value of 1024, but some servers simply drop
4379 larger lines while others do log them. If a server supports long
4380 lines, it may make sense to set this value here in order to avoid
4381 truncating long lines. Similarly, if a server drops long lines,
4382 it is preferable to truncate them before sending them. Accepted
4383 values are 80 to 65535 inclusive. The default value of 1024 is
4384 generally fine for all standard usages. Some specific cases of
4385 long captures or JSON-formatted logs may require larger values.
4386
4387 <facility> must be one of the 24 standard syslog facilities :
4388
4389 kern user mail daemon auth syslog lpr news
4390 uucp cron auth2 ftp ntp audit alert cron2
4391 local0 local1 local2 local3 local4 local5 local6 local7
4392
4393 <level> is optional and can be specified to filter outgoing messages. By
4394 default, all messages are sent. If a level is specified, only
4395 messages with a severity at least as important as this level
4396 will be sent. An optional minimum level can be specified. If it
4397 is set, logs emitted with a more severe level than this one will
4398 be capped to this level. This is used to avoid sending "emerg"
4399 messages on all terminals on some default syslog configurations.
4400 Eight levels are known :
4401
4402 emerg alert crit err warning notice info debug
4403
4404 It is important to keep in mind that it is the frontend which decides what to
4405 log from a connection, and that in case of content switching, the log entries
4406 from the backend will be ignored. Connections are logged at level "info".
4407
4408 However, backend log declaration define how and where servers status changes
4409 will be logged. Level "notice" will be used to indicate a server going up,
4410 "warning" will be used for termination signals and definitive service
4411 termination, and "alert" will be used for when a server goes down.
4412
4413 Note : According to RFC3164, messages are truncated to 1024 bytes before
4414 being emitted.
4415
4416 Example :
4417 log global
4418 log 127.0.0.1:514 local0 notice # only send important events
4419 log 127.0.0.1:514 local0 notice # same but limit output level
4420 log "${LOCAL_SYSLDG}:514" local0 notice # send to local server
```

```

4421
4422
4423 log-format <string>
4424 Specifies the log format string to use for traffic logs
4425 May be used in sections: defaults | frontend | listen | backend
4426 yes | yes | yes | no
4427
4428 This directive specifies the log format string that will be used for all logs
4429 resulting from traffic passing through the frontend using this line. If the
4430 directive is used in a defaults section, all subsequent frontends will use
4431 the same log format. Please see section 8.2.4 which covers the log format
4432 string in depth.
4433
4434 log-format-sd <string>
4435 Specifies the RFC5424 structured-data log format string
4436 May be used in sections: defaults | frontend | listen | backend
4437 yes | yes | yes | no
4438
4439 This directive specifies the RFC5424 structured-data log format string that
4440 will be used for all logs resulting from traffic passing through the frontend
4441 using this line. If the directive is used in a defaults section, all
4442 subsequent frontends will use the same log format. Please see section 8.2.4
4443 which covers the log format string in depth.
4444
4445 See https://tools.ietf.org/html/rfc5424#section-6.3 for more information
4446 about the RFC5424 structured-data part.
4447
4448 Note : This log format string will be used only for loggers that have set
4449 log format to "rfc5424".
4450
4451 Example :
4452 log-format-sd [exampleSDID@1234\ bytes="\%B\\"" status="\%ST\"]
4453
4454
4455 log-tag <string>
4456 Specifies the log tag to use for all outgoing logs
4457 May be used in sections: defaults | frontend | listen | backend
4458 yes | yes | yes | yes
4459
4460 Sets the tag field in the syslog header to this string. It defaults to the
4461 log-tag set in the global section, otherwise the program name as launched
4462 from the command line, which usually is "haproxy". Sometimes it can be useful
4463 to differentiate between multiple processes running on the same host, or to
4464 differentiate customer instances running in the same process. In the backend,
4465 logs about servers up/down will use this tag. As a hint, it can be convenient
4466 to set a log-tag related to a hosted customer in a defaults section then put
4467 all the frontends and backends for that customer, then start another customer
4468 in a new defaults section. See also the global "log-tag" directive.
4469
4470 max-keep-alive-queue <value>
4471 Set the maximum server queue size for maintaining keep-alive connections
4472 May be used in sections: defaults | frontend | listen | backend
4473 yes | no | yes | yes
4474
4475 HTTP keep-alive tries to reuse the same server connection whenever possible,
4476 but sometimes it can be counter-productive, for example if a server has a lot
4477 of connections while other ones are idle. This is especially true for static
4478 servers.
4479
4480 The purpose of this setting is to set a threshold on the number of queued
4481 connections at which haproxy stops trying to reuse the same server and prefers
4482 to find another one. The default value, -1, means there is no limit. A value
4483 of zero means that keep-alive requests will never be queued. For very close
4484 servers which can be reached with a low latency and which are not sensible to
4485 breaking keep-alive, a low value is recommended (eg: local static server can

```

```

4486 use a value of 10 or less). For remote servers suffering from a high latency,
4487 higher values might be needed to cover for the latency and/or the cost of
4488 picking a different server.
4489
4490 Note that this has no impact on responses which are maintained to the same
4491 server consecutively to a 401 response. They will still go to the same server
4492 even if they have to be queued.
4493
4494 See also : "option http-server-close", "option prefer-last-server", server
4495 "maxconn" and cookie persistence.
4496
4497 maxconn <n>
4498 Fix the maximum number of concurrent connections on a frontend
4499 May be used in sections: defaults | frontend | listen | backend
4500 yes | yes | yes | no
4501
4502 Arguments :
4503 <n> is the maximum number of concurrent connections the frontend will
4504 accept to serve. Excess connections will be queued by the system
4505 in the socket's listen queue and will be served once a connection
4506 closes.
4507
4508 If the system supports it, it can be useful on big sites to raise this limit
4509 very high so that haproxy manages connection queues, instead of leaving the
4510 clients with unanswered connection attempts. This value should not exceed the
4511 global maxconn. Also, keep in mind that a connection contains two buffers
4512 of 8kB each, as well as some other data resulting in about 17 kB of RAM being
4513 consumed per established connection. That means that a medium system equipped
4514 with 1GB of RAM can withstand around 40000-50000 concurrent connections if
4515 properly tuned.
4516
4517 Also, when <n> is set to large values, it is possible that the servers
4518 are not sized to accept such loads, and for this reason it is generally wise
4519 to assign them some reasonable connection limits.
4520
4521 By default, this value is set to 2000.
4522
4523 See also : "server", global section's "maxconn", "fullconn"
4524
4525 mode { tcp|http|health }
4526 Set the running mode or protocol of the instance
4527 May be used in sections: defaults | frontend | listen | backend
4528 yes | yes | yes | yes
4529
4530 Arguments :
4531 tcp The instance will work in pure TCP mode. A full-duplex connection
4532 will be established between clients and servers, and no layer 7
4533 examination will be performed. This is the default mode. It
4534 should be used for SSL, SSH, SMTP, ...
4535
4536 http The instance will work in HTTP mode. The client request will be
4537 analyzed in depth before connecting to any server. Any request
4538 which is not RFC-compliant will be rejected. Layer 7 filtering,
4539 processing and switching will be possible. This is the mode which
4540 brings HAProxy most of its value.
4541
4542 health The instance will work in "health" mode. It will just reply "OK"
4543 to incoming connections and close the connection. Alternatively,
4544 If the "htpchk" option is set, "HTTP/1.0 200 OK" will be sent
4545 instead. Nothing will be logged in either case. This mode is used
4546 to reply to external components health checks. This mode is
4547 deprecated and should not be used anymore as it is possible to do
4548 the same and even better by combining TCP or HTTP modes with the
4549 "monitor" keyword.
4550

```

When doing content switching, it is mandatory that the frontend and the backend are in the same mode (generally HTTP), otherwise the configuration will be refused.

Example :
defaults http_instances
mode http

See also : "monitor", "monitor-net"

```
monitor fail { if | unless } <condition>
```

Add a condition to report a failure to a monitor HTTP request.

May be used in sections : defaults | frontend | listen | backend
no | yes | yes | no

Arguments :

if <cond> the monitor request will fail if the condition is satisfied, and will succeed otherwise. The condition should describe a combined test which must induce a failure if all conditions are met, for instance a low number of servers both in a backend and its backup.

unless <cond> the monitor request will succeed only if the condition is satisfied, and will fail otherwise. Such a condition may be based on a test on the presence of a minimum number of active servers in a list of backends.

This statement adds a condition which can force the response to a monitor request to report a failure. By default, when an external component queries the URI dedicated to monitoring, a 200 response is returned. When one of the conditions above is met, haproxy will return 503 instead of 200. This is very useful to report a site failure to an external component which may base routing advertisements between multiple sites on the availability reported by haproxy. In this case, one would rely on an ACL involving the "nbsrv" criterion. Note that "monitor fail" only works in HTTP mode. Both status messages may be tweaked using "errorfile" or "errorloc" if needed.

Example:

```
frontend ww
mode http
acl site_dead nbsrv(dynamic) lt 2
acl site_dead nbsrv(static) lt 2
monitor-uri /site.alive
monitor fail if site_dead
```

See also : "monitor-net", "monitor-uri", "errorfile", "errorloc"

monitor-net <source>

Declare a source network which is limited to monitor requests

May be used in sections : defaults | frontend | listen | backend
yes | yes | yes | no

Arguments :

<source> is the source IPv4 address or network which will only be able to get monitor responses to any request. It can be either an IPv4 address, a host name, or an address followed by a slash ('/') followed by a mask.

In TCP mode, any connection coming from a source matching <source> will cause the connection to be immediately closed without any log. This allows another equipment to probe the port and verify that it is still listening, without forwarding the connection to a remote server.

In HTTP mode, a connection coming from a source matching <source> will be accepted, the following response will be sent without waiting for a request,

then the connection will be closed : "HTTP/1.0 200 OK". This is normally enough for any front-end HTTP probe to detect that the service is up and running without forwarding the request to a backend server. Note that this response is sent in raw format, without any transformation. This is important as it means that it will not be SSL-encrypted on SSL listeners.

Monitor requests are processed very early, just after tcp-request connection ACLs which are the only ones able to block them. These connections are short lived and never wait for any data from the client. They cannot be logged, and it is the intended purpose. They are only used to report HAProxy's health to an upper component, nothing more. Please note that "monitor fail" rules do not apply to connections intercepted by "monitor-net".

Last, please note that only one "monitor-net" statement can be specified in a frontend. If more than one is found, only the last one will be considered.

Example :

```
# addresses .252 and .253 are just probing us.
frontend ww
monitor-net 192.168.0.252/31
```

See also : "monitor fail", "monitor-uri"

monitor-uri <uri>

Intercept a URI used by external components' monitor requests

May be used in sections : defaults | frontend | listen | backend
yes | yes | yes | no

Arguments :
<uri> is the exact URI which we want to intercept to return HAProxy's health status instead of forwarding the request.

When an HTTP request referencing <uri> will be received on a frontend, HAProxy will not forward it nor log it, but instead will return either "HTTP/1.0 200 OK" or "HTTP/1.0 503 Service unavailable", depending on failure conditions defined with "monitor fail". This is normally enough for any front-end HTTP probe to detect that the service is up and running without forwarding the request to a backend server. Note that the HTTP method, the version and all headers are ignored, but the request must at least be valid at the HTTP level. This keyword may only be used with an HTTP-mode frontend.

Monitor requests are processed very early. It is not possible to block nor divert them using ACLs. They cannot be logged either, and it is the intended purpose. They are only used to report HAProxy's health to an upper component, nothing more. However, it is possible to add any number of conditions using "monitor fail" and ACLs so that the result can be adjusted to whatever check can be imagined (most often the number of available servers in a backend).

Example :

```
# Use /haproxy_test to report haproxy's status
frontend ww
mode http
monitor-uri /haproxy_test
```

See also : "monitor fail", "monitor-net"

option abortonclose

no option abortonclose

Enable or disable early dropping of aborted requests pending in queues. May be used in sections : defaults | frontend | listen | backend
yes | no | yes | yes

Arguments : none

In presence of very high loads, the servers will take some time to respond.

The per-instance connection queue will inflate, and the response time will increase respective to the size of the queue times the average per-session response time. When clients will wait for more than a few seconds, they will often hit the "STOP" button on their browser, leaving a useless request in the queue, and slowing down other users, and the servers as well, because the request will eventually be served, then aborted at the first error encountered while delivering the response.

As there is no way to distinguish between a full STOP and a simple output close on the client side, HTTP agents should be conservative and consider that the client might only have closed its output channel while waiting for the response. However, this introduces risks of congestion when lots of users do the same, and is completely useless nowadays because probably no client at all will close the session while waiting for the response. Some HTTP agents support this behaviour (Squid, Apache, HAPROXY), and others do not (TUX, most hardware-based load balancers). So the probability for a closed input channel to represent a user hitting the "STOP" button is close to 100%, and the risk of being the single component to break rare but valid traffic is extremely low, which adds to the temptation to be able to abort a session early while still not served and not pollute the servers.

In HAPROXY, the user can choose the desired behaviour using the option "abortonclose". By default (without the option) the behaviour is HTTP compliant and aborted requests will be served. But when the option is specified, a session with an incoming channel closed will be aborted while it is still possible, either pending in the queue for a connection slot, or during the connection establishment if the server has not yet acknowledged the connection request. This considerably reduces the queue size and the load on saturated servers when users are tempted to click on STOP, which in turn reduces the response time for other users.

If this option has been enabled in a "defaults" section, it can be disabled in a specific instance by prepending the "no" keyword before it.

See also : "timeout queue" and server's "maxconn" and "maxqueue" parameters

option accept-invalid-http-request

no option accept-invalid-http-request

Enable or disable relaxing of HTTP request parsing

May be used in sections : defaults | frontend | listen | backend
yes | yes | yes | no

Arguments : none

By default, HAPROXY complies with RFC7230 in terms of message parsing. This means that invalid characters in header names are not permitted and cause an error to be returned to the client. This is the desired behaviour as such forbidden characters are essentially used to build attacks exploiting server weaknesses, and bypass security filtering. Sometimes, a buggy browser or server will emit invalid header names for whatever reason (configuration, implementation) and the issue will not be immediately fixed. In such a case, it is possible to relax HAPROXY's header name parser to accept any character even if that does not make sense, by specifying this option. Similarly, the list of characters allowed to appear in a URI is well defined by RFC3986, and chars 0-31, 32 (space), 34 ('\"), 60 ('<'), 62 ('>'), 92 ('\\'), 94 ('^'), 96 ('`'), 123 ('{'), 124 ('|'), 125 ('}'), 127 (delete) and anything above are not allowed at all. HAPROXY always blocks a number of them (0..32, 127). The remaining ones are blocked by default unless this option is enabled. This option also relaxes the test on the HTTP version, it allows HTTP/0.9 requests to pass through (no version specified) and multiple digits for both the major and the minor version.

This option should never be enabled by default as it hides application bugs and open security breaches. It should only be deployed after a problem has been confirmed.

When this option is enabled, erroneous header names will still be accepted in requests, but the complete request will be captured in order to permit later analysis using the "show errors" request on the UNIX stats socket. Similarly, requests containing invalid chars in the URI part will be logged. Doing this also helps confirming that the issue has been solved.

If this option has been enabled in a "defaults" section, it can be disabled in a specific instance by prepending the "no" keyword before it.

See also : "option accept-invalid-http-response" and "show errors" on the stats socket.

option accept-invalid-http-response

no option accept-invalid-http-response

Enable or disable relaxing of HTTP response parsing

May be used in sections : defaults | frontend | listen | backend
yes | no | yes | yes

Arguments : none

By default, HAPROXY complies with RFC7230 in terms of message parsing. This means that invalid characters in header names are not permitted and cause an error to be returned to the client. This is the desired behaviour as such forbidden characters are essentially used to build attacks exploiting server weaknesses, and bypass security filtering. Sometimes, a buggy browser or server will emit invalid header names for whatever reason (configuration, implementation) and the issue will not be immediately fixed. In such a case, it is possible to relax HAPROXY's header name parser to accept any character even if that does not make sense, by specifying this option. This option also relaxes the test on the HTTP version format, it allows multiple digits for both the major and the minor version.

This option should never be enabled by default as it hides application bugs and open security breaches. It should only be deployed after a problem has been confirmed.

When this option is enabled, erroneous header names will still be accepted in responses, but the complete response will be captured in order to permit later analysis using the "show errors" request on the UNIX stats socket. Doing this also helps confirming that the issue has been solved.

If this option has been enabled in a "defaults" section, it can be disabled in a specific instance by prepending the "no" keyword before it.

See also : "option accept-invalid-http-request" and "show errors" on the stats socket.

option allbackups

no option allbackups

Use either all backup servers at a time or only the first one

May be used in sections : defaults | frontend | listen | backend
yes | no | yes | yes

Arguments : none

By default, the first operational backup server gets all traffic when normal servers are all down. Sometimes, it may be preferred to use multiple backups at once, because one will not be enough. When "option allbackups" is enabled, the load balancing will be performed among all backup servers when all normal ones are unavailable. The same load balancing algorithm will be used and the servers' weights will be respected. Thus, there will not be any priority order between the backup servers anymore.

This option is mostly used with static server farms dedicated to return a

"sorry" page when an application is completely offline.

If this option has been enabled in a "defaults" section, it can be disabled in a specific instance by prepending the "no" keyword before it.

option checkcache

no option checkcache

Analyze all server responses and block responses with cacheable cookies

May be used in sections : defaults | frontend | listen | backend

yes | no | yes | yes

Arguments : none

Some high-level frameworks set application cookies everywhere and do not always let enough control to the developer to manage how the responses should be cached. When a session cookie is returned on a cacheable object, there is a high risk of session crossing or stealing between users traversing the same caches. In some situations, it is better to block the response than to let some sensitive session information go in the wild.

The option "checkcache" enables deep inspection of all server responses for strict compliance with HTTP specification in terms of cacheability. It carefully checks "Cache-control", "Pragma" and "Set-cookie" headers in server response to check if there's a risk of caching a cookie on a client-side proxy. When this option is enabled, the only responses which can be delivered to the client are :

- all those without "Set-Cookie" header ;
- all those with a return code other than 200, 203, 206, 300, 301, 410, provided that the server has not set a "Cache-control: public" header ;
- all those that come from a POST request, provided that the server has not set a 'Cache-Control: public' header ;
- those with a 'Pragma: no-cache' header
- those with a 'Cache-control: private' header
- those with a 'Cache-control: no-store' header
- those with a 'Cache-control: max-age=0' header
- those with a 'Cache-control: s-maxage=0' header
- those with a 'Cache-control: no-cache' header
- those with a 'Cache-control: no-cache="set-cookie"' header
- those with a 'Cache-control: no-cache="set-cookie, ' header (allowing other fields after set-cookie)

If a response doesn't respect these requirements, then it will be blocked just as if it was from an "rspsdeny" filter, with an "HTTP 502 bad gateway". The session state shows "PH--" meaning that the proxy blocked the response during headers processing. Additionally, an alert will be sent in the logs so that admins are informed that there's something to be fixed.

Due to the high impact on the application, the application should be tested in depth with the option enabled before going to production. It is also a good practice to always activate it during tests, even if it is not used in production, as it will report potentially dangerous application behaviours.

If this option has been enabled in a "defaults" section, it can be disabled in a specific instance by prepending the "no" keyword before it.

option cliticpka

no option cliticpka

Enable or disable the sending of TCP keepalive packets on the client side

May be used in sections : defaults | frontend | listen | backend

yes | yes | yes | no

Arguments : none

When there is a firewall or any session-aware component between a client and a server, and when the protocol involves very long sessions with long idle

periods (eg: remote desktops), there is a risk that one of the intermediate components decides to expire a session which has remained idle for too long.

Enabling socket-level TCP keep-alives makes the system regularly send packets to the other end of the connection, leaving it active. The delay between keep-alive probes is controlled by the system only and depends both on the operating system and its tuning parameters.

It is important to understand that keep-alive packets are neither emitted nor received at the application level. It is only the network stacks which sees them. For this reason, even if one side of the proxy already uses keep-alives to maintain its connection alive, those keep-alive packets will not be forwarded to the other side of the proxy.

Please note that this has nothing to do with HTTP keep-alive.

Using option "cliticpka" enables the emission of TCP keep-alive probes on the client side of a connection, which should help when session expirations are noticed between HAProxy and a client.

If this option has been enabled in a "defaults" section, it can be disabled in a specific instance by prepending the "no" keyword before it.

See also : "option srvtcpka", "option tcpka"

option contstats

Enable continuous traffic statistics updates

May be used in sections : defaults | frontend | listen | backend

yes | yes | yes | no

Arguments : none

By default, counters used for statistics calculation are incremented only when a session finishes. It works quite well when serving small objects, but with big ones (for example large images or archives) or with A/V streaming, a graph generated from haproxy counters looks like a hedgehog. With this option enabled counters get incremented continuously, during a whole session. Recounting touches a hotpath directly so it is not enabled by default, as it has small performance impact (~0.5%).

option dontlog-normal

no option dontlog-normal

Enable or disable logging of normal, successful connections

May be used in sections : defaults | frontend | listen | backend

yes | yes | yes | no

Arguments : none

There are large sites dealing with several thousand connections per second and for which logging is a major pain. Some of them are even forced to turn logs off and cannot debug production issues. Setting this option ensures that normal connections, those which experience no error, no timeout, no retry nor redispatch, will not be logged. This leaves disk space for anomalies. In HTTP mode, the response status code is checked and return codes 5xx will still be logged.

It is strongly discouraged to use this option as most of the time, the key to complex issues is in the normal logs which will not be logged here. If you need to separate logs, see the "log-separate-errors" option instead.

See also : "log", "dontlognull", "log-separate-errors" and section 8 about logging.

option dontlognull

4941 no option dontlognull
4942 Enable or disable logging of null connections
4943 May be used in sections : defaults | frontend | listen | backend
4944 yes | yes | yes | no
4945 Arguments : none
4946
4947 In certain environments, there are components which will regularly connect to
4948 various systems to ensure that they are still alive. It can be the case from
4949 another load balancer as well as from monitoring systems. By default, even a
4950 simple port probe or scan will produce a log. If those connections pollute
4951 the logs too much, it is possible to enable option "dontlognull" to indicate
4952 that a connection on which no data has been transferred will not be logged,
4953 which typically corresponds to those probes. Note that errors will still be
4954 returned to the client and accounted for in the stats. If this is not what is
4955 desired, option http-ignore-probes can be used instead.
4956
4957 It is generally recommended not to use this option in uncontrolled
4958 environments (eg: internet), otherwise scans and other malicious activities
4959 would not be logged.
4960
4961 If this option has been enabled in a "defaults" section, it can be disabled
4962 in a specific instance by prepending the "no" keyword before it.
4963
4964 See also : "log", "http-ignore-probes", "monitor-net", "monitor-uri", and
4965 section 8 about logging.
4966
4967
4968 option forceclose
4969 no option forceclose
4970 Enable or disable active connection closing after response is transferred.
4971 May be used in sections : defaults | frontend | listen | backend
4972 yes | yes | yes | yes
4973 Arguments : none
4974
4975 Some HTTP servers do not necessarily close the connections when they receive
4976 the "Connection: close" set by "option httpclose", and if the client does not
4977 close either, then the connection remains open till the timeout expires. This
4978 causes high number of simultaneous connections on the servers and shows high
4979 global session times in the logs.
4980
4981 When this happens, it is possible to use "option forceclose". It will
4982 actively close the outgoing server channel as soon as the server has finished
4983 to respond and release some resources earlier than with "option httpclose".
4984
4985 This option may also be combined with "option http-pretend-keepalive", which
4986 will disable sending of the "Connection: close" header, but will still cause
4987 the connection to be closed once the whole response is received.
4988
4989 This option disables and replaces any previous "option httpclose", "option
4990 http-server-close", "option http-keep-alive", or "option http-tunnel".
4991
4992 If this option has been enabled in a "defaults" section, it can be disabled
4993 in a specific instance by prepending the "no" keyword before it.
4994
4995 See also : "option httpclose" and "option http-pretend-keepalive"
4996
4997
4998 option forwardfor [except <network>] [header <name>] [if-none]
4999 Enable insertion of the X-Forwarded-For header to requests sent to servers
5000 May be used in sections : defaults | frontend | listen | backend
5001 yes | yes | yes | yes
5002 Arguments :
5003 <network> is an optional argument used to disable this option for sources
5004 matching <network>
5005 <name> an optional argument to specify a different "X-Forwarded-For"

5006 header name.
5007
5008 Since HAProxy works in reverse-proxy mode, the servers see its IP address as
5009 their client address. This is sometimes annoying when the client's IP address
5010 is expected in server logs. To solve this problem, the well-known HTTP header
5011 "X-Forwarded-For" may be added by HAProxy to all requests sent to the server.
5012 This header contains a value representing the client's IP address. Since this
5013 header is always appended at the end of the existing header list, the server
5014 must be configured to always use the last occurrence of this header only. See
5015 the server's manual to find how to enable use of this standard header. Note
5016 that only the last occurrence of the header must be used, since it is really
5017 possible that the client has already brought one.
5018
5019 The keyword "header" may be used to supply a different header name to replace
5020 the default "X-Forwarded-For". This can be useful where you might already
5021 have a "X-Forwarded-For" header from a different application (eg: stunnel),
5022 and you need preserve it. Also if your backend server doesn't use the
5023 "X-Forwarded-For" header and requires different one (eg: Zeus Web Servers
5024 require "X-Cluster-Client-IP").
5025
5026 Sometimes, a same HAProxy instance may be shared between a direct client
5027 access and a reverse-proxy access (for instance when an SSL reverse-proxy is
5028 used to decrypt HTTPS traffic). It is possible to disable the addition of the
5029 header for a known source address or network by adding the "except" keyword
5030 followed by the network address. In this case, any source IP matching the
5031 network will not cause an addition of this header. Most common uses are with
5032 private networks or 127.0.0.1.
5033
5034 Alternatively, the keyword "if-none" states that the header will only be
5035 added if it is not present. This should only be used in perfectly trusted
5036 environment, as this might cause a security issue if headers reaching haproxy
5037 are under the control of the end-user.
5038
5039 This option may be specified either in the frontend or in the backend. If at
5040 least one of them uses it, the header will be added. Note that the backend's
5041 setting of the header subargument takes precedence over the frontend's if
5042 both are defined. In the case of the "if-none" argument, if at least one of
5043 the frontend or the backend does not specify it, it wants the addition to be
5044 mandatory, so it wins.
5045
5046 Examples :
5047 # Public HTTP address also used by stunnel on the same machine
5048 frontend www
5049 mode http
5050 option forwardfor except 127.0.0.1 # stunnel already adds the header
5051 # Those servers want the IP Address in X-Client
5052 backend www
5053 mode http
5054 option forwardfor header X-Client
5055
5056 See also : "option httpclose", "option http-server-close",
5057 "option forceclose", "option http-keep-alive"
5058
5059
5060 option http-buffer-request
5061 no option http-buffer-request
5062 Enable or disable waiting for whole HTTP request body before proceeding
5063 May be used in sections : defaults | frontend | listen | backend
5064 yes | yes | yes | yes
5065 Arguments : none
5066
5067 It is sometimes desirable to wait for the body of an HTTP request before
5068 taking a decision. This is what is being done by "balance url_param" for
5069 example. The first use case is to buffer requests from slow clients before
5070

connecting to the server. Another use case consists in taking the routing decision based on the request body's contents. This option placed in a frontend or backend forces the HTTP processing to wait until either the whole body is received, or the request buffer is full, or the first chunk is complete in case of chunked encoding. It can have undesired side effects with some applications abusing HTTP by expecting unbuffered transmissions between the frontend and the backend, so this should definitely not be used by default.

See also : "option http-no-delay", "timeout http-request"

option http-ignore-probes

no option http-ignore-probes

Enable or disable logging of null connections and request timeouts

May be used in sections : defaults | frontend | listen | backend
yes | yes | yes | no

Arguments : none

Recently some browsers started to implement a "pre-connect" feature consisting in speculatively connecting to some recently visited web sites just in case the user would like to visit them. This results in many connections being established to web sites, which end up in 408 Request Timeout if the timeout strikes first, or 400 Bad Request when the browser decides to close them first. These ones pollute the log and feed the error counters. There was already "option dontlognull" but it's insufficient in this case. Instead, this option does the following things :

- prevent any 400/408 message from being sent to the client if nothing was received over a connection before it was closed ;
- prevent any log from being emitted in this situation ;
- prevent any error counter from being incremented

That way the empty connection is silently ignored. Note that it is better not to use this unless it is clear that it is needed, because it will hide real problems. The most common reason for not receiving a request and seeing a 408 is due to an MTU inconsistency between the client and an intermediary element such as a VPN, which blocks too large packets. These issues are generally seen with POST requests as well as GET with large cookies. The logs are often the only way to detect them.

If this option has been enabled in a "defaults" section, it can be disabled in a specific instance by prepending the "no" keyword before it.

See also : "log", "dontlognull", "errorfile", and section 8 about logging.

option http-keep-alive

no option http-keep-alive

Enable or disable HTTP keep-alive from client to server

May be used in sections : defaults | frontend | listen | backend
yes | yes | yes | yes

Arguments : none

By default HAProxy operates in keep-alive mode with regards to persistent connections: for each connection it processes each request and response, and leaves the connection idle on both sides between the end of a response and the start of a new request. This mode may be changed by several options such as "option http-server-close", "option forceclose", "option httpclose" or "option http-tunnel". This option allows to set back the keep-alive mode, which can be useful when another mode was used in a defaults section.

Setting "option http-keep-alive" enables HTTP keep-alive mode on the client- and server- sides. This provides the lowest latency on the client side (slow network) and the fastest session reuse on the server side at the expense of maintaining idle connections to the servers. In general, it is possible

with this option to achieve approximately twice the request rate that the "http-server-close" option achieves on small objects. There are mainly two situations where this option may be useful :

- when the server is non-HTTP compliant and authenticates the connection instead of requests (eg: NTLM authentication)
- when the cost of establishing the connection to the server is significant compared to the cost of retrieving the associated object from the server.

This last case can happen when the server is a fast static server of cache. In this case, the server will need to be properly tuned to support high enough connection counts because connections will last until the client sends another request.

If the client request has to go to another backend or another server due to content switching or the load balancing algorithm, the idle connection will immediately be closed and a new one re-opened. Option "prefer-last-server" is available to try optimize server selection so that if the server currently attached to an idle connection is usable, it will be used.

In general it is preferred to use "option http-server-close" with application servers, and some static servers might benefit from "option http-keep-alive".

At the moment, logs will not indicate whether requests came from the same session or not. The accept date reported in the logs corresponds to the end of the previous request, and the request time corresponds to the time spent waiting for a new request. The keep-alive request time is still bound to the timeout defined by "timeout http-keep-alive" or "timeout http-request" if not set.

This option disables and replaces any previous "option httpclose", "option http-server-close", "option forceclose" or "option http-tunnel". When backend and frontend options differ, all of these 4 options have precedence over "option http-keep-alive".

See also : "option forceclose", "option http-server-close",

"option prefer-last-server", "option http-pretend-keepalive",
"option httpclose", and "l.i. The HTTP transaction model".

option http-no-delay

no option http-no-delay

Instruct the system to favor low interactive delays over performance in HTTP

May be used in sections : defaults | frontend | listen | backend
yes | yes | yes | yes

Arguments : none

In HTTP, each payload is unidirectional and has no notion of interactivity. Any agent is expected to queue data somewhat for a reasonably low delay.

There are some very rare server-to-server applications that abuse the HTTP protocol and expect the payload phase to be highly interactive, with many interleaved data chunks in both directions within a single request. This is absolutely not supported by the HTTP specification and will not work across most proxies or servers. When such applications attempt to do this through haproxy, it works but they will experience high delays due to the network optimizations which favor performance by instructing the system to wait for enough data to be available in order to only send full packets. Typical delays are around 200 ms per round trip. Note that this only happens with abnormal uses. Normal uses such as CONNECT requests nor WebSockets are not affected.

When "option http-no-delay" is present in either the frontend or the backend used by a connection, all such optimizations will be disabled in order to make the exchanges as fast as possible. Of course this offers no guarantee on

the functionality, as it may break at any other place. But if it works via HAProxy, it will work as fast as possible. This option should never be used by default, and should never be used at all unless such a buggy application is discovered. The impact of using this option is an increase of bandwidth usage and CPU usage, which may significantly lower performance in high latency environments.

See also : "option http-buffer-request"

option http-pretend-keepalive

no option http-pretend-keepalive

Define whether haproxy will announce keepalive to the server or not

May be used in sections : defaults | frontend | listen | backend
yes | yes | yes | yes

Arguments : none

When running with "option http-server-close" or "option forceclose", haproxy adds a "Connection: close" header to the request forwarded to the server. Unfortunately, when some servers see this header, they automatically refrain from using the chunked encoding for responses of unknown length, while this is totally unrelated. The immediate effect is that this prevents haproxy from maintaining the client connection alive. A second effect is that a client or a cache could receive an incomplete response without being aware of it, and consider the response complete.

By setting "option http-pretend-keepalive", haproxy will make the server believe it will keep the connection alive. The server will then not fall back to the abnormal undesired above. When haproxy gets the whole response, it will close the connection with the server just as it would do with the "forceclose" option. That way the client gets a normal response and the connection is correctly closed on the server side.

It is recommended not to enable this option by default, because most servers will more efficiently close the connection themselves after the last packet, and release its buffers slightly earlier. Also, the added packet on the network could slightly reduce the overall peak performance. However it is worth noting that when this option is enabled, haproxy will have slightly less work to do. So if haproxy is the bottleneck on the whole architecture, enabling this option might save a few CPU cycles.

This option may be set both in a frontend and in a backend. It is enabled if at least one of the frontend or backend holding a connection has it enabled. This option may be combined with "option httpclose", which will cause keepalive to be announced to the server and close to be announced to the client. This practice is discouraged though.

If this option has been enabled in a "defaults" section, it can be disabled in a specific instance by prepending the "no" keyword before it.

See also : "option forceclose", "option http-server-close", and
"option http-keep-alive"

option http-server-close

no option http-server-close

Enable or disable HTTP connection closing on the server side

May be used in sections : defaults | frontend | listen | backend
yes | yes | yes | yes

Arguments : none

By default HAProxy operates in keep-alive mode with regards to persistent connections: for each connection it processes each request and response, and leaves the connection idle on both sides between the end of a response and the start of a new request. This mode may be changed by several options such

as "option http-server-close", "option forceclose", "option httpclose" or "option http-tunnel". Setting "option http-server-close" enables HTTP connection-close mode on the server side while keeping the ability to support HTTP keep-alive and pipelining on the client side. This provides the lowest latency on the client side (slow network) and the fastest session reuse on the server side to save server resources, similarly to "option forceclose".

It also permits non-keepalive capable servers to be served in keep-alive mode to the clients if they conform to the requirements of RFC2616. Please note that some servers do not always conform to those requirements when they see "Connection: close" in the request. The effect will be that keep-alive will never be used. A workaround consists in enabling "option http-pretend-keepalive".

At the moment, logs will not indicate whether requests came from the same session or not. The accept date reported in the logs corresponds to the end of the previous request, and the request time corresponds to the time spent waiting for a new request. The keep-alive request time is still bound to the timeout defined by "timeout http-keep-alive" or "timeout http-request" if not set.

This option may be set both in a frontend and in a backend. It is enabled if at least one of the frontend or backend holding a connection has it enabled. It disables and replaces any previous "option httpclose", "option forceclose", "option http-tunnel" or "option http-keep-alive". Please check section 4 ("Proxies") to see how this option combines with others when frontend and backend options differ.

If this option has been enabled in a "defaults" section, it can be disabled in a specific instance by prepending the "no" keyword before it.

See also : "option forceclose", "option http-pretend-keepalive",
"option httpclose", "option http-keep-alive", and
"1.1. The HTTP transaction model".

option http-tunnel

no option http-tunnel

Disable or enable HTTP connection processing after first transaction

May be used in sections : defaults | frontend | listen | backend
yes | yes | yes | yes

Arguments : none

By default HAProxy operates in keep-alive mode with regards to persistent connections: for each connection it processes each request and response, and leaves the connection idle on both sides between the end of a response and the start of a new request. This mode may be changed by several options such as "option http-server-close", "option forceclose", "option httpclose" or "option http-tunnel".

Option "http-tunnel" disables any HTTP processing past the first request and the first response. This is the mode which was used by default in versions 1.0 to 1.5-dev21. It is the mode with the lowest processing overhead, which is normally not needed anymore unless in very specific cases such as when using an in-house protocol that looks like HTTP but is not compatible, or just to log one request per client in order to reduce log size. Note that everything which works at the HTTP level, including header parsing/addition, cookie processing or content switching will only work for the first request and will be ignored after the first response.

If this option has been enabled in a "defaults" section, it can be disabled in a specific instance by prepending the "no" keyword before it.

See also : "option forceclose", "option http-server-close",
"option httpclose", "option http-keep-alive", and
"1.1. The HTTP transaction model".

```

5331
5332
5333 option http-use-proxy-header
5334 no option http-use-proxy-header
5335 Make use of non-standard Proxy-Connection header instead of Connection
5336 May be used in sections : defaults | frontend | listen | backend
5337 yes | yes | yes | no
5338 Arguments : none
5339
5340 While RFC2616 explicitly states that HTTP/1.1 agents must use the
5341 Connection header to indicate their wish of persistent or non-persistent
5342 connections, both browsers and proxies ignore this header for proxied
5343 connections and make use of the undocumented, non-standard Proxy-Connection
5344 header instead. The issue begins when trying to put a load balancer between
5345 browsers and such proxies, because there will be a difference between what
5346 haproxy understands and what the client and the proxy agree on.
5347
5348 By setting this option in a frontend, haproxy can automatically switch to use
5349 that non-standard header if it sees proxied requests. A proxied request is
5350 defined here as one where the URI begins with neither a '/' nor a '*'. The
5351 choice of header only affects requests passing through proxies making use of
5352 one of the "httpclose", "forceclose" and "http-server-close" options. Note
5353 that this option can only be specified in a frontend and will affect the
5354 request along its whole life.
5355
5356 Also, when this option is set, a request which requires authentication will
5357 automatically switch to use proxy authentication headers if it is itself a
5358 proxied request. That makes it possible to check or enforce authentication in
5359 front of an existing proxy.
5360
5361 This option should normally never be used, except in front of a proxy.
5362
5363 See also : "option httpclose", "option forceclose" and "option
5364 http-server-close".
5365
5366
5367 option httpchk
5368 option httpchk <uri>
5369 option httpchk <method> <uri>
5370 option httpchk <method> <uri> <version>
5371 Enable HTTP protocol to check on the servers health
5372 May be used in sections : defaults | frontend | listen | backend
5373 yes | no | yes | yes
5374 Arguments :
5375 <method> is the optional HTTP method used with the requests. When not set,
5376 the "OPTIONS" method is used, as it generally requires low server
5377 processing and is easy to filter out from the logs. Any method
5378 may be used, though it is not recommended to invent non-standard
5379 ones.
5380
5381 <uri> is the URI referenced in the HTTP requests. It defaults to " / "
5382 which is accessible by default on almost any server, but may be
5383 changed to any other URI. Query strings are permitted.
5384
5385 <version> is the optional HTTP version string. It defaults to "HTTP/1.0"
5386 but some servers might behave incorrectly in HTTP 1.0, so turning
5387 it to HTTP/1.1 may sometimes help. Note that the Host field is
5388 mandatory in HTTP/1.1, and as a trick, it is possible to pass it
5389 after "\r\n" following the version string.
5390
5391 By default, server health checks only consist in trying to establish a TCP
5392 connection. When "option httpchk" is specified, a complete HTTP request is
5393 sent once the TCP connection is established, and responses 2xx and 3xx are
5394 considered valid, while all other ones indicate a server failure, including
5395 the lack of any response.

```

```

5396
5397
5398 The port and interval are specified in the server configuration.
5399
5400 This option does not necessarily require an HTTP backend, it also works with
5401 plain TCP backends. This is particularly useful to check simple scripts bound
5402 to some dedicated ports using the inetd daemon.
5403
5404 Examples :
5405 # Relay HTTPS traffic to Apache instance and check service availability
5406 # using HTTP request "OPTIONS * HTTP/1.1" on port 80.
5407 backend https_relay
5408 mode tcp
5409 option httpchk OPTIONS * HTTP/1.1\r\nHost:\ www
5410 server apache1 192.168.1.1:443 check port 80
5411
5412 See also : "option ssl-hello-chk", "option smtpchk", "option mysql-check",
5413 "option pgsql-check", "http-check" and the "check", "port" and
5414 "inter" server options.
5415
5416 option httpclose
5417 no option httpclose
5418 Enable or disable passive HTTP connection closing
5419 May be used in sections : defaults | frontend | listen | backend
5420 yes | yes | yes | yes
5421 Arguments : none
5422
5423 By default HAProxy operates in keep-alive mode with regards to persistent
5424 connections: for each connection it processes each request and response, and
5425 leaves the connection idle on both sides between the end of a response and
5426 the start of a new request. This mode may be changed by several options such
5427 as "option http-server-close", "option forceclose", "option httpclose" or
5428 "option http-tunnel".
5429
5430 If "option httpclose" is set, HAProxy will work in HTTP tunnel mode and check
5431 if a "Connection: close" header is already set in each direction, and will
5432 add one if missing. Each end should react to this by actively closing the TCP
5433 connection after each transfer, thus resulting in a switch to the HTTP close
5434 mode. Any "Connection" header different from "close" will also be removed.
5435 Note that this option is deprecated since what it does is very cheap but not
5436 reliable. Using "option http-server-close" or "option forceclose" is strongly
5437 recommended instead.
5438
5439 It seldom happens that some servers incorrectly ignore this header and do not
5440 close the connection even though they reply "Connection: close". For this
5441 reason, they are not compatible with older HTTP 1.0 browsers. If this happens
5442 it is possible to use the "option forceclose" which actively closes the
5443 request connection once the server responds. Option "forceclose" also
5444 releases the server connection earlier because it does not have to wait for
5445 the client to acknowledge it.
5446
5447 This option may be set both in a frontend and in a backend. It is enabled if
5448 at least one of the frontend or backend holding a connection has it enabled.
5449 It disables and replaces any previous "option http-server-close",
5450 "option forceclose", "option http-keep-alive" or "option http-tunnel". Please
5451 check section 4 ("Proxies") to see how this option combines with others when
5452 frontend and backend options differ.
5453
5454 If this option has been enabled in a "defaults" section, it can be disabled
5455 in a specific instance by prepending the "no" keyword before it.
5456
5457 See also : "option forceclose", "option http-server-close" and
5458 "1.1. The HTTP transaction model".
5459
5460

```

5461 option httplog [clf]
 5462 Enable logging of HTTP request, session state and timers
 5463 May be used in sections : defaults | frontend | listen | backend
 5464 yes | yes | yes | yes
 5465 Arguments :
 5466 clf if the "clf" argument is added, then the output format will be
 5467 the CLF format instead of HAProxy's default HTTP format. You can
 5468 use this when you need to feed HAProxy's logs through a specific
 5469 log analyser which only support the CLF format and which is not
 5470 extensible.

5471
 5472 By default, the log output format is very poor, as it only contains the
 5473 source and destination addresses, and the instance name. By specifying
 5474 "option httplog", each log line turns into a much richer format including,
 5475 but not limited to, the HTTP request, the connection timers, the session
 5476 status, the connections numbers, the captured headers and cookies, the
 5477 frontend, backend and server name, and of course the source address and
 5478 ports.

5479 This option may be set either in the frontend or the backend.

5480
 5481 Specifying only "option httplog" will automatically clear the 'clf' mode
 5482 if it was set by default.

5483 See also : section 8 about logging.

option http_proxy

5485 no option http_proxy

5486 Enable or disable plain HTTP proxy mode

5487 May be used in sections : defaults | frontend | listen | backend
 5488 yes | yes | yes | yes

5489 Arguments : none

5490 It sometimes happens that people need a pure HTTP proxy which understands
 5491 basic proxy requests without caching nor any fancy feature. In this case,
 5492 it may be worth setting up an HAProxy instance with the "option http_proxy"
 5493 set. In this mode, no server is declared, and the connection is forwarded to
 5494 the IP address and port found in the URL after the "http://" scheme.

5495 No host address resolution is performed, so this only works when pure IP
 5496 addresses are passed. Since this option's usage perimeter is rather limited,
 5497 it will probably be used only by experts who know they need exactly it. Last,
 5498 if the clients are susceptible of sending keep-alive requests, it will be
 5499 needed to add "option httpclose" to ensure that all requests will correctly
 5500 be analyzed.

5501 If this option has been enabled in a "defaults" section, it can be disabled
 5502 in a specific instance by prepending the "no" keyword before it.

5503 Example :

5504 # this backend understands HTTP proxy requests and forwards them directly.
 5505 backend direct_forward
 5506 option httpclose
 5507 option http_proxy

5508 See also : "option httpclose"

option independent-streams

5509 no option independent-streams

5510 Enable or disable independent timeout processing for both directions

5511 May be used in sections : defaults | frontend | listen | backend
 5512 yes | yes | yes | yes

5513 Arguments : none

5514 By default, when data is sent over a socket, both the write timeout and the
 5515 read timeout for that socket are refreshed, because we consider that there is
 5516 activity on that socket, and we have no other means of guessing if we should
 5517 receive data or not.

5518 While this default behaviour is desirable for almost all applications, there
 5519 exists a situation where it is desirable to disable it, and only refresh the
 5520 read timeout if there are incoming data. This happens on sessions with large
 5521 timeouts and low amounts of exchanged data such as telnet session. If the
 5522 server suddenly disappears, the output data accumulates in the system's
 5523 socket buffers, both timeouts are correctly refreshed, and there is no way
 5524 to know the server does not receive them, so we don't timeout. However, when
 5525 the underlying protocol always echoes sent data, it would be enough by itself
 5526 to detect the issue using the read timeout. Note that this problem does not
 5527 happen with more verbose protocols because data won't accumulate long in the
 5528 socket buffers.

5529 When this option is set on the frontend, it will disable read timeout updates
 5530 on data sent to the client. There probably is little use of this case. When
 5531 the option is set on the backend, it will disable read timeout updates on
 5532 data sent to the server. Doing so will typically break large HTTP posts from
 5533 slow lines, so use it with caution.

5534 Note: older versions used to call this setting "option independent-streams"
 5535 with a spelling mistake. This spelling is still supported but
 5536 deprecated.

5537 See also : "timeout client", "timeout server" and "timeout tunnel"

option ldap-check

5538 Use LDAPv3 health checks for server testing

5539 May be used in sections : defaults | frontend | listen | backend
 5540 yes | no | yes | yes

5541 Arguments : none

5542 It is possible to test that the server correctly talks LDAPv3 instead of just
 5543 testing that it accepts the TCP connection. When this option is set, an
 5544 LDAPv3 anonymous simple bind message is sent to the server, and the response
 5545 is analyzed to find an LDAPv3 bind response message.

5546 The server is considered valid only when the LDAP response contains success
 5547 resultCode (<http://tools.ietf.org/html/rfc4511#section-4.1.9>).

5548 Logging of bind requests is server dependent see your documentation how to
 5549 configure it.

5550 Example :

5551 option ldap-check

5552 See also : "option httpchk"

option external-check

5553 Use external processes for server health checks

5554 May be used in sections : defaults | frontend | listen | backend
 5555 yes | no | yes | yes

5556 It is possible to test the health of a server using an external command.
 5557 This is achieved by running the executable set using "external-check
 5558 command".

5559 Requires the "external-check" global to be set.

5560

5591 See also : "external-check", "external-check command", "external-check path"
 5592
 5593
 5594 option log-health-checks
 5595 no option log-health-checks
 5596 Enable or disable logging of health checks status updates
 5597 May be used in sections : defaults | frontend | listen | backend
 5598 yes | no | yes | yes
 5599 Arguments : none
 5600
 5601 By default, failed health check are logged if server is UP and successful
 5602 health checks are logged if server is DOWN, so the amount of additional
 5603 information is limited.
 5604
 5605 When this option is enabled, any change of the health check status or to
 5606 the server's health will be logged, so that it becomes possible to know
 5607 that a server was failing occasional checks before crashing, or exactly when
 5608 it failed to respond a valid HTTP status, then when the port started to
 5609 reject connections, then when the server stopped responding at all.
 5610
 5611 Note that status changes not caused by health checks (eg: enable/disable on
 5612 the CLI) are intentionally not logged by this option.
 5613
 5614 See also: "option httpchk", "option ldap-check", "option mysql-check",
 5615 "option pgsql-check", "option redis-check", "option smtpchk",
 5616 "option tcp-check", "log" and section 8 about logging.
 5617
 5618 option log-separate-errors
 5619 no option log-separate-errors
 5620 Change log level for non-completely successful connections
 5621 May be used in sections : defaults | frontend | listen | backend
 5622 yes | yes | yes | no
 5623
 5624 Arguments : none
 5625
 5626 Sometimes looking for errors in logs is not easy. This option makes haproxy
 5627 raise the level of logs containing potentially interesting information such
 5628 as errors, timeouts, retries, redispatches, or HTTP status codes 5xx. The
 5629 level changes from "info" to "err". This makes it possible to log them
 5630 separately to a different file with most syslog daemons. Be careful not to
 5631 remove them from the original file, otherwise you would lose ordering which
 5632 provides very important information.
 5633
 5634 Using this option, large sites dealing with several thousand connections per
 5635 second may log normal traffic to a rotating buffer and only archive smaller
 5636 error logs.
 5637
 5638 See also : "log", "dontlognull", "dontlog-normal" and section 8 about
 5639 logging.
 5640
 5641 option logasap
 5642 no option logasap
 5643 Enable or disable early logging of HTTP requests
 5644 May be used in sections : defaults | frontend | listen | backend
 5645 yes | yes | yes | no
 5646
 5647 Arguments : none
 5648
 5649 By default, HTTP requests are logged upon termination so that the total
 5650 transfer time and the number of bytes appear in the logs. When large objects
 5651 are being transferred, it may take a while before the request appears in the
 5652 logs. Using "option logasap" the request gets logged as soon as the server
 5653 sends the complete headers. The only missing information in the logs will be
 5654 the total number of bytes which will indicate everything except the amount
 5655 of data transferred, and the total time which will not take the transfer

5656 time into account. In such a situation, it's a good practice to capture the
 5657 "Content-Length" response header so that the logs at least indicate how many
 5658 bytes are expected to be transferred.
 5659
 5660 Examples :
 5661 listen http_proxy 0.0.0.0:80
 5662 mode http
 5663 option httplog
 5664 option logasap
 5665 log 192.168.2.200 local3
 5666
 5667 >>> Feb 6 12:14:14 localhost \
 5668 haproxy[14389]: 10.0.1.2:33317 [06/Feb/2009:12:14:14.655] http-in \
 5669 static/srv1 9/10/7/14/+30 200 +243 - - - - - 3/1/1/0 1/0 \
 5670 "GET /image.iso HTTP/1.0"
 5671
 5672 See also : "option httplog", "capture response header", and section 8 about
 5673 logging.
 5674
 5675 option mysql-check [user <username> [post-41]]
 5676 Use MySQL health checks for server testing
 5677 May be used in sections : defaults | frontend | listen | backend
 5678 yes | no | yes | yes
 5679 Arguments :
 5680 <username> This is the username which will be used when connecting to MySQL
 5681 server.
 5682
 5683 post-41 Send post v4.1 client compatible checks
 5684
 5685 If you specify a username, the check consists of sending two MySQL packet,
 5686 one Client Authentication packet, and one QUIT packet, to correctly close
 5687 MySQL session. We then parse the MySQL Handshake Initialisation packet and/or
 5688 Error packet. It is a basic but useful test which does not produce error nor
 5689 aborted connect on the server. However, it requires adding an authorization
 5690 in the MySQL table, like this :
 5691
 5692 USE mysql;
 5693 INSERT INTO user (Host,User) values ('<ip_of_haproxy>','<username>');
 5694 FLUSH PRIVILEGES;
 5695
 5696 If you don't specify a username (it is deprecated and not recommended), the
 5697 check only consists in parsing the MySQL Handshake Initialisation packet or
 5698 Error packet, we don't send anything in this mode. It was reported that it
 5699 can generate lockout if check is too frequent and/or if there is not enough
 5700 traffic. In fact, you need in this case to check MySQL "max_connect_errors"
 5701 value as if a connection is established successfully within fewer than MySQL
 5702 "max_connect_errors" attempts after a previous connection was interrupted,
 5703 the error count for the host is cleared to zero. If HAProxy's server get
 5704 blocked, the "FLUSH HOSTS" statement is the only way to unblock it.
 5705
 5706 Remember that this does not check database presence nor database consistency.
 5707 To do this, you can use an external check with xinetd for example.
 5708
 5709 The check requires MySQL >=3.22, for older version, please use TCP check.
 5710
 5711 Most often, an incoming MySQL server needs to see the client's IP address for
 5712 various purposes, including IP privilege matching and connection logging.
 5713 When possible, it is often wise to masquerade the client's IP address when
 5714 connecting to the server using the "usersrc" argument of the "source" keyword,
 5715 which requires the transparent proxy feature to be compiled in, and the MySQL
 5716 server to route the client via the machine hosting haproxy.
 5717
 5718 See also: "option httpchk"
 5719
 5720

5721 option nolinger
5722 no option nolinger
5723 Enable or disable immediate session resource cleaning after close
5724 May be used in sections: defaults | frontend | listen | yes | yes | yes
5725 Arguments : none
5726
5727 When clients or servers abort connections in a dirty way (eg: they are
5728 physically disconnected), the session timeouts triggers and the session is
5729 closed. But it will remain in FIN_WAIT1 state for some time in the system,
5730 using some resources and possibly limiting the ability to establish newer
5731 connections.
5732
5733 When this happens, it is possible to activate "option nolinger" which forces
5734 the system to immediately remove any socket's pending data on close. Thus,
5735 the session is instantly purged from the system's tables. This usually has
5736 side effects such as increased number of TCP resets due to old retransmits
5737 getting immediately rejected. Some firewalls may sometimes complain about
5738 this too.
5739
5740 For this reason, it is not recommended to use this option when not absolutely
5741 needed. You know that you need it when you have thousands of FIN_WAIT1
5742 sessions on your system (TIME_WAIT ones do not count).
5743
5744 This option may be used both on frontends and backends, depending on the side
5745 where it is required. Use it on the frontend for clients, and on the backend
5746 for servers.
5747
5748 If this option has been enabled in a "defaults" section, it can be disabled
5749 in a specific instance by prepending the "no" keyword before it.
5750
5751 option originalto [except <network>] [header <name>]
5752 Enable insertion of the X-Original-To header to requests sent to servers
5753 May be used in sections : defaults | frontend | listen | yes | yes | yes
5754 Arguments :
5755 <network> is an optional argument used to disable this option for sources
5756 matching <network>
5757 an optional argument to specify a different "X-Original-To"
5758 <name> header name.
5759
5760 Since HAProxy can work in transparent mode, every request from a client can
5761 be redirected to the proxy and HAProxy itself can proxy every request to a
5762 complex SQUID environment and the destination host from SO_ORIGINAL_DST will
5763 be lost. This is annoying when you want access rules based on destination IP
5764 addresses. To solve this problem, a new HTTP header "X-Original-To" may be
5765 added by HAProxy to all requests sent to the server. This header contains a
5766 value representing the original destination IP address. Since this must be
5767 configured to always use the last occurrence of this header only. Note that
5768 only the last occurrence of the header must be used, since it is really
5769 possible that the client has already brought one.
5770
5771 The keyword "header" may be used to supply a different header name to replace
5772 the default "X-Original-To". This can be useful where you might already
5773 have a "X-Original-To" header from a different application, and you need
5774 preserve it. Also if your backend server doesn't use the "X-Original-To"
5775 header and requires different one.
5776
5777 Sometimes, a same HAProxy instance may be shared between a direct client
5778 access and a reverse-proxy access (for instance when an SSL reverse-proxy is
5779 used to decrypt HTTPS traffic). It is possible to disable the addition of the
5780 header for a known source address or network by adding the "except" keyword
5781 followed by the network address. In this case, any source IP matching the
5782 network will not cause an addition of this header. Most common uses are with
5783

5786 private networks or 127.0.0.1.
5787
5788 This option may be specified either in the frontend or in the backend. If at
5789 least one of them uses it, the header will be added. Note that the backend's
5790 setting of the header subargument takes precedence over the frontend's if
5791 both are defined.
5792
5793 Examples :
5794 # Original Destination address
5795 frontend www
5796 mode http
5797 option originalto except 127.0.0.1
5798
5799 # Those servers want the IP Address in X-Client-Dst
5800 backend www
5801 mode http
5802 option originalto header X-Client-Dst
5803
5804 See also : "option httpclose", "option http-server-close",
5805 "option forceclose"
5806
5807 option persist
5808 no option persist
5809 Enable or disable forced persistence on down servers
5810 May be used in sections: defaults | frontend | listen | yes | yes | yes
5811 Arguments : none
5812
5813 When an HTTP request reaches a backend with a cookie which references a dead
5814 server, by default it is redispached to another server. It is possible to
5815 force the request to be sent to the dead server first using "option persist"
5816 if absolutely needed. A common use case is when servers are under extreme
5817 load and spend their time flapping. In this case, the users would still be
5818 directed to the server they opened the session on, in the hope they would be
5819 correctly served. It is recommended to use "option redispach" in conjunction
5820 with this option so that in the event it would not be possible to connect to
5821 the server at all (server definitely dead), the client would finally be
5822 redirected to another valid server.
5823
5824 If this option has been enabled in a "defaults" section, it can be disabled
5825 in a specific instance by prepending the "no" keyword before it.
5826
5827 See also : "option redispach", "retries", "force-persist"
5828
5829 option pgsq1-check [user <username>]
5830 Use PostgreSQL health checks for server testing
5831 May be used in sections : defaults | frontend | listen | yes | yes | yes
5832 Arguments :
5833 <username> This is the username which will be used when connecting to
5834 PostgreSQL server.
5835
5836 The check sends a PostgreSQL StartupMessage and waits for either
5837 Authentication request or ErrorResponse message. It is a basic but useful
5838 test which does not produce error nor aborted connect on the server.
5839 This check is identical with the "mysql-check".
5840
5841 See also: "option httpchk"
5842
5843 option prefer-last-server
5844 no option prefer-last-server
5845 Allow multiple load balanced requests to remain on the same server
5846

5851 May be used in sections: defaults | frontend | listen | backend
 5852 yes | no | yes | yes
 5853 Arguments : none
 5854

5855 When the load balancing algorithm in use is not deterministic, and a previous
 5856 request was sent to a server to which haproxy still holds a connection, it is
 5857 sometimes desirable that subsequent requests on a same session go to the same
 5858 server as much as possible. Note that this is different from persistence, as
 5859 we only indicate a preference which haproxy tries to apply without any form
 5860 of warranty. The real use is for keep-alive connections sent to servers. When
 5861 this option is used, haproxy will try to reuse the same connection that is
 5862 attached to the server instead of rebalancing to another server, causing a
 5863 close of the connection. This can make sense for static file servers. It does
 5864 not make much sense to use this in combination with hashing algorithms. Note,
 5865 haproxy already automatically tries to stick to a server which sends a 401 or
 5866 to a proxy which sends a 407 (authentication required). This is mandatory for
 5867 use with the broken NTLM authentication challenge, and significantly helps in
 5868 troubleshooting some faulty applications. Option prefer-last-server might be
 5869 desirable in these environments as well, to avoid redistributing the traffic
 5870 after every other response.

5871 If this option has been enabled in a "defaults" section, it can be disabled
 5872 in a specific instance by prepending the "no" keyword before it.

5873 See also: "option http-keep-alive"

5874 option redispatch
 5875 option redispatch <interval>
 5876 no option redispatch
 5877 Enable or disable session redistribution in case of connection failure
 5878 May be used in sections: defaults | frontend | listen | backend
 5879 yes | no | yes | yes

5880 Arguments :
 5881 <interval> The optional integer value that controls how often redispatches
 5882 occur when retrying connections. Positive value P indicates a
 5883 redispatch is desired on every Pth retry, and negative value
 5884 N indicate a redispatch is desired on the Nth retry prior to the
 5885 last retry. For example, the default of -1 preserves the
 5886 historical behaviour of redispatching on the last retry, a
 5887 positive value of 1 would indicate a redispatch on every retry,
 5888 and a positive value of 3 would indicate a redispatch on every
 5889 third retry. You can disable redispatches with a value of 0.

5890 In HTTP mode, if a server designated by a cookie is down, clients may
 5891 definitely stick to it because they cannot flush the cookie, so they will not
 5892 be able to access the service anymore.

5893 Specifying "option redispatch" will allow the proxy to break their
 5894 persistence and redistribute them to a working server.
 5895
 5896 It also allows to retry connections to another server in case of multiple
 5897 connection failures. Of course, it requires having "retries" set to a nonzero
 5898 value.

5900 This form is the preferred form, which replaces both the "redispatch" and
 5901 "redispatch" keywords.

5902 If this option has been enabled in a "defaults" section, it can be disabled
 5903 in a specific instance by prepending the "no" keyword before it.

5904 See also : "redispatch", "retries", "force-persist"

5916 option redis-check
 5917 Use redis health checks for server testing
 5918 May be used in sections : defaults | frontend | listen | backend
 5919 yes | no | yes | yes
 5920 Arguments : none
 5921

5922 It is possible to test that the server correctly talks REDIS protocol instead
 5923 of just testing that it accepts the TCP connection. When this option is set,
 5924 a PING redis command is sent to the server, and the response is analyzed to
 5925 find the "+PONG" response message.

5926 Example :
 5927 option redis-check

5928 See also : "option httpchk"

5929 option smtpchk
 5930 option smtpchk <hello> <domain>
 5931 Use SMTP health checks for server testing
 5932 May be used in sections : defaults | frontend | listen | backend
 5933 yes | no | yes | yes

5934 Arguments :
 5935 <hello> is an optional argument. It is the "hello" command to use. It can
 5936 be either "HELO" (for SMTP) or "EHLO" (for ESTMP). All other
 5937 values will be turned into the default command ("HELO").

5938 <domain> is the domain name to present to the server. It may only be
 5939 specified (and is mandatory) if the hello command has been
 5940 specified. By default, "localhost" is used.

5941 When "option smtpchk" is set, the health checks will consist in TCP
 5942 connections followed by an SMTP command. By default, this command is
 5943 "HELO localhost". The server's return code is analyzed and only return codes
 5944 starting with a "2" will be considered as valid. All other responses,
 5945 including a lack of response will constitute an error and will indicate a
 5946 dead server.

5947 This test is meant to be used with SMTP servers or relays. Depending on the
 5948 request, it is possible that some servers do not log each connection attempt,
 5949 so you may want to experiment to improve the behaviour. Using telnet on port
 5950 25 is often easier than adjusting the configuration.

5951 Most often, an incoming SMTP server needs to see the client's IP address for
 5952 various purposes, including spam filtering, anti-spoofing and logging. When
 5953 possible, it is often wise to masquerade the client's IP address when
 5954 connecting to the server using the "usesrc" argument of the "source" keyword,
 5955 which requires the transparent proxy feature to be compiled in.

5956 Example :
 5957 option smtpchk HELO mydomain.org

5958 See also : "option httpchk", "source"

5959 option socket-stats
 5960 no option socket-stats

5961 Enable or disable collecting & providing separate statistics for each socket.
 5962 May be used in sections : defaults | frontend | listen | backend
 5963 yes | yes | yes | no

5964 Arguments : none

5965

5966

5967

5968

5969

5970

5971

5972

5973

5974

5975

5976

5977

5978

5979

5980

5981 option splice-auto
5982 no option splice-response
5983 Enable or disable automatic kernel acceleration on sockets in both directions
5984 May be used in sections : defaults | frontend | listen | backend
5985 yes | yes | yes | yes
5986 Arguments : none
5987
5988 When this option is enabled either on a frontend or on a backend, haproxy
5989 will automatically evaluate the opportunity to use kernel tcp splicing to
5990 forward data between the client and the server, in either direction. Haproxy
5991 uses heuristics to estimate if kernel splicing might improve performance or
5992 not. Both directions are handled independently. Note that the heuristics used
5993 are not much aggressive in order to limit excessive use of splicing. This
5994 option requires splicing to be enabled at compile time, and may be globally
5995 disabled with the global option "nossplice". Since splice uses pipes, using it
5996 requires that there are enough spare pipes.
5997
5998 Important note: kernel-based TCP splicing is a Linux-specific feature which
5999 first appeared in kernel 2.6.25. It offers kernel-based acceleration to
6000 transfer data between sockets without copying these data to user-space, thus
6001 providing noticeable performance gains and CPU cycles savings. Since many
6002 early implementations are buggy, corrupt data and/or are inefficient, this
6003 feature is not enabled by default, and it should be used with extreme care.
6004 While it is not possible to detect the correctness of an implementation,
6005 2.6.29 is the first version offering a properly working implementation. In
6006 case of doubt, splicing may be globally disabled using the global "nossplice"
6007 keyword.
6008
6009 Example :
6010 option splice-auto
6011
6012 If this option has been enabled in a "defaults" section, it can be disabled
6013 in a specific instance by prepending the "no" keyword before it.
6014
6015 See also : "option splice-request", "option splice-response", and global
6016 options "nossplice" and "maxpipes"
6017
6018
6019 option splice-request
6020 no option splice-request
6021 Enable or disable automatic kernel acceleration on sockets for requests
6022 May be used in sections : defaults | frontend | listen | backend
6023 yes | yes | yes | yes
6024 Arguments : none
6025
6026 When this option is enabled either on a frontend or on a backend, haproxy
6027 will use kernel tcp splicing whenever possible to forward data going from
6028 the client to the server. It might still use the recv/send scheme if there
6029 are no spare pipes left. This option requires splicing to be enabled at
6030 compile time, and may be globally disabled with the global option "nossplice".
6031 Since splice uses pipes, using it requires that there are enough spare pipes.
6032
6033 Important note: see "option splice-auto" for usage limitations.
6034
6035 Example :
6036 option splice-request
6037
6038 If this option has been enabled in a "defaults" section, it can be disabled
6039 in a specific instance by prepending the "no" keyword before it.
6040
6041 See also : "option splice-auto", "option splice-response", and global options
6042 "nossplice" and "maxpipes"
6043
6044 option splice-response
6045

6046 no option splice-response
6047 Enable or disable automatic kernel acceleration on sockets for responses
6048 May be used in sections : defaults | frontend | listen | backend
6049 yes | yes | yes | yes
6050 Arguments : none
6051
6052 When this option is enabled either on a frontend or on a backend, haproxy
6053 will use kernel tcp splicing whenever possible to forward data going from
6054 the server to the client. It might still use the recv/send scheme if there
6055 are no spare pipes left. This option requires splicing to be enabled at
6056 compile time, and may be globally disabled with the global option "nossplice".
6057 Since splice uses pipes, using it requires that there are enough spare pipes.
6058
6059 Important note: see "option splice-auto" for usage limitations.
6060
6061 Example :
6062 option splice-response
6063
6064 If this option has been enabled in a "defaults" section, it can be disabled
6065 in a specific instance by prepending the "no" keyword before it.
6066
6067 See also : "option splice-auto", "option splice-request", and global options
6068 "nossplice" and "maxpipes"
6069
6070 option srvtcpka
6071 no option srvtcpka
6072 Enable or disable the sending of TCP keepalive packets on the server side
6073 May be used in sections : defaults | frontend | listen | backend
6074 yes | no | yes | yes
6075 Arguments : none
6076
6077 When there is a firewall or any session-aware component between a client and
6078 a server, and when the protocol involves very long sessions with long idle
6079 periods (eg: remote desktops), there is a risk that one of the intermediate
6080 components decides to expire a session which has remained idle for too long.
6081
6082 Enabling socket-level TCP keep-alives makes the system regularly send packets
6083 to the other end of the connection, leaving it active. The delay between
6084 keep-alive probes is controlled by the system only and depends both on the
6085 operating system and its tuning parameters.
6086
6087 It is important to understand that keep-alive packets are neither emitted nor
6088 received at the application level. It is only the network stacks which sees
6089 them. For this reason, even if one side of the proxy already uses keep-alives
6090 to maintain its connection alive, those keep-alive packets will not be
6091 forwarded to the other side of the proxy.
6092
6093 Please note that this has nothing to do with HTTP keep-alive.
6094
6095 Using option "srvtcpka" enables the emission of TCP keep-alive probes on the
6096 server side of a connection, which should help when session expirations are
6097 noticed between HAProxy and a server.
6098
6099 If this option has been enabled in a "defaults" section, it can be disabled
6100 in a specific instance by prepending the "no" keyword before it.
6101
6102 See also : "option cliitcpka", "option tcpka"
6103
6104 option ssl-hello-chk
6105 Use SSLv3 client hello health checks for server testing
6106 May be used in sections : defaults | frontend | listen | backend
6107 yes | no | yes | yes
6108 Arguments : none
6109
6110

When some SSL-based protocols are relayed in TCP mode through HAProxy, it is possible to test that the server correctly talks SSL instead of just testing that it accepts the TCP connection. When "option ssl-hello-chk" is set, pure SSLv3 client hello messages are sent once the connection is established to the server, and the response is analyzed to find an SSL server hello message. The server is considered valid only when the response contains this server hello message.

All servers tested till there correctly reply to SSLv3 client hello messages, and most servers tested do not even log the requests containing only hello messages, which is appreciable.

Note that this check works even when SSL support was not built into haproxy because it forges the SSL message. When SSL support is available, it is best to use native SSL health checks instead of this one.

See also: "option httpchk", "check-ssl"

option tcp-check

Perform health checks using tcp-check send/expect sequences

May be used in sections: defaults | frontend | listen | backend
yes | no | yes | yes

This health check method is intended to be combined with "tcp-check" command lists in order to support send/expect types of health check sequences.

TCP checks currently support 4 modes of operations :

- no "tcp-check" directive : the health check only consists in a connection attempt, which remains the default mode.
- "tcp-check send" or "tcp-check send-binary" only is mentioned : this is used to send a string along with a connection opening. With some protocols, it helps sending a "QUIT" message for example that prevents the server from logging a connection error for each health check. The check result will still be based on the ability to open the connection only.

- "tcp-check expect" only is mentioned : this is used to test a banner. The connection is opened and haproxy waits for the server to present some contents which must validate some rules. The check result will be based on the matching between the contents and the rules. This is suited for POP, IMAP, SMTP, FTP, SSH, TELNET.

- both "tcp-check send" and "tcp-check expect" are mentioned : this is used to test a hello-type protocol. Haproxy sends a message, the server responds and its response is analysed. the check result will be based on the matching between the response contents and the rules. This is often suited for protocols which require a binding or a request/response model. LDAP, MySQL, Redis are example of such protocols, though they already all have their dedicated checks with a deeper understanding of the respective protocols.
- In this mode, many questions may be sent and many answers may be analysed.

A fifth mode can be used to insert comments in different steps of the script.

For each tcp-check rule you create, you can add a "comment" directive, followed by a string. This string will be reported in the log and stderr in debug mode. It is useful to make user-friendly error reporting.

The "comment" is of course optional.

Examples :

```
# perform a POP check (analyse only server's banner)
option tcp-check
tcp-check expect string +OK POP3 ready comment POP protocol

# perform an IMAP check (analyse only server's banner)
option tcp-check
tcp-check expect string *OK IMAP4 ready comment IMAP protocol

# look for the redis master server after ensuring it speaks well
# redis protocol, then it exits properly.
# (send a command then analyse the response 3 times)
option tcp-check
tcp-check comment PING phase
tcp-check send PING\r\n
tcp-check expect string +PONG
tcp-check comment role check
tcp-check send info replication\r\n
tcp-check expect string role:master
tcp-check comment QUIT phase
tcp-check send QUIT\r\n
tcp-check expect string +OK
```

forge a HTTP request, then analyse the response
(send many headers before analyzing)

```
option tcp-check
tcp-check comment forge and send HTTP request
tcp-check send HEAD / HTTP/1.1\r\n
tcp-check send Host: www.mydomain.com\r\n
tcp-check send User-Agent: HAProxy tcpcheck\r\n
tcp-check send \r\n
tcp-check expect rstring HTTP/1..\ (2..|3..) comment check HTTP response
```

See also : "tcp-check expect", "tcp-check send"

option tcp-smart-accept

no option tcp-smart-accept

Enable or disable the saving of one ACK packet during the accept sequence

May be used in sections: defaults | frontend | listen | backend
yes | yes | yes | no

Arguments : none

When an HTTP connection request comes in, the system acknowledges it on behalf of HAProxy, then the client immediately sends its request, and the system acknowledges it too while it is notifying HAProxy about the new connection. HAProxy then reads the request and responds. This means that we have one TCP ACK sent by the system for nothing, because the request could very well be acknowledged by HAProxy when it sends its response.

For this reason, in HTTP mode, HAProxy automatically asks the system to avoid sending this useless ACK on platforms which support it (currently at least Linux). It must not cause any problem, because the system will send it anyway after 40 ms if the response takes more time than expected to come.

During complex network debugging sessions, it may be desirable to disable this optimization because delayed ACKs can make troubleshooting more complex when trying to identify where packets are delayed. It is then possible to fall back to normal behaviour by specifying "no option tcp-smart-accept".

It is also possible to force it for non-HTTP proxies by simply specifying "option tcp-smart-accept". For instance, it can make sense with some services such as SMTP where the server speaks first.

6241 It is recommended to avoid forcing this option in a defaults section. In case
6242 of doubt, consider setting it back to automatic values by prepending the
6243 "default" keyword before it, or disabling it using the "no" keyword.

6244 See also : "option tcp-smart-connect"

6245 option tcp-smart-connect

6246 no option tcp-smart-connect

6247 Enable or disable the saving of one ACK packet during the connect sequence
6248 May be used in sections : defaults | frontend | listen | backend

6249 yes | no | yes | yes

6250 Arguments : none

6251 On certain systems (at least Linux), HAProxy can ask the kernel not to
6252 immediately send an empty ACK upon a connection request, but to directly
6253 send the buffer request instead. This saves one packet on the network and
6254 thus boosts performance. It can also be useful for some servers, because they
6255 immediately get the request along with the incoming connection.

6256 This feature is enabled when "option tcp-smart-connect" is set in a backend.
6257 It is not enabled by default because it makes network troubleshooting more
6258 complex.

6259 It only makes sense to enable it with protocols where the client speaks first
6260 such as HTTP. In other situations, if there is no data to send in place of
6261 the ACK, a normal ACK is sent.

6262 If this option has been enabled in a "defaults" section, it can be disabled
6263 in a specific instance by prepending the "no" keyword before it.

6264 See also : "option tcp-smart-accept"

6265 option tcpka

6266 Enable or disable the sending of TCP keepalive packets on both sides

6267 May be used in sections : defaults | frontend | listen | backend

6268 yes | yes | yes | yes

6269 Arguments : none

6270 When there is a firewall or any session-aware component between a client and
6271 a server, and when the protocol involves very long sessions with long idle
6272 periods (eg: remote desktops), there is a risk that one of the intermediate
6273 components decides to expire a session which has remained idle for too long.

6274 Enabling socket-level TCP keep-alives makes the system regularly send packets
6275 to the other end of the connection, leaving it active. The delay between
6276 keep-alive probes is controlled by the system only and depends both on the
6277 operating system and its tuning parameters.

6278 It is important to understand that keep-alive packets are neither emitted nor
6279 received at the application level. It is only the network stacks which sees
6280 them. For this reason, even if one side of the proxy already uses keep-alives
6281 to maintain its connection alive, those keep-alive packets will not be
6282 forwarded to the other side of the proxy.

6283 Please note that this has nothing to do with HTTP keep-alive.

6284 Using option "tcpka" enables the emission of TCP keep-alive probes on both
6285 the client and server sides of a connection. Note that this is meaningful
6286 only in "defaults" or "listen" sections. If this option is used in a
6287 frontend, only the client side will get keep-alives, and if this option is
6288 used in a backend, only the server side will get keep-alives. For this
6289 reason, it is strongly recommended to explicitly use "option cliptcpka" and
6290 "option srvtcpka" when the configuration is split between frontends and
6291 backends.

6306 backends.

6307 See also : "option cliptcpka", "option srvtcpka"

6308 option tcplog

6309 Enable advanced logging of TCP connections with session state and timers

6310 May be used in sections : defaults | frontend | listen | backend

6311 yes | yes | yes | yes

6312 Arguments : none

6313 By default, the log output format is very poor, as it only contains the
6314 source and destination addresses, and the instance name. By specifying
6315 "option tcplog", each log line turns into a much richer format including, but
6316 not limited to, the connection timers, the session status, the connections
6317 numbers, the frontend, backend and server name, and of course the source
6318 address and ports. This option is useful for pure TCP proxies in order to
6319 find which of the client or server disconnects or times out. For normal HTTP
6320 proxies, it's better to use "option httplog" which is even more complete.

6321 This option may be set either in the frontend or the backend.

6322 See also : "option httplog", and section 8 about logging.

6323 option transparent

6324 no option transparent

6325 Enable client-side transparent proxying

6326 May be used in sections : defaults | frontend | listen | backend

6327 yes | no | yes | yes

6328 Arguments : none

6329 This option was introduced in order to provide layer 7 persistence to layer 3
6330 load balancers. The idea is to use the OS's ability to redirect an incoming
6331 connection for a remote address to a local process (here HAProxy), and let
6332 this process know what address was initially requested. When this option is
6333 used, sessions without cookies will be forwarded to the original destination
6334 IP address of the incoming request (which should match that of another
6335 equipment), while requests with cookies will still be forwarded to the
6336 appropriate server.

6337 Note that contrary to a common belief, this option does NOT make HAProxy
6338 present the client's IP to the server when establishing the connection.

6339 See also: the "usesrc" argument of the "source" keyword, and the
6340 "transparent" option of the "bind" keyword.

6341 external-check command <command>

6342 Executable to run when performing an external-check

6343 May be used in sections : defaults | frontend | listen | backend

6344 yes | no | yes | yes

6345 Arguments :

6346 <command> is the external command to run

6347 The arguments passed to the to the command are:

6348 <proxy_address> <proxy_port> <server_address> <server_port>

6349 The <proxy_address> and <proxy_port> are derived from the first listener
6350 that is either IPv4, IPv6 or a UNIX socket. In the case of a UNIX socket
6351 listener the proxy_address will be the path of the socket and the
6352 <proxy_port> will be the string "NOT_USED". In a backend section, it's not
6353 possible to determine a listener, and both <proxy_address> and <proxy_port>

6371 will have the string value "NOT_USED".
6372
6373 Some values are also provided through environment variables.
6374

6375 Environment variables :
6376 HAPROXY_PROXY_ADDR The first bind address if available (or empty if not
6377 applicable, for example in a "backend" section).
6378

6379 HAPROXY_PROXY_ID The backend id.
6380

6381 HAPROXY_PROXY_NAME The backend name.
6382

6383 HAPROXY_PROXY_PORT The first bind port if available (or empty if not
6384 applicable, for example in a "backend" section or
6385 for a UNIX socket).
6386

6387 HAPROXY_SERVER_ADDR The server address.
6388

6389 HAPROXY_SERVER_CURCONN The current number of connections on the server.
6390

6391 HAPROXY_SERVER_ID The server id.
6392

6393 HAPROXY_SERVER_MAXCONN The server max connections.
6394

6395 HAPROXY_SERVER_NAME The server name.
6396

6397 HAPROXY_SERVER_PORT The server port if available (or empty for a UNIX
6398 socket).
6399

6400 PATH The PATH environment variable used when executing
6401 the command may be set using "external-check path".
6402

6403 If the command executed and exits with a zero status then the check is
6404 considered to have passed, otherwise the check is considered to have
6405 failed.
6406

6407 Example :
6408 external-check command /bin/true
6409

6410 See also : "external-check", "option external-check", "external-check path"
6411
6412

6413 external-check path <path>
6414

6415 The value of the PATH environment variable used when running an external-check
6416 May be used in sections : defaults | frontend | listen | backend
6417 yes | no | yes | yes
6418

6419 Arguments :
6420 <path> is the path used when executing external command to run
6421
6422 The default path is "".

6423 Example :
6424 external-check path "/usr/bin/bin"
6425
6426 See also : "external-check", "option external-check",
6427 "external-check command"
6428
6429

6430 persist rdp-cookie
6431 persist rdp-cookie(<name>)
6432 Enable RDP cookie-based persistence
6433 May be used in sections : defaults | frontend | listen | backend
6434 yes | no | yes | yes
6435

Arguments :

6436 <name> is the optional name of the RDP cookie to check. If omitted, the
6437 default cookie name "msts" will be used. There currently is no
6438 valid reason to change this name.
6439

6440 This statement enables persistence based on an RDP cookie. The RDP cookie
6441 contains all information required to find the server in the list of known
6442 servers. So when this option is set in the backend, the request is analysed
6443 and if an RDP cookie is found, it is decoded. If it matches a known server
6444 which is still UP (or if "option persist" is set), then the connection is
6445 forwarded to this server.
6446

6447 Note that this only makes sense in a TCP backend, but for this to work, the
6448 frontend must have waited long enough to ensure that an RDP cookie is present
6449 in the request buffer. This is the same requirement as with the "rdp-cookie"
6450 load-balancing method. Thus it is highly recommended to put all statements in
6451 a single "listen" section.
6452

6453 Also, it is important to understand that the terminal server will emit this
6454 RDP cookie only if it is configured for "token redirection mode", which means
6455 that the "IP address redirection" option is disabled.
6456

6457 Example :

```
6458     listen tse-farm  
6459     bind :3389  
6460     # wait up to 5s for an RDP cookie in the request  
6461     tcp-request inspect-delay 5s  
6462     tcp-request content accept if RDP_COOKIE  
6463     # apply RDP cookie persistence  
6464     persist rdp-cookie  
6465     # if server is unknown, let's balance on the same cookie.  
6466     # alternatively, "balance leastconn" may be useful too.  
6467     balance rdp-cookie  
6468     server srv1 1.1.1.1:3389  
6469     server srv2 1.1.1.2:3389  
6470  
6471  
6472  
6473
```

6474 See also : "balance rdp-cookie", "tcp-request", the "req_rdp_cookie" ACL and
6475 the rdp_cookie pattern fetch function.
6476

6477 rate-limit sessions <rate>
6478

6479 Set a limit on the number of new sessions accepted per second on a frontend
6480 May be used in sections : defaults | frontend | listen | backend
6481 yes | yes | yes | no
6482

6483 Arguments :
6484 The <rate> parameter is an integer designating the maximum number
6485 of new sessions per second to accept on the frontend.
6486
6487 When the frontend reaches the specified number of new sessions per second, it
6488 stops accepting new connections until the rate drops below the limit again.
6489 During this time, the pending sessions will be kept in the socket's backlog
6490 (in system buffers) and haproxy will not even be aware that sessions are
6491 pending. When applying very low limit on a highly loaded service, it may make
6492 sense to increase the socket's backlog using the "backlog" keyword.
6493

6494 This feature is particularly efficient at blocking connection-based attacks
6495 or service abuse on fragile servers. Since the session rate is measured every
6496 millisecond, it is extremely accurate. Also, the limit applies immediately,
6497 no delay is needed at all to detect the threshold.
6498
6499 Example : limit the connection rate on SMTP to 10 per second max
6500 listen smtp
6501 mode tcp
6502 bind :25
6503 rate-limit sessions 10
6504 server 127.0.0.1:1025

6501
6502 **Note** : When the maximum rate is reached, the frontend's status is not changed
6503 but its sockets appear as "WAITING" in the statistics if the
6504 "socket-stats" option is enabled.
6505

6506 See also : the "backlog" keyword and the "fe_sess_rate" ACL criterion.
6507

6508 redirect location <loc> [code <code>] <option> [{if | unless} <condition>]
6509 redirect prefix <px> [code <code>] <option> [{if | unless} <condition>]
6510 redirect scheme <sch> [code <code>] <option> [{if | unless} <condition>]
6511 Return an HTTP redirection if/unless a condition is matched
6512 May be used in sections : defaults | frontend | listen | backend
6513 no | yes | yes | yes
6514
6515

6516 If/unless the condition is matched, the HTTP request will lead to a redirect
6517 response. If no condition is specified, the redirect applies unconditionally.
6518

6519 **Arguments** :
6520 <loc> With "redirect location", the exact value in <loc> is placed into
6521 the HTTP "Location" header. When used in an "http-request" rule,
6522 <loc> value follows the log-format rules and can include some
6523 dynamic values (see Custom Log Format in section 8.2.4).
6524

6525 <px> With "redirect prefix", the "Location" header is built from the
6526 concatenation of <px> and the complete URI path, including the
6527 query string, unless the "drop-query" option is specified (see
6528 below). As a special case, if <px> equals exactly "/", then
6529 nothing is inserted before the original URI. It allows one to
6530 redirect to the same URL (for instance, to insert a cookie). When
6531 used in an "http-request" rule, <px> value follows the log-format
6532 rules and can include some dynamic values (see Custom Log Format
6533 in section 8.2.4).
6534

6535 <sch> With "redirect scheme", then the "Location" header is built by
6536 concatenating <sch> with "://" then the first occurrence of the
6537 "Host" header, and then the URI path, including the query string
6538 unless the "drop-query" option is specified (see below). If no
6539 path is found or if the path is "*", then "/" is used instead. If
6540 no "Host" header is found, then an empty host component will be
6541 returned, which most recent browsers interpret as redirecting to
6542 the same host. This directive is mostly used to redirect HTTP to
6543 HTTPS. When used in an "http-request" rule, <sch> value follows
6544 the log-format rules and can include some dynamic values (see
6545 Custom Log Format in section 8.2.4).
6546

6547 <code> The code is optional. It indicates which type of HTTP redirection
6548 is desired. Only codes 301, 302, 303, 307 and 308 are supported,
6549 with 302 used by default if no code is specified. 301 means
6550 "Moved permanently", and a browser may cache the Location. 302
6551 means "Moved temporarily" and means that the browser should not
6552 cache the redirection. 303 is equivalent to 302 except that the
6553 browser will fetch the location with a GET method. 307 is just
6554 like 302 but makes it clear that the same method must be reused.
6555 Likewise, 308 replaces 301 if the same method must be used.
6556

6557 <option> There are several options which can be specified to adjust the
6558 expected behaviour of a redirection :
6559

6560 - "drop-query"
6561 When this keyword is used in a prefix-based redirection, then the
6562 location will be set without any possible query-string, which is useful
6563 for directing users to a non-secure page for instance. It has no effect
6564 with a location-type redirect.
6565

6566 - "append-slash"
6567 This keyword may be used in conjunction with "drop-query" to redirect
6568 users who use a URL not ending with a '/' to the same one with the '/'.
6569 It can be useful to ensure that search engines will only see one URL.
6570 For this, a return code 301 is preferred.
6571

6572 - "set-cookie NAME[=value]"
6573 A "Set-Cookie" header will be added with NAME (and optionally "=value")
6574 to the response. This is sometimes used to indicate that a user has
6575 been seen, for instance to protect against some types of DoS. No other
6576 cookie option is added, so the cookie will be a session cookie. Note
6577 that for a browser, a sole cookie name without an equal sign is
6578 different from a cookie with an equal sign.
6579

6580 - "clear-cookie NAME[=]"
6581 A "Set-Cookie" header will be added with NAME (and optionally "=",) but
6582 with the "Max-Age" attribute set to zero. This will tell the browser to
6583 delete this cookie. It is useful for instance on logout pages. It is
6584 important to note that clearing the cookie "NAME" will not remove a
6585 cookie set with "NAME=value". You have to clear the cookie "NAME=" for
6586 that, because the browser makes the difference.
6587

6588 Example: move the login URL only to HTTPS.

```
6589 acl clear      dst_port 80
6590 acl secure     dst_port 8080
6591 acl login_page url_beg /login
6592 acl logout     url_beg /logout
6593 acl uid_given url_reg /login?userid=[^&]+
6594 acl cookie_set hdr_sub(cookie) SEEN=1
```

```
6595 redirect prefix https://mysite.com set-cookie SEEN=1 if !cookie_set
6596 redirect prefix https://mysite.com if login_page !secure
6597 redirect prefix http://mysite.com drop-query if login_page !uid_given
6598 redirect location http://mysite.com/ if !login_page secure
6599 redirect location / clear-cookie USERID=
6600
```

6601 Example: send redirects for request for articles without a '/'.
6602

```
6603 acl missing_slash path_reg ^/article/[^\/*]*$
6604 redirect code 301 prefix / drop-query append-slash if missing_slash
6605
```

6606 Example: redirect all HTTP traffic to HTTPS when SSL is handled by haproxy.
6607 redirect scheme https if !{ ssl_fc }

6608 Example: append 'www.' prefix in front of all hosts not having it
6609 http-request redirect code 301 location www. %[hdr(host)] %[req.uri] \\
6610 unless { hdr_beg(host) -i www }

6611 See section 7 about ACL usage.

6612 redispatch (deprecated)

6613 redispatch (deprecated)
6614 Enable or disable session redistribution in case of connection failure
6615 May be used in sections: defaults | frontend | listen | backend
6616 yes | no | yes | yes
6617 Arguments : none
6618

6619 In HTTP mode, if a server designated by a cookie is down, clients may
6620 definitely stick to it because they cannot flush the cookie, so they will not
6621 be able to access the service anymore.
6622

6623 Specifying "redispatch" will allow the proxy to break their persistence and
6624 redistribute them to a working server.
6625

6626 It also allows to retry last connection to another server in case of multiple
6627
6628
6629
6630

connection failures. Of course, it requires having "retries" set to a nonzero value.

This form is deprecated, do not use it in any new configuration, use the new "option redispatch" instead.

See also : "option redispatch"

```
reqadd <string> [(if | unless) <cond>]
```

Add a header at the end of the HTTP request

May be used in sections : defaults | frontend | listen | backend
no | yes | yes | yes

Arguments :

<string> is the complete line to be added. Any space or known delimiter must be escaped using a backslash ('\'). Please refer to section 6 about HTTP header manipulation for more information.

<cond> is an optional matching condition built from ACLs. It makes it possible to ignore this rule when other conditions are not met.

A new line consisting in <string> followed by a line feed will be added after the last header of an HTTP request.

Header transformations only apply to traffic which passes through HAProxy, and not to traffic generated by HAProxy, such as health-checks or error responses.

Example : add "X-Proto: SSL" to requests coming via port 81

```
acl is-ssl dst_port 81
reqadd X-Proto:\ SSL if is-ssl
```

See also: "rspadd", "http-request", section 6 about HTTP header manipulation, and section 7 about ACLs.

```
reqallow <search> [(if | unless) <cond>]
```

reqallow <search> [(if | unless) <cond>] (ignore case)

Definitely allow an HTTP request if a line matches a regular expression

May be used in sections : defaults | frontend | listen | backend
no | yes | yes | yes

Arguments :

<search> is the regular expression applied to HTTP headers and to the request line. This is an extended regular expression. Parenthesis grouping is supported and no preliminary backslash is required.

Any space or known delimiter must be escaped using a backslash ('\'). The pattern applies to a full line at a time. The "reqallow" keyword strictly matches case while "reqiallow" ignores case.

<cond> is an optional matching condition built from ACLs. It makes it possible to ignore this rule when other conditions are not met.

A request containing any line which matches extended regular expression <search> will mark the request as allowed, even if any later test would result in a deny. The test applies both to the request line and to request headers. Keep in mind that URLs in request line are case-sensitive while header names are not.

It is easier, faster and more powerful to use ACLs to write access policies. Reqdeny, reqallow and reqpass should be avoided in new designs.

Example :

```
# allow www.* but refuse *.local
reqiallow ^Host:\ www\.
```

```
reqdeny ^Host:\ *\.
```

See also: "reqdeny", "block", "http-request", section 6 about HTTP header manipulation, and section 7 about ACLs.

```
reqdel <search> [(if | unless) <cond>]
reqidel <search> [(if | unless) <cond>] (ignore case)
```

Delete all headers matching a regular expression in an HTTP request

May be used in sections : defaults | frontend | listen | backend
no | yes | yes | yes

Arguments :

<search> is the regular expression applied to HTTP headers and to the request line. This is an extended regular expression. Parenthesis grouping is supported and no preliminary backslash is required. Any space or known delimiter must be escaped using a backslash ('\'). The pattern applies to a full line at a time. The "reqdel" keyword strictly matches case while "reqidel" ignores case.

<cond> is an optional matching condition built from ACLs. It makes it possible to ignore this rule when other conditions are not met.

Any header line matching extended regular expression <search> in the request will be completely deleted. Most common use of this is to remove unwanted and/or dangerous headers or cookies from a request before passing it to the next servers.

Header transformations only apply to traffic which passes through HAProxy, and not to traffic generated by HAProxy, such as health-checks or error responses. Keep in mind that header names are not case-sensitive.

Example :

```
# remove X-Forwarded-For header and SERVER cookie
reqidel ^X-Forwarded-For:.
reqidel ^Cookie:. *SERVER=
```

See also: "reqadd", "reqrep", "rspdel", "http-request", section 6 about HTTP header manipulation, and section 7 about ACLs.

```
reqdeny <search> [(if | unless) <cond>]
reqideny <search> [(if | unless) <cond>] (ignore case)
```

Deny an HTTP request if a line matches a regular expression

May be used in sections : defaults | frontend | listen | backend
no | yes | yes | yes

Arguments :

<search> is the regular expression applied to HTTP headers and to the request line. This is an extended regular expression. Parenthesis grouping is supported and no preliminary backslash is required. Any space or known delimiter must be escaped using a backslash ('\'). The pattern applies to a full line at a time. The "reqdeny" keyword strictly matches case while "reqideny" ignores case.

<cond> is an optional matching condition built from ACLs. It makes it possible to ignore this rule when other conditions are not met.

A request containing any line which matches extended regular expression <search> will mark the request as denied, even if any later test would result in an allow. The test applies both to the request line and to request headers. Keep in mind that URLs in request line are case-sensitive while header names are not.

A denied request will generate an "HTTP 403 forbidden" response once the complete request has been parsed. This is consistent with what is practiced

6761 using ACLs.
 6762
 6763 It is easier, faster and more powerful to use ACLs to write access policies.
 6764 Reqdeny, reqallow and reqpass should be avoided in new designs.
 6765

6766 Example :
 6767 # refuse *.local, then allow www.*
 6768 reqdeny ^Host:\.*\.local
 6769 reqallow ^Host:\ www\
 6770

6771 See also: "reqallow", "rspdeny", "block", "http-request", section 6 about
 6772 HTTP header manipulation, and section 7 about ACLs.
 6773

6774 reqpass <search> [if | unless] <cond>] (ignore case)
 6775 reqpass <search> [if | unless] <cond>] (ignore case)
 6776 Ignore any HTTP request line matching a regular expression in next rules
 6777 May be used in sections : defaults | frontend | listen | backend
 6778 no | yes | yes | yes | yes

6779 Arguments :
 6780 <search> is the regular expression applied to HTTP headers and to the
 6781 request line. This is an extended regular expression. Parenthesis
 6782 grouping is supported and no preliminary backslash is required.
 6783 Any space or known delimiter must be escaped using a backslash
 6784 (\.). The pattern applies to a full line at a time. The
 6785 "reqpass" keyword strictly matches case while "reqipass" ignores
 6786 case.
 6787

6788 <cond> is an optional matching condition built from ACLs. It makes it
 6789 possible to ignore this rule when other conditions are not met.
 6790

6791 A request containing any line which matches extended regular expression
 6792 <search> will skip next rules, without assigning any deny or allow verdict.
 6793 The test applies both to the request line and to request headers. Keep in
 6794 mind that URLs in request line are case-sensitive while header names are not.
 6795

6796 It is easier, faster and more powerful to use ACLs to write access policies.
 6797 Reqdeny, reqallow and reqpass should be avoided in new designs.
 6798

6799 Example :
 6800 # refuse *.local, then allow www.*, but ignore "www.private.local"
 6801 reqpass ^Host:\ www.private\.local
 6802 reqdeny ^Host:\.*\.local
 6803 reqallow ^Host:\ www\
 6804

6805 See also: "reqallow", "reqdeny", "block", "http-request", section 6 about
 6806 HTTP header manipulation, and section 7 about ACLs.
 6807

6808 reqrep <search> <string> [if | unless] <cond>] (ignore case)
 6809 reqrep <search> <string> [if | unless] <cond>] (ignore case)
 6810 Replace a regular expression with a string in an HTTP request line
 6811 May be used in sections : defaults | frontend | listen | backend
 6812 no | yes | yes | yes | yes
 6813 Arguments :
 6814 <search> is the regular expression applied to HTTP headers and to the
 6815 request line. This is an extended regular expression. Parenthesis
 6816 grouping is supported and no preliminary backslash is required.
 6817 Any space or known delimiter must be escaped using a backslash
 6818 (\.). The pattern applies to a full line at a time. The "reqrep"
 6819 keyword strictly matches case while "reqirep" ignores case.
 6820

6821 <string> is the complete line to be added. Any space or known delimiter
 6822 must be escaped using a backslash (\.). References to matched
 6823 pattern groups are possible using the common \N form, with N
 6824
 6825

6826 being a single digit between 0 and 9. Please refer to section
 6827 6 about HTTP header manipulation for more information.
 6828

6829 <cond> is an optional matching condition built from ACLs. It makes it
 6830 possible to ignore this rule when other conditions are not met.
 6831

6832 Any line matching extended regular expression <search> in the request (both
 6833 the request line and header lines) will be completely replaced with <string>.
 6834 Most common use of this is to rewrite URLs or domain names in "Host" headers.
 6835

6836 Header transformations only apply to traffic which passes through HAProxy,
 6837 and not to traffic generated by HAProxy, such as health-checks or error
 6838 responses. Note that for increased readability, it is suggested to add enough
 6839 spaces between the request and the response. Keep in mind that URLs in
 6840 request line are case-sensitive while header names are not.
 6841

6842 Example :
 6843 # replace "/static/" with "/" at the beginning of any request path.
 6844 reqrep ^([^\:]*)\ /static/(.*) \1 /&2
 6845 # replace "www.mydomain.com" with "www" in the host name.
 6846 reqrep ^Host:\ www.mydomain.com Host:\ www
 6847

6848 See also: "reqadd", "reqdel", "rsprep", "tune.bufsize", "http-request",
 6849 section 6 about HTTP header manipulation, and section 7 about ACLs.
 6850

6851 reqtarpit <search> [if | unless] <cond>]
 6852 reqtarpit <search> [if | unless] <cond>] (ignore case)
 6853 Tarpit an HTTP request containing a line matching a regular expression
 6854 May be used in sections : defaults | frontend | listen | backend
 6855 no | yes | yes | yes | yes
 6856 Arguments :
 6857 <search> is the regular expression applied to HTTP headers and to the
 6858 request line. This is an extended regular expression. Parenthesis
 6859 grouping is supported and no preliminary backslash is required.
 6860 Any space or known delimiter must be escaped using a backslash
 6861 (\.). The pattern applies to a full line at a time. The
 6862 "reqtarpit" keyword strictly matches case while "reqitarpit"
 6863 ignores case.
 6864

6865 <cond> is an optional matching condition built from ACLs. It makes it
 6866 possible to ignore this rule when other conditions are not met.
 6867

6868 A request containing any line which matches extended regular expression
 6869 <search> will be tarpit, which means that it will connect to nowhere, will
 6870 be kept open for a pre-defined time, then will return an HTTP error 500 so
 6871 that the attacker does not suspect it has been tarpitted. The status 500 will
 6872 be reported in the logs, but the completion flags will indicate "PT". The
 6873 delay is defined by "timeout tarpit", or "timeout connect" if the former is
 6874 not set.
 6875

6876 The goal of the tarpit is to slow down robots attacking servers with
 6877 identifiable requests. Many robots limit their outgoing number of connections
 6878 and stay connected waiting for a reply which can take several minutes to
 6879 come. Depending on the environment and attack, it may be particularly
 6880 efficient at reducing the load on the network and firewalls.
 6881

6882 Examples :
 6883 # ignore user-agents reporting any flavour of "Mozilla" or "MSIE", but
 6884 # block all others.
 6885 reqipass ^User-Agent:.*(Mozilla|MSIE)
 6886 reqitarpit ^User-Agent:
 6887
 6888 # block bad guys
 6889 acl badguys src 10.1.0.3 172.16.13.20/28
 6890

6891 reqitarpit . if badguys
6892
6893 See also: "reqallow", "reqdeny", "reqpass", "http-request", section 6
6894 about HTTP header manipulation, and section 7 about ACLs.
6895
6896 retries <value>
6897 Set the number of retries to perform on a server after a connection failure
6898 May be used in sections: defaults | frontend | listen | backend
6899 yes | no | yes | yes
6900 Arguments :
6901 <value> is the number of times a connection attempt should be retried on
6902 a server when a connection either is refused or times out. The
6903 default value is 3.
6904
6905 It is important to understand that this value applies to the number of
6906 connection attempts, not full requests. When a connection has effectively
6907 been established to a server, there will be no more retry.
6908
6909 In order to avoid immediate reconnections to a server which is restarting,
6910 a turn-around timer of min("timeout connect", one second) is applied before
6911 a retry occurs.
6912
6913 When "option redispatch" is set, the last retry may be performed on another
6914 server even if a cookie references a different server.
6915
6916 See also : "option redispatch"
6917
6918 rspadd <string> [{if | unless} <cond>]
6919 Add a header at the end of the HTTP response
6920 May be used in sections : defaults | frontend | listen | backend
6921 no | yes | yes | yes
6922 Arguments :
6923 <string> is the complete line to be added. Any space or known delimiter
6924 must be escaped using a backslash ('\'). Please refer to section
6925 6 about HTTP header manipulation for more information.
6926
6927 <cond>
6928 is an optional matching condition built from ACLs. It makes it
6929 possible to ignore this rule when other conditions are not met.
6930
6931 A new line consisting in <string> followed by a line feed will be added after
6932 the last header of an HTTP response.
6933
6934 Header transformations only apply to traffic which passes through HAProxy,
6935 and not to traffic generated by HAProxy, such as health-checks or error
6936 responses.
6937
6938 See also: "rspdel" "reqadd", "http-response", section 6 about HTTP header
6939 manipulation, and section 7 about ACLs.
6940
6941 rspdel <search> [{if | unless} <cond>]
6942 Delete all headers matching a regular expression in an HTTP response
6943 May be used in sections : defaults | frontend | listen | backend
6944 no | yes | yes | yes
6945 Arguments :
6946 <search> is the regular expression applied to HTTP headers and to the
6947 response line. This is an extended regular expression, so
6948 parenthesis grouping is supported and no preliminary backslash
6949 is required. Any space or known delimiter must be escaped using
6950 a backslash ('\'). The pattern applies to a full line at a time.
6951 The "rspdel" keyword strictly matches case while "rspidel"
6952 ignores case.
6953
6954
6955

6956 <cond> is an optional matching condition built from ACLs. It makes it
6957 possible to ignore this rule when other conditions are not met.
6958
6959 Any header line matching extended regular expression <search> in the response
6960 will be completely deleted. Most common use of this is to remove unwanted
6961 and/or sensitive headers or cookies from a response before passing it to the
6962 client.
6963
6964 Header transformations only apply to traffic which passes through HAProxy,
6965 and not to traffic generated by HAProxy, such as health-checks or error
6966 responses. Keep in mind that header names are not case-sensitive.
6967
6968 Example :
6969 # remove the Server header from responses
6970 rspidel ^Server:.*
6971
6972 See also: "rspadd", "rsprep", "reqdel", "http-response", section 6 about
6973 HTTP header manipulation, and section 7 about ACLs.
6974
6975 rspdeny <search> [{if | unless} <cond>]
6976 rspdeny <search> [{if | unless} <cond>] (ignore case)
6977 Block an HTTP response if a line matches a regular expression
6978 May be used in sections : defaults | frontend | listen | backend
6979 no | yes | yes | yes
6980 Arguments :
6981 <search> is the regular expression applied to HTTP headers and to the
6982 response line. This is an extended regular expression, so
6983 parenthesis grouping is supported and no preliminary backslash
6984 is required. Any space or known delimiter must be escaped using
6985 a backslash ('\'). The pattern applies to a full line at a time.
6986 The "rspdeny" keyword strictly matches case while "rspideny"
6987 ignores case.
6988
6989 <cond> is an optional matching condition built from ACLs. It makes it
6990 possible to ignore this rule when other conditions are not met.
6991
6992 A response containing any line which matches extended regular expression
6993 <search> will mark the request as denied. The test applies both to the
6994 response line and to response headers. Keep in mind that header names are not
6995 case-sensitive.
6996
6997 Main use of this keyword is to prevent sensitive information leak and to
6998 block the response before it reaches the client. If a response is denied, it
6999 will be replaced with an HTTP 502 error so that the client never retrieves
7000 any sensitive data.
7001
7002 It is easier, faster and more powerful to use ACLs to write access policies.
7003 Rspdeny should be avoided in new designs.
7004
7005 Example :
7006 # Ensure that no content type matching ms-word will leak
7007 rspdeny ^Content-type:.*\/ms-word
7008
7009 See also: "reqdeny", "acl", "block", "http-response", section 6 about
7010 HTTP header manipulation and section 7 about ACLs.
7011
7012 rsprep <search> <string> [{if | unless} <cond>]
7013 rsprep <search> <string> [{if | unless} <cond>] (ignore case)
7014 Replace a regular expression with a string in an HTTP response line
7015 May be used in sections : defaults | frontend | listen | backend
7016 no | yes | yes | yes
7017 Arguments :
7018
7019
7020

7021 **<search>** is the regular expression applied to HTTP headers and to the
 7022 response line. This is an extended regular expression, so
 7023 parenthesis grouping is supported and no preliminary backslash
 7024 is required. Any space or known delimiter must be escaped using
 7025 a backslash ('\'). The pattern applies to a full line at a time.
 7026 The "rsprep" keyword strictly matches case while "rsprep"
 7027 ignores case.

7028 **<string>** is the complete line to be added. Any space or known delimiter
 7029 must be escaped using a backslash ('\'). References to matched
 7030 pattern groups are possible using the common \N form, with N
 7031 being a single digit between 0 and 9. Please refer to section
 7032 6 about HTTP header manipulation for more information.
 7033

7034 **<cond>** is an optional matching condition built from ACLs. It makes it
 7035 possible to ignore this rule when other conditions are not met.
 7036

7037 Any line matching extended regular expression <search> in the response (both
 7038 the response line and header lines) will be completely replaced with
 7039 <string>. Most common use of this is to rewrite Location headers.
 7040

7041 Header transformations only apply to traffic which passes through HAProxy,
 7042 and not to traffic generated by HAProxy, such as health-checks or error
 7043 responses. Note that for increased readability, it is suggested to add enough
 7044 spaces between the request and the response. Keep in mind that header names
 7045 are not case-sensitive.
 7046

7047 **Example :**
 7048 # replace "Location: 127.0.0.1:8080" with "Location: www.mydomain.com"
 7049 rsprep ^Location:\ 127.0.0.1:8080 Location:\ www.mydomain.com
 7050

7051 See also: "rspadd", "rspsdel", "rsprep", "http-response", section 6 about
 7052 HTTP header manipulation, and section 7 about ACLs.
 7053

7054 **server <name> <address>[:<port>] [<param>]**

7055 Declare a server in a backend

7056 May be used in sections : defaults | frontend | listen | backend

7057 no | no | yes | yes

7058 **Arguments :**

7059 **<name>** is the internal name assigned to this server. This name will

7060 appear in logs and alerts. If "http-send-name-header" is

7061 set, it will be added to the request header sent to the server.
 7062

7063 **<address>** is the IPv4 or IPv6 address of the server. Alternatively, a
 7064 resolvable hostname is supported, but this name will be resolved
 7065 during start-up. Address "0.0.0.0" or "*" has a special meaning.
 7066 It indicates that the connection will be forwarded to the same IP
 7067 address as the one from the client connection. This is useful in
 7068 transparent proxy architectures where the client's connection is
 7069 intercepted and haproxy must forward to the original destination
 7070 address. This is more or less what the "transparent" keyword does
 7071 except that with a server it's possible to limit concurrency and
 7072 to report statistics. Optionally, an address family prefix may be
 7073 used before the address to force the family regardless of the
 7074 address format, which can be useful to specify a path to a unix
 7075 socket with no slash ('/'). Currently supported prefixes are :
 7076 - 'ipv4q' -> address is always IPv4
 7077 - 'ipv6q' -> address is always IPv6
 7078 - 'unixq' -> address is a path to a local unix socket
 7079 - 'abnsq' -> address is in abstract namespace (Linux only)
 7080 You may want to reference some environment variables in the
 7081 address parameter, see section 2.3 about environment
 7082 variables.
 7083

7086 **<port>** is an optional port specification. If set, all connections will
 7087 be sent to this port. If unset, the same port the client
 7088 connected to will be used. The port may also be prefixed by a "+"
 7089 or a "-". In this case, the server's port will be determined by
 7090 adding this value to the client's port.
 7091

7092 **<param>** is a list of parameters for this server. The "server" keywords
 7093 accepts an important number of options and has a complete section
 7094 dedicated to it. Please refer to section 5 for more details.
 7095

7096 **Examples :**

7097 server first 10.1.1.1:1080 cookie first check inter 1000
 7098 server second 10.1.1.2:1080 cookie second check inter 1000
 7099 server transp ipv4q
 7100 server backup "\${SRV_BACKUP}:1080" backup
 7101 server www1 dc1 "\${LAN_DC1}.101:80"
 7102 server www1_dc2 "\${LAN_DC2}.101:80"
 7103

7104 **Note:** regarding Linux's abstract namespace sockets, HAProxy uses the whole
 7105 sun_path length is used for the address length. Some other programs
 7106 such as socat use the string length only by default. Pass the option
 7107 "unix-tightsocketlen=0" to any abstract socket definition in socat to
 7108 make it compatible with HAProxy's.
 7109

7110 See also: "default-server", "http-send-name-header" and section 5 about
 7111 server options
 7112

7113 **server-state-file-name [<file>]**

7114 Set the server state file to read, load and apply to servers available in
 7115 this backend. It only applies when the directive "load-server-state-from-file"
 7116 is set to "local". When <file> is not provided or if this directive is not
 7117 set, then backend name is used. If <file> starts with a slash '/', then it is
 7118 considered as an absolute path. Otherwise, <file> is concatenated to the
 7119 global directive "server-state-file-base".
 7120

7121 **Example:** the minimal configuration below would make HAProxy look for the
 7122 state server file '/etc/haproxy/status/bk':
 7123

7124 global

7125 server-state-file-base /etc/haproxy/status

7126 backend bk

7127 load-server-state-from-file

7128 See also: "server-state-file-base", "load-server-state-from-file", and
 7129 "show servers state"

7130 **source <addr>[:<port>] [usesrc { <addr2>[:<port2>] | client | cliptip } }]**
 7131 **source <addr>[:<port>] [usesrc { <addr2>[:<port2>] | hdr_ip[<hdr>[:<occ>]] }]**
 7132 **source <addr>[:<port>] [interface <name>]**

7133 Set the source address for outgoing connections

7134 May be used in sections : defaults | frontend | listen | backend
 7135 yes | no | yes | yes

7136 **Arguments :**

7137 **<addr>** is the IPv4 address HAProxy will bind to before connecting to a
 7138 server. This address is also used as a source for health checks.
 7139

7140 The default value of 0.0.0.0 means that the system will select
 7141 the most appropriate address to reach its destination. Optionally
 7142 an address family prefix may be used before the address to force
 7143 the family regardless of the address format, which can be useful
 7144 to specify a path to a unix socket with no slash ('/'). Currently
 7145 supported prefixes are :
 7146 - 'ipv4q' -> address is always IPv4
 7147 - 'ipv6q' -> address is always IPv6
 7148

7151 - 'unix@' -> address is a path to a local unix socket
7152 - 'absn@' -> address is in abstract namespace (Linux only)
7153 You may want to reference some environment variables in the
7154 address parameter, see section 2.3 about environment variables.
7155
7156
7157 <port> is an optional port. It is normally not needed but may be useful
7158 in some very specific contexts. The default value of zero means
7159 the system will select a free port. Note that port ranges are not
7160 supported in the backend. If you want to force port ranges, you
7161 have to specify them on each "server" line.
7162
7163 <addr2> is the IP address to present to the server when connections are
7164 forwarded in full transparent proxy mode. This is currently only
7165 supported on some patched Linux kernels. When this address is
7166 specified, clients connecting to the server will be presented
7167 with this address, while health checks will still use the address
7168 <addr>.
7169
7170 <port2> is the optional port to present to the server when connections
7171 are forwarded in full transparent proxy mode (see <addr2> above).
7172 The default value of zero means the system will select a free
7173 port.
7174
7175 <hdr> is the name of a HTTP header in which to fetch the IP to bind to.
7176 This is the name of a comma-separated header list which can
7177 contain multiple IP addresses. By default, the last occurrence is
7178 used. This is designed to work with the X-Forwarded-For header
7179 and to automatically bind to the client's IP address as seen
7180 by previous proxy, typically Stunnel. In order to use another
7181 occurrence from the last one, please see the <occ> parameter
7182 below. When the header (or occurrence) is not found, no binding
7183 is performed so that the proxy's default IP address is used. Also
7184 keep in mind that the header name is case insensitive, as for any
7185 HTTP header.
7186
7187 <occ> is the occurrence number of a value to be used in a multi-value
7188 header. This is to be used in conjunction with "hdr_ip(<hdr>)",
7189 in order to specify which occurrence to use for the source IP
7190 address. Positive values indicate a position from the first
7191 occurrence, 1 being the first one. Negative values indicate
7192 positions relative to the last one, -1 being the last one. This
7193 is helpful for situations where an X-Forwarded-For header is set
7194 at the entry point of an infrastructure and must be used several
7195 proxy layers away. When this value is not specified, -1 is
7196 assumed. Passing a zero here disables the feature.

7197 <name> is an optional interface name to which to bind to for outgoing
7198 traffic. On systems supporting this features (currently, only
7199 Linux), this allows one to bind all traffic to the server to
7200 this interface even if it is not the one the system would select
7201 based on routing tables. This should be used with extreme care.
7202 Note that using this option requires root privileges.

7203 The "source" keyword is useful in complex environments where a specific
7204 address only is allowed to connect to the servers. It may be needed when a
7205 private address must be used through a public gateway for instance, and it is
7206 known that the system cannot determine the adequate source address by itself.

7207 An extension which is available on certain patched Linux kernels may be used
7208 through the "usrc" optional keyword. It makes it possible to connect to the
7209 servers with an IP address which does not belong to the system itself. This
7210 is called "full transparent proxy mode". For this to work, the destination
7211 servers have to route their traffic back to this address through the machine
7212 running HAProxy, and IP forwarding must generally be enabled on this machine.
7213
7214
7215

7216 In this "full transparent proxy" mode, it is possible to force a specific IP
7217 address to be presented to the servers. This is not much used in fact. A more
7218 common use is to tell HAProxy to present the client's IP address. For this,
7219 there are two methods :

7220
7221 - present the client's IP and port addresses. This is the most transparent
7222 mode, but it can cause problems when IP connection tracking is enabled on
7223 the machine, because a same connection may be seen twice with different
7224 states. However, this solution presents the huge advantage of not
7225 limiting the system to the 64k outgoing address+port couples, because all
7226 of the client ranges may be used.

7227
7228 - present only the client's IP address and select a spare port. This
7229 solution is still quite elegant but slightly less transparent (downstream
7230 firewalls logs will not match upstream's). It also presents the downside
7231 of limiting the number of concurrent connections to the usual 64k ports.
7232 However, since the upstream and downstream ports are different, local IP
7233 connection tracking on the machine will not be upset by the reuse of the
7234 same session.

7235 This option sets the default source for all servers in the backend. It may
7236 also be specified in a "defaults" section. Finer source address specification
7237 is possible at the server level using the "source" server option. Refer to
7238 section 5 for more information.

7239 In order to work, "us-src" requires root privileges.

7240 Examples :

7241 backend private
7242 # Connect to the servers using our 192.168.1.200 source address
7243 source 192.168.1.200

7244 backend transparent_ssl1

7245 # Connect to the SSL farm from the client's source address
7246 source 192.168.1.200 us-src clientip

7247 backend transparent_ssl2

7248 # Connect to the SSL farm from the client's source address and port
7249 # not recommended if IP conntrack is present on the local machine.
7250 source 192.168.1.200 us-src client

7251 backend transparent_ssl3

7252 # Connect to the SSL farm from the client's source address. It
7253 # is more conntrack-friendly.
7254 source 192.168.1.200 us-src clientip

7255 backend transparent_smtp

7256 # Connect to the SMTP farm from the client's source address/port
7257 # with Tproxy version 4.

7258 source 0.0.0.0 us-src clientip

7259 backend transparent_http

7260 # Connect to the servers using the client's IP as seen by previous
7261 # proxy.
7262 source 0.0.0.0 us-src hdr_ip(x-forwarded-for,-1)

7263 See also : the "source" server option in section 5, the Tproxy patches for
7264 the Linux kernel on www.balabit.com, the "bind" keyword.

7265 srvtimerout <timeout> (deprecated)

7266 Set the maximum inactivity time on the server side.

7267 May be used in sections : defaults | frontend | listen | backend
7268 yes | no | yes | yes

7269 Arguments :

7270

7281 <timeout> is the timeout value specified in milliseconds by default, but
 7282 can be in any other unit if the number is suffixed by the unit,
 7283 as explained at the top of this document.
 7284

7285 The inactivity timeout applies when the server is expected to acknowledge or
 7286 send data. In HTTP mode, this timeout is particularly important to consider
 7287 during the first phase of the server's response, when it has to send the
 7288 headers, as it directly represents the server's processing time for the
 7289 request. To find out what value to put there, it's often good to start with
 7290 what would be considered as unacceptable response times, then check the logs
 7291 to observe the response time distribution, and adjust the value accordingly.
 7292

7293 The value is specified in milliseconds by default, but can be in any other
 7294 unit if the number is suffixed by the unit, as specified at the top of this
 7295 document. In TCP mode (and to a lesser extent, in HTTP mode), it is highly
 7296 recommended that the client timeout remains equal to the server timeout in
 7297 order to avoid complex situations to debug. Whatever the expected server
 7298 response times, it is a good practice to cover at least one or several TCP
 7299 packet losses by specifying timeouts that are slightly above multiples of 3
 7300 seconds (eg: 4 or 5 seconds minimum).
 7301

7302 This parameter is specific to backends, but can be specified once for all in
 7303 "defaults" sections. This is in fact one of the easiest solutions not to
 7304 forget about it. An unspecified timeout results in an infinite timeout, which
 7305 is not recommended. Such a usage is accepted and works but reports a warning
 7306 during startup because it may results in accumulation of expired sessions in
 7307 the system if the system's timeouts are not configured either.
 7308

7309 This parameter is provided for compatibility but is currently deprecated.
 7310 Please use "timeout server" instead.
 7311

7312 See also : "timeout server", "timeout tunnel", "timeout client" and
 7313 "clitimeout".
 7314

7315 stats admin { if | unless } <cond>

7316 Enable statistics admin level if/unless a condition is matched
 7317 May be used in sections : defaults | frontend | listen | backend
 7318 no | yes | yes | yes
 7319 no | yes | yes | yes
 7320

7321 This statement enables the statistics admin level if/unless a condition is
 7322 matched.
 7323

7324 The admin level allows to enable/disable servers from the web interface. By
 7325 default, statistics page is read-only for security reasons.
 7326

7327 Note : Consider not using this feature in multi-process mode (nbproc > 1)
 7328 unless you know what you do : memory is not shared between the
 7329 processes, which can result in random behaviours.
 7330

7331 Currently, the POST request is limited to the buffer size minus the reserved
 7332 buffer space, which means that if the list of servers is too long, the
 7333 request won't be processed. It is recommended to alter few servers at a
 7334 time.
 7335

7336 Example :
 7337 # statistics admin level only for localhost
 7338 backend stats.localhost
 7339 stats enable
 7340 stats admin if LOCALHOST
 7341

7342 Example :
 7343 # statistics admin level always enabled because of the authentication
 7344 backend stats.auth
 7345 stats enable

7346 stats auth admin:AdMiN123
 7347 stats admin if TRUE
 7348

7349 Example :

7350 # statistics admin level depends on the authenticated user
 7351 userlist stats-auth
 7352 group admin users admin
 7353 user admin insecure-password AdMiN123
 7354 group readonly users haproxy
 7355 user haproxy insecure-password haproxy
 7356
 7357 backend stats_auth
 7358 stats enable
 7359 acl AUTH http_auth http_auth(stats-auth) admin
 7360 stats http-request auth unless AUTH
 7361 stats admin if AUTH_ADMIN
 7362
 7363
 7364
 7365
 7366

7367 See also : "stats enable", "stats auth", "stats http-request", "nbproc",
 7368 "bind-process", section 3.4 about userlists and section 7 about
 7369 ACL usage.
 7370

7371 stats auth <user>:<passwd>

7372 Enable statistics with authentication and grant access to an account
 7373 May be used in sections : defaults | frontend | listen | backend
 7374 defaults | yes | yes | yes | yes
 7375 yes | yes | yes | yes
 7376

7377 Arguments :
 7378 <user> is a user name to grant access to
 7379
 7380 <passwd> is the cleartext password associated to this user
 7381

7382 This statement enables statistics with default settings, and restricts access
 7383 to declared users only. It may be repeated as many times as necessary to
 7384 allow as many users as desired. When a user tries to access the statistics
 7385 without a valid account, a "401 Forbidden" response will be returned so that
 7386 the browser asks the user to provide a valid user and password. The real
 7387 which will be returned to the browser is configurable using "stats realm".
 7388

7389 Since the authentication method is HTTP Basic Authentication, the passwords
 7390 circulate in cleartext on the network. Thus, it was decided that the
 7391 configuration file would also use cleartext passwords to remind the users
 7392 that those ones should not be sensitive and not shared with any other account.
 7393

7394 It is also possible to reduce the scope of the proxies which appear in the
 7395 report using "stats scope".
 7396

7397 Though this statement alone is enough to enable statistics reporting, it is
 7398 recommended to set all other settings in order to avoid relying on default
 7399 unobvious parameters.
 7400

7401 Example :
 7402 # public access (limited to this backend only)
 7403 backend public.www
 7404 server srv1 192.168.0.1:80
 7405 stats enable
 7406 stats hide-version
 7407 stats scope .
 7408 stats uri /admin?stats
 7409 stats realm Haproxy\ Statistics
 7410 stats auth admin1:AdMiN123
 7411 stats auth admin2:AdMiN321
 7412

7413 # internal monitoring access (unlimited)
 7414 backend private_monitoring

```
7411 stats enable
7412 stats uri /admin?stats
7413 stats refresh 5s
7414
7415 See also : "stats enable", "stats realm", "stats scope", "stats uri"
7416
7417 stats enable
7418 Enable statistics reporting with default settings
7419 May be used in sections : defaults | frontend | listen | backend
7420 Arguments : none
7421
7422 This statement enables statistics reporting with default settings defined
7423 at build time. Unless stated otherwise, these settings are used :
7424 - stats uri : /haproxy?stats
7425 - stats realm : "HAProxy Statistics"
7426 - stats auth : no authentication
7427 - stats scope : no restriction
7428
7429 Though this statement alone is enough to enable statistics reporting, it is
7430 recommended to set all other settings in order to avoid relying on default
7431 unobvious parameters.
7432
7433 Example :
7434 # public access (limited to this backend only)
7435 backend public_www
7436 server srv1 192.168.0.1:80
7437 stats enable
7438 stats hide-version
7439 stats scope /admin?stats
7440 stats realm Haproxy\ Statistics
7441 stats auth admin1:AdMiN123
7442 stats auth admin2:AdMiN321
7443
7444 # internal monitoring access (unlimited)
7445 backend private_monitoring
7446 stats enable
7447 stats uri /admin?stats
7448 stats refresh 5s
7449
7450 See also : "stats auth", "stats realm", "stats uri"
7451
7452 stats hide-version
7453 Enable statistics and hide HAProxy version reporting
7454 May be used in sections : defaults | frontend | listen | backend
7455 Arguments : none
7456
7457 By default, the stats page reports some useful status information along with
7458 the statistics. Among them is HAProxy's version. However, it is generally
7459 considered dangerous to report precise version to anyone, as it can help them
7460 target known weaknesses with specific attacks. The "stats hide-version"
7461 statement removes the version from the statistics report. This is recommended
7462 for public sites or any site with a weak login/password.
7463
7464 Though this statement alone is enough to enable statistics reporting, it is
7465 recommended to set all other settings in order to avoid relying on default
7466 unobvious parameters.
7467
7468 Example :
7469 # public access (limited to this backend only)
7470 backend public_www
7471 server srv1 192.168.0.1:80
7472 stats enable
7473 stats hide-version
7474 stats scope .
```

```
7476 server srv1 192.168.0.1:80
7477 stats enable
7478 stats hide-version
7479 stats scope /admin?stats
7480 stats uri Haproxy\ Statistics
7481 stats realm admin1:AdMiN123
7482 stats auth admin2:AdMiN321
7483
7484 # internal monitoring access (unlimited)
7485 backend private_monitoring
7486 stats enable
7487 stats uri /admin?stats
7488 stats refresh 5s
7489
7490 See also : "stats auth", "stats enable", "stats realm", "stats uri"
7491
7492 stats http-request { allow | deny | auth [realm <realm>] }
7493 [ { if | unless } <condition> ]
7494
7495 Access control for statistics
7496
7497 May be used in sections: defaults | frontend | listen | backend
7498 no | yes | yes | yes
7499
7500 As "http-request", these set of options allow to fine control access to
7501 statistics. Each option may be followed by if/unless and acl.
7502 First option with matched condition (or option without condition) is final.
7503 For "deny" a 403 error will be returned, for "allow" normal processing is
7504 performed, for "auth" a 401/407 error code is returned so the client
7505 should be asked to enter a username and password.
7506
7507 There is no fixed limit to the number of http-request statements per
7508 instance.
7509
7510 See also : "http-request", section 3.4 about userlists and section 7
7511 about ACL usage.
7512
7513 stats realm <realm>
7514 Enable statistics and set authentication realm
7515 May be used in sections : defaults | frontend | listen | backend
7516 yes | yes | yes | yes
7517
7518 Arguments :
7519 <realm> is the name of the HTTP Basic Authentication realm reported to
7520 the browser. The browser uses it to display it in the pop-up
7521 inviting the user to enter a valid username and password.
7522
7523 The realm is read as a single word, so any spaces in it should be escaped
7524 using a backslash ('\').
7525
7526 This statement is useful only in conjunction with "stats auth" since it is
7527 only related to authentication.
7528
7529 Though this statement alone is enough to enable statistics reporting, it is
7530 recommended to set all other settings in order to avoid relying on default
7531 unobvious parameters.
7532
7533 Example :
7534 # public access (limited to this backend only)
7535 backend public_www
7536 server srv1 192.168.0.1:80
7537 stats enable
7538 stats hide-version
7539 stats scope .
```

```
7541 stats uri /admin?stats
7542 stats realm Haproxy\ Statistics
7543 stats auth admin1:AdMiN123
7544 stats auth admin2:AdMiN321
7545
7546 # internal monitoring access (unlimited)
7547 backend private_monitoring
7548 stats enable
7549 stats uri /admin?stats
7550 stats refresh 5s
7551
7552 See also : "stats auth", "stats enable", "stats uri"
7553
7554 stats refresh <delay>
7555 Enable statistics with automatic refresh
7556 May be used in sections : defaults | frontend | listen | backend
7557 yes | yes | yes | yes
7558 Arguments :
7559 <delay> is the suggested refresh delay, specified in seconds, which will
7560 be returned to the browser consulting the report page. While the
7561 browser is free to apply any delay, it will generally respect it
7562 and refresh the page this every seconds. The refresh interval may
7563 be specified in any other non-default time unit, by suffixing the
7564 unit after the value, as explained at the top of this document.
7565
7566 This statement is useful on monitoring displays with a permanent page
7567 reporting the load balancer's activity. When set, the HTML report page will
7568 include a link "refresh"/"stop refresh" so that the user can select whether
7569 he wants automatic refresh of the page or not.
7570
7571 Though this statement alone is enough to enable statistics reporting, it is
7572 recommended to set all other settings in order to avoid relying on default
7573 unobvious parameters.
7574
7575 Example :
7576 # public access (limited to this backend only)
7577 backend public_www
7578 server srv1 192.168.0.1:80
7579 stats enable
7580 stats hide-version
7581 stats scope . /admin?stats
7582 stats uri Haproxy\ Statistics
7583 stats realm Haproxy\ Statistics
7584 stats auth admin1:AdMiN123
7585 stats auth admin2:AdMiN321
7586
7587 # internal monitoring access (unlimited)
7588 backend private_monitoring
7589 stats enable
7590 stats uri /admin?stats
7591 stats refresh 5s
7592
7593 See also : "stats auth", "stats enable", "stats realm", "stats uri"
7594
7595 stats scope { <name> | "." }
7596 Enable statistics and limit access scope
7597 May be used in sections : defaults | frontend | listen | backend
7598 yes | yes | yes | yes
7599 Arguments :
7600 <name> is the name of a listen, frontend or backend section to be
7601 reported. The special name "." (a single dot) designates the
7602 section in which the statement appears.
7603
7604
7605
```

```
7606 When this statement is specified, only the sections enumerated with this
7607 statement will appear in the report. All other ones will be hidden. This
7608 statement may appear as many times as needed if multiple sections need to be
7609 reported. Please note that the name checking is performed as simple string
7610 comparisons, and that it is never checked that a give section name really
7611 exists.
7612
7613 Though this statement alone is enough to enable statistics reporting, it is
7614 recommended to set all other settings in order to avoid relying on default
7615 unobvious parameters.
7616
7617 Example :
7618 # public access (limited to this backend only)
7619 backend public_www
7620 server srv1 192.168.0.1:80
7621 stats enable
7622 stats hide-version
7623 stats scope . /admin?stats
7624 stats uri Haproxy\ Statistics
7625 stats realm Haproxy\ Statistics
7626 stats auth admin1:AdMiN123
7627 stats auth admin2:AdMiN321
7628
7629 # internal monitoring access (unlimited)
7630 backend private_monitoring
7631 stats enable
7632 stats uri /admin?stats
7633 stats refresh 5s
7634
7635 See also : "stats auth", "stats enable", "stats realm", "stats uri"
7636
7637 stats show-desc [ <desc> ]
7638 Enable reporting of a description on the statistics page.
7639 May be used in sections : defaults | frontend | listen | backend
7640 yes | yes | yes | yes
7641
7642 <desc> is an optional description to be reported. If unspecified, the
7643 description from global section is automatically used instead.
7644
7645 This statement is useful for users that offer shared services to their
7646 customers, where node or description should be different for each customer.
7647
7648 Though this statement alone is enough to enable statistics reporting, it is
7649 recommended to set all other settings in order to avoid relying on default
7650 unobvious parameters. By default description is not shown.
7651
7652 Example :
7653 # internal monitoring access (unlimited)
7654 backend private_monitoring
7655 stats enable
7656 stats show-desc Master node for Europe, Asia, Africa
7657 stats uri /admin?stats
7658 stats refresh 5s
7659
7660 See also: "show-node", "stats enable", "stats uri" and "description" in
7661 global section.
7662
7663 stats show-legends
7664 Enable reporting additional information on the statistics page
7665 May be used in sections : defaults | frontend | listen | backend
7666 yes | yes | yes | yes
7667 Arguments : none
7668
7669
7670
```

7671 Enable reporting additional information on the statistics page :
7672 - cap: capabilities (proxy)
7673 - mode: one of tcp, http or health (proxy)
7674 - id: SNMP ID (proxy, socket, server)
7675 - IP (socket, server)
7676 - cookie (backend, server)
7677

7678 Though this statement alone is enough to enable statistics reporting, it is
7679 recommended to set all other settings in order to avoid relying on default
7680 unobvious parameters. Default behaviour is not to show this information.
7681

7682 See also: "stats enable", "stats uri".
7683

7684 stats show-node [<name>]
7685 Enable reporting of a host name on the statistics page.
7686 May be used in sections : defaults | frontend | listen | backend
7687 yes | yes | yes | yes | yes
7688

7689 Arguments:
7690 <name> is an optional name to be reported. If unspecified, the
7691 node name from global section is automatically used instead.
7692

7693 This statement is useful for users that offer shared services to their
7694 customers, where node or description might be different on a stats page
7695 provided for each customer. Default behaviour is not to show host name.
7696

7697 Though this statement alone is enough to enable statistics reporting, it is
7698 recommended to set all other settings in order to avoid relying on default
7699 unobvious parameters.
7700

7701 Example:
7702 # internal monitoring access (unlimited)
7703 backend private_monitoring
7704 stats enable
7705 stats show-node Europe-1
7706 stats uri /admin?stats
7707 stats refresh 5s
7708

7709 See also: "show-desc", "stats enable", "stats uri", and "node" in global
7710 section.
7711

7712 stats uri <prefix>
7713 Enable statistics and define the URI prefix to access them
7714 May be used in sections : defaults | frontend | listen | backend
7715 yes | yes | yes | yes | yes
7716

7717 Arguments :
7718 <prefix> is the prefix of any URI which will be redirected to stats. This
7719 prefix may contain a question mark (?) to indicate part of a
7720 query string.
7721

7722 The statistics URI is intercepted on the relayed traffic, so it appears as a
7723 page within the normal application. It is strongly advised to ensure that the
7724 selected URI will never appear in the application, otherwise it will never be
7725 possible to reach it in the application.
7726

7727 The default URI compiled in haproxy is "/haproxy?stats", but this may be
7728 changed at build time, so it's better to always explicitly specify it here.
7729 It is generally a good idea to include a question mark in the URI so that
7730 intermediate proxies refrain from caching the results. Also, since any string
7731 beginning with the prefix will be accepted as a stats request, the question
7732 mark helps ensuring that no valid URI will begin with the same words.
7733

7734 It is sometimes very convenient to use "/" as the URI prefix, and put that
7735 statement in a "listen" instance of its own. That makes it easy to dedicate

7736 an address or a port to statistics only.
7737

7738 Though this statement alone is enough to enable statistics reporting, it is
7739 recommended to set all other settings in order to avoid relying on default
7740 unobvious parameters.
7741

7742 Example :
7743 # public access (limited to this backend only)
7744 backend public_www
7745 server srv1 192.168.0.1:80
7746 stats enable
7747 stats hide-version
7748 stats scope /admin?stats
7749 stats uri /admin?stats
7750 stats realm Haproxy\ Statistics
7751 stats auth admin1:AdMiN123
7752 stats auth admin2:AdMiN321
7753

7754 # internal monitoring access (unlimited)
7755 backend private_monitoring
7756 stats enable
7757 stats uri /admin?stats
7758 stats refresh 5s
7759

7760 See also : "stats auth", "stats enable", "stats realm"
7761

7762 stick match <pattern> [table <table>] [{if | unless} <cond>]
7763 Define a request pattern matching condition to stick a user to a server
7764 May be used in sections : defaults | frontend | listen | backend
7765 no | no | yes | yes
7766

7767 Arguments :
7768 is a sample expression rule as described in section 7.3. It
7769 <pattern> describes what elements of the incoming request or connection
7770 will be analysed in the hope to find a matching entry in a
7771 stickiness table. This rule is mandatory.
7772

7773 <table> is an optional stickiness table name. If unspecified, the same
7774 backend's table is used. A stickiness table is declared using
7775 the "stick-table" statement.
7776

7777 <cond> is an optional matching condition. It makes it possible to match
7778 on a certain criterion only when other conditions are met (or
7779 not met). For instance, it could be used to match on a source IP
7780 address except when a request passes through a known proxy, in
7781 which case we'd match on a header containing that IP address.
7782

7783 Some protocols or applications require complex stickiness rules and cannot
7784 always simply rely on cookies nor hashing. The "stick match" statement
7785 describes a rule to extract the stickiness criterion from an incoming request
7786 or connection. See section 7 for a complete list of possible patterns and
7787 transformation rules.
7788

7789 The table has to be declared using the "stick-table" statement. It must be of
7790 a type compatible with the pattern. By default it is the one which is present
7791 in the same backend. It is possible to share a table with other backends by
7792 referencing it using the "table" keyword. If another table is referenced,
7793 the server's ID inside the backends are used. By default, all server IDs
7794 start at 1 in each backend, so the server ordering is enough. But in case of
7795 doubt, it is highly recommended to force server IDs using their "id" setting.
7796

7797 It is possible to restrict the conditions where a "stick match" statement
7798 will apply, using "if" or "unless" followed by a condition. See section 7 for
7799 ACL based conditions.
7800

There is no limit on the number of "stick match" statements. The first that applies and matches will cause the request to be directed to the same server as was used for the request which created the entry. That way, multiple matches can be used as fallbacks.

The stick rules are checked after the persistence cookies, so they will not affect stickiness if a cookie has already been used to select a server. That way, it becomes very easy to insert cookies and match on IP addresses in order to maintain stickiness between HTTP and HTTPS.

Note : Consider not using this feature in multi-process mode (nbproc > 1) unless you know what you do : memory is not shared between the processes, which can result in random behaviours.

Example :

```
# forward SMTP users to the same server they just used for POP in the
# last 30 minutes
backend pop
mode tcp
balance roundrobin
stick store-request src
stick-table type ip size 200k expire 30m
server s1 192.168.1.1:110
server s2 192.168.1.1:110
```

backend smtp

```
mode tcp
balance roundrobin
stick match src table pop
server s1 192.168.1.1:25
server s2 192.168.1.1:25
```

See also : "stick-table", "stick on", "nbproc", "bind-process" and section 7 about ACLs and samples fetching.

stick on <pattern> [table <table>] [(if | unless) <condition>]

Define a request pattern to associate a user to a server

May be used in sections : defaults | frontend | listen | backend
no | no | yes | yes

Note : This form is exactly equivalent to "stick match" followed by "stick store-request", all with the same arguments. Please refer to both keywords for details. It is only provided as a convenience for writing more maintainable configurations.

Note : Consider not using this feature in multi-process mode (nbproc > 1) unless you know what you do : memory is not shared between the processes, which can result in random behaviours.

Examples :

```
# The following form ...
stick on src table pop if !localhost
```

```
# ...is strictly equivalent to this one :
stick match src table pop if !localhost
stick store-request src table pop if !localhost
```

Use cookie persistence for HTTP, and stick on source address for HTTPS as well as HTTP without cookie. Share the same table between both accesses.

```
backend http
mode http
balance roundrobin
```

```
stick on src table https
cookie SRV insert indirect nocache
server s1 192.168.1.1:80 cookie s1
server s2 192.168.1.1:80 cookie s2

backend https
mode tcp
balance roundrobin
stick-table type ip size 200k expire 30m
stick on src
server s1 192.168.1.1:443
server s2 192.168.1.1:443
```

See also : "stick match", "stick store-request", "nbproc" and "bind-process".

stick store-request <pattern> [table <table>] [(if | unless) <condition>]

Define a request pattern used to create an entry in a stickiness table

May be used in sections : defaults | frontend | listen | backend
no | no | yes | yes

Arguments :

<pattern> is a sample expression rule as described in section 7.3. It describes what elements of the incoming request or connection will be analysed, extracted and stored in the table once a server is selected.

<table> is an optional stickiness table name. If unspecified, the same backend's table is used. A stickiness table is declared using the "stick-table" statement.

<cond> is an optional storage condition. It makes it possible to store certain criteria only when some conditions are met (or not met). For instance, it could be used to store the source IP address except when the request passes through a known proxy, in which case we'd store a converted form of a header containing that IP address.

Some protocols or applications require complex stickiness rules and cannot always simply rely on cookies nor hashing. The "stick store-request" statement describes a rule to decide what to extract from the request and when to do it, in order to store it into a stickiness table for further requests to match it using the "stick match" statement. Obviously the extracted part must make sense and have a chance to be matched in a further request. Storing a client's IP address for instance often makes sense. Storing an ID found in a URL parameter also makes sense. Storing a source port will almost never make any sense because it will be randomly matched. See section 7 for a complete list of possible patterns and transformation rules.

The table has to be declared using the "stick-table" statement. It must be of a type compatible with the pattern. By default it is the one which is present in the same backend. It is possible to share a table with other backends by referencing it using the "table" keyword. If another table is referenced, the server's ID inside the backends are used. By default, all server IDs start at 1 in each backend, so the server ordering is enough. But in case of doubt, it is highly recommended to force server IDs using their "id" setting.

It is possible to restrict the conditions where a "stick store-request" statement will apply, using "if" or "unless" followed by a condition. This condition will be evaluated while parsing the request, so any criteria can be used. See section 7 for ACL based conditions.

There is no limit on the number of "stick store-request" statements, but there is a limit of 8 simultaneous stores per request or response. This makes it possible to store up to 8 criteria, all extracted from either the

request or the response, regardless of the number of rules. Only the 8 first ones which match will be kept. Using this, it is possible to feed multiple tables at once in the hope to increase the chance to recognize a user on another protocol or access method. Using multiple store-request rules with the same table is possible and may be used to find the best criterion to rely on, by arranging the rules by decreasing preference order. Only the first extracted criterion for a given table will be stored. All subsequent store-request rules referencing the same table will be skipped and their ACLs will not be evaluated.

The "store-request" rules are evaluated once the server connection has been established, so that the table will contain the real server that processed the request.

Note : Consider not using this feature in multi-process mode (nbproc > 1) unless you know what you do : memory is not shared between the processes, which can result in random behaviours.

Example :

```
# forward SMTP users to the same server they just used for POP in the
# last 30 minutes
backend pop
    mode tcp
    balance roundrobin
```

```
stick store-request src
stick-table type ip size 200k expire 30m
server s1 192.168.1.1:110
server s2 192.168.1.1:110
```

```
backend smtp
    mode tcp
    balance roundrobin
    stick match src table pop
server s1 192.168.1.1:25
server s2 192.168.1.1:25
```

See also : "stick-table", "stick on", "nbproc", "bind-process" and section 7 about ACLs and sample fetching.

```
stick-table type {ip | integer | string [len <length>] | binary [len <length>]}
size <size> [expire <expire>] [nopurge] [peers <peersect>]
[store <data_type>]*
```

Configure the stickiness table for the current section

May be used in sections : defaults | frontend | listen | backend
no | yes | yes | yes | yes

Arguments :

ip a table declared with "type ip" will only store IPv4 addresses.
This form is very compact (about 50 bytes per entry) and allows very fast entry lookup and stores with almost no overhead. This is mainly used to store client source IP addresses.

ipv6 a table declared with "type ipv6" will only store IPv6 addresses.
This form is very compact (about 60 bytes per entry) and allows very fast entry lookup and stores with almost no overhead. This is mainly used to store client source IP addresses.

integer a table declared with "type integer" will store 32bit integers which can represent a client identifier found in a request for instance.

string a table declared with "type string" will store substrings of up to <len> characters. If the string provided by the pattern extractor is larger than <len>, it will be truncated before

being stored. During matching, at most <len> characters will be compared between the string in the table and the extracted pattern. When not specified, the string is automatically limited to 32 characters.

binary a table declared with "type binary" will store binary blocks of <len> bytes. If the block provided by the pattern extractor is larger than <len>, it will be truncated before being stored. If the block provided by the sample expression is shorter than <len>, it will be padded by 0. When not specified, the block is automatically limited to 32 bytes.

<length> is the maximum number of characters that will be stored in a "string" type table (See type "string" above). Or the number of bytes of the block in "binary" type table. Be careful when changing this parameter as memory usage will proportionally increase.

<size> is the maximum number of entries that can fit in the table. This value directly impacts memory usage. Count approximately 50 bytes per entry, plus the size of a string if any. The size supports suffixes "k", "m", "g" for 2¹⁰, 2²⁰ and 2³⁰ factors.

[nopurge]

indicates that we refuse to purge older entries when the table is full. When not specified and the table is full when haproxy wants to store an entry in it, it will flush a few of the oldest entries in order to release some space for the new ones. This is most often the desired behaviour. In some specific cases, it be desirable to refuse new entries instead of purging the older ones. That may be the case when the amount of data to store is far above the hardware limits and we prefer not to offer access to new clients than to reject the ones already connected. When using this parameter, be sure to properly set the "expire" parameter (see below).

<peersect> is the name of the peers section to use for replication. Entries which associate keys to server IDs are kept synchronized with the remote peers declared in this section. All entries are also automatically learned from the local peer (old process) during a soft restart.

NOTE : each peers section may be referenced only by tables belonging to the same unique process.

<expire> defines the maximum duration of an entry in the table since it was last created, refreshed or matched. The expiration delay is defined using the standard time format, similarly as the various timeouts. The maximum duration is slightly above 24 days. See section 2.2 for more information. If this delay is not specified, the session won't automatically expire, but older entries will be removed once full. Be sure not to use the "nopurge" parameter if not expiration delay is specified.

<data_type> is used to store additional information in the stick-table. This may be used by ACLs in order to control various criteria related to the activity of the client matching the stick-table. For each item specified here, the size of each entry will be inflated so that the additional data can fit. Several data types may be stored with an entry. Multiple data types may be specified after the "store" keyword, as a comma-separated list. Alternatively, it is possible to repeat the "store" keyword followed by one or several data types. Except for the "server_id" type which is automatically detected and enabled, all data types must be explicitly declared to be stored. If an ACL references a data type which is not stored, the ACL will simply not match. Some

data types require an argument which must be passed just after the type between parenthesis. See below for the supported data types and their arguments.

The data types that can be stored with an entry are the following :

- `server_id` : this is an integer which holds the numeric ID of the server a request was assigned to. It is used by the "stick match", "stick store", and "stick on" rules. It is automatically enabled when referenced.
- `gpc0` : first General Purpose Counter. It is a positive 32-bit integer which may be used for anything. Most of the time it will be used to put a special tag on some entries, for instance to note that a specific behaviour was detected and must be known for future matches.
- `gpc0_rate(<period>)` : increment rate of the first General Purpose Counter over a period. It is a positive 32-bit integer which may be used for anything. Just like `<gpc0>`, it counts events, but instead of keeping a cumulative count, it maintains the rate at which the counter is incremented. Most of the time it will be used to measure the frequency of occurrence of certain events (eg: requests to a specific URL).
- `conn_cnt` : Connection Count. It is a positive 32-bit integer which counts the absolute number of connections received from clients which matched this entry. It does not mean the connections were accepted, just that they were received.
- `conn_cur` : Current Connections. It is a positive 32-bit integer which stores the concurrent connection counts for the entry. It is incremented once an incoming connection matches the entry, and decremented once the connection leaves. That way it is possible to know at any time the exact number of concurrent connections for an entry.
- `conn_rate(<period>)` : frequency counter (takes 12 bytes). It takes an integer parameter `<period>` which indicates in milliseconds the length of the period over which the average is measured. It reports the average incoming connection rate over that period, in connections per period. The result is an integer which can be matched using ACLs.
- `sess_cnt` : Session Count. It is a positive 32-bit integer which counts the absolute number of sessions received from clients which matched this entry. A session is a connection that was accepted by the layer 4 rules.
- `sess_rate(<period>)` : frequency counter (takes 12 bytes). It takes an integer parameter `<period>` which indicates in milliseconds the length of the period over which the average is measured. It reports the average incoming session rate over that period, in sessions per period. The result is an integer which can be matched using ACLs.
- `http_req_cnt` : HTTP request Count. It is a positive 32-bit integer which counts the absolute number of HTTP requests received from clients which matched this entry. It does not matter whether they are valid requests or not. Note that this is different from sessions when keep-alive is used on the client side.
- `http_req_rate(<period>)` : frequency counter (takes 12 bytes). It takes an integer parameter `<period>` which indicates in milliseconds the length of the period over which the average is measured. It reports the average HTTP request rate over that period, in requests per period. The result is an integer which can be matched using ACLs. It does not matter whether they are valid requests or not. Note that this is different from sessions when keep-alive is used on the client side.
- `http_err_cnt` : HTTP Error Count. It is a positive 32-bit integer which counts the absolute number of HTTP requests errors induced by clients which matched this entry. Errors are counted on invalid and truncated

requests, as well as on denied or tarptitted requests, and on failed authentications. If the server responds with 4xx, then the request is also counted as an error since it's an error triggered by the client (eg: vulnerability scan).

- `http_err_rate(<period>)` : frequency counter (takes 12 bytes). It takes an integer parameter `<period>` which indicates in milliseconds the length of the period over which the average is measured. It reports the average HTTP request error rate over that period, in requests per period (see `http_err_cnt` above for what is accounted as an error). The result is an integer which can be matched using ACLs.
- `bytes_in_cnt` : client to server byte count. It is a positive 64-bit integer which counts the cumulated amount of bytes received from clients which matched this entry. Headers are included in the count. This may be used to limit abuse of upload features on photo or video servers.
- `bytes_in_rate(<period>)` : frequency counter (takes 12 bytes). It takes an integer parameter `<period>` which indicates in milliseconds the length of the period over which the average is measured. It reports the average incoming bytes rate over that period, in bytes per period. It may be used to detect users which upload too much and too fast. Warning: with large uploads, it is possible that the amount of uploaded data will be counted once upon termination, thus causing spikes in the average transfer speed instead of having a smooth one. This may partially be smoothed with "option contstats" though this is not perfect yet. Use of `byte_in_cnt` is recommended for better fairness.
- `bytes_out_cnt` : server to client byte count. It is a positive 64-bit integer which counts the cumulated amount of bytes sent to clients which matched this entry. Headers are included in the count. This may be used to limit abuse of bots sucking the whole site.
- `bytes_out_rate(<period>)` : frequency counter (takes 12 bytes). It takes an integer parameter `<period>` which indicates in milliseconds the length of the period over which the average is measured. It reports the average outgoing bytes rate over that period, in bytes per period. It may be used to detect users which download too much and too fast. Warning: with large transfers, it is possible that the amount of transferred data will be counted once upon termination, thus causing spikes in the average transfer speed instead of having a smooth one. This may partially be smoothed with "option contstats" though this is not perfect yet. Use of `byte_out_cnt` is recommended for better fairness.

There is only one stick-table per proxy. At the moment of writing this doc, it does not seem useful to have multiple tables per proxy. If this happens to be required, simply create a dummy backend with a stick-table in it and reference it.

It is important to understand that stickiness based on learning information has some limitations, including the fact that all learned associations are lost upon restart. In general it can be good as a complement but not always as an exclusive stickiness.

Last, memory requirements may be important when storing many data types. Indeed, storing all indicators above at once in each entry requires 116 bytes per entry, or 116 MB for a 1-million entries table. This is definitely not something that can be ignored.

Example:

- # Keep track of counters of up to 1 million IP addresses over 5 minutes
- # and store a general purpose counter and the average connection rate
- # computed over a sliding window of 30 seconds.
- stick-table type ip size 1m expire 5m store gpc0,conn_rate(30s)

See also : "stick match", "stick on", "stick store-request", section 2.2 about time format and section 7 about ACLs.

stick store-response <pattern> [table <table>] [(if | unless) <condition>]
Define a request pattern used to create an entry in a stickiness table
May be used in sections : defaults | frontend | listen | backend
no | no | yes | yes

Arguments :
<pattern> is a sample expression rule as described in section 7.3. It describes what elements of the response or connection will be analysed, extracted and stored in the table once a server is selected.

<table> is an optional stickiness table name. If unspecified, the same backend's table is used. A stickiness table is declared using the "stick-table" statement.

<cond> is an optional storage condition. It makes it possible to store certain criteria only when some conditions are met (or not met). For instance, it could be used to store the SSL session ID only when the response is a SSL server hello.

Some protocols or applications require complex stickiness rules and cannot always simply rely on cookies nor hashing. The "stick store-response" statement describes a rule to decide what to extract from the response and when to do it, in order to store it into a stickiness table for further requests to match it using the "stick match" statement. Obviously the extracted part must make sense and have a chance to be matched in a further request. Storing an ID found in a header of a response makes sense. See section 7 for a complete list of possible patterns and transformation rules.

The table has to be declared using the "stick-table" statement. It must be of a type compatible with the pattern. By default it is the one which is present in the same backend. It is possible to share a table with other backends by referencing it using the "table" keyword. If another table is referenced, the server's ID inside the backends are used. By default, all server IDs start at 1 in each backend, so the server ordering is enough. But in case of doubt, it is highly recommended to force server IDs using their "id" setting. It is possible to restrict the conditions where a "stick store-response" statement will apply, using "if" or "unless" followed by a condition. This condition will be evaluated while parsing the response, so any criteria can be used. See section 7 for ACL based conditions.

There is no limit on the number of "stick store-response" statements, but there is a limit of 8 simultaneous stores per request or response. This makes it possible to store up to 8 criteria, all extracted from either the request or the response, regardless of the number of rules. Only the 8 first ones which match will be kept. Using this, it is possible to feed multiple tables at once in the hope to increase the chance to recognize a user on another protocol or access method. Using multiple store-response rules with the same table is possible and may be used to find the best criterion to rely on, by arranging the rules by decreasing preference order. Only the first extracted criterion for a given table will be stored. All subsequent store-response rules referencing the same table will be skipped and their ACLs will not be evaluated. However, even if a store-request rule references a table, a store-response rule may also use the same table. This means that each table may learn exactly one element from the request and one element from the response at once.

The table will contain the real server that processed the request.

Example :
Learn SSL session ID from both request and response and create affinity.
backend https
mode tcp
balance roundrobin
maximum SSL session ID length is 32 bytes.
stick-table type binary len 32 size 30k expire 30m

acl clienthello req_ssl_hello_type 1
acl serverhello rep_ssl_hello_type 2
use tcp content accepts to detects ssl client and server hello.
tcp-request inspect-delay 5s
tcp-request content accept if clienthello

no timeout on response inspect delay by default.
tcp-response content accept if serverhello

SSL session ID (SSLID) may be present on a client or server hello.
Its length is coded on 1 byte at offset 43 and its value starts
at offset 44.

Match and learn on request if client hello.
stick on payload_lv(43,1) if clienthello

Learn on response if server hello.
stick store-response payload_lv(43,1) if serverhello

server s1 192.168.1.1:443
server s2 192.168.1.1:443

See also : "stick-table", "stick on", and section 7 about ACLs and pattern extraction.

tcp-check connect [params*]

Opens a new connection

May be used in sections: defaults | frontend | listen | backend
no | no | yes | yes

When an application lies on more than a single TCP port or when HAProxy load-balance many services in a single backend, it makes sense to probe all the services individually before considering a server as operational.

When there are no TCP port configured on the server line neither server port directive, then the 'tcp-check connect port <port>' must be the first step of the sequence.

In a tcp-check ruleset a 'connect' is required, it is also mandatory to start the ruleset with a 'connect' rule. Purpose is to ensure admin know what they do.

Parameters :

They are optional and can be used to describe how HAProxy should open and use the TCP connection.

port if not set, check port or server port is used.
It tells HAProxy where to open the connection to.

<port> must be a valid TCP port source integer, from 1 to 65535.

send-proxy send a PROXY protocol string

ssl opens a ciphered connection

Examples:


```
8321 # check HTTP and HTTPs services on a server.
8322 # first open port 80 thanks to server line port directive, then
8323 # tcp-check opens port 443, ciphered and run a request on it:
8324 option tcp-check
8325 tcp-check connect
8326 tcp-check send GET\ /\ HTTP/1.0\r\n
8327 tcp-check send Host:\ haproxy.lwt.eu\r\n
8328 tcp-check send \r\n
8329 tcp-check expect rstring (2..[3..])
8330 tcp-check connect port 443 ssl
8331 tcp-check send GET\ /\ HTTP/1.0\r\n
8332 tcp-check send Host:\ haproxy.lwt.eu\r\n
8333 tcp-check send \r\n
8334 tcp-check expect rstring (2..[3..])
8335 server www 10.0.0.1 check port 80
8336
8337 # check both POP and IMAP from a single server:
8338 option tcp-check
8339 tcp-check connect port 110
8340 tcp-check expect string +OK\ POP3\ ready
8341 tcp-check connect port 143
8342 tcp-check expect string *\ OK\ IMAP4\ ready
8343 server mail 10.0.0.1 check
8344
8345 See also : "option tcp-check", "tcp-check send", "tcp-check expect"
8346
8347
8348 tcp-check expect [!] <match> <pattern>
8349 Specify data to be collected and analysed during a generic health check
8350 May be used in sections: defaults | frontend | listen | backend
8351 no | no | yes | yes
8352
8353 Arguments :
8354 <match> is a keyword indicating how to look for a specific pattern in the
8355 response. The keyword may be one of "string", "rstring" or
8356 binary.
8357 The keyword may be preceded by an exclamation mark ("!") to negate
8358 the match. Spaces are allowed between the exclamation mark and the
8359 keyword. See below for more details on the supported keywords.
8360
8361 <pattern> is the pattern to look for. It may be a string or a regular
8362 expression. If the pattern contains spaces, they must be escaped
8363 with the usual backslash ('\').
8364 If the match is set to binary, then the pattern must be passed as
8365 a serie of hexadecimal digits in an even number. Each sequence of
8366 two digits will represent a byte. The hexadecimal digits may be
8367 used upper or lower case.
8368
8369 The available matches are intentionally similar to their http-check cousins :
```

```
8370 string <string> : test the exact string matches in the response buffer.
8371 A health check response will be considered valid if the
8372 response's buffer contains this exact string. If the
8373 "string" keyword is prefixed with "i", then the response
8374 will be considered invalid if the body contains this
8375 string. This can be used to look for a mandatory pattern
8376 in a protocol response, or to detect a failure when a
8377 specific error appears in a protocol banner.
8378
8379 rstring <regex> : test a regular expression on the response buffer.
8380 A health check response will be considered valid if the
8381 response's buffer matches this expression. If the
8382 "rstring" keyword is prefixed with "i", then the response
8383 will be considered invalid if the body matches the
```

```
8386 expression.
8387
8388 binary <hexstring> : test the exact string in its hexadecimal form matches
8389 in the response buffer. A health check response will
8390 be considered valid if the response's buffer contains
8391 this exact hexadecimal string.
8392 Purpose is to match data on binary protocols.
8393
8394 It is important to note that the responses will be limited to a certain size
8395 defined by the global "tune.chksize" option, which defaults to 16384 bytes.
8396 Thus, too large responses may not contain the mandatory pattern when using
8397 "string", "rstring" or binary. If a large response is absolutely required, it
8398 is possible to change the default max size by setting the global variable.
8399 However, it is worth keeping in mind that parsing very large responses can
8400 waste some CPU cycles, especially when regular expressions are used, and that
8401 it is always better to focus the checks on smaller resources. Also, in its
8402 current state, the check will not find any string nor regex past a null
8403 character in the response. Similarly it is not possible to request matching
8404 the null character.
8405
8406 Examples :
8407 # perform a POP check
8408 option tcp-check
8409 tcp-check expect string +OK\ POP3\ ready
8410
8411 # perform an IMAP check
8412 option tcp-check
8413 tcp-check expect string *\ OK\ IMAP4\ ready
8414
8415 # look for the redis master server
8416 option tcp-check
8417 tcp-check send PING\r\n
8418 tcp-check expect string +PONG
8419 tcp-check send info\ replication\r\n
8420 tcp-check expect string role:master
8421 tcp-check send QUIT\r\n
8422 tcp-check expect string +OK
8423
8424 See also : "option tcp-check", "tcp-check connect", "tcp-check send",
8425 "tcp-check send-binary", "http-check expect", tune.chksize
```

```
8426 tcp-check send <data>
8427 Specify a string to be sent as a question during a generic health check
8428 May be used in sections: defaults | frontend | listen | backend
8429 no | no | yes | yes
```

```
8430 <data> : the data to be sent as a question during a generic health check
8431 session. For now, <data> must be a string.
```

```
8432 Examples :
8433 # look for the redis master server
8434 option tcp-check
8435 tcp-check send info\ replication\r\n
8436 tcp-check expect string role:master
8437
8438 See also : "option tcp-check", "tcp-check connect", "tcp-check expect",
8439 "tcp-check send-binary", tune.chksize
```

```
8440 tcp-check send-binary <hexastring>
8441 Specify an hexa digits string to be sent as a binary question during a raw
8442 tcp health check
8443 May be used in sections: defaults | frontend | listen | backend
8444
8445
```

8451 no | no | yes | yes | yes
8452
8453 <data> : the data to be sent as a question during a generic health check
8454 session. For now, <data> must be a string.
8455 <hexastring> : test the exact string in its hexadecimal form matches in the
8456 response buffer. A health check response will be considered
8457 valid if the response's buffer contains this exact
8458 hexadecimal string.
8459 Purpose is to send binary data to ask on binary protocols.

Examples :

8460 # redis check in binary
8461 option tcp-check
8462 tcp-check send-binary 50494e470d0a # PING\r\n
8463 tcp-check expect binary 2b504f4e47 # +PONG
8464
8465
8466
8467

8468 See also : "option tcp-check", "tcp-check connect", "tcp-check expect",
8469 "tcp-check send", tune.chksize
8470
8471

tcp-request connection <action> [(if | unless) <condition>]

8472 Perform an action on an incoming connection depending on a layer 4 condition
8473 May be used in sections : defaults | frontend | listen | backend
8474 no | yes | yes | no

8475 Arguments : defines the action to perform if the condition applies. See
8476 below.
8477
8478

8479 <condition> is a standard layer4-only ACL-based condition (see section 7).
8480
8481

8482 Immediately after acceptance of a new incoming connection, it is possible to
8483 evaluate some conditions to decide whether this connection must be accepted
8484 or dropped or have its counters tracked. Those conditions cannot make use of
8485 any data contents because the connection has not been read from yet, and the
8486 buffers are not yet allocated. This is used to selectively and very quickly
8487 accept or drop connections from various sources with a very low overhead. If
8488 some contents need to be inspected in order to take the decision, the
8489 "tcp-request content" statements must be used instead.

8490 The "tcp-request connection" rules are evaluated in their exact declaration
8491 order. If no rule matches or if there is no rule, the default action is to
8492 accept the incoming connection. There is no specific limit to the number of
8493 rules which may be inserted.
8494
8495

8496 Four types of actions are supported :

8497 - accept :
8498 accepts the connection if the condition is true (when used with "if")
8499 or false (when used with "unless"). The first such rule executed ends
8500 the rules evaluation.
8501

8502 - reject :
8503 rejects the connection if the condition is true (when used with "if")
8504 or false (when used with "unless"). The first such rule executed ends
8505 the rules evaluation. Rejected connections do not even become a
8506 session, which is why they are accounted separately for in the stats,
8507 as "denied connections". They are not considered for the session
8508 rate-limit and are not logged either. The reason is that these rules
8509 should only be used to filter extremely high connection rates such as
8510 the ones encountered during a massive DoS attack. Under these extreme
8511 conditions, the simple action of logging each event would make the
8512 system collapse and would considerably lower the filtering capacity. If
8513 logging is absolutely desired, then "tcp-request content" rules should
8514 be used instead.
8515

8516 - expect-proxy layer4 :
8517 configures the client-facing connection to receive a PROXY protocol
8518 header before any byte is read from the socket. This is equivalent to
8519 having the "accept-proxy" keyword on the "bind" line, except that using
8520 the TCP rule allows the PROXY protocol to be accepted only for certain
8521 IP address ranges using an ACL. This is convenient when multiple layers
8522 of load balancers are passed through by traffic coming from public
8523 hosts.
8524

8525 - capture <sample> len <length> :

8526 This only applies to "tcp-request content" rules. It captures sample
8527 expression <sample> from the request buffer, and converts it to a
8528 string of at most <len> characters. The resulting string is stored into
8529 the next request "capture" slot, so it will possibly appear next to
8530 some captured HTTP headers. It will then automatically appear in the
8531 logs, and it will be possible to extract it using sample fetch rules to
8532 feed it into headers or anything. The length should be limited given
8533 that this size will be allocated for each capture during the whole
8534 session life. Please check section 7.3 (Fetching samples) and "capture
8535 request header" for more information.
8536

8537 - { track-sc0 | track-sc1 | track-sc2 } <key> [table <table>] :

8538 enables tracking of sticky counters from current connection. These
8539 rules do not stop evaluation and do not change default action. 3 sets
8540 of counters may be simultaneously tracked by the same connection. The
8541 first "track-sc0" rule executed enables tracking of the counters of the
8542 specified table as the first set. The first "track-sc1" rule executed
8543 enables tracking of the counters of the specified table as the second
8544 set. The first "track-sc2" rule executed enables tracking of the
8545 counters of the specified table as the third set. It is a recommended
8546 practice to use the first set of counters for the per-frontend counters
8547 and the second set for the per-backend ones. But this is just a
8548 guideline, all may be used everywhere.
8549

8550 These actions take one or two arguments :

8551 <key> is mandatory, and is a sample expression rule as described
8552 in section 7.3. It describes what elements of the incoming
8553 request or connection will be analysed, extracted, combined,
8554 and used to select which table entry to update the counters.
8555 Note that "tcp-request connection" cannot use content-based
8556 fetches.
8557

8558 <table> is an optional table to be used instead of the default one,
8559 which is the stick-table declared in the current proxy. All
8560 the counters for the matches and updates for the key will
8561 then be performed in that table until the session ends.

8562 Once a "track-sc*" rule is executed, the key is looked up in the table
8563 and if it is not found, an entry is allocated for it. Then a pointer to
8564 that entry is kept during all the session's life, and this entry's
8565 counters are updated as often as possible, every time the session's
8566 counters are updated, and also systematically when the session ends.
8567 Counters are only updated for events that happen after the tracking has
8568 been started. For example, connection counters will not be updated when
8569 tracking layer 7 information, since the connection event happens before
8570 layer7 information is extracted.
8571

8572 If the entry tracks concurrent connection counters, one connection is
8573 counted for as long as the entry is tracked, and the entry will not
8574 expire during that time. Tracking counters also provides a performance
8575 advantage over just checking the keys, because only one table lookup is
8576 performed for all ACL checks that make use of it.
8577

8578 - sc-inc-gpc0(<sc-id>) :

8579 The "sc-inc-gpc0" increments the GPC0 counter according to the sticky
8580

counter designated by <sc-id>. If an error occurs, this action silently fails and the actions evaluation continues.

- sc-set-gpt0(<sc-id>) <int>:

This action sets the GPT0 tag according to the sticky counter designated by <sc-id> and the value of <int>. The expected result is a boolean. If an error occurs, this action silently fails and the actions evaluation continues.

- "silent-drop" :

This stops the evaluation of the rules and makes the client-facing connection suddenly disappear using a system-dependent way that tries to prevent the client from being notified. The effect is then that the client still sees an established connection while there's none on HAProxy. The purpose is to achieve a comparable effect to "tarpit" except that it doesn't use any local resource at all on the machine running HAProxy. It can resist much higher loads than "tarpit", and slow down stronger attackers. It is important to understand the impact of using this mechanism. All stateful equipments placed between the client and HAProxy (firewalls, proxies, load balancers) will also keep the established connection for a long time and may suffer from this action. On modern Linux systems running with enough privileges, the TCP REPAIR socket option is used to block the emission of a TCP reset. On other systems, the socket's TTL is reduced to 1 so that the TCP reset doesn't pass the first router, though it's still delivered to local networks. Do not use it unless you fully understand how it works.

Note that the "if/unless" condition is optional. If no condition is set on the action, it is simply performed unconditionally. That can be useful for "track-sc*" actions as well as for changing the default action to a reject.

Example: accept all connections from white-listed hosts, reject too fast connection without counting them, and track accepted connections. This results in connection rate being capped from abusive sources.

```
tcp-request connection accept if { src -f /etc/haproxy/whitelist.lst }
tcp-request connection reject if { src_conn_rate gt 10 }
tcp-request connection track-sc0 src
```

Example: accept all connections from white-listed hosts, count all other connections and reject too fast ones. This results in abusive ones being blocked as long as they don't slow down.

```
tcp-request connection accept if { src -f /etc/haproxy/whitelist.lst }
tcp-request connection track-sc0 src
tcp-request connection reject if { sc0_conn_rate gt 10 }
```

Example: enable the PROXY protocol for traffic coming from all known proxies.

```
tcp-request connection expect-proxy layer4 if { src -f proxies.lst }
```

See section 7 about ACL usage.

See also : "tcp-request content", "stick-table"

tcp-request content <action> [(if | unless) <condition>]

Perform an action on a new session depending on a layer 4-7 condition

May be used in sections : defaults | frontend | listen | backend

```
no | yes | yes | yes
```

Arguments :

<action> defines the action to perform if the condition applies. See below.

<condition> is a standard layer 4-7 ACL-based condition (see section 7).

A request's contents can be analysed at an early stage of request processing called "TCP content inspection". During this stage, ACL-based rules are evaluated every time the request contents are updated, until either an "accept" or a "reject" rule matches, or the TCP request inspection delay expires with no matching rule.

The first difference between these rules and "tcp-request connection" rules is that "tcp-request content" rules can make use of contents to take a decision. Most often, these decisions will consider a protocol recognition or validity. The second difference is that content-based rules can be used in both frontends and backends. In case of HTTP keep-alive with the client, all tcp-request content rules are evaluated again, so haproxy keeps a record of what sticky counters were assigned by a "tcp-request connection" versus a "tcp-request content" rule, and flushes all the content-related ones after processing an HTTP request, so that they may be evaluated again by the rules being evaluated again for the next request. This is of particular importance when the rule tracks some L7 information or when it is conditioned by an L7-based ACL, since tracking may change between requests.

Content-based rules are evaluated in their exact declaration order. If no rule matches or if there is no rule, the default action is to accept the contents. There is no specific limit to the number of rules which may be inserted.

Several types of actions are supported :

- accept : the request is accepted
- reject : the request is rejected and the connection is closed
- capture : the specified sample expression is captured
 - { track-sc0 | track-scl | track-sc2 } <key> [table <table>]
 - sc-inc-gp0(<sc-id>)
 - set-gpt0(<sc-id>) <int>
 - set-var(<var-name>) <expr>
 - silent-drop

They have the same meaning as their counter-parts in "tcp-request connection" so please refer to that section for a complete description.

While there is nothing mandatory about it, it is recommended to use the track-sc0 in "tcp-request connection" rules, track-scl for "tcp-request content" rules in the frontend, and track-sc2 for "tcp-request content" rules in the backend, because that makes the configuration more readable and easier to troubleshoot, but this is just a guideline and all counters may be used everywhere.

Note that the "if/unless" condition is optional. If no condition is set on the action, it is simply performed unconditionally. That can be useful for "track-sc*" actions as well as for changing the default action to a reject.

It is perfectly possible to match layer 7 contents with "tcp-request content" rules, since HTTP-specific ACL matches are able to preliminarily parse the contents of a buffer before extracting the required data. If the buffered contents do not parse as a valid HTTP message, then the ACL does not match. The parser which is involved there is exactly the same as for all other HTTP processing, so there is no risk of parsing something differently. In an HTTP backend connected to from an HTTP frontend, it is guaranteed that HTTP contents will always be immediately present when the rule is evaluated first.

Tracking layer7 information is also possible provided that the information are present when the rule is processed. The rule processing engine is able to wait until the inspect delay expires when the data to be tracked is not yet available.

The "set-var" is used to set the content of a variable. The variable is declared inline.

<var-name> The name of the variable starts by an indication about its scope.

The allowed scopes are:

"sess" : the variable is shared with all the session,
"txn" : the variable is shared with all the transaction
(request and response)

"req" : the variable is shared only during the request
processing

"res" : the variable is shared only during the response
processing..

This prefix is followed by a name. The separator is a '.'.
The name may only contain characters 'a-z', 'A-Z', '0-9' and '_'.

<expr> Is a standard HAProxy expression formed by a sample-fetch
followed by some converters.

Example:

```
tcp-request content set-var(sess.my_var) src
```

Example:

```
# Accept HTTP requests containing a Host header saying "example.com"  
# and reject everything else.
```

```
acl is_host_com hdr(Host) -i example.com
```

```
tcp-request inspect-delay 30s
```

```
tcp-request content accept if is_host_com
```

```
tcp-request content reject
```

Example:

```
# reject SMTP connection if client speaks first
```

```
tcp-request inspect-delay 30s
```

```
acl content_present req_len gt 0
```

```
tcp-request content reject if content_present
```

```
# Forward HTTPS connection only if client speaks
```

```
tcp-request inspect-delay 30s
```

```
acl content_present req_len gt 0
```

```
tcp-request content accept if content_present
```

```
tcp-request content reject
```

Example:

```
# Track the last IP from X-Forwarded-For
```

```
tcp-request inspect-delay 10s
```

```
tcp-request content track-sc0 hdr(x-forwarded-for,-1)
```

Example:

```
# track request counts per "base" (concatenation of Host+URL)
```

```
tcp-request inspect-delay 10s
```

```
tcp-request content track-sc0 base table req-rate
```

Example: track per-frontend and per-backend counters, block abusers at the
frontend when the backend detects abuse.

```
frontend http
```

```
# Use General Purpose Counter 0 in SC0 as a global abuse counter
```

```
# protecting all our sites
```

```
stick-table type ip size 1m expire 5m store gpc0
```

```
tcp-request connection track-sc0 src
```

```
tcp-request connection reject if { sc0_get_gpc0 gt 0 }
```

```
...  
use_backend http_dynamic if { path_end .php }
```

```
backend http_dynamic
```

```
# if a source makes too fast requests to this dynamic site (tracked
```

```
# by SC1), block it globally in the frontend.
```

```
8776 stick-table type ip size 1m expire 5m store http_req_rate(10s)  
8777 acl click_too_fast scl_http_req_rate gt 10  
8778 acl mark_as_abuser sc0_inc_gpc0 gt 0  
8779 tcp-request content track-sc1 src  
8780 tcp-request content reject if click_too_fast mark_as_abuser  
8781  
8782 See section 7 about ACL usage.  
8783  
8784 See also : "tcp-request connection", "tcp-request inspect-delay"  
8785  
8786  
8787  
8788 tcp-request inspect-delay <timeout>  
8789 Set the maximum allowed time to wait for data during content inspection  
8790 May be used in sections : defaults | frontend | listen | backend  
8791 no | yes | yes | yes  
8792 Arguments :  
8793 <timeout> is the timeout value specified in milliseconds by default, but  
8794 can be in any other unit if the number is suffixed by the unit,  
8795 as explained at the top of this document.
```

People using haproxy primarily as a TCP relay are often worried about the risk of passing any type of protocol to a server without any analysis. In order to be able to analyze the request contents, we must first withhold the data then analyze them. This statement simply enables withholding of data for at most the specified amount of time.

TCP content inspection applies very early when a connection reaches a frontend, then very early when the connection is forwarded to a backend. This means that a connection may experience a first delay in the frontend and a second delay in the backend if both have tcp-request rules.

Note that when performing content inspection, haproxy will evaluate the whole rules for every new chunk which gets in, taking into account the fact that those data are partial. If no rule matches before the aforementioned delay, a last check is performed upon expiration, this time considering that the contents are definitive. If no delay is set, haproxy will not wait at all and will immediately apply a verdict based on the available information. Obviously this is unlikely to be very useful and might even be racy, so such setups are not recommended.

As soon as a rule matches, the request is released and continues as usual. If the timeout is reached and no rule matches, the default policy will be to let it pass through unaffected.

For most protocols, it is enough to set it to a few seconds, as most clients send the full request immediately upon connection. Add 3 or more seconds to cover TCP retransmits but that's all. For some protocols, it may make sense to use large values, for instance to ensure that the client never talks before the server (eg: SMTP), or to wait for a client to talk before passing data to the server (eg: SSL). Note that the client timeout must cover at least the inspection delay, otherwise it will expire first. If the client closes the connection or if the buffer is full, the delay immediately expires since the contents will not be able to change anymore.

See also : "tcp-request content accept", "tcp-request content reject",
"timeout client".

tcp-response content <action> [if | unless] <condition>]

Perform an action on a session response depending on a layer 4-7 condition
May be used in sections : defaults | frontend | listen | backend
no | no | yes | yes

Arguments :
<action> defines the action to perform if the condition applies. See below.

<condition> is a standard layer 4-7 ACL-based condition (see section 7).

Response contents can be analysed at an early stage of response processing called "TCP content inspection". During this stage, ACL-based rules are evaluated every time the response contents are updated, until either an "accept", "close" or a "reject" rule matches, or a TCP response inspection delay is set and expires with no matching rule.

Most often, these decisions will consider a protocol recognition or validity.

Content-based rules are evaluated in their exact declaration order. If no rule matches or if there is no rule, the default action is to accept the contents. There is no specific limit to the number of rules which may be inserted.

Several types of actions are supported :

- accept :
 - accepts the response if the condition is true (when used with "if") or false (when used with "unless"). The first such rule executed ends the rules evaluation.

- close :
 - immediately closes the connection with the server if the condition is true (when used with "if"), or false (when used with "unless"). The first such rule executed ends the rules evaluation. The main purpose of this action is to force a connection to be finished between a client and a server after an exchange when the application protocol expects some long time outs to elapse first. The goal is to eliminate idle connections which take significant resources on servers with certain protocols.

- reject :
 - rejects the response if the condition is true (when used with "if") or false (when used with "unless"). The first such rule executed ends the rules evaluation. Rejected session are immediately closed.

- set-var(<var-name>) <expr>
 - Sets a variable.

- sc-inc-gpc0(<sc-id>) :

This action increments the GPC0 counter according to the sticky counter designated by <sc-id>. If an error occurs, this action fails silently and the actions evaluation continues.

- sc-set-gpt0(<sc-id>) <int> :

This action sets the GPT0 tag according to the sticky counter designated by <sc-id> and the value of <int>. The expected result is a boolean. If an error occurs, this action silently fails and the actions evaluation continues.

- "silent-drop" :

This stops the evaluation of the rules and makes the client-facing connection suddenly disappear using a system-dependant way that tries to prevent the client from being notified. The effect it then that the client still sees an established connection while there's none on HAProxy. The purpose is to achieve a comparable effect to "tarpit" except that it doesn't use any local resource at all on the machine running HAProxy. It can resist much higher loads than "tarpit", and slow down stronger attackers. It is important to understand the impact of using this mechanism. All stateful equipments placed between the client and HAProxy (firewalls, proxies, load balancers) will also keep the established connection for a long time and may suffer from this action. On modern Linux systems running with enough privileges, the TCP_REPAIR socket option is used to block the emission of a TCP

8905

reset. On other systems, the socket's TTL is reduced to 1 so that the TCP reset doesn't pass the first router, though it's still delivered to local networks. Do not use it unless you fully understand how it works.

Note that the "if/unless" condition is optional. If no condition is set on the action, it is simply performed unconditionally. That can be useful for changing the default action to a reject.

It is perfectly possible to match layer 7 contents with "tcp-response content" rules, but then it is important to ensure that a full response has been buffered, otherwise no contents will match. In order to achieve this, the best solution involves detecting the HTTP protocol during the inspection period.

The "set-var" is used to set the content of a variable. The variable is declared inline.

<var-name> The name of the variable starts by an indication about its scope. The allowed scopes are:

"sess" : the variable is shared with all the session.

"txn" : the variable is shared with all the transaction (request and response)

"req" : the variable is shared only during the request processing
 "res" : the variable is shared only during the response processing.

This prefix is followed by a name. The separator is a '.'.

The name may only contain characters 'a-z', 'A-Z', '0-9' and '_'.

<expr> Is a standard HAProxy expression formed by a sample-fetch followed by some converters.

Example:

tcp-request content set-var(sess.my_var) src

See section 7 about ACL usage.

See also : "tcp-request content", "tcp-response inspect-delay"

tcp-response inspect-delay <timeout>

Set the maximum allowed time to wait for a response during content inspection. May be used in sections :

	defaults	no	yes	yes
frontend				
listen				
backend				

no | no | yes | yes

Arguments :

<timeout> is the timeout value specified in milliseconds by default, but can be in any other unit if the number is suffixed by the unit, as explained at the top of this document.

See also : "tcp-response content", "tcp-request inspect-delay".

timeout check <timeout>

Set additional check timeout, but only after a connection has been already established.

May be used in sections:

	defaults	yes	no	yes	yes
frontend					
listen					
backend					

Arguments:

<timeout> is the timeout value specified in milliseconds by default, but can be in any other unit if the number is suffixed by the unit, as explained at the top of this document.

If set, haproxy uses min("timeout connect", "inter") as a connect timeout

8970

for check and "timeout check" as an additional read timeout. The "min" is used so that people running with *very* long "timeout connect" (eg. those who needed this due to the queue or tarpit) do not slow down their checks. (Please also note that there is no valid reason to have such long connect timeouts, because "timeout queue" and "timeout tarpit" can always be used to avoid that).

If "timeout check" is not set haproxy uses "inter" for complete check timeout (connect + read) exactly like all <1.3.15 version.

In most cases check request is much simpler and faster to handle than normal requests and people may want to kick out laggy servers so this timeout should be smaller than "timeout server".

This parameter is specific to backends, but can be specified once for all in "defaults" sections. This is in fact one of the easiest solutions not to forget about it.

See also: "timeout connect", "timeout queue", "timeout server", "timeout tarpit".

timeout client <timeout>

timeout cliptimeout <timeout> (deprecated)

Set the maximum inactivity time on the client side.

May be used in sections : defaults | frontend | listen | backend
yes | yes | yes | no

Arguments :

<timeout> is the timeout value specified in milliseconds by default, but can be in any other unit if the number is suffixed by the unit, as explained at the top of this document.

The inactivity timeout applies when the client is expected to acknowledge or send data. In HTTP mode, this timeout is particularly important to consider during the first phase, when the client sends the request, and during the response while it is reading data sent by the server. The value is specified in milliseconds by default, but can be in any other unit if the number is suffixed by the unit, as specified at the top of this document. In TCP mode (and to a lesser extent, in HTTP mode), it is highly recommended that the client timeout remains equal to the server timeout in order to avoid complex situations to debug. It is a good practice to cover one or several TCP packet losses by specifying timeouts that are slightly above multiples of 3 seconds (eg: 4 or 5 seconds). If some long-lived sessions are mixed with short-lived sessions (eg: WebSocket and HTTP), it's worth considering "timeout tunnel", which overrides "timeout client" and "timeout server" for tunnels, as well as "timeout client-fin" for half-closed connections.

This parameter is specific to frontends, but can be specified once for all in "defaults" sections. This is in fact one of the easiest solutions not to forget about it. An unspecified timeout results in an infinite timeout, which is not recommended. Such a usage is accepted and works but reports a warning during startup because it may results in accumulation of expired sessions in the system if the system's timeouts are not configured either.

This parameter replaces the old, deprecated "cliptimeout". It is recommended to use it to write new configurations. The form "timeout cliptimeout" is provided only by backwards compatibility but its use is strongly discouraged.

See also : "cliptimeout", "timeout server", "timeout tunnel".

timeout client-fin <timeout>

Set the inactivity timeout on the client side for half-closed connections.

May be used in sections : defaults | frontend | listen | backend
yes | yes | yes | no

Arguments :
<timeout> is the timeout value specified in milliseconds by default, but can be in any other unit if the number is suffixed by the unit, as explained at the top of this document.

The inactivity timeout applies when the client is expected to acknowledge or send data while one direction is already shut down. This timeout is different from "timeout client" in that it only applies to connections which are closed in one direction. This is particularly useful to avoid keeping connections in FIN WAIT state for too long when clients do not disconnect cleanly. This problem is particularly common long connections such as RDP or WebSocket. Note that this timeout can override "timeout tunnel" when a connection shuts down in one direction.

This parameter is specific to frontends, but can be specified once for all in "defaults" sections. By default it is not set, so half-closed connections will use the other timeouts (timeout.client or timeout.tunnel).

See also : "timeout client", "timeout server-fin", and "timeout tunnel".

timeout connect <timeout>

timeout conntimeout <timeout> (deprecated)

Set the maximum time to wait for a connection attempt to a server to succeed.

May be used in sections : defaults | frontend | listen | backend
yes | no | yes | yes

Arguments :

<timeout> is the timeout value specified in milliseconds by default, but can be in any other unit if the number is suffixed by the unit, as explained at the top of this document.

If the server is located on the same LAN as haproxy, the connection should be immediate (less than a few milliseconds). Anyway, it is a good practice to cover one or several TCP packet losses by specifying timeouts that are slightly above multiples of 3 seconds (eg: 4 or 5 seconds). By default, the connect timeout also presets both queue and tarpit timeouts to the same value if these have not been specified.

This parameter is specific to backends, but can be specified once for all in "defaults" sections. This is in fact one of the easiest solutions not to forget about it. An unspecified timeout results in an infinite timeout, which is not recommended. Such a usage is accepted and works but reports a warning during startup because it may results in accumulation of failed sessions in the system if the system's timeouts are not configured either.

This parameter replaces the old, deprecated "conntimeout". It is recommended to use it to write new configurations. The form "timeout conntimeout" is provided only by backwards compatibility but its use is strongly discouraged.

See also: "timeout check", "timeout queue", "timeout server", "conntimeout", "timeout tarpit".

timeout http-keep-alive <timeout>

Set the maximum allowed time to wait for a new HTTP request to appear

May be used in sections : defaults | frontend | listen | backend
yes | yes | yes | yes

Arguments :

<timeout> is the timeout value specified in milliseconds by default, but can be in any other unit if the number is suffixed by the unit, as explained at the top of this document.

By default, the time to wait for a new request in case of keep-alive is set by "timeout http-request". However this is not always convenient because some people want very short keep-alive timeouts in order to release connections

faster, and others prefer to have larger ones but still have short timeouts once the request has started to present itself.

The "http-keep-alive" timeout covers these needs. It will define how long to wait for a new HTTP request to start coming after a response was sent. Once the first byte of request has been seen, the "http-request" timeout is used to wait for the complete request to come. Note that empty lines prior to a new request do not refresh the timeout and are not counted as a new request.

There is also another difference between the two timeouts : when a connection expires during timeout http-keep-alive, no error is returned, the connection just closes. If the connection expires in "http-request" while waiting for a connection to complete, a HTTP 408 error is returned.

In general it is optimal to set this value to a few tens to hundreds of milliseconds, to allow users to fetch all objects of a page at once but without waiting for further clicks. Also, if set to a very small value (eg: 1 millisecond) it will probably only accept pipelined requests but not the non-pipelined ones. It may be a nice trade-off for very large sites running with tens to hundreds of thousands of clients.

If this parameter is not set, the "http-request" timeout applies, and if both are not set, "timeout client" still applies at the lower level. It should be set in the frontend to take effect, unless the frontend is in TCP mode, in which case the HTTP backend's timeout will be used.

See also : "timeout http-request", "timeout client".

timeout http-request <timeout>

Set the maximum allowed time to wait for a complete HTTP request

May be used in sections : defaults | frontend | listen | backend
yes | yes | yes | yes

Arguments :

<timeout> is the timeout value specified in milliseconds by default, but can be in any other unit if the number is suffixed by the unit, as explained at the top of this document.

In order to offer DoS protection, it may be required to lower the maximum accepted time to receive a complete HTTP request without affecting the client timeout. This helps protecting against established connections on which nothing is sent. The client timeout cannot offer a good protection against this abuse because it is an inactivity timeout, which means that if the attacker sends one character every now and then, the timeout will not trigger. With the HTTP request timeout, no matter what speed the client types, the request will be aborted if it does not complete in time. When the timeout expires, an HTTP 408 response is sent to the client to inform it about the problem, and the connection is closed. The logs will report termination codes "cR". Some recent browsers are having problems with this standard, well-documented behaviour, so it might be needed to hide the 408 code using "option http-ignore-probes" or "errorfile 408/dev/null". See more details in the explanations of the "cR" termination code in section 8.5.

By default, this timeout only applies to the header part of the request, and not to any data. As soon as the empty line is received, this timeout is not used anymore. When combined with "option http-buffer-request", this timeout also applies to the body of the request.. It is used again on keep-alive connections to wait for a second request if "timeout http-keep-alive" is not set.

Generally it is enough to set it to a few seconds, as most clients send the full request immediately upon connection. Add 3 or more seconds to cover TCP retransmits but that's all. Setting it to very low values (eg: 50 ms) will generally work on local networks as long as there are no packet losses. This will prevent people from sending bare HTTP requests using telnet.

If this parameter is not set, the client timeout still applies between each chunk of the incoming request. It should be set in the frontend to take effect, unless the frontend is in TCP mode, in which case the HTTP backend's timeout will be used.

See also : "errorfile", "http-ignore-probes", "timeout http-keep-alive", and "timeout client", "option http-buffer-request".

timeout queue <timeout>

Set the maximum time to wait in the queue for a connection slot to be free

May be used in sections : defaults | frontend | listen | backend
yes | no | yes | yes

Arguments :

<timeout> is the timeout value specified in milliseconds by default, but can be in any other unit if the number is suffixed by the unit, as explained at the top of this document.

When a server's maxconn is reached, connections are left pending in a queue which may be server-specific or global to the backend. In order not to wait indefinitely, a timeout is applied to requests pending in the queue. If the timeout is reached, it is considered that the request will almost never be served, so it is dropped and a 503 error is returned to the client.

The "timeout queue" statement allows to fix the maximum time for a request to be left pending in a queue. If unspecified, the same value as the backend's connection timeout ("timeout connect") is used, for backwards compatibility with older versions with no "timeout queue" parameter.

See also : "timeout connect", "contimeout".

timeout server <timeout>

Set the maximum inactivity time on the server side.

May be used in sections : defaults | frontend | listen | backend
yes | no | yes | yes

Arguments :

<timeout> is the timeout value specified in milliseconds by default, but can be in any other unit if the number is suffixed by the unit, as explained at the top of this document.

The inactivity timeout applies when the server is expected to acknowledge or send data. In HTTP mode, this timeout is particularly important to consider during the first phase of the server's response, when it has to send the headers, as it directly represents the server's processing time for the request. To find out what value to put there, it's often good to start with what would be considered as unacceptable response times, then check the logs to observe the response time distribution, and adjust the value accordingly.

The value is specified in milliseconds by default, but can be in any other unit if the number is suffixed by the unit, as specified at the top of this document. In TCP mode (and to a lesser extent, in HTTP mode), it is highly recommended that the client timeout remains equal to the server timeout in order to avoid complex situations to debug. Whatever the expected server response times, it is a good practice to cover at least one or several TCP packet losses by specifying timeouts that are slightly above multiples of 3 seconds (eg: 4 or 5 seconds minimum). If some long-lived sessions are mixed with short-lived sessions (eg: WebSocket and HTTP), it's worth considering "timeout tunnel", which overrides "timeout client" and "timeout server" for tunnels.

This parameter is specific to backends, but can be specified once for all in "defaults" sections. This is in fact one of the easiest solutions not to

forget about it. An unspecified timeout results in an infinite timeout, which is not recommended. Such a usage is accepted and works but reports a warning during startup because it may results in accumulation of expired sessions in the system if the system's timeouts are not configured either.

This parameter replaces the old, deprecated "srvtimeout". It is recommended to use it to write new configurations. The form "timeout srvtimeout" is provided only by backwards compatibility but its use is strongly discouraged.

See also : "srvtimeout", "timeout client" and "timeout tunnel".

timeout server-fin <timeout>

Set the inactivity timeout on the server side for half-closed connections.

May be used in sections : defaults | frontend | listen | backend
yes | no | yes | yes

Arguments :

<timeout> is the timeout value specified in milliseconds by default, but can be in any other unit if the number is suffixed by the unit, as explained at the top of this document.

The inactivity timeout applies when the server is expected to acknowledge or send data while one direction is already shut down. This timeout is different from "timeout server" in that it only applies to connections which are closed in one direction. This is particularly useful to avoid keeping connections in FIN_WAIT state for too long when a remote server does not disconnect cleanly. This problem is particularly common long connections such as RDP or WebSocket. Note that this timeout can override "timeout tunnel" when a connection shuts down in one direction. This setting was provided for completeness, but in most situations, it should not be needed.

This parameter is specific to backends, but can be specified once for all in "defaults" sections. By default it is not set, so half-closed connections will use the other timeouts (timeout.server or timeout.tunnel).

See also : "timeout client-fin", "timeout server", and "timeout tunnel".

timeout tarpit <timeout>

Set the duration for which tarpitted connections will be maintained

May be used in sections : defaults | frontend | listen | backend
yes | yes | yes | yes

Arguments :

<timeout> is the tarpit duration specified in milliseconds by default, but can be in any other unit if the number is suffixed by the unit, as explained at the top of this document.

When a connection is tarpitted using "reqtarpit", it is maintained open with no activity for a certain amount of time, then closed. "timeout tarpit" defines how long it will be maintained open.

The value is specified in milliseconds by default, but can be in any other unit if the number is suffixed by the unit, as specified at the top of this document. If unspecified, the same value as the backend's connection timeout ("timeout connect") is used, for backwards compatibility with older versions with no "timeout tarpit" parameter.

See also : "timeout connect", "contimeout".

timeout tunnel <timeout>

Set the maximum inactivity time on the client and server side for tunnels.

May be used in sections : defaults | frontend | listen | backend
yes | no | yes | yes

Arguments :

<timeout> is the timeout value specified in milliseconds by default, but can be in any other unit if the number is suffixed by the unit, as explained at the top of this document.

The tunnel timeout applies when a bidirectional connection is established between a client and a server, and the connection remains inactive in both directions. This timeout supersedes both the client and server timeouts once the connection becomes a tunnel. In TCP, this timeout is used as soon as no analyser remains attached to either connection (eg: tcp content rules are accepted). In HTTP, this timeout is used when a connection is upgraded (eg: when switching to the WebSocket protocol, or forwarding a CONNECT request to a proxy), or after the first response when no keepalive/close option is specified.

Since this timeout is usually used in conjunction with long-lived connections, it usually is a good idea to also set "timeout client-fin" to handle the situation where a client suddenly disappears from the net and does not acknowledge a close, or sends a shutdown and does not acknowledge pending data anymore. This can happen in lossy networks where firewalls are present, and is detected by the presence of large amounts of sessions in a FIN_WAIT state.

The value is specified in milliseconds by default, but can be in any other unit if the number is suffixed by the unit, as specified at the top of this document. Whatever the expected normal idle time, it is a good practice to cover at least one or several TCP packet losses by specifying timeouts that are slightly above multiples of 3 seconds (eg: 4 or 5 seconds minimum).

This parameter is specific to backends, but can be specified once for all in "defaults" sections. This is in fact one of the easiest solutions not to forget about it.

Example :

```
defaults http
    option http-server-close
    timeout connect 5s
    timeout client 30s
    timeout client-fin 30s
    timeout server 30s
    timeout tunnel 1h # timeout to use with WebSocket and CONNECT
```

See also : "timeout client", "timeout client-fin", "timeout server".

transparent (deprecated)

Enable client-side transparent proxying

May be used in sections : defaults | frontend | listen | backend
yes | no | yes | yes

Arguments : none

This keyword was introduced in order to provide layer 7 persistence to layer 3 load balancers. The idea is to use the OS's ability to redirect an incoming connection for a remote address to a local process (here HAProxy), and let this process know what address was initially requested. When this option is used, sessions without cookies will be forwarded to the original destination IP address of the incoming request (which should match that of another equipment), while requests with cookies will still be forwarded to the appropriate server.

The "transparent" keyword is deprecated, use "option transparent" instead.

Note that contrary to a common belief, this option does NOT make HAProxy present the client's IP to the server when establishing the connection.

See also: "option transparent"


```
9361 unique-id-format <string>
9362 Generate a unique ID for each request.
9363 May be used in sections : defaults | frontend | listen | backend
9364 yes | yes | yes | no
9365 Arguments :
9366 <string> is a log-format string.
9367
9368 This keyword creates a ID for each request using the custom log format. A
9369 unique ID is useful to trace a request passing through many components of
9370 a complex infrastructure. The newly created ID may also be logged using the
9371 %ID tag the log-format string.
9372
9373 The format should be composed from elements that are guaranteed to be
9374 unique when combined together. For instance, if multiple haproxy instances
9375 are involved, it might be important to include the node name. It is often
9376 needed to log the incoming connection's source and destination addresses
9377 and ports. Note that since multiple requests may be performed over the same
9378 connection, including a request counter may help differentiate them.
9379 Similarly, a timestamp may protect against a rollover of the counter.
9380 Logging the process ID will avoid collisions after a service restart.
9381
9382 It is recommended to use hexadecimal notation for many fields since it
9383 makes them more compact and saves space in logs.
9384
9385 Example:
9386
9387 unique-id-format {%+X}%\ %ci:%cp_%fi:%fp_%Ts_%rt:%pid
9388 will generate:
9389
9390 7F000001:8296_7F00001E:1F90_4F7B0A69_0003:790A
9391
9392 See also: "unique-id-header"
9393
9394 unique-id-header <name>
9395 Add a unique ID header in the HTTP request.
9396 May be used in sections : defaults | frontend | listen | backend
9397 yes | yes | yes | no
9398 Arguments :
9399 <name> is the name of the header.
9400
9401 Add a unique-id header in the HTTP request sent to the server, using the
9402 unique-id-format. It can't work if the unique-id-format doesn't exist.
9403
9404 Example:
9405
9406 unique-id-format {%+X}%\ %ci:%cp_%fi:%fp_%Ts_%rt:%pid
9407 unique-id-header X-Unique-ID
9408 will generate:
9409
9410 X-Unique-ID: 7F000001:8296_7F00001E:1F90_4F7B0A69_0003:790A
9411
9412 See also: "unique-id-format"
9413
9414 use backend <backend> [{if | unless} <condition>]
9415 Switch to a specific backend if/unless an ACL-based condition is matched.
9416 May be used in sections : defaults | frontend | listen | backend
9417 no | yes | yes | no
9418 Arguments :
9419 <backend> is the name of a valid backend or "listen" section, or a
9420 "log-format" string resolving to a backend name.
9421
9422 <condition> is a condition composed of ACLs, as described in section 7. If
```

```
9426 it is omitted, the rule is unconditionally applied.
9427
9428 When doing content-switching, connections arrive on a frontend and are then
9429 dispatched to various backends depending on a number of conditions. The
9430 relation between the conditions and the backends is described with the
9431 "use backend" keyword. While it is normally used with HTTP processing, it can
9432 also be used in pure TCP, either without content using stateless ACLs (eg:
9433 source address validation) or combined with a "tcp-request" rule to wait for
9434 some payload.
9435
9436 There may be as many "use backend" rules as desired. All of these rules are
9437 evaluated in their declaration order, and the first one which matches will
9438 assign the backend.
9439
9440 In the first form, the backend will be used if the condition is met. In the
9441 second form, the backend will be used if the condition is not met. If no
9442 condition is valid, the backend defined with "default_backend" will be used.
9443 If no default backend is defined, either the servers in the same section are
9444 used (in case of a "listen" section) or, in case of a frontend, no server is
9445 used and a 503 service unavailable response is returned.
9446
9447 Note that it is possible to switch from a TCP frontend to an HTTP backend. In
9448 this case, either the frontend has already checked that the protocol is HTTP,
9449 and backend processing will immediately follow, or the backend will wait for
9450 a complete HTTP request to get in. This feature is useful when a frontend
9451 must decode several protocols on a unique port, one of them being HTTP.
9452
9453 When <backend> is a simple name, it is resolved at configuration time, and an
9454 error is reported if the specified backend does not exist. If <backend> is
9455 a log-format string instead, no check may be done at configuration time, so
9456 the backend name is resolved dynamically at run time. If the resulting
9457 backend name does not correspond to any valid backend, no other rule is
9458 evaluated, and the default backend directive is applied instead. Note that
9459 when using dynamic backend names, it is highly recommended to use a prefix
9460 that no other backend uses in order to ensure that an unauthorized backend
9461 cannot be forced from the request.
9462
9463 It is worth mentioning that "use_backend" rules with an explicit name are
9464 used to detect the association between frontends and backends to compute the
9465 backend's "fullconn" setting. This cannot be done for dynamic names.
9466
9467 See also: "default_backend", "tcp-request", "fullconn", "log-format", and
9468 section 7 about ACLs.
9469
9470 use-server <server> if <condition>
9471 use-server <server> unless <condition>
9472 Only use a specific server if/unless an ACL-based condition is matched.
9473 May be used in sections : defaults | frontend | listen | backend
9474 no | no | yes | yes
9475 Arguments :
9476 <server> is the name of a valid server in the same backend section.
9477 <condition> is a condition composed of ACLs, as described in section 7.
9478
9479 By default, connections which arrive to a backend are load-balanced across
9480 the available servers according to the configured algorithm, unless a
9481 persistence mechanism such as a cookie is used and found in the request.
9482
9483 Sometimes it is desirable to forward a particular request to a specific
9484 server without having to declare a dedicated backend for this server. This
9485 can be achieved using the "use-server" rules. These rules are evaluated after
9486 the "redirect" rules and before evaluating cookies, and they have precedence
9487 on them. There may be as many "use-server" rules as desired. All of these
9488 rules are evaluated in their declaration order, and the first one which
```

matches will assign the server.

If a rule designates a server which is down, and "option persist" is not used and no force-persist rule was validated, it is ignored and evaluation goes on with the next rules until one matches.

In the first form, the server will be used if the condition is met. In the second form, the server will be used if the condition is not met. If no condition is valid, the processing continues and the server will be assigned according to other persistence mechanisms.

Note that even if a rule is matched, cookie processing is still performed but does not assign the server. This allows prefixed cookies to have their prefix stripped.

The "use-server" statement works both in HTTP and TCP mode. This makes it suitable for use with content-based inspection. For instance, a server could be selected in a farm according to the TLS SNI field. And if these servers have their weight set to zero, they will not be used for other traffic.

```
Example :
# intercept incoming TLS requests based on the SNI field
use-server www if { req_ssl_sni -i www.example.com }
server www 192.168.0.1:443 weight 0
use-server mail if { req_ssl_sni -i mail.example.com }
server mail 192.168.0.1:587 weight 0
use-server imap if { req_ssl_sni -i imap.example.com }
server mail 192.168.0.1:993 weight 0
# all the rest is forwarded to this server
server default 192.168.0.2:443 check
```

See also: "use_backend", section 5 about server and section 7 about ACLs.

5. Bind and Server options

The "bind", "server" and "default-server" keywords support a number of settings depending on some build options and on the system HAProxy was built on. These settings generally each consist in one word sometimes followed by a value, written on the same line as the "bind" or "server" line. All these options are described in this section.

5.1. Bind options

The "bind" keyword supports a certain number of settings which are all passed as arguments on the same line. The order in which those arguments appear makes no importance, provided that they appear after the bind address. All of these parameters are optional. Some of them consist in a single words (booleans), while other ones expect a value after them. In this case, the value must be provided immediately after the setting name.

The currently supported settings are the following ones.

accept-proxy

Enforces the use of the PROXY protocol over any connection accepted by any of the sockets declared on the same line. Versions 1 and 2 of the PROXY protocol are supported and correctly detected. The PROXY protocol dictates the layer 3/4 addresses of the incoming connection to be used everywhere an address is used, with the only exception of "tcp-request connection" rules which will only see the real connection address. Logs will reflect the addresses indicated in the protocol, unless it is violated, in which case the real address will still be used. This keyword combined with support from external

components can be used as an efficient and reliable alternative to the X-Forwarded-For mechanism which is not always reliable and not even always usable. See also "tcp-request connection expect-proxy" for a finer-grained setting of which client is allowed to use the protocol.

alpn <protocols>

This enables the TLS ALPN extension and advertises the specified protocol list as supported on top of ALPN. The protocol list consists in a comma-delimited list of protocol names, for instance: "http/1.1,http/1.0" (without quotes). This requires that the SSL library is built with support for TLS extensions enabled (check with haproxy -vv). The ALPN extension replaces the initial NPN extension.

backlog <backlog>

Sets the socket's backlog to this value. If unspecified, the frontend's backlog is used instead, which generally defaults to the maxconn value.

edhe <named curve>

This setting is only available when support for OpenSSL was built in. It sets the named curve (RFC 4492) used to generate ECDH ephemeral keys. By default, used named curve is prime256v1.

ca-file <cafile>

This setting is only available when support for OpenSSL was built in. It designates a PEM file from which to load CA certificates used to verify client's certificate.

ca-ignore-err [all<errorID>,...]

This setting is only available when support for OpenSSL was built in. Sets a comma separated list of errorIDs to ignore during verify at depth > 0. If set to 'all', all errors are ignored. SSL handshake is not aborted if an error is ignored.

ca-sign-file <cafile>

This setting is only available when support for OpenSSL was built in. It designates a PEM file containing both the CA certificate and the CA private key used to create and sign server's certificates. This is a mandatory setting when the dynamic generation of certificates is enabled. See 'generate-certificates' for details.

ca-sign-passphrase <passphrase>

This setting is only available when support for OpenSSL was built in. It is the CA private key passphrase. This setting is optional and used only when the dynamic generation of certificates is enabled. See 'generate-certificates' for details.

ciphers <ciphers>

This setting is only available when support for OpenSSL was built in. It sets the string describing the list of cipher algorithms ("cipher suite") that are negotiated during the SSL/TLS handshake. The format of the string is defined in "man 1 ciphers" from OpenSSL man pages, and can be for instance a string such as "AES:ALL:!aNULL:!eNULL:!RC4:@STRENGTH" (without quotes).

crl-file <crlfile>

This setting is only available when support for OpenSSL was built in. It designates a PEM file from which to load certificate revocation list used to verify client's certificate.

crt <cert>

This setting is only available when support for OpenSSL was built in. It designates a PEM file containing both the required certificates and any associated private keys. This file can be built by concatenating multiple PEM files into one (e.g. cat cert.pem key.pem > combined.pem). If your CA requires an intermediate certificate, this can also be concatenated into this file.

If the OpenSSL used supports Diffie-Hellman, parameters present in this file are loaded.

If a directory name is used instead of a PEM file, then all files found in that directory will be loaded in alphabetic order unless their name ends with '.issuer', '.ocsp' or '.sctl' (reserved extensions). This directive may be specified multiple times in order to load certificates from multiple files or directories. The certificates will be presented to clients who provide a valid TLS Server Name Indication field matching one of their CN or alt subjects. Wildcards are supported, where a wildcard character '*' is used instead of the first hostname component (eg: *.example.org matches www.example.org but not www.sub.example.org).

If no SNI is provided by the client or if the SSL library does not support TLS extensions, or if the client provides an SNI hostname which does not match any certificate, then the first loaded certificate will be presented. This means that when loading certificates from a directory, it is highly recommended to load the default one first as a file or to ensure that it will always be the first one in the directory.

Note that the same cert may be loaded multiple times without side effects.

Some CAs (such as Godaddy) offer a drop down list of server types that do not include HAProxy when obtaining a certificate. If this happens be sure to choose a webserver that the CA believes requires an intermediate CA (for Godaddy, selection Apache Tomcat will get the correct bundle. but many others, e.g. nginx, result in a wrong bundle that will not work for some clients).

For each PEM file, haproxy checks for the presence of file at the same path suffixed by ".ocsp". If such file is found, support for the TLS Certificate Status Request extension (also known as "OCSP stapling") is automatically enabled. The content of this file is optional. If not empty, it must contain a valid OCSP Response in DER format. In order to be valid an OCSP Response must comply with the following rules: it has to indicate a good status, it has to be a single response for the certificate of the PEM file, and it has to be valid at the moment of addition. If these rules are not respected the OCSP Response is ignored and a warning is emitted. In order to identify which certificate an OCSP Response applies to, the issuer's certificate is necessary. If the issuer's certificate is not found in the PEM file, it will be loaded from a file at the same path as the PEM file suffixed by ".issuer" if it exists otherwise it will fail with an error.

For each PEM file, haproxy also checks for the presence of file at the same path suffixed by ".sctl". If such file is found, support for Certificate Transparency (RFC6962) TLS extension is enabled. The file must contain a valid Signed Certificate Timestamp List, as described in RFC. File is parsed to check basic syntax, but no signatures are verified.

There are cases where it is desirable support multiple key types (RSA/ECDSA) in the cipher suites offered to the clients. This allows clients that support EC certificates to be able to use EC ciphers, while simultaneously supporting older, RSA only clients.

In order to provide this functionality, multiple PEM files, each with a different key type, are required. To associate these PEM files into a "cert bundle" that is recognized by haproxy, they must be named in the following way: All PEM files that are to be bundled must have the same base name, with a suffix indicating the key type. Currently, three suffixes are supported: rsa, dsa and ecdsa. For example, if www.example.com has two PEM files, an RSA file and an ECDSA file, they must be named: "example.pem.rsa" and "example.pem.ecdsa". The first part of the filename is arbitrary; only the suffix matters. To load this bundle into haproxy, specify the base name only:

Example : bind :8443 ssl crt example.pem

Note that the suffix is not given to haproxy, this tells haproxy to look for a cert bundle.

Haproxy will load all PEM files in the bundle at the same time to try to support multiple key types. PEM files are combined based on Common Name (CN) and Subject Alternative Name (SAN) to support SNI lookups. This means that even if you give haproxy a cert bundle, if there are no shared CN/SAN entries in the certificates in that bundle, haproxy will not be able to provide multi-cert support.

Assuming bundle in the example above contained the following:

Filename	CN	SAN
example.pem.rsa	www.example.com	rsa.example.com
example.pem.ecdsa	www.example.com	ecdsa.example.com

Users connecting with an SNI of "www.example.com" will be able to use both RSA and ECDSA cipher suites. Users connecting with an SNI of "rsa.example.com" will only be able to use RSA cipher suites, and users connecting with "ecdsa.example.com" will only be able to use ECDSA cipher suites.

If a directory name is given as the <cert> argument, haproxy will automatically search and load bundled files in that directory.

OCSP files (.ocsp) and issuer files (.issuer) are supported with multi-cert bundling. Each certificate can have its own .ocsp and .issuer file. At this time, sctl is not supported in multi-certificate bundling.

crt-ignore-err <errors>

This setting is only available when support for OpenSSL was built in. Sets a comma separated list of errorIDs to ignore during verify at depth == 0. If set to 'all', all errors are ignored. SSL handshake is not aborted if an error is ignored.

crt-list <file>

This setting is only available when support for OpenSSL was built in. It designates a list of PEM file with an optional list of SNI filter per certificate, with the following format for each line :

<certfile> [[]<snifilter> ...]

Wildcards are supported in the SNI filter. Negative filter are also supported, only useful in combination with a wildcard filter to exclude a particular SNI. The certificates will be presented to clients who provide a valid TLS Server Name Indication field matching one of the SNI filters. If no SNI filter is specified, the CN and alt subjects are used. This directive may be specified multiple times. See the "crt" option for more information. The default certificate is still needed to meet OpenSSL expectations. If it is not used, the 'strict-sni' option may be used.

Multi-cert bundling (see "crt") is support with crt-list, as long as only the base name is given in the crt-list. Due to the nature of bundling, all SNI filters given to a multi-cert bundle entry are ignored.

defer-accept

Is an optional keyword which is supported only on certain Linux kernels. It states that a connection will only be accepted once some data arrive on it, or at worst after the first retransmit. This should be used only on protocols for which the client talks first (eg: HTTP). It can slightly improve

performance by ensuring that most of the request is already available when the connection is accepted. On the other hand, it will not be able to detect connections which don't talk. It is important to note that this option is broken in all kernels up to 2.6.31, as the connection is never accepted until the client talks. This can cause issues with front firewalls which would see an established connection while the proxy will only see it in SYN_RECV. This option is only supported on TCPv4/TCPv6 sockets and ignored by other ones.

force-sslvs3

This option enforces use of SSLv3 only on SSL connections instantiated from this listener. SSLv3 is generally less expensive than the TLS counterparts for high connection rates. This option is also available on global statement "ssl-default-bind-options". See also "no-tlsv*" and "no-sslvs3".

force-tlsv10

This option enforces use of TLSv1.0 only on SSL connections instantiated from this listener. This option is also available on global statement "ssl-default-bind-options". See also "no-tlsv*" and "no-sslvs3".

force-tlsv11

This option enforces use of TLSv1.1 only on SSL connections instantiated from this listener. This option is also available on global statement "ssl-default-bind-options". See also "no-tlsv*", and "no-sslvs3".

force-tlsv12

This option enforces use of TLSv1.2 only on SSL connections instantiated from this listener. This option is also available on global statement "ssl-default-bind-options". See also "no-tlsv*", and "no-sslvs3".

generate-certificates

This setting is only available when support for OpenSSL was built in. It enables the dynamic SSL certificates generation. A CA certificate and its private key are necessary (see 'ca-sign-file'). When HAProxy is configured as a transparent forward proxy, SSL requests generate errors because of a common name mismatch on the certificate presented to the client. With this option enabled, HAProxy will try to forge a certificate using the SNI hostname indicated by the client. This is done only if no certificate matches the SNI hostname (see 'crt-list'). If an error occurs, the default certificate is used, else the 'strict-sni' option is set.

It can also be used when HAProxy is configured as a reverse proxy to ease the deployment of an architecture with many backends.

Creating a SSL certificate is an expensive operation, so a LRU cache is used to store forged certificates (see 'tune.ssl.ssl-ctx-cache-size'). It increases the HAProxy's memory footprint to reduce latency when the same certificate is used many times.

gid <gid>

Sets the group of the UNIX sockets to the designated system gid. It can also be set by default in the global section's "unix-bind" statement. Note that some platforms simply ignore this. This setting is equivalent to the "group" setting except that the group name is used instead of its gid. This setting is ignored by non UNIX sockets.

group <group>

Sets the group of the UNIX sockets to the designated system group. It can also be set by default in the global section's "unix-bind" statement. Note that some platforms simply ignore this. This setting is equivalent to the "gid" setting except that the group name is used instead of its gid. This setting is ignored by non UNIX sockets.

id <id>

Fixes the socket ID. By default, socket IDs are automatically assigned, but sometimes it is more convenient to fix them to ease monitoring. This value must be strictly positive and unique within the listener/frontend. This

option can only be used when defining only a single socket.

interface <interface>

Restricts the socket to a specific interface. When specified, only packets received from that particular interface are processed by the socket. This is currently only supported on Linux. The interface must be a primary system interface, not an aliased interface. It is also possible to bind multiple frontends to the same address if they are bound to different interfaces. Note that binding to a network interface requires root privileges. This parameter is only compatible with TCPv4/TCPv6 sockets.

level <level>

This setting is used with the stats sockets only to restrict the nature of the commands that can be issued on the socket. It is ignored by other sockets. <level> can be one of :

- "user" is the least privileged level ; only non-sensitive stats can be read, and no change is allowed. It would make sense on systems where it is not easy to restrict access to the socket.
- "operator" is the default level and fits most common uses. All data can be read, and only non-sensitive changes are permitted (eg: clear max counters).
- "admin" should be used with care, as everything is permitted (eg: clear all counters).

maxconn <maxconn>

Limits the sockets to this number of concurrent connections. Extraneous connections will remain in the system's backlog until a connection is released. If unspecified, the limit will be the same as the frontend's maxconn. Note that in case of port ranges or multiple addresses, the same value will be applied to each socket. This setting enables different limitations on expensive sockets, for instance SSL entries which may easily eat all memory.

mode <mode>

Sets the octal mode used to define access permissions on the UNIX socket. It can also be set by default in the global section's "unix-bind" statement. Note that some platforms simply ignore this. This setting is ignored by non UNIX sockets.

mss <maxseg>

Sets the TCP Maximum Segment Size (MSS) value to be advertised on incoming connections. This can be used to force a lower MSS for certain specific ports, for instance for connections passing through a VPN. Note that this relies on a kernel feature which is theoretically supported under Linux but was buggy in all versions prior to 2.6.28. It may or may not work on other operating systems. It may also not change the advertised value but change the effective size of outgoing segments. The commonly advertised value for TCPv4 over Ethernet networks is 1460 = 1500(MTU) - 40(IP+TCP). If this value is positive, it will be used as the advertised MSS. If it is negative, it will indicate by how much to reduce the incoming connection's advertised MSS for outgoing segments. This parameter is only compatible with TCP v4/v6 sockets.

name <name>

Sets an optional name for these sockets, which will be reported on the stats page.

namespace <name>

On Linux, it is possible to specify which network namespace a socket will belong to. This directive makes it possible to explicitly bind a listener to a namespace different from the default one. Please refer to your operating system's documentation to find more details about network namespaces.

nice <nice>

Sets the 'niceness' of connections initiated from the socket. Value must be in the range -1024..1024 inclusive, and defaults to zero. Positive values

means that such connections are more friendly to others and easily offer their place in the scheduler. On the opposite, negative values mean that connections want to run with a higher priority than others. The difference only happens under high loads when the system is close to saturation. Negative values are appropriate for low-latency or administration services, and high values are generally recommended for CPU intensive tasks such as SSL processing or bulk transfers which are less sensible to latency. For example, it may make sense to use a positive value for an SMTP socket and a negative one for an RDP socket.

no-sslv3

This setting is only available when support for OpenSSL was built in. It disables support for SSLv3 on any sockets instantiated from the listener when SSL is supported. Note that SSLv2 is forced disabled in the code and cannot be enabled using any configuration option. This option is also available on global statement "ssl-default-bind-options". See also "force-tls*", and "force-sslv3".

no-tls-tickets

This setting is only available when support for OpenSSL was built in. It disables the stateless session resumption (RFC 5077 TLS Ticket extension) and force to use stateful session resumption. Stateless session resumption is more expensive in CPU usage. This option is also available on global statement "ssl-default-bind-options".

no-tlsv10

This setting is only available when support for OpenSSL was built in. It disables support for TLSv1.0 on any sockets instantiated from the listener when SSL is supported. Note that SSLv2 is forced disabled in the code and cannot be enabled using any configuration option. This option is also available on global statement "ssl-default-bind-options". See also "force-tlsv*", and "force-sslv3".

no-tlsv11

This setting is only available when support for OpenSSL was built in. It disables support for TLSv1.1 on any sockets instantiated from the listener when SSL is supported. Note that SSLv2 is forced disabled in the code and cannot be enabled using any configuration option. This option is also available on global statement "ssl-default-bind-options". See also "force-tlsv*", and "force-sslv3".

no-tlsv12

This setting is only available when support for OpenSSL was built in. It disables support for TLSv1.2 on any sockets instantiated from the listener when SSL is supported. Note that SSLv2 is forced disabled in the code and cannot be enabled using any configuration option. This option is also available on global statement "ssl-default-bind-options". See also "force-tlsv*", and "force-sslv3".

npn <protocols>

This enables the NPN TLS extension and advertises the specified protocol list as supported on top of NPN. The protocol list consists in a comma-delimited list of protocol names, for instance: "http/1.1,http/1.0" (without quotes). This requires that the SSL library is built with support for TLS extensions enabled (check with haproxy -vv). Note that the NPN extension has been replaced with the ALPN extension (see the "alpn" keyword).

process [all | odd | even | <number 1-64> [-<number 1-64>]]

This restricts the list of processes on which this listener is allowed to run. It does not enforce any process but eliminates those which do not match. If the frontend uses a "bind-process" setting, the intersection between the two is applied. If in the end the listener is not allowed to run on any remaining process, a warning is emitted, and the listener will either run the first process of the listener if a single process was specified, or on all of its processes if multiple processes were specified. For the unlikely

case where several ranges are needed, this directive may be repeated. The main purpose of this directive is to be used with the stats sockets and have one different socket per process. The second purpose is to have multiple bind lines sharing the same IP:port but not the same process in a listener, so that the system can distribute the incoming connections into multiple queues and allow a smoother inter-process load balancing. Currently Linux 3.9 and above is known for supporting this. See also "bind-process" and "nproc".

ssl

This setting is only available when support for OpenSSL was built in. It enables SSL deciphering on connections instantiated from this listener. A certificate is necessary (see "crt" above). All contents in the buffers will appear in clear text, so that ACLs and HTTP processing will only have access to deciphered contents.

strict-sni

This setting is only available when support for OpenSSL was built in. The SSL/TLS negotiation is allow only if the client provided an SNI which match a certificate. The default certificate is not used. See the "crt" option for more information.

tcp-ut <delay>

Sets the TCP User Timeout for all incoming connections instantiated from this listening socket. This option is available on Linux since version 2.6.37. It allows haproxy to configure a timeout for sockets which contain data not receiving an acknowledgement for the configured delay. This is especially useful on long-lived connections experiencing long idle periods such as remote terminals or database connection pools, where the client and server timeouts must remain high to allow a long period of idle, but where it is important to detect that the client has disappeared in order to release all resources associated with its connection (and the server's session). The argument is a delay expressed in milliseconds by default. This only works for regular TCP connections, and is ignored for other protocols.

tfo

Is an optional keyword which is supported only on Linux kernels ≥ 3.7 . It enables TCP Fast Open on the listening socket, which means that clients which support this feature will be able to send a request and receive a response during the 3-way handshake starting from second connection, thus saving one round-trip after the first connection. This only makes sense with protocols that use high connection rates and where each round trip matters. This can possibly cause issues with many firewalls which do not accept data on SYN packets, so this option should only be enabled once well tested. This option is only supported on TCPv4/TCPv6 sockets and ignored by other ones. You may need to build HAProxy with USE_TFO=1 if your libc doesn't define TCP_FASTOPEN.

tls-ticket-keys <keyfile>

Sets the TLS ticket keys file to load the keys from. The keys need to be 48 bytes long, encoded with base64 (ex. openssl rand -base64 48). Number of keys is specified by the TLS_TICKETS_NO build option (default 3) and at least as many keys need to be present in the file. Last TLS_TICKETS_NO keys will be used for decryption and the penultimate one for encryption. This enables easy key rotation by just appending new key to the file and reloading the process. Keys must be periodically rotated (ex. every 12h) or Perfect Forward Secrecy is compromised. It is also a good idea to keep the keys off any permanent storage such as hard drives (hint: use tmpfs and don't swap those files). Lifetime hint can be changed using tune.ssl.timeout.

transparent

Is an optional keyword which is supported only on certain Linux kernels. It indicates that the addresses will be bound even if they do not belong to the local machine, and that packets targeting any of these addresses will be intercepted just as if the addresses were locally configured. This normally requires that IP forwarding is enabled. Caution! do not use this with the

default address '*', as it would redirect any traffic for the specified port. This keyword is available only when HAProxy is built with USE_LINUX_TPROXY=1. This parameter is only compatible with TCPv4 and TCPv6 sockets, depending on kernel version. Some distribution kernels include backports of the feature, so check for support with your vendor.

v4v6

Is an optional keyword which is supported only on most recent systems including Linux kernels >= 2.4.21. It is used to bind a socket to both IPv4 and IPv6 when it uses the default address. Doing so is sometimes necessary on systems which bind to IPv6 only by default. It has no effect on non-IPv6 sockets, and is overridden by the "v6only" option.

v6only

Is an optional keyword which is supported only on most recent systems including Linux kernels >= 2.4.21. It is used to bind a socket to IPv6 only when it uses the default address. Doing so is sometimes preferred to doing it system-wide as it is per-listener. It has no effect on non-IPv6 sockets and has precedence over the "v4v6" option.

uid <uid>

Sets the owner of the UNIX sockets to the designated system uid. It can also be set by default in the global section's "unix-bind" statement. Note that some platforms simply ignore this. This setting is equivalent to the "user" setting except that the user numeric ID is used instead of its name. This setting is ignored by non UNIX sockets.

user <user>

Sets the owner of the UNIX sockets to the designated system user. It can also be set by default in the global section's "unix-bind" statement. Note that some platforms simply ignore this. This setting is equivalent to the "uid" setting except that the user name is used instead of its uid. This setting is ignored by non UNIX sockets.

verify [none|optional|required]

This setting is only available when support for OpenSSL was built in. If set to 'none', client certificate is not requested. This is the default. In other cases, a client certificate is requested. If the client does not provide a certificate after the request and if 'verify' is set to 'required', then the handshake is aborted, while it would have succeeded if set to 'optional'. The certificate provided by the client is always verified using CA's from 'ca-file' and optional CRLs from 'crl-file'. On verify failure the handshake is aborted, regardless of the 'verify' option, unless the error code exactly matches one of those listed with 'ca-ignore-err' or 'crt-ignore-err'.

5.2. Server and default-server options

The "server" and "default-server" keywords support a certain number of settings which are all passed as arguments on the server line. The order in which those arguments appear does not count, and they are all optional. Some of those settings are single words (booleans) while others expect one or several values after them. In this case, the values must immediately follow the setting name. Except default-server, all those settings must be specified after the server's address if they are used:

```
server <name> <address>[:port] [settings ...]
default-server [settings ...]
```

The currently supported settings are the following ones.

addr <ipv4|ipv6>

Using the "addr" parameter, it becomes possible to use a different IP address to send health-checks. On some servers, it may be desirable to dedicate an IP address to specific component able to perform complex tests which are more

suitable to health-checks than the application. This parameter is ignored if the "check" parameter is not set. See also the "port" parameter.

Supported in default-server: No

agent-check

Enable an auxiliary agent check which is run independently of a regular health check. An agent health check is performed by making a TCP connection to the port set by the "agent-port" parameter and reading an ASCII string. The string is made of a series of words delimited by spaces, tabs or commas in any order, optionally terminated by '\r' and/or '\n', each consisting of :

- An ASCII representation of a positive integer percentage, e.g. "75%". Values in this format will set the weight proportional to the initial weight of a server as configured when haproxy starts. Note that a zero weight is reported on the stats page as "DRAIN" since it has the same effect on the server (it's removed from the LB farm).

- The word "ready". This will turn the server's administrative state to the READY mode, thus cancelling any DRAIN or MAINT state

- The word "drain". This will turn the server's administrative state to the DRAIN mode, thus it will not accept any new connections other than those that are accepted via persistence.

- The word "maint". This will turn the server's administrative state to the MAINT mode, thus it will not accept any new connections at all, and health checks will be stopped.

- The words "down", "failed", or "stopped", optionally followed by a description string after a sharp ('#'). All of these mark the server's operating state as DOWN, but since the word itself is reported on the stats page, the difference allows an administrator to know if the situation was expected or not : the service may intentionally be stopped, may appear up but fail some validity tests, or may be seen as down (eg: missing process, or port not responding).

- The word "up" sets back the server's operating state as UP if health checks also report that the service is accessible.

Parameters which are not advertised by the agent are not changed. For example, an agent might be designed to monitor CPU usage and only report a relative weight and never interact with the operating status. Similarly, an agent could be designed as an end-user interface with 3 radio buttons allowing an administrator to change only the administrative state. However, it is important to consider that only the agent may revert its own actions, so if a server is set to DRAIN mode or to DOWN state using the agent, the agent must implement the other equivalent actions to bring the service into operations again.

Failure to connect to the agent is not considered an error as connectivity is tested by the regular health check which is enabled by the "check" parameter. Warning though, it is not a good idea to stop an agent after it reports "down", since only an agent reporting "up" will be able to turn the server up again. Note that the CLI on the Unix stats socket is also able to force an agent's result in order to workaround a bogus agent if needed.

Requires the "agent-port" parameter to be set. See also the "agent-inter" parameter.

Supported in default-server: No

agent-send <string>

If this option is specified, haproxy will send the given string (verbatim) to the agent server upon connection. You could, for example, encode

the backend name into this string, which would enable your agent to send different responses based on the backend. Make sure to include a '\n' if you want to terminate your request with a newline.

agent-inter <delay>

The "agent-inter" parameter sets the interval between two agent checks to <delay> milliseconds. If left unspecified, the delay defaults to 2000 ms.

Just as with every other time-based parameter, it may be entered in any other explicit unit among { us, ms, s, m, h, d }. The "agent-inter" parameter also serves as a timeout for agent checks "timeout check" is not set. In order to reduce "resonance" effects when multiple servers are hosted on the same hardware, the agent and health checks of all servers are started with a small time offset between them. It is also possible to add some random noise in the agent and health checks interval using the global "spread-checks" keyword. This makes sense for instance when a lot of backends use the same servers.

See also the "agent-check" and "agent-port" parameters.

Supported in default-server: Yes

agent-port <port>

The "agent-port" parameter sets the TCP port used for agent checks.

See also the "agent-check" and "agent-inter" parameters.

Supported in default-server: Yes

backup

When "backup" is present on a server line, the server is only used in load balancing when all other non-backup servers are unavailable. Requests coming with a persistence cookie referencing the server will always be served though. By default, only the first operational backup server is used, unless the "allbackups" option is set in the backend. See also the "allbackups" option.

Supported in default-server: No

ca-file <cafile>

This setting is only available when support for OpenSSL was built in. It designates a PEM file from which to load CA certificates used to verify server's certificate.

Supported in default-server: No

check

This option enables health checks on the server. By default, a server is always considered available. If "check" is set, the server is available when accepting periodic TCP connections, to ensure that it is really able to serve requests. The default address and port to send the tests to are those of the server, and the default source is the same as the one defined in the backend. It is possible to change the address using the "addr" parameter, the port using the "port" parameter, the source address using the "source" address, and the interval and timers using the "inter", "rise" and "fall" parameters. The request method is define in the backend using the "httpchk", "smtpchk", "mysql-check", "pgsql-check" and "ssl-hello-chk" options. Please refer to those options and parameters for more information.

Supported in default-server: No

check-send-proxy

This option forces emission of a PROXY protocol line with outgoing health checks, regardless of whether the server uses send-proxy or not for the normal traffic. By default, the PROXY protocol is enabled for health checks

if it is already enabled for normal traffic and if no "port" nor "addr" directive is present. However, if such a directive is present, the "check-send-proxy" option needs to be used to force the use of the protocol. See also the "send-proxy" option for more information.

Supported in default-server: No

check-ssl

This option forces encryption of all health checks over SSL, regardless of whether the server uses SSL or not for the normal traffic. This is generally used when an explicit "port" or "addr" directive is specified and SSL health checks are not inherited. It is important to understand that this option inserts an SSL transport layer below the checks, so that a simple TCP connect check becomes an SSL connect, which replaces the old ssl-hello-chk. The most common use is to send HTTPS checks by combining "httpchk" with SSL checks. All SSL settings are common to health checks and traffic (eg: ciphers). See the "ssl" option for more information.

Supported in default-server: No

ciphers <ciphers>

This option sets the string describing the list of cipher algorithms that is negotiated during the SSL/TLS handshake with the server. The format of the string is defined in "man 1 ciphers". When SSL is used to communicate with servers on the local network, it is common to see a weaker set of algorithms than what is used over the internet. Doing so reduces CPU usage on both the server and haproxy while still keeping it compatible with deployed software. Some algorithms such as RC4-SHA1 are reasonably cheap. If no security at all is needed and just connectivity, using DES can be appropriate.

Supported in default-server: No

cookie <value>

The "cookie" parameter sets the cookie value assigned to the server to <value>. This value will be checked in incoming requests, and the first operational server possessing the same value will be selected. In return, in cookie insertion or rewrite modes, this value will be assigned to the cookie sent to the client. There is nothing wrong in having several servers sharing the same cookie value, and it is in fact somewhat common between normal and backup servers. See also the "cookie" keyword in backend section.

Supported in default-server: No

crl-file <crlfile>

This setting is only available when support for OpenSSL was built in. It designates a PEM file from which to load certificate revocation list used to verify server's certificate.

Supported in default-server: No

crt <cert>

This setting is only available when support for OpenSSL was built in. It designates a PEM file from which to load both a certificate and the associated private key. This file can be built by concatenating both PEM files into one. This certificate will be sent if the server send a client certificate request.

Supported in default-server: No

disabled

The "disabled" keyword starts the server in the "disabled" state. That means that it is marked down in maintenance mode, and no connection other than the ones allowed by persist mode will reach it. It is very well suited to setup new servers, because normal traffic will never reach them, while it is still possible to test the service by making use of the force-persist mechanism.

10271 Supported in default-server: No

10272

10273 error-limit <count>

10274 If health observing is enabled, the "error-limit" parameter specifies the

10275 number of consecutive errors that triggers event selected by the "on-error"

10276 option. By default it is set to 10 consecutive errors.

10277

10278 Supported in default-server: Yes

10279

10280

10281 See also the "check", "error-limit" and "on-error".

10282

10283 fall <count>

10284 The "fall" parameter states that a server will be considered as dead after

10285 <count> consecutive unsuccessful health checks. This value defaults to 3 if

10286 unspecified. See also the "check", "inter", and "rise" parameters.

10287

10288 Supported in default-server: Yes

10289

10290 force-ssl v3

10291 This option enforces use of SSLv3 only when SSL is used to communicate with

10292 the server. SSLv3 is generally less expensive than the TLS counterparts for

10293 high connection rates. This option is also available on global statement

10294 "ssl-default-server-options". See also "no-tls*", "no-ssl v3".

10295

10296 Supported in default-server: No

10297

10298 force-tls v10

10299 This option enforces use of TLSv1.0 only when SSL is used to communicate with

10300 the server. This option is also available on global statement

10301 "ssl-default-server-options". See also "no-tls*", "no-ssl v3".

10302

10303 Supported in default-server: No

10304

10305 force-tls v11

10306 This option enforces use of TLSv1.1 only when SSL is used to communicate with

10307 the server. This option is also available on global statement

10308 "ssl-default-server-options". See also "no-tls*", "no-ssl v3".

10309

10310 Supported in default-server: No

10311

10312 force-tls v12

10313 This option enforces use of TLSv1.2 only when SSL is used to communicate with

10314 the server. This option is also available on global statement

10315 "ssl-default-server-options". See also "no-tls*", "no-ssl v3".

10316

10317 Supported in default-server: No

10318

10319 id <value>

10320 Set a persistent ID for the server. This ID must be positive and unique for

10321 the proxy. An unused ID will automatically be assigned if unset. The first

10322 assigned value will be 1. This ID is currently only returned in statistics.

10323

10324 Supported in default-server: No

10325

10326 inter <delay>

10327 fastinter <delay>

10328 downinter <delay>

10329 The "inter" parameter sets the interval between two consecutive health checks

10330 to <delay> milliseconds. If left unspecified, the delay defaults to 2000 ms.

10331 It is also possible to use "fastinter" and "downinter" to optimize delays

10332 between checks depending on the server state :

10333

10334 Server state | Interval used

10335 -----+-----

10336 UP 100% (non-transitional) | "inter"

10337 -----+-----

10338 Transitionally UP (going down "fall"), | "fastinter" if set,

10339 Transitionally DOWN (going up "rise"), | "inter" otherwise.

10340 or yet unchecked. |

10341 -----+-----

10342 DOWN 100% (non-transitional) | "downinter" if set,

10343 | "inter" otherwise.

10344 -----+-----

10345

10346 Just as with every other time-based parameter, they can be entered in any

10347 other explicit unit among { us, ms, s, m, h, d }. The "inter" parameter also

10348 serves as a timeout for health checks sent to servers if "timeout check" is

10349 not set. In order to reduce "resonance" effects when multiple servers are

10350 hosted on the same hardware, the agent and health checks of all servers

10351 are started with a small time offset between them. It is also possible to

10352 add some random noise in the agent and health checks interval using the

10353 global "spread-checks" keyword. This makes sense for instance when a lot

10354 of backends use the same servers.

10355

10356 Supported in default-server: Yes

10357

10358 maxconn <maxconn>

10359 The "maxconn" parameter specifies the maximal number of concurrent

10360 connections that will be sent to this server. If the number of incoming

10361 concurrent requests goes higher than this value, they will be queued, waiting

10362 for a connection to be released. This parameter is very important as it can

10363 save fragile servers from going down under extreme loads. If a "minconn"

10364 parameter is specified, the limit becomes dynamic. The default value is "0"

10365 which means unlimited. See also the "minconn" and "maxqueue" parameters, and

10366 the backend's "fullconn" keyword.

10367

10368 Supported in default-server: Yes

10369

10370 maxqueue <maxqueue>

10371 The "maxqueue" parameter specifies the maximal number of connections which

10372 will wait in the queue for this server. If this limit is reached, next

10373 requests will be redispached to other servers instead of indefinitely

10374 waiting to be served. This will break persistence but may allow people to

10375 quickly re-log in when the server they try to connect to is dying. The

10376 default value is "0" which means the queue is unlimited. See also the

10377 "maxconn" and "minconn" parameters.

10378

10379 Supported in default-server: Yes

10380

10381 minconn <minconn>

10382 When the "minconn" parameter is set, the maxconn limit becomes a dynamic

10383 limit following the backend's load. The server will always accept at least

10384 <minconn> connections, never more than <maxconn>, and the limit will be on

10385 the ramp between both values when the backend has less than <fullconn>

10386 concurrent connections. This makes it possible to limit the load on the

10387 server during normal loads, but push it further for important loads without

10388 overloading the server during exceptional loads. See also the "maxconn"

10389 and "maxqueue" parameters, as well as the "fullconn" backend keyword.

10390

10391 Supported in default-server: Yes

10392

10393 namespace <name>

10394 On Linux, it is possible to specify which network namespace a socket will

10395 belong to. This directive makes it possible to explicitly bind a server to

10396 a namespace different from the default one. Please refer to your operating

10397 system's documentation to find more details about network namespaces.

10398

10399 no-ssl-reuse

10400 This option disables SSL session reuse when SSL is used to communicate with

the server. It will force the server to perform a full handshake for every new connection. It's probably only useful for benchmarking, troubleshooting, and for paranoid users.

Supported in default-server: No

no-sslV3

This option disables support for SSLv3 when SSL is used to communicate with the server. Note that SSLv2 is disabled in the code and cannot be enabled using any configuration option. See also "force-sslV3", "force-tlsV*".

Supported in default-server: No

no-tls-tickets

This setting is only available when support for OpenSSL was built in. It disables the stateless session resumption (RFC 5077 TLS Ticket extension) and force to use stateful session resumption. Stateless session resumption is more expensive in CPU usage for servers. This option is also available on global statement "ssl-default-server-options".

Supported in default-server: No

no-tlsV10

This option disables support for TLSv1.0 when SSL is used to communicate with the server. Note that SSLv2 is disabled in the code and cannot be enabled using any configuration option. TLSv1 is more expensive than SSLv3 so it often makes sense to disable it when communicating with local servers. This option is also available on global statement "ssl-default-server-options". See also "force-sslV3", "force-tlsV*".

Supported in default-server: No

no-tlsV11

This option disables support for TLSv1.1 when SSL is used to communicate with the server. Note that SSLv2 is disabled in the code and cannot be enabled using any configuration option. TLSv1 is more expensive than SSLv3 so it often makes sense to disable it when communicating with local servers. This option is also available on global statement "ssl-default-server-options". See also "force-sslV3", "force-tlsV*".

Supported in default-server: No

no-tlsV12

This option disables support for TLSv1.2 when SSL is used to communicate with the server. Note that SSLv2 is disabled in the code and cannot be enabled using any configuration option. TLSv1 is more expensive than SSLv3 so it often makes sense to disable it when communicating with local servers. This option is also available on global statement "ssl-default-server-options". See also "force-sslV3", "force-tlsV*".

Supported in default-server: No

non-stick

Never add connections allocated to this server to a stick-table. This may be used in conjunction with backup to ensure that stick-table persistence is disabled for backup servers.

Supported in default-server: No

observe <mode>

This option enables health adjusting based on observing communication with the server. By default this functionality is disabled and enabling it also requires to enable health checks. There are two supported modes: "layer4" and "layer7". In layer4 mode, only successful/unsuccessful tcp connections are significant. In layer7, which is only allowed for http proxies, responses

received from server are verified, like valid/wrong http code, unparsable headers, a timeout, etc. Valid status codes include 100 to 499, 501 and 505.

Supported in default-server: No

See also the "check", "on-error" and "error-limit".

on-error <mode>

Select what should happen when enough consecutive errors are detected.

Currently, four modes are available:

- fastinter: force fastinter
- fail-check: simulate a failed check, also forces fastinter (default)
- sudden-death: simulate a pre-fatal failed health check, one more failed check will mark a server down, forces fastinter
- mark-down: mark the server immediately down and force fastinter

Supported in default-server: Yes

See also the "check", "observe" and "error-limit".

on-marked-down <action>

Modify what occurs when a server is marked down.

Currently one action is available:

- shutdown-sessions: Shutdown peer sessions. When this setting is enabled, all connections to the server are immediately terminated when the server goes down. It might be used if the health check detects more complex cases than a simple connection status, and long timeouts would cause the service to remain unresponsive for too long a time. For instance, a health check might detect that a database is stuck and that there's no chance to reuse existing connections anymore. Connections killed this way are logged with a 'D' termination code (for "Down").

Actions are disabled by default

Supported in default-server: Yes

on-marked-up <action>

Modify what occurs when a server is marked up.

Currently one action is available:

- shutdown-backup-sessions: Shutdown sessions on all backup servers. This is done only if the server is not in backup state and if it is not disabled (it must have an effective weight > 0). This can be used sometimes to force an active server to take all the traffic back after recovery when dealing with long sessions (eg: LDAP, SQL, ...). Doing this can cause more trouble than it tries to solve (eg: incomplete transactions), so use this feature with extreme care. Sessions killed because a server comes up are logged with an 'U' termination code (for "Up").

Actions are disabled by default

Supported in default-server: Yes

port <port>

Using the "port" parameter, it becomes possible to use a different port to send health-checks. On some servers, it may be desirable to dedicate a port to a specific component able to perform complex tests which are more suitable to health-checks than the application. It is common to run a simple script in inetd for instance. This parameter is ignored if the "check" parameter is not set. See also the "addr" parameter.

Supported in default-server: Yes

redir <prefix>

The "redir" parameter enables the redirection mode for all GET and HEAD requests addressing this server. This means that instead of having HAProxy

forward the request to the server, it will send an "HTTP 302" response with the "Location" header composed of this prefix immediately followed by the requested URI beginning at the leading '/' of the path component. That means that no trailing slash should be used after <prefix>. All invalid requests will be rejected, and all non-GET or HEAD requests will be normally served by the server. Note that since the response is completely forged, no header mangling nor cookie insertion is possible in the response. However, cookies in requests are still analysed, making this solution completely usable to direct users to a remote location in case of local disaster. Main use consists in increasing bandwidth for static servers by having the clients directly connect to them. Note: never use a relative location here, it would cause a loop between the client and HAProxy!

Example : server srv1 192.168.1.1:80 redir http://image1.mydomain.com check

Supported in default-server: No

rise <count>

The "rise" parameter states that a server will be considered as operational after <count> consecutive successful health checks. This value defaults to 2 if unspecified. See also the "check", "inter" and "fall" parameters.

Supported in default-server: Yes

resolve-prefer <family>

When DNS resolution is enabled for a server and multiple IP addresses from different families are returned, HAProxy will prefer using an IP address from the family mentioned in the "resolve-prefer" parameter. Available families: "ipv4" and "ipv6"

Default value: ipv6

Supported in default-server: Yes

Example: server s1 appl.domain.com:80 resolvers mydns resolve-prefer ipv6

resolvers <id>

Points to an existing "resolvers" section to resolve current server's hostname.

In order to be operational, DNS resolution requires that health check is enabled on the server. Actually, health checks triggers the DNS resolution. You must precise one 'resolvers' parameter on each server line where DNS resolution is required.

Supported in default-server: No

Example: server s1 appl.domain.com:80 check resolvers mydns

See also chapter 5.3

send-proxy

The "send-proxy" parameter enforces use of the PROXY protocol over any connection established to this server. The PROXY protocol informs the other end about the layer 3/4 addresses of the incoming connection, so that it can know the client's address or the public address it accessed to, whatever the upper layer protocol. For connections accepted by an "accept-proxy" listener, the advertised address will be used. Only TCPv4 and TCPv6 address families are supported. Other families such as Unix sockets, will report an UNKNOWN family. Servers using this option can fully be chained to another instance of haproxy listening with an "accept-proxy" setting. This setting must not be used if the server isn't aware of the protocol. When health checks are sent to the server, the PROXY protocol is automatically used when this option is set, unless there is an explicit "port" or "addr" directive, in which case an explicit "check-send-proxy" directive would also be needed to use the PROXY protocol. See also the "accept-proxy" option of the "bind" keyword.

10595

Supported in default-server: No

send-proxy-v2

The "send-proxy-v2" parameter enforces use of the PROXY protocol version 2 over any connection established to this server. The PROXY protocol informs the other end about the layer 3/4 addresses of the incoming connection, so that it can know the client's address or the public address it accessed to, whatever the upper layer protocol. This setting must not be used if the server isn't aware of this version of the protocol. See also the "send-proxy" option of the "bind" keyword.

Supported in default-server: No

send-proxy-v2-ssl

The "send-proxy-v2-ssl" parameter enforces use of the PROXY protocol version 2 over any connection established to this server. The PROXY protocol informs the other end about the layer 3/4 addresses of the incoming connection, so that it can know the client's address or the public address it accessed to, whatever the upper layer protocol. In addition, the SSL information extension of the PROXY protocol is added to the PROXY protocol header. This setting must not be used if the server isn't aware of this version of the protocol. See also the "send-proxy-v2" option of the "bind" keyword.

Supported in default-server: No

send-proxy-v2-ssl-cn

The "send-proxy-v2-ssl" parameter enforces use of the PROXY protocol version 2 over any connection established to this server. The PROXY protocol informs the other end about the layer 3/4 addresses of the incoming connection, so that it can know the client's address or the public address it accessed to, whatever the upper layer protocol. In addition, the SSL information extension of the PROXY protocol, along with the Common Name from the subject of the client certificate (if any), is added to the PROXY protocol header. This setting must not be used if the server isn't aware of this version of the protocol. See also the "send-proxy-v2" option of the "bind" keyword.

Supported in default-server: No

slowstart <start time in ms>

The "slowstart" parameter for a server accepts a value in milliseconds which indicates after how long a server which has just come back up will run at full speed. Just as with every other time-based parameter, it can be entered in any other explicit unit among { us, ms, s, m, h, d }. The speed grows linearly from 0 to 100% during this time. The limitation applies to two parameters :

- maxconn: the number of connections accepted by the server will grow from 1 to 100% of the usual dynamic limit defined by (minconn,maxconn,fullconn).

- weight: when the backend uses a dynamic weighted algorithm, the weight grows linearly from 1 to 100%. In this case, the weight is updated at every health-check. For this reason, it is important that the "inter" parameter is smaller than the "slowstart", in order to maximize the number of steps.

The slowstart never applies when haproxy starts, otherwise it would cause trouble to running servers. It only applies when a server has been previously seen as failed.

Supported in default-server: Yes

sni <expression>

The "sni" parameter evaluates the sample fetch expression, converts it to a string and uses the result as the host name sent in the SNI TLS extension to the server. A typical use case is to send the SNI received from the client in

10660

a bridged HTTPS scenario, using the "ssl_fc_sni" sample fetch for the expression, though alternatives such as req.hdr(host) can also make sense.

Supported in default-server: no

source <addr>[:<sp>[:<-ph>]] [usesrc { <addr2>[:<port2>] | client | clientip }]
source <addr>[:<port>] [usesrc { <addr2>[:<port2>] | hdr_ip(<hdr>[:<occ>]) }]
source <addr>[:<sp>[:<-sph>]] [interface <name>] ...
The "source" parameter sets the source address which will be used when connecting to the server. It follows the exact same parameters and principle as the backend "source" keyword, except that it only applies to the server referencing it. Please consult the "source" keyword for details.

Additionally, the "source" statement on a server line allows one to specify a source port range by indicating the lower and higher bounds delimited by a dash ('-'). Some operating systems might require a valid IP address when a source port range is specified. It is permitted to have the same IP/range for several servers. Doing so makes it possible to bypass the maximum of 64k total concurrent connections. The limit will then reach 64k connections per server.

Supported in default-server: No

ssl

This option enables SSL ciphering on outgoing connections to the server. It is critical to verify server certificates using "verify" when using SSL to connect to servers, otherwise the communication is prone to trivial man-in-the-middle attacks rendering SSL useless. When this option is used, health checks are automatically sent in SSL too unless there is a "port" or an "addr" directive indicating the check should be sent to a different location. See the "check-ssl" option to force SSL health checks.

Supported in default-server: No

tcp-out <delay>

Sets the TCP User Timeout for all outgoing connections to this server. This option is available on Linux since version 2.6.37. It allows haproxy to configure a timeout for sockets which contain data not receiving an acknowledgement for the configured delay. This is especially useful on long-lived connections experiencing long idle periods such as remote terminals or database connection pools, where the client and server timeouts must remain high to allow a long period of idle, but where it is important to detect that the server has disappeared in order to release all resources associated with its connection (and the client's session). One typical use case is also to force dead server connections to die when health checks are too slow or during a soft reload since health checks are then disabled. The argument is a delay expressed in milliseconds by default. This only works for regular TCP connections, and is ignored for other protocols.

track [<proxy>]/<server>

This option enables ability to set the current state of the server by tracking another one. It is possible to track a server which itself tracks another server, provided that at the end of the chain, a server has health checks enabled. If <proxy> is omitted the current one is used. If disable-on-404 is used, it has to be enabled on both proxies.

Supported in default-server: No

verify [none|required]

This setting is only available when support for OpenSSL was built in. If set to 'none', server certificate is not verified. In the other case, The certificate provided by the server is verified using CAs from 'ca-file' and optional CRLs from 'crl-file'. If 'ssl_server_verify' is not specified in global section, this is the default. On verify failure the handshake is aborted. It is critically important to verify server certificates when

using SSL to connect to servers, otherwise the communication is prone to trivial man-in-the-middle attacks rendering SSL totally useless.

Supported in default-server: No

verifyhost <hostname>

This setting is only available when support for OpenSSL was built in, and only takes effect if 'verify required' is also specified. When set, the hostnames in the subject and subjectAltNameNames of the certificate provided by the server are checked. If none of the hostnames in the certificate match the specified hostname, the handshake is aborted. The hostnames in the server-provided certificate may include wildcards.

Supported in default-server: No

weight <weight>

The "weight" parameter is used to adjust the server's weight relative to other servers. All servers will receive a load proportional to their weight relative to the sum of all weights, so the higher the weight, the higher the load. The default weight is 1, and the maximal value is 256. A value of 0 means the server will not participate in load-balancing but will still accept persistent connections. If this parameter is used to distribute the load according to server's capacity, it is recommended to start with values which can both grow and shrink, for instance between 10 and 100 to leave enough room above and below for later adjustments.

Supported in default-server: Yes

5.3. Server IP address resolution using DNS

HAProxy allows using a host name on the server line to retrieve its IP address using name servers. By default, HAProxy resolves the name when parsing the configuration file, at startup and cache the result for the process' life. This is not sufficient in some cases, such as in Amazon where a server's IP can change after a reboot or an ELB Virtual IP can change based on current workload.

This chapter describes how HAProxy can be configured to process server's name resolution at run time.

Whether run time server name resolution has been enable or not, HAProxy will carry on doing the first resolution when parsing the configuration.

Bear in mind that DNS resolution is triggered by health checks. This makes health checks mandatory to allow DNS resolution.

5.3.1. Global overview

As we've seen in introduction, name resolution in HAProxy occurs at two different steps of the process life:

1. when starting up, HAProxy parses the server line definition and matches a host name. It uses libc functions to get the host name resolved. This resolution relies on /etc/resolv.conf file.

2. at run time, when HAProxy gets prepared to run a health check on a server, it verifies if the current name resolution is still considered as valid. If not, it processes a new resolution, in parallel of the health check.

A few other events can trigger a name resolution at run time:

- when a server's health check ends up in a connection timeout: this may be because the server has a new IP address. So we need to trigger a name resolution to know this new IP.

10791
10792 A few things important to notice:
10793 - all the name servers are queried in the mean time. HAProxy will process the
10794 first valid response.
10795
10796 - a resolution is considered as invalid (NX, timeout, refused), when all the
10797 servers return an error.
10798
10799
10800 5.3.2. The resolvers section
10801 -----
10802
10803 This section is dedicated to host information related to name resolution in
10804 HAProxy.
10805 There can be as many as resolvers section as needed. Each section can contain
10806 many name servers.
10807
10808 When multiple name servers are configured in a resolvers section, then HAProxy
10809 uses the first valid response. In case of invalid responses, only the last one
10810 is treated. Purpose is to give the chance to a slow server to deliver a valid
10811 answer after a fast faulty or outdated server.
10812
10813 When each server returns a different error type, then only the last error is
10814 used by HAProxy to decide what type of behavior to apply.
10815
10816 Two types of behavior can be applied:
10817 1. stop DNS resolution
10818 2. replay the DNS query with a new query type
10819 In such case, the following types are applied in this exact order:
10820 1. ANY query type
10821 2. query type corresponding to family pointed by resolve-prefer
10822 server's parameter
10823 3. remaining family type
10824
10825 HAProxy stops DNS resolution when the following errors occur:
10826 - invalid DNS response packet
10827 - wrong name in the query section of the response
10828 - NX domain
10829 - Query refused by server
10830 - CNAME not pointing to an IP address
10831
10832 HAProxy tries a new query type when the following errors occur:
10833 - no Answer records in the response
10834 - DNS response truncated
10835 - Error in DNS response
10836 - No expected DNS records found in the response
10837 - name server timeout
10838
10839 For example, with 2 name servers configured in a resolvers section:
10840 - first response is valid and is applied directly, second response is ignored
10841 - first response is invalid and second one is valid, then second response is
10842 applied;
10843 - first response is a NX domain and second one a truncated response, then
10844 HAProxy replays the query with a new type;
10845 - first response is truncated and second one is a NX Domain, then HAProxy
10846 stops resolution.
10847
10848 resolvers <resolvers id>
10849 Creates a new name server list labelled <resolvers id>
10850
10851 A resolvers section accept the following parameters:
10852
10853 nameserver <id> <ip>:<port>
10854 DNS server description:
10855

10856 <id> : label of the server, should be unique
10857 <ip> : IP address of the server
10858 <port> : port where the DNS service actually runs
10859
10860 hold <status> <period>
10861 Defines <period> during which the last name resolution should be kept based
10862 on last resolution <status>
10863 <status> : last name resolution status. Only "valid" is accepted for now.
10864 <period> : interval between two successive name resolution when the last
10865 answer was in <status>. It follows the HAProxy time format.
10866 <period> is in milliseconds by default.
10867
10868 Default value is 10s for "valid".
10869
10870 Note: since the name resolution is triggered by the health checks, a new
10871 resolution is triggered after <period> modulo the <inter> parameter of
10872 the health check.
10873
10874 resolve_retries <nb>
10875 Defines the number <nb> of queries to send to resolve a server name before
10876 giving up.
10877 Default value: 3
10878
10879 A retry occurs on name server timeout or when the full sequence of DNS query
10880 type failover is over and we need to start up from the default ANY query
10881 type.
10882
10883 timeout <event> <time>
10884 Defines timeouts related to name resolution
10885 <event> : the event on which the <time> timeout period applies to.
10886 events available are:
10887 - retry: time between two DNS queries, when no response have
10888 been received.
10889 Default value: 1s
10890 <time> : time related to the event. It follows the HAProxy time format.
10891 <time> is expressed in milliseconds.
10892
10893 Example of a resolvers section (with default values):
10894
10895 resolvers mydns
10896 nameserver dns1 10.0.0.1:53
10897 nameserver dns2 10.0.0.2:53
10898 resolve_retries 3
10899 timeout retry 1s
10900 hold valid 10s
10901
10902
10903 6. HTTP header manipulation
10904 -----
10905
10906 In HTTP mode, it is possible to rewrite, add or delete some of the request and
10907 response headers based on regular expressions. It is also possible to block a
10908 request or a response if a particular header matches a regular expression,
10909 which is enough to stop most elementary protocol attacks, and to protect
10910 against information leak from the internal network.
10911
10912 If HAProxy encounters an "Informational Response" (status code 1xx), it is able
10913 to process all rsp* rules which can allow, deny, rewrite or delete a header,
10914 but it will refuse to add a header to any such messages as this is not
10915 HTTP-compliant. The reason for still processing headers in such responses is to
10916 stop and/or fix any possible information leak which may happen, for instance
10917 because another downstream equipment would unconditionally add a header, or if
10918 a server name appears there. When such messages are seen, normal processing
10919 still occurs on the next non-informational messages.
10920

This section covers common usage of the following keywords, described in detail in section 4.2 :

```
- reqadd <string>
- reqallow <search>
- reqallow <search>
- reqdel <search>
- reqdel <search>
- reqdeny <search>
- reqdeny <search>
- reqpass <search>
- reqpass <search>
- reqrep <search> <replace>
- reqrep <search> <replace>
- reqtarpit <search>
- reqtarpit <search>
- rspadd <string>
- rspdel <search>
- rspdeny <search>
- rspdeny <search>
- rsprep <search> <replace>
- rsprep <search> <replace>
```

With all these keywords, the same conventions are used. The <search> parameter is a POSIX extended regular expression (regex) which supports grouping through parenthesis (without the backslash). Spaces and other delimiters must be prefixed with a backslash ('\') to avoid confusion with a field delimiter.

Other characters may be prefixed with a backslash to change their meaning :

```
\t for a tab
\r for a carriage return (CR)
\n for a new line (LF)
\ to mark a space and differentiate it from a delimiter
\# to mark a sharp and differentiate it from a comment
\\ to use a backslash in a regex
\\\\ to use a backslash in the text (*2 for regex, *2 for haproxy)
\\XX to write the ASCII hex code XX as in the C language
```

The <replace> parameter contains the string to be used to replace the largest portion of text matching the regex. It can make use of the special characters above, and can reference a substring which is delimited by parenthesis in the regex, by writing a backslash ('\') immediately followed by one digit from 0 to 9 indicating the group position (0 designating the entire line). This practice is very common to users of the "sed" program.

The <string> parameter represents the string which will systematically be added after the last header line. It can also use special character sequences above.

Notes related to these keywords :

- these keywords are not always convenient to allow/deny based on header contents. It is strongly recommended to use ACLs with the "block" keyword instead, resulting in far more flexible and manageable rules.
- lines are always considered as a whole. It is not possible to reference a header name only or a value only. This is important because of the way headers are written (notably the number of spaces after the colon).
- the first line is always considered as a header, which makes it possible to rewrite or filter HTTP requests URIs or response codes, but in turn makes it harder to distinguish between headers and request line. The regex prefix ^[\ \t]*[\ \t] matches any HTTP method followed by a space, and the prefix ^[^\ \t]*: matches any header name followed by a colon.

- for performances reasons, the number of characters added to a request or to a response is limited at build time to values between 1 and 4 kB. This should normally be far more than enough for most usages. If it is too short on occasional usages, it is possible to gain some space by removing some useless headers before adding new ones.

- keywords beginning with "reqi" and "rspi" are the same as their counterpart without the 'i' letter except that they ignore case when matching patterns.
- when a request passes through a frontend then a backend, all req* rules from the frontend will be evaluated, then all req* rules from the backend will be evaluated. The reverse path is applied to responses.

- req* statements are applied after "block" statements, so that "block" is always the first one, but before "use_backend" in order to permit rewriting before switching.

7. Using ACLs and fetching samples

Haproxy is capable of extracting data from request or response streams, from client or server information, from tables, environmental information etc... The action of extracting such data is called fetching a sample. Once retrieved, these samples may be used for various purposes such as a key to a stick-table, but most common usages consist in matching them against predefined constant data called patterns.

7.1. ACL basics

The use of Access Control Lists (ACL) provides a flexible solution to perform content switching and generally to take decisions based on content extracted from the request, the response or any environmental status. The principle is simple :

- extract a data sample from a stream, table or the environment
- optionally apply some format conversion to the extracted sample
- apply one or multiple pattern matching methods on this sample
- perform actions only when a pattern matches the sample

The actions generally consist in blocking a request, selecting a backend, or adding a header.

In order to define a test, the "acl" keyword is used. The syntax is :

```
acl <aclname> <criterion> [flags] [operator] [<value>] ...
```

This creates a new ACL <aclname> or completes an existing one with new tests. Those tests apply to the portion of request/response specified in <criterion> and may be adjusted with optional flags [flags]. Some criteria also support an operator which may be specified before the set of values. Optionally some conversion operators may be applied to the sample, and they will be specified as a comma-delimited list of keywords just after the first keyword. The values are of the type supported by the criterion, and are separated by spaces.

ACL names must be formed from upper and lower case letters, digits, '-' (dash), '.' (underscore), '.' (dot) and ':' (colon). ACL names are case-sensitive, which means that "my_acl" and "My_Acl" are two different ACLs.

There is no enforced limit to the number of ACLs. The unused ones do not affect performance, they just consume a small amount of memory.

The criterion generally is the name of a sample fetch method, or one of its ACL

Specific declinations. The default test method is implied by the output type of this sample fetch method. The ACL declinations can describe alternate matching methods of a same sample fetch method. The sample fetch methods are the only ones supporting a conversion.

Sample fetch methods return data which can be of the following types :

- boolean
- integer (signed or unsigned)
- IPv4 or IPv6 address
- string
- data block

Converters transform any of these data into any of these. For example, some converters might convert a string to a lower-case string while other ones would turn a string to an IPv4 address, or apply a netmask to an IP address. The resulting sample is of the type of the last converter applied to the list, which defaults to the type of the sample fetch method.

Each sample or converter returns data of a specific type, specified with its keyword in this documentation. When an ACL is declared using a standard sample fetch method, certain types automatically involved a default matching method which are summarized in the table below :

Sample or converter	Default output type	matching method
boolean	bool	
integer	int	
ip	ip	
string	str	
binary	none, use "-m"	

Note that in order to match a binary samples, it is mandatory to specify a matching method, see below.

The ACL engine can match these types against patterns of the following types :

- boolean
- integer or integer range
- IP address / network
- string (exact, substring, suffix, prefix, subdir, domain)
- regular expression
- hex block

The following ACL flags are currently supported :

- i : ignore case during matching of all subsequent patterns.
- f : load patterns from a file.
- m : use a specific pattern matching method
- n : forbid the DNS resolutions
- M : load the file pointed by -f like a map file.
- u : force the unique id of the ACL
- : force end of flags. Useful when a string looks like one of the flags.

The "-f" flag is followed by the name of a file from which all lines will be read as individual values. It is even possible to pass multiple "-f" arguments if the patterns are to be loaded from multiple files. Empty lines as well as lines beginning with a sharp (#) will be ignored. All leading spaces and tabs will be stripped. If it is absolutely necessary to insert a valid pattern beginning with a sharp, just prefix it with a space so that it is not taken for

a comment. Depending on the data type and match method, haproxy may load the lines into a binary tree, allowing very fast lookups. This is true for IPv4 and exact string matching. In this case, duplicates will automatically be removed.

The "-M" flag allows an ACL to use a map file. If this flag is set, the file is parsed as two column file. The first column contains the patterns used by the ACL, and the second column contain the samples. The sample can be used later by a map. This can be useful in some rare cases where an ACL would just be used to check for the existence of a pattern in a map before a mapping is applied.

The "-u" flag forces the unique id of the ACL. This unique id is used with the socket interface to identify ACL and dynamically change its values. Note that a file is always identified by its name even if an id is set.

Also, note that the "-i" flag applies to subsequent entries and not to entries loaded from files preceding it. For instance :

```
acl valid-ua hdr(user-agent) -f exact-ua.lst -i -f generic-ua.lst test
```

In this example, each line of "exact-ua.lst" will be exactly matched against the "user-agent" header of the request. Then each line of "generic-ua" will be case-insensitively matched. Then the word "test" will be insensitively matched as well.

The "-m" flag is used to select a specific pattern matching method on the input sample. All ACL-specific criteria imply a pattern matching method and generally do not need this flag. However, this flag is useful with generic sample fetch methods to describe how they're going to be matched against the patterns. This is required for sample fetches which return data type for which there is no obvious matching method (eg: string or binary). When "-m" is specified and followed by a pattern matching method name, this method is used instead of the default one for the criterion. This makes it possible to match contents in ways that were not initially planned, or with sample fetch methods which return a string. The matching method also affects the way the patterns are parsed.

The "-n" flag forbids the dns resolutions. It is used with the load of ip files. By default, if the parser cannot parse ip address it considers that the parsed string is maybe a domain name and try dns resolution. The flag "-n" disable this resolution. It is useful for detecting malformed ip lists. Note that if the DNS server is not reachable, the haproxy configuration parsing may last many minutes waiting fir the timeout. During this time no error messages are displayed. The flag "-n" disable this behavior. Note also that during the runtime, this function is disabled for the dynamic acl modifications.

There are some restrictions however. Not all methods can be used with all sample fetch methods. Also, if "-m" is used in conjunction with "-f", it must be placed first. The pattern matching method must be one of the following :

- "found" : only check if the requested sample could be found in the stream, but do not compare it against any pattern. It is recommended not to pass any pattern to avoid confusion. This matching method is particularly useful to detect presence of certain contents such as headers, cookies, etc... even if they are empty and without comparing them to anything nor counting them.

- "bool" : check the value as a boolean. It can only be applied to fetches which return a boolean or integer value, and takes no pattern. Value zero or false does not match, all other values do match.

- "int" : match the value as an integer. It can be used with integer and boolean samples. Boolean false is integer 0, true is integer 1.

- "ip" : match the value as an IPv4 or IPv6 address. It is compatible with IP address samples only, so it is implied and never needed.

- "bin" : match the contents against an hexadecimal string representing a binary sequence. This may be used with binary or string samples.
- "len" : match the sample's length as an integer. This may be used with binary or string samples.
- "str" : exact match : match the contents against a string. This may be used with binary or string samples.
- "sub" : substring match : check that the contents contain at least one of the provided string patterns. This may be used with binary or string samples.
- "reg" : regex match : match the contents against a list of regular expressions. This may be used with binary or string samples.
- "beg" : prefix match : check that the contents begin like the provided string patterns. This may be used with binary or string samples.
- "end" : suffix match : check that the contents end like the provided string patterns. This may be used with binary or string samples.
- "dir" : subdir match : check that a slash-delimited portion of the contents exactly matches one of the provided string patterns. This may be used with binary or string samples.
- "dom" : domain match : check that a dot-delimited portion of the contents exactly match one of the provided string patterns. This may be used with binary or string samples.

For example, to quickly detect the presence of cookie "JSESSIONID" in an HTTP request, it is possible to do :

```
aclickess present cook(JSESSIONID) -m found
```

In order to apply a regular expression on the 500 first bytes of data in the buffer, one would use the following acl :

```
acl script tag payload(0,500) -m reg -i <script>
```

On systems where the regex library is much slower when using "-i", it is possible to convert the sample to lowercase before matching, like this :

```
acl script tag payload(0,500),lower -m req <script>
```

All ACL-specific criteria imply a default matching method. Most often, these criteria are composed by concatenating the name of the original sample fetch method and the matching method. For example, "hdr_beg" applies the "beg" match to samples retrieved using the "hdr" fetch method. Since all ACL-specific criteria rely on a sample fetch method, it is always possible instead to use the original sample fetch method and the explicit matching method using "-m".

If an alternate match is specified using "-m" on an ACL-specific criterion, the matching method is simply applied to the underlying sample fetch method. For example, all ACLs below are exact equivalent :

acl short_form	hdr_beg(host)	www.
acl alternate1	hdr_beg(host)	-m beg www.
acl alternate2	hdr_dom(host)	-m beg www.
acl alternate3	hdr(host)	-m beg www.

The table below summarizes the compatibility matrix between sample or converter types and the pattern types to fetch against. It indicates for each compatible combination the name of the matching method to be used, surrounded with angle

brackets ">" and "<" when the method is the default one and will work by default without "-m".

	Input sample type									
pattern type	boolean	integer	ip	string	binary					
none (presence only)	found	found	found	found	found					
none (boolean value)	> bool <	bool	bool	bool	bool					
integer (value)	int	> int <	int	int	int					
integer (length)	len	len	len	len	len					
IP address			> ip <	ip	ip					
exact string	str	str	str	> str <	str					
prefix	beg	beg	beg	beg	beg					
suffix	end	end	end	end	end					
substring	sub	sub	sub	sub	sub					
subdir	dir	dir	dir	dir	dir					
domain	dom	dom	dom	dom	dom					
regex	reg	reg	reg	reg	reg					
hex block				bin	bin					

7.1.1. Matching booleans

In order to match a boolean, no value is needed and all values are ignored. Boolean matching is used by default for all fetch methods of type "boolean". When boolean matching is used, the fetched value is returned as-is, which means that a boolean "true" will always match and a boolean "false" will never match.

Boolean matching may also be enforced using "-m bool" on fetch methods which return an integer value. Then, integer value 0 is converted to the boolean "false" and all other values are converted to "true".

7.1.2. Matching integers

Integer matching applies by default to integer fetch methods. It can also be enforced on boolean fetches using "-m int". In this case, "false" is converted to the integer 0, and "true" is converted to the integer 1.

Integer matching also supports integer ranges and operators. Note that integer matching only applies to positive values. A range is a value expressed with a lower and an upper bound separated with a colon, both of which may be omitted.

For instance, "1024:65535" is a valid range to represent a range of unprivileged ports, and "1024:" would also work. "0:1023" is a valid representation of privileged ports, and ":1023" would also work.

As a special case, some ACL functions support decimal numbers which are in fact

two integers separated by a dot. This is used with some version checks for instance. All integer properties apply to those decimal numbers, including ranges and operators.

For an easier usage, comparison operators are also supported. Note that using operators with ranges does not make much sense and is strongly discouraged. Similarly, it does not make much sense to perform order comparisons with a set of values.

Available operators for integer matching are :

```
eq : true if the tested value equals at least one value
ge : true if the tested value is greater than or equal to at least one value
gt : true if the tested value is greater than at least one value
le : true if the tested value is less than or equal to at least one value
lt : true if the tested value is less than at least one value
```

For instance, the following ACL matches any negative Content-Length header :

```
acl negative-length hdr_val(content-length) lt 0
```

This one matches SSL versions between 3.0 and 3.1 (inclusive) :

```
acl sslv3 req_ssl_ver 3:3.1
```

7.1.3. Matching strings

String matching applies to string or binary fetch methods, and exists in 6 different forms :

- exact match (-m str) : the extracted string must exactly match the patterns ;
- substring match (-m sub) : the patterns are looked up inside the extracted string, and the ACL matches if any of them is found inside ;
- prefix match (-m beg) : the patterns are compared with the beginning of the extracted string, and the ACL matches if any of them matches.
- suffix match (-m end) : the patterns are compared with the end of the extracted string, and the ACL matches if any of them matches.
- subdir match (-m sub) : the patterns are looked up inside the extracted string, delimited with slashes ("/"), and the ACL matches if any of them matches.
- domain match (-m dom) : the patterns are looked up inside the extracted string, delimited with dots ("."), and the ACL matches if any of them matches.

String matching applies to verbatim strings as they are passed, with the exception of the backslash ("\") which makes it possible to escape some characters such as the space. If the "-i" flag is passed before the first string, then the matching will be performed ignoring the case. In order to match the string "-i", either set it second, or pass the "--" flag before the first string. Same applies of course to match the string "--".

7.1.4. Matching regular expressions (regexes)

Just like with string matching, regex matching applies to verbatim strings as they are passed, with the exception of the backslash ("\") which makes it

possible to escape some characters such as the space. If the "-i" flag is passed before the first regex, then the matching will be performed ignoring the case. In order to match the string "-i", either set it second, or pass the "--" flag before the first string. Same principle applies of course to match the string "--".

7.1.5. Matching arbitrary data blocks

It is possible to match some extracted samples against a binary block which may not safely be represented as a string. For this, the patterns must be passed as a series of hexadecimal digits in an even number, when the match method is set to binary. Each sequence of two digits will represent a byte. The hexadecimal digits may be used upper or lower case.

Example :

```
# match "Hello\n" in the input stream (\x48 \x65 \x6c \x6f \x0a)
acl hello payload(0,6) -m bin 48656c6cf0a
```

7.1.6. Matching IPv4 and IPv6 addresses

IPv4 addresses values can be specified either as plain addresses or with a netmask appended, in which case the IPv4 address matches whenever it is within the network. Plain addresses may also be replaced with a resolvable host name, but this practice is generally discouraged as it makes it more difficult to read and debug configurations. If hostnames are used, you should at least ensure that they are present in /etc/hosts so that the configuration does not depend on any random DNS match at the moment the configuration is parsed.

IPv6 may be entered in their usual form, with or without a netmask appended. Only bit counts are accepted for IPv6 netmasks. In order to avoid any risk of trouble with randomly resolved IP addresses, host names are never allowed in IPv6 patterns.

HAProxy is also able to match IPv4 addresses with IPv6 addresses in the following situations :

- tested address is IPv4, pattern address is IPv4, the match applies in IPv4 using the supplied mask if any.
- tested address is IPv6, pattern address is IPv6, the match applies in IPv6 using the supplied mask if any.
- tested address is IPv6, pattern address is IPv4, the match applies in IPv4 using the pattern's mask if the IPv6 address matches with 2002::IPV4:::IPV4 or ::ffff:IPV4, otherwise it fails.
- tested address is IPv4, pattern address is IPv6, the IPv4 address is first converted to IPv6 by prefixing ::ffff: in front of it, then the match is applied in IPv6 using the supplied IPv6 mask.

7.2. Using ACLs to form conditions

Some actions are only performed upon a valid condition. A condition is a combination of ACLs with operators. 3 operators are supported :

- AND (implicit)
- OR (explicit with the "or" keyword or the "||" operator)
- Negation with the exclamation mark ("!")

A condition is formed as a disjunctive form:

```
[!]acl1 [!]acl2 ... [!]acln { or [!]acl1 [!]acl2 ... [!]acln } ...
```



```
11441 Such conditions are generally used after an "if" or "unless" statement,
11442 indicating when the condition will trigger the action.
11443
11444 For instance, to block HTTP requests to the "*" URL with methods other than
11445 "OPTIONS", as well as POST requests without content-length, and GET or HEAD
11446 requests with a content-length greater than 0, and finally every request which
11447 is not either GET/HEAD/POST/OPTIONS !
11448
11449     acl missing_cl hdr_cnt(Content-length) eq 0
11450     block if HTTP_URL_STAR !METH_OPTIONS || METH_POST missing_cl
11451     block if METH_GET HTTP_CONTENT
11452     block unless METH_GET or METH_POST or METH_OPTIONS
11453
11454 To select a different backend for requests to static contents on the "www" site
11455 and to every request on the "img", "video", "download" and "ftp" hosts :
11456
11457     acl url_static path_beg          /static/images /img /css
11458     acl url_static path_end          .gif .png .jpg .css .js
11459     acl host_www    hdr_beg(host) -i www
11460     acl host_static hdr_beg(host) -i img. video. download. ftp.
11461
11462     # now use backend "static" for all static-only hosts, and for static urls
11463     # of host "www". Use backend "www" for the rest.
11464     use_backend static if host_static or host_www url_static
11465     use_backend www    if host_www
11466
11467 It is also possible to form rules using "anonymous ACLs". Those are unnamed ACL
11468 expressions that are built on the fly without needing to be declared. They must
11469 be enclosed between braces, with a space before and after each brace (because
11470 the braces must be seen as independent words). Example :
11471
11472     The following rule :
11473
11474         acl missing_cl hdr_cnt(Content-length) eq 0
11475         block if METH_POST missing_cl
11476
11477     Can also be written that way :
11478
11479         block if METH_POST { hdr_cnt(Content-length) eq 0 }
11480
11481 It is generally not recommended to use this construct because it's a lot easier
11482 to leave errors in the configuration when written that way. However, for very
11483 simple rules matching only one source IP address for instance, it can make more
11484 sense to use them than to declare ACLs with random names. Another example of
11485 good use is the following :
11486
11487     With named ACLs :
11488
11489         acl site_dead nbsrv(dynamic) lt 2
11490         acl site_dead nbsrv(static) lt 2
11491         monitor fail if site_dead
11492
11493     With anonymous ACLs :
11494
11495         monitor fail if { nbsrv(dynamic) lt 2 } || { nbsrv(static) lt 2 }
11496
11497 See section 4.2 for detailed help on the "block" and "use_backend" keywords.
```

7.3. Fetching samples

Historically, sample fetch methods were only used to retrieve data to match against patterns using ACLs. With the arrival of stick-tables, a new class of

sample fetch methods was created, most often sharing the same syntax as their ACL counterpart. These sample fetch methods are also known as "fetches". As of now, ACLs and fetches have converged. All ACL fetch methods have been made available as fetch methods, and ACLs may use any sample fetch method as well.

This section details all available sample fetch methods and their output type. Some sample fetch methods have deprecated aliases that are used to maintain compatibility with existing configurations. They are then explicitly marked as deprecated and should not be used in new setups.

The ACL derivatives are also indicated when available, with their respective matching methods. These ones all have a well defined default pattern matching method, so it is never necessary (though allowed) to pass the "-m" option to indicate how the sample will be matched using ACLs.

As indicated in the sample type versus matching compatibility matrix above, when using a generic sample fetch method in an ACL, the "-m" option is mandatory unless the sample type is one of boolean, integer, IPv4 or IPv6. When the same keyword exists as an ACL keyword and as a standard fetch method, the ACL engine will automatically pick the ACL-only one by default.

Some of these keywords support one or multiple mandatory arguments, and one or multiple optional arguments. These arguments are strongly typed and are checked when the configuration is parsed so that there is no risk of running with an incorrect argument (eg: an unresolved backend name). Fetch function arguments are passed between parenthesis and are delimited by commas. When an argument is optional, it will be indicated below between square brackets (['']). When all arguments are optional, the parenthesis may be omitted.

Thus, the syntax of a standard sample fetch method is one of the following :

```
- name
- name(arg1)
- name(arg1,arg2)
```

7.3.1. Converters

Sample fetch methods may be combined with transformations to be applied on top of the fetched sample (also called "converters"). These combinations form what is called "sample expressions" and the result is a "sample". Initially this was only supported by "stick on" and "stick store-request" directives but this has now be extended to all places where samples may be used (acls, log-format, unique-id-format, add-header, ...).

These transformations are enumerated as a series of specific keywords after the sample fetch method. These keywords may equally be appended immediately after the fetch keyword's argument, delimited by a comma. These keywords can also support some arguments (eg: a netmask) which must be passed in parenthesis.

A certain category of converters are bitwise and arithmetic operators which support performing basic operations on integers. Some bitwise operations are supported (and, or, xor, cpl) and some arithmetic operations are supported (add, sub, mul, div, mod, neg). Some comparators are provided (odd, even, not, bool) which make it possible to report a match without having to write an ACL.

The currently available list of transformation keywords include :

add(<values>)

Adds <values> to the input value of type signed integer, and returns the result as a signed integer. <values> can be a numeric value or a variable name. The name of the variable starts by an indication about its scope. The allowed scopes are:

"sess" : the variable is shared with all the session,

"txn" : the variable is shared with all the transaction (request and

```
11571         response),
11572         "req" : the variable is shared only during the request processing,
11573         "res" : the variable is shared only during the response processing.
11574         This prefix is followed by a name. The separator is a '.'. The name may only
11575         contain characters 'a-z', 'A-Z', '0-9' and '-'.
11576
11577     and(<value>)
11578     Performs a bitwise "AND" between <value> and the input value of type signed
11579     integer, and returns the result as an signed integer. <value> can be a
11580     numeric value or a variable name. The name of the variable starts by an
11581     indication about its scope. The allowed scopes are:
11582     "sess" : the variable is shared with all the session,
11583     "txn" : the variable is shared with all the transaction (request and
11584             response),
11585     "req" : the variable is shared only during the request processing,
11586     "res" : the variable is shared only during the response processing.
11587     This prefix is followed by a name. The separator is a '.'. The name may only
11588     contain characters 'a-z', 'A-Z', '0-9' and '-'.
11589
11590     base64
11591     Converts a binary input sample to a base64 string. It is used to log or
11592     transfer binary content in a way that can be reliably transferred (eg:
11593     an SSL ID can be copied in a header).
11594
11595     bool
11596     Returns a boolean TRUE if the input value of type signed integer is
11597     non-null, otherwise returns FALSE. Used in conjunction with and(), it can be
11598     used to report true/false for bit testing on input values (eg: verify the
11599     presence of a flag).
11600
11601     bytes(<offset>[,<length>])
11602     Extracts some bytes from an input binary sample. The result is a binary
11603     sample starting at an offset (in bytes) of the original sample and
11604     optionally truncated at the given length.
11605
11606     cpl
11607     Takes the input value of type signed integer, applies a ones-complement
11608     (flips all bits) and returns the result as an signed integer.
11609
11610     crc32([<avalanche>])
11611     Hashes a binary input sample into an unsigned 32-bit quantity using the CRC32
11612     hash function. Optionally, it is possible to apply a full avalanche hash
11613     function to the output if the optional <avalanche> argument equals 1. This
11614     converter uses the same functions as used by the various hash-based load
11615     balancing algorithms, so it will provide exactly the same results. It is
11616     provided for compatibility with other software which want a CRC32 to be
11617     computed on some input keys, so it follows the most common implementation as
11618     found in Ethernet, Gzip, PNG, etc... It is slower than the other algorithms
11619     but may provide a better or at least less predictable distribution. It must
11620     not be used for security purposes as a 32-bit hash is trivial to break. See
11621     also "djb2", "sdbm", "wt6" and the "hash-type" directive.
11622
11623     da-csv-conv(<prop>[,<prop>*])
11624     Asks the DeviceAtlas converter to identify the User Agent string passed on
11625     input, and to emit a string made of the concatenation of the properties
11626     enumerated in argument, delimited by the separator defined by the global
11627     keyword "deviceatlas-property-separator", or by default the pipe character
11628     ('|'). There's a limit of 5 different properties imposed by the haproxy
11629     configuration language.
11630
11631     Example:
11632     frontend www
11633     bind *:8881
11634     default backend servers
11635     http-request set-header X-DeviceAtlas-Data %[req,fnhdr(User-Agent),da-csv(prima...
```

```
11636
11637     debug
11638     This converter is used as debug tool. It dumps on screen the content and the
11639     type of the input sample. The sample is returned as is on its output. This
11640     converter only exists when haproxy was built with debugging enabled.
11641
11642     div(<value>)
11643     Divides the input value of type signed integer by <value>, and returns the
11644     result as an signed integer. If <value> is null, the largest unsigned
11645     integer is returned (typically 2^63-1). <value> can be a numeric value or a
11646     variable name. The name of the variable starts by an indication about it
11647     scope. The scope allowed are:
11648     "sess" : the variable is shared with all the session,
11649     "txn" : the variable is shared with all the transaction (request and
11650             response),
11651     "req" : the variable is shared only during the request processing,
11652     "res" : the variable is shared only during the response processing.
11653     This prefix is followed by a name. The separator is a '.'. The name may only
11654     contain characters 'a-z', 'A-Z', '0-9' and '-'.
11655
11656     djb2([<avalanches>])
11657     Hashes a binary input sample into an unsigned 32-bit quantity using the DJB2
11658     hash function. Optionally, it is possible to apply a full avalanche hash
11659     function to the output if the optional <avalanches> argument equals 1. This
11660     converter uses the same functions as used by the various hash-based load
11661     balancing algorithms, so it will provide exactly the same results. It is
11662     mostly intended for debugging, but can be used as a stick-table entry to
11663     collect rough statistics. It must not be used for security purposes as a
11664     32-bit hash is trivial to break. See also "crc32", "sdbm", "wt6" and the
11665     "hash-type" directive.
11666
11667     even
11668     Returns a boolean TRUE if the input value of type signed integer is even
11669     otherwise returns FALSE. It is functionally equivalent to "not,and(1),bool".
11670
11671     field(<index>,<delimiters>)
11672     Extracts the substring at the given index considering given delimiters from
11673     an input string. Indexes start at 1 and delimiters are a string formatted
11674     list of chars.
11675
11676     hex
11677     Converts a binary input sample to an hex string containing two hex digits per
11678     input byte. It is used to log or transfer hex dumps of some binary input data
11679     in a way that can be reliably transferred (eg: an SSL ID can be copied in a
11680     header).
11681
11682     http-date([<offset>])
11683     Converts an integer supposed to contain a date since epoch to a string
11684     representing this date in a format suitable for use in HTTP header fields. If
11685     an offset value is specified, then it is a number of seconds that is added to
11686     the date before the conversion is operated. This is particularly useful to
11687     emit Date header fields, Expires values in responses when combined with a
11688     positive offset, or Last-Modified values when the offset is negative.
11689
11690     in table(<table>)
11691     Uses the string representation of the input sample to perform a look up in
11692     the specified table. If the key is not found in the table, a boolean false
11693     is returned. Otherwise a boolean true is returned. This can be used to verify
11694     the presence of a certain key in a table tracking some elements (eg: whether
11695     or not a source IP address or an Authorization header was already seen).
11696
11697     ipmask(<mask>)
11698     Apply a mask to an IPv4 address, and use the result for lookups and storage.
11699     This can be used to make all hosts within a certain mask to share the same
11700     table entries and as such use the same server. The mask can be passed in
```

```
11701 dotted form (eg: 255.255.255.0) or in CIDR form (eg: 24).
11702
11703 json[<input-code>])
11704 Escapes the input string and produces an ASCII ouput string ready to use as a
11705 JSON string. The converter tries to decode the input string according to the
11706 <input-code> parameter. It can be "ascii", "utf8", "utf8s", "utf8" or
11707 "utf8ps". The "ascii" decoder never fails. The "utf8" decoder detects 3 types
11708 of errors:
11709 - bad UTF-8 sequence (lone continuation byte, bad number of continuation
11710 bytes, ...)
11711 - invalid range (the decoded value is within a UTF-8 prohibited range),
11712 - code overlong (the value is encoded with more bytes than necessary).
11713
11714 The UTF-8 JSON encoding can produce a "too long value" error when the UTF-8
11715 character is greater than 0xffff because the JSON string escape specification
11716 only authorizes 4 hex digits for the value encoding. The UTF-8 decoder exists
11717 in 4 variants designated by a combination of two suffix letters : "p" for
11718 "permissive" and "s" for "silently ignore". The behaviors of the decoders
11719 are :
11720 - "ascii" : never fails ;
11721 - "utf8" : fails on any detected errors ;
11722 - "utf8s" : never fails, but removes characters corresponding to errors ;
11723 - "utf8p" : accepts and fixes the overlong errors, but fails on any other
11724 error ;
11725 - "utf8ps" : never fails, accepts and fixes the overlong errors, but removes
11726 characters corresponding to the other errors.
11727
11728 This converter is particularly useful for building properly escaped JSON for
11729 logging to servers which consume JSON-formatted traffic logs.
11730
11731 Example:
11732 capture request header user-agent len 150
11733 capture request header Host len 15
11734 log-format {"ip":%[src],"user-agent":%[capture.req.hdr(1),json]}
11735
11736
11737 Input request from client 127.0.0.1:
11738 GET / HTTP/1.0
11739 User-Agent: Very "Ugly" UA 1/2
11740
11741 Output log:
11742 {"ip":"127.0.0.1","user-agent":"Very \"Ugly\" UA 1\2"}
11743
11744 Language(<values>[,<default>])
11745 Returns the value with the highest q-factor from a list as extracted from the
11746 "accept-language" header using "req.fhdr". Values with no q-factor have a
11747 q-factor of 1. Values with a q-factor of 0 are dropped. Only values which
11748 belong to the list of semi-colon delimited <values> will be considered. The
11749 argument <values> syntax is "lang[:lang[:lang[:...]]]". If no value matches the
11750 given list and a default value is provided, it is returned. Note that language
11751 names may have a variant after a dash ('-'). If this variant is present in the
11752 list, it will be matched, but if it is not, only the base language is checked.
11753 The match is case-sensitive, and the output string is always one of those
11754 provided in arguments. The ordering of arguments is meaningless, only the
11755 ordering of the values in the request counts, as the first value among
11756 multiple sharing the same q-factor is used.
11757
11758 Example :
11759 # this configuration switches to the backend matching a
11760 # given language based on the request :
11761
11762 acl es req.fhdr(accept-language), language(es;fr;en) -m str es
11763 acl fr req.fhdr(accept-language), language(es;fr;en) -m str fr
11764 acl en req.fhdr(accept-language), language(es;fr;en) -m str en
11765 use_backend spanish if es
```

```
11766 use_backend french if fr
11767 use_backend english if en
11768 default_backend choose_your_language
11769
11770 lower
11771 Convert a string sample to lower case. This can only be placed after a string
11772 sample fetch function or after a transformation keyword returning a string
11773 type. The result is of type string.
11774
11775 time(<format>[,<offset>])
11776 Converts an integer supposed to contain a date since epoch to a string
11777 representing this date in local time using a format defined by the <format>
11778 string using strftime(3). The purpose is to allow any date format to be used
11779 in logs. An optional <offset> in seconds may be applied to the input date
11780 (positive or negative). See the strftime() man page for the format supported
11781 by your operating system. See also the utime converter.
11782
11783 Example :
11784 # Emit two colons, one with the local time and another with ip:port
11785 # Eg: 20140710162350 127.0.0.1:57325
11786 log-format %[date,ttime(%Y%m%d%H%M%S)]\ %ci:%cp
11787
11788 map(<map_file>[,<default_values>])
11789 map_<map_file>(<map_file>[,<default_values>])
11790 map_<match_type> <output_type>(<map_file>[,<default_values>])
11791 Search the input value from <map_file> using the <match_type> matching method,
11792 and return the associated value converted to the type <output_type>. If the
11793 input value cannot be found in the <map_file>, the converter returns the
11794 <default_values>. If the <default_values> is not set, the converter fails and
11795 acts as if no input value could be fetched. If the <match_type> is not set, it
11796 defaults to "str". Likewise, if the <output_type> is not set, it defaults to
11797 "str". For convenience, the "map" keyword is an alias for "map_str" and maps a
11798 string to another string.
11799
11800 It is important to avoid overlapping between the keys : IP addresses and
11801 strings are stored in trees, so the first of the finest match will be used.
11802 Other keys are stored in lists, so the first matching occurrence will be used.
11803
11804 The following array contains the list of all map functions available sorted by
11805 input type, match type and output type.
11806
11807 input type | match method | output type str | output type int | output type ip
11808 -----+-----+-----+-----+-----+
11809 str | str | map_str | map_str_int | map_str_ip
11810 -----+-----+-----+-----+-----+
11811 str | beg | map_beg | map_beg_int | map_end_ip
11812 -----+-----+-----+-----+-----+
11813 str | sub | map_sub | map_sub_int | map_sub_ip
11814 -----+-----+-----+-----+-----+
11815 str | dir | map_dir | map_dir_int | map_dir_ip
11816 -----+-----+-----+-----+-----+
11817 str | dom | map_dom | map_dom_int | map_dom_ip
11818 -----+-----+-----+-----+-----+
11819 str | end | map_end | map_end_int | map_end_ip
11820 -----+-----+-----+-----+-----+
11821 str | reg | map_reg | map_reg_int | map_reg_ip
11822 -----+-----+-----+-----+-----+
11823 int | int | map_int | map_int_int | map_int_ip
11824 -----+-----+-----+-----+-----+
11825 ip | ip | map_ip | map_ip_int | map_ip_ip
11826 -----+-----+-----+-----+-----+
11827
11828 The file contains one key + value per line. Lines which start with '#' are
11829 ignored, just like empty lines. Leading tabs and spaces are stripped. The key
```


11961 contain characters 'a-z', 'A-Z', '0-9' and '._'.

11962 sub(<value>)

11963 Subtracts <value> from the input value of type signed integer, and returns

11964 the result as an signed integer. Note: in order to subtract the input from

11965 a constant, simply perform a "neg.add(value)". <value> can be a numeric value

11966 or a variable name. The name of the variable starts by an indication about its

11967 scope. The allowed scopes are:

11968 "sess" : the variable is shared with all the session,

11969 "txn" : the variable is shared with all the transaction (request and

11970 response),

11971 "req" : the variable is shared only during the request processing,

11972 "res" : the variable is shared only during the response processing.

11973 This prefix is followed by a name. The separator is a '.'. The name may only

11974 contain characters 'a-z', 'A-Z', '0-9' and '._'.

11975

11976 table_bytes_in_rate(<table>)

11977 Uses the string representation of the input sample to perform a look up in

11978 the specified table. If the key is not found in the table, integer value zero

11979 is returned. Otherwise the converter returns the average client-to-server

11980 bytes rate associated with the input sample in the designated table, measured

11981 in amount of bytes over the period configured in the table. See also the

11982 sc_bytes_in_rate sample fetch keyword.

11983

11984

11985

11986 table_bytes_out_rate(<table>)

11987 Uses the string representation of the input sample to perform a look up in

11988 the specified table. If the key is not found in the table, integer value zero

11989 is returned. Otherwise the converter returns the average server-to-client

11990 bytes rate associated with the input sample in the designated table, measured

11991 in amount of bytes over the period configured in the table. See also the

11992 sc_bytes_out_rate sample fetch keyword.

11993

11994

11995 table_conn_cnt(<table>)

11996 Uses the string representation of the input sample to perform a look up in

11997 the specified table. If the key is not found in the table, integer value zero

11998 is returned. Otherwise the converter returns the cumulated amount of incoming

11999 connections associated with the input sample in the designated table. See

12000 also the sc_conn_cnt sample fetch keyword.

12001

12002 table_conn_cur(<table>)

12003 Uses the string representation of the input sample to perform a look up in

12004 the specified table. If the key is not found in the table, integer value zero

12005 is returned. Otherwise the converter returns the current amount of concurrent

12006 tracked connections associated with the input sample in the designated table.

12007 See also the sc_conn_cur sample fetch keyword.

12008

12009 table_conn_rate(<table>)

12010 Uses the string representation of the input sample to perform a look up in

12011 the specified table. If the key is not found in the table, integer value zero

12012 is returned. Otherwise the converter returns the average incoming connection

12013 rate associated with the input sample in the designated table. See also the

12014 sc_conn_rate sample fetch keyword.

12015

12016 table_gpt0(<table>)

12017 Uses the string representation of the input sample to perform a look up in

12018 the specified table. If the key is not found in the table, boolean value zero

12019 is returned. Otherwise the converter returns the current value of the first

12020 general purpose tag associated with the input sample in the designated table.

12021 See also the sc_get_gpt0 sample fetch keyword.

12022

12023 table_gpc0(<table>)

12024 Uses the string representation of the input sample to perform a look up in

12025 the specified table. If the key is not found in the table, integer value zero

12026 general purpose counter associated with the input sample in the designated

12027 table. See also the sc_get_gpc0 sample fetch keyword.

12028

12029 table_gpc0_rate(<table>)

12030 Uses the string representation of the input sample to perform a look up in

12031 the specified table. If the key is not found in the table, integer value zero

12032 is returned. Otherwise the converter returns the frequency which the gpc0

12033 counter was incremented over the configured period in the table, associated

12034 with the input sample in the designated table. See also the sc_get_gpc0_rate

12035 sample fetch keyword.

12036

12037 table_http_err_cnt(<table>)

12038 Uses the string representation of the input sample to perform a look up in

12039 the specified table. If the key is not found in the table, integer value zero

12040 is returned. Otherwise the converter returns the cumulated amount of HTTP

12041 errors associated with the input sample in the designated table. See also the

12042 sc_http_err_cnt sample fetch keyword.

12043

12044 table_http_err_rate(<table>)

12045 Uses the string representation of the input sample to perform a look up in

12046 the specified table. If the key is not found in the table, integer value zero

12047 is returned. Otherwise the average rate of HTTP errors associated with the

12048 input sample in the designated table, measured in amount of errors over the

12049 period configured in the table. See also the sc_http_err_rate sample fetch

12050 keyword.

12051

12052 table_http_req_cnt(<table>)

12053 Uses the string representation of the input sample to perform a look up in

12054 the specified table. If the key is not found in the table, integer value zero

12055 is returned. Otherwise the converter returns the cumulated amount of HTTP

12056 requests associated with the input sample in the designated table. See also

12057 the sc_http_req_cnt sample fetch keyword.

12058

12059 table_http_req_rate(<table>)

12060 Uses the string representation of the input sample to perform a look up in

12061 the specified table. If the key is not found in the table, integer value zero

12062 is returned. Otherwise the average rate of HTTP requests associated with the

12063 input sample in the designated table, measured in amount of requests over the

12064 period configured in the table. See also the sc_http_req_rate sample fetch

12065 keyword.

12066

12067 table_kbytes_in(<table>)

12068 Uses the string representation of the input sample to perform a look up in

12069 the specified table. If the key is not found in the table, integer value zero

12070 is returned. Otherwise the converter returns the cumulated amount of client-

12071 to-server data associated with the input sample in the designated table,

12072 measured in kilobytes. The test is currently performed on 32-bit integers,

12073 which limits values to 4 terabytes. See also the sc_kbytes_in sample fetch

12074 keyword.

12075

12076 table_kbytes_out(<table>)

12077 Uses the string representation of the input sample to perform a look up in

12078 the specified table. If the key is not found in the table, integer value zero

12079 is returned. Otherwise the converter returns the cumulated amount of server-

12080 to-client data associated with the input sample in the designated table,

12081 measured in kilobytes. The test is currently performed on 32-bit integers,

12082 which limits values to 4 terabytes. See also the sc_kbytes_out sample fetch

12083 keyword.

12084

12085 table_server_id(<table>)

12086 Uses the string representation of the input sample to perform a look up in

12087 the specified table. If the key is not found in the table, integer value zero

12088 is returned. Otherwise the converter returns the server ID associated with

12089 the input sample in the designated table. A server ID is associated to a

12090 sample by a "stick" rule when a connection to a server succeeds. A server ID

zero means that no server is associated with this key.

table_sess_cnt(<table>)

Uses the string representation of the input sample to perform a look up in the specified table. If the key is not found in the table, integer value zero is returned. Otherwise the converter returns the cumulated amount of incoming sessions associated with the input sample in the designated table. Note that a session here refers to an incoming connection being accepted by the "tcp-request connection" rulesets. See also the `sc_sess_cnt` sample fetch keyword.

table_sess_rate(<table>)

Uses the string representation of the input sample to perform a look up in the specified table. If the key is not found in the table, integer value zero is returned. Otherwise the converter returns the average incoming session rate associated with the input sample in the designated table. Note that a session here refers to an incoming connection being accepted by the "tcp-request connection" rulesets. See also the `sc_sess_rate` sample fetch keyword.

table_trackers(<table>)

Uses the string representation of the input sample to perform a look up in the specified table. If the key is not found in the table, integer value zero is returned. Otherwise the converter returns the current amount of concurrent connections tracking the same key as the input sample in the designated table. It differs from `table_conn_cur` in that it does not rely on any stored information but on the table's reference count (the "use" value which is returned by "show table" on the CLI). This may sometimes be more suited for layer7 tracking. It can be used to tell a server how many concurrent connections there are from a given address for example. See also the `sc_trackers` sample fetch keyword.

upper

Convert a string sample to upper case. This can only be placed after a string sample fetch function or after a transformation keyword returning a string type. The result is of type string.

url_dec

Takes an url-encoded string provided as input and returns the decoded version as output. The input and the output are of type string.

utime(<format>[,<offset>])

Converts an integer supposed to contain a date since epoch to a string representing this date in UTC time using a format defined by the `<format>` string using `strftime(3)`. The purpose is to allow any date format to be used in logs. An optional `<offset>` in seconds may be applied to the input date (positive or negative). See the `strftime()` man page for the format supported by your operating system. See also the `[time converter]`.

Example :

```
# Emit two colons, one with the UTC time and another with ip:port
# Eg: 20140710162350 127.0.0.1:57325
log-format %[date,utime(%Y%m%d%H%M%S)]\ %ci:%cp
```

word(<index>,<delimiters>)

Extracts the nth word considering given delimiters from an input string. Indexes start at 1 and delimiters are a string formatted list of chars.

wt6([<avalanches>])

Hashes a binary input sample into an unsigned 32-bit quantity using the WT6 hash function. Optionally, it is possible to apply a full avalanche hash function to the output if the optional `<avalanches>` argument equals 1. This converter uses the same functions as used by the various hash-based load balancing algorithms, so it will provide exactly the same results. It is

mostly intended for debugging, but can be used as a stick-table entry to collect rough statistics. It must not be used for security purposes as a 32-bit hash is trivial to break. See also "crc32", "djb2", "sdbm", and the "hash-type" directive.

xor(<values>)

Performs a bitwise "XOR" (exclusive OR) between `<values>` and the input value of type signed integer, and returns the result as an signed integer. `<value>` can be a numeric value or a variable name. The name of the variable starts by an indication about its scope. The allowed scopes are: "sess" : the variable is shared with all the session, "txn" : the variable is shared with all the transaction (request and response).

"req" : the variable is shared only during the request processing,

"res" : the variable is shared only during the response processing.

This prefix is followed by a name. The separator is a '.'. The name may only contain characters 'a-z', 'A-Z', '0-9' and '_'.

7.3.2. Fetching samples from internal states

A first set of sample fetch methods applies to internal information which does not even relate to any client information. These ones are sometimes used with "monitor-fail" directives to report an internal status to external watchers. The sample fetch methods described in this section are usable anywhere.

always_false : boolean

Always returns the boolean "false" value. It may be used with ACLs as a temporary replacement for another one when adjusting configurations.

always_true : boolean

Always returns the boolean "true" value. It may be used with ACLs as a temporary replacement for another one when adjusting configurations.

avg_queue(<backends>]) : integer

Returns the total number of queued connections of the designated backend divided by the number of active servers. The current backend is used if no backend is specified. This is very similar to "queue" except that the size of the farm is considered, in order to give a more accurate measurement of the time it may take for a new connection to be processed. The main usage is with ACL to return a sorry page to new users when it becomes certain they will get a degraded service, or to pass to the backend servers in a header so that they decide to work in degraded mode or to disable some functions to speed up the processing a bit. Note that in the event there would not be any active server anymore, twice the number of queued connections would be considered as the measured value. This is a fair estimate, as we expect one server to get back soon anyway, but we still prefer to send new traffic to another backend if in better shape. See also the "queue", "be_conn", and "be_sess_rate" sample fetches.

be_conn(<backends>]) : integer

Applies to the number of currently established connections on the backend, possibly including the connection being evaluated. If no backend name is specified, the current one is used. But it is also possible to check another backend. It can be used to use a specific farm when the nominal one is full. See also the "fe_conn", "queue" and "be_sess_rate" criteria.

be_sess_rate(<backends>]) : integer

Returns an integer value corresponding to the sessions creation rate on the backend, in number of new sessions per second. This is used with ACLs to switch to an alternate backend when an expensive or fragile one reaches too high a session rate, or to limit abuse of service (eg. prevent sucking of an online dictionary). It can also be useful to add this element to logs using a log-format directive.

```
12221
12222
12223
12224
12225
12226
12227
12228
12229
12230
12231
12232
12233
12234
12235
12236
12237
12238
12239
12240
12241
12242
12243
12244
12245
12246
12247
12248
12249
12250
12251
12252
12253
12254
12255
12256
12257
12258
12259
12260
12261
12262
12263
12264
12265
12266
12267
12268
12269
12270
12271
12272
12273
12274
12275
12276
12277
12278
12279
12280
12281
12282
12283
12284
12285
```

Example :

Redirect to an error page if the dictionary is requested too often

```
backend dynamic
acl being_scanned be_sess_rate gt 100
redirect location /denied.html if being_scanned
```

bin(<hexa>) : bin

Returns a binary chain. The input is the hexadecimal representation of the string.

bool(<bool>) : bool

Returns a boolean value. <bool> can be 'true', 'false', '1' or '0'.

'false' and '0' are the same. 'true' and '1' are the same.

connslots(<[<backend>]) : integer

Returns an integer value corresponding to the number of connection slots still available in the backend, by totaling the maximum amount of connections on all servers and the maximum queue size. This is probably only used with ACLs.

The basic idea here is to be able to measure the number of connection "slots" still available (connection + queue), so that anything beyond that (intended usage; see "use_backend" keyword) can be redirected to a different backend.

'connslots' = number of available server connection slots, + number of available server queue slots.

Note that while "fe_conn" may be used, "connslots" comes in especially useful when you have a case of traffic going to one single ip, splitting into multiple backends (perhaps using ACLs to do name-based load balancing) and you want to be able to differentiate between different backends, and their available "connslots". Also, whereas "nbrsrv" only measures servers that are actually *down*, this fetch is more fine-grained and looks into the number of available connection slots as well. See also "queue" and "avg_queue".

OTHER CAVEATS AND NOTES: at this point in time, the code does not take care of dynamic connections. Also, if any of the server maxconn, or maxqueue is 0, then this fetch clearly does not make sense, in which case the value returned will be -1.

date([<offset>]) : integer

Returns the current date as the epoch (number of seconds since 01/01/1970).

If an offset value is specified, then it is a number of seconds that is added to the current date before returning the value. This is particularly useful to compute relative dates, as both positive and negative offsets are allowed. It is useful combined with the http_date converter.

Example :

set an expires header to now+1 hour in every response

```
http-response set-header Expires %[date(3600)].http_date]
```

env(<name>) : string

Returns a string containing the value of environment variable <name>. As a reminder, environment variables are per-process and are sampled when the process starts. This can be useful to pass some information to a next hop server, or with ACLs to take specific action when the process is started a certain way.

Examples :

Pass the Via header to next hop with the local hostname in it

```
http-request add-header Via 1.1\ %[env(HOSTNAME)]
```

```
12286
12287
12288
12289
12290
12291
12292
12293
12294
12295
12296
12297
12298
12299
12300
12301
12302
12303
12304
12305
12306
12307
12308
12309
12310
12311
12312
12313
12314
12315
12316
12317
12318
12319
12320
12321
12322
12323
12324
12325
12326
12327
12328
12329
12330
12331
12332
12333
12334
12335
12336
12337
12338
12339
12340
12341
12342
12343
12344
12345
12346
12347
12348
12349
12350
```

reject cookie-less requests when the STOP environment variable is set

```
http-request deny if { cook(SESSIONID) -m found } { env(STOP) -m found }
```

fe_conn(<[<frontend>]) : integer

Returns the number of currently established connections on the frontend, possibly including the connection being evaluated. If no frontend name is specified, the current one is used. But it is also possible to check another frontend. It can be used to return a sorry page before hard-blocking, or to use a specific backend to drain new requests when the farm is considered full. This is mostly used with ACLs but can also be used to pass some statistics to servers in HTTP headers. See also the "dst_conn", "be_conn", "fe_sess_rate" fetches.

fe_sess_rate(<[<frontend>]) : integer

Returns an integer value corresponding to the sessions creation rate on the frontend, in number of new sessions per second. This is used with ACLs to limit the incoming session rate to an acceptable range in order to prevent abuse of service at the earliest moment, for example when combined with other layer 4 ACLs in order to force the clients to wait a bit for the rate to go down below the limit. It can also be useful to add this element to logs using a log-format directive. See also the "rate-limit sessions" directive for use in frontends.

Example :

This frontend limits incoming mails to 10/s with a max of 100 concurrent connections. We accept any connection below 10/s, and # force excess clients to wait for 100 ms. Since clients are limited to # 100 max, there cannot be more than 10 incoming mails per second.

```
frontend mail
bind :25
mode tcp
maxconn 100
acl too_fast fe_sess_rate ge 10
tcp-request inspect-delay 100ms
tcp-request content accept if ! too_fast
tcp-request content accept if WAIT_END
```

int(<integer>) : signed integer

Returns a signed integer.

ip4(<[<ip4>]) : ip4

Returns an ip4.

ip6(<[<ip6>]) : ip6

Returns an ip6.

meth(<[<method>]) : method

Returns a method.

nbproc : integer

Returns an integer value corresponding to the number of processes that were started (it equals the global "nbproc" setting). This is useful for logging and debugging purposes.

nbrsrv(<[<backend>]) : integer

Returns an integer value corresponding to the number of usable servers of either the current backend or the named backend. This is mostly used with ACLs but can also be useful when added to logs. This is normally used to switch to an alternate backend when the number of servers is too low to handle some load. It is useful to report a failure when combined with "monitor fail".

proc : integer

Returns an integer value corresponding to the position of the process calling the function, between 1 and global.nbproc. This is useful for logging and

debugging purposes.

`queue([<backend>])` : integer

Returns the total number of queued connections of the designated backend, including all the connections in server queues. If no backend name is specified, the current one is used, but it is also possible to check another one. This is useful with ACLs or to pass statistics to backend servers. This can be used to take actions when queuing goes above a known level, generally indicating a surge of traffic or a massive slowdown on the servers. One possible action could be to reject new users but still accept old ones. See also the "avg_queue", "be_conn", and "be_sess_rate" fetches.

`rand([<range>])` : integer

Returns a random integer value within a range of <range> possible values, starting at zero. If the range is not specified, it defaults to 2^32, which gives numbers between 0 and 4294967295. It can be useful to pass some values needed to take some routing decisions for example, or just for debugging purposes. This random must not be used for security purposes.

`srv_conn([<backend>][<server>])` : integer

Returns an integer value corresponding to the number of currently established connections on the designated server, possibly including the connection being evaluated. If <backend> is omitted, then the server is looked up in the current backend. It can be used to use a specific farm when one server is full, or to inform the server about our view of the number of active connections with it. See also the "fe_conn", "be_conn" and "queue" fetch methods.

`srv_is_up([<backend>][<server>])` : boolean

Returns true when the designated server is UP, and false when it is either DOWN or in maintenance mode. If <backend> is omitted, then the server is looked up in the current backend. It is mainly used to take action based on an external status reported via a health check (eg: a geographical site's availability). Another possible use which is more of a hack consists in using dummy servers as boolean variables that can be enabled or disabled from the CLI, so that rules depending on those ACLs can be tweaked in realtime.

`srv_sess_rate([<backend>][<server>])` : integer

Returns an integer corresponding to the sessions creation rate on the designated server, in number of new sessions per second. If <backend> is omitted, then the server is looked up in the current backend. This is mostly used with ACLs but can make sense with logs too. This is used to switch to an alternate backend when an expensive or fragile one reaches too high a session rate, or to limit abuse of service (eg. prevent latent requests from overloading servers).

Example :

```
# Redirect to a separate back
acl srv1_full srv_sess_rate(be1/srv1) gt 50
acl srv2_full srv_sess_rate(be1/srv2) gt 50
use_backend be2 if srv1_full or srv2_full
```

`stopping` : boolean

Returns TRUE if the process calling the function is currently stopping. This can be useful for logging, or for relaxing certain checks or helping close certain connections upon graceful shutdown.

`str(<string>)` : string

Returns a string.

`table_avl([<table>])` : integer

Returns the total number of available entries in the current proxy's stick-table or in the designated stick-table. See also `table_cnt`.

`table_cnt([<table>])` : integer

Returns the total number of entries currently in use in the current proxy's stick-table or in the designated stick-table. See also `src_conn_cnt` and `table_avl` for other entry counting methods.

`var(<var-name>)` : undefined

Returns a variable with the stored type. If the variable is not set, the sample fetch fails. The name of the variable starts by an indication about its scope. The scope allowed are:

"sess" : the variable is shared with all the session,

"txn" : the variable is shared with all the transaction (request and response),

"req" : the variable is shared only during the request processing,

"res" : the variable is shared only during the response processing.

This prefix is followed by a name. The separator is a '.'. The name may only contain characters 'a-z', 'A-Z', '0-9' and '_'.

7.3.3. Fetching samples at Layer 4

The layer 4 usually describes just the transport layer which in haproxy is closest to the connection, where no content is yet made available. The fetch methods described here are usable as low as the "tcp-request connection" rule sets unless they require some future information. Those generally include TCP/IP addresses and ports, as well as elements from stick-tables related to the incoming connection. For retrieving a value from a sticky counters, the counter number can be explicitly set as 0, 1, or 2 using the pre-defined "sc0_", "sc1_", or "sc2_" prefix, or it can be specified as the first integer argument when using the "sc_" prefix. An optional table may be specified with the "sc#" form, in which case the currently tracked key will be looked up into this alternate table instead of the table currently being tracked.

`be_id` : integer

Returns an integer containing the current backend's id. It can be used in frontends with responses to check which backend processed the request.

`dst` : ip

This is the destination IPv4 address of the connection on the client side, which is the address the client connected to. It can be useful when running in transparent mode. It is of type IP and works on both IPv4 and IPv6 tables. On IPv6 tables, IPv4 address is mapped to its IPv6 equivalent, according to RFC 4291.

`dst_conn` : integer

Returns an integer value corresponding to the number of currently established connections on the same socket including the one being evaluated. It is normally used with ACLs but can as well be used to pass the information to servers in an HTTP header or in logs. It can be used to either return a sorry page before hard-blocking, or to use a specific backend to drain new requests when the socket is considered saturated. This offers the ability to assign different limits to different listening ports or addresses. See also the "fe_conn" and "be_conn" fetches.

`dst_port` : integer

Returns an integer value corresponding to the destination TCP port of the connection on the client side, which is the port the client connected to. This might be used when running in transparent mode, when assigning dynamic ports to some clients for a whole application session, to stick all users to a same server, or to pass the destination port information to a server using an HTTP header.

`fe_id` : integer

Returns an integer containing the current frontend's id. It can be used in backends to check from which backend it was called, or to stick all users coming via a same frontend to the same server.


```
12481      sc_bytes_in_rate(<ctr>[,<table>]) : integer
12482      sc0_bytes_in_rate([<table>]) : integer
12483      sc1_bytes_in_rate([<table>]) : integer
12484      sc2_bytes_in_rate([<table>]) : integer
12485      Returns the average client-to-server bytes rate from the currently tracked
12486      counters, measured in amount of bytes over the period configured in the
12487      table. See also src_bytes_in_rate.
12488
12489      sc_bytes_out_rate(<ctr>[,<table>]) : integer
12490      sc0_bytes_out_rate([<table>]) : integer
12491      sc1_bytes_out_rate([<table>]) : integer
12492      sc2_bytes_out_rate([<table>]) : integer
12493      Returns the average server-to-client bytes rate from the currently tracked
12494      counters, measured in amount of bytes over the period configured in the
12495      table. See also src_bytes_out_rate.
12496
12497      sc_clr_gpc0(<ctr>[,<table>]) : integer
12498      sc0_clr_gpc0([<table>]) : integer
12499      sc1_clr_gpc0([<table>]) : integer
12500      sc2_clr_gpc0([<table>]) : integer
12501      Clears the first General Purpose Counter associated to the currently tracked
12502      counters, and returns its previous value. Before the first invocation, the
12503      stored value is zero, so first invocation will always return zero. This is
12504      typically used as a second ACL in an expression in order to mark a connection
12505      when a first ACL was verified :
12506      # block if 5 consecutive requests continue to come faster than 10 sess
12507      # per second, and reset the counter as soon as the traffic slows down.
12508      acl abuse sc0_http_req_rate gt 10
12509      acl kill sc0_inc_gpc0 gt 5
12510      acl save sc0_clr_gpc0 ge 0
12511      tcp-request connection accept if !abuse save
12512      tcp-request connection reject if abuse kill
12513
12514      sc_conn_cnt(<ctr>[,<table>]) : integer
12515      sc0_conn_cnt([<table>]) : integer
12516      sc1_conn_cnt([<table>]) : integer
12517      sc2_conn_cnt([<table>]) : integer
12518      Returns the cumulated number of incoming connections from currently tracked
12519      counters. See also src_conn_cnt.
12520
12521      sc_conn_cur(<ctr>[,<table>]) : integer
12522      sc0_conn_cur([<table>]) : integer
12523      sc1_conn_cur([<table>]) : integer
12524      sc2_conn_cur([<table>]) : integer
12525      Returns the current amount of concurrent connections tracking the same
12526      tracked counters. This number is automatically incremented when tracking
12527      begins and decremented when tracking stops. See also src_conn_cur.
12528
12529      sc_conn_rate(<ctr>[,<table>]) : integer
12530      sc0_conn_rate([<table>]) : integer
12531      sc1_conn_rate([<table>]) : integer
12532      sc2_conn_rate([<table>]) : integer
12533      Returns the average connection rate from the currently tracked counters,
12534      measured in amount of connections over the period configured in the table.
12535      See also src_conn_rate.
12536
12537      sc_get_gpc0(<ctr>[,<table>]) : integer
12538      sc0_get_gpc0([<table>]) : integer
12539      sc1_get_gpc0([<table>]) : integer
12540      sc2_get_gpc0([<table>]) : integer
12541      Returns the value of the first General Purpose Counter associated to the
12542      currently tracked counters. See also src_get_gpc0 and sc/sc0/sc1/sc2_inc_gpc0.
```

```
12546      sc_get_gpt0(<ctr>[,<table>]) : integer
12547      sc0_get_gpt0([<table>]) : integer
12548      sc1_get_gpt0([<table>]) : integer
12549      sc2_get_gpt0([<table>]) : integer
12550      Returns the value of the first General Purpose Tag associated to the
12551      currently tracked counters. See also src_get_gpt0.
12552
12553      sc_gpc0_rate(<ctr>[,<table>]) : integer
12554      sc0_gpc0_rate([<table>]) : integer
12555      sc1_gpc0_rate([<table>]) : integer
12556      sc2_gpc0_rate([<table>]) : integer
12557      Returns the average increment rate of the first General Purpose Counter
12558      associated to the currently tracked counters. It reports the frequency
12559      which the gpc0 counter was incremented over the configured period. See also
12560      src_gpc0_rate, sc/sc0/sc1/sc2_get_gpc0, and sc/sc0/sc1/sc2_inc_gpc0. Note
12561      that the "gpc0_rate" counter must be stored in the stick-table for a value to
12562      be returned, as "gpc0" only holds the event count.
12563
12564      sc_http_err_cnt(<ctr>[,<table>]) : integer
12565      sc0_http_err_cnt([<table>]) : integer
12566      sc1_http_err_cnt([<table>]) : integer
12567      sc2_http_err_cnt([<table>]) : integer
12568      Returns the cumulated number of HTTP errors from the currently tracked
12569      counters. This includes the both request errors and 4xx error responses.
12570      See also src_http_err_cnt.
12571
12572      sc_http_err_rate(<ctr>[,<table>]) : integer
12573      sc0_http_err_rate([<table>]) : integer
12574      sc1_http_err_rate([<table>]) : integer
12575      sc2_http_err_rate([<table>]) : integer
12576      Returns the average rate of HTTP errors from the currently tracked counters,
12577      measured in amount of errors over the period configured in the table. This
12578      includes the both request errors and 4xx error responses. See also
12579      src_http_err_rate.
12580
12581      sc_http_req_cnt(<ctr>[,<table>]) : integer
12582      sc0_http_req_cnt([<table>]) : integer
12583      sc1_http_req_cnt([<table>]) : integer
12584      sc2_http_req_cnt([<table>]) : integer
12585      Returns the cumulated number of HTTP requests from the currently tracked
12586      counters. This includes every started request, valid or not. See also
12587      src_http_req_cnt.
12588
12589      sc_http_req_rate(<ctr>[,<table>]) : integer
12590      sc0_http_req_rate([<table>]) : integer
12591      sc1_http_req_rate([<table>]) : integer
12592      sc2_http_req_rate([<table>]) : integer
12593      Returns the average rate of HTTP requests from the currently tracked
12594      counters, measured in amount of requests over the period configured in
12595      the table. This includes every started request, valid or not. See also
12596      src_http_req_rate.
12597
12598      sc_inc_gpc0(<ctr>[,<table>]) : integer
12599      sc0_inc_gpc0([<table>]) : integer
12600      sc1_inc_gpc0([<table>]) : integer
12601      sc2_inc_gpc0([<table>]) : integer
12602      Increments the first General Purpose Counter associated to the currently
12603      tracked counters, and returns its new value. Before the first invocation,
12604      the stored value is zero, so first invocation will increase it to 1 and will
12605      return 1. This is typically used as a second ACL in an expression in order
12606      to mark a connection when a first ACL was verified :
12607      acl abuse sc0_http_req_rate gt 10
12608      acl kill sc0_inc_gpc0 gt 0
12609      tcp-request connection reject if abuse kill
12610
```

```
12611 sc_kbytes_in(<ctr>[,<table>]) : integer
12612 sc0_kbytes_in(<table>) : integer
12613 sc1_kbytes_in(<table>) : integer
12614 sc2_kbytes_in(<table>) : integer
12615 Returns the total amount of client-to-server data from the currently tracked
12616 counters, measured in kilobytes. The test is currently performed on 32-bit
12617 integers, which limits values to 4 terabytes. See also src_kbytes_in.
12618
12619
12620 sc_kbytes_out(<ctr>[,<table>]) : integer
12621 sc0_kbytes_out(<table>) : integer
12622 sc1_kbytes_out(<table>) : integer
12623 sc2_kbytes_out(<table>) : integer
12624 Returns the total amount of server-to-client data from the currently tracked
12625 counters, measured in kilobytes. The test is currently performed on 32-bit
12626 integers, which limits values to 4 terabytes. See also src_kbytes_out.
12627
12628 sc_sess_cnt(<ctr>[,<table>]) : integer
12629 sc0_sess_cnt(<table>) : integer
12630 sc1_sess_cnt(<table>) : integer
12631 sc2_sess_cnt(<table>) : integer
12632 Returns the cumulated number of incoming connections that were transformed
12633 into sessions, which means that they were accepted by a "tcp-request
12634 connection" rule, from the currently tracked counters. A backend may count
12635 more sessions than connections because each connection could result in many
12636 backend sessions if some HTTP keep-alive is performed over the connection
12637 with the client. See also src_sess_cnt.
12638
12639 sc_sess_rate(<ctr>[,<table>]) : integer
12640 sc0_sess_rate(<table>) : integer
12641 sc1_sess_rate(<table>) : integer
12642 sc2_sess_rate(<table>) : integer
12643 Returns the average session rate from the currently tracked counters,
12644 measured in amount of sessions over the period configured in the table. A
12645 session is a connection that got past the early "tcp-request connection"
12646 rules. A backend may count more sessions than connections because each
12647 connection could result in many backend sessions if some HTTP keep-alive is
12648 performed over the connection with the client. See also src_sess_rate.
12649
12650
12651 sc_tracked(<ctr>[,<table>]) : boolean
12652 sc0_tracked(<table>) : boolean
12653 sc1_tracked(<table>) : boolean
12654 sc2_tracked(<table>) : boolean
12655 Returns true if the designated session counter is currently being tracked by
12656 the current session. This can be useful when deciding whether or not we want
12657 to set some values in a header passed to the server.
12658
12659
12660 sc_trackers(<ctr>[,<table>]) : integer
12661 sc0_trackers(<table>) : integer
12662 sc1_trackers(<table>) : integer
12663 sc2_trackers(<table>) : integer
12664 Returns the current amount of concurrent connections tracking the same
12665 tracked counters. This number is automatically incremented when tracking
12666 begins and decremented when tracking stops. It differs from sc0_conn_cur in
12667 that it does not rely on any stored information but on the table's reference
12668 count (the "use" value which is returned by "show table" on the CLI). This
12669 may sometimes be more suited for layer7 tracking. It can be used to tell a
12670 server how many concurrent connections there are from a given address for
12671 example.
12672
12673
12674 so_id : integer
12675 Returns an integer containing the current listening socket's id. It is useful
12676 in frontends involving many "bind" lines, or to stick all users coming via a
12677 same socket to the same server.
```

```
12676 src : ip
12677 This is the source IPv4 address of the client of the session. It is of type
12678 IP and works on both IPv4 and IPv6 tables. On IPv6 tables, IPv4 addresses are
12679 mapped to their IPv6 equivalent, according to RFC 4291. Note that it is the
12680 TCP-level source address which is used, and not the address of a client
12681 behind a proxy. However if the "accept-proxy" bind directive is used, it can
12682 be the address of a client behind another PROXY-protocol compatible component
12683 for all rule sets except "tcp-request connection" which sees the real address.
12684
12685
12686 Example:
12687 # add an HTTP header in requests with the originating address' country
12688 http-request set-header X-Country %[src,map_ip(geoip.lst)]
12689
12690
12691 src_bytes_in_rate(<table>) : integer
12692 Returns the average bytes rate from the incoming connection's source address
12693 in the current proxy's stick-table or in the designated stick-table, measured
12694 in amount of bytes over the period configured in the table. If the address is
12695 not found, zero is returned. See also sc/sc0/scl/sc2_bytes_in_rate.
12696
12697
12698 src_bytes_out_rate(<table>) : integer
12699 Returns the average bytes rate to the incoming connection's source address in
12700 the current proxy's stick-table or in the designated stick-table, measured in
12701 amount of bytes over the period configured in the table. If the address is
12702 not found, zero is returned. See also sc/sc0/scl/sc2_bytes_out_rate.
12703
12704
12705 src_clr_gpc0(<table>) : integer
12706 Clears the first General Purpose Counter associated to the incoming
12707 connection's source address in the current proxy's stick-table or in the
12708 designated stick-table, and returns its previous value. If the address is not
12709 found, an entry is created and 0 is returned. This is typically used as a
12710 second ACL in an expression in order to mark a connection when a first ACL
12711 was verified :
12712
12713 # block if 5 consecutive requests continue to come faster than 10 sess
12714 # per second, and reset the counter as soon as the traffic slows down.
12715 acl abuse src_http_req_rate gt 10
12716 acl kill src_inc_gpc0 gt 5
12717 acl save src_clr_gpc0 ge 0
12718 tcp-request connection accept if !abuse save
12719 tcp-request connection reject if abuse kill
12720
12721
12722 src_conn_cnt(<table>) : integer
12723 Returns the cumulated number of connections initiated from the current
12724 incoming connection's source address in the current proxy's stick-table or in
12725 the designated stick-table. If the address is not found, zero is returned.
12726 See also sc/sc0/scl/sc2_conn_cnt.
12727
12728
12729 src_conn_cur(<table>) : integer
12730 Returns the current amount of concurrent connections initiated from the
12731 current incoming connection's source address in the current proxy's
12732 stick-table or in the designated stick-table. If the address is not found,
12733 zero is returned. See also sc/sc0/scl/sc2_conn_cur.
12734
12735
12736 src_conn_rate(<table>) : integer
12737 Returns the average connection rate from the incoming connection's source
12738 address in the current proxy's stick-table or in the designated stick-table,
12739 measured in amount of connections over the period configured in the table. If
12740 the address is not found, zero is returned. See also sc/sc0/scl/sc2_conn_rate.
12741
12742
12743 src_get_gpc0(<table>) : integer
12744 Returns the value of the first General Purpose Counter associated to the
12745 incoming connection's source address in the current proxy's stick-table or in
12746 the designated stick-table. If the address is not found, zero is returned.
12747 See also sc/sc0/scl/sc2_get_gpc0 and src_inc_gpc0.
12748
12749
```

```
12741 src_get_gpt0([<table>]) : integer
12742 Returns the value of the first General Purpose Tag associated to the
12743 incoming connection's source address in the current proxy's stick-table or in
12744 the designated stick-table. If the address is not found, zero is returned.
12745 See also sc/sc0/sc1/sc2_get_gpt0.
12746
12747 src_gpc0_rate([<table>]) : integer
12748 Returns the average increment rate of the first General Purpose Counter
12749 associated to the incoming connection's source address in the current proxy's
12750 stick-table or in the designated stick-table. It reports the frequency
12751 which the gpc0 counter was incremented over the configured period. See also
12752 sc/sc0/sc1/sc2_gpc0_rate, src_get_gpc0, and sc/sc0/sc1/sc2_inc_gpc0. Note
12753 that the "gpc0_rate" counter must be stored in the stick-table for a value to
12754 be returned, as "gpc0" only holds the event count.
12755
12756 src_http_err_cnt([<table>]) : integer
12757 Returns the cumulated number of HTTP errors from the incoming connection's
12758 source address in the current proxy's stick-table or in the designated
12759 stick-table. This includes the both request errors and 4xx error responses.
12760 See also sc/sc0/sc1/sc2_http_err_cnt. If the address is not found, zero is
12761 returned.
12762
12763 src_http_err_rate([<table>]) : integer
12764 Returns the average rate of HTTP errors from the incoming connection's source
12765 address in the current proxy's stick-table or in the designated stick-table,
12766 measured in amount of errors over the period configured in the table. This
12767 includes the both request errors and 4xx error responses. If the address is
12768 not found, zero is returned. See also sc/sc0/sc1/sc2_http_err_rate.
12769
12770 src_http_req_cnt([<table>]) : integer
12771 Returns the cumulated number of HTTP requests from the incoming connection's
12772 source address in the current proxy's stick-table or in the designated stick-
12773 table. This includes every started request, valid or not. If the address is
12774 not found, zero is returned. See also sc/sc0/sc1/sc2_http_req_cnt.
12775
12776 src_http_req_rate([<table>]) : integer
12777 Returns the average rate of HTTP requests from the incoming connection's
12778 source address in the current proxy's stick-table or in the designated stick-
12779 table, measured in amount of requests over the period configured in the
12780 table. This includes every started request, valid or not. If the address is
12781 not found, zero is returned. See also sc/sc0/sc1/sc2_http_req_rate.
12782
12783 src_inc_gpc0([<table>]) : integer
12784 Increments the first General Purpose Counter associated to the incoming
12785 connection's source address in the current proxy's stick-table or in the
12786 designated stick-table, and returns its new value. If the address is not
12787 found, an entry is created and 1 is returned. See also sc0/sc2/sc2_inc_gpc0.
12788 This is typically used as a second ACL in an expression in order to mark a
12789 connection when a first ACL was verified :
12790
12791 acl abuse src_http_req_rate gt 10
12792 acl kill src_inc_gpc0 gt 0
12793 tcp-request connection reject if abuse kill
12794
12795 src_kbytes_in([<table>]) : integer
12796 Returns the total amount of data received from the incoming connection's
12797 source address in the current proxy's stick-table or in the designated
12798 stick-table, measured in kilobytes. If the address is not found, zero is
12799 returned. The test is currently performed on 32-bit integers, which limits
12800 values to 4 terabytes. See also sc/sc0/sc1/sc2_kbytes_in.
12801
12802 src_kbytes_out([<table>]) : integer
12803 Returns the total amount of data sent to the incoming connection's source
12804 address in the current proxy's stick-table or in the designated stick-table,
12805 measured in kilobytes. If the address is not found, zero is returned. The
```

```
12806 test is currently performed on 32-bit integers, which limits values to 4
12807 terabytes. See also sc/sc0/sc1/sc2_kbytes_out.
12808
12809 src_port : integer
12810 Returns an integer value corresponding to the TCP source port of the
12811 connection on the client side, which is the port the client connected from.
12812 Usage of this function is very limited as modern protocols do not care much
12813 about source ports nowadays.
12814
12815 src_sess_cnt([<table>]) : integer
12816 Returns the cumulated number of connections initiated from the incoming
12817 connection's source IPv4 address in the current proxy's stick-table or in the
12818 designated stick-table, that were transformed into sessions, which means that
12819 they were accepted by "tcp-request" rules. If the address is not found, zero
12820 is returned. See also sc/sc0/sc1/sc2_sess_cnt.
12821
12822 src_sess_rate([<table>]) : integer
12823 Returns the average session rate from the incoming connection's source
12824 address in the current proxy's stick-table or in the designated stick-table,
12825 measured in amount of sessions over the period configured in the table. A
12826 session is a connection that went past the early "tcp-request" rules. If the
12827 address is not found, zero is returned. See also sc/sc0/sc1/sc2_sess_rate.
12828
12829 src_updt_conn_cnt([<table>]) : integer
12830 Creates or updates the entry associated to the incoming connection's source
12831 address in the current proxy's stick-table or in the designated stick-table.
12832 This table must be configured to store the "conn_cnt" data type, otherwise
12833 the match will be ignored. The current count is incremented by one, and the
12834 expiration timer refreshed. The updated count is returned, so this match
12835 can't return zero. This was used to reject service abusers based on their
12836 source address. Note: it is recommended to use the more complete "track-sc*"
12837 actions in "tcp-request" rules instead.
12838
12839 Example :
12840 # This frontend limits incoming SSH connections to 3 per 10 second for
12841 # each source address, and rejects excess connections until a 10 second
12842 # silence is observed. At most 20 addresses are tracked.
12843 listen ssh
12844 bind :22
12845 mode tcp
12846 maxconn 100
12847 stick-table type ip size 20 expire 10s store conn_cnt
12848 tcp-request content reject if { src_updt_conn_cnt gt 3 }
12849 server local 127.0.0.1:22
12850
12851 srv_id : integer
12852 Returns an integer containing the server's id when processing the response.
12853 While it's almost only used with ACLs, it may be used for logging or
12854 debugging.
12855
12856
12857 7.3.4. Fetching samples at Layer 5
12858 -----
12859
12860 The layer 5 usually describes just the session layer which in haproxy is
12861 closest to the session once all the connection handshakes are finished, but
12862 when no content is yet made available. The fetch methods described here are
12863 usable as low as the "tcp-request content" rule sets unless they require some
12864 future information. Those generally include the results of SSL negotiations.
12865
12866 ssl_bc : boolean
12867 Returns true when the back connection was made via an SSL/TLS transport
12868 layer and is locally deciphered. This means the outgoing connection was made
12869 other a server with the "ssl" option.
12870
```

12871 **ssl_bc_alg_keysize** : integer
12872 Returns the symmetric cipher key size supported in bits when the outgoing
12873 connection was made over an SSL/TLS transport layer.
12874
12875 **ssl_bc_cipher** : string
12876 Returns the name of the used cipher when the outgoing connection was made
12877 over an SSL/TLS transport layer.
12878
12879 **ssl_bc_protocol** : string
12880 Returns the name of the used protocol when the outgoing connection was made
12881 over an SSL/TLS transport layer.
12882
12883 **ssl_bc_unique_id** : binary
12884 When the outgoing connection was made over an SSL/TLS transport layer,
12885 returns the TLS unique ID as defined in RFC5929 section 3. The unique id
12886 can be encoded to base64 using the converter: "ssl_bc_unique_id,base64".
12887
12888 **ssl_bc_session_id** : binary
12889 Returns the SSL ID of the back connection when the outgoing connection was
12890 made over an SSL/TLS transport layer. It is useful to log if we want to know
12891 if session was reused or not.
12892
12893 **ssl_bc_use_keysize** : integer
12894 Returns the symmetric cipher key size used in bits when the outgoing
12895 connection was made over an SSL/TLS transport layer.
12896
12897 **ssl_c_ca_err** : integer
12898 When the incoming connection was made over an SSL/TLS transport layer,
12899 returns the ID of the first error detected during verification of the client
12900 certificate at depth > 0, or 0 if no error was encountered during this
12901 verification process. Please refer to your SSL library's documentation to
12902 find the exhaustive list of error codes.
12903
12904 **ssl_c_ca_err_depth** : integer
12905 When the incoming connection was made over an SSL/TLS transport layer,
12906 returns the depth in the CA chain of the first error detected during the
12907 verification of the client certificate. If no error is encountered, 0 is
12908 returned.
12909
12910 **ssl_c_der** : binary
12911 Returns the DER formatted certificate presented by the client when the
12912 incoming connection was made over an SSL/TLS transport layer. When used for
12913 an ACL, the value(s) to match against can be passed in hexadecimal form.
12914
12915 **ssl_c_err** : integer
12916 When the incoming connection was made over an SSL/TLS transport layer,
12917 returns the ID of the first error detected during verification at depth 0, or
12918 0 if no error was encountered during this verification process. Please refer
12919 to your SSL library's documentation to find the exhaustive list of error
12920 codes.
12921
12922 **ssl_c_i_dn([<entry>[,<occ>]])** : string
12923 When the incoming connection was made over an SSL/TLS transport layer,
12924 returns the full distinguished name of the issuer of the certificate
12925 presented by the client when no <entry> is specified, or the value of the
12926 first given entry found from the beginning of the DN. If a positive/negative
12927 occurrence number is specified as the optional second argument, it returns
12928 the value of the nth given entry value from the beginning/end of the DN.
12929 For instance, "ssl_c_i_dn(OU,2)" the second organization unit, and
12930 "ssl_c_i_dn(CN)" retrieves the common name.
12931
12932 **ssl_c_key_alg** : string
12933 Returns the name of the algorithm used to generate the key of the certificate
12934 presented by the client when the incoming connection was made over an SSL/TLS
12935 transport layer.

12936
12937 **ssl_c_notafter** : string
12938 Returns the end date presented by the client as a formatted string
12939 YMMDDhhmmss[Z] when the incoming connection was made over an SSL/TLS
12940 transport layer.
12941
12942 **ssl_c_notbefore** : string
12943 Returns the start date presented by the client as a formatted string
12944 YMMDDhhmmss[Z] when the incoming connection was made over an SSL/TLS
12945 transport layer.
12946
12947 **ssl_c_s_dn([<entry>[,<occ>]])** : string
12948 When the incoming connection was made over an SSL/TLS transport layer,
12949 returns the full distinguished name of the subject of the certificate
12950 presented by the client when no <entry> is specified, or the value of the
12951 first given entry found from the beginning of the DN. If a positive/negative
12952 occurrence number is specified as the optional second argument, it returns
12953 the value of the nth given entry value from the beginning/end of the DN.
12954 For instance, "ssl_c_s_dn(OU,2)" the second organization unit, and
12955 "ssl_c_s_dn(CN)" retrieves the common name.
12956
12957 **ssl_c_serial** : binary
12958 Returns the serial of the certificate presented by the client when the
12959 incoming connection was made over an SSL/TLS transport layer. When used for
12960 an ACL, the value(s) to match against can be passed in hexadecimal form.
12961
12962 **ssl_c_sha1** : binary
12963 Returns the SHA-1 fingerprint of the certificate presented by the client when
12964 the incoming connection was made over an SSL/TLS transport layer. This can be
12965 used to stick a client to a server, or to pass this information to a server.
12966 Note that the output is binary, so if you want to pass that signature to the
12967 server, you need to encode it in hex or base64, such as in the example below:
12968 http-request set-header X-SSL-Client-SHA1 %[ssl_c_sha1,hex]
12969
12970 **ssl_c_sig_alg** : string
12971 Returns the name of the algorithm used to sign the certificate presented by
12972 the client when the incoming connection was made over an SSL/TLS transport
12973 layer.
12974
12975 **ssl_c_used** : boolean
12976 Returns true if current SSL session uses a client certificate even if current
12977 connection uses SSL session resumption. See also "ssl_fc_has_crt".
12978
12979 **ssl_c_verify** : integer
12980 Returns the verify result error ID when the incoming connection was made over
12981 an SSL/TLS transport layer, otherwise zero if no error is encountered. Please
12982 refer to your SSL library's documentation for an exhaustive list of error
12983 codes.
12984
12985 **ssl_c_version** : integer
12986 Returns the version of the certificate presented by the client when the
12987 incoming connection was made over an SSL/TLS transport layer.
12988
12989 **ssl_f_der** : binary
12990 Returns the DER formatted certificate presented by the frontend when the
12991 incoming connection was made over an SSL/TLS transport layer. When used for
12992 an ACL, the value(s) to match against can be passed in hexadecimal form.
12993
12994 **ssl_f_i_dn([<entry>[,<occ>]])** : string
12995 When the incoming connection was made over an SSL/TLS transport layer,
12996 returns the full distinguished name of the issuer of the certificate
12997 presented by the frontend when no <entry> is specified, or the value of the
12998 first given entry found from the beginning of the DN. If a positive/negative
12999 occurrence number is specified as the optional second argument, it returns
13000

the value of the nth given entry value from the beginning/end of the DN.
For instance, "ssl_f_i_dn(0U,2)" the second organization unit, and
"ssl_f_i_dn(CN)" retrieves the common name.

ssl_f_key_alg : string

Returns the name of the algorithm used to generate the key of the certificate presented by the frontend when the incoming connection was made over an SSL/TLS transport layer.

ssl_f_notafter : string

Returns the end date presented by the frontend as a formatted string
YYMMDDhhmmss[Z] when the incoming connection was made over an SSL/TLS transport layer.

ssl_f_notbefore : string

Returns the start date presented by the frontend as a formatted string
YYMMDDhhmmss[Z] when the incoming connection was made over an SSL/TLS transport layer.

ssl_f_s_dn([<entry>[,<occ>]]) : string

When the incoming connection was made over an SSL/TLS transport layer, returns the full distinguished name of the subject of the certificate presented by the frontend when no <entry> is specified, or the value of the first given entry found from the beginning of the DN. If a positive/negative occurrence number is specified as the optional second argument, it returns the value of the nth given entry value from the beginning/end of the DN.
For instance, "ssl_f_s_dn(0U,2)" the second organization unit, and
"ssl_f_s_dn(CN)" retrieves the common name.

ssl_f_serial : binary

Returns the serial of the certificate presented by the frontend when the incoming connection was made over an SSL/TLS transport layer. When used for an ACL, the value(s) to match against can be passed in hexadecimal form.

ssl_f_sha1 : binary

Returns the SHA-1 fingerprint of the certificate presented by the frontend when the incoming connection was made over an SSL/TLS transport layer. This can be used to know which certificate was chosen using SNI.

ssl_f_sig_alg : string

Returns the name of the algorithm used to sign the certificate presented by the frontend when the incoming connection was made over an SSL/TLS transport layer.

ssl_f_version : integer

Returns the version of the certificate presented by the frontend when the incoming connection was made over an SSL/TLS transport layer.

ssl_fc : boolean

Returns true when the front connection was made via an SSL/TLS transport layer and is locally deciphered. This means it has matched a socket declared with a "bind" line having the "ssl" option.

Example :

```
# This passes "X-Proto: https" to servers when client connects over SSL
listen http-https
    bind :80
    bind :443 ssl crt /etc/haproxy.pem
    http-request add-header X-Proto https if { ssl_fc }
```

ssl_fc_alg_keysize : integer

Returns the symmetric cipher key size supported in bits when the incoming connection was made over an SSL/TLS transport layer.

ssl_fc_alpn : string

This extracts the Application Layer Protocol Negotiation field from an incoming connection made via a TLS transport layer and locally deciphered by haproxy. The result is a string containing the protocol name advertised by the client. The SSL library must have been built with support for TLS extensions enabled (check haproxy -vv). Note that the TLS ALPN extension is not advertised unless the "alpn" keyword on the "bind" line specifies a protocol list. Also, nothing forces the client to pick a protocol from this list, any other one may be requested. The TLS ALPN extension is meant to replace the TLS NPN extension. See also "ssl_fc_npn".

ssl_fc_cipher : string

Returns the name of the used cipher when the incoming connection was made over an SSL/TLS transport layer.

ssl_fc_has_crt : boolean

Returns true if a client certificate is present in an incoming connection over SSL/TLS transport layer. Useful if 'verify' statement is set to 'optional'.
Note: on SSL session resumption with Session ID or TLS ticket, client certificate is not present in the current connection but may be retrieved from the cache or the ticket. So prefer "ssl_c_used" if you want to check if current SSL session uses a client certificate.

ssl_fc_has_sni : boolean

This checks for the presence of a Server Name Indication (SNI) in an incoming connection was made over an SSL/TLS transport layer. Returns true when the incoming connection presents a TLS SNI field. This requires that the SSL library is build with support for TLS extensions enabled (check haproxy -vv).

ssl_fc_is_resumed: boolean

Returns true if the SSL/TLS session has been resumed through the use of SSL session cache or TLS tickets.

ssl_fc_npn : string

This extracts the Next Protocol Negotiation field from an incoming connection made via a TLS transport layer and locally deciphered by haproxy. The result is a string containing the protocol name advertised by the client. The SSL library must have been built with support for TLS extensions enabled (check haproxy -vv). Note that the TLS NPN extension is not advertised unless the "npn" keyword on the "bind" line specifies a protocol list. Also, nothing forces the client to pick a protocol from this list, any other one may be requested. Please note that the TLS NPN extension was replaced with ALPN.

ssl_fc_protocol : string

Returns the name of the used protocol when the incoming connection was made over an SSL/TLS transport layer.

ssl_fc_unique_id : binary

When the incoming connection was made over an SSL/TLS transport layer, returns the TLS unique ID as defined in RFC5929 section 3. The unique ID can be encoded to base64 using the converter: "ssl_bc_unique_id,base64".

ssl_fc_session_id : binary

Returns the SSL ID of the front connection when the incoming connection was made over an SSL/TLS transport layer. It is useful to stick a given client to a server. It is important to note that some browsers refresh their session ID every few minutes.

ssl_fc_sni : string

This extracts the Server Name Indication (SNI) field from an incoming connection made via an SSL/TLS transport layer and locally deciphered by haproxy. The result (when present) typically is a string matching the HTTPS host name (253 chars or less). The SSL library must have been built with support for TLS extensions enabled (check haproxy -vv).

13130

This fetch is different from "req_ssl_sni" above in that it applies to the connection being deciphered by haproxy and not to SSL contents being blindly forwarded. See also "ssl_fc_sni_end" and "ssl_fc_sni_reg" below. This requires that the SSL library is build with support for TLS extensions enabled (check haproxy -vv).

ACL derivatives :

ssl_fc_sni_end : suffix match
ssl_fc_sni_reg : regex match

ssl_fc_use_keysize : integer

Returns the symmetric cipher key size used in bits when the incoming connection was made over an SSL/TLS transport layer.

7.3.5. Fetching samples from buffer contents (Layer 6)

Fetching samples from buffer contents is a bit different from the previous sample fetches above because the sampled data are ephemeral. These data can only be used when they're available and will be lost when they're forwarded. For this reason, samples fetched from buffer contents during a request cannot be used in a response for example. Even while the data are being fetched, they can change. Sometimes it is necessary to set some delays or combine multiple sample fetch methods to ensure that the expected data are complete and usable, for example through TCP request content inspection. Please see the "tcp-request content" keyword for more detailed information on the subject.

payload(<offset>,<length>) : binary (deprecated)

This is an alias for "req.payload" when used in the context of a request (eg: "stick on", "stick match"), and for "res.payload" when used in the context of a response such as in "stick store response".

payload_lv(<offset>,<length>[,<offset2>]) : binary (deprecated)

This is an alias for "req.payload_lv" when used in the context of a request (eg: "stick on", "stick match"), and for "res.payload_lv" when used in the context of a response such as in "stick store response".

req.len : integer

req.len : integer (deprecated)

Returns an integer value corresponding to the number of bytes present in the request buffer. This is mostly used in ACL. It is important to understand that this test does not return false as long as the buffer is changing. This means that a check with equality to zero will almost always immediately match at the beginning of the session, while a test for more data will wait for that data to come in and return false only when haproxy is certain that no more data will come in. This test was designed to be used with TCP request content inspection.

req.payload(<offset>,<length>) : binary

This extracts a binary block of <length> bytes and starting at byte <offset> in the request buffer. As a special case, if the <length> argument is zero, the whole buffer from <offset> to the end is extracted. This can be used with ACLs in order to check for the presence of some content in a buffer at any location.

ACL alternatives :

payload(<offset>,<length>) : hex binary match

req.payload_lv(<offset>,<length>[,<offset2>]) : binary

This extracts a binary block whose size is specified at <offset> for <length> bytes, and which starts at <offset2> if specified or just after the length in the request buffer. The <offset2> parameter also supports relative offsets if prepended with a '+' or '-' sign.

ACL alternatives :

payload_lv(<offset>,<length>[,<offset2>]) : hex binary match

Example : please consult the example from the "stick store-response" keyword.
req.proto_http : boolean
req.proto_http : boolean (deprecated)
Returns true when data in the request buffer look like HTTP and correctly parses as such. It is the same parser as the common HTTP request parser which is used so there should be no surprises. The test does not match until the request is complete, failed or timed out. This test may be used to report the protocol in TCP logs, but the biggest use is to block TCP request analysis until a complete HTTP request is present in the buffer, for example to track a header.

Example:

```
# track request counts per "base" (concatenation of Host+URL)
tcp-request inspect-delay 10s
tcp-request content reject if !HTTP
tcp-request content track-sc0 base table req-rate
```

req_rdp_cookie(<name>) : string

req_rdp_cookie(<name>) : string (deprecated)

When the request buffer looks like the RDP protocol, extracts the RDP cookie <name>, or any cookie if unspecified. The parser only checks for the first cookie, as illustrated in the RDP protocol specification. The cookie name is case insensitive. Generally the "MSTS" cookie name will be used, as it can contain the user name of the client connecting to the server if properly configured on the client. The "MSTHASH" cookie is often used as well for session stickiness to servers.

This differs from "balance rdp-cookie" in that any balancing algorithm may be used and thus the distribution of clients to backend servers is not linked to a hash of the RDP cookie. It is envisaged that using a balancing algorithm such as "balance roundrobin" or "balance leastconn" will lead to a more even distribution of clients to backend servers than the hash used by "balance rdp-cookie".

ACL derivatives :

req_rdp_cookie(<name>) : exact string match

Example :

```
listen tse-farm
bind 0.0.0.0:3389
# wait up to 5s for an RDP cookie in the request
tcp-request inspect-delay 5s
tcp-request content accept if RDP_COOKIE
# apply RDP cookie persistence
persist rdp-cookie
# Persist based on the msthash cookie
# This is only useful makes sense if
# balance rdp-cookie is not used
stick-table type string size 204800
stick on req_rdp_cookie(msthash)
server srv1 1.1.1.1:3389
server srv2 1.1.1.2:3389
```

See also : "balance rdp-cookie", "persist rdp-cookie", "tcp-request" and the "req_rdp_cookie" ACL.

req_rdp_cookie_cnt(<name>) : integer

req_rdp_cookie_cnt(<name>) : integer (deprecated)

Tries to parse the request buffer as RDP protocol, then returns an integer corresponding to the number of RDP cookies found. If an optional cookie name is passed, only cookies matching this name are considered. This is mostly

```
13261 used in ACL.
13262
13263 ACL derivatives :
13264     req_rdp_cookie_cnt(<name>[ ] : integer match
13265
13266 req_ssl_ec_ext : boolean
13267 Returns a boolean identifying if client sent the Supported Elliptic Curves
13268 Extension as defined in RFC4492, section 5.1. within the SSL ClientHello
13269 message. This can be used to present ECC compatible clients with EC
13270 certificate and to use RSA for all others, on the same IP address. Note that
13271 this only applies to raw contents found in the request buffer and not to
13272 contents deciphered via an SSL data layer, so this will not work with "bind"
13273 lines having the "ssl" option.
13274
13275 req_ssl_hello_type : integer
13276 req_ssl_hello_type : integer (deprecated)
13277 Returns an integer value containing the type of the SSL hello message found
13278 in the request buffer if the buffer contains data that parse as a complete
13279 SSL (v3 or superior) client hello message. Note that this only applies to raw
13280 contents found in the request buffer and not to contents deciphered via an
13281 SSL data layer, so this will not work with "bind" lines having the "ssl"
13282 option. This is mostly used in ACL to detect presence of an SSL hello message
13283 that is supposed to contain an SSL session ID usable for stickiness.
13284
13285 req_ssl_sni : string
13286 req_ssl_sni : string (deprecated)
13287 Returns a string containing the value of the Server Name TLS extension sent
13288 by a client in a TLS stream passing through the request buffer if the buffer
13289 contains data that parse as a complete SSL (v3 or superior) client hello
13290 message. Note that this only applies to raw contents found in the request
13291 buffer and not to contents deciphered via an SSL data layer, so this will not
13292 work with "bind" lines having the "ssl" option. SNI normally contains the
13293 name of the host the client tries to connect to (for recent browsers). SNI is
13294 useful for allowing or denying access to certain hosts when SSL/TLS is used
13295 by the client. This test was designed to be used with TCP request content
13296 inspection. If content switching is needed, it is recommended to first wait
13297 for a complete client hello (type 1), like in the example below. See also
13298 "ssl_fc_sni".
13299
13300 ACL derivatives :
13301     req_ssl_sni : exact string match
13302
13303 Examples :
13304     # Wait for a client hello for at most 5 seconds
13305     tcp-request inspect-delay 5s
13306     use backend bk allow if { req_ssl_sni -f allowed_sites }
13307     default_backend bk_sorry_page
13308
13309 req_ssl_st_ext : integer
13310 Returns 0 if the client didn't send a SessionTicket TLS Extension (RFC5077)
13311 Returns 1 if the client sent SessionTicket TLS Extension
13312 Returns 2 if the client also sent non-zero length TLS SessionTicket
13313 Note that this only applies to raw contents found in the request buffer and
13314 not to contents deciphered via an SSL data layer, so this will not work with
13315 "bind" lines having the "ssl" option. This can for example be used to detect
13316 whether the client sent a SessionTicket or not and stick it accordingly, if
13317 no SessionTicket then stick on SessionID or don't stick as there's no server
13318 side state is there when SessionTickets are in use.
13319
13320 req_ssl_ver : integer
13321 req_ssl_ver : integer (deprecated)
13322 Returns an integer value containing the version of the SSL/TLS protocol of a
13323 stream present in the request buffer. Both SSLv2 hello messages and SSLv3
13324 messages are supported. TLSv1 is announced as SSL version 3.1. The value is
13325
```

```
13326 composed of the major version multiplied by 65536, added to the minor
13327 version. Note that this only applies to raw contents found in the request
13328 buffer and not to contents deciphered via an SSL data layer, so this will not
13329 work with "bind" lines having the "ssl" option. The ACL version of the test
13330 matches against a decimal notation in the form MAJOR.MINOR (eg: 3.1). This
13331 fetch is mostly used in ACL.
13332
13333 ACL derivatives :
13334     req_ssl_ver : decimal match
13335
13336 res.len : integer
13337 Returns an integer value corresponding to the number of bytes present in the
13338 response buffer. This is mostly used in ACL. It is important to understand
13339 that this test does not return false as long as the buffer is changing. This
13340 means that a check with equality to zero will almost always immediately match
13341 at the beginning of the session, while a test for more data will wait for
13342 that data to come in and return false only when haproxy is certain that no
13343 more data will come in. This test was designed to be used with TCP response
13344 content inspection.
13345
13346 res.payload(<offset>,<length>) : binary
13347 This extracts a binary block of <length> bytes and starting at byte <offset>
13348 in the response buffer. As a special case, if the <length> argument is zero,
13349 the whole buffer from <offset> to the end is extracted. This can be used
13350 with ACLs in order to check for the presence of some content in a buffer at
13351 any location.
13352
13353 res.payload_lv(<offset>,<length>[,<offset2>]) : binary
13354 This extracts a binary block whose size is specified at <offset1> for <length>
13355 bytes, and which starts at <offset2> if specified or just after the length in
13356 the response buffer. The <offset2> parameter also supports relative offsets
13357 if prepended with a '+' or '-' sign.
13358
13359 Example : please consult the example from the "stick store-response" keyword.
13360
13361 res_ssl_hello_type : integer
13362 rep_ssl_hello_type : integer (deprecated)
13363 Returns an integer value containing the type of the SSL hello message found
13364 in the response buffer if the buffer contains data that parses as a complete
13365 SSL (v3 or superior) hello message. Note that this only applies to raw
13366 contents found in the response buffer and not to contents deciphered via an
13367 SSL data layer, so this will not work with "server" lines having the "ssl"
13368 option. This is mostly used in ACL to detect presence of an SSL hello message
13369 that is supposed to contain an SSL session ID usable for stickiness.
13370
13371 wait_end : boolean
13372 This fetch either returns true when the inspection period is over, or does
13373 not fetch. It is only used in ACLs, in conjunction with content analysis to
13374 avoid returning a wrong verdict early. It may also be used to delay some
13375 actions, such as a delayed reject for some special addresses. Since it either
13376 stops the rules evaluation or immediately returns true, it is recommended to
13377 use this acl as the last one in a rule. Please note that the default ACL
13378 "WAIT_END" is always usable without prior declaration. This test was designed
13379 to be used with TCP request content inspection.
13380
13381 Examples :
13382     # delay every incoming request by 2 seconds
13383     tcp-request inspect-delay 2s
13384     tcp-request content accept if WAIT_END
13385
13386     # don't immediately tell bad guys they are rejected
13387     tcp-request inspect-delay 10s
13388     acl goodguys src 10.0.0.0/24
13389     acl badguys src 10.0.1.0/24
13390     tcp-request content accept if goodguys
```

```
13391 tcp-request content reject if badguys WAIT_END
13392 tcp-request content reject
13393
```

7.3.6. Fetching HTTP samples (Layer 7)

It is possible to fetch samples from HTTP contents, requests and responses. This application layer is also called layer 7. It is only possible to fetch the data in this section when a full HTTP request or response has been parsed from its respective request or response buffer. This is always the case with all HTTP specific rules and for sections running with "mode http". When using TCP content inspection, it may be necessary to support an inspection delay in order to let the request or response come in first. These fetches may require a bit more CPU resources than the layer 4 ones, but not much since the request and response are indexed.

base : string
This returns the concatenation of the first Host header and the path part of the request, which starts at the first slash and ends before the question mark. It can be useful in virtual hosted environments to detect URL abuses as well as to improve shared caches efficiency. Using this with a limited size stick table also allows one to collect statistics about most commonly requested objects by host/path. With ACLs it can allow simple content switching rules involving the host and the path at the same time, such as "www.example.com/favicon.ico". See also "path" and "uri".

ACL derivatives :
base : exact string match
base_beg : prefix match
base_dir : subdir match
base_end : domain match
base_len : length match
base_reg : regex match
base_sub : substring match

base32 : integer
This returns a 32-bit hash of the value returned by the "base" fetch method above. This is useful to track per-URL activity on high traffic sites without having to store all URLs. Instead a shorter hash is stored, saving a lot of memory. The output type is an unsigned integer. The hash function used is SDBM with full avalanche on the output. Technically, base32 is exactly equal to "base,sdbm(1)".

base32+src : binary
This returns the concatenation of the base32 fetch above and the src fetch below. The resulting type is of type binary, with a size of 8 or 20 bytes depending on the source address family. This can be used to track per-IP, per-URL counters.

capture.req.hdr(<idx>) : string

This extracts the content of the header captured by the "capture request header", idx is the position of the capture keyword in the configuration. The first entry is an index of 0. See also: "capture request header".

capture.req.method : string

This extracts the METHOD of an HTTP request. It can be used in both request and response. Unlike "method", it can be used in both request and response because it's allocated.

capture.req.uri : string

This extracts the request's URI, which starts at the first slash and ends before the first space in the request (without the host part). Unlike "path" and "url", it can be used in both request and response because it's

allocated.

capture.req.ver : string
This extracts the request's HTTP version and returns either "HTTP/1.0" or "HTTP/1.1". Unlike "req.ver", it can be used in both request, response, and logs because it relies on a persistent flag.

capture.res.hdr(<idx>) : string
This extracts the content of the header captured by the "capture response header", idx is the position of the capture keyword in the configuration. The first entry is an index of 0. See also: "capture response header"

capture.res.ver : string
This extracts the response's HTTP version and returns either "HTTP/1.0" or "HTTP/1.1". Unlike "res.ver", it can be used in logs because it relies on a persistent flag.

req.body : binary
This returns the HTTP request's available body as a block of data. It requires that the request body has been buffered made available using "option http-buffer-request". In case of chunked-encoded body, currently only the first chunk is analyzed.

req.body_param(<[<name>]) : string
This fetch assumes that the body of the POST request is url-encoded. The user can check if the "content-type" contains the value "application/x-www-form-urlencoded". This extracts the first occurrence of the parameter <name> in the body, which ends before '&'. The parameter name is case-sensitive. If no name is given, any parameter will match, and the first one will be returned. The result is a string corresponding to the value of the parameter <name> as presented in the request body (no URL decoding is performed). Note that the ACL version of this fetch iterates over multiple parameters and will iteratively report all parameters values if no name is given.

req.body.len : integer

This returns the length of the HTTP request's available body in bytes. It may be lower than the advertised length if the body is larger than the buffer. It requires that the request body has been buffered made available using "option http-buffer-request".

req.body.size : integer

This returns the advertised length of the HTTP request's body in bytes. It will represent the advertised Content-Length header, or the size of the first chunk in case of chunked encoding. In order to parse the chunks, it requires that the request body has been buffered made available using "option http-buffer-request".

req.cook(<[<name>]) : string

cook(<[<name>]) : string (deprecated)

This extracts the last occurrence of the cookie name <name> on a "Cookie" header line from the request, and returns its value as string. If no name is specified, the first cookie value is returned. When used with ACLs, all matching cookies are evaluated. Spaces around the name and the value are ignored as requested by the Cookie header specification (RFC6265). The cookie name is case-sensitive. Empty cookies are valid, so an empty cookie may very well return an empty value if it is present. Use the "found" match to detect presence. Use the res.cook() variant for response cookies sent by the server.

ACL derivatives :

cook(<[<name>]) : exact string match

cook_beg(<[<name>]) : prefix match

cook_dir(<[<name>]) : subdir match

cook_dom(<[<name>]) : domain match

13520


```
13521 cook_end([<name>]) : suffix match
13522 cook_len([<name>]) : length match
13523 cook_reg([<name>]) : regex match
13524 cook_sub([<name>]) : substring match
13525
13526 req.cook_cnt([<name>]) : integer
13527 req.cook_cnt([<name>]) : integer (deprecated)
13528 Returns an integer value representing the number of occurrences of the cookie
13529 <name> in the request, or all cookies if <name> is not specified.
13530
13531 req.cook_val([<name>]) : integer
13532 req_val([<name>]) : integer (deprecated)
13533 This extracts the last occurrence of the cookie name <name> on a "Cookie"
13534 header line from the request, and converts its value to an integer which is
13535 returned. If no name is specified, the first cookie value is returned. When
13536 used in ACLs, all matching names are iterated over until a value matches.
13537
13538 cookie([<name>]) : string (deprecated)
13539 This extracts the last occurrence of the cookie name <name> on a "Cookie"
13540 header line from the request, or a "Set-Cookie" header from the response, and
13541 returns its value as a string. A typical use is to get multiple clients
13542 sharing a same profile use the same server. This can be similar to what
13543 "appsession" did with the "request-learn" statement, but with support for
13544 multi-peer synchronization and state keeping across restarts. If no name is
13545 specified, the first cookie value is returned. This fetch should not be used
13546 anymore and should be replaced by req.cook() or res.cook() instead as it
13547 ambiguously uses the direction based on the context where it is used.
13548
13549 hdr([<name>[,<occ>]]) : string
13550 This is equivalent to req.hdr() when used on requests, and to res.hdr() when
13551 used on responses. Please refer to these respective fetches for more details.
13552 In case of doubt about the fetch direction, please use the explicit ones.
13553 Note that contrary to the hdr() sample fetch method, the hdr_* ACL keywords
13554 unambiguously apply to the request headers.
13555
13556 req.fhdr([<name>[,<occ>]]) : string
13557 This extracts the last occurrence of header <name> in an HTTP request. When
13558 used from an ACL, all occurrences are iterated over until a match is found.
13559 Optionally, a specific occurrence might be specified as a position number.
13560 Positive values indicate a position from the first occurrence, with 1 being
13561 the first one. Negative values indicate positions relative to the last one,
13562 with -1 being the last one. It differs from req.hdr() in that any commas
13563 present in the value are returned and are not used as delimiters. This is
13564 sometimes useful with headers such as User-Agent.
13565
13566 req.fhdr_cnt([<name>]) : integer
13567 Returns an integer value representing the number of occurrences of request
13568 header field name <name>, or the total number of header fields if <name> is
13569 not specified. Contrary to its req.hdr_cnt() cousin, this function returns
13570 the number of full line headers and does not stop on commas.
13571
13572 req.hdr([<name>[,<occ>]]) : string
13573 This extracts the last occurrence of header <name> in an HTTP request. When
13574 used from an ACL, all occurrences are iterated over until a match is found.
13575 Optionally, a specific occurrence might be specified as a position number.
13576 Positive values indicate a position from the first occurrence, with 1 being
13577 the first one. Negative values indicate positions relative to the last one,
13578 with -1 being the last one. A typical use is with the X-Forwarded-For header
13579 once converted to IP, associated with an IP stick-table. The function
13580 considers any comma as a delimiter for distinct values. If full-line headers
13581 are desired instead, use req.fhdr(). Please carefully check RFC2616 to know
13582 how certain headers are supposed to be parsed. Also, some of them are case
13583 insensitive (eg: Connection).
13584
13585 ACL derivatives :
```

```
13586 hdr([<name>[,<occ>]]) : exact string match
13587 hdr_beg([<name>[,<occ>]]) : prefix match
13588 hdr_dir([<name>[,<occ>]]) : subdir match
13589 hdr_dom([<name>[,<occ>]]) : domain match
13590 hdr_end([<name>[,<occ>]]) : suffix match
13591 hdr_len([<name>[,<occ>]]) : length match
13592 hdr_reg([<name>[,<occ>]]) : regex match
13593 hdr_sub([<name>[,<occ>]]) : substring match
13594
13595 req.hdr_cnt([<name>]) : integer
13596 req_cnt([<header>]) : integer (deprecated)
13597 Returns an integer value representing the number of occurrences of request
13598 header field name <name>, or the total number of header field values if
13599 <name> is not specified. It is important to remember that one header line may
13600 contain several headers if it has several values. The function considers any
13601 comma as a delimiter for distinct values. If full-line headers are desired
13602 instead, req.fhdr_cnt() should be used instead. With ACLs, it can be used to
13603 detect presence, absence or abuse of a specific header, as well as to block
13604 request smuggling attacks by rejecting requests which contain more than one
13605 of certain headers. See "req.hdr" for more information on header matching.
13606
13607 req.hdr_ip([<name>[,<occ>]]) : ip
13608 This extracts the last occurrence of header <name> in an HTTP request,
13609 converts it to an IPv4 or IPv6 address and returns this address. When used
13610 with ACLs, all occurrences are checked, and if <name> is omitted, every value
13611 of every header is checked. Optionally, a specific occurrence might be
13612 specified as a position number. Positive values indicate a position from the
13613 first occurrence, with 1 being the first one. Negative values indicate
13614 positions relative to the last one, with -1 being the last one. A typical use
13615 is with the X-Forwarded-For and X-Client-IP headers.
13616
13617 req.hdr_val([<name>[,<occ>]]) : integer
13618 hdr_val([<name>[,<occ>]]) : integer (deprecated)
13619 This extracts the last occurrence of header <name> in an HTTP request, and
13620 converts it to an integer value. When used with ACLs, all occurrences are
13621 checked, and if <name> is omitted, every value of every header is checked.
13622 Optionally, a specific occurrence might be specified as a position number.
13623 Positive values indicate a position from the first occurrence, with 1 being
13624 the first one. Negative values indicate positions relative to the last one,
13625 with -1 being the last one. A typical use is with the X-Forwarded-For header.
13626
13627 http.auth(<userlist>) : boolean
13628 Returns a boolean indicating whether the authentication data received from
13629 the client match a username & password stored in the specified userlist. This
13630 fetch function is not really useful outside of ACLs. Currently only http
13631 basic auth is supported.
13632
13633 http.auth_group(<userlist>) : string
13634 Returns a string corresponding to the user name found in the authentication
13635 data received from the client if both the user name and password are valid
13636 according to the specified userlist. The main purpose is to use it in ACLs
13637 where it is then checked whether the user belongs to any group within a list.
13638 This fetch function is not really useful outside of ACLs. Currently only http
13639 basic auth is supported.
13640
13641 ACL derivatives :
13642 http.auth_group(<userlist>) : group ...
13643 Returns true when the user extracted from the request and whose password is
13644 valid according to the specified userlist belongs to at least one of the
13645 groups.
13646
13647 http.first_req : boolean
13648 Returns true when the request being processed is the first one of the
13649 connection. This can be used to add or remove headers that may be missing
13650
```

from some requests when a request is not the first one, or to help grouping requests in the logs.

method : integer + string

Returns an integer value corresponding to the method in the HTTP request. For example, "GET" equals 1 (check sources to establish the matching). Value 9 means "other method" and may be converted to a string extracted from the stream. This should not be used directly as a sample, this is only meant to be used from ACLs, which transparently convert methods from patterns to these integer + string values. Some predefined ACL already check for most common methods.

ACL derivatives :

method : case insensitive method match

Example :

```
# only accept GET and HEAD requests
acl valid_method method GET HEAD
http-request deny if ! valid_method
```

path : string

This extracts the request's URL path, which starts at the first slash and ends before the question mark (without the host part). A typical use is with prefetch-capable caches, and with portals which need to aggregate multiple information from databases and keep them in caches. Note that with outgoing caches, it would be wiser to use "url" instead. With ACLs, it's typically used to match exact file names (eg: "/login.php"), or directory parts using the derivative forms. See also the "url" and "base" fetch methods.

ACL derivatives :

```
path : exact string match
path_beg : prefix match
path_dir : subdir match
path_dom : domain match
path_end : suffix match
path_len : length match
path_reg : regex match
path_sub : substring match
```

query : string

This extracts the request's query string, which starts after the first question mark. If no question mark is present, this fetch returns nothing. If a question mark is present but nothing follows, it returns an empty string. This means it's possible to easily know whether a query string is present using the "found" matching method. This fetch is the complement of "path" which stops before the question mark.

req_hdr_names([<delim>]) : string

This builds a string made from the concatenation of all header names as they appear in the request when the rule is evaluated. The default delimiter is the comma (',') but it may be overridden as an optional argument <delim>. In this case, only the first character of <delim> is considered.

req_ver : string

req_ver : string (deprecated)

Returns the version string from the HTTP request, for example "1.1". This can be useful for logs, but is mostly there for ACL. Some predefined ACL already check for versions 1.0 and 1.1.

ACL derivatives :

```
req_ver : exact string match
```

res.comp : boolean

Returns the boolean "true" value if the response has been compressed by HAProxy, otherwise returns boolean "false". This may be used to add

information in the logs.

res.comp_algo : string

Returns a string containing the name of the algorithm used if the response was compressed by HAProxy, for example : "deflate". This may be used to add some information in the logs.

res.cookie([<name>]) : string

scookie([<name>]) : string (deprecated)

This extracts the last occurrence of the cookie name <name> on a "Set-Cookie" header line from the response, and returns its value as string. If no name is specified, the first cookie value is returned.

ACL derivatives :

```
scookie([<name>] : exact string match
```

res.cookie_cnt([<name>]) : integer

scookie_cnt([<name>]) : integer (deprecated)

Returns an integer value representing the number of occurrences of the cookie <name> in the response, or all cookies if <name> is not specified. This is mostly useful when combined with ACLs to detect suspicious responses.

res.cookie_val([<name>]) : integer

scookie_val([<name>]) : integer (deprecated)

This extracts the last occurrence of the cookie name <name> on a "Set-Cookie" header line from the response, and converts its value to an integer which is returned. If no name is specified, the first cookie value is returned.

res.fhdr([<name>[,<occ>]]) : string

This extracts the last occurrence of header <name> in an HTTP response, or of the last header if no <name> is specified. Optionally, a specific occurrence might be specified as a position number. Positive values indicate a position from the first occurrence, with 1 being the first one. Negative values indicate positions relative to the last one, with -1 being the last one. It differs from res.hdr() in that any commas present in the value are returned and are not used as delimiters. If this is not desired, the res.hdr() fetch should be used instead. This is sometimes useful with headers such as Date or Expires.

res.fhdr_cnt([<name>]) : integer

Returns an integer value representing the number of occurrences of response header field name <name>, or the total number of header fields if <name> is not specified. Contrary to its res.hdr_cnt() cousin, this function returns the number of full line headers and does not stop on commas. If this is not desired, the res.hdr_cnt() fetch should be used instead.

res.hdr([<name>[,<occ>]]) : string

shdr([<name>[,<occ>]]) : string (deprecated)

This extracts the last occurrence of header <name> in an HTTP response, or of the last header if no <name> is specified. Optionally, a specific occurrence might be specified as a position number. Positive values indicate a position from the first occurrence, with 1 being the first one. Negative values indicate positions relative to the last one, with -1 being the last one. This can be useful to learn some data into a stick-table. The function considers any comma as a delimiter for distinct values. If this is not desired, the res.fhdr() fetch should be used instead.

ACL derivatives :

```
shdr([<name>[,<occ>]]) : exact string match
shdr_beg([<name>[,<occ>]]) : prefix match
shdr_dir([<name>[,<occ>]]) : subdir match
shdr_dom([<name>[,<occ>]]) : domain match
shdr_end([<name>[,<occ>]]) : suffix match
shdr_len([<name>[,<occ>]]) : length match
shdr_reg([<name>[,<occ>]]) : regex match
```

```
13781      shdr_sub([<name>[,<occ>]]) : substring match
13782
13783      res.hdr_cnt([<name>]) : integer
13784      shdr_cnt([<name>]) : integer (deprecated)
13785      Returns an integer value representing the number of occurrences of response
13786      header field name <name>, or the total number of header fields if <name> is
13787      not specified. The function considers any comma as a delimiter for distinct
13788      values. If this is not desired, the res.fhdr_cnt() fetch should be used
13789      instead.
13790
13791      res.hdr_ip([<name>[,<occ>]]) : ip
13792      shdr_ip([<name>[,<occ>]]) : ip (deprecated)
13793      This extracts the last occurrence of header <name> in an HTTP response,
13794      convert it to an IPv4 or IPv6 address and returns this address. Optionally, a
13795      specific occurrence might be specified as a position number. Positive values
13796      indicate a position from the first occurrence, with 1 being the first one.
13797      Negative values indicate positions relative to the last one, with -1 being
13798      the last one. This can be useful to learn some data into a stick table.
13799
13800      res.hdr_names([<delim>]) : string
13801      This builds a string made from the concatenation of all header names as they
13802      appear in the response when the rule is evaluated. The default delimiter is
13803      the comma (',') but it may be overridden as an optional argument <delim>. In
13804      this case, only the first character of <delim> is considered.
13805
13806      res.hdr_val([<names>[,<occ>]]) : integer
13807      shdr_val([<names>[,<occ>]]) : integer (deprecated)
13808      This extracts the last occurrence of header <names> in an HTTP response, and
13809      converts it to an integer value. Optionally, a specific occurrence might be
13810      specified as a position number. Positive values indicate a position from the
13811      first occurrence, with 1 being the first one. Negative values indicate
13812      positions relative to the last one, with -1 being the last one. This can be
13813      useful to learn some data into a stick table.
13814
13815      res.ver : string
13816      resp.ver : string (deprecated)
13817      Returns the version string from the HTTP response, for example "1.1". This
13818      can be useful for logs, but is mostly there for ACL.
13819
13820      ACL derivatives :
13821      resp_ver : exact string match
13822
13823      set-cookie([<name>]) : string (deprecated)
13824      This extracts the last occurrence of the cookie name <name> on a "Set-Cookie"
13825      header line from the response and uses the corresponding value to match. This
13826      can be comparable to what "appsession" did with default options, but with
13827      support for multi-peer synchronization and state keeping across restarts.
13828
13829      This fetch function is deprecated and has been superseded by the "res.cook"
13830      fetch. This keyword will disappear soon.
13831
13832      status : integer
13833      Returns an integer containing the HTTP status code in the HTTP response, for
13834      example, 302. It is mostly used within ACLs and integer ranges, for example,
13835      to remove any Location header if the response is not a 3xx.
13836
13837      url : string
13838      This extracts the request's URL as presented in the request. A typical use is
13839      with prefetch-capable caches, and with portals which need to aggregate
13840      multiple information from databases and keep them in caches. With ACLs, using
13841      "path" is preferred over using "url", because clients may send a full URL as
13842      is normally done with proxies. The only real use is to match "*" which does
13843      not match in "path", and for which there is already a predefined ACL. See
13844      also "path" and "base".
13845
```

```
13846      ACL derivatives :
13847      url      : exact string match
13848      url_beg  : prefix match
13849      url_dir  : subdir match
13850      url_dom  : domain match
13851      url_end  : suffix match
13852      url_len  : length match
13853      url_reg  : regex match
13854      url_sub  : substring match
13855
13856      url_ip : ip
13857      This extracts the IP address from the request's URL when the host part is
13858      presented as an IP address. Its use is very limited. For instance, a
13859      monitoring system might use this field as an alternative for the source IP in
13860      order to test what path a given source address would follow, or to force an
13861      entry in a table for a given source address. With ACLs it can be used to
13862      restrict access to certain systems through a proxy, for example when combined
13863      with option "http_proxy".
13864
13865      url_port : integer
13866      This extracts the port part from the request's URL. Note that if the port is
13867      not specified in the request, port 80 is assumed. With ACLs it can be used to
13868      restrict access to certain systems through a proxy, for example when combined
13869      with option "http_proxy".
13870
13871      url_param([<name>[,<delim>]]) : string
13872      url_param([<name>[,<delim>]]) : string
13873      This extracts the first occurrence of the parameter <name> in the query
13874      string, which begins after either '?' or <delim>, and which ends before '&',
13875      ';' or <delim>. The parameter name is case-sensitive. If no name is given,
13876      any parameter will match, and the first one will be returned. The result is
13877      a string corresponding to the value of the parameter <names> as presented in
13878      the request (no URL decoding is performed). This can be used for session
13879      stickiness based on a client ID, to extract an application cookie passed as a
13880      URL parameter, or in ACLs to apply some checks. Note that the ACL version of
13881      this fetch iterates over multiple parameters and will iteratively report all
13882      parameters values if no name is given
13883
13884      ACL derivatives :
13885      urlp(<name>[,<delim>]) : exact string match
13886      urlp_beg(<name>[,<delim>]) : prefix match
13887      urlp_dir(<name>[,<delim>]) : subdir match
13888      urlp_dom(<name>[,<delim>]) : domain match
13889      urlp_end(<name>[,<delim>]) : suffix match
13890      urlp_len(<name>[,<delim>]) : length match
13891      urlp_reg(<name>[,<delim>]) : regex match
13892      urlp_sub(<name>[,<delim>]) : substring match
13893
13894      Example :
13895      # match http://example.com/foo?PHPSESSIONID=some_id
13896      stick on urlp(PHPSESSIONID)
13897      # match http://example.com/foo;JSESSIONID=some_id
13898      stick on urlp(JSESSIONID,;)
13899
13900      urlp_val([<name>[,<delim>]]) : integer
13901      See "urlp" above. This one extracts the URL parameter <name> in the request
13902      and converts it to an integer value. This can be used for session stickiness
13903      based on a user ID for example, or with ACLs to match a page number or price.
13904
13905      7.4. Pre-defined ACLs
13906      -----
13907      Some predefined ACLs are hard-coded so that they do not have to be declared in
13908
```

every frontend which needs them. They all have their names in upper case in order to avoid confusion. Their equivalence is provided below.

ACL name	Equivalent to	Usage
FALSE	always_false	never match
HTTP_1_0	req_proto_http	match if protocol is valid HTTP
HTTP_1_1	req_ver 1.0	match HTTP version 1.0
HTTP_CONTENT	req_ver 1.1	match HTTP version 1.1
HTTP_URL_ABS	hdr_val(content-length) gt 0	match an existing content-length
HTTP_URL_SLASH	url_reg ^[/\.:]*://	match absolute URL with scheme
HTTP_URL_STAR	url_beg /	match URL beginning with "/"
LOCALHOST	url_equal *	match URL equal to "*"
METH_CONNECT	src 127.0.0.1/8	match connection from local host
METH_GET	method CONNECT	match HTTP CONNECT method
METH_HEAD	method GET HEAD	match HTTP GET or HEAD method
METH_OPTIONS	method HEAD	match HTTP HEAD method
METH_POST	method OPTIONS	match HTTP OPTIONS method
METH_TRACE	method POST	match HTTP POST method
RDP_COOKIE	method TRACE	match HTTP TRACE method
REQ_CONTENT	req_rdp_cookie_cnt gt 0	match presence of an RDP cookie
TRUE	req_len gt 0	match data in the request buffer
WAIT_END	always_true	always match
	wait_end	wait for end of content analysis

8. Logging

One of HAProxy's strong points certainly lies in its precise logs. It probably provides the finest level of information available for such a product, which is very important for troubleshooting complex environments. Standard information provided in logs include client ports, TCP/HTTP state timers, precise session state at termination and precise termination cause, information about decisions to direct traffic to a server, and of course the ability to capture arbitrary headers.

In order to improve administrators reactivity, it offers a great transparency about encountered problems, both internal and external, and it is possible to send logs to different sources at the same time with different level filters :

- global process-level logs (system errors, start/stop, etc..)
- per-instance system and internal errors (lack of resource, bugs, ...)
- per-instance external troubles (servers up/down, max connections)
- per-instance activity (client connections), either at the establishment or at the termination.
- per-request control of log-level, eg:
 - http-request set-log-level silent if sensitive_request

The ability to distribute different levels of logs to different log servers allow several production teams to interact and to fix their problems as soon as possible. For example, the system team might monitor system-wide errors, while the application team might be monitoring the up/down for their servers in real time, and the security team might analyze the activity logs with one hour delay.

8.1. Log levels

TCP and HTTP connections can be logged with information such as the date, time, source IP address, destination address, connection duration, response times, HTTP request, HTTP return code, number of bytes transmitted, conditions in which the session ended, and even exchanged cookies values. For example

track a particular user's problems. All messages may be sent to up to two syslog servers. Check the "log" keyword in section 4.2 for more information about log facilities.

8.2. Log formats

HAProxy supports 5 log formats. Several fields are common between these formats and will be detailed in the following sections. A few of them may vary slightly with the configuration, due to indicators specific to certain options. The supported formats are as follows :

- the default format, which is very basic and very rarely used. It only provides very basic information about the incoming connection at the moment it is accepted : source IP:port, destination IP:port, and frontend-name. This mode will eventually disappear so it will not be described to great extents.

- the TCP format, which is more advanced. This format is enabled when "option tcplog" is set on the frontend. HAProxy will then usually wait for the connection to terminate before logging. This format provides much richer information, such as timers, connection counts, queue size, etc... This format is recommended for pure TCP proxies.

- the HTTP format, which is the most advanced for HTTP proxying. This format is enabled when "option httplog" is set on the frontend. It provides the same information as the TCP format with some HTTP-specific fields such as the request, the status code, and captures of headers and cookies. This format is recommended for HTTP proxies.

- the CLF HTTP format, which is equivalent to the HTTP format, but with the fields arranged in the same order as the CLF format. In this mode, all timers, captures, flags, etc... appear one per field after the end of the common fields, in the same order they appear in the standard HTTP format.

- the custom log format, allows you to make your own log line.

Next sections will go deeper into details for each of these formats. Format specification will be performed on a "field" basis. Unless stated otherwise, a field is a portion of text delimited by any number of spaces. Since syslog servers are susceptible of inserting fields at the beginning of a line, it is always assumed that the first field is the one containing the process name and identifier.

Note : Since log lines may be quite long, the log examples in sections below might be broken into multiple lines. The example log lines will be prefixed with 3 closing angle brackets ('>>>') and each time a log is broken into multiple lines, each non-final line will end with a backslash ('\') and the next line will start indented by two characters.

8.2.1. Default log format

This format is used when no specific option is set. The log is emitted as soon as the connection is accepted. One should note that this currently is the only format which logs the request's destination IP and ports.

Example :

```
listen www
mode http
log global
server srv1 127.0.0.1:8000
```

```
>>> Feb 6 12:12:09 localhost \
haproxy[14385]: Connect from 10.0.1.2:33312 to 10.0.3.31:8012 \
(www/HTTP)
```

Field	Format	Extract from the example above
1	process_name '[' pid ']:'	haproxy[14385]:
2	'Connect from'	Connect from
3	source_ip ':' source_port	10.0.1.2:33312
4	'to'	to
5	destination_ip ':' destination_port	10.0.3.31:8012
6	(' frontend_name '/' mode ')	(www/HTTP)

Detailed fields description :

- "source_ip" is the IP address of the client which initiated the connection.
- "source_port" is the TCP port of the client which initiated the connection.
- "destination_ip" is the IP address the client connected to.
- "destination_port" is the TCP port the client connected to.
- "frontend_name" is the name of the frontend (or listener) which received and processed the connection.
- "mode" is the mode the frontend is operating (TCP or HTTP).

In case of a UNIX socket, the source and destination addresses are marked as "unix:" and the ports reflect the internal ID of the socket which accepted the connection (the same ID as reported in the stats).

It is advised not to use this deprecated format for newer installations as it will eventually disappear.

8.2.2. TCP log format

The TCP format is used when "option tcplog" is specified in the frontend, and is the recommended format for pure TCP proxies. It provides a lot of precious information for troubleshooting. Since this format includes timers and byte counts, the log is normally emitted at the end of the session. It can be emitted earlier if "option logasap" is specified, which makes sense in most environments with long sessions such as remote terminals. Sessions which match the "monitor" rules are never logged. It is also possible not to emit logs for sessions for which no data were exchanged between the client and the server, by specifying "option dontlognull" in the frontend. Successful connections will not be logged if "option dontlog-normal" is specified in the frontend. A few fields may slightly vary depending on some configuration options, those are marked with a star (*) after the field name below.

Example :

```
frontend fnt
mode tcp
option tcplog
log global
default_backend bck

backend bck
server srv1 127.0.0.1:8000
```

```
>>> Feb 6 12:12:56 localhost \
haproxy[14387]: 10.0.1.2:33313 [06/Feb/2009:12:12:51.443] fnt \
bck/srv1 0/0/5007 212 -- 0/0/0/0/3 0/0
```

Field	Format	Extract from the example above
1	process_name '[' pid ']:'	haproxy[14387]:
2	client_ip ':' client_port	10.0.1.2:33313
3	(' accept_date ')	[06/Feb/2009:12:12:51.443]
4	frontend_name	fnt
5	backend_name '/' server_name	bck/srv1

6	Tw '/' Tc '/' Tt*	0/0/5007
7	bytes_read*	212
8	termination_state	--
9	actconn '/' feconn '/' beconn '/' srv_conn '/' retries*	0/0/0/0/3
10	srv_queue '/' backend_queue	0/0

Detailed fields description :

- "client_ip" is the IP address of the client which initiated the TCP connection to haproxy. If the connection was accepted on a UNIX socket instead, the IP address would be replaced with the word "unix". Note that when the connection is accepted on a socket configured with "accept-proxy" and the PROXY protocol is correctly used, then the logs will reflect the forwarded connection's information.
- "client_port" is the TCP port of the client which initiated the connection. If the connection was accepted on a UNIX socket instead, the port would be replaced with the ID of the accepting socket, which is also reported in the stats interface.
- "accept_date" is the exact date when the connection was received by haproxy (which might be very slightly different from the date observed on the network if there was some queuing in the system's backlog). This is usually the same date which may appear in any upstream firewall's log.
- "frontend_name" is the name of the frontend (or listener) which received and processed the connection.
- "backend_name" is the name of the backend (or listener) which was selected to manage the connection to the server. This will be the same as the frontend if no switching rule has been applied, which is common for TCP applications.
- "server_name" is the name of the last server to which the connection was sent, which might differ from the first one if there were connection errors and a redispatch occurred. Note that this server belongs to the backend which processed the request. If the connection was aborted before reaching a server, "<NOSRV>" is indicated instead of a server name.
- "Tw" is the total time in milliseconds spent waiting in the various queues. It can be "-1" if the connection was aborted before reaching the queue. See "Timers" below for more details.
- "Tc" is the total time in milliseconds spent waiting for the connection to establish to the final server, including retries. It can be "-1" if the connection was aborted before a connection could be established. See "Timers" below for more details.
- "Tt" is the total time in milliseconds elapsed between the accept and the last close. It covers all possible processing. There is one exception, if "option logasap" was specified, then the time counting stops at the moment the log is emitted. In this case, a '+' sign is prepended before the value, indicating that the final one will be larger. See "Timers" below for more details.
- "bytes_read" is the total number of bytes transmitted from the server to the client when the log is emitted. If "option logasap" is specified, the this value will be prefixed with a '+' sign indicating that the final one may be larger. Please note that this value is a 64-bit counter, so log analysis tools must be able to handle it without overflowing.
- "termination_state" is the condition the session was in when the session ended. This indicates the session state, which side caused the end of session to happen, and for what reason (timeout, error, ...). The normal flags should be "--", indicating the session was closed by either end with no data remaining in buffers. See below "Session state at disconnection"

for more details.

- "actconn" is the total number of concurrent connections on the process when the session was logged. It is useful to detect when some per-process system limits have been reached. For instance, if actconn is close to 512 when multiple connection errors occur, chances are high that the system limits the process to use a maximum of 1024 file descriptors and that all of them are used. See section 3 "Global parameters" to find how to tune the system.

- "feconn" is the total number of concurrent connections on the frontend when the session was logged. It is useful to estimate the amount of resource required to sustain high loads, and to detect when the frontend's "maxconn" has been reached. Most often when this value increases by huge jumps, it is because there is congestion on the backend servers, but sometimes it can be caused by a denial of service attack.

- "beconn" is the total number of concurrent connections handled by the backend when the session was logged. It includes the total number of concurrent connections active on servers as well as the number of connections pending in queues. It is useful to estimate the amount of additional servers needed to support high loads for a given application. Most often when this value increases by huge jumps, it is because there is congestion on the backend servers, but sometimes it can be caused by a denial of service attack.

- "srv_conn" is the total number of concurrent connections still active on the Server when the session was logged. It can never exceed the server's configured "maxconn" parameter. If this value is very often close or equal to the server's "maxconn", it means that traffic regulation is involved a lot, meaning that either the server's maxconn value is too low, or that there aren't enough servers to process the load with an optimal response time. When only one of the server's "srv_conn" is high, it usually means that this server has some trouble causing the connections to take longer to be processed than on other servers.

- "retries" is the number of connection retries experienced by this session when trying to connect to the server. It must normally be zero, unless a server is being stopped at the same moment the connection was attempted. Several retries generally indicate either a network problem between haproxy and the server, or a misconfigured system backlog on the server preventing new connections from being queued. This field may optionally be prefixed with a '+' sign, indicating that the session has experienced a redispach after the maximal retry count has been reached on the initial server. In this case, the server name appearing in the log is the one the connection was redispached to, and not the first one, though both may sometimes be the same in case of hashing for instance. So as a general rule of thumb, when a '+' is present in front of the retry count, this count should not be attributed to the logged server.

- "srv_queue" is the total number of requests which were processed before this one in the server queue. It is zero when the request has not gone through the server queue. It makes it possible to estimate the approximate server's response time by dividing the time spent in queue by the number of requests in the queue. It is worth noting that if a session experiences a redispach and passes through two server queues, their positions will be cumulated. A request should not pass through both the server queue and the backend queue unless a redispach occurs.

- "backend_queue" is the total number of requests which were processed before this one in the backend's global queue. It is zero when the request has not gone through the global queue. It makes it possible to estimate the average queue length, which easily translates into a number of missing servers when divided by a server's "maxconn" parameter. It is worth noting that if a session experiences a redispach, it may pass twice in the backend's queue, and then both positions will be cumulated. A request should not pass

through both the server queue and the backend queue unless a redispach occurs.

8.2.3. HTTP log format

The HTTP format is the most complete and the best suited for HTTP proxies. It is enabled by when "option httplog" is specified in the frontend. It provides the same level of information as the TCP format with additional features which are specific to the HTTP protocol. Just like the TCP format, the log is usually emitted at the end of the session, unless "option logasap" is specified, which generally only makes sense for download sites. A session which matches the "monitor" rules will never be logged. It is also possible not to log sessions for which no data were sent by the client by specifying "option dontlognull" in the frontend. Successful connections will not be logged if "option dontlog-normal" is specified in the frontend.

Most fields are shared with the TCP log, some being different. A few fields may slightly vary depending on some configuration options. Those ones are marked with a star (*) after the field name below.

Example :

```
frontend http-in
mode http
option httplog
log global
default_backend bck

backend static
server srv1 127.0.0.1:8000
```

```
>>> Feb  6 12:14:14 localhost \
haproxy[14389]: 10.0.1.2:33317 [06/Feb/2009:12:14:14.655] http-in \
static/srv1 10/0/30/69/109 200 2750 - - ---- 1/1/1/0 0/0 {lwt.eu} \
{} "GET /index.html HTTP/1.1"
```

Field	Format	Extract from the example above
1	process_name '[pid]:'	haproxy[14389]:
2	client_ip ':' client_port	10.0.1.2:33317
3	'[' accept_date ']'	[06/Feb/2009:12:14:14.655]
4	frontend_name	http-in
5	backend_name '/' server_name	static/srv1
6	Tq '/' Tw '/' Tc '/' Tr '/' Tt*	10/0/30/69/109
7	status_code	200
8	bytes_read*	2750
9	captured_request_cookie	-
10	captured_response_cookie	-
11	termination_state	----
12	actconn '/' feconn '/' beconn '/' retries*	1/1/1/0
13	srv_queue '/' backend_queue	0/0
14	{' captured_request_headers* '}	{haproxy.lwt.eu}
15	{' captured_response_headers* '}	{}
16	,,, http_request ',,,'	"GET /index.html HTTP/1.1"

Detailed fields description :

- "client_ip" is the IP address of the client which initiated the TCP connection to haproxy. If the connection was accepted on a UNIX socket instead, the IP address would be replaced with the word "unix". Note that when the connection is accepted on a socket configured with "accept-proxy" and the PROXY protocol is correctly used, then the logs will reflect the forwarded connection's information.

- "client_port" is the TCP port of the client which initiated the connection.

If the connection was accepted on a UNIX socket instead, the port would be replaced with the ID of the accepting socket, which is also reported in the stats interface.

- "accept_date" is the exact date when the TCP connection was received by haproxy (which might be very slightly different from the date observed on the network if there was some queuing in the system's backlog). This is usually the same date which may appear in any upstream firewall's log. This does not depend on the fact that the client has sent the request or not.

- "frontend_name" is the name of the frontend (or listener) which received and processed the connection.

- "backend_name" is the name of the backend (or listener) which was selected to manage the connection to the server. This will be the same as the frontend if no switching rule has been applied.

- "server_name" is the name of the last server to which the connection was sent, which might differ from the first one if there were connection errors and a redispatch occurred. Note that this server belongs to the backend which processed the request. If the request was aborted before reaching a server, "<NOSRV>" is indicated instead of a server name. If the request was intercepted by the stats subsystem, "<STATS>" is indicated instead.

- "Tq" is the total time in milliseconds spent waiting for the client to send a full HTTP request, not counting data. It can be "-1" if the connection was aborted before a complete request could be received. It should always be very small because a request generally fits in one single packet. Large times here generally indicate network trouble between the client and haproxy. See "Timers" below for more details.

- "Tw" is the total time in milliseconds spent waiting in the various queues. It can be "-1" if the connection was aborted before reaching the queue. See "Timers" below for more details.

- "Tc" is the total time in milliseconds spent waiting for the connection to establish to the final server, including retries. It can be "-1" if the request was aborted before a connection could be established. See "Timers" below for more details.

- "Tr" is the total time in milliseconds spent waiting for the server to send a full HTTP response, not counting data. It can be "-1" if the request was aborted before a complete response could be received. It generally matches the server's processing time for the request, though it may be altered by the amount of data sent by the client to the server. Large times here on "GET" requests generally indicate an overloaded server. See "Timers" below for more details.

- "Tt" is the total time in milliseconds elapsed between the accept and the last close. It covers all possible processing. There is one exception, if "option logasap" was specified, then the time counting stops at the moment the log is emitted. In this case, a '+' sign is prepended before the value, indicating that the final one will be larger. See "Timers" below for more details.

- "status_code" is the HTTP status code returned to the client. This status is generally set by the server, but it might also be set by haproxy when the server cannot be reached or when its response is blocked by haproxy.

- "bytes_read" is the total number of bytes transmitted to the client when the log is emitted. This does include HTTP headers. If "option logasap" is specified, the this value will be prefixed with a '+' sign indicating that the final one may be larger. Please note that this value is a 64-bit counter, so log analysis tools must be able to handle it without overflowing.

- "captured_request_cookie" is an optional "name=value" entry indicating that the client had this cookie in the request. The cookie name and its maximum length are defined by the "capture cookie" statement in the frontend configuration. The field is a single dash ('-') when the option is not set. Only one cookie may be captured, it is generally used to track session ID exchanges between a client and a server to detect session crossing between clients due to application bugs. For more details, please consult the section "Capturing HTTP headers and cookies" below.

- "captured_response_cookie" is an optional "name=value" entry indicating that the server has returned a cookie with its response. The cookie name and its maximum length are defined by the "capture cookie" statement in the frontend configuration. The field is a single dash ('-') when the option is not set. Only one cookie may be captured, it is generally used to track session ID exchanges between a client and a server to detect session crossing between clients due to application bugs. For more details, please consult the section "Capturing HTTP headers and cookies" below.

- "termination_state" is the condition the session was in when the session ended. This indicates the session state, which side caused the end of session to happen, for what reason (timeout, error, ...), just like in TCP logs, and information about persistence operations on cookies in the last two characters. The normal flags should begin with "...", indicating the session was closed by either end with no data remaining in buffers. See below "Session state at disconnection" for more details.

- "actconn" is the total number of concurrent connections on the process when the session was logged. It is useful to detect when some per-process system limits have been reached. For instance, if actconn is close to 512 or 1024 when multiple connection errors occur, chances are high that the system limits the process to use a maximum of 1024 file descriptors and that all of them are used. See section 3 "Global parameters" to find how to tune the system.

- "feconn" is the total number of concurrent connections on the frontend when the session was logged. It is useful to estimate the amount of resource required to sustain high loads, and to detect when the frontend's "maxconn" has been reached. Most often when this value increases by huge jumps, it is because there is congestion on the backend servers, but sometimes it can be caused by a denial of service attack.

- "beconn" is the total number of concurrent connections handled by the backend when the session was logged. It includes the total number of concurrent connections active on servers as well as the number of connections pending in queues. It is useful to estimate the amount of additional servers needed to support high loads for a given application. Most often when this value increases by huge jumps, it is because there is congestion on the backend servers, but sometimes it can be caused by a denial of service attack.

- "srv_conn" is the total number of concurrent connections still active on the server when the session was logged. It can never exceed the server's configured "maxconn" parameter. If this value is very often close or equal to the server's "maxconn", it means that traffic regulation is involved a lot, meaning that either the server's maxconn value is too low, or that there aren't enough servers to process the load with an optimal response time. When only one of the server's "srv_conn" is high, it usually means that this server has some trouble causing the requests to take longer to be processed than on other servers.

- "retries" is the number of connection retries experienced by this session when trying to connect to the server. It must normally be zero, unless a server is being stopped at the same moment the connection was attempted. Frequent retries generally indicate either a network problem between

14366
14367
14368
14369
14370
14371
14372
14373
14374
14375
14376
14377
14378
14379
14380
14381
14382
14383
14384
14385
14386
14387
14388
14389
14390
14391
14392
14393
14394
14395
14396
14397
14398
14399
14400
14401
14402
14403
14404
14405
14406
14407
14408
14409
14410
14411
14412
14413
14414
14415
14416
14417
14418
14419
14420
14421
14422
14423
14424
14425
14426
14427
14428
14429
14430

haproxy and the server, or a misconfigured system backlog on the server preventing new connections from being queued. This field may optionally be prefixed with a '+' sign, indicating that the session has experienced a redispatch after the maximal retry count has been reached on the initial server. In this case, the server name appearing in the log is the one the connection was redispatched to, and not the first one, though both may sometimes be the same in case of hashing for instance. So as a general rule of thumb, when a '+' is present in front of the retry count, this count should not be attributed to the logged server.

- "srv_queue" is the total number of requests which were processed before this one in the server queue. It is zero when the request has not gone through the server queue. It makes it possible to estimate the approximate server's response time by dividing the time spent in queue by the number of requests in the queue. It is worth noting that if a session experiences a redispatch and passes through two server queues, their positions will be cumulated. A request should not pass through both the server queue and the backend queue unless a redispatch occurs.

- "backend_queue" is the total number of requests which were processed before this one in the backend's global queue. It is zero when the request has not gone through the global queue. It makes it possible to estimate the average queue length, which easily translates into a number of missing servers when divided by a server's "maxconn" parameter. It is worth noting that if a session experiences a redispatch, it may pass twice in the backend's queue, and then both positions will be cumulated. A request should not pass through both the server queue and the backend queue unless a redispatch occurs.

- "captured_request_headers" is a list of headers captured in the request due to the presence of the "capture request header" statement in the frontend. Multiple headers can be captured, they will be delimited by a vertical bar (|). When no capture is enabled, the braces do not appear, causing a shift of remaining fields. It is important to note that this field may contain spaces, and that using it requires a smarter log parser than when it's not used. Please consult the section "Capturing HTTP headers and cookies" below for more details.

- "captured_response_headers" is a list of headers captured in the response due to the presence of the "capture response header" statement in the frontend. Multiple headers can be captured, they will be delimited by a vertical bar (|). When no capture is enabled, the braces do not appear, causing a shift of remaining fields. It is important to note that this field may contain spaces, and that using it requires a smarter log parser than when it's not used. Please consult the section "Capturing HTTP headers and cookies" below for more details.

- "http_request" is the complete HTTP request line, including the method, request and HTTP version string. Non-printable characters are encoded (see below the section "Non-printable characters"). This is always the last field, and it is always delimited by quotes and is the only one which can contain quotes. If new fields are added to the log format, they will be added before this field. This field might be truncated if the request is huge and does not fit in the standard syslog buffer (1024 characters). This is the reason why this field must always remain the last one.

8.2.4. Custom log format

The directive log-format allows you to customize the logs in http mode and tcp mode. It takes a string as argument.

HAproxy understands some log format variables. % precedes log format variables. Variables can take arguments using braces ({}), and multiple arguments are

separated by commas within the braces. Flags may be added or removed by prefixing them with a '+' or '-' sign.

Special variable "%o" may be used to propagate its flags to all other variables on the same format string. This is particularly handy with quoted string formats ("%Q").

If a variable is named between square brackets ('[' .. ']') then it is used as a sample expression rule (see section 7.3). This it useful to add some less common information such as the client's SSL certificate's DN, or to log the key that would be used to store an entry into a stick table.

Note: spaces must be escaped. A space character is considered as a separator. In order to emit a verbatim ' ', it must be preceded by another ' ', resulting in '%%'. HAproxy will automatically merge consecutive separators.

Flags are :
* Q: quote a string
* X: hexadecimal representation (IPs, Ports, %Ts, %rt, %pid)

Example:

log-format %T\ %t\ Some\ Text
log-format %(+Q)o\ %t\ %s\ %{-Q}r

At the moment, the default HTTP format is defined this way :

log-format %ci:%cp\ [%t]\ %ft\ %b/%s\ %Tq/%Tw/%Tc/%Tr/%Tt\ %ST\ %B\ %CC\ \
%CS\ %tsc\ %ac/%fc/%bc/%sc/%rc\ %sq/%bq\ %hr\ %hs\ %{-Q}r

the default CLF format is defined this way :

log-format %(+Q)o\ %{-Q}ci\ -\ -\ [%T]\ %r\ %ST\ %B\ \"%\" \"%r\" \"%cp\" \
%ms\ %ft\ %b\ %s\ %Tq\ %Tw\ %Tc\ %Tr\ %Tt\ %tsc\ %ac\ %fc\ \
%bc\ %sc\ %rc\ %sq\ %bq\ %CC\ %CS\ %hr\ %hs\

and the default TCP format is defined this way :

log-format %ci:%cp\ [%t]\ %ft\ %b/%s\ %Tw/%Tc/%Tt\ %B\ %ts\ \
%ac/%fc/%bc/%sc/%rc\ %sq/%bq

Please refer to the table below for currently defined variables :

+-----+ R var field name (8.2.2 and 8.2.3 for description) type +-----+			
%o special variable, apply flags on all next var +-----+			
%B bytes_read (from server to client) numeric +-----+			
%CC captured_request_cookie string +-----+			
%CS captured_response_cookie string +-----+			
%H hostname string +-----+			
%HM HTTP method (ex: POST) string +-----+			
%HP HTTP request URI without query string (path) string +-----+			
%HQ HTTP request URI query string (ex: ?bar=baz) string +-----+			
%HU HTTP request URI (ex: /foo?bar=baz) string +-----+			
%ID unique-id string +-----+			
%ST status_code numeric +-----+			
%T gmt_date_time date +-----+			
%Tc Tc numeric +-----+			
%Tl Local_date_time date +-----+			
%Tq Tq numeric +-----+			
%Tr Tr numeric +-----+			
%Ts timestamp numeric +-----+			


```
14561 | | %Tt | | numeric |
14562 | | %Tw | | numeric |
14563 | | %U | | numeric |
14564 | | %a | | numeric |
14565 | | %b | | string |
14566 | | %bc | | numeric |
14567 | | %bi | | IP |
14568 | | %bp | | numeric |
14569 | | %bq | | numeric |
14570 | | %ci | | IP |
14571 | | %cp | | numeric |
14572 | | %f | | string |
14573 | | %fc | | numeric |
14574 | | %fi | | IP |
14575 | | %fp | | numeric |
14576 | | %ft | | string |
14577 | | %lc | | numeric |
14578 | | %lr | | string |
14579 | | %hl | | string list |
14580 | | %s | | string |
14581 | | %sl | | string list |
14582 | | %ms | | numeric |
14583 | | %pid | | numeric |
14584 | | %r | | string |
14585 | | %rc | | numeric |
14586 | | %rt | | numeric |
14587 | | %s | | string |
14588 | | %sc | | numeric |
14589 | | %sl | | IP |
14590 | | %sp | | numeric |
14591 | | %sq | | numeric |
14592 | | %ssl | | string |
14593 | | %sslv | | string |
14594 | | %t | | date |
14595 | | %ts | | date (with millisecond resolution) |
14596 | | %tsc | | string |
14597 | | %tsc | | string |
+-----+
R = Restrictions : H = mode http only ; S = SSL only

8.2.5. Error log format
-----
When an incoming connection fails due to an SSL handshake or an invalid PROXY
protocol header, haproxy will log the event using a shorter, fixed line format.
By default, logs are emitted at the LOG.INFO level, unless the option
"log-separate-errors" is set in the backend, in which case the LOG_ERR level
will be used. Connections on which no data are exchanged (eg: probes) are not
logged if the "dontlognull" option is set.

The format looks like this :
>>> Dec 3 18:27:14 localhost \
haproxy[6103]: 127.0.0.1:56059 [03/Dec/2012:17:35:10.380] frt/f1: \
Connection error during SSL handshake

Field Format Extract from the example above
1 process_name '[' pid ']:' haproxy[6103]:
2 client_ip ':' client_port 127.0.0.1:56059
3 '[' accept_date ']' [03/Dec/2012:17:35:10.380]
4 frontend_name "/" bind_name ":" frt/f1:
5 message Connection error during SSL handshake

These fields just provide minimal information to help debugging connection
```

```
14626 failures.
14627
14628
14629 8.3. Advanced logging options
14630 -----
14631
14632 Some advanced logging options are often looked for but are not easy to find out
14633 just by looking at the various options. Here is an entry point for the few
14634 options which can enable better logging. Please refer to the keywords reference
14635 for more information about their usage.
14636
14637
14638 8.3.1. Disabling logging of external tests
14639 -----
14640
14641 It is quite common to have some monitoring tools perform health checks on
14642 haproxy. Sometimes it will be a layer 3 load-balancer such as LVS or any
14643 commercial load-balancer, and sometimes it will simply be a more complete
14644 monitoring system such as Nagios. When the tests are very frequent, users often
14645 ask how to disable logging for those checks. There are three possibilities :
14646
14647 - if connections come from everywhere and are just TCP probes, it is often
14648 desired to simply disable logging of connections without data exchange, by
14649 setting "option dontlognull" in the frontend. It also disables logging of
14650 port scans, which may or may not be desired.
14651
14652 - if the connection come from a known source network, use "monitor-net" to
14653 declare this network as monitoring only. Any host in this network will then
14654 only be able to perform health checks, and their requests will not be
14655 logged. This is generally appropriate to designate a list of equipment
14656 such as other load-balancers.
14657
14658 - if the tests are performed on a known URI, use "monitor-uri" to declare
14659 this URI as dedicated to monitoring. Any host sending this request will
14660 only get the result of a health-check, and the request will not be logged.
14661
14662
14663 8.3.2. Logging before waiting for the session to terminate
14664 -----
14665
14666 The problem with logging at end of connection is that you have no clue about
14667 what is happening during very long sessions, such as remote terminal sessions
14668 or large file downloads. This problem can be worked around by specifying
14669 "option logasap" in the frontend. Haproxy will then log as soon as possible,
14670 just before data transfer begins. This means that in case of TCP, it will still
14671 log the connection status to the server, and in case of HTTP, it will log just
14672 after processing the server headers. In this case, the number of bytes reported
14673 is the number of header bytes sent to the client. In order to avoid confusion
14674 with normal logs, the total time field and the number of bytes are prefixed
14675 with a '+' sign which means that real numbers are certainly larger.
14676
14677
14678 8.3.3. Raising log level upon errors
14679 -----
14680
14681 Sometimes it is more convenient to separate normal traffic from errors logs,
14682 for instance in order to ease error monitoring from log files. When the option
14683 "log-separate-errors" is used, connections which experience errors, timeouts,
14684 retries, redispaches or HTTP status codes 5xx will see their syslog level
14685 raised from "info" to "err". This will help a syslog daemon store the log in
14686 a separate file. It is very important to keep the errors in the normal traffic
14687 file too, so that log ordering is not altered. You should also be careful if
14688 you already have configured your syslog daemon to store all logs higher than
14689 "notice" in an "admin" file, because the "err" level is higher than "notice".
14690
```

8.3.4. Disabling logging of successful connections

Although this may sound strange at first, some large sites have to deal with multiple thousands of logs per second and are experiencing difficulties keeping them intact for a long time or detecting errors within them. If the option "dontlog-normal" is set on the frontend, all normal connections will not be logged. In this regard, a normal connection is defined as one without any error, timeout, retry nor redispach. In HTTP, the status code is checked too, and a response with a status 5xx is not considered normal and will be logged too. Of course, doing is is really discouraged as it will remove most of the useful information from the logs. Do this only if you have no other alternative.

8.4. Timing events

Timers provide a great help in troubleshooting network problems. All values are reported in milliseconds (ms). These timers should be used in conjunction with the session termination flags. In TCP mode with "option tcplog" set on the frontend, 3 control points are reported under the form "Tw/Tc/Tt", and in HTTP mode, 5 control points are reported under the form "Tq/Tw/Tc/Tr/Tt" :

- Tq: total time to get the client request (HTTP mode only). It's the time elapsed between the moment the client connection was accepted and the moment the proxy received the last HTTP header. The value "-1" indicates that the end of headers (empty line) has never been seen. This happens when the client closes prematurely or times out.

- Tw: total time spent in the queues waiting for a connection slot. It accounts for backend queue as well as the server queues, and depends on the queue size, and the time needed for the server to complete previous requests. The value "-1" means that the request was killed before reaching the queue, which is generally what happens with invalid or denied requests.

- Tc: total time to establish the TCP connection to the server. It's the time elapsed between the moment the proxy sent the connection request, and the moment it was acknowledged by the server, or between the TCP SYN packet and the matching SYN/ACK packet in return. The value "-1" means that the connection never established.

- Tr: server response time (HTTP mode only). It's the time elapsed between the moment the TCP connection was established to the server and the moment the server sent its complete response headers. It purely shows its request processing time, without the network overhead due to the data transmission. It is worth noting that when the client has data to send to the server, for instance during a POST request, the time already runs, and this can distort apparent response time. For this reason, it's generally wise not to trust too much this field for POST requests initiated from clients behind an untrusted network. A value of "-1" here means that the last the response header (empty line) was never seen, most likely because the server timeout stroke before the server managed to process the request.

- Tt: total session duration time, between the moment the proxy accepted it and the moment both ends were closed. The exception is when the "logasap" option is specified. In this case, it only equals (Tq+Tw+Tc+Tr), and is prefixed with a '+' sign. From this field, we can deduce "Td", the data transmission time, by subtracting other timers when valid :

$$Td = Tt - (Tq + Tw + Tc + Tr)$$

Timers with "-1" values have to be excluded from this equation. In TCP mode, "Tq" and "Tr" have to be excluded too. Note that "Tt" can never be

negative.

These timers provide precious indications on trouble causes. Since the TCP protocol defines retransmit delays of 3, 6, 12... seconds, we know for sure that timers close to multiples of 3s are nearly always related to lost packets due to network problems (wires, negotiation, congestion). Moreover, if "Tt" is close to a timeout value specified in the configuration, it often means that a session has been aborted on timeout.

Most common cases :

- If "Tq" is close to 3000, a packet has probably been lost between the client and the proxy. This is very rare on local networks but might happen when clients are on far remote networks and send large requests. It may happen that values larger than usual appear here without any network cause. Sometimes, during an attack or just after a resource starvation has ended, haproxy may accept thousands of connections in a few milliseconds. The time spent accepting these connections will inevitably slightly delay processing of other connections, and it can happen that request times in the order of a few tens of milliseconds are measured after a few thousands of new connections have been accepted at once. Setting "option http-server-close" may display larger request times since "Tq" also measures the time spent waiting for additional requests.

- If "Tc" is close to 3000, a packet has probably been lost between the server and the proxy during the server connection phase. This value should always be very low, such as 1 ms on local networks and less than a few tens of ms on remote networks.

- If "Tr" is nearly always lower than 3000 except some rare values which seem to be the average majored by 3000, there are probably some packets lost between the proxy and the server.

- If "Tt" is large even for small byte counts, it generally is because neither the client nor the server decides to close the connection, for instance because both have agreed on a keep-alive connection mode. In order to solve this issue, it will be needed to specify "option httpclose" on either the frontend or the backend. If the problem persists, it means that the server ignores the "close" connection mode and expects the client to close. Then it will be required to use "option forceclose". Having the smallest possible 'Tt' is important when connection regulation is used with the "maxconn" option on the servers, since no new connection will be sent to the server until another one is released.

Other noticeable HTTP log cases ('xx' means any value to be ignored) :

Tq/Tw/Tc/Tr/+Tt The "option logasap" is present on the frontend and the log was emitted before the data phase. All the timers are valid except "Tt" which is shorter than reality.

-1/xx/xx/xx/Tt The client was not able to send a complete request in time or it aborted too early. Check the session termination flags then "timeout http-request" and "timeout client" settings.

Tq/-1/xx/xx/Tt It was not possible to process the request, maybe because servers were out of order, because the request was invalid or forbidden by ACL rules. Check the session termination flags.

Tq/Tw/-1/xx/Tt The connection could not establish on the server. Either it actively refused it or it timed out after Tt-(Tq+Tw) ms. Check the session termination flags, then check the "timeout connect" setting. Note that the tarpit action might return similar-looking patterns, with "Tw" equal to the time the client connection was maintained open.

14821 Tq/Tw/Tc/-/Tt The server has accepted the connection but did not return
14822 a complete response in time, or it closed its connection
14823 unexpectedly after $t-(tq+tw+tc)$ ms. Check the session
14824 termination flags, then check the "timeout server" setting.
14825
14826
14827
14828
14829
14830
14831
14832
14833
14834
14835
14836
14837
14838
14839
14840
14841
14842
14843
14844
14845
14846
14847
14848
14849
14850
14851
14852
14853
14854
14855
14856
14857
14858
14859
14860
14861
14862
14863
14864
14865
14866
14867
14868
14869
14870
14871
14872
14873
14874
14875
14876
14877
14878
14879
14880
14881
14882
14883
14884
14885

8.5. Session state at disconnection

TCP and HTTP logs provide a session termination indicator in the
"termination state" field, just before the number of active connections. It is
2-characters long in TCP mode, and is extended to 4 characters in HTTP mode,
each of which has a special meaning :

- On the first character, a code reporting the first event which caused the session to terminate :
- C : the TCP session was unexpectedly aborted by the client.
- S : the TCP session was unexpectedly aborted by the server, or the server explicitly refused it.
- P : the session was prematurely aborted by the proxy, because of a connection limit enforcement, because a DENY filter was matched, because of a security check which detected and blocked a dangerous error in server response which might have caused information leak (eg: cacheable cookie).
- L : the session was locally processed by haproxy and was not passed to a server. This is what happens for stats and redirects.
- R : a resource on the proxy has been exhausted (memory, sockets, source ports, ...). Usually, this appears during the connection phase, and system logs should contain a copy of the precise error. If this happens, it must be considered as a very serious anomaly which should be fixed as soon as possible by any means.
- I : an internal error was identified by the proxy during a self-check. This should NEVER happen, and you are encouraged to report any log containing this, because this would almost certainly be a bug. It would be wise to preventively restart the process after such an event too, in case it would be caused by memory corruption.
- D : the session was killed by haproxy because the server was detected as down and was configured to kill all connections when going down.
- U : the session was killed by haproxy on this backup server because an active server was detected as up and was configured to kill all backup connections when going up.
- K : the session was actively killed by an admin operating on haproxy.
- c : the client-side timeout expired while waiting for the client to send or receive data.
- s : the server-side timeout expired while waiting for the server to send or receive data.
- : normal session completion, both the client and the server closed with nothing left in the buffers.
- on the second character, the TCP or HTTP session state when it was closed :
- R : the proxy was waiting for a complete, valid REQUEST from the client

14886
14887
14888
14889
14890
14891
14892
14893
14894
14895
14896
14897
14898
14899
14900
14901
14902
14903
14904
14905
14906
14907
14908
14909
14910
14911
14912
14913
14914
14915
14916
14917
14918
14919
14920
14921
14922
14923
14924
14925
14926
14927
14928
14929
14930
14931
14932
14933
14934
14935
14936
14937
14938
14939
14940
14941
14942
14943
14944
14945
14946
14947
14948
14949
14950

(HTTP mode only). Nothing was sent to any server.

Q : the proxy was waiting in the QUEUE for a connection slot. This can only happen when servers have a 'maxconn' parameter set. It can also happen in the global queue after a redispatch consecutive to a failed attempt to connect to a dying server. If no redispatch is reported, then no connection attempt was made to any server.

C : the proxy was waiting for the CONNECTION to establish on the server. The server might at most have noticed a connection attempt.

H : the proxy was waiting for complete, valid response HEADERS from the server (HTTP only).

D : the session was in the DATA phase.

L : the proxy was still transmitting LAST data to the client while the server had already finished. This one is very rare as it can only happen when the client dies while receiving the last packets.

T : the request was tarptitted. It has been held open with the client during the whole "timeout tarpit" duration or until the client closed, both of which will be reported in the "Tw" timer.

- normal session completion after end of data transfer.
- the third character tells whether the persistence cookie was provided by the client (only in HTTP mode) :
- N : the client provided NO cookie. This is usually the case for new visitors, so counting the number of occurrences of this flag in the logs generally indicate a valid trend for the site frequentation.
- I : the client provided an INVALID cookie matching no known server. This might be caused by a recent configuration change, mixed cookies between HTTP/HTTPS sites, persistence conditionally ignored, or an attack.
- D : the client provided a cookie designating a server which was DOWN, so either "option persist" was used and the client was sent to this server, or it was not set and the client was redispatched to another server.
- V : the client provided a VALID cookie, and was sent to the associated server.
- E : the client provided a valid cookie, but with a last date which was older than what is allowed by the "maxidle" cookie parameter, so the cookie is consider EXPIRED and is ignored. The request will be redispatched just as if there was no cookie.
- O : the client provided a valid cookie, but with a first date which was older than what is allowed by the "maxlife" cookie parameter, so the cookie is consider too OLD and is ignored. The request will be redispatched just as if there was no cookie.
- U : a cookie was present but was not used to select the server because some other server selection mechanism was used instead (typically a "use-server" rule).
- : does not apply (no cookie set in configuration).
- the last character reports what operations were performed on the persistence cookie returned by the server (only in HTTP mode) :

N : NO cookie was provided by the server, and none was inserted either.
I : no cookie was provided by the server, and the proxy INSERTED one.
Note that in "cookie insert" mode, if the server provides a cookie, it will still be overwritten and reported as "I" here.

U : the proxy UPDATED the last date in the cookie that was presented by the client. This can only happen in insert mode with "maxidle". It happens every time there is activity at a different date than the date indicated in the cookie. If any other change happens, such as a redispatch, then the cookie will be marked as inserted instead.

P : a cookie was PROVIDED by the server and transmitted as-is.

R : the cookie provided by the server was REMITTEN by the proxy, which happens in "cookie rewrite" or "cookie prefix" modes.

D : the cookie provided by the server was DELETED by the proxy.

- : does not apply (no cookie set in configuration).

The combination of the two first flags gives a lot of information about what was happening when the session terminated, and why it did terminate. It can be helpful to detect server saturation, network troubles, local system resource starvation, attacks, etc...

The most common termination flags combinations are indicated below. They are alphabetically sorted, with the lowercase set just after the upper case for easier finding and understanding.

Flags	Reason
-------	--------

--	Normal termination.
----	---------------------

CC	The client aborted before the connection could be established to the server. This can happen when haproxy tries to connect to a recently dead (or unchecked) server, and the client aborts while haproxy is waiting for the server to respond or for "timeout connect" to expire.
----	-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------

CD	The client unexpectedly aborted during data transfer. This can be caused by a browser crash, by an intermediate equipment between the client and haproxy which decided to actively break the connection, by network routing issues between the client and haproxy, or by a keep-alive session between the server and the client terminated first by the client.
----	-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------

cd	The client did not send nor acknowledge any data for as long as the "timeout client" delay. This is often caused by network failures on the client side, or the client simply leaving the net uncetainly.
----	-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------

CH	The client aborted while waiting for the server to start responding. It might be the server taking too long to respond or the client clicking the 'Stop' button too fast.
----	---------------------------------------------------------------------------------------------------------------------------------------------------------------------------

CH	The "timeout client" stroke while waiting for client data during a POST request. This is sometimes caused by too large TCP MSS values for PPoE networks which cannot transport full-sized packets. It can also happen when client timeout is smaller than server timeout and the server takes too long to respond.
----	--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------

CQ	The client aborted while its session was queued, waiting for a server with enough empty slots to accept it. It might be that either all the servers were saturated or that the assigned server was taking too long a time to respond.
----	---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------

CR	The client aborted before sending a full HTTP request. Most likely the request was typed by hand using a telnet client, and aborted too early. The HTTP status code is likely a 400 here. Sometimes this might also be caused by an IDS killing the connection between haproxy and the client. "option http-ignore-probes" can be used to ignore connections without any data transfer.
----	-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------

cR	The "timeout http-request" stroke before the client sent a full HTTP request. This is sometimes caused by too large TCP MSS values on the client side for PPoE networks which cannot transport full-sized packets, or by clients sending requests by hand and not typing fast enough, or forgetting to enter the empty line at the end of the request. The HTTP status code is likely a 408 here. Note: recently, some browsers started to implement a "pre-connect" feature consisting in speculatively connecting to some recently visited web sites just in case the user would like to visit them. This results in many connections being established to web sites, which end up in 408 Request timeout if the timeout strikes first, or 400 Bad Request when the browser decides to close them first. These ones pollute the log and feed the error counters. Some versions of some browsers have even been reported to display the error code. It is possible to work around the undesirable effects of this behaviour by adding "option http-ignore-probes" in the frontend, resulting in connections with zero data transfer to be totally ignored. This will definitely hide the errors of people experiencing connectivity issues though.
----	-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------

CT	The client aborted while its session was tarpitted. It is important to check if this happens on valid requests, in order to be sure that no wrong tarpit rules have been written. If a lot of them happen, it might make sense to lower the "timeout tarpit" value to something closer to the average reported "tw" timer, in order not to consume resources for just a few attackers.
----	----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------

LR	The request was intercepted and locally handled by haproxy. Generally it means that this was a redirect or a stats request.
----	-----------------------------------------------------------------------------------------------------------------------------

SC	The server or an equipment between it and haproxy explicitly refused the TCP connection (the proxy received a TCP RST or an ICMP message in return). Under some circumstances, it can also be the network stack telling the proxy that the server is unreachable (eg: no route, or no ARP response on local network). When this happens in HTTP mode, the status code is likely a 502 or 503 here.
----	----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------

sc	The "timeout connect" stroke before a connection to the server could complete. When this happens in HTTP mode, the status code is likely a 503 or 504 here.
----	-------------------------------------------------------------------------------------------------------------------------------------------------------------

SD	The connection to the server died with an error during the data transfer. This usually means that haproxy has received an RST from the server or an ICMP message from an intermediate equipment while exchanging data with the server. This can be caused by a server crash or by a network issue on an intermediate equipment.
----	---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------

sd	The server did not send nor acknowledge any data for as long as the "timeout server" setting during the data phase. This is often caused by too short timeouts on L4 equipments before the server (firewalls, load-balancers, ...), as well as keep-alive sessions maintained between the client and the server expiring first on haproxy.
----	--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------

SH	The server aborted before sending its full HTTP response headers, or it crashed while processing the request. Since a server aborting at this moment is very rare, it would be wise to inspect its logs to control whether it crashed and why. The logged request may indicate a small set of faulty requests, demonstrating bugs in the application. Sometimes this might also be caused by an IDS killing the connection
----	----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------

15081 between haproxy and the server.

15082

15083 SH The "timeout server" stroke before the server could return its

15084 response headers. This is the most common anomaly, indicating too

15085 long transactions, probably caused by server or database saturation.

15086 The immediate workaround consists in increasing the "timeout server"

15087 setting, but it is important to keep in mind that the user experience

15088 will suffer from these long response times. The only long term

15089 solution is to fix the application.

15090

15091 sQ The session spent too much time in queue and has been expired. See

15092 the "timeout queue" and "timeout connect" settings to find out how to

15093 fix this if it happens too often. If it often happens massively in

15094 short periods, it may indicate general problems on the affected

15095 servers due to I/O or database congestion, or saturation caused by

15096 external attacks.

15097

15098 PC The proxy refused to establish a connection to the server because the

15099 process' socket limit has been reached while attempting to connect.

15100 The global "maxconn" parameter may be increased in the configuration

15101 so that it does not happen anymore. This status is very rare and

15102 might happen when the global "ulimit-n" parameter is forced by hand.

15103

15104 PD The proxy blocked an incorrectly formatted chunked encoded message in

15105 a request or a response, after the server has emitted its headers. In

15106 most cases, this will indicate an invalid message from the server to

15107 the client. Haproxy supports chunk sizes of up to 2GB - 1 (2147483647

15108 bytes). Any larger size will be considered as an error.

15109

15110 PH The proxy blocked the server's response, because it was invalid,

15111 incomplete, dangerous (cache control), or matched a security filter.

15112 In any case, an HTTP 502 error is sent to the client. One possible

15113 cause for this error is an invalid syntax in an HTTP header name

15114 containing unauthorized characters. It is also possible but quite

15115 rare, that the proxy blocked a chunked-encoding request from the

15116 client due to an invalid syntax, before the server responded. In this

15117 case, an HTTP 400 error is sent to the client and reported in the

15118 logs.

15119

15120 PR The proxy blocked the client's HTTP request, either because of an

15121 invalid HTTP syntax, in which case it returned an HTTP 400 error to

15122 the client, or because a deny filter matched, in which case it

15123 returned an HTTP 403 error.

15124

15125 PT The proxy blocked the client's request and has tarptitted its

15126 connection before returning it a 500 server error. Nothing was sent

15127 to the server. The connection was maintained open for as long as

15128 reported by the "tw" timer field.

15129

15130 RC A local resource has been exhausted (memory, sockets, source ports)

15131 preventing the connection to the server from establishing. The error

15132 logs will tell precisely what was missing. This is very rare and can

15133 only be solved by proper system tuning.

15134

15135 The combination of the two last flags gives a lot of information about how

15136 persistence was handled by the client, the server and by haproxy. This is very

15137 important to troubleshoot disconnections, when users complain they have to

15138 re-authenticate. The commonly encountered flags are :

15139

15140 -- Persistence cookie is not enabled.

15141

15142 NN No cookie was provided by the client, none was inserted in the

15143 response. For instance, this can be in insert mode with "postonly"

15144 set on a GET request.

15145

15146 II A cookie designating an invalid server was provided by the client,

15147 a valid one was inserted in the response. This typically happens when

15148 a "server" entry is removed from the configuration, since its cookie

15149 value can be presented by a client when no other server knows it.

15150

15151 NI No cookie was provided by the client, one was inserted in the

15152 response. This typically happens for first requests from every user

15153 in "insert" mode, which makes it an easy way to count real users.

15154

15155 VN A cookie was provided by the client, none was inserted in the

15156 response. This happens for most responses for which the client has

15157 already got a cookie.

15158

15159 VU A cookie was provided by the client, with a last visit date which is

15160 not completely up-to-date, so an updated cookie was provided in

15161 response. This can also happen if there was no date at all, or if

15162 there was a date but the "maxidle" parameter was not set, so that the

15163 cookie can be switched to unlimited time.

15164

15165 EI A cookie was provided by the client, with a last visit date which is

15166 too old for the "maxidle" parameter, so the cookie was ignored and a

15167 new cookie was inserted in the response.

15168

15169 OI A cookie was provided by the client, with a first visit date which is

15170 too old for the "maxlife" parameter, so the cookie was ignored and a

15171 new cookie was inserted in the response.

15172

15173 DI The server designated by the cookie was down, a new server was

15174 selected and a new cookie was emitted in the response.

15175

15176 VI The server designated by the cookie was not marked dead but could not

15177 be reached. A redispatch happened and selected another one, which was

15178 then advertised in the response.

15179

15180

15181

15182

15183

15184

15185

15186

15187

15188

15189

15190

15191

15192

15193

15194

15195

15196

15197

15198

15199

15200

15201

15202

15203

15204

15205

15206

15207

15208

15209

15210

8.6. Non-printable characters

In order not to cause trouble to log analysis tools or terminals during log consulting, non-printable characters are not sent as-is into log files, but are converted to the two-digits hexadecimal representation of their ASCII code, prefixed by the character '#'. The only characters that can be logged without being escaped are comprised between 32 and 126 (inclusive). Obviously, the escape character '#' itself is also encoded to avoid any ambiguity ("23"). It is the same for the character '.' which becomes "#22", as well as '{', '|' and '}' when logging headers.

Note that the space character (' ') is not encoded in headers, which can cause issues for tools relying on space count to locate fields. A typical header containing spaces is "User-Agent".

Last, it has been observed that some syslog daemons such as syslog-ng escape the quote ('"') with a backslash ('\'). The reverse operation can safely be performed since no quote may appear anywhere else in the logs.

8.7. Capturing HTTP cookies

Cookie capture simplifies the tracking a complete user session. This can be achieved using the "capture cookie" statement in the frontend. Please refer to section 4.2 for more details. Only one cookie can be captured, and the same cookie will simultaneously be checked in the request ("Cookie:" header) and in the response ("Set-Cookie:" header). The respective values will be reported in the HTTP logs at the "captured_request_cookie" and "captured_response_cookie"

locations (see section 8.2.3 about HTTP log format). When either cookie is not seen, a dash ('-') replaces the value. This way, it's easy to detect when a user switches to a new session for example, because the server will reassign it a new cookie. It is also possible to detect if a server unexpectedly sets a wrong cookie to a client, leading to session crossing.

Examples :

```
# capture the first cookie whose name starts with "ASPSESSION"
capture cookie ASPSESSION len 32

# capture the first cookie whose name is exactly "vgnvisitor"
capture cookie vgnvisitor= len 32
```

8.8. Capturing HTTP headers

Header captures are useful to track unique request identifiers set by an upper proxy, virtual host names, user-agents, POST content-length, referrers, etc. In the response, one can search for information about the response length, how the server asked the cache to behave, or an object location during a redirection.

Header captures are performed using the "capture request header" and "capture response header" statements in the frontend. Please consult their definition in section 4.2 for more details.

It is possible to include both request headers and response headers at the same time. Non-existent headers are logged as empty strings, and if one header appears more than once, only its last occurrence will be logged. Request headers are grouped within braces '{' and '}' in the same order as they were declared, and delimited with a vertical bar '|' without any space. Response headers follow the same representation, but are displayed after a space following the request headers block. These blocks are displayed just before the HTTP request in the logs.

As a special case, it is possible to specify an HTTP header capture in a TCP frontend. The purpose is to enable logging of headers which will be parsed in an HTTP backend if the request is then switched to this HTTP backend.

Example :

```
# This instance chains to the outgoing proxy
listen proxy-out
mode http
option httplog
option logasap
log global
server cache1 192.168.1.1:3128
```

```
# log the name of the virtual server
capture request header Host len 20

# log the amount of data uploaded during a POST
capture request header Content-Length len 10

# log the beginning of the referrer
capture request header Referer len 20

# server name (useful for outgoing proxies only)
capture response header Server len 20

# logging the content-length is useful with "option logasap"
capture response header Content-Length len 10

# log the expected cache behaviour on the response
capture response header Cache-Control len 8
```

```
15276
15277 # the Via header will report the next proxy's name
15278 capture response header Via len 20
15279
15280 # log the URL location during a redirection
15281 capture response header Location len 20
15282
15283
15284
15285
15286
15287
15288
15289
15290
15291
15292
15293
15294
15295
15296
15297
15298
15299
15300
15301
15302
15303
15304
15305
15306
15307
15308
15309
15310
15311
15312
15313
15314
15315
15316
15317
15318
15319
15320
15321
15322
15323
15324
15325
15326
15327
15328
15329
15330
15331
15332
15333
15334
15335
15336
15337
15338
15339
15340
```

```
>>> Aug 9 20:26:09 localhost \
haproxy[2022]: 127.0.0.1:34014 [09/Aug/2004:20:26:09] proxy-out \
proxy-out/cachel 0/0/0/162+162 200 +350 - - ---- 0/0/0/0 0/0 \
{fr.adserver.yahoo.co||http://fr.f416.mail.} {1864|private|} \
"GET http://fr.adserver.yahoo.com/"
```

```
>>> Aug 9 20:30:46 localhost \
haproxy[2022]: 127.0.0.1:34020 [09/Aug/2004:20:30:46] proxy-out \
proxy-out/cachel 0/0/0/182+182 200 +279 - - ---- 0/0/0/0 0/0 \
{w.ods.org|} {Formilux/0.1.8|3495|}| \
"GET http://trafic.lwt.eu/ HTTP/1.1"
```

```
>>> Aug 9 20:30:46 localhost \
haproxy[2022]: 127.0.0.1:34028 [09/Aug/2004:20:30:46] proxy-out \
proxy-out/cachel 0/0/2/126+128 301 +223 - - ---- 0/0/0/0 0/0 \
{www.sytadin.equipement.gouv.fr||http://trafic.lwt.eu/} \
{Apache|230|}|http://www.sytadin.} \
"GET http://www.sytadin.equipement.gouv.fr/ HTTP/1.1"
```

8.9. Examples of logs

These are real-world examples of logs accompanied with an explanation. Some of them have been made up by hand. The syslog part has been removed for better reading. Their sole purpose is to explain how to decipher them.

```
>>> haproxy[674]: 127.0.0.1:33318 [15/Oct/2003:08:31:57.130] px-http \
px-http/srv1 6559/0/7/147/6723 200 243 - - ---- 5/3/3/1/0 0/0 \
"HEAD / HTTP/1.0"
```

=> long request (6.5s) entered by hand through 'telnet'. The server replied in 147 ms, and the session ended normally ('----')

```
>>> haproxy[674]: 127.0.0.1:33319 [15/Oct/2003:08:31:57.149] px-http \
px-http/srv1 6559/1230/7/147/6870 200 243 - - ---- 324/239/239/99/0 \
0/9 "HEAD / HTTP/1.0"
```

=> Idem, but the request was queued in the global queue behind 9 other requests, and waited there for 1230 ms.

```
>>> haproxy[674]: 127.0.0.1:33320 [15/Oct/2003:08:32:17.654] px-http \
px-http/srv1 9/0/7/14/+30 200 +243 - - ---- 3/3/3/1/0 0/0 \
"GET /image.iso HTTP/1.0"
```

=> request for a long data transfer. The "logasap" option was specified, so the log was produced just before transferring data. The server replied in 14 ms, 243 bytes of headers were sent to the client, and total time from accept to first data byte is 30 ms.

```
>>> haproxy[674]: 127.0.0.1:33320 [15/Oct/2003:08:32:17.925] px-http \
px-http/srv1 9/0/7/14/30 502 243 - - PH-- 3/2/2/0/0 0/0 \
"GET /cgi-bin/bug.cgi? HTTP/1.0"
```

=> the proxy blocked a server response either because of an "rspdeny" or "rspideny" filter, or because the response was improperly formatted and not HTTP-compliant, or because it blocked sensitive information which risked being cached. In this case, the response is replaced with a "502

```
15341 bad gateway". The flags ("PH--") tell us that it was haproxy who decided
15342 to return the 502 and not the server.
15343
15344 >>> haproxy[18113]: 127.0.0.1:34548 [15/Oct/2003:15:18:55.798] px-http \
15345 px-http/<NOSRV> -1/-1/-1/-1/8490 -1 0 - - CR-- 2/2/2/0/0 0/0 ""
15346
15347 => the client never completed its request and aborted itself ("C---") after
15348 8.5s, while the proxy was waiting for the request headers ("R--").
15349 Nothing was sent to any server.
15350
15351 >>> haproxy[18113]: 127.0.0.1:34549 [15/Oct/2003:15:19:06.103] px-http \
15352 px-http/<NOSRV> -1/-1/-1/-1/50001 408 0 - - cR-- 2/2/2/0/0 0/0 ""
15353
15354 => The client never completed its request, which was aborted by the
15355 time-out ("C---") after 50s, while the proxy was waiting for the request
15356 headers ("R--"). Nothing was sent to any server, but the proxy could
15357 send a 408 return code to the client.
15358
15359 >>> haproxy[18989]: 127.0.0.1:34550 [15/Oct/2003:15:24:28.312] px-tcp \
15360 px-tcp/srv1 0/0/5007 0 cD 0/0/0/0/0 0/0
15361
15362 => This log was produced with "option tcplog". The client timed out after
15363 5 seconds ("c---").
15364
15365 >>> haproxy[18989]: 10.0.0.1:34552 [15/Oct/2003:15:26:31.462] px-http \
15366 px-http/srv1 3183/-1/-1/-1/11215 503 0 - - SC-- 205/202/202/115/3 \
15367 0/0 "HEAD / HTTP/1.0"
15368
15369 => The request took 3s to complete (probably a network problem), and the
15370 connection to the server failed ('SC--') after 4 attempts of 2 seconds
15371 (config says 'retries 3'), and no redispatch (otherwise we would have
15372 seen '+3'). Status code 503 was returned to the client. There were 115
15373 connections on this server, 202 connections on this proxy, and 205 on
15374 the global process. It is possible that the server refused the
15375 connection because of too many already established.
15376
15377 /*
15378 * Local variables:
15379 * fill-column: 79
15380 * End:
15381 */
15382
```