

```
1
2
3
4
5
6
7
8
9
10
11
12
13
14
15
16
17
18
19
20
21
22
23
24
25
26
27
28
29
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48
49
50
51
52
53
54
55
56
57
58
59
60
61
62
63
64
65
```

```
-----
haproxy Starter Guide
-----
version 1.6
```

This document is an introduction to HAProxy for all those who don't know it, as well as for those who want to re-discover it when they know older versions. Its primary focus is to provide users with all the elements to decide if HAProxy is the product they're looking for or not. Advanced users may find here some parts of solutions to some ideas they had just because they were not aware of a given new feature. Some sizing information are also provided, the product's lifecycle is explained, and comparisons with partially overlapping products are provided.

This document doesn't provide any configuration help nor hint, but it explains where to find the relevant documents. The summary below is meant to help you search sections by name and navigate through the document.

Note to documentation contributors :

This document is formatted with 80 columns per line, with even number of spaces for indentation and without tabs. Please follow these rules strictly so that it remains easily printable everywhere. If you add sections, please update the summary below for easier searching.

Summary

1. Available documentation

2. Quick introduction to load balancing and load balancers

3. Introduction to HAProxy What HAProxy is and is not

3.1. How HAProxy works

3.2. Basic features

3.3. Proxying

3.3.1. SSL

3.3.2. Monitoring

3.3.3. High availability

3.3.4. Load balancing

3.3.5. Stickiness

3.3.6. Sampling and converting information

3.3.7. Maps

3.3.8. ACLs and conditions

3.3.9. Content switching

3.3.10. Stick-tables

3.3.11. Formatted strings

3.3.12. HTTP rewriting and redirection

3.3.13. Server protection

3.3.14. Logging

3.3.15. Statistics

3.3.16. Advanced features

3.4. Management

3.4.1. System-specific capabilities

3.4.2. Scripting

3.4.3. Sizing

3.5. How to get HAProxy

3.6.

4. Companion products and alternatives

4.1. Apache HTTP server

4.2. NGINX

4.3. Varnish

4.4. Alternatives

1. Available documentation

The complete HAProxy documentation is contained in the following documents. Please ensure to consult the relevant documentation to save time and to get the most accurate response to your needs. Also please refrain from sending questions to the mailing list whose responses are present in these documents.

- intro.txt (this document) : it presents the basics of load balancing, HAProxy as a product, what it does, what it doesn't do, some known traps to avoid, some OS-specific limitations, how to get it, how it evolves, how to ensure you're running with all known fixes how to update it, complements and alternatives.

- management.txt : it explains how to start haproxy, how to manage it at runtime, how to manage it on multiple nodes, how to proceed with seamless upgrades.

- configuration.txt : the reference manual details all configuration keywords and their options. It is used when a configuration change is needed.

- architecture.txt : the architecture manual explains how to best architect a load-balanced infrastructure and how to interact with third party products.

- coding-style.txt : this is for developers who want to propose some code to the project. It explains the style to adopt for the code. It's not very strict and not all the code base completely respects it but contributions which diverge too much from it will be rejected.

- proxy-protocol.txt : this is the de-facto specification of the PROXY protocol which is implemented by HAProxy and a number of third party products.

- README : how to build haproxy from sources

2. Quick introduction to load balancing and load balancers

Load balancing consists in aggregating multiple components in order to achieve a total processing capacity above each component's individual capacity, without any intervention from the end user and in a scalable way. This results in more operations being performed simultaneously by the time it takes a component to perform only one. A single operation however will still be performed on a single component at a time and will not get faster than without load balancing. It always requires at least as many operations as available components and an efficient load balancing mechanism to make use of all components and to fully benefit from the load balancing. A good example of this is the number of lanes on a highway which allows as many cars to pass during the same time frame without increasing their individual speed.

Examples of load balancing :

- Process scheduling in multi-processor systems
- Link load balancing (eg: EtherChannel, Bonding)
- IP address load balancing (eg: ECMP, DNS roundrobin)
- Server load balancing (via load balancers)

The mechanism or component which performs the load balancing operation is called a load balancer. In web environments these components are called a "network load balancer", and more commonly a "load balancer" given that this activity is by far the best known case of load balancing.

A load balancer may act :

- at the link level : this is called link load balancing, and it consists in choosing what network link to send a packet to;
- at the network level : this is called network load balancing, and it consists in choosing what route a series of packets will follow;
- at the server level : this is called server load balancing and it consists in deciding what server will process a connection or request.

Two distinct technologies exist and address different needs, though with some overlapping. In each case it is important to keep in mind that load balancing consists in diverting the traffic from its natural flow and that doing so always requires a minimum of care to maintain the required level of consistency between all routing decisions.

The first one acts at the packet level and processes packets more or less individually. There is a 1-to-1 relation between input and output packets, so it is possible to follow the traffic on both sides of the load balancer using a regular network sniffer. This technology can be very cheap and extremely fast. It is usually implemented in hardware (ASICs) allowing to reach line rate, such as switches doing ECMP. Usually stateless, it can also be stateful (consider the session a packet belongs to and called layer4-LB or L4), may support DSR (direct server return, without passing through the LB again) if the packets were not modified, but provides almost no content awareness. This technology is very well suited to network-level load balancing, though it is sometimes used for very basic server load balancing at high speed.

The second one acts on session contents. It requires that the input streams is reassembled and processed as a whole. The contents may be modified, and the output stream is segmented into new packets. For this reason it is generally performed by proxies and they're often called layer 7 load balancers or L7. This implies that there are two distinct connections on each side, and that there is no relation between input and output packets sizes nor counts. Clients and servers are not required to use the same protocol (for example IPv4 vs IPv6, clear vs SSL). The operations are always stateful, and the return traffic must pass through the load balancer. The extra processing comes with a cost so it's not always possible to achieve line rate, especially with small packets. On the other hand, it offers wide possibilities and is generally achieved by pure software, even if embedded into hardware appliances. This technology is very well suited for server load balancing.

Packet-based load balancers are generally deployed in cut-through mode, so they are installed on the normal path of the traffic and divert it according to the configuration. The return traffic doesn't necessarily pass through the load balancer. Some modifications may be applied to the network destination address in order to direct the traffic to the proper destination. In this case, it is mandatory that the return traffic passes through the load balancer. If the routes doesn't make this possible, the load balancer may also replace the packets' source address with its own in order to force the return traffic to pass through it.

Proxy-based load balancers are deployed as a server with their own IP address and ports, without architecture changes. Sometimes this requires to perform some adaptations to the applications so that clients are properly directed to the load balancer's IP address and not directly to the server's. Some load balancers may have to adjust some servers' responses to make this possible (eg: the HTTP Location header field used in HTTP redirects). Some proxy-based load balancers may intercept traffic for an address they don't own, and spoof the client's address when connecting to the server. This allows them to be deployed as if they were a regular router or firewall, in a cut-through mode very similar to the packet based load balancers. This is particularly appreciated for products which combine both packet mode and proxy mode. In this case DSR is obviously still not possible and the return traffic still has to be routed back to the load balancer.

A very scalable layered approach would consist in having a front router which receives traffic from multiple load balanced links, and uses ECMP to distribute this traffic to a first layer of multiple stateful packet-based load balancers (L4). These L4 load balancers in turn pass the traffic to an even larger number of proxy-based load balancers (L7), which have to parse the contents to decide what server will ultimately receive the traffic.

The number of components and possible paths for the traffic increases the risk of failure; in very large environments, it is even normal to permanently have a few faulty components being fixed or replaced. Load balancing done without awareness of the whole stack's health significantly degrades availability. For this reason, any sane load balancer will verify that the components it intends to deliver the traffic to are still alive and reachable, and it will stop delivering traffic to faulty ones. This can be achieved using various methods.

The most common one consists in periodically sending probes to ensure the component is still operational. These probes are called "health checks". They must be representative of the type of failure to address. For example a ping-based check will not detect that a web server has crashed and doesn't listen to a port anymore, while a connection to the port will verify this, and a more advanced request may even validate that the server still works and that the database it relies on is still accessible. Health checks often involve a few retries to cover for occasional measuring errors. The period between checks must be small enough to ensure the faulty component is not used for too long after an error occurs.

Other methods consist in sampling the production traffic sent to a destination to observe if it is processed correctly or not, and to evince the components which return inappropriate responses. However this requires to sacrifice a part of the production traffic and this is not always acceptable. A combination of these two mechanisms provides the best of both worlds, with both of them being used to detect a fault, and only health checks to detect the end of the fault. A last method involves centralized reporting : a central monitoring agent periodically updates all load balancers about all components' state. This gives a global view of the infrastructure to all components, though sometimes with less accuracy or responsiveness. It's best suited for environments with many load balancers and many servers.

Layer 7 load balancers also face another challenge known as stickiness or persistence. The principle is that they generally have to direct multiple subsequent requests or connections from a same origin (such as an end user) to the same target. The best known example is the shopping cart on an online store. If each click leads to a new connection, the user must always be sent to the server which holds his shopping cart. Content-awareness makes it easier to spot some elements in the request to identify the server to deliver it to, but that's not always enough. For example if the source address is used as a key to pick a server, it can be decided that a hash-based algorithm will be used and that a given IP address will always be sent to the same server based on a divide of the address by the number of available servers. But if one server fails, the result changes and all users are suddenly sent to a different server and lose their shopping cart. The solution against this issue consists in memorizing the chosen target so that each time the same visitor is seen, he's directed to the same server regardless of the number of available servers. The information may be stored in the load balancer's memory, in which case it may have to be replicated to other load balancers if it's not alone, or it may be stored in the client's memory using various methods provided that the client is able to present this information back with every request (cookie insertion, redirection to a sub-domain, etc). This mechanism provides the extra benefit of not having to rely on unstable or unevenly distributed information (such as the source IP address). This is in fact the strongest reason to adopt a layer 7 load balancer instead of a layer 4 one.

In order to extract information such as a cookie, a host header field, a URL or whatever, a load balancer may need to decrypt SSL/TLS traffic and even

possibly to reencrypt it when passing it to the server. This expensive task explains why in some high-traffic infrastructures, sometimes there may be a lot of load balancers.

Since a layer 7 load balancer may perform a number of complex operations on the traffic (decrypt, parse, modify, match cookies, decide what server to send to, etc), it can definitely cause some trouble and will very commonly be accused of being responsible for a lot of trouble that it only revealed. Often it will be discovered that servers are unstable and periodically go up and down, or for web servers, that they deliver pages with some hard-coded links forcing the clients to connect directly to one specific server without passing via the load balancer, or that they take ages to respond under high load causing timeouts. That's why logging is an extremely important aspect of layer 7 load balancing. Once a trouble is reported, it is important to figure if the load balancer took a wrong decision and if so why so that it doesn't happen anymore.

3. Introduction to HAProxy

HAProxy is written "HAProxy" to designate the product, "haproxy" to designate the executable program, software package or a process, though both are commonly used for both purposes, and is pronounced H-A-Proxy. Very early it used to stand for "high availability proxy" and the name was written in two separate words, though by now it means nothing else than "HAProxy".

3.1. What HAProxy is and is not

HAProxy is :

- a TCP proxy : it can accept a TCP connection from a listening socket, connect to a server and attach these sockets together allowing traffic to flow in both directions;
- an HTTP reverse-proxy (called a "gateway" in HTTP terminology) : it presents itself as a server, receives HTTP requests over connections accepted on a listening TCP socket, and passes the requests from these connections to servers using different connections.
- an SSL terminator / initiator / offloader : SSL/TLS may be used on the connection coming from the client, on the connection going to the server, or even on both connections.
- a TCP normalizer : since connections are locally terminated by the operating system, there is no relation between both sides, so abnormal traffic such as invalid packets, flag combinations, window advertisements, sequence numbers, incomplete connections (SYN floods), or so will not be passed to the other side. This protects fragile TCP stacks from protocol attacks, and also allows to optimize the connection parameters with the client without having to modify the servers' TCP stack settings.
- an HTTP normalizer : when configured to process HTTP traffic, only valid complete requests are passed. This protects against a lot of protocol-based attacks. Additionally, protocol deviations for which there is a tolerance in the specification are fixed so that they don't cause problem on the servers (eg: multiple-line headers).
- an HTTP fixing tool : it can modify / fix / add / remove / rewrite the URL or any request or response header. This helps fixing interoperability issues in complex environments.
- a content-based switch : it can consider any element from the request to decide what server to pass the request or connection to. Thus it is possible

to handle multiple protocols over a same port (eg: http, https, ssh).

- a server load balancer : it can load balance TCP connections and HTTP requests. In TCP mode, load balancing decisions are taken for the whole connection. In HTTP mode, decisions are taken per request.

- a traffic regulator : it can apply some rate limiting at various points, protect the servers against overloading, adjust traffic priorities based on the contents, and even pass such information to lower layers and outer network components by marking packets.

- a protection against DDoS and service abuse : it can maintain a wide number of statistics per IP address, URL, cookie, etc and detect when an abuse is happening, then take action (slow down the offenders, block them, send them to outdated contents, etc).

- an observation point for network troubleshooting : due to the precision of the information reported in logs, it is often used to narrow down some network-related issues.

- an HTTP compression offloader : it can compress responses which were not compressed by the server, thus reducing the page load time for clients with poor connectivity or using high-latency, mobile networks.

HAProxy is not :

- an explicit HTTP proxy, ie, the proxy that browsers use to reach the internet. There are excellent open-source software dedicated for this task, such as Squid. However HAProxy can be installed in front of such a proxy to provide load balancing and high availability.
- a caching proxy : it will return as-is the contents its received from the server and will not interfere with any caching policy. There are excellent open-source software for this task such as Varnish. HAProxy can be installed in front of such a cache to provide SSL offloading, and scalability through smart load balancing.
- a data scrubber : it will not modify the body of requests nor responses.
- a web server : during startup, it isolates itself inside a chroot jail and drops its privileges, so that it will not perform any single file-system access once started. As such it cannot be turned into a web server. There are excellent open-source software for this such as Apache or Nginx, and HAProxy can be installed in front of them to provide load balancing and high availability.
- a packet-based load balancer : it will not see IP packets nor UDP datagrams, will not perform NAT or even less DSR. These are tasks for lower layers. Some kernel-based components such as IPVS (Linux Virtual Server) already do this pretty well and complement perfectly with HAProxy.

3.2. How HAProxy works

HAProxy is a single-threaded, event-driven, non-blocking engine combining a very fast I/O layer with a priority-based scheduler. As it is designed with a data forwarding goal in mind, its architecture is optimized to move data as fast as possible with the least possible operations. As such it implements a layered model offering bypass mechanisms at each level ensuring data don't reach higher levels when not needed. Most of the processing is performed in the kernel, and HAProxy does its best to help the kernel do the work as fast as possible by giving some hints or by avoiding certain operation when it guesses they could be grouped later. As a result, typical figures show 15% of the processing time spent in HAProxy versus 85% in the kernel in TCP or HTTP close mode, and about

391 30% for HAProxy versus 70% for the kernel in HTTP keep-alive mode.

392
393 A single process can run many proxy instances; configurations as large as
394 300000 distinct proxies in a single process were reported to run fine. Thus
395 there is usually no need to start more than one process for all instances.

396
397 It is possible to make HAProxy run over multiple processes, but it comes with
398 a few limitations. In general it doesn't make sense in HTTP close or TCP modes
399 because the kernel-side doesn't scale very well with some operations such as
400 connect(). It scales pretty well for HTTP keep-alive mode but the performance
401 that can be achieved out of a single process generally outperforms common needs
402 by an order of magnitude. It does however make sense when used as an SSL
403 offloader, and this feature is well supported in multi-process mode.

404
405 HAProxy only requires the haproxy executable and a configuration file to run.
406 For logging it is highly recommended to have a properly configured syslog daemon
407 and log rotations in place. The configuration files are parsed before starting,
408 then HAProxy tries to bind all listening sockets, and refuses to start if
409 anything fails. Past this point it cannot fail anymore. This means that there
410 are no runtime failures and that if it accepts to start, it will work until it
411 is stopped.

412
413 Once HAProxy is started, it does exactly 3 things :

- 414 - process incoming connections;
- 415 - periodically check the servers' status (known as health checks);
- 416 - exchange information with other haproxy nodes.

417
418 Processing incoming connections is by far the most complex task as it depends
419 on a lot of configuration possibilities, but it can be summarized as the 9 steps
420 below :

- 421 - accept incoming connections from listening sockets that belong to a
422 configuration entity known as a "frontend", which references one or multiple
423 listening addresses;
- 424 - apply the frontend-specific processing rules to these connections that may
425 result in blocking them, modifying some headers, or intercepting them to
426 execute some internal applets such as the statistics page or the CLI;
- 427 - pass these incoming connections to another configuration entity representing
428 a server farm known as a "backend", which contains the list of servers and
429 the load balancing strategy for this server farm;
- 430 - apply the backend-specific processing rules to these connections;
- 431 - decide which server to forward the connection to according to the load
432 balancing strategy;
- 433 - apply the backend-specific processing rules to the response data;
- 434 - apply the frontend-specific processing rules to the response data;
- 435 - emit a log to report what happened in fine details;
- 436 - in HTTP, loop back to the second step to wait for a new request, otherwise
437 close the connection.

438
439 Frontends and backends are sometimes considered as half-proxies, since they only
440 look at one side of an end-to-end connection; the frontend only cares about the
441 clients while the backend only cares about the servers. HAProxy also supports
442 full proxies which are exactly the union of a frontend and a backend. When HTTP
443 processing is desired, the configuration will generally be split into frontends
444
445

456 and backends as they open a lot of possibilities since any frontend may pass a
457 connection to any backend. With TCP-only proxies, using frontends and backends
458 rarely provides a benefit and the configuration can be more readable with full
459 proxies.

461 3.3. Basic features

462 -----

463
464 This section will enumerate a number of features that HAProxy implements, some
465 of which are generally expected from any modern load balancer, and some of
466 which are a direct benefit of HAProxy's architecture. More advanced features
467 will be detailed in the next section.

470 3.3.1. Basic features : Proxying

471 -----

472
473 Proxying is the action of transferring data between a client and a server over
474 two independent connections. The following basic features are supported by
475 HAProxy regarding proxying and connection management :

- 476 - Provide the server with a clean connection to protect them against any
477 client-side defect or attack;
- 478 - Listen to multiple IP address and/or ports, even port ranges;
- 479 - Transparent accept : intercept traffic targetting any arbitrary IP address
480 that doesn't even belong to the local system;
- 481 - Server port doesn't need to be related to listening port, and may even be
482 translated by a fixed offset (useful with ranges);
- 483 - Transparent connect : spoof the client's (or any) IP address if needed
484 when connecting to the server;
- 485 - Provide a reliable return IP address to the servers in multi-site LBs;
- 486 - Offload the server thanks to buffers and possibly short-lived connections
487 to reduce their concurrent connection count and their memory footprint;
- 488 - Optimize TCP stacks (eg: SACK), congestion control, and reduce RTT impacts;
- 489 - Support different protocol families on both sides (eg: IPv4/IPv6/Unix);
- 490 - Timeout enforcement : HAProxy supports multiple levels of timeouts depending
491 on the stage the connection is, so that a dead client or server, or an
492 attacker cannot be granted resources for too long;
- 493 - Protocol validation: HTTP, SSL, or payload are inspected and invalid
494 protocol elements are rejected, unless instructed to accept them anyway;
- 495 - Policy enforcement : ensure that only what is allowed may be forwarded;
- 496 - Both incoming and outgoing connections may be limited to certain network
497 namespaces (linux only), making it easy to build a cross-container,
498 multi-tenant load balancer;
- 499 - PROXY protocol presents the client's IP address to the server even for
500 non-HTTP traffic. This is an HAProxy extension that was adopted by a number
501 of third-party products by now, at least these ones at the time of writing :
 - 502 - client : haproxy, stud, stunnel, exaproxy, ELB, squid
 - 503 - server : haproxy, stud, postfix, exim, nginx, squid, node.js, varnish

3.3.2. Basic features : SSL

HAProxy's SSL stack is recognized as one of the most featureful according to Google's engineers (<http://istisfastyet.com/>). The most commonly used features making it quite complete are :

- SNI-based multi-hosting with no limit on sites count and focus on performance. At least one deployment is known for running 50000 domains with their respective certificates;
- support for wildcard certificates reduces the need for many certificates ;
- certificate-based client authentication with configurable policies on failure to present a valid certificate. This allows to present a different server farm to regenerate the client certificate for example;
- authentication of the backend server ensures the backend server is the real one and not a man in the middle;
- authentication with the backend server lets the backend server it's really the expected haproxy node that is connecting to it;
- TLS NPN and ALPN extensions make it possible to reliably offload SPDY/HTTP2 connections and pass them in clear text to backend servers;
- OSCP stapling further reduces first page load time by delivering inline an OSCP response when the client requests a Certificate Status Request;
- Dynamic record sizing provides both high performance and low latency, and significantly reduces page load time by letting the browser start to fetch new objects while packets are still in flight;
- permanent access to all relevant SSL/TLS layer information for logging, access control, reporting etc... These elements can be embedded into HTTP header or even as a PROXY protocol extension so that the offloaded server gets all the information it would have had if it performed the SSL termination itself.
- Detect, log and block certain known attacks even on vulnerable SSL libs, such as the Heartbleed attack affecting certain versions of OpenSSL.
- support for stateless session resumption (RFC 5077 TLS Ticket extension). TLS tickets can be updated from CLI which provides them means to implement Perfect Forward Secrecy by frequently rotating the tickets.

3.3.3. Basic features : Monitoring

HAProxy focuses a lot on availability. As such it cares about servers state, and about reporting its own state to other network components :

- Servers state is continuously monitored using per-server parameters. This ensures the path to the server is operational for regular traffic;
- Health checks support two hysteresis for up and down transitions in order to protect against state flapping;
- Checks can be sent to a different address/port/protocol : this makes it easy to check a single service that is considered representative of multiple ones, for example the HTTPS port for an HTTP+HTTPS server.
- Servers can track other servers and go down simultaneously : this ensures that servers hosting multiple services can fail atomically and that noone

will be sent to a partially failed server;

- Agents may be deployed on the server to monitor load and health : a server may be interested in reporting its load, operational status, administrative status independently from what health checks can see. By running a simple agent on the server, it's possible to consider the server's view of its own health in addition to the health checks validating the whole path;
- Various check methods are available : TCP connect, HTTP request, SMTP hello, SSL hello, LDAP, SQL, Redis, send/expect scripts, all with/without SSL;
- State change is notified in the logs and stats page with the failure reason (eg: the HTTP response received at the moment the failure was detected). An e-mail can also be sent to a configurable address upon such a change ;
- Server state is also reported on the stats interface and can be used to take routing decisions so that traffic may be sent to different farms depending on their sizes and/or health (eg: loss of an inter-DC link);
- HAProxy can use health check requests to pass information to the servers, such as their names, weight, the number of other servers in the farm etc... so that servers can adjust their response and decisions based on this knowledge (eg: postpone backups to keep more CPU available);
- Servers can use health checks to report more detailed state than just on/off (eg: I would like to stop, please stop sending new visitors);
- HAProxy itself can report its state to external components such as routers or other load balancers, allowing to build very complete multi-path and multi-layer infrastructures.

3.3.4. Basic features : High availability

Just like any serious load balancer, HAProxy cares a lot about availability to ensure the best global service continuity :

- Only valid servers are used ; the other ones are automatically evinced from load balancing farms ; under certain conditions it is still possible to force to use them though;
- Support for a graceful shutdown so that it is possible to take servers out of a farm without affecting any connection;
- Backup servers are automatically used when active servers are down and replace them so that sessions are not lost when possible. This also allows to build multiple paths to reach the same server (eg: multiple interfaces);
- Ability to return a global failed status for a farm when too many servers are down. This, combined with the monitoring capabilities makes it possible for an upstream component to choose a different LB node for a given service;
- Stateless design makes it easy to build clusters : by design, HAProxy does its best to ensure the highest service continuity without having to store information that could be lost in the event of a failure. This ensures that a takeover is the most seamless possible;
- Integrates well with standard VRRP daemon keepalived : HAProxy easily tells keepalived about its state and copes very well with floating virtual IP addresses. Note: only use IP redundancy protocols (VRRP/CARP) over cluster-based solutions (Heartbeat, ...) as they're the ones offering the fastest, most seamless, and most reliable switchover.

3.3.5. Basic features : Load balancing

HAProxy offers a fairly complete set of load balancing features, most of which are unfortunately not available in a number of other load balancing products :

- no less than 9 load balancing algorithms are supported, some of which apply to input data to offer an infinite list of possibilities. The most common ones are round-robin (for short connections, pick each server in turn), leastconn (for long connections, pick the least recently used of the servers with the lowest connection count), source (for SSL farms or terminal server farms, the server directly depends on the client's source address), uri (for HTTP caches, the server directly depends on the HTTP URI), hdr (the server directly depends on the contents of a specific HTTP header field), first (for short-lived virtual machines, all connections are packed on the smallest possible subset of servers so that unused ones can be powered down);
- all algorithms above support per-server weights so that it is possible to accommodate from different server generations in a farm, or direct a small fraction of the traffic to specific servers (debug mode, running the next version of the software, etc);
- dynamic weights are supported for round-robin, leastconn and consistent hashing ; this allows server weights to be modified on the fly from the CLI or even by an agent running on the server;
- slow-start is supported whenever a dynamic weight is supported; this allows a server to progressively take the traffic. This is an important feature for fragile application servers which require to compile classes at runtime as well as cold caches which need to fill up before being run at full throttle;
- hashing can apply to various elements such as client's source address, URL components, query string element, header field values, POST parameter, RDP cookie;
- consistent hashing protects server farms against massive redistribution when adding or removing servers in a farm. That's very important in large cache farms and it allows slow-start to be used to refill cold caches;
- a number of internal metrics such as the number of connections per server, per backend, the amount of available connection slots in a backend etc makes it possible to build very advanced load balancing strategies.

3.3.6. Basic features : Stickiness

Application load balancing would be useless without stickiness. HAProxy provides a fairly comprehensive set of possibilities to maintain a visitor on the same server even across various events such as server addition/removal, down/up cycles, and some methods are designed to be resistant to the distance between multiple load balancing nodes in that they don't require any replication :

- stickiness information can be individually matched and learned from different places if desired. For example a JSESSIONID cookie may be matched both in a cookie and in the URL. Up to 8 parallel sources can be learned at the same time and each of them may point to a different stick-table;
- stickiness information can come from anything that can be seen within a request or response, including source address, TCP payload offset and length, HTTP query string elements, header field values, cookies, and so on...

- stick-tables are replicated between all nodes in a multi-master fashion ;
- commonly used elements such as SSL-ID or RDP cookies (for TSE farms) are directly accessible to ease manipulation;
- all sticking rules may be dynamically conditioned by ACLs;
- it is possible to decide not to stick to certain servers, such as backup servers, so that when the nominal server comes back, it automatically takes the load back. This is often used in multi-path environments;
- in HTTP it is often preferred not to learn anything and instead manipulate a cookie dedicated to stickiness. For this, it's possible to detect, rewrite, insert or prefix such a cookie to let the client remember what server was assigned;
- the server may decide to change or clean the stickiness cookie on logout, so that leaving visitors are automatically unbound from the server;
- using ACL-based rules it is also possible to selectively ignore or enforce stickiness regardless of the server's state; combined with advanced health checks, that helps admins verify that the server they're installing is up and running before presenting it to the whole world;
- an innovative mechanism to set a maximum idle time and duration on cookies ensures that stickiness can be smoothly stopped on devices which are never closed (smartphones, TVs, home appliances) without having to store them on persistent storage;
- multiple server entries may share the same stickiness keys so that stickiness is not lost in multi-path environments when one path goes down;
- soft-stop ensures that only users with stickiness information will continue to reach the server they've been assigned to but no new users will go there.

3.3.7. Basic features : Sampling and converting information

HAProxy supports information sampling using a wide set of "sample fetch functions". The principle is to extract pieces of information known as samples, for immediate use. This is used for stickiness, to build conditions, to produce information in logs or to enrich HTTP headers.

Samples can be fetched from various sources :

- constants : integers, strings, IP addresses, binary blocks;
- the process : date, environment variables, server/frontend/backend/process state, byte/connection counts/rates, queue length, random generator, ...
- variables : per-session, per-request, per-response variables;
- the client connection : source and destination addresses and ports, and all related statistics counters;
- the SSL client session : protocol, version, algorithm, cipher, key size, session ID, all client and server certificate fields, certificate serial, SNI, ALPN, NPN, client support for certain extensions;
- request and response buffers contents : arbitrary payload at offset/length, data length, RDP cookie, decoding of SSL hello type, decoding of TLS SNI;
- HTTP (request and response) : method, URI, path, query string arguments, status code, headers values, positional header value, cookies, captures,

authentication, body elements;

A sample may then pass through a number of operators known as "converters" to experience some transformation. A converter consumes a sample and produces a new one, possibly of a completely different type. For example, a converter may be used to return only the integer length of the input string, or could turn a string to upper case. Any arbitrary number of converters may be applied in series to a sample before final use. Among all available sample converters, the following ones are the most commonly used :

- arithmetic and logic operators : they make it possible to perform advanced computation on input data, such as computing ratios, percentages or simply converting from one unit to another one;
- IP address masks are useful when some addresses need to be grouped by larger networks;
- data representation : url-decode, base64, hex, JSON strings, hashing;
- string conversion : extract substrings at fixed positions, fixed length, extract specific fields around certain delimiters, extract certain words, change case, apply regex-based substitution ;
- date conversion : convert to http date format, convert local to UTC and conversely, add or remove offset;
- lookup an entry in a stick table to find statistics or assigned server;
- map-based key-to-value conversion from a file (mostly used for geolocation).

3.3.8. Basic features : Maps

Maps are a powerful type of converter consisting in loading a two-columns file into memory at boot time, then looking up each input sample from the first column and either returning the corresponding pattern on the second column if the entry was found, or returning a default value. The output information also being a sample, it can in turn experience other transformations including other map lookups. Maps are most commonly used to translate the client's IP address to an AS number or country code since they support a longest match for network addresses but they can be used for various other purposes.

Part of their strength comes from being updatable on the fly either from the CLI or from certain actions using other samples, making them capable of storing and retrieving information between subsequent accesses. Another strength comes from the binary tree based indexation which makes them extremely fast even when they contain hundreds of thousands of entries, making geolocation very cheap and easy to set up.

3.3.9. Basic features : ACLs and conditions

Most operations in HAProxy can be made conditional. Conditions are built by combining multiple ACLs using logic operators (AND, OR, NOT). Each ACL is a series of tests based on the following elements :

- a sample fetch method to retrieve the element to test ;
- an optional series of converters to transform the element ;
- a list of patterns to match against ;
- a matching method to indicate how to compare the patterns with the sample

For example, the sample may be taken from the HTTP "Host" header, it could then be converted to lower case, then matched against a number of regex patterns using the regex matching method.

Technically, ACLs are built on the same core as the maps, they share the exact same internal structure, pattern matching methods and performance. The only real difference is that instead of returning a sample, they only return "found" or "not found". In terms of usage, ACL patterns may be declared inline in the configuration file and do not require their own file. ACLs may be named for ease of use or to make configurations understandable. A named ACL may be declared multiple times and it will evaluate all definitions in turn until one matches.

About 13 different pattern matching methods are provided, among which IP address mask, integer ranges, substrings, regex. They work like functions, and just like with any programming language, only what is needed is evaluated, so when a condition involving an OR is already true, next ones are not evaluated, and similarly when a condition involving an AND is already false, the rest of the condition is not evaluated.

There is no practical limit to the number of declared ACLs, and a handful of commonly used ones are provided. However experience has shown that setups using a lot of named ACLs are quite hard to troubleshoot and that sometimes using anonymous ACLs inline is easier as it requires less references out of the scope being analysed.

3.3.10. Basic features : Content switching

HAProxy implements a mechanism known as content-based switching. The principle is that a connection or request arrives on a frontend, then the information carried with this request or connection are processed, and at this point it is possible to write ACLs-based conditions making use of these information to decide what backend will process the request. Thus the traffic is directed to one backend or another based on the request's contents. The most common example consists in using the Host header and/or elements from the path (sub-directories or file-name extensions) to decide whether an HTTP request targets a static object or the application, and to route static objects traffic to a backend made of fast and light servers, and all the remaining traffic to a more complex application server, thus constituting a fine-grained virtual hosting solution. This is quite convenient to make multiple technologies coexist as a more global solution.

Another use case of content-switching consists in using different load balancing algorithms depending on various criteria. A cache may use a URI hash while an application would use round robin.

Last but not least, it allows multiple customers to use a small share of a common resource by enforcing per-backend (thus per-customer connection limits).

Content switching rules scale very well, though their performance may depend on the number and complexity of the ACLs in use. But it is also possible to write dynamic content switching rules where a sample value directly turns into a backend name and without making use of ACLs at all. Such configurations have been reported to work fine at least with 300000 backends in production.

3.3.11. Basic features : Stick-tables

Stick-tables are commonly used to store stickiness information, that is, to keep a reference to the server a certain visitor was directed to. The key is then the identifier associated with the visitor (its source address, the SSL ID of the connection, an HTTP or RDP cookie, the customer number extracted from the URL or

911 from the payload, ...) and the stored value is then the server's identifier.
912
913 Stick tables may use 3 different types of samples for their keys : integers,
914 strings and addresses. Only one stick-table may be referenced in a proxy, and it
915 is designated everywhere with the proxy name. Up to 8 key may be tracked in
916 parallel. The server identifier is committed during request or response
917 processing once both the key and the server are known.
918

919 Stick-table contents may be replicated in active-active mode with other HAProxy
920 nodes known as "peers" as well as with the new process during a reload operation
921 so that all load balancing nodes share the same information and take the same
922 routing decision if a client's requests are spread over multiple nodes.

923
924 Since stick-tables are indexed on what allows to recognize a client, they are
925 often also used to store extra information such as per-client statistics. The
926 extra statistics take some extra space and need to be explicitly declared. The
927 type of statistics that may be stored includes the input and output bandwidth,
928 the number of concurrent connections, the connection rate and count over a
929 period, the amount and frequency of errors, some specific tags and counters,
930 etc... In order to support keeping such information without being forced to
931 stick to a given server, a special "tracking" feature is implemented and allows
932 to track up to 3 simultaneous keys from different tables at the same time
933 regardless of stickiness rules. Each stored statistics may be searched, dumped
934 and cleared from the CLI and adds to the live troubleshooting capabilities.
935

936 While this mechanism can be used to surclass a returning visitor or to adjust
937 the delivered quality of service depending on good or bad behaviour, it is
938 mostly used to fight against service abuse and more generally DDoS as it allows
939 to build complex models to detect certain bad behaviours at a high processing
940 speed.
941

942 3.3.12. Basic features : Formated strings

943 -----
944
945 There are many places where HAProxy needs to manipulate character strings, such
946 as logs, redirects, header additions, and so on. In order to provide the
947 greatest flexibility, the notion of formated strings was introduced, initially
948 for logging purposes, which explains why it's still called "log-format". These
949 strings contain escape characters allowing to introduce various dynamic data
950 including variables and sample fetch expressions into strings, and even to
951 adjust the encoding while the result is being turned into a string (for example,
952 adding quotes). This provides a powerful way to build header contents or to
953 customize log lines. Additionally, in order to remain simple to build most
954 common strings, about 50 special tags are provided as shortcuts for information
955 commonly used in logs.
956

957 3.3.13. Basic features : HTTP rewriting and redirection

958 -----
959
960 Installing a load balancer in front of an application that was never designed
961 for this can be a challenging task without the proper tools. One of the most
962 commonly requested operation in this case is to adjust requests and response
963 headers to make the load balancer appear as the origin server and to fix hard
964 coded information. This comes with changing the path in requests (which is
965 strongly advised against), modifying Host header field, modifying the Location
966 response header field for redirects, modifying the path and domain attribute
967 for cookies, and so on. It also happens that a number of servers are somewhat
968 verbose and tend to leak too much information in the response, making them more
969 vulnerable to targetted attacks. While it's theoretically not the role of a load
970 balancer to clean this up, in practice it's located at the best place in the
971 infrastructure to guarantee that everything is cleaned up.
972
973 Similarly, sometimes the load balancer will have to intercept some requests and
974
975

976 respond with a redirect to a new target URL. While some people tend to confuse
977 redirects and rewriting, these are two completely different concepts, since the
978 rewriting makes the client and the server see different things (and disagree on
979 the location of the page being visited) while redirects ask the client to visit
980 the new URL so that it sees the same location as the server.
981

982 In order to do this, HAProxy supports various possibilities for rewriting and
983 redirect, among which :

- 984 - regex-based URL and header rewriting in requests and responses. Regex are
985 the most commonly used tool to modify header values since they're easy to
986 manipulate and well understood;
- 987 - headers may also be appended, deleted or replaced based on formatted strings
988 so that it is possible to pass information there (eg: client side TLS
989 algorithm and cipher);
- 990 - HTTP redirects can use any 3xx code to a relative, absolute, or completely
991 dynamic (formatted string) URI;
- 992 - HTTP redirects also support some extra options such as setting or clearing
993 a specific cookie, dropping the query string, appending a slash if missing,
994 and so on;
- 995 - all operations support ACL-based conditions;

996 3.3.14. Basic features : Server protection

997 -----
998
999 HAProxy does a lot to maximize service availability, and for this it deploys
1000 large efforts to protect servers against overloading and attacks. The first
1001 and most important point is that only complete and valid requests are forwarded
1002 to the servers. The initial reason is that HAProxy needs to find the protocol
1003 elements it needs to stay synchronized with the byte stream, and the second
1004 reason is that until the request is complete, there is no way to know if some
1005 elements will change its semantics. The direct benefit from this is that servers
1006 are not exposed to invalid or incomplete requests. This is a very effective
1007 protection against slowloris attacks, which have almost no impact on HAProxy.

1008 Another important point is that HAProxy contains buffers to store requests and
1009 responses, and that by only sending a request to a server when it's complete and
1010 by reading the whole response very quickly from the local network, the server
1011 side connection is used for a very short time and this preserves server
1012 resources as much as possible.

1013 A direct extension to this is that HAProxy can artificially limit the number of
1014 concurrent connections or outstanding requests to a server, which guarantees
1015 that the server will never be overloaded even if it continuously runs at 100% of
1016 its capacity during traffic spikes. All excess requests will simply be queued to
1017 be processed when one slot is released. In the end, this huge resource savings
1018 most often ensures so much better server response times that it ends up actually
1019 being faster than by overloading the server. Queued requests may be redispached
1020 to other servers, or even aborted in queue when the client aborts, which also
1021 protects the servers against the "reload effect", where each click on "reload"
1022 by a visitor on a slow-loading page usually induces a new request and maintains
1023 the server in an overloaded state.

1024 The slow-start mechanism also protects restarting servers against high traffic
1025 levels while they're still finalizing their startup or compiling some classes.

1026 Regarding the protocol-level protection, it is possible to relax the HTTP parser
1027 to accept non standard-compliant but harmless requests or responses and even to
1028 fix them. This allows bogus applications to be accessible while a fix is being
1029 developed. In parallel, offending messages are completely captured with a
1030

1041 detailed report that help developers spot the issue in the application. The most
1042 dangerous protocol violations are properly detected and dealt with and fixed.
1043 For example malformed requests or responses with two Content-length headers are
1044 either fixed if the values are exactly the same, or rejected if they differ,
1045 since it becomes a security problem. Protocol inspection is not limited to HTTP,
1046 it is also available for other protocols like TLS or RDP.

1047
1048 When a protocol violation or attack is detected, there are various options to
1049 respond to the user, such as returning the common "HTTP 400 bad request",
1050 closing the connection with a TCP reset, faking an error after a long delay
1051 ("tarpit") to confuse the attacker. All of these contribute to protecting the
1052 servers by discouraging the offending client from pursuing an attack that
1053 becomes very expensive to maintain.

1054 HAProxy also proposes some more advanced options to protect against accidental
1055 data leaks and session crossing. Not only it can log suspicious server responses
1056 but it will also log and optionally block a response which might affect a given
1057 visitors' confidentiality. One such example is a cacheable cookie appearing in a
1058 cacheable response and which may result in an intermediary cache to deliver it
1059 to another visitor, causing an accidental session sharing.

3.3.15. Basic features : Logging

1061
1062
1063
1064
1065
1066 Logging is an extremely important feature for a load balancer, first because a
1067 load balancer is often accused of the trouble it reveals, and second because it
1068 is placed at a critical point in an infrastructure where all normal and abnormal
1069 activity needs to be analysed and correlated with other components.

1070 HAProxy provides very detailed logs, with millisecond accuracy and the exact
1071 connection accept time that can be searched in firewalls logs (eg: for NAT
1072 correlation). By default, TCP and HTTP logs are quite detailed an contain
1073 everything needed for troubleshooting, such as source IP address and port,
1074 frontend, backend, server, timers (request receipt duration, queue duration,
1075 connection setup time, response headers time, data transfer time), global
1076 process state, connection counts, queue status, retries count, detailed
1077 stickiness actions and disconnect reasons, header captures with a safe output
1078 encoding. It is then possible to extend or replace this format to include any
1079 sampled data, variables, captures, resulting in very detailed information. For
1080 example it is possible to log the number cumulated requests for this client or
1081 the number of different URLs for the client.

1082
1083
1084 The log level may be adjusted per request using standard ACLs, so it is possible
1085 to automatically silent some logs considered as pollution and instead raise
1086 warnings when some abnormal behaviour happen for a small part of the traffic
1087 (eg: too many URLs or HTTP errors for a source address). Administrative logs are
1088 also emitted with their own levels to inform about the loss or recovery of a
1089 server for example.

1090 Each frontend and backend may use multiple independant log outputs, which eases
1091 multi-tenancy. Logs are preferably sent over UDP, maybe JSON-encoded, and are
1092 truncated after a configurable line length in order to guarantee delivery.

3.3.16. Basic features : Statistics

1093
1094
1095
1096
1097
1098
1099 HAProxy provides a web-based statistics reporting interface with authentication,
1100 security levels and scopes. It is thus possible to provide each hosted customer
1101 with his own page showing only his own instances. This page can be located in a
1102 hidden URL part of the regular web site so that no new port needs to be opened.
1103 This page may also report the availability of other HAProxy nodes so that it is
1104 easy to spot if everything works as expected at a glance. The view is synthetic
1105 with a lot of details accessible (such as error causes, last access and last

1106 change duration, etc), which are also accessible as a CSV table that other tools
1107 may import to draw graphs. The page may self-refresh to be used as a monitoring
1108 page on a large display. In administration mode, the page also allows to change
1109 server state to ease maintenance operations.

3.4. Advanced features

3.4.1. Advanced features : Management

1110
1111
1112
1113
1114
1115
1116
1117
1118 HAProxy is designed to remain extremely stable and safe to manage in a regular
1119 production environment. It is provided as a single executable file which doesn't
1120 require any installation process. Multiple versions can easily coexist, meaning
1121 that it's possible (and recommended) to upgrade instances progressively by
1122 order of criticality instead of migrating all of them at once. Configuration files
1123 are easily versionned. Configuration checking is done off-line so it doesn't
1124 require to restart a service that will possibly fail. During configuration
1125 checks, a number of advanced mistakes may be detected (eg: for example, a rule
1126 hiding another one, or stickiness that will not work) and detailed warnings and
1127 configuration hints are proposed to fix them. Backwards configuration file
1128 compatibility goes very far away in time, with version 1.5 still fully
1129 supporting configurations for versions 1.1 written 13 years before, and 1.6
1130 only dropping support for almost unused, obsolete keywords that can be done
1131 differently. The configuration and software upgrade mechanism is smooth and non
1132 disruptive in that it allows old and new processes to coexist on the system,
1133 each handling its own connections. System status, build options and library
1134 compatibility are reported on startup.

1135
1136
1137
1138
1139
1140
1141
1142
1143
1144
1145
1146
1147
1148
1149
1150
1151
1152
1153
1154
1155
1156
1157
1158
1159
1160
1161
1162
1163
1164
1165
1166
1167
1168
1169
1170
1171
1172
1173
1174
1175
1176
1177
1178
1179
1180
1181
1182
1183
1184
1185
1186
1187
1188
1189
1190
1191
1192
1193
1194
1195
1196
1197
1198
1199
1200
1201
1202
1203
1204
1205
1206
1207
1208
1209
1210
1211
1212
1213
1214
1215
1216
1217
1218
1219
1220
1221
1222
1223
1224
1225
1226
1227
1228
1229
1230
1231
1232
1233
1234
1235
1236
1237
1238
1239
1240
1241
1242
1243
1244
1245
1246
1247
1248
1249
1250
1251
1252
1253
1254
1255
1256
1257
1258
1259
1260
1261
1262
1263
1264
1265
1266
1267
1268
1269
1270
1271
1272
1273
1274
1275
1276
1277
1278
1279
1280
1281
1282
1283
1284
1285
1286
1287
1288
1289
1290
1291
1292
1293
1294
1295
1296
1297
1298
1299
1300
1301
1302
1303
1304
1305
1306
1307
1308
1309
1310
1311
1312
1313
1314
1315
1316
1317
1318
1319
1320
1321
1322
1323
1324
1325
1326
1327
1328
1329
1330
1331
1332
1333
1334
1335
1336
1337
1338
1339
1340
1341
1342
1343
1344
1345
1346
1347
1348
1349
1350
1351
1352
1353
1354
1355
1356
1357
1358
1359
1360
1361
1362
1363
1364
1365
1366
1367
1368
1369
1370
1371
1372
1373
1374
1375
1376
1377
1378
1379
1380
1381
1382
1383
1384
1385
1386
1387
1388
1389
1390
1391
1392
1393
1394
1395
1396
1397
1398
1399
1400
1401
1402
1403
1404
1405
1406
1407
1408
1409
1410
1411
1412
1413
1414
1415
1416
1417
1418
1419
1420
1421
1422
1423
1424
1425
1426
1427
1428
1429
1430
1431
1432
1433
1434
1435
1436
1437
1438
1439
1440
1441
1442
1443
1444
1445
1446
1447
1448
1449
1450
1451
1452
1453
1454
1455
1456
1457
1458
1459
1460
1461
1462
1463
1464
1465
1466
1467
1468
1469
1470
1471
1472
1473
1474
1475
1476
1477
1478
1479
1480
1481
1482
1483
1484
1485
1486
1487
1488
1489
1490
1491
1492
1493
1494
1495
1496
1497
1498
1499
1500
1501
1502
1503
1504
1505
1506
1507
1508
1509
1510
1511
1512
1513
1514
1515
1516
1517
1518
1519
1520
1521
1522
1523
1524
1525
1526
1527
1528
1529
1530
1531
1532
1533
1534
1535
1536
1537
1538
1539
1540
1541
1542
1543
1544
1545
1546
1547
1548
1549
1550
1551
1552
1553
1554
1555
1556
1557
1558
1559
1560
1561
1562
1563
1564
1565
1566
1567
1568
1569
1570
1571
1572
1573
1574
1575
1576
1577
1578
1579
1580
1581
1582
1583
1584
1585
1586
1587
1588
1589
1590
1591
1592
1593
1594
1595
1596
1597
1598
1599
1600
1601
1602
1603
1604
1605
1606
1607
1608
1609
1610
1611
1612
1613
1614
1615
1616
1617
1618
1619
1620
1621
1622
1623
1624
1625
1626
1627
1628
1629
1630
1631
1632
1633
1634
1635
1636
1637
1638
1639
1640
1641
1642
1643
1644
1645
1646
1647
1648
1649
1650
1651
1652
1653
1654
1655
1656
1657
1658
1659
1660
1661
1662
1663
1664
1665
1666
1667
1668
1669
1670
1671
1672
1673
1674
1675
1676
1677
1678
1679
1680
1681
1682
1683
1684
1685
1686
1687
1688
1689
1690
1691
1692
1693
1694
1695
1696
1697
1698
1699
1700
1701
1702
1703
1704
1705
1706
1707
1708
1709
1710
1711
1712
1713
1714
1715
1716
1717
1718
1719
1720
1721
1722
1723
1724
1725
1726
1727
1728
1729
1730
1731
1732
1733
1734
1735
1736
1737
1738
1739
1740
1741
1742
1743
1744
1745
1746
1747
1748
1749
1750
1751
1752
1753
1754
1755
1756
1757
1758
1759
1760
1761
1762
1763
1764
1765
1766
1767
1768
1769
1770
1771
1772
1773
1774
1775
1776
1777
1778
1779
1780
1781
1782
1783
1784
1785
1786
1787
1788
1789
1790
1791
1792
1793
1794
1795
1796
1797
1798
1799
1800
1801
1802
1803
1804
1805
1806
1807
1808
1809
1810
1811
1812
1813
1814
1815
1816
1817
1818
1819
1820
1821
1822
1823
1824
1825
1826
1827
1828
1829
1830
1831
1832
1833
1834
1835
1836
1837
1838
1839
1840
1841
1842
1843
1844
1845
1846
1847
1848
1849
1850
1851
1852
1853
1854
1855
1856
1857
1858
1859
1860
1861
1862
1863
1864
1865
1866
1867
1868
1869
1870
1871
1872
1873
1874
1875
1876
1877
1878
1879
1880
1881
1882
1883
1884
1885
1886
1887
1888
1889
1890
1891
1892
1893
1894
1895
1896
1897
1898
1899
1900
1901
1902
1903
1904
1905
1906
1907
1908
1909
1910
1911
1912
1913
1914
1915
1916
1917
1918
1919
1920
1921
1922
1923
1924
1925
1926
1927
1928
1929
1930
1931
1932
1933
1934
1935
1936
1937
1938
1939
1940
1941
1942
1943
1944
1945
1946
1947
1948
1949
1950
1951
1952
1953
1954
1955
1956
1957
1958
1959
1960
1961
1962
1963
1964
1965
1966
1967
1968
1969
1970
1971
1972
1973
1974
1975
1976
1977
1978
1979
1980
1981
1982
1983
1984
1985
1986
1987
1988
1989
1990
1991
1992
1993
1994
1995
1996
1997
1998
1999
2000
2001
2002
2003
2004
2005
2006
2007
2008
2009
2010
2011
2012
2013
2014
2015
2016
2017
2018
2019
2020
2021
2022
2023
2024
2025
2026
2027
2028
2029
2030
2031
2032
2033
2034
2035
2036
2037
2038
2039
2040
2041
2042
2043
2044
2045
2046
2047
2048
2049
2050
2051
2052
2053
2054
2055
2056
2057
2058
2059
2060
2061
2062
2063
2064
2065
2066
2067
2068
2069
2070
2071
2072
2073
2074
2075
2076
2077
2078
2079
2080
2081
2082
2083
2084
2085
2086
2087
2088
2089
2090
2091
2092
2093
2094
2095
2096
2097
2098
2099
2100
2101
2102
2103
2104
2105
2106
2107
2108
2109
2110
2111
2112
2113
2114
2115
2116
2117
2118
2119
2120
2121
2122
2123
2124
2125
2126
2127
2128
2129
2130
2131
2132
2133
2134
2135
2136
2137
2138
2139
2140
2141
2142
2143
2144
2145
2146
2147
2148
2149
2150
2151
2152
2153
2154
2155
2156
2157
2158
2159
2160
2161
2162
2163
2164
2165
2166
2167
2168
2169
2170
2171
2172
2173
2174
2175
2176
2177
2178
2179
2180
2181
2182
2183
2184
2185
2186
2187
2188
2189
2190
2191
2192
2193
2194
2195
2196
2197
2198
2199
2200
2201
2202
2203
2204
2205
2206
2207
2208
2209
2210
2211
2212
2213
2214
2215
2216
2217
2218
2219
2220
2221
2222
2223
2224
2225
2226
2227
2228
2229
2230
2231
2232
2233
2234
2235
2236
2237
2238
2239
2240
2241
2242
2243
2244
2245
2246
2247
2248
2249
2250
2251
2252
2253
2254
2255
2256
2257
2258
2259
2260
2261
2262
2263
2264
2265
2266
2267
2268
2269
2270
2271
2272
2273
2274
2275
2276
2277
2278
2279
2280
2281
2282
2283
2284
2285
2286
2287
2288
2289
2290
2291
2292
2293
2294
2295
2296
2297
2298
2299
2300
2301
2302
2303
2304
2305
2306
2307
2308
2309
2310
2311
2312
2313
2314
2315
2316
2317
2318
2319
2320
2321
2322
2323
2324
2325
2326
2327
2328
2329
2330
2331
2332
2333
2334
2335
2336
2337
2338
2339
2340
2341
2342
2343
2344
2345
2346
2347
2348
2349
2350
2351
2352
2353
2354
2355
2356
2357
2358
2359
2360
2361
2362
2363
2364
2365
2366
2367
2368
2369
2370
2371
2372
2373
2374
2375
2376
2377
2378
2379
2380
2381
2382
2383
2384
2385
2386
2387
2388
2389
2390
2391
2392
2393
2394
2395
2396
2397
2398
2399
2400
2401
2402
2403
2404
2405
2406
2407
2408
2409
2410
2411
2412
2413
2414
2415
2416
2417
2418
2419
2420
2421
2422
2423
2424
2425
2426
2427
2428
2429
2430
2431
2432
2433
2434
2435
2436
2437
2438
2439
2440
2441
2442
2443
2444
2445
2446
2447
2448
2449
2450
2451
2452
2453
2454
2455
2456
2457
2458
2459
2460
2461
2462
2463
2464
2465
2466
2467
2468
2469
2470
2471
2472
2473
2474
2475
2476
2477
2478
2479
2480
2481
2482
2483
2484
2485
2486
2487
2488
2489
2490
2491
2492
2493
2494
2495
2496
2497
2498
2499
2500

- socat (in order to connect to the CLI, though certain forks of netcat can also do it to some extents);
- halog from the latest HAProxy version : this is the log analysis tool, it parses native TCP and HTTP logs extremely fast (1 to 2 GB per second) and extracts useful information and statistics such as requests per URL, per source address, URLs sorted by response time or error rate, termination codes etc... It was designed to be deployed on the production servers to help troubleshoot live issues so it has to be there ready to be used;
- tcpdump : this is highly recommended to take the network traces needed to troubleshoot an issue that was made visible in the logs. There is a moment where application and haproxy's analysis will diverge and the network traces are the only way to say who's right and who's wrong. It's also fairly common to detect bugs in network stacks and hypervisors thanks to tcpdump;
- strace : it is tcpdump's companion. It will report what HAProxy really sees and will help sort out the issues the operating system is responsible for from the ones HAProxy is responsible for. Strace is often requested when a bug in HAProxy is suspected;

3.4.2. Advanced features : System-specific capabilities

Depending on the operating system HAProxy is deployed on, certain extra features may be available or needed. While it is supported on a number of platforms, HAProxy is primarily developed on Linux, which explains why some features are only available on this platform.

The transparent bind and connect features, the support for binding connections to a specific network interface, as well as the ability to bind multiple processes to the same IP address and ports are only available on Linux and BSD systems, though only Linux performs a kernel-side load balancing of the incoming requests between the available processes.

On Linux, there are also a number of extra features and optimizations including support for network namespaces (also known as "containers") allowing HAProxy to be a gateway between all containers, the ability to set the MSS, Netfilter marks and IP TOS field on the client side connection, support for TCP FastOpen on the listening side, TCP user timeouts to let the kernel quickly kill connections when it detects the client has disappeared before the configured timeouts, TCP splicing to let the kernel forward data between the two sides of a connections thus avoiding multiple memory copies, the ability to enable the "defer-accept" bind option to only get notified of an incoming connection once data become available in the kernel buffers, and the ability to send the request with the ACK confirming a connect (sometimes called "biggy-back") which is enabled with the "tcp-smart-connect" option. On Linux, HAProxy also takes great care of manipulating the TCP delayed ACKs to save as many packets as possible on the network.

Some systems have an unreliable clock which jumps back and forth in the past and in the future. This used to happen with some NUMA systems where multiple processors didn't see the exact same time of day, and recently it became more common in virtualized environments where the virtual clock has no relation with the real clock, resulting in huge time jumps (sometimes up to 30 seconds have been observed). This causes a lot of trouble with respect to timeout enforcement in general. Due to this flaw of these systems, HAProxy maintains its own monotonic clock which is based on the system's clock but where drift is measured and compensated for. This ensures that even with a very bad system clock, timers remain reasonably accurate and timeouts continue to work. Note that this problem affects all the software running on such systems and is not specific to HAProxy. The common effects are spurious timeouts or application freezes. Thus if this behaviour is detected on a system, it must be fixed, regardless of the fact that

HAProxy protects itself against it.

3.4.3. Advanced features : Scripting

HAProxy can be built with support for the Lua embedded language, which opens a wide area of new possibilities related to complex manipulation of requests or responses, routing decisions, statistics processing and so on. Using Lua it is even possible to establish parallel connections to other servers to exchange information. This way it becomes possible (though complex) to develop an authentication system for example. Please refer to the documentation in the file "doc/lua-api/index.rst" for more information on how to use Lua.

3.5. Sizing

Typical CPU usage figures show 15% of the processing time spent in HAProxy versus 85% in the kernel in TCP or HTTP close mode, and about 30% for HAProxy versus 70% for the kernel in HTTP keep-alive mode. This means that the operating system and its tuning have a strong impact on the global performance.

Usages vary a lot between users, some focus on bandwidth, other ones on request rate, others on connection concurrency, others on SSL performance. This section aims at providing a few elements to help in this task.

It is important to keep in mind that every operation comes with a cost, so each individual operation adds its overhead on top of the other ones, which may be negligible in certain circumstances, and which may dominate in other cases.

When processing the requests from a connection, we can say that :

- forwarding data costs less than parsing request or response headers;
- parsing request or response headers cost less than establishing then closing a connection to a server;
- establishing an closing a connection costs less than a TLS resume operation;
- a TLS resume operation costs less than a full TLS handshake with a key computation;
- an idle connection costs less CPU than a connection whose buffers hold data;
- a TLS context costs even more memory than a connection with data;

So in practice, it is cheaper to process payload bytes than header bytes, thus it is easier to achieve high network bandwidth with large objects (few requests per volume unit) than with small objects (many requests per volume unit). This explains why maximum bandwidth is always measured with large objects, while request rate or connection rates are measured with small objects.

Some operations scale well on multiple process spread over multiple processors, and others don't scale as well. Network bandwidth doesn't scale very far because the CPU is rarely the bottleneck for large objects, it's mostly the network bandwidth and data busses to reach the network interfaces. The connection rate doesn't scale well over multiple processors due to a few locks in the system when dealing with the local ports table. The request rate over persistent connections scales very well as it doesn't involve much memory nor network bandwidth and doesn't require to access locked structures. TLS key computation scales very well as it's totally CPU-bound. TLS resume scales moderately well, but reaches its limits around 4 processes where the overhead of accessing the shared table offsets the small gains expected from more power.

The performance numbers one can expect from a very well tuned system are in the following range. It is important to take them as orders of magnitude and to expect significant variations in any direction based on the processor, IRQ setting, memory type, network interface type, operating system tuning and so on.

The following numbers were found on a Core i7 running at 3.7 GHz equipped with a dual-port 10 Gbps NICs running Linux kernel 3.10, HAProxy 1.6 and OpenSSL 1.0.2. HAProxy was running as a single process on a single dedicated CPU core, and two extra cores were dedicated to network interrupts :

- 20 Gbps of maximum network bandwidth in clear text for objects 256 kB or higher, 10 Gbps for 41kB or higher;

- 4.6 Gbps of TLS traffic using AES256-GCM cipher with large objects;

- 83000 TCP connections per second from client to server;

- 82000 HTTP connections per second from client to server;

- 97000 HTTP requests per second in server-close mode (keep-alive with the client, close with the server);

- 243000 HTTP requests per second in end-to-end keep-alive mode;

- 300000 filtered TCP connections per second (anti-DDoS)

- 160000 HTTPS requests per second in keep-alive mode over persistent TLS connections;

- 13100 HTTPS requests per second using TLS resumed connections;

- 1300 HTTPS connections per second using TLS connections renegotiated with RSA2048;

- 20000 concurrent saturated connections per GB of RAM, including the memory required for system buffers; it is possible to do better with careful tuning but this setting it easy to achieve.

- about 8000 concurrent TLS connections (client-side only) per GB of RAM, including the memory required for system buffers;

- about 5000 concurrent end-to-end TLS connections (both sides) per GB of RAM including the memory required for system buffers;

Thus a good rule of thumb to keep in mind is that the request rate is divided by 10 between TLS keep-alive and TLS resume, and between TLS resume and TLS renegotiation, while it's only divided by 3 between HTTP keep-alive and HTTP close. Another good rule of thumb is to remember that a high frequency core with AES instructions can do around 5 Gbps of AES-GCM per core.

Having more core rarely helps (except for TLS) and is even counter-productive due to the lower frequency. In general a small number of high frequency cores is better.

Another good rule of thumb is to consider that on the same server, HAProxy will be able to saturate :

- about 5-10 static file servers or caching proxies;

- about 100 anti-virus proxies;

- and about 100-1000 application servers depending on the technology in use.

3.6. How to get HAProxy

HAProxy is an opensource project covered by the GPLv2 license, meaning that everyone is allowed to redistribute it provided that access to the sources is also provided upon request, especially if any modifications were made.

HAProxy evolves as a main development branch called "master" or "mainline", from which new branches are derived once the code is considered stable. A lot of web sites run some development branches in production on a voluntarily basis, either to participate to the project or because they need a bleeding edge feature, and their feedback is highly valuable to fix bugs and judge the overall quality and stability of the version being developed.

The new branches that are created when the code is stable enough constitute a stable version and are generally maintained for several years, so that there is no emergency to migrate to a newer branch even when you're not on the latest. Once a stable branch is issued, it may only receive bug fixes, and very rarely minor feature updates when that makes users' life easier. All fixes that go into a stable branch necessarily come from the master branch. This guarantees that no fix will be lost after an upgrade. For this reason, if you fix a bug, please make the patch against the master branch, not the stable branch. You may even discover it was already fixed. This process also ensures that regressions in a stable branch are extremely rare, so there is never any excuse for not upgrading to the latest version in your current branch.

Branches are numbered with two digits delimited with a dot, such as "1.6". A complete version includes one or two sub-version numbers indicating the level of fix. For example, version 1.5.14 is the 14th fix release in branch 1.5 after version 1.5.0 was issued. It contains 126 fixes for individual bugs, 24 updates on the documentation, and 75 other backported patches, most of which were needed to fix the aforementioned 126 bugs. An existing feature may never be modified nor removed in a stable branch, in order to guarantee that upgrades within the same branch will always be harmless.

HAProxy is available from multiple sources, at different release rhythms :

- The official community web site : <http://www.haproxy.org/> : this site provides the sources of the latest development release, all stable releases, as well as nightly snapshots for each branch. The release cycle is not fast, several months between stable releases, or between development snapshots. Very old versions are still supported there. Everything is provided as sources only, so whatever comes from there needs to be rebuilt and/or repackaged;

- A number of operating systems such as Linux distributions and BSD ports. These systems generally provide long-term maintained versions which do not always contain all the fixes from the official ones, but which at least contain the critical fixes. It often is a good option for most users who do not seek advanced configurations and just want to keep updates easy;

- Commercial versions from <http://www.haproxy.com/> : these are supported professional packages built for various operating systems or provided as appliances, based on the latest stable versions and including a number of features backported from the next release for which there is a strong demand. It is the best option for users seeking the latest features with the reliability of a stable branch, the fastest response time to fix bugs, or simply support contracts on top of an opensource product;

In order to ensure that the version you're using is the latest one in your branch, you need to proceed this way :

- verify which HAProxy executable you're running : some systems ship it by default and administrators install their versions somewhere else on the system, so it is important to verify in the startup scripts which one is

1431 used;

1432

1433 - determine which source your HAProxy version comes from. For this, it's

1434 generally sufficient to type "haproxy -v". A development version will

1435 appear like this, with the "dev" word after the branch number :

1436

1437 HA-Proxy version 1.6-dev3-385ecc-68 2015/08/18

1438

1439 A stable version will appear like this, as well as unmodified stable

1440 versions provided by operating system vendors :

1441

1442 HA-Proxy version 1.5.14 2015/07/02

1443

1444 And a nightly snapshot of a stable version will appear like this with an

1445 hexadecimal sequence after the version, and with the date of the snapshot

1446 instead of the date of the release :

1447

1448 HA-Proxy version 1.5.14-e4766ba 2015/07/29

1449

1450 Any other format may indicate a system-specific package with its own

1451 patch set. For example HAProxy Enterprise versions will appear with the

1452 following format (<branch>-<latest commit>-<revision>) :

1453

1454 HA-Proxy version 1.5.0-994126-357 2015/07/02

1455

1456 - for system-specific packages, you have to check with your vendor's package

1457 repository or update system to ensure that your system is still supported,

1458 and that fixes are still provided for your branch. For community versions,

1459 coming from haproxy.org, just visit the site, verify the status of your

1460 branch and compare the latest version with yours to see if you're on the

1461 latest one. If not you can upgrade. If your branch is not maintained

1462 anymore, you're definitely very late and will have to consider an upgrade

1463 to a more recent branch (carefully read the README when doing so).

1464

1465 HAProxy will have to be updated according to the source it came from. Usually it

1466 follows the system vendor's way of upgrading a package. If it was taken from

1467 sources, please read the README file in the sources directory after extracting

1468 the sources and follow the instructions for your operating system.

1469

1470

1471 4. Companion products and alternatives

1472 -----

1473

1474 HAProxy integrates fairly well with certain products listed below, which is why

1475 they are mentioned here even if not directly related to HAProxy.

1476

1477

1478 4.1. Apache HTTP server

1479 -----

1480

1481 Apache is the de-facto standard HTTP server. It's a very complete and modular

1482 project supporting both file serving and dynamic contents. It can serve as a

1483 frontend for some application servers. In can even proxy requests and cache

1484 responses. In all of these use cases, a front load balancer is commonly needed.

1485 Apache can work in various modes, certain being heavier than other ones. Certain

1486 modules still require the heavier pre-forked model and will prevent Apache from

1487 scaling well with a high number of connections. In this case HAProxy can provide

1488 a tremendous help by enforcing the per-server connection limits to a safe value

1489 and will significantly speed up the server and preserve its resources that will

1490 be better used by the application.

1491

1492 Apache can extract the client's address from the X-Forwarded-For header by using

1493 the "mod_rpaf" extension. HAProxy will automatically feed this header when

1494 "option forwardfor" is specified in its configuration. HAProxy may also offer a

1495 nice protection to Apache when exposed to the internet, where it will better

1496 resist to a wide number of types of DoS.

1497

1498

1499 4.2. NGINX

1500 -----

1501

1502 NGINX is the second de-facto standard HTTP server. Just like Apache, it covers a

1503 wide range of features. NGINX is built on a similar model as HAProxy so it has

1504 no problem dealing with tens of thousands of concurrent connections. When used

1505 as a gateway to some applications (eg: using the included PHP FPM), it can often

1506 be beneficial to set up some frontend connection limiting to reduce the load

1507 on the PHP application. HAProxy will clearly be useful there both as a regular

1508 load balancer and as the traffic regulator to speed up PHP by decongestionning

1509 it. Also since both products use very little CPU thanks to their event-driven

1510 architecture, it's often easy to install both of them on the same system. NGINX

1511 implements HAProxy's PROXY protocol, thus it is easy for HAProxy to pass the

1512 client's connection information to NGINX so that the application gets all the

1513 relevant information. Some benchmarks have also shown that for large static

1514 file serving, implementing consistent hash on HAProxy in front of NGINX can be

1515 beneficial by optimizing the OS' cache hit ratio, which is basically multiplied

1516 by the number of server nodes.

1517

1518 4.3. Varnish

1519 -----

1520

1521 Varnish is a smart caching reverse-proxy, probably best described as a web

1522 application accelerator. Varnish doesn't implement SSL/TLS and wants to dedicate

1523 all of its CPU cycles to what it does best. Varnish also implements HAProxy's

1524 PROXY protocol so that HAProxy can very easily be deployed in front of Varnish

1525 as an SSL offloader as well as a load balancer and pass it all relevant client

1526 information. Also, Varnish naturally supports decompression from the cache when

1527 a server has provided a compressed object, but doesn't compress however. HAProxy

1528 can then be used to compress outgoing data when backend servers do not implement

1529 compression, though it's rarely a good idea to compress on the load balancer

1530 unless the traffic is low.

1531

1532 When building large caching farms across multiple nodes, HAProxy can make use of

1533 consistent URL hashing to intelligently distribute the load to the caching nodes

1534 and avoid cache duplication, resulting in a total cache size which is the sum of

1535 all caching nodes.

1536

1537

1538 4.4. Alternatives

1539 -----

1540

1541 Linux Virtual Server (LVS or IPVS) is the layer 4 load balancer included within

1542 the Linux kernel. It works at the packet level and handles TCP and UDP. In most

1543 cases it's more a complement than an alternative since it doesn't have layer 7

1544 knowledge at all.

1545

1546 Pound is another well-known load balancer. It's much simpler and has much less

1547 features than HAProxy but for many very basic setups both can be used. Its

1548 author has always focused on code auditability first and wants to maintain the

1549 set of features low. Its thread-based architecture scales less well with high

1550 connection counts, but it's a good product.

1551

1552 Pen is a quite light load balancer. It supports SSL, maintains persistence using

1553 a fixed-size table of its clients' IP addresses. It supports a packet-oriented

1554 mode allowing it to support direct server return and UDP to some extents. It is

1555 meant for small loads (the persistence table only has 2048 entries).

1556

1557 NGINX can do some load balancing to some extents, though it's clearly not its

1558 primary function. Production traffic is used to detect server failures, the

1559 load balancing algorithms are more limited, and the stickiness is very limited.

1560

1561 But it can make sense in some simple deployment scenarios where it is already
1562 present. The good thing is that since it integrates very well with HAProxy,
1563 there's nothing wrong with adding HAProxy later when its limits have been faced.
1564

1565 Varnish also does some load balancing of its backend servers and does support
1566 real health checks. It doesn't implement stickiness however, so just like with
1567 NGINX, as long as stickiness is not needed that can be enough to start with.
1568 And similarly, since HAProxy and Varnish integrate so well together, it's easy
1569 to add it later into the mix to complement the feature set.
1570