

Lua: Architecture and first steps
~~~~~  
version 1.0  
  
author: Thierry FOURNIER  
contact: tfournier at arpalert dot org

HAProxy is a powerful load balancer. It embeds many options and many configuration styles in order to give a solution to many load balancing problems. However, HAProxy is not universal and some special or specific problems doesn't have solution with the native software.

This text is not a full explanation of the Lua syntax.

This text is not a replacement of the HAProxy Lua API documentation. The API documentation can be found at the project root, in the documentation directory. The goal of this text is to discover how Lua is implemented in HAProxy and using it efficiently.

However, this can be read by Lua beginners. Some examples are detailed.

#### Why a scripting language in HAProxy

=====

HAProxy 1.5 makes at possible to do many things using samples, but some people wants to more combining results of samples fetches, programming conditions and loops which is not possible. Sometimes people implement these functionalities in patches which have no meaning outside their network. These people must maintain these patches, or worse we must integrate them in the HAProxy mainstream.

Their need is to have an embedded programming language in order to no longer modify the HAProxy source code, but to write their own control code. Lua is encountered very often in the software industry, and in some open source projects. It is easy to understand, efficient, light without external dependencies, and leaves the resource control to the implementation. Its design is close to the HAProxy philosophy which uses components for what they do perfectly.

The HAProxy control block allows one to take a decision based on the comparison between samples and patterns. The samples are extracted using fetch functions easily extensible, and are used by actions which are also extensible. It seems natural to allow Lua to give samples, modify them, and to be an action target. So, Lua uses the same entities as the configuration language. This is the most natural and reliable way fir the Lua integration. So, the Lua engine allow one to add new sample fetch functions, new converter functions and new actions. These new entities can access the existing samples fetches and converters allowing to extend them without rewriting them.

The writing of the first Lua functions shows that implementing complex concepts like protocol analysers is easy and can be extended to full services. It appears that these services are not easy to implement with the HAProxy configuration model which is base on four steps: fetch, convert, compare and action. HAProxy is extended with a notion of services which are a formalisation of the existing services like stats, cli and peers. The service is an autonomous entity with a behaviour pattern close to that of an external client or server. The Lua engine inherits from this new service and offers new possibilities for writing services.

This scripting language is useful for testing new features as proof of concept. Later, if there is general interest, the proof of concept could be integrated with C language in the HAProxy core.

The HAProxy Lua integration also provides also a simple way for distributing Lua packages. The final user needs only to install the Lua file, load it in HAProxy and follow the attached documentation.

#### Design and technical things

=====

Lua is integrated into the HAProxy event driven core. We want to preserve the fast processing of HAProxy. To ensure this, we implement some technical concepts between HAProxy and the Lua library.

The following paragraph also describes the interactions between Lua and HAProxy from a technical point of view.

#### Prerequisite

-----

Reading the following documentation links is required to understand the current paragraph:

HAProxy doc: <http://cbonte.github.io/haproxy-dconv/configuration-1.6.html>  
Lua API: <http://www.lua.org/manual/5.3/>  
HAProxy API: <http://www.arpalert.org/src/haproxy-lua-api/1.6/index.html>  
Lua guide: <http://www.lua.org/pil/>

#### more about Lua choice

-----

Lua language is very simple to extend. It is easy to add new functions written in C in the core language. It not require to embed very intrusive libraries, and we do not change compilation processes.

The amount of memory consumed can be controlled, and the issues due to lack of memory are perfectly caught. The maximum amount of memory allowed for the Lua processes is configurable. If some memory is missing, the current Lua action fails, and the HAProxy processing flow continues.

Lua provides a way for implementing event driven design. When the Lua code wants to do a blocking action, the action is started, it executes non blocking operations, and returns control to the HAProxy scheduler when it needs to wait for some external event.

The Lua process can be interrupted after a number of instructions executed. The Lua execution will resume later. This is a useful way for controlling the execution time. This system also keeps HAProxy responsive. When the Lua execution is interrupted, HAProxy accepts some connections or transfers pending data. The Lua execution does not block the main HAProxy processing, except in some cases which we will see later.

#### Lua function integration

-----

The Lua actions, sample fetches, converters and services are integrated in HAProxy with "register\_\*" functions. The register system is a choice for providing HAProxy Lua packages easily. The register system adds new sample fetches, converters, actions or services usable in the HAProxy configuration file.

The register system is defined in the "core" functions collection. This collection is provided by HAProxy and is always available. Below, the list of these functions:

- core.register\_action()
- core.register\_converters()
- core.register\_fetches()

```
131 - core.register_init()
132 - core.register_service()
133 - core.register_task()
134
```

135 These functions are the execution entry points.

136  
137 HTTP action must be used for manipulating HTTP request headers. This action  
138 can not manipulates HTTP content. It is dangerous to use the channel  
139 manipulation object with an HTTP request in an HTTP action. The channel  
140 manipulation can transform a valid request in an invalid request. In this case,  
141 the action will never resume and the processing will be frozen. HAProxy  
142 discards the request after the reception timeout.

```
143 Non blocking design
144 -----
145
```

146 HAProxy is an event driven software, so blocking system calls are absolutely  
147 forbidden. However, the Lua allows to do blocking actions. When an action  
148 blocks, HAProxy is waiting and do nothing, so the basic functionalities like  
149 accepting connections or forwarding data are blocked while the end of the system  
150 call. In this case HAProxy will be less responsive.

151  
152 This is very insidious because when the developer tries to execute its Lua code  
153 with only one stream, HAProxy seems to run fine. When the code is used with  
154 production stream, HAProxy encounters some slow processing, and it cannot  
155 hold the load.

156  
157 However, during the initialisation state, you can obviously using blocking  
158 functions. There are typically used for loading files.

159  
160 The list of prohibited standard Lua functions during the runtime contains all  
161 that do filesystem access:

```
162 - os.remove()
163 - os.rename()
164 - os.tmpname()
165 - package.*()
166 - io.*()
167 - file.*()
168
```

169 Some other functions are prohibited:

```
170 - os.execute(), waits for the end of the required execution blocking HAProxy.
171
172 - os.exit(), is not really dangerous for the process, but its not the good way
173 for exiting the HAProxy process.
174
```

```
175 - print(), writes data on stdout. In some cases these writes are blocking, the
176 best practice is reserving this call for debugging. We must prefer
177 to use core.log() or TXN.log() for sending messages.
178
```

179  
180 Some HAProxy functions have a blocking behaviour pattern in the Lua code, but  
181 there are compatible with the non blocking design. These functions are:

```
182 - All the socket class
183 - core.sleep()
184
```

```
185 Responsive design
186 -----
187
```

188 HAProxy must process connexions accept, forwarding data and processing timeouts  
189 as soon as possible. The first thing is to believe that a Lua script with a long  
190 execution time should impact the expected responsive behaviour.

191 It is not the case, the Lua script execution are regularly interrupted, and

196 HAProxy can process other things. These interruptions are exprimed in number of  
197 Lua instructions. The number of interruptions between two interrupts is  
198 configured with the following "tune" option:  
199

```
200 tune.lua.forced-yield <nb>
201
```

202 The default value is 10 000. For determining it, I ran benchmark on my laptop.  
203 I executed a Lua loop between 10 seconds with differents values for the  
204 "tune.lua.forced-yield" option, and I noted the results:

| configured<br>instructions<br>between two<br>forced yields | Number of<br>loops executed<br>in millions |
|------------------------------------------------------------|--------------------------------------------|
| 10                                                         | 160                                        |
| 500                                                        | 670                                        |
| 1000                                                       | 680                                        |
| 5000                                                       | 700                                        |
| 7000                                                       | 700                                        |
| 8000                                                       | 700                                        |
| 9000                                                       | 710 <- ceil                                |
| 10000                                                      | 710                                        |
| 100000                                                     | 710                                        |
| 1000000                                                    | 710                                        |

210  
211  
212  
213  
214  
215  
216  
217  
218  
219  
220  
221  
222 The result showed that from 9000 instructions between two interrupt, we reached  
223 a ceil, so the default parameter is 10 000.

224 When HAProxy interrupts the Lua processing, we have two states possible:

- Lua is resumable, and it returns control to the HAProxy scheduler,
- Lua is not resumable, and we just check the execution timeout.

225  
226 The second case occurs if it is required by the HAProxy core. This state is  
227 forced if the Lua is processed in a non resumable HAProxy part, like sample  
228 fetches or converters.

229  
230 It occurs also if the Lua is non resumable. For example, if some code is  
231 executed through the Lua pcall() function, the execution is not resumable. This  
232 is explained later.

233  
234 So, the Lua code must be fast and simple when is executed as sample fetches and  
235 converters, it could be slow and complex when is executed as actions and  
236 services.

```
237 Execution time
238 -----
239
```

240 The Lua execution time is measured and limited. Each group of functions have its  
241 own timeout configured. The time measured is the real Lua execution time, and  
242 not the difference between the end time and the start time. The groups are:

- main code and init are not submitted to the timeout,
- fetches, converters and action have a default timeout of 4s,
- task, by default does not have timeout,
- service have a default timeout of 4s.

243 The corresponding tune option are:

- tune.lua.session-timeout (fetches, converters and action)
- tune.lua.task-timeout (task)
- tune.lua.service-timeout (services)

244 The tasks does not have a timeout because it runs in background along the

## 261 HAProxy process life.

262 For example, if an Lua script is executed during 1,1s and the script executes a  
 263 sleep of 1 second, the effective measured running time is 0,1s.

264  
 265 This timeout is useful for preventing infinite loops. During the runtime, it  
 266 should never triggered.

## 267 The stack and the coprocess

268 -----

269 The Lua execution is organized around a stack. Each Lua action, even out of the  
 270 effective execution, affects the stack. HAProxy integration uses one main stack,  
 271 which is common for all the process, and a secondary one used as coprocess.  
 272 After the initialization, the main stack is no longer used by HAProxy, except  
 273 for global storage. The second type of stack is used by all the Lua functions  
 274 called from different Lua actions declared in HAProxy. The main stack permits  
 275 to store coroutines pointers, and some global variables.

276 Do you want to see an example of how seems Lua C development around a stack ?  
 277 Some examples follows. This first one, is a simple addition:

```
278 lua_pushnumber(L, 1)
279 lua_pushnumber(L, 2)
280 lua_arith(L, LUA_OPADD)
```

281 Its easy, we push 1 on the stack, after, we push 2, and finally, we perform an  
 282 addition. The two top entries of the stack are added, popped, and the result is  
 283 pushed. It is a classic way with a stack.

284 Now an example for constructing array and objects. Its little bit more  
 285 complicated. The difficult consist to keep in mind the state of the stack while  
 286 we write the code. The goal is to create the entity described below. Note that  
 287 the notation "\*1" is a metatable reference. The metatable will be explained  
 288 later.

```
289 name*1 = {
290   [0] = <userdata>,
291 }
292
293 *1 = {
294   "__index" = {
295     "method1" = <function>,
296     "method2" = <function>
297   }
298   "__gc" = <function>
299 }
```

300 Let's go:

```
301 lua_newtable()           // The "name" table
302 lua_newtable()           // The metatable *1
303 lua_pushstring("__index")
304 lua_newtable()           // The "__index" table
305 lua_pushstring("method1")
306 lua_pushfunction(function)
307 lua_settable(-3)
308 lua_pushstring("method2")
309 lua_pushfunction(function)
310 lua_settable(-3)
311 lua_settable(-3)
312 lua_pushstring("__gc")
313 lua_pushfunction(function)
314 lua_settable(-1)
315 // insert "gc"
316 // attach metatable to "name"
```

```
326 lua_pushnumber(0)
327 lua_pushuserdata(userdata)
328 lua_settable(-3)
329 lua_setglobal("name")
330
```

331 So, coding for Lua in C, is not complex, but it needs some mental gymnastic.

332 The object concept and the HAProxy format

333 -----

334 The objects seems to not be a native concept. An Lua object is a table. We can  
 335 note that the table notation accept three forms:

- 336 1. mytable["entry"](mytable, "param")
- 337 2. mytable.entry(mytable, "param")
- 338 3. mytable:entry("param")

339 These three notation have the same behaviour pattern: a function is executed  
 340 with the itself table as first parameter and string "param" as second parameter  
 341 The notation with [] is commonly used for storing data in a hash table, and the  
 342 dotted notation is used for objects. The notation with ":" indicates that the  
 343 first parameter is the element at the left of the symbol ":".

344 So, an object is a table and each entry of the table is a variable. A variable  
 345 can be a function. These are the first concepts of the object notation in the  
 346 Lua, but it is not the end.

347 With the objects, we usually expect classes and inheritance. This is the role of  
 348 the metable. A metable is a table with predefined entries. These entries modify  
 349 the default behaviour of the table. The simplest example is the "\_\_index" entry.  
 350 If this entry exists, it is called when a value is requested in the table. The  
 351 behaviour is the following:

- 352 1 - looks in the table if the entry exists, and if it the case, return it
- 353 2 - looks if a metatable exists, and if the "\_\_index" entry exists
- 354 3 - if "\_\_index" is a function, execute it with the key as parameter, and  
 355 returns the result of the function.
- 356 4 - if "\_\_index" is a table, looks if the requested entry exists, and if  
 357 exists, return it.
- 358 5 - if not exists, return to step 2

359 The behaviour of the point 5 represents the inheritance.

360 In HAProxy all the provided objects are tables, the entry "[0]" contains private  
 361 data, there are often userdata or userdata. The metatable is registered in  
 362 the global part of the main Lua stack, and it is called with the case sensitive  
 363 class name. A great part of these class must not be used directly because it  
 364 requires an initialisation using the HAProxy internal structs.

365 The HAProxy objects uses unified conventions. An Lua object is always a table.  
 366 In most cases, an HAProxy Lua object need some private data. These are always  
 367 set in the index [0] of the array. The metatable entry "\_\_tostring" returns the  
 368 object name.

369 The Lua developer can add entries to the HAProxy object. He just works carefully  
 370 and prevent to modify the index [0].

371 Common HAProxy objects are:

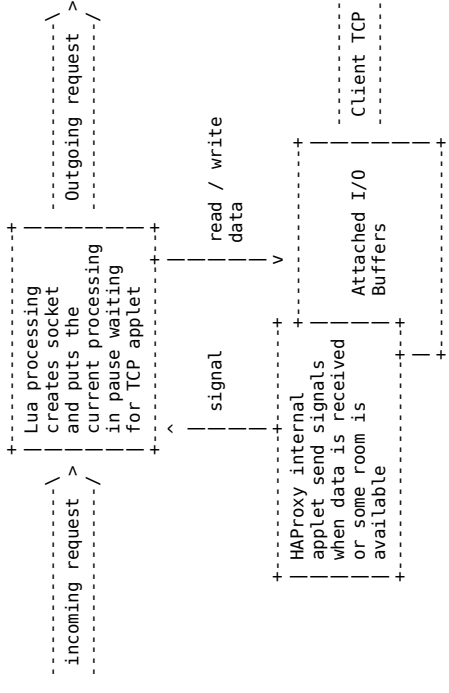
- 372 - TXN : manipulates the transaction between the client and the server
- 373 - Channel : manipulates proxified data between the client and the server

```
391 - HTTP      : manipulates HTTP between the client and the server
392 - Map       : manipulates HAProxy maps.
393 - Fetches   : access to all HAProxy sample fetches
394 - Converters : access to all HAProxy sample converters
395 - AppletTCP : process client request like a TCP server
396 - AppletHTTP : process client request like an HTTP server
397 - Socket    : establish tcp connection to a server (ipv4/ipv6/socket/ssl/...)
398
399 The garbage collector and the memory allocation
400 -----
401
402 Lua doesn't really have a global memory limit, but HAProxy implements it. This
403 permits to control the amount of memory dedicated to the Lua processes. It is
404 specially useful with embedded environments.
405
406 When the memory limit is reached, HAProxy refuses to give more memory to the Lua
407 scripts. The current Lua execution is terminated with an error and HAProxy
408 continue its processing.
409
410 The max amount of memory is configured with the option:
411
412 tune.lua.maxmem
413
414 As many other script languages, Lua uses a garbage collector for reusing its
415 memory. The Lua developer can work without memory preoccupation. Usually, the
416 garbage collector is controlled by the Lua core, but sometimes it will be useful
417 to run when the user/developer requires. So the garbage collector can be called
418 from C part or Lua part.
419
420 Sometimes, objects using userdata or userdata requires to free some memory
421 block or close filedescriptor not controlled by the Lua. A dedicated garbage
422 collection function is provided through the metatable. It is referenced with the
423 special entry " __gc".
424
425 Generally, in HAProxy, the garbage collector does this job without any
426 intervention. However some object uses a great amount of memory, and we want to
427 release as quick as possible. The problem is that only the GC knows if the object
428 is in use or not. The reason is simple variable containing objects can be shared
429 between coroutines and the main thread, so an object can used everywhere in
430 HAProxy.
431
432 The only one example is the HAProxy sockets. These are explained later, just for
433 understanding the GC issues, a quick overview of the socket follows. The HAProxy
434 socket uses an internal session and stream, these sessions uses resources like
435 memory and file descriptor and in some cases keeps a socket open while it is no
436 longer used by Lua.
437
438 If the HAProxy socket is used, we forcing a garbage collector cycle after the
439 end of each function using HAProxy socket. The reason is simple: if the socket
440 is no longer used, we want to close the connection quickly.
441
442 A special flag is used in HAProxy indicating that a HAProxy socket is created.
443 If this flag is set, a full GC cycle is started after each Lua action. This is
444 not free, we loose about 10% of performances, but it is the only way for closing
445 sockets quickly.
446
447 The yield concept / longjmp issues
448 -----
449
450 The "yield" is an action which does some Lua processing in pause and give back
451 the hand to the HAProxy core. This action is do when the Lua needs to wait about
452 data or other things. The most basically example is the sleep() function. In a
453 event driven software the code must not process blocking systems call, so the
454 sleep blocks the software between a lot of time. In HAProxy, an Lua sleep does a
455 yield, and ask to the scheduler to be waked up in a required sleep time.
```

```
456 Meanwhile, the HAProxy scheduler dos other things, like accepting new connection
457 or forwarding data.
458
459 A yield is also executed regularly, after a lot of Lua instruction processed.
460 This yield permits to control the effective execution time, and also give back
461 the hand to the haproxy core. When HAProxy finish to process the pending jobs,
462 the Lua execution continue.
463
464 This special "yield" uses the Lua "debug" functions. Lua provides a debug method
465 called "lua_sethook()" which permits to interrupt the execution after some
466 configured condition and call a function. This condition used in HAProxy is
467 a number of instruction processed and when a function returns. The function
468 called controls the effective execution time, and if it is possible send a
469 "yield".
470
471 The yield system is based on a couple setjmp/longjmp. In brief, the setjmp()
472 stores a stack state, and the longjmp restores the stack in its state which had
473 before the last Lua execution.
474
475 Lua can immediately stop is execution if an error occurs. This system uses also
476 the longjmp system. In HAProxy, we try to use this sytem only for unrecoverable
477 errors. Maybe some trivial errors targets an exception, but we try to remove it.
478
479 It seems that Lua uses the longjmp system for having a behaviour like the java
480 try / catch. We can use the function pcall() to executes some code. The function
481 pcall() run a setjmp(). So, if any error occurs while the Lua code execution,
482 the flow immediately return from the pcall() with an error.
483
484 The big issue of this behaviour is that we cannot do a yield. So if some Lua code
485 executes a library using pcall for catching errors, HAProxy must be wait for the
486 end of execution without processing any accept or any stream. The cause is the
487 yield must be jump to the root of execution. The intermediate setjmp() avoid
488 this behaviour.
489
490 HAProxy start Lua execution
491 + Lua puts a setjmp()
492 + Lua executes code
493 + Some code is executed in a pcall()
494 + pcall() puts a setjmp()
495 + Lua executes code
496 + A yield is require for a sleep function
497 it cannot be jumps to the Lua root execution.
498
499 Another issue with the processing of strong errors is the manipulation of the
500 Lua stack outside of an Lua processing. If one of the functions called occurs a
501 strong error, the default behaviour is an abort(). It is not acceptable when
502 HAProxy is in runtime mode. The Lua documentation propose to use another
503 setjmp/longjmp to avoid the abort(). The goal is to puts a setjmp between
504 manipulating the Lua stack and using an alternative "panic" function which jumps
505 to the setjmp() in error case.
506
507 All of these behaviours are very dangerous for the stability, and the internal
508 HAProxy code must be modified with many precautions.
509
510 For preserving a good behaviour of HAProxy, the yield is mandatory.
511 Unfortunately, some HAProxy part are not adapted for resuming an execution after
512 a yield. These part are the sample fetches and the sample converters. So, the
513 Lua code written in these parts of HAProxy must be quickly executed, and can not
514 do actions which require yield like TCP connection or simple sleep.
515
516 HAProxy socket object
517 -----
518
519 -----
520
```

The HAProxy design is optimized for the data transfers between a client and a server, and processing the many errors which can occurs during these exchanges. HAProxy is not designed for having a third connection established to a third party server.

The solution consist to puts the main stream in pause waiting for the end of the exchanges with the third connection. This is completed by a signal between internal tasks. The following graph shows the HAProxy Lua socket:



A more detailed graph is available in the "doc/internals" directory.

The HAProxy Lua socket uses a full HAProxy session / stream for establishing the connection. This mechanism provides all the facilities and HAProxy features, like the SSL stack, many socket type, and support for namespaces. Technically it support the proxy protocol, but there are no way to enable it.

How compiling HAProxy with Lua

=====

HAProxy 1.6 requires Lua 5.3. Lua 5.3 offers some features which makes easy the integration. Lua 5.3 is young, and some distros do not distribute it. Luckily, Lua is a great product because it does not require exotic dependencies, and its build process is really easy.

The compilation process for linux is easy:

```
- download the source tarball
wget http://www.lua.org/ftp/luatxt-5.3.1.1.tar.gz
```

```
- untar it
tar xf luatxt-5.3.1.1.tar.gz
```

```
- enter the directory
cd luatxt-5.3.1
```

```
- build the library for linux
make linux
```

```
- install it:
sudo make INSTALL_TOP=/opt/luatxt-5.3.1
```

HAProxy builds with your favourite options, plus the following options for embedding the Lua script language:

```
- download the source tarball
wget http://www.haproxy.org/download/1.6/src/haproxy-1.6.2.tar.gz
```

```
- untar it
tar xf haproxy-1.6.2.tar.gz
```

```
- enter the directory
cd haproxy-1.6.2
```

```
- build HAProxy:
make TARGET=linux \
  USE_DL=1 \
  USE_LUA=1 \
  LUA_LIB=/opt/luatxt-5.3.1/lib \
  LUA_INC=/opt/luatxt-5.3.1/include
```

```
- install it:
sudo make PREFIX=/opt/haproxy-1.6.2 install
```

```
First steps with Lua
=====
```

Now, its time to using Lua in HAProxy.

```
Start point
-----
```

The HAProxy global directive "lua-load <file>" allow to load an lua file. This is the entry point. This load become during the configuration parsing, and the Lua file is immediately executed.

All the register\_\*() function must be called at this time because there are used just after the processing of the global section, in the frontend/backend/listen sections.

The most simple "Hello world !" is the following line a loaded Lua file:

```
core.Alert("Hello World !");
```

It display a log during the HAProxy startup:

```
[alert] 285/083533 (14465) : Hello World !
```

Default path and libraries

```
-----
```

Lua can embed some libraries. These libraries can be included from different paths. It seems that Lua doesn't like subdirectories. In the following example, I try to load a compiled library, so the first line is Lua code, the second line is an 'strace' extract proving that the library was opened. The next lines are the associated error.

```
require("luac/concat")
```

```
open("./luac/concat.so", O_RDONLY|O_CLOEXEC) = 4
```

```
[ALERT] 293/175822 (22806) : parsing [commonstats.conf:15] : lua runtime
error: error loading module 'luac/concat' from file './luac/concat.so':
        ./luac/concat.so: undefined symbol: luaopen_luac/concat
```

Lua tries to load the C symbol 'luaopen\_luac/concat'. When Lua tries to open a

```

651 library, it tries to execute the function associated to the symbol
652 "luaopen_<libname>".
653
654 The variable "<libname>" is defined using the content of the variable
655 "package.cpath" and/or "package.path". The default definition of the
656 "package.cpath" (on my computer is ) variable is:
657
658     /usr/local/lib/lua/5.3/?so:/usr/local/lib/lua/5.3/loadall.so:./?.so
659
660 The "<libname>" is the content which replaces the symbol "<?>". In th previous
661 example, its "luac/concat", and obviously the Lua core try to load the function
662 associated with the symbol "luaopen_luac/concat".
663
664 My conclusion is that Lua doesn't support subdirectories. So, for loading
665 libraries in subdirectory, it must fill the variable with the name of this
666 subdirectory. The extension .so must disappear, otherwise Lua try to execute the
667 function associated with the symbol "luaopen_concat.so". The following syntax is
668 correct:
669
670     package.cpath = package.cpath .. ";/luac/?so"
671     require("concat")
672
673 First useful example
674 -----
675
676     core.register_fetches("my-hash", function(txn, salt)
677         return txn.sc:sdbm(salt .. txn.sf:req_fhdr("host") .. txn.sf:path() .. txn.sf:src..
678         end)
679
680 You will see that these 3 line can generate a lot of explanations :)
681
682 Core.register_fetches() is executed during the processing of the global section
683 by the HAProxy configuration parser. A new sample fetch is declared with name
684 "my-hash", this name is always prefixed by "lua.". So this new declared
685 sample fetch will be used calling "lua.my-hash" in the HAProxy configuration
686 file.
687
688 The second parameter is an inline declared anonymous function. Note the closed
689 parenthesis after the keyword "end" which end the function. The first parameter
690 of these anonymous function is "txn". It an object of class TXN. It provides
691 access functions. The second parameter is an arbitrary value provided by the
692 HAProxy configuration file. This parameter is optional, the developer must
693 check if its present.
694
695 The anonymous function registration is executed when the HAProxy backend or
696 frontend configuration references the sample fetch "lua.my-hash".
697
698 This example can writed with an other style, like below:
699
700     function my_hash(txn, salt)
701         return txn.sc:sdbm(salt .. txn.sf:req_fhdr("host") .. txn.sf:path() .. txn.sf:src..
702         end
703
704     core.register_fetches("my-hash", my_hash)
705
706 This second form is clearer, but the first one is compact.
707
708 The operator "..." is a string concatenation. If one of the two operands are not a
709 string, an error occurs and the execution is immediately stopped. This is
710 important to keep in mind for the following things.
711
712 Now I write the example on more than one line. Its an easiest way for commenting
713 the code:
714
715     1. function my_hash(txn, salt)

```

```

716     2. local str = ""
717     3. str = str .. salt
718     4. str = str .. txn.sf:req_fhdr("host")
719     5. str = str .. txn.sf:path()
720     6. str = str .. txn.sf:src()
721     7. local result = txn.sc:sdbm(str, 1)
722     8. return result
723     9. end
724    10.
725    11. core.register_fetches("my-hash", my_hash)
726
727 local
728 ~~~~
729
730 The first keyword is "local". This is a really important keyword. You must
731 understand that the function "my_hash" will be called for each HAProxy request
732 using the declared sample fetch. So, this function can be executed many times in
733 parallel.
734
735 By default, Lua uses global variables. so in this example, il the variable "str"
736 is declared without the keyword "local", it will be shared by all the parallel
737 executions of the function and obviously, the content of the requests will be
738 shared.
739
740 This warning is very important. I tried to write useful Lua code like a rewrite
741 of the statistics page, and its very hard to thing to declare each variable as
742 "local".
743
744 I guess than this behaviour will be the cause of many trouble on the mailing
745 list.
746
747 str = str ..
748 ~~~~~
749
750 Now a parenthesis about the form "str = str ..". This form allow to do string
751 concatenations. Remember that Lua uses a garbage collector, so what happens when
752 we do "str = str .. 'another string'" ?
753
754     str = str .. "another string"
755     ^   ^   ^   ^
756     1   2   3   4
757
758 Lua execute first the concatenation operator (3), it allocates memory for the
759 resulting string and fill this memory with the concatenation of the operands 2
760 and 4. Next, it free the variable 1, now the old content of 1 can be garbage
761 collected. and finally, the new content of 1 is the concatenation.
762
763 what the matter ? when we do this operation many times, we consume a lot of
764 memory, and the string data is duplicated and move many times. So, this practice
765 is expensive in execution time and memory consumption.
766
767 There are easy ways to prevent this behaviour. I guess that a C binding for
768 concatenation with chunks will be available ASAP (it is already written). I do
769 some benchmarks. I compare the execution time of 1 000 times, 1 000
770 concatenation of 10 bytes written in pure Lua and with a C library. The result is
771 10 times faster in C (1s in Lua, and 0.1s in C).
772
773 txn
774 ~~~~
775
776 txn is an HAProxy object of class TXN. The documentation is available in the
777 HAProxy Lua API reference. This class allow the access to the native HAProxy
778 sample_fetches and converters. The object txn contains 2 members dedicated to
779 the sample_fetches and 2 members dedicated to the converters.
780

```

The sample fetches members are "f" (as sample-Fetch) and "sf" (as String sample-Fetch). These two members contains exactly the same functions. All the HAProxy native sample fetches are available, obviously, the Lua registered sample fetches are not available. Unfortunately, HAProxy sample fetches names are not compatible with the Lua function names, and they are renamed. The rename convention is simple, we replace all the '.', '+', and '-' by '\_' . The '.' is the object member separator, and the "-" and "+" is math operator.

Now, that I'm writing this article, I know the Lua better than I wrote the sample-fetches wrapper. The original HAProxy sample-fetches name should be used using alternative manner to call an object member, so the sample-fetch "req.fhdr" (actually renamed req.fhdr") is should be used like this:

```
txn.f["req.fhdr"](txn.f, ...)
```

However, I think that this form is not elegant.

The "s" collection return a data with a type near to the original returned type. A string return an Lua string, an integer returns an Lua integer and an IP address returns an Lua string. Sometime the data is not or not yet available, in this case it returns the Lua nil value.

The "sf" collection guarantee that a string will be always returned. If the data is not available, an empty string is returned. The main usage of these collection is to concatenate the returned sample-fetches without testing each function.

The parameters of the sample-fetches are according with the haproxy documentation.

The converters runs exactly with the same manner as the sample fetches. The only one difference is that the first parameter is the converter entry element. The "c" collection returns a precise result, and the "sc" collection returns always a string.

The sample-fetches used in the example function are "txn.sf:req.fhdr()", "txn.sf:path()" and "txn.sf:src()". The converter are "txn.sc:sdbm()". The same function with the "s" collection of sample-fetches and the "c" collection of converter should be written like this:

```
1. function my_hash(txn, salt)
2.   local str = ""
3.   str = str .. salt
4.   str = str .. tostring(txn.f:req.fhdr("host"))
5.   str = str .. tostring(txn.f:path())
6.   str = str .. tostring(txn.f:src())
7.   local result = tostring(txn.c:sdbm(str, 1))
8.   return result
9. end
10.
11. core.register_fetches("my-hash", my_hash)
```

```
tostring
~~~~~
```

The function tostring ensure that its parameter is returned as a string. If the parameter is a table or a thread or anything that will not have any sense as a string, a form like the typename followed by a pointer is returned. For example:

```
t = {}
print(tostring(t))

returns:
table: 0x15facc0
```

For objects, if the special function \_\_tostring() is registered in the attached metatable, it will be called with the table itself as first argument. The HAProxy objects returns its own type.

About the converters entry point

In HAProxy, a converter is a stateless function that takes a data as entry and returns a transformation of this data as output. In Lua it is exactly the same behaviour.

So, the registered Lua function doesn't have any special parameters, juste a variable as input which contains the value to convert, and it must return data.

The data required as input by the Lua converter is a string. So HAProxy will always provide a string as input. If the native sample fetch is not a string it will ne converted in best effort.

The returned value will have anything type, it will be converted as sample of the near HAProxy type. The conversion rules from Lua variables to HAProxy samples are:

| Lua        | HAProxy sample types |
|------------|----------------------|
| "number"   | "sint"               |
| "boolean"  | "bool"               |
| "string"   | "str"                |
| "userdata" | "bool" (false)       |
| "nil"      | "bool" (false)       |
| "table"    | "bool" (false)       |
| "function" | "bool" (false)       |
| "thread"   | "bool" (false)       |

The function used for registering a converter is:

```
core.register_converters()
```

The task entry point

The function "core.register\_task(fcn)" executes once the function "fcn" when the scheduler starts. This way is used for executing background task. For example, you can use this functionality for periodically checking the health of an other service, and giving the result to each proxy needing it.

The task is started once, if you want periodic actions, you can use the "core.sleep()" or "core.msleap()" for waiting the next runtime.

Storing Lua variable between function in the same session

All the functions registered as action or sample fetch can share an Lua context. This context is a memory zone in the stack. sample fetch and action uses the same stack, so both can access to the context.

The context is accessible via the function get\_priv and set\_priv provided by an object of class TXN. The value given to set\_priv replaces the current stored value. This value can be a table, it is useful if a lot of data can be shared.

If the value stored is a table, you can add or remove entries from the table without storing again the new table. Maybe an example will be clearer:

```
local t = {}
txn:set_priv(t)
```

```
911 t["entry1"] = "foo"
912 t["entry2"] = "bar"
913
914 -- this will display "foo"
915 print(txn:get_priv()["entry1"])
916
917 HTTP actions
918 =====
919
920 ... coming soon ...
921
922 Lua is fast, but my service require more execution speed
923 =====
924
925 We can wrote C modules for Lua. These modules must run with HAProxy while they
926 are compliant with the HAProxy Lua version. A simple example is the "concat"
927 module.
928
929 It is very easy to write and compile a C Lua library, however, I don't see
930 documentation about this process. So the current chapter is a quick howto.
931
932 The entry point
933 -----
934
935 The entry point is called "luaopen_<name>", where <name> is the name of the ".so"
936 file. An hello world is like this:
937
938 #include <stdio.h>
939 #include <lua.h>
940 #include <lauxlib.h>
941
942 int luaopen_mymod(lua_State *L)
943 {
944 printf("Hello world\n");
945 return 0;
946 }
947
948 The build
949 -----
950
951 The compilation of the source file requires the Lua "include" directory. The
952 compilation and the link of the object file requires the -fpIC option. That's
953 all.
954
955 cc -I/opt/lua/include -fpIC -shared -o mymod.so mymod.c
956
957 Usage
958 ----
959
960 You can load this module with the following Lua syntax:
961
962 require("mymod")
963
964 When you start HAProxy, this module just print "Hello world" when its loaded.
965 Please, remember that HAProxy doesn't allow blocking method, so if you write a
966 function doing filesystem access or synchronous network access, all the HAProxy
967 process will fail.
```